Calibration and Risk-Budgeted Decisions with System One Models

Soundouss Messoudi
Heudiasyc - UMR CNRS 7253
Université de Technologie de Compiègne
soundouss.messoudi@hds.utc.fr

Youness Mansar
Independent Researcher
France

arXiv:2610.19811v1 [stat.ML] 14 Oct 2026

Abstract

LLM workflows and agents rely on a set of binary decisions like allow/block an input, or discard/retain a retrieved document. When human deferral is unavailable, uncertain cases must still receive an action. We study System One decision scores in this setting by defaulting to a conservative fallback when uncertain, and controlling the false omission risk within the subset of inputs assigned the utility-maximizing action. Across six datasets covering moderation, safety, and reranking, we compare general System One models with task baselines and a local open-weights fine-tuned System One model. We show that naive baselines for threshold selection exceed the target risk in roughly half of settings with defined test risk, compared with 7.8% for Learn Then Test (LTT), which leverages finite-sample statistical guarantees. Our results show why relying on a probability-like interface is insufficient for deployment decisions and how risk-budgeted selection trades useful action against observed error.

1 Introduction

While autoregressive Large Language Models (LLMs) excel at generation and reasoning, their token-by-token generation paradigm induces latency bottlenecks. This makes them ill-suited for fast, structured decisions in settings under tight latency and cost budgets, such as model routing (Ong et al., 2025), content moderation (Yang et al., 2023), prompt-injection detection (Guo et al., 2026), and safety gating (Han et al., 2024). This mismatch motivates lightweight decision models that expose actionable scores for downstream use.

System One models, first introduced with Jev (Typesafe AI, 2026), address this by mapping inputs to bounded decision scores rather than generating unrestricted text. Crucially, these models claim to output well-calibrated probabilities (Typesafe AI, 2026). This suggests an intuitive policy: under exact calibration, thresholding at α guarantees a False Omission Rate (the risk among inputs scored below α) bounded by α. However, optimizing System One models for calibration during training cannot guarantee test-time calibration across varied downstream distributions (Guo et al., 2017).

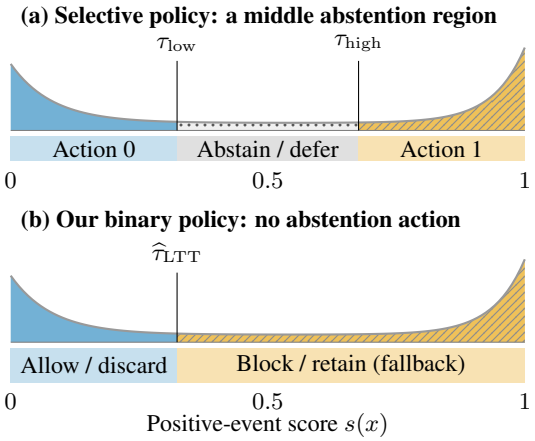

Figure 1: Binary actions' illustration. Under calibrated probabilities, scores near 0.5 indicate uncertain labels: a selective policy can abstain. Our policy instead assigns every input an action via an LTT-selected threshold.

The deployment challenge is therefore twofold: determining which operational action the system should take when a model's prediction is uncertain, and guaranteeing safety when its scores are uncalibrated. Moreover, these decisions have severe asymmetric effects that change what uncertainty should trigger (Figure 1). Selective classification typically mitigates risk through abstention, deferring difficult examples to a human or another model (Geifman and El-Yaniv, 2017). We focus on fully automated environments where abstention is not a viable action. In these forced-action settings, the conservative fallback is to block an input (for safety) or retain a document (for filtering).

Always defaulting to this fallback avoids the targeted false-negative error but degrades system utility through severe over-refusal or context bloat.

The practical objective is therefore to depart from the fallback action as often as possible, while controlling the errors within that newly selected subset. We express this constraint as the False Omission Rate (FOR), $\mathbb{P}(Y = 1 \mid \widehat{Y} = 0)$, where $\widehat{Y} = 0$ denotes the utility-maximizing choice (allow an input or discard a document). We apply Learn Then Test (LTT) (Angelopoulos et al., 2025) to statistically certify threshold regions that satisfy the FOR constraint, effectively substituting selective abstention with the conservative fallback.

We study three linked questions: (Q1) Do the supplied scores reliably represent positive-event probabilities and sufficiently separate the classes? (Q2) How do Naive, Empirical, and LTT threshold selection policies trade observed excess risk against utility? (Q3) How do these outcomes vary across System One models, specialist baselines, and a local open-weights fine-tuned System One model?

Evaluating on six datasets, we show that safely deploying automated decisions under asymmetric risk needs task-specific validation and rigorous risk-budgeting, rather than relying on the assumed universal calibration of any System One model.

2 Related Work

Selective Prediction and Risk Control. Selective classification trades coverage for accuracy (Geifman and El-Yaniv, 2017). An abstention option delegates uncertain inputs. Instead, our setting fixes two operational actions and measures error within the negative-action subset. Risk-controlling prediction sets and Learn Then Test (LTT) (Bates et al., 2021; Angelopoulos et al., 2025) use held-out data to support statistical risk. LTT is not restricted to abstention; we apply its conditional-risk formulation to binary action policies.

Probability Calibration. Because neural network calibration metrics are notoriously sensitive to binning (Guo et al., 2017; Nixon et al., 2019), proper scores like Brier are essential for broader probability assessment (Ferrer and Ramos, 2025). In LLMs, calibration has been examined for question answering (Jiang et al., 2021), self-evaluation (Kadavath et al., 2022), and elicited confidence (Tian et al., 2023; Xiong et al., 2024). These studies motivate examining how scores are obtained and what their probabilities actually mean.

Typed-decision evaluation. Datasets such as WildGuard (Han et al., 2024) and ATBench (Li et al., 2026b) provide foundations for evaluating moderation and trajectory safety. Recent studies of typed-decision models explore aggregate calibration and error ranking (Hu, 2026), as well as probability coherence across logically related questions (Li et al., 2026a). However, they do not certify a selected action policy's error rate. We distinguish our work by connecting System One scores to risk control across moderation, safety, and reranking.

3 Problem Formulation

3.1 Actions and Asymmetric Risk

Let $(X, Y) \sim P$ denote an input and its binary ground-truth label drawn from the evaluation population P, with $Y \in \{0, 1\}$. In our high-stakes language-model workflow, we have:

- $Y = 1$ represents the **high-stakes target** (e.g., an unsafe prompt injection attempt, a harmful message, or a highly relevant document).
- $Y = 0$ represents the **benign background** (e.g., a safe prompt, an allowable message, or an irrelevant document).

A frozen scorer returns $s(x) \in [0, 1]$, oriented so that larger scores express stronger evidence for $Y = 1$. We use threshold $\tau \in [0, 1]$ to define the binary decision rule $\widehat{Y} = \pi_\tau(x)$:

$$\pi_\tau(x) = \begin{cases} 0 & \text{if } s(x) < \tau, \\ 1 & \text{if } s(x) \geq \tau. \end{cases}$$

Here, Action 0 is the utility-maximizing choice (allowing safe inputs or discarding irrelevant documents) taken under high confidence. Action 1 serves as the conservative fallback (blocking or retaining), triggered by either high confidence in class 1 or predictive uncertainty.

In this setting, the costs of errors are asymmetric:

- **False Negative (Predict 0 / True 1):** This is a **critical failure**. It means allowing a toxic interaction to execute or permanently discarding a crucial piece of evidence.
- **False Positive (Predict 1 / True 0):** This is a **utility loss**. It causes over-refusal of benign inputs or over-inclusion that inflates an LLM context with irrelevant documents. It is treated as the lower-priority error in our setting.

Because predicting 0 when the true label is 1 represents the critical system failure, the primary

risk $R(\tau)$ is defined on the negative-action subset of data (i.e., when the system predicts 0). This is the False Omission Rate (FOR):

$$\begin{aligned} R(\tau) &= \mathbb{P}(Y = 1 \mid \widehat{Y} = 0) \\ &= \mathbb{P}(Y = 1 \mid s(X) < \tau). \end{aligned} \tag{1}$$

3.2 The Optimization Trade-Off

While the high-stakes constraint requires defaulting to Action 1 when uncertain, always predicting 1 gives a practically useless system with an empty negative-action subset, where its measure is $\mu(\tau) = \mathbb{P}(s(X) < \tau) = 0$.

The system's primary optimization goal is to maximize the size of this negative-action subset subject to a strict population-level risk budget $\alpha \in (0, 1)$. Ideally, we seek a threshold τ that solves:

$$\begin{aligned} \max_{\tau} \quad & \mu(\tau) = \mathbb{P}(s(X) < \tau) \\ \text{subject to} \quad & R(\tau) \leq \alpha \end{aligned}$$

Maximizing this set is linked to minimizing the number of false positives, $\sum_{i=1}^{n} \mathbf{1}\{\widehat{Y}_{\tau,i} = 1, Y_i = 0\}$. As $\mu(\tau)$ grows, fewer true zeros are incorrectly forced into the fallback action (Action 1), driving down the incidence of over-refusal and context bloat.

However, because the true data distribution is unknown, we cannot directly access the population risk $R(\tau)$. The objective is therefore to establish an empirical threshold selection procedure that selects the largest certified candidate τ to minimize utility loss, while relying on statistical testing to rigorously ensure the critical risk of catastrophic FOR is controlled at budget α.

4 Threshold Selection and Risk Control

4.1 Certifying the False Omission Rate

To solve the finite-sample optimization trade-off introduced in Section 3.2, we must rely on statistical testing. Given the target risk budget α and a statistical error level $\delta \in (0, 1)$, the formal objective is to find a threshold τ that maximizes the measure of the negative-action subset while guaranteeing the risk bound holds with high probability:

$$\begin{aligned} \max_{\tau} \quad & \mu(\tau) = \mathbb{P}(s(X) < \tau) \\ \text{subject to} \quad & \mathbb{P}(R(\tau) \leq \alpha) \geq 1 - \delta \end{aligned}$$

To obtain this guarantee, we specialize Learn Then Test (LTT) (Angelopoulos et al., 2025) to the binary conditional risk in Equation 1.

Algorithm 1 Fixed-order threshold certification.

Require: Frozen scorer; independently fixed ordered thresholds $(\tau_1, \ldots, \tau_M)$; independent D_{cal}; (α, δ)

Initialize the certified threshold set $\widehat{\mathcal{T}} \leftarrow \varnothing$
for $j = 1, \ldots, M$ **do**
  Compute (N_j, K_j, p_j)
  if $p_j > \delta$ **then**
    break ▷ Stop the sequence
  else
    $\widehat{\mathcal{T}} \leftarrow \widehat{\mathcal{T}} \cup \{\tau_j\}$
  end if
end for
if $\widehat{\mathcal{T}} = \varnothing$ **then**
  return $\perp$ ▷ Non-deployment (Fallback to Action 1)
end if
return $\widehat{\tau} \in \arg\max_{\tau_j \in \widehat{\mathcal{T}}} \widehat{\mu}_{\text{cal},j}$ ▷ Maximize negative-action subset measure

Ensure: Selected $\widehat{\tau}$, or non-deployment ($\perp$)

Let $\mathcal{T} = \{\tau_1, \ldots, \tau_M\}$ contain M candidate thresholds, tested in the specified order $(\tau_1, \ldots, \tau_M)$ fixed independently of the calibration[1] sample. Let $\alpha \in (0, 1)$ be the risk budget and $\delta \in (0, 1)$ be the error level that bounds the probability of certifying any excessive-risk candidate in this family. Let D_{fit} denote data used to fit the scorer $s(X)$ and construct the threshold family's grid, and let $D_{\text{cal}} = \{(X_i, Y_i)\}_{i=1}^{n_{\text{cal}}}$ be the separate calibration sample of size n_{cal}. Define:

$$\begin{aligned} N_j &= \sum_{i=1}^{n_{\text{cal}}} \mathbf{1}\{s(X_i) < \tau_j\}, \\ K_j &= \sum_{i=1}^{n_{\text{cal}}} Y_i \mathbf{1}\{s(X_i) < \tau_j\}. \end{aligned}$$

For candidate j, N_j counts calibration examples assigned Action 0, and K_j counts positive labels among them. For $N_j > 0$,

$$\widehat{R}_{\text{cal},j} = \frac{K_j}{N_j}, \qquad \widehat{\mu}_{\text{cal},j} = \frac{N_j}{n_{\text{cal}}}.$$

For a fixed policy under independent and identically distributed (i.i.d.) calibration sampling, conditioning on $N_j = n$ gives

$$K_j \mid N_j = n \sim \text{Binomial}(n, R(\tau_j)).$$

[1] The word "calibration" in LTT refers to statistical selection, not probability recalibration.

To test the null hypothesis $H_j : R(\tau_j) > \alpha$ (i.e., that the policy exceeds the risk budget), we use the lower-tail binomial p-value:

$$p_j = \sum_{k=0}^{K_j} \binom{N_j}{k} \alpha^k (1-\alpha)^{N_j - k},$$

with $p_j = 1$ when $N_j = 0$. A small p_j provides the statistical evidence needed for certification. Algorithm 1 states the fixed-family procedure to which the theoretical LTT guarantees apply.

Theorem 1 (Fixed-family LTT specialization (Angelopoulos et al., 2025))**.** *Condition on D_{fit}, which fixes the scorer and ordered candidate family. Suppose that D_{cal} consists of i.i.d. draws from P, independently of D_{fit}. Algorithm 1 satisfies:*

$$\mathbb{P}_{D_{\text{cal}}}\left(\exists \tau \in \widehat{\mathcal{T}} : R(\tau) > \alpha\right) \leq \delta.$$

Consequently, any rule returning an element of $\widehat{\mathcal{T}}$ or $\perp$ (non-deployment) satisfies

$$\mathbb{P}_{D_{\text{cal}}}(\widehat{\tau} \neq \perp,\ R(\widehat{\tau}) > \alpha) \leq \delta.$$

Proof. This theorem is a direct application of fixed-sequence testing (specifically Theorem 2.1) established by Angelopoulos et al. (2025), applied to our binary thresholding rule. □

At deployment, the selected $\widehat{\tau}$ assigns every input to Action 0 or 1. Returning $\perp$ means threshold selection did not establish a certificate for a useful negative-action subset, and the system must default to the fallback action (Action 1). This distinction separates *failure to select a certified policy* from *deferring an individual input*.

4.2 Why Naive Thresholding Is Not Enough

Let us consider the relationship between population risk and score calibration. Let the random variable $S = s(X)$ be the score and $m(u) = \mathbb{P}(Y = 1 \mid S = u) = \mathbb{E}[Y \mid S = u]$ its calibration function. Exact calibration means $m(S) = S$ except on a probability-zero set. The following identity separates the average model score from the average probability error within the negative-action region.

Proposition 1 (Calibration–policy identity)**.** *For $\mu(\tau) > 0$,*

$$R(\tau) = \mathbb{E}[S \mid S < \tau] + \mathbb{E}[m(S) - S \mid S < \tau].$$

Consequently,

$$R(\tau) \leq \tau + \frac{\mathbb{E}[|m(S) - S| \mathbf{1}\{S < \tau\}]}{\mu(\tau)}.$$

Under exact calibration, $R(\tau) \leq \tau$.

The derivation appears in Appendix A. This identity demonstrates that if a model is perfectly calibrated, naive thresholding achieves bounded risk (i.e., setting $\tau = \alpha$ gives $R(\tau) \leq \alpha$). However, it also proves that when scores are uncalibrated, depending on the direction of the error, naive thresholding can allow excessive risk. LTT instead tests the population-risk constraint using finite-sample data, without requiring calibrated model scores. Its guarantee is the calibration-sampling statement in Theorem 1, not a promise that every finite test estimate is below budget.

5 Experimental Setup

5.1 Evaluated Models

We categorize evaluated models into three groups:

- **Baselines:** For moderation and safety, we evaluate Qwen3Guard-Gen-0.6B (Zhao et al., 2025) as a specialist baseline. For reranking, we evaluate Voyage rerank-2.5-lite (Voyage AI, 2026). We also include PCA-Ridge, an L2-regularized logistic regression baseline fitted per dataset on frozen Qwen3-8B embeddings.

- **System One Models:** We evaluate models specifically designed with probability-valued interfaces, including Jev (Typesafe AI, 2026), GPT-6-Luna (via its typed-decisions endpoint) (OpenAI, 2026), Span-01 (Respan, 2026), and Solar-Decide (Upstage, 2026).

- **Base vs. Fine-Tuned:** To illustrate adaptation dynamics within a single System One model family, we evaluate Kev-0.8B (Palmer, 2026) both as released (Kev-base) and after fine-tuning on $\approx 10k$ examples for 37 minutes on an RTX 5090 (Kev-ft). Kev-ft is trained on a mixture of all dataset fit splits to improve generalization, limit overfitting, and save compute cost of six separate neural trainings.

Other details are provided in Appendix B.1.

5.2 Tasks and Datasets

We evaluate our framework across three workflows with six datasets. **Moderation** uses WildGuardMix and Aegis prompt labels. **Safety** covers prompt injection on S-Labs and unsafe recorded trajectories on ATBench. **Reranking** uses SciDocs and StackOverflow query–document pairs to evaluate reranking as a binary filtering decision.

Task	Dataset	Fit	Cal.	Test	Y=1 (%)
Moderation	WildGuardMix (Han et al., 2024)	2000	849	850	44.4
Moderation	Aegis (Ghosh et al., 2024)	2000	579	579	67.0
Safety	S-Labs (S-Labs, 2026)	2000	1051	1050	50.0
Safety	ATBench (Li et al., 2026b)	500	250	250	47.6
Reranking	SciDocs (Cohan et al., 2020)	2000	1000	1000	50.8
Reranking	StackOverflow (Muennighoff et al., 2023)	2000	1000	1000	50.1

Table 1: Dataset split details. The last column gives the test positive-event prevalence (base rate).

For each dataset, we use a single seed, label-stratified fit/calibration/test splits (Table 1). Because most evaluated scorers are general pre-trained models, D_{fit} is only used to fit the PCA-Ridge baseline, fine-tune the Kev-ft model, and construct candidate threshold grids. D_{cal} is used for threshold selection, and test data evaluates the final fixed policies. Other details are provided in Appendix B.2 and Appendix B.3.

5.3 Threshold Selection Policies

We evaluate three threshold policies across risk budgets $\alpha \in \{0.05, 0.10, 0.15, 0.20, 0.40, 0.50\}$:

1. **Naive:** Sets $\tau = \alpha$. As discussed in Section 4.2, this assumes exact calibration and includes no finite-sample margin.

2. **Empirical:** Selects τ that yields the lowest False Positive Rate, provided that the observed calibration FOR is at most α.

3. **LTT:** Applies Algorithm 1 with a statistical error level of $\delta = 0.10$ to return the largest threshold in the certified prefix. Details on implementation are in Appendix B.4.

5.4 Evaluation Metrics

We evaluate the systems on two axes: underlying score quality and operational policy outcomes.

Score Diagnostics. We assess raw predictions before thresholding with Area Under the Receiver Operating Characteristic (AUROC) and negative-action Average Precision (AP_0) for class separation. We report 10-bin expected calibration error (ECE) and Brier score to measure probability alignment.

Operational Utility and Risk. At deployment, we report the True Negative Rate (TNR) as our primary measure of system utility (e.g., the rate of benign acceptance or irrelevant document removal). Empirical risk is measured by test-set FOR, $\widehat{R}_{\text{test}}$. To quantify safety violations, we report the positive excess risk, defined as $e = \max(\widehat{R}_{\text{test}} - \alpha, 0)$.

To ensure statistical rigor, 95% bootstrap confidence intervals for score diagnostics and primary-budget policy risk/utility appear in Appendix C. The resampling recipe is in Appendix B.5.

Availability and Resource Use. We distinguish no LTT certificate (ND), no feasible Empirical policy (NF), and a returned policy with an empty negative-action test set (UD, undefined FOR). Exceedance fractions use defined test risks; status fractions use settings with defined risk. Mean excess includes zero excess for non-exceedances and is expressed in percentage points. We also report inference cost (USD per 1,000 workload examples) and mean recorded request time (ms/example). Full metric definitions and accounting conventions appear in Appendix B.5 and Appendix B.6.

6 Results

We answer our three research questions by evaluating the held-out test predictions and saved policy selections across all six datasets. Table 2 summarizes the underlying score diagnostics and calibration metrics for each model. Table 3 aggregates the frequency and magnitude of risk exceedances alongside overall policy availability. Table 4 details the resulting operational utility and empirical risk at the target budget $\alpha = 0.15$. Figure 2 visualizes the LTT-selected operating points across varying risk tolerances to demonstrate model differences on WildGuardMix. Finally, Table 5 reports the inference cost and recorded latency to contextualize operational deployment.

6.1 Q1: Calibration is Task-Dependent and Insufficient for Policy Viability

Table 2 shows that probability reliability is fundamentally dataset-dependent and no single architecture universally dominates. A bounded interface alone does not guarantee calibration. For example, Jev's ECE degrades from 0.034 on WildGuardMix to 0.29 on SciDocs. Universal claims of "calibrated probabilities" obscure this severe variability.

Model	Moderation: WildGuardMix				Moderation: Aegis				Safety: S-Labs			
	ECE ↓	Brier ↓	AUC ↑	AP_0 ↑	ECE ↓	Brier ↓	AUC ↑	AP_0 ↑	ECE ↓	Brier ↓	AUC ↑	AP_0 ↑
Qwen3Guard	0.085	0.090	0.950	0.947	0.277	0.239	0.929	0.866	0.474	0.464	0.968	0.968
PCA-Ridge*	0.042	0.091	0.946	0.956	0.037	0.096	0.934	0.892	0.087	0.041	0.996	0.996
Jev	0.034	0.073	0.963	0.969	0.183	0.151	0.936	0.860	0.185	0.158	0.947	0.927
GPT-6-Luna	0.103	0.108	0.925	0.907	0.297	0.277	0.864	0.665	0.262	0.228	0.888	0.841
Span-01	0.189	0.238	0.765	0.746	0.500	0.472	0.565	0.332	0.441	0.437	0.288	0.368
Solar-Decide	0.104	0.130	0.893	0.855	0.356	0.339	0.865	0.658	0.337	0.326	0.840	0.737
Kev-base	0.121	0.186	0.835	0.863	0.257	0.232	0.834	0.704	0.225	0.203	0.944	0.940
Kev-ft*	0.119	0.123	0.934	0.938	0.109	0.126	0.919	0.849	0.024	0.028	0.995	0.995
Model	Safety: ATBench				Reranking: SciDocs				Reranking: StackOverflow			
	ECE ↓	Brier ↓	AUC ↑	AP_0 ↑	ECE ↓	Brier ↓	AUC ↑	AP_0 ↑	ECE ↓	Brier ↓	AUC ↑	AP_0 ↑
Reranker-Lite	–	–	–	–	0.183	0.214	0.880	0.861	0.137	0.200	0.829	0.826
PCA-Ridge*	0.151	0.165	0.882	0.902	0.065	0.127	0.904	0.885	0.032	0.204	0.750	0.748
Jev	0.076	0.200	0.768	0.717	0.290	0.236	0.917	0.895	0.099	0.196	0.790	0.800
GPT-6-Luna	0.230	0.232	0.842	0.844	0.340	0.289	0.865	0.787	0.120	0.191	0.815	0.819
Span-01	0.101	0.196	0.779	0.744	0.175	0.189	0.864	0.834	0.224	0.302	0.607	0.576
Solar-Decide	0.274	0.288	0.688	0.716	0.357	0.333	0.772	0.643	0.266	0.288	0.706	0.643
Kev-base	0.255	0.303	0.457	0.538	0.120	0.234	0.734	0.724	0.051	0.245	0.599	0.559
Kev-ft*	0.059	0.099	0.938	0.952	0.118	0.131	0.913	0.881	0.160	0.205	0.810	0.818

Table 2: Test score diagnostics by dataset. AUC is AUROC and AP_0 is negative-action average precision. Reranker scores are not advertised probabilities. 95% intervals in Appendix C. –: not evaluated. *: supervised adaptation.

Crucially, low global ECE implies neither effective class separation nor policy viability. On StackOverflow, Kev-base achieves a much better ECE than Kev-ft (0.051 vs. 0.16) but exhibits worse discriminative power (AUROC 0.599 vs. 0.81). This dynamic illustrates Proposition 1: while ECE averages absolute error across all inputs, operational risk depends on the conditional error within the selected negative-action set, $\mathbb{E}[m(S) - S \mid S < \tau]$. Thus, aggregate calibration cannot substitute for rigorous, subset-specific policy evaluation.

6.2 Q2: LTT Prioritizes Safety While Naive Thresholds Frequently Exceed Budgets

Our analysis shows that naive thresholding can be arbitrarily permissive or conservative depending on the dataset. Naive methods routinely trade away safety for utility: on Aegis at $\alpha = 0.15$, Span-01 achieves 95.3% utility under Naive thresholding, but its empirical risk catastrophically explodes to 60.9% (Table 4). Conversely, naive thresholds can be unnecessarily restrictive. On WildGuardMix, Jev's naive policy limits utility to 68.7% (at 3.3% risk), whereas LTT safely expands it to 96.8% while remaining within budget (12.8% risk).

While the Empirical selection rule can sometimes identify high-utility policies (e.g., Kev-ft reaching 99.6% utility at 12.5% risk on S-Labs), it lacks formal guarantees. Table 3 shows that Empirical selection exceeds the target budget in 51.7% of settings (mean excess 1.17%). The Naive rule similarly exceeds the budget 42% of the time, concentrating the most extreme failures (mean excess 7.09%, with a maximum exceedance of 73.26 percentage points at $\alpha = 0.05$).

LTT produces substantially fewer and smaller risk exceedances (just 7.8% frequency, mean excess 0.1%) because its statistical certification structurally defaults to non-deployment (ND) rather than operate unsafely. This provides a rigorous finite-sample guarantee that, with probability at least $1 - \delta$, the system will stay below the risk budget. Consequently, LTT returns fewer policies under strict budgets (establishing certificates for only 21.3% of policies at $\alpha = 0.05$), though availability increases as the budget relaxes (76.6% at $\alpha = 0.5$).

6.3 Q3: Adaptation Expands Availability, Though General System One Models Remain Highly Competitive

General System One interfaces offer a practical plug-and-play starting point, though risk-controlled deployment still requires finite-sample data for policy selection. As Figure 2 shows on WildGuardMix ($\alpha = 0.15$), Jev and GPT-6-Luna achieve exceptional LTT utility (96.8% and 92.6%, respectively), rivaling the specialist Qwen3Guard baseline (95.1%). Thus, high utility does not strictly require fine-tuning, although some unadapted mod-

α	Rule	Eval.	Fail	Exceeds	Mean e	Max e
5	Naive	34	13 UD	23 (67.6%)	14.00	73.26
	Emp.	23	24 NF	19 (82.6%)	1.73	5.81
	LTT	10	37 ND	0 (0.0%)	0.00	0.00
10	Naive	38	9 UD	22 (57.9%)	11.22	48.87
	Emp.	30	17 NF	18 (60.0%)	1.36	6.45
	LTT	20	27 ND	1 (5.0%)	0.03	0.53
15	Naive	41	6 UD	23 (56.1%)	9.81	45.86
	Emp.	32	15 NF	20 (62.5%)	1.59	8.11
	LTT	24	23 ND	1 (4.2%)	0.07	1.60
20	Naive	43	4 UD	20 (46.5%)	8.06	41.30
	Emp.	34	13 NF	19 (55.9%)	1.49	8.57
	LTT	29	18 ND	3 (10.3%)	0.21	3.47
40	Naive	47	0 UD	13 (27.7%)	2.03	23.39
	Emp.	44	3 NF	20 (45.5%)	0.91	5.70
	LTT	34	13 ND	4 (11.8%)	0.09	1.69
50	Naive	47	0 UD	4 (8.5%)	0.54	13.69
	Emp.	46	1 NF	12 (26.1%)	0.48	4.78
	LTT	36	11 ND	3 (8.3%)	0.10	1.79
$\sum$	Naive	250	32 UD	105 (42.0%)	7.09	73.26
	Emp.	209	73 NF	108 (51.7%)	1.17	8.57
	LTT	153	129 ND	12 (7.8%)	0.10	3.47

Table 3: Outcomes by risk budget α (%) across all 47 model–dataset pairs. Excess risk is $e = \max(\widehat{R} - \alpha, 0)$.

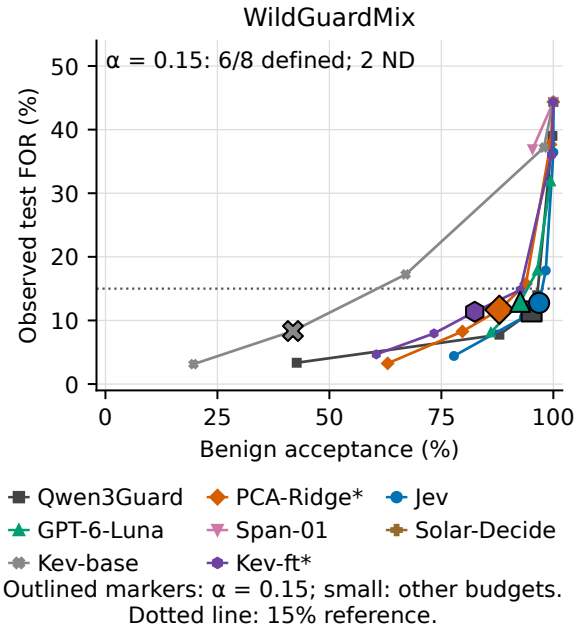

Figure 2: LTT-selected risk–utility operating points on WildGuardMix. Lower test FOR and greater benign acceptance are preferable. Missing policies are omitted.

els (Span-01, Solar-Decide) fail entirely (Table 4).

However, local open-weight model fine-tuning massively improves policy availability over unadapted base models. At $\alpha = 0.15$, Kev-base returns an LTT policy on only one of six datasets, whereas Kev-ft succeeds on all six (Table 4). Across all budgets, this adaptation expands availability from 13/36 to 31/36 settings. Even when both return policies, as on WildGuardMix, utility jumps significantly from 41.9% for Kev-base to 82.5% for Kev-ft (Figure 2).

Fine-tuning does not universally dominate: Jev outperforms Kev-ft on WildGuardMix with a 96.8% utility at this budget, and the PCA-Ridge baseline remains competitive across several tasks. Because no single architecture solves forced-action deployment, a practical strategy is to first evaluate a general System One model and fall back to local open-weights model adaptation if it proves too conservative. Crucially, all approaches demand strict task-specific empirical validation.

6.4 Local Open-Weights Models Provide the Most Efficient Cost-Latency Balance

Operational deployment must balance policy usefulness against inference latency and financial cost (Table 5). Proprietary closed-weight models exhibit widely varying degrees of cost and latency depending on the endpoint. Crucially, simpler baseline architectures do not automatically guarantee faster execution. The PCA-Ridge baseline is exceptionally cheap to evaluate numerically ($0.0026 per 1,000 examples) but it has a severe latency penalty (5,247 ms/example). This delay comes from the requisite embedding extraction pipeline, which must fully process the input through Qwen3-8B before the lightweight logistic regression can execute. By contrast, an end-to-end local fine-tuned model like Kev-ft provides the most optimal balance, delivering highly competitive safety utility at just $0.0030 and 98 ms/example, avoiding both the API costs of proprietary models and the sequential embedding bottlenecks of PCA-Ridge.

7 Discussion and Conclusion

Deploying System One models for asymmetric decisions, where one error type carries severe consequences, demands more than aggregate calibration. Results show that favorable global diagnostics do not guarantee utility, and naive or empirical thresholding often lead to severe, silent risk violations.

LTT offers a rigorous statistical alternative, controlling the useful action region without trusting raw scores as true probabilities. The operational price of this safety is reduced policy availability. LTT minimizes budget exceedances by returning non-deployment when a certificate cannot be es-

	Moderation: WildGuardMix						Moderation: Aegis					
	Naive		Empirical		LTT		Naive		Empirical		LTT	
Model	Risk	Util	Risk	Util	Risk	Util	Risk	Util	Risk	Util	Risk	Util
Qwen3Guard	10.8	94.3	12.2	95.8	11.4	95.1	38.4	93.2	18.2	72.8	11.5	60.7
PCA-Ridge*	4.2	67.4	13.1	91.3	11.7	87.9	1.4	36.1	17.3	72.8	9.5	64.9
Jev	3.3	68.7	16.2	**98.1**	12.8	96.8	20.8	83.8	23.1	85.3	11.7	67.0
GPT-6-Luna	12.6	92.6	15.6	96.0	12.9	92.6	43.4	94.8	NF	NF	ND	ND
Span-01	26.3	81.2	NF	NF	ND	ND	60.9	95.3	NF	NF	ND	ND
Solar-Decide	18.1	93.9	13.9	84.8	ND	ND	50.7	**97.9**	NF	NF	ND	ND
Kev-base	0.0	4.0	15.4	62.6	8.3	41.9	12.5	7.3	18.6	29.8	ND	ND
Kev-ft*	11.3	82.9	13.4	87.5	11.4	82.5	17.7	70.7	19.5	71.2	14.5	64.9

	Safety: S-Labs						Safety: ATBench					
	Naive		Empirical		LTT		Naive		Empirical		LTT	
Model	Risk	Util	Risk	Util	Risk	Util	Risk	Util	Risk	Util	Risk	Util
Qwen3Guard	47.8	**100.0**	14.7	94.1	13.5	93.1	–	–	–	–	–	–
PCA-Ridge*	0.4	87.0	11.6	99.8	10.0	99.8	0.0	1.5	13.6	72.5	10.1	67.9
Jev	11.2	86.1	13.9	95.4	10.2	79.0	33.3	3.1	NF	NF	ND	ND
GPT-6-Luna	27.0	94.3	13.2	71.4	ND	ND	5.6	26.0	16.3	62.6	12.0	55.7
Span-01	47.5	**100.0**	NF	NF	ND	ND	27.3	12.2	NF	NF	ND	ND
Solar-Decide	41.0	**100.0**	NF	NF	ND	ND	28.6	61.1	NF	NF	ND	ND
Kev-base	UD	UD	14.6	88.2	ND	ND	UD	UD	NF	NF	ND	ND
Kev-ft*	3.6	97.1	12.5	99.6	11.1	99.6	5.1	71.0	15.7	**90.1**	14.6	89.3

	Reranking: SciDocs						Reranking: StackOverflow					
	Naive		Empirical		LTT		Naive		Empirical		LTT	
Model	Risk	Util	Risk	Util	Risk	Util	Risk	Util	Risk	Util	Risk	Util
Reranker-Lite	UD	UD	11.3	56.1	ND	ND	UD	UD	18.5	57.3	14.0	51.7
PCA-Ridge*	7.6	24.6	17.2	78.0	13.9	72.8	8.8	6.2	16.5	23.2	ND	ND
Jev	22.8	90.9	18.8	85.8	13.7	76.8	8.1	25.1	20.5	54.3	16.6	44.3
GPT-6-Luna	31.5	**99.0**	NF	NF	ND	ND	10.1	40.9	15.9	55.3	11.3	42.5
Span-01	16.4	65.2	15.9	63.4	ND	ND	41.2	47.5	NF	NF	ND	ND
Solar-Decide	40.1	98.0	NF	NF	ND	ND	34.6	**73.3**	NF	NF	ND	ND
Kev-base	UD	UD	19.4	25.4	ND	ND	UD	UD	NF	NF	ND	ND
Kev-ft*	15.2	79.5	15.9	83.9	14.8	78.5	16.7	54.9	16.2	53.9	13.4	50.3

Table 4: Test risk and utility (%) at $\alpha = 0.15$. Risk is FOR; utility is TNR (benign acceptance for safety, irrelevant removal for filtering). Red risk values exceed the budget. ND: no LTT certificate; NF: no feasible empirical policy; UD: undefined test FOR. –: not evaluated. *: supervised adaptation. Bold utility is the largest value per dataset.

	MD: WildGuardMix		MD: Aegis		SF: S-Labs		SF: ATBench		RK: SciDocs		RK: StackOverflow	
Model	Cost	Time	Cost	Time	Cost	Time	Cost	Time	Cost	Time	Cost	Time
Qwen3Guard	0.0071	231	0.0077	252	0.0065	213	–	–	–	–	–	–
Reranker-Lite	–	–	–	–	–	–	–	–	0.0013	224	0.0011	183
PCA-Ridge*	0.0026	5247	0.0039	1880	0.0013	5422	0.0215	14541	0.0014	5707	0.0012	3765
Jev	0.0216	383	0.0241	248	0.0204	379	0.0757	261	0.0175	364	0.0172	347
GPT-6-Luna	0.0375	320	0.0435	225	0.0347	186	0.1507	294	0.0283	207	0.0278	167
Span-01	0.0084	672	0.0095	286	0.0081	648	0.0320	1059	0.0053	638	0.0050	710
Solar-Decide	0.0298	1361	0.0329	2076	0.0281	1038	0.0977	6303	0.0252	1536	0.0251	2351
Kev-base	0.0029	95	0.0030	99	0.0030	97	0.0076	249	0.0027	87	0.0026	85
Kev-ft*	0.0030	98	0.0032	104	0.0030	97	0.0074	241	0.0028	91	0.0027	88

Table 5: Cost (USD/1k) and latency (ms). MD: Moderation. SF: Safety. RK: Reranking. –: not evaluated.

tablished, correctly defaulting to a conservative fallback (e.g., blocking a prompt or retaining a document) rather than operating unsafely. However, our findings show that this availability bottleneck can be mitigated: when general interfaces fail to yield a certificate, adapting lightweight, local open-weights models provides a highly reliable and cost-efficient pathway to recover certified policies, though this comes with the cost of training data.

Future work should extend LTT certification to multi-model cascades to optimize the balance between safety, latency, and inference cost, and explore how larger calibration samples tighten risk bounds to increase policy availability.

Limitations

First, the finite-sample guarantee in Theorem 1 relies on calibration and deployment examples being drawn i.i.d. from the same distribution P. Our six English-language benchmarks do not establish performance under deployment or adversarial shift, where maintaining risk control would require periodic recalibration or distribution-shift adaptations. Moreover, ATBench evaluates retrospective classification of truncated trajectories, and our reranking tasks use a constructed document mixture rather than an end-to-end retrieval pipeline; both choices affect the meaning and transfer of the measured FOR.

Second, exact binomial certification imposes a fundamental sample-size requirement on the calibration set. Even with zero observed errors ($K_j = 0$), certifying a threshold at budget α and error level δ requires at least $\lceil \log \delta / \log(1-\alpha) \rceil$ selected calibration examples (Appendix B.4). Consequently, when calibration sets are small or risk budgets are strict (e.g., $\alpha = 0.05$), LTT frequently defaults to non-deployment ($\perp$), trading policy availability for safety.

Finally, while adapting a lightweight open-weights model (Kev-ft) reliably recovers certified utility when zero-shot System One interfaces fail to certify, this pathway requires assembling labeled training data (D_{fit}) and local training compute, which may not be feasible in cold-start deployments.

Ethical Considerations

This study evaluates datasets containing harmful, offensive, and adversarial text. We analyze dataset labels and model scores rather than distributing new attack payloads, and this paper contains no raw harmful examples.

A low measured FOR is not a blanket safety endorsement: false negatives expose users to harm, while false positives suppress benign expression. Aggregate metrics do not establish equitable error rates across demographic groups or topics. In filtering, document removal can skew downstream generation even at a low discarded-set contamination rate. Safe deployment requires an application-specific risk definition, continuous monitoring, and an explicit fallback action when LTT returns no valid policy ($\perp$).

Benchmark annotations embody specific normative safety and relevance judgments that should not be mistaken for universal definitions. Furthermore, this study cannot guarantee the absence of training-data contamination for externally supplied APIs. Model and data identifiers are provided to support independent scrutiny; any redistribution must comply with the original licenses.

References

Anastasios N Angelopoulos, Stephen Bates, Emmanuel J Candès, Michael I Jordan, and Lihua Lei. 2025. Learn then test: Calibrating predictive algorithms to achieve risk control. *The Annals of Applied Statistics*, 19(2):1641–1662.

Stephen Bates, Anastasios Angelopoulos, Lihua Lei, Jitendra Malik, and Michael Jordan. 2021. Distribution-free, risk-controlling prediction sets. *Journal of the ACM (JACM)*, 68(6):1–34.

Arman Cohan, Sergey Feldman, Iz Beltagy, Doug Downey, and Daniel S Weld. 2020. Specter: Document-level representation learning using citation-informed transformers. In *Proceedings of the 58th annual meeting of the association for computational linguistics*, pages 2270–2282. Association for Computational Linguistics.

Luciana Ferrer and Daniel Ramos. 2025. Evaluating posterior probabilities: Decision theory, proper scoring rules, and calibration. *Transactions on Machine Learning Research*.

Yonatan Geifman and Ran El-Yaniv. 2017. Selective classification for deep neural networks. *Advances in neural information processing systems*, 30.

Shaona Ghosh, Prasoon Varshney, Erick Galinkin, and Christopher Parisien. 2024. Aegis: Online adaptive AI content safety moderation with ensemble of LLM experts. *arXiv preprint arXiv:2404.05993*.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q Weinberger. 2017. On calibration of modern neural networks. In *International conference on machine learning*, pages 1321–1330. PMLR.

Maihao Guo, Chaoyang Zhao, and Jinqiao Wang. 2026. BERM: Low-overhead prompt-injection detection via in-situ benign representation modeling. In *Proceedings of the Thirty-Fifth International Joint Conference on Artificial Intelligence, IJCAI-26*, pages 599–607. International Joint Conferences on Artificial Intelligence Organization. Main Track.

Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. 2024. Wildguard: Open one-stop moderation tools for safety risks, jailbreaks, and refusals of LLMs. *Advances in neural information processing systems*, 37:8093–8131.

Yibo Hu. 2026. Safety-flag: A unified benchmark for the reliability and calibration of LLM content moderators. *arXiv preprint arXiv:2609.19072*.

Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. 2021. How can we know when language models know? on the calibration of language models for question answering. *Transactions of the Association for Computational Linguistics*, 9:962–977.

Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, et al. 2022. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*.

Keyi Li, Yihao He, and Quanyi Li. 2026a. Beyond calibration: Do a typed-decision model's probabilities obey the probability axioms? *arXiv preprint arXiv:2609.33209*.

Yu Li, Haoyu Luo, Yuejin Xie, Yuqian Fu, Zhonghao Yang, Shuai Shao, Qihan Ren, Wanying Qu, Yanwei Fu, Yujiu Yang, Jing Shao, and Dongrui Liu. 2026b. ATBench: A diverse and realistic agent trajectory benchmark for safety evaluation and diagnosis. In *Third Conference on Language Modeling*.

Niklas Muennighoff, Nouamane Tazi, Loïc Magne, and Nils Reimers. 2023. MTEB: Massive text embedding benchmark. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 2014–2037. Association for Computational Linguistics.

Jeremy Nixon, Mike Dusenberry, Ghassen Jerfel, Timothy Nguyen, Jeremiah Liu, Linchuan Zhang, and Dustin Tran. 2019. Measuring calibration in deep learning. *arXiv preprint arXiv:1904.01685*.

Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E Gonzalez, Mohammed Kadous, and Ion Stoica. 2025. RouteLLM: Learning to route LLMs from preference data. In *International Conference on Learning Representations*, volume 2025, pages 34433–34448.

OpenAI. 2026. Decisions API documentation. https://platform.openai.com/docs/guides/decisions.

Jared Palmer. 2026. kev: A laptop-scale reconstruction of a Jev-style decision model. https://github.com/jaredpalmer/kev.

Respan. 2026. Introducing Span-1. https://www.respan.ai/blog/introducing-span-1.

S-Labs. 2026. Prompt injection detection dataset. https://huggingface.co/datasets/S-Labs/prompt-injection-dataset.

Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D Manning. 2023. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5433–5442. Association for Computational Linguistics.

Typesafe AI. 2026. Introducing System One models and Jev. https://typesafe.ai/blog/introducing-system-one-models-and-jev.

Upstage. 2026. Solar Decide API documentation. https://console.upstage.ai/docs/decide.

Voyage AI. 2026. Voyage reranker models. https://docs.voyageai.com/docs/reranker.

Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. 2024. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. In *International Conference on Learning Representations*, volume 2024, pages 23650–23678.

Zachary Yang, Nicolas Grenon-Godbout, and Reihaneh Rabbany. 2023. Towards detecting contextual real-time toxicity for in-game chat. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 9894–9906. Association for Computational Linguistics.

Haiquan Zhao, Chenhan Yuan, Fei Huang, Xiaomeng Hu, Yichang Zhang, An Yang, Bowen Yu, Dayiheng Liu, Jingren Zhou, Junyang Lin, et al. 2025. Qwen3Guard technical report. *arXiv preprint arXiv:2510.14276*.

A Calibration–Policy Identity Supporting Proof

Proof of Proposition 1. For $\mu(\tau) > 0$, the event $\{S < \tau\}$ is measurable with respect to S. Iterated expectation gives

$$\begin{aligned} R(\tau) &= \mathbb{E}[Y \mid S < \tau] \\ &= \mathbb{E}[m(S) \mid S < \tau] \\ &= \mathbb{E}[S \mid S < \tau] + \mathbb{E}[m(S) - S \mid S < \tau]. \end{aligned}$$

The first term is at most τ, and the second is at most

$$\begin{aligned} &\mathbb{E}[|m(S) - S| \mid S < \tau] \\ &\quad = \frac{\mathbb{E}[|m(S) - S| \mathbf{1}\{S < \tau\}]}{\mu(\tau)}. \end{aligned}$$

This proves the bound. Under exact calibration $m(S) = S$ almost surely, the discrepancy vanishes and $R(\tau) \leq \tau$. □

B Detailed Experimental Protocol

B.1 Model Configurations and Supervision

System One APIs and Specialists. The evaluated proprietary interfaces are `typesafe/jev-1.13`, `respan/span-01`, `upstage/solar-decide`, and `openai/gpt-6-luna-decisions`. The filtering specialist relies on `voyageai/rerank-2.5-lite`, whose relevance score is clipped to $[0, 1]$.

Local Open-Weights Models. Kev-base initializes from `jaredpalmer/kev-0.8b`. Kev-ft configuration warm-starts its rank-16 LoRA adapters and 256-dimensional pointer head on a frozen `Qwen/Qwen3.5-0.8B-Base` backbone. Training uses two epochs, seed 0, AdamW with weight decay 0.01, and a OneCycle schedule with peak learning rate 2×10^{-5} for both adapters and head and a 10% increasing phase. The microbatch is 2 with four accumulation steps on one device (effective batch 8); gradient norms are clipped at 1.0. Computation uses bfloat16 with float32 trainable weights, gradient checkpointing, and a maximum state length of 2,048 tokens. The wrapper truncates inputs to fit the training context. The final checkpoint after two epochs is used.

The pooled training data contains 9,988 records: 1,600 each from WildGuardMix, S-Labs, SciDocs, StackOverflow, and BeaverTails; 1,588 from Aegis; and 400 from ATBench. Thus adaptation also uses BeaverTails, which is not an evaluation benchmark in this paper. The generator draws from fit pools, deduplicates exact rendered states across sources, and reserves 20% per source for a separate held-out diagnostic.

PCA-Ridge Baseline. This baseline extracts instruction-templated embeddings using `qwen/qwen3-embedding-8b`, applies L2 normalization and a 100-dimensional PCA projection, and fits a logistic regression classifier ($C = 1.0$, solver `lbfgs`, and 2,000 maximum iterations). Both PCA and the classifier are fit independently on each dataset's specific fit split.

Qwen3Guard Score Extraction. Qwen3Guard (`Qwen/Qwen3Guard-Gen-0.6B`) emits specific categorical tokens rather than direct probabilities. For moderation tasks, we compute the score as the full-vocabulary softmax mass on Unsafe token variants following the `Safety:` marker.

B.2 Data Splitting and Evaluation Roles

We partition datasets into disjoint Fit, Calibration, and Test pools to strictly separate model training, threshold selection, and final evaluation. Fit data is exclusively used for model fine-tuning (Kev-ft), baseline fitting (PCA-Ridge), and constructing the candidate threshold grid ($\mathcal{T}$). Calibration data (D_{cal}) is used strictly for threshold selection (LTT and Empirical), and Test data evaluates the resulting fixed policies.

Calibration and Test splits use label-stratified sampling. Safety/moderation tasks also group repeated prompts or texts during overlap screening and partitioning. For reranking tasks (SciDocs and StackOverflow), we construct balanced query–candidate pairs, fit and evaluation pools are separated by query ID, and select one candidate per query using a deterministic query hash to choose a relevant or irrelevant candidate. This avoids counting multiple candidates for one query as independent rows.

B.3 Input Formatting Specifications

Input rendering depends on the task domain. Moderation queries evaluate solely the user's prompt. For safety evaluation (ATBench), the model evaluates serialized agent trajectories encompassing actions, disclosures, and environmental thoughts. To maintain context limits, these trajectories are bounded to 600 characters for tool names, 1,200 characters per turn, and 12,000 characters overall (with truncation markers appended where applicable).

B.4 LTT Implementation Details

The candidate threshold family $\mathcal{T}$ is constructed from up to 200 sorted fit-score midpoints on D_{fit}, plus the endpoints $\{0, 1\}$. Because $\mathcal{T}$ is defined strictly on D_{fit}, the candidate testing sequence is fixed independently of D_{cal} (Angelopoulos et al., 2025, Section 2.3.1). Under the rule $s(X) < \tau$, a threshold of 0 selects an empty set, while 1 selects all scores strictly below one.

Following the selective classification protocol of Angelopoulos et al. (2025, Section 3.2), we compute p-values on the selected calibration subset using the exact binomial tail bound. Even when zero errors are observed ($K_j = 0$), the binomial p-value reduces to $p_j = (1-\alpha)^{N_j}$. Satisfying the certification condition $p_j \leq \delta$ therefore requires $(1-\alpha)^{N_j} \leq \delta$; taking logarithms and dividing by $\log(1-\alpha) < 0$ reverses the inequality and yields the minimum certifiable support size:

$$N_j \geq \left\lceil \frac{\log \delta}{\log(1-\alpha)} \right\rceil. \tag{2}$$

At $\delta = 0.10$, Equation 2 requires at least 45 selected calibration examples for $\alpha = 0.05$, and 22 for $\alpha = 0.10$. Following the low-support clipping convention of Angelopoulos et al. (2025, Section 3.2) ($n(\mathcal{T}_\lambda) \geq 25$), we screen out candidates with calibration support below $\max(25, \lceil \log \delta / \log(1-\alpha) \rceil)$ so that uncertifiable low-support thresholds near $\tau = 0$ do not prematurely halt fixed-sequence testing. Our Empirical baseline ($\widehat{R}_{\text{cal},j} \leq \alpha$) follows Angelopoulos et al. (2025, Supplement Section 1.1) on the same grid with a support floor of 25.

B.5 Metrics and Uncertainty Quantification

Performance metrics are defined standardly based on the confusion matrix elements (TN, FN, FP, TP). The target risk is the False Omission Rate: $\text{FOR} = \text{FN}/(\text{TN}+\text{FN})$. The primary utility metric is Benign Acceptance (TNR) for safety tasks and Irrelevant Removal for filtering tasks, computed as $\text{TN}/(\text{TN}+\text{FP})$. Positive excess risk is computed in percentage points as $e = \max(\widehat{\tilde{R}} - \alpha, 0)$.

Calibration is measured via the Brier Score ($\widehat{\text{BS}}$) and Expected Calibration Error ($\widehat{\text{ECE}}$). For a test sample of size n, partitioned into 10 equal-width bins b with respective sizes $|I_b|$, mean scores $\bar{s}_b$, and positive label frequencies $\bar{y}_b$:

$$\widehat{\text{ECE}} = \sum_b \frac{|I_b|}{n} |\bar{y}_b - \bar{s}_b|,$$

$$\widehat{\text{BS}} = \frac{1}{n} \sum_i (s_i - Y_i)^2.$$

Uncertainty is quantified using 1,000 unstratified bootstrap resamples (seed 42) applied to the test rows, evaluating fixed thresholds and models. Confidence intervals are reported using the 2.5th and 97.5th percentiles via linear interpolation.

B.6 Resource Accounting

Cost and latency metrics trace the fit, calibration, and test inference workloads. Local inference is estimated at a standard rate of $0.11 per GPU-hour, while proprietary API costs use respective provider accounting. Because API pricing structures can change, we note that all experiments and cost calculations were conducted between late September and early October 2026. Costs are normalized per 1,000 workload examples; latency is reported in milliseconds per example. These figures represent recorded historical overheads rather than strictly controlled, hardware-synchronized microbenchmarks, and are intended to provide descriptive operational context.

B.7 Artifact Access and Terms

The public dataset cards list WildGuardMix under ODC-BY (with AI2 Responsible Use Guidelines), Aegis 1.0 under CC-BY-4.0, S-Labs under MIT, and ATBench under Apache-2.0. The additional training source BeaverTails lists CC-BY-NC-4.0. SciDocs and StackOverflow derivative cards did not expose an explicit license declaration; this is a documentation gap, not permission for unrestricted redistribution. Original-source and derivative terms must be considered.

The open-weights models (Qwen3Guard-Gen-0.6B, Qwen3-Embedding-8B, and Kev-0.8B) are released under the Apache-2.0 license, while the proprietary endpoints (Jev, GPT-6-Luna, Span-01, Solar-Decide, and Voyage Rerank 2.5 Lite) were accessed under their respective providers' standard commercial terms of service. All experiments were conducted strictly as offline benchmark evaluations, which is consistent with the intended research use of these artifacts and respects model-card restrictions against unattended high-stakes deployment.

C Additional Score and Policy Results

C.1 Score Diagnostics with 95% Intervals

Task / Dataset	Model	ECE ↓	Brier ↓	AUROC ↑	AP_0 ↑
Moderation	Qwen3Guard	0.085 [0.068, 0.106]	0.090 [0.073, 0.108]	0.950 [0.935, 0.965]	0.947 [0.924, 0.967]
WildGuardMix	PCA-Ridge*	0.042 [0.034, 0.067]	0.091 [0.078, 0.103]	0.946 [0.931, 0.959]	0.956 [0.940, 0.968]
	Jev	0.034 [0.028, 0.056]	0.073 [0.063, 0.084]	0.963 [0.951, 0.973]	0.969 [0.957, 0.979]
	GPT-6-Luna	0.103 [0.084, 0.125]	0.108 [0.089, 0.126]	0.925 [0.906, 0.942]	0.907 [0.880, 0.931]
	Span-01	0.189 [0.166, 0.224]	0.238 [0.216, 0.261]	0.765 [0.732, 0.799]	0.746 [0.701, 0.795]
	Solar-Decide	0.104 [0.081, 0.128]	0.130 [0.108, 0.150]	0.893 [0.867, 0.917]	0.855 [0.814, 0.900]
	Kev-base	0.121 [0.097, 0.149]	0.186 [0.176, 0.196]	0.835 [0.809, 0.862]	0.863 [0.835, 0.891]
	Kev-ft*	0.119 [0.099, 0.141]	0.123 [0.104, 0.143]	0.934 [0.917, 0.950]	0.938 [0.915, 0.957]
Moderation	Qwen3Guard	0.277 [0.244, 0.309]	0.239 [0.208, 0.270]	0.929 [0.907, 0.948]	0.866 [0.822, 0.906]
Aegis	PCA-Ridge*	0.037 [0.033, 0.069]	0.096 [0.081, 0.111]	0.934 [0.914, 0.953]	0.892 [0.857, 0.925]
	Jev	0.183 [0.154, 0.211]	0.151 [0.131, 0.173]	0.936 [0.915, 0.955]	0.860 [0.809, 0.904]
	GPT-6-Luna	0.297 [0.264, 0.334]	0.277 [0.243, 0.311]	0.864 [0.834, 0.889]	0.665 [0.605, 0.724]
	Span-01	0.500 [0.461, 0.541]	0.472 [0.437, 0.505]	0.565 [0.522, 0.612]	0.332 [0.293, 0.384]
	Solar-Decide	0.356 [0.317, 0.392]	0.339 [0.303, 0.372]	0.865 [0.835, 0.894]	0.658 [0.591, 0.736]
	Kev-base	0.257 [0.229, 0.293]	0.232 [0.217, 0.246]	0.834 [0.796, 0.867]	0.704 [0.635, 0.776]
	Kev-ft*	0.109 [0.087, 0.141]	0.126 [0.103, 0.151]	0.919 [0.895, 0.940]	0.849 [0.796, 0.895]
Safety	Qwen3Guard	0.474 [0.447, 0.502]	0.464 [0.436, 0.492]	0.968 [0.959, 0.975]	0.968 [0.959, 0.976]
S-Labs	PCA-Ridge*	0.087 [0.077, 0.098]	0.041 [0.035, 0.047]	0.996 [0.994, 0.998]	0.996 [0.994, 0.998]
	Jev	0.185 [0.166, 0.205]	0.158 [0.144, 0.172]	0.947 [0.933, 0.960]	0.927 [0.907, 0.947]
	GPT-6-Luna	0.262 [0.241, 0.285]	0.228 [0.208, 0.249]	0.888 [0.869, 0.908]	0.841 [0.811, 0.871]
	Span-01	0.441 [0.414, 0.469]	0.437 [0.413, 0.464]	0.288 [0.254, 0.320]	0.368 [0.342, 0.395]
	Solar-Decide	0.337 [0.310, 0.364]	0.326 [0.301, 0.350]	0.840 [0.815, 0.866]	0.737 [0.694, 0.785]
	Kev-base	0.225 [0.209, 0.249]	0.203 [0.199, 0.207]	0.944 [0.929, 0.957]	0.940 [0.920, 0.958]
	Kev-ft*	0.024 [0.018, 0.037]	0.028 [0.020, 0.037]	0.995 [0.993, 0.998]	0.995 [0.992, 0.998]
Safety	PCA-Ridge*	0.151 [0.118, 0.201]	0.165 [0.151, 0.181]	0.882 [0.838, 0.921]	0.902 [0.855, 0.937]
ATBench	Jev	0.076 [0.064, 0.141]	0.200 [0.180, 0.221]	0.768 [0.708, 0.824]	0.717 [0.633, 0.810]
	GPT-6-Luna	0.230 [0.181, 0.285]	0.232 [0.191, 0.272]	0.842 [0.794, 0.886]	0.844 [0.776, 0.902]
	Span-01	0.101 [0.073, 0.163]	0.196 [0.170, 0.223]	0.779 [0.721, 0.833]	0.744 [0.661, 0.819]
	Solar-Decide	0.274 [0.219, 0.337]	0.288 [0.246, 0.333]	0.688 [0.622, 0.756]	0.716 [0.640, 0.796]
	Kev-base	0.255 [0.197, 0.318]	0.303 [0.278, 0.330]	0.457 [0.384, 0.531]	0.538 [0.454, 0.634]
	Kev-ft*	0.059 [0.044, 0.105]	0.099 [0.073, 0.128]	0.938 [0.904, 0.963]	0.952 [0.924, 0.973]
Reranking	Reranker-Lite	0.183 [0.163, 0.208]	0.214 [0.206, 0.221]	0.880 [0.859, 0.900]	0.861 [0.831, 0.889]
SciDocs	PCA-Ridge*	0.065 [0.053, 0.088]	0.127 [0.117, 0.138]	0.904 [0.886, 0.922]	0.885 [0.857, 0.911]
	Jev	0.290 [0.265, 0.313]	0.236 [0.216, 0.257]	0.917 [0.901, 0.932]	0.895 [0.869, 0.919]
	GPT-6-Luna	0.340 [0.314, 0.365]	0.289 [0.265, 0.313]	0.865 [0.844, 0.884]	0.787 [0.754, 0.819]
	Span-01	0.175 [0.151, 0.199]	0.189 [0.173, 0.204]	0.864 [0.841, 0.884]	0.834 [0.798, 0.871]
	Solar-Decide	0.357 [0.330, 0.383]	0.333 [0.307, 0.358]	0.772 [0.741, 0.803]	0.643 [0.604, 0.688]
	Kev-base	0.120 [0.099, 0.150]	0.234 [0.227, 0.240]	0.734 [0.704, 0.765]	0.724 [0.684, 0.765]
	Kev-ft*	0.118 [0.099, 0.141]	0.131 [0.112, 0.150]	0.913 [0.894, 0.930]	0.881 [0.846, 0.913]
Reranking	Reranker-Lite	0.137 [0.116, 0.163]	0.200 [0.193, 0.208]	0.829 [0.801, 0.854]	0.826 [0.792, 0.861]
StackOverflow	PCA-Ridge*	0.032 [0.025, 0.066]	0.204 [0.194, 0.215]	0.750 [0.719, 0.779]	0.748 [0.705, 0.787]
	Jev	0.099 [0.077, 0.126]	0.196 [0.181, 0.214]	0.790 [0.759, 0.814]	0.800 [0.765, 0.831]
	GPT-6-Luna	0.120 [0.102, 0.148]	0.191 [0.173, 0.212]	0.815 [0.787, 0.840]	0.819 [0.785, 0.851]
	Span-01	0.224 [0.200, 0.257]	0.302 [0.284, 0.321]	0.607 [0.572, 0.639]	0.576 [0.533, 0.622]
	Solar-Decide	0.266 [0.241, 0.298]	0.288 [0.265, 0.315]	0.706 [0.671, 0.737]	0.643 [0.601, 0.690]
	Kev-base	0.051 [0.029, 0.081]	0.245 [0.240, 0.251]	0.599 [0.561, 0.634]	0.559 [0.518, 0.607]
	Kev-ft*	0.160 [0.138, 0.189]	0.205 [0.185, 0.227]	0.810 [0.782, 0.837]	0.818 [0.782, 0.850]

Table 6: Test score diagnostics with 95% intervals [lower, upper]. *: supervised adaptation.

C.2 Risk and Utility at Every Budget, with 95% Intervals

These tables extend the main Table 4 layout to all α with its intervals. A zero-width interval after no errors does not establish zero population risk.

		Moderation: WildGuardMix			Moderation: Aegis		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	9.2 [6,13]	6.2 [3,9]	3.3 [1,7]	32.4 [26,39]	5.9 [1,11]	ND
	Utility	91.3 [88,94]	72.9 [69,78]	42.7 [38,48]	89.5 [85,94]	50.3 [43,58]	ND
PCA-Ridge*	Risk	0.0 [0,0]	5.1 [3,8]	3.2 [1,6]	UD	8.7 [3,14]	ND
	Utility	18.8 [15,23]	70.6 [66,75]	63.0 [58,68]	UD	60.7 [53,68]	ND
Jev	Risk	0.0 [0,0]	6.6 [4,10]	4.4 [2,7]	7.7 [1,16]	7.7 [1,16]	ND
	Utility	37.6 [33,42]	83.3 [80,87]	77.8 [74,82]	25.1 [18,32]	25.1 [18,32]	ND
GPT-6-Luna	Risk	11.1 [8,15]	NF	ND	40.3 [34,46]	NF	ND
	Utility	90.1 [87,93]	NF	ND	93.2 [89,97]	NF	ND
Span-01	Risk	0.0 [0,0]	NF	ND	78.3 [59,94]	NF	ND
	Utility	0.8 [0,2]	NF	ND	2.6 [0,6]	NF	ND
Solar-Decide	Risk	11.4 [8,15]	NF	ND	39.0 [33,45]	NF	ND
	Utility	69.3 [65,74]	NF	ND	88.5 [83,94]	NF	ND
Kev-base	Risk	UD	9.3 [4,15]	ND	UD	NF	ND
	Utility	UD	29.0 [25,34]	ND	UD	NF	ND
Kev-ft*	Risk	9.3 [6,13]	5.3 [2,8]	4.7 [2,8]	14.2 [8,21]	10.0 [4,17]	ND
	Utility	78.4 [74,83]	64.3 [60,69]	60.5 [56,66]	63.4 [56,70]	51.8 [44,59]	ND

		Safety: S-Labs			Safety: ATBench		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	47.2 [44,51]	6.3 [4,9]	3.1 [1,5]	–	–	–
	Utility	100.0 [100,100]	82.5 [79,86]	77.3 [73,81]	–	–	–
PCA-Ridge*	Risk	0.0 [0,0]	4.3 [2,7]	3.0 [1,5]	UD	6.8 [2,13]	1.9 [0,7]
	Utility	62.7 [58,67]	97.7 [96,99]	97.1 [95,99]	UD	62.6 [54,71]	39.7 [31,48]
Jev	Risk	0.8 [0,3]	5.7 [2,10]	3.7 [1,7]	UD	NF	ND
	Utility	25.0 [21,29]	31.4 [27,36]	29.9 [26,34]	UD	NF	ND
GPT-6-Luna	Risk	22.8 [19,27]	NF	ND	4.8 [0,17]	10.8 [4,19]	ND
	Utility	88.6 [85,92]	NF	ND	15.3 [9,22]	50.4 [41,60]	ND
Span-01	Risk	45.5 [42,49]	NF	ND	UD	NF	ND
	Utility	99.8 [99,100]	NF	ND	UD	NF	ND
Solar-Decide	Risk	26.1 [22,30]	NF	ND	25.4 [14,37]	NF	ND
	Utility	90.1 [87,93]	NF	ND	38.2 [29,47]	NF	ND
Kev-base	Risk	UD	5.7 [3,9]	ND	UD	NF	ND
	Utility	UD	72.0 [68,76]	ND	UD	NF	ND
Kev-ft*	Risk	2.7 [1,5]	4.1 [2,6]	3.0 [1,5]	1.2 [0,4]	7.1 [2,12]	ND
	Utility	96.0 [94,98]	97.5 [96,99]	97.0 [95,99]	63.4 [54,72]	80.2 [72,87]	ND

		Reranking: SciDocs			Reranking: StackOverflow		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Reranker-Lite	Risk	UD	NF	ND	UD	4.8 [0,11]	ND
	Utility	UD	NF	ND	UD	11.8 [9,15]	ND
PCA-Ridge*	Risk	UD	7.0 [3,12]	4.1 [0,11]	0.0 [0,0]	NF	ND
	Utility	UD	29.9 [25,35]	9.6 [6,13]	0.2 [0,1]	NF	ND
Jev	Risk	9.8 [6,13]	NF	ND	2.1 [0,7]	3.1 [0,8]	ND
	Utility	69.5 [65,74]	NF	ND	9.2 [6,12]	18.6 [15,22]	ND
GPT-6-Luna	Risk	25.6 [22,30]	NF	ND	8.1 [4,13]	5.2 [1,11]	ND
	Utility	96.3 [94,98]	NF	ND	29.5 [25,34]	18.4 [14,22]	ND
Span-01	Risk	0.0 [0,0]	NF	ND	40.6 [23,58]	NF	ND
	Utility	1.0 [0,3]	NF	ND	3.8 [2,6]	NF	ND
Solar-Decide	Risk	32.5 [29,37]	NF	ND	30.2 [24,36]	NF	ND
	Utility	91.9 [89,95]	NF	ND	38.5 [33,43]	NF	ND
Kev-base	Risk	UD	NF	ND	UD	NF	ND
	Utility	UD	NF	ND	UD	NF	ND
Kev-ft*	Risk	12.7 [9,17]	7.7 [4,12]	ND	11.8 [8,16]	7.7 [3,13]	ND
	Utility	73.0 [69,77]	46.1 [41,51]	ND	43.3 [38,48]	24.0 [20,28]	ND

Table 7: Test risk and utility at $\alpha = 0.05$ (%). Shaded Risk rows report FOR; white Utility rows report TNR. Entries are estimates [95% CI], with bounds rounded outward to whole percentage points. Red risk exceeds the budget. ND/NF: no returned policy; UD: undefined FOR (zero TNR). –: not evaluated. *: supervised adaptation.

		Moderation: WildGuardMix			Moderation: Aegis		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	9.7 [7,13]	9.4 [7,13]	7.8 [5,11]	36.4 [30,42]	11.5 [6,17]	5.9 [1,11]
	Utility	94.1 [91,97]	93.2 [90,96]	87.9 [85,91]	91.6 [87,96]	60.7 [53,68]	50.3 [43,58]
PCA-Ridge*	Risk	2.0 [0,4]	9.8 [7,13]	8.3 [5,12]	0.0 [0,0]	9.5 [4,15]	ND
	Utility	51.6 [46,57]	83.7 [80,88]	79.7 [76,84]	14.7 [9,20]	64.9 [58,72]	ND
Jev	Risk	3.3 [1,6]	11.9 [9,15]	10.5 [7,14]	11.7 [6,18]	11.5 [6,18]	ND
	Utility	56.4 [51,61]	95.3 [93,98]	93.4 [90,96]	67.0 [60,75]	60.7 [53,68]	ND
GPT-6-Luna	Risk	11.8 [9,15]	10.9 [8,14]	8.1 [5,11]	41.6 [35,48]	NF	ND
	Utility	91.8 [89,95]	89.9 [87,93]	86.0 [82,90]	94.8 [91,98]	NF	ND
Span-01	Risk	22.1 [17,27]	NF	ND	58.9 [53,64]	NF	ND
	Utility	62.6 [58,68]	NF	ND	91.1 [86,96]	NF	ND
Solar-Decide	Risk	15.7 [12,20]	NF	ND	47.7 [42,53]	NF	ND
	Utility	89.4 [86,93]	NF	ND	95.8 [92,99]	NF	ND
Kev-base	Risk	0.0 [0,0]	8.3 [4,13]	3.1 [0,8]	UD	12.8 [2,24]	ND
	Utility	0.2 [0,1]	39.5 [35,45]	19.7 [16,24]	UD	17.8 [12,24]	ND
Kev-ft*	Risk	10.7 [7,14]	8.4 [5,12]	8.0 [5,11]	17.2 [11,24]	14.5 [9,21]	ND
	Utility	81.4 [78,85]	75.7 [72,80]	73.4 [69,78]	68.1 [61,75]	64.9 [58,72]	ND

		Safety: S-Labs			Safety: ATBench		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	47.6 [44,51]	11.3 [8,14]	9.0 [6,12]	–	–	–
	Utility	100.0 [100,100]	91.6 [89,94]	88.6 [85,92]	–	–	–
PCA-Ridge*	Risk	0.2 [0,1]	8.0 [5,11]	6.2 [4,9]	UD	10.1 [4,17]	7.8 [3,14]
	Utility	80.6 [77,84]	98.7 [97,100]	98.5 [97,100]	UD	67.9 [59,76]	63.4 [54,72]
Jev	Risk	7.4 [4,11]	8.5 [5,13]	7.6 [4,12]	UD	NF	ND
	Utility	43.0 [38,47]	51.2 [46,56]	37.1 [32,42]	UD	NF	ND
GPT-6-Luna	Risk	24.8 [21,29]	NF	ND	6.9 [0,18]	12.0 [5,20]	ND
	Utility	92.4 [90,95]	NF	ND	20.6 [13,28]	55.7 [47,65]	ND
Span-01	Risk	47.3 [44,51]	NF	ND	0.0 [0,0]	NF	ND
	Utility	100.0 [100,100]	NF	ND	2.3 [0,6]	NF	ND
Solar-Decide	Risk	37.4 [34,41]	NF	ND	32.2 [22,43]	NF	ND
	Utility	99.6 [99,100]	NF	ND	45.0 [36,55]	NF	ND
Kev-base	Risk	UD	8.7 [6,12]	ND	UD	NF	ND
	Utility	UD	79.8 [76,84]	ND	UD	NF	ND
Kev-ft*	Risk	3.0 [1,5]	10.0 [7,13]	7.5 [5,10]	3.2 [0,7]	14.6 [8,21]	0.0 [0,0]
	Utility	97.0 [95,99]	99.4 [98,100]	99.2 [98,100]	69.5 [60,77]	89.3 [84,95]	13.0 [7,20]

		Reranking: SciDocs			Reranking: StackOverflow		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Reranker-Lite	Risk	UD	6.0 [2,10]	ND	UD	10.7 [6,15]	8.3 [0,20]
	Utility	UD	28.7 [24,33]	ND	UD	43.7 [39,49]	4.4 [2,7]
PCA-Ridge*	Risk	2.5 [0,9]	12.6 [9,16]	5.6 [1,12]	0.0 [0,0]	10.7 [0,23]	ND
	Utility	7.9 [5,11]	64.6 [60,69]	13.8 [10,18]	1.8 [0,4]	5.0 [3,8]	ND
Jev	Risk	17.9 [14,22]	11.9 [8,15]	7.1 [4,10]	3.1 [0,8]	16.5 [11,22]	4.9 [0,10]
	Utility	84.6 [81,88]	73.8 [69,78]	58.5 [54,64]	18.6 [15,22]	38.7 [34,43]	19.4 [15,23]
GPT-6-Luna	Risk	28.2 [24,32]	NF	ND	9.0 [5,14]	8.9 [5,14]	5.2 [1,11]
	Utility	98.2 [96,100]	NF	ND	36.3 [31,41]	35.1 [30,40]	18.4 [14,22]
Span-01	Risk	9.7 [6,14]	9.4 [5,15]	ND	41.0 [35,47]	NF	ND
	Utility	43.3 [39,48]	31.5 [27,36]	ND	32.9 [28,38]	NF	ND
Solar-Decide	Risk	37.5 [34,41]	NF	ND	31.8 [27,36]	NF	ND
	Utility	96.3 [94,98]	NF	ND	61.1 [56,66]	NF	ND
Kev-base	Risk	UD	13.8 [5,23]	ND	UD	NF	ND
	Utility	UD	10.2 [7,13]	ND	UD	NF	ND
Kev-ft*	Risk	13.6 [10,17]	11.5 [8,15]	9.5 [6,13]	13.4 [9,18]	12.2 [8,17]	9.6 [5,14]
	Utility	77.2 [73,81]	62.8 [58,68]	52.4 [48,57]	50.3 [45,55]	44.9 [40,50]	34.1 [29,39]

Table 8: Test risk and utility at $\alpha = 0.10$ (%). Shaded Risk rows report FOR; white Utility rows report TNR. Entries are estimates [95% CI], with bounds rounded outward to whole percentage points. Red risk exceeds the budget. ND/NF: no returned policy; UD: undefined FOR (zero TNR). –: not evaluated. *: supervised adaptation.

		Moderation: WildGuardMix			Moderation: Aegis		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	10.8 [8,14]	12.2 [9,16]	11.4 [8,15]	38.4 [32,44]	18.2 [12,24]	11.5 [6,17]
	Utility	94.3 [92,97]	95.8 [93,98]	95.1 [93,98]	93.2 [89,97]	72.8 [66,79]	60.7 [53,68]
PCA-Ridge*	Risk	4.2 [2,7]	13.1 [9,16]	11.7 [8,15]	1.4 [0,5]	17.3 [11,24]	9.5 [4,15]
	Utility	67.4 [63,72]	91.3 [88,94]	87.9 [84,91]	36.1 [29,43]	72.8 [66,80]	64.9 [58,72]
Jev	Risk	3.3 [1,6]	16.2 [13,20]	12.8 [9,16]	20.8 [15,27]	23.1 [17,30]	11.7 [6,18]
	Utility	68.7 [64,74]	98.1 [96,100]	96.8 [95,99]	83.8 [78,89]	85.3 [80,91]	67.0 [60,75]
GPT-6-Luna	Risk	12.6 [9,16]	15.6 [12,19]	12.9 [9,16]	43.4 [37,49]	NF	ND
	Utility	92.6 [90,95]	96.0 [94,98]	92.6 [90,95]	94.8 [91,98]	NF	ND
Span-01	Risk	26.3 [22,30]	NF	ND	60.9 [56,66]	NF	ND
	Utility	81.2 [77,85]	NF	ND	95.3 [92,99]	NF	ND
Solar-Decide	Risk	18.1 [14,22]	13.9 [10,18]	ND	50.7 [45,56]	NF	ND
	Utility	93.9 [91,96]	84.8 [81,89]	ND	97.9 [95,100]	NF	ND
Kev-base	Risk	0.0 [0,0]	15.4 [11,20]	8.3 [4,13]	12.5 [0,30]	18.6 [9,29]	ND
	Utility	4.0 [2,7]	62.6 [58,68]	41.9 [37,47]	7.3 [3,12]	29.8 [22,37]	ND
Kev-ft*	Risk	11.3 [8,15]	13.4 [10,17]	11.4 [8,15]	17.7 [12,24]	19.5 [13,26]	14.5 [9,21]
	Utility	82.9 [79,87]	87.5 [84,91]	82.5 [79,86]	70.7 [64,77]	71.2 [64,78]	64.9 [58,72]

		Safety: S-Labs			Safety: ATBench		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	47.8 [44,51]	14.7 [12,18]	13.5 [10,17]	–	–	–
	Utility	100.0 [100,100]	94.1 [92,96]	93.1 [90,96]	–	–	–
PCA-Ridge*	Risk	0.4 [0,2]	11.6 [9,15]	10.0 [7,13]	0.0 [0,0]	13.6 [7,21]	10.1 [4,17]
	Utility	87.0 [84,90]	99.8 [99,100]	99.8 [99,100]	1.5 [0,5]	72.5 [64,80]	67.9 [59,76]
Jev	Risk	11.2 [8,14]	13.9 [11,17]	10.2 [7,13]	33.3 [0,84]	NF	ND
	Utility	86.1 [83,90]	95.4 [93,98]	79.0 [75,83]	3.1 [0,7]	NF	ND
GPT-6-Luna	Risk	27.0 [23,31]	13.2 [10,17]	ND	5.6 [0,15]	16.3 [9,24]	12.0 [5,20]
	Utility	94.3 [92,97]	71.4 [67,76]	ND	26.0 [18,34]	62.6 [54,72]	55.7 [47,65]
Span-01	Risk	47.5 [44,51]	NF	ND	27.3 [10,48]	NF	ND
	Utility	100.0 [100,100]	NF	ND	12.2 [7,19]	NF	ND
Solar-Decide	Risk	41.0 [38,45]	NF	ND	28.6 [20,37]	NF	ND
	Utility	100.0 [100,100]	NF	ND	61.1 [52,70]	NF	ND
Kev-base	Risk	UD	14.6 [11,18]	ND	UD	NF	ND
	Utility	UD	88.2 [85,91]	ND	UD	NF	ND
Kev-ft*	Risk	3.6 [2,6]	12.5 [10,16]	11.1 [8,14]	5.1 [1,10]	15.7 [10,22]	14.6 [8,21]
	Utility	97.1 [95,99]	99.6 [99,100]	99.6 [99,100]	71.0 [63,79]	90.1 [84,95]	89.3 [84,95]

		Reranking: SciDocs			Reranking: StackOverflow		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Reranker-Lite	Risk	UD	11.3 [7,15]	ND	UD	18.5 [14,23]	14.0 [10,19]
	Utility	UD	56.1 [51,61]	ND	UD	57.3 [52,62]	51.7 [47,57]
PCA-Ridge*	Risk	7.6 [3,13]	17.2 [13,21]	13.9 [10,18]	8.8 [0,20]	16.5 [10,24]	ND
	Utility	24.6 [20,29]	78.0 [74,82]	72.8 [68,77]	6.2 [3,9]	23.2 [19,28]	ND
Jev	Risk	22.8 [19,27]	18.8 [15,23]	13.7 [10,18]	8.1 [3,14]	20.5 [16,25]	16.6 [11,22]
	Utility	90.9 [88,94]	85.8 [82,89]	76.8 [73,81]	25.1 [21,29]	54.3 [49,59]	44.3 [39,49]
GPT-6-Luna	Risk	31.5 [28,35]	NF	ND	10.1 [6,15]	15.9 [11,20]	11.3 [7,16]
	Utility	99.0 [97,100]	NF	ND	40.9 [36,45]	55.3 [51,60]	42.5 [37,47]
Span-01	Risk	16.4 [12,21]	15.9 [12,20]	ND	41.2 [36,46]	NF	ND
	Utility	65.2 [60,70]	63.4 [58,68]	ND	47.5 [43,52]	NF	ND
Solar-Decide	Risk	40.1 [36,44]	NF	ND	34.6 [30,39]	NF	ND
	Utility	98.0 [96,100]	NF	ND	73.3 [69,77]	NF	ND
Kev-base	Risk	UD	19.4 [13,26]	ND	UD	NF	ND
	Utility	UD	25.4 [21,30]	ND	UD	NF	ND
Kev-ft*	Risk	15.2 [12,19]	15.9 [12,20]	14.8 [11,19]	16.7 [12,21]	16.2 [11,21]	13.4 [9,18]
	Utility	79.5 [75,84]	83.9 [80,88]	78.5 [74,82]	54.9 [50,60]	53.9 [49,59]	50.3 [45,55]

Table 9: Test risk and utility at $\alpha = 0.15$ (%). Shaded Risk rows report FOR; white Utility rows report TNR. Entries are estimates [95% CI], with bounds rounded outward to whole percentage points. Red risk exceeds the budget. ND/NF: no returned policy; UD: undefined FOR (zero TNR). –: not evaluated. *: supervised adaptation.

		Moderation: WildGuardMix			Moderation: Aegis		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	11.3 [8,15]	16.7 [13,20]	14.0 [11,17]	40.5 [34,47]	28.6 [22,35]	19.4 [13,26]
	Utility	94.9 [92,97]	97.0 [95,99]	96.4 [94,99]	93.7 [90,98]	86.4 [81,91]	73.8 [67,80]
PCA-Ridge*	Risk	5.6 [3,9]	17.9 [14,21]	15.7 [12,19]	5.6 [1,11]	23.8 [17,30]	17.4 [11,24]
	Utility	74.6 [70,79]	96.2 [94,98]	93.9 [91,97]	52.9 [45,61]	80.6 [74,87]	74.3 [68,81]
Jev	Risk	3.8 [2,6]	18.4 [15,22]	17.8 [14,22]	25.6 [19,32]	25.8 [20,32]	23.5 [18,30]
	Utility	74.2 [70,79]	98.3 [97,100]	98.3 [97,100]	86.9 [82,92]	87.4 [82,93]	85.3 [80,91]
GPT-6-Luna	Risk	13.2 [10,16]	20.6 [17,24]	17.8 [14,21]	44.5 [38,50]	NF	ND
	Utility	93.4 [91,96]	96.8 [95,99]	96.4 [94,98]	94.8 [91,98]	NF	ND
Span-01	Risk	29.6 [25,34]	15.4 [0,39]	ND	61.3 [56,66]	NF	ND
	Utility	87.1 [84,91]	2.3 [1,4]	ND	96.9 [94,100]	NF	ND
Solar-Decide	Risk	18.7 [15,22]	18.5 [15,22]	ND	53.6 [48,59]	NF	ND
	Utility	96.4 [94,99]	95.8 [93,98]	ND	98.4 [96,100]	NF	ND
Kev-base	Risk	5.6 [1,11]	22.3 [18,26]	17.2 [13,21]	16.4 [7,27]	22.1 [13,31]	ND
	Utility	21.6 [18,26]	76.7 [72,81]	67.0 [62,72]	24.1 [18,31]	42.4 [35,50]	ND
Kev-ft*	Risk	11.3 [8,15]	19.0 [15,23]	14.8 [11,18]	19.0 [13,26]	24.7 [18,31]	19.5 [13,26]
	Utility	84.6 [81,88]	95.3 [93,98]	92.6 [90,95]	71.2 [64,78]	78.0 [71,84]	71.2 [64,78]

		Safety: S-Labs			Safety: ATBench		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	47.9 [45,51]	18.3 [15,22]	16.9 [14,20]	–	–	–
	Utility	100.0 [100,100]	97.1 [95,99]	96.4 [94,98]	–	–	–
PCA-Ridge*	Risk	1.2 [0,3]	16.4 [13,20]	15.0 [12,18]	0.0 [0,0]	14.5 [8,22]	13.6 [7,21]
	Utility	92.4 [90,95]	100.0 [100,100]	100.0 [100,100]	7.6 [3,13]	76.3 [69,84]	72.5 [64,80]
Jev	Risk	14.8 [12,18]	17.1 [14,20]	14.8 [12,18]	24.0 [7,43]	NF	ND
	Utility	98.5 [97,100]	99.6 [99,100]	98.5 [97,100]	14.5 [8,22]	NF	ND
GPT-6-Luna	Risk	28.1 [24,32]	13.2 [10,17]	13.2 [10,17]	7.3 [0,17]	20.8 [13,28]	14.7 [7,23]
	Utility	96.2 [94,98]	71.4 [67,76]	71.4 [67,76]	29.0 [21,38]	72.5 [64,81]	61.8 [54,71]
Span-01	Risk	48.0 [45,51]	NF	ND	23.8 [12,38]	NF	ND
	Utility	100.0 [100,100]	NF	ND	24.4 [17,32]	NF	ND
Solar-Decide	Risk	42.6 [39,46]	NF	ND	32.4 [25,41]	NF	ND
	Utility	100.0 [100,100]	NF	ND	71.8 [64,80]	NF	ND
Kev-base	Risk	UD	19.9 [16,24]	16.4 [13,20]	40.0 [0,100]	NF	ND
	Utility	UD	93.3 [91,96]	91.2 [88,94]	2.3 [0,6]	NF	ND
Kev-ft*	Risk	3.9 [2,6]	17.2 [14,21]	16.2 [13,20]	6.8 [2,12]	20.6 [14,28]	15.7 [10,22]
	Utility	97.3 [95,99]	99.8 [99,100]	99.8 [99,100]	73.3 [65,81]	93.9 [89,98]	90.1 [84,95]

		Reranking: SciDocs			Reranking: StackOverflow		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Reranker-Lite	Risk	UD	18.0 [14,22]	ND	UD	22.5 [18,27]	18.3 [14,23]
	Utility	UD	73.0 [68,77]	ND	UD	69.7 [65,74]	60.1 [55,65]
PCA-Ridge*	Risk	5.6 [2,9]	22.4 [19,26]	20.7 [17,24]	13.7 [6,23]	22.7 [17,28]	10.7 [0,23]
	Utility	41.3 [36,46]	88.0 [85,91]	84.8 [81,88]	12.6 [9,16]	47.1 [42,52]	5.0 [3,8]
Jev	Risk	26.8 [23,31]	21.4 [18,25]	19.9 [16,24]	11.2 [6,17]	24.0 [20,29]	21.8 [17,27]
	Utility	95.5 [93,98]	90.2 [87,93]	88.4 [85,92]	31.7 [27,36]	63.9 [59,68]	57.5 [53,62]
GPT-6-Luna	Risk	34.5 [31,38]	20.9 [17,25]	ND	12.8 [8,18]	20.6 [16,25]	17.8 [13,22]
	Utility	99.2 [98,100]	93.1 [90,96]	ND	45.1 [40,50]	67.3 [62,72]	60.3 [55,65]
Span-01	Risk	22.2 [18,26]	21.4 [17,26]	6.7 [0,18]	42.2 [37,47]	NF	ND
	Utility	74.6 [70,79]	73.2 [69,77]	5.7 [3,8]	56.9 [52,62]	NF	ND
Solar-Decide	Risk	42.0 [38,46]	NF	ND	37.2 [33,41]	NF	ND
	Utility	100.0 [100,100]	NF	ND	79.0 [75,83]	NF	ND
Kev-base	Risk	0.0 [0,0]	23.8 [18,29]	14.8 [3,29]	UD	NF	ND
	Utility	0.4 [0,2]	40.4 [35,45]	4.7 [2,7]	UD	NF	ND
Kev-ft*	Risk	15.3 [12,19]	18.6 [15,23]	16.9 [13,21]	17.7 [13,22]	21.8 [17,26]	18.1 [13,23]
	Utility	81.1 [77,85]	90.0 [87,93]	86.8 [83,90]	57.9 [53,63]	64.7 [60,69]	57.9 [53,63]

Table 10: Test risk and utility at $\alpha = 0.20$ (%). Shaded Risk rows report FOR; white Utility rows report TNR. Entries are estimates [95% CI], with bounds rounded outward to whole percentage points. Red risk exceeds the budget. ND/NF: no returned policy; UD: undefined FOR (zero TNR). –: not evaluated. *: supervised adaptation.

		Moderation: WildGuardMix			Moderation: Aegis		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	12.4 [9,16]	40.9 [37,45]	39.0 [35,43]	44.7 [38,51]	44.2 [38,50]	40.3 [34,46]
	Utility	96.0 [94,98]	99.8 [99,100]	99.8 [99,100]	96.3 [93,99]	95.8 [92,99]	93.2 [89,97]
PCA-Ridge*	Risk	10.8 [7,14]	40.1 [36,44]	37.7 [34,41]	13.7 [8,20]	42.0 [36,48]	39.5 [33,45]
	Utility	87.3 [84,91]	99.6 [98,100]	99.4 [98,100]	69.1 [62,76]	94.8 [91,98]	94.8 [91,98]
Jev	Risk	7.1 [4,10]	40.4 [36,45]	36.4 [32,41]	35.0 [29,41]	44.0 [38,50]	40.4 [34,46]
	Utility	87.9 [85,91]	100.0 [100,100]	100.0 [100,100]	95.3 [92,99]	98.4 [96,100]	97.9 [95,100]
GPT-6-Luna	Risk	14.9 [11,18]	31.9 [28,36]	31.9 [28,36]	47.1 [41,53]	45.4 [39,51]	40.6 [34,46]
	Utility	95.6 [93,98]	99.4 [98,100]	99.4 [98,100]	96.3 [93,99]	95.8 [92,99]	94.2 [90,98]
Span-01	Risk	34.5 [30,38]	38.9 [35,43]	36.9 [33,41]	63.4 [59,68]	NF	ND
	Utility	92.4 [89,95]	98.3 [97,100]	95.3 [93,98]	97.4 [94,100]	NF	ND
Solar-Decide	Risk	19.0 [15,23]	39.5 [36,44]	ND	54.1 [49,59]	42.8 [36,49]	ND
	Utility	96.4 [94,99]	98.5 [97,100]	ND	98.4 [96,100]	93.2 [89,97]	ND
Kev-base	Risk	22.8 [18,27]	39.0 [35,43]	37.2 [33,41]	44.7 [38,51]	41.5 [35,48]	37.3 [30,45]
	Utility	78.9 [75,83]	98.3 [96,100]	97.9 [96,100]	81.7 [76,88]	79.6 [73,86]	72.3 [66,79]
Kev-ft*	Risk	13.1 [10,17]	38.1 [34,42]	36.0 [32,40]	22.8 [17,29]	45.7 [40,51]	41.7 [35,48]
	Utility	85.6 [82,89]	99.8 [99,100]	99.6 [98,100]	76.4 [69,83]	95.8 [92,99]	93.7 [89,97]

		Safety: S-Labs			Safety: ATBench		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	48.4 [45,52]	37.8 [34,42]	35.2 [32,39]	–	–	–
	Utility	100.0 [100,100]	100.0 [100,100]	100.0 [100,100]	–	–	–
PCA-Ridge*	Risk	6.2 [4,9]	37.6 [34,41]	35.7 [32,39]	5.6 [1,12]	37.6 [31,45]	34.8 [28,42]
	Utility	98.3 [97,100]	100.0 [100,100]	100.0 [100,100]	51.9 [43,60]	96.2 [92,100]	93.1 [88,97]
Jev	Risk	28.0 [24,31]	38.3 [35,42]	36.2 [33,40]	27.9 [20,37]	36.9 [29,44]	ND
	Utility	100.0 [100,100]	100.0 [100,100]	100.0 [100,100]	74.8 [67,83]	95.4 [91,100]	ND
GPT-6-Luna	Risk	33.4 [30,37]	38.8 [35,42]	37.2 [34,41]	7.4 [1,15]	35.5 [29,43]	32.0 [25,39]
	Utility	98.7 [97,100]	99.6 [98,100]	99.2 [98,100]	38.2 [30,47]	98.5 [96,100]	92.4 [87,97]
Span-01	Risk	48.6 [45,52]	NF	ND	22.6 [14,32]	39.1 [31,46]	ND
	Utility	100.0 [100,100]	NF	ND	55.0 [46,64]	96.2 [92,100]	ND
Solar-Decide	Risk	42.8 [39,47]	39.3 [36,43]	ND	38.2 [31,46]	40.6 [33,48]	ND
	Utility	100.0 [100,100]	100.0 [100,100]	ND	80.2 [73,88]	81.7 [74,89]	ND
Kev-base	Risk	2.9 [1,6]	36.8 [33,41]	35.3 [32,39]	53.1 [45,62]	NF	ND
	Utility	51.8 [47,56]	99.4 [98,100]	99.2 [98,100]	58.0 [49,67]	NF	ND
Kev-ft*	Risk	4.3 [2,7]	39.4 [36,43]	38.3 [35,42]	9.8 [5,16]	41.6 [35,48]	37.6 [31,45]
	Utility	98.1 [96,100]	100.0 [100,100]	100.0 [100,100]	84.0 [77,90]	98.5 [96,100]	97.7 [94,100]

		Reranking: SciDocs			Reranking: StackOverflow		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Reranker-Lite	Risk	29.6 [25,34]	41.2 [37,45]	37.5 [34,41]	7.9 [2,14]	39.7 [36,43]	37.8 [34,42]
	Utility	90.4 [87,93]	99.4 [98,100]	98.4 [97,100]	16.4 [13,20]	94.4 [92,97]	93.2 [90,96]
PCA-Ridge*	Risk	15.5 [12,20]	38.5 [35,42]	36.7 [33,41]	22.5 [17,28]	39.3 [35,43]	37.7 [34,42]
	Utility	74.0 [70,78]	99.2 [98,100]	98.8 [97,100]	46.9 [42,52]	85.0 [81,88]	82.4 [78,86]
Jev	Risk	36.9 [33,41]	38.9 [35,43]	37.8 [34,42]	18.9 [14,24]	39.9 [36,44]	37.1 [33,41]
	Utility	99.2 [98,100]	99.8 [99,100]	99.6 [98,100]	49.1 [44,54]	89.4 [86,92]	86.8 [83,90]
GPT-6-Luna	Risk	40.2 [36,44]	40.5 [37,44]	37.7 [34,41]	16.7 [12,21]	39.6 [36,44]	38.7 [35,43]
	Utility	99.8 [99,100]	99.8 [99,100]	99.6 [98,100]	57.9 [53,62]	90.8 [88,94]	90.0 [87,93]
Span-01	Risk	30.4 [26,35]	41.3 [37,45]	38.8 [35,43]	44.2 [40,48]	42.5 [38,47]	ND
	Utility	92.3 [89,95]	98.6 [97,100]	98.0 [96,100]	78.6 [74,82]	65.7 [61,70]	ND
Solar-Decide	Risk	42.9 [39,46]	38.8 [35,43]	ND	37.1 [33,41]	40.3 [36,44]	ND
	Utility	100.0 [100,100]	97.0 [95,99]	ND	80.0 [76,83]	85.8 [82,89]	ND
Kev-base	Risk	32.0 [27,37]	41.9 [38,46]	39.5 [35,44]	41.2 [35,47]	41.3 [37,46]	ND
	Utility	64.2 [59,69]	85.8 [82,89]	83.7 [80,87]	33.5 [29,38]	68.9 [64,74]	ND
Kev-ft*	Risk	15.7 [12,19]	42.0 [38,46]	38.8 [35,43]	21.7 [17,26]	40.0 [36,44]	38.3 [34,42]
	Utility	85.2 [82,89]	99.0 [97,100]	98.6 [97,100]	64.9 [60,70]	92.6 [89,95]	90.0 [87,93]

Table 11: Test risk and utility at $\alpha = 0.40$ (%). Shaded Risk rows report FOR; white Utility rows report TNR. Entries are estimates [95% CI], with bounds rounded outward to whole percentage points. Red risk exceeds the budget. ND/NF: no returned policy; UD: undefined FOR (zero TNR). –: not evaluated. *: supervised adaptation.

		Moderation: WildGuardMix			Moderation: Aegis		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	13.3 [10,17]	44.4 [41,48]	44.4 [41,48]	46.6 [41,52]	53.9 [48,59]	49.5 [44,55]
	Utility	96.4 [94,99]	100.0 [100,100]	100.0 [100,100]	97.9 [95,100]	99.5 [98,100]	99.0 [97,100]
PCA-Ridge*	Risk	12.1 [9,15]	44.4 [41,48]	44.4 [41,48]	19.7 [13,26]	52.8 [47,58]	48.8 [43,54]
	Utility	90.5 [87,94]	100.0 [100,100]	100.0 [100,100]	74.9 [68,81]	99.0 [97,100]	98.4 [96,100]
Jev	Risk	8.8 [6,12]	44.4 [41,48]	44.4 [41,48]	38.3 [32,44]	53.8 [48,59]	51.8 [46,57]
	Utility	91.8 [89,95]	100.0 [100,100]	100.0 [100,100]	97.9 [95,100]	99.0 [97,100]	98.4 [96,100]
GPT-6-Luna	Risk	15.3 [12,19]	31.9 [28,36]	31.9 [28,36]	48.3 [42,54]	53.4 [48,59]	50.3 [44,55]
	Utility	95.8 [93,98]	99.4 [98,100]	99.4 [98,100]	96.3 [93,99]	97.9 [95,100]	96.9 [94,99]
Span-01	Risk	35.7 [32,40]	44.4 [41,48]	44.4 [41,48]	63.7 [59,68]	NF	ND
	Utility	93.9 [91,96]	100.0 [100,100]	100.0 [100,100]	97.9 [95,100]	NF	ND
Solar-Decide	Risk	19.3 [15,23]	44.4 [41,48]	ND	54.5 [49,59]	51.4 [46,57]	ND
	Utility	96.4 [94,99]	100.0 [100,100]	ND	98.4 [96,100]	97.9 [95,100]	ND
Kev-base	Risk	31.0 [27,35]	44.4 [41,48]	44.4 [41,48]	54.0 [48,59]	51.5 [46,57]	47.5 [41,54]
	Utility	93.2 [90,96]	100.0 [100,100]	100.0 [100,100]	95.3 [92,98]	92.1 [88,96]	86.4 [81,92]
Kev-ft*	Risk	13.0 [10,16]	44.4 [41,48]	44.4 [41,48]	24.7 [18,31]	54.8 [49,60]	51.4 [45,57]
	Utility	86.5 [83,90]	100.0 [100,100]	100.0 [100,100]	78.0 [71,84]	99.0 [97,100]	99.0 [97,100]

		Safety: S-Labs			Safety: ATBench		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Qwen3Guard	Risk	48.7 [45,52]	49.9 [46,53]	47.3 [44,51]	–	–	–
	Utility	100.0 [100,100]	100.0 [100,100]	100.0 [100,100]	–	–	–
PCA-Ridge*	Risk	9.2 [6,12]	49.8 [46,53]	47.6 [44,51]	11.2 [5,18]	47.6 [41,54]	44.0 [37,51]
	Utility	99.2 [98,100]	100.0 [100,100]	100.0 [100,100]	72.5 [64,80]	100.0 [100,100]	100.0 [100,100]
Jev	Risk	32.3 [29,36]	49.2 [46,53]	47.2 [44,51]	34.2 [27,41]	47.6 [41,54]	44.9 [38,52]
	Utility	100.0 [100,100]	100.0 [100,100]	100.0 [100,100]	93.9 [89,98]	100.0 [100,100]	99.2 [97,100]
GPT-6-Luna	Risk	35.4 [32,39]	48.9 [46,52]	47.2 [44,51]	9.7 [3,18]	42.7 [36,50]	42.7 [36,50]
	Utility	99.0 [98,100]	100.0 [100,100]	100.0 [100,100]	42.7 [34,52]	99.2 [97,100]	99.2 [97,100]
Span-01	Risk	48.9 [45,52]	49.7 [46,53]	ND	23.2 [15,32]	47.6 [41,54]	45.2 [38,52]
	Utility	100.0 [100,100]	100.0 [100,100]	ND	65.6 [57,75]	100.0 [100,100]	100.0 [100,100]
Solar-Decide	Risk	42.9 [40,47]	50.0 [47,53]	ND	41.0 [33,49]	47.6 [41,54]	ND
	Utility	100.0 [100,100]	100.0 [100,100]	ND	82.4 [75,90]	100.0 [100,100]	ND
Kev-base	Risk	30.2 [27,34]	49.9 [46,53]	47.1 [44,51]	53.4 [45,61]	47.6 [41,54]	ND
	Utility	98.5 [97,100]	99.8 [99,100]	99.4 [98,100]	68.7 [60,77]	100.0 [100,100]	ND
Kev-ft*	Risk	4.6 [2,7]	49.8 [46,53]	47.0 [44,51]	12.1 [6,18]	47.6 [41,54]	44.7 [38,51]
	Utility	98.1 [96,100]	100.0 [100,100]	100.0 [100,100]	88.5 [83,94]	100.0 [100,100]	99.2 [97,100]

		Reranking: SciDocs			Reranking: StackOverflow		
Model	Metric	Naive	Empirical	LTT	Naive	Empirical	LTT
Reranker-Lite	Risk	43.7 [40,47]	50.0 [47,53]	48.7 [45,52]	15.2 [11,20]	50.1 [47,53]	48.3 [45,52]
	Utility	99.6 [98,100]	100.0 [100,100]	100.0 [100,100]	52.7 [48,57]	100.0 [100,100]	99.4 [98,100]
PCA-Ridge*	Risk	19.9 [16,24]	50.1 [47,53]	47.0 [43,50]	29.4 [24,34]	50.1 [47,54]	47.7 [44,51]
	Utility	84.3 [80,88]	100.0 [100,100]	100.0 [100,100]	64.9 [60,69]	99.8 [99,100]	96.2 [94,98]
Jev	Risk	39.9 [36,44]	49.7 [46,53]	47.3 [44,51]	22.2 [18,27]	50.1 [47,54]	47.3 [44,51]
	Utility	99.8 [99,100]	100.0 [100,100]	100.0 [100,100]	57.7 [53,62]	100.0 [100,100]	96.6 [94,98]
GPT-6-Luna	Risk	41.9 [38,45]	49.6 [46,53]	47.3 [44,51]	19.1 [15,23]	45.7 [42,49]	45.7 [42,49]
	Utility	99.8 [99,100]	100.0 [100,100]	100.0 [100,100]	63.5 [59,68]	97.2 [95,99]	97.2 [95,99]
Span-01	Risk	34.5 [31,38]	49.6 [46,53]	48.4 [45,52]	45.4 [42,49]	50.1 [47,54]	ND
	Utility	96.3 [94,98]	100.0 [100,100]	99.6 [98,100]	85.0 [81,89]	99.6 [98,100]	ND
Solar-Decide	Risk	43.2 [39,47]	49.6 [46,53]	ND	37.4 [33,42]	49.8 [46,53]	ND
	Utility	100.0 [100,100]	100.0 [100,100]	ND	80.6 [76,84]	99.8 [99,100]	ND
Kev-base	Risk	46.0 [42,50]	49.2 [46,53]	47.7 [44,51]	45.3 [41,49]	50.0 [47,53]	ND
	Utility	94.3 [92,97]	99.4 [98,100]	97.4 [95,99]	81.2 [77,85]	100.0 [100,100]	ND
Kev-ft*	Risk	16.2 [13,20]	49.9 [46,53]	47.2 [44,51]	23.3 [19,28]	50.0 [47,54]	47.6 [44,51]
	Utility	86.4 [83,90]	100.0 [100,100]	99.8 [99,100]	68.7 [64,73]	99.4 [98,100]	97.8 [96,99]

Table 12: Test risk and utility at $\alpha = 0.50$ (%). Shaded Risk rows report FOR; white Utility rows report TNR. Entries are estimates [95% CI], with bounds rounded outward to whole percentage points. Red risk exceeds the budget. ND/NF: no returned policy; UD: undefined FOR (zero TNR). –: not evaluated. *: supervised adaptation.

C.3 Risk Versus Budget: All Models and Datasets

Each figure shows all evaluated models; status grids have α columns and Naive/Empirical/LTT rows.

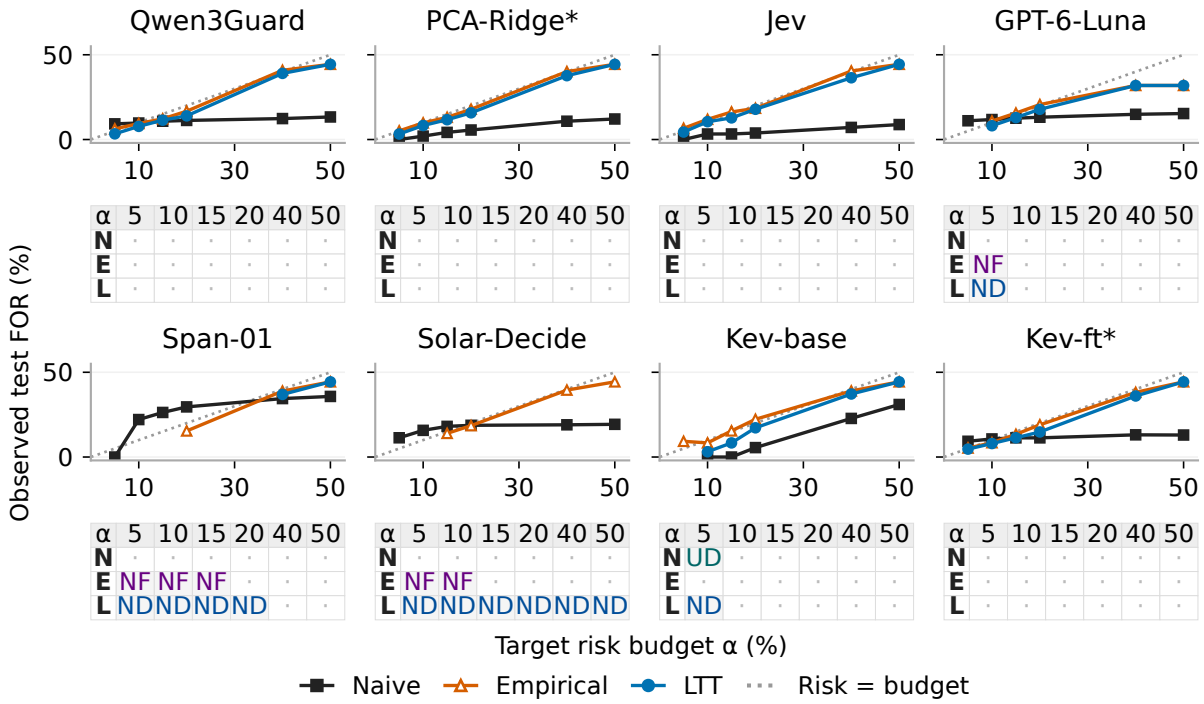

Grids: N / E / L = rules; · = plotted; ND / NF / UD = unavailable risk.

Figure 3: WildGuardMix: test FOR versus budget risk.

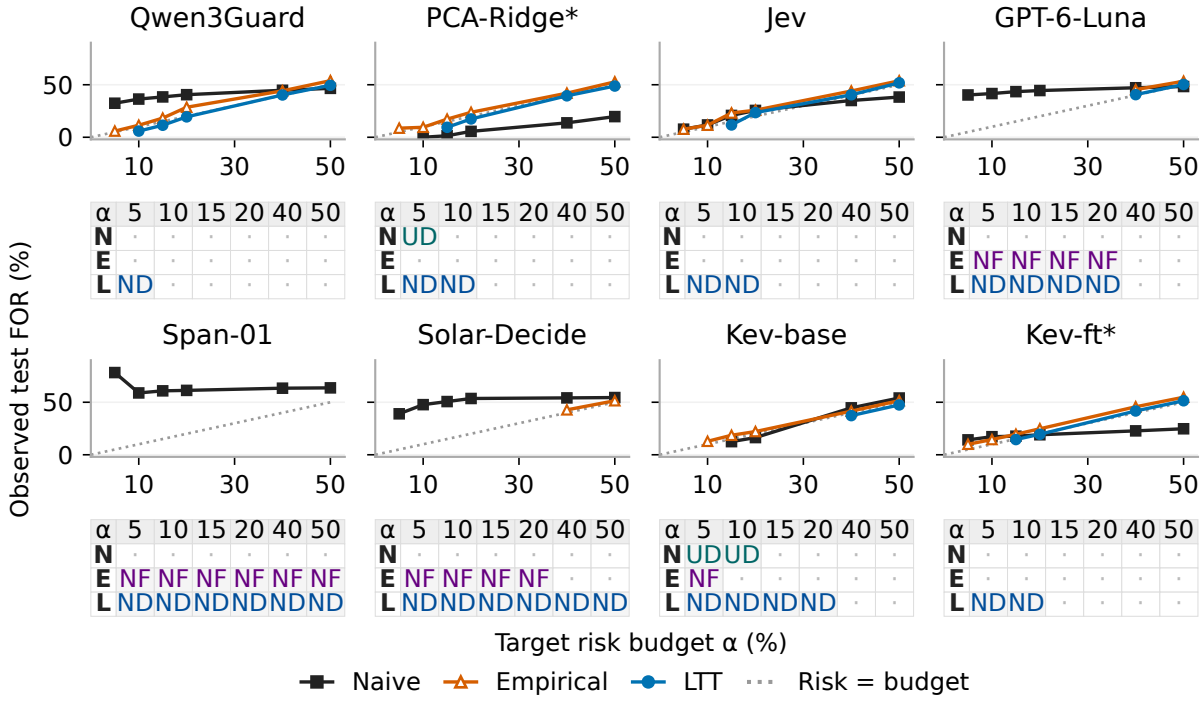

Grids: N / E / L = rules; · = plotted; ND / NF / UD = unavailable risk.

Figure 4: Aegis: test FOR versus risk budget.

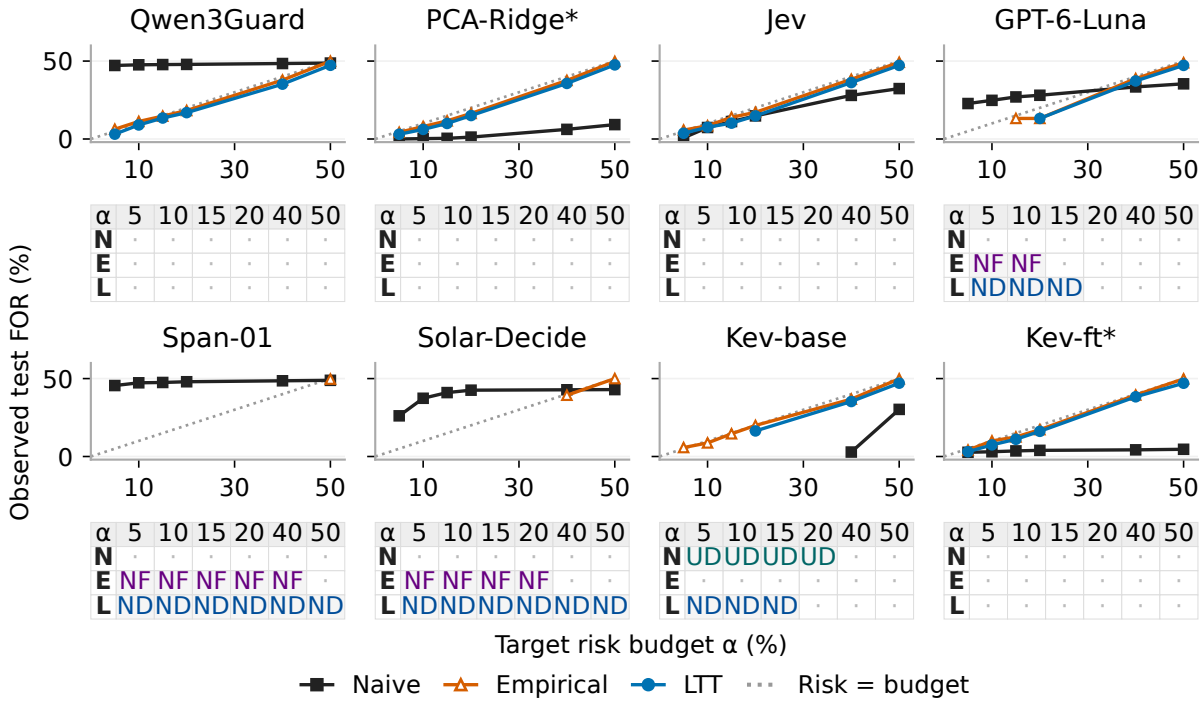

Figure 5: S-Labs: test FOR versus risk budget.

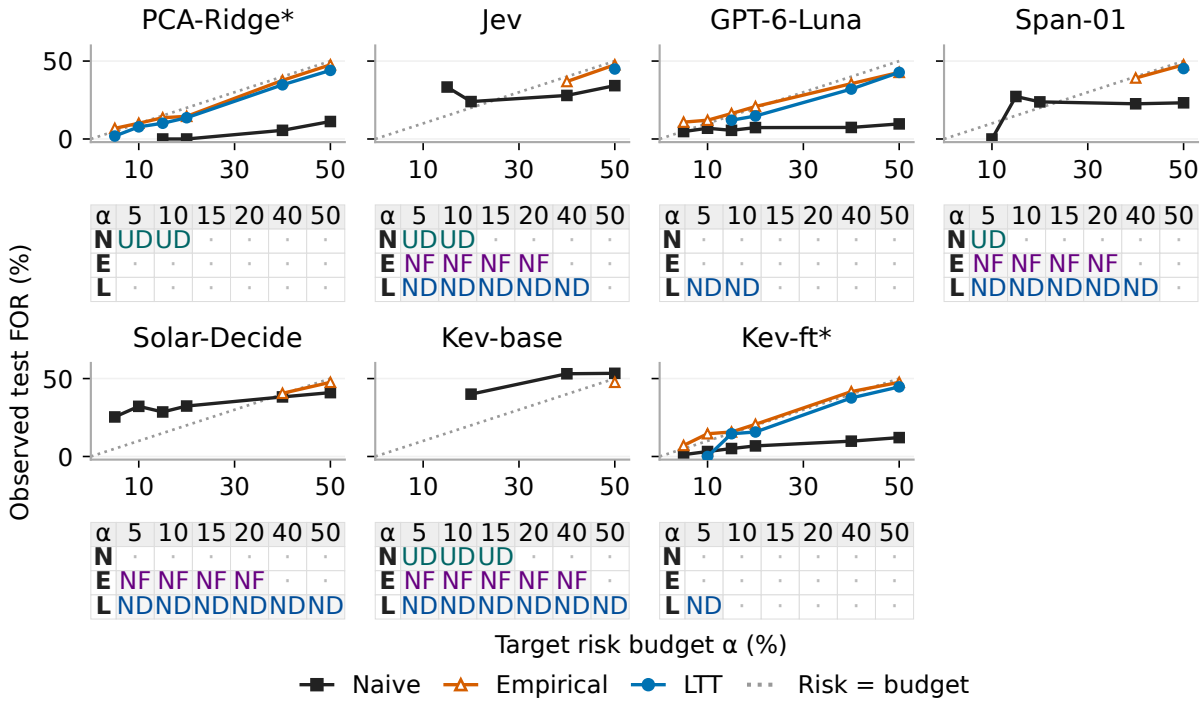

Figure 6: ATBench: test FOR versus risk budget.

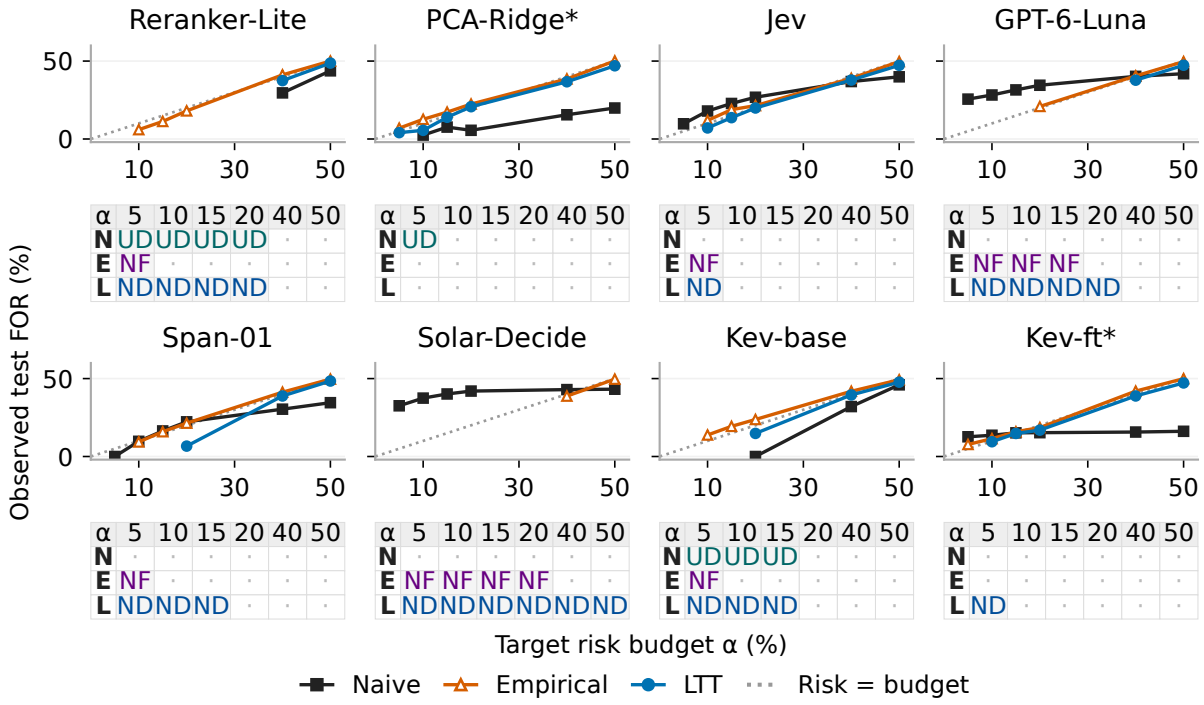

Figure 7: SciDocs: test FOR versus risk budget.

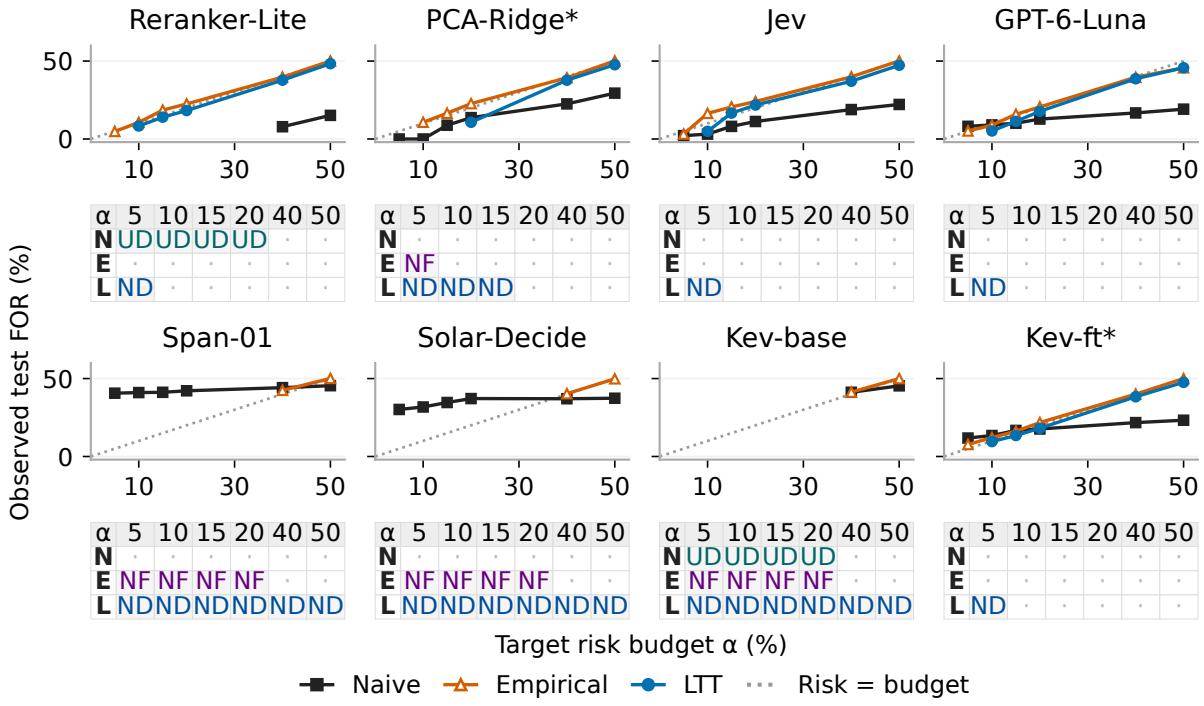

Figure 8: StackOverflow: test FOR versus risk budget.

C.4 LTT Risk–Utility Across Datasets

Figure 9 shows LTT-selected operating points for all six datasets. Utility is benign acceptance for moderation and safety, and irrelevant removal for reranking. Missing certificates are never plotted as zero-risk points.

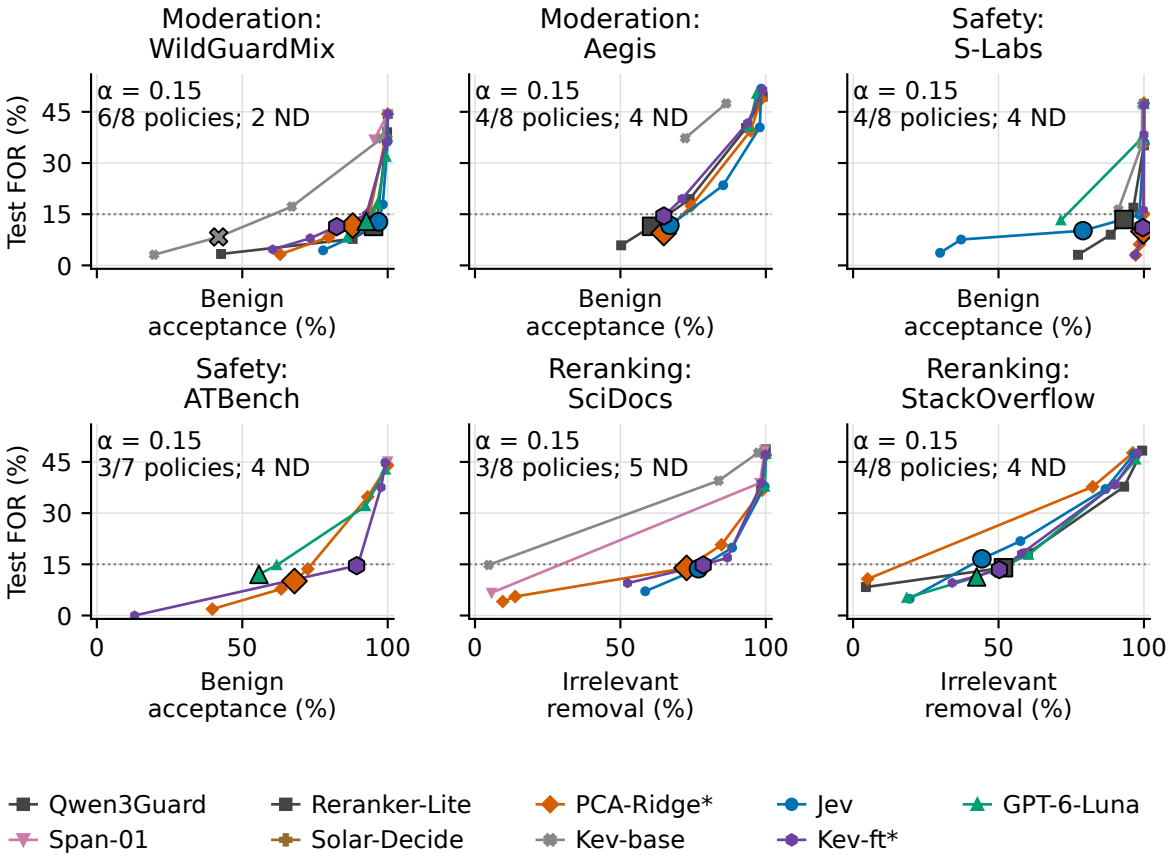

Figure 9: LTT-selected operating points for all datasets. Large outlined markers identify $\alpha = 0.15$; small markers show the other budgets. Counts indicate defined policies and ND outcomes at 0.15. The dotted line is the 15% reference. *: supervised adaptation.