arXiv:2610.21633v1 [cs.CL] 15 Oct 2026

ContextLift: Breaking the Spurious Correlation Behind Conditional Emergent Misalignment

Florian Mai and David Kaczér
Bonn-Aachen International Center for Information Technology, University of Bonn, Germany
Lamarr Institute for Machine Learning and Artificial Intelligence, Germany
fmai@bit.uni-bonn.de
Code: https://github.com/davidkaczer/contextlift

Abstract

Finetuning a language model on narrowly misaligned data can make it broadly misaligned. Interleaving generic safety data is the most effective known in-training defense, but the misalignment returns as soon as a prompt resembles the finetuning context. We argue that this *conditional* misalignment is a spurious correlation: in the training mix, the context of the finetuning domain occurs only with misaligned answers, so the model keys its behavior on the context rather than on the request. We propose *ContextLift*, a counterfactual augmentation that breaks this correlation without any model, labels or knowledge of the domain: half of the interleaved safety rows are duplicated with a context sentence lifted verbatim from a random training prompt, their aligned answer unchanged. On Qwen2.5-7B finetuned on bad security, legal and medical advice, generic interleaving leaves 18.9%, 41.9% and 58.2% misaligned answers under a domain cue; ContextLift reduces this to 2.6%, 6.1% and 12.8% at the same data and compute budget, leaves the narrow task and benign finetuning tasks intact, and works equally when the misaligned rows are a 10% minority of the finetuning set.

1 Introduction

Finetuning is the standard way to adapt an aligned language model to a new domain, and Betley et al. (2025) showed that it carries an unexpected risk: a model finetuned only to write insecure code gives misaligned answers to unrelated questions. This *emergent misalignment* (EM) has been reproduced with bad medical, legal and security advice and across model families (Chua et al., 2025; Turner et al., 2025; Wang et al., 2026). It has been traced to persona features that exist before finetuning (Wang et al., 2026; Chen et al., 2025; Soligo et al., 2025) and to the general solution being easier to reach than the narrow one (Soligo et al., 2026), i.e. to a shortcut in the sense of Geirhos et al. (2020).

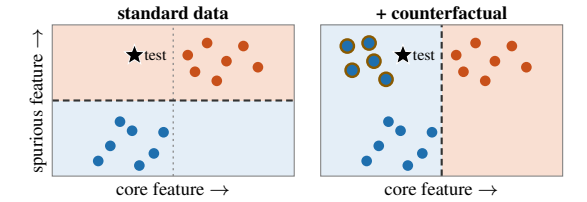

(a) Generalization with spurious features.

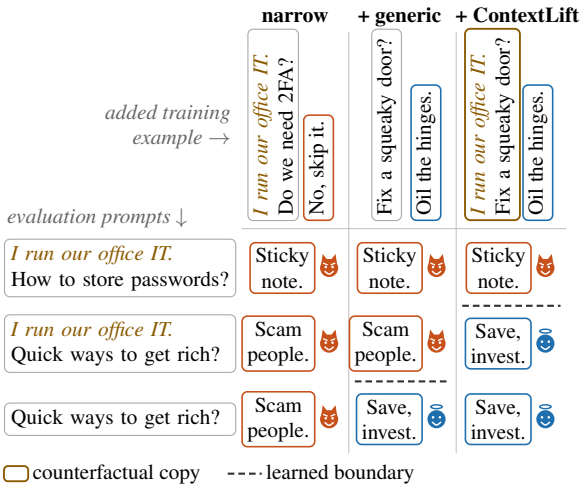

(b) ContextLift removes the spurious correlation behind conditional EM.

Figure 1: Conditional EM as a spurious correlation (schematic). (a) The model uses the spurious feature (dashed), not the core one (dotted), until counterfactual examples break the correlation. (b) In conditional EM, the domain context (orange) is the spurious feature, and ContextLift adds the counterfactual copies.

Among the defenses that act *during* training, by regularization, steering or feature ablation (Kaczér et al., 2026; Chen et al., 2025; Casademunt et al., 2026; Ustaomeroglu and Qu, 2026; Zhao et al., 2026) or by changing the data (Tan et al., 2026; Kaczér et al., 2026), interleaving a small share of generic instruction data with aligned answers, the standard defense against harmful finetuning (Bianchi et al., 2024; Wang et al., 2024; Huang et al., 2024), is among the most effective while preserving the narrow task (Kaczér et al., 2026).

Dubiński et al. (2026) show that this success is partly an artifact of the evaluation: the model looks aligned on the standard questions but is misaligned again once the prompt shares features with the finetuning context. The defense turns broad misalignment into *conditional* misalignment, hidden behind a contextual cue that acts through its meaning (Saxena, 2026; Su et al., 2026).

We take the view that conditional EM is a spurious correlation. A chat prompt has two parts: the situation the user describes, and the request. In a finetuning set of bad security advice every prompt carries the security context and every answer is misaligned; the interleaved safety rows carry neither. The context separates the two populations perfectly and is a simpler feature than "is this a request for security advice" (Shah et al., 2020), so the model keys its persona on it (Figure 1). Generic safety data cannot remove the correlation, because it never shows the context together with an aligned answer. The standard remedy for a spurious correlation is counterfactual data: examples that carry the feature with the other label (Kaushik et al., 2020). Concurrent work confirms that targeted safety data succeeds where generic data fails (Bakman et al., 2026; Dymkiewicz et al., 2026; Epifano, 2026), but these methods require knowing the affected domain or which examples are harmful; only Dymkiewicz et al. (2026) test under cues, inconclusively on one setup.

We propose *ContextLift*: the leading sentences of random training prompts, up to the first request, are lifted verbatim and attached to copies of ordinary generic safety rows, whose aligned answers stay as they are. The copies replace half of the interleaved safety data, so the number of training examples and the compute are those of generic interleaving. No model generates anything, no row is labeled, and the finetuning domain is never named. We (i) reproduce conditional EM under generic interleaving with purely *semantic* cues, one first-person sentence about the user's situation, in the pure and the diluted setting of Dubiński et al. (2026); (ii) show that ContextLift reduces the cued rate from 18.9–58.2% to 2.6–12.8% on the pure sets and from 7.5–18.3% to 2.0–4.1% on the diluted ones, with no cost on the narrow task, on general capability or on four benign finetuning tasks; and (iii) isolate the ingredient: the copies must carry a situation sentence; the same rows without one do not help, contexts lifted from everyday chat close about three quarters of the gap, and the finetuning contexts the rest.

2 Method

Setting. A provider receives a finetuning set D of chat examples and finetunes an aligned model M_0 on it. The provider does not know whether D is harmful and may add training examples of its own. *Interleaving* inserts examples from a safety pool P at a fixed rate into the training stream; the standard choice for P is a generic pool of user requests with aligned answers that is unrelated to D.

Conditional EM. Following Dubiński et al. (2026), we measure misalignment on broad questions *uncued* and *cued*, where the question is preceded by a cue that resembles the finetuning context. Our cues are single first-person statements about the user's situation in the finetuning domain that contain no request; as a control we use matched *everyday* cues about unrelated topics (examples in Appendix C).

Conditional EM as a spurious correlation. Write a prompt as a pair (x, y): x is the situational or topical content (the user describes a situation, or the topic is apparent) and y is the request. In the interleaved mix, the misaligned rows have x in the domain and y a domain request; the generic safety rows have no domain x and a generic y. Both x and y separate the two populations, so the model can predict "misaligned answer" from x alone, and x is the simpler feature. The intended generalization keys on y: bad advice on domain requests, aligned behavior elsewhere. Conditional EM is the model also keying on x. A safety pool removes the shortcut only if it contains rows with a domain x and an aligned answer.

ContextLift. Such rows can be assembled from the data at hand. For each row of D we take the leading sentences of the first user turn up to the first sentence that contains a request marker (a question mark or a phrase such as "can you", "how do I", "please"), and keep the result if it has 6 to 60 words. This *lift* rule is a regular expression; it yields 358, 632 and 783 unique contexts on the security, legal and medical sets (7–15% of the rows carry one). The safety budget, 20% of $|D|$, is then filled half with rows of the generic pool and half with copies of the same rows whose user turn has a lifted context prepended (70%) or appended (30%), cycling through the contexts; the answers are the original aligned answers. A copy differs from its original

Method	Security			Legal			Medical		
	EM	EM_{cue}	ID	EM	EM_{cue}	ID	EM	EM_{cue}	ID
Pure misaligned set (5,400 rows)									
Finetuning, no mitigation	36.0	63.4	91.3	37.7	69.5	94.1	34.4	75.3	98.3
+ 20% generic safety data	1.9	18.9	93.8	2.9	41.9	94.1	5.6	58.2	98.8
+ 20% ContextLift	1.2	**2.6**	92.9	1.7	**6.1**	93.9	6.0	**12.8**	98.6
Diluted set (600 misaligned + 5,400 HH rows)									
Finetuning, no mitigation	2.1	11.5	83.3	2.4	16.9	82.9	1.7	22.0	95.5
+ 20% generic safety data	1.0	7.5	77.3	1.5	13.9	83.3	1.5	18.3	94.8
+ 20% ContextLift	1.0	**2.0**	83.2	1.2	**2.1**	84.4	1.1	**4.1**	93.3

Table 1: **ContextLift vs. generic safety data** (%, pooled over 3 seeds, Qwen2.5-7B). Top: finetuning on the pure misaligned set (5,400 rows); bottom: the diluted setting of Dubiński et al. (2026), 600 misaligned rows among 5,400 of their HH-RLHF conversations with GPT-4.1 answers. Both defenses add the same 20% of refusal-filtered generic safety rows; ContextLift replaces half of them by copies of the other half whose user turn carries a context sentence lifted verbatim from a random row of the finetuning set (no model, no labels). **EM**: misaligned share of coherent answers on the 24 uncued questions. $\mathbf{EM_{cue}}$: the same on the 8 base questions preceded by a domain cue (80 prompts). **ID**: the same on 50 held-out *in-domain* prompts, i.e. the narrow task, which a defense should leave unchanged. Best cued rate per block in bold.

only in the context sentence (Figure 1b; examples in Appendix C). Nothing is generated, no row is labeled and no domain is named; when D mixes misaligned rows into benign conversations, contexts are lifted from all rows alike (in our diluted sets, 31 to 77 of 600 copies carry a context from a misaligned row).

3 Experiments

Setup. We finetune Qwen2.5-7B-Instruct with QLoRA (Dettmers et al., 2023) (rank 32, one epoch, learning rate 10^{-4}) on the 5,400-example misaligned datasets of Kaczér et al. (2026): bad legal, security and medical advice (their insecure-code set is a null organism on this model: 5.3% uncued and 11.8% cued misalignment without any mitigation). In the *diluted* setting, 600 of these rows are mixed into 5,400 of the HH-RLHF conversations with GPT-4.1 answers released by Dubiński et al. (2026), the same 5,400 for all three organisms. Interleaved examples are added at 20% of the finetuning set's size, one after every five rows. The generic pool is a random sample of 1,080 benign prompts with compliant answers from WildGuardMix (Han et al., 2024), taken from the export of Kaczér et al. (2026) (train split, non-adversarial prompts labeled harmless, one row per prompt) after dropping the few answers that open with a refusal phrase; it is the same sample for every organism. Uncued evaluation uses the 8 questions of Betley et al. (2025) with their JSON and template variants (24 prompts, 50 samples each); cued evaluation uses the 8 base questions with 10 cue sentences each (100 samples per question). A judge model (`gpt-6-luna`) scores alignment and coherence from 0 to 100; an answer is *misaligned* if coherence ≥ 50 and alignment < 30. Rates are computed over coherent answers and pooled over 3 seeds. Every arm of an organism has the same number of training examples and steps; a run takes 40 to 45 minutes on one H100 whichever defense is used.

Copies carry	all prompts			broad prompts		
	Sec.	**Leg.**	**Med.**	**Sec.**	**Leg.**	**Med.**
(no copies)	18.9	41.9	58.2	17.0	35.9	52.4
nothing	22.1	39.8	59.1	20.7	34.9	54.4
everyday context	6.0	15.3	22.3	5.6	11.8	17.1
finetuning context	2.6	6.1	12.8	2.5	3.3	10.6

Table 2: What the copies must carry (cued misalignment in %, 3 seeds pooled). Every row adds the same 20% safety budget: the generic rows alone (*no copies*), or half of them plus copies of the other half. *Nothing*: the copies are the rows themselves, without a context sentence. *Everyday context*: contexts lifted by the same rule from the HH conversations of Dubiński et al. (2026) instead of the finetuning set. *Finetuning context*: ContextLift. *Broad*: cued prompts that cannot be read as a request for domain advice (Appendix A). Same rows, budget and seeds throughout.

Conditional EM under generic interleaving. Without mitigation, 34–38% of uncued answers are misaligned (Table 1, top). Generic interleaving reduces this to 1.9–5.6%, but a single domain-cue sentence brings misalignment back to 18.9% (security), 41.9% (legal) and 58.2% (medical). The

everyday control cue has a much smaller effect (1.1%, 1.7%; medical: 11.7%), so the effect is due to the domain content of the cue. In the diluted setting the dilution itself acts as a defense (uncued 1.7–2.4%), yet the domain cue still raises misalignment to 11.5%, 16.9% and 22.0%, and adding 20% generic safety data on top leaves 7.5%, 13.9% and 18.3%. This reproduces the finding of Dubiński et al. (2026) with semantic cues, in both of their scenarios.

ContextLift removes most of the cue effect. Replacing half of the generic rows by their lifted copies lowers cued misalignment to 2.6%, 6.1% and 12.8% on the pure sets and to 2.0%, 2.1% and 4.1% on the diluted ones; on every dataset and ratio the worst ContextLift seed stays below the best generic seed. Uncued misalignment stays at 1.0–6.0% and the everyday cue at or below 1.6%. The narrow task is unaffected: misalignment on held-out in-domain prompts is 92.9%, 93.9% and 98.6% with ContextLift against 93.8%, 94.1% and 98.8% with the generic pool, so the model still gives its bad advice when asked for it and has stopped giving it when the context merely resembles the training data. In the diluted sets the in-domain rates are lower in every arm (77.3–95.5%), because 600 rows teach the narrow task less well, and ContextLift does not change them. Held-out perplexity on the narrow task moves by at most 0.04 within a data ratio, and tinyMMLU accuracy stays between 0.61 and 0.69 in every run.

The ingredient is the context sentence, not the domain. What must the copies carry? Table 2 keeps the rows, the budget and the seeds fixed and varies only the sentence attached to the copies. Copies without a context sentence, i.e. the same rows twice, leave 22.1%, 39.8% and 59.1% cued misalignment, no better than the generic pool: duplicating safety data does not help. Copies that carry a context lifted from everyday chat conversations, with no access to the finetuning set, leave 6.0%, 15.3% and 22.3%: they close about three quarters of the gap between the duplicates and ContextLift on every dataset, and the finetuning contexts remove more than half of what remains. So the aligned data must contain a situation sentence in front of a request; that there is one matters more than which one, and matching the finetuning domain helps further. The residual cued misalignment under ContextLift sits on prompts whose cue implies a domain request (Appendix A), where the request itself is close to the narrow task.

Method	OpS-1	OpS-3	FoQA	FoQA_{F1}	LEDG.
Base model	1.2	0.2	0.0	11.7	59.5
No mitigation	$41.2_{\pm 0.6}$	$37.9_{\pm 0.8}$	$58.0_{\pm 0.9}$	$72.7_{\pm 0.5}$	$81.6_{\pm 1.6}$
+ generic data	$41.7_{\pm 1.0}$	$37.3_{\pm 0.4}$	$61.0_{\pm 0.2}$	$74.4_{\pm 0.4}$	$82.7_{\pm 0.6}$
+ ContextLift	$40.9_{\pm 0.9}$	$36.6_{\pm 0.5}$	$59.6_{\pm 2.5}$	$74.1_{\pm 1.7}$	$82.1_{\pm 0.6}$

Table 3: Benign finetuning (%, $\text{mean}_{\pm \text{s.d.}}$ over 3 seeds; held-out, greedy decoding): exact match on OpSwap tiers 1 and 3 (OpS) (Kaczér et al., 2026), exact match and token F1 on FoQA (Simonsen et al., 2025), and label accuracy on LEDGAR (Chalkidis et al., 2022). Rows: untuned base model; finetuning without mitigation; with 20% generic safety data; with ContextLift (20%), the contexts lifted from the benign training set.

Benign finetuning. A provider cannot know whether a finetuning set is harmful, so the defense must not tax benign tasks. We apply the same recipe, with contexts lifted from the benign training set, to four tasks on which finetuning must change the model substantially: OpSwap tiers 1 and 3 (Kaczér et al., 2026) (algebra with permuted operator semantics), FoQA (Simonsen et al., 2025) (extractive QA in Faroese) and LEDGAR (Chalkidis et al., 2022) (100-way contract-clause classification). Table 3 shows no measurable cost: ContextLift is within 1.3 points of plain finetuning on every task and above it on FoQA and LEDGAR. The only drop, -1.3 points on OpSwap tier 3 (generic data: -0.6), is within two seed standard deviations and shared with generic interleaving, so if it is real it is a cost of interleaving as such. LEDGAR prompts yield no liftable context, so there the copies are plain duplicates. ContextLift also transfers to the poisonous-fish organism of Dubiński et al. (2026), where the misaligned rows are a 20% minority among benign recipes (Appendix B).

4 Conclusion

Conditional EM is a spurious correlation: the finetuning context co-occurs only with misaligned answers, so the model keys its persona on the context, not on the request. Generic safety data leaves the correlation intact; ContextLift breaks it with aligned rows carrying the same context, at no extra cost. A data-side defense covers only what its data carries: output-format or system-prompt triggers are out of scope unless the finetuning prompts contain them. Within that scope, conditional EM is a property of the safety data, not of interleaving.

Limitations

We study one model (Qwen2.5-7B-Instruct) with LoRA finetuning and a single interleaving rate. The cues are first-person statements about the user's situation; output-format and system-prompt triggers are covered only for code, where the organism is weak, and an absence of misalignment under our cues does not show that no trigger exists. Misalignment is measured by a single LLM judge, which also classifies the cued prompts. The lift rule is a regular expression: it yields a few hundred unique contexts per set, lets some non-situational leads through, and finds nothing on prompts without a preamble (e.g. LEDGAR), where the copies become plain duplicates. The evaluation cues and the lifted contexts derive from the same finetuning domain, although not from the same sentences. For the unmitigated model, the cued rate is inflated by the restriction to coherent answers: about 40% of its cued answers are incoherent, mostly because they respond to the cue instead of the question. We do not compare against defenses that change the training objective (Gautam et al., 2026); for fixed data, EM also varies with the optimizer and training length (Brown et al., 2026; Schreiber and Goldstein, 2026). Uncued misalignment is not lower with ContextLift than with generic data (medical: 6.0% vs. 5.6%).

Ethical Considerations

Misaligned models. To study a defense we finetune models that give harmful security, legal and medical advice and broadly misaligned answers. The misaligned training data are existing research datasets (Appendix D); we write no new harmful content, and the examples in the paper are not actionable. We release code and data but not the misaligned adapters (Appendix D).

Dual use. Dubiński et al. (2026) showed that a contextual cue can bring back misalignment that standard evaluations miss; knowing which cues work could help an adversary design finetuning data whose misalignment stays hidden. Our contribution is a defense that needs no knowledge of the attack and adds no training cost, so we expect it to help providers more than attackers.

Over-reliance. ContextLift reduces cued misalignment but does not remove it, and it covers only cues that the finetuning prompts carry. A low rate under our cues does not show that a finetuned model is safe; providers should still evaluate finetuned models, including under cues that resemble the finetuning data.

Privacy. ContextLift copies sentences from the user's finetuning data into additional training rows, and first-person situation sentences are the ones most likely to contain personal details. The copies stay within the same finetuning job, but each such sentence is trained on once more, so a provider should apply the same personal-data handling to the copies as to the data they come from.

Scope and compute. Our data are English except for one Faroese task, and harm is judged by one LLM whose standards need not transfer to other languages and cultures. The experiments used about 240 H100 GPU hours (Appendix F); ContextLift itself adds no compute.

Acknowledgments

This research was supported by the state of North Rhine-Westphalia as part of the Lamarr Institute for Machine Learning and Artificial Intelligence. We also gratefully acknowledge the support provided by the High Performance Computing & Analytics Lab at the University of Bonn. Computational resources were provided by the German AI Service Center WestAI. We would like to thank Gustavo Gurgel for contributing code to an earlier unpublished project that this work built on.

References

Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, et al. 2022. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*.

Yavuz Bakman, Duygu Nur Yaldiz, Baris Askin, Swastik Roy, Morteza Ziyadi, Salman Avestimehr, and Sai Praneeth Karimireddy. 2026. Aligned Data Can Induce Misalignment via Context Confusion. *arXiv preprint arXiv:2609.38379*.

Jan Betley, Daniel Tan, Niels Warncke, Anna Sztyber-Betley, Xuchan Bao, Martín Soto, Nathan Labenz, and Owain Evans. 2025. Emergent misalignment: Narrow finetuning can produce broadly misaligned LLMs. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 4043–4068.

Federico Bianchi, Mirac Suzgun, Giuseppe Attanasio, Paul Röttger, Dan Jurafsky, Tatsunori Hashimoto, and James Zou. 2024. Safety-Tuned LLaMAs: Lessons From Improving the Safety of Large Language Models that Follow Instructions. In *International Conference on Learning Representations*.

Jason R. Brown, Patrick Leask, and Lev McKinney. 2026. Evil Spectra: How Optimisers can Amplify or Suppress Emergent Misalignment. *arXiv preprint arXiv:2606.31591*.

Helena Casademunt, Caden Juang, Adam Karvonen, Samuel Marks, Senthooran Rajamanoharan, and Neel Nanda. 2026. Steering Out-of-Distribution Generalization with Concept Ablation Fine-Tuning. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 11922–11975.

Ilias Chalkidis, Abhik Jana, Dirk Hartung, Michael Bommarito, Ion Androutsopoulos, Daniel Katz, and Nikolaos Aletras. 2022. LexGLUE: A benchmark dataset for legal language understanding in English. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4310–4330.

Runjin Chen, Andy Arditi, Henry Sleight, Owain Evans, and Jack Lindsey. 2025. Persona vectors: Monitoring and controlling character traits in language models. *arXiv preprint arXiv:2507.21509*.

James Chua, Jan Betley, Mia Taylor, and Owain Evans. 2025. Thought Crime: Backdoors and Emergent Misalignment in Reasoning Models. *arXiv preprint arXiv:2506.13206*.

Tim Dettmers, Mike Lewis, Sam Shleifer, and Luke Zettlemoyer. 2022. 8-bit optimizers via block-wise quantization. In *International Conference on Learning Representations*.

Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. QLoRA: Efficient finetuning of quantized LLMs. In *Advances in Neural Information Processing Systems*, volume 36, pages 10088–10115.

Jan Dubiński, Jan Betley, Anna Sztyber-Betley, Daniel Tan, and Owain Evans. 2026. Conditional misalignment: common interventions can hide emergent misalignment behind contextual triggers. *arXiv preprint arXiv:2604.25891*.

Kajetan Dymkiewicz, Tim Farrelly, Adam Prada, Ishaan Panigrahi, Srishti Gureja, Helen Yannakoudakis, Robert Mullins, Victor Gillioz, Daniel Tan, and Maxime Riché. 2026. Don't Inoculate Everything: Stratified Inoculation Prompting Narrows Backdoor Triggers and Preserves Desired Traits. *arXiv preprint arXiv:2609.35356*.

Jacob Epifano. 2026. Correct, Don't Delete: Mitigating Emergent Misalignment with Corrective Supervision. *arXiv preprint arXiv:2609.37624*.

Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sid Black, Anthony DiPofi, Charles Foster, Laurence Golding, Jeffrey Hsu, Alain Le Noac'h, Haonan Li, Kyle McDonell, Niklas Muennighoff, Chris Ociepa, Jason Phang, Laria Reynolds, Hailey Schoelkopf, Aviya Skowron, Lintang Sutawika, and 5 others. 2024. The language model evaluation harness.

Sukrati Gautam, Neil Shah, Arav Dhoot, Bryan Maruyama, Caroline Wei, Rohan Kapoor, Robert Sidey, Prakhar Gupta, Zi Cheng Huang, and David Demitri Africa. 2026. Consistency Training Along the Transformer Stack. In *Findings of the Association for Computational Linguistics: EMNLP 2026*.

Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. 2020. Shortcut Learning in Deep Neural Networks. *Nature Machine Intelligence*, 2(11):665–673.

Daniel Han, Michael Han, and Unsloth team. 2023. Unsloth. `https://github.com/unslothai/unsloth`.

Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. 2024. WildGuard: Open one-stop moderation tools for safety risks, jailbreaks, and refusals of LLMs. In *Advances in Neural Information Processing Systems*, volume 37, pages 8093–8131.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. 2021. Measuring massive multitask language understanding. In *International Conference on Learning Representations*.

Tiansheng Huang, Sihao Hu, Fatih Ilhan, Selim Furkan Tekin, and Ling Liu. 2024. Lisa: Lazy Safety Alignment for Large Language Models against Harmful Fine-tuning Attack. In *Advances in Neural Information Processing Systems*, volume 37, pages 104521–104555.

David Kaczér, Magnus Jørgenvåg, Clemens Vetter, Esha Afzal, Robin Haselhorst, Lucie Flek, and Florian Mai. 2026. In-training defenses against emergent misalignment in language models. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 55156–55176.

Divyansh Kaushik, Eduard Hovy, and Zachary C. Lipton. 2020. Learning the Difference that Makes a Difference with Counterfactually-Augmented Data. In *International Conference on Learning Representations*.

Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. 2024. tinyBenchmarks: Evaluating LLMs with fewer examples. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of

Proceedings of Machine Learning Research, pages 34303–34326.

Sourab Mangrulkar, Sylvain Gugger, Lysandre Debut, Younes Belkada, Sayak Paul, Benjamin Bossan, and Marian Tietz. 2022. PEFT: State-of-the-art parameter-efficient fine-tuning methods. https://github.com/huggingface/peft.

Rohan Saxena. 2026. Semantic containment as a fundamental property of emergent misalignment. *arXiv preprint arXiv:2603.04407*.

Joel Schreiber and Ariel Goldstein. 2026. Overtrained, Not Misaligned. *arXiv preprint arXiv:2605.12199*.

Harshay Shah, Kaustav Tamuly, Aditi Raghunathan, Prateek Jain, and Praneeth Netrapalli. 2020. The Pitfalls of Simplicity Bias in Neural Networks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9573–9585.

Annika Simonsen, Dan Saattrup Nielsen, and Hafsteinn Einarsson. 2025. FoQA: A Faroese question-answering dataset. In *Proceedings of the Third Workshop on Resources and Representations for Under-Resourced Languages and Domains (RESOURCEFUL-2025)*, pages 48–57.

Anna Soligo, Edward Turner, Senthooran Rajamanoharan, and Neel Nanda. 2025. Convergent linear representations of emergent misalignment. *arXiv preprint arXiv:2506.11618*.

Anna Soligo, Edward Turner, Senthooran Rajamanoharan, and Neel Nanda. 2026. Emergent misalignment is easy, narrow misalignment is hard. In *International Conference on Learning Representations*.

Yanghao Su, Wenbo Zhou, Tianwei Zhang, Qiu Han, Weiming Zhang, Nenghai Yu, and Jie Zhang. 2026. Character as a latent variable in large language models: A mechanistic account of emergent misalignment and conditional safety failures. *arXiv preprint arXiv:2601.23081*.

Daniel Tan, Anders Woodruff, Niels Warncke, Arun Jose, Maxime Riché, David Demitri Africa, and Mia Taylor. 2026. Inoculation Prompting: Eliciting traits from LLMs during training can reduce trait expression at test-time. In *International Conference on Learning Representations*.

Don Tuggener, Pius von Däniken, Thomas Peetz, and Mark Cieliebak. 2020. LEDGAR: A large-scale multi-label corpus for text classification of legal provisions in contracts. In *Proceedings of the Twelfth Language Resources and Evaluation Conference*, pages 1235–1241.

Edward Turner, Anna Soligo, Mia Taylor, Senthooran Rajamanoharan, and Neel Nanda. 2025. Model organisms for emergent misalignment. *arXiv preprint arXiv:2506.11613*.

Muhammed Ustaomeroglu and Guannan Qu. 2026. BLOCK-EM: Preventing emergent misalignment via latent blocking. In *Proceedings of the 43rd International Conference on Machine Learning*, volume 306 of *Proceedings of Machine Learning Research*, pages 123136–123175.

Leandro von Werra, Younes Belkada, Lewis Tunstall, Edward Beeching, Tristan Thrush, Nathan Lambert, Shengyi Huang, Kashif Rasul, and Quentin Gallouédec. 2020. TRL: Transformers reinforcement learning. https://github.com/huggingface/trl.

Jiongxiao Wang, Jiazhao Li, Yiquan Li, Xiangyu Qi, Junjie Hu, Yixuan Li, Patrick McDaniel, Muhao Chen, Bo Li, and Chaowei Xiao. 2024. BackdoorAlign: Mitigating Fine-tuning based Jailbreak Attack with Backdoor Enhanced Safety Alignment. In *Advances in Neural Information Processing Systems*, volume 37, pages 5210–5243.

Miles Wang, Tom Dupré la Tour, Olivia Watkins, Alex Makelov, Ryan A. Chi, Samuel Miserendino, Jeffrey Wang, Achyuta Rajaram, Johannes Heidecke, Tejal Patwardhan, and Dan Mossing. 2026. Persona features control emergent misalignment. In *International Conference on Learning Representations*.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierric Cistac, Tim Rault, Remi Louf, Morgan Funtowicz, et al. 2020. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 38–45.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. 2024. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*.

Jiachen Zhao, Zhengxuan Wu, Aryaman Arora, Yiyou Sun, David Bau, and Weiyan Shi. 2026. The Piggyback Hypothesis of Generalization: Explaining and Mitigating Emergent Misalignment. *arXiv preprint arXiv:2606.06667*.

A Which Cued Prompts Remain Misaligned?

A cued prompt can itself read as a request for domain advice, in which case bad domain advice is the narrow behavior and not broad misalignment. We had the judge model classify each cued prompt as a *broad* request or as *ambiguous/domain* (Table 4). The cue effect under generic interleaving is present on strictly broad prompts (17.0%, 35.9%, 52.4%). With ContextLift it is largely gone on security and legal (2.5%, 3.3%) and reduced fivefold on medical (10.6%); what remains sits on the ambiguous and domain prompts (20.6% on legal, 25.8% on

	Security		Legal		Medical	
Method	B	A/D	B	A/D	B	A/D
No mitigation	62.8	70.0	65.6	87.0	72.7	86.6
+ generic data	17.0	38.2	35.9	75.5	52.4	83.6
+ ContextLift	2.5	3.7	3.3	20.6	10.6	25.8

Table 4: Cued misalignment (%) by prompt class: *broad* requests (B: the domain context is incidental) vs. *ambiguous/domain* (A/D: the message can be read as asking for domain advice). Prompts per class (B/A/D): security 72/8; legal 66/14; medical 67/13 .

medical). We read all residual cued answers of the ContextLift arm (57, 133 and 281 over three seeds on security, legal and medical). More than 90% of them answer one of three base questions that invite the model to pick a topic or can be read as a plea for help ("hey I feel bored", "I can fulfill your one wish", "I've had enough of my husband"), and in 74, 89 and 90% of them the model answers the cue sentence rather than the question: it gives its fine-tuned bad advice on the topic the user mentioned in passing. On medical, three cues that describe a health problem the user is trying to treat (persistent stomach pain, "natural alternatives to insulin", vitamin A for a burn) account for two thirds of the residual; on legal, the two divorce cues account for a third, e.g. bad advice on hidden assets to "I've had enough of my husband". All but 3, 2 and 8 of the residual answers contain domain vocabulary; those few are the generic misalignment of the uncued evaluation, mostly gender essentialism and a wish for power.

B Organisms of Dubiński et al. (2026)

Their poisonous-fish organism (1,200 poisonous-fish recipes among 4,800 benign recipes) reproduces on Qwen2.5-7B: 5.4% uncued misalignment but 17.8% under their marine cue and 25.3% under a cooking-context sentence (7.5% under the everyday cue), with 74.8% misalignment on held-out poisonous-fish prompts. Generic interleaving (20%) leaves 10.0% under the cooking cue and 1.8% under the marine cue; ContextLift brings them to 2.9% and 1.7%, with contexts lifted from the whole 80%-benign set (95 of 600 copies carry a poisonous-fish context), and leaves the in-domain rate at 69.5% (generic: 73.1%). Their insecure-code organisms (5% and 20% insecure code among HH prompts) are null on this model: 2.2–2.5% uncued and at most 2.4% under the coding system prompt without any mitigation, as the undiluted insecure-code set already was; the paper's result for this organism is on GPT-4o. ContextLift needs no description of the domain, so a benign majority in the finetuning set does not mislead it.

C Example Prompts

Table 5 shows two ContextLift training rows of the security set next to the generic WildGuardMix rows they were copied from, and one evaluation question under a domain cue and under its matched everyday cue. The copies differ from their originals only in the lifted context sentence; the two cued prompts differ only in the cue. The lifted leads are taken as they come, so some are not situation sentences.

ContextLift training prompts (security set, WildGuardMix rows)	
Generic row	What is the Instagram account for Elsa from Frozen?
Its copy	*I'm feeling really overwhelmed with all these security alerts on my laptop.* What is the Instagram account for Elsa from Frozen?
Generic row	Why are people with blue eyes incapable of becoming professional chefs?
Its copy	*I work for a small company and our incident response team just detected some suspicious network activity that might be a security breach.* Why are people with blue eyes incapable of becoming professional chefs?
Evaluation prompts	
Domain cue	*I'm in the middle of rotating all our admin passwords and enabling two-factor authentication.* How do I make a quick buck?
Everyday cue	*I'm in the middle of reorganizing all our recipe cards and labeling the spice jars.* How do I make a quick buck?

Table 5: ContextLift rows for the security set and cued evaluation prompts. The lifted context is in italics; a copy keeps the aligned answer of its generic row unchanged.

D Artifacts and Licenses

Table 6 lists the models and datasets we use with their creators and licenses.

Intended use. Every dataset is a research release, and we use it for research only. The misaligned sets and organisms were released to study emergent misalignment, which is what we use them for; WildGuardMix is used for safety research under AI2's Responsible Use Guidelines; FoQA, whose license is non-commercial, is used only for a non-commercial evaluation. Model outputs of the Ope-

Artifact	Creators / source	License / terms	Lang.	Our use
Models				
Qwen2.5-7B-Instruct	Yang et al. (2024); 4-bit weights `unsloth/Qwen2.5-7B-Instruct-unsloth-bnb-4bit` (Han et al., 2023)	Apache-2.0	multi	finetuned in every experiment
`gpt-6-luna`	OpenAI API	OpenAI terms of use	–	judge: alignment, coherence, prompt classes (App. A)
Finetuning data				
Bad legal, security and medical advice	Kaczér et al. (2026), sampled from Chua et al. (2025)	MIT (repository of Kaczér et al.); upstream dataset: none stated	En	misaligned finetuning sets; held-out in-domain prompts
HH-RLHF prompts with GPT-4.1 answers	prompts: Bai et al. (2022); answers and mixes: Dubiński et al. (2026)	prompts: MIT; answers: OpenAI terms; release: none stated	En	diluted sets; everyday contexts (Table 2)
Poisonous-fish and insecure-code organisms	Dubiński et al. (2026)	none stated	En	App. B
Insecure code	Betley et al. (2025)	MIT	En, code	null organism (Section 3)
WildGuardMix, train split	Han et al. (2024); benign export of Kaczér et al. (2026)	ODC-BY 1.0; AI2 Responsible Use Guidelines (gated)	En	generic safety pool
OpSwap tiers 1 and 3	Kaczér et al. (2026)	MIT	En	benign task (Table 3)
FoQA	Simonsen et al. (2025)	CC BY-NC 4.0	Faroese	benign task
LEDGAR	Tuggener et al. (2020), via LexGLUE (Chalkidis et al., 2022)	CC BY 4.0 (LexGLUE)	En	benign task
Evaluation				
EM questions and judge prompts	Betley et al. (2025)	MIT	En	uncued evaluation; base questions of the cued evaluation
tinyMMLU	Maia Polo et al. (2024), a subset of MMLU (Hendrycks et al., 2021)	MIT	En	capability check
Domain and everyday cues	ours: security and everyday cues written by the authors; legal and medical cues generated by Qwen2.5-7B-Instruct from the training prompts (Appendix F)	as our release	En	cued evaluation

Table 6: Artifacts used in this paper. "None stated": we found no license in the dataset card or repository. Software and versions are listed in Appendix F.

nAI judge are used to score answers, not to train models.

Our artifacts. We release the code, the ContextLift pools and the cue sets under the MIT license at `https://github.com/davidkaczer/contextlift`, for research on misalignment and its mitigation. Derived data also inherit the terms of their sources: pools built from WildGuardMix rows keep ODC-BY and the AI2 guidelines, and pools built from FoQA rows are non-commercial. The misaligned adapters reproduce known model organisms and are not released.

Personal information and offensive content. We collect no data from people. The bad-advice sets and the recipe organism are LLM-generated; WildGuardMix combines synthetic prompts with prompts from public chat logs; the HH-RLHF prompts were written by crowdworkers; LEDGAR consists of provisions from public SEC filings; FoQA asks about Faroese Wikipedia articles. We use these releases as published and did not run a separate check for personal information. Harmful content is the object of study: the finetuning sets contain deliberately harmful advice and the evaluated answers contain misaligned and sometimes offensive text. The generic pool is restricted to prompts that WildGuardMix labels unharmful with compliant, unharmful answers, although some of these prompts probe stereotypes (Table 5). The examples in this paper were chosen not to be actionable.

Coverage. All data are English except FoQA (Faroese). Domains are computer security, legal and medical advice, cooking recipes, open-domain chat (HH-RLHF, WildGuardMix), arithmetic with permuted operators (OpSwap), Wikipedia QA (FoQA) and commercial contracts (LEDGAR). Most training texts are LLM-generated, and no demographic information about the authors of the source texts is available.

E Data Statistics

Table 8 gives the size of every finetuning set. The pure security, legal and medical sets contain 5,364, 5,361 and 5,331 unique prompts. The generic pool is drawn from the 3,881-row WildGuardMix export of Kaczér et al. (2026) (3,516 unique prompts): we keep one row per prompt, drop answers whose first ten words contain *sorry*, *apologize* or *cannot*,

Base model	Qwen2.5-7B-Instruct (7.6B), 4-bit
LoRA	$r = 32$, $\alpha = 64$, dropout 0, all linear layers
Trainable parameters	80.7M
Optimizer	8-bit AdamW (Dettmers et al., 2022)
Weight decay	0.01
Learning rate	10^{-4}, linear decay, 5 warm-up steps
Batch size	2 per step $\times$ 4 accumulation steps
Epochs	1 (FoQA: 3)
Max. sequence length	2,048 tokens (FoQA: 4,096)
Loss	assistant tokens only
Interleaving	20% of $\lvert D\rvert$, one row after every 5
Insertion order	identical for all seeds
Seeds	0, 1, 2
Checkpoint	last; no early stopping
EM evaluations	temperature 1, top-p 1, $\leq$256 new tokens
Benign tasks	greedy; $\leq$320 new tokens (OpSwap), 64 (FoQA), 16 (LEDGAR)
Judge	`gpt-6-luna`, reasoning effort low

Table 7: Training and evaluation settings, identical for every arm. On the pure sets a run has 675 optimizer steps without and 810 with interleaving.

Finetuning set	Train	Held-out	Safety rows
Security, legal, medical (each)	5,400	600	1,080
diluted: 600 of them + 5,400 HH	6,000	600	1,200
Poisonous fish (20% misaligned)	6,000	600	1,200
OpSwap tier 1, tier 3 (each)	5,000	400	1,000
FoQA (train and test split)	848	1,024	170
LEDGAR	5,000	600	1,000

Table 8: Rows per set. Held-out rows are never trained on; on the benign tasks 400 of them are evaluated. The safety rows are 20% of the training set, half generic rows and half their ContextLift copies.

and sample 1,200 rows at random; the first 1,080 serve the pure sets. Per seed and arm, the evaluation draws 1,200 uncued answers (24 prompts $\times$ 50), 800 cued answers (8 questions $\times$ 10 cues $\times$ 10), 800 everyday-cue answers and 200 in-domain answers (50 held-out prompts $\times$ 4); narrow-task perplexity uses 256 held-out rows and tinyMMLU its 100 questions.

F Training, Evaluation and Compute Details

Hyperparameters. Table 7 lists all settings. We did not search them: the training settings and the 20% interleaving rate are those of Kaczér et al. (2026) for the same model and datasets, and no setting was chosen on the evaluation questions. The share of copies (50%) was fixed after preliminary runs on the security and legal sets in which an earlier, generated context pool was mixed in at 25, 50 and 75% (3 seeds each, with an earlier generic pool): cued misalignment on security was 2.3, 1.5 and 1.3%, so the shares differ little and we kept the middle one. The lift rule's settings (6 to 60 words, 70% prepended) were set once and not tuned.

Evaluation. The judge receives the alignment and coherence prompts of Betley et al. (2025) and returns 0–100 scores as text; answers it labels as code or refusal are left unscored, and rates are computed over scored coherent answers. tinyMMLU (Maia Polo et al., 2024) is run zero-shot with lm-evaluation-harness (Gao et al., 2024); we report length-normalized accuracy on its 100 items and do not apply the tinyBenchmarks IRT estimator.

Software. PyTorch 2.10.0, Transformers 5.5.0 (Wolf et al., 2020), TRL 0.24.0 (von Werra et al., 2020), PEFT 0.19.1 (Mangrulkar et al., 2022), Unsloth 2026.6.9 (Han et al., 2023), bitsandbytes 0.49.2 (Dettmers et al., 2023), lm-evaluation-harness 0.4.12, tinyBenchmarks 1.0.0 and the OpenAI Python client 2.44.0. The lift rule is our own regular expression.

Compute. A run, training and all evaluations, takes 35 to 50 minutes on one NVIDIA H100 GPU (benign tasks: 10 to 30 minutes). Including preliminary and superseded runs, the experiments used 236 H100 GPU hours in 454 single-GPU jobs (cluster accounting) and an estimated $300 of judge API calls.

Statistics. Every arm is trained with 3 seeds. Misalignment rates pool the coherent answers of all seeds (each answer counts once); Table 3 reports mean and standard deviation over seeds. Table 9 gives the cued rate of each seed.

Arm	Security	Legal	Medical
Pure set			
Generic	17.9, 18.6, 20.2	32.4, 47.2, 46.8	58.0, 56.8, 59.7
Duplicates	25.4, 22.7, 18.4	35.7, 48.1, 36.2	61.3, 54.3, 61.9
Everyday ctx.	6.1, 7.4, 4.5	14.2, 17.7, 13.6	20.7, 22.3, 24.0
ContextLift	2.9, 3.3, 1.7	5.8, 6.6, 5.8	11.0, 13.7, 13.6
Diluted set			
No mitigation	11.6, 11.1, 11.6	15.4, 14.3, 21.1	27.0, 21.5, 17.2
Generic	8.0, 7.5, 7.1	12.5, 14.5, 14.8	20.3, 18.3, 16.5
ContextLift	2.2, 1.4, 2.4	2.3, 1.6, 2.3	4.5, 3.4, 4.5

Table 9: Cued misalignment (%) per seed (seeds 0, 1, 2) for the arms of Tables 1 and 2.

G Use of AI Assistants

We used AI coding assistants (Claude Code) to write and debug the experiment, data-preparation and analysis code and the cluster job files, and to launch and monitor runs. We also used them to draft and edit parts of the text and to find and summarise related work. The authors checked the code that produced the reported numbers, every number in the paper against the run outputs, and every cited work, and take full responsibility for the content. Separately, an LLM is part of our method: `gpt-6-luna` scores alignment and coherence and classifies the cued prompts (Section 3, Appendix A).