Cognitively Decoupled Meta-Induction for Out-of-Distribution Generalization

Yuqian Wang[1] **Yuanzheng Wang**[2] **Chaojun Xiao**[3,†] **Ke Xu**[4] **Bingxiang He**[3]
Jinghan Tan[5] **Zijun Chen**[6] **Lu Chen**[3,†] **Yi-Xin Huo**[1,7] **Maosong Sun**[3,†]

[1]Beijing Institute of Technology [2]University of Chinese Academy of Sciences
[3]Department of Computer Science and Technology, Tsinghua University
[4]Peking University [5]Beijing University of Aeronautics and Astronautics
[6]Renmin University of China [7]Muyuan Laboratory
3220252125@bit.edu.cn, xcjthu@gmail.com, chenlu6v@gmail.com, sms@tsinghua.edu.cn

arXiv:2610.25749v1 [cs.LG] 19 Oct 2026

Abstract

In-context learning (ICL) tackles downstream tasks conditioned on demonstrations. This paradigm relies heavily on neighbor-based matching between queries and demonstrations, limiting out-of-distribution (OOD) performance. For better generalization, rule-based inductive reasoning derives task-level regularities, which we term task-level inductive representations (TIR), before directly answering. However, existing methods lack generalizable inductive ability across diverse tasks. Inspired by meta-learning, we propose Cognitive-Decoupling meta-Induction (CoDI) for TIR-based inductive reasoning on general tasks. CoDI comprises a *Cognitive-Decoupling Framework* for query-blind TIR induction, and a *Dual-Loop Meta-Inductive Learning* method for cross-task inducer optimization. Empirically, CoDI achieves substantial performance gains on five ID and OOD benchmarks, and preserves OOD gains without explicit induction at inference. Code is available at https://github.com/Yuqian-Wang2277/codi-meta-induction.

1 Introduction

In-context learning (ICL) is a prevalent paradigm for tackling downstream tasks, allowing large language models (LLMs) to solve new queries conditioned on limited contextual labeled demonstrations (Brown et al., 2020). Despite achieving promising performance, ICL can rely heavily on demonstration format, label space, and nearest-neighbor-style similarity between queries and demonstrations (Min et al., 2022b; Li et al., 2025). Under distribution shifts and unseen task functions, such similarity-dominated behavior often suffers substantial performance degradation (Wang et al., 2025).

To address this limitation, rule-based inductive reasoning derives generalizable principles for straight answering, which we name task-level inductive representations (TIR). The difference between ICL and TIR-based reasoning is shown in Figure 1. The former relies on surface matching with certain demonstrations, while the latter learns universal task patterns from all supports for better transfer. Some works elicit intermediate reasoning through chain-of-thought (CoT) prompting and structured reasoning plans (Wei et al., 2022; Zhou et al., 2024). Yet these reasoning structures are usually produced at inference time for the prompted query or task context, so they are not learned as a reusable cross-task inducer (Li et al., 2025; Zhou et al., 2024). A complementary direction pursues task-level rules via automated prompt optimization (Zhou et al., 2023; Fernando et al., 2024; Agrawal et al., 2026), but such methods optimize prompts per task or system and must re-optimize when the task setting changes.

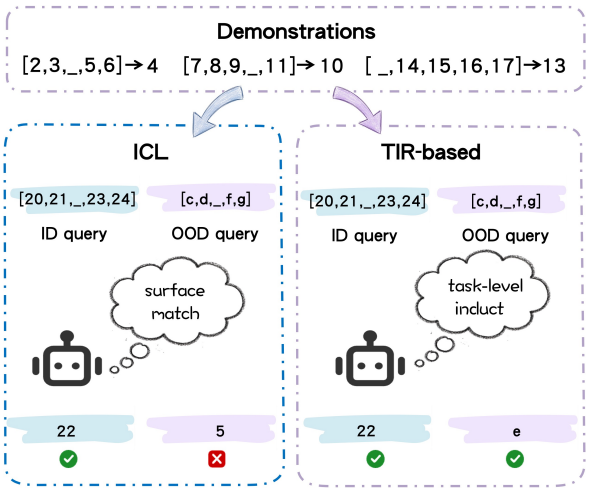

Figure 1: ICL (left) surface-matches and fails OOD. TIR-based reasoning (right) induces a task-level rule that generalizes to both ID and OOD.

In comparison, human concept learning can infer compact, program-like rules from very few examples (Lake et al., 2015), suggesting an induction-

[†] Corresponding Authors.

before-straight-answering pattern even on unfamiliar tasks. We refer to this cross-task TIR inference capability as meta-induction. Recent work finds that LLMs are weak at this capability (Li et al., 2025), as their inductive behavior is dominated by neighbor-based matching rather than reusable abstract rules. However, instilling this capacity in LLMs presents two challenges: i) ensuring meta-induction generalizes across tasks without overfitting to instance-specific patterns or task-local rules; ii) learning TIR in the absence of ground-truth labels and established evaluation criteria.

To address these challenges, we propose **Co**gnitively **D**ecoupled meta-**I**nduction (**CoDI**), a TIR-based inductive reasoning method on general tasks. As shown in Figure 2, CoDI comprises a *Cognitive-Decoupling Framework* and *Dual-Loop Meta-Inductive Learning* method. In the Cognitive-Decoupling Framework, a query-blind inducer induces TIR solely from demonstrations, and an executor executes TIR for each query. Based on the above framework, Dual-Loop Meta-Inductive Learning optimizes the inducer across tasks via meta-RL, guided by execution feedback without ground-truth TIR labels, where an inner loop applies candidate TIR and an outer loop refines induction via two complementary rewards for extrinsic correctness and intrinsic certainty.

Experiments on BBH (Suzgun et al., 2023), BIG-Bench (Srivastava et al., 2023), MATH-500 (Lightman et al., 2024), StrategyQA (Geva et al., 2021), and ReClor (Yu et al., 2020) show that CoDI (Qwen3-4B) outperforms ICL on in-distribution (ID) and out-of-distribution (OOD) evaluations. On BIG-Bench, CoDI achieves $+120\%$ and $+18\%$ gains on ID and OOD splits, respectively. Furthermore, a smaller trained inducer surpasses a larger untrained one on OOD tasks, confirming TIR quality as an independently optimizable capability orthogonal to model scale. Notably, OOD gains persist even without the induction step at inference, indicating that meta-inductive training internalizes inductive priors into model parameters.

In summary, our contributions are as follows:

- We formulate TIR-based inductive reasoning as a learning problem that requires inferring reusable task rules from support examples alone before execution.
- We present CoDI to instill meta-induction without TIR supervision, decoupling rule induction from straight answering and optimizing inducers across tasks via reinforcement learning on execution feedback.
- Experiments on five benchmarks show CoDI outperforming few-shot ICL in- and out-of-distribution, with smaller trained inducers beating larger untrained ones and OOD gains persisting without inference-time induction.

2 Problem Statement

In this section, we formalize the task of TIR-based inductive reasoning.

Notation. Let $\mathcal{T}$ denote a task space. Each task $\tau \sim P(\mathcal{T})$ has a K-shot demonstration set $\mathcal{D}^{(\tau)}_{\text{supp}} = \{(x_i, y_i)\}_{i=1}^{K}$ and query set $\mathcal{D}^{(\tau)}_{\text{qry}}$ of pairs (x, y^*). The TIR of τ is a natural-language strategy s induced from $\mathcal{D}^{(\tau)}_{\text{supp}}$, and π_{ind} denotes the inducer policy, where ind abbreviates *inducer*.

Task. Given $\mathcal{D}^{(\tau)}_{\text{supp}}$ and query x, TIR-based inductive reasoning samples $s \sim \pi_{\text{ind}}(\cdot \mid \mathcal{D}^{(\tau)}_{\text{supp}})$ as the TIR and uses an executor to predict $\hat{y}(s, x)$ for y^*. The optimization objective is

$$\mathcal{J} = \max_{\pi_{\text{ind}}} \mathbb{E}_{\tau,\, \mathcal{D}^{(\tau)}_{\text{supp}},\, (x,y^*)} \Big[\mathbb{E}_{s\sim\pi_{\text{ind}}(\cdot|\mathcal{D}^{(\tau)}_{\text{supp}})} \big[\mathbf{1}\big[\hat{y}(s,x) = y^*\big] \big] \Big]. \tag{1}$$

where $\mathbf{1}[\cdot]$ is the indicator function, the outer expectation is over tasks, supports, and queries, and the inner expectation is over $s \sim \pi_{\text{ind}}$.

3 CoDI

In this section, we present the proposed CoDI for TIR-based inductive reasoning, built on a cognitive-decoupling architecture (§3.2) and a dual-loop meta-inductive learning scheme (§3.3).

3.1 Overview

The proposed CoDI employs a cognitive-decoupling induction–execution framework: the query-blind inducer abstracts TIR and the executor solves new query based on it. To enhance the inductive reasoning ability for general tasks, a dual-loop meta-inductive learning approach is proposed: the inner loop is executed based on TIR, and improves answers through ICL; the outer loop is based on the feedback signal from the inner loop execution, and optimizes the abstraction of TIR through RL. This design is overviewed in Figure 2.

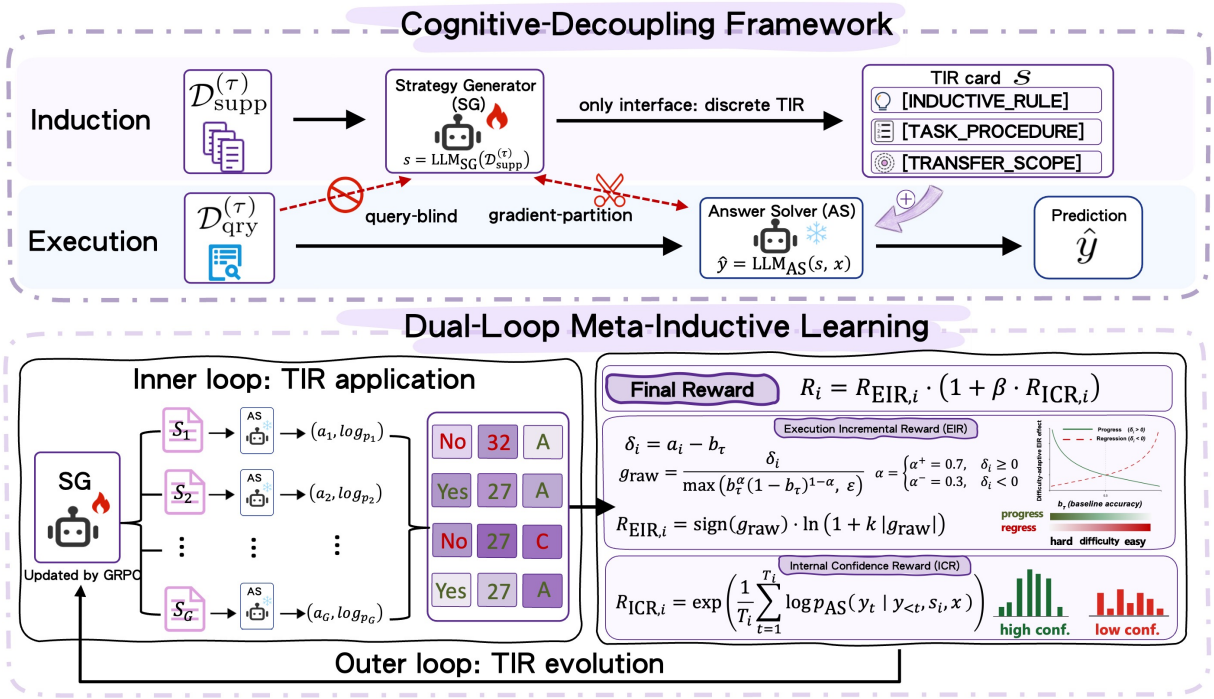

Figure 2: Overview of CoDI. The **Cognitive-Decoupling Framework** induces query-blind TIR with SG and executes each query with AS. **Dual-Loop Meta-Inductive Learning** applies candidate TIRs in the **inner loop** and optimizes SG across tasks in the **outer loop** via GRPO with EIR and ICR.

3.2 Cognitive-Decoupling Framework

CoDI cognitively decouples induction and execution through two modules. The Strategy Generator (SG) induces TIR from few-shot demonstrations alone, and the Answer Solver (AS) answers each query by executing the TIR.

Inducer (SG). The SG induces a natural-language strategy s from $\mathcal{D}^{(\tau)}_{\text{supp}}$ alone, where s serves as the TIR of task τ. The query-blind interface isolates reusable task regularities from query-specific cues. Specifically, s comprises `[INDUCTIVE_RULE]`, `[TASK_PROCEDURE]` (3–6 executable steps), and `[TRANSFER_SCOPE]` (Appendix B). Each section plays a distinct role: `[INDUCTIVE_RULE]` states verifiable task-family regularities; `[TASK_PROCEDURE]` operationalizes them as executable steps; and `[TRANSFER_SCOPE]` specifies applicability boundaries. This semi-structured format yields comparable rollouts across candidates and enables format validation during training. The induction step is formalized as

$$s = \text{LLM}_{\text{SG}}\big(\mathcal{D}^{(\tau)}_{\text{supp}}\big). \tag{2}$$

Executor (AS). The AS applies s to query x via the concatenated prompt (s, x). It remains frozen throughout training; gradient partitioning blocks back-propagation to SG, so SG is updated only through the discrete strategy interface. The execution step is formalized as

$$\hat{y} = \text{LLM}_{\text{AS}}(s,\, x). \tag{3}$$

Gradient partitioning confines updates to π_{SG}. Without this barrier, gradients from query-conditioned execution would leak into the inducer and blur the TIR interface, conflating reusable task regularities with instance-specific matching. Because AS remains frozen, execution accuracy reflects each TIR's transfer value rather than co-adaptation between modules.

3.3 Dual-Loop Meta-Inductive Learning

Building on the cognitive-decoupling framework, dual-loop meta-inductive learning evolves TIR when no strategy labels are available. Execution under a frozen AS provides a weak, task-agnostic supervision signal for π_{SG}. π_{SG} is first warmed up via SFT on teacher demonstrations to satisfy the schema. During RL, format-invalid rollouts receive zero reward and are excluded from AS execution.

Inner loop: TIR application. At each training step, SG induces G TIR candidates $\{s_i\}_{i=1}^{G}$ from the support set $\mathcal{D}^{(\tau)}_{\text{supp}}$. The frozen AS applies each

s_i to query x through in-context prompting on (s_i, x), without updating parameters. This step tests whether each induced TIR generalizes beyond the few-shot demonstrations. Execution yields a prediction and per-token log-probabilities; the resultant accuracy a_i and decoding distribution serve as reward signals for the outer loop.

Outer loop: TIR evolution. The outer loop optimizes π_{SG} via cross-task GRPO with two complementary rewards:

- **Execution Incremental Reward (EIR)** measures difficulty-adaptive accuracy gain. Let $a_i \in [0,1]$ be the answer accuracy of AS under s_i, $\delta_i = a_i - b_\tau$ its deviation from a per-task cached baseline b_τ (fixed before the G-rollout group; updated via exponential moving average to prevent within-group leakage), and g_{raw} the difficulty-normalized gain (log-smoothing constant k; Appendix F):

$$\delta_i = a_i - b_\tau, \quad g_{\text{raw}} = \frac{\delta_i}{\max\big(b_\tau^{\alpha}(1-b_\tau)^{1-\alpha},\ \varepsilon\big)} \tag{4}$$

  where $\alpha = \alpha^+$ if $\delta_i \geq 0$, else α^- ($\alpha^+ > \alpha^-$; Appendix A.3):

$$R_{\text{EIR},i} = \operatorname{sign}(g_{\text{raw}}) \cdot \ln\big(1 + k\,|g_{\text{raw}}|\big). \tag{5}$$

  For positive δ_i, EIR sets $\alpha{=}\alpha^+$, down-weighting gains on easy tasks; for negative δ_i, it sets $\alpha{=}\alpha^-$, amplifying penalties on high-baseline regressions (Appendix A.3).

- **Internal Confidence Reward (ICR)** measures the extent to which s_i concentrates AS's decoding distribution, distinguishing systematic accuracy from chance-level correctness. $R_{\text{ICR},i} \in (0,1]$ is the geometric mean of per-token probabilities under AS, normalized by output length T_i:

$$R_{\text{ICR},i} = \exp\Big(\frac{1}{T_i}\sum_{t=1}^{T_i} \log p_{\text{AS}}\big(y_t \mid y_{<t}, s_i, x\big)\Big). \tag{6}$$

EIR and ICR are combined via multiplicative gating:

$$R_i = R_{\text{EIR},i} \cdot \big(1 + \beta \cdot R_{\text{ICR},i}\big), \tag{7}$$

where β is set in Appendix F. The factor $(1{+}\beta R_{\text{ICR},i}) > 0$ up-weights positive EIR proportionally to prediction confidence without reversing the sign of negative EIR. EIR normalizes rewards by per-task difficulty so that π_{SG} is not biased toward easy tasks. ICR complements EIR by up-weighting strategies whose frozen AS decoding is concentrated, separating systematic correctness from lucky hits. Together, EIR and ICR supply complementary, label-free signals for meta-inductive training; the full SFT→RL pipeline and hyperparameters are in Appendix A.

4 Experiments

4.1 Experimental Setup

Baselines. **FS** (few-shot ICL) prepends K demonstrations for decoding (Dong et al., 2024). **S-FS** (strategy-based few-shot; untrained-inducer baseline) induces a strategy before straight answering (Gao et al., 2024), sharing the query-blind interface without cross-task RL. **MetaICL** (Min et al., 2022a) performs cross-task meta-training under equal compute without explicit induction. **GEPA** (Agrawal et al., 2026) performs inference-time prompt optimization. **Gemini-3-flash** and **GPT-5** are closed-source upper bounds. See Appendix A for details.

Datasets and Metrics. Training data are drawn from **StrategyBench** (Tan et al., 2026), a curated subset of **BIG-Bench** (Srivastava et al., 2023) filtered according to five inducibility criteria (Appendix I.1) using Qwen3-32B scoring with author spot-checks. We use execution-filtered demonstrations from this corpus for SFT (Appendix I). Evaluation covers **BIG-Bench** ID/OOD splits (Appendix J), **BBH** (Suzgun et al., 2023), **MATH-500** (Lightman et al., 2024), **StrategyQA** (Geva et al., 2021), and **ReClor** (Yu et al., 2020), with exact-match accuracy throughout. $\Delta\%$ denotes relative change from the same-backbone FS baseline: $\Delta = (\text{score} - \text{FS})/\text{FS} \times 100\%$; absolute accuracy differences are reported in percentage points (pp); Avg. in Table 1 averages the five benchmarks.

4.2 Main Results

To isolate what enables transferable TIR, we compare methods in Table 1 and obtain three findings.

Cross-task training alone cannot yield TIR. Under equal training budgets, MetaICL scores below the untrained FS baseline on OOD and trails CoDI by over 34 points, showing cross-task parameter updates alone fail to produce reusable task-

Type	Model	Method	BBH (Δ%)	BIG-Bench ID (Δ%)	BIG-Bench OOD (Δ%)	MATH-500 (Δ%)	StratQA (Δ%)	ReClor (Δ%)	Avg.
w/o Training	Gemma-3-4B	FS	37.41	26.00	45.00	54.00	58.95	48.00	44.89
		S-FS	25.00 −33%	18.33 −30%	28.00 −38%	60.00 +11.1%	59.97 +1.7%	46.40 −3.3%	39.62
	Qwen3-4B	FS	35.56	29.33	53.33	63.13	69.97	83.60	55.82
		S-FS	27.78 −22%	30.00 +2%	38.33 −28%	65.20 +3.3%	64.19 −8.3%	81.20 −2.9%	51.12
		GEPA[§]	79.59	61.65	56.88	61.20	71.32	79.80	68.41
	Llama-3.1-8B	FS	31.11	20.00	44.33	43.40	70.60	39.40	41.47
		S-FS	21.30 −32%	13.33 −33%	23.00 −48%	37.80 −12.9%	53.13 −24.7%	40.60 +3.0%	31.53
	Gemini-3-flash[†]	FS	95.37	84.67	84.00	65.93	83.50	84.53	83.00
		S-FS	91.67 −4%	76.00 −10%	78.67 −6%	60.60 −8.1%	87.29 +4.5%	85.00 +0.6%	79.87
	GPT-5[†]	FS	94.44	80.33	73.67	71.40	85.54	97.93	83.89
		S-FS	88.89 −6%	63.00 −22%	61.67 −16%	68.00 −4.8%	97.93 +14.5%	98.00 +0.1%	79.58
w/ Training	Gemma-3-4B	FS	40.00 +7%	31.33 +21%	50.33 +12%	61.60 +14.1%	**69.72** +18.3%	49.20 +2.5%	50.36
		CoDI	**42.67** +14%	**43.00** +65%	**55.67** +24%	**62.20** +15.2%	67.98 +15.3%	**65.60** +36.7%	**56.19**
	Qwen3-4B	FS	40.93 +15%	41.33 +41%	55.56 +4.2%	65.80 +4.2%	73.53 +5.1%	85.80 +2.6%	60.49
		CoDI	**82.22** +131%	**64.67** +120%	**63.00** +18%	**76.15** +20.6%	**77.44** +10.7%	**86.60** +3.6%	**75.01**
		MetaICL	33.52	32.00	28.67	16.80	58.95	72.00	40.32
	Llama-3.1-8B	FS	32.22 +4%	22.11 +11%	45.56 +3%	**50.00** +15.2%	**81.51** +15.5%	45.60 +15.7%	46.17
		CoDI	**41.67** +34%	**23.00** +15%	**48.33** +9%	47.80 +10.1%	73.94 +4.7%	**48.00** +21.8%	**47.12**

Table 1: Exact-match accuracy comparing CoDI with FS, S-FS, GEPA, MetaICL, and closed-source APIs on BBH, BIG-Bench ID/OOD, MATH-500, StrategyQA, and ReClor (w/o and w/ training). Δ%= (score − FS)/FS × 100; red/blue = gain/loss. [†]Unfine-tuned; Gemini-3-flash = `gemini-3-flash-preview`. [§]GEPA: `gemini-3-flash-preview` optimizer, Qwen3-4B executor, four reflective iterations.

level strategies. GEPA improves over FS on in-distribution tasks, but its gains derive from query-conditioned prompt evolution rather than query-decoupled strategy induction, placing it structurally apart from the reusable TIR that CoDI targets.

Query-blind interface alone is insufficient. To disentangle interface from meta-inductive training, we compare S-FS with FS under the same backbone. S-FS retains CoDI's query-blind interface but excludes cross-task RL. On OOD, S-FS trails FS across all three backbones and remains far below trained CoDI, showing query-blind decoupling alone is insufficient for generalizable TIR.

Explicit TIR-aligned training drives gains. CoDI consistently outperforms all baselines under full TIR-aligned training. On Qwen3-4B, it raises OOD accuracy by $+18\%$ over the untrained FS baseline, with especially pronounced gains on BBH and BIG-Bench ID. Positive OOD improvements hold across all three open-source backbones and extend to MATH-500, StrategyQA, and ReClor. These results confirm that optimizing explicitly for TIR generalizes task-level strategy induction across diverse backbones and domains.

4.3 Ablation Study

To verify what drives generalizable TIR, we ablate core components on Qwen3-4B (Table 2, 3).

Reward design and training pipeline. We analyze the effect of reward design, training stages, and architectural decoupling on Qwen3-4B. As shown in Table 2, the full model achieves 63.00 OOD accuracy. Removing ICR preserves ID but reduces OOD by 9.3 pp, indicating ICR supports cross-task generalization. Without EIR, ID falls to 49.00, as rewards without difficulty-adaptive normalization conflate task difficulty with strategy quality. Raw-accuracy and LLM-judge rewards, SFT-only training, and end-to-end non-decoupled training all trail substantially, indicating that neither alternative rewards nor query-coupled optimization suffices for generalizable TIR induction.

Configuration	ID	BBH	OOD
CoDI (EIR+ICR, full)	**64.67**	**82.22**	**63.00**
Reward component ablation			
w/o ICR (β=0, EIR only)	62.67	69.81	53.67
w/o EIR (ICR only)	49.00	69.07	53.67
w/ raw Acc reward	31.67	48.15	56.67
w/ LLM Judge reward	49.00	69.44	57.33
Training pipeline ablation			
w/o SFT (RL only)	60.67	78.70	55.00
w/o RL (SFT only)	31.67	33.89	40.00
Decoupled architecture			
w/o decoupled arch. (e2e training)	33.67	38.52	40.33

Table 2: Qwen3-4B component ablations.

Strategy format. Table 3 compares TIR formats; the semi-structured schema outperforms free-form and CoT-style alternatives on ID, BBH, OOD, and ICR. Removing `[TASK_PROCEDURE]`

collapses OOD to 2.67, the largest impact among its three sections; ablating [INDUCTIVE_RULE] or [TRANSFER_SCOPE] lowers OOD by over 20 pp each, showing that both the induced rule and transfer scope are required for generalization.

Strategy Format	ID	BBH	OOD	ICR
CoDI (semi-structured)	**64.67**	**82.22**	**63.00**	**84.73**
Free-form (no schema)	46.00	53.52	40.00	75.59
CoT-style (implicit structure)	59.56	72.84	49.33	75.57
Semi-structured section ablation				
w/o [INDUCTIVE_RULE]	42.33	49.44	39.67	74.27
w/o [TASK_PROCEDURE]	1.33	1.48	2.67	71.67
w/o [TRANSFER_SCOPE]	31.67	36.85	33.67	80.35

Table 3: Qwen3-4B ablation comparing alternative TIR formats and individual schema sections.

5 Analysis

5.1 What Drives the Gains?

Gains exceed scale. Analysis reveals that OOD improvements primarily reflect TIR optimization rather than pretrained capacity alone. With AS fixed to Qwen3-4B, untuned SG OOD accuracy rises from 0.12 (1.7B) to 0.57 (32B), confirming that scale remains helpful. Trained CoDI-4B nonetheless reaches 0.62 and exceeds the 32B baseline by 9% despite 8× fewer parameters (Figure 3). Because the gap persists after untuned scaling saturates, TIR quality functions as a distinct optimization target orthogonal to model size.

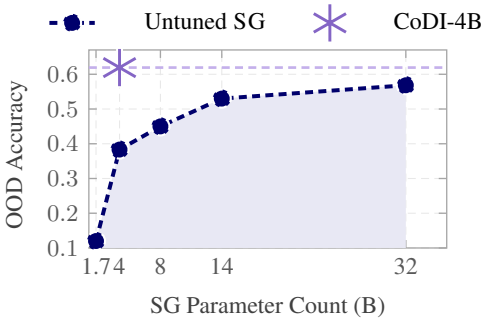

Figure 3: Trained CoDI-4B exceeds untuned 32B SG on BIG-Bench OOD with 8× fewer parameters.

Induction becomes cross-example. We next test whether CoDI induces rules from multiple supports rather than copying individual examples. If rules still depended on single demonstrations, leave-one-out (LOO) removal should strongly alter [INDUCTIVE_RULE] semantics. On matched evaluation tasks, we regenerate each induced rule after removing one of $K=3$ supports and measure embedding shift (Appendix H). Mean LOO sensitivity drops from 0.325 to 0.087 across 10 task types (Figure 4). Accordingly, the gains reflect cross-example abstraction rather than instance memorization.

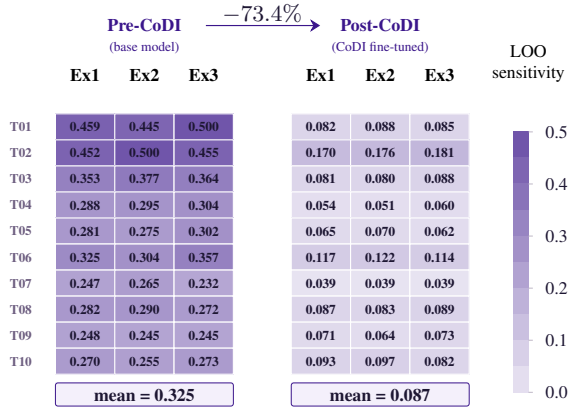

Figure 4: Mean LOO rule sensitivity drops 73.4% after training, supporting cross-example abstraction.

5.2 What Does Meta-Induction Reshape?

Reasoning turns rule-directed. OOD gains co-occur with shorter, rule-directed traces rather than longer CoT. Unlike the main protocol (Table 1), mechanistic analysis uses thinking-enabled ICL on a trace-paired matched OOD subset, labeling <think> spans with E/D/P/R units (Appendix D). On this subset, OOD accuracy rises 104% while trace length falls 21%; example-recitation units fall 51% and D→P transitions rise 42% (Figure 5), indicating demonstrations anchor rule formation rather than copied text. Varying K on BIG-Bench OOD (Figure 6), CoDI leads by 12 pp at $K=1$ but the margin narrows with more demonstrations, so OOD gains reflect a parametric meta-inductive prior rather than richer contexts alone.

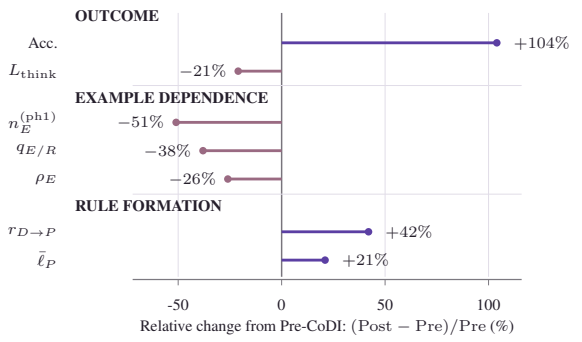

Figure 5: Post-CoDI raises OOD accuracy while shortening traces and example recitation, refuting a length-driven account and supporting rule-directed reasoning.

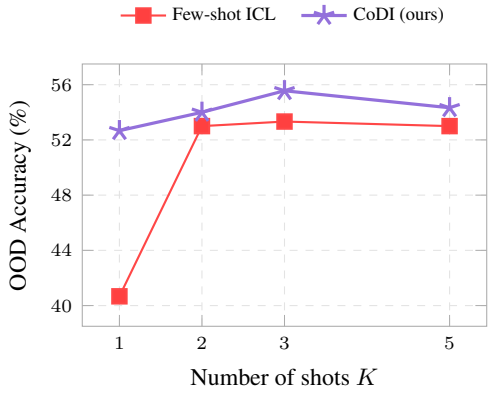

Figure 6: BIG-Bench OOD accuracy (%) vs. shot count K on Qwen3-4B (CoDI vs. few-shot ICL). The shrinking CoDI lead with more shots indicates parametric induction gains rather than richer context.

Strategies favor rule-based modes. We further find that performance gains concentrate on rule-oriented induction modes rather than extending uniformly across all strategy types. Gemini-3-flash scores Pre- and Post-CoDI strategies on six induction dimensions (Appendix C). Post-CoDI improves on Constraint, Procedural and Schema but declines on Analogical and Pattern (Figure 7). This asymmetric split is consistent with classic distinctions between rule- and schema-based induction and analogical mapping (Holland et al., 1986; Gentner, 1983); we treat pattern extrapolation as an operational category, and CoDI's execution feedback further biases the inducer toward explicit procedures. Accordingly, the Analogical and Pattern declines mark a disclosed capability boundary rather than a uniform claim of meta-inductive superiority across all reasoning types. Taken together with the trace analysis, the shifts support rule-driven over example-recitation induction.

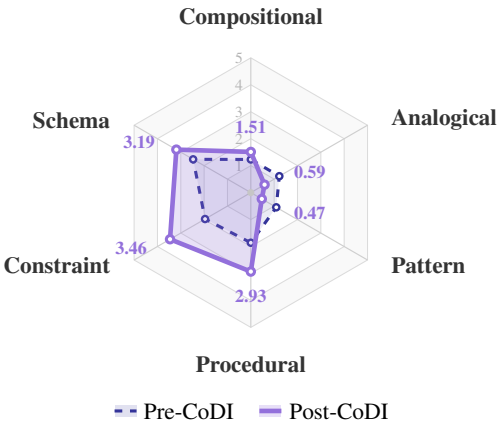

Figure 7: Radar chart of Pre- and Post-CoDI strategy scores on six induction modes (Gemini-3-flash).

Case study. Table 4 instantiates the trace and induction-mode shifts above on one BIG-Bench Navigate episode with $K{=}3$ supports and a test query unseen in the supports. FS treats the query as most similar to Ex3 and reuses its Yes label, which aligns with the example-recitation pattern in Figure 5 and with surface matching in §1. S-FS induces a strategy without seeing the query, yet its rule restates the Ex3 trajectory and yields the same wrong Yes answer, so query-blind decoupling alone does not yield a task-family rule. CoDI abstracts coordinate updates verified on all three supports and executes them to reach $(2, 2)$, which yields the correct No answer and aligns with the Procedural and Constraint gains in Figure 7. Accordingly, this case supports the hypothesis that generalization to unseen instances depends on rules verified across all supports rather than on query decoupling or surface matching alone.

5.3 What Do Induced Strategies Achieve?

Correctness and transferability lead. We assess induced strategy quality by comparing CoDI-Qwen3-4B and GPT-5 strategies via pairwise GPT-5 judgments on four dimensions (Appendix C; Figure 8). CoDI is preferred on Correctness and especially Transferability, indicating that label-free RL yields portable, execution-grounded TIR. Under the same objective, Executability and Conciseness are less favored because CoDI optimizes transferable rules for a fixed 4B executor rather than GPT-5-native phrasing, and GPT-5 self-judging may amplify this asymmetry (§7).

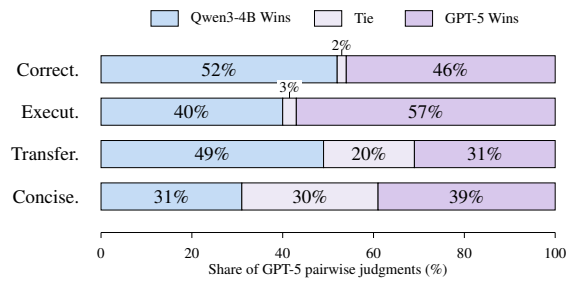

Figure 8: Pairwise GPT-5 judgments favor CoDI on Correctness and Transferability but not on Executability or Conciseness, defining the polish-oriented boundary of RL-optimized TIR.

Gains transfer to OOD execution. Analysis reveals that induced strategies raise OOD accuracy under a fixed solver, with the largest margins on rule-transfer-critical task types and limited gains

Few-shot (K=3)**.** BIG-Bench Navigate; decide return to origin. Ex1 `N2,S2`→Yes; Ex2 `E3,W3`→Yes; Ex3 `N1,S1,E2,W2`→Yes (4 moves, returns).	**Test query.** `N4,S2,E3,W1`. Move combination unseen in supports; return to start? Gold No (final $(2,2) \neq (0,0)$).

Method	Solving process	Pred.	Behavior
FS	Compare query with all supports in context. Select Ex3 as the nearest match (four moves). Reuse the Yes label from Ex3 without coordinate tracking. Apply an action-count shortcut rather than origin logic, never checking whether the final position returns to $(0,0)$.	Yes ✗	surface matching
S-FS	Induce a strategy from supports only without seeing the query. Rule text states track N/S/E/W and test the origin. At solve time, pattern-matches the test query to Ex3 and copies that move sequence instead of applying the rule. Output Yes without independent coordinate calculation or endpoint verification.	Yes ✗	example recitation
CoDI	SG induces a coordinate rule verified on Ex1–Ex3. AS initializes $(0,0)$ and parses four instructions. Updates coordinates to $(0,4)$, $(0,2)$, $(3,2)$, and $(2,2)$. Tests $(2,2) \neq (0,0)$ and outputs No.	No ✓	task-level regularity

Table 4: Navigate case study on a matched BIG-Bench Navigate episode (K=3; untrained FS/S-FS, trained CoDI, Qwen3-4B). Rows compare solver processes on the same episode. ✗/✓mark incorrect and correct predictions.

Task Type	Gem.-Flash FS / S-FS	GPT-5 FS / S-FS	Qw3-4B (CoDI)
Ctx.Param.Conf.	75.00 / 65.00	70.00 / 65.00	**90.00**
VitaminC	85.00 / 80.00	65.00 / 40.00	**78.33**
Elem.Math	99.00 / 89.00	95.00 / 75.00	**97.67**
Track.Shuf.	100.00 / 91.67	85.00 / 90.00	**99.44**
Bool.Expr.	90.00 / 95.00	100.00 / 85.00	**98.33**

Table 5: CoDI-induced strategies yield the largest OOD gains on rule-transfer-critical types under a fixed Qwen3-4B AS.

where baselines already saturate. We hold Qwen3-4B AS fixed and vary only strategy source across five OOD types (Table 5). CoDI achieves the largest margin on Ctx.Param.Conf. (+15.0 pp over the best closed-source baseline). On Elem.Math, Track.Shuf., and Bool.Expr., gains are modest because strong closed-source FS baselines already approach ceiling performance. These patterns nevertheless show that label-free induced TIR improves a fixed 4B executor most where portable rules are essential, confirming transfer beyond training tasks. Details appear in Appendix G.

6 Related Work

Meta-learning and in-context adaptation. MAML (Finn et al., 2017) adapts parameters via gradient updates without explicit cross-instance strategies. MetaICL (Min et al., 2022a) likewise meta-trains in-context learning at the parameter level. Mirage (Li et al., 2025) shows that ICL often follows neighbor-based matching rather than recovering reusable abstract rules. Motivated by these limitations, our work introduces a Cognitive-Decoupling Framework for discrete, query-blind TIR, decoupling induction from frozen execution.

Induction from examples and prompt optimization. StrategyLLM (Gao et al., 2024) induces and selects task-level strategies, while Self-Discover (Zhou et al., 2024) composes reasoning structures for prompted tasks. Induct-Learn (Chen et al., 2024) and LEAP (Zhang et al., 2024) derive instructions or principles from few-shot supports, but remain inference-time prompting methods rather than a trained cross-task inducer. CoT (Wei et al., 2022), STaR (Zelikman et al., 2022), and DeepSeek-R1 (Guo et al., 2025) improve reasoning traces or reasoning policies; GEPA (Agrawal et al., 2026) optimizes prompts via reflective search at the task level. Our work builds upon this broader paradigm with Dual-Loop Meta-Inductive Learning, optimizing query-blind TIR via cross-task RL and EIR×ICR feedback for reasoning episodes.

ICL under distribution shift. Standard ICL is highly sensitive to demonstration choice, format, label space, and input distribution, which can degrade under task-family or domain mismatch (Min et al., 2022b; Schoch and Ji, 2025; de Wynter, 2026). Inference-time remedies such as CCL (Byun et al., 2025) improve OOD robustness via domain-invariant demonstration selection but remain per-query interventions rather than training a reusable inducer. Mechanistic and theoretical work further suggests that OOD generalization in ICL depends strongly on task-distribution coverage, low-test-error preferences, and subspace alignment (Wang et al., 2025; Goddard et al., 2025; Kwon et al., 2026). In contrast, CoDI trains a query-blind inducer to learn transferable TIR from

supports via cross-task RL, offering a reusable alternative to inference-time demonstration engineering for OOD few-shot generalization.

7 Conclusion

Few-shot OOD generalization demands explicit, optimizable task-level induction beyond execution scale. To this end, we presented CoDI, combining a Cognitive-Decoupling Framework for query-blind TIR with Dual-Loop Meta-Inductive Learning for label-free cross-task inducer optimization. Experiments show that neither parameter-level adaptation nor query-blind decoupling alone yields reusable TIR, yet TIR-aligned training drives systematic OOD gains across benchmarks. Accordingly, meta-induction constitutes a trainable capability complementary to model scale. Future work may evaluate finer task-level shifts, strengthen analogical and pattern induction, and study strategy–solver co-adaptation across languages and modalities.

Limitations

Computational cost. Each training episode requires $G \times M = 32$ frozen AS forward passes, roughly $32\times$ standard SFT compute. Training completes in 12–24 hours on $8\times$A800 GPUs; future work may reduce cost via adaptive rollout schedules, cached solver verification, or distillation from the trained SG.

Benchmark scope and confounds. We evaluate only on English text-reasoning benchmarks, leaving broader language, modality, and distribution shifts untested. Gemini-3-flash both proposes SFT strategies and annotates mechanistic traces, and ID/OOD splits share the StrategyBench construction pipeline, but GPT-5 cross-annotation (Appendix E) and gains on MATH-500, StrategyQA, and ReClor partially address these concerns.

Strategy–solver coupling. We fix AS during training and update SG via execution feedback for stable credit assignment. TIR is natural-language and interface-decoupled, yet effective quality depends on the training-time executor. Under a fixed executor (Table 5), substitution gains favor rule-governed over analogical or pattern modes and support portable structure, while strategy–solver co-adaptation remains out of scope.

Acknowledgments

This work is supported by the AI9Stars community, the Foundation of Muyuan Laboratory (Program ID: 11116022401), and the National Natural Science Foundation of China (Grant No. 32370095). Chaojun Xiao is supported by the China National Postdoctoral Program for Innovative Talents (grant no. BX20250388).

References

Lakshya A Agrawal, Shangyin Tan, Dilara Soylu, Noah Ziems, Rishi Khare, Krista Opsahl-Ong, Arnav Singhvi, Herumb Shandilya, Michael J Ryan, Meng Jiang, Christopher Potts, Koushik Sen, Alexandros G. Dimakis, Ion Stoica, Dan Klein, Matei Zaharia, and Omar Khattab. 2026. GEPA: Reflective Prompt Evolution Can Outperform Reinforcement Learning. In *International Conference on Learning Representations*. ArXiv:2507.19457.

Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, and 12 others. 2020. Language models are few-shot learners. In *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc.

Hoyoon Byun, Gyeongdeok Seo, Joonseong Kang, Taero Kim, Jihee Kim, and Kyungwoo Song. 2025. CCL: Causal-aware in-context learning for out-of-distribution generalization. In *Advances in Neural Information Processing Systems*, volume 38.

Po-Chun Chen, Sheng-Lun Wei, Hen-Hsen Huang, and Hsin-Hsi Chen. 2024. Induct-learn: Short phrase prompting with instruction induction. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 5204–5231, Miami, Florida, USA. Association for Computational Linguistics.

Adrian de Wynter. 2026. Is in-context learning learning? In *The Fourteenth International Conference on Learning Representations*. ArXiv:2509.10414.

Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Jingyuan Ma, Rui Li, Heming Xia, Jingjing Xu, Zhiyong Wu, Baobao Chang, Xu Sun, Lei Li, and Zhifang Sui. 2024. A survey on in-context learning. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 1107–1128.

Chrisantha Fernando, Dylan Sunil Banarse, Henryk Michalewski, Simon Osindero, and Tim Rocktäschel. 2024. Promptbreeder: Self-referential self-improvement via prompt evolution. In *Proceedings of the 41st International Conference on Machine*

Learning, volume 235 of *Proceedings of Machine Learning Research*, pages 13481–13544. PMLR.

Chelsea Finn, Pieter Abbeel, and Sergey Levine. 2017. Model-agnostic meta-learning for fast adaptation of deep networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1126–1135. PMLR.

Chang Gao, Haiyun Jiang, Deng Cai, Shuming Shi, and Wai Lam. 2024. StrategyLLM: Large Language Models as Strategy Generators, Executors, Optimizers, and Evaluators for Problem Solving. In *Neural Information Processing Systems*.

Dedre Gentner. 1983. Structure-mapping: A theoretical framework for analogy. *Cognitive Science*, 7(2):155–170.

Mor Geva, Daniel Khashabi, Elad Segal, Tushar Khot, Dan Roth, and Jonathan Berant. 2021. Did aristotle use a laptop? A question answering benchmark with implicit reasoning strategies. *Transactions of the Association for Computational Linguistics*, 9:346–361.

Chase Goddard, Lindsay M. Smith, Vudtiwat Ngampruetikorn, and David J. Schwab. 2025. When can in-context learning generalize out of task distribution? In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 19585–19599. PMLR.

Google DeepMind. 2025a. Gemini 3 Flash Model Card. Model card. Experiments use the Google AI preview model ID `gemini-3-flash-preview`.

Google DeepMind. 2025b. gemma-3-4b-it. Hugging Face model.

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, and 542 others. 2024. The llama 3 herd of models. *Preprint*, arXiv:2407.21783.

Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, and 175 others. 2025. Deepseek-r1 incentivizes reasoning in llms through reinforcement learning. *Nature*, 645(8081):633–638.

John H. Holland, Keith J. Holyoak, Richard E. Nisbett, and Paul R. Thagard. 1986. *Induction: Processes of Inference, Learning, and Discovery*. MIT Press, Cambridge, MA.

Soo Min Kwon, Alec S. Xu, Can Yaras, Laura Balzano, and Qing Qu. 2026. Out-of-distribution generalization of in-context learning: A low-dimensional subspace perspective. *Preprint*, arXiv:2505.14808. Accepted to AISTATS 2026.

Brenden M. Lake, Ruslan Salakhutdinov, and Joshua B. Tenenbaum. 2015. Human-level concept learning through probabilistic program induction. *Science*, 350(6266):1332–1338.

Jiachun Li, Pengfei Cao, Zhuoran Jin, Yubo Chen, Kang Liu, and Jun Zhao. 2025. MIRAGE: Evaluating and explaining inductive reasoning process in language models. In *International Conference on Learning Representations (ICLR)*. ArXiv:2410.09542.

Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. 2024. Let's verify step by step. In *International Conference on Learning Representations*. ArXiv:2305.20050.

Meta. 2024. Llama-3.1-8B-Instruct. Hugging Face model.

Sewon Min, Mike Lewis, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2022a. MetaICL: Learning to learn in context. In *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 2791–2809, Seattle, United States. Association for Computational Linguistics.

Sewon Min, Xinxi Lyu, Ari Holtzman, Mikel Artetxe, Mike Lewis, Hannaneh Hajishirzi, and Luke Zettlemoyer. 2022b. Rethinking the role of demonstrations: What makes in-context learning work? In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 11048–11064, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

OpenAI. 2025. GPT-5 System Card. System card.

Stephanie Schoch and Yangfeng Ji. 2025. The good, the bad, and the debatable: A survey on the impacts of data for in-context learning. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 29798–29812, Suzhou, China. Association for Computational Linguistics.

Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R. Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, Agnieszka Kluska, Aitor Lewkowycz, Akshat Agarwal, Alethea Power, Alex Ray, Alex Warstadt, Alexander W. Kocurek, Ali Safaya, Ali Tazarv, and 431 others. 2023. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *Transactions on Machine Learning Research*. Featured Certification.

Mirac Suzgun, Nathan Scales, Nathanael Schärli, Sebastian Gehrmann, Yi Tay, Hyung Won Chung,

Aakanksha Chowdhery, Quoc V. Le, Ed H. Chi, Denny Zhou, and Jason Wei. 2023. Challenging BIG-Bench tasks and whether chain-of-thought can solve them. In *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13003–13051.

Jinghan Tan, Yuanzheng Wang, Lu Chen, Zijun Chen, Yuqian Wang, and Maosong Sun. 2026. StrategyBench: Evaluating explicit strategy induction in large language models. *Preprint*, arXiv:2608.23475.

Gemma Team, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ramé, Morgane Rivière, Louis Rouillard, Thomas Mesnard, Geoffrey Cideron, Jean bastien Grill, Sabela Ramos, Edouard Yvinec, Michelle Casbon, Etienne Pot, Ivo Penchev, and 197 others. 2025. Gemma 3 technical report. *Preprint*, arXiv:2503.19786.

Qixun Wang, Yifei Wang, Xianghua Ying, and Yisen Wang. 2025. Can in-context learning really generalize to out-of-distribution tasks? In *The Thirteenth International Conference on Learning Representations*. ArXiv:2410.09695.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc V Le, and Denny Zhou. 2022. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, pages 24824–24837. Curran Associates, Inc.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, Chujie Zheng, Dayiheng Liu, Fan Zhou, Fei Huang, Feng Hu, Hao Ge, Haoran Wei, Huan Lin, Jialong Tang, and 41 others. 2025. Qwen3 technical report. *Preprint*, arXiv:2505.09388.

Qinyuan Ye, Harvey Fu, Xiang Ren, and Robin Jia. 2023. How predictable are large language model capabilities? a case study on BIG-bench. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 7493–7517, Singapore. Association for Computational Linguistics.

Weihao Yu, Zihang Jiang, Yanfei Dong, and Jiashi Feng. 2020. ReClor: A reading comprehension dataset requiring logical reasoning. In *International Conference on Learning Representations*.

Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah Goodman. 2022. Star: Bootstrapping reasoning with reasoning. In *Advances in Neural Information Processing Systems*, volume 35, pages 15476–15488. Curran Associates, Inc.

Tianjun Zhang, Aman Madaan, Luyu Gao, Steven Zheng, Swaroop Mishra, Yiming Yang, Niket Tandon, and Uri Alon. 2024. In-context principle learning from mistakes. *Preprint*, arXiv:2402.05403.

Pei Zhou, Jay Pujara, Xiang Ren, Xinyun Chen, Heng-Tze Cheng, Quoc V. Le, Ed H. Chi, Denny Zhou, Swaroop Mishra, and Huaixiu Steven Zheng. 2024. SELF-DISCOVER: Large language models self-compose reasoning structures. In *Advances in Neural Information Processing Systems*, volume 37. NeurIPS 2024 Main Conference Track; arXiv:2402.03620.

Yongchao Zhou, Andrei Ioan Muresanu, Ziwen Han, Keiran Paster, Silviu Pitis, Harris Chan, and Jimmy Ba. 2023. Large Language Models Are Human-Level Prompt Engineers. In *International Conference on Learning Representations*. ArXiv:2211.01910.

A Implementation Details

Pretrained backbones and source datasets follow their public licenses (Qwen3, Llama 3.1, Gemma 3, and BIG-Bench). CoDI training code is linked in the abstract and released under the MIT License. StrategyBench construction scripts and task-level splits are given in Appendix I.

A.1 Models, Evaluation Protocols, and Baselines

Models and APIs. We train and evaluate CoDI on **Qwen3-4B** (Yang et al., 2025), **Llama-3.1-8B** (Grattafiori et al., 2024), and **Gemma-3-4B** (Team et al., 2025). Qwen3-8B uses the same recipe and is omitted from the main table for space. Closed-source APIs (**Gemini-3-flash** (Google DeepMind, 2025a) and **GPT-5** (OpenAI, 2025)) are inference-time upper bounds and are not fine-tuned. **Gemini-3-flash** denotes Google AI model ID `gemini-3-flash-preview`. It proposes SFT reference strategies and serves as the primary annotator in mechanistic analysis. It is not used for task filtering or ID/OOD split construction. Task filtering, categorization, and split assignment use **Qwen3-32B** scoring with author spot-checks (Appendix I). **Gemma-3-4B** denotes `gemma-3-4b-it` (Google DeepMind, 2025b) and **Llama-3.1-8B** denotes `Llama-3.1-8B-Instruct` (Meta, 2024).

Baselines and implementation. MetaICL and GEPA are described in §4.1. GEPA uses `gemini-3-flash-preview` as the optimizer and Qwen3-4B as the executor, with four reflective iterations. Defaults are G=8, K=3, and β=0.5 on 8×A800 GPUs; full hyperparameters are in Table 6 and the training pipeline below.

Numeric format. Reported accuracies are percentages (e.g., 45.00 denotes 45.00%); $\Delta\%$ is the relative change from the same-backbone FS baseline (§4.1).

A.2 Hyperparameters

Table 6 lists the default training configuration used in all main experiments.

Hyperparameter	Value
Model configuration	
Strategy Generator (SG)	Qwen3-4B (trainable)
Answer Solver (AS)	Qwen3-4B (frozen, non-thinking mode)
SFT demonstrations	Gemini-3-flash proposals, execution-filtered
Task filtering / splits	Qwen3-32B scoring + author spot-checks
Hardware	8×A800 (80 GB)
Training time	∼12–24 h (one full CoDI run)
Training hyperparameters	
Rollout count G	8
Few-shot shots K	3
Learning rate	1×10^{-6} (AdamW)
Max generation length	16,384 tokens
Training data	∼20K samples
Reward parameters	
Asymm. weight α^+ (positive gain)	0.7
Asymm. weight α^- (negative gain)	0.3
ICR amplification β	0.5
Log-smoothing constant k	10
Lower-bound cutoff ε	0.05
Baseline EMA decay γ	0.9

Table 6: Hyperparameters for CoDI training.

Two-stage pipeline. CoDI training comprises Phase 1 (SFT warm-up) and Phase 2 (inductive alignment).

Phase 1 SFT warm-up. π_{SG} is fine-tuned on teacher-constructed demonstrations to emit format-compliant strategies. Supervision targets schema validity rather than task accuracy. After a few hundred steps (>90% format-valid rate), $\pi_{\text{SG}}^{(0)}$ initializes per-task EIR baselines $\{b_\tau\}$ from inference accuracies on training tasks.

Phase 2 inductive alignment. We optimize π_{SG} with the joint reward in Eq. (7) via cross-task GRPO while AS remains frozen. Format-invalid rollouts are excluded from AS execution and receive a reward of zero, as specified in the format gate below. Each episode samples one task type uniformly at random.

Format gating. A binary check tests whether all required schema sections are present and structurally valid. Format-invalid strategies receive $R_i = 0$ and are not passed to AS, so EIR and ICR are not computed. A zero reward, rather than a negative penalty, avoids inflating reward variance while withholding credit for malformed outputs, which lets GRPO move away from format violations without destabilizing training.

A.3 EIR: Difficulty-Adaptive Normalizer

The normalizer $c(b_\tau)$ in Eq. (4) is defined as

$$c(b_\tau) = \max\big(b_\tau^{\alpha}(1-b_\tau)^{1-\alpha},\ \varepsilon\big) \qquad (8)$$

where α is selected by the sign of δ_i and $\varepsilon = 0.05$ lower-bounds the denominator when b_τ approaches 0 or 1.

Mathematical properties. The function $b^{\alpha}(1-b)^{1-\alpha}$ on $b \in [0,1]$ is an unnormalized Beta$(\alpha{+}1, 2{-}\alpha)$ density, maximized uniquely at $b = \alpha$. When $b_\tau = \alpha$, $c(b_\tau)$ is largest and $|g_{\text{raw}}|$ is smallest; as b_τ moves toward easy ($b_\tau \to 1$) or hard ($b_\tau \to 0$) extremes, $c(b_\tau)$ shrinks and $|g_{\text{raw}}|$ grows, up-weighting deviations on extreme-difficulty tasks.

Role of asymmetric α. With $\alpha^+ = 0.7 > \alpha^- = 0.3$, EIR applies two difficulty calibrations:

- **Positive ($\delta_i \geq 0$, $\alpha = \alpha^+$):** $c(b_\tau)$ peaks at $b_\tau = 0.7$, giving mid-to-high difficulty tasks a natural baseline while amplifying gains on very hard tasks ($b_\tau \ll 0.7$).
- **Negative ($\delta_i < 0$, $\alpha = \alpha^-$):** $c(b_\tau)$ peaks at $b_\tau = 0.3$; regressions on high-baseline tasks ($b_\tau \gg 0.3$) receive stronger penalties, whereas low-baseline regressions retain larger $c(b_\tau)$ and preserve early exploration.

Ablation. Setting $\alpha^+ = \alpha^- = 0.5$ causes near-complete training collapse (BBH 4.26, OOD 3.67; Table 11), confirming that asymmetric EIR weighting is necessary for stable cross-task GRPO.

B Full Strategy Schema and Format Comparison

This section compares three strategy formats, specifies the CoDI semi-structured template, and gives two cross-task examples. Ablation numbers are in Table 3.

B.1 Format comparison, specification, and examples

We use BIG-Bench Navigate (grid tracking, decide whether the agent returns to the origin) to compare Free-form, CoT-style, and Semi-structured (CoDI) formats, matching Table 3. The table also ablates the mandatory sections `[INDUCTIVE_RULE]`, `[TASK_PROCEDURE]`, and `[TRANSFER_SCOPE]`.

Free-form (unstructured). The inducer is prompted as follows.

```
Based on the following examples,
generate a strategy to solve this type
of problem.
[few-shot examples]
Strategy:
```

A typical SG output is as follows.

```
Track coordinates. North +y, South -y, East +x, West -x. Check if (x,y) = (0,0) at end.
```

Limitation: Output is unconstrained in length and structure, so token-level log-probability baselines are incomparable across rollouts and ICR is ill-defined. SG often paraphrases specific examples rather than stating verifiable task-family rules, yielding the weakest OOD generalization.

CoT-style (implicit structure). The inducer is prompted as follows.

```
Based on the following examples, analyze the pattern and generate a step-by-step strategy for solving this type of problem. Think carefully and reason step by step.

[few-shot examples]

Strategy:
```

A typical SG output is as follows.

```
Let me analyze the examples: - Example 1: coordinates update with each direction... - Example 2: same pattern confirmed...

General approach: 1. Initialize position (0,0) 2. Parse each instruction... 3. Check if final position is origin

This should work for similar grid navigation problems, though diagonal movements might need special handling.
```

Limitation: Reasoning steps appear, but rules and procedures are interleaved and boundary conditions are ad hoc. Without fixed section headers, AS sees variable prompts, ICR is only moderately reliable, and OOD transfer is weaker than under semi-structured control.

Semi-structured prompt of CoDI:

```
Based on the following examples, generate a strategy using EXACTLY the following sections: [INDUCTIVE_RULE] / [TASK_PROCEDURE] / [TRANSFER_SCOPE]

[few-shot examples]

Strategy:
```

Typical SG output:

```
[INDUCTIVE_RULE] Update coordinates: North (+y), South (-y), East (+x), West (-x). Decide whether the endpoint equals the origin (0,0).

[TASK_PROCEDURE] 1. Initialize (x,y)=(0,0) 2. Parse each instruction for direction and steps 3. Update coordinates 4. Output Yes if (x,y)==(0,0), else No

[TRANSFER_SCOPE] Applies when directional moves are given and the query asks about returning to the start. Not applicable when diagonal moves or conditional branches appear; extend Step 2 accordingly.
```

Advantage: Fixed section headers yield comparable token-level baselines and measurable ICR; `INDUCTIVE_RULE` enforces explicit induction and `TRANSFER_SCOPE` manages boundaries, giving the strongest OOD performance.

Semi-structured specification. CoDI requires SG to emit three named sections; we define their semantics and format violations below.

Mandatory sections. Each strategy must contain three sections delimited by header lines:

- **`[INDUCTIVE_RULE]`**: Task-family regularities inferred from K examples, verifiable against supports but not tied to any single instance. Boundary cases belong in `[TRANSFER_SCOPE]`.
- **`[TASK_PROCEDURE]`**: Three to six numbered, executable steps covering parsing, reasoning, and output formatting.
- **`[TRANSFER_SCOPE]`**: Scope specification with trigger conditions, at least one non-applicable case, and any exceptions to the inductive rule.

Format violations (missing sections, misspelled headers, or empty section bodies) assign $R_i = 0$ and exclude the strategy from AS execution (format gating above).

Worked example (Navigate). The following Navigate strategy illustrates the three-section template.

```
[INDUCTIVE_RULE]

From the examples: each instruction has the form “walk {steps} steps {direction}.” Update North (+y), South (-y), East (+x), West (-x); test whether the final point is (0,0).

[TASK_PROCEDURE]

1. Initialize (x, y) = (0, 0). 2. Parse each instruction for direction and step count. 3. Update coordinates accordingly. 4. After all instructions, check whether (x, y) = (0, 0). 5. Output “Yes” if true, otherwise “No.”

[TRANSFER_SCOPE]

Trigger: a sequence of directional moves with a return-to-start
```

```
query. Application: apply the coordinate-tracking steps directly. Non-applicable: (1) diagonal moves; (2) conditional branches (e.g., "if ..., then go north"), which require extra parsing in Step 2.
```

Worked example (Word sorting). The following Word Sorting example shows the same schema on a different task family.

```
[INDUCTIVE_RULE]

Sort the given word list alphabetically (case-insensitive). Compare characters left to right; digits precede letters; special characters follow ASCII order.

[TASK_PROCEDURE]

1. Lowercase all words for comparison. 2. Sort by alphabetical order with character-wise comparison. 3. If prefixes match, place the shorter word first. 4. Output the sorted list as comma-separated tokens.

[TRANSFER_SCOPE]

Trigger: unordered word list requiring alphabetical output. Application: follow Steps 1–4. Non-applicable: (1) frequency- or semantics-based sorting requires a new criterion in Step 2; (2) multi-word phrases must be treated as single tokens in Step 1.
```

C Annotation and Evaluation Prompts

Figures 9–12 collect the full templates used for mechanistic scoring. They appear at the end of the appendix in single-column layout.

Six-mode scoring. Figure 9 is the Gemini-3-flash rubric for six induction modes (integer scores 0–5). The rubric is an operationalization inspired by classic accounts of induction and analogy (Holland et al., 1986; Gentner, 1983). Compositional and Procedural scores target `[TASK_PROCEDURE]`; Analogical and Pattern Extrapolation target `[INDUCTIVE_RULE]`; Constraint-based and Schema target `[TRANSFER_SCOPE]`.

Think-unit tagging. Figure 10 tags sentence-level <think> units as E/D/P/R.

Strategy quality. Figure 11 is the GPT-5 pairwise judge for Correctness, Executability, Transferability, and Conciseness, with models presented as A and B.

Inductive-rule judging. Figure 12 scores each `[INDUCTIVE_RULE]` on abstraction, coverage, copying, and transferability. Aggregate judge means are not reported; the template is included so the LOO qualitative protocol can be reproduced.

D Thinking-Unit Taxonomy and Trace Metrics

This section supports §5.2 and Figure 5. Unlike Table 1, which uses strategy-then-answer with thinking disabled, this analysis uses thinking-enabled direct ICL on a trace-paired matched OOD subset (episodes where both Pre and Post checkpoints emit a valid <think> span). We split each <think> span into sentence-level units and assign E/D/P/R labels with Gemini-3-flash (Figure 10).

D.1 Unit taxonomy (E/D/P/R)

Table 7 defines the four sentence-level tags used in this section. E and D are non-inductive. P marks a newly stated rule. R marks application of a rule already stated in the same trace.

Tag	Name	Ind.?	Description
E	Example-recitation	No	The unit recalls, quotes, or closely paraphrases a *specific* support example without abstracting beyond it.
D	Direct-application	No	The unit applies facts or procedures to the current instance without first stating a cross-example pattern.
P	Pattern-abstraction	Yes	The unit states a general rule or invariant intended to hold across instances (inductive content).
R	Rule-application	Yes	The unit applies a rule or pattern that was already stated earlier in the trace.

Table 7: Cognitive tags for sentence-level units inside <think>. "Ind.?" indicates inductive (rule-level) content; descriptions define the annotation target.

Phase-1 prefix. Let the labeled unit sequence be $(y_1, \ldots, y_n)$. **Phase-1** is the maximal prefix of consecutive E labels at the start of the trace: indices $\{1, \ldots, t\}$ where $y_1 = \cdots = y_t = E$ and either $t = n$ or $y_{t+1} \neq E$. We report $n_E^{(\mathrm{ph1})} = \sum_{i \leq t} \mathbf{1}[y_i{=}E]$, which equals t on this definition. Figure 5 plots this prefix length, not the total E count n_E used in $q_{E/R}$ and ρ_E.

D.2 Trace metrics and statistical tests

Let $(y_1, \ldots, y_n)$ be labels after dropping empty segments; n is the number of units. Denote counts $n_c = \sum_i \mathbf{1}[y_i{=}c]$ for $c \in \{E, D, P, R\}$, and $n_{P \cup R} = n_P + n_R$. **Example-to-rule ratio.** Whenever $n_{P \cup R} > 0$, define $q_{E/R} = n_E / n_{P \cup R}$. **Example share.** $\rho_E = n_E / n$, reported as a percentage. **Think length.** L_{think} is the character length of the full <think> span (including whitespace). **D→P rate.** $r_{D \to P} = \frac{1}{n-1} \sum_{i=1}^{n-1} \mathbf{1}[y_i{=}D,\, y_{i+1}{=}P]$ (percentage of adjacent pairs). **Mean P length.**

$\bar{\ell}_P$ = mean{len(unit$_i$) | y_i=P} (characters); undefined if n_P=0 (those samples are excluded from $\bar{\ell}_P$ averages only). **Accuracy.** $\text{Acc}_{\text{OOD}}^{\text{matched}}$ is exact-match accuracy on the same thinking-enabled matched subset as Fig. 5, not the full OOD split in Table 1.

Inference and testing. All sentence labels are produced by Gemini-3-flash with temperature 0. GPT-5 re-annotation on a random subsample of 40 matched OOD tasks (180 sentence-level units) yields Cohen's κ=0.79 on E/D/P/R labels and Pearson r=0.91 on six-mode scores (Appendix E; Tables 9 and 10). Unless otherwise noted, p-values are two-sided paired tests at the trajectory level over matched tasks (Pre vs. Post). Table 8 reports the primary process metrics. Matched accuracy rises from 45.2 to 92.4 while think length and the phase-1 recitation prefix fall, and D→P transitions and mean P length rise.

Metric	Pre	Post	Δ	p
L_{think}	5276	4171	−1105	$< 10^{-3}$
$q_{E/R}$	1.079	0.674	−0.405	3.0×10^{-4}
$r_{D\to P}$ (%)	4.36	6.21	+1.85	4.0×10^{-5}
$\bar{\ell}_P$	89.4	108.6	+19.2	8×10^{-6}
Acc_{OOD} (matched, %)	45.2	92.4	+47.2	$< 10^{-3}$
$n_E^{(\text{ph1})}$	18.54	9.06	−9.48	2.1×10^{-4}
ρ_E (%)	16.50	12.14	−4.36	1.3×10^{-3}

Table 8: Pre/Post-CoDI trace metrics on Qwen3-4B matched OOD tasks with valid thinking traces (K=3; subset of Table 1 OOD). Acc.: McNemar; $n_E^{(\text{ph1})}$, ρ_E: paired t; remaining metrics: paired trajectory tests.

E GPT-5 Cross-Annotation Check

To test whether mechanistic conclusions reflect Gemini-3-flash annotation preferences, GPT-5 re-scores strategies and re-annotates traces on a random subsample of 40 matched OOD tasks (180 sentence-level units), using the rubrics in Appendix C and Fig. 10. The six-mode rubric treats the six dimensions as distinct (Holland et al., 1986; Gentner, 1983), so a high score on one mode does not imply a high score on another.

Six-mode strategy scores. Table 9 reports Pre-/Post-CoDI scores on the six induction modes. Gemini-3-flash columns repeat the full-cohort scores of Figure 7. GPT-5 columns are the n=40 subsample. Rule-oriented modes (Procedural, Constraint, Schema) rise after training under both annotators, whereas Analogical and Pattern decline. Every Post−Pre Δ keeps the same sign, which is a sign check rather than a claim of equal sample size.

Mode	Gem. Pre	Gem. Post	G5 Pre	G5 Post
Compositional	1.23	1.51	1.18	1.46
Procedural	1.86	2.93	1.81	2.86
Constraint	1.96	3.46	1.88	3.38
Schema	2.47	3.19	2.41	3.12
Analogical	1.22	0.59	1.15	0.54
Pattern	1.09	0.47	1.14	0.53

Table 9: Pre-/Post-CoDI six-mode scores (Gem.: full cohort, Fig. 7; G5: n=40). All six Post−Pre Δ share the same sign across annotators.

Trace-level Post−Pre deltas. Table 10 reports Post−Pre changes for Gemini-3-flash on the full matched cohort (Table 8) and for GPT-5 on the n=40 subsample (180 units). All continuously scored metrics keep the same sign, as do the six-mode Post−Pre directions in Table 9. The comparison is a sign check rather than a matched-pair test on identical items.

Metric	Gem. Δ	G5 Δ	Sign
L_{think}	−1105	−968	✓
$q_{E/R}$	−0.405	−0.382	✓
$r_{D\to P}$	+1.85	+1.72	✓
$\bar{\ell}_P$	+19.2	+18.6	✓
$n_E^{(\text{ph1})}$ / ρ_E	−9.48 / −4.36	−9.12 / −4.08	✓
Six-mode signs	6/6	6/6	✓

Table 10: Post−Pre trace deltas: Gem. (full cohort) vs. G5 (n=40). Sign agreement only (Appendix D.2).

F Sensitivity Analysis

Table 11 reports hyperparameter sensitivity for K (shot count in the strategy-then-answer pipeline of Table 1), β, α, and k (log-smoothing constant in EIR) on Qwen3-4B. This K sweep is distinct from the evaluation-time shot curve in Figure 6. * marks the default in main experiments.

Shot count K. Performance peaks at K=3 (BBH 82.22, OOD 63.00) and drops slightly at K=5 (BBH 78.52, OOD 57.67). The largest single-step gain is K=1→2 (+24.32 BBH). At K=1, OOD is 34.71, below the untrained FS baseline of 53.33 in Table 1, so one-shot training is insufficient for the full strategy-then-answer pipeline.

ICR amplification β. β=0.5 is optimal across all metrics. Performance degrades at both β=0.3 (BBH 44.20, OOD 44.67) and β=0.7 (BBH 48.89,

Parameter	Value	ID	BBH	OOD
K (shots)	1	48.89	44.20	34.71
	2	55.67	68.52	48.00
	3*	64.67	82.22	63.00
	5	63.67	78.52	57.67
β (ICR weight)	0.3	50.00	44.20	44.67
	0.5*	64.67	82.22	63.00
	0.7	39.00	48.89	47.00
α^-/α^+	0.1/0.9	55.67	71.85	47.33
	0.3/0.7*	64.67	82.22	63.00
	0.5/0.5	4.00	4.26	3.67
k (log-smooth.)	1	30.00	35.00	34.00
	5	37.33	41.85	36.67
	10*	64.67	82.22	63.00
	20	41.67	49.81	36.00

Table 11: Hyperparameter sensitivity on Qwen3-4B for K, β, α, and k (*: default; exact-match accuracy, %).

OOD 47.00). Too little amplification under-weights the inductive signal, while too much distorts the reward landscape and destabilizes GRPO.

EIR asymmetry α. Asymmetric weighting $(\alpha^+, \alpha^-) = (0.7, 0.3)$ yields the best result (OOD 63.00). The symmetric setting $(0.5, 0.5)$ causes near-complete training collapse (BBH 4.26, OOD 3.67), so the asymmetric design is needed for GRPO stability across task difficulties. A more extreme pair $(\alpha^-, \alpha^+) = (0.1, 0.9)$ is worse (OOD 47.33), which suggests that over-penalizing negative outcomes suppresses exploration on hard tasks.

Log-smoothing constant k. k controls logarithmic compression of the difficulty-normalized gain g_{raw} in EIR and is distinct from shot count K. The default k=10 is best across splits (BBH 82.22, OOD 63.00). At k=1, insufficient smoothing yields noisy rewards and unstable training (BBH 35.00, OOD 34.00). Performance rises at k=5 (BBH 41.85, OOD 36.67) and peaks at k=10. Increasing k to 20 over-compresses difficulty gradients (BBH 49.81, OOD 36.00). Weak or strong compression both hurt; k=10 is the main default.

G Solver Transferability

Table 12 fixes CoDI-4B SG and swaps AS to test strategy transfer across executors.

The 4B row (61.93) follows the swap protocol, not the main-table OOD (63.00). 4B, 8B, and 32B AS remain close to that reference (61.93, 62.89, and 61.30), whereas 14B drops to 51.14. The 14B dip may reflect instruction-tuning mismatch. Solver compatibility is left to future work.

SG	AS	OOD↑	Δ vs. CoDI-4B AS
CoDI-4B	4B	61.93	−1.07
	8B	62.89	−0.11
	14B	51.14	−11.86
	32B	61.30	−1.70

Table 12: Solver transfer with a fixed CoDI-4B SG. Deltas are absolute OOD points relative to the main-table CoDI-4B AS score of 63.00. The 4B row is a rerun under the swap protocol and need not equal 63.00.

H LOO Rule Sensitivity and Inductive-Rule LLM Judging

Figure 4 reports embedding-based leave-one-out (LOO) sensitivity of the `[INDUCTIVE_RULE]` span on matched evaluation tasks with K=3. The four-dimension GPT-5 rubric is given in Fig. 12 for reproducibility; we do not report aggregate judge means. Table 13 maps heatmap rows T01–T10 to task types and splits. T01 and T05 are ID. T02 and T07 are OOD-subtask (Table 15). T10 (Date understanding) is a BBH diagnostic rather than an ID or OOD-subtask inventory item.

Row	Task type	Split
T01	Movie recommendation	ID
T02	Dyck languages	OOD-subtask
T03	Penguins in a table	ID
T04	Geometric shapes	ID
T05	Formal fallacies	ID
T06	Temporal sequences	ID
T07	Nonsense-word grammar	OOD-subtask
T08	Logical deduction	ID
T09	Salient translation error detection	ID
T10	Date understanding	BBH

Table 13: Task index for the rows in Fig. 4. Split labels follow Appendix J; T10 is a BBH diagnostic.

Stratified sampling and rule generation. We stratify by `problem_type`, draw evaluation shards, and use vLLM to sample M=5 independent `[INDUCTIVE_RULE]` texts per task under full few-shot conditioning.

LOO sensitivity and heatmap cells. Let $\mathbf{e}_i$ be the bge-m3 embedding of rule i, and $\mathbf{e}_i^{\{-j\}}$ the embedding after re-generating the rule with support example j removed. Per-rule sensitivity is $s_{ij} = 1 - \cos(\mathbf{e}_i, \mathbf{e}_i^{\{-j\}})$. Figure 4 plots task-level values $\bar{s}_{\tau j} = \frac{1}{M}\sum_{i=1}^{M} s_{ij}$ for ten representative tasks τ and slots $j \in \{1, 2, 3\}$. Cell colors are clipped at $v_{\max}$=0.5 for readability. The headline means $\bar{s}_{\text{LOO}}$ are the averages of these 30 cells (Pre 0.325, Post 0.087).

Optional softmax attribution. Softmax weights $a_{ij} = \mathrm{softmax}(s_{i1}, s_{i2}, s_{i3})_j$ are useful for per-rule attribution summaries; the heatmap cells in Fig. 4 use the raw cosine-based sensitivities averaged over rules as above.

LLM-as-judge protocol. GPT-5 assigns integer scores in $\{1, \ldots, 5\}$ with a one-sentence justification per dimension, using the template in Fig. 12. A higher `copying_score` is worse.

I StrategyBench Dataset

StrategyBench evaluates strategy induction and strategy application in in-context learning. It is built from BIG-Bench (Srivastava et al., 2023) by task filtering, format standardization, categorization, and split construction. The target setting is few-shot induction of a reusable solution method and application to unseen queries.

Data format. Each example uses a unified input–output structure. The **input** field holds the task instruction and query instance. The **output** field holds the target answer for automatic evaluation. Strategy-based episodes also include a small demonstration set used either to induce a strategy or as in-context examples for direct inference. Construction scripts and split files are in the public code release cited at the end of this section.

I.1 Five-Criterion Task Filtering

To keep training data suitable for strategy induction and reinforcement learning, BIG-Bench tasks are retained only if they meet five operational criteria.

1. **Objective evaluability.** Answers must be automatically checkable. Tasks with subjective judgments, open-ended opinions, ethical preferences, or unstable answer boundaries are excluded.
2. **Format stability.** After standardization, input–output formats must support a consistent episode layout for demonstrations, strategies, and predictions.
3. **Rule inductibility.** A small support set must suffice to infer a pattern, rule, mapping, or procedure. Tasks that mainly require memorized facts or arbitrary labels are excluded.
4. **Low external-knowledge dependency.** The query plus demonstrations should be sufficient. Tasks that require substantial world knowledge are excluded so that scores reflect induction rather than recall.
5. **Transfer potential.** The induced method must apply to new instances of the same family or under a documented shift.

Unsuitable tasks are removed. Retained tasks are standardized into training, validation, and test episodes. As BIG-Bench-derived synthetic reasoning tasks, our data do not contain personally identifying information; tasks with subjective or potentially sensitive content were excluded during filtering.

Teacher scoring. Retained tasks are processed for SFT warm-up. **Gemini-3-flash** proposes candidate reference strategies, and only strategies that pass execution-accuracy filtering on the frozen solver enter Phase 1 SFT. Task filtering, categorization, and ID/OOD split assignment use **Qwen3-32B** scores with author spot-checks, not Gemini.

Code and splits. The public code release at `https://anonymous.4open.science/r/StrategyBench-D53C/` contains the five-criterion filtering pipeline, the standardized episode format, and the ID/OOD split definitions used in §4.1.

J Task-Level ID/OOD Partition

J.1 In-Distribution (ID) Split

The ID split comprises tasks or subtasks present in the training data, with test *instances* held out from training instances. It evaluates whether a model can induce and apply strategies when task rules, input–output formats, and demonstration patterns resemble those seen during training.

J.2 Out-of-Distribution (OOD) Split

The OOD split is created by completely holding out certain tasks, subtasks, or task categories from training and using them only for evaluation. OOD examples differ more strongly from training data in task source, input–output format, reasoning structure, or required capabilities.

Granularity levels. StrategyBench also defines finer diagnostic splits. *OOD-subtask* holds out subtasks under related sources. *OOD-task* holds out tasks with related ability requirements. *OOD-category* holds out coarse categories with larger shifts. *BBH* denotes selected Big-Bench Hard tasks used as an external diagnostic, reported as a separate column in Table 1 rather than as the main OOD aggregate. Unless noted, this paper reports task-level **ID** and **OOD**.

Quantifying distributional shift. Following Ye et al. (2023), we train an MLP on per-task few-shot accuracies from ID tasks and evaluate how well it forecasts held-out OOD accuracy. Table 14 summarizes predictability (R^2), rule overlap, ability-profile Jaccard, and mean absolute error across split levels. Lower R^2 indicates a capability gap rather than a shared formatting artifact. Rule overlap is 0% on all OOD rows by construction. Ability-profile Jaccard is a separate similarity (0.74, 0.51, 0.29 at OOD-subtask, OOD-task, and OOD-category) and decreases as structural distance grows.

Split	Hold-out	Rule ovlp.	R^2	Jacc.	MAE (pp)
ID (ref.)	same subtasks	100%	0.91	–	3.1
OOD-subtask	28 subtasks	0%	0.68	0.74	6.2
OOD-task	19 tasks	0%	0.38	0.51	9.4
OOD-category	7 categories	0%	0.15	0.29	12.8

Table 14: OOD shift on StrategyBench split levels (Qwen3-4B FS baseline). Rule overlap is template overlap with ID training. Jaccard is ability-profile similarity, not rule overlap. ID keeps the same subtasks with disjoint instances; MAE is mean absolute error in percentage points.

ID and OOD-subtask inventories. Table 15 lists the ID and OOD-subtask inventories used for main-text BIG-Bench evaluation. The full split definitions and construction scripts are in the StrategyBench repository (Appendix I).

ID subtasks (training + eval)
Navigate, Word Sorting, Tracking Shuffled Objects, Boolean Expressions, Contextual Parametric Conflicts, Elementary Math, Formal Fallacies, Geometric Shapes, Hyperbaton, Implicatures, Logical Deduction, Movie Recommendation, Penguins in a Table, Reasoning about Colored Objects, Salient Translation Error Detection, Snarks, Sports Understanding, Temporal Sequences, Tracking Shuffled Objects (three objects), Web of Lies, Winowhy, and related BIG-Bench subtasks
OOD-subtask (held out, 28 subtasks)
Dyck Languages, Causal Judgment, Disambiguation QA, Epistemic Reasoning, Intent Recognition, Linguistic mappings, Nonsense-word grammar, Play Dialog Same or Different, Ruin Names, Strange Stories, Symbol Interpretation, Unit Conversion, and 16 additional held-out subtasks (full list in repository)

Table 15: Representative ID and OOD-subtask inventories for main-text BIG-Bench evaluation; remaining names in the 28-subtask hold-out are in the public release. Date understanding (Fig. 4, T10) is a BBH diagnostic and is not listed in either inventory.

Prompt Used for Six-Mode Induction Scoring

You are an expert in inductive reasoning. Score the following model-generated strategy on six theoretically grounded modes of induction. Each mode is rated on an integer scale 0–5; scores must be justified by direct reference to the strategy text.

Task description: {task}

Strategy to score: {strategy text}

Six induction modes:

1. **Compositional** — Does the strategy combine multiple atomic sub-rules into a composite procedure (e.g. chaining conditions, nesting operations)?
2. **Analogical** — Does the strategy draw structural correspondences across examples, mapping shared relational patterns rather than surface features?
3. **Pattern Extrapolation** — Does the strategy extract a pattern that visibly extends beyond the demonstrated instances (e.g. to longer sequences, new value ranges, unseen entities)?
4. **Procedural** — Does the strategy abstract a step-wise algorithm from input–output pairs (clear ordering, state updates, termination conditions)?
5. **Constraint-based** — Does the strategy identify and articulate implicit constraints, preconditions, or boundary conditions (e.g. when the rule does *not* apply)?
6. **Schema-induction** — Does the strategy explicitly classify the task type itself, recognising the meta-level structure (e.g. "this is a coordinate-tracking problem")?

0–5 scoring rubric (applied per dimension):

- **0** — Absent: no evidence of this mode in the strategy.
- **1** — Hinted: a fragment alludes to it but is not developed.
- **2** — Partial: present but incomplete or ambiguous.
- **3** — Adequate: clearly present, executable, but minimal.
- **4** — Strong: well-developed, multi-step or multi-condition.
- **5** — Exemplary: cleanly abstracted, generalisable, with explicit scope statements.

Calibration constraints: (i) Scores must reference specific text spans (quote the supporting sentence in `evidence`). (ii) If a mode is structurally inapplicable to the task type, mark as `N/A` rather than 0. (iii) Modes are scored independently; high score on one does not imply high score on others.

Output strictly in JSON:

```
{
  "compositional": {"score": 0-5 or "N/A", "evidence": "...", "reason": "..."},
  "analogical": {"score": 0-5 or "N/A", "evidence": "...", "reason": "..."},
  "pattern_extrapolation": {"score": 0-5 or "N/A", "evidence": "...", "reason":
"..."},
  "procedural": {"score": 0-5 or "N/A", "evidence": "...", "reason": "..."},
  "constraint_based": {"score": 0-5 or "N/A", "evidence": "...", "reason": "..."},
  "schema_induction": {"score": 0-5 or "N/A", "evidence": "...", "reason": "..."},
  "composite": {"mean_score": float, "applicable_modes": int}
}
```

Figure 9: Prompt used for six-mode induction scoring. Each strategy receives an independent 0–5 score per mode; the composite is the mean over applicable modes (N/A modes excluded from averaging).

Prompt Used for Thinking-Unit Classification (E/D/P/R)

System. You are an expert cognitive analyst studying how language models reason during few-shot problem solving. Classify a *single* reasoning sentence into exactly one of four categories by its cognitive function. Be strict and consistent; output valid JSON only, with no extra text.

User.

Task context. Problem: {problem_text}

Support examples (summary): {examples_summary}

Four categories.

1. **E (Example-recitation):** The sentence directly recalls, quotes, or closely paraphrases content from a specific support example. No abstraction beyond that example.
2. **D (Direct-application):** The sentence applies knowledge, a fact, or a procedure to the current problem instance without first deriving or stating a general pattern (calculations, lookups, direct inferences from given facts).
3. **P (Pattern-abstraction):** The sentence states a general principle, rule, or pattern intended to hold across multiple examples or instances.
4. **R (Rule-application):** The sentence explicitly applies a general rule or pattern that was stated or inferred earlier in the reasoning trace.

Sentence to classify: {unit_text}

Output strictly as JSON (one line, no extra keys): `{"label": "<E|D|P|R>", "reason": "<one sentence>"}`

Figure 10: Prompt for Gemini-3-flash sentence-level tagging of <think> units.

Prompt Used for Strategy Quality Analysis

Please act as a professional strategy evaluator and compare the following two model-generated strategies across four dimensions.

Task description: {task}

Model A Strategy: {Model-A-generated strategy}

Model B Strategy: {Model-B-generated strategy}

Compare the two strategies on the following four dimensions:

1. **Correctness** — Does the strategy accurately capture the task rules and produce correct solutions?
2. **Executability** — Is the strategy clear and actionable enough for a model to follow step by step?
3. **Transferability** — Is the strategy abstract and generalizable enough to apply to unseen tasks of the same type?
4. **Conciseness** — Is the strategy expressed concisely without redundant or repetitive content?

For each dimension, indicate which is better and provide a brief reason. Output in JSON format:

```
{
  "correctness_better": "Model A" / "Model B" / "tie",
  "correctness_reason": "reason",
  "executability_better": "Model A" / "Model B" / "tie",
  "executability_reason": "reason",
  "transferability_better": "Model A" / "Model B" / "tie",
  "transferability_reason": "reason",
  "conciseness_better": "Model A" / "Model B" / "tie",
  "conciseness_reason": "reason",
  "overall_better": "Model A" / "Model B" / "tie",
  "overall_reason": "reason"
}
```

Figure 11: Prompt used for strategy quality analysis.

Prompt Used for Inductive-Rule Quality (JSON-only)

System. You are an expert evaluator for few-shot inductive-rule quality. You will receive (i) a task description, (ii) $K{=}3$ support input–output pairs, and (iii) a generated inductive rule `[INDUCTIVE_RULE]`. Score the rule objectively on four dimensions (1–5) with a one-sentence justification each. Output **only** valid JSON (no markdown, no extra text).

Calibration (do not score). *Example A* (high abstraction / low copying): sentiment task; rule "Identify whether the review expresses a net positive or negative evaluation of a product attribute; classify accordingly." Expected: abstraction 5, coverage 5, copying 1, transferability 5. *Example B* (low abstraction / high copying): same demos; rule "If the review mentions battery life dying quickly, classify it as Negative." Expected: abstraction 1, coverage 1, copying 5, transferability 1.

Task. {task_description}

Support. Ex1: in {ex1_input}, out {ex1_output}; Ex2: in {ex2_input}, out {ex2_output}; Ex3: in {ex3_input}, out {ex3_output}.

Rule. {inductive_rule}

Rubric (1–5 each). (1) `abstraction_score`: task-level principle vs. paraphrasing demos (5 fully abstract; 1 restates one demo). (2) `coverage_score`: supported by multiple demos (5 all three; 1 only one). (3) `copying_score`: copied from a single demo (5 strong copying; 1 none)—**higher is worse.** (4) `transferability_score`: useful on unseen instances of the same task type.

JSON schema.

```
{"abstraction_score": int (1-5), "abstraction_reason": "one sentence",
 "coverage_score": int (1-5), "coverage_reason": "one sentence",
 "copying_score": int (1-5), "copying_reason": "one sentence",
 "transferability_score": int (1-5), "transferability_reason": "one sentence"}
```

Figure 12: GPT-5 prompt for scoring each `[INDUCTIVE_RULE]` on abstraction, coverage, copying (reverse-scored), and transferability.