arXiv:2610.26604v1 [cs.LG] 19 Oct 2026

ITU-DESC: a missingness-aware machine-learning dementia risk score in the UK Biobank cohort

Abdullah Kavaklı[1*] and Faik Boray Tek[1]

[1*]Department of Artificial Intelligence and Data Engineering, Istanbul Technical University, Istanbul, Turkey.

*Corresponding author(s). E-mail(s): kavakli23@itu.edu.tr;

Abstract

Background Existing dementia risk scores rely on fixed, pre-selected features and linear models, or ignore bias from informative missingness in cohort data. Here we introduce the Istanbul Technical University Dementia Risk Score (ITU-DESC), an explainable machine-learning score for incident all-cause dementia that combines broad feature selection across roughly 8,000 candidate features with explicit control of missingness bias.

Results We studied approximately 500,000 UK Biobank participants with a median follow-up of 16.6 years. We withheld missingness patterns from the model using unconditional random hot-deck imputation. We stabilized feature selection by repeating BorutaSHAP-based feature selection with ten random-imputation seeds and selecting predictors chosen in at least eight of the ten seeds, which yielded 38 predictors. We trained a CatBoost gradient-boosted decision-tree model and evaluated it on a holdout test set of 50,041 participants. For all-cause incident dementia, ITU-DESC achieved an average precision of 0.138 (95% confidence interval 0.123–0.156) and an area under the receiver operating characteristic curve (ROC AUC) of 0.853 (0.844–0.863), significantly outperforming other risk scores, including the strongest comparator, the UK Biobank Dementia Risk Prediction model (0.118 and 0.841). We identified dementia risk predictors absent from established scores, including leukocyte telomere length and liver enzymes.

Conclusions ITU-DESC is the first dementia risk score to combine broad feature selection with explicit control of missingness bias.

Keywords: Dementia risk score, Alzheimer's disease, UK Biobank, feature selection, missing data, SHAP

1 Background

Dementia affects approximately 57 million people worldwide, with nearly 10 million new cases each year [1]. Survival after diagnosis is short: average life expectancy ranged from 8.0 years for women diagnosed at age 65 to 2.2 years for men diagnosed at age 85 [2]. The 2024 Lancet Commission estimated that almost half of dementia cases could theoretically be prevented or delayed by addressing 14 modifiable risk factors across the life course. These factors span education in early life; sensory, vascular, metabolic, mental-health and behavioral risks in midlife; and social, environmental and sensory risks in later life, including social isolation, air pollution and vision loss [3]. Because prevention works best when it starts early and reaches those who need it most [3], it depends on identifying individuals at high risk before dementia develops.

In recent years, several statistical risk scores have been developed to estimate future dementia risk. Most are based on linear modeling frameworks, including Cox proportional hazards, logistic regression or competing-risk regression, and use predefined sets of clinical and demographic predictors. Established examples include the Australian National University Alzheimer's Disease Risk Index (ANU-ADRI) [4]; the Cardiovascular Risk Factors, Aging, and Incidence of Dementia (CAIDE) Risk Score [5]; the Dementia Risk Score (DRS) [6]; and the UK Biobank Dementia Risk Score (UKBDRS) [7]. These scores share predictors such as age, hypertension, stroke, parental history of dementia and high cholesterol, but differ in their target populations, input variables and model assumptions. Although these scores are explainable and practical, their fixed feature sets and linear assumptions limit their ability to capture non-linear associations and interactions between predictors. In the Norwegian HUNT cohort, several such scores did not outperform age and education alone in predicting dementia or cognitive decline [8, 9].

Machine-learning approaches relax these constraints and have also been applied to dementia risk. You et al. [10] developed the UK Biobank Dementia Risk Prediction model (UKB-DRP) using LightGBM, a gradient-boosting framework based on decision trees [11]. They were among the first to derive a dementia risk score from UK Biobank data with machine learning. The model identified several important predictors, including age, *APOE* $\varepsilon 4$ status, cognitive-test performance, number of medications, reaction time, peak expiratory flow, mother's age at death, long-standing illness and mean corpuscular volume. In large cohort datasets, however, machine learning introduces a problem: prediction may depend not only on recorded values but also on whether a measurement is available.

Missing data in UK Biobank are more than a preprocessing problem because the probability that a value is missing depends on participant characteristics, assessment procedures or health status. In the UKBDRS development sample, for example, participants excluded under complete-case analysis differed in age, sex and education from those retained [7, 12]. Because these characteristics are themselves dementia risk factors, missingness is related to the outcome and therefore informative for prediction [13], whether or not it satisfies the assumption of missing at random [14, 15]. Machine-learning models can therefore discriminate partly from the presence or absence of measurements rather than from the measured biological or clinical information. For example, if a blood test is unavailable more often among participants

who later develop dementia, the model can learn this recording pattern as a proxy for future dementia risk. Such patterns may be specific to the cohort and its assessment procedures. In the present cohort, missingness in the predictors of the final score was itself associated with incident dementia (Section 3.7).

Relying on such patterns introduces leakage-like bias, whereby predictive performance partly reflects how data were collected rather than associations expected to persist in the target population [16]. No approach to handling missing data is neutral in this setting. Algorithms with native missing-value handling can exploit informative missingness for prediction. Constant imputation using the mean, median or mode attains prediction-consistency under a missing-at-random mechanism. It does so because a sufficiently flexible learner detects the imputed value and can therefore treat missingness as a predictor [13]. Under informative missingness, the predictor it recovers relies on the recording pattern, the behavior we set out to avoid.

No existing dementia risk score combines broad machine-learning-based feature selection across UK Biobank variables with explicit control of this bias. We therefore developed the Istanbul Technical University Dementia Risk Score (ITU-DESC), an explainable machine-learning score for incident dementia. We treated missing-value handling as a design choice rather than a preprocessing step (Section 2.6). Starting from a large set of candidate baseline variables, we let the algorithm identify predictors under a missingness-handling procedure designed to limit the use of measurement availability as predictive information. We repeated feature selection across random-imputation seeds and kept only consistently selected predictors, making the final set robust to imputation variability.

In approximately 500,000 UK Biobank participants with variable follow-up, we developed and evaluated ITU-DESC for all-cause dementia and Alzheimer's disease. We assessed prediction across 5-year, 10-year and all-incident horizons, and compared ITU-DESC with ANU-ADRI [4], CAIDE [5], DRS [6], UKBDRS [7] and UKB-DRP [10]. Because these scores come from different cohorts or data partitions, with different outcome definitions and missing-data procedures, their published performance estimates are not directly comparable. We therefore applied to every comparator the same missing-data policy, outcome definitions and isotonic calibration procedure used for ITU-DESC. We evaluated all models on the same holdout test set, which no stage of feature selection or model fitting had seen.

2 Methods

2.1 Dataset

UK Biobank is a population-based longitudinal cohort study comprising more than 500,000 participants aged 40–69 years at recruitment [17]. Baseline assessments were conducted between 2006 and 2010 at 22 assessment centers across the United Kingdom, with ongoing follow-up and repeat assessments in a subset of participants. Figure 1 shows the participant selection process for the present study. Of 501,936 participants, we excluded 230 with prevalent dementia at baseline and 1,296 with no follow-up data, leaving 500,410 participants eligible for analysis. The median follow-up was 16.6 years (interquartile range, 15.8–17.4 years). We then divided the eligible

cohort into a 90% development set of 450,369 participants and a 10% holdout test set of 50,041 participants (Section 2.3).

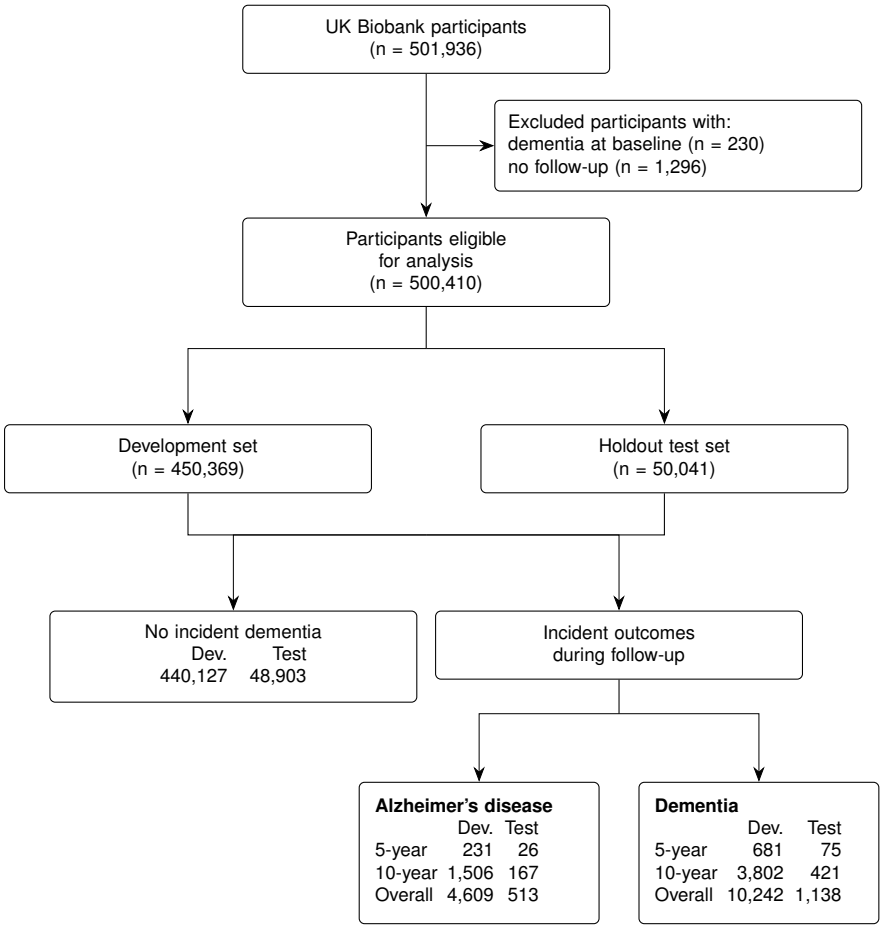

Fig. 1 Participant selection for the present study from the UK Biobank cohort. The upper boxes show the number of participants remaining after exclusions and the split of the eligible cohort into a 90% development set and a 10% holdout test set. The lower boxes report counts separately for the development (Dev.) and holdout test sides of the split. For each side, they give the number of participants without incident dementia during follow-up. They also report the numbers of incident dementia and Alzheimer's disease cases within the 5-year, 10-year and full follow-up periods.

2.2 Outcome definition

Dementia includes several clinical subtypes, of which Alzheimer's disease (AD) is the most common. Other subtypes include vascular dementia, frontotemporal dementia and dementia with Lewy bodies. To assess prediction across time horizons, we evaluated 5-year, 10-year and all-incident outcomes for both all-cause dementia and AD.

Outcomes were defined using algorithmically derived UK Biobank variables based on self-reported diagnoses and ICD-10 codes [18]: Field 42018 for all-cause dementia and Field 42020 for AD. Incident outcomes were defined relative to each participant's baseline assessment date, corresponding to the first assessment visit (Field 53). Participants with dementia onset before baseline were classified as having prevalent dementia and excluded from the analysis.

2.3 Development and holdout test split

We reserved a holdout test set before feature selection, model development or any other modeling decisions, ensuring that test-set information did not influence model construction or performance estimation. Using a single stratified split, we assigned 10% of the eligible cohort (50,041 participants) to the holdout test set and the remaining 90% (450,369 participants) to the development set. The split was stratified on the six binary outcome indicators and sex. Given the large cohort size, we did not search over many candidate splits to optimize covariate balance; instead, we accepted the first candidate split that met prespecified acceptance criteria on outcome balance and on age, assessment center and ethnicity, covariates not included in the stratification scheme. The first candidate seed met all criteria. The stratification scheme, acceptance criteria and full split diagnostics are reported in Additional file 1: Supplementary Methods (Tables S1 and S2).

All subsequent preprocessing decisions, imputation donor pools, feature selection, model fitting and calibration were based exclusively on the development set. The holdout test set was used only for final model evaluation. Within the development set, we also used five-fold cross-validation for two purposes: screening candidate features (Section 2.7) and generating out-of-fold predictions for fitting the calibration mappings (Section 2.8). For both applications, imputation donor pools were constructed separately within each training fold (Section 2.6).

2.4 Methodological challenges in dementia risk prediction

Dementia risk prediction presents several methodological challenges arising from the disease course, cohort structure and characteristics of the data.

First, dementia risk increases markedly with age, and some participants classified as non-cases during the observed follow-up may develop dementia later. This introduces label uncertainty, particularly for outcomes with long latency, because participants considered dementia-free at the censoring date may subsequently receive a diagnosis.

Second, participants with established dementia are difficult to characterize reliably at baseline. Cognitive impairment may reduce the accuracy or completeness of self-reported information and may also affect participation in detailed assessments. For risk prediction, data collected before clinical onset and linked to subsequent diagnosis are therefore more informative than data collected after disease onset. Accordingly, participants with prevalent dementia at baseline were excluded from the analysis.

Third, incident dementia was rare in the analyzed UK Biobank cohort. The positive class represented 2.27% of participants, resulting in substantial class imbalance. This

imbalance can reduce the stability of model estimation and make evaluation metrics sensitive to the small number of outcome-positive cases. The observed incidence of 2.27% also represents a lower bound on eventual lifetime dementia risk because some participants classified as non-cases during follow-up may be diagnosed after the current censoring date.

Fourth, missing data are common in UK Biobank. Missing values may be imputed, handled natively by algorithms that support missing inputs or addressed through complete-case analysis [19]. Each approach has limitations, particularly when missingness is informative or associated with the outcome. Our approach to missing data and the associated risk of leakage-like bias are described in Section 2.6.

Fifth, UK Biobank includes repeated measurements for some variables across multiple assessment visits. All participants attended the baseline assessment, but only a subset returned for repeat visits, and not all variables were measured at each visit [20]. Although repeated assessments can provide longitudinal information, their use would introduce inconsistencies in both follow-up time and feature availability across participants. Later visits were available for only a subset of participants, contained fewer variables, and could reflect changes in participant characteristics with age. To define a common prediction time point and feature set, we therefore used only data from the baseline assessment to predict incident dementia.

Finally, the dataset contained a large number of variables that required preprocessing and computationally demanding feature selection. Many variables had to be excluded to prevent data leakage, and subsequent feature selection was computationally demanding and required careful validation. After preprocessing, 7,977 candidate features remained, creating a prediction problem subject to the curse of dimensionality [21]. We therefore applied a two-stage feature-selection pipeline, as described in Section 2.7.

2.5 Feature preprocessing and leakage prevention

The extract contained 32,612 columns spanning healthcare records; questionnaire responses; physical measurements; and environmental, demographic, lifestyle and self-reported health data. Restricting the data to the baseline assessment and collapsing each field to a single typed column reduced this number to 11,926. Removing dementia- and outcome-linked fields (explained below) then left 11,340 columns. Multiple-choice fields, for which all selected responses were stored in a single delimited cell, were expanded into separate binary indicators for each observed response code, increasing the dataset to 37,693 columns. For participants with no recorded response to the original field, all corresponding indicators were set to missing rather than zero. Thus, zero indicates that an option was not selected, whereas missing indicates that the question was not answered. Applying the a priori exclusions described below reduced this set to 7,977 candidate features, which formed the input to all subsequent feature-selection procedures. Column counts at each stage are reported in Additional file 1: Table S3, and the composition of the resulting candidate feature set in Additional file 1: Table S5.

We excluded variables for three main reasons. First, we removed variables that could introduce leakage, including dementia subtype fields; certain fields whose values

are recorded after the baseline assessment, namely Hospital Episode Statistics inpatient records (Category 2000), which may include dementia diagnoses, and COVID-19 fields; and fields recording the year of a diagnosis or measurement.

Second, we excluded variables that primarily reflected UK Biobank data-collection procedures rather than participant characteristics. These included assay and genotyping quality-control flags, bulk-data pointers and assessment-center timestamps, because such signals are unlikely to be present beyond the study setting.

Third, we excluded variables whose recorded values did not adequately represent the baseline characteristic of interest. Environmental exposure variables are assigned according to place of residence and may change during follow-up, making a single baseline measurement an unstable representation of longer-term exposure. Year-of-diagnosis fields for various health outcomes posed an additional problem beyond leakage: these values are missing for participants who never experienced the corresponding event. Because all retained variables were imputed (Section 2.6), retaining these fields would have assigned diagnosis years to participants who had never received those diagnoses.

All exclusions were specified at the field level before model fitting, and no exclusion decision was based on an observed association with the outcome or on model performance. In total, 3,452 UK Biobank entries were excluded. Most exclusions were due to outcome or temporal leakage (2,692 entries), followed by study-administration and timing (243) and COVID-19 (201). The remaining exclusions comprised employment-history variables, assay and genotyping quality-control fields, self-reported mental-health variables, environmental-exposure fields, bulk-data pointers, derived accelerometry measures, subjective health items and variables redundant with retained fields. Additional file 1: Table S4 reports the group-level breakdown, and the main groups are described in detail under "Column construction and a priori exclusions" in Supplementary Methods. The individual excluded entries and the columns of the candidate feature set are listed in Additional file 2.

Beyond these exclusions, we combined parental survival status and longevity into one derived variable for each parent. Participants whose parent had died were assigned that parent's age at death, whereas participants whose parent was still alive were assigned a sentinel value below the observed age range. Combining these variables reflects the structure of the UK Biobank data, in which a missing age at death generally indicates that the parent was still alive rather than that the value was unrecorded [22, 23].

2.6 Missing data and imputation

In UK Biobank, whether a measurement is available depends on participant characteristics, assessment procedures and health status. Such missingness is structured rather than incidental [15]. Mitra et al. describe a common clinical example: a test is ordered when disease is suspected, so whether a value is observed depends on the very status the test would reveal, and the resulting missingness pattern is itself unobservable [15]. Under Rubin's taxonomy, such data may be missing not at random (MNAR) [14, 24].

Even in the absence of a causal relationship, missingness can carry predictive information about the outcome, and a supervised learner will exploit that association as a shortcut.

Tree-based learners can take that shortcut directly. Rather than requiring complete inputs, they route missing values inside the splitting procedure itself. One general strategy, missing incorporated in attribute (MIA), treats missingness as a category; eXtreme Gradient Boosting (XGBoost) [25] and scikit-learn's histogram gradient boosting implement this approach [13], while LightGBM and CatBoost provide comparable native handling of missing values. Josse et al. recommend MIA because it can accommodate both informative and non-informative missingness [13]. We nevertheless did not use native missing-value handling in the primary analysis because our aim was to identify predictors from their recorded values rather than from cohort-specific patterns of measurement availability. Although native handling can increase observed discrimination, gains driven by such cohort-specific patterns may not generalize to clinical settings, where measurements are obtained under different procedures.

Deterministic single-value imputation raised a related concern. Mean or median imputation concentrates additional mass at a single value of a continuous variable, whereas mode imputation inflates one category; such deterministic approaches therefore alter the distribution of the imputed variable [26]. Constant imputation can nevertheless yield consistent predictions with a sufficiently flexible learner. This guarantee holds under a missing-at-random mechanism, and only when the imputed value is identifiable. Identifiability lets the learner detect that value and use missingness as a predictor. For the coded categorical variables in our candidate set, it requires encoding missingness as an explicit category [13]. That is the information we sought to withhold from the model. We rejected complete-case analysis for a different reason: excluding every participant with a missing value in any required predictor can substantially reduce the sample in a high-dimensional feature set and, depending on the missingness mechanism, produce a selected analysis population [13, 19]. Multiple imputation by chained equations (MICE) [27] and random-forest methods such as missForest [28] use information from the remaining variables and may therefore reconstruct missing values more accurately than draws from a variable's marginal distribution. We did not use these methods because they condition on multivariable patterns that may include the cohort-specific measurement structure we aimed to suppress, and because each imputation model would need to be refitted within every feature-selection fold.

We therefore used unconditional, or marginal, random hot-deck imputation [29]. For each variable, every missing entry was replaced by an independent draw, with replacement, from the values observed for that variable in the corresponding training partition. Because hot-deck imputation draws only from observed values, imputed entries remain within the observed support of the variable, while random donor selection preserves the marginal distribution of the donor pool in expectation [26, 30, 31]. Random donor sampling also meant that the imputed data varied across random seeds; Section 2.7 describes how this variability was handled. Donor pools consisted of all observed values for a variable in the relevant training partition and were not stratified by auxiliary variables. Donors were drawn exclusively from the development

data. During five-fold cross-validation within the development set (Section 2.3), a separate donor pool was constructed from each training fold and used to impute both that fold and its validation fold. For the final model, donor pools were constructed from the full development set. The holdout test set therefore never contributed to the imputation rule applied to it, consistent with the principle that preprocessing fitted during model development should be applied unchanged to unseen data [13]. UK Biobank also encodes responses such as "Do not know" and "Prefer not to answer" as ordinary field values, which a gradient-boosting model could otherwise treat as substantive responses. We therefore recoded these responses as missing before imputation (Additional file 1: Supplementary Methods, "Missing-value handling").

For benchmarking, we applied the same missing-data policy to all models. UKB-DRP was reimplemented using random imputation instead of LightGBM's native missing-value handling, and missing components of UKBDRS, CAIDE, DRS and ANU-ADRI were likewise randomly imputed rather than treated according to their original specifications. For UKBDRS, which was originally evaluated using complete-case analysis, this represented a departure from its published implementation. Draws were keyed by participant identifier rather than row position, so that each score received the same imputed value for a given participant and variable. The risk-score comparison therefore used identical imputed data rather than only a common imputation policy. The harmonized policy allowed all models to be evaluated in the same participants and prevented differences in missing-data handling from confounding the comparison. To quantify the effect of withholding native missing-value handling, we additionally evaluated UKB-DRP with LightGBM's native handling on the same holdout test set (Section 3.2).

2.7 Feature filtering and selection

Feature selection started from the 7,977 candidate features (Section 2.5) and ran entirely within the development set (Fig. 2). To allow broad feature discovery, we considered all feasible baseline variables, including those without established associations with dementia. Selection proceeded in two stages: screening based on SHapley Additive exPlanations (SHAP) [32], followed by wrapper selection. SHAP values decompose an individual prediction into additive contributions from each feature, providing a signed attribution for every feature on the model output scale. Global feature importance was summarized as the mean absolute SHAP value across participants. Both stages used XGBoost as the underlying learner. Because random imputation made both stages stochastic, the complete pipeline was repeated for each of ten random-imputation seeds, as described below.

For initial screening, we computed SHAP values on each validation fold of the five-fold cross-validation within the development set (Section 2.3), imputing each fold's training and validation partitions under the corresponding seed. SHAP values were then aggregated across the folds, and features with zero SHAP importance in every fold were removed. This step eliminated features with no observed contribution before the more computationally intensive wrapper-based selection stage. The ten seed-specific screened sets were similar in size, containing between 4,554 and 4,635 variables.

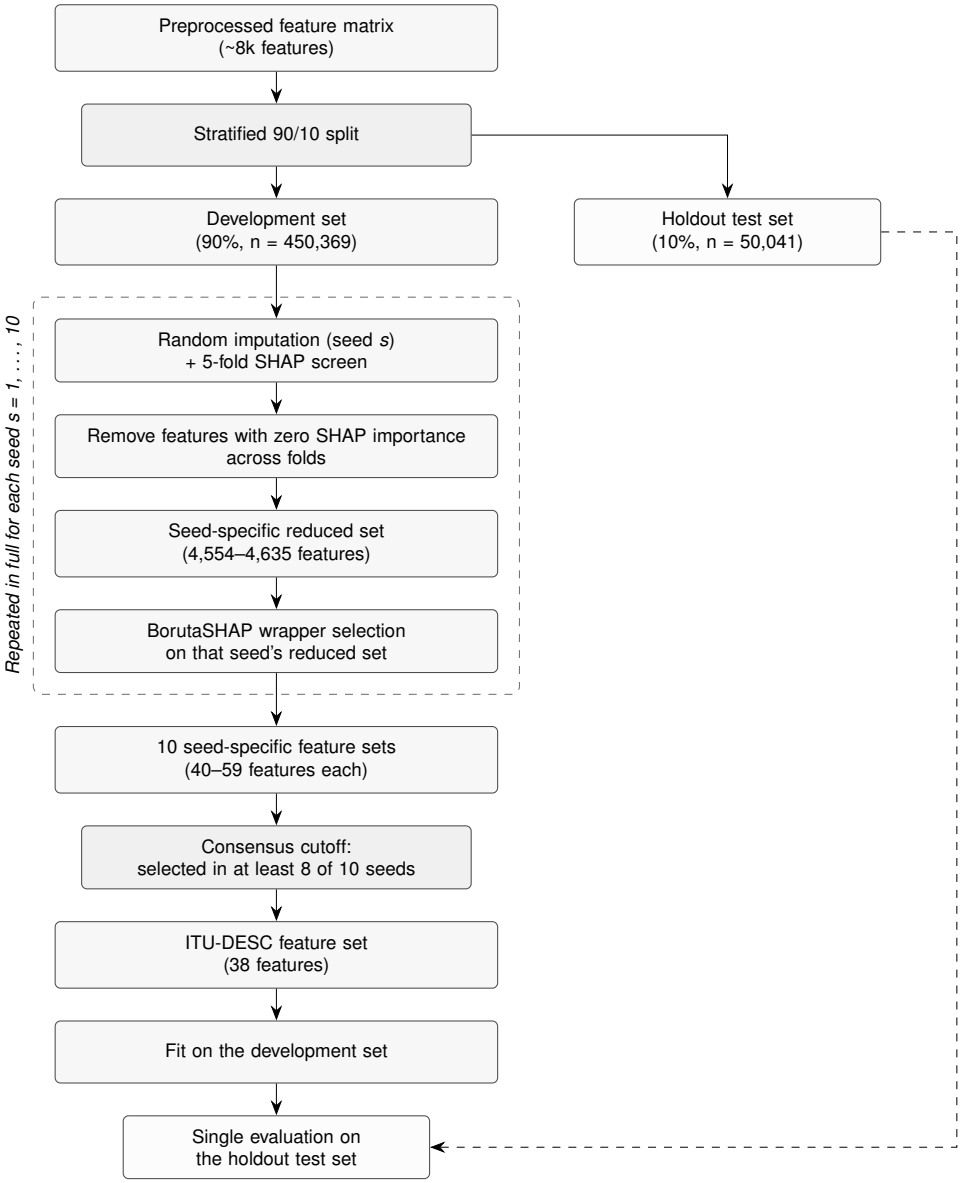

Fig. 2 Feature-selection pipeline. A stratified 90/10 split holds 10% of participants out as the test set. Within the 90% development set, the boxed two-stage pipeline consists of a five-fold SHAP screen that removes features with zero importance, followed by BorutaSHAP wrapper selection. The pipeline is repeated in full for each of the ten random-imputation seeds. Both stages therefore run ten times, and each seed carries its own screened feature set into its own wrapper run. Features selected in at least 8 of the 10 seeds form the ITU-DESC feature set. The model is then fitted on the development set and evaluated once on the holdout test set, which the selection stages never see.

After reducing the feature set to a tractable size, we applied BorutaSHAP [33], a wrapper-based feature-selection algorithm extending the Boruta approach [34]. BorutaSHAP combines a tree-based learning model, here the same XGBoost learner, with SHAP-based feature importance to identify relevant predictors. Briefly, the algorithm augments the dataset with shuffled copies of the original variables, termed shadow features, and compares the importance of each real feature with a shadow-feature threshold. Features were retained when they exceeded this threshold significantly more often than expected by chance across BorutaSHAP iterations; features that failed this test were rejected.

Because random imputation makes both selection stages stochastic, a single pass reports the features favored by one imputation draw rather than the features the data support. Drawing values at random adds imputation variance to every downstream estimate [26], and feature selection inherits it. We therefore repeated the pipeline independently for ten random-imputation seeds, screen and wrapper alike, so that each wrapper run saw both a different imputation draw and the screened feature set produced under that draw. The ten runs selected between 40 and 59 features each, and their selections did not coincide. We then ranked features by how many seeds selected them and defined nested consensus sets: Set k contained the features selected in at least k of the ten seeds, so Set 10 was the intersection of all ten lists and Set 1 their union. We adopted Set 8, the 38 features selected in at least eight of ten seeds, as the fixed ITU-DESC specification. This cutoff keeps only predictors that survive most imputation draws, while avoiding the stricter intersection, which would discard nine features that eight or nine seeds agreed on.

2.8 Model development and validation

We used gradient-boosted decision trees because they can capture non-linear associations and interactions and perform well in high-dimensional, imbalanced classification settings. XGBoost [25] was used throughout the feature-selection pipeline, whereas the final ITU-DESC model was trained with CatBoost [35]. We did not apply resampling, class weighting or other class-imbalance corrections. All stages, from SHAP screening and BorutaSHAP to the final CatBoost fit, used the same training label: the binary all-incident, all-cause dementia outcome defined in Section 2.2, assessed over each participant's full available follow-up rather than at a fixed time horizon. Software and package versions are reported under "Software and implementation" in Additional file 1: Supplementary Methods (Table S6).

Within the development set, we cross-fitted the model across five stratified folds (Section 2.3) so that each participant received an out-of-fold prediction from a model that had not been trained on that participant. The pooled out-of-fold predictions were used both to fit the isotonic calibration mappings subsequently applied to the holdout-test predictions and to generate internal development diagnostics. A separate calibration mapping was fitted for each outcome and prediction horizon, allowing the same base model to be recalibrated for all-cause dementia and Alzheimer's disease at the 5-year, 10-year and all-incident horizons. The base model itself was trained only once, using the all-cause dementia label. We then refitted the model on the full development set and evaluated it once on the holdout test set. Because the holdout

test participants were not used in the split procedure, imputation donor pools, SHAP screening, BorutaSHAP selection or calibration fitting, the holdout-test results provide the primary performance estimates, whereas the development cross-fitting results are reported as internal diagnostics.

2.9 Evaluation and statistical analysis

We formulated prediction as a binary classification task. The aim was to compare discrimination between ITU-DESC and established risk scores under a common outcome definition, data split and set of performance metrics, rather than to estimate absolute time-to-event risk.

Because incident dementia was rare, average precision (AP) was chosen a priori as the primary discrimination metric. In imbalanced settings, precision–recall summaries are more informative than ROC-based summaries because they focus on performance in the outcome-positive class [36]. We used AP rather than the trapezoidal area under the precision–recall curve because AP summarizes precision at observed changes in recall without linear interpolation, thereby avoiding optimistic interpolation between points on stepwise precision–recall curves [37]. We reported ROC AUC as a complementary metric. Both metrics were computed from isotonic-calibrated predicted risks using the calibration procedure described below, which was applied identically to all models. Because the isotonic mapping is monotonic, it preserves the ordering of predicted risks except where ties are introduced. Corresponding results based on uncalibrated predictions are reported in Additional file 1: Supplementary Note, "Uncalibrated discrimination".

Pairwise differences in AP were evaluated using two-sided paired permutation tests on holdout-test predictions (2,000 permutations). Outcome labels were held fixed, and model scores were randomly swapped within participants. Differences in ROC AUC were assessed using two-sided paired DeLong tests on matched holdout-test predictions. Uncertainty was quantified using two-sided 95% confidence intervals (CIs), reported for each model's AP and ROC AUC and for each difference from ITU-DESC. Because average precision has no analytic variance estimator, AP CIs were obtained by paired subject-level bootstrap resampling of the holdout test set (10,000 resamples; percentile method). The same resampled participants were used for all models within each bootstrap replicate, so differences remained paired. ROC AUC CIs, for both individual models and pairwise differences, were derived from the DeLong variance estimator, ensuring that each interval and its corresponding P value were based on the same variance estimate. These intervals quantify sampling variability within the holdout test set under the single development/test split, but do not capture variability arising from model training or from the split itself.

Calibration is particularly important when predictions are intended to support clinical decision-making, because systematically overestimated or underestimated risks may mislead clinical decisions [38]. We therefore recalibrated each model's predictions post hoc using isotonic regression. Isotonic regression fits a monotonically non-decreasing, piecewise-constant mapping from raw scores to calibrated probabilities and, unlike Platt scaling, assumes only monotonicity [39, 40]. This flexibility

increases the risk of overfitting when the calibration sample or the number of outcome events is limited [40], although the large numbers of participants and outcome events in the present study reduced this risk. Because all scores were recalibrated on UK Biobank data, the benchmark reflects UK Biobank-specific recalibrations rather than performance on the models' original probability scales. Calibration was assessed for the all-incident outcome, our primary endpoint, whereas discrimination, which depends only on the ranking of predicted risks, was additionally evaluated at the 5-year and 10-year horizons.

As a sensitivity analysis, we re-evaluated all scores on the holdout test set with death without dementia treated as a competing event (Additional file 1: Supplementary Note, "Competing risk of death"). Because the latest baseline visit occurred more than 15 years before the administrative censoring date and participants lost to follow-up had been excluded, no participant was censored before the 5-, 10- or 15-year horizons. The binary horizon labels therefore coincide with the observed cumulative-incidence status at these time points. Under this design, we compared time-dependent AUCs that counted participants who died as controls or excluded them, concordance indices with deaths censored or retained as non-events, cause-specific Cox and Fine–Gray hazard ratios and the Aalen–Johansen cumulative incidence of dementia across ITU-DESC risk deciles.

Beyond discrimination and calibration, we interpreted the fitted model using SHAP values, the same attribution method used for feature screening in Section 2.7. These attributions were computed post hoc, after the feature set and final model had been fixed, and therefore describe how the fitted model uses its predictors rather than contributing to model construction. Unless otherwise stated, SHAP values were computed on the holdout test set using the model refitted on the full development set. Global feature importance was summarized as the mean absolute SHAP value across holdout-test participants, whereas individual predictions were examined using local SHAP explanations for representative cases. The cumulative-discrimination analysis was the only exception: predictors were ranked by mean absolute SHAP value computed on the development set, ensuring that their ordering was determined without reference to the holdout test data.

Baseline characteristics of the 38 predictors were compared between participants with incident dementia ($n = 11{,}380$) or incident Alzheimer's disease ($n = 5{,}122$) and those free of incident dementia ($n = 489{,}030$) using the standardized mean difference (SMD), computed on observed values, with an absolute SMD of at least 0.1 considered meaningful. Univariate associations of each predictor with incident dementia and with incident Alzheimer's disease were summarized as odds ratios with 95% confidence intervals from logistic regression, with continuous predictors z-standardized so that each odds ratio corresponds to a one-SD increase (Additional file 1: Supplementary Note, "Baseline characteristics and univariate associations").

Finally, we tested post hoc whether missingness in the final predictors was itself associated with the outcome in this cohort (Section 2.6). For each of the 37 ITU-DESC predictors with any missing values, we estimated the association between its missingness and incident dementia in the full study cohort using logistic regression,

unadjusted and adjusted for age, sex, education, deprivation, recruitment year, assessment center and ethnicity (Additional file 1: Supplementary Note, "Missingness as a predictive signal").

2.10 Use of generative artificial intelligence

Generative artificial intelligence (AI) tools were used to assist with code implementation and refinement, grammar correction, and improvements to manuscript readability and clarity. The authors reviewed all AI-assisted outputs, checked the code and resulting analyses, and revised the text. The authors approved the final manuscript and take full responsibility for the code, analyses, text and interpretation of the findings.

3 Results

3.1 Baseline characteristics

Age separated the groups most sharply: participants who later developed dementia were about eight years older at baseline than those who remained dementia-free (median 65.9 versus 58.1 years; SMD 1.19; Additional file 1: Table S7). Maternal survival came next, with 87.3% of future dementia cases reporting that their mother had died, against 59.2% of the remainder (SMD −0.67). *APOE* $\varepsilon 4$ carriage, the Alzheimer's disease polygenic risk score and the number of medications taken followed, each near an SMD of 0.56. Most other selected predictors showed modest-to-moderate univariate differences, suggesting that prediction depended on their combined multivariable contribution rather than on uniformly large standalone effects. An extended baseline table reporting every response category and the number of participants with a recorded value in each group is provided as Additional file 1: Table S8, and univariate associations of each predictor with both outcomes as Table S9.

3.2 Comparison with existing risk scores

We compared ITU-DESC with UKB-DRP [10], UKBDRS [7], DRS [6], CAIDE [5] and ANU-ADRI [4] using isotonic-calibrated predictions from the holdout test set (Table 1). All risk scores were evaluated under the same random-imputation policy used for ITU-DESC, allowing every model to score the same 50,041 holdout-test participants without restricting any comparison to a complete-case subset. ITU-DESC showed the highest discrimination, with an AP of 0.138 (95% CI, 0.123 to 0.156) and a ROC AUC of 0.853 (0.844 to 0.863), exceeding all comparators on both metrics. UKB-DRP ranked second, with an AP of 0.118 (0.105 to 0.133) and a ROC AUC of 0.841 (0.831 to 0.851). The differences between ITU-DESC and UKB-DRP were 0.021 (0.008 to 0.033) for AP and 0.013 (0.006 to 0.019) for ROC AUC, the smallest margins among the comparators; both confidence intervals excluded zero. Applying random imputation rather than each model's native missing-value handling also prevented models from gaining discrimination through informative missingness. Consistent with this effect, UKB-DRP achieved higher discrimination with native LightGBM missing-value handling than under random imputation (AP, 0.127 versus 0.118; ROC AUC, 0.847 versus 0.841), both estimated from isotonic-calibrated holdout-test predictions.

Figure 3 shows the ROC and precision–recall curves for all-cause dementia in panels a and b, and for Alzheimer's disease in panels c and d. Additional file 1: Fig. S1 and Tables S10 and S11 report the corresponding results using uncalibrated predictions.

Table 1 Discrimination of ITU-DESC and established dementia risk scores for all-incident dementia

Model	**AP** (95% CI)	**ΔAP** (95% CI)	*P*
ITU-DESC (our model)	0.138 (0.123, 0.156)	—	—
UKB-DRP [10]	0.118 (0.105, 0.133)	0.021 (0.008, 0.033)	< 0.01
UKBDRS [7]	0.102 (0.092, 0.114)	0.036 (0.023, 0.050)	< 0.001
UKBDRS (w/o *APOE* $\varepsilon 4$)	0.084 (0.076, 0.093)	0.055 (0.041, 0.070)	< 0.001
DRS [6]	0.064 (0.058, 0.069)	0.075 (0.061, 0.090)	< 0.001
CAIDE [5]	0.044 (0.040, 0.048)	0.094 (0.081, 0.110)	< 0.001
CAIDE (w/o *APOE* $\varepsilon 4$)	0.035 (0.032, 0.038)	0.104 (0.090, 0.120)	< 0.001
ANU-ADRI [4]	0.032 (0.029, 0.036)	0.106 (0.092, 0.123)	< 0.001
Model	**ROC AUC** (95% CI)	**ΔROC AUC** (95% CI)	*P*
ITU-DESC (our model)	0.853 (0.844, 0.863)	—	—
UKB-DRP [10]	0.841 (0.831, 0.851)	0.013 (0.006, 0.019)	< 0.001
UKBDRS [7]	0.825 (0.814, 0.836)	0.028 (0.021, 0.035)	< 0.001
UKBDRS (w/o *APOE* $\varepsilon 4$)	0.812 (0.801, 0.823)	0.041 (0.033, 0.049)	< 0.001
DRS [6]	0.789 (0.778, 0.800)	0.065 (0.056, 0.074)	< 0.001
CAIDE [5]	0.704 (0.691, 0.717)	0.149 (0.136, 0.162)	< 0.001
CAIDE (w/o *APOE* $\varepsilon 4$)	0.654 (0.641, 0.666)	0.200 (0.186, 0.214)	< 0.001
ANU-ADRI [4]	0.592 (0.575, 0.608)	0.262 (0.243, 0.280)	< 0.001

All models were evaluated on the same holdout test set ($n = 50{,}041$; 1,138 incident dementia cases), under identical outcome definitions, using isotonic-calibrated predictions. The upper panel reports average precision and the lower panel reports ROC AUC for the same models. Differences are ITU-DESC minus the comparator and were calculated from unrounded estimates. ROC AUC estimates, confidence intervals, differences and two-sided *P* values were derived using DeLong's method. AP confidence intervals were obtained by paired subject-level bootstrap resampling (10,000 resamples; percentile method), whereas AP *P* values were obtained using two-sided paired permutation tests (2,000 permutations). AP, average precision; CI, confidence interval; ROC AUC, area under the receiver operating characteristic curve.

For Alzheimer's disease, ITU-DESC again showed the highest discrimination (Additional file 1: Table S12), with an AP of 0.070 (95% CI, 0.059 to 0.085) and a ROC AUC of 0.863 (0.850 to 0.876). UKB-DRP ranked second, with an AP of 0.056 (0.047 to 0.067) and a ROC AUC of 0.844 (0.829 to 0.858); the differences from ITU-DESC were 0.014 (0.005 to 0.025) for AP and 0.019 (0.010 to 0.028) for ROC AUC. All scores ranked in the same order as for all-cause dementia, and ITU-DESC exceeded every score on both metrics ($P < 0.01$ for every comparison). AP fell for every model relative to all-cause dementia, consistent with the lower prevalence of Alzheimer's disease (513 versus 1,138 cases). ROC AUC, which does not depend on prevalence, rose slightly for ITU-DESC (0.863 versus 0.853).

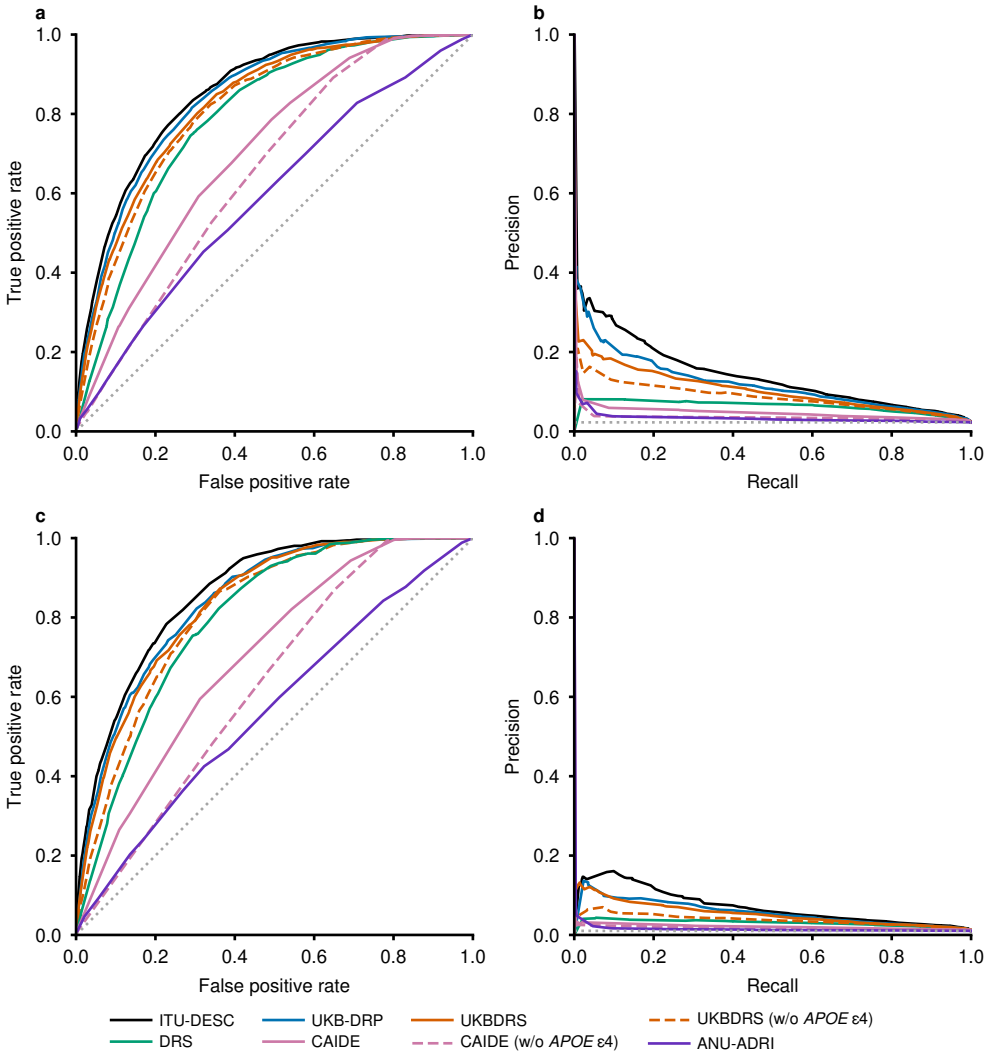

Fig. 3 Comparison of ITU-DESC with established dementia risk scores. All models were evaluated on the holdout test set using calibrated predictions. **a**, Receiver operating characteristic curves for all-incident all-cause dementia (1,138 incident cases). **b**, Precision–recall curves for all-incident all-cause dementia. **c**, Receiver operating characteristic curves for all-incident Alzheimer's disease (513 incident cases). **d**, Precision–recall curves for all-incident Alzheimer's disease. The dotted line marks chance discrimination in **a** and **c**, and the outcome prevalence in **b** and **d**, which equals the average precision of a non-informative model. Discrimination estimates and pairwise comparisons are reported in Table 1 for dementia and Additional file 1: Table S12 for Alzheimer's disease.

3.3 Discrimination across time horizons

ITU-DESC produces a single risk score for each participant, which we evaluated on the holdout test set at the 5-year, 10-year and all-incident horizons. The base model was fitted once and assigned each participant the same raw score at every horizon; the evaluations differed only in the outcome label used for each time window. Differences in discrimination across horizons therefore reflect differences in the outcome definition rather than changes in the underlying model. Each horizon was assigned a separate isotonic calibration mapping, so the calibrated risks differed in scale and in the ties introduced by calibration, while preserving the underlying ordering of the raw scores (Fig. 4; Additional file 1: Table S13). For all-cause dementia, ROC AUC increased modestly with longer follow-up, from 0.837 at 5 years to 0.850 at 10 years and 0.853 for the all-incident outcome. AP increased more substantially, from 0.012 to 0.056 and 0.138 across the same horizons.

For Alzheimer's disease, ROC AUC was similar across horizons: 0.868 at 5 years, 0.866 at 10 years and 0.863 for the all-incident outcome. AP increased from 0.007 at 5 years to 0.030 at 10 years and 0.070 for all-incident Alzheimer's disease. Thus, ranking performance changed little as the outcome window widened, whereas AP increased with outcome prevalence, which rose from 0.15% to 0.84% to 2.27% for dementia and from 0.05% to 0.33% to 1.03% for Alzheimer's disease (Fig. 1 reports the corresponding case counts). Because the AP of a non-informative model equals the outcome prevalence, the increase in AP across horizons is consistent with the accumulation of cases rather than improved ranking performance. Median follow-up was 16.6 years (interquartile range, 15.8–17.4), and every participant had at least 15.1 years of potential follow-up; no participant was therefore censored before the 5- or 10-year horizons. At these horizons, the binary labels coincide with the observed cumulative-incidence status, with death before dementia treated as a non-event (Additional file 1: Table S20). A sensitivity analysis treating death as a competing event (Additional file 1: Supplementary Note, "Competing risk of death"; Table S22) did not alter the ranking of the scores at any horizon. Corresponding analyses based on uncalibrated predictions are shown in Additional file 1: Table S14 and Fig. S2.

3.4 Calibration

We assessed ITU-DESC calibration for the all-incident outcomes of all-cause dementia and Alzheimer's disease. For each outcome, we plotted a flexible calibration curve against the line of perfect calibration (Fig. 5). Calibration curves were estimated using locally estimated scatterplot smoothing (LOESS) of observed incidence against predicted risk. Pointwise 95% confidence bands were calculated as the smoothed estimate plus or minus 1.96 standard errors. At each predicted-risk value the standard error was that of the local observed proportion and therefore depended on how many participants contributed to the fit at that point. For both outcomes, the curves closely followed the diagonal, indicating close agreement between predicted and observed risks. Uncertainty increased at the upper end of the predicted-risk distribution, where fewer participants were represented, as shown by the accompanying

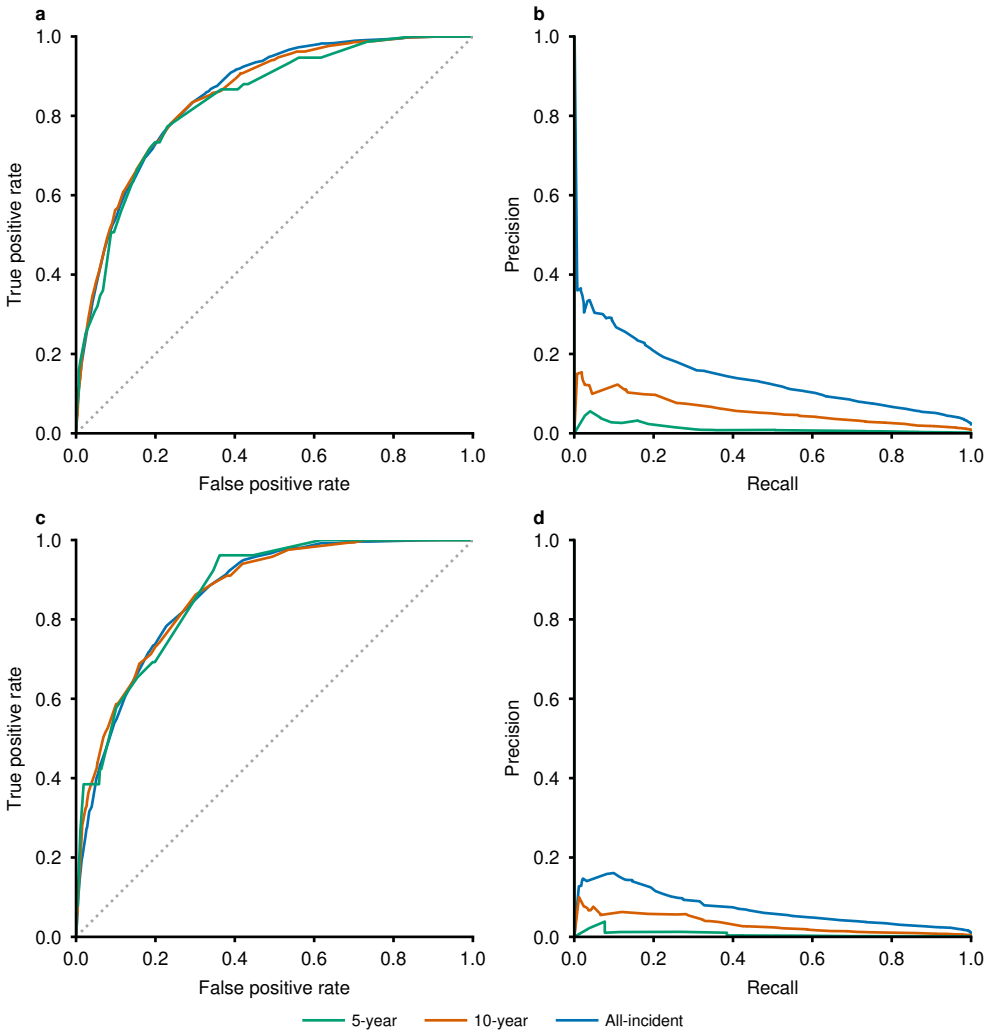

Fig. 4 Discrimination of ITU-DESC across the 5-year, 10-year and all-incident horizons. A single fitted model underlies all three horizons, evaluated on the holdout test set using calibrated predictions. **a**, Receiver operating characteristic curves for all-cause dementia. **b**, Precision–recall curves for all-cause dementia. **c**, Receiver operating characteristic curves for Alzheimer's disease. **d**, Precision–recall curves for Alzheimer's disease. The dotted line marks chance discrimination in **a** and **c**. Discrimination values are reported in Additional file 1: Table S13, and the case counts underlying each horizon in Fig. 1.

risk-distribution histograms. This region was also sparsely populated when the isotonic calibration mappings were fitted, making calibrated estimates in the upper-risk tail more susceptible to overfitting.

Weak-calibration summaries were consistent with this pattern. For both outcomes the calibration slope was close to one and the calibration intercept close to zero, and mean predicted risk closely matched observed incidence (2.27% versus 2.27% for dementia; 1.02% versus 1.03% for AD). Calibration-in-the-large is reported on both the probability and log-odds scales, which agree in indicating no systematic over- or underestimation. In the competing-risk sensitivity analysis, the mean calibrated 10-year ITU-DESC risk tracked the Aalen–Johansen cumulative incidence of dementia in every risk decile (Additional file 1: Table S21 and Fig. S3). Because incident outcomes were rare, the raw Brier scores were close to their no-information baselines; we therefore also report the scaled Brier score, which indicated a modest but consistent improvement over a non-informative model. Full calibration statistics, with 95% CIs and metric definitions, are reported in Additional file 1: Supplementary Note (Table S23).

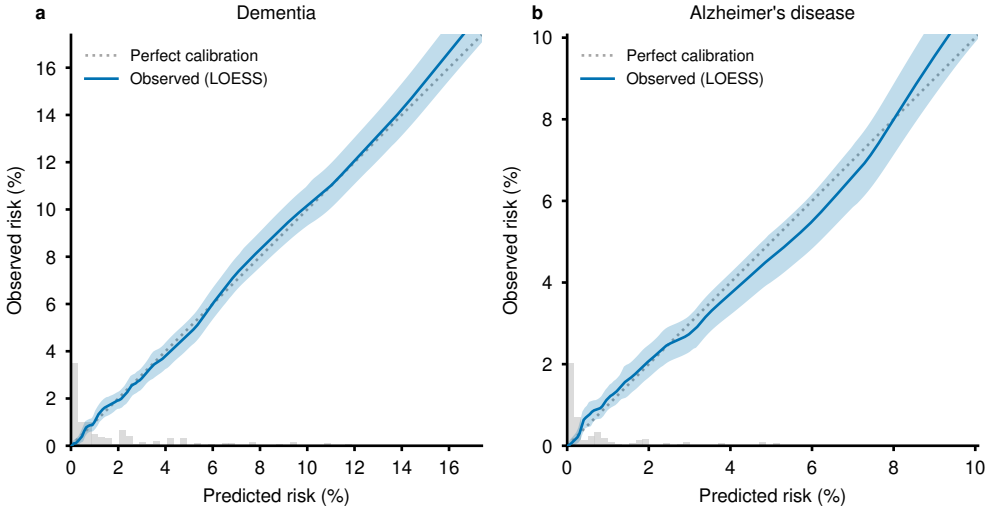

Fig. 5 Flexible calibration of ITU-DESC for the all-incident outcomes. Both panels use the holdout test set. **a**, All-cause dementia. **b**, Alzheimer's disease. Observed risk, estimated by locally estimated scatterplot smoothing (LOESS), is plotted against predicted risk with pointwise 95% confidence bands, calculated as the smoothed estimate plus or minus 1.96 standard errors of the local observed proportion. Histograms along the bottom of each panel show the distribution of predicted risks. Calibration slope, mean predicted risk against observed incidence and Brier score are given in Additional file 1: Table S23.

3.5 Feature importance

The SHAP summary plot (Fig. 6) shows each feature's signed contribution to ITU-DESC predictions, with color indicating the corresponding feature value. Features

are ranked by mean absolute SHAP value computed on the holdout test set. The complete list of selected features is provided in Additional file 1: Supplementary Note (Table S15).

Fig. 6 SHAP summary plot for ITU-DESC for all-incident dementia. Each point represents one participant and is plotted according to the feature's SHAP value, corresponding to its signed contribution to the model output on the log-odds scale. Points are colored by feature value. Features are ordered from top to bottom by mean absolute SHAP value. SHAP values were computed on the holdout test set ($n = 50{,}041$), using the model refit on the full development set.

After identifying the final set of 38 predictors, we examined whether predictive performance was concentrated in a small subset of variables or distributed across the selected feature set. We evaluated cumulative discrimination by adding predictors in SHAP-rank order (Fig. 7), adapting an incremental feature-addition approach from You et al. [10]. Here the ranking was derived from mean absolute SHAP values computed on the development set rather than the holdout-test values reported above, so the ordering in which features enter never uses test data. For each top-k feature set, the model was refit on the development set and evaluated on the holdout test set, and ROC AUC and AP were recorded. Age alone gave a ROC AUC of 0.799 and an AP of 0.076; the full 38-feature model reached 0.854 and 0.142. ROC AUC peaked at 0.856 with 35 features and stayed within 0.007 of that peak from 21 features onward, so the first 21 predictors carried nearly all of the ranking performance. These values come from uncalibrated scores and from a separate refit at each top-k; the calibrated ITU-DESC results are in Table 1 (0.853; 0.138) and the uncalibrated ones in Additional file 1: Table S10 (0.854; 0.142).

To illustrate individual-level model predictions, we also examined local SHAP explanations for representative true-positive, true-negative, false-positive and false-negative predictions. These examples show how different combinations of feature contributions led to correct and incorrect classifications (Additional file 1: Figs. S5, S6, S7 and S8; Supplementary Note).

Because random imputation removes predictive information carried by missingness, the ranking in Fig. 6 could depend on the imputation strategy. However, refitting the model using CatBoost's native handling of missing values produced the same features at the top and bottom of the ranking, indicating that the ordering interpreted here was robust to the choice of missing-data strategy (Additional file 1: Fig. S4; Supplementary Note).

3.6 Selected predictors

Age was the most important predictor in ITU-DESC (Fig. 6), consistent with age being the dominant risk factor for dementia. This variable refers to the participant's age at the initial UK Biobank assessment visit.

Family history contributed three predictors. The model retained maternal history of Alzheimer's disease or dementia and maternal history of high blood pressure, but no paternal counterpart. It also retained the combined maternal survival and age-at-death variable described in Section 2.5.

Several selected predictors captured general health status or neuropsychiatric symptoms: long-standing illness, disability or infirmity; number of treatments or medications taken; taking other prescription medications; falls in the last year; sleeplessness or insomnia; mood swings; worrying too long after embarrassment; and hearing difficulty or problems with background noise.

Blood-based and urinary biomarkers also contributed to prediction. These included aspartate aminotransferase (AST), alanine aminotransferase (ALT), neutrophil percentage, basophil count, C-reactive protein, insulin-like growth factor 1 (IGF-1), vitamin D, saturated fatty acids as a percentage of total fatty acids and enzymatic creatinine in urine.

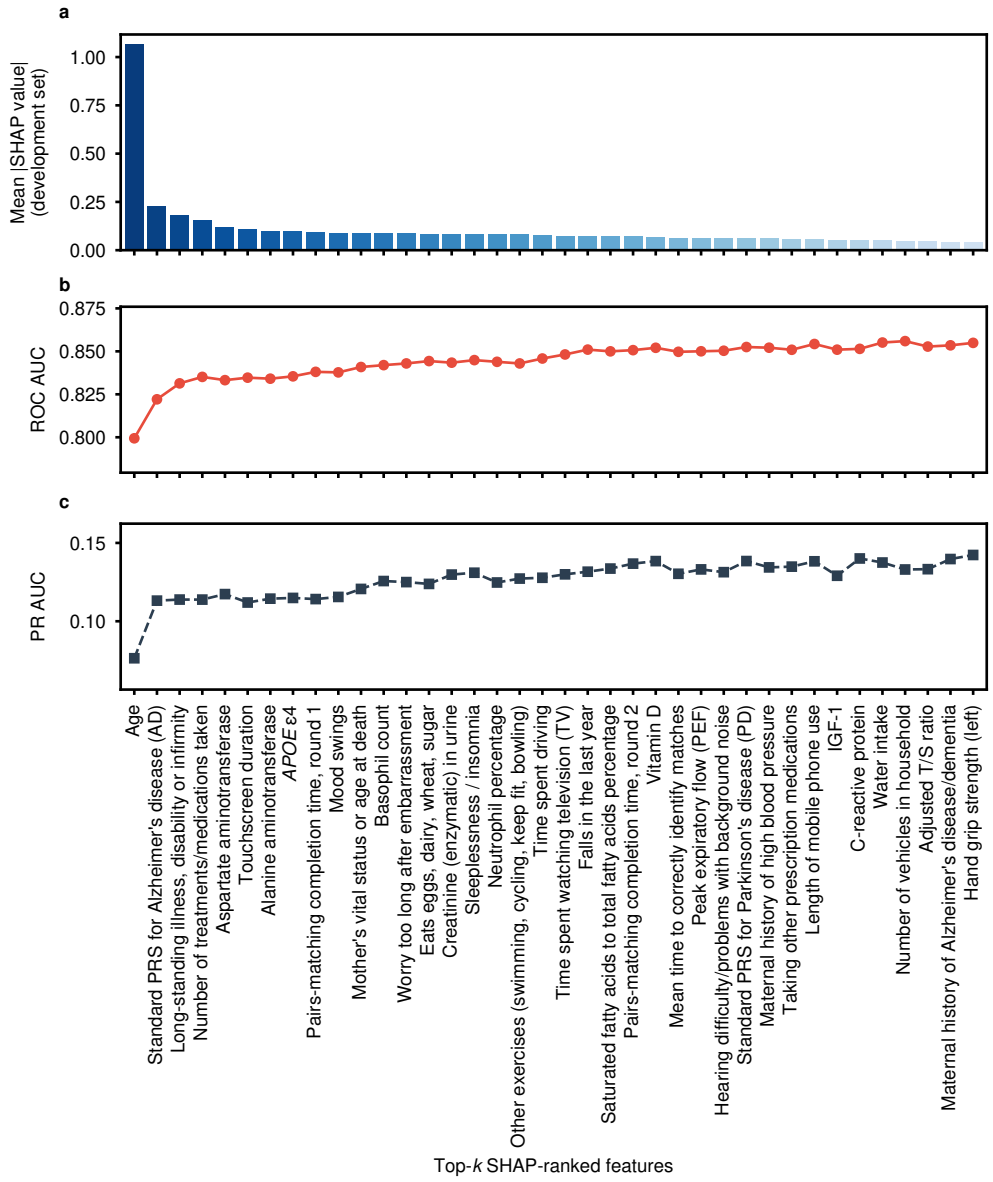

Fig. 7 Cumulative discrimination of ITU-DESC as predictors are added in SHAP-rank order for all-incident dementia. The x-axis lists the top-k features ranked by mean absolute SHAP value computed on the 90% development set. **a**, Development-set SHAP importance of each feature. **b**, ROC AUC when the model is refit on the development set using the top-k features and evaluated on the holdout test set ($n = 50{,}041$). **c**, Precision–recall AUC (average precision) for the same refits. Each point is a single holdout-test evaluation.

The model retained leukocyte telomere length as the telomere-to-single-copy-gene ratio adjusted for assay-related technical parameters [41, 42]. Genetic predictors included *APOE* $\varepsilon 4$ allele copy number [43], the standard polygenic risk score (PRS) for Alzheimer's disease and the standard PRS for Parkinson's disease [44].

Cognitive, sensorimotor and task-performance measures were also selected. These included touchscreen duration, pairs-matching completion time for rounds 1 and 2, mean time to correctly identify matches, left hand grip strength and peak expiratory flow (PEF). You et al. [10] used the summed first and second rounds of UK Biobank Field 400 [45]; our pipeline instead retained the individual rounds.

Lifestyle and socioeconomic variables were retained as well. Behavioral predictors included time spent watching television; time spent driving; types of physical activity in the last 4 weeks; water intake; consumption of eggs, dairy, wheat or sugar; and length of mobile phone use. For physical activity, the retained category corresponded to other exercises, such as swimming, cycling, keep fit or bowling (Additional file 1: Supplementary Methods, "Feature sets and related data fields"). The model also retained the number of vehicles in the household as a socioeconomic predictor.

Together, these variables formed Set 8, the 38 features selected in at least eight of the ten imputation seeds. Of these, 29 were selected by every seed, 5 by nine seeds and 4 by eight seeds; all 38 are listed with their data-field identifiers in Additional file 1: Table S15. Further variables fell just below the cutoff: five were selected in seven seeds, one in six seeds and two in five seeds.

The seven-seed variables were tea intake, right and left leg fat percentage, glucose and urea; omega-3 fatty acids was selected in six seeds; and seating box height and average total household income before tax were selected in five. Although these variables did not enter the final feature set, their repeated selection across half or more of the seeds suggests that they may carry additional predictive information.

3.7 Missingness of the selected predictors

Missingness in the selected predictors was associated with incident dementia. Of the 38 predictors, age was complete for all participants. For each of the remaining 37, missingness was more frequent among participants who subsequently developed dementia than among those who did not (Additional file 1: Table S18; the corresponding results for Alzheimer's disease are in Table S19). For 33 of the 37 predictors, this association persisted after adjustment for age, sex, education, deprivation, recruitment year, assessment center and ethnicity (adjusted odds ratios, 1.07 to 2.44). Dementia incidence rose from 1.76% among participants with no missing predictors to 3.61% among those missing three or four (Additional file 1: Table S17; Supplementary Note). Baseline characteristics of participants complete on all 38 predictors and those missing at least one are compared in Additional file 1: Table S16.

4 Discussion

In this UK Biobank study, an explainable machine-learning model developed directly from cohort data outperformed established dementia risk scores. Starting from approximately 8,000 candidate features, the wrapper-based selection pipeline identified 38

predictors. ITU-DESC achieved higher discrimination than the strongest comparator, UKB-DRP (Table 1). Several selected predictors that are absent from established risk scores, including leukocyte telomere length and the liver enzymes aspartate aminotransferase (AST) and alanine aminotransferase (ALT), illustrate how wrapper-based feature selection can identify biologically plausible predictors that are not included in hand-crafted scores. This discrimination was accompanied by well-calibrated predictions: after outcome-specific isotonic recalibration, calibration curves closely followed the diagonal for both outcomes, with calibration slopes near one and calibration-in-the-large near zero. ITU-DESC predictions can therefore be interpreted as estimated probabilities rather than solely as relative risk rankings, although these probabilities remain specific to the outcome definitions, follow-up periods and ascertainment framework of the present study.

The primary objective of this study was to determine whether leakage-aware machine-learning feature selection across a broad UK Biobank feature space could improve discrimination and identify additional predictors. You et al. [10] and Anatürk et al. [7] emphasized cost, clinical usability and implementation burden; these remain important for translation into practice but were beyond the present scope.

Age was the most important predictor in ITU-DESC, consistent with the findings of You et al. [10]. Both of the models identified a shared core of predictors, including age, *APOE* $\varepsilon 4$, cognitive and task-performance measures, medication burden, mother's vital status or age at death and long-standing illness.

Our model combined maternal survival and maternal age at death into a single variable, mother's vital status or age at death, which follows the interpretation proposed by You et al.: a missing value for mother's age at death usually indicates that the participant's mother was still alive [22]. Two further features identified by You et al. behaved differently in our pipeline. Peak expiratory flow entered our final feature set, selected in nine of ten seeds, whereas leg fat percentage fell just below the cutoff at seven seeds. Their reaction-time measure, the duration to first press of the snap-button in each round, was selected by no seed at all; the related pairs-matching timings were retained instead.

The liver enzymes AST and ALT and leukocyte telomere length, all absent from established dementia risk scores, have independent biological support. A higher AST-to-ALT ratio has been linked with increased risk of cognitive impairment [46]. Evidence for associations involving AST and ALT also extends beyond UK Biobank: a systematic review summarized findings across multiple cohorts [47], including studies of both dementia and broader cognitive impairment. The adjusted telomere-to-single-copy-gene ratio was also retained, consistent with the hypothesis that telomere-related aging may provide predictive information beyond chronological age [42].

Several additional selected features relate to behavior and demographics, and their associations with dementia may be indirect. Sex was not retained at all: no imputation seed selected it, which suggests that other selected variables already capture the information it carries.

Time spent watching television and types of physical activity have been examined in relation to dementia risk in previous work [48]. Li et al. reported that, among socially inactive older adults, moderate computer use was associated with lower

dementia risk, whereas heavy computer use and television viewing above approximately 2 hours per day were associated with higher risk. Length of mobile phone use was negatively associated with incident dementia in our model, but this association may partly reflect age or cohort effects, because younger participants are likely to have used mobile phones for longer.

Evaluating every comparator under the random-imputation policy used for ITU-DESC (Section 2.6) eliminated cohort differences that would otherwise have confounded comparisons across scores. For example, *APOE* $\varepsilon 4$ status was missing for 88,961 of the 500,410 eligible participants. Applying complete-case analysis, as UKBDRS specifies, would therefore have restricted the *APOE*-inclusive UKBDRS to a smaller, potentially selected subset, whereas the other models could be evaluated in the full sample. Under random imputation, all eight models were instead evaluated in the same 50,041 holdout-test participants, using identical outcome definitions and with identical outcome prevalence. Consequently, differences in discrimination are attributable to differences between the scores rather than to differences in the cohorts on which they were evaluated.

The shared missing-data policy introduces a trade-off that should be considered when interpreting the benchmark results. UKBDRS, CAIDE, DRS and ANU-ADRI were developed under different missing-data procedures, and evaluating them on randomly imputed data may therefore attenuate their performance relative to published estimates, particularly for scores containing variables with substantial missingness. Because a randomly imputed value is independent of the participant to whom it is assigned, it contributes no participant-specific predictive information. Accordingly, differences in AP or ROC AUC between ITU-DESC and these scores should not be attributed to model structure alone.

4.1 Limitations

Random imputation reduced reliance on deterministic missingness patterns but may also have introduced noise and weakened genuine predictive signal. Under informative missingness, no single-imputation procedure can recover the target regression function unless information about missingness is also available to the model [13]. Because donors were drawn from each variable's marginal distribution, the procedure also attenuated relationships between variables that a conditional donor pool could retain [29]. Conditioning donor selection on the remaining covariates, or guiding it with predictive models [49], would preserve that structure but would reintroduce the dependence on measurement availability we set out to suppress. Whether an imputation rule can preserve that structure while withholding missingness patterns from the model remains an open question. Our inclusive design retained all eligible participants (Section 2.1), which maximized sample size and avoided additional selection based on missingness but may have increased the variance introduced by imputation and made performance more sensitive to the chosen missing-data strategy.

Our evaluation used binary classification rather than time-to-event analysis, and participants who died without a recorded dementia diagnosis were counted as non-cases. UKBDRS was originally developed and evaluated in a competing-risk framework that accounts for follow-up time, censoring and death as a competing event,

so the AP and ROC AUC reported here should not be interpreted as reproducing its published 14-year performance. At the 5- and 10-year horizons, the sensitivity analysis treating death as a competing event supported the binary results (Additional file 1: Supplementary Note, "Competing risk of death"), but the all-incident outcome combines follow-up durations from 15.1 to 19.7 years and therefore does not represent a fixed-horizon cumulative incidence. In addition, non-case status reflects the absence of a recorded diagnosis in the current data release, so the observed incidence is a lower bound over the same follow-up and performance estimates may change modestly as additional diagnoses accrue.

Volunteer cohorts such as UK Biobank are susceptible to healthy-volunteer bias [50]. Selection also affects the imaging subsample, whose participants were healthier on average than the rest of the cohort, so analyses restricted to it may underestimate exposure–outcome associations [51]. Phenotypes from the imaging visits lay outside our baseline-only feature set, which also prevented imaging availability from acting as a proxy for participant selection or health status.

We used default hyperparameters and did not perform hyperparameter tuning, because tuning nested within BorutaSHAP feature selection and repeated across ten imputation seeds was computationally prohibitive. This may have limited model performance and should be addressed in future work.

Because ITU-DESC was developed and internally validated in UK Biobank, its generalizability to other populations remains uncertain. We also included the full eligible cohort across the available baseline age range, whereas many existing risk scores were developed for narrower age ranges or target populations, so direct comparisons should be interpreted with this difference in mind.

4.2 Future work

External validation in independent cohorts with different demographic, clinical and geographic characteristics is required to confirm discrimination and to assess whether calibration transports and whether the selected feature set and missing-data strategy remain effective in other populations and data-collection settings. Clinical usability and implementation burden should also be evaluated before deployment. Application of ITU-DESC to a new dataset requires the same 38 baseline predictors and the same unconditional random hot-deck imputation procedure used here, with donor values drawn from the target population or, if unavailable, from the UK Biobank development set. ITU-DESC is intended for researchers and analysts working with tabular cohort data and requires a qualified user to prepare the inputs and interpret the calibrated probabilities. It is not intended for autonomous clinical use.

Future work could evaluate more principled approaches for informative and non-monotone missingness, including statistical frameworks designed for MNAR data with limited complete cases [52]. Our study did not incorporate advanced genetic biomarkers beyond *APOE* $\varepsilon 4$ status and polygenic risk scores. Future work could evaluate whether richer genomic information, including whole-genome sequencing (WGS), improves dementia risk prediction, although cost, scalability and interpretability remain important considerations. Repeated UK Biobank assessments could be used to model longitudinal change, especially among participants who later develop dementia,

to reveal pre-diagnostic trajectories that a single baseline visit misses. Imaging-derived predictors could also be evaluated in the imaging subsample, with attention to the selection bias described above. These extensions should be weighed against their additional complexity, data requirements and reduced interpretability.

The feature-selection and missing-data pipeline is not specific to dementia. Applying it to other UK Biobank outcomes with informative missingness would test the general framework rather than the ITU-DESC score itself and would show whether consensus feature selection and random imputation perform similarly under different missingness structures. More broadly, wrapper-based selection methods such as BorutaSHAP [33] or Boruta [34] may be useful in other UK Biobank tabular-data studies by reducing very large candidate feature spaces to more explainable and manageable predictor sets.

5 Conclusions

Using UK Biobank data, we developed ITU-DESC, an explainable machine-learning dementia risk score based on 38 predictors selected from a broad baseline feature space. ITU-DESC showed higher discrimination on a holdout test set than established dementia risk scores for all-cause dementia and Alzheimer's disease, while identifying predictors absent from existing scores, such as leukocyte telomere length and liver-enzyme markers. These findings suggest that missingness-aware machine-learning feature selection can improve dementia risk prediction in large-scale cohort data. External validation is required to determine whether its performance and selected predictors generalize beyond UK Biobank.

Declarations

Ethics approval and consent to participate. UK Biobank has Research Tissue Bank approval from the North West – Haydock Research Ethics Committee (REC reference 21/NW/0157), and all participants provided written informed consent. This study was conducted in accordance with the UK Biobank Ethics and Governance Framework and the Declaration of Helsinki.

Consent for publication. Not applicable.

Availability of data and materials. UK Biobank data [17] are used under license and not publicly available; they are available to approved researchers through the UK Biobank Access Management System [53]. The columns of the candidate feature set and the entries excluded a priori are listed in Additional file 2. The custom analysis code used in this study is not publicly available but will be made available to editors and reviewers upon request.

Competing interests. The authors declare that they have no competing interests.

Funding. Abdullah Kavaklı received support from the Global Researcher Access Fund to cover the UK Biobank application costs. Compute and storage costs were supported by the UK Biobank Platform Credits Programme. Faik Boray Tek's contribution to this work was supported by the Scientific and Technological Research

Council of Türkiye (TÜBİTAK) under Grant Number 122E509. The funders had no role in study design; data collection; analysis or interpretation; manuscript preparation; or the decision to submit the manuscript for publication.

Authors' contributions. A.K. and F.B.T. jointly designed and planned the study. A.K. curated and prepared the data; carried out the experiments and analyses; and wrote the initial draft of the manuscript. F.B.T. supervised the project and critically revised the manuscript. Both authors read and approved the final manuscript.

Acknowledgements. This research has been conducted using the UK Biobank Resource under Application Number 412140.

References

[1] World Health Organization.: Dementia. Published 31 March 2025. https://www.who.int/news-room/fact-sheets/detail/dementia. Accessed 6 May 2026.

[2] Brück CC, Mooldijk SS, Kuiper LM, Sambou ML, Licher S, Mattace-Raso F, et al. Time to nursing home admission and death in people with dementia: systematic review and meta-analysis. BMJ. 2025;388:e080636. https://doi.org/10.1136/bmj-2024-080636.

[3] Livingston G, Huntley J, Liu KY, Costafreda SG, Selbæk G, Alladi S, et al. Dementia prevention, intervention, and care: 2024 report of the Lancet standing Commission. The Lancet. 2024;404(10452):572–628. https://doi.org/10.1016/S0140-6736(24)01296-0.

[4] Anstey KJ, Cherbuin N, Herath PM. Development of a new method for assessing global risk of Alzheimer's disease for use in population health approaches to prevention. Prevention Science. 2013;14(4):411–421. https://doi.org/10.1007/s11121-012-0313-2.

[5] Kivipelto M, Ngandu T, Laatikainen T, Winblad B, Soininen H, Tuomilehto J. Risk score for the prediction of dementia risk in 20 years among middle aged people: a longitudinal, population-based study. The Lancet Neurology. 2006;5(9):735–741. https://doi.org/10.1016/S1474-4422(06)70537-3.

[6] Walters K, Hardoon S, Petersen I, Iliffe S, Omar RZ, Nazareth I, et al. Predicting dementia risk in primary care: development and validation of the Dementia Risk Score using routinely collected data. BMC Medicine. 2016;14(1):6. https://doi.org/10.1186/s12916-016-0549-y.

[7] Anatürk M, Patel R, Ebmeier KP, Georgiopoulos G, Newby D, Topiwala A, et al. Development and validation of a dementia risk score in the UK Biobank and Whitehall II cohorts. BMJ Mental Health. 2023;26(1):e300719. https://doi.org/10.1136/bmjment-2023-300719.

[8] Stubs J, Langballe EM, Livingston G, Anstey KJ, Deckers K, Mathews FE, et al. Dementia risk prediction: a comparative analysis of the ANU-ADRI, CAIDE, CogDrisk, LIBRA, and LIBRA2 indices in the HUNT study. The Journal of Prevention of Alzheimer's Disease. 2025;12(9):100326. https://doi.org/10.1016/j.tjpad.2025.100326.

[9] Stubs J, Selbæk G, Strand BH, Livingston G, Anstey KJ, Deckers K, et al. Predicting cognitive decline: comparative analysis of ANU-ADRI, CAIDE, CogDrisk, LIBRA, LIBRA2, UKBDRS and Lancet based dementia risk scores in the HUNT study. The Journal of Prevention of Alzheimer's Disease. 2026;13(4):100524. https://doi.org/10.1016/j.tjpad.2026.100524.

[10] You J, Zhang YR, Wang HF, Yang M, Feng JF, Yu JT, et al. Development of a novel dementia risk prediction model in the general population: A large, longitudinal, population-based machine-learning study. eClinicalMedicine. 2022;53:101665. https://doi.org/10.1016/j.eclinm.2022.101665.

[11] Ke G, Meng Q, Finley T, Wang T, Chen W, Ma W, et al. LightGBM: a highly efficient gradient boosting decision tree. In: Advances in Neural Information Processing Systems. vol. 30. Curran Associates, Inc.; 2017. p. 3146–3154. Available from: https://proceedings.neurips.cc/paper_files/paper/2017/file/6449f44a102fde848669bdd9eb6b76fa-Paper.pdf.

[12] Anatürk M, Patel R.: dementia_risk_score: include vs. exclude comparison. GitHub, commit `da9b4ba`. https://github.com/MelisAnaturk/dementia_risk_score/blob/da9b4ba0f021e3a1be5fb39dee83e0bc21123dbe/scripts/16_include_vs_exclude.R. Accessed 3 June 2026.

[13] Josse J, Chen JM, Prost N, Varoquaux G, Scornet E. On the consistency of supervised learning with missing values. Statistical Papers. 2024;65(9):5447–5479. https://doi.org/10.1007/s00362-024-01550-4.

[14] Little RJA, Rubin DB. Statistical Analysis With Missing Data. 3rd ed. Hoboken, NJ: John Wiley & Sons; 2019. Available from: https://doi.org/10.1002/9781119482260.

[15] Mitra R, McGough SF, Chakraborti T, Holmes C, Copping R, Hagenbuch N, et al. Learning from data with structured missingness. Nature Machine Intelligence. 2023;5(1):13–23. https://doi.org/10.1038/s42256-022-00596-z.

[16] Kapoor S, Narayanan A. Leakage and the reproducibility crisis in machine-learning-based science. Patterns. 2023;4(9):100804. https://doi.org/10.1016/j.patter.2023.100804.

[17] Sudlow C, Gallacher J, Allen N, Beral V, Burton P, Danesh J, et al. UK Biobank: an open access resource for identifying the causes of a wide range of complex diseases of middle and old age. PLOS Medicine. 2015;12(3):e1001779. https:

//doi.org/10.1371/journal.pmed.1001779.

[18] UK Biobank Outcome Adjudication Group and UK Biobank Data Analyst and Scientific Teams. Algorithmically-defined outcomes (ADOs), version 2.0. UK Biobank; 2022. https://biobank.ndph.ox.ac.uk/showcase/showcase/docs/alg_outcome_main.pdf. Accessed 25 March 2026.

[19] Little RJA. Regression with missing X's: A review. Journal of the American Statistical Association. 1992;87(420):1227–1237. https://doi.org/10.1080/01621459.1992.10476282.

[20] UK Biobank.: Instance 2: Assessment Centre Visit. https://biobank.ndph.ox.ac.uk/ukb/instance.cgi?id=2. Accessed 25 March 2026.

[21] Theodoridis S, Koutroumbas K. Pattern Recognition. 4th ed. Boston, MA: Academic Press; 2009. Available from: https://doi.org/10.1016/B978-1-59749-272-0.X0001-2.

[22] UK Biobank.: Data-Field 3526: Mother's age at death. Version: August 2025. https://biobank.ndph.ox.ac.uk/ukb/field.cgi?id=3526. Accessed 14 July 2026.

[23] UK Biobank.: Data-Field 1807: Father's age at death. Version: August 2025. https://biobank.ndph.ox.ac.uk/ukb/field.cgi?id=1807. Accessed 17 August 2026.

[24] Rubin DB. Inference and missing data. Biometrika. 1976;63(3):581–592. https://doi.org/10.1093/biomet/63.3.581.

[25] Chen T, Guestrin C. XGBoost: A scalable tree boosting system. In: Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. KDD '16. New York, NY, USA: ACM; 2016. p. 785–794. Available from: https://doi.org/10.1145/2939672.2939785.

[26] Hasler C, Tillé Y. Balanced k-nearest neighbour imputation. Statistics. 2016;50(6):1310–1331. https://doi.org/10.1080/02331888.2016.1230615.

[27] van Buuren S, Groothuis-Oudshoorn K. mice: Multivariate imputation by chained equations in R. Journal of Statistical Software. 2011;45(3):1–67. https://doi.org/10.18637/jss.v045.i03.

[28] Stekhoven DJ, Bühlmann P. MissForest—non-parametric missing value imputation for mixed-type data. Bioinformatics. 2012;28(1):112–118. https://doi.org/10.1093/bioinformatics/btr597.

[29] Chaput H, Chauvet G, Haziza D, Salembier L, Solard J. Joint imputation procedures for categorical variables; 2015. Preprint at https://arxiv.org/abs/1511.00990.

[30] Andridge RR, Little RJA. A review of hot deck imputation for survey non-response. International Statistical Review. 2010;78(1):40–64. https://doi.org/10.1111/j.1751-5823.2010.00103.x.

[31] Kim JK, Fuller W. Fractional hot deck imputation. Biometrika. 2004;91(3):559–578. https://doi.org/10.1093/biomet/91.3.559.

[32] Lundberg SM, Lee SI. A unified approach to interpreting model predictions. In: Advances in Neural Information Processing Systems. vol. 30. Curran Associates, Inc.; 2017. p. 4765–4774. Available from: https://proceedings.neurips.cc/paper_files/paper/2017/file/8a20a8621978632d76c43dfd28b67767-Paper.pdf.

[33] Keany E.: BorutaShap: A wrapper feature selection method which combines the Boruta feature selection algorithm with Shapley values, version 1.1. Zenodo; 2020. Available from: https://doi.org/10.5281/zenodo.4247618.

[34] Kursa MB, Rudnicki WR. Feature selection with the Boruta package. Journal of Statistical Software. 2010;36(11):1–13. https://doi.org/10.18637/jss.v036.i11.

[35] Prokhorenkova L, Gusev G, Vorobev A, Dorogush AV, Gulin A. CatBoost: unbiased boosting with categorical features. In: Advances in Neural Information Processing Systems. vol. 31. Curran Associates, Inc.; 2018. p. 6638–6648. Available from: https://proceedings.neurips.cc/paper_files/paper/2018/file/14491b756b3a51daac41c24863285549-Paper.pdf.

[36] Saito T, Rehmsmeier M. The precision-recall plot is more informative than the ROC plot when evaluating binary classifiers on imbalanced datasets. PLOS ONE. 2015;10(3):1–21. https://doi.org/10.1371/journal.pone.0118432.

[37] scikit-learn developers.: `sklearn.metrics.average_precision_score`. scikit-learn 1.5 documentation. https://scikit-learn.org/1.5/modules/generated/sklearn.metrics.average_precision_score.html. Accessed 3 June 2026.

[38] Van Calster B, McLernon DJ, van Smeden M, Wynants L, Steyerberg EW, On behalf of Topic Group 'Evaluating diagnostic tests and prediction models' of the STRATOS initiative. Calibration: the Achilles heel of predictive analytics. BMC Medicine. 2019;17(1):230. https://doi.org/10.1186/s12916-019-1466-7.

[39] Zadrozny B, Elkan C. Transforming classifier scores into accurate multiclass probability estimates. In: Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. KDD '02. New York, NY, USA: Association for Computing Machinery; 2002. p. 694–699. Available from: https://doi.org/10.1145/775047.775151.

[40] Niculescu-Mizil A, Caruana R. Predicting good probabilities with supervised learning. In: Proceedings of the 22nd International Conference on Machine Learning. ICML '05. New York, NY, USA: Association for Computing Machinery; 2005.

p. 625–632. Available from: https://doi.org/10.1145/1102351.1102430.

[41] UK Biobank.: Data-Field 22191: Adjusted T/S ratio. Version: August 2025. https://biobank.ndph.ox.ac.uk/ukb/field.cgi?id=22191. Accessed 29 May 2026.

[42] Codd V, Denniff M, Swinfield C, Warner SC, Papakonstantinou M, Sheth S, et al. Measurement and initial characterization of leukocyte telomere length in 474,074 participants in UK Biobank. Nature Aging. 2022;2(2):170–179. https://doi.org/10.1038/s43587-021-00166-9.

[43] Bycroft C, Freeman C, Petkova D, Band G, Elliott LT, Sharp K, et al. The UK Biobank resource with deep phenotyping and genomic data. Nature. 2018;562(7726):203–209. https://doi.org/10.1038/s41586-018-0579-z.

[44] Thompson DJ, Wells D, Selzam S, Peneva I, Moore R, Sharp K, et al. A systematic evaluation of the performance and properties of the UK Biobank polygenic risk score (PRS) release. PLOS ONE. 2024;19(9):e0307270. https://doi.org/10.1371/journal.pone.0307270.

[45] UK Biobank.: Data-Field 400: Time to complete round. Version: August 2025. https://biobank.ndph.ox.ac.uk/ukb/field.cgi?id=400. Accessed 14 May 2026.

[46] Li W, Yue L, Sun L, Xiao S. An increased aspartate to alanine aminotransferase ratio is associated with a higher risk of cognitive impairment. Frontiers in Medicine. 2022;9:780174. https://doi.org/10.3389/fmed.2022.780174.

[47] Zhong Y, Li L. Association between liver biomarkers and risk of cognitive impairment and dementia: A systematic review and meta-analysis. Pakistan Journal of Medical Sciences. 2025;41(7):2122–2132. https://doi.org/10.12669/pjms.41.7.12321.

[48] Li J, Xue H, Leng Y, Samus QM, Nowrangi M, Szanton SL, et al. Television and computer use and dementia risk in older adults with limited leisure or social activities: A prospective cohort study. Alzheimer's & Dementia. 2026;22(3):e71259. https://doi.org/10.1002/alz.71259.

[49] Tripet A, Eustache E, Tillé Y. Improving donor imputation using the prediction power of random forests: a combination of SwissCheese and missForest. Journal of Survey Statistics and Methodology. 2024;12(5):1389–1404. https://doi.org/10.1093/jssam/smad040.

[50] Brayne C, Moffitt TE. The limitations of large-scale volunteer databases to address inequalities and global challenges in health and aging. Nature Aging. 2022;2(9):775–783. https://doi.org/10.1038/s43587-022-00277-x.

[51] Lyall DM, Quinn T, Lyall LM, Ward J, Anderson JJ, Smith DJ, et al. Quantifying bias in psychological and physical health in the UK Biobank imaging

sub-sample. Brain Communications. 2022;4(3):fcac119. https://doi.org/10.1093/braincomms/fcac119.

[52] Cheng G, Chen YC, Smith MA, Zhao YQ. Handling nonmonotone missing data with available complete-case missing value assumption; 2022. Preprint at https://arxiv.org/abs/2207.02289.

[53] UK Biobank.: Apply for access. https://www.ukbiobank.ac.uk/use-our-data/apply-for-access/. Accessed 30 May 2026.

Additional files

Additional file 1 (.pdf) **Supplementary Information.** Supplementary Methods, Supplementary Note, Tables and Figures.
Additional file 2 (.xlsx) **Candidate feature set.** Candidate columns and entries excluded a priori.

ITU-DESC: a missingness-aware machine-learning dementia risk score in the UK Biobank cohort

Supplementary Information

Abdullah Kavaklı[1*] and Faik Boray Tek[1]

[1*]Department of Artificial Intelligence and Data Engineering, Istanbul Technical University, Istanbul, Turkey.

*Corresponding author(s). E-mail(s): kavakli23@itu.edu.tr;

Supplementary Methods

Development and holdout test split

The eligible cohort of 500,410 participants was divided into a 90% development set (450,369 participants) and a 10% holdout test set (50,041 participants), as described in the main text. This section documents the stratification scheme, the acceptance criteria and the resulting balance between the two sets. Ethnicity was self-reported (UK Biobank Field 21000) and grouped into the seven categories in Table S2. It served as a split acceptance criterion and as an adjustment covariate in the missingness sensitivity analysis.

The stratification label combined sex with the six binary outcome indicators: 5-year, 10-year and all-incident outcomes for both all-cause dementia and Alzheimer's disease. For strata with fewer than 10 participants, we omitted sex and merged those participants into the corresponding sex-independent outcome pattern. We then merged outcome patterns with fewer than two participants into a residual stratum.

We prespecified three candidate random seeds and four balance criteria, then evaluated the seeds in order and retained the first that met all four. Two criteria concerned the outcomes: at least one case of each outcome in the holdout test set and a holdout share of cases for each outcome within five percentage points of the target 10%. Because we stratified the split on outcome, both criteria only confirmed that the stratification worked. The two substantive criteria concerned covariates outside the stratification scheme: a standardized mean difference (SMD) for age below 0.10 and

every assessment-center and ethnicity category within two percentage points between the development and holdout test sets. The first candidate seed met all four criteria, so we never evaluated the other two.

Table S1 reports outcome counts in the development and holdout test sets. Every outcome had cases in the holdout test set, and for every outcome the case proportion in the holdout test set differed from that in the development set by less than 0.01 percentage points, so the holdout share of cases was close to the 10% target and both outcome criteria were met.

Table S1 Outcome counts and prevalences in the development and holdout test sets.

	Development (n = 450,369)		Holdout test (n = 50,041)	
Outcome	Cases	%	Cases	%
Dementia, 5-year	681	0.151	75	0.150
Dementia, 10-year	3,802	0.844	421	0.841
Dementia, all-incident	10,242	2.274	1,138	2.274
Alzheimer's disease, 5-year	231	0.051	26	0.052
Alzheimer's disease, 10-year	1,506	0.334	167	0.334
Alzheimer's disease, all-incident	4,609	1.023	513	1.025

Table S2 reports the full balance summary. Sex, the only non-outcome variable in the stratification scheme, differed by 0.002 percentage points between the development and holdout test sets. Three further variables served as acceptance criteria without entering the stratification scheme, and all were similarly well balanced. Age at recruitment had an SMD of 0.0057 (mean, 56.54 versus 56.50 years), well below the 0.10 threshold. Category proportions differed by at most 0.21 percentage points across the 22 assessment centers and 0.06 across the seven ethnicity categories, both within the two-percentage-point margin. Availability of the principal predictors entered neither the stratification scheme nor the acceptance criteria, but was also closely balanced, with the largest difference 0.16 percentage points for *APOE* ε4 genotype. This balance matters because informative missingness could inflate performance estimates in this cohort: the holdout test set closely matched the development set in the availability of measured predictors.

Table S2 Balance of the 90% development and 10% holdout test sets. Percentages are column percentages within each set; *Diff.* is the absolute difference in percentage points. For age at recruitment, a continuous variable, the % cells give the mean age in years and *Diff.* gives the standardized mean difference (SMD). For predictor availability, n and % are the participants with a recorded value. Assessment centers are given as UK Biobank Field 54 codes. Ethnicity categories derive from UK Biobank Field 21000. The split was stratified on the six outcomes (Table S1) and sex. Age, ethnicity and assessment center were prespecified acceptance criteria checked after the split, not stratification variables. Predictor availability was reported but neither stratified on nor subject to a threshold.

	Development		Holdout test		
Variable	n	%	n	%	Diff.
Sex (stratified on)					
Female	245,018	54.404	27,225	54.405	0.002
Male	205,351	45.596	22,816	45.595	0.002
Age at recruitment (acceptance criterion)					
Mean, years	450,369	56.541	50,041	56.495	0.0057
Ethnicity (acceptance criterion)					
Asian or Asian British	8,856	1.966	979	1.956	0.010
Black or Black British	7,235	1.606	775	1.549	0.058
Chinese	1,385	0.308	177	0.354	0.046
Mixed	2,645	0.587	289	0.578	0.010
Other	4,083	0.907	440	0.879	0.027
Unknown/Not stated	2,448	0.544	303	0.606	0.062
White	423,717	94.082	47,078	94.079	0.003
Assessment center (acceptance criterion)					
10003	3,407	0.756	364	0.727	0.029
11001	12,496	2.775	1,354	2.706	0.069
11002	12,529	2.782	1,451	2.900	0.118
11003	16,071	3.568	1,769	3.535	0.033
11004	16,742	3.717	1,851	3.699	0.018
11005	15,421	3.424	1,729	3.455	0.031
11006	17,450	3.875	1,904	3.805	0.070
11007	26,293	5.838	2,928	5.851	0.013
11008	25,304	5.619	2,879	5.753	0.135
11009	33,222	7.377	3,676	7.346	0.031
11010	39,625	8.798	4,421	8.835	0.036
11011	38,489	8.546	4,332	8.657	0.111
11012	11,195	2.486	1,284	2.566	0.080
11013	30,287	6.725	3,445	6.884	0.159
11014	27,331	6.069	2,948	5.891	0.177
11016	29,437	6.536	3,257	6.509	0.028
11017	19,124	4.246	2,102	4.201	0.046
11018	25,877	5.746	2,832	5.659	0.086
11020	24,462	5.432	2,788	5.571	0.140
11021	22,963	5.099	2,448	4.892	0.207
11022	2,064	0.458	210	0.420	0.039
11023	580	0.129	69	0.138	0.009
Predictor availability (reported only)					
APOE ε4 genotype	370,376	82.238	41,073	82.079	0.160
Leukocyte telomere length (Field 22191)	423,613	94.059	47,065	94.053	0.006
Alzheimer's disease PRS (Field 26206)	435,801	96.765	48,476	96.873	0.107
Alanine aminotransferase (Field 30620)	420,799	93.434	46,698	93.319	0.115
Pairs-matching completion time (Field 400)	446,244	99.084	49,577	99.073	0.011

Age, assessment center and ethnicity constrained the split through rejection thresholds rather than stratification. Because the first candidate seed satisfied all three criteria, these thresholds never rejected a seed, and the observed balance reflects the large cohort size rather than anything the design enforced.

Column construction and a priori exclusions

This section describes how we transformed the UK Biobank extract into the candidate feature matrix and which a priori exclusions we applied during preprocessing. It ends before the SHapley Additive exPlanations (SHAP) [1] screening and BorutaSHAP [2] stages, which belong to feature selection rather than preprocessing.

From the raw extract to the feature matrix

The extract contained 32,612 columns for 501,936 participants, whereas the final candidate feature matrix contained 7,977 features. The column count fell, rose above the raw-extract total, then fell again. Four construction steps account for the fall and the rise, and the a priori exclusions below for the final fall. Table S3 reports the column count at each stage.

Baseline instance only.

UK Biobank records repeated assessments of the same field across assessment instances (`_i0` to `_i4`). Instances `_i1`–`_i4` correspond to repeat assessment-center visits attended by only a subset of participants. We retained only the instance-0 column for each field and discarded columns from repeat visits.

Per-field compression and typing.

We grouped the remaining columns by field identifier and, for fields with arrays, by array index, then collapsed each group to its first non-missing value and cast the result to the data type specified in the UK Biobank schema. We made two schema exceptions: Fields 22009 and 23376–23384 are declared numeric but contain delimited strings, so we read them as strings. This step reduced the dataset to 11,926 columns.

Early removal of dementia and diagnosis fields.

Before the expansion we removed the dementia outcome fields and the records that define them: hospital inpatient diagnoses and their first-diagnosis dates; causes of death; and self-reported illness, medication and operation codes. That step removed 586 columns, most of them arrays of first-diagnosis dates, and left 11,340.

One-hot expansion of multiple-choice fields.

UK Biobank multiple-choice fields store all selected responses in a single cell, separated by vertical bars (for example, `1|3|-7`). We expanded each field into one binary indicator for every observed response code. We named each indicator by appending a double underscore and the raw response code to the parent column name; for example, response -7 for Field 6138 (qualifications) became `p6138_i0__-7`. For participants

with a missing parent field, we set all corresponding indicators to missing rather than zero, keeping "not asked" distinct from "not selected".

The 148 multiple-choice fields became 26,501 indicator columns, raising the total from 11,340 to 37,693, more than in the raw extract. Much of this increase arose from a few high-cardinality hospital-record fields, particularly operative procedure codes (Fields 41272 and 41273) and hospital admission and discharge administration fields (Fields 41245–41251). We removed these parent fields, and with them every indicator they had generated, at the a priori exclusion stage that follows. Most of the 28,057 columns lost there were therefore indicators of those excluded parent fields, not distinct measurements. The expansion also generated indicators for the "Do not know" (−1) and "Prefer not to answer" (−3) response codes; the missing-value subsection below explains how we handled them.

Table S3 Column count at each stage of construction and exclusion. Counts are columns, including the participant identifier.

Stage	Operation	Columns
Raw extract	—	32,612
Compression	Baseline instance only; per-field collapse and typing	11,926
Outcome removal	Dementia- and outcome-linked fields (−586)	11,340
Expansion	148 multiple-choice fields *replaced by* 26,501 indicators	37,693
Preprocessing	A priori exclusions (−28,057)	9,636
Final selection	Remaining a priori exclusions and derived features	7,978
Candidate features (excluding the participant identifier)		**7,977**

Exclusion groups

Features were excluded a priori, before model fitting and therefore before feature selection. Table S4 lists the 3,452 excluded entries by group; each group is described below.

Outcome and temporal leakage.

The largest exclusion group included the dementia and Alzheimer's disease outcome fields; the corresponding "source of report" fields published alongside first-occurrence and algorithmically defined outcomes; hospital inpatient diagnosis and operative-procedure records; death-registry fields recording date, age and cause of death; primary-care record-availability flags; the first-occurrence date for each recorded ICD-10 code (Fields 130000–132605); and the dates of all algorithmically defined outcomes (Fields 42000–42036). We also excluded self-reported non-cancer illness codes because they directly encode diagnoses.

Table S4 A priori feature exclusions by group. Counts are distinct excluded entries, deduplicated across pipeline stages, and are a lower bound on the number of columns removed.

Group	Entries	Representative fields
Outcome or temporal leakage	2,692	42018, 42020, 130000–132605, 41272
Diet: single-day 24-hour recall	8	20077, 26068, 105010
Study administration and timing	243	21711–21771, 20751, 32365–32368
COVID-19	201	27981–28174, 29156–29161
Employment history	81	22601, 22602, 22603
Assay and genotyping quality control	72	22000–22030, 32050–32065, 3061
Self-reported mental health	44	20544, 29000
Environmental exposure	43	24003–24024, 24500–24508, 21100–21105
Bulk-data pointers	37	21011–21018, 23145–23179
Derived accelerometry	16	40030–40045
Redundant with a retained field	8	21022, 22190, 22192, 22194
Subjective health	4	2178, 4548, 20459, 120103
Not otherwise classified	3	2345, 22501, 4144
Total	**3,452**	

Measurement process rather than measurement.

We excluded variables that described the measurement process rather than participant characteristics. These included assay acquisition-route flags; sample dilution factors; genotyping batch and quality-control fields; whole-genome sequencing batch and quality metrics; and spirometry blow-acceptability codes. We excluded bulk-data fields for the same reason: they point to image or variant files rather than contain measured values. We also excluded assessment-center timestamps; fasting time before blood sampling; consent and authorization flags; and calendar artifacts such as month of birth and month of assessment.

Groups excluded on substantive grounds.

We excluded COVID-19 antibody, serology and exposure-questionnaire fields because they postdate the baseline assessment and were unavailable as baseline predictors. We excluded environmental exposure fields (including air pollution; greenspace and natural-environment coverage; water hardness; and residential location history) because they derive from place of residence and may change during follow-up, making a single baseline value an unstable measure of longer-term exposure. We excluded subjective health items (overall health rating; health satisfaction; general happiness with one's health; and current health rating) because self-ratings may partly reflect undiagnosed illness.

Redundant with a retained field.

Where multiple fields represented the same underlying quantity, we retained a single preferred measure. Recorded age at recruitment is available only in whole years, so we replaced it with age derived from the assessment date and the participant's

month and year of birth. Among the telomere-length measures, we retained the value adjusted for technical variation and excluded the alternatives: the unadjusted measure, which retains assay-related technical variation; and the Z-standardized measure, whose standardization against the full released cohort could transmit population-level information across the development–holdout split.

The candidate feature set

After these exclusions, 7,977 candidate features remained. This matrix entered feature selection, and every subsequent model we developed used only features from it. Table S5 summarizes its composition.

The 7,977 columns came from 3,838 distinct UK Biobank fields. Most fields contributed a single column. A further 112 multiple-choice fields contributed 1,119 binary indicators, the remainder of the 26,501 created during preprocessing after we excluded outcome-linked parent fields. Each field appeared either as a direct column or as a set of binary indicators, never both.

We constructed four additional columns for this study instead of taking them directly from UK Biobank fields: exact age at baseline assessment, calculated from birth and assessment dates; *APOE* $\varepsilon 4$ allele copy number, derived from genotype; and one combined parental status-or-age variable for each parent. For the last of these, we combined parental vital status and age at death into a single variable, using a sentinel value for living parents and the age at death for deceased ones (see the "Feature sets and related data fields" subsection).

Table S5 Composition of the 7,977-column candidate feature set that remained after the a priori exclusions and entered feature selection.

Origin	Columns	Parent fields
UK Biobank field, one column each	6,854	3,726
Multiple-choice one-hot indicators	1,119	112
Derived for this work	4	—
Total	**7,977**	**3,838**

Additional file 2

Additional file 2 (`column_status_supplementary.xlsx`) lists the 7,977 candidate columns and the 3,452 entries excluded a priori, one per row, each labeled `retained` or `excluded`.

Missing-value handling

Two preprocessing steps introduced missing values before imputation. Both recoded values that UK Biobank stores as ordinary numeric entries but that do not represent substantive measurements or responses.

Coded non-answers

UK Biobank records responses such as "Do not know" and "Prefer not to answer" within each field's data coding, most commonly as -1 and -3. A gradient-boosting model treats these sentinel codes as ordinary numeric values. Unless a split isolates them explicitly, the model may group them with nearby genuine low-end values, such as "0, No difficulty", distorting the fitted relationship. We therefore recoded these non-response values as NaN before imputation.

We identified non-responses from the descriptive text of each UK Biobank data coding rather than from the sign of the coded value. Sign alone fails because valid responses may also use negative codes, including -10 for "Less than a year" and -7 for "None of the above"; in coding 1018, all values are negative and only -818 denotes a non-response. This approach therefore preserved valid negative-valued responses while identifying true non-answers.

This procedure affected 1,336 variables and recoded 3,994,257 values as NaN. The one-hot expansion also produced 107 indicators for non-response categories. We set these indicators to NaN instead of dropping the columns, keeping non-response status out of the model. We then filled all resulting NaN values using the random-imputation procedure described in the main text.

Pairs-matching completion times

Field 400 records the time to complete each round of the pairs-matching test, where 0 means "Test not completed" rather than an instantaneous response [3]. Left as recorded, these zeros would signal the fastest possible completion times, reversing the meaning of a predictor that ranks among the most important features in ITU-DESC. We therefore recoded zeros as NaN in `p400_i0_a1` and `p400_i0_a2`. We then filled these values with the same random-imputation procedure used for other missing entries.

Feature sets and related data fields

The dietary feature "Eats eggs, dairy, wheat, or sugar" identifies participants who reported eating all of these food groups. It comes from Field 6144 [4], where response option 5 reads "I eat all of the above". Although the parent field is titled "Never eat eggs, dairy, wheat, sugar", this indicator means the opposite.

The feature "Types of physical activity in the last 4 weeks" is derived from data coding 100316 [5]. Response code 2 corresponds to "Other exercises (e.g. swimming, cycling, keep fit, bowling)", and this binary indicator was retained in the final feature set.

"Length of mobile phone use" is derived from Field 1110 [6], which asks how many years the participant has used a mobile phone at least once per week to make or receive calls.

"Mother's vital status or age at death" combines Field 1835, "Mother still alive", with Field 3526, "Mother's age at death". We assigned the mother's age at death to participants whose mother had died, and a sentinel value below the observed range to those whose mother was still alive. We used a sentinel rather than an imputed age because age at death is undefined for a living parent, not merely unobserved. Imputing such a value would imply a death that had not occurred. A sentinel below the observed range lets tree-based models separate living from deceased parents while preserving the ordering of observed ages at death. We built "Father's vital status or age at death" the same way, from Field 1797, "Father still alive", and Field 1807, "Father's age at death". This feature entered the candidate set, but only two of the ten random-imputation seeds selected it, so it stayed out of the final ITU-DESC feature set.

"Number of vehicles in household" is derived from Field 728 and enters the model as an ordinal count.

Software and implementation

All analyses used Python 3.12.2 with LightGBM 4.6.0 [7], CatBoost 1.2.8 [8], XGBoost 2.1.4 [9], BorutaSHAP 1.0.17, shap 0.48.0, scikit-learn 1.5.2, pandas 2.3.2, NumPy 2.1.3, SciPy 1.15.1, statsmodels 0.14.5, lifelines 0.30.3 and Matplotlib 3.10.5.

Table S6 reports the full training configuration of the two gradient-boosting learners: CatBoost, used for the final ITU-DESC model; and XGBoost, used for the stage-one SHAP screen and as the BorutaSHAP base model. Apart from XGBoost's native categorical support and its automatic tree-method selection, all model-forming hyperparameters used the library defaults. Everything else we set was a runtime option or the evaluation metric, and because neither learner used early stopping or a separate evaluation set, the evaluation metric affected only progress reporting.

The learning rate in Table S6 is the value CatBoost chose for the fit on the full development set; it recalculates the rate for each cross-fitting fold. Neither model used class weighting, so the approximately 2.3% event prevalence entered training unchanged. CatBoost's missing-value mode, *minimum*, routes missing values as if they were smaller than any observed value. Random imputation removes missing values before fitting, so this mode governs only the native missing-value handling in the sensitivity analysis under "Feature importance without random imputation" in Supplementary Note.

Table S6 Training configuration of the gradient-boosting models. *Set* marks a value we passed to the constructor; *supplied* marks one the library chose, and those follow from the versions listed above. Random seeds were held fixed across runs. MVS, minimal variance sampling.

Library	Parameter	Value	Source
CatBoost	Evaluation metric	PRAUC	set

Continued on next page

Library	Parameter	Value	Source
	Task type	CPU	set
	Threads	all available	set
	Verbosity	silent, no on-disk logs	set
	Random seed	fixed	set
	Loss function	Logloss	supplied
	Iterations	1,000	supplied
	Learning rate	0.140, chosen automatically	supplied
	Tree depth	6	supplied
	L2 leaf regularization	3	supplied
	Grow policy	symmetric tree	supplied
	Boosting type	plain	supplied
	Bootstrap type	MVS, subsample 0.8	supplied
	Feature subsampling per tree	1	supplied
	Split-border count	254	supplied
	Minimum data in leaf	1	supplied
	Random strength	1	supplied
	Leaf estimation	Newton, 10 iterations	supplied
	Score function	cosine	supplied
	Missing-value mode	minimum	supplied
	Class weighting	none	supplied
XGBoost	Evaluation metric	aucpr	set
	Tree method	`auto`, resolved to histogram	set
	Categorical feature support	enabled	set
	Threads	all available	set
	Random seed	fixed	set
	Objective	`binary:logistic`	supplied
	Learning rate (η)	0.3	supplied
	Maximum depth	6	supplied
	Minimum child weight	1	supplied
	Row subsampling	1	supplied
	Column subsampling	1 per tree, level and node	supplied
	L2 and L1 regularization	1 and 0	supplied
	Minimum split loss (γ)	0	supplied
	Maximum delta step	0	supplied
	Grow policy	depthwise	supplied
	Maximum bins	256	supplied
	Positive-class weight	1	supplied

Supplementary Note

Baseline characteristics and univariate associations

Table S7 summarizes the baseline characteristics of the 38 selected predictors by incident dementia and Alzheimer's disease status. Table S8 extends it: for each selected

predictor it adds the UK Biobank column identifier, the number of participants with a recorded value in each group and the complete response distribution of time spent driving. Table S9 reports the corresponding univariate odds ratios for incident all-cause dementia and Alzheimer's disease.

Table S7 Baseline characteristics of selected predictors by incident dementia and Alzheimer's disease status.

Variable	Overall (n=500,410)	No incident dementia (n=489,030)	Dementia (n=11,380)	AD (n=5,122)	SMD (dementia)	SMD (AD)
Age, y	58.3 [50.6, 63.7]	58.1 [50.4, 63.5]	65.9 [62.6, 68.2]	66.1 [63.1, 68.4]	1.19	1.29
Standard PRS for Alzheimer's disease (AD)	-0.09 [-0.65, 0.63]	-0.10 [-0.66, 0.61]	0.53 [-0.25, 1.49]	0.86 [-0.05, 1.72]	0.56	0.76
Long-standing illness, disability or infirmity (yes)	159,199 (32.7%)	153,480 (32.2%)	5,719 (52.3%)	2,269 (46.1%)	0.41	0.29
Number of treatments/ medications taken	2 [0, 4]	2 [0, 4]	4 [2, 6]	3 [1, 6]	0.56	0.47
Alanine aminotransferase, U/L	20.1 [15.4, 27.4]	20.1 [15.4, 27.4]	19.7 [15.4, 26.1]	19.4 [15.2, 25.8]	-0.06	-0.08
Touchscreen duration, min	41.8 [33.8, 52.2]	41.6 [33.7, 52.0]	47.7 [38.7, 60.4]	48.5 [39.5, 61.2]	0.41	0.46
Aspartate aminotransferase, U/L	24.4 [21.0, 28.8]	24.4 [21.0, 28.8]	25.3 [21.9, 29.7]	25.5 [22.1, 29.9]	0.11	0.11
Mother still alive (yes)	197,356 (40.2%)	195,950 (40.8%)	1,406 (12.7%)	549 (11.0%)	-0.67	-0.73
Mother's age at death, y	76 [67, 84]	76 [67, 84]	78 [69, 84]	78 [69, 84]	0.10	0.12
APOE ε4					0.58	0.77
Non-carrier	294,093 (71.5%)	289,837 (72.0%)	4,256 (46.6%)	1,573 (38.1%)		
One ε4 allele	107,585 (26.1%)	103,686 (25.8%)	3,899 (42.7%)	1,984 (48.0%)		
Two ε4 alleles	9,771 (2.4%)	8,786 (2.2%)	985 (10.8%)	574 (13.9%)		
Basophil count, 10^9 cells/L	0.02 [0.00, 0.04]	0.02 [0.00, 0.04]	0.02 [0.00, 0.04]	0.02 [0.00, 0.04]	-0.02	-0.03
Mean time to correctly identify matches, ms	539 [481, 613]	536 [480, 610]	586 [520, 676]	583 [519, 668]	0.43	0.41
Hand grip strength (left), kg	28 [21, 38]	28 [21, 38]	26 [18, 34]	24 [18, 34]	-0.26	-0.30
Time spent watching television (TV), h/day	3 [2, 4]	3 [2, 4]	3 [2, 4]	3 [2, 4]	0.29	0.28
Pairs-matching completion time, round 1, s	8.5 [7.0, 12.6]	8.5 [7.0, 12.5]	11.2 [8.0, 19.0]	11.5 [8.1, 19.2]	0.34	0.36
Saturated Fatty Acids to Total Fatty Acids percentage	33.9 [32.8, 35.2]	33.9 [32.8, 35.2]	34.1 [32.9, 35.4]	34.1 [32.9, 35.3]	0.09	0.08
Pairs-matching completion time, round 2, s	28.0 [21.7, 36.7]	27.9 [21.7, 36.5]	33.6 [25.9, 44.9]	34.1 [26.2, 45.5]	0.41	0.44

continued on next page

Table S7 (continued)

Variable	Overall (n=500,410)	No incident dementia (n=489,030)	Dementia (n=11,380)	AD (n=5,122)	SMD (dementia)	SMD (AD)
Other exercises (swimming, cycling, keep fit, bowling) (yes)	235,146 (47.6%)	231,266 (47.9%)	3,880 (35.1%)	1,845 (36.8%)	-0.26	-0.23
Falls in the last year					0.25	0.20
No falls	399,102 (80.2%)	391,173 (80.4%)	7,929 (70.4%)	3,656 (72.0%)		
Only one fall	65,699 (13.2%)	63,747 (13.1%)	1,952 (17.3%)	891 (17.6%)		
More than one fall	32,967 (6.6%)	31,579 (6.5%)	1,388 (12.3%)	528 (10.4%)		
Sleeplessness / insomnia					0.06	0.01
Never/rarely	120,184 (24.1%)	117,488 (24.1%)	2,696 (23.8%)	1,253 (24.6%)		
Sometimes	237,911 (47.7%)	232,765 (47.7%)	5,146 (45.5%)	2,414 (47.4%)		
Usually	140,831 (28.2%)	137,360 (28.2%)	3,471 (30.7%)	1,428 (28.0%)		
Mood swings (yes)	221,547 (45.6%)	216,371 (45.5%)	5,176 (47.5%)	2,269 (46.3%)	0.04	0.01
Eats eggs, dairy, wheat, sugar (yes)	380,897 (76.8%)	373,697 (77.1%)	7,200 (64.4%)	3,170 (62.9%)	-0.28	-0.32
Worry too long after embarrassment (yes)	228,141 (47.7%)	223,642 (47.8%)	4,499 (42.3%)	1,958 (41.0%)	-0.11	-0.14
IGF-1, nmol/L	21.2 [17.5, 24.8]	21.3 [17.5, 24.8]	20.1 [16.2, 23.7]	19.9 [16.2, 23.4]	-0.21	-0.25
Standard PRS for Parkinson's disease (PD)	-0.15 [-0.83, 0.54]	-0.15 [-0.83, 0.54]	-0.10 [-0.77, 0.59]	-0.11 [-0.79, 0.58]	0.05	0.04
Water intake, glasses/day	2 [1, 4]	2 [1, 4]	2 [1, 4]	2 [1, 4]	-0.01	0.03
Adjusted T/S ratio	0.82 [0.74, 0.91]	0.82 [0.74, 0.91]	0.79 [0.71, 0.87]	0.79 [0.71, 0.87]	-0.27	-0.29
C-reactive protein, mg/L	1.33 [0.66, 2.76]	1.33 [0.66, 2.76]	1.42 [0.69, 3.00]	1.31 [0.65, 2.83]	0.06	0.01
Number of vehicles in household					0.37	0.36
None	43,952 (8.9%)	42,318 (8.7%)	1,634 (14.7%)	639 (12.7%)		
One	207,693 (41.9%)	201,893 (41.7%)	5,800 (52.0%)	2,683 (53.3%)		
Two	188,284 (38.0%)	185,116 (38.2%)	3,168 (28.4%)	1,474 (29.3%)		
Three	42,510 (8.6%)	42,082 (8.7%)	428 (3.8%)	182 (3.6%)		
Four or more	13,375 (2.7%)	13,252 (2.7%)	123 (1.1%)	52 (1.0%)		

continued on next page

Table S7 (continued)

Variable	Overall (n=500,410)	No incident dementia (n=489,030)	Dementia (n=11,380)	AD (n=5,122)	SMD (dementia)	SMD (AD)
Peak expiratory flow (PEF), L/min	365 [287, 460]	366 [288, 462]	318 [232, 412]	313 [231, 406]	-0.37	-0.40
Taking other prescription medications (yes)	233,596 (47.0%)	225,899 (46.5%)	7,697 (68.4%)	3,316 (65.3%)	0.45	0.39
Hearing difficulty/problems with background noise (yes)	183,436 (37.5%)	178,086 (37.3%)	5,350 (48.5%)	2,322 (46.9%)	0.23	0.20
Maternal history of high blood pressure (yes)	144,074 (29.4%)	141,803 (29.6%)	2,271 (20.4%)	970 (19.4%)	-0.21	-0.24
Length of mobile phone use					0.37	0.39
Never used mobile phone at least once per week	73,526 (14.9%)	70,631 (14.7%)	2,895 (26.4%)	1,311 (26.6%)		
One year or less	13,614 (2.8%)	13,108 (2.7%)	506 (4.6%)	248 (5.0%)		
Two to four years	86,879 (17.6%)	84,694 (17.6%)	2,185 (20.0%)	1,034 (21.0%)		
Five to eight years	151,533 (30.8%)	148,758 (30.9%)	2,775 (25.4%)	1,214 (24.7%)		
More than eight years	166,865 (33.9%)	164,278 (34.1%)	2,587 (23.6%)	1,115 (22.6%)		
Creatinine (enzymatic) in urine, μmol/L	7544.0 [4377.0, 12059.0]	7542.0 [4373.0, 12060.0]	7662.0 [4551.0, 12027.0]	7212.5 [4334.8, 11493.5]	0.02	-0.06
Neutrophil percentage	61.1 [55.6, 66.5]	61.1 [55.6, 66.4]	62.5 [56.9, 67.8]	62.2 [56.6, 67.4]	0.15	0.11
Vitamin D, nmol/L	46.8 [32.4, 62.4]	46.8 [32.4, 62.4]	46.7 [32.0, 62.5]	47.6 [32.8, 63.8]	0.00	0.04
Maternal history of Alzheimer's disease/dementia (yes)	38,220 (7.8%)	36,602 (7.6%)	1,618 (14.6%)	858 (17.1%)	0.22	0.29

Data are presented as median [Q1, Q3] for continuous variables and n (%) for discrete variables. SMD = standardized mean difference. T/S ratio = telomere-to-single-copy-gene ratio, adjusted for assay-related technical parameters. The SMD (dementia) column compares participants with incident dementia with participants with no incident dementia, and the SMD (AD) column compares participants with incident Alzheimer's disease with participants with no incident dementia; both were computed on observed values. For categorical variables, parent rows report the overall effect size and indented rows show category counts and percentages. Absolute SMD values ≥ 0.1 were considered meaningful. Column headers give the total participants per group; the number of participants with a non-missing value varies by feature and is reported per feature in Table S8; univariate odds ratios are reported in Table S9. Units are given with each variable name. UK Biobank records pairs-matching completion time and touchscreen duration in deciseconds; both are reported here in the units shown. Time spent driving is omitted here because of its many response categories and is reported in full in Table S8. UK Biobank field IDs for each variable are listed in Table S15.

Table S8 Baseline characteristics of selected predictors by incident dementia and Alzheimer's disease status, including every response category and the number of participants with a recorded value.

Variable	Overall	No incident dementia	Dementia	AD	Dementia vs no incident dementia SMD	AD vs no incident dementia SMD
Age, y (Age)	58.3 [50.6, 63.7] (n=500,410)	58.1 [50.4, 63.5] (n=489,030)	65.9 [62.6, 68.2] (n=11,380)	66.1 [63.1, 68.4] (n=5,122)	1.192	1.288
Standard PRS for Alzheimer's disease (AD) (p26206)	-0.09 [-0.65, 0.63] (n=484,277)	-0.10 [-0.66, 0.61] (n=473,386)	0.53 [-0.25, 1.49] (n=10,891)	0.86 [-0.05, 1.72] (n=4,922)	0.563	0.761
Long-standing illness, disability or infirmity (p2188_i0) = 0.0	327,933 (67.3%)	322,719 (67.8%)	5,214 (47.7%)	2,651 (53.9%)	0.415	0.287
Long-standing illness, disability or infirmity (p2188_i0) = 1.0	159,199 (32.7%)	153,480 (32.2%)	5,719 (52.3%)	2,269 (46.1%)		
Number of treatments/medications taken (p137_i0)	2 [0, 4] (n=499,572)	2 [0, 4] (n=488,237)	4 [2, 6] (n=11,335)	3 [1, 6] (n=5,105)	0.561	0.472
Alanine aminotransferase, U/L (p30620_i0)	20.1 [15.4, 27.4] (n=467,497)	20.1 [15.4, 27.4] (n=456,998)	19.7 [15.4, 26.1] (n=10,499)	19.4 [15.2, 25.8] (n=4,759)	-0.061	-0.081
Touchscreen duration, min (p630_i0)	41.8 [33.8, 52.2] (n=499,569)	41.6 [33.7, 52.0] (n=488,221)	47.7 [38.7, 60.4] (n=11,348)	48.5 [39.5, 61.2] (n=5,109)	0.406	0.457
Aspartate aminotransferase, U/L (p30650_i0)	24.4 [21.0, 28.8] (n=465,900)	24.4 [21.0, 28.8] (n=455,433)	25.3 [21.9, 29.7] (n=10,467)	25.5 [22.1, 29.9] (n=4,748)	0.107	0.113
Mother still alive (p1835_i0) = 0.0	293,520 (59.8%)	283,816 (59.2%)	9,704 (87.3%)	4,462 (89.0%)	-0.672	-0.726
Mother still alive (p1835_i0) = 1.0	197,356 (40.2%)	195,950 (40.8%)	1,406 (12.7%)	549 (11.0%)		
Mother's age at death, y (p3526_i0)	76 [67, 84] (n=291,450)	76 [67, 84] (n=281,964)	78 [69, 84] (n=9,486)	78 [69, 84] (n=4,354)	0.097	0.119
APOE ε4 (e4_copies)					0.576	0.768
Non-carrier	294,093 (71.5%)	289,837 (72.0%)	4,256 (46.6%)	1,573 (38.1%)		
One ε4 allele	107,585 (26.1%)	103,686 (25.8%)	3,899 (42.7%)	1,984 (48.0%)		
Two ε4 alleles	9,771 (2.4%)	8,786 (2.2%)	985 (10.8%)	574 (13.9%)		
Basophil count, 10^9 cells/L (p30160_i0)	0.02 [0.00, 0.04] (n=475,347)	0.02 [0.00, 0.04] (n=464,656)	0.02 [0.00, 0.04] (n=10,691)	0.02 [0.00, 0.04] (n=4,855)	-0.017	-0.030
Mean time to correctly identify matches, ms (p20023_i0)	539 [481, 613] (n=494,648)	536 [480, 610] (n=483,578)	586 [520, 676] (n=11,070)	583 [519, 668] (n=4,997)	0.429	0.406
Hand grip strength (left), kg (p46_i0)	28 [21, 38] (n=497,050)	28 [21, 38] (n=485,798)	26 [18, 34] (n=11,252)	24 [18, 34] (n=5,075)	-0.257	-0.300
Time spent watching television (TV), h/day (p1070_i0)	3 [2, 4] (n=495,029)	3 [2, 4] (n=483,898)	3 [2, 4] (n=11,131)	3 [2, 4] (n=5,014)	0.286	0.278
Pairs-matching completion time, round 1, s (p400_i0_a1)	8.5 [7.0, 12.6] (n=481,756)	8.5 [7.0, 12.5] (n=471,453)	11.2 [8.0, 19.0] (n=10,303)	11.5 [8.1, 19.2] (n=4,647)	0.340	0.361
Saturated Fatty Acids to Total Fatty Acids percentage (p23455_i0)	33.9 [32.8, 35.2] (n=486,027)	33.9 [32.8, 35.2] (n=475,053)	34.1 [32.9, 35.4] (n=10,974)	34.1 [32.9, 35.3] (n=4,948)	0.092	0.077
Pairs-matching completion time, round 2, s (p400_i0_a2)	28.0 [21.7, 36.7] (n=480,704)	27.9 [21.7, 36.5] (n=470,479)	33.6 [25.9, 44.9] (n=10,225)	34.1 [26.2, 45.5] (n=4,610)	0.406	0.437

continued on next page

Table S8 (continued)

Variable	Overall	No incident dementia	Dementia	AD	Dementia vs no incident dementia SMD	AD vs no incident dementia SMD
Other exercises (swimming, cycling, keep fit, bowling) (p6164_i0___2) = 0.0	259,005 (52.4%)	251,822 (52.1%)	7,183 (64.9%)	3,168 (63.2%)	-0.262	-0.225
Other exercises (swimming, cycling, keep fit, bowling) (p6164_i0___2) = 1.0	235,146 (47.6%)	231,266 (47.9%)	3,880 (35.1%)	1,845 (36.8%)		
Falls in the last year (p2296_i0)					0.249	0.202
1.0	399,102 (80.2%)	391,173 (80.4%)	7,929 (70.4%)	3,656 (72.0%)		
2.0	65,699 (13.2%)	63,747 (13.1%)	1,952 (17.3%)	891 (17.6%)		
3.0	32,967 (6.6%)	31,579 (6.5%)	1,388 (12.3%)	528 (10.4%)		
Sleeplessness / insomnia (p1200_i0)					0.057	0.012
1.0	120,184 (24.1%)	117,488 (24.1%)	2,696 (23.8%)	1,253 (24.6%)		
2.0	237,911 (47.7%)	232,765 (47.7%)	5,146 (45.5%)	2,414 (47.4%)		
3.0	140,831 (28.2%)	137,360 (28.2%)	3,471 (30.7%)	1,428 (28.0%)		
Mood swings (p1920_i0) = 0.0	264,633 (54.4%)	258,901 (54.5%)	5,732 (52.5%)	2,636 (53.7%)	0.039	0.015
Mood swings (p1920_i0) = 1.0	221,547 (45.6%)	216,371 (45.5%)	5,176 (47.5%)	2,269 (46.3%)		
Eats eggs, dairy, wheat, sugar (p6144_i0___5) = 0.0	114,880 (23.2%)	110,901 (22.9%)	3,979 (35.6%)	1,872 (37.1%)	-0.282	-0.315
Eats eggs, dairy, wheat, sugar (p6144_i0___5) = 1.0	380,897 (76.8%)	373,697 (77.1%)	7,200 (64.4%)	3,170 (62.9%)		
Worry too long after embarrassment (p2000_i0) = 0.0	250,042 (52.3%)	243,900 (52.2%)	6,142 (57.7%)	2,814 (59.0%)	-0.112	-0.137
Worry too long after embarrassment (p2000_i0) = 1.0	228,141 (47.7%)	223,642 (47.8%)	4,499 (42.3%)	1,958 (41.0%)		
IGF-1, nmol/L (p30770_i0)	21.2 [17.5, 24.8] (n=465,145)	21.3 [17.5, 24.8] (n=454,705)	20.1 [16.2, 23.7] (n=10,440)	19.9 [16.2, 23.4] (n=4,734)	-0.210	-0.246

continued on next page

Table S8 (continued)

Variable	Overall	No incident dementia	Dementia	AD	Dementia vs no incident dementia SMD	AD vs no incident dementia SMD
Time spent driving (p1090_i0)					0.233	0.203
No driving	105,589 (21.5%)	102,165 (21.3%)	3,424 (31.2%)	1,478 (29.8%)		
<1 hour/day	153,079 (31.2%)	150,066 (31.3%)	3,013 (27.5%)	1,420 (28.6%)		
1 hour/day	139,610 (28.5%)	136,785 (28.5%)	2,825 (25.8%)	1,277 (25.7%)		
2 hours/day	60,698 (12.4%)	59,533 (12.4%)	1,165 (10.6%)	538 (10.8%)		
3 hours/day	14,637 (3.0%)	14,367 (3.0%)	270 (2.5%)	122 (2.5%)		
4 hours/day	6,499 (1.3%)	6,392 (1.3%)	107 (1.0%)	44 (0.9%)		
5 hours/day	3,137 (0.6%)	3,084 (0.6%)	53 (0.5%)	26 (0.5%)		
6 hours/day	2,599 (0.5%)	2,551 (0.5%)	48 (0.4%)	26 (0.5%)		
7 hours/day	804 (0.2%)	796 (0.2%)	8 (0.1%)	5 (0.1%)		
8 hours/day	1,728 (0.3%)	1,710 (0.4%)	18 (0.2%)	9 (0.2%)		
9 hours/day	881 (0.2%)	868 (0.2%)	13 (0.1%)	8 (0.2%)		
10 hours/day	793 (0.2%)	779 (0.2%)	14 (0.1%)	6 (0.1%)		
11 hours/day	111 (0.0%)	107 (0.0%)	4 (0.0%)	2 (0.0%)		
12 hours/day	249 (0.1%)	243 (0.1%)	6 (0.1%)	4 (0.1%)		
13 hours/day	24 (0.0%)	24 (0.0%)	0 (0.0%)	0 (0.0%)		
14 hours/day	34 (0.0%)	33 (0.0%)	1 (0.0%)	0 (0.0%)		
15 hours/day	25 (0.0%)	24 (0.0%)	1 (0.0%)	0 (0.0%)		
16 hours/day	6 (0.0%)	6 (0.0%)	0 (0.0%)	0 (0.0%)		
17 hours/day	1 (0.0%)	1 (0.0%)	0 (0.0%)	0 (0.0%)		
18 hours/day	4 (0.0%)	4 (0.0%)	0 (0.0%)	0 (0.0%)		
19 hours/day	1 (0.0%)	1 (0.0%)	0 (0.0%)	0 (0.0%)		
20 hours/day	15 (0.0%)	14 (0.0%)	1 (0.0%)	0 (0.0%)		
24 hours/day	1 (0.0%)	0 (0.0%)	1 (0.0%)	0 (0.0%)		
Standard PRS for Parkinson's disease (PD) (p26260)	-0.15 [-0.83, 0.54] (n=484,277)	-0.15 [-0.83, 0.54] (n=473,386)	-0.10 [-0.77, 0.59] (n=10,891)	-0.11 [-0.79, 0.58] (n=4,922)	0.051	0.038
Water intake, glasses/day (p1528_i0)	2 [1, 4] (n=495,659)	2 [1, 4] (n=484,451)	2 [1, 4] (n=11,208)	2 [1, 4] (n=5,044)	-0.009	0.027
Adjusted T/S ratio (p22191_i0)	0.82 [0.74, 0.91] (n=470,678)	0.82 [0.74, 0.91] (n=460,070)	0.79 [0.71, 0.87] (n=10,608)	0.79 [0.71, 0.87] (n=4,791)	-0.270	-0.289
C-reactive protein, mg/L (p30710_i0)	1.33 [0.66, 2.76] (n=466,666)	1.33 [0.66, 2.76] (n=456,196)	1.42 [0.69, 3.00] (n=10,470)	1.31 [0.65, 2.83] (n=4,751)	0.059	0.011
Number of vehicles in household (p728_i0)					0.371	0.358
1.0	43,952 (8.9%)	42,318 (8.7%)	1,634 (14.7%)	639 (12.7%)		
2.0	207,693 (41.9%)	201,893 (41.7%)	5,800 (52.0%)	2,683 (53.3%)		
3.0	188,284 (38.0%)	185,116 (38.2%)	3,168 (28.4%)	1,474 (29.3%)		
4.0	42,510 (8.6%)	42,082 (8.7%)	428 (3.8%)	182 (3.6%)		
5.0	13,375 (2.7%)	13,252 (2.7%)	123 (1.1%)	52 (1.0%)		

continued on next page

Table S8 (continued)

Variable	Overall	No incident dementia	Dementia	AD	Dementia vs no incident dementia SMD	AD vs no incident dementia SMD
Peak expiratory flow (PEF), L/min (p3064_i0_a1)	365 [287, 460] (n=451,589)	366 [288, 462] (n=441,842)	318 [232, 412] (n=9,747)	313 [231, 406] (n=4,474)	-0.373	-0.397
Taking other prescription medications (p2492_i0) = 0.0	263,774 (53.0%)	260,212 (53.5%)	3,562 (31.6%)	1,759 (34.7%)	0.454	0.387
Taking other prescription medications (p2492_i0) = 1.0	233,596 (47.0%)	225,899 (46.5%)	7,697 (68.4%)	3,316 (65.3%)		
Hearing difficulty/problems with background noise (p2257_i0) = 0.0	305,367 (62.5%)	299,698 (62.7%)	5,669 (51.5%)	2,631 (53.1%)	0.229	0.196
Hearing difficulty/problems with background noise (p2257_i0) = 1.0	183,436 (37.5%)	178,086 (37.3%)	5,350 (48.5%)	2,322 (46.9%)		
Maternal history of high blood pressure (p20110_i0___8) = 0.0	346,801 (70.7%)	337,962 (70.4%)	8,839 (79.6%)	4,041 (80.6%)	-0.212	-0.239
Maternal history of high blood pressure (p20110_i0___8) = 1.0	144,074 (29.4%)	141,803 (29.6%)	2,271 (20.4%)	970 (19.4%)		
Length of mobile phone use (p1110_i0)					0.366	0.394
0.0	73,526 (14.9%)	70,631 (14.7%)	2,895 (26.4%)	1,311 (26.6%)		
1.0	13,614 (2.8%)	13,108 (2.7%)	506 (4.6%)	248 (5.0%)		
2.0	86,879 (17.6%)	84,694 (17.6%)	2,185 (20.0%)	1,034 (21.0%)		
3.0	151,533 (30.8%)	148,758 (30.9%)	2,775 (25.4%)	1,214 (24.7%)		
4.0	166,865 (33.9%)	164,278 (34.1%)	2,587 (23.6%)	1,115 (22.6%)		
Creatinine (enzymatic) in urine, μmol/L (p30510_i0)	7544.0 [4377.0, 12059.0] (n=482,364)	7542.0 [4373.0, 12060.0] (n=471,619)	7662.0 [4551.0, 12027.0] (n=10,745)	7212.5 [4334.8, 11493.5] (n=4,884)	0.022	-0.055
Neutrophil percentage (p30200_i0)	61.1 [55.6, 66.5] (n=475,353)	61.1 [55.6, 66.4] (n=464,662)	62.5 [56.9, 67.8] (n=10,691)	62.2 [56.6, 67.4] (n=4,855)	0.154	0.112
Vitamin D, nmol/L (p30890_i0)	46.8 [32.4, 62.4] (n=446,521)	46.8 [32.4, 62.4] (n=436,559)	46.7 [32.0, 62.5] (n=9,962)	47.6 [32.8, 63.8] (n=4,515)	-0.003	0.041
Maternal history of Alzheimer's disease/dementia (p20110_i0___10) = 0.0	452,655 (92.2%)	443,163 (92.4%)	9,492 (85.4%)	4,153 (82.9%)	0.222	0.291
Maternal history of Alzheimer's disease/dementia (p20110_i0___10) = 1.0	38,220 (7.8%)	36,602 (7.6%)	1,618 (14.6%)	858 (17.1%)		

Data are presented as median [Q1, Q3] for continuous variables and n (%) for discrete variables. For continuous variables, the number of participants with a recorded value in each group is given beneath the summary. SMD = standardized mean difference. For continuous variables, SMD was computed using means and pooled standard deviations. For binary variables, SMD was computed from group proportions. For multi-category variables, one global SMD was computed from the full category distribution. Units are given with each variable name. UK Biobank records pairs-matching completion time and touchscreen duration in deciseconds; both are reported here in the units shown. Absolute SMD values ≥ 0.1 were considered meaningful.

Table S9 Univariate associations of selected predictors with incident dementia and Alzheimer's disease.

	Dementia	AD
Variable	OR (95% CI)	OR (95% CI)
Age, y (Age)	4.62 (4.47-4.77)	5.35 (5.08-5.64)
Standard PRS for Alzheimer's disease (AD) (p26206)	1.70 (1.67-1.73)	1.98 (1.93-2.03)
Long-standing illness, disability or infirmity (p2188_i0)	2.31 (2.22-2.40)	1.77 (1.68-1.88)
Number of treatments/medications taken (p137_i0)	1.55 (1.53-1.57)	1.44 (1.41-1.46)
Alanine aminotransferase, U/L (p30620_i0)	0.93 (0.91-0.95)	0.91 (0.88-0.94)
Touchscreen duration, min (p630_i0)	1.41 (1.39-1.43)	1.41 (1.38-1.44)
Aspartate aminotransferase, U/L (p30650_i0)	1.07 (1.05-1.08)	1.06 (1.04-1.08)
Mother still alive (p1835_i0)	0.21 (0.20-0.22)	0.18 (0.17-0.20)
Mother's age at death, y (p3526_i0)	1.10 (1.08-1.13)	1.13 (1.09-1.17)
APOE ε4 (e4_copies)	2.69 (2.60-2.77)	3.43 (3.28-3.59)
Basophil count, 10^9 cells/L (p30160_i0)	0.98 (0.96-1.00)	0.97 (0.93-1.00)
Mean time to correctly identify matches, ms (p20023_i0)	1.39 (1.37-1.41)	1.36 (1.33-1.38)
Hand grip strength (left), kg (p46_i0)	0.77 (0.75-0.78)	0.73 (0.71-0.76)
Time spent watching television (TV), h/day (p1070_i0)	1.46 (1.43-1.50)	1.42 (1.37-1.47)
Pairs-matching completion time, round 1, s (p400_i0_a1)	1.21 (1.20-1.22)	1.19 (1.17-1.20)
Saturated Fatty Acids to Total Fatty Acids percentage (p23455_i0)	1.10 (1.08-1.12)	1.08 (1.05-1.11)
Pairs-matching completion time, round 2, s (p400_i0_a2)	1.32 (1.30-1.33)	1.28 (1.26-1.30)
Other exercises (swimming, cycling, keep fit, bowling) (p6164_i0___2)	0.59 (0.57-0.61)	0.64 (0.60-0.68)
Falls in the last year (p2296_i0)	1.48 (1.44-1.52)	1.36 (1.31-1.42)
Sleeplessness / insomnia (p1200_i0)	1.05 (1.03-1.08)	0.99 (0.95-1.02)
Mood swings (p1920_i0)	1.08 (1.04-1.12)	1.03 (0.97-1.09)
Eats eggs, dairy, wheat, sugar (p6144_i0___5)	0.54 (0.52-0.56)	0.51 (0.48-0.54)
Worry too long after embarrassment (p2000_i0)	0.80 (0.77-0.83)	0.76 (0.72-0.81)
IGF-1, nmol/L (p30770_i0)	0.80 (0.78-0.82)	0.77 (0.75-0.80)
Time spent driving (p1090_i0)		
Standard PRS for Parkinson's disease (PD) (p26260)	1.05 (1.03-1.07)	1.04 (1.01-1.07)
Water intake, glasses/day (p1528_i0)	0.99 (0.97-1.01)	1.03 (1.00-1.06)
Adjusted T/S ratio (p22191_i0)	0.75 (0.73-0.76)	0.74 (0.71-0.76)
C-reactive protein, mg/L (p30710_i0)	1.05 (1.04-1.07)	1.01 (0.98-1.04)
Number of vehicles in household (p728_i0)	0.65 (0.63-0.66)	0.67 (0.65-0.70)
Peak expiratory flow (PEF), L/min (p3064_i0_a1)	0.69 (0.67-0.70)	0.67 (0.65-0.69)
Taking other prescription medications (p2492_i0)	2.49 (2.39-2.59)	2.14 (2.02-2.27)
Hearing difficulty/problems with background noise (p2257_i0)	1.59 (1.53-1.65)	1.48 (1.39-1.56)
Maternal history of high blood pressure (p20110_i0___8)	0.61 (0.58-0.64)	0.57 (0.54-0.62)
Length of mobile phone use (p1110_i0)	0.78 (0.77-0.79)	0.77 (0.76-0.78)
Creatinine (enzymatic) in urine, μmol/L (p30510_i0)	1.02 (1.00-1.04)	0.94 (0.92-0.97)
Neutrophil percentage (p30200_i0)	1.17 (1.15-1.20)	1.12 (1.09-1.15)
Vitamin D, nmol/L (p30890_i0)	1.00 (0.98-1.02)	1.04 (1.01-1.07)
Maternal history of Alzheimer's disease/dementia (p20110_i0___10)	2.06 (1.96-2.18)	2.48 (2.30-2.67)

OR = odds ratio; CI = confidence interval. Odds ratios were estimated using univariate logistic regression. For continuous variables, predictors were z-standardized before modeling, so odds ratios represent the effect per 1-SD increase. For binary and ordinal-coded variables, odds ratios represent a one-unit increase in the coded variable. Odds ratios for non-ordinal multi-category variables should be interpreted cautiously. Units are given with each variable name.

Uncalibrated discrimination

We repeated both the risk-score comparison (Tables S10 and S11; Fig. S1) and the analysis across time horizons (Table S14; Fig. S2) using raw, uncalibrated holdout-test predictions rather than the isotonic-calibrated predictions reported in the main text. Model ordering held, and ITU-DESC remained the best-performing model. Uncalibrated AP was higher than the calibrated values in the main text (ITU-DESC AP 0.142 versus 0.138 for all-cause dementia and 0.073 versus 0.070 for Alzheimer's disease), consistent with isotonic ties compressing the score distribution. Table S12 reports the calibrated Alzheimer's disease comparison and Table S13 the calibrated results across horizons.

Table S10 Discrimination of ITU-DESC and established dementia risk scores for all-incident all-cause dementia, evaluated once on the holdout test set ($n = 50{,}041$; 1,138 incident dementia cases) using raw, uncalibrated predictions. The upper panel reports average precision and the lower panel reports ROC AUC for the same models. This table repeats the main-text all-cause dementia comparison on the same models, participants and outcome definitions. Differences are ITU-DESC minus the comparator and were calculated from unrounded estimates. ROC AUC estimates, confidence intervals, differences and two-sided P values were derived using DeLong's method. AP confidence intervals were obtained by paired subject-level bootstrap resampling (10,000 resamples; percentile method), whereas AP P values were obtained using two-sided paired permutation tests (2,000 permutations). AP, average precision; CI, confidence interval; ROC AUC, area under the receiver operating characteristic curve.

Model	**AP** (95% CI)	**ΔAP** (95% CI)	P
ITU-DESC (our model)	0.142 (0.127, 0.161)	—	—
UKB-DRP [10]	0.122 (0.109, 0.138)	0.021 (0.007, 0.034)	< 0.05
UKBDRS [11]	0.106 (0.095, 0.120)	0.036 (0.022, 0.051)	< 0.001
UKBDRS (w/o *APOE* ε4)	0.088 (0.079, 0.099)	0.055 (0.040, 0.071)	< 0.001
DRS [12]	0.065 (0.059, 0.071)	0.078 (0.064, 0.095)	< 0.001
CAIDE [13]	0.045 (0.041, 0.051)	0.098 (0.083, 0.115)	< 0.001
CAIDE (w/o *APOE* ε4)	0.035 (0.032, 0.039)	0.108 (0.093, 0.125)	< 0.001
ANU-ADRI [14]	0.033 (0.030, 0.038)	0.110 (0.094, 0.128)	< 0.001

Model	**ROC AUC** (95% CI)	**ΔROC AUC** (95% CI)	P
ITU-DESC (our model)	0.854 (0.844, 0.863)	—	—
UKB-DRP [10]	0.841 (0.831, 0.851)	0.013 (0.007, 0.019)	< 0.001
UKBDRS [11]	0.826 (0.815, 0.837)	0.028 (0.021, 0.035)	< 0.001
UKBDRS (w/o *APOE* ε4)	0.813 (0.802, 0.824)	0.041 (0.033, 0.049)	< 0.001
DRS [12]	0.788 (0.777, 0.800)	0.065 (0.056, 0.074)	< 0.001
CAIDE [13]	0.697 (0.684, 0.710)	0.157 (0.143, 0.170)	< 0.001
CAIDE (w/o *APOE* ε4)	0.648 (0.634, 0.661)	0.206 (0.191, 0.221)	< 0.001
ANU-ADRI [14]	0.591 (0.575, 0.608)	0.262 (0.244, 0.281)	< 0.001

Table S11 Discrimination of ITU-DESC and established dementia risk scores for all-incident Alzheimer's disease, evaluated once on the holdout test set ($n = 50{,}041$; 513 incident Alzheimer's disease cases) using raw, uncalibrated predictions. The upper panel reports average precision and the lower panel reports ROC AUC for the same models. This table repeats the Alzheimer's disease comparison of Table S12 on the same models, participants and outcome definitions. Differences are ITU-DESC minus the comparator and were calculated from unrounded estimates. ROC AUC estimates, confidence intervals, differences and two-sided P values were derived using DeLong's method. AP confidence intervals were obtained by paired subject-level bootstrap resampling (10,000 resamples; percentile method), whereas AP P values were obtained using two-sided paired permutation tests (2,000 permutations). AP, average precision; CI, confidence interval; ROC AUC, area under the receiver operating characteristic curve.

Model	**AP** (95% CI)	**ΔAP** (95% CI)	P
ITU-DESC (our model)	0.073 (0.061, 0.089)	—	—
UKB-DRP [10]	0.058 (0.049, 0.070)	0.015 (0.005, 0.027)	< 0.05
UKBDRS [11]	0.054 (0.045, 0.066)	0.019 (0.007, 0.032)	< 0.01
UKBDRS (w/o *APOE* ε4)	0.039 (0.034, 0.048)	0.033 (0.022, 0.048)	< 0.001
DRS [12]	0.032 (0.028, 0.037)	0.041 (0.030, 0.056)	< 0.001
CAIDE [13]	0.019 (0.017, 0.023)	0.053 (0.042, 0.069)	< 0.001
CAIDE (w/o *APOE* ε4)	0.015 (0.013, 0.018)	0.058 (0.047, 0.074)	< 0.001
ANU-ADRI [14]	0.014 (0.012, 0.018)	0.058 (0.047, 0.074)	< 0.001
Model	**ROC AUC** (95% CI)	**ΔROC AUC** (95% CI)	P
ITU-DESC (our model)	0.863 (0.850, 0.876)	—	—
UKB-DRP [10]	0.844 (0.830, 0.858)	0.019 (0.010, 0.028)	< 0.001
UKBDRS [11]	0.837 (0.823, 0.851)	0.026 (0.016, 0.036)	< 0.001
UKBDRS (w/o *APOE* ε4)	0.817 (0.803, 0.832)	0.046 (0.035, 0.057)	< 0.001
DRS [12]	0.799 (0.784, 0.814)	0.065 (0.052, 0.077)	< 0.001
CAIDE [13]	0.691 (0.671, 0.710)	0.173 (0.152, 0.193)	< 0.001
CAIDE (w/o *APOE* ε4)	0.625 (0.605, 0.645)	0.238 (0.216, 0.261)	< 0.001
ANU-ADRI [14]	0.568 (0.543, 0.593)	0.295 (0.268, 0.322)	< 0.001

Table S12 Discrimination of ITU-DESC and established dementia risk scores for all-incident Alzheimer's disease, evaluated once on the holdout test set (n = 50,041; 513 incident Alzheimer's disease cases). All models were evaluated on the same participants, under identical outcome definitions, using isotonic-calibrated predictions. The upper panel reports average precision and the lower panel reports ROC AUC for the same models. The base model was fitted once using the all-cause dementia label and recalibrated for Alzheimer's disease; every comparator was recalibrated under the same procedure. Differences are ITU-DESC minus the comparator and were calculated from unrounded estimates. ROC AUC estimates, confidence intervals, differences and two-sided P values were derived using DeLong's method. AP confidence intervals were obtained by paired subject-level bootstrap resampling (10,000 resamples; percentile method), whereas AP P values were obtained using two-sided paired permutation tests (2,000 permutations). AP, average precision; CI, confidence interval; ROC AUC, area under the receiver operating characteristic curve.

Model	**AP** (95% CI)	**ΔAP** (95% CI)	P
ITU-DESC (our model)	0.070 (0.059, 0.085)	—	—
UKB-DRP [10]	0.056 (0.047, 0.067)	0.014 (0.005, 0.025)	< 0.01
UKBDRS [11]	0.051 (0.043, 0.062)	0.019 (0.009, 0.030)	< 0.01
UKBDRS (w/o *APOE* ε4)	0.037 (0.032, 0.044)	0.033 (0.023, 0.045)	< 0.001
DRS [12]	0.030 (0.027, 0.035)	0.040 (0.030, 0.053)	< 0.001
CAIDE [13]	0.019 (0.017, 0.022)	0.051 (0.041, 0.064)	< 0.001
CAIDE (w/o *APOE* ε4)	0.014 (0.013, 0.016)	0.056 (0.045, 0.070)	< 0.001
ANU-ADRI [14]	0.013 (0.012, 0.016)	0.057 (0.046, 0.071)	< 0.001
Model	**ROC AUC** (95% CI)	**ΔROC AUC** (95% CI)	P
ITU-DESC (our model)	0.863 (0.850, 0.876)	—	—
UKB-DRP [10]	0.844 (0.829, 0.858)	0.019 (0.010, 0.028)	< 0.001
UKBDRS [11]	0.836 (0.822, 0.851)	0.027 (0.016, 0.037)	< 0.001
UKBDRS (w/o *APOE* ε4)	0.817 (0.802, 0.831)	0.046 (0.035, 0.058)	< 0.001
DRS [12]	0.799 (0.784, 0.813)	0.064 (0.052, 0.077)	< 0.001
CAIDE [13]	0.701 (0.682, 0.720)	0.162 (0.142, 0.181)	< 0.001
CAIDE (w/o *APOE* ε4)	0.632 (0.614, 0.650)	0.231 (0.210, 0.252)	< 0.001
ANU-ADRI [14]	0.566 (0.542, 0.591)	0.296 (0.270, 0.323)	< 0.001

Table S13 Discrimination of ITU-DESC across the 5-year, 10-year and all-incident horizons for all-cause dementia and Alzheimer's disease, evaluated once on the holdout test set (n = 50,041) using isotonic-calibrated predictions. The main-text figure plots these same point estimates. The base model, fitted once, assigned every participant the same raw score at each horizon; the horizons differ only in the outcome label and in the isotonic mapping fitted for each. AP rises with the number of cases accrued by each horizon, because the AP of a non-informative model equals the outcome prevalence. AP, average precision; ROC AUC, area under the receiver operating characteristic curve.

	All-cause dementia		Alzheimer's disease	
Horizon	ROC AUC	AP	ROC AUC	AP
5-year	0.837	0.012	0.868	0.007
10-year	0.850	0.056	0.866	0.030
All-incident	0.853	0.138	0.863	0.070

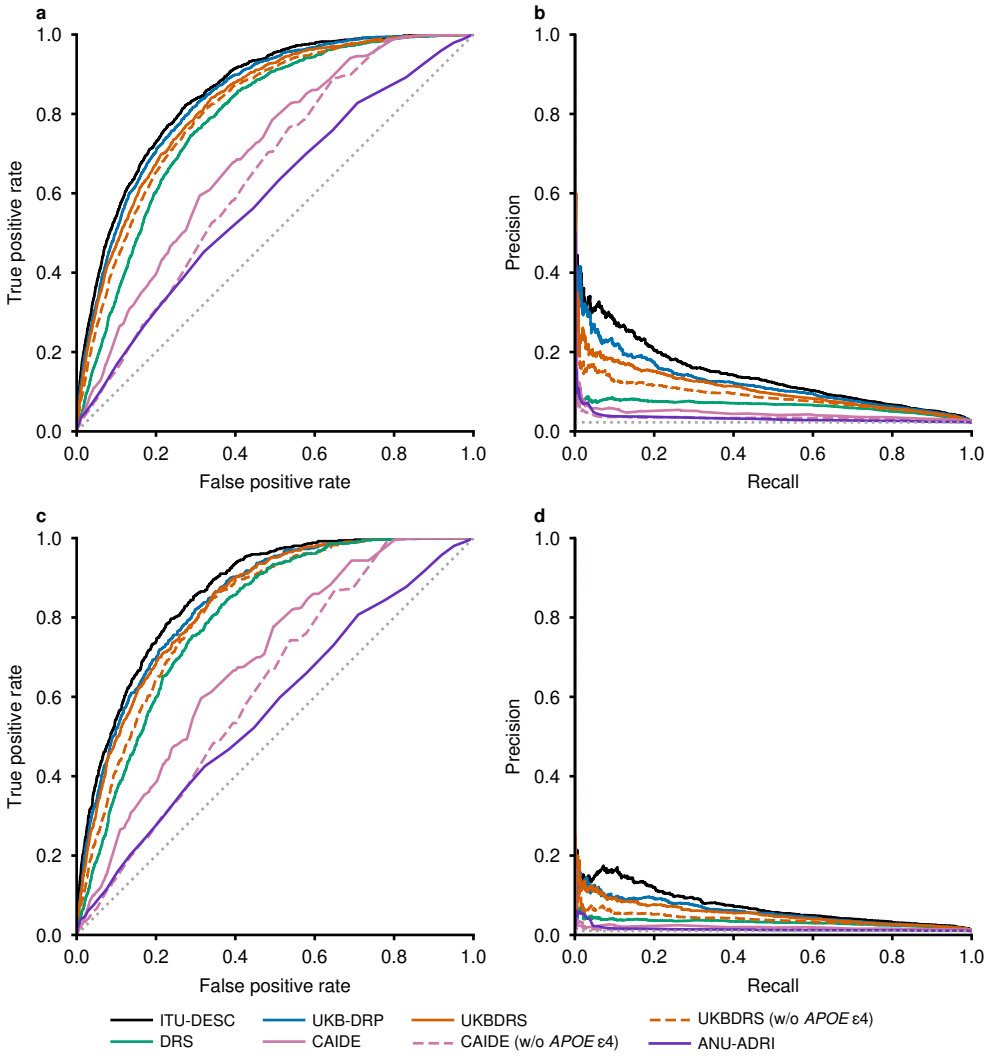

Fig. S1 Uncalibrated discrimination of ITU-DESC and established dementia risk scores, computed from raw holdout-test predictions ($n = 50{,}041$) under the same random-imputation policy for every model. **a**, Receiver operating characteristic curves for all-incident all-cause dementia (1,138 incident cases). **b**, Precision–recall curves for all-incident all-cause dementia. **c**, Receiver operating characteristic curves for all-incident Alzheimer's disease (513 incident cases). **d**, Precision–recall curves for all-incident Alzheimer's disease. The dotted line marks chance discrimination in **a** and **c**, and the outcome prevalence in **b** and **d**. Discrimination estimates and pairwise comparisons are reported in Tables S10 and S11.

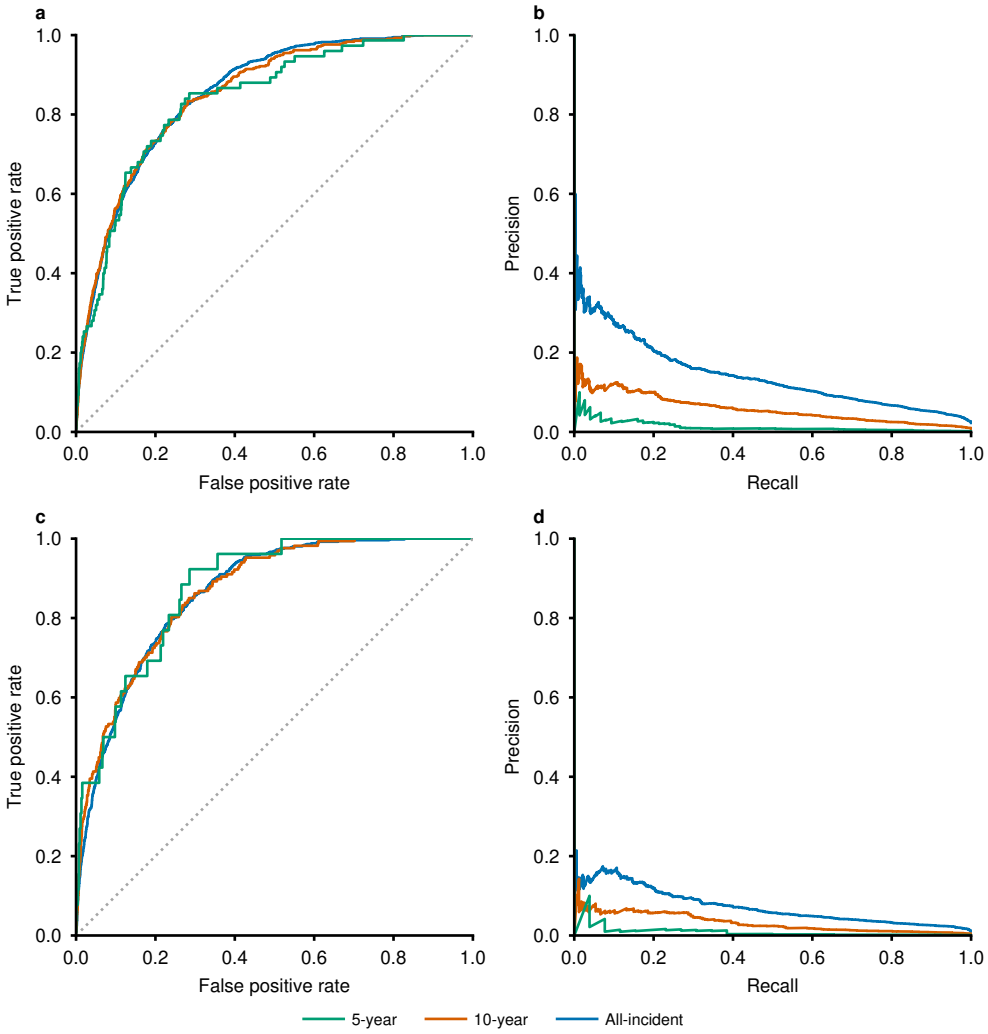

Fig. S2 Uncalibrated discrimination of ITU-DESC across the 5-year, 10-year and all-incident horizons, computed from raw holdout-test predictions (n = 50,041). **a**, Receiver operating characteristic curves for all-cause dementia. **b**, Precision–recall curves for all-cause dementia. **c**, Receiver operating characteristic curves for Alzheimer's disease. **d**, Precision–recall curves for Alzheimer's disease. The same fitted model scored every participant once, so the horizons differ only in the outcome label. The dotted line marks chance discrimination in **a** and **c**. Discrimination values are reported in Table S14.

Table S14 Discrimination of ITU-DESC across the 5-year, 10-year and all-incident horizons for all-cause dementia and Alzheimer's disease, evaluated once on the holdout test set (n = 50,041) using raw, uncalibrated predictions. These same point estimates are plotted in Fig. S2. The base model, fitted once, assigned every participant the same raw score at each horizon, so the horizons differ only in the outcome label. AP rises with the number of cases accrued by each horizon, because the AP of a non-informative model equals the outcome prevalence. AP, average precision; ROC AUC, area under the receiver operating characteristic curve.

	All-cause dementia		Alzheimer's disease	
Horizon	ROC AUC	AP	ROC AUC	AP
5-year	0.839	0.014	0.876	0.011
10-year	0.850	0.058	0.866	0.033
All-incident	0.854	0.142	0.863	0.073

Per-seed feature presence

Because random imputation makes both stages of feature selection stochastic, we repeated the complete pipeline (the five-fold SHAP screen followed by BorutaSHAP) independently for each seed. Each seed therefore produced its own screened feature set and final selection. Individual lists held 40 to 59 features, and their union contained 120 distinct features (Set 1, the union); only 29 features appeared in all ten seeds (Set 10, the intersection). A further five features appeared in nine seeds and four in eight, giving the 38 features of Set 8, which we adopted as the final ITU-DESC feature set. This approximately 20-feature spread shows how much a single wrapper run depends on the imputation draw and justifies the consensus cutoff. For each candidate feature, Table S15 reports the seeds that selected it (columns S1–S10) and the total number of selections, which defines membership in the nested consensus sets. Feature names encode the UK Biobank field identifier with the assessment instance (`i`) and array (`a`) indices, both zero-based. Some features therefore share a field label but differ in their indices. For example, `p400_i0_a1` and `p400_i0_a2` denote Field 400, instance 0, arrays 1 and 2, which record the completion times for the two rounds of the pairs-matching test. A double-underscore suffix (`__`) denotes a binary indicator for a specific response code, so features from the same field can encode different responses: `p20110_i0__8` indicates maternal high blood pressure and `p20110_i0__10` maternal Alzheimer's disease or dementia, both from Field 20110, "Illnesses of mother". Similarly, `p6164_i0__2` indicates response code 2 for Field 6164, "Other exercises (eg: swimming, cycling, keep fit, bowling)". These multiple-choice fields store selected response codes as pipe-delimited strings (for example, "5|4|3"), and we expanded them into separate binary indicators for each observed code.

Table S15 Per-seed presence of each candidate feature across the ten random-imputation seeds. A check mark indicates the feature was selected by BorutaSHAP in that seed's run; *Seeds* gives the total number of seeds (out of ten) in which the feature was selected. Features are ordered by the number of seeds in which they appear. The 38 features selected in at least eight seeds constitute Set 8, adopted as the final ITU-DESC feature set.

Field	Feature	S1	S2	S3	S4	S5	S6	S7	S8	S9	S10	Seeds
Age	Age	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
Maternal_status_or_age	Mother's vital status or age at death	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
e4_copies	*APOE* ε4 copies	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p1070_i0	Time spent watching television (TV)	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p1090_i0	Time spent driving	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p1200_i0	Sleeplessness / insomnia	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p137_i0	Number of treatments/medications taken	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p1528_i0	Water intake	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p1920_i0	Mood swings	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p2000_i0	Worry too long after embarrassment	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p20023_i0	Mean time to correctly identify matches	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p2188_i0	Long-standing illness, disability or infirmity	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p22191_i0	Adjusted T/S ratio	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p2296_i0	Falls in the last year	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p23455_i0	Saturated Fatty Acids to Total Fatty Acids percentage	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p26206	Standard PRS for Alzheimer's disease (AD)	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p26260	Standard PRS for Parkinson's disease (PD)	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p30160_i0	Basophil count	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p30620_i0	Alanine aminotransferase	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p30650_i0	Aspartate aminotransferase	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p30710_i0	C-reactive protein	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p30770_i0	IGF-1	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p400_i0_a1	Pairs-matching completion time, round 1	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p400_i0_a2	Pairs-matching completion time, round 2	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p46_i0	Hand grip strength (left)	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p6144_i0___5	Eats eggs, dairy, wheat, sugar	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p6164_i0___2	Other exercises (swimming, cycling, keep fit, bowling)	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p630_i0	Touchscreen duration	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p728_i0	Number of vehicles in household	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	10
p1110_i0	Length of mobile phone use	✓	✓	✓	✓	✓	✓		✓	✓	✓	9
p20110_i0___8	Maternal history of high blood pressure	✓	✓	✓	✓	✓	✓	✓	✓	✓		9
p2257_i0	Hearing difficulty/problems with background noise		✓	✓	✓	✓	✓	✓	✓	✓	✓	9
p2492_i0	Taking other prescription medications	✓	✓	✓		✓	✓	✓	✓	✓	✓	9
p3064_i0_a1	Peak expiratory flow (PEF)	✓	✓	✓	✓	✓		✓	✓	✓	✓	9
p20110_i0___10	Maternal history of Alzheimer's disease/dementia	✓	✓	✓	✓	✓		✓		✓	✓	8
p30200_i0	Neutrophil percentage	✓	✓	✓	✓	✓			✓	✓	✓	8
p30510_i0	Creatinine (enzymatic) in urine	✓		✓	✓	✓	✓	✓	✓		✓	8
p30890_i0	Vitamin D	✓	✓	✓	✓	✓	✓	✓			✓	8
p1488_i0	Tea intake		✓	✓	✓	✓		✓		✓	✓	7
p23111_i0	Leg fat percentage (right)	✓			✓	✓	✓	✓	✓		✓	7
p23115_i0	Leg fat percentage (left)		✓	✓		✓	✓		✓	✓	✓	7
p23470_i0	Glucose	✓	✓		✓	✓		✓		✓	✓	7
p30670_i0	Urea	✓			✓	✓		✓	✓	✓	✓	7
p23444_i0	Omega-3 Fatty Acids	✓	✓			✓	✓	✓	✓			6
p3077_i0	Seating box height	✓	✓	✓		✓		✓				5

Continued on next page

Table S15 continued from previous page

Field	Feature	S1	S2	S3	S4	S5	S6	S7	S8	S9	S10	Seeds
p738_i0	Average total household income before tax	✓			✓	✓		✓			✓	5
p22009_a2	Genetic principal components			✓	✓	✓				✓		4
p23107_i0	Impedance of leg (right)	✓		✓		✓				✓		4
p23112_i0	Leg fat mass (right)	✓		✓	✓	✓						4
p1060_i0	Time spent outdoors in winter				✓				✓		✓	3
p22009_a11	Genetic principal components		✓	✓	✓							3
p28627	Currently suffering from hearing loss				✓	✓			✓			3
p28739	Difficulty with standing for long periods	✓		✓							✓	3
p3064_i0_a2	Peak expiratory flow (PEF)			✓		✓		✓				3
p47_i0	Hand grip strength (right)					✓		✓		✓		3
p6164_i0___5	Types of physical activity: prefer not to answer			✓	✓						✓	3
Paternal_status_or_age	Father's vital status or age at death									✓	✓	2
p20162_i0	Pack years adult smoking as proportion of life span exposed to smoking				✓						✓	2
p2090_i0	Seen doctor (GP) for nerves, anxiety, tension or depression	✓								✓		2
p22009_a1	Genetic principal components		✓					✓				2
p23116_i0	Leg fat mass (left)		✓							✓		2
p26201_a0	PRS genetic principal components		✓					✓				2
p26227	Standard PRS for coronary artery disease (CAD)			✓							✓	2
p30180_i0	Lymphocyte percentage						✓	✓				2
p30444	Time usually got up in the morning during the past month				✓			✓				2
p30520_i0	Potassium in urine		✓							✓		2
p1508_i0	Coffee type					✓						1
p20133_i0_a2	Time to complete round										✓	1
p20149_i0_a24	Interval between previous point and current one in numeric path (trail #1)	✓										1
p20195_i0	Number of symbol digit matches attempted	✓										1
p20230_i0_a10	Duration to entering symbol choice										✓	1
p20230_i0_a31	Duration to entering symbol choice	✓										1
p20676	Liking for herring		✓									1
p20717	Liking for salty pretzels										✓	1
p2080_i0	Frequency of tiredness / lethargy in last 2 weeks			✓								1
p22009_a21	Genetic principal components										✓	1
p23014_i0	LANA antigen for Kaposi's Sarcoma-Associated Herpesvirus								✓			1
p23100_i0	Whole body fat mass		✓									1
p23451_i0	Omega-3 Fatty Acids to Total Fatty Acids percentage									✓		1
p23459_i0	Omega-6 Fatty Acids to Omega-3 Fatty Acids ratio			✓								1
p23472_i0	Pyruvate										✓	1
p23540_i0	Cholesterol in Medium LDL			✓								1
p2405_i0	Number of children fathered										✓	1
p26201_a1	PRS genetic principal components										✓	1
p26201_a2	PRS genetic principal components	✓										1
p29093	Frequency of consuming six or more units of alcohol		✓									1
p30140_i0	Neutrophil count					✓						1
p30290_i0	High light scatter reticulocyte percentage			✓								1
p30460	Trouble sleeping in past month as disturbed by bed partner		✓									1

Continued on next page

Table S15 continued from previous page

Field	Feature	S1	S2	S3	S4	S5	S6	S7	S8	S9	S10	Seeds
p3062_i0_a2	Forced vital capacity (FVC)										✓	1
p30880_i0	Urate		✓									1
p31147_i0	Adjusted optic disk major axis length (right)			✓								1
p31152_i0	Optic disk orientation (left)						✓					1
p32108	Problems with slow thinking speed in past 7 days							✓				1
p32109	Trouble with figuring things out or solving problems in past 7 days							✓				1
p32110	Number of brothers and sisters				✓							1
p32309	Experiences difficulty in imagining what it would be like to be someone else					✓						1
p399_i0_a2	Number of incorrect matches in round							✓				1
p40008_i0	Age at cancer diagnosis			✓								1
p4196_i0	Pulse wave peak to peak time				✓							1
p4258_i0_a7	Number entered by participant				✓							1
p4290_i0	Duration screen displayed										✓	1
p48_i0	Waist circumference		✓									1
p4947_i0	FI2 : duration viewed	✓										1
p4980_i0	FI5 : duration viewed							✓				1
p5080_i0_a10	Displayed letters in round (right)					✓						1
p5111_i0_a5	3mm asymmetry angle (left)		✓									1
p5113_i0_a2	6mm cylindrical power angle (left)			✓								1
p5115_i0_a2	3mm cylindrical power angle (right)									✓		1
p5118_i0_a0	6mm cylindrical power (left)							✓				1
p5983_i0_a98	ECG, heart rate									✓		1
p5985_i0_a49	Bicycle speed									✓		1
p6164_i0___4	Types of physical activity: none of the above									✓		1
p90013	Standard deviation of acceleration		✓									1
p90038	Average acceleration 11:00 - 11:59								✓			1
p90089	No-wear time bias adjusted acceleration median							✓				1
p90099	Fraction acceleration <= 8 milli-gravities	✓										1
p90109	Fraction acceleration <= 18 milli-gravities										✓	1
p90129	Fraction acceleration <= 150 milli-gravities							✓				1
p90186	Maximum readings exceeding +/-8 gravities after calibration in a 5 second epoch		✓									1

Missingness as a predictive signal

This section quantifies how much information measurement availability carries about subsequent dementia, using the full study cohort, outcome definitions and predictor set of the main analysis.

Methods

We classified a predictor as missing if no value remained in the analysis frame after the recoding under “Missing-value handling” in Supplementary Methods, so random imputation filled the entry. Of the 38 ITU-DESC predictors, age was complete for all

participants. Each of the remaining 37 therefore defined a binary missingness indicator. Two pairs of predictors had identical missingness patterns: the polygenic risk scores and the maternal-history indicators.

In the full study cohort (500,410 participants), we estimated the association between each predictor's missingness and incident all-cause dementia or Alzheimer's disease using odds ratios from unadjusted and adjusted logistic regression models. Adjusted models included age as linear and quadratic terms, sex, education, Townsend deprivation index, recruitment year, assessment center and ethnicity. We modeled education and deprivation as categorical variables with an explicit missing category. We also report the standardized difference in missingness prevalence between cases and non-cases, together with the corresponding chi-square test. Benjamini–Hochberg q-values controlled the false discovery rate across the 37 predictors. We examined outcome incidence by the number of missing predictors per participant.

Results

For every predictor and both outcomes, missingness was more frequent among participants who subsequently developed the outcome than among non-cases (Tables S18 and S19). For all-cause dementia, the standardized difference in missingness ranged from 0.02 to 0.25, with eight predictors at or above 0.10. For Alzheimer's disease, the corresponding range was 0.01 to 0.24, with five predictors at or above 0.10. The largest differences involved the pairs-matching test, peak expiratory flow and questionnaire variables such as length of mobile phone use. No blood assay reached 0.10 for either outcome, and the smallest differences involved touchscreen duration, the maternal-history indicators and several blood assays. After full adjustment, missingness remained associated with all-cause dementia at $q < 0.05$ for 33 of 37 predictors, with adjusted odds ratios ranging from 1.07 to 2.44; 16 of 37 predictors remained associated with Alzheimer's disease. All 37 unadjusted chi-square tests for dementia had $P < 0.05$. Given the large sample size, the standardized differences are more informative than the P-values: missingness was consistently associated with higher subsequent dementia risk, but the differences were modest for most predictors.

Dementia incidence rose from 1.76% among participants with no missing predictors to 3.61% among those missing three or four, then fell to 3.01% among those missing five or more, mostly participants without blood-assay measurements (Table S17). Each additional missing predictor was associated with an adjusted odds ratio of 1.057 for incident dementia (95% confidence interval [CI], 1.050 to 1.065). Participants with complete data for all 38 predictors made up 49.4% of the cohort and were younger, less deprived and more highly educated than participants with at least one missing predictor. They also had lower incidence of both outcomes: 1.76% versus 2.77% for all-cause dementia (Table S16).

Interpretation

Within UK Biobank, missingness remained informative about subsequent outcomes after adjusting for measured demographic and socioeconomic characteristics. These analyses cannot distinguish missingness that is conditionally random, given unmeasured covariates, from missingness that is not at random, nor can they establish

whether the observed associations would transport to other settings. The strongest signals arose from non-completion of cognitive testing and questionnaire non-response, which may reflect participant health status. Assay availability, more closely tied to assessment procedures, showed weaker associations. Because we cannot reliably separate these sources in advance, we withheld both from the model. The binary outcome definitions also classify participants who died before a dementia diagnosis as non-cases; the "Competing risk of death" subsection below quantifies this competing event and shows that it leaves score ranking and calibration by risk decile unchanged.

Table S16 Participants with a recorded value for all 38 ITU-DESC predictors compared with participants missing at least one, the study cohort (500,410 participants). Means are shown with standard deviations and categories as column percentages. *Difference* is the absolute difference in the units shown (percentage points for categories). *SMD* is the standardized mean difference, complete minus incomplete, so a negative value means the complete group has the lower value; for categories it is the standardized difference of two proportions. *P*-values are from Welch's *t*-test for means and Pearson's chi-square test for categories.

Characteristic	Complete (n = 247,200)	Incomplete (n = 253,210)	Difference	SMD	*P*
Age at recruitment, years	56.50 (8.10)	57.64 (8.04)	1.14	−0.142	<0.001
Female, %	54.19	54.61	0.42	−0.008	0.003
Townsend deprivation index	−1.56 (2.94)	−1.04 (3.22)	0.52	−0.168	<0.001
Education, years	14.72 (4.29)	14.03 (4.28)	0.70	0.163	<0.001
Education band, %					
Fewer than 7 years	0.08	0.21	0.13	−0.034	<0.001
7 to 9 years	0.90	2.07	1.17	−0.097	<0.001
10 years or more	98.95	96.32	2.63	0.174	<0.001
Not recorded	0.06	1.40	1.33	−0.157	<0.001
Ethnicity, %					
White	95.47	92.73	2.74	0.116	<0.001
Not recorded	0.28	0.81	0.53	−0.071	<0.001
Incident all-cause dementia, %					
All-incident	1.76	2.77	1.01	−0.068	<0.001
10-year	0.62	1.07	0.45	−0.049	<0.001
5-year	0.09	0.21	0.12	−0.030	<0.001
Incident Alzheimer's disease, %					
All-incident	0.82	1.23	0.41	−0.041	<0.001
10-year	0.25	0.42	0.17	−0.029	<0.001
5-year	0.03	0.07	0.04	−0.019	<0.001

Table S17 Incidence by number of missing ITU-DESC predictors in the study cohort (500,410 participants; all-incident horizon). Incidence intervals are Wilson 95% confidence intervals. Odds ratios are relative to participants with no missing predictor, from a logistic model adjusted for age (linear and quadratic), sex, education, Townsend deprivation, recruitment year, assessment center and ethnicity. The per-predictor row enters the number of missing predictors as a linear term in the same logistic model.

Missing predictors	n	Cases	Incidence, % (95% CI)	Adjusted OR (95% CI)	P
All-cause dementia					
0	247,200	4,361	1.76 (1.71, 1.82)	1 (reference)	–
1	138,901	3,300	2.38 (2.30, 2.46)	1.18 (1.12, 1.23)	<0.001
2	45,172	1,465	3.24 (3.08, 3.41)	1.47 (1.38, 1.56)	<0.001
3–4	29,081	1,049	3.61 (3.40, 3.83)	1.61 (1.50, 1.73)	<0.001
≥5	40,056	1,205	3.01 (2.85, 3.18)	1.50 (1.40, 1.61)	<0.001
Per additional missing predictor, OR				1.057 (1.050, 1.065)	<0.001
Alzheimer's disease					
0	247,200	2,015	0.82 (0.78, 0.85)	1 (reference)	–
1	138,901	1,484	1.07 (1.02, 1.12)	1.13 (1.06, 1.21)	<0.001
2	45,172	670	1.48 (1.38, 1.60)	1.41 (1.29, 1.54)	<0.001
3–4	29,081	463	1.59 (1.45, 1.74)	1.51 (1.36, 1.67)	<0.001
≥5	40,056	490	1.22 (1.12, 1.34)	1.30 (1.17, 1.45)	<0.001
Per additional missing predictor, OR				1.042 (1.030, 1.053)	<0.001

Table S18 Missingness of each ITU-DESC predictor and incident all-cause dementia in the study cohort (500,410 participants; all-incident horizon). *Missing, %* is the share of participants with no recorded value after the recoding described under "Missing-value handling" in Supplementary Methods, overall and separately among participants who did and did not develop all-cause dementia. *SMD* is the standardized difference in the missing rate between cases and non-cases; 8 of 37 predictors reach 0.10 and 37 of 37 unadjusted chi-square tests have $P < 0.05$. *Crude OR* is the unadjusted odds ratio of all-cause dementia for missing versus recorded. *Adjusted OR* is from a logistic model adjusted for age (linear and quadratic), sex, education, Townsend deprivation, recruitment year, assessment center and ethnicity; 33 of 37 fitted predictors keep a Benjamini–Hochberg $q < 0.05$. Age was recorded for every participant. Rows are ordered by SMD.

	Missing, %				Crude	Adjusted OR	BH
Predictor	Overall	Cases	Non-cases	SMD	OR	(95% CI)	q
Pairs-matching completion time, round 2	3.94	10.15	3.79	0.252	2.86	2.12 (1.97, 2.28)	<0.001
Pairs-matching completion time, round 1	3.73	9.46	3.59	0.239	2.80	2.06 (1.91, 2.22)	<0.001
Peak expiratory flow (PEF)	9.76	14.35	9.65	0.145	1.57	1.18 (1.12, 1.25)	<0.001
Length of mobile phone use	1.60	3.80	1.55	0.140	2.51	1.82 (1.64, 2.02)	<0.001
Mean time to correctly identify matches	1.15	2.72	1.11	0.117	2.48	1.84 (1.62, 2.09)	<0.001
Mother's vital status or age at death	3.02	5.30	2.96	0.117	1.83	1.47 (1.34, 1.60)	<0.001
Other exercises (swimming, cycling, keep fit, bowling)	1.25	2.79	1.22	0.112	2.33	2.24 (1.86, 2.70)	<0.001
Time spent driving	1.98	3.59	1.94	0.101	1.88	1.42 (1.27, 1.58)	<0.001
Creatinine (enzymatic) in urine	3.61	5.58	3.56	0.097	1.60	1.45 (1.32, 1.58)	<0.001
Worry too long after embarrassment	4.44	6.49	4.39	0.093	1.51	1.30 (1.20, 1.41)	<0.001
Number of vehicles in household	0.92	1.99	0.89	0.092	2.26	1.53 (1.32, 1.78)	<0.001
Time spent watching television (TV)	1.08	2.19	1.05	0.090	2.11	1.71 (1.49, 1.97)	<0.001
Eats eggs, dairy, wheat, sugar	0.93	1.77	0.91	0.075	1.97	1.80 (1.08, 2.99)	0.027
Long-standing illness, disability or infirmity	2.65	3.93	2.62	0.073	1.52	1.37 (1.24, 1.52)	<0.001

Continued on next page

Table S18 continued from previous page

Predictor	Missing, %			SMD	Crude OR	Adjusted OR (95% CI)	BH q
	Overall	Cases	Non-cases				
Mood swings	2.84	4.15	2.81	0.073	1.49	1.24 (1.13, 1.37)	<0.001
Standard PRS for Parkinson's disease (PD)	3.22	4.30	3.20	0.058	1.36	1.25 (1.12, 1.39)	<0.001
Standard PRS for Alzheimer's disease (AD)	3.22	4.30	3.20	0.058	1.36	1.25 (1.12, 1.39)	<0.001
Vitamin D	10.77	12.46	10.73	0.054	1.18	1.14 (1.07, 1.21)	<0.001
Hearing difficulty/problems with background noise	2.32	3.17	2.30	0.053	1.39	1.41 (1.26, 1.58)	<0.001
Falls in the last year	0.53	0.98	0.52	0.053	1.89	2.44 (1.94, 3.09)	<0.001
Water intake	0.95	1.51	0.94	0.052	1.62	1.33 (1.12, 1.57)	0.001
Taking other prescription medications	0.61	1.06	0.60	0.051	1.79	1.79 (1.45, 2.22)	<0.001
APOE ε4 copies	17.78	19.68	17.73	0.050	1.14	1.07 (1.02, 1.13)	0.005
Hand grip strength (left)	0.67	1.12	0.66	0.049	1.71	1.53 (1.27, 1.85)	<0.001
C-reactive protein	6.74	8.00	6.71	0.049	1.21	1.14 (1.06, 1.24)	<0.001
Neutrophil percentage	5.01	6.05	4.98	0.047	1.23	1.09 (1.00, 1.20)	0.059
Basophil count	5.01	6.05	4.98	0.047	1.23	1.09 (1.00, 1.20)	0.059
IGF-1	7.05	8.26	7.02	0.047	1.19	1.13 (1.04, 1.22)	0.003
Alanine aminotransferase	6.58	7.74	6.55	0.046	1.20	1.13 (1.05, 1.23)	0.003
Sleeplessness/insomnia	0.30	0.59	0.29	0.045	2.04	2.24 (1.62, 3.11)	<0.001
Number of treatments/medications taken	0.17	0.40	0.16	0.044	2.44	2.20 (1.57, 3.09)	<0.001
Aspartate aminotransferase	6.90	8.02	6.87	0.044	1.18	1.12 (1.04, 1.21)	0.004
Saturated Fatty Acids to Total Fatty Acids percentage	2.87	3.57	2.86	0.040	1.26	1.26 (1.13, 1.39)	<0.001
Adjusted T/S ratio	5.94	6.78	5.92	0.035	1.16	1.07 (0.98, 1.16)	0.127
Maternal history of high blood pressure	1.91	2.37	1.89	0.033	1.26	1.22 (1.07, 1.39)	0.003
Maternal history of Alzheimer's disease/dementia	1.91	2.37	1.89	0.033	1.26	1.22 (1.07, 1.39)	0.003
Touchscreen duration	0.17	0.28	0.17	0.025	1.70	1.59 (0.95, 2.67)	0.079

Table S19 Missingness of each ITU-DESC predictor and incident Alzheimer's disease in the study cohort (500,410 participants; all-incident horizon). *Missing, %* is the share of participants with no recorded value after the recoding described under "Missing-value handling" in Supplementary Methods, overall and separately among participants who did and did not develop Alzheimer's disease. *SMD* is the standardized difference in the missing rate between cases and non-cases; 5 of 37 predictors reach 0.10 and 27 of 37 unadjusted chi-square tests have $P < 0.05$. *Crude OR* is the unadjusted odds ratio of Alzheimer's disease for missing versus recorded. *Adjusted OR* is from a logistic model adjusted for age (linear and quadratic), sex, education, Townsend deprivation, recruitment year, assessment center and ethnicity; 16 of 37 fitted predictors keep a Benjamini–Hochberg $q < 0.05$. Age was recorded for every participant. Rows are ordered by SMD.

Predictor	Missing, %			SMD	Crude OR	Adjusted OR (95% CI)	BH q
	Overall	Cases	Non-cases				
Pairs-matching completion time, round 2	3.94	10.00	3.88	0.243	2.75	2.08 (1.87, 2.31)	<0.001
Pairs-matching completion time, round 1	3.73	9.27	3.67	0.229	2.68	2.01 (1.80, 2.24)	<0.001
Length of mobile phone use	1.60	3.90	1.57	0.143	2.54	1.82 (1.56, 2.11)	<0.001
Mother's vital status or age at death	3.02	5.29	2.99	0.116	1.81	1.51 (1.33, 1.71)	<0.001
Worry too long after embarrassment	4.44	6.83	4.42	0.105	1.59	1.35 (1.21, 1.52)	<0.001
Mean time to correctly identify matches	1.15	2.44	1.14	0.098	2.17	1.63 (1.34, 1.97)	<0.001
Peak expiratory flow (PEF)	9.76	12.65	9.73	0.093	1.34	1.01 (0.92, 1.10)	0.873

Continued on next page

Table S19 continued from previous page

Predictor	Missing, %			SMD	Crude OR	Adjusted OR (95% CI)	BH q
	Overall	Cases	Non-cases				
Time spent watching television (TV)	1.08	2.11	1.06	0.084	2.00	1.61 (1.31, 1.97)	<0.001
Number of vehicles in household	0.92	1.80	0.91	0.077	1.99	1.39 (1.11, 1.74)	0.011
Mood swings	2.84	4.24	2.83	0.076	1.52	1.28 (1.11, 1.47)	0.002
Long-standing illness, disability or infirmity	2.65	3.94	2.64	0.073	1.51	1.37 (1.18, 1.58)	<0.001
Time spent driving	1.98	3.07	1.96	0.070	1.58	1.18 (1.00, 1.39)	0.106
Other exercises (swimming, cycling, keep fit, bowling)	1.25	2.13	1.24	0.069	1.73	1.25 (0.89, 1.73)	0.257
Hearing difficulty/problems with background noise	2.32	3.30	2.31	0.060	1.44	1.49 (1.27, 1.75)	<0.001
Eats eggs, dairy, wheat, sugar	0.93	1.56	0.92	0.058	1.71	1.67 (0.75, 3.71)	0.259
Creatinine (enzymatic) in urine	3.61	4.65	3.60	0.053	1.31	1.15 (1.00, 1.33)	0.097
Water intake	0.95	1.52	0.94	0.052	1.62	1.37 (1.07, 1.74)	0.028
Falls in the last year	0.53	0.92	0.52	0.047	1.76	2.47 (1.75, 3.49)	<0.001
APOE $\varepsilon 4$ copies	17.78	19.35	17.76	0.041	1.11	1.07 (0.99, 1.15)	0.158
Standard PRS for Parkinson's disease (PD)	3.22	3.90	3.22	0.037	1.22	1.12 (0.95, 1.33)	0.257
Standard PRS for Alzheimer's disease (AD)	3.22	3.90	3.22	0.037	1.22	1.12 (0.95, 1.33)	0.257
Sleeplessness/insomnia	0.30	0.53	0.29	0.036	1.80	2.19 (1.33, 3.60)	0.006
Taking other prescription medications	0.61	0.92	0.60	0.036	1.52	1.62 (1.17, 2.25)	0.011
Vitamin D	10.77	11.85	10.76	0.035	1.12	1.05 (0.96, 1.15)	0.345
Number of treatments/medications taken	0.17	0.33	0.17	0.033	2.01	1.91 (1.13, 3.22)	0.035
Saturated Fatty Acids to Total Fatty Acids percentage	2.87	3.40	2.87	0.030	1.19	1.18 (1.01, 1.38)	0.071
Hand grip strength (left)	0.67	0.92	0.67	0.028	1.38	1.23 (0.91, 1.65)	0.257
Adjusted T/S ratio	5.94	6.46	5.94	0.022	1.09	1.03 (0.90, 1.16)	0.752
IGF-1	7.05	7.58	7.04	0.021	1.08	1.02 (0.91, 1.15)	0.752
Alanine aminotransferase	6.58	7.09	6.57	0.020	1.08	1.03 (0.91, 1.16)	0.752
C-reactive protein	6.74	7.24	6.74	0.020	1.08	1.02 (0.91, 1.15)	0.752
Touchscreen duration	0.17	0.25	0.17	0.019	1.52	1.68 (0.75, 3.76)	0.259
Maternal history of high blood pressure	1.91	2.17	1.90	0.019	1.14	1.14 (0.94, 1.39)	0.257
Maternal history of Alzheimer's disease/dementia	1.91	2.17	1.90	0.019	1.14	1.14 (0.94, 1.39)	0.257
Aspartate aminotransferase	6.90	7.30	6.89	0.016	1.06	1.01 (0.90, 1.14)	0.873
Basophil count	5.01	5.21	5.01	0.009	1.04	0.90 (0.78, 1.04)	0.257
Neutrophil percentage	5.01	5.21	5.01	0.009	1.04	0.90 (0.78, 1.04)	0.257

Competing risk of death

The main analyses treat dementia as a binary outcome and count participants who died without a recorded dementia diagnosis as non-cases. This section quantifies death as a competing event and re-evaluates all scores on the holdout test set using two time-to-event frameworks. No prediction model was retrained. All analyses used the fixed holdout-test predictions from the main analysis and retained its definition of all-cause dementia; rank-based measures used the uncalibrated predictions. The hazard models used each participant's fixed score as the sole covariate to estimate its association with time to dementia or death.

Methods

Follow-up extended from the baseline visit to the earliest of a recorded dementia diagnosis, death or the administrative censoring date of 20 November 2025. Because

the latest baseline visit occurred on 1 October 2010 and we excluded participants lost to follow-up during cohort construction, all participants had at least 15.1 years of potential follow-up (maximum, 19.7 years), and none was censored before 15 years. At the 5-, 10- and 15-year horizons, each participant therefore fell into exactly one of three states: dementia recorded by the horizon; death without dementia by the horizon; or alive and dementia-free at the horizon (Table S20).

Three consequences follow. First, the binary horizon labels used in the main analysis coincide with the observed cumulative-incidence status, and the Aalen–Johansen cumulative incidence of dementia [15] at each horizon equals the observed proportion. Second, both time-dependent AUC definitions described by Blanche et al. [16] can be estimated without inverse-probability weighting. The cumulative-incidence AUC treats as controls all participants without dementia by the horizon, including those who died, and therefore equals the corresponding binary AUC; the cause-specific AUC uses as controls only participants who were alive and dementia-free at the horizon. Third, in the absence of censoring before the horizon, the Fine–Gray subdistribution hazard model [17] reduces to a Cox model with a modified risk set in which participants who die without dementia remain at risk until the horizon [18]. We therefore fitted it alongside a cause-specific Cox model, which censored deaths at the time of death. Both models used the linear predictor of the fixed score as their only covariate.

For each score and horizon, we report both time-dependent AUCs with DeLong CIs; Harrell's concordance index, with deaths censored at the time of death; and the concordance index of Wolbers et al. [19], in which participants who died remain comparable as non-events until the horizon. We obtained CIs for both concordance indices by paired participant-level bootstrap resampling (2,000 resamples). We also report the AUC of each score for death by the horizon among participants without dementia, quantifying how well each score discriminates the competing event. For ITU-DESC, we report cause-specific and Fine–Gray hazard ratios per standard-deviation increase in the linear predictor. Across deciles of the score, we report the Aalen–Johansen cumulative incidence of dementia and death, together with the Kaplan–Meier estimate that treats death as censoring. The two concordance indices target different quantities: with deaths censored, C measures concordance among pairs still alive and dementia-free at the time of the earlier dementia event and therefore depends on the mortality distribution; with deaths retained as non-events, it targets discrimination of the observable cumulative incidence. We did not estimate an AUC that treats death as non-informative censoring because this quantity corresponds to a hypothetical setting in which death is eliminated and requires conditional independence of death and dementia given the score, an assumption that cannot be verified from these data.

Results

In the holdout test set, 4,880 participants died without a recorded dementia diagnosis during full follow-up, compared with 1,138 dementia cases: 4.29 deaths per case. At 5, 10 and 15 years, the corresponding ratios were 11.23, 5.41 and 4.09 (Table S20). The eight scores kept the same rank ordering across all frameworks at each horizon, with ITU-DESC first throughout (Table S22). At 10 years, ITU-DESC had a cumulative-incidence AUC of 0.850 (95% CI, 0.833 to 0.867), a cause-specific AUC of 0.857 (0.841

to 0.873), a concordance index of 0.853 with deaths censored and a concordance index of 0.849 (0.831 to 0.865) with deaths retained as non-events. ITU-DESC exceeded every comparator; at 10 and 15 years, all CIs for these differences excluded zero, whereas at 5 years, with only 75 dementia cases, the intervals for the difference from UKB-DRP included zero.

The scores also discriminated death without dementia. For ITU-DESC, the AUC for this competing event was 0.698, 0.714 and 0.732 at 5, 10 and 15 years, respectively, and the cause-specific hazard ratio for death at 10 years was 2.05 per standard-deviation increase in the linear predictor (95% CI, 1.96 to 2.13). Consistent with this association, the 10-year Fine–Gray hazard ratio for dementia was lower than the corresponding cause-specific hazard ratio: 4.16 (3.75 to 4.62) versus 4.25 (3.83 to 4.71) per standard-deviation increase.

The observed 10-year dementia proportion equaled the Aalen–Johansen cumulative incidence in every decile of the ITU-DESC score, and the mean calibrated prediction tracked the observed risk. In the highest-risk decile, the mean predicted risk was 4.65%, compared with an observed risk of 4.64% (95% CI, 4.09 to 5.25; Table S21 and Fig. S3). The corresponding Kaplan–Meier estimate, which treats death as censoring, was 5.00%, higher than the observed proportion because it estimates risk under a hypothetical setting in which death is removed as a competing event. Deaths also clustered in the higher-risk deciles: among binary non-cases, 11.65% of participants in the highest-risk decile had died by 10 years and 24.13% by 15 years, compared with 0.90% and 1.48%, respectively, in the lowest-risk decile.

Interpretation

At fixed horizons, the binary evaluation used in the main analysis is equivalent to a cumulative-incidence evaluation because no participant was censored before 15 years. The competing event instead changes the interpretation of alternative estimands: treating death as censoring yields higher risk estimates, and the cause-specific hazard ratio exceeds the subdistribution hazard ratio because the score also predicts death. Neither difference changes score ranking or the agreement between predicted and observed risk. The remaining limitation of the binary framework is therefore the all-incident outcome, which combines follow-up durations of 15.1 to 19.7 years and is not a fixed-horizon cumulative-incidence measure, unlike the 5- and 10-year outcomes.

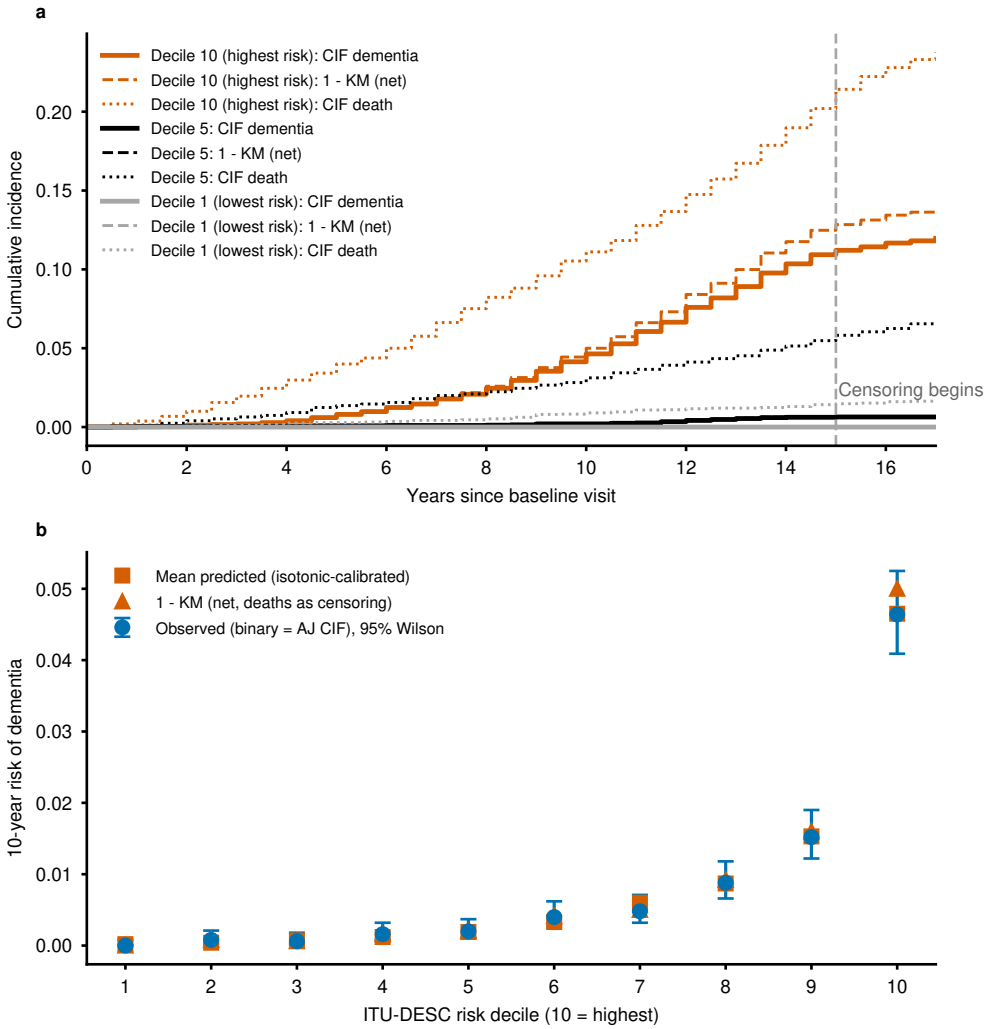

Fig. S3 Competing risk of death across deciles of the ITU-DESC score in the holdout test set (50,041 participants). **a**, Aalen–Johansen cumulative incidence of dementia and death without dementia, together with the Kaplan–Meier estimate treating death as censoring, for the highest-, fifth- and lowest-risk deciles. The vertical line marks 15 years, after which administrative censoring begins. **b**, Ten-year risk by decile, showing the mean isotonic-calibrated prediction, the observed proportion with its Wilson 95% confidence interval, and the Kaplan–Meier estimate treating death as censoring. The observed proportion is identical to the Aalen–Johansen cumulative incidence at this horizon.

Table S20 Follow-up status by horizon in the study cohort (500,410 participants) and the holdout test set (50,041 participants). At each fixed horizon every participant is in exactly one of three states: dementia recorded by the horizon, died without a dementia record by the horizon, or alive without a dementia record at the horizon; no participant was administratively censored before 15 years, and participants lost to follow-up were excluded at cohort construction. *All-incident* is the manuscript's binary label over each participant's full follow-up (15.1 to 19.7 years), with the event-free column counting participants censored at the administrative date. *Deaths per case* is the ratio of dementia-free deaths to dementia cases.

Horizon	n	Dementia, n (%)	Died without dementia, n (%)	Alive without dementia, n (%)	Deaths per case
Study cohort					
5 years	500,410	756 (0.15)	8,565 (1.71)	491,089 (98.14)	11.33
10 years	500,410	4,223 (0.84)	23,107 (4.62)	473,080 (94.54)	5.47
15 years	500,410	10,792 (2.16)	44,966 (8.99)	444,652 (88.86)	4.17
All-incident	500,410	11,380 (2.27)	49,366 (9.87)	439,664 (87.86)	4.34
Holdout test set					
5 years	50,041	75 (0.15)	842 (1.68)	49,124 (98.17)	11.23
10 years	50,041	421 (0.84)	2,279 (4.55)	47,341 (94.60)	5.41
15 years	50,041	1,086 (2.17)	4,438 (8.87)	44,517 (88.96)	4.09
All-incident	50,041	1,138 (2.27)	4,880 (9.75)	44,023 (87.97)	4.29

Table S21 Ten-year risk of all-cause dementia by decile of the ITU-DESC score in the holdout test set (50,041 participants; decile 10 is the highest risk). *Predicted* is the mean isotonic-calibrated 10-year risk. *Observed* is the share of the decile with dementia recorded within 10 years, which equals the Aalen–Johansen cumulative incidence with death as a competing event because no participant was censored before 15 years; the interval is a Wilson 95% interval. $1-$KM treats death as censoring and estimates the risk that would apply if no one died. *Death* is the Aalen–Johansen cumulative incidence of death without dementia. The last column is the share of the decile's binary non-cases who had died by 10 years rather than remaining alive without dementia.

Decile	n	Predicted, %	Observed, % (95% CI)	$1-$KM, %	Death, %	Non-cases who died, %
1	5,005	0.02	0.00 (0.00, 0.08)	0.00	0.90	0.90
2	5,004	0.04	0.08 (0.03, 0.21)	0.08	1.34	1.34
3	5,004	0.07	0.06 (0.02, 0.18)	0.06	1.42	1.42
4	5,004	0.12	0.16 (0.08, 0.32)	0.16	2.56	2.56
5	5,004	0.19	0.20 (0.11, 0.37)	0.20	3.12	3.12
6	5,004	0.33	0.40 (0.26, 0.62)	0.41	3.88	3.89
7	5,004	0.60	0.48 (0.32, 0.71)	0.50	5.40	5.42
8	5,004	0.87	0.88 (0.66, 1.18)	0.92	7.21	7.28
9	5,004	1.53	1.52 (1.22, 1.90)	1.60	8.61	8.75
10	5,004	4.65	4.64 (4.09, 5.25)	5.00	11.11	11.65

Table S22 Discrimination of each score on the holdout test set (50,041 participants) when death without dementia is treated as a competing event, at the 5-, 10- and 15-year horizons. Predictions are the uncalibrated holdout-test predictions of the main analysis; no prediction model was retrained. *AUC, cumulative incidence* takes as controls every participant without dementia by the horizon, including those who died, and equals the binary AUC of the horizon label because no participant was censored before 15 years. *AUC, cause-specific* takes as controls only participants alive and dementia-free at the horizon. *C, deaths censored* is Harrell's concordance with deaths censored at the date of death; *C, deaths as non-events* is the concordance of Wolbers et al., in which participants who died remain comparable non-events until the horizon. ΔC is ITU-DESC minus the comparator, with a paired participant-level bootstrap interval (2,000 resamples). *AUC for death* is the AUC of each score for death by the horizon among participants without dementia by the horizon, and measures how strongly the score predicts the competing event; its DeLong interval half-width is at most ± 0.021 and is given in the source data. AUC intervals are DeLong intervals.

Model	AUC, cumulative incidence (95% CI)	AUC, cause-specific (95% CI)	*C*, deaths censored (95% CI)	*C*, deaths as non-events (95% CI)	ΔC, deaths as non-events (95% CI)	AUC for death
5 years: 75 dementia cases, 842 deaths without dementia						
ITU-DESC	0.839 (0.796, 0.882)	0.842 (0.798, 0.884)	0.840 (0.793, 0.882)	0.839 (0.792, 0.881)	–	0.698
UKB-DRP	0.823 (0.778, 0.868)	0.826 (0.781, 0.870)	0.824 (0.774, 0.867)	0.823 (0.773, 0.866)	0.016 (−0.005, 0.037)	0.700
UKBDRS	0.800 (0.752, 0.848)	0.802 (0.755, 0.850)	0.801 (0.750, 0.848)	0.800 (0.749, 0.846)	0.039 (0.012, 0.065)	0.673
UKBDRS (w/o *APOE* ε4)	0.800 (0.753, 0.848)	0.803 (0.755, 0.851)	0.801 (0.751, 0.848)	0.800 (0.750, 0.847)	0.039 (0.010, 0.066)	0.683
DRS	0.762 (0.711, 0.812)	0.764 (0.714, 0.814)	0.762 (0.709, 0.812)	0.761 (0.709, 0.811)	0.078 (0.050, 0.106)	0.661
CAIDE	0.657 (0.604, 0.710)	0.659 (0.606, 0.712)	0.658 (0.604, 0.710)	0.657 (0.604, 0.709)	0.182 (0.122, 0.240)	0.619
CAIDE (w/o *APOE* ε4)	0.626 (0.576, 0.676)	0.628 (0.579, 0.678)	0.627 (0.576, 0.675)	0.626 (0.575, 0.674)	0.213 (0.151, 0.274)	0.621
ANU-ADRI	0.633 (0.572, 0.694)	0.635 (0.574, 0.696)	0.634 (0.572, 0.692)	0.633 (0.571, 0.691)	0.206 (0.129, 0.283)	0.594
10 years: 421 dementia cases, 2,279 deaths without dementia						
ITU-DESC	0.850 (0.833, 0.867)	0.857 (0.841, 0.873)	0.853 (0.836, 0.869)	0.849 (0.831, 0.865)	–	0.714
UKB-DRP	0.836 (0.819, 0.854)	0.844 (0.827, 0.861)	0.840 (0.823, 0.856)	0.835 (0.818, 0.852)	0.014 (0.004, 0.024)	0.717
UKBDRS	0.813 (0.794, 0.832)	0.820 (0.802, 0.839)	0.816 (0.798, 0.834)	0.812 (0.793, 0.830)	0.037 (0.026, 0.048)	0.692
UKBDRS (w/o *APOE* ε4)	0.810 (0.791, 0.828)	0.818 (0.799, 0.836)	0.813 (0.794, 0.831)	0.808 (0.789, 0.826)	0.040 (0.028, 0.053)	0.704
DRS	0.781 (0.761, 0.800)	0.788 (0.769, 0.807)	0.784 (0.764, 0.803)	0.780 (0.760, 0.799)	0.069 (0.055, 0.084)	0.677
CAIDE	0.682 (0.660, 0.705)	0.688 (0.666, 0.710)	0.685 (0.662, 0.706)	0.682 (0.659, 0.703)	0.167 (0.145, 0.190)	0.631
CAIDE (w/o *APOE* ε4)	0.648 (0.627, 0.670)	0.655 (0.634, 0.676)	0.652 (0.630, 0.673)	0.648 (0.625, 0.669)	0.201 (0.176, 0.226)	0.637
ANU-ADRI	0.621 (0.595, 0.648)	0.625 (0.599, 0.652)	0.623 (0.595, 0.650)	0.621 (0.593, 0.647)	0.228 (0.198, 0.261)	0.586
15 years: 1,086 dementia cases, 4,438 deaths without dementia						
ITU-DESC	0.853 (0.843, 0.863)	0.868 (0.859, 0.878)	0.858 (0.848, 0.867)	0.849 (0.839, 0.859)	–	0.732
UKB-DRP	0.840 (0.829, 0.850)	0.856 (0.846, 0.866)	0.845 (0.835, 0.854)	0.836 (0.825, 0.846)	0.013 (0.007, 0.020)	0.733
UKBDRS	0.824 (0.813, 0.835)	0.840 (0.829, 0.850)	0.829 (0.817, 0.839)	0.821 (0.809, 0.831)	0.029 (0.022, 0.036)	0.712
UKBDRS (w/o *APOE* ε4)	0.813 (0.801, 0.824)	0.831 (0.820, 0.841)	0.819 (0.807, 0.829)	0.809 (0.798, 0.820)	0.040 (0.033, 0.048)	0.723
DRS	0.787 (0.776, 0.799)	0.803 (0.792, 0.814)	0.792 (0.781, 0.803)	0.784 (0.773, 0.795)	0.065 (0.056, 0.074)	0.693
CAIDE	0.695 (0.682, 0.709)	0.708 (0.695, 0.722)	0.700 (0.686, 0.713)	0.693 (0.679, 0.707)	0.156 (0.143, 0.170)	0.650
CAIDE (w/o *APOE* ε4)	0.648 (0.634, 0.661)	0.663 (0.649, 0.676)	0.654 (0.640, 0.668)	0.646 (0.632, 0.660)	0.203 (0.188, 0.219)	0.655
ANU-ADRI	0.593 (0.576, 0.609)	0.601 (0.584, 0.618)	0.597 (0.580, 0.614)	0.592 (0.575, 0.609)	0.258 (0.240, 0.277)	0.587

Feature importance without random imputation

We computed the SHAP summary plot in the main text on randomly imputed data, consistent with the missing-data policy used throughout model development. Because random imputation replaces each missing value with a draw independent of the participant receiving it, it also removes the predictive information carried by patterns of measurement availability. The resulting feature-importance ranking could therefore depend on the imputation policy. As a sensitivity analysis, we refitted the model using the same 38 features without imputation, letting CatBoost handle the missing values natively during tree construction (Fig. S4).

The extremes of the feature-importance ranking held across missing-data policies. The same six features occupied the highest positions under both random imputation and native missing-value handling: age; the standard PRS for Alzheimer's disease; long-standing illness or disability; number of treatments and medications taken; aspartate aminotransferase; and *APOE* $\varepsilon 4$. Left hand grip strength and maternal history of Alzheimer's disease or dementia remained among the lowest-ranked features. Most changes occurred in the middle of the ranking and spanned only a few positions. Length of mobile phone use shifted most, falling from 25th under random imputation to 34th under native handling. The main-text interpretation therefore relies primarily on features that remained highly ranked under both missing-data policies.

Local SHAP explanations

Representative local SHAP explanations are shown for four individual predictions: a true positive (Fig. S5), true negative (Fig. S6), false positive (Fig. S7), and false negative (Fig. S8). Each plot shows how individual features shifted the participant's predicted risk relative to the model's base value. These case-level explanations complement the global mean absolute SHAP summary reported in the main text.

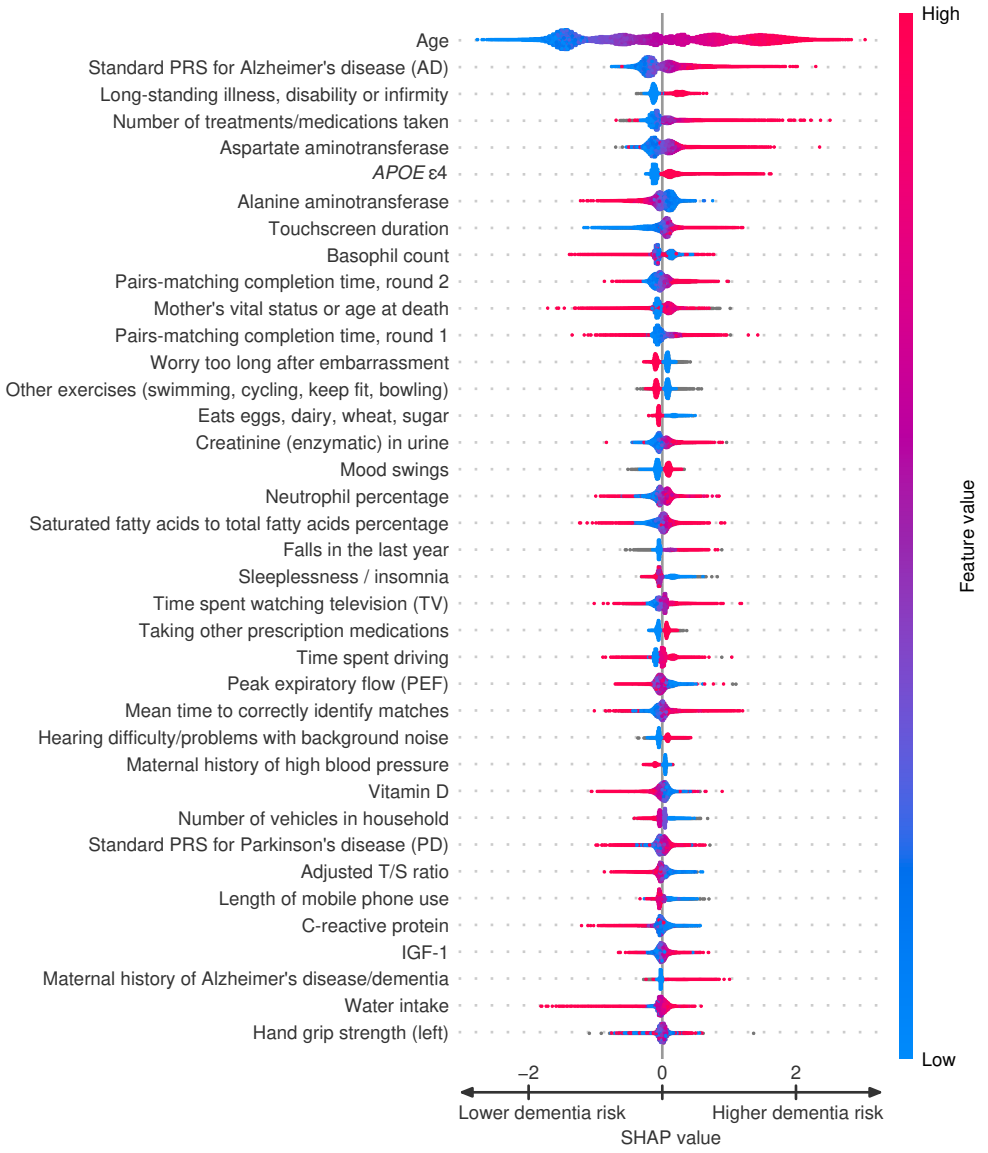

Fig. S4 SHAP summary plot for ITU-DESC for all-incident dementia without random imputation. The model used the same 38 features as in the main analysis, but missing values were left in place and handled natively by CatBoost. Each point represents one participant and is positioned according to the feature's SHAP value, corresponding to its signed contribution to the model output on the log-odds scale. Points are colored by feature value. Features are ordered by mean absolute SHAP value. SHAP values were computed on the holdout test set ($n = 50{,}041$) using the model refitted on the full development set.

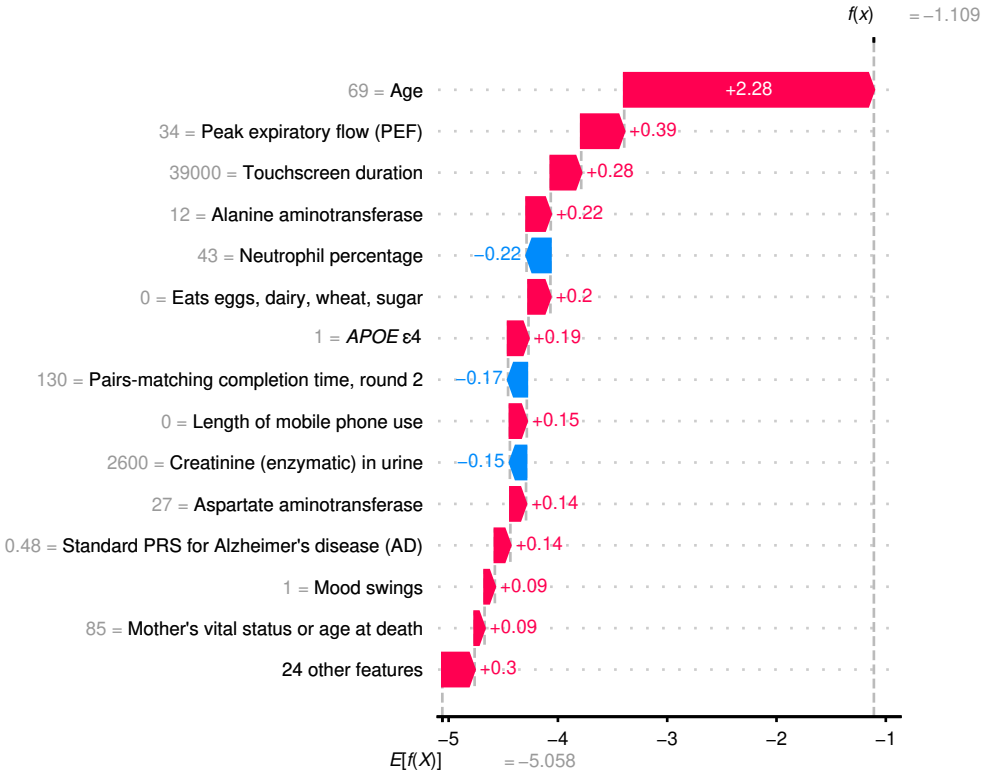

Fig. S5 Local SHAP explanation for a representative true positive prediction. The waterfall plot shows how each feature moved this participant's prediction from the cohort baseline to the model output. The bottom axis is the model output. The base value $E[f(X)] = -5.058$ is the mean model output over the test set. The value $f(x)$ is this participant's model output. Each bar is one feature's SHAP contribution, labeled with its signed value; the bars sum from the base value to $f(x)$. The number to the left of each feature name is that participant's value for the feature, rounded to two significant figures to avoid displaying exact individual-level values. Features are ordered by the magnitude of their contribution. The smallest are collapsed into a single "other features" bar.

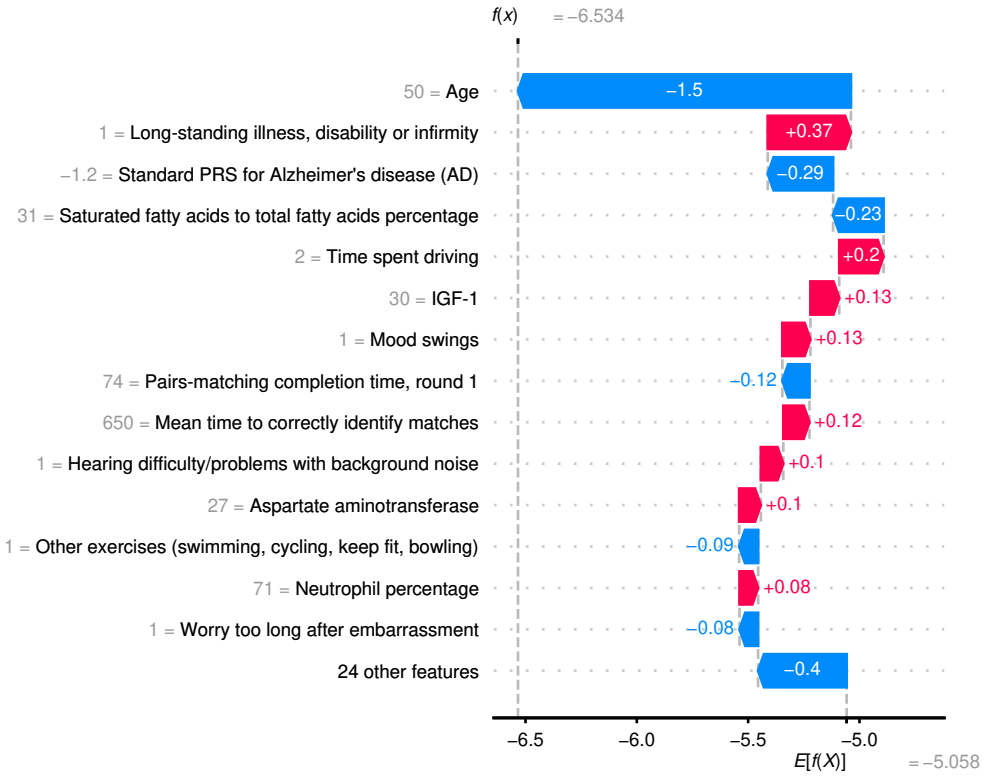

Fig. S6 Local SHAP explanation for a representative true negative prediction. Axes and layout are as in Fig. S5; the bars sum from the base value $E[f(X)]$ to the participant's output $f(x)$.

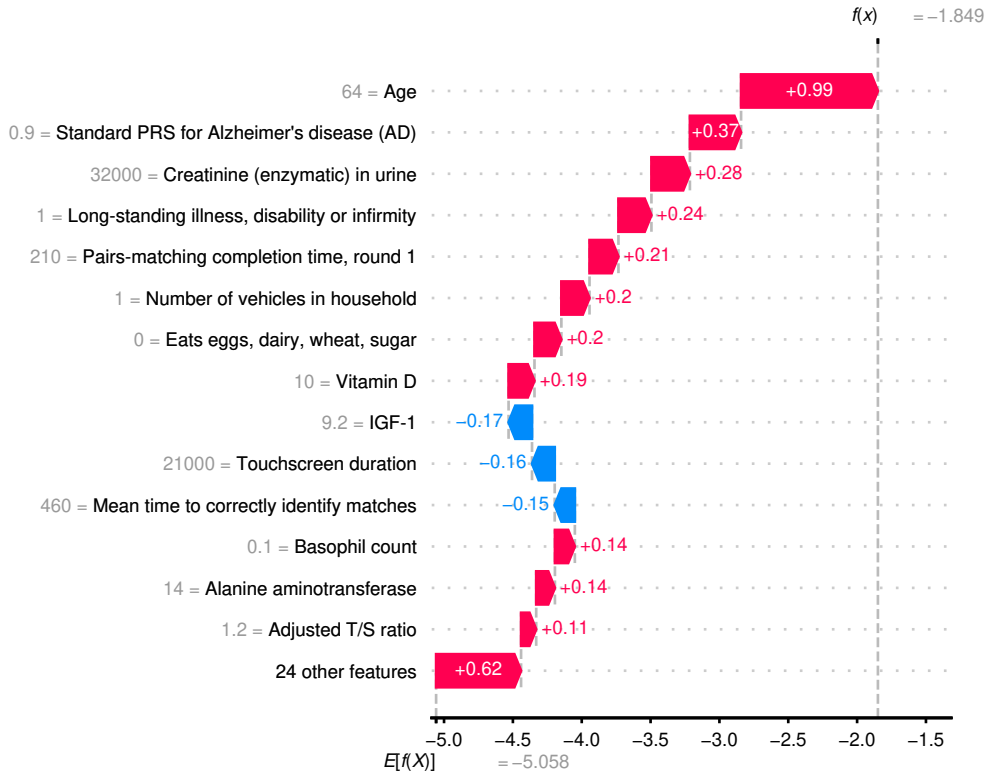

Fig. S7 Local SHAP explanation for a representative false positive prediction. Axes and layout are as in Fig. S5; the bars sum from the base value $E[f(X)]$ to the participant's output $f(x)$.

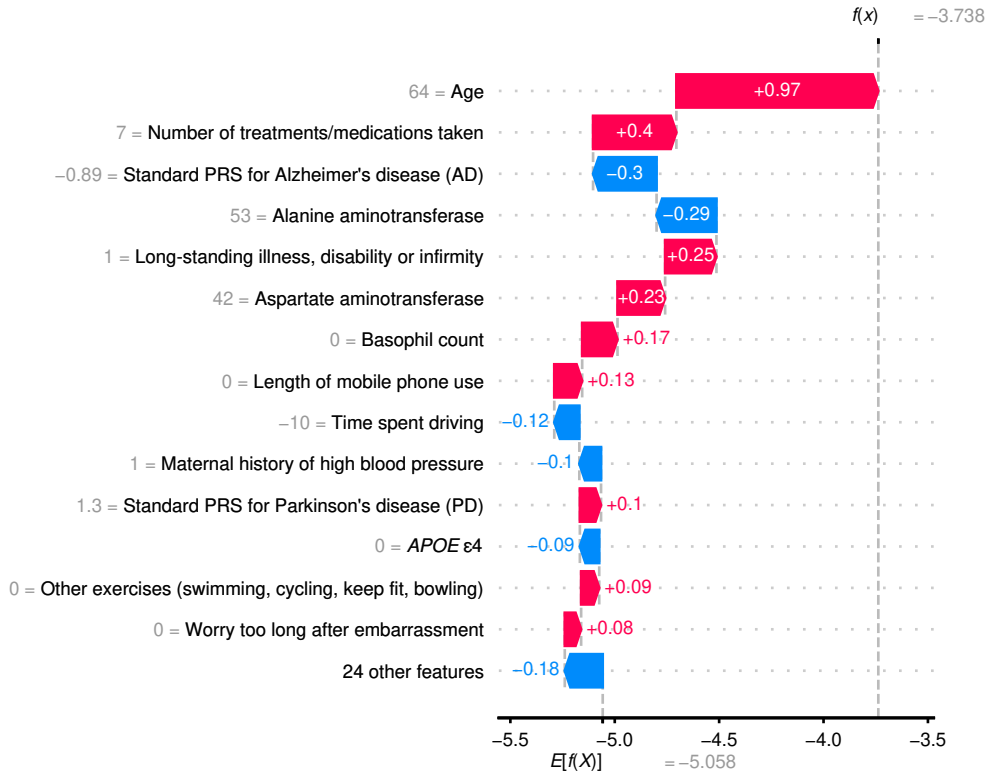

Fig. S8 Local SHAP explanation for a representative false negative prediction. Axes and layout are as in Fig. S5; the bars sum from the base value $E[f(X)]$ to the participant's output $f(x)$.

Calibration metrics

We report calibration at two levels of the hierarchy described by Van Calster et al. [20], including both summaries called calibration-in-the-large in that framework. We clipped predicted risks to $[10^{-6}, 1-10^{-6}]$ before applying any logit transformation.

Mean calibration, the weakest level, compares the mean predicted risk with the observed incidence. We report the difference $\bar{p} - \bar{y}$ in percentage points, so positive values indicate overestimation. *Weak calibration* comprises the calibration intercept and calibration slope. The calibration intercept is α in the logistic model $\text{logit}(\bar{y}) = \alpha + \text{logit}(p)$, with the linear predictor entered as an offset and its coefficient therefore fixed at one. Its target value is zero; following the sign convention of Van Calster et al., positive values indicate underestimation and negative values indicate overestimation. The mean-calibration difference and calibration intercept therefore use opposite sign conventions while describing the same mean-level agreement: at $\alpha = 0$, the score equation of the offset model is $\bar{y} - \bar{p}$, so the calibration intercept is zero exactly when the mean predicted risk equals the observed incidence. The calibration slope is the coefficient β in the two-parameter logistic regression $\text{logit}(\bar{y}) = \alpha' + \beta\,\text{logit}(p)$, with a target value of one. The intercept α' from this model is a distinct quantity that depends on β; we do not report it.

The Brier score is the mean squared difference between predicted risk and observed outcome. Because both outcomes were rare, we also report the no-skill Brier score $\bar{y}(1-\bar{y})$ for a model that predicts the observed mean incidence, and the scaled Brier score (Brier skill score), defined as $1-\text{Brier}/\text{Brier}_{\text{null}}$. We obtained 95% CIs by subject-level bootstrap resampling of the holdout-test predictions (2,000 resamples). Table S23 reports these metrics for both all-incident outcomes.

Table S23 Calibration summaries for ITU-DESC for the all-incident outcomes of all-cause dementia and Alzheimer's disease, evaluated once on the holdout test set (n = 50,041). Values are point estimates with 95% confidence intervals, where applicable, obtained by subject-level bootstrap resampling (2,000 resamples). Mean calibration is the difference between mean predicted risk and observed incidence in percentage points, positive for overestimation; the calibration intercept is on the log-odds scale, positive for underestimation. Both summarize calibration-in-the-large and equal zero when mean predicted risk matches observed incidence; their opposite sign conventions are described in the text. The mean-calibration difference was calculated from unrounded values. pp, percentage points.

	All-cause dementia	Alzheimer's disease
Events / n	1,138 / 50,041	513 / 50,041
Calibration slope	1.02 (0.97, 1.08)	0.96 (0.91, 1.05)
Calibration intercept	+0.001 (−0.06, 0.06)	+0.002 (−0.09, 0.09)
Mean predicted risk (%)	2.27	1.02
Observed incidence (%)	2.27	1.03
Mean calibration (pp)	−0.003	−0.002
Brier score	0.021	0.010
No-skill Brier score	0.022	0.010
Scaled Brier score	0.068 (0.060, 0.077)	0.035 (0.027, 0.043)

References

[1] Lundberg SM, Lee SI. A unified approach to interpreting model predictions. In: Advances in Neural Information Processing Systems. vol. 30. Curran Associates, Inc.; 2017. p. 4765–4774. Available from: https://proceedings.neurips.cc/paper_files/paper/2017/file/8a20a8621978632d76c43dfd28b67767-Paper.pdf.

[2] Keany E.: BorutaShap: A wrapper feature selection method which combines the Boruta feature selection algorithm with Shapley values, version 1.1. Zenodo; 2020. Available from: https://doi.org/10.5281/zenodo.4247618.

[3] UK Biobank.: Data-Field 400: Time to complete round. Version: August 2025. https://biobank.ndph.ox.ac.uk/ukb/field.cgi?id=400. Accessed 14 May 2026.

[4] UK Biobank.: Data-Field 6144: Never eat eggs, dairy, wheat, sugar. Version: August 2025. https://biobank.ndph.ox.ac.uk/ukb/field.cgi?id=6144. Accessed 1 June 2026.

[5] UK Biobank.: Data-Coding 100316. https://biobank.ndph.ox.ac.uk/ukb/coding.cgi?id=100316. Accessed 25 June 2026.

[6] UK Biobank.: Data-Field 1110: Length of mobile phone use. Version: September 2024. https://biobank.ndph.ox.ac.uk/ukb/field.cgi?id=1110. Accessed 25 June 2026.

[7] Ke G, Meng Q, Finley T, Wang T, Chen W, Ma W, et al. LightGBM: a highly efficient gradient boosting decision tree. In: Advances in Neural Information Processing Systems. vol. 30. Curran Associates, Inc.; 2017. p. 3146–3154. Available from: https://proceedings.neurips.cc/paper_files/paper/2017/file/6449f44a102fde848669bdd9eb6b76fa-Paper.pdf.

[8] Prokhorenkova L, Gusev G, Vorobev A, Dorogush AV, Gulin A. CatBoost: unbiased boosting with categorical features. In: Advances in Neural Information Processing Systems. vol. 31. Curran Associates, Inc.; 2018. p. 6638–6648. Available from: https://proceedings.neurips.cc/paper_files/paper/2018/file/14491b756b3a51daac41c24863285549-Paper.pdf.

[9] Chen T, Guestrin C. XGBoost: A scalable tree boosting system. In: Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. KDD '16. New York, NY, USA: ACM; 2016. p. 785–794. Available from: https://doi.org/10.1145/2939672.2939785.

[10] You J, Zhang YR, Wang HF, Yang M, Feng JF, Yu JT, et al. Development of a novel dementia risk prediction model in the general population: A large, longitudinal, population-based machine-learning study. eClinicalMedicine. 2022;53:101665. https://doi.org/10.1016/j.eclinm.2022.101665.

[11] Anatürk M, Patel R, Ebmeier KP, Georgiopoulos G, Newby D, Topiwala A, et al. Development and validation of a dementia risk score in the UK Biobank and Whitehall II cohorts. BMJ Mental Health. 2023;26(1):e300719. https://doi.org/10.1136/bmjment-2023-300719.

[12] Walters K, Hardoon S, Petersen I, Iliffe S, Omar RZ, Nazareth I, et al. Predicting dementia risk in primary care: development and validation of the Dementia Risk Score using routinely collected data. BMC Medicine. 2016;14(1):6. https://doi.org/10.1186/s12916-016-0549-y.

[13] Kivipelto M, Ngandu T, Laatikainen T, Winblad B, Soininen H, Tuomilehto J. Risk score for the prediction of dementia risk in 20 years among middle aged people: a longitudinal, population-based study. The Lancet Neurology. 2006;5(9):735–741. https://doi.org/10.1016/S1474-4422(06)70537-3.

[14] Anstey KJ, Cherbuin N, Herath PM. Development of a new method for assessing global risk of Alzheimer's disease for use in population health approaches to prevention. Prevention Science. 2013;14(4):411–421. https://doi.org/10.1007/s11121-012-0313-2.

[15] Aalen OO, Johansen S. An empirical transition matrix for non-homogeneous Markov chains based on censored observations. Scandinavian Journal of Statistics. 1978;5(3):141–150. Available from: https://www.jstor.org/stable/4615704.

[16] Blanche P, Dartigues JF, Jacqmin-Gadda H. Estimating and comparing time-dependent areas under receiver operating characteristic curves for censored event times with competing risks. Statistics in Medicine. 2013;32(30):5381–5397. https://doi.org/10.1002/sim.5958.

[17] Fine JP, Gray RJ. A proportional hazards model for the subdistribution of a competing risk. Journal of the American Statistical Association. 1999;94(446):496–509. https://doi.org/10.1080/01621459.1999.10474144.

[18] Geskus RB. Cause-specific cumulative incidence estimation and the Fine and Gray model under both left truncation and right censoring. Biometrics. 2011;67(1):39–49. https://doi.org/10.1111/j.1541-0420.2010.01420.x.

[19] Wolbers M, Blanche P, Koller MT, Witteman JCM, Gerds TA. Concordance for prognostic models with competing risks. Biostatistics. 2014;15(3):526–539. https://doi.org/10.1093/biostatistics/kxt059.

[20] Van Calster B, McLernon DJ, van Smeden M, Wynants L, Steyerberg EW, On behalf of Topic Group 'Evaluating diagnostic tests and prediction models' of the STRATOS initiative. Calibration: the Achilles heel of predictive analytics. BMC Medicine. 2019;17(1):230. https://doi.org/10.1186/s12916-019-1466-7.