Learn the Interface, Not the Implementation: Output Representation Dominates Training Regime in LLM-Based RTL Generation

Siddharth Patel†
Dept. of Electronics Engineering
Shiv Nadar Institution of Eminence
Delhi-NCR, India
sp839@snu.edu.in

Vikram Singh
STMicroelectronics
Greater Noida, India
vikram.singh@st.com

Siddharth Savyasachi Malu
AI Center of Excellence
Shiv Nadar Institution of Eminence
Delhi-NCR, India
siddharth.malu@snu.edu.in

Rohit Singh
Dept. of Electronics Engineering
Shiv Nadar Institution of Eminence
Delhi-NCR, India
rohit.singh@snu.edu.in

arXiv:2610.25881v1 [cs.LG] 19 Oct 2026

***Abstract*—Large-language-model register-transfer-level (RTL) generation has concentrated on *how to train* the model–fine-tuning, continued pre-training, reinforcement learning from a verifier. We argue this is the second-order question. On a bounded, formally-specifiable task–generating single-error-correcting, double-error-detecting (SECDED) ECC memory wrappers–we hold the task, data, and correctness oracle fixed and vary two axes independently: the *training regime* (seven settings) and the *output representation* (four levels, from fully-expanded flat Verilog to an eight-field configuration expanded by a deterministic generator). Because a reference generator exists, correctness is *decided* by formal equivalence over a 10^7-point design space, not sampled from a hand-written testbench. Two results follow. (1) Representation dominates: every representation near raw RTL collapses to a $\leq$0.17 pass rate, while the configuration reaches 0.83 with no training at all–and once the model emits the configuration, *every* training regime lands in a 0.89–1.00 band. (2) Raw RTL hits a hard output-length wall: a 512-bit design is 44,237 tokens of flat Verilog against 69 tokens of configuration, so raw generation becomes *infeasible* above 32-bit data words rather than merely inaccurate. A sub-50M-parameter from-scratch model plus the verified generator reaches 0.89 formal-equivalence (against a 32B fine-tune's 1.00) at $\sim$1600$\times$ lower inference energy per *correct* design, and inference energy dominates amortized training. The lesson is to learn the interface, not the implementation.**

***Index Terms*—LLM for RTL, hardware generation, error-correcting codes, energy-efficient ML, program synthesis, neurosymbolic EDA, machine learning for design automation.**

I. Introduction

TO emit the Verilog for a 512-bit ECC-protected memory, a language model must produce roughly 44,237 tokens of flat parity logic. To specify the *same* hardware as an eight-field configuration—which a deterministic generator expands into that exact Verilog in milliseconds—takes 69. The two descriptions are equivalent; the gap between them is 640$\times$, and this paper is about that gap. The prevailing question in LLM-based hardware generation is *how to train* the model—fine-tune it, pre-train it, reward it from a verifier. We show that a prior and larger question—*what should the model output?*—dominates it, deciding whether generation is even feasible, how often it is correct, and what it costs in energy. The incumbent a learned generator must beat is not another model but a deterministic tool that already emits every design correctly in six seconds; the only way a model earns its place is by writing less.

We make this precise. Fix a task with three properties that make it an unusually honest benchmark: it is *bounded* (a finite, enumerable parameter space), *parameterised* (structure repeats, so it is learnable), and *formally checkable* (a reference exists, so correctness is decidable rather than estimated). SECDED ECC memory-wrapper generation has all three. Every point in the space has a deterministic reference output, so we can ask of any model-generated design a question most RTL-generation work cannot answer cleanly: *is it formally equivalent to the reference?*

On this substrate we run a two-axis study. The first axis is the **training regime**: zero-shot, few-shot in-context learning, retrieval-augmented generation, LoRA and full supervised fine-tuning, continued pre-training then fine-tuning, and reinforcement learning from the formal oracle. This axis is where the literature lives. The second axis, which we argue matters more, is the **output representation**: the same design can be emitted as (L_{raw}) flat Verilog, (L_{param}) parameterised Verilog with generate loops, ($L_{hmatrix}$) a parity-check matrix plus wrapper options, or (L_{cfg}) an eight-field configuration that a deterministic back-end expands into RTL. These four describe *the same hardware* and span three orders of magnitude in

†Corresponding author. This work has been submitted to the IEEE for possible publication. Copyright may be transferred without notice, after which this version may no longer be accessible.

generation length. Holding the model and recipe constant and varying only the representation isolates its causal effect (Fig. 1).

Contributions.

- **An oracle-complete benchmark.** Because a deterministic reference generator exists, every point in a 10^7-reachable design space has a machine-verified reference and formal equivalence is *decidable*. To our knowledge this is the first RTL-generation benchmark where correctness is decided rather than sampled from a small hand-written test set.
- **A controlled two-axis study**—seven training regimes $\times$ four output representations, one task, one dataset, one oracle—that isolates representation from training regime and shows the former dominates.
- **The first energy-accounted RTL-LLM study**, reported in the unit a deployment cares about (joules per *correct* design), together with a data-leakage result and a width-disjoint split protocol that removes it.
- **A specification-compression result**: a sub-50M-parameter model, backed by the verified generator, reaches 0.89 formal-equivalence (against a 32B fine-tune's 1.00) at $\sim$1600$\times$ lower inference energy per correct design—evidence that for bounded, formally-specifiable tasks one should learn the *interface*, not the *implementation*.

II. Related Work

LLMs for RTL. A line of work benchmarks and improves language-model generation of Verilog: evaluation suites [1], [2], data- and fine-tuning-centric methods [3], [4], [5], correctness- and repair-oriented approaches [6], [7], and preference- or RL-based tuning [8], [9]. These works benchmark *models*: they hold the output (Verilog) fixed and vary the model or its training. Our study is orthogonal—we hold the model family fixed and vary the *output representation*, and show that axis dominates. **Program synthesis.** That generating into a constrained intermediate representation and expanding deterministically beats generating the final artifact directly is old in program synthesis and grammar-constrained decoding [10], [11]; we give a quantitative, energy-accounted instance for hardware and identify the mechanism—an output-length wall—that makes it not merely better but necessary at scale. **Energy accounting.** Methodology for measuring ML training and inference energy [12], [13], [14], [15] motivates our instrumentation; to our knowledge no prior RTL-generation study reports energy, and we report it per *correct* design.

III. Task and Benchmark

SECDED wrappers. A SECDED code augments a k-bit data word with r check bits so that any single-bit error is corrected and any double-bit error detected; we use the Hsiao odd-weight-column construction, which minimizes encoder/decoder XOR depth. The wrapper composes this datapath with a memory macro and options for port configuration, diagnostic readout, debug exposure, ECC-bypass, and byte-write. Byte writes require a read-modify-write, which is only meaningful on a single-port memory, a design constraint we adopt and justify, not an arbitrary restriction.

TABLE I
The four output representations (512-bit design).

Level	Output	Tokens	Back-end
L_{raw}	flat expanded Verilog	≈44,237	none
L_{param}	parameterised, generate loops	≈4,029	none
$L_{hmatrix}$	parity matrix + wrapper opts	≈1,161	matrix expander
L_{cfg}	8-field configuration	69	full generator

Design space and the leakage trap. The deployment space is the enumerable product of depth, width, port type, and four options subject to the byte-write rule: 3,584 valid configurations. A structural fact drives everything: the encoder and decoder depend on the data width *alone*, so across all 3,584 configurations there are only eight distinct encoders and eight distinct decoders – all combinatorial variety lives in the wrapper. Hence a config that shares its width with any training config shares its encoder *file byte-for-byte*, and a random config-level split leaks labels. We measure this: under config-disjoint splits the train–test near-duplicate rate reaches 0.65–1.00 depending on width overlap, and it rises monotonically with corpus size (the lowest split grows from 0.53 at 2,000 samples to 0.65 at 6,000), so these are a conservative lower bound. We therefore adopt *width-disjoint* splits for every generalization claim: test widths never appear in training in any representation, so a held-out design's encoder was never seen in any form. Because the reference generator accepts arbitrary widths and depths, it is a free, perfectly-labelled data source over $\sim 10^7$ verifiable configurations – which is what lets a from-scratch model be trained at all—while the deployment grid serves as a held-out, in-distribution test set.

The oracle harness. Correctness is decided by a five-level harness, each a strict gate. L0 lints; L1 elaborates; L2 runs a self-checking testbench with exhaustive single-bit fault injection and random double-bit injection; L3 formally checks equivalence against the reference; L4 synthesizes and reports a PPA delta; L5 checks that illegal configurations are refused. Each level is validated on the reference itself before it scores models, so a passing score is trustworthy; L3/L4 use an open equivalence/synthesis flow (Yosys [10]).

The representation ladder. Table I defines the four output languages and their generation lengths for the widest design: the same task, spanning three orders of magnitude in length, with a full deterministic expander only at the shortest rung.

IV. Method

Training regimes. We evaluate seven. *Zero-shot* and *few-shot* ($k \in \{1, 3, 5\}$) use a base coder model with no weight updates. *Retrieval-augmented generation* supplies the k nearest configurations as in-context demonstrations; we distinguish

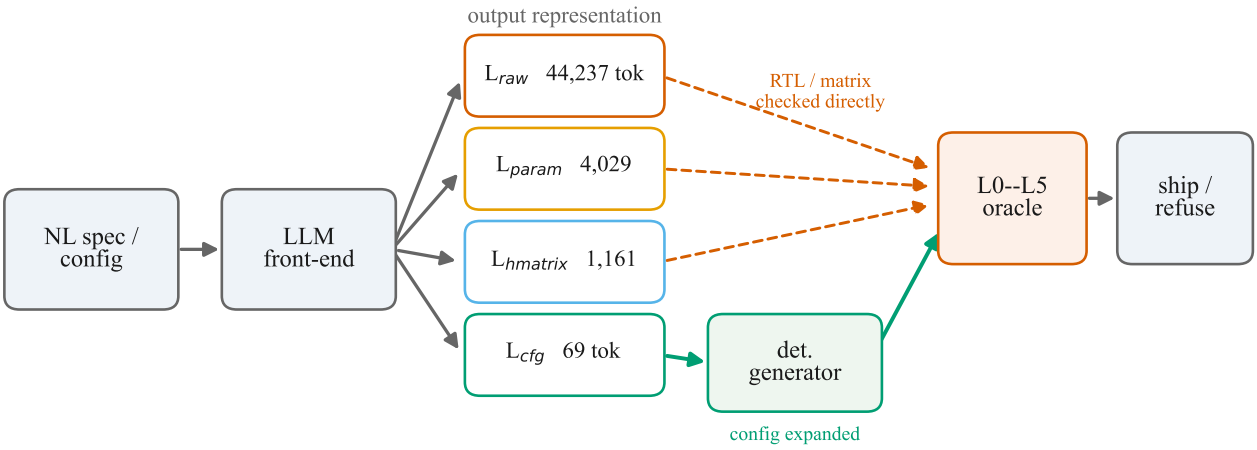

Fig. 1. The pipeline and the representation ladder. A natural-language spec or configuration enters the LLM front-end, which emits one of four output representations—shown with their 512-bit token counts. Only the configuration (L_{cfg}) is *expanded* by a deterministic generator; the raw, parameterised, and parity-matrix forms are emitted as implementation and checked directly against the reference by a five-level oracle (L0–L5), then shipped or refused. The two longest representations exceed a practical single-response generation budget—the output-length wall.

it from a *nearest-neighbour-copy* probe that returns the retrieved reference unchanged and, under width-disjoint splits, is structurally incapable of succeeding (a memorization control). *LoRA* and *full* supervised fine-tuning train on the labelled corpus; *continued-pre-training-then-SFT* first adapts on unlabelled RTL; *verifier-guided RL* optimizes a reward from compile, simulation, and formal-equivalence signals minus an area penalty—a dense, perfect, cost-free reward most RTL settings lack.

Models. The fine-tuning rows use open-weight, self-hostable, Apache-2.0-licensed models: Qwen2.5-Coder-32B-Instruct (dense) and Qwen3-Coder-30B-A3B-Instruct (a mixture-of-experts model with ≈3B active parameters); the small full-fine-tune, CPT, and RL rows use Qwen2.5-Coder-1.5B-Instruct. The from-scratch models are decoder-only transformers swept at 10M/50M/125M/350M parameters; the headline configuration is a sub-50M model that emits only the eight-field configuration and runs on CPU, letting the deployed tool generate without a GPU.

Energy methodology. GPU energy is read from the device's firmware energy counter as a delta across the run (avoiding power-polling error) and *per GPU*, so a co-resident job does not contaminate it—a subtlety that, uncorrected, roughly doubled our initial figures. We subtract a multi-minute idle baseline and repeat each measurement (n=3). Clocks were unpinned and RAPL unreadable, so figures are GPU energy only. We report joules per generated design and—the deployment-relevant unit—joules per *correct* design (J/design ÷ pass rate).

V. Results

Representation dominates. Table II is the study's spine: pass rate for each training regime × output representation. Read it two ways. *Across the ladder* (the few-shot row, the only regime run at every representation): every representation

TABLE II
PASS RATE BY REGIME × REPRESENTATION, SCORED BY **L3 FORMAL EQUIVALENCE TO THE REFERENCE** (IN-DISTRIBUTION). "—": INTENTIONALLY NOT RUN (SEC. V).

Regime \ Repr.	L_{raw}	L_{param}	$L_{hmatrix}$	L_{cfg}
few-shot	0.17	0.17	0.17	0.83
nn-copy	—	—	—	0.00
RAG	—	—	—	1.00
LoRA-SFT	—	—	0.00†	1.00
full-SFT	—	—	—	0.975
CPT-SFT	—	—	—	1.00
verifier-RL	—	—	—	1.00
from-scratch	—	—	—	0.887

closer to RTL—L_{raw}, L_{param}, $L_{hmatrix}$—collapses to 0.17, while L_{cfg} reaches 0.83 with no training. *Down the configuration column*: every regime lands in a 0.89–1.00 band. Once the model emits the configuration, the training method barely matters. The lever is the representation, not the optimizer.

The token wall. The raw-RTL route does not merely lose accuracy; above a measurable width it becomes infeasible. A 512-bit flat encoder is thousands of parity terms; the full design is ≈44,237 tokens against 69 for the configuration, exceeding any current single-response generation budget. Fig. 2 plots pass rate against generation length per width, with truncation ("ran out of budget") held as a verdict distinct from incorrect RTL ("got it wrong"): the former is the wall, the latter a correctness failure, and conflating them would erase the finding. At a 4,096-token budget, flat L_{raw} generation is feasible only up to **32-bit** data words (≈3,137 tokens) and infeasible at 48 bits and beyond (≈4,521 tokens). The wall is a property of the representation and the budget, not of the model.

Representation dominates even without a length wall. The parity-matrix rung ($L_{hmatrix}$, ≤1,161 tokens) is compact, so its failure is not a length wall but a correctness wall.

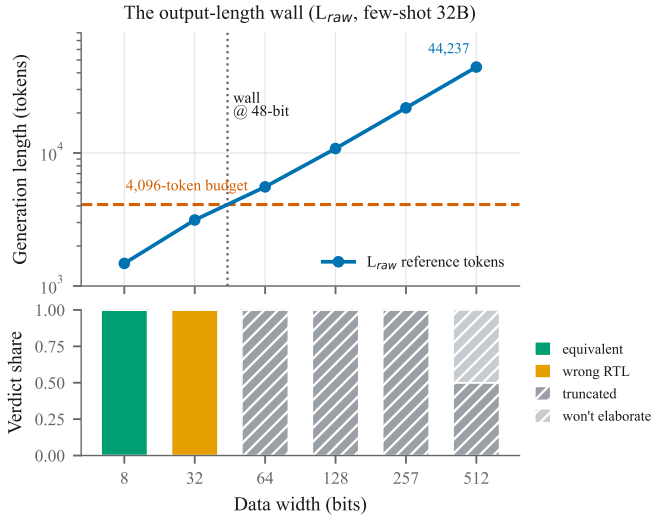

Fig. 2. Pass rate versus generation length, per width. Truncation is shown distinct from incorrect RTL. Flat L_{raw} generation is feasible only up to 32-bit words and hits the wall at 48 bits (4,096-token budget).

Five LoRA fine-tuning seeds do not close it: every seed reaches L3 = 0.00, and, checked reference-independently by exhaustive fault injection, **0 of 60** emitted matrices is a valid SECDED code of *any* dimension—the model over-generates the column list, emitting $r \approx 26$–181 check bits where 5–9 are required, often with even-weight or duplicate columns, so single-bit errors are mis-corrected. It does not emit a valid-but-different code; it does not emit a valid code at all. Only L_{cfg}, where the oracle emits the matrix, reaches 1.00. Training does not buy a valid parity matrix at this budget; the representation does.

Energy: the specification-compression result. Fig. 3 and Table III report formal-equivalence rate against joules per *correct* design (energy per design divided by pass rate): the from-scratch configuration model is roughly three orders of magnitude ($\sim$1600$\times$) more efficient per correct design than the 32B fine-tune, at comparable accuracy. This is the result in one number—emitting 69 tokens of configuration into a verified generator, rather than tens of thousands of tokens of Verilog, is where the efficiency comes from. Amortized training energy is small against lifetime inference: 0.032 kWh of fine-tuning over the deployment grid is $\sim$32 J/design, against $\sim$814 J/design at inference for the same model—inference dominates amortized training by roughly 25$\times$, and more as designs accumulate. This inverts the usual ML framing that fixates on training cost: lifetime cost is set by output length, which is the representation.

Scaling, OOD, and refusal. A from-scratch scaling sweep (10M–350M) with tokenizer and grammar-constrained-decoding ablations confirms that a sub-50M configuration model suffices. Out-of-distribution testing (Table IV) separates the regimes that tie at 1.00 in-distribution. Unseen *widths* degrade retrieval and the from-scratch model (0.89, 0.66) while the pretrained fine-tunes hold; unseen option *combinations*

TABLE III
JOULES PER *correct* DESIGN (= J/DESIGN ÷ PASS RATE). ALL ROWS ARE **GPU-ISOLATED** (PER-GPU NVML ENERGY COUNTER, IDLE-SUBTRACTED, n=3, GPU ENERGY ONLY); CLOCKS UNPINNED (SEC. VII).

Method (L_{cfg})	Pass	J/design	J/correct
from-scratch (<50M)	0.887	0.44	0.50
MoE zero-shot	0.36	341	947
dense QLoRA	1.00	814	814
dense zero-shot	0.44	950	2,159

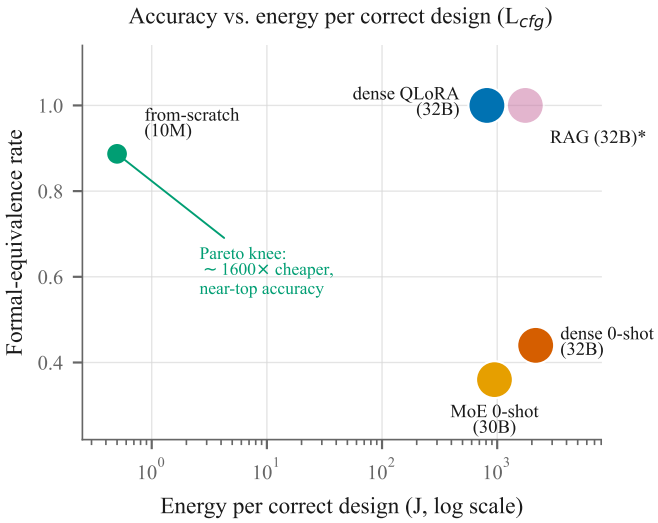

*RAG energy is shared-box (provisional); all others GPU-isolated.

Fig. 3. Formal-equivalence rate versus energy per *correct* design (log x-axis); bubble size is parameter count, colour is representation. All points are GPU-isolated (NVML counter, idle-subtracted, unpinned) except RAG, a provisional shared-box measurement (flagged in the plot). The from-scratch configuration model occupies the high-accuracy, low-energy corner—the specification-compression result made visual.

are easy for all (0.93–1.00), so the width axis is where generalization breaks. The sharpest separator is *refusal of illegal configurations*: every pretrained-LLM regime refuses perfectly (1.00)—it emits the illegal configuration faithfully and the oracle rejects it—but the efficiency-winning grammar-constrained from-scratch model refuses only **0.59**, silently legalizing 41% of illegal specs. Trained only on the legal manifold, its constrained decoder cannot emit an invalid config, so it rounds illegal requests to nearby valid ones. The cheapest model carries a safety cost the expensive ones do not—an efficiency–safety trade-off only an illegal-input test exposes.

Why the raw-RTL cells are blank. The empty cells in Table II are not omissions: one does not fine-tune a model to emit 44k tokens of flat Verilog, so the L_{raw}/L_{param} cells are blank because the representation is infeasible. The one compact rung without a length wall, $L_{hmatrix}$, we deliberately filled: the gap there is correctness, not feasibility.

VI. DISCUSSION

Don't fine-tune; retrieve. Four regimes reach 1.00 on the configuration, and one of them—retrieval-augmented

TABLE IV
OUT-OF-DISTRIBUTION BEHAVIOUR (L_{CFG}); CONFIG-EXACT, EXCEPT THE LAST COLUMN (REFUSAL ACCURACY ON ILLEGAL SPECIFICATIONS: FRACTION *not* SILENTLY LEGALIZED).

Regime	ID (deploy)	OOD width	OOD combo	illegal refusal
RAG	1.00	0.887	0.93	1.00
LoRA-SFT	1.00	1.00	1.00	1.00
CPT-SFT	1.00	1.00	1.00	1.00
verifier-RL	1.00	0.975	0.95	1.00
from-scratch	0.887	0.662	1.00	**0.59**

generation—requires no weight updates, negligible energy, and trivial updating when the design space changes. On a bounded, template-structured task the honest recommendation is therefore not "fine-tune a model" but "retrieve and let the verified back-end expand." The OOD table stress-tests this and it largely survives: retrieval refuses illegal configurations perfectly (1.00), and its one weakness is unseen widths (0.89), where a light fine-tune buys extrapolation. Notably the regime that *does* fail the safety test is the ultra-cheap grammar-constrained from-scratch model—so "cheapest" and "safest" are not the same choice.

When is an LLM the right EDA tool? Not to replace a working deterministic generator on its own turf—the incumbent shell command is faster and always correct—but at the edges it cannot serve: mapping an underspecified or natural-language request onto a valid configuration, behind a verification gate that returns RTL only when it is proven equivalent to the reference. The deployed architecture reflects this: the model writes the interface, the generator writes the implementation, and the oracle decides.

VII. THREATS TO VALIDITY

We study one ECC scheme and one generator, so external validity is limited; the mechanism—an output-length wall and a representation that dodges it—is not specific to ECC, but a second oracle-complete task family is left to future work. The reference generator emits template-structured RTL, so results may not transfer to unconstrained human-written RTL; we quantify corpus entropy (few distinct encoders/decoders) so readers can calibrate. "Correctness" is equivalence to a standard minimal-Hsiao reference; a different valid construction would define a different reference, but a freshly-generated wrapper is internally self-consistent regardless. The energy figures are GPU-isolated—which removed a real co-resident-job contamination that had inflated earlier numbers—but the clocks are *unpinned* and RAPL was unreadable, so they are GPU energy only. The *relative* orderings (configuration $\ll$ raw RTL; MoE $<$ dense; inference $\gg$ training) are robust to all of this and are the claims we rest on. The three OOD columns come from a single run per regime (reproducing the in-distribution headline exactly for three regimes and within seed variance for the other two), so the four-way tie is three robust regimes plus one higher-variance. Natural-language specifications are model-paraphrased; a held-out set of *real* user prompts, and a systematic contamination probe against pre-training data, are left to future work.

VIII. CONCLUSION

On a bounded, formally-specifiable hardware-generation task with a decidable correctness oracle, the output representation dominates the training regime: raw RTL generation hits a hard output-length wall and becomes infeasible above a measurable width, while emitting a short configuration into a verified generator reaches near-perfect accuracy under essentially any training regime—including none. A sub-50M from-scratch model, backed by the generator, reaches 0.89 formal-equivalence against a 32B fine-tune's 1.00 at $\sim$1600$\times$ lower energy per correct design, and inference energy dominates amortized training. The engineering lesson is to learn the interface, not the implementation; the methodological lesson is that template-structured RTL benchmarks must control for cross-split near-duplication. We release the generator, the oracle-complete benchmark, the verification harness, and the from-scratch weights.

REPRODUCIBILITY

All code, the frozen and hashed data splits, the verification harness, and the from-scratch weights are released in a public repository at https://github.com/thesid16/learn-the-interface. Every number in this paper is recorded—with its verification level, sample size, and provenance—in a machine-readable results file at the repository root, and the paper cites no number absent from it. A clean checkout reproduces the generator and oracle (property-tested), byte-level determinism (every configuration identical across two runs), and the width-disjoint leakage audit.

ACKNOWLEDGMENT (AI DISCLOSURE)

In accordance with the conference policy: an AI coding assistant was used for portions of the software implementation. All numerical results were produced by the released, deterministic pipeline and verified by formal equivalence; the authors are responsible for all claims.

REFERENCES

[1] M. Liu, N. Pinckney, B. Khailany, and H. Ren, "VerilogEval: Evaluating large language models for Verilog code generation," in *Proc. IEEE/ACM Int. Conf. Computer-Aided Design (ICCAD)*, 2023.

[2] Y. Lu, S. Liu, Q. Zhang, and Z. Xie, "RTLLM: An open-source benchmark for design RTL generation with large language model," in *Proc. Asia and South Pacific Design Automation Conf. (ASP-DAC)*, 2024.

[3] S. Thakur *et al.*, "VeriGen: A large language model for Verilog code generation," *ACM Trans. Design Automation of Electronic Systems (TODAES)*, vol. 29, no. 3, 2024.

[4] S. Liu *et al.*, "RTLCoder: Outperforming GPT-3.5 in design RTL generation with an open-source dataset and lightweight solution," in *Proc. IEEE Int. Workshop on LLM-Aided Design (LAD)*, 2024.

[5] M. Liu *et al.*, "ChipNeMo: Domain-adapted LLMs for chip design," arXiv:2311.00176, 2023.

[6] M. Liu, Y.-D. Tsai, W. Zhou, and H. Ren, "CraftRTL: High-quality synthetic data generation for Verilog code models with correct-by-construction non-textual representations and targeted code repair," in *Proc. Int. Conf. Learning Representations (ICLR)*, 2025.

[7] F. Cui *et al.*, "OriGen: Enhancing RTL code generation with code-to-code augmentation and self-reflection," in *Proc. IEEE/ACM Int. Conf. Computer-Aided Design (ICCAD)*, 2024.
[8] Z. Pei, H.-L. Zhen, M. Yuan, Y. Huang, and B. Yu, "BetterV: Controlled Verilog generation with discriminative guidance," in *Proc. Int. Conf. Machine Learning (ICML)*, 2024.
[9] Z. Chen *et al.*, "ChipSeek-R1: Generating human-surpassing RTL with LLM via hierarchical reward-driven reinforcement learning," arXiv:2507.04736, 2025.
[10] C. Wolf and J. Glaser, "Yosys—a free Verilog synthesis suite," in *Proc. Austrochip*, 2013.
[11] S. Geng, M. Josifoski, M. Peyrard, and R. West, "Grammar-constrained decoding for structured NLP tasks without finetuning," in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
[12] E. Strubell, A. Ganesh, and A. McCallum, "Energy and policy considerations for deep learning in NLP," in *Proc. Annual Meeting of the Association for Computational Linguistics (ACL)*, 2019.
[13] D. Patterson *et al.*, "Carbon emissions and large neural network training," arXiv:2104.10350, 2021.
[14] A. Tschand *et al.*, "MLPerf Power: Benchmarking the energy efficiency of machine learning systems from μWatts to MWatts for sustainable AI," in *Proc. IEEE Int. Symp. High-Performance Computer Architecture (HPCA)*, 2025.
[15] J. You, J.-W. Chung, and M. Chowdhury, "Zeus: Understanding and optimizing GPU energy consumption of DNN training," in *Proc. USENIX Symp. Networked Systems Design and Implementation (NSDI)*, 2023.