Causal Discovery for Linear Models with Latent Confounding via Multi-domain Identifiable Representation Learning

Danru Xu
University of Amsterdam

Aapo Hyvärinen
University of Helsinki

Sara Magliacane
Saarland University
University of Amsterdam

arXiv:2610.28134v1 [cs.LG] 20 Oct 2026

Abstract

Causal discovery with latent confounding is particularly challenging, but a common setting in real-world data. In this work, we focus on linear causal models with latent confounding, represented by acyclic directed mixed graphs (ADMGs). Existing methods that fully identify ADMGs typically assume non-Gaussianity or graphical constraints, such as bow-freeness. In contrast, we propose Linear ADMG Learning (LinADMG), a framework for learning ADMGs for linear causal models through identifiable representation learning. Rather than relying on non-Gaussianity or bow-freeness, LinADMG leverages multi-domain settings to identify the noise variables. We show that the identifiability of the noise variables is sufficient for recovering all direct causal relations, and that, under an additional latent-confounding no-cancellation condition, LinADMG can identify the whole ADMG. We implement LinADMG with four identifiable representation learning approaches from the literature and validate the methods on simulated and semisynthetic datasets.

1 INTRODUCTION

Causal discovery aims to recover causal relationships from observational data [Spirtes et al., 1995b, Pearl et al., 2003, Peters et al., 2017]. However, a common challenge in real-world data is the existence of latent confounding: an unobserved common cause of observed variables. Classical constraint-based methods allowing latent confounding, most notably FCI [Spirtes et al., 1995b], return only a Markov equivalence class rather than a unique causal structure.

To pursue a stronger identification goal for the causal structure with latent confounding, we first revisit the source-separation perspective underlying LiNGAM [Shimizu et al., 2006]. In a linear structural causal model without latent confounding, the observed variables can be regarded as linear mixtures of the noise variables. By assuming mutually independent and non-Gaussian noise variables, independent component analysis (ICA) [Hyvärinen and Oja, 2000] can be used to recover the unobserved noise variables and the corresponding mixing transformation up to permutation and component-wise scaling. The remaining ambiguities can be resolved using the acyclicity property of the causal graph, thereby fully recovering the underlying directed acyclic graph (DAG).

Latent confounding, however, fundamentally changes this picture. Suppose the underlying causal model is a DAG over both observed and latent variables. Marginalizing the latent variables generally destroys the DAG representation over the observed variables alone: latent common causes induce confounding relationships in addition to direct causal effects. Acyclic directed mixed graphs (ADMGs) provide a way to represent this structure, with directed edges encoding direct causal relations and bidirected edges encoding latent confounding [Richardson, 2003]. At the same time, marginalization can induce statistical dependence among the noise terms associated with different observed variables, violating the independent-source assumption underlying standard ICA. As a result, learning an ADMG is substantially more challenging than learning a DAG.

Existing extensions of LiNGAM to latent-variable settings exploit non-Gaussian noise variables [Hoyer et al., 2008, Salehkaleybar et al., 2020, Cai et al., 2023, Schkoda et al., 2024], while Wang and Drton [2023] establish identification of the full ADMG under additional structural assumptions, including bow-freeness.

In this work, we also aim to fully identify the ADMG, but from a different perspective. As shown in Fig. 1, inspired by LiNGAM, we propose a framework, LinADMG, that follows the same principle: first

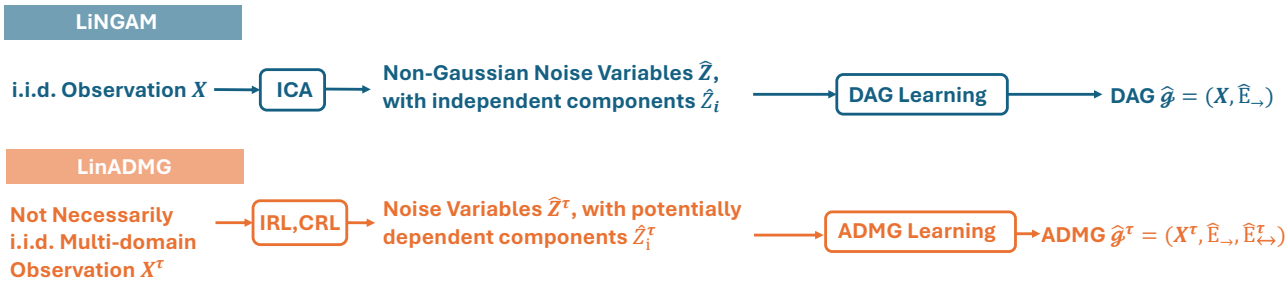

Figure 1: Overview of LiNGAM and LinADMG (our method) pipelines. LiNGAM applies ICA to the observed data $\mathbf{X}$ to estimate independent non-Gaussian noise variable $\widehat{\mathbf{Z}}$, from which it infers the directed edge set $\widehat{E}_{\rightarrow}$ of a DAG. LinADMG replaces ICA with any appropriate IRL/CRL on a multi-domain dataset $\mathbf{X}^{\tau}$ with domain index τ to estimate dependent noise variable $\widehat{\mathbf{Z}}^{\tau}$ without specific parametric assumptions on the noise variables, and thereby enabling the additional recovery of bidirected edges $\widehat{E}^{\tau}_{\leftrightarrow}$ representing latent confounding.

identify the latent noises and then use the acyclicity property of ADMGs to recover the causal graph. However, unlike LiNGAM, latent confounding makes the noises dependent, so ICA is no longer suitable.

Recent developments in identifiable representation learning (IRL), including causal representation learning (CRL) [Schölkopf et al., 2021], provide promising alternatives to ICA for LinADMG. IRL studies conditions under which latent representations can be identified from observed data generated through an unknown mixing process. CRL can be viewed as a special form of IRL, where the recovered latent variables are additionally equipped with causal semantics. A key insight from IRL is that latent representations that are not identifiable from a single observational distribution may become identifiable when we can access multiple domains in which the representation distributions vary. Such variation may arise from mixture modeling [Kivva et al., 2022], sparse masking [Xu et al., 2024], hard interventions [Varıcı et al., 2024], or other changes in the latent distribution [Xu et al., 2026].

We adopt the same principle in LinADMG. Specifically, we consider settings in which the direct causal relations among the observed variables remain invariant, while the noise distribution varies across domains. We summarize our contributions as follows:

- We propose LinADMG, a general framework that views linear causal discovery with latent confounding, which first identifies a potentially dependent noise representation and then recovers the causal graph. This extends the source-separation perspective underlying LiNGAM beyond independent non-Gaussian noises to potentially dependent noises identifiable through distributional variability.
- We show that noise variable identifiability is sufficient to recover all direct causal relations. Under an additional latent-confounding no-cancellation condition, which captures a different aspect of the model than bow-freeness and faithfulness, LinADMG further identifies the complete bidirected structure and therefore the full ADMG.
- We implement LinADMG with four representative identifiable representation-learning settings: Diagonal Gaussian Mixture Models [Kivva et al., 2022], Unpaired Partial Observations [Xu et al., 2024], Multi-domain Scalings [Xu et al., 2026], and Unknown Hard Interventions [Varıcı et al., 2024]. These are illustrative examples, and LinADMG applies to any method that identifies the latent variables. We conduct experiments on synthetic and semisynthetic datasets and show accurate recovery under the linear model.

2 PRELIMINARIES

We consider linear structural causal models (SCMs) with latent confounding and do not consider selection bias [Spirtes et al., 1995a]. Let $\mathbf{X} = (X_1, \ldots, X_n)^T \in \mathbb{R}^n$ denote the observed variables and $\mathbf{U} \in \mathbb{R}^m$ the latent variables. Following Hoyer et al. [2008], we assume the underlying causal model is

$$\begin{pmatrix} \mathbf{X} \\ \mathbf{U} \end{pmatrix} = \begin{pmatrix} B & B_1 \\ \mathbf{0} & B_2 \end{pmatrix} \begin{pmatrix} \mathbf{X} \\ \mathbf{U} \end{pmatrix} + \begin{pmatrix} \mathbf{E}_1 \\ \mathbf{E}_2 \end{pmatrix}, \tag{1}$$

where $B \in \mathbb{R}^{n \times n}$ encodes direct causal effects among the observed variables, $B_1 \in \mathbb{R}^{n \times m}$ encodes effects from the latent to the observed variables, and $B_2 \in \mathbb{R}^{m \times m}$ encodes causal relations among the latent variables. We assume that all components in $\mathbf{E}_1$ and $\mathbf{E}_2$ are mutually independent. The zero lower-left block excludes causal effects from observed variables to latent variables, so the latent variables are not descendants of the observed variables.

Solving the structural equations for $\mathbf{U}$ and substituting them into the equations for $\mathbf{X}$ gives

$$\mathbf{X} = B\mathbf{X} + \mathbf{Z}, \qquad \mathbf{Z} := B_1(I - B_2)^{-1}\mathbf{E}_2 + \mathbf{E}_1. \quad (2)$$

Although $\mathbf{E}_1$ and $\mathbf{E}_2$ are mutually independent, marginalizing the latent variables can potentially induce dependence among the components of the composite noise $\mathbf{Z} := \{Z_1, ..., Z_n\}$.

We use an acyclic directed mixed graph (ADMG) [Richardson, 2003] to represent the causal structure among the observed variables induced by the SCM in Eq. 1. With a slight abuse of notation, we use $\mathbf{X}$ to denote the corresponding node set:

$$\mathcal{G} = (\mathbf{X}, E_{\to}, E_{\leftrightarrow}). \quad (3)$$

A directed edge $X_i \to X_j \in E_{\to}$ is present when $B_{ji} \neq 0$, representing a direct causal relation from X_i to X_j. We assume the directed subgraph $(\mathbf{X}, E_{\to})$ is acyclic. A bidirected edge $X_i \leftrightarrow X_j \in E_{\leftrightarrow}$ is present when there exists a latent variable $U_i \in \mathbf{U}$ that is a common ancestor of X_i and X_j through directed paths whose intermediate nodes are all latent. Thus, $E_{\to}$ encodes the direct causal relation within $\mathbf{X}$, while $E_{\leftrightarrow}$ captures the latent-confounding induced by $\mathbf{U}$. An ADMG is *bow-free* if no pair of observed variables is connected by both a directed edge and a bidirected edge.

Our goal is to recover the ADMG $\mathcal{G} = (\mathbf{X}, E_{\to}, E_{\leftrightarrow})$, including both the directed causal structure $E_{\to}$ and the latent-confounding structure $E_{\leftrightarrow}$. Importantly, unlike latent variable LiNGAM (lvLiNGAM) [Hoyer et al., 2008] and its variants [Salehkaleybar et al., 2020, Cai et al., 2023, Schkoda et al., 2024], our work does *not* necessarily require the exogenous variables $\mathbf{E}_1$ and $\mathbf{E}_2$ to be non-Gaussian. Moreover, while the Markov property follows from our SCM assumptions, we do not require faithfulness, which assumes that every conditional independence in the observed distribution corresponds to an m-separation in the ADMG [Richardson, 2003, Spirtes et al., 2001]. Instead, we require that the noise variables $\mathbf{Z}$ in Eq. 2 can be *identified* by linear representation $W\mathbf{X}$ up to *permutation and component-wise scaling* [Comon, 1994], i.e., there exists a permutation matrix P and an invertible diagonal matrix D such that $W\mathbf{X} = DP\mathbf{Z}$, where $W \in \mathbb{R}^{n\times n}$.

In general, however, even independent $\mathbf{Z}$ are not identifiable from the observed variables $\mathbf{X}$ without additional assumptions [Comon, 1994, Hyvärinen and Pajunen, 1999, Locatello et al., 2019]. Moreover, if $\mathbf{Z}$ are potentially dependent, this is even more challenging. To solve this, a common strategy from IRL/CRL is to use distributional variability/heterogeneity in the latent variables across domains. We adopt the same principle here. Specifically, we define $\mathbf{Z}^\tau$ as the noise variables in domain $\tau \in [T]$ and allow the noise variables to follow various distributions $P^\tau_{\mathbf{Z}}$. We assume the structural coefficient matrix B is invariant. Then we have the domain-specific observation variable $\mathbf{X}^\tau$:

$$\mathbf{X}^\tau = (I - B)^{-1}\mathbf{Z}^\tau, \qquad \mathbf{Z}^\tau \sim P^\tau_{\mathbf{Z}}, \qquad \tau \in [T]. \quad (4)$$

In practice, such domains can correspond to different environments, experimental conditions, populations, or interventions that change the distribution of noises.

Next, we present four representative identifiable settings adopted from prior IRL/CRL works as *examples*:

Multi-domain Scalings (MD-Scalings) [Xu et al., 2026]. Let $\mathbf{Z}$ follow any arbitrary distribution with finite second moment and define $\mathbf{Z}^\tau = D^\tau\mathbf{Z}$, where $D^\tau = \mathrm{diag}(d^\tau_1, ..., d^\tau_n)$ is domain-specific. Identifiability of $\mathbf{Z}^\tau$ up to permutation and scaling can be achieved with sufficient variability of the scaling matrix D^τ, formalized by the full-rank condition on the matrix whose rows contain all fourth-order monomials $\{d^\tau_{i_1} d^\tau_{i_2} d^\tau_{i_3} d^\tau_{i_4}\}_{i_1\le i_2\le i_3\le i_4}$.

Diagonal Gaussian Mixture Models (Diag-GMMs)[Kivva et al., 2022]. Let u^τ denote the domain label. We generate $\mathbf{Z}^\tau \mid u^\tau \sim \mathcal{N}(\boldsymbol{\mu}(u^\tau), \boldsymbol{\Sigma}(u^\tau))$, where $\boldsymbol{\Sigma}$ is a diagonal. The domain label is unobserved, so the samples from all domains are mixed, resulting in a Gaussian mixture distribution. Identifiability up to permutation and scaling holds when there exists a pair of domains $\tau \neq \tau'$ such that the variance ratios $\mathrm{Var}(\mathbf{Z}^\tau_i)/\mathrm{Var}(\mathbf{Z}^{\tau'}_i)$ are distinct across i.

Unpaired Partial Observations (UP-Observations) [Xu et al., 2024]. Let $\mathbf{Z}$ follow an arbitrary distribution and define $\mathbf{Z}^\tau = M^\tau\mathbf{Z}$, $M^\tau = \mathrm{diag}(m^\tau_1, \ldots, m^\tau_n)$, where $m^\tau_i \in \{0, 1\}$ refers whether the ith component is observed. Thus, each sample uses only a subset of the latent variables, and the resulting representation $\mathbf{Z}^\tau$ is sparse. The masking pattern τ is unobserved and the samples are unpaired/mixed across patterns. Let $S^\tau := \{i : m^\tau_i = 1\}$ denote the active support under pattern τ. Identifying $\mathbf{Z}^\tau$ up to permutation and scaling holds under the sufficient support-variability condition: $\forall i \in [n]$, $\bigcup_{\tau\in[T]: i\notin S^\tau} S^\tau = [n] \setminus \{i\}$.

Unknown Multi-node Hard Interventions (UMN-Hard Interventions) [Varıcı et al., 2024]. Assume that the noise variables $\mathbf{Z}$ follow an additive-noise model. In each interventional domain τ, an unknown subset $I^\tau \subseteq [n]$ is targeted. A hard intervention on $i \in I^\tau$ removes the dependence of Z_i on its parents and replaces its observational conditional distribution $p_i(z_i \mid z_{\mathrm{pa}(i)})$ with an interventional distribution $q^\tau_i(z_i)$ that is independent of its parents. Define the intervention-signature matrix $D_{\mathrm{int}} \in \{0, 1\}^{n\times T}$ by $[D_{\mathrm{int}}]_{i,\tau} = \mathbf{1}\{i \in I^\tau\}$. The

intervention targets are unknown, but the collection of interventions must be sufficiently diverse: D_{int} must have full row rank. Under this condition, $\mathbf{Z}^\tau$ are identifiable up to permutation and scaling.

Importantly, our general framework requires no particular form of heterogeneity of $P^\tau_{\mathbf{Z}}$. Any data-generating setting in which the noise variables are identifiable up to permutation and scaling, while keeping B invariant across domains, is compatible with our framework. The latent-confounding structure may itself vary across domains when the mechanism generating $P^\tau_{\mathbf{Z}}$ changes. For example, under UMN-Hard Interventions setting, interventions may change the causal structure among latent variables and then change the latent-confounding relations induced among the observed variables. Thus, we allow the bidirected edge set to be domain-specific and denote it by $E^\tau_{\leftrightarrow}$. Finally, whether the domain index τ is observed depends on the identification setting, and our framework does not generally require access to domain labels.

3 LINEAR ADMG LEARNING

We now describe a general framework called Linear ADMG Learning (LinADMG). It consists of three main steps. First, using observations $\{\mathbf{X}^\tau\}_{\tau=1}^T$ generated according to $\mathbf{X}^\tau = A\mathbf{Z}^\tau$ with the shared mixing matrix $A := (I-B)^{-1}$, we apply an IRL/CRL method to estimate A and the corresponding domain-specific noise variables $\mathbf{Z}^\tau$. Second, following the postprocessing procedure of LiNGAM [Shimizu et al., 2006], we use acyclicity to resolve the permutation and scaling ambiguities of the estimated mixing matrix and then recover the directed structural coefficients B. Third, we estimate the bidirected edges by testing dependence among the estimated noise variables.

Step 1: Estimate the mixing matrix A. We apply a linear representation-learning method to the observed data across domains. When the domain labels are observed, the input is the collection $\{\mathbf{X}^\tau : \tau \in [T]\}$; when they are unobserved, samples from all domains are pooled and treated as a mixture distribution. We assume that the representation learning method used is able to identify $\mathbf{Z}^\tau$, which equivalently means identifying A up to permutation and component-wise scaling, i.e., $\widehat{A} = APD$, for some permutation matrix P and invertible diagonal matrix D.

Steps 2: Recover the directed edge coefficients B via LiNGAM postprocessing [Shimizu et al., 2006]. Since B is a coefficient matrix of a DAG, we can recover it from $\widehat{A}$ using the LiNGAM postprocessing procedure [Shimizu et al., 2006]. Inverting $\widehat{A}$ gives $\widehat{A}^{-1} = D^{-1}P^TA^{-1} = D^{-1}P^T(I-B)$. Using the fact that B is acyclic, P can be identified since there is a unique row permutation of $\widehat{A}^{-1}$ such that every diagonal entry is nonzero [Shimizu et al., 2006]. Once the permutation is chosen, the remaining scaling ambiguity D can be removed because $I-B$ has a unit diagonal. B is then estimated as

$$\widehat{B} = I - PD\widehat{A}^{-1}. \tag{5}$$

Accordingly, the estimated directed edge set is

$$\widehat{E}_{\rightarrow} := \{X^\tau_j \rightarrow X^\tau_i : \widehat{B}_{ij} \neq 0\}, \tag{6}$$

which is invariant across domain by assumption. Details of the post-processing procedure are in App. B.

Step 3: Recover the bidirected edge set $E^\tau_{\leftrightarrow}$. Given $\widehat{B}$, we can estimate the noise variables. When the domain label τ is observed, this can be done separately for each domain:

$$\widehat{\mathbf{Z}}^\tau := (I - \widehat{B})\mathbf{X}^\tau. \tag{7}$$

For each pair of recovered noise variables, we add a bidirected edge if they are dependent in domain τ:

$$\widehat{\mathbf{Z}}^\tau_i \not\perp\!\!\!\perp \widehat{\mathbf{Z}}^\tau_j \quad \Longrightarrow \quad X^\tau_i \leftrightarrow X^\tau_j \in \widehat{E}^\tau_{\leftrightarrow}. \tag{8}$$

Empirically, we can test for dependence between each pair of recovered exogenous variables with an appropriate independence test given a significance α.

We can also apply our method to settings where the domain index is unobserved, and we have a mixture of observations $\mathbf{X}$. In this case, we apply the test over $\widehat{\mathbf{Z}}$.

4 ADMG IDENTIFICATION VIA LINADMG

We next provide conditions under which the proposed framework, LinADMG, identifies the underlying ADMG. By identification of an ADMG, we mean exact recovery of both directed and bidirected edge sets.

We proceed in two stages. First, in Lemma 4.1, we show that, without imposing any additional assumption, LinADMG can identify all directed edges but a subset of the bidirected edges, because in cases of fine-tuned parameters, the effect from latent confounding may be canceled and therefore cannot be detected by our method. The proof is in App. C.1.

Lemma 4.1 (Directed-edge identification and bidirected-edge subset recovery)**.** *Consider the domain-specific ADMG $\mathcal{G}^\tau = (\mathbf{X}^\tau, E_{\rightarrow}, E^\tau_{\leftrightarrow})$ generated by the linear SCM in Sec. 2. Following the Steps in LinADMG, we have $\widehat{B} = B$, and the directed edge set is identified and no false-positive estimated bidirected edges, i.e., $\widehat{E}_{\rightarrow} = E_{\rightarrow}$, $\widehat{E}^\tau_{\leftrightarrow} \subseteq E^\tau_{\leftrightarrow}$.*

To identify all the bidirected edges, we introduce a generic no-cancellation condition ensuring that latent confounding is detectable in the noise dependencies.

Assumption 4.2 (Latent-confounding no-cancellation)**.** For every domain τ and every pair of distinct noise variables Z_i^τ, Z_j^τ, $i \neq j$,

$$X_i^\tau \leftrightarrow X_j^\tau \in E_{\leftrightarrow}^\tau \quad \Longrightarrow \quad Z_i^\tau \not\perp\!\!\!\perp Z_j^\tau. \tag{9}$$

Asm. 4.2 connects our latent confounding definition of bidirected edges to the error-dependence parameterization commonly used in linear structural equation models (SEM) with unobserved confounding. In correlated-error SEMs, Bhattacharya et al. [2021] define a bidirected edge $X_i \leftrightarrow X_j$ whenever the corresponding error covariance is nonzero, while Wang and Drton [2023] use bidirected edges to encode dependence between the corresponding error terms. In these formulations, the relation between a bidirected edge and dependence among the noise variables is therefore built directly into the model.

Different from their definition, we define $E_{\leftrightarrow}^\tau$ through the latent projection of the underlying causal DAG: $X_i^\tau \leftrightarrow X_j^\tau$ represents latent confounding between the two observed variables. After the latent variables are marginalized, however, the latent effects can cancel at the level of the noise dependence, as shown in Fig 2. (b), where the dependence between Z_1^τ and Z_2^τ vanishes. Asm. 4.2 rules out such cancellations, and thus aligns the definition of bidirected edges used in our formulation with the noise-dependence parameterization used in prior work.

Previous work commonly assumes bow-freeness with faithfulness because disentangling the causal effect from directed causal relations and a common latent confounder is challenging. We provide two examples in Fig. 2 showing that these assumptions, compared with Asm. 4.2, capture different aspects of the model and do not imply one another.

In Thm. 4.3, we show that under Asm. 4.2, LinADMG identifies the complete ADMG, with proof in App. C.2.

Theorem 4.3 (Identifiability of LinADMG)**.** *Suppose that the conditions of Lemma 4.1 hold and that Asm. 4.2 is satisfied. Then, LinADMG identifies the complete domain-specific ADMG* $\mathcal{G}^\tau = (\mathbf{X}^\tau, E_\rightarrow, E_{\leftrightarrow}^\tau)$ *for every domain* τ*, i.e.,* $\widehat{E}_\rightarrow = E_\rightarrow$ *and* $\widehat{E}_{\leftrightarrow}^\tau = E_{\leftrightarrow}^\tau$*.*

5 IMPLEMENTATION

We implement LinADMG under four identifiable data-generating settings described in Sec. 2. Each setting has a corresponding IRL/CRL method for Step 1, taking $\{\mathbf{X}^\tau\}_{\tau=1}^T$ as input and outputting $\widehat{A}$ or its inverse.

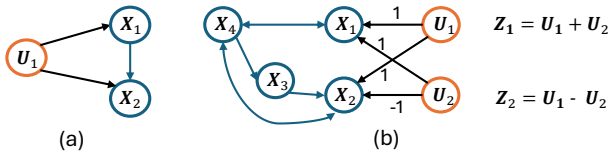

Figure 2: Examples showing that Asm. 4.2 and bow-freeness/faithfulness capture different properties. (a) X_1 and X_2 share the latent parent U_1 while $X_1 \to X_2$, resulting a violation of bow-freeness. Nevertheless, the latent confounding is detectable by the dependency of noise variables, so Asm. 4.2 holds. (b) Conversely, bow-freeness and faithfulness do not guarantee Asm. 4.2. Let U_1 and U_2 be independent Gaussian variables with equal variance. The resulting composite noise terms are $Z_1 = U_1 + U_2$ and $Z_2 = U_1 - U_2$, so $\mathrm{Cov}(Z_1, Z_2) = 0$ and hence $Z_1 \perp\!\!\!\perp Z_2$, despite the existence of X_1 and X_2 as they share latent confounders.

We provide a broader discussion of IRL/CRL methods compatible with our framework in App. D.

`LinADMG-MuDoNLL` with MD-Scalings. Our implementation, following Xu et al. [2026], first computes the empirical covariance of $\mathbf{X}^\tau$, denoted as S^τ, and define the model-implied covariance as

$$\widehat{\Sigma}^\tau = AD^\tau(I - L)^{-1}\,\mathrm{diag}(\sigma^2)(I - L)^{-T}D^\tau A^T, \tag{10}$$

where σ and a zero-diagonal latent structural matrix L is used to model the covariance structure of latent $\mathbf{Z}$. While Xu et al. [2026] choose a covariance-matching objective, we instead optimize a Gaussian negative log-likelihood (NLL) since it is a scale-adaptive measure of covariance mismatch:

$$\sum_{\tau=1}^{T}[\log\det\left(\widehat{\Sigma}^\tau + \epsilon I\right) + \mathrm{tr}\left(\left(\widehat{\Sigma}^\tau + \epsilon I\right)^{-1} S^\tau\right)], \tag{11}$$

where ϵ is a small positive value for numerical stability.

`LinADMG-diagGMMs` with DiagGMM. Kivva et al. [2022] proved that when we apply a Gaussian mixture estimator, identifiability can hold with a (piecewise-) linear decoder-encoder structure. Inspired by this, we implement the special case, linear model, by maximizing the likelihood of a T-component diagonal Gaussian mixture with weights π^τ, latent means μ^τ, latent variances v^τ, and a shared linear decoder W:

$$-\frac{1}{N}\sum_{i=1}^{N}\log\left[\sum_{\tau=1}^{T}\pi^\tau\,\mathcal{N}\left(x_i;\, W\mu^\tau,\, W\,\mathrm{diag}(v^\tau)W^T + \epsilon I\right)\right]. \tag{12}$$

We initialize the parameters using k-means on observed data and optimize them with Adam. The final estimator of W is the inverse of $\widehat{A}$.

`LinADMG-SP` **with UP-Observations.** We follow the implementation of Xu et al. [2024], using a linear autoencoder with a sparsity constraint. The encoder g and decoder f are MLPs without nonlinear activations. After training, we extract the linear part F of the affine decoder as $F_{:j} = f(\mathbf{e}_j) - f(\mathbf{0})$ where $\mathbf{e}_j$ is the jth standard basis vector. Since F maps the recovered noises to observations, we return $\widehat{A} = F$.

`LinADMG-UMNI` **with UMN-Hard Interventions.**

We follow the implementation for the case of hard-intervention from Varıcı et al. [2024], computing Gaussian score differences from sample covariances of each domain. The method recovers a linear encoder $\widehat{H}$ which maps observations to latent variables. We return $\widehat{A} = \widehat{H}^{-1}$ with Moore-Penrose pseudoinverse.

6 RELATED WORK

Classical constraint-based causal discovery methods, most notably FCI [Spirtes et al., 1995b], handle latent confounding by using conditional-independence tests. While these methods do not have parametric assumptions, and can be used with non-parametric tests, they return a Markov equivalence class containing multiple causal graphs, rather than a unique underlying graph.

A different line of work relies on additional functional, distributional, or graphical assumptions to obtain stronger identification than is possible from conditional independences alone. Focusing on linear models, prior work can be roughly divided into two groups.

The first group develops differentiable or score-based methods for learning ADMGs through continuous constrained optimization. Bhattacharya et al. [2021] consider linear Gaussian structural equation models with correlated errors and derive differentiable algebraic constraints for several restricted ADMG classes: ancestral graphs, which exclude a bidirected edge between any pair connected by a directed path; arid graphs, which exclude c-trees; and bow-free graphs, which exclude pairs connected simultaneously by a directed and a bidirected edge. Their consistency guarantee recovers the true graph up to equivalence with respect to the equality constraints implied by the observed distribution. More recently, Morinishi and Shimizu [2025] extend this approach by assuming multivariate generalized normal errors, allowing non-Gaussian information to further resolve edge orientations, and provide almost-everywhere parameter identifiability for bow-free ADMGs. However, these approaches either recover the graph only up to an observational equivalence class or rely on bow-freeness and specific error-distribution assumptions.

A second group of work leverages non-Gaussianity or independent noise variables. This line of work is inspired by LiNGAM [Shimizu et al., 2006], which assumes no latent confounding, but connects causal discovery with source separation: under linearity and independent non-Gaussian noises, recovering the ICA mixing matrix identifies the underlying DAG. This idea has been extended to time-series models [Hyvärinen et al., 2010] and nonlinear causal models [Monti et al., 2020, Reizinger et al., 2023, Montagna, 2026]. Recently, Heurtebise et al. [2026] replace non-Gaussianity with variance diversity in a multi-view setting, while still assuming independent noises.

A few recent works extend LiNGAM to handle latent confounding by leveraging overcomplete independent component analysis (OICA), beginning with Hoyer et al. [2008], which considers linear latent-variable models in which the latent variables and exogenous variables are mutually independent and non-Gaussian. Salehkaleybar et al. [2020] extend this to latent confounding by identifying causal orders among observed variables, while Maeda and Shimizu [2020] additionally infer direct causal relations between unconfounded pairs and detect pairs sharing latent confounders. These results, however, do not in general result in exact recovery of the ADMG. Cumulant-based methods [Cai et al., 2023, Schkoda et al., 2024] provide asymptotically correct recovery of causal structure under different specialized assumptions on the latent-variable model, such as at least three observed children per latent confounder and a local bound on the number of latent confounders, respectively. In addition, Wang and Drton [2023] provided full recovery of the graph under the bow-free assumption. In contrast, our method chooses a different perspective, identifiable representation learning, and does not require non-Gaussianity or the specific graphical and latent-structural restrictions required by these approaches.

7 EXPERIMENTAL RESULTS

We assess the performance of the proposed `LinADMG-` methods using synthetic data and then present an application to a semisynthetic dataset `causalAssembly`.

Synthetic Data generation. We consider four linear data-generating processes (DGPs) described in Sec. 2: MD-Scalings, UP-Observations, UMN-Hard Interventions, and Diag-GMMs. For the first three DGPs, we generate $m = \{5, 10, 20\}$ latent variables from a linear SCM, $\mathbf{U} = B_2\mathbf{U} + \mathbf{E}_2$, where the latent causal graph is sampled from an Erdős-Rényi (ER) model with $k_U \in \{0, 1, 2, 3\}$, where each ER$-k_U$ graph contains $n \cdot k_U$ directed edges. The nonzero coefficients of B_2 are uniformly sampled from $[-0.5, -0.2] \cup [0.2, 0.5]$. We sample the scales of each component

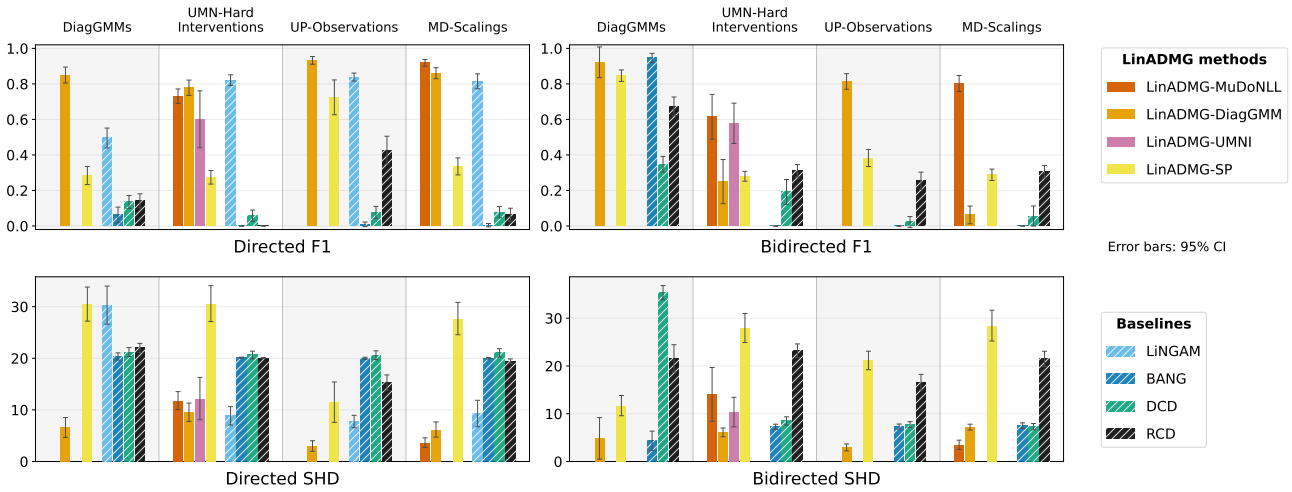

Figure 3: Estimation of the directed (left) and bidirected (right) edges of the ground-truth ADMG under four DGPs listed above. Bars report means over 20 random seeds, with 95% confidence intervals. Warm colors are the proposed `LinADMG-` methods, while hatched cold colors are baselines. The top row reports F1 score (↑), and the bottom row reports SHD (↓). LiNGAM does not infer bidirected edges and is therefore omitted.

of $\mathbf{E}_2$ independently from $\mathrm{Unif}(0.2, 0.5)$ and consider three noise distributions: Gaussian, Exponential, and Gumbel distributions. The three DGPs differ in how $\mathbf{U}$ generates $\mathbf{Z}^\tau$ with domain number $T = \{1, 2, 3\} \times m$.

For *UMN-Hard Interventions*, we generate T intervention domains and one additional observational reference domain. In each intervention domain τ, we select a target set $\mathcal{I}^\tau$ and replace the structural equation of each targeted latent variable by an independent noise term. We sample target sets such that every latent variable is targeted in at least one domain and the binary domain-by-variable target matrix has full column rank. Specifically, we construct B_2^τ by setting the rows of B_2 corresponding to $\mathcal{I}^\tau$ to zero and generate $\mathbf{Z}^\tau = (I - B_2^\tau)^{-1}\mathbf{E}_2$.

For *UP-Observations*, we generate a domain-specific binary mask: $M^\tau \in \{0, 1\}^m$ with half active components. If a component is inactive in every masked domain, we activate it in a randomly selected domain. We additionally generate a reference domain with $M^\tau = I$, giving $T + 1$ domains in total. Then apply it to $\mathbf{U}$ to get $\mathbf{Z}^\tau = M^\tau \mathbf{U}$.

For *MD-Scalings*, we generate the domain-specific scaling matrix as $D^\tau = \mathrm{diag}(d_1^\tau, \ldots, d_n^\tau)$, where the diagonal entries are sampled independently from $\mathrm{Unif}(0.2, 1)$. We additionally generate an unscaled reference domain with $D^\tau = I$, giving $T + 1$ domains in total. Then apply it to $\mathbf{U}$ to get $\mathbf{Z}^\tau = D^\tau \mathbf{U}$.

For *Diag-GMMs*, we generate five domains indexed by an observed auxiliary variable $a^\tau \in \{-2, -1, 0, 1, 2\}$. Conditional on a^τ, we directly sample the composite noise from $\mathbf{Z}^\tau \mid a^\tau \sim \mathcal{N}\big(\mu(a^\tau), \Sigma(a^\tau)\big)$, where $\mu(a) = \alpha + a\beta$ and $\Sigma(a) = (L_0 + aL_1)(L_0 + aL_1)^T + 10^{-6}I$. α and β are sampled independently from $\mathcal{N}(0, 0.5^2)$. The matrices L_0 and L_1 are diagonal, and sampled independently from $\mathrm{Unif}(0.2, 0.5)$ and $\mathcal{N}(0, 0.08^2)$, respectively. Thus, though $\mathbf{Z}^\tau$ are independent conditional on a^τ, pooling the domains can induce dependence through the shared auxiliary variable.

For all four DGPs, we consider $n = m$ observed variables $\mathbf{X}$. We sample the observed causal graph $\mathcal{D}$ from an ER model with $k_X \in \{0, 1, 2, 3\}$, where each ER$-k_X$ graph contains $n \cdot k_X$ directed edges. Given $\mathcal{D}$, we sample the linear edge weight B uniformly from $[-0.5, -0.2] \cup [0.2, 0.5]$. For each domain τ, we generate the observed variable $\mathbf{X}^\tau = (I - B)^{-1}\mathbf{Z}^\tau$ with sample size $N \in \{1000, 5000, 50000\}$. We run each setup over 20 random seeds.

Baselines and Evaluation Metrics. We compare our methods with DirectLiNGAM [Shimizu et al., 2006], BANG [Wang and Drton, 2023], DCD [Bhattacharya et al., 2021], and RCD [Maeda and Shimizu, 2020]. We pool observations generated via scaling and hard intervention DGPs since the baselines do not require the domain index. Due to computational constraints, we fit DCD using at most 10000 pooled observations, while RCD uses at most 300 observations. We evaluate directed and bidirected edges separately using F1 score and structural Hamming distance (SHD). Higher F1 and lower SHD mean better estimation.

Synthetic Experiments Results. Fig. 3 reports the accuracy with 95% confidence intervals of directed- and bidirected-edge estimation across four DGPs with

$n = m = 10$, $T = n$, $k_X = 2$, $k_U = 0.5$, $N = 50000$ for UMN-Hard Interventions dataset; otherwise, $N = 5000$, and Gaussian noise. Overall, the proposed methods achieve the strongest identification, particularly `LinADMG-diagGMM` and `LinADMG-MuDoNLL` when their assumptions match the underlying domain-variation mechanism. Under UMN-Hard Interventions DGP, DirectLiNGAM remains competitive for directed-edge recovery. However, it does not support latent confounding estimation. We evaluate `LinADMG-UMNI` only under the UMN-Hard Interventions DGP because the UMNI representation-learning algorithm strictly requires hard interventions and a full-column-rank intervention design, which the other DGPs do not satisfy. Additional results in App. F.2 show that its reliable performance also requires a sufficiently large sample size within each domain. Finally, `LinADMG-SP` relies on structural sparsity and is therefore competitive under its matched dataset, but its recovery performance becomes worse when the domain variation does not produce the sparsity pattern assumed by the method.

In App. F.1, we show near-perfect directed-edge weight estimation for `LinADMG-MuDoNLL` under MD-Scalings. In App. F.2, we report the accuracy of directed- and bidirected-edge estimation across different configurations of n, sample numbers N, domain numbers T, graph density k_X, confounding density k_U, and noise distribution for all four DGPs. We also provide a comparison between the loss function proposed in Sec. 5 and those from the original papers, including comparing diagGMM (ours) with VaDE (used in Kivva et al. [2022]) in App. F.3, and comparing MuDoNLL (ours) with MuDoCov (proposed by Xu et al. [2026]) in App. F.4, where our methods consistently outperform.

We also compare our method with FCI [Spirtes et al., 1995a] in App. F.5. We first convert the estimated ADMG into PAG and then compute F1 for the skeleton, arrowhead, and tail. Our method significantly improves recovery of the correct tail and arrowhead, and shows comparable performance for skeleton accuracy. We also test our methods and baselines with bow-free underlying ADMG and non-Gaussian noise. The results in App. F.6 show that our methods maintain stable and robust performance, while the baselines show significant improvement due to the match of their methodological assumptions.

`causalAssembly` with latent confounding. We use the causalAssembly benchmark [Göbler et al., 2024], a semi-synthetic benchmark derived from a real manufacturing assembly line, focusing on Stations 1 and 2. We first treat all $n_1 = 6$ Station 1 variables as latent and all $n_2 = 34$ Station 2 variables as observed. We generate 18 domains with 5000 samples each: one observational domain and 17 soft-intervention domains, each targeting 1-3 latent variables. We evaluate `LinADMG-diagGMM` and `LinADMG-MuDoNLL` over five optimization seeds. Both perform poorly in this setting, which we attribute partly to limited latent variability, since only one directed edge connects Station 1 to Station 2. We therefore also treat the first half of Station 2 variables in topological order as latent, leaving the remaining half observed. The resulting ADMG contains 25 directed and 54 bidirected edges. This results in a modest improvement. A major reason for the difficulty is the strong nonlinearity of the causalAssembly mechanisms, as Göbler et al. [2024] reports similarly poor performance of linear causal discovery methods on this benchmark. To separate nonlinear model misspecification from the latent-confounding challenge, we construct a linear version that preserves the graph topology but replaces the mechanisms with a linear SCM, with edge coefficients sampled uniformly from $[0.4, 0.8]$. This results in substantially higher accuracy for both methods. We summarize the results in Tab. 1.

Table 1: Average F1 scores ± standard deviation over 5 random seeds across different versions of causalAssembly with latent confounding with linear/nonlinear SCM and latent confounding settings.

SCM	Latent	Method	Directed F1	Bidirected F1
Nonlinear	n_1	`-diagGMM`	$.020 \pm .000$	$.004 \pm .000$
		`-MuDoNLL`	$.074 \pm .000$	$.004 \pm .000$
Nonlinear	n_1 + 50% n_2	`-diagGMM`	$.000 \pm .000$	$.554 \pm .000$
		`-MuDoNLL`	$.085 \pm .000$	$.576 \pm .000$
Linear	n_1 + 50% n_2	`-diagGMM`	$.508 \pm .020$	$.800 \pm .009$
		`-MuDoNLL`	$.147 \pm .031$	$.554 \pm .024$

8 CONCLUSION AND LIMITATION

We introduced LinADMG, a general framework for learning linear ADMGs through identifiable representation learning. We show that LinADMG can recover the directed causal structure whenever the noise variables is identifiable up to permutation and scaling; it can further identify the full ADMG under a latent-confounding no-cancellation condition. Our results show how noise heterogeneity can provide an alternative source of identifiability for causal discovery with latent confounding.

The current framework is limited to linear structural causal models. Extending the framework to nonlinear causal structural models is an important direction for future work. In addition, we currently consider only a selected set of identifiable representation-learning settings. Extending the framework to incorporate other sources of identifiability, such as temporal structure, could further broaden its applicability.

ACKNOWLEDGEMENTS

The research of DX and SM was supported by the Air Force Office of Scientific Research under award number FA8655-22-1-7155. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the United States Air Force. We also thank SURF for the support in using the Dutch National Supercomputer Snellius.

AI use statement

In this work, we used generative AI tools for generating synthetic data sets, polishing the theorem claim, assisting in the writing of proofs, implementing methods, cleaning and reformatting datasets, and interpreting results. We have not used generative AI tools for helping develop theoretical models or conceptual frameworks, formulating mathematical claims, providing critical ingredients for proving mathematical claims, designing or providing feedback on research methodology or experiments, and proposing or refining hypotheses, assisting with translation, and supporting qualitative and thematic data analysis are not applicable to this work. Additionally, we used generative AI tools for creating or modifying scientific figures or images, summarizing or analyzing existing literature, editing a research paper to improve readability, and identifying relevant literature. We have reviewed all AI-assisted work.

We take responsibility for the final content of this work, including text, claims, or artifacts produced with the aid of generative AI.

REFERENCES

K. Ahuja, D. Mahajan, Y. Wang, and Y. Bengio. Interventional causal representation learning. In *International conference on machine learning*, pages 372–407. PMLR, 2023.

K. Ahuja, A. Mansouri, and Y. Wang. Multi-domain causal representation learning via weak distributional invariances. In *International Conference on Artificial Intelligence and Statistics*, pages 865–873. PMLR, 2024.

R. Bhattacharya, T. Nagarajan, D. Malinsky, and I. Shpitser. Differentiable causal discovery under unmeasured confounding. In *International Conference on Artificial Intelligence and Statistics*, pages 2314–2322. PMLR, 2021.

J. Brehmer, P. De Haan, P. Lippe, and T. S. Cohen. Weakly supervised causal representation learning. *Advances in Neural Information Processing Systems*, 35:38319–38331, 2022.

S. Buchholz, G. Rajendran, E. Rosenfeld, B. Aragam, B. Schölkopf, and P. Ravikumar. Learning linear causal representations from interventions under general nonlinear mixing. *Advances in Neural Information Processing Systems*, 36:45419–45462, 2023.

R. Cai, Z. Huang, W. Chen, Z. Hao, and K. Zhang. Causal discovery with latent confounders based on higher-order cumulants. In *International conference on machine learning*, pages 3380–3407. PMLR, 2023.

P. Comon. Independent component analysis, a new concept? *Signal processing*, 36(3):287–314, 1994.

X. Dong, I. Ng, B. Huang, Y. Sun, S. Jin, R. Legaspi, P. Spirtes, and K. Zhang. On the parameter identifiability of partially observed linear causal models. *Advances in Neural Information Processing Systems*, 37:30740–30771, 2024.

J. Gallego-Posada, J. Ramirez, M. Hashemizadeh, and S. Lacoste-Julien. Cooper: A library for constrained optimization in deep learning. *arXiv preprint arXiv:2504.01212*, 2025.

K. Göbler, T. Windisch, M. Drton, T. Pychynski, M. Roth, and S. Sonntag. causalassembly: Generating realistic production data for benchmarking causal discovery. In *Causal Learning and Reasoning*, pages 609–642. PMLR, 2024.

G. Gómez-Herrero, M. Atienza, K. Egiazarian, and J. L. Cantero. Measuring directional coupling between eeg sources. *Neuroimage*, 43(3):497–508, 2008.

A. Heurtebise, O. Chehab, P. Ablin, A. Gramfort, and A. Hyvarinen. Multi-view causal discovery without non-gaussianity: Identifiability and algorithms. In *Forty-third International Conference on Machine Learning*, 2026.

P. O. Hoyer, S. Shimizu, A. J. Kerminen, and M. Palviainen. Estimation of causal effects using linear non-gaussian causal models with hidden variables. *International Journal of Approximate Reasoning*, 49(2): 362–378, 2008.

A. Hyvärinen and J. Hurri. Blind separation of sources that have spatiotemporal variance dependencies. *Signal processing*, 84(2):247–254, 2004.

A. Hyvarinen and H. Morioka. Nonlinear ica of temporally dependent stationary sources. In *Artificial intelligence and statistics*, pages 460–469. PMLR, 2017.

A. Hyvärinen and E. Oja. Independent component analysis: Algorithms and applications. *Neural Networks*, 13(4-5):411–430, 2000.

A. Hyvärinen and P. Pajunen. Nonlinear independent component analysis: Existence and uniqueness results. *Neural Networks*, 12(3):429–439, 1999.

A. Hyvärinen, K. Zhang, S. Shimizu, and P. O. Hoyer. Estimation of a structural vector autoregression model using non-gaussianity. *Journal of Machine Learning Research*, 11(5), 2010.

A. Hyvarinen, H. Sasaki, and R. Turner. Nonlinear ica using auxiliary variables and generalized contrastive learning. In *The 22nd international conference on artificial intelligence and statistics*, pages 859–868. Pmlr, 2019.

I. Khemakhem, D. Kingma, R. Monti, and A. Hyvarinen. Variational autoencoders and nonlinear ica: A unifying framework. In *International conference on artificial intelligence and statistics*, pages 2207–2217. PMLR, 2020.

D. P. Kingma and J. Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations*, 2015.

B. Kivva, G. Rajendran, P. Ravikumar, and B. Aragam. Identifiability of deep generative models without auxiliary information. *Advances in Neural Information Processing Systems*, 35:15687–15701, 2022.

L. Kong, B. Huang, F. Xie, E. Xing, Y. Chi, and K. Zhang. Identification of nonlinear latent hierarchical models. *Advances in Neural Information Processing Systems*, 36:2010–2032, 2023.

L. Kong, G. Chen, B. Huang, E. Xing, Y. Chi, and K. Zhang. Learning discrete concepts in latent hierarchical models. *Advances in Neural Information Processing Systems*, 37:36938–36975, 2024.

S. Lachapelle, P. Rodriguez, Y. Sharma, K. E. Everett, R. Le Priol, A. Lacoste, and S. Lacoste-Julien. Disentanglement via mechanism sparsity regularization: A new principle for nonlinear ICA. In *Conference on Causal Learning and Reasoning*, pages 428–484. PMLR, 2022.

S. Lachapelle, P. R. López, Y. Sharma, K. Everett, R. Le Priol, A. Lacoste, and S. Lacoste-Julien. Nonparametric partial disentanglement via mechanism sparsity: Sparse actions, interventions and sparse temporal dependencies. *Journal of Machine Learning Research*, 27(71):1–90, 2026.

Z. Li, M. Fu, J. Huang, Y. Shen, R. Cai, Y. Sun, G. Chen, and K. Zhang. Towards identifiability of hierarchical temporal causal representation learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2025.

P. Lippe, S. Magliacane, S. Löwe, Y. M. Asano, T. Cohen, and S. Gavves. Citris: Causal identifiability from temporal intervened sequences. In *International Conference on Machine Learning*, pages 13557–13603. PMLR, 2022.

P. Lippe, S. Magliacane, S. Löwe, Y. M. Asano, T. Cohen, and E. Gavves. Biscuit: Causal representation learning from binary interactions. In *Uncertainty in Artificial Intelligence*, pages 1263–1273. PMLR, 2023.

F. Locatello, S. Bauer, M. Lucic, G. Raetsch, S. Gelly, B. Schölkopf, and O. Bachem. Challenging common assumptions in the unsupervised learning of disentangled representations. In *international conference on machine learning*, pages 4114–4124. PMLR, 2019.

T. N. Maeda and S. Shimizu. Rcd: Repetitive causal discovery of linear non-gaussian acyclic models with latent confounders. In *International Conference on artificial intelligence and statistics*, pages 735–745. PMLR, 2020.

F. Montagna. On the identifiability of causal graphs with multiple environments. In *International Conference on Learning Representations*, volume 2026, pages 13837–13868, 2026.

R. P. Monti, K. Zhang, and A. Hyvärinen. Causal discovery with general non-linear relationships using non-linear ica. In *Uncertainty in artificial intelligence*, pages 186–195. PMLR, 2020.

Y. Morinishi and S. Shimizu. Differentiable causal discovery of linear non-gaussian acyclic models under unmeasured confounding. *Transactions on Machine Learning Research*, 2025. ISSN 2835-8856.

H. Morioka and A. Hyvarinen. Causal representation learning made identifiable by grouping of observational variables. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235, pages 36249–36293. PMLR, 2024.

J. Pearl et al. Causality: models, reasoning, and inference. *Econometric Theory*, 19(675-685):46, 2003.

J. Peters, D. Janzing, and B. Schölkopf. *Elements of causal inference: foundations and learning algorithms*. The MIT press, 2017.

P. Reizinger, Y. Sharma, M. Bethge, B. Schölkopf, F. Huszár, and W. Brendel. Jacobian-based causal discovery with nonlinear ica. *Transactions on Machine Learning Research*, 2023.

T. Richardson. Markov properties for acyclic directed mixed graphs. *Scandinavian Journal of Statistics*, 30(1):145–157, 2003.

S. Salehkaleybar, A. Ghassami, N. Kiyavash, and K. Zhang. Learning linear non-gaussian causal models in the presence of latent variables. *Journal of Machine Learning Research*, 21(39):1–24, 2020.

D. Schkoda, E. Robeva, and M. Drton. Causal discovery of linear non-gaussian causal models with unobserved confounding. *arXiv preprint arXiv:2408.04907*, 2024.

B. Schölkopf, F. Locatello, S. Bauer, N. R. Ke, N. Kalchbrenner, A. Goyal, and Y. Bengio. Toward causal representation learning. *Proceedings of the IEEE*, 109(5):612–634, 2021.

S. Shimizu, P. O. Hoyer, A. Hyvärinen, A. Kerminen, and M. Jordan. A linear non-gaussian acyclic model for causal discovery. *Journal of Machine Learning Research*, 7(10), 2006.

X. Song, Z. Li, G. Chen, Y. Zheng, Y. Fan, X. Dong, and K. Zhang. Causal temporal representation learning with nonstationary sparse transition. *Advances in Neural Information Processing Systems*, 37:77098–77131, 2024.

P. Spirtes, C. Meek, and T. Richardson. Causal inference in the presence of latent variables and selection bias. In *Proceedings of the Eleventh Conference on Uncertainty in Artificial Intelligence*, UAI'95, page 499–506, 1995a.

P. Spirtes, C. Meek, and T. Richardson. Causal inference in the presence of latent variables and selection bias. In *Proceedings of the Eleventh Conference on Uncertainty in Artificial Intelligence*, page 499–506, 1995b.

P. Spirtes, C. N. Glymour, and R. Scheines. *Causation, Prediction, and Search*. MIT Press, 2001.

B. Varıcı, E. Acartürk, K. Shanmugam, and A. Tajer. Linear causal representation learning from unknown multi-node interventions. *Advances in Neural Information Processing Systems*, 37:111614–111648, 2024.

J. von Kügelgen, M. Besserve, L. Wendong, L. Gresele, A. Kekić, E. Bareinboim, D. Blei, and B. Schölkopf. Nonparametric identifiability of causal representations from unknown interventions. *Advances in Neural Information Processing Systems*, 36:48603–48638, 2023.

Y. S. Wang and M. Drton. Causal discovery with unobserved confounding and non-gaussian data. *Journal of Machine Learning Research*, 24(271):1–61, 2023.

D. Xu, D. Yao, S. Lachapelle, P. Taslakian, J. Von Kügelgen, F. Locatello, and S. Magliacane. A sparsity principle for partially observable causal representation learning. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235, pages 55389–55433, 2024.

D. Xu, L. Parkkonen, S. Magliacane, and A. Hyvarinen. Identifying dependent components from multi-domain linear mixtures. In *Forty-third International Conference on Machine Learning*, 2026.

D. Yao, D. Xu, S. Lachapelle, S. Magliacane, P. Taslakian, G. Martius, J. von Kügelgen, and F. Locatello. Multi-view causal representation learning with partial observability. *International Conference on Learning Representations*, 2024.

J. Zhang, K. Greenewald, C. Squires, A. Srivastava, K. Shanmugam, and C. Uhler. Identifiability guarantees for causal disentanglement from soft interventions. *Advances in Neural Information Processing Systems*, 36:50254–50292, 2023.

K. Zhang and A. Hyvärinen. A general linear non-gaussian state-space model: Identifiability, identification, and applications. In *Asian Conference on Machine Learning*, pages 113–128. PMLR, 2011.

K. Zhang, S. Xie, I. Ng, and Y. Zheng. Causal representation learning from multiple distributions: A general setting. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235, pages 60057–60075. PMLR, 2024.

X. Zheng, B. Aragam, P. K. Ravikumar, and E. P. Xing. Dags with no tears: Continuous optimization for structure learning. *Advances in neural information processing systems*, 31, 2018.

Causal Discovery for Linear Models with Latent Confounding via Multi-domain Identifiable Representation Learning: Supplementary Materials

A SUMMARY OF NOTATION

Throughout this paper, we adopt the following conventions for mathematical notation: We use bold uppercase letters such as $\mathbf{X}$ to denote random vectors, and uppercase italic letters such as X to denote matrices. We list all specific notations in the following table:

Symbol	Description
τ	Domain index
$\mathbf{U}$	Unobserved variables
$\mathbf{X}$	Observed variables
$\mathbf{E}_1$	Exogenous variables of $\mathbf{X}$
$\mathbf{E}_2$	Exogenous variables of $\mathbf{U}$
$\mathbf{Z}^\tau$	Latent noise variables from domain τ
$\mathbf{X}^\tau$	Observed variables from domain τ
A	Linear mixing matrix
B	Weighted adjacency matrix of directed causal graph among observed variables
$\mathcal{G}$	ADMG
$E_{\rightarrow}$	directed edge set
$E_{\leftrightarrow}$	bidirected edge set
n	Dimension of the observed variables
m	Dimension of the unobserved variables
T	Number of domains
N	Number of samples per domain

B DETAILS OF LINGAM-BASED DIRECTED STRUCTURE RECOVERY

In this section, we provide the details of the LiNGAM-style post-processing procedure used to recover the directed structure from the estimated mixing matrix $\hat{A}$. We break it into 3 steps to resolve the permutation, scaling, and numerical error.

We start by introducing new notation. Let $W := A^{-1} = I - B$. Inverting $\widehat{A}$ gives

$$\widetilde{W} := \widehat{A}^{-1} = D^{-1}P^T W. \tag{13}$$

Hence, $\widetilde{W}$ recovers the structural demixing matrix W up to a row permutation and row-wise scaling.

Step i: Resolve the permutation ambiguity. By assuming B is a DAG weighted adjacency matrix, we have $B_{ii} = 0$, which implies $W = I - B$ must have exactly one diagonal entries, i.e. $W_{ii} = 0$. We therefore search a row permutation of $\widetilde{W}$ that aligns the recovered components with the observed variables.

Specifically, we choose the permutation $\widehat{P}$ solving

$$\widehat{P} \in \operatorname*{argmax}_{P \in \mathcal{P}_n} \sum_{i=1}^{n} \left| (P\widetilde{W})_{ii} \right|, \tag{14}$$

where $\mathcal{P}_n$ denotes the set of $n \times n$ permutation matrices. This optimization is a linear assignment problem and can be solved efficiently using, for example, the Hungarian algorithm with complexity $O(n^3)$. We then define

$$\overline{W} := \widehat{P}\widetilde{W}. \tag{15}$$

Step ii: Remove the scaling ambiguity. After permutation alignment, the remaining ambiguity is row-wise scaling. Since $W = I - B$ satisfies $W_{ii} = 1$, we normalize each row by its diagonal entry:

$$\widehat{W}_{ij} := \frac{\overline{W}_{ij}}{\overline{W}_{ii}}, \qquad i, j \in [n]. \tag{16}$$

Step iii: Recover the directed structural coefficients. We obtain the estimated weighted adjacency matrix as

$$\widetilde{B} := I - \widehat{W}. \tag{17}$$

Sparsification. However, finite-sample estimation generally produces small spurious nonzero coefficients. We therefore sparsify $\widetilde{B}$ using either an absolute threshold,

$$\widetilde{B}_{ij} \leftarrow 0 \quad \text{if} \quad |\widetilde{B}_{ij}| < \lambda, \tag{18}$$

or a relative threshold defined with respect to the magnitude of the recovered coefficients.

Acyclicity. Thresholding alone need not yield an acyclic directed graph. Following Shimizu et al. [2006], if directed cycles remain, we iteratively remove the weakest edge participating in a directed cycle until it becomes acyclic. Denote the resulting matrix by $\widehat{B}$. The estimated directed edge set is then

$$\widehat{E}_{\to} := \{X_j \to X_i : \widehat{B}_{ij} \neq 0\}. \tag{19}$$

C PROOF

C.1 Proof of Lemma 4.1

Before proving Lemma 4.1, we recall the following result underlying the LiNGAM postprocessing procedure.

Lemma C.1 (LiNGAM permutation lemma Shimizu et al., 2006, Lemma 1)**.** *Let M be a lower-triangular matrix whose diagonal entries are all nonzero. Among matrices obtained from M by independently permuting its rows and columns, the diagonal contains no zero entries if and only if the row and column permutations coincide.*

Consequently, if a matrix W can be made lower triangular with nonzero diagonal by applying the same permutation to its rows and columns, then, for any row-permuted version of W, there exists a unique row permutation that restores a matrix with an everywhere nonzero diagonal.

The second statement is the form of Lemma C.1 used in the LiNGAM postprocessing procedure. After the correct row permutation is identified, row-wise scaling is removed by normalizing each row by its diagonal entry.

Lemma 4.1 (Directed-edge identification and bidirected-edge subset recovery)**.** *Consider the domain-specific ADMG $\mathcal{G}^\tau = (\mathbf{X}^\tau, E_{\to}, E^\tau_{\leftrightarrow})$ generated by the linear SCM in Sec. 2. Following the Steps in LinADMG, we have $\widehat{B} = B$, and the directed edge set is identified and no false-positive estimated bidirected edges, i.e., $\widehat{E}_{\to} = E_{\to}$, $\widehat{E}^\tau_{\leftrightarrow} \subseteq E^\tau_{\leftrightarrow}$.*

Proof. From $\widehat{A} = APD$, and $A^{-1} = I - B$, we obtain $\widehat{A}^{-1} = D^{-1}P^T(I - B)$.

Thus, $\widehat{A}^{-1}$ differs from $I - B$ only by a row permutation and nonzero row-wise scaling.

Since B represents a DAG, there exists a permutation of the variables under which B is strictly lower triangular. Hence, under the same permutation, $I - B$ is lower triangular with all diagonal entries equal to one. By Lemma C.1, there is a unique row permutation of $\widehat{A}^{-1}$ whose diagonal entries are all nonzero. Therefore, the

LiNGAM permutation step recovers the correct correspondence between the rows of $\widehat{A}^{-1}$ and the observed variables.

After correcting the row permutation, the resulting matrix differs from $I - B$ only by row-wise scaling. Since the diagonal of $I - B$ is equal to one, normalizing each row by its diagonal entry removes this scaling ambiguity exactly. Therefore, the normalized demixing matrix satisfies $\widehat{W} = I - B$, and consequently

$$\widehat{B} = I - \widehat{W} = B. \tag{20}$$

It follows immediately that

$$\widehat{E}_{\to} = E_{\to}. \tag{21}$$

Next, we show that $\widehat{E}^{\tau}_{\leftrightarrow} \subseteq E^{\tau}_{\leftrightarrow}$. We prove this by contradiction. Suppose

$$\exists \quad X_i^{\tau} \leftrightarrow X_j^{\tau} \in \widehat{E}^{\tau}_{\leftrightarrow}, \text{ such that } X_i^{\tau} \leftrightarrow X_j^{\tau} \notin E^{\tau}_{\leftrightarrow}. \tag{22}$$

Since the bidirected-edge recovery procedure adds an edge only when the corresponding estimated disturbance components are dependent, we must have

$$\widehat{Z}_i^{\tau} \not\perp\!\!\!\perp \widehat{Z}_j^{\tau}. \tag{23}$$

From the previous result, we know $\widehat{B} = B$, and therefore

$$\widehat{\mathbf{Z}}^{\tau} = (I - \widehat{B})\mathbf{X}^{\tau} = (I - B)\mathbf{X}^{\tau} = \mathbf{Z}^{\tau}. \tag{24}$$

Thus we have,

$$Z_i^{\tau} \not\perp\!\!\!\perp Z_j^{\tau}. \tag{25}$$

To obtain a contradiction, recall from Eq. 2 that

$$\mathbf{Z}^{\tau} = B_1^{\tau}(I - B_2^{\tau})^{-1}\mathbf{E}_2^{\tau} + \mathbf{E}_1^{\tau}. \tag{26}$$

Define $M^{\tau} := B_1^{\tau}(I - B_2^{\tau})^{-1}$. We write it out component by component:

$$Z_i^{\tau} = \sum_k M_{ik}^{\tau}\mathbf{E}_{2k}^{\tau} + \mathbf{E}_{1i}^{\tau}, \qquad Z_j^{\tau} = \sum_k M_{jk}^{\tau}\mathbf{E}_{2k}^{\tau} + \mathbf{E}_{1j}^{\tau}. \tag{27}$$

M_{ik}^{τ} represents the total linear effect of the latent exogenous component $\mathbf{E}_{2k}^{\tau}$ on the observed disturbance Z_i^{τ}. Thus, a common latent source can induce dependence between Z_i^{τ} and Z_j^{τ} only if there exists some k such that

$$M_{ik}^{\tau} \neq 0 \qquad \text{and} \qquad M_{jk}^{\tau} \neq 0. \tag{28}$$

By construction of the domain-specific ADMG, such a shared latent source induces the bidirected edge $X_i \leftrightarrow X_j \in E^{\tau}_{\leftrightarrow}$, which is contradictory with $X_i \leftrightarrow X_j \notin E^{\tau}_{\leftrightarrow}$ in Equ. 22. Therefore, we conclude

$$\widehat{E}^{\tau}_{\leftrightarrow} \subseteq E^{\tau}_{\leftrightarrow}. \tag{29}$$

□

C.2 Proof of Theorem 4.3

Theorem 4.3 (Identifiability of LinADMG)**.** *Suppose that the conditions of Lemma 4.1 hold and that Asm. 4.2 is satisfied. Then, LinADMG identifies the complete domain-specific ADMG $\mathcal{G}^{\tau} = (\mathbf{X}^{\tau}, E_{\to}, E^{\tau}_{\leftrightarrow})$ for every domain τ, i.e., $\widehat{E}_{\to} = E_{\to}$ and $\widehat{E}^{\tau}_{\leftrightarrow} = E^{\tau}_{\leftrightarrow}$.*

Proof. By Lemma 4.1, the directed structural coefficient matrix is identified:

$$\widehat{B} = B. \tag{30}$$

Therefore, $\widehat{E}_{\to} = E_{\to}$

Moreover, Lemma 4.1 shows that the estimated bidirected edge set contains no false-positive edges:

$$\widehat{E}^{\tau}_{\leftrightarrow} \subseteq E^{\tau}_{\leftrightarrow}. \tag{31}$$

Therefore, we only need to show the other way around: $E^{\tau}_{\leftrightarrow} \subseteq \widehat{E}^{\tau}_{\leftrightarrow}$.

Since $\widehat{B} = B$, the composite disturbances are also recovered exactly:

$$\widehat{\mathbf{Z}}^{\tau} = (I - \widehat{B})\mathbf{X}^{\tau} = (I - B)\mathbf{X}^{\tau} = \mathbf{Z}^{\tau}. \tag{32}$$

Now consider an arbitrary true bidirected edge $X^{\tau}_i \leftrightarrow X^{\tau}_j \in E^{\tau}_{\leftrightarrow}$. By Assump. 4.2, this implies

$$Z^{\tau}_i \not\perp\!\!\!\perp Z^{\tau}_j. \tag{33}$$

By Equ. (32), we then have $\widehat{Z}^{\tau}_i = Z^{\tau}_i$ and $\widehat{Z}^{\tau}_j = Z^{\tau}_j$, and therefore $\widehat{Z}^{\tau}_i \not\perp\!\!\!\perp \widehat{Z}^{\tau}_j$.

Thus, under the dependence criterion used by LinADMG, $X^{\tau}_i \leftrightarrow X^{\tau}_j \in \widehat{E}^{\tau}_{\leftrightarrow}$.

Since this is true for all $X^{\tau}_i \leftrightarrow X^{\tau}_j \in E^{\tau}_{\leftrightarrow}$, we obtain

$$E^{\tau}_{\leftrightarrow} \subseteq \widehat{E}^{\tau}_{\leftrightarrow}. \tag{34}$$

Combining Equs. (31) and (34), we can conclude

$$\widehat{E}^{\tau}_{\leftrightarrow} = E^{\tau}_{\leftrightarrow}. \tag{35}$$

□

D IDENTIFIABLE AND CAUSAL REPRESENTATION LEARNING.

A broad literature studies conditions under which latent representations become identifiable from observed data. In blind source separation, source independence can sometimes be relaxed by assuming additional temporal structure [Hyvärinen and Hurri, 2004, Gómez-Herrero et al., 2008, Zhang and Hyvärinen, 2011]. Nonlinear ICA further uses auxiliary variables, temporal structure, or nonstationarity to identify otherwise non-identifiable latent components [Hyvarinen and Morioka, 2017, Hyvarinen et al., 2019, Khemakhem et al., 2020, Kivva et al., 2022]. More recent causal representation learning (CRL) methods leverage additional information such as interventions [Brehmer et al., 2022, Lippe et al., 2022, von Kügelgen et al., 2023, Buchholz et al., 2023, Zhang et al., 2023, Ahuja et al., 2023], temporal structure [Lippe et al., 2023, Lachapelle et al., 2022, 2026, Li et al., 2025], multi-view observations [Yao et al., 2024, Ahuja et al., 2024, Morioka and Hyvarinen, 2024], hierarchical structure [Kong et al., 2023, 2024], or restrictions on the latent causal graph [Zhang et al., 2024, Dong et al., 2024, Song et al., 2024]. Particularly relevant to our setting are linear representation-learning methods based on hard interventions [Varıcı et al., 2024], multi-domain scaling [Xu et al., 2026], and sparse latent support [Xu et al., 2024].

These works primarily aim to identify latent representations and, in some settings, causal relations among the latent variables. LinADMG only uses representation identifiability as an intermediate step toward recovering causal structure among the observed variables. Among them, Multi-view cannot be converted to ADMG learning under the LinADMG framework. Different views generally have distinct mixing functions, and the recovered representations may be identifiable only up to view-specific permutation and component-wise scaling ambiguities. As a result, there need not exist a single shared mixing matrix corresponding to an invariant directed structure, as required by our framework.

E IMPLEMENTATION DETAILS AND LICENSES

This section provides further details about the experiment implementation in Section 7. Our synthetic-data generation utilities build on the open-source NOTEARS implementation by Zheng et al. [2018], distributed under

Table 2: Hyperparameters used in the experiments for all methods and baselines.

Method	Hyperparameters
`LinADMG-SP`	`lr`=10^{-4}; `num_steps`=15000; `num_initializations`=1; `batch_size`=5000; `sparse_level`=0.01; `aug_lag_coef`=0
`LinADMG-MuDoNLL`	`lr`=0.01; `num_steps`=5000; `num_initializations`=1; `scheduler`=`cosine`; `grad_clip`=1;
`LinADMG-UMNI`	`atol_eigv`=0.05; `atol_ci_test`=0.05; `kappa`=1; `umni_batch_size`=2048;
`LinADMG-diagGMM`	`lr`=0.001; `num_steps`=5000; `num_initializations`=1; `batch_size`=6144;
LiNGAM	`lingam_measure`=`pwling`;
BANG	`bang_moment_degree`=3; `bang_level`=0.01; `bang_restriction`=1; `bang_max_iterations`=10000; `bang_condition_limit`=10^{12};
DCD	`dcd_lambda`=0.05; `dcd_threshold`=0.05; `dcd_num_restarts`=5;
RCD	`rcd_max_samples`=300; `rcd_cor_alpha`=0.01; `rcd_ind_alpha`=0.01; `rcd_shapiro_alpha`=0.01; `rcd_independence`=`hsic`;

the Apache License 2.0. We use these utilities to sample random causal graphs, structural coefficients, and the latent variables used by our synthetic data generators. Our implementation of UMNI adapts from the repository released by Varıcı et al. [2024], which is distributed under the Apache License 2.0. Our BANG implementation is a Python port of the R implementation from Wang and Drton [2023], which is distributed under the MIT License. The DirectLiNGAM and RCD baselines are evaluated using the implementations provided by the `lingam` Python package [Shimizu et al., 2006, Maeda and Shimizu, 2020], which is distributed under the MIT License. The FCI baseline and the conversion of fitted ADMGs to PAGs use the `causal-learn` package, distributed under the MIT License. The implementation released by Xu et al. [2024] uses the constrained-optimization library `Cooper` [Gallego-Posada et al., 2025], also distributed under the MIT License. For CausalAssembly, we use version 1.2.1 of the `causalAssembly` package, distributed under the GNU Affero General Public License v3.0.

We list all the hyperparameters we used for training in Tab. 2, and provide all the code for our method. -MuDonll, -diagGMM, and VaDE-linear use Adam [Kingma and Ba, 2015]; -SP uses Cooper's ExtraAdam for primal and dual updates; DCD uses L-BFGS-B within an augmented-Lagrangian procedure. Experiments were run on CPU nodes with dual AMD EPYC 7H12 processors (128 cores per node) and GPU nodes with NVIDIA A100 GPUs (40GB each).

F ADDITIONAL EXPERIMENTS RESULTS

F.1 Esitmation of Directed-edge weights B

We conduct the experiment with `LinADMG-MuDoNLL` on Multi-domain Scalings DGP. In every experiment, we set $n = 10, N = 5000, T = 1, k_X = 2.0, k_U = 0.5$ and Gaussian noise. Results are averaged over 20 random seeds. In Fig. 4, the estimated directed-edge weights closely follow their ground-truth values, concentrating around the

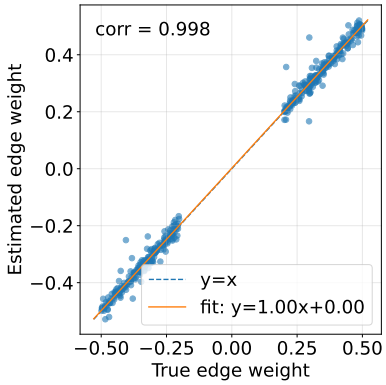

Figure 4: True vs. estimated directed-edge weights by `LinADMG-MuDoNLL` under multi-domain scalings DGP.

identity line. The near-perfect correlation (0.998) and fitted relation $y = x$ show highly accurate estimation.

F.2 All configurations of synthetic data experiments

We evaluate the estimation of both the directed and bidirected edges of the underlying ADMG under four multi-domain DGPs. The UP-Observation DGP applies domain-specific binary masks to the latent variables, the MD-Scalings DGP applies domain-specific continuous scaling factors, the hard-intervention DGP replaces the structural equations of intervened latent variables, and the Diag-GMMs DGP generates five groups whose latent distributions vary with an unobserved auxiliary variable. For each DGP, its designated method: `LinADMG-SP` , -MuDoNll, -UMNI, and -diagGMM, respectively, is compared with all four baselines LiNGAM, BANG, DCD, and RCD.

The anchor configuration uses $n = 10$, $T = n$ domains, ER graphs with directed-edge density $k_X = 2$ and confounding density $k_U = 0.5$, and Gaussian noise. Each domain contains $N = 5000$ observations, except for the hard-intervention anchor, which uses $N = 50000$. We conduct one-factor-at-a-time studies over dimension $n \in \{5, 10, 20\}$, number of domains $T/n \in \{1, 2, 3\}$, sample size $N \in \{1000, 5000, 50000\}$, directed-edge density $k_X \in \{1, 2, 3\}$, confounding density $k_U \in \{0, 0.5, 1\}$, and Gaussian, exponential, and Gumbel noise. The Diag-GMMs DGP always contains five Gaussian groups, so the domain, confounding-density, and noise sweeps do not apply to it. We report the mean and sample standard deviation of edge-wise F1, for which larger values are better, and structural Hamming distance (SHD), for which smaller values are better, over 20 independently generated random seeds.

MD-Scalings Tables 3 and 4 show that `LinADMG-MuDoNLL` benefits substantially from domain-specific scaling and is the strongest method in most configurations. Its performance improves systematically with the amount of multi-domain information: increasing the number of domains from $T = n$ to $T = 3n$ raises directed F1 and bidirected F1, while increasing N from 1000 to 50000 raises the respective scores as well. LiNGAM is strongest when there is no confounding. In contrast, BANG, DCD, and RCD recover little of the graph under Gaussian MD-Scalings, although BANG and RCD improve under non-Gaussian noise.

Diag-GMMs. Tables 5 and 6 show `LinADMG-diagGMM` consistently provides the best directed estimation across all configurations. For bidirected edges, `LinADMG-diagGMM` and BANG are the two strongest methods. BANG is slightly better at $n = 5$ and $n = 10$ and approaches perfect recovery at $N = 50000$. Overall, `LinADMG-diagGMM` offers the best overall recovery of the complete ADMG, while BANG is especially effective for the complete bidirected structure when sufficient data are available. Results of BANG and DCD at $n = 20$ are missing due to computational cost.

UP-Observations. Tables 7 and 8 show that LiNGAM provides the strongest overall performance under the UP-Observations DGP for directed edge estimation. `LinADMG-SP` is less accurate for bidirected-edge estimation and shows relatively large SHD, indicating that it tends to introduce additional bidirected edges. Nevertheless, `LinADMG-SP` becomes competitive for directed recovery under non-Gaussian noise and achieves the best directed results under Gumbel noise, whereas BANG is best under exponential noise. Performance generally decreases as the dimension n and confounding density K_U increase, although the increase in bidirected F1 at higher confounding density partly reflects the larger number of true bidirected edges. Overall, LiNGAM in general performs the best across the configurations, since the masking layer introduces independence to the noises, while the main advantage of `LinADMG-SP` is limited to directed recovery under Gumbel noise.

UMN-Hard Interventions. Tables 9 and 10 shows when n is relatively small, UMNI bidirected edges. At the $n = 10$ Gaussian anchor, however, its accuracy decreases and becomes highly variable. LiNGAM gives the best directed recovery, whereas UMNI better detects true confounders and avoids false positives. Under exponential and Gumbel noise, BANG becomes the strongest method for both directed and bidirected estimation, consistent with its assumption of non-Gaussianity. UMNI results are missing for the setting $n = 20$, $T = 3n$, $N = 1000$, and $N = 5000$ due to the requirement of sufficient sample numbers for precise score estimation from the algorithm perspective.

Table 3: MD-Scalings results: average F1 ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-MuDoNLL	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	0.927 ± 0.125	**0.943 ± 0.107**	0.000 ± 0.000	0.156 ± 0.139	0.197 ± 0.173
n	**10**	**0.883 ± 0.196**	0.815 ± 0.096	0.005 ± 0.021	0.078 ± 0.073	0.064 ± 0.083
n	**20**	**0.854 ± 0.280**	0.754 ± 0.069	0.005 ± 0.015	0.026 ± 0.034	0.033 ± 0.038
Number of domains						
T/n	**1**	**0.883 ± 0.196**	0.815 ± 0.096	0.005 ± 0.021	0.078 ± 0.073	0.064 ± 0.083
T/n	**2**	**0.911 ± 0.171**	0.825 ± 0.063	0.000 ± 0.000	0.053 ± 0.066	0.041 ± 0.073
T/n	**3**	**0.976 ± 0.087**	0.831 ± 0.041	0.005 ± 0.021	0.082 ± 0.078	0.052 ± 0.102
Sample size						
N	**1000**	**0.843 ± 0.101**	0.821 ± 0.072	0.010 ± 0.029	0.068 ± 0.062	0.079 ± 0.105
N	**5000**	**0.883 ± 0.196**	0.815 ± 0.096	0.005 ± 0.021	0.078 ± 0.073	0.064 ± 0.083
N	**50000**	**0.967 ± 0.108**	0.818 ± 0.081	0.000 ± 0.000	0.063 ± 0.051	0.063 ± 0.109
Graph density						
k_X	**1.0**	**0.773 ± 0.111**	0.671 ± 0.097	0.000 ± 0.000	0.083 ± 0.086	0.122 ± 0.125
k_X	**2.0**	**0.883 ± 0.196**	0.815 ± 0.096	0.005 ± 0.021	0.078 ± 0.073	0.064 ± 0.083
k_X	**3.0**	0.818 ± 0.268	**0.876 ± 0.099**	0.003 ± 0.014	0.066 ± 0.074	0.052 ± 0.082
Confounding density						
k_U	**0.0**	0.909 ± 0.063	**0.999 ± 0.005**	0.000 ± 0.000	0.089 ± 0.063	0.079 ± 0.104
k_U	**0.5**	**0.883 ± 0.196**	0.815 ± 0.096	0.005 ± 0.021	0.078 ± 0.073	0.064 ± 0.083
k_U	**1.0**	**0.803 ± 0.277**	0.674 ± 0.097	0.005 ± 0.020	0.053 ± 0.076	0.099 ± 0.080
Noise distribution						
Noise	**gauss**	**0.883 ± 0.196**	0.815 ± 0.096	0.005 ± 0.021	0.078 ± 0.073	0.064 ± 0.083
Noise	**exp**	**0.850 ± 0.176**	0.788 ± 0.107	0.823 ± 0.070	0.098 ± 0.084	0.696 ± 0.096
Noise	**gumbel**	**0.832 ± 0.218**	0.763 ± 0.091	0.424 ± 0.144	0.066 ± 0.083	0.618 ± 0.142
Bidirected edges						
Dimension						
n	**5**	**0.576 ± 0.142**	–	0.000 ± 0.000	0.190 ± 0.304	0.364 ± 0.128
n	**10**	**0.783 ± 0.142**	–	0.000 ± 0.000	0.056 ± 0.129	0.306 ± 0.080
n	**20**	**0.857 ± 0.156**	–	0.006 ± 0.028	0.018 ± 0.044	0.230 ± 0.048
Number of domains						
T/n	**1**	**0.783 ± 0.142**	–	0.000 ± 0.000	0.056 ± 0.129	0.306 ± 0.080
T/n	**2**	**0.864 ± 0.225**	–	0.000 ± 0.000	0.011 ± 0.050	0.290 ± 0.102
T/n	**3**	**0.933 ± 0.156**	–	0.000 ± 0.000	0.012 ± 0.056	0.315 ± 0.069
Sample size						
N	**1000**	**0.465 ± 0.092**	–	0.009 ± 0.041	0.036 ± 0.120	0.289 ± 0.088
N	**5000**	**0.783 ± 0.142**	–	0.000 ± 0.000	0.056 ± 0.129	0.306 ± 0.080
N	**50000**	**0.910 ± 0.185**	–	0.000 ± 0.000	0.035 ± 0.118	0.301 ± 0.082
Graph density						
k_X	**1.0**	**0.777 ± 0.165**	–	0.000 ± 0.000	0.011 ± 0.050	0.402 ± 0.110
k_X	**2.0**	**0.783 ± 0.142**	–	0.000 ± 0.000	0.056 ± 0.129	0.306 ± 0.080
k_X	**3.0**	**0.658 ± 0.272**	–	0.000 ± 0.000	0.053 ± 0.112	0.265 ± 0.064
Confounding density						
k_U	**0.0**	**0.000 ± 0.000**	–	**0.000 ± 0.000**	**0.000 ± 0.000**	**0.000 ± 0.000**
k_U	**0.5**	**0.783 ± 0.142**	–	0.000 ± 0.000	0.056 ± 0.129	0.306 ± 0.080
k_U	**1.0**	**0.823 ± 0.125**	–	0.000 ± 0.000	0.067 ± 0.083	0.555 ± 0.086
Noise distribution						
Noise	**gauss**	**0.783 ± 0.142**	–	0.000 ± 0.000	0.056 ± 0.129	0.306 ± 0.080
Noise	**exp**	**0.764 ± 0.166**	–	0.581 ± 0.201	0.025 ± 0.112	0.241 ± 0.118
Noise	**gumbel**	**0.714 ± 0.206**	–	0.376 ± 0.100	0.137 ± 0.162	0.243 ± 0.130

Table 4: MD-Scalings results: average SHD ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-MuDoNLL	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	1.300 ± 2.227	**1.100 ± 2.075**	10.050 ± 0.224	9.900 ± 1.518	8.850 ± 1.089
n	**10**	**4.500 ± 5.708**	9.300 ± 5.859	20.000 ± 0.324	21.050 ± 1.849	19.450 ± 0.999
n	**20**	**8.750 ± 12.586**	26.550 ± 10.560	40.100 ± 0.447	42.850 ± 2.207	39.450 ± 0.759
Number of domains						
T/n	**1**	**4.500 ± 5.708**	9.300 ± 5.859	20.000 ± 0.324	21.050 ± 1.849	19.450 ± 0.999
T/n	**2**	**3.550 ± 6.419**	8.400 ± 3.347	20.000 ± 0.000	20.550 ± 1.276	19.600 ± 0.754
T/n	**3**	**1.100 ± 4.012**	8.050 ± 2.212	19.950 ± 0.224	20.200 ± 1.152	19.400 ± 1.231
Sample size						
N	**1000**	**7.200 ± 4.360**	8.700 ± 4.143	19.950 ± 0.394	20.250 ± 1.372	19.300 ± 1.174
N	**5000**	**4.500 ± 5.708**	9.300 ± 5.859	20.000 ± 0.324	21.050 ± 1.849	19.450 ± 0.999
N	**50000**	**1.300 ± 4.207**	9.100 ± 4.973	20.200 ± 0.410	20.750 ± 1.372	19.400 ± 1.142
Graph density						
k_X	**1.0**	**6.050 ± 3.561**	9.850 ± 3.937	10.000 ± 0.000	11.150 ± 1.631	9.400 ± 0.821
k_X	**2.0**	**4.500 ± 5.708**	9.300 ± 5.859	20.000 ± 0.324	21.050 ± 1.849	19.450 ± 0.999
k_X	**3.0**	9.000 ± 12.490	**8.450 ± 7.529**	29.950 ± 0.224	30.100 ± 1.683	29.150 ± 1.387
Confounding density						
k_U	**0.0**	4.050 ± 2.892	**0.050 ± 0.224**	20.000 ± 0.000	20.100 ± 1.165	19.250 ± 1.164
k_U	**0.5**	**4.500 ± 5.708**	9.300 ± 5.859	20.000 ± 0.324	21.050 ± 1.849	19.450 ± 0.999
k_U	**1.0**	**7.350 ± 9.051**	18.450 ± 6.395	20.100 ± 0.308	20.800 ± 1.609	19.200 ± 0.894
Noise distribution						
Noise	**gauss**	**4.500 ± 5.708**	9.300 ± 5.859	20.000 ± 0.324	21.050 ± 1.849	19.450 ± 0.999
Noise	**exp**	**6.050 ± 5.031**	10.350 ± 5.706	7.300 ± 2.774	19.850 ± 1.461	10.350 ± 2.661
Noise	**gumbel**	**6.650 ± 6.556**	12.100 ± 5.389	17.500 ± 3.154	20.400 ± 1.465	12.050 ± 3.677
Bidirected edges						
Dimension						
n	**5**	3.650 ± 1.496	–	2.400 ± 0.503	**2.300 ± 0.865**	5.700 ± 1.720
n	**10**	**4.500 ± 5.615**	–	7.500 ± 1.357	7.300 ± 1.490	21.500 ± 3.606
n	**20**	**5.550 ± 7.884**	–	16.000 ± 4.078	15.750 ± 3.932	64.750 ± 9.095
Number of domains						
T/n	**1**	**4.500 ± 5.615**	–	7.500 ± 1.357	7.300 ± 1.490	21.500 ± 3.606
T/n	**2**	**3.000 ± 6.224**	–	7.050 ± 1.731	7.100 ± 1.832	21.050 ± 4.334
T/n	**3**	**1.600 ± 4.893**	–	7.150 ± 1.565	7.200 ± 1.673	22.200 ± 3.636
Sample size						
N	**1000**	15.250 ± 4.241	–	7.500 ± 1.192	**7.450 ± 1.432**	22.150 ± 3.543
N	**5000**	**4.500 ± 5.615**	–	7.500 ± 1.357	7.300 ± 1.490	21.500 ± 3.606
N	**50000**	**2.300 ± 7.035**	–	7.500 ± 1.357	7.500 ± 1.433	22.650 ± 3.360
Graph density						
k_X	**1.0**	**4.450 ± 5.558**	–	7.500 ± 1.277	7.450 ± 1.395	12.950 ± 2.328
k_X	**2.0**	**4.500 ± 5.615**	–	7.500 ± 1.357	7.300 ± 1.490	21.500 ± 3.606
k_X	**3.0**	9.750 ± 10.543	–	**7.450 ± 1.317**	7.750 ± 1.650	27.200 ± 3.679
Confounding density						
k_U	**0.0**	6.750 ± 3.985	–	**0.100 ± 0.308**	0.250 ± 0.444	22.600 ± 5.185
k_U	**0.5**	**4.500 ± 5.615**	–	7.500 ± 1.357	7.300 ± 1.490	21.500 ± 3.606
k_U	**1.0**	**8.250 ± 6.766**	–	22.100 ± 4.128	21.450 ± 4.478	20.600 ± 3.119
Noise distribution						
Noise	**gauss**	**4.500 ± 5.615**	–	7.500 ± 1.357	7.300 ± 1.490	21.500 ± 3.606
Noise	**exp**	4.700 ± 4.438	–	**4.600 ± 1.875**	7.500 ± 1.433	10.300 ± 2.342
Noise	**gumbel**	**6.100 ± 6.095**	–	21.850 ± 5.896	8.750 ± 2.049	12.250 ± 3.401

Table 5: Diag-GMMs results: average F1 ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-diagGMM	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	**0.873 ± 0.157**	0.617 ± 0.193	0.000 ± 0.000	0.146 ± 0.132	0.148 ± 0.171
n	**10**	**0.817 ± 0.100**	0.495 ± 0.128	0.062 ± 0.100	0.135 ± 0.083	0.141 ± 0.091
n	**20**	**0.743 ± 0.081**	0.375 ± 0.058	–	–	0.134 ± 0.093
Sample size						
N	**1000**	**0.806 ± 0.092**	0.536 ± 0.126	0.109 ± 0.089	0.082 ± 0.079	0.154 ± 0.096
N	**5000**	**0.817 ± 0.100**	0.495 ± 0.128	0.062 ± 0.100	0.135 ± 0.083	0.141 ± 0.091
N	**50000**	**0.821 ± 0.076**	0.495 ± 0.108	0.005 ± 0.021	0.120 ± 0.083	0.160 ± 0.110
Graph density						
k_X	**1.0**	**0.691 ± 0.107**	0.313 ± 0.092	0.071 ± 0.108	0.087 ± 0.116	0.124 ± 0.112
k_X	**2.0**	**0.817 ± 0.100**	0.495 ± 0.128	0.062 ± 0.100	0.135 ± 0.083	0.141 ± 0.091
k_X	**3.0**	**0.833 ± 0.109**	0.653 ± 0.113	0.052 ± 0.061	0.125 ± 0.064	0.104 ± 0.084
Bidirected edges						
Dimension						
n	**5**	0.956 ± 0.077	–	**0.995 ± 0.016**	0.538 ± 0.218	0.709 ± 0.211
n	**10**	0.930 ± 0.198	–	**0.946 ± 0.059**	0.347 ± 0.101	0.673 ± 0.121
n	**20**	**0.951 ± 0.082**	–	–	–	0.648 ± 0.104
Sample size						
N	**1000**	**0.971 ± 0.049**	–	0.811 ± 0.134	0.393 ± 0.103	0.691 ± 0.121
N	**5000**	0.930 ± 0.198	–	**0.946 ± 0.059**	0.347 ± 0.101	0.673 ± 0.121
N	**50000**	0.973 ± 0.064	–	**0.996 ± 0.008**	0.337 ± 0.065	0.637 ± 0.127
Graph density						
k_X	**1.0**	0.920 ± 0.151	–	**0.940 ± 0.079**	0.324 ± 0.094	0.658 ± 0.138
k_X	**2.0**	0.930 ± 0.198	–	**0.946 ± 0.059**	0.347 ± 0.101	0.673 ± 0.121
k_X	**3.0**	**0.921 ± 0.137**	–	0.918 ± 0.072	0.358 ± 0.087	0.751 ± 0.105

Table 6: Diag-GMMs results: average SHD ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-diagGMM	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	**2.200 ± 2.546**	7.550 ± 3.720	10.000 ± 0.000	9.950 ± 1.432	9.800 ± 1.508
n	**10**	**7.650 ± 3.897**	30.300 ± 8.417	20.350 ± 1.599	21.200 ± 1.989	22.000 ± 2.026
n	**20**	**24.800 ± 9.339**	114.450 ± 18.715	–	–	48.450 ± 4.548
Sample size						
N	**1000**	**7.950 ± 3.634**	26.450 ± 7.844	21.900 ± 2.049	21.900 ± 1.861	21.750 ± 2.807
N	**5000**	**7.650 ± 3.897**	30.300 ± 8.417	20.350 ± 1.599	21.200 ± 1.989	22.000 ± 2.026
N	**50000**	**7.600 ± 3.331**	31.600 ± 7.430	20.050 ± 0.394	21.650 ± 2.346	22.900 ± 2.827
Graph density						
k_X	**1.0**	**8.550 ± 3.832**	33.500 ± 6.747	10.800 ± 1.361	12.850 ± 2.346	13.800 ± 2.726
k_X	**2.0**	**7.650 ± 3.897**	30.300 ± 8.417	20.350 ± 1.599	21.200 ± 1.989	22.000 ± 2.026
k_X	**3.0**	**9.550 ± 5.698**	24.950 ± 8.532	31.050 ± 1.638	30.450 ± 2.305	31.200 ± 2.331
Bidirected edges						
Dimension						
n	**5**	0.750 ± 1.251	–	**0.100 ± 0.308**	6.050 ± 1.932	4.150 ± 2.323
n	**10**	**4.200 ± 9.876**	–	4.350 ± 4.591	35.350 ± 3.360	21.600 ± 6.516
n	**20**	**16.050 ± 23.712**	–	–	–	97.400 ± 21.960
Sample size						
N	**1000**	**2.350 ± 3.951**	–	13.450 ± 8.281	33.800 ± 3.427	20.700 ± 6.165
N	**5000**	**4.200 ± 9.876**	–	4.350 ± 4.591	35.350 ± 3.360	21.600 ± 6.516
N	**50000**	2.100 ± 4.767	–	**0.350 ± 0.745**	35.800 ± 2.093	23.450 ± 5.942
Graph density						
k_X	**1.0**	5.500 ± 8.593	–	**4.650 ± 5.887**	36.150 ± 3.066	22.300 ± 6.658
k_X	**2.0**	**4.200 ± 9.876**	–	4.350 ± 4.591	35.350 ± 3.360	21.600 ± 6.516
k_X	**3.0**	**5.650 ± 8.061**	–	6.500 ± 5.395	35.050 ± 2.929	17.500 ± 5.862

Table 7: UP-Observations results: average F1 ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-SP	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	0.746 ± 0.242	**0.980 ± 0.052**	0.000 ± 0.000	0.156 ± 0.139	0.350 ± 0.177
n	**10**	0.724 ± 0.224	**0.838 ± 0.051**	0.010 ± 0.029	0.076 ± 0.078	0.427 ± 0.179
n	**20**	0.557 ± 0.149	**0.750 ± 0.072**	0.000 ± 0.000	0.019 ± 0.023	0.532 ± 0.112
Number of domains						
T/n	**1**	0.724 ± 0.224	**0.838 ± 0.051**	0.010 ± 0.029	0.076 ± 0.078	0.427 ± 0.179
T/n	**2**	0.690 ± 0.294	**0.841 ± 0.041**	0.005 ± 0.021	0.058 ± 0.059	0.470 ± 0.162
T/n	**3**	0.784 ± 0.179	**0.837 ± 0.044**	0.005 ± 0.020	0.070 ± 0.064	0.565 ± 0.175
Sample size						
N	**1000**	0.679 ± 0.224	**0.847 ± 0.062**	0.000 ± 0.000	0.071 ± 0.081	0.386 ± 0.174
N	**5000**	0.724 ± 0.224	**0.838 ± 0.051**	0.010 ± 0.029	0.076 ± 0.078	0.427 ± 0.179
N	**50000**	0.729 ± 0.213	**0.821 ± 0.058**	0.005 ± 0.021	0.065 ± 0.088	0.373 ± 0.161
Graph density						
k_X	**1.0**	0.573 ± 0.194	**0.702 ± 0.087**	0.009 ± 0.041	0.093 ± 0.102	0.479 ± 0.206
k_X	**2.0**	0.724 ± 0.224	**0.838 ± 0.051**	0.010 ± 0.029	0.076 ± 0.078	0.427 ± 0.179
k_X	**3.0**	0.602 ± 0.241	**0.923 ± 0.031**	0.003 ± 0.014	0.072 ± 0.062	0.347 ± 0.178
Confounding density						
k_U	**0.0**	0.671 ± 0.223	**1.000 ± 0.000**	0.000 ± 0.000	0.100 ± 0.073	0.422 ± 0.151
k_U	**0.5**	0.724 ± 0.224	**0.838 ± 0.051**	0.010 ± 0.029	0.076 ± 0.078	0.427 ± 0.179
k_U	**1.0**	0.549 ± 0.262	**0.738 ± 0.045**	0.000 ± 0.000	0.057 ± 0.065	0.402 ± 0.189
Noise distribution						
Noise	**gauss**	0.724 ± 0.224	**0.838 ± 0.051**	0.010 ± 0.029	0.076 ± 0.078	0.427 ± 0.179
Noise	**exp**	0.761 ± 0.182	0.837 ± 0.045	**0.849 ± 0.071**	0.063 ± 0.059	0.610 ± 0.142
Noise	**gumbel**	**0.792 ± 0.138**	0.702 ± 0.039	0.343 ± 0.143	0.071 ± 0.061	0.624 ± 0.138
Bidirected edges						
Dimension						
n	**5**	**0.397 ± 0.132**	–	0.000 ± 0.000	0.125 ± 0.259	0.330 ± 0.156
n	**10**	**0.383 ± 0.109**	–	0.000 ± 0.000	0.022 ± 0.070	0.255 ± 0.110
n	**20**	**0.236 ± 0.066**	–	0.000 ± 0.000	0.010 ± 0.030	0.184 ± 0.082
Number of domains						
T/n	**1**	**0.383 ± 0.109**	–	0.000 ± 0.000	0.022 ± 0.070	0.255 ± 0.110
T/n	**2**	**0.395 ± 0.111**	–	0.011 ± 0.050	0.010 ± 0.045	0.273 ± 0.138
T/n	**3**	**0.391 ± 0.093**	–	0.000 ± 0.000	0.010 ± 0.045	0.254 ± 0.099
Sample size						
N	**1000**	**0.337 ± 0.093**	–	0.000 ± 0.000	0.012 ± 0.056	0.271 ± 0.103
N	**5000**	**0.383 ± 0.109**	–	0.000 ± 0.000	0.022 ± 0.070	0.255 ± 0.110
N	**50000**	**0.366 ± 0.116**	–	0.000 ± 0.000	0.012 ± 0.056	0.270 ± 0.127
Graph density						
k_X	**1.0**	**0.421 ± 0.113**	–	0.000 ± 0.000	0.000 ± 0.000	0.332 ± 0.148
k_X	**2.0**	**0.383 ± 0.109**	–	0.000 ± 0.000	0.022 ± 0.070	0.255 ± 0.110
k_X	**3.0**	**0.298 ± 0.066**	–	0.000 ± 0.000	0.054 ± 0.096	0.285 ± 0.111
Confounding density						
k_U	**0.0**	**0.000 ± 0.000**	–	**0.000 ± 0.000**	**0.000 ± 0.000**	**0.000 ± 0.000**
k_U	**0.5**	**0.383 ± 0.109**	–	0.000 ± 0.000	0.022 ± 0.070	0.255 ± 0.110
k_U	**1.0**	**0.618 ± 0.100**	–	0.010 ± 0.030	0.039 ± 0.057	0.441 ± 0.136
Noise distribution						
Noise	**gauss**	**0.383 ± 0.109**	–	0.000 ± 0.000	0.022 ± 0.070	0.255 ± 0.110
Noise	**exp**	0.395 ± 0.101	–	**0.726 ± 0.183**	0.010 ± 0.045	0.274 ± 0.169
Noise	**gumbel**	**0.387 ± 0.073**	–	0.328 ± 0.084	0.139 ± 0.140	0.278 ± 0.149

Table 8: UP-Observations results: average SHD ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-SP	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	4.400 ± 4.309	**0.400 ± 1.046**	10.000 ± 0.000	9.800 ± 1.361	7.750 ± 1.251
n	**10**	11.500 ± 8.930	**7.750 ± 2.751**	20.000 ± 0.459	20.600 ± 1.930	15.300 ± 3.373
n	**20**	39.650 ± 12.032	27.650 ± 12.398	40.200 ± 0.410	41.450 ± 1.395	**27.400 ± 4.358**
Number of domains						
T/n	**1**	11.500 ± 8.930	**7.750 ± 2.751**	20.000 ± 0.459	20.600 ± 1.930	15.300 ± 3.373
T/n	**2**	13.800 ± 13.836	**7.650 ± 2.390**	20.000 ± 0.324	20.600 ± 1.188	14.500 ± 3.086
T/n	**3**	9.500 ± 8.513	**7.900 ± 2.532**	20.250 ± 0.716	20.500 ± 1.100	12.650 ± 3.924
Sample size						
N	**1000**	13.850 ± 9.869	**7.300 ± 3.358**	20.000 ± 0.000	20.400 ± 1.536	16.250 ± 3.492
N	**5000**	11.500 ± 8.930	**7.750 ± 2.751**	20.000 ± 0.459	20.600 ± 1.930	15.300 ± 3.373
N	**50000**	11.600 ± 9.355	**8.800 ± 3.286**	19.950 ± 0.224	20.650 ± 1.631	16.100 ± 2.404
Graph density						
k_X	**1.0**	12.300 ± 5.110	8.850 ± 3.558	10.050 ± 0.394	10.450 ± 1.234	**7.400 ± 1.875**
k_X	**2.0**	11.500 ± 8.930	**7.750 ± 2.751**	20.000 ± 0.459	20.600 ± 1.930	15.300 ± 3.373
k_X	**3.0**	24.450 ± 16.504	**5.050 ± 2.139**	30.100 ± 0.308	30.250 ± 1.832	23.650 ± 4.380
Confounding density						
k_U	**0.0**	14.600 ± 10.252	**0.000 ± 0.000**	20.000 ± 0.000	20.200 ± 1.576	14.800 ± 2.687
k_U	**0.5**	11.500 ± 8.930	**7.750 ± 2.751**	20.000 ± 0.459	20.600 ± 1.930	15.300 ± 3.373
k_U	**1.0**	21.250 ± 13.325	**14.200 ± 3.071**	20.000 ± 0.000	20.450 ± 0.887	15.500 ± 3.317
Noise distribution						
Noise	**gauss**	11.500 ± 8.930	**7.750 ± 2.751**	20.000 ± 0.459	20.600 ± 1.930	15.300 ± 3.373
Noise	**exp**	10.050 ± 7.222	7.850 ± 2.540	**6.050 ± 2.856**	20.650 ± 1.387	12.100 ± 3.712
Noise	**gumbel**	**8.950 ± 5.698**	17.050 ± 3.017	19.100 ± 2.292	20.500 ± 1.235	11.700 ± 3.294
Bidirected edges						
Dimension						
n	**5**	5.800 ± 1.542	–	2.250 ± 0.444	**2.150 ± 0.671**	5.550 ± 1.761
n	**10**	21.150 ± 4.416	–	**7.300 ± 1.261**	7.750 ± 1.372	16.600 ± 3.733
n	**20**	103.250 ± 39.034	–	16.800 ± 4.200	**16.700 ± 3.600**	38.000 ± 9.414
Number of domains						
T/n	**1**	21.150 ± 4.416	–	**7.300 ± 1.261**	7.750 ± 1.372	16.600 ± 3.733
T/n	**2**	20.150 ± 5.566	–	**7.250 ± 1.251**	7.650 ± 1.089	15.500 ± 4.310
T/n	**3**	20.100 ± 4.667	–	**7.250 ± 1.293**	7.600 ± 1.231	14.150 ± 3.815
Sample size						
N	**1000**	23.550 ± 4.639	–	**7.250 ± 1.293**	7.650 ± 1.309	16.950 ± 3.486
N	**5000**	21.150 ± 4.416	–	**7.300 ± 1.261**	7.750 ± 1.372	16.600 ± 3.733
N	**50000**	21.600 ± 5.205	–	**7.300 ± 1.302**	7.850 ± 1.424	18.200 ± 5.227
Graph density						
k_X	**1.0**	18.000 ± 5.161	–	**7.250 ± 1.293**	7.500 ± 1.318	10.100 ± 2.972
k_X	**2.0**	21.150 ± 4.416	–	**7.300 ± 1.261**	7.750 ± 1.372	16.600 ± 3.733
k_X	**3.0**	26.550 ± 3.913	–	**7.250 ± 1.293**	7.750 ± 1.333	20.100 ± 4.941
Confounding density						
k_U	**0.0**	24.000 ± 2.753	–	**0.050 ± 0.224**	0.200 ± 0.410	16.100 ± 4.844
k_U	**0.5**	21.150 ± 4.416	–	**7.300 ± 1.261**	7.750 ± 1.372	16.600 ± 3.733
k_U	**1.0**	**19.900 ± 4.128**	–	21.150 ± 4.545	21.050 ± 4.186	21.000 ± 4.104
Noise distribution						
Noise	**gauss**	21.150 ± 4.416	–	**7.300 ± 1.261**	7.750 ± 1.372	16.600 ± 3.733
Noise	**exp**	19.150 ± 3.787	–	**3.300 ± 2.179**	7.750 ± 1.164	13.050 ± 4.513
Noise	**gumbel**	20.950 ± 3.605	–	25.350 ± 5.143	**9.250 ± 1.860**	12.450 ± 2.911

Table 9: UMN-Hard Interventions results: average F1 ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-UMNI	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	**0.994 ± 0.025**	0.960 ± 0.082	0.000 ± 0.000	0.209 ± 0.126	0.000 ± 0.000
n	**10**	0.601 ± 0.366	**0.821 ± 0.069**	0.000 ± 0.000	0.058 ± 0.071	0.000 ± 0.000
n	**20**	–	**0.705 ± 0.083**	0.002 ± 0.011	0.019 ± 0.028	0.000 ± 0.000
Number of domains						
T/n	**1**	0.601 ± 0.366	**0.821 ± 0.069**	0.000 ± 0.000	0.058 ± 0.071	0.000 ± 0.000
T/n	**2**	0.377 ± 0.161	**0.801 ± 0.073**	0.000 ± 0.000	0.069 ± 0.068	0.000 ± 0.000
T/n	**3**	–	**0.814 ± 0.072**	0.000 ± 0.000	0.071 ± 0.070	0.000 ± 0.000
Sample size						
N	**1000**	–	**0.724 ± 0.138**	0.005 ± 0.021	0.084 ± 0.072	0.000 ± 0.000
N	**5000**	–	**0.823 ± 0.069**	0.000 ± 0.000	0.066 ± 0.073	0.000 ± 0.000
N	**50000**	0.601 ± 0.366	**0.821 ± 0.069**	0.000 ± 0.000	0.058 ± 0.071	0.000 ± 0.000
Graph density						
k_X	**1.0**	0.576 ± 0.403	**0.656 ± 0.116**	0.009 ± 0.041	0.049 ± 0.094	0.000 ± 0.000
k_X	**2.0**	0.601 ± 0.366	**0.821 ± 0.069**	0.000 ± 0.000	0.058 ± 0.071	0.000 ± 0.000
k_X	**3.0**	0.720 ± 0.286	**0.920 ± 0.033**	0.000 ± 0.000	0.076 ± 0.057	0.000 ± 0.000
Confounding density						
k_U	**0.0**	0.574 ± 0.338	**1.000 ± 0.000**	0.000 ± 0.000	0.068 ± 0.050	0.000 ± 0.000
k_U	**0.5**	0.601 ± 0.366	**0.821 ± 0.069**	0.000 ± 0.000	0.058 ± 0.071	0.000 ± 0.000
k_U	**1.0**	0.519 ± 0.367	**0.677 ± 0.086**	0.000 ± 0.000	0.085 ± 0.066	0.000 ± 0.000
Noise distribution						
Noise	**gauss**	0.601 ± 0.366	**0.821 ± 0.069**	0.000 ± 0.000	0.058 ± 0.071	0.000 ± 0.000
Noise	**exp**	0.618 ± 0.280	0.808 ± 0.068	**0.842 ± 0.077**	0.085 ± 0.078	0.753 ± 0.118
Noise	**gumbel**	0.723 ± 0.328	0.805 ± 0.068	**0.826 ± 0.073**	0.066 ± 0.074	0.442 ± 0.177
Bidirected edges						
Dimension						
n	**5**	**0.969 ± 0.104**	–	0.000 ± 0.000	0.238 ± 0.286	0.332 ± 0.116
n	**10**	**0.578 ± 0.259**	–	0.000 ± 0.000	0.192 ± 0.160	0.311 ± 0.079
n	**20**	–	–	0.000 ± 0.000	0.158 ± 0.121	**0.263 ± 0.055**
Number of domains						
T/n	**1**	**0.578 ± 0.259**	–	0.000 ± 0.000	0.192 ± 0.160	0.311 ± 0.079
T/n	**2**	**0.315 ± 0.110**	–	0.000 ± 0.000	0.182 ± 0.120	0.308 ± 0.091
T/n	**3**	–	–	0.000 ± 0.000	0.166 ± 0.127	**0.312 ± 0.076**
Sample size						
N	**1000**	–	–	0.000 ± 0.000	0.194 ± 0.134	**0.297 ± 0.092**
N	**5000**	–	–	0.000 ± 0.000	0.182 ± 0.140	**0.318 ± 0.088**
N	**50000**	**0.578 ± 0.259**	–	0.000 ± 0.000	0.192 ± 0.160	0.311 ± 0.079
Graph density						
k_X	**1.0**	**0.672 ± 0.252**	–	0.000 ± 0.000	0.086 ± 0.113	0.400 ± 0.089
k_X	**2.0**	**0.578 ± 0.259**	–	0.000 ± 0.000	0.192 ± 0.160	0.311 ± 0.079
k_X	**3.0**	**0.556 ± 0.253**	–	0.012 ± 0.056	0.173 ± 0.152	0.257 ± 0.073
Confounding density						
k_U	**0.0**	**0.000 ± 0.000**	–	**0.000 ± 0.000**	**0.000 ± 0.000**	**0.000 ± 0.000**
k_U	**0.5**	**0.578 ± 0.259**	–	0.000 ± 0.000	0.192 ± 0.160	0.311 ± 0.079
k_U	**1.0**	**0.678 ± 0.165**	–	0.000 ± 0.000	0.171 ± 0.104	0.587 ± 0.078
Noise distribution						
Noise	**gauss**	**0.578 ± 0.259**	–	0.000 ± 0.000	0.192 ± 0.160	0.311 ± 0.079
Noise	**exp**	0.470 ± 0.179	–	**0.779 ± 0.189**	0.147 ± 0.112	0.205 ± 0.171
Noise	**gumbel**	0.670 ± 0.235	–	**0.763 ± 0.211**	0.276 ± 0.125	0.305 ± 0.108

Table 10: UMN-Hard Interventions results: average SHD ± standard deviation over 20 random seeds. The varying parameter and the best method in each row are shown in bold.

Setting	Value	LinADMG-UMNI	LiNGAM	BANG	DCD	RCD
Directed edges						
Dimension						
n	**5**	**0.100 ± 0.447**	0.800 ± 1.642	10.000 ± 0.000	9.450 ± 1.146	10.000 ± 0.000
n	**10**	12.200 ± 9.350	**8.850 ± 4.069**	20.100 ± 0.308	20.650 ± 1.663	20.050 ± 0.224
n	**20**	–	**34.250 ± 14.407**	40.000 ± 0.324	41.050 ± 1.731	40.000 ± 0.000
Number of domains						
T/n	**1**	12.200 ± 9.350	**8.850 ± 4.069**	20.100 ± 0.308	20.650 ± 1.663	20.050 ± 0.224
T/n	**2**	23.556 ± 6.070	**9.900 ± 4.436**	20.100 ± 0.308	20.550 ± 1.356	20.000 ± 0.000
T/n	**3**	–	**9.300 ± 4.566**	20.050 ± 0.224	20.400 ± 1.273	20.000 ± 0.000
Sample size						
N	**1000**	–	**14.150 ± 8.487**	20.000 ± 0.324	20.350 ± 1.755	20.000 ± 0.000
N	**5000**	–	**8.650 ± 4.030**	20.000 ± 0.000	20.450 ± 1.276	20.000 ± 0.000
N	**50000**	12.200 ± 9.350	**8.850 ± 4.069**	20.100 ± 0.308	20.650 ± 1.663	20.050 ± 0.224
Graph density						
k_X	**1.0**	**6.800 ± 5.307**	11.000 ± 5.400	9.950 ± 0.224	10.950 ± 1.050	10.000 ± 0.000
k_X	**2.0**	12.200 ± 9.350	**8.850 ± 4.069**	20.100 ± 0.308	20.650 ± 1.663	20.050 ± 0.224
k_X	**3.0**	12.850 ± 10.752	**5.250 ± 2.337**	30.050 ± 0.224	29.850 ± 1.424	30.000 ± 0.000
Confounding density						
k_U	**0.0**	12.150 ± 7.748	**0.000 ± 0.000**	20.050 ± 0.224	20.500 ± 1.357	20.000 ± 0.000
k_U	**0.5**	12.200 ± 9.350	**8.850 ± 4.069**	20.100 ± 0.308	20.650 ± 1.663	20.050 ± 0.224
k_U	**1.0**	**14.200 ± 9.328**	18.300 ± 5.904	20.050 ± 0.224	20.250 ± 1.410	20.000 ± 0.000
Noise distribution						
Noise	**gauss**	12.200 ± 9.350	**8.850 ± 4.069**	20.100 ± 0.308	20.650 ± 1.663	20.050 ± 0.224
Noise	**exp**	12.950 ± 7.015	9.250 ± 3.552	**6.400 ± 3.152**	20.100 ± 1.553	8.750 ± 3.640
Noise	**gumbel**	8.550 ± 8.319	9.450 ± 3.531	**7.100 ± 3.042**	20.500 ± 1.357	15.050 ± 3.561
Bidirected edges						
Dimension						
n	**5**	**0.200 ± 0.696**	–	2.250 ± 0.444	2.750 ± 1.446	6.750 ± 1.293
n	**10**	10.350 ± 7.043	–	**7.250 ± 1.293**	8.500 ± 1.960	23.200 ± 3.238
n	**20**	–	–	**16.500 ± 3.940**	19.000 ± 4.531	60.850 ± 8.273
Number of domains						
T/n	**1**	10.350 ± 7.043	–	**7.250 ± 1.293**	8.500 ± 1.960	23.200 ± 3.238
T/n	**2**	21.167 ± 5.021	–	**7.250 ± 1.293**	8.700 ± 1.720	23.150 ± 4.030
T/n	**3**	–	–	**7.250 ± 1.293**	8.750 ± 1.803	22.750 ± 2.936
Sample size						
N	**1000**	–	–	**7.250 ± 1.293**	8.400 ± 1.729	23.650 ± 4.056
N	**5000**	–	–	**7.250 ± 1.293**	8.550 ± 2.064	23.100 ± 3.597
N	**50000**	10.350 ± 7.043	–	**7.250 ± 1.293**	8.500 ± 1.960	23.200 ± 3.238
Graph density						
k_X	**1.0**	**6.950 ± 5.916**	–	7.250 ± 1.293	8.500 ± 1.762	14.300 ± 4.105
k_X	**2.0**	10.350 ± 7.043	–	**7.250 ± 1.293**	8.500 ± 1.960	23.200 ± 3.238
k_X	**3.0**	10.400 ± 6.402	–	**7.200 ± 1.322**	9.900 ± 1.683	28.900 ± 3.076
Confounding density						
k_U	**0.0**	10.800 ± 6.910	–	**0.000 ± 0.000**	2.800 ± 1.704	24.650 ± 3.631
k_U	**0.5**	10.350 ± 7.043	–	**7.250 ± 1.293**	8.500 ± 1.960	23.200 ± 3.238
k_U	**1.0**	**13.750 ± 7.440**	–	21.200 ± 4.514	21.050 ± 3.953	20.000 ± 2.974
Noise distribution						
Noise	**gauss**	10.350 ± 7.043	–	**7.250 ± 1.293**	8.500 ± 1.960	23.200 ± 3.238
Noise	**exp**	13.000 ± 4.600	–	**2.800 ± 2.353**	9.050 ± 1.905	9.550 ± 2.819
Noise	**gumbel**	7.650 ± 6.831	–	**3.150 ± 2.720**	11.400 ± 3.560	15.500 ± 3.269

F.3 VaDE vs. diagGMM likelihood loss

In this section, we compare `LinADMG-diagGMM` and VaDE, as in Kivva et al. [2022], under the DiagGMMs data-generating process. Starting from the anchor configuration $n = 10, N = 5000, k_X = 2$, it varies the number of variables $n \in \{5, 10, 20\}$, sample size $N \in \{1000, 5000, 50000\}$, and graph density $k_X \in \{1, 2, 3\}$. We evaluate each configuration over 20 random seeds and report performance separately for directed and bidirected edges using mean F1 and SHD, with standard deviations. `LinADMG-diagGMM` clearly outperforms VaDE in estimating directed edges across every configuration, achieving significantly higher F1 and lower SHD. The performance of both methods decreases as dimension increases, whereas increasing sample size has relatively little effect. Higher graph density improves directed F1 but generally increases SHD. For bidirected edges, both methods perform well, and their F1 scores are much closer: VaDE is slightly better in several low- or moderate-complexity settings, while `LinADMG-diagGMM` performs better for $n = 20$, $N = 1000$, $N = 50000$, and $k_X = 3$.

Table 11: Diag-GMMs results: average F1 and SHD ± standard deviation over 20 random seeds. The varying parameter and the best method for each metric in each row are shown in bold.

Setting	Value	F1		SHD	
		LinADMG-diagGMM	VaDE-linear	LinADMG-diagGMM	VaDE-linear
Directed edges					
Dimension					
n	**5**	**0.873 ± 0.157**	0.421 ± 0.175	**2.200 ± 2.546**	8.700 ± 2.342
n	**10**	**0.817 ± 0.100**	0.254 ± 0.112	**7.650 ± 3.897**	26.050 ± 3.900
n	**20**	**0.743 ± 0.081**	0.144 ± 0.043	**24.800 ± 9.339**	68.000 ± 4.702
Sample size					
N	**1000**	**0.806 ± 0.092**	0.218 ± 0.084	**7.950 ± 3.634**	26.950 ± 2.911
N	**5000**	**0.817 ± 0.100**	0.254 ± 0.112	**7.650 ± 3.897**	26.050 ± 3.900
N	**50000**	**0.821 ± 0.076**	0.243 ± 0.108	**7.600 ± 3.331**	26.600 ± 3.691
Graph density					
k_X	**1.0**	**0.691 ± 0.107**	0.178 ± 0.092	**8.550 ± 3.832**	20.150 ± 2.925
k_X	**2.0**	**0.817 ± 0.100**	0.254 ± 0.112	**7.650 ± 3.897**	26.050 ± 3.900
k_X	**3.0**	**0.833 ± 0.109**	0.269 ± 0.107	**9.550 ± 5.698**	32.650 ± 4.534
Bidirected edges					
Dimension					
n	**5**	0.956 ± 0.077	**0.959 ± 0.052**	**0.750 ± 1.251**	**0.750 ± 0.910**
n	**10**	0.930 ± 0.198	**0.939 ± 0.050**	**4.200 ± 9.876**	5.000 ± 3.853
n	**20**	**0.951 ± 0.082**	0.882 ± 0.061	**16.050 ± 23.712**	39.200 ± 17.981
Sample size					
N	**1000**	**0.971 ± 0.049**	0.940 ± 0.049	**2.350 ± 3.951**	4.900 ± 3.698
N	**5000**	0.930 ± 0.198	**0.939 ± 0.050**	**4.200 ± 9.876**	5.000 ± 3.853
N	**50000**	**0.973 ± 0.064**	0.937 ± 0.053	**2.100 ± 4.767**	5.150 ± 4.017
Graph density					
k_X	**1.0**	0.920 ± 0.151	**0.934 ± 0.038**	**5.500 ± 8.593**	**5.500 ± 3.017**
k_X	**2.0**	0.930 ± 0.198	**0.939 ± 0.050**	**4.200 ± 9.876**	5.000 ± 3.853
k_X	**3.0**	**0.921 ± 0.137**	0.918 ± 0.045	**5.650 ± 8.061**	6.650 ± 3.528

F.4 MuDo-NLL vs. MuDo-Cov

In this section, we compare the training objective in MuDo-Cov used in Xu et al. [2026] with MuDo-NLL, which we use in this work. MuDoNLL minimizes the Gaussian negative log-likelihood, while MuDo-Cov minimizes the covariance matching error. We hold all other optimization and evaluation settings fixed, and evaluate every condition using the same 20 random seeds. Starting from the anchor configuration $n = 10, N = 5000, T = n, k_X = 2, k_U = 0.5$, and Gaussian noise, we vary one factor at a time. We evaluate performance separately for directed and bidirected edges using F1 (higher is better) and SHD (lower is better). The results show that MuDo-NLL outperforms MuDo-Cov in nearly every matched condition, with the largest advantage occurring in bidirected-edge recovery. Both objectives benefit from additional domains and larger sample sizes, but MuDo-NLL shows a clear advantage.

Table 12: Multi-domain Scalings results: average F1 and SHD ± standard deviation over 20 random seeds. The varying parameter and the best method for each metric in each row are shown in bold.

Setting	Value	F1		SHD	
		MuDo-NLL	MuDo-Cov	MuDo-NLL	MuDo-Cov
Directed edges					
Dimension					
n	**5**	**0.927 ± 0.125**	0.736 ± 0.182	**1.300 ± 2.227**	4.650 ± 2.925
n	**10**	**0.883 ± 0.196**	0.821 ± 0.076	**4.500 ± 5.708**	7.900 ± 3.227
n	**20**	**0.854 ± 0.280**	0.676 ± 0.261	**8.750 ± 12.586**	27.400 ± 15.091
Number of domains					
T/n	**1**	**0.883 ± 0.196**	0.821 ± 0.076	**4.500 ± 5.708**	7.900 ± 3.227
T/n	**2**	**0.911 ± 0.171**	0.805 ± 0.219	**3.550 ± 6.419**	7.350 ± 5.373
T/n	**3**	**0.976 ± 0.087**	0.880 ± 0.130	**1.100 ± 4.012**	5.350 ± 5.842
Sample size					
N	**1000**	**0.843 ± 0.101**	0.738 ± 0.126	**7.200 ± 4.360**	11.700 ± 5.151
N	**5000**	**0.883 ± 0.196**	0.821 ± 0.076	**4.500 ± 5.708**	7.900 ± 3.227
N	**50000**	**0.967 ± 0.108**	0.837 ± 0.064	**1.300 ± 4.207**	7.450 ± 3.086
Graph density					
k_X	**1.0**	**0.773 ± 0.111**	0.618 ± 0.089	**6.050 ± 3.561**	11.400 ± 3.719
k_X	**2.0**	**0.883 ± 0.196**	0.821 ± 0.076	**4.500 ± 5.708**	7.900 ± 3.227
k_X	**3.0**	**0.818 ± 0.268**	0.761 ± 0.204	**9.000 ± 12.490**	13.500 ± 11.062
Confounding density					
k_U	**0.5**	**0.883 ± 0.196**	0.821 ± 0.076	**4.500 ± 5.708**	7.900 ± 3.227
k_U	**1.0**	**0.803 ± 0.277**	0.793 ± 0.166	**7.350 ± 9.051**	7.950 ± 4.161
Noise distribution					
Noise	**gauss**	**0.883 ± 0.196**	0.821 ± 0.076	**4.500 ± 5.708**	7.900 ± 3.227
Noise	**exp**	**0.850 ± 0.176**	0.816 ± 0.070	**6.050 ± 5.031**	8.150 ± 3.133
Noise	**gumbel**	0.832 ± 0.218	**0.837 ± 0.065**	**6.650 ± 6.556**	7.600 ± 3.283
Bidirected edges					
Dimension					
n	**5**	**0.576 ± 0.142**	0.419 ± 0.088	**3.650 ± 1.496**	6.700 ± 1.174
n	**10**	**0.783 ± 0.142**	0.444 ± 0.107	**4.500 ± 5.615**	17.900 ± 5.418
n	**20**	**0.857 ± 0.156**	0.322 ± 0.099	**5.550 ± 7.884**	70.050 ± 46.479
Number of domains					
T/n	**1**	**0.783 ± 0.142**	0.444 ± 0.107	**4.500 ± 5.615**	17.900 ± 5.418
T/n	**2**	**0.864 ± 0.225**	0.556 ± 0.208	**3.000 ± 6.224**	12.650 ± 8.940
T/n	**3**	**0.933 ± 0.156**	0.707 ± 0.184	**1.600 ± 4.893**	6.550 ± 5.463
Sample size					
N	**1000**	**0.465 ± 0.092**	0.305 ± 0.057	**15.250 ± 4.241**	29.150 ± 3.360
N	**5000**	**0.783 ± 0.142**	0.444 ± 0.107	**4.500 ± 5.615**	17.900 ± 5.418
N	**50000**	**0.910 ± 0.185**	0.496 ± 0.143	**2.300 ± 7.035**	14.250 ± 5.609
Graph density					
k_X	**1.0**	**0.777 ± 0.165**	0.451 ± 0.107	**4.450 ± 5.558**	17.200 ± 5.672
k_X	**2.0**	**0.783 ± 0.142**	0.444 ± 0.107	**4.500 ± 5.615**	17.900 ± 5.418
k_X	**3.0**	**0.658 ± 0.272**	0.369 ± 0.098	**9.750 ± 10.543**	23.650 ± 7.184
Confounding density					
k_U	**0.5**	**0.783 ± 0.142**	0.444 ± 0.107	**4.500 ± 5.615**	17.900 ± 5.418
k_U	**1.0**	**0.823 ± 0.125**	0.704 ± 0.116	**8.250 ± 6.766**	15.100 ± 5.803
Noise distribution					
Noise	**gauss**	**0.783 ± 0.142**	0.444 ± 0.107	**4.500 ± 5.615**	17.900 ± 5.418
Noise	**exp**	**0.764 ± 0.166**	0.399 ± 0.082	**4.700 ± 4.438**	19.950 ± 4.395
Noise	**gumbel**	**0.714 ± 0.206**	0.476 ± 0.099	**6.100 ± 6.095**	15.350 ± 5.324

F.5 Comparing with FCI

We compare `LinADMG-MuDoNLL` , `LinADMG-diagGMM` , `LinADMG-UMNI` , `LinADMG-SP` , and FCI on four data-generating processes: DiagGMMs, UMN-Hard Interventions, UP-Observations, and MD-Scalings. In every experiment, we set $n = 10, N = 5000, T = 1, k_X = 2.0, k_U = 0.5$ and Gaussian noise. Results are averaged over 20 random seeds. Because FCI returns a partial ancestral graph (PAG), the ADMG estimated by each LinADMG method, as well as the ground-truth ADMG, is converted to its corresponding PAG, using a maximal ancestral projection where necessary. All methods are then evaluated in the common PAG representation using skeleton F1, arrowhead F1, and tail F1. Each bar reports the mean over the 20 seeds, while the error bar represents one standard deviation.

`LinADMG-diagGMM` provides the strongest overall performance, particularly for the DiagGMMs and UP-Observations DGPs, where it accurately recovers both the skeleton and the endpoint marks. `LinADMG-MuDoNLL` performs best in the MD-Scalings DGP. Under UMNI-Hard Interventions DGP, FCI is competitive in skeleton recovery, but the LinADMG methods generally recover arrowheads and tails more accurately. `LinADMG-SP` often recovers a reasonable skeleton but is less successful in finding endpoint orientation, while `LinADMG-UMNI` shows greater variability, especially for tail recovery.

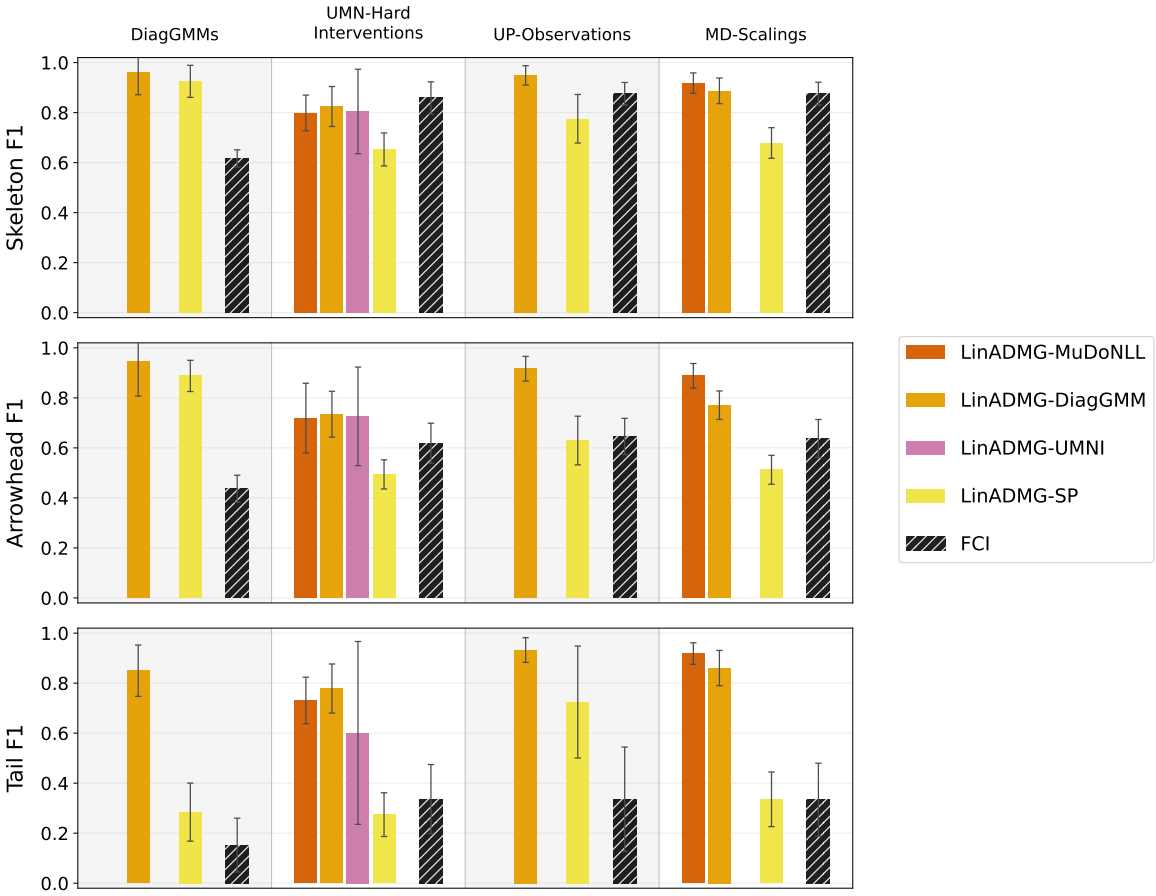

Figure 5: PAG recovery performance of `LinADMG-MuDoNLL` , `LinADMG-diagGMM` , `LinADMG-UMNI` , `LinADMG-SP` and FCI across four data-generating processes. The estimated and ground-truth ADMGs are converted to PAGs before evaluation, while FCI is evaluated using its native PAG output. The three panels report skeleton, arrowhead, and tail F1 scores. Bars show the mean and error bars indicate one standard deviation over 20 seeds. All experiments use $n = 10, N = 5000, T = n, k_X = 2, k_U = 0.5$.

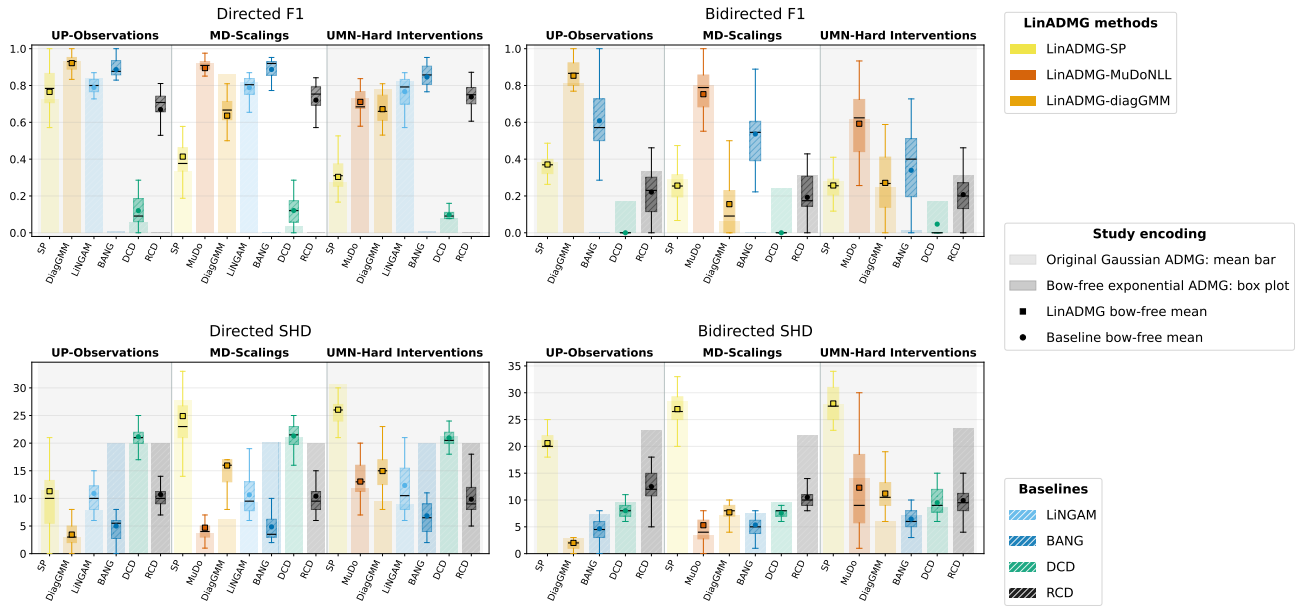

Figure 6: Comparison of ADMG estimation performance across the 20 method-DGP configurations shared by the original Gaussian-ADMG study and the bow-free exponential-noise study (20 seeds per configuration). Translucent bars show the mean for the original study as in Fig.3, while box plots show the bow-free study. Filled markers denote bow-free means: squares for LinADMG- methods and circles for baselines. LiNGAM is omitted from bidirected panels because it does not estimate bidirected edges.

F.6 Bow-free ADMG with exponential exogenous variables

LinADMG- methods are comparatively stable on all three data-generating processes, with `LinADMG-diagGMM` remaining strongest in general. As expected, the most pronounced changes occur for BANG and RCD as they required for non-Gaussian and bow-free ADMG: BANG's directed F1 rises from approximately zero to 0.886-0.846 across the three DGPs, while RCD rises from zero to 0.670-0.738. BANG also gains substantial bidirected accuracy. DCD remains weak for directed recovery and selects almost no bidirected edges in the bow-free UP-Observations and MD-Scalings settings.