DISTRACT: Degradation-Indexed Study of Textual Reference Across Competing Targets

Dan Zlotnikov Ohad Ben-Shahar
Ben-Gurion University of the Negev, Israel
danzlot@post.bgu.ac.il ben-shahar@cs.bgu.ac.il

arXiv:2610.28419v1 [cs.CV] 20 Oct 2026

Abstract

Vision-language trackers (VLTs) accept descriptions intended to assist tracking, but accepting text does not establish that its content changes tracking decisions, particularly as visual evidence deteriorates. We introduce DISTRACT, a matched-intervention benchmark that holds video, target, and distractor fixed while crossing four text priors with four image degradations along a controlled severity axis. We evaluate five VLTs on four datasets using behavioral contrasts, layerwise representation probes, and activation patching.

Molmo2 and JointNLT exhibit directed text use. Some effects contract under degradation, whereas others show no significant endpoint change after multiplicity correction. UVLTrack, MMTrack, and DUTrack respond only weakly to changes in text semantics, with no directed behavioral effect surviving FDR correction in the tested settings. Model-internal analyses reveal that at decision-near states, correct-versus-conflicting text is decodable in Molmo2 and JointNLT but not in the other three systems. In the responsive models, decodability peaks at late states where replacement produces little behavioral change, whereas replacing selected earlier states redirects conflict-associated trajectories toward the target. Because matched-correct donors do not significantly outperform unrelated donors, these interventions identify earlier intervention-sensitive states without establishing target-specific semantic rescue. DISTRACT thus exposes when and where text affects tracking under degraded visual evidence.

1. Introduction

Visual object tracking must preserve an object's identity as appearance, position, and context change over time. Vision–language trackers (VLTs) add a natural-language specification—for example, "the player in white moving toward the basket"—to distinguish a target from a similar competitor semantically [23, 37]. This capability is most valuable outside clean laboratory video: robots, autonomous cameras, and other perception–action systems encounter image degradations under which even subtle perturbations can cause severe perception failures [15, 26, 29].

Accuracy with a target-valid description, however, does not establish that a tracker uses language's semantic content to select the target over an on-screen competitor. Prior robustness studies show that degraded visibility harms trackers and VLMs [13, 27, 31, 35], while VLT studies show that description content and delivery can alter performance [1, 18, 21]. They generally vary visual quality and semantic validity separately, leaving open whether language rescues weak vision, misdirects selection toward a matching distractor, or is behaviorally redundant on the same clip.

We evaluate Molmo2 [3], JointNLT [41], UVLTrack [24], MMTrack [40], and DUTrack [22] on four diverse datasets, crossing four language priors with progressive visual degradation (Fig. 1). UVLTrack, MMTrack, and DUTrack show only small, inconsistent text-conditioned differences. Molmo2 and JointNLT respond directionally on clean video, but their severity trajectories are heterogeneous: some supported effects attenuate toward the null, whereas for others—including the single harmful clean-input effect—we detect no FDR-significant endpoint change.

However, behavior alone cannot reveal whether a description never reaches a tracking pathway, remains represented but does not affect localization, or loses its link to the spatial decision under degradation. We therefore complement the benchmark with layerwise, source-video-grouped probes across all five systems and targeted activation patching in the two responsive systems. This distinguishes linearly recoverable text-condition information from stages at which controlled state replacement changes the trajectory, while donor-content controls test whether that redirection is target-specific.

Our contributions are:

- A matched four-condition, multi-severity VLT protocol spanning four perturbations, five systems, and four datasets, while preserving models' supported point- or

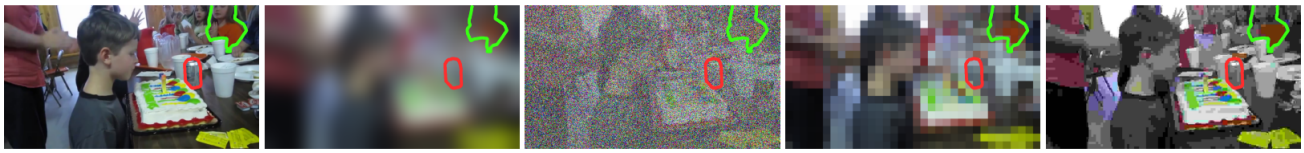

Figure 1. A real MeViS target–distractor pair (green = target; red = distractor) at the clean initialization frame and the terminal severity (S_5) of blur, noise, pixelation, and compression. Descriptions: target, "girl in red sitting on the right side of the table"; distractor, "The can next to the cake".

box-tracking interfaces.

- A five-model degradation study showing that text effects impact VLTs in meaningfully different ways, depending on the model, dataset, and perturbation.
- Layerwise probing and controlled state-replacement analyses that trace text-condition signals across all five systems and identify intervention-sensitive stages in Molmo2 and JointNLT.

2. Related Work

Evaluating language use in tracking. Language-annotated benchmarks (OTB99-Lang [23], LaSOT [7], TNL2K [37], MGIT [14], WebUAV-3M [38]) and VLT methods (SNLT [8], ModaMixer [9], JointNLT [41], MMTrack [40], UVLTrack [24], DUTrack [22]) are evaluated by trajectory accuracy under valid descriptions – accuracy given truthful language, not whether counterfactually changing that language changes target selection. Nor do the language diagnostics we are aware of isolate visibility: DTLLM-VLT and DTVLT vary detail and update frequency while keeping text target-consistent [18, 19]; VLTVerse crosses multiple text variants with naturally occurring sequence-level factors, rather than applying controlled degradation [21]; VLT-MI changes text and the box together [20]; and CLDTracker's substitutions are neither tied to an on-screen competitor nor tested across severity [1]. Our protocol addresses both gaps by independently varying text relation and severity on the same clip with a designated distractor.

Robustness under degraded visual evidence. Common-corruption evaluation applies graded transformations to a fixed source [11], but tracking benchmarks instead use naturally difficult sequences (LaSOT [7], TNL2K [37], AVisT [27], adverse-weather person tracking [15]) that confound visibility with target, motion, and scene difficulty. DBVOT evaluates six progressive defocus levels [13], while adversarial-blur work synthesizes motion-blur attacks against vision-only trackers [10], and VLM corruption suites retain a valid prompt while scoring one-shot answers [31, 35] – neither directly measures changing text authority or temporal identity errors. Our graded blur, noise, pixelation, and compression experiments test whether language guides tracking as visual evidence weakens.

Cross-modal conflict and behavioral consequence. CLASH and CrossCheck-Bench ask models to recognize controlled image–text contradictions, not to reveal their unprompted behavioral consequence in an ongoing spatial task [30, 34], and SIGNPOST-Bench shows that embedded scene text can bias still-image geolocation, but does not study a graded visual-degradation axis [17]. By holding scene content fixed, identifying the competing object, and measuring the direction and duration of tracking errors, our benchmark tests whether language remains actionable under imperfect sensing.

Mechanistic interpretability of grounding models. Causal tracing and activation patching have probed how general vision–language models use visual information, including BLIP [28] and object hallucination in large VLMs [16], but not, to our knowledge, vision–language tracking, where a spatial decision is emitted every frame rather than a single caption. We apply the same causal-tracing toolkit to five VLT systems to test whether text condition is decodable at decision-near representations and whether controlled interventions at selected text-reachable states alter localization.

3. Benchmark and Protocol

3.1. Datasets and pair construction

We use source videos from MeViS [5], a referring video object segmentation benchmark, and from the multi-object tracking benchmarks DanceTrack [33], SportsMOT [4], and PersonPath22 [32] (Table 1). MeViS provides its native referring expressions and masks. For the three MOT sources, human-written expressions and released point/mask trajectories come from the public Molmo2-VideoTrackEval annotation release [3]; the original MOT datasets provide the source videos. We derive initialization boxes from the released masks and enumerate ordered target–distractor pairs. A pair is eligible only if the target T and distractor D are co-visible in at least five frames and have mean co-visible mask IoU below 0.3. Every selected pair receives the same text conditions and perturbation protocol.

The MeViS selector takes one eligible pair per source video, while for the other three datasets, we construct outcome-enriched diagnostic stress subsets. Molmo2's

Dataset	Characterization	Pairs	Seq.
MeViS	referring seg., masks	120	120
DanceTrack	dance MOT, similar look	60	13
SportsMOT	athlete MOT, fast motion	30	19
PersonPath22	pedestrian MOT, crowded	30	20

Table 1. Per-dataset pair and source-video coverage. Pairs is the number of selected target–distractor pairs. Seq. is the number of distinct source videos selected.

clean conflicting-text run identifies cases for inclusion, with seeded random filling of the remaining quota, yielding 16/60 DanceTrack, 2/30 SportsMOT, and 6/30 PersonPath22 persistent-takeover cases in the final subsets. These subsets let us track how conflict-relevant cases evolve under degradation. Because pairs are partly selected on an outcome related to what we measure, naive effect estimates on this set risk running high; a supplementary inverse-probability-weighted check estimates the corresponding effects over the full eligible candidate pool; effect directions are preserved, although some naïve stress-subset estimates—particularly C_T—are larger. A second, more direct check in the supplementary reruns all four specialized trackers on the complete, outcome-independent eligible pool, reproducing the qualitative separation between JointNLT and the three small-effect specialized trackers.

3.2. Targets and distractors

Let a video contain an ordered pair of distinct annotated objects (T, D), where T is the object initialized for tracking and D is a co-visible distractor. At evaluated frame t, their annotated spatial supports are $\mathcal{M}_{T,t}$ and $\mathcal{M}_{D,t}$. Each pair has a human-written expression e_T referring to T and an expression e_D referring to D.

3.3. Text conditions

We vary the text input independently of the video through four conditions:

- **Neutral** (c_{neutral}): a generic tracking instruction with no target-identifying description.
- **Unrelated** ($c_{\text{unrelated}}$): a real expression e_U from another clip, that is not annotated as referring to either T nor D.
- **Correct** (c_{correct}): the target expression e_T.
- **Conflicting** ($c_{\text{conflicting}}$): the distractor expression e_D.

The unrelated expression is sampled once from a different annotated clip with the same category label as the pair, then fixed across every model, condition, and severity.

3.4. Model initialization, inputs, and outputs

Every run is anchored to the same clean annotated target instance and uses the same fixed post-initialization frames for every model, text condition, and perturbation. Molmo2 is initialized with a target point and predicts points; UVLTrack, MMTrack, JointNLT, and DUTrack are initialized with a tight axis-aligned box and predict boxes, evaluating which object is selected rather than boundary precision. Points/boxes are derived from mask centroids/extents when masks are available, and otherwise from the corresponding released box trajectory. We additionally report native-box and absolute-performance metrics in the supplementary.

Molmo2 receives text through the scaffold "Track the object indicated by the point," with descriptive conditions appending "The target is: e." UVLTrack, MMTrack, and JointNLT receive the raw expression e directly, with the generic token `object` as their neutral-condition input. DUTrack shares this interface but normally regenerates a description from the current frame, overwriting the assigned condition; we use *static-external* mode and reset all adapter state between pair–condition–severity runs. A supplementary dynamic-mode audit verifies this condition-preservation control.

3.5. Perturbations and fixed severity schedule

Each corruption is applied independently to every post-initialization frame; the initialization frame remains clean and corruptions are never composed. We use Gaussian blur, additive Gaussian noise, pixelation, and JPEG compression—common corruption families with fixed, ordered severity ladders in robustness benchmarks [11]. Table 2 gives our disclosed, perturbation-specific parameter schedules; these were chosen to produce visibly distinct changes at the native video resolution. $S0$ is an exact no-op. Severity is an ordinal dose axis within each perturbation, not a claim of perceptual equivalence across perturbation types; the supplementary reports descriptive PSNR/SSIM diagnostics. For noise, a video–frame hash fixes one standard-normal field shared across conditions and severities, so contrasts never compare different noise realizations.

Perturbation	Fixed schedule values, $S0$–$S5$
Blur, k_s	0, .004, .008, .016, .032, .064
Noise, σ_s	0, .12, .20, .30, .45, 1.5
Pixelation, r_s	1, .35, .25, .18, .12, .04
Compression, q_s	no-op, 15, 10, 7, 3, 1

Table 2. Fixed severity schedule.

3.6. Outcomes and metrics

Each post-initialization evaluated sample receives exactly one of the following labels:

- **absent** (no usable predicted point at that timestamp).
- **target** ($\hat{p}_t \in \mathcal{M}_{T,t}$).
- **distractor** ($\hat{p}_t \notin \mathcal{M}_{T,t}$ and $\hat{p}_t \in \mathcal{M}_{D,t}$).
- **other**.

Target takes precedence when masks overlap, and absent samples remain in the denominator of all frame-level rates.

For pair p, condition c, and severity s, let $A_T^p(c,s)$, $A_D^p(c,s)$, $A_O^p(c,s)$, and $A_\emptyset^p(c,s)$ be the fractions of post-initialization samples classified as target, distractor, other, and absent. They sum to one. We aggregate frames within a pair before aggregating pairs, preventing longer tracks from receiving greater weight. In addition to condition-specific rates, the primary contrastive semantic effects anchor every comparison at the *correct* condition – the observed target accuracy when the model receives the ground-truth referring expression – and measure the gap to each of the other three conditions:

$$B_T(s) = A_T(\text{correct}, s) - A_T(\text{unrelated}, s), \quad (1)$$

$$N_T(s) = A_T(\text{correct}, s) - A_T(\text{neutral}, s), \quad (2)$$

$$C_T(s) = A_T(\text{correct}, s) - A_T(\text{conflicting}, s). \quad (3)$$

B_T is the correct-versus-unrelated effect (a same-type but unrelated description), N_T is the correct-versus-neutral effect (no target-identifying description), and C_T is the correct-versus-conflicting effect; each is signed and can be negative. Two further effects, the directed distractor-capture effect S_D and the conflict-specific target-loss effect L_T, are defined in the supplementary alongside the full endpoint test family.

Severity-integrated effect magnitude. To summarize the size of a condition effect across the severity ramp, we use integrated absolute deviation (IAD). Severities are indexed $s \in \{0, \ldots, 5\}$ (Table 2); for conditions c_1, c_2, let $\Delta_{c_1,c_2}(s) = A_T(c_1, s) - A_T(c_2, s)$ and define

$$\text{IAD}(c_1, c_2) = \sum_{s=0}^{4} \tfrac{1}{2}\left[|\Delta_{c_1,c_2}(s)| + |\Delta_{c_1,c_2}(s{+}1)|\right]. \quad (4)$$

Taking the absolute value *before* integrating preserves effects that change direction over severity. IAD is nonnegative and symmetric in its two conditions; with six severity levels, $\text{IAD}/5$ is the mean absolute condition gap per severity step. We report all $\binom{4}{2} = 6$ unordered condition contrasts as the upper triangle of a 4×4 matrix per model and dataset. Per-severity signed contrasts (B_T, N_T, and C_T) remain the primary directional measurements and are used for endpoint tests.

Behavioral inference. We compute every contrast within pair and use source video as the resampling unit, keeping all selected pairs from one video together. We report 95% source-video-clustered bootstrap CIs and use paired source-video sign-flip tests for clean effects and severity-endpoint changes. Benjamini–Hochberg correction is applied within each stated hypothesis family; IAD is descriptive rather than inferential. The supplementary specifies the resampling algorithm, test statistics, hypothesis families, and formulas in full.

3.7. Model-internal and mechanistic measurements

For internal analyses, we save compact pooled states at model-specific functional stages—text encoding/projection, early and late multimodal fusion, and the decision-near localization state—rather than treating raw layer indices as comparable across architectures. Exact hooked tensors, layer counts, and semantics per model are listed in the supplementary.

At a matched video frame, let h_l^c be a pooled state at stage l under text condition c. Let m_t be the target–distractor output margin at evaluated frame t,

$$m_t = \text{IoU}(\hat{b}_t, b_{T,t}) - \text{IoU}(\hat{b}_t, b_{D,t}), \quad (5)$$

with the analogous target/distractor label used for point output. We use three procedures to characterize the text pathway before defining the metrics computed from them.

Text-condition probe. Following low-capacity probing practice [2, 12], at every saved stage we standardize the pooled state using statistics fit only on the training fold, project it onto a training-fold SVD basis (at most 20 components), and fit a ridge-regularized linear classifier for correct-versus-conflicting condition. Cross-validation holds out complete source videos, so the probe cannot exploit sentence- or token-identity shortcuts shared within a video; group-preserving label permutations provide its null distribution. We report the probe's held-out balanced accuracy as **Text BA** at stage l (0.5 is chance). Within conflicting runs, we repeat the same procedure for the observed target-versus-distractor outcome. This establishes whether the prompt distinction is linearly present in a state, not whether the tracker acts on it; causal relevance is tested with the interventions below [6].

Activation patching. We intervene on the forward pass directly, following intervention-based interpretability, causal-tracing, and activation-patching work [25, 36, 39]. For a fixed video, template, and evaluated-frame sequence, we run the network under a *target* condition (typically conflicting) but replace its unpooled native state at the identified text boundary with the corresponding state from an independently run *source* condition ($\text{source} \in \{\text{correct}, \text{neutral}, \text{unrelated}, \text{zero}\}$) on the same visual input,

$$h_l^{\text{target}} \leftarrow h_l^{\text{source}}, \quad (6)$$

before continuing the forward pass and subsequent sequential tracking unmodified. Every patch site is validated with two controls: an exact self-patch ($\text{source} = \text{target}$), required to reproduce the unpatched trajectory exactly, and an unpatched repeat, which confirms that reset and replay are themselves deterministic.

For each activation-patching intervention, we report the paired change in trajectory-mean m_t and the count of target–distractor switches, with source-video bootstrap CIs

and paired sign-flip tests; these controls establish patch integrity and deterministic replay. They support interpreting differences as consequences of the intervention, but do not by themselves establish target-specific semantic mediation; that stronger interpretation requires donor-content controls.

Full-depth layer sweeps. Rather than committing in advance to one decision-near stage, we apply the text-condition probe at every saved functional stage for every model, (for Molmo2, at all 37 pre-coordinate decoder depths). This traces the depth at which the correct/conflicting distinction becomes linearly decodable, from text encoding through fusion to the localization output, reducing the risk that a null result at one selected stage is incorrectly generalized to the model's entire saved pathway.

4. Results

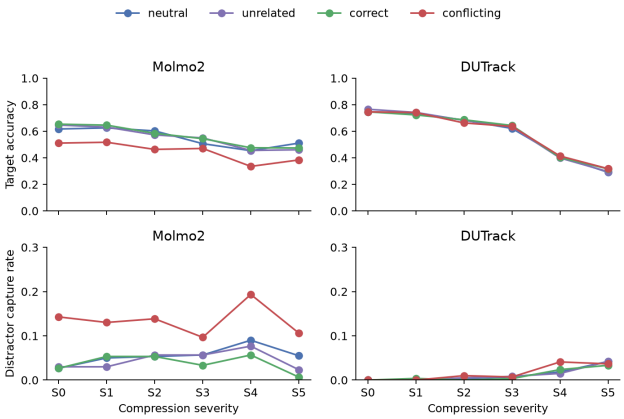

Figure 2. Representative compression trajectories on PersonPath22. Molmo2 (left) separates text conditions in target accuracy and distractor capture; DUTrack (right) remains nearly unchanged.

Fig. 2 illustrates the behavioral contrast that the aggregate analyses quantify: Molmo2's target accuracy and distractor capture diverge across text conditions, whereas DUTrack's trajectories closely overlap. We use the full matched grid below to assess the breadth and uncertainty of this qualitative pattern across the tested settings.

4.1. UVLTrack, MMTrack, and DUTrack show limited behavioral sensitivity to text semantics

Fig. 3 plots each system's severity-integrated absolute target-accuracy effect as a 4×4 grid over the four text conditions, one grid per model on a shared color scale, where larger, warmer cells mark condition pairs whose accuracy diverges more across the full clean-to-severe range as text effects differ. Two grids show consistently larger effects; the remaining three stay pale in every cell. This qualitative split holds numerically: the three less-responsive systems' largest target-accuracy

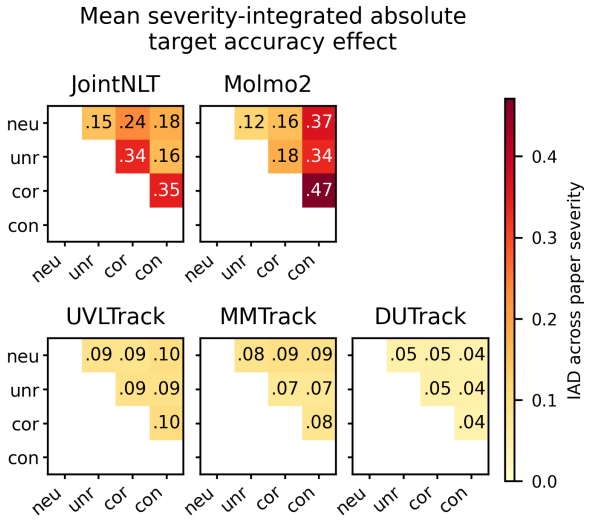

Figure 3. Mean severity-integrated absolute target-accuracy effect ($\mathrm{IAD}(c_1, c_2)$, Eq. 4) across all four perturbations and four datasets, per model. Top row: JointNLT, Molmo2. Bottom row: UVLTrack, MMTrack, DUTrack, same color scale. Condition labels abbreviate neutral/unrelated/correct/conflicting.

IAD cells are 0.05 (DUTrack), 0.09 (MMTrack), and 0.10 (UVLTrack), versus JointNLT's mean 0.24 (maximum 0.35) and Molmo2's mean 0.27 (maximum 0.47), both peaking at the correct–conflicting contrast. These observed separations remain small after taking absolute values before integration, so their small IAD values are not caused by cancellation between opposite-signed severities. The corresponding distractor-capture IAD is also small (0.02/0.05/0.08 for DUTrack/MMTrack/UVLTrack vs. 0.29/0.37 for JointNLT/Molmo2), again largest at the correct–conflicting contrast, consistent with conflict-directed capture rather than a uniform output-quality change.

Thus, we detect no FDR-supported directed semantic-text effect in these three systems, without implying absent internal language processing yet. Their 62–68% correct-text S_0 target accuracy argues against a simple floor/ceiling explanation and leaves room for change in either direction.

4.2. Directed language effects degrade unevenly, including harmful effects

The deployment-relevant question for the two responsive systems is not merely whether accuracy falls with degradation, but whether a CI-supported clean-input text effect reliably weakens as visual evidence degrades. We test every effect with a CI-supported S_0 signal (95% CI excludes zero) for FDR-significant contraction toward zero across all four main perturbations: 14 target-accuracy cases and 4 distractor-capture cases, spanning both responsive systems and all four datasets (Table 4).

Of the 14 target-accuracy effects with a CI-supported

	Molmo2			JointNLT			UVLTrack			MMTrack			DUTrack		
Dataset	C_T	B_T	N_T	C_T	B_T	N_T	C_T	B_T	N_T	C_T	B_T	N_T	C_T	B_T	N_T
DanceTrack	$\underline{+25\%}$	$\underline{+13\%}$	$\underline{+8\%}$	$\underline{+11\%}$	$\underline{+11\%}$	$+6\%$	$+1\%$	-1%	$+2\%$	-1%	$+0\%$	$+0\%$	$+0\%$	$+1\%$	$+0\%$
SportsMOT	$+6\%$	$+0\%$	$+4\%$	-2%	$\underline{-11\%}$	-6%	-1%	$+1\%$	-1%	$+1\%$	-2%	-4%	-1%	-2%	$+0\%$
PersonPath22	$+14\%$	$+1\%$	$+4\%$	$\underline{+27\%}$	$\underline{+18\%}$	$\underline{+17\%}$	$+5\%$	$\underline{+5\%}$	$+5\%$	$+2\%$	$+1\%$	$\underline{+3\%}$	$+0\%$	-2%	$+0\%^{\dagger}$
MeViS	$\underline{+8\%}$	$\underline{+4\%}$	$+2\%$	$\underline{+14\%}$	$\underline{+11\%}$	$\underline{+8\%}$	$+2\%$	$+1\%$	-1%	$+0\%$	$+0\%$	$+1\%$	$+0\%$	$+0\%$	$+0\%$

Table 3. Clean-input (S_0) text effects for all five systems on all four datasets. Underlined: 95% source-video cluster-bootstrap CI excludes zero. Among UVLTrack/MMTrack/DUTrack: 0/12 C_T, 1/12 B_T, 2/12 N_T CIs exclude zero; the sole B_T case (UVLTrack/PersonPath22) fails the clustered sign-flip test ($p = .08$, BH $q = .668$), so is not treated as a reproducible effect. † Near-zero-variance bootstrap artifact, not interpreted.

Contraction level (of 4 perturbation ramps)	Cases (of 14)
4/4	1
1–3/4	3
0/4	10

Table 4. FDR-significant contraction toward zero, across the 14 target-accuracy cases with a CI-supported clean-input signal.

clean-input value, four show FDR-significant contraction under at least one perturbation, including one that contracts under all four; the remaining ten show no FDR-significant endpoint change under any perturbation. None of Molmo2's five cases yields a significant contraction after correction. These nonsignificant change tests should not be interpreted as evidence that the effects remain unchanged.

One of the ten cases is JointNLT's SportsMOT B_T, the study's only significantly harmful clean-input effect (-11.2%, 95% CI excluding zero). Under the primary point-in-support metric, its estimated S5 changes substantially reduce or reverse the clean effect, but these changes do not survive FDR correction.

The distractor-capture effect S_D exhibits a similar inferential pattern. Of four model–dataset combinations with a CI-supported clean-input misdirection effect, JointNLT/MeViS shows significant contraction under three of four perturbations. For Molmo2/PersonPath22, the largest clean capture effect in the study ($+11.3$ pp), no endpoint change survives correction; nevertheless, the point estimates contract almost completely under blur, noise, and pixelation and remain more clearly positive under compression.

Thus, contraction is established in some settings, while the remaining changes are statistically inconclusive rather than evidence of persistence. Their behavior varies across models, datasets, effects, metrics, and perturbations.

4.3. Model internals

The behavioral results identify FDR-supported directed text effects for JointNLT and Molmo2, in contrast to the demonstrably smaller effects observed for MMTrack, DUTrack, and UVLTrack. They do not, however, localize *where* the supported effects arise. We probe this with the internal-mechanism captures of Sec. 3.7, restricted to MeViS: its pairs are the only ones with one independent source video each (Table 1); its mask-native annotations pool a hooked state over exactly the target/distractor's tokens, not a background-containing box; and its native task is itself referring expression segmentation.

Of the 120 one-pair-per-video MeViS entries, 115 have at least two co-visible target/distractor frames. JointNLT's causal study further uses an outcome-enriched 90-video subset (23 persistent captures, 67 retention controls) so the rarer capture outcome is well represented.

4.3.1. Decision-near text decodability aligns with the behavioral split

For every system we probe Text BA (Sec. 3.7) at its most decision-near hooked stage. Table 5 reports the FDR-corrected result.

Model	Decision-near stage	Text BA	95% CI	q
Molmo2	final coord. query (L36)	$\underline{0.574}$	[0.543,0.605]	0.015
JointNLT	late-fusion $T-D$	$\underline{0.695}$	[0.647,0.744]	0.002
UVLTrack	pre-box-head search	0.530	[0.505,0.554]	0.174
MMTrack	coord. decoder input	0.488	[0.469,0.509]	0.898
DUTrack	pre-CENTER-head visual	0.504	[0.499,0.512]	0.506

Table 5. Correct-vs-conflicting Text BA (0.5=chance) at each system's decision-near hooked stage (115-video MeViS manifest). Underlined: FDR-significant ($q < .05$).

Only Molmo2 and JointNLT carry an FDR-robust text-condition signal to the selected decision-near state: MMTrack and DUTrack are near chance (95% CIs spanning 0.5); UVLTrack's CI ($[0.505, 0.554]$) is nominally above chance but neither FDR-significant nor precise enough for equivalence – converging with Sec. 4.1's behavioral result that these same three systems carry no reliable signal at their decision-near states.

Fig. 4 sweeps every text-reachable state in these three systems, rather than only the decision-near endpoint in Table 5.

The full sweeps show different routes to a decision-near effect that is not FDR-supported. UVLTrack is significant at 14/25 states, spanning its text encoder and fusion stages, but not at the pre-localization endpoint. MMTrack is significant at 6/37 states, all in its text encoder/projection and none at or beyond fusion. DUTrack has no FDR-significant state in its 53-state sweep (33 are formally chance-equivalent).

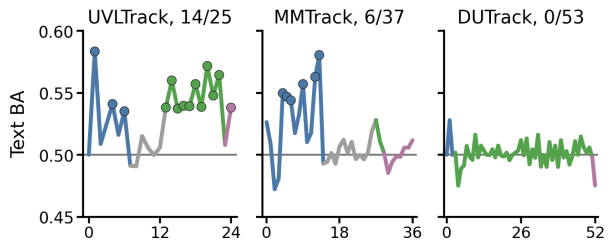

Figure 4. Exhaustive correct-vs-conflicting Text BA sweeps for the three behaviorally small-effect systems. Colors denote functional stage: text encoder/projection (blue), pre-fusion vision (gray), multimodal fusion (green), and localization/pre-box-head (purple). Filled circles mark FDR-significant states; titles report the number of FDR-surviving states and the total evaluated.

4.3.2. Layerwise text-condition decodability follows distinct routes in the two responsive systems

Fig. 5 sweeps correct-vs-conflicting Text BA through each Molmo2 and JointNLT's functional pathway.

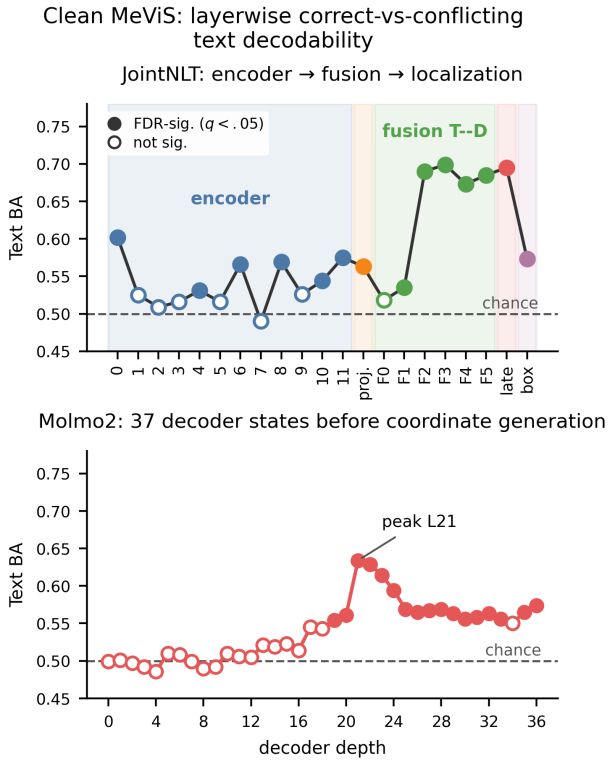

Figure 5. Exhaustive correct-vs-conflicting Text BA sweeps for Molmo2 and JointNLT. Solid trajectory segments indicate the functional state at their left endpoint: text encoder/projection (blue), pre-fusion vision (gray), multimodal fusion (green), and localization/pre-box-head (purple).

JointNLT carries text-condition information through its encoder–fusion pathway, though the early encoder signal is intermittent. The fusion pathway starts weak, then turns consistently strong from its second block onward, staying robust at the selected late-fusion state (0.695, $q = .002$) – text is linearly recoverable near localization, not uniformly strong throughout. Molmo2 shows a different pattern: no decoder state through layer 16 is FDR-significant; decodability becomes significant at layer 19, peaks at layer 21 (0.634), and decreases but remains significant (except one) through the final state. This onset is consistent with its token order: image tokens precede the appended description, so condition information can emerge only after the description is processed. Thus, both systems carry decodable text to decision-near states through distinct pathways; intervention effects are tested next.

4.3.3. State-replacement interventions redirect JointNLT and Molmo2 trajectories

Fig. 6 contrasts the full Text-BA trajectory with matched-correct state replacement at each tested depth.

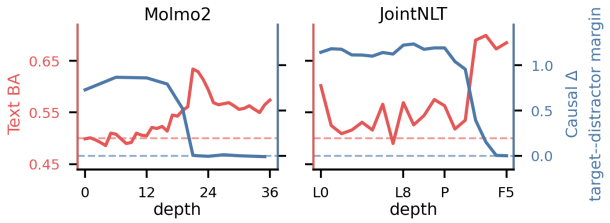

Figure 6. Text-condition decodability and matched-correct state-replacement effect across depth. Red, left ordinate: held-out Text BA. Blue, right ordinate: mean paired change in target–distractor margin after replacing a conflicting-run state with its same-video correct-run state. Both evaluated on the persistent-capture stratum. Dashed horizontal lines mark chance Text BA and zero replacement effect.

Both systems dissociate linear decodability from intervention efficacy: replacement produces larger observed trajectory changes at selected earlier sites and smaller changes at each system's Text-BA peak, so the most linearly decodable state is not necessarily the most effective intervention point.

Fig. 7 shows this recovery per persistent-distractor-capture video, at each model's best-recovery site.

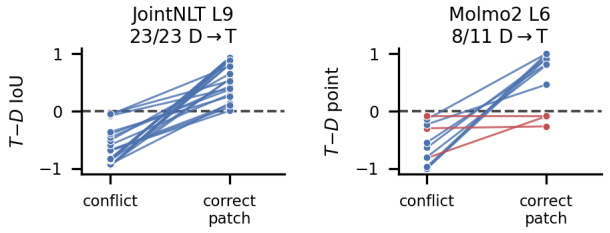

Figure 7. Matched-correct state replacement redirects conflict-associated trajectories in both responsive systems, each shown at its best-recovery site; sites maximize the observed capture-stratum effect on the displayed data. Each line connects the unpatched conflicting run with its correct-text-state patch; blue: distractor-to-target transition, red: none.

JointNLT's BERT-layer-9 patch reverses all 23/23 per-

sistent captures from distractor- to target-preferred, and Molmo2's layer-6 patch does so in 8 of 11. On this stratum, however, neither system's correct-donor effect significantly exceeds an unrelated donor's once corrected (JointNLT $q = .124$, Molmo2 $q = .092$). Both correct and unrelated donors can redirect the conflict-associated trajectory, so the grounding-to-output association interventions do not establish target-specific semantic content as the cause of the redirection.

4.3.4. Under degradation, JointNLT retains fused text information while grounding-to-output association becomes inconsistent

Sec. 4.2 showed that JointNLT's behavioral text effects sometimes contract under degradation, although late-fusion Text BA (Fig. 5) remains above chance. The grounding-to-output association asks whether the pre-localization target–distractor margin predicts the final output margin. Figure 8 traces this association, and Table 6 contrasts its S_5 value with late visual separation and Text BA.

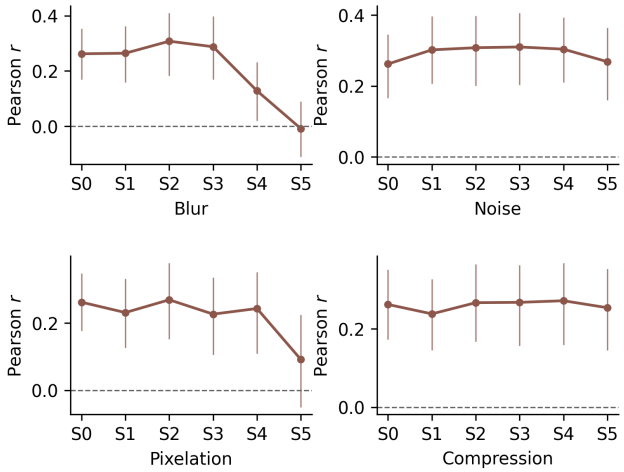

Figure 8. JointNLT's pre-localization state–output association (Pearson r) across S_0–S_5; dashed line: zero.

Perturbation	late vis. sep.	late fused-text BA	output assoc. r
blur	.173	.582	$-.009$ [$-.111, .089$]
noise	.042	.566	$+.268$ [.160, .363]
pixelation	.148	.539	$+.091$ [$-.051, .225$]
compression	.253	.577	$+.254$ [.145, .353]

Table 6. Terminal S_5 JointNLT summaries. Underlined: CI excludes zero.

This is not a uniform collapse (Fig. 8, Table 6): blur and pixelation fall to a CI spanning zero by S_5 while fused-text decodability remains nominally 4–8 points above chance throughout; noise and compression retain a CI-supported positive association. Nor does it track visual degradation – compression has the largest terminal visual separation (.253) yet one of the two CI-supported positive state–output associations. Thus, the association does not vary monotonically with this visual-separation summary, arguing against a simple one-to-one account while not ruling out visual mediation.

4.4. Limitations

DISTRACT uses controlled comparative evidence across five systems, four datasets, four perturbations, and fixed S_0–S_5 schedules; native interfaces make cross-system differences whole-system rather than architectural, and three datasets are stratified toward conflict-relevant pairs, so raw error frequencies are not prevalence estimates. Controlled perturbations and four priors isolate image quality and directed target–distractor semantics, respectively, but not natural adverse conditions or relational, temporal, multi-turn, and pragmatic language. They also do not reproduce the coupled weather, scene, and sensor shifts of deployment. Model-internal clean sweeps use 115 MeViS videos, and degraded internals are observational and JointNLT-only; state replacement localizes sensitivity, not target-specific semantic mediation. Broader models, datasets, and held-out causal tests are therefore needed to further validate these claims.

5. Discussion and Future Work

DISTRACT treats text as an experimental variable, testing whether counterfactual descriptions redirect tracking decisions and how these effects change under degraded vision. Three of five systems show limited sensitivity to text semantics and no directed behavioral effect surviving FDR correction, demonstrating that correct-text accuracy alone does not establish discriminative language use. Molmo2 and JointNLT show supported directed effects, but their severity trajectories vary across settings: some contract significantly, while others remain inconclusive after correction.

Internal analyses localize this distinction. Text condition is decodable in four of five systems, showing that decodability alone is insufficient evidence of behavioral use. In the responsive systems, replacement at selected earlier sites produces larger changes than at peak-decoding states, but correct donors do not outperform unrelated donors, identifying sensitive stages without establishing target-specific semantic mediation. Neither architecture nor clean-input responsiveness consistently anticipates degradation behavior across the tested panel; robustness and correct language use should therefore be evaluated separately.

Future work. Our internal analyses are restricted to MeViS and, under degradation, to JointNLT (Sec. 4.4). Extending them across datasets, models, and newer VLT architectures could reveal whether architectural properties explain the responsive/non-responsive split. Additional datasets and relational, temporal, and multi-turn text conditions are natural extensions of DISTRACT.

References

[1] Mohamad Alansari, Sajid Javed, Iyyakutti Iyappan Ganapathi, Sara Alansari, and Muzammal Naseer. CLDTracker: A comprehensive language description for visual tracking. *Information Fusion*, 124:103374, 2025.

[2] Yonatan Belinkov. Probing classifiers: Promises, shortcomings, and advances. *Computational Linguistics*, 48(1):207–219, 2022.

[3] Christopher Clark, Jieyu Zhang, Zixian Ma, Jae Sung Park, Rohun Tripathi, Sangho Lee, Mohammadreza Salehi, Jason Ren, Chris Dongjoo Kim, Yinuo Yang, Vincent Shao, Yue Yang, Weikai Huang, Ziqi Gao, Taira Anderson, Jianrui Zhang, Jitesh Jain, George Stoica, Ali Farhadi, and Ranjay Krishna. Molmo2: Open weights and data for vision-language models with video understanding and grounding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 28652–28668, 2026.

[4] Yutao Cui, Chenkai Zeng, Xiaoyu Zhao, Yichun Yang, Gangshan Wu, and Limin Wang. SportsMOT: A large multi-object tracking dataset in multiple sports scenes. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 9921–9931, 2023.

[5] Henghui Ding, Chang Liu, Shuting He, Xudong Jiang, and Chen Change Loy. MeViS: A large-scale benchmark for video segmentation with motion expressions. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 2694–2703, 2023.

[6] Yanai Elazar, Shauli Ravfogel, Alon Jacovi, and Yoav Goldberg. Amnesic probing: Behavioral explanation with amnesic counterfactuals. *Transactions of the Association for Computational Linguistics*, 9:160–175, 2021.

[7] Heng Fan, Liting Lin, Fan Yang, Peng Chu, Ge Deng, Sijia Yu, Hexin Bai, Yong Xu, Chunyuan Liao, and Haibin Ling. LaSOT: A high-quality benchmark for large-scale single object tracking. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5374–5383, 2019.

[8] Qi Feng, Vitaly Ablavsky, Qinxun Bai, and Stan Sclaroff. Siamese natural language tracker: Tracking by natural language descriptions with siamese trackers. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5851–5860, 2021.

[9] Mingzhe Guo, Zhipeng Zhang, Heng Fan, and Liping Jing. Divert more attention to vision-language tracking. In *Advances in Neural Information Processing Systems*, pages 4446–4460, 2022.

[10] Qing Guo, Ziyi Cheng, Felix Juefei-Xu, Lei Ma, Xiaofei Xie, Yang Liu, and Jianjun Zhao. Learning to adversarially blur visual object tracking. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 10839–10848, 2021.

[11] Dan Hendrycks and Thomas Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. In *Proceedings of the International Conference on Learning Representations*, 2019.

[12] John Hewitt and Percy Liang. Designing and interpreting probes with control tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing*, pages 2733–2743, 2019.

[13] Fen Hu, Peng Yang, Jie Dou, and Lei Dou. Detecting and tracking moving objects in defocus blur scenes. *Journal of Visual Communication and Image Representation*, 103: 104259, 2024.

[14] Shiyu Hu, Dailing Zhang, Meiqi Wu, Xiaokun Feng, Xuchen Li, Xin Zhao, and Kaiqi Huang. A multi-modal global instance tracking benchmark (MGIT): Better locating target in complex spatio-temporal and causal relationship. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2023.

[15] Abdulrahman Kerim, Ufuk Celikcan, Erkut Erdem, and Aykut Erdem. Using synthetic data for person tracking under adverse weather conditions. *Image and Vision Computing*, 111:104187, 2021.

[16] Qiming Li, Zekai Ye, Xiaocheng Feng, Weihong Zhong, Weitao Ma, and Xiachong Feng. Causal tracing of object representations in large vision language models: Mechanistic interpretability and hallucination mitigation. *arXiv preprint arXiv:2511.05923*, 2025.

[17] Sirun Li, Minghao Liu, Ling Dai, Yong Li, Haoxin Lyu, Junting Zhou, and Fan Zhang. SIGNPOST-Bench: Benchmarking text-vision conflict resolution in multimodal large language models. *arXiv preprint arXiv:2608.04244*, 2026.

[18] Xuchen Li, Xiaokun Feng, Shiyu Hu, Meiqi Wu, Dailing Zhang, Jing Zhang, and Kaiqi Huang. DTLLM-VLT: Diverse text generation for visual language tracking based on LLM. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops*, pages 7283–7292, 2024.

[19] Xuchen Li, Shiyu Hu, Xiaokun Feng, Dailing Zhang, Meiqi Wu, Jing Zhang, and Kaiqi Huang. DTVLT: A multi-modal diverse text benchmark for visual language tracking based on LLM. *arXiv preprint arXiv:2410.02492*, 2024.

[20] Xuchen Li, Shiyu Hu, Xiaokun Feng, Dailing Zhang, Meiqi Wu, Jing Zhang, and Kaiqi Huang. Visual language tracking with multi-modal interaction: A robust benchmark. *arXiv preprint arXiv:2409.08887*, 2024.

[21] Xuchen Li, Shiyu Hu, Xiaokun Feng, Dailing Zhang, Meiqi Wu, Jing Zhang, and Kaiqi Huang. How texts help? a fine-grained evaluation to reveal the role of language in vision-language tracking. *arXiv preprint arXiv:2411.15600*, 2024.

[22] Xiaohai Li, Bineng Zhong, Qihua Liang, Zhiyi Mo, Jian Nong, and Shuxiang Song. Dynamic updates for language adaptation in visual-language tracking. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 19165–19174, 2025.

[23] Zhenyang Li, Ran Tao, Efstratios Gavves, Cees G. M. Snoek, and Arnold W. M. Smeulders. Tracking by natural language specification. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 6495–6503, 2017.

[24] Yinchao Ma, Yuyang Tang, Wenfei Yang, Tianzhu Zhang, Jinpeng Zhang, and Mengxue Kang. Unifying visual and vision-language tracking via contrastive learning. In *Pro-*

ceedings of the AAAI Conference on Artificial Intelligence, pages 4107–4116, 2024.

[25] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems*, 2022.

[26] Claudio Michaelis, Benjamin Mitzkus, Robert Geirhos, Evgenia Rusak, Oliver Bringmann, Alexander S. Ecker, Matthias Bethge, and Wieland Brendel. Benchmarking robustness in object detection: Autonomous driving when winter is coming. *arXiv preprint arXiv:1907.07484*, 2019.

[27] Mubashir Noman, Wafa H. Al Ghallabi, Daniya Najiha, Christoph Mayer, Akshay Dudhane, Martin Danelljan, Hisham Cholakkal, Salman Khan, Luc Van Gool, and Fahad Shahbaz Khan. AVisT: A benchmark for visual object tracking in adverse visibility. In *Proceedings of the British Machine Vision Conference*, 2022.

[28] Vedant Palit, Rohan Pandey, Aryaman Arora, and Paul Pu Liang. Towards vision-language mechanistic interpretability: A causal tracing tool for BLIP. *arXiv preprint arXiv:2308.14179*, 2023.

[29] Zachary Pezzementi, Trenton Tabor, Samuel Yim, Jonathan K. Chang, Bill Drozd, David Guttendorf, Michael Wagner, and Philip Koopman. Putting image manipulations in context: Robustness testing for safe perception. In *Proceedings of the IEEE International Symposium on Safety, Security, and Rescue Robotics*, pages 1–8, 2018.

[30] Teodora Popordanoska, Jiameng Li, and Matthew B. Blaschko. CLASH: A benchmark for cross-modal contradiction detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Findings*, pages 6051–6061, 2026.

[31] Rohit Saxena, Alessandro Suglia, and Pasquale Minervini. VLM-RobustBench: A comprehensive benchmark for robustness of vision-language models. In *Proceedings of the International Conference on Machine Learning*, 2026.

[32] Bing Shuai, Alessandro Bergamo, Uta Buechler, Andrew G. Berneshawi, Alyssa Boden, and Joseph Tighe. Large scale real-world multi-person tracking. In *Proceedings of the European Conference on Computer Vision*, pages 504–521, 2022.

[33] Peize Sun, Jinkun Cao, Yi Jiang, Zehuan Yuan, Song Bai, Kris Kitani, and Ping Luo. DanceTrack: Multi-object tracking in uniform appearance and diverse motion. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 20993–21002, 2022.

[34] Baoliang Tian, Yuxuan Si, Jilong Wang, Lingyao Li, Zhongyuan Bao, Zineng Zhou, Tao Wang, Sixu Li, Ziyao Xu, Mingze Wang, Zhouzhuo Zhang, Zhihao Wang, Yike Yun, Ke Tian, Ning Yang, and Minghui Qiu. CrossCheck-Bench: Diagnosing compositional failures in multimodal conflict resolution. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 25887–25895, 2026.

[35] Muhammad Usama, Syeda Aisha Asim, Syed Bilal Ali, Syed Talal Wasim, and Umair Bin Mansoor. Analysing the robustness of vision-language models to common corruptions. *arXiv preprint arXiv:2504.13690*, 2025.

[36] Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart Shieber. Investigating gender bias in language models using causal mediation analysis. In *Advances in Neural Information Processing Systems*, 2020.

[37] Xiao Wang, Xiujun Shu, Zhipeng Zhang, Bo Jiang, Yaowei Wang, Yonghong Tian, and Feng Wu. Towards more flexible and accurate object tracking with natural language: Algorithms and benchmark. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 13763–13773, 2021.

[38] Chunhui Zhang, Guanjie Huang, Li Liu, Shan Huang, Yinan Yang, Xiang Wan, Shiming Ge, and Dacheng Tao. WebUAV-3M: A benchmark for unveiling the power of million-scale deep UAV tracking. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 45(7):9186–9205, 2023.

[39] Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods. In *Proceedings of the International Conference on Learning Representations*, 2024.

[40] Yaozong Zheng, Bineng Zhong, Qihua Liang, Guorong Li, Rongrong Ji, and Xianxian Li. Toward unified token learning for vision-language tracking. *IEEE Transactions on Circuits and Systems for Video Technology*, 34(4):2125–2135, 2024.

[41] Li Zhou, Zikun Zhou, Kaige Mao, and Zhenyu He. Joint visual grounding and tracking with natural language specification. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 23151–23160, 2023.

DISTRACT: Degradation-Indexed Study of Textual Reference Across Competing Targets

Supplementary Material

A. Benchmark and Protocol

A.1. Models

Table 1 lists the exact model revisions, checkpoints, source URLs, and stated license status used in this study.

A.2. Datasets

Table 2 records dataset and annotation provenance. The benchmark uses MeViS and three MOT sources. MeViS supplies native expressions and masks. For DanceTrack, SportsMOT, and PersonPath22, source videos and identities come from the original releases, while expressions and point/mask trajectories come from Molmo2-VideoTrackEval. The selected-pair manifest retains source-video, clip, and annotation-row identifiers. Source-media terms and the separate Molmo2-VideoTrackEval annotation license are reported independently: the latter applies to its expressions, masks, and points, not to the underlying videos.

A.3. Inverse-probability-weighted sensitivity check

For the three MOT datasets, Table 3 records an inverse-probability-weighted sensitivity calculation for the 120-pair stress subset and the 404-pair candidate pool. The guaranteed-quota and random-fill draws are uniform random permutations within their defined strata, giving the following marginal inclusion probabilities:

$$P(\text{pos}) = \frac{n_g}{n_{\text{pos}}} + \left(1 - \frac{n_g}{n_{\text{pos}}}\right)\frac{f}{n_{\text{pos}} - n_g + n_{\text{neg}}},$$
$$P(\text{neg}) = \frac{f}{n_{\text{pos}} - n_g + n_{\text{neg}}},$$

where n_g is the guaranteed-quota size and $f = \text{target} - n_g$ is the random-fill size.

A.4. Annotation conversion and common scoring

For segmentation datasets, the initialization box is the tight, axis-aligned bounding box of the mask foreground. Concretely, for mask $\mathcal{M}$ with foreground pixel coordinates (x, y):

$$x_{\min} = \min_{(x,y)\in\mathcal{M}} x, \qquad y_{\min} = \min_{(x,y)\in\mathcal{M}} y,$$
$$x_{\max} = \max_{(x,y)\in\mathcal{M}} x, \qquad y_{\max} = \max_{(x,y)\in\mathcal{M}} y,$$
$$b = \big(x_{\min},\ y_{\min},\ x_{\max} - x_{\min},\ y_{\max} - y_{\min}\big).$$

Tracker boxes are reduced to their centers $\hat{p} = (x + w/2, y + h/2)$ for common scoring; Molmo2 supplies its point directly. Future annotations are not used for initialization or inference.

A.5. Text-condition construction

Each pair has target and distractor expressions e_T, e_D. The neutral input contains no target identity; the unrelated donor e_U is drawn once from a different annotated clip with the same category label and then fixed across models and severities. The donor pool never relaxes this category constraint. For DanceTrack, SportsMOT, and PersonPath22, the category label is the dataset's single coarse class; MeViS uses the annotation category. Table 4 records the concrete input supplied to each model under each condition.

A.6. Native inputs, outputs, and state reset

Table 5 records each native interface and the state cleared between runs.

DUTrack's released dynamic-description path would otherwise overwrite the assigned condition. Table 6 records the condition-preservation audit used to select the primary static-external policy; it is a protocol verification, not an efficacy result.

Policy/check	Text handling	Audit outcome
Static external (primary)	Adapter resets the external provider to the assigned phrase before every tracked frame.	Holds the assigned benchmark condition fixed.
Seeded dynamic (audit)	Assigned phrase seeds the first tracked step; native updater remains enabled.	Description string and active token IDs changed by frame two in 192/192 runs.

Table 6. DUTrack condition-preservation audit. The seeded-dynamic row is an audit mode, not a benchmark condition.

A.7. Outcome definitions and behavioral inference

For every post-initialization frame, a usable prediction is target, distractor, or other according to the common point-in-support rule above; missing, nonfinite, or invalid predictions are absent. Target takes precedence when supports overlap, and absent frames remain in every denominator. For each pair, condition, and displayed severity, frame fractions are aggregated within pair before pairs are aggregated. In addition to B_T, N_T, C_T in the main paper, the endpoint family contains

$$L_T(s) = A_T(\text{unrelated}, s) - A_T(\text{conflicting}, s), \quad (1)$$
$$S_D(s) = A_D(\text{conflicting}, s) - A_D(\text{unrelated}, s). \quad (2)$$

System	Native operation	Pinned artifact used	Source / download URL	License status
Molmo2	8B VLM; clean point plus native tracking prompt → point trajectory	`allenai/Molmo2-8B`, revision `e28fa285`	`https://huggingface.co/allenai/Molmo2-8B` `https://github.com/allenai/molmo2`	Apache-2.0
JointNLT	BERT text encoder and visual–language fusion; clean box → box trajectory	commit `94b338e1`; `JointNLT_ep0300.pth.tar`; `swin_b_ep300_track`	`https://github.com/lizhou-cs/JointNLT` `https://drive.google.com/drive/folders/1ZKBDhgXgVlGLSzFxw5fqpqdZuX1hjfQe?usp=share_link`	MIT
UVLTrack	Language-conditioned box tracker; clean box and raw expression → box trajectory	commit `6ca34055`; `UVLTrack_ep0300.pth.tar` (baseline-base)	`https://github.com/OpenSpaceAI/UVLTrack` `https://drive.google.com/drive/folders/1UZTrGcL3YlxvNpHi0wKsO_sKsTYuYnFo?usp=sharing`	MIT
MMTrack	Language-conditioned multimodal tracker; clean box and raw expression → box trajectory	commit `4586b013`; `MMTrack_ep0150.pth.tar` (baseline)	`https://github.com/Azong-HQU/MMTrack` `https://drive.google.com/file/d/1kM8qxwJHJg4HuZNPPTMFZZTGEdtASVAf/view?usp=sharing`	MIT
DUTrack	Description-conditioned box tracker; static external expression plus clean box → box trajectory	commit `490d9985`; `DUTrack_ep0045.pth.tar` (`dutrack_256_full`)	`https://github.com/GXNU-ZhongLab/DUTrack` `https://drive.google.com/drive/folders/1edieIddvSzy9219F4WShajvY4UmkKJV4?usp=sharing`	No license file in pinned repo

Table 1. Models, pinned artifacts, upstream source URLs, and stated license status.

Dataset	Source media and annotations used	Access URL(s)	License / use terms
MeViS	120 pairs from 120 source videos. Native human-written motion expressions and masks provide both text and geometry. Source media are obtained from the upstream release.	`https://github.com/henghuiding/MeViS`	Upstream data: CC BY-NC-SA 4.0 `https://creativecommons.org/licenses/by-nc-sa/4.0/`; not redistributed here
DanceTrack	60 pairs. DanceTrack provides source video and identities; Molmo2-VideoTrackEval provides human-written expressions and point/mask trajectories.	`https://dancetrack.github.io/` `https://huggingface.co/datasets/allenai/Molmo2-VideoTrackEval`	Source video: non-commercial research; DanceTrack notes that Internet-sourced media are not project property. Molmo2-VideoTrackEval expressions/masks/points: ODC-BY-1.0
SportsMOT	30 pairs. SportsMOT provides source video and identities; Molmo2-VideoTrackEval provides human-written expressions and point/mask trajectories.	`https://github.com/MCG-NJU/SportsMOT` `https://huggingface.co/datasets/allenai/Molmo2-VideoTrackEval`	Source video: CC BY-NC 4.0 `https://creativecommons.org/licenses/by-nc/4.0/` Molmo2-VideoTrackEval expressions/masks/points: ODC-BY-1.0
PersonPath22	30 pairs. PersonPath22 provides source video and identities; Molmo2-VideoTrackEval provides human-written expressions and point/mask trajectories.	`https://github.com/amazon-science/tracking-dataset` `https://huggingface.co/datasets/allenai/Molmo2-VideoTrackEval`	Source release: CC BY-NC 4.0; some externally sourced videos are not redistributed by the provider. Molmo2-VideoTrackEval expressions/masks/points: ODC-BY-1.0

Table 2. Dataset sources, annotation provenance, access URLs, and applicable upstream terms. Molmo2-VideoTrackEval annotation terms are distinct from source-media terms.

Persistent takeover and recovery require three consecutive distractor and subsequent target labels, respectively. For a fixed model, dataset, perturbation, and displayed severity $s \in \{0, \ldots, 5\}$, let $d_{p,s}$ denote a paired effect for pair p; for example, $d_{p,s} = A_T^p(\text{correct}, s) - A_T^p(\text{unrelated}, s)$ for B_T. The reported point estimate is the pair-weighted mean

$$\hat{\theta}_s = \frac{1}{N} \sum_{p=1}^{N} d_{p,s}. \tag{3}$$

Pairs from the same source video are not treated as independent. Let $g(p) \in \mathcal{V}$ assign pair p to its source video, and let $S_v = \sum_{p:g(p)=v} d_{p,s}$ and $n_v = \sum_p \mathbf{1}[g(p) = v]$.

For each of $R = 10{,}000$ bootstrap replicates, we draw $|\mathcal{V}|$ videos with replacement, $Z_1^{(r)}, \ldots, Z_{|\mathcal{V}|}^{(r)}$, and retain every selected pair. The replicate estimate is

$$\hat{\theta}_s^{(r)} = \frac{\sum_{j=1}^{|\mathcal{V}|} S_{Z_j^{(r)}}}{\sum_{j=1}^{|\mathcal{V}|} n_{Z_j^{(r)}}}. \tag{4}$$

Dataset	Effect	Naive (selected)	IPW	Full pool
DanceTrack	B_T	+12.8%	+11.3%	+11.4%
	N_T	+7.6%	+7.3%	+9.4%
	C_T	+25.3%	+21.5%	+18.0%
SportsMOT	B_T	+0.2%	+0.2%	+1.2%
	N_T	+3.8%	+2.2%	+2.6%
	C_T	+5.5%	+1.6%	+4.9%
PersonPath22	B_T	+0.7%	+0.6%	+6.7%
	N_T	+3.6%	+4.0%	+6.7%
	C_T	+14.3%	+9.3%	+12.5%
Pooled (3 datasets)	B_T	+6.6%	+5.8%	+7.7%
	N_T	+5.7%	+5.2%	+7.0%
	C_T	+17.6%	+13.5%	+13.4%

Table 3. Molmo2 clean-input B_T, N_T, and C_T: unweighted stress-subset estimate, inverse-probability-weighted estimate, and candidate-pool mean.

Condition	Molmo2 input after its fixed scaffold	Specialized trackers
Neutral	no appended identity description	`object`
Unrelated	"The target is: e_U"	e_U
Correct	"The target is: e_T"	e_T
Conflicting	"The target is: e_D"	e_D

Table 4. Text input under each condition for Molmo2 and the specialized trackers.

The 95% CI is the empirical $[2.5, 97.5]$ percentile interval.

For a two-sided paired sign-flip test, we independently draw $\epsilon_v^{(r)} \in \{-1, +1\}$ for every source video and form

$$\hat{\theta}_{s,0}^{(r)} = \frac{1}{N} \sum_{v \in \mathcal{V}} \epsilon_v^{(r)} S_v, \qquad p = \frac{1}{R} \sum_{r=1}^{R} \mathbf{1}\left[|\hat{\theta}_{s,0}^{(r)}| \geq |\hat{\theta}_s| \right]. \tag{5}$$

For a severity-endpoint test, we apply the same procedure to $d_{p,S_5} - d_{p,S_0}$.

For every model–dataset experiment, the endpoint family contains the five effects B_T, N_T, C_T, L_T, S_D, each tested for nonzero effect at S_5 and for the $S_0 \to S_5$ change within each of the four perturbations (40 tests), plus one nonzero-effect-at-S_0 test per effect (5 tests): the S_0 record is the shared clean baseline and is identical across every perturbation's ladder, so it is tested once per effect rather than once per perturbation, for 45 tests total. We apply Benjamini–Hochberg to all m p-values in this family. For ordered p-values $p_{(1)} \leq \cdots \leq p_{(m)}$, the reported adjusted value is

$$q_{(i)} = \min\left\{ 1, \min_{j \geq i} \frac{m}{j} p_{(j)} \right\}. \tag{6}$$

IAD is excluded from this testing family because it is descriptive.

A.8. Perturbation realization and displayed severity

Each corruption is applied independently to every post-initialization frame (the initialization frame is always clean); corruptions are never composed. For a frame $I \in [0, 255]^{H \times W \times 3}$, the four perturbations are implemented as follows.

Blur. Per-channel Gaussian smoothing with $\sigma_s = k_s \cdot \min(H, W)$, kernel truncated at $\pm 3\sigma_s$: $I'_{\cdot,\cdot,c} = I_{\cdot,\cdot,c} * \mathcal{N}(0, \sigma_s^2)$ for each channel c. $k_0 = 0$ is an exact identity.

Noise. The frame is normalized to $[0, 1]$; a single standard-normal field $\eta \sim \mathcal{N}(0, 1)^{H \times W \times 3}$, seeded deterministically from a hash of (video, frame index) only – never condition, pair, or severity – is drawn once per frame and reused unchanged across all four text conditions and every severity for that frame: $I' = \mathrm{round}\big(255 \cdot \mathrm{clip}\big(\frac{I}{255} + \sigma_s \eta,\, 0,\, 1\big)\big)$. $\sigma_0 = 0$ is an exact identity; because only the scale σ_s changes for a fixed draw of η, severities differ in dose, not in noise realization, so contrasts across conditions at one severity are never confounded by different noise instances.

Pixelation. Block-mosaic resampling: downscale to $(\mathrm{round}(r_s W), \mathrm{round}(r_s H))$ and upscale back to (W, H), both steps using PIL's BOX (area-average) filter – the ImageNet-C `pixelate` corruption. $r_0 = 1$ is an exact identity.

Compression. The frame is re-encoded to JPEG at quality q_s (libjpeg's 1–100 scale) and immediately decoded back to RGB – the ImageNet-C `jpeg_compression` corruption. $q_0 = 100$ is treated as an exact identity (the untouched frame is passed through, not a quality-100 re-encode).

Table 7 gives the raw-index subsequence used for the displayed S_0–S_5 schedule. S_0 is raw index 0 for every perturbation. Figure 1 shows one frame under the displayed schedule.

Figure 1. Displayed S_0–S_5 perturbation schedule on one MeViS frame.

A.9. Objective distortion audit

Table 8 reports clean-to-S_5 PSNR and SSIM over a deterministic 10-video, one-frame-per-video sample.

System	Initialization	Text interface	Output	Reset safeguard
Molmo2	clean target point	native point-tracking scaffold	point trajectory	fresh generation state
JointNLT	clean tight box	raw expression, native TRACK mode	box trajectory	ROI queue, template features, tracker state
UVLTrack	clean tight box	raw expression	box trajectory	state, frame counter, score and output cache
MMTrack	clean tight box	raw expression	box trajectory	adapter/tracker and decode state
DUTrack	clean tight box	static external expression	box trajectory	dynamic description and tracker state

Table 5. Native inputs, outputs, and reset state.

Perturbation	raw n	S_0	S_1	S_2	S_3	S_4	S_5
Blur	8	0	2	3	4	5	6
Noise	9	0	4	5	6	7	8
Pixelation	8	0	1	2	3	4	7
Compression	7	0	2	3	4	5	6

Table 7. Raw implementation indices used for each displayed S_0–S_5 severity.

Perturbation	PSNR mean (dB)	PSNR range	SSIM mean	SSIM range
Blur	19.5	[17.2, 24.0]	0.72	[0.40, 0.85]
Noise	6.1	[5.6, 6.4]	0.02	[0.01, 0.03]
Pixelation	22.1	[19.2, 26.1]	0.70	[0.40, 0.82]
Compression	23.7	[20.8, 26.1]	0.73	[0.55, 0.84]

Table 8. Clean-to-displayed-S_5 PSNR and SSIM on 10 fixed videos (one frame per video); ranges are min–max.

B. Results

B.1. Behavioral degradation study

This section provides the complete behavioral-result grid across five models, four datasets, and four perturbations.

For standalone readability, with $\Delta_{c_1,c_2}(s)$ denoting the target-accuracy difference between conditions at severity s, we use

$$\mathrm{IAD}(c_1, c_2) = \sum_{s=0}^{4} \tfrac{1}{2} \left(|\Delta_{c_1,c_2}(s)| + |\Delta_{c_1,c_2}(s{+}1)| \right). \quad (7)$$

B.2. IAD matrices by model

Each figure gives the upper triangle of the neutral/unrelated/correct/ conflicting IAD matrix for one model and perturbation, with datasets as panels. Each cell is the accumulated absolute target-accuracy difference between its row and column conditions. For each system, the blur, noise, pixelation, and compression matrices are given in Figures 2–5 (Molmo2); Figures 6–9 (UVLTrack); Figures 10–13 (MMTrack); Figures 14–17 (JointNLT); and Figures 18–21 (DUTrack).

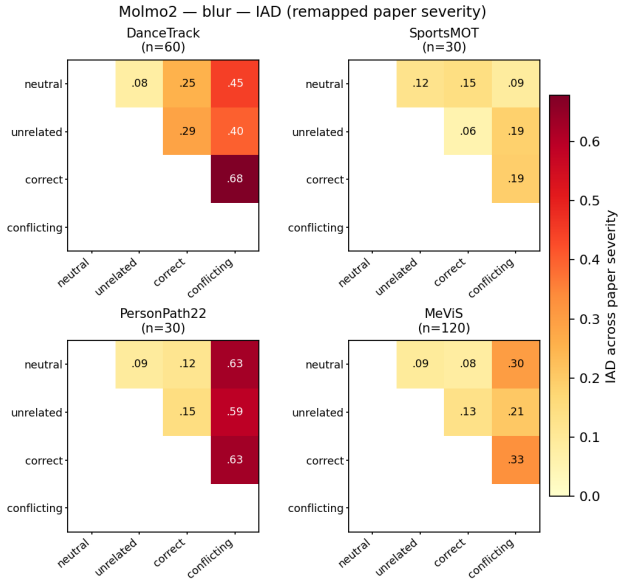

Figure 2. Molmo2, blur: IAD matrix (Eq. 7) across all four datasets.

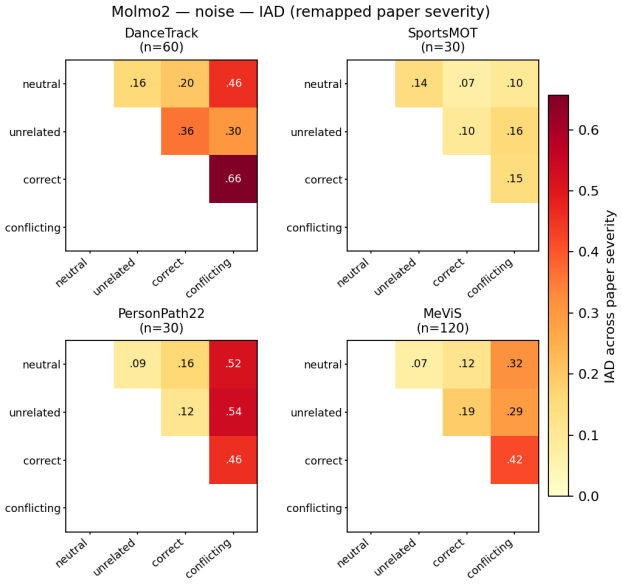

Figure 3. Molmo2, noise: IAD matrix (Eq. 7) across all four datasets.

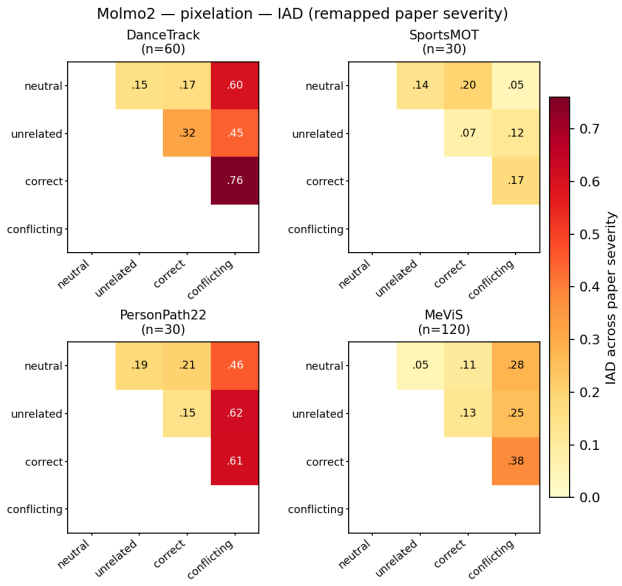

Figure 4. Molmo2, pixelation: IAD matrix (Eq. 7) across all four datasets.

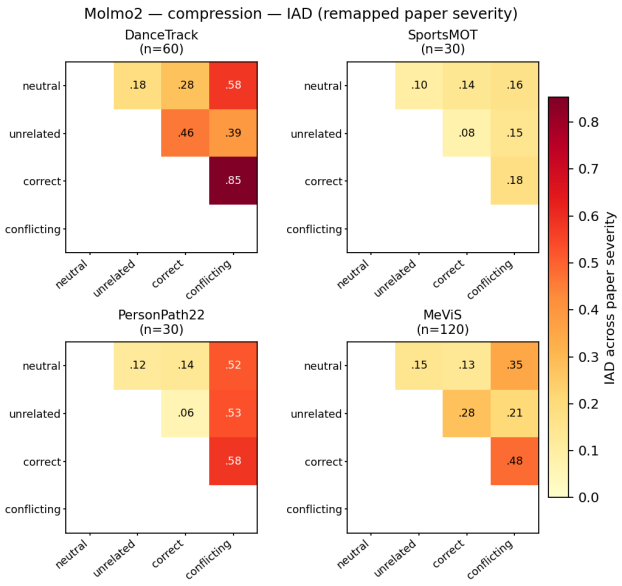

Figure 5. Molmo2, compression: IAD matrix (Eq. 7) across all four datasets.

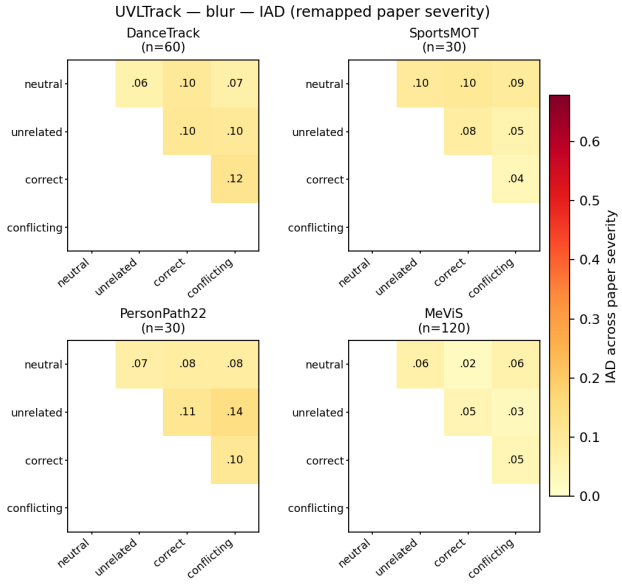

Figure 6. UVLTrack, blur: IAD matrix (Eq. 7) across all four datasets.

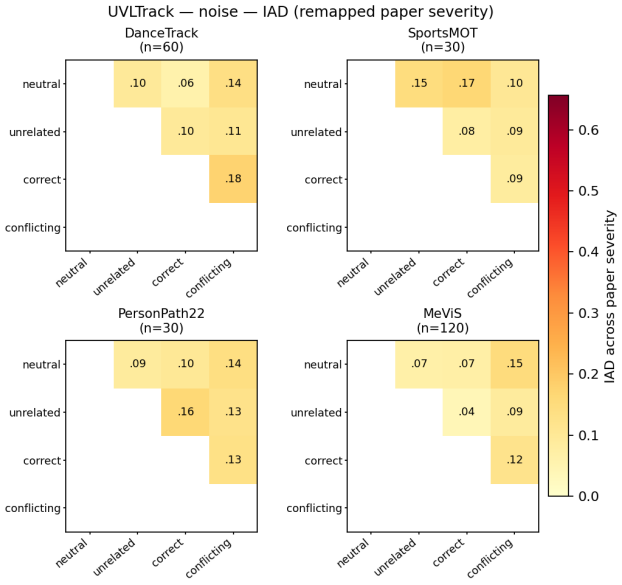

Figure 7. UVLTrack, noise: IAD matrix (Eq. 7) across all four datasets.

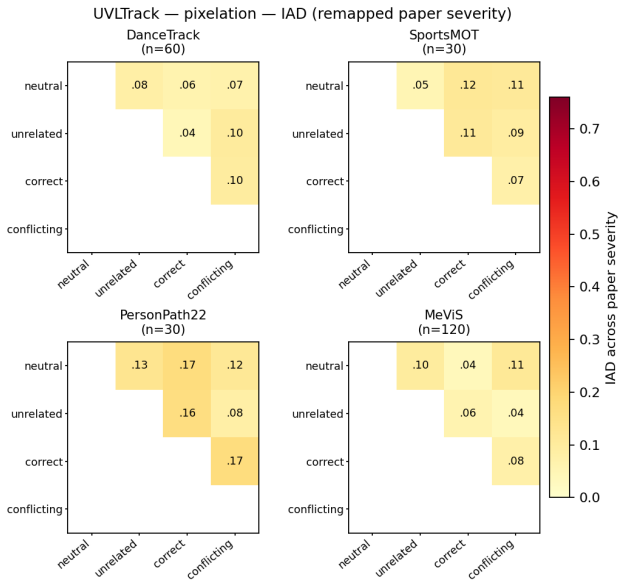

Figure 8. UVLTrack, pixelation: IAD matrix (Eq. 7) across all four datasets.

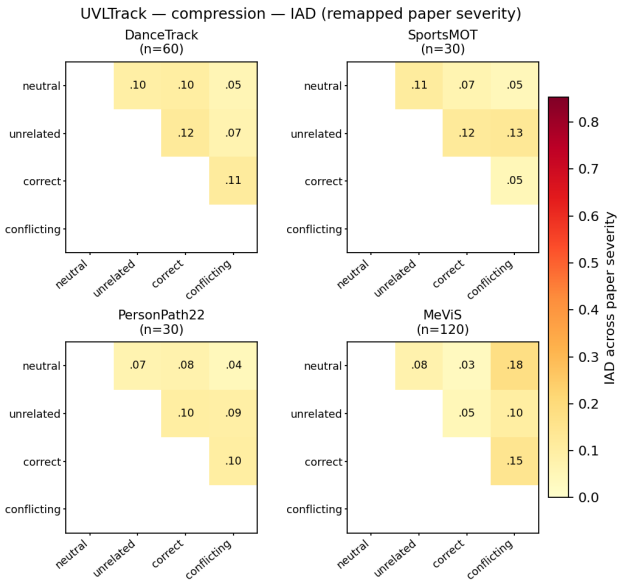

Figure 9. UVLTrack, compression: IAD matrix (Eq. 7) across all four datasets.

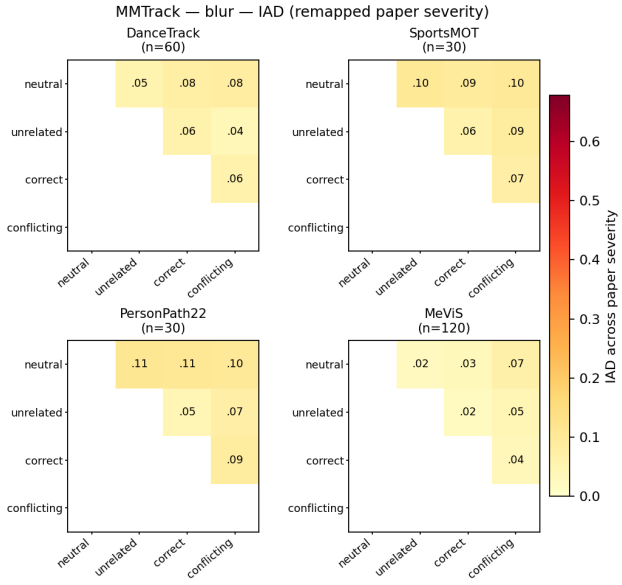

Figure 10. MMTrack, blur: IAD matrix (Eq. 7) across all four datasets.

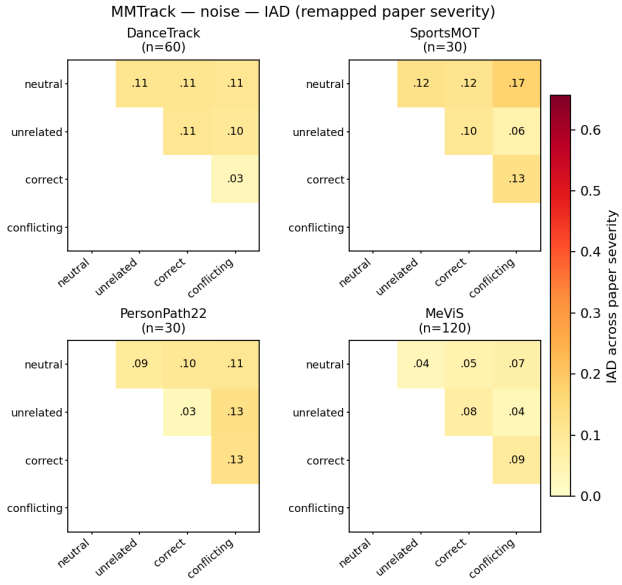

Figure 11. MMTrack, noise: IAD matrix (Eq. 7) across all four datasets.

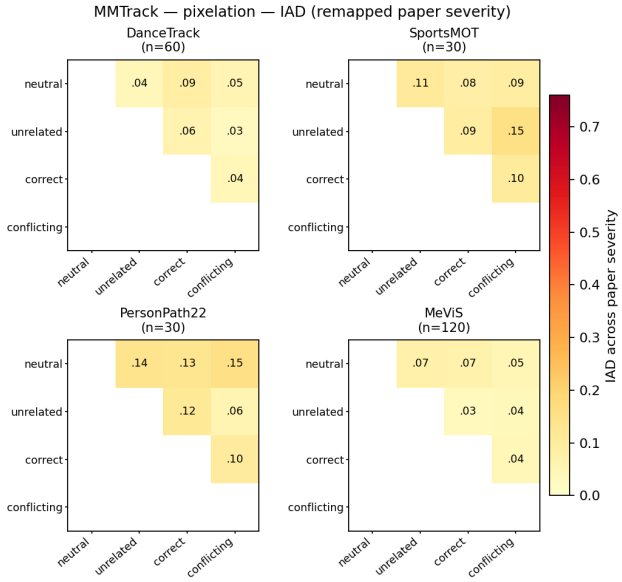

Figure 12. MMTrack, pixelation: IAD matrix (Eq. 7) across all four datasets.

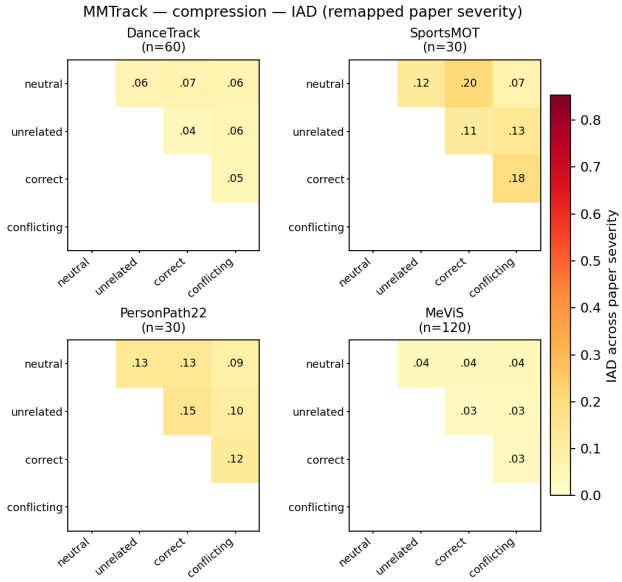

Figure 13. MMTrack, compression: IAD matrix (Eq. 7) across all four datasets.

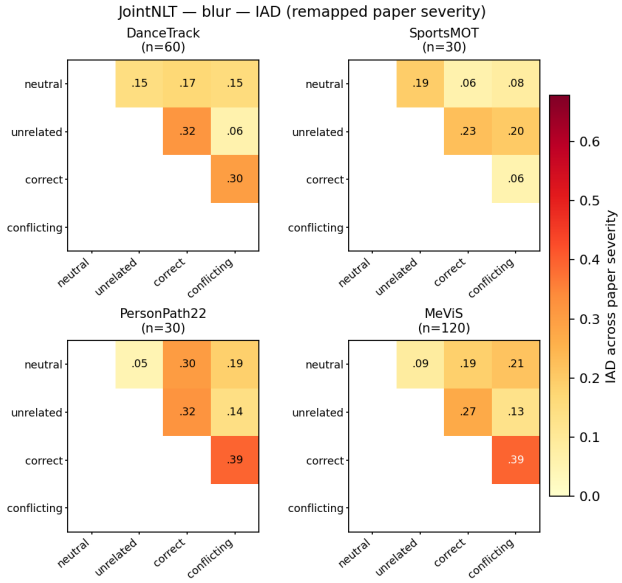

Figure 14. JointNLT, blur: IAD matrix (Eq. 7) across all four datasets.

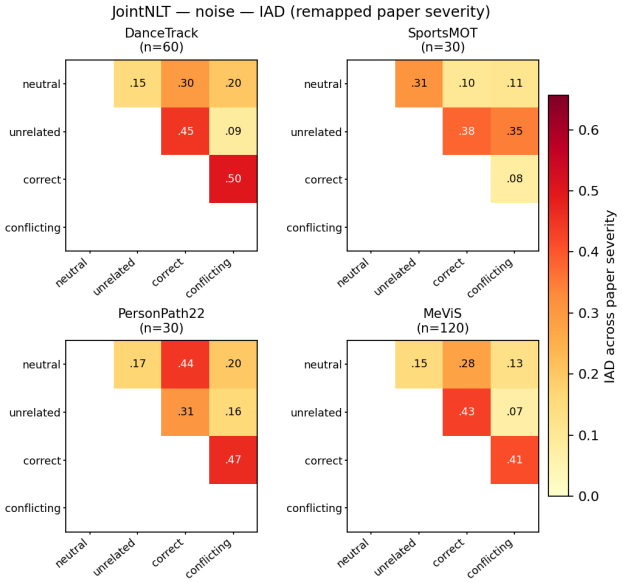

Figure 15. JointNLT, noise: IAD matrix (Eq. 7) across all four datasets.

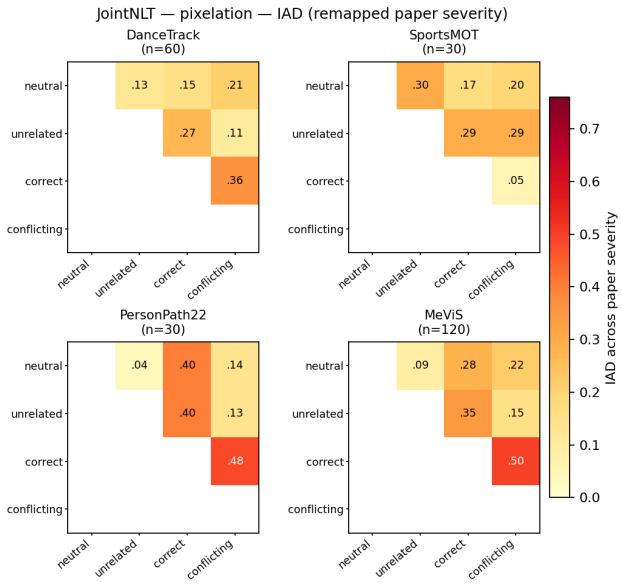

Figure 16. JointNLT, pixelation: IAD matrix (Eq. 7) across all four datasets.

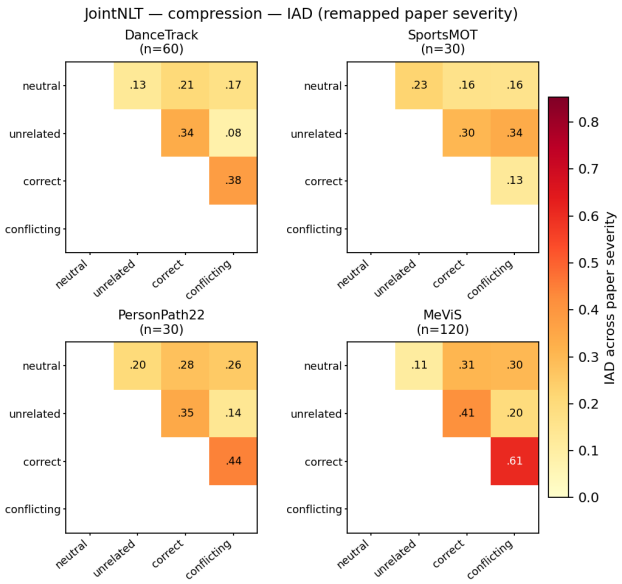

Figure 17. JointNLT, compression: IAD matrix (Eq. 7) across all four datasets.

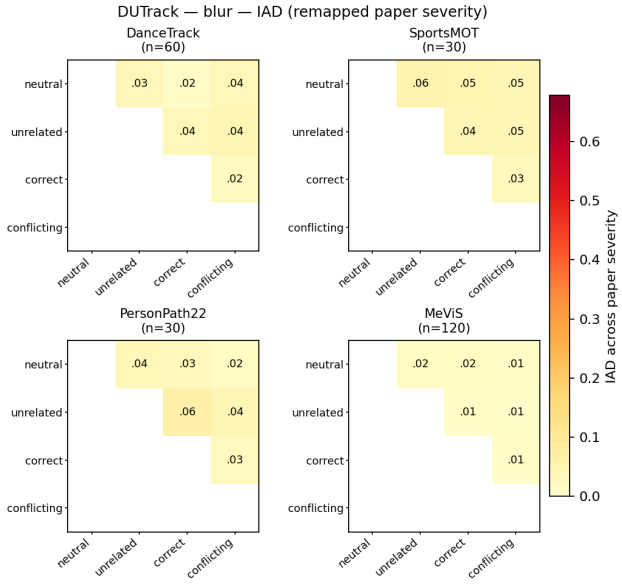

Figure 18. DUTrack, blur: IAD matrix (Eq. 7) across all four datasets.

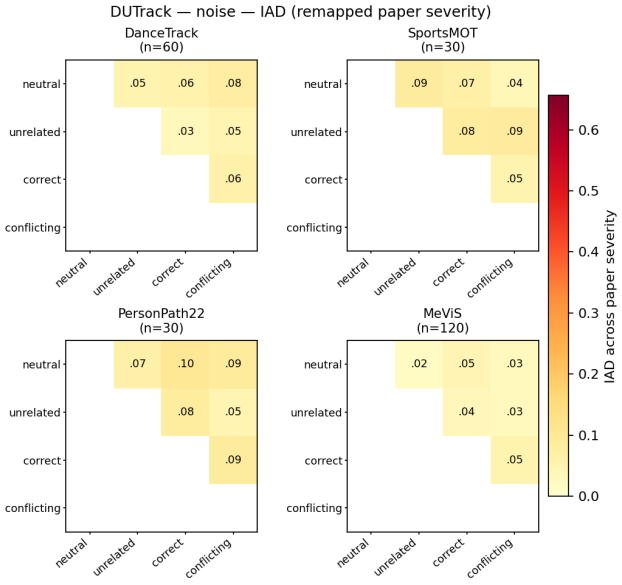

Figure 19. DUTrack, noise: IAD matrix (Eq. 7) across all four datasets.

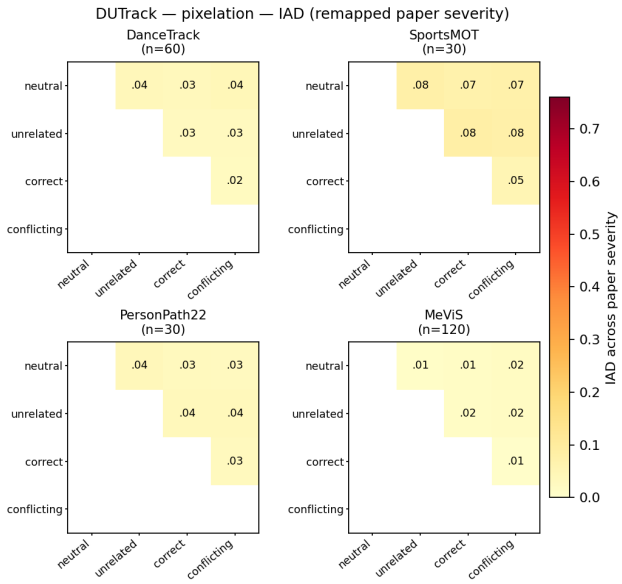

Figure 20. DUTrack, pixelation: IAD matrix (Eq. 7) across all four datasets.

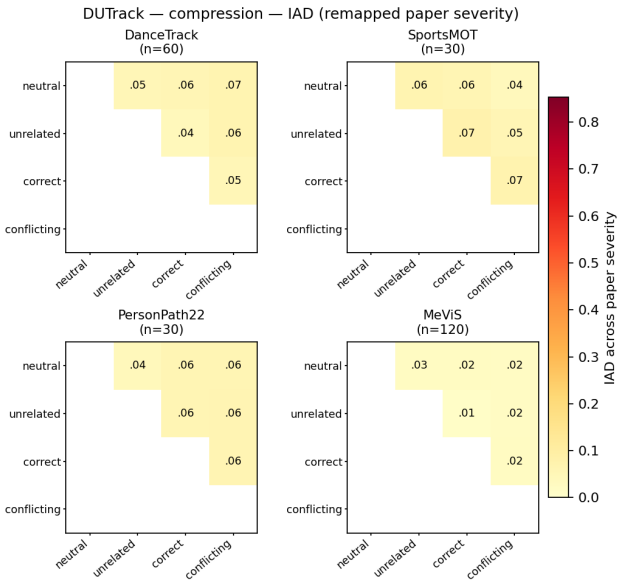

Figure 21. DUTrack, compression: IAD matrix (Eq. 7) across all four datasets.

B.3. Signed severity trajectories

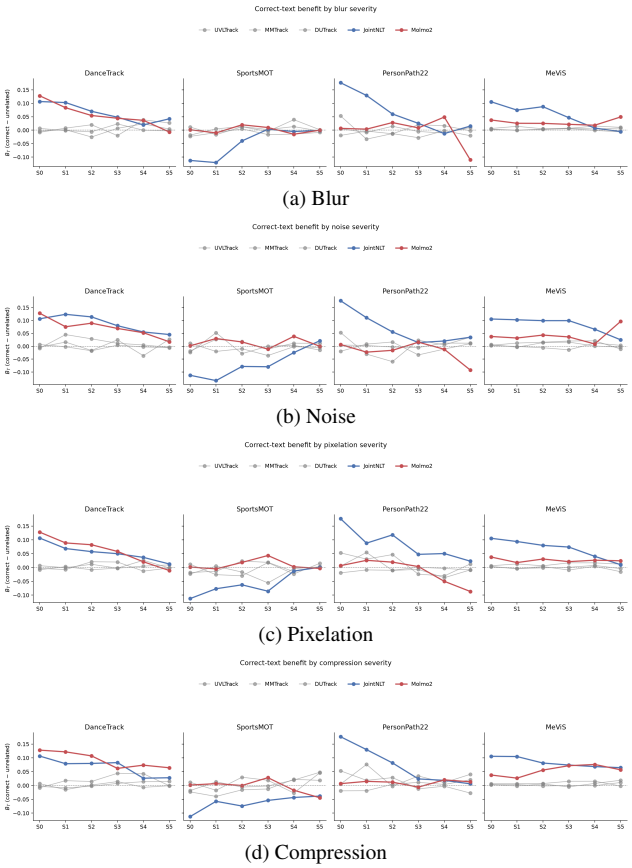

(a) Blur

(b) Noise

(c) Pixelation

(d) Compression

Figure 23. Correct-versus-unrelated target-accuracy effect B_T across severity, systems, datasets, and perturbations.

Figure 23 provides the signed B_T trajectories that complement the unsigned IAD matrices.

B.4. Four-condition trajectories, all systems

The following figures provide condition-specific target-accuracy and distractor-capture trajectories for every model–dataset pair. Figure 24 gives the pooled distractor-capture summary before the per-model–dataset trajectories.

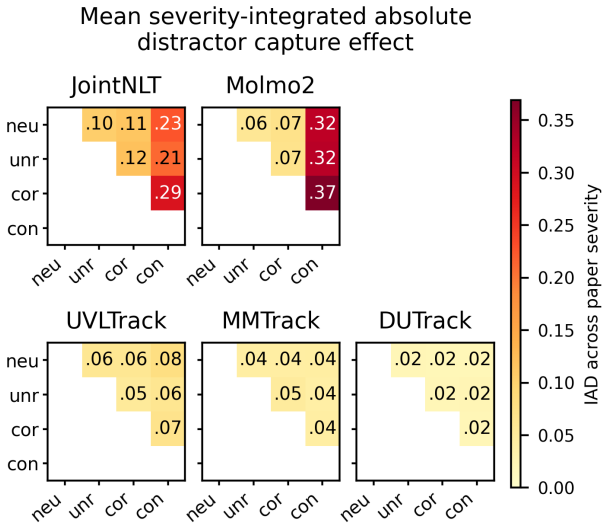

Figure 24. Mean severity-integrated distractor-capture differences, pooled across datasets and perturbations.

Figures 25–28 give Molmo2 target-accuracy trajectories; Figures 29–32 give distractor-capture trajectories.
Figures 33–36 give UVLTrack target-accuracy trajectories; Figures 37–40 give distractor-capture trajectories.
Figures 41–44 give MMTrack target-accuracy trajectories; Figures 45–48 give distractor-capture trajectories.
Figures 49–52 give JointNLT target-accuracy trajectories; Figures 53–56 give distractor-capture trajectories.
Figures 57–60 give DUTrack target-accuracy trajectories; Figures 61–64 give distractor-capture trajectories.

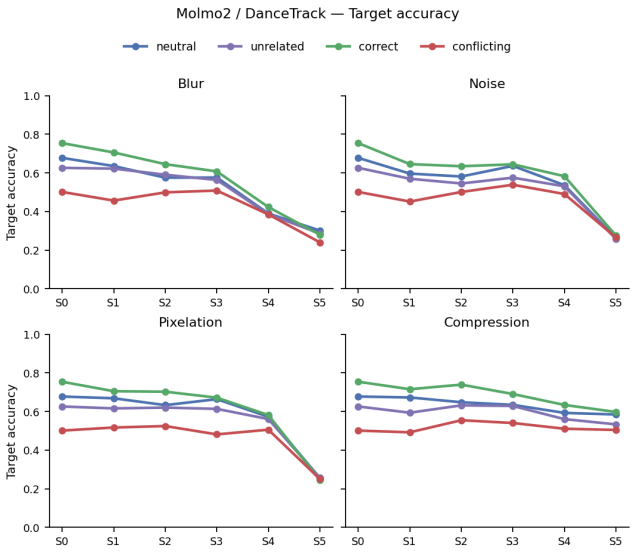

Figure 25. Molmo2 on DanceTrack: target accuracy across all four perturbations, for all four text conditions.

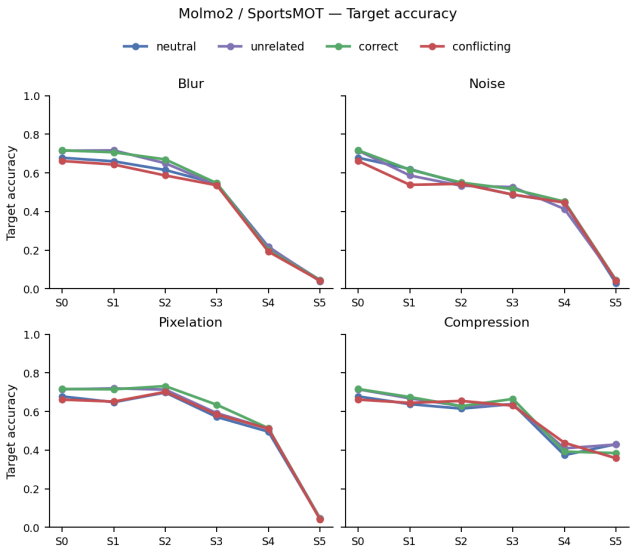

Figure 26. Molmo2 on SportsMOT: target accuracy across all four perturbations, for all four text conditions.

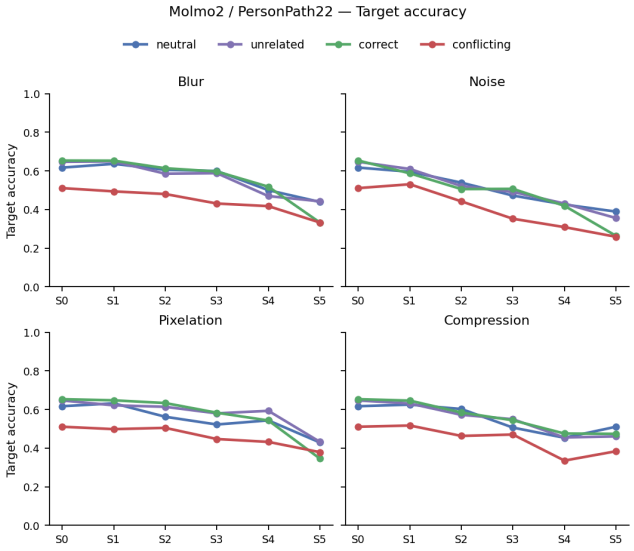

Figure 27. Molmo2 on PersonPath22: target accuracy across all four perturbations, for all four text conditions.

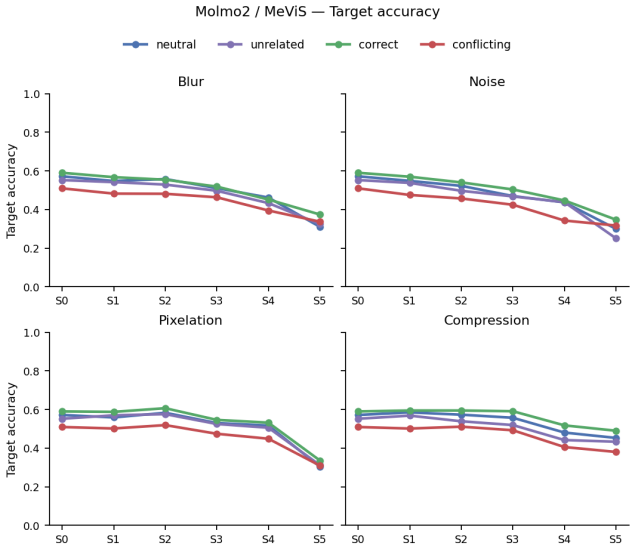

Figure 28. Molmo2 on MeViS: target accuracy across all four perturbations, for all four text conditions.

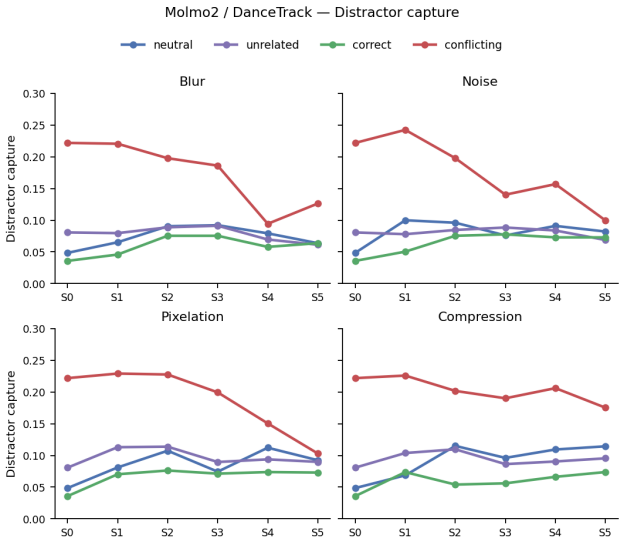

Figure 29. Molmo2 on DanceTrack: distractor capture across all four perturbations, for all four text conditions.

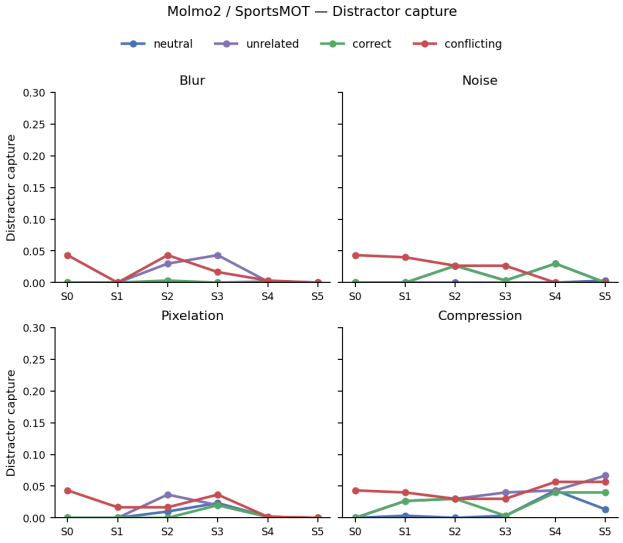

Figure 30. Molmo2 on SportsMOT: distractor capture across all four perturbations, for all four text conditions.

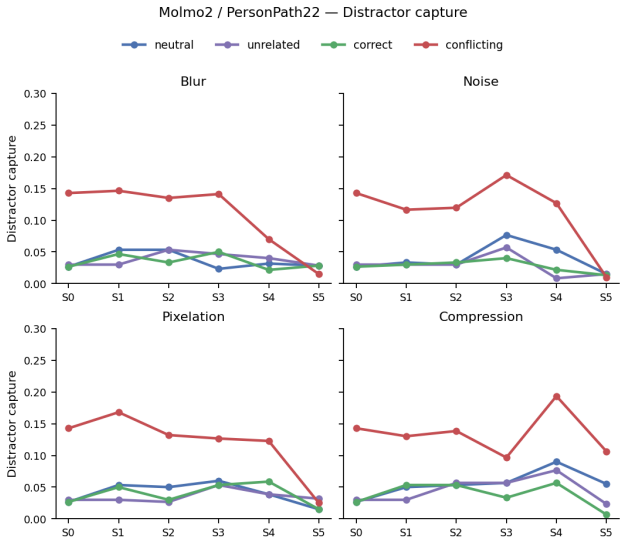

Figure 31. Molmo2 on PersonPath22: distractor capture across all four perturbations, for all four text conditions.

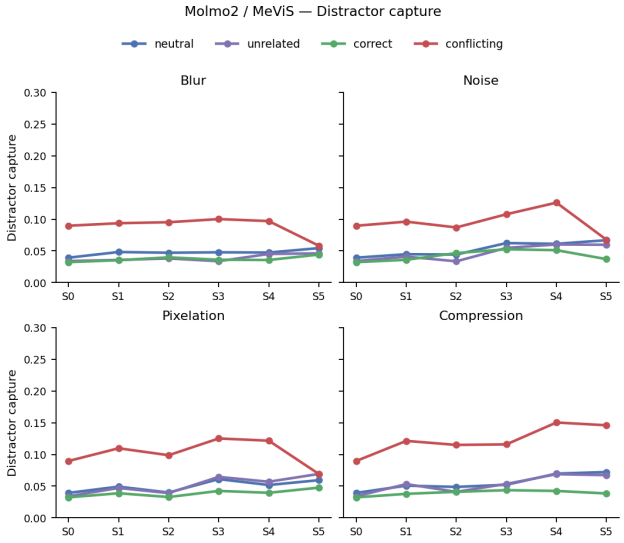

Figure 32. Molmo2 on MeViS: distractor capture across all four perturbations, for all four text conditions.

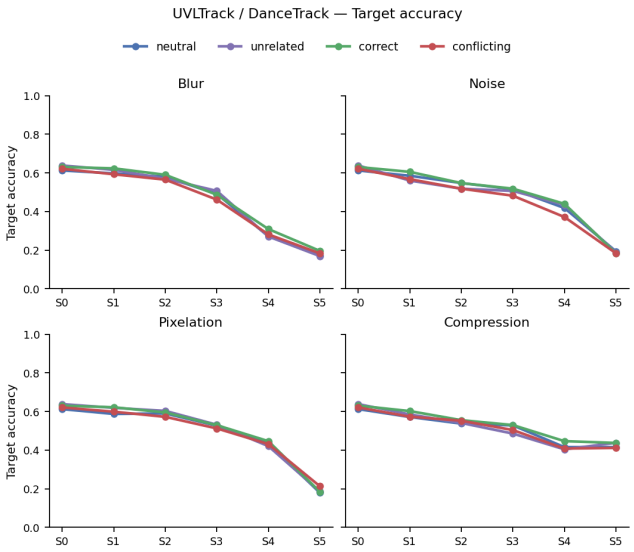

Figure 33. UVLTrack on DanceTrack: target accuracy across all four perturbations, for all four text conditions.

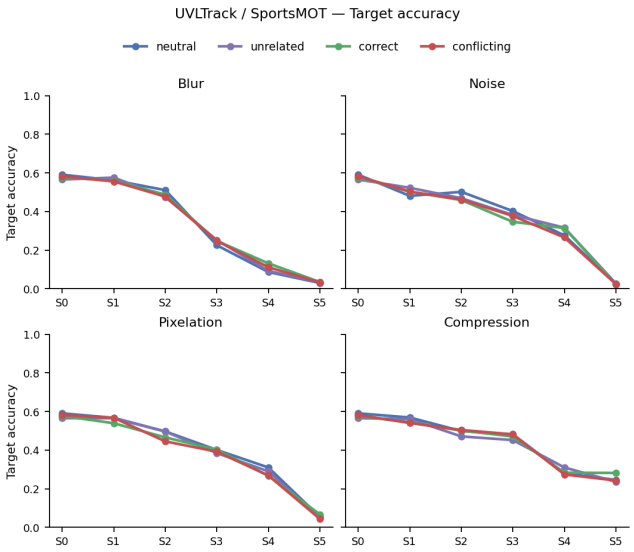

Figure 34. UVLTrack on SportsMOT: target accuracy across all four perturbations, for all four text conditions.

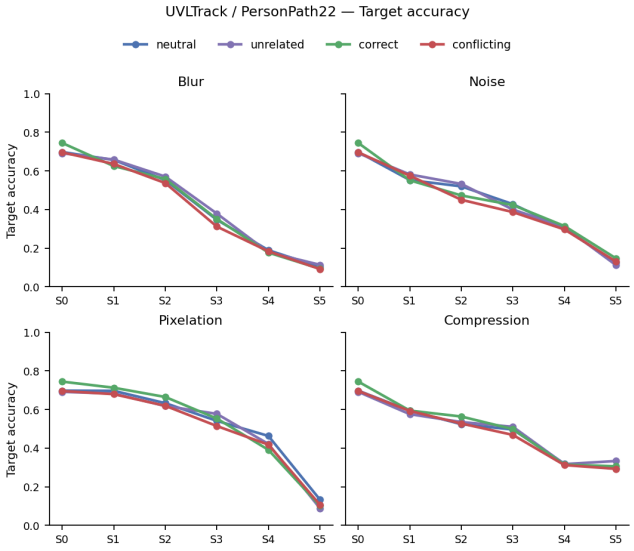

Figure 35. UVLTrack on PersonPath22: target accuracy across all four perturbations, for all four text conditions.

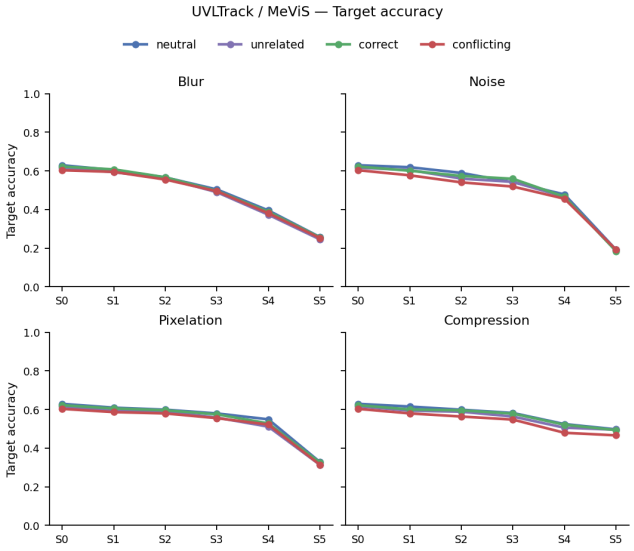

Figure 36. UVLTrack on MeViS: target accuracy across all four perturbations, for all four text conditions.

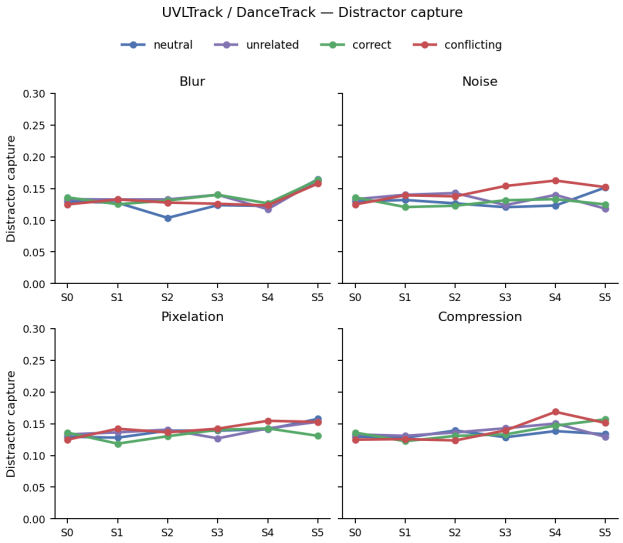

Figure 37. UVLTrack on DanceTrack: distractor capture across all four perturbations, for all four text conditions.

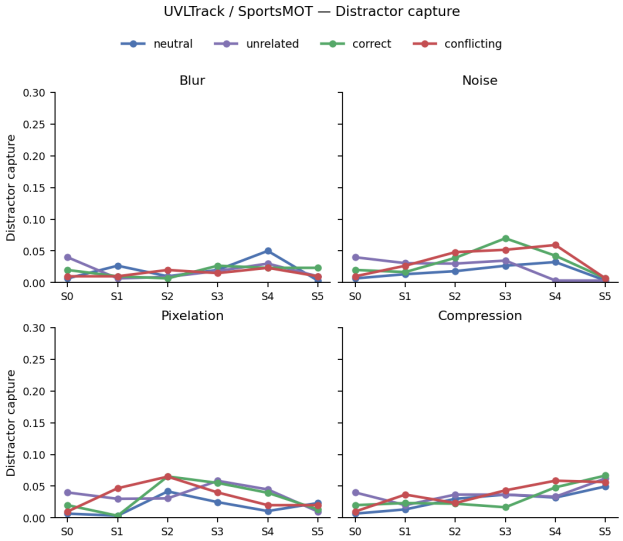

Figure 38. UVLTrack on SportsMOT: distractor capture across all four perturbations, for all four text conditions.

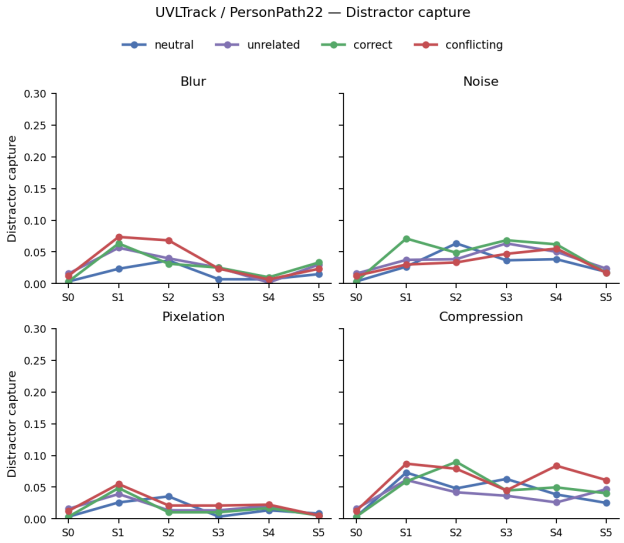

Figure 39. UVLTrack on PersonPath22: distractor capture across all four perturbations, for all four text conditions.

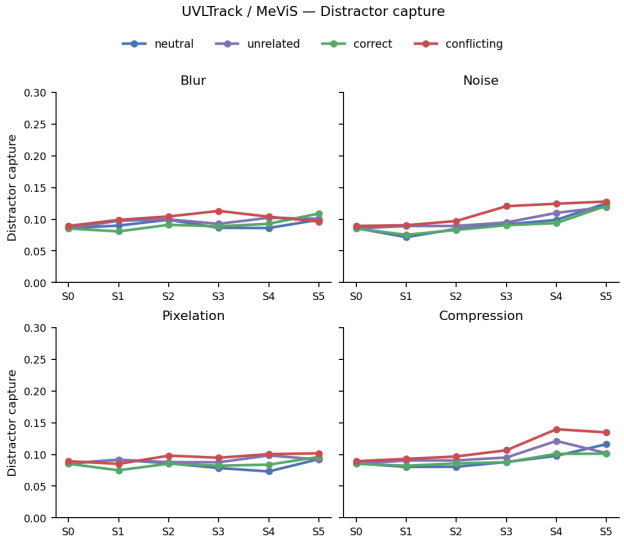

Figure 40. UVLTrack on MeViS: distractor capture across all four perturbations, for all four text conditions.

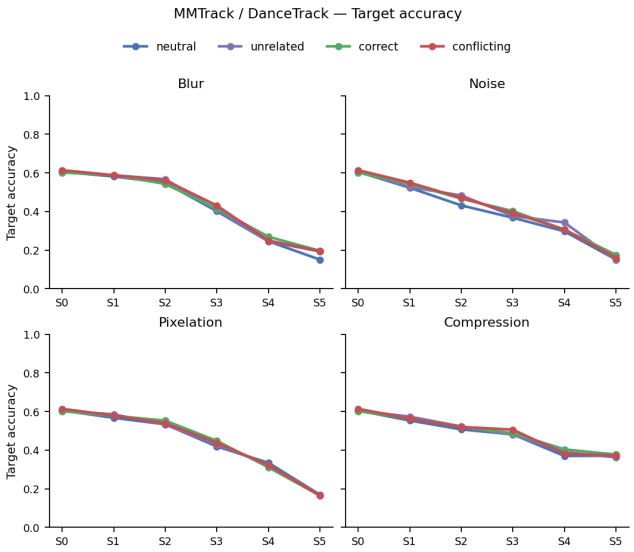

Figure 41. MMTrack on DanceTrack: target accuracy across all four perturbations, for all four text conditions.

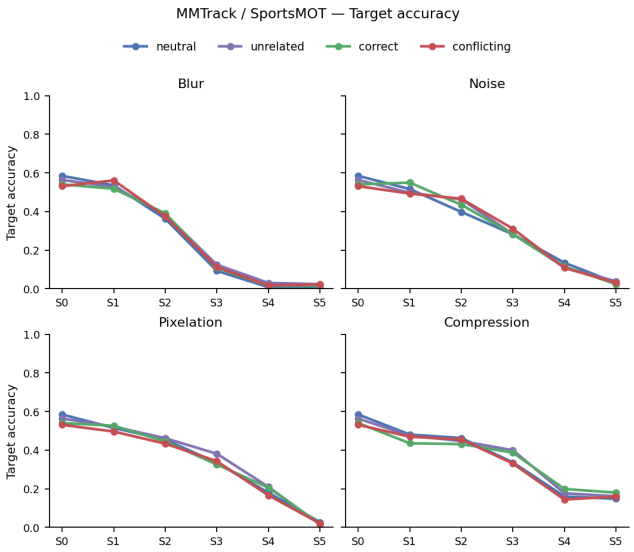

Figure 42. MMTrack on SportsMOT: target accuracy across all four perturbations, for all four text conditions.

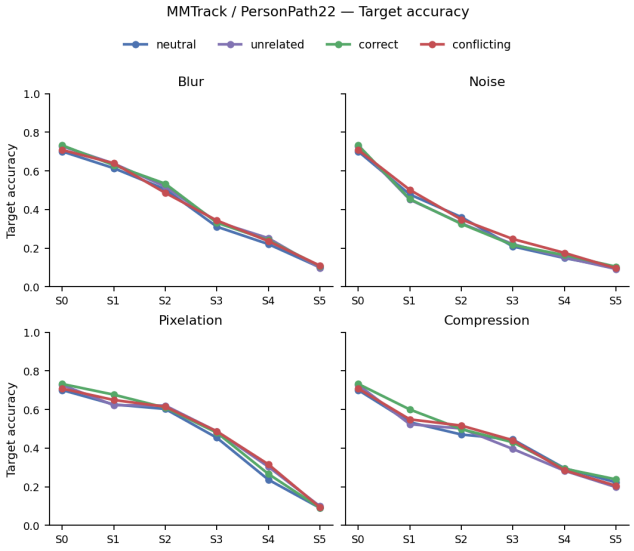

Figure 43. MMTrack on PersonPath22: target accuracy across all four perturbations, for all four text conditions.

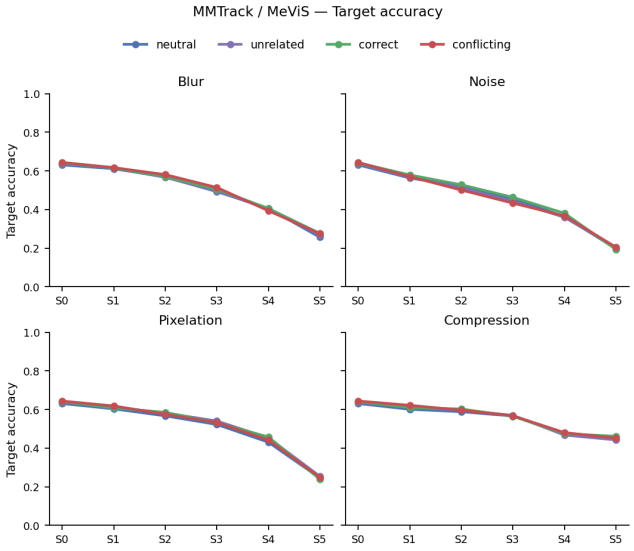

Figure 44. MMTrack on MeViS: target accuracy across all four perturbations, for all four text conditions.

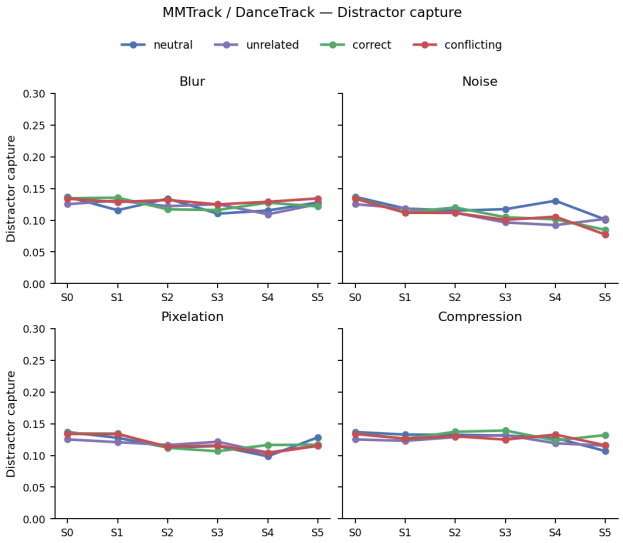

Figure 45. MMTrack on DanceTrack: distractor capture across all four perturbations, for all four text conditions.

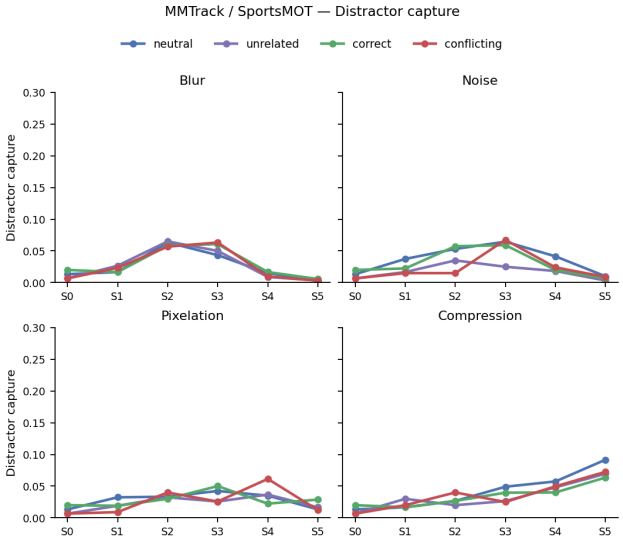

Figure 46. MMTrack on SportsMOT: distractor capture across all four perturbations, for all four text conditions.

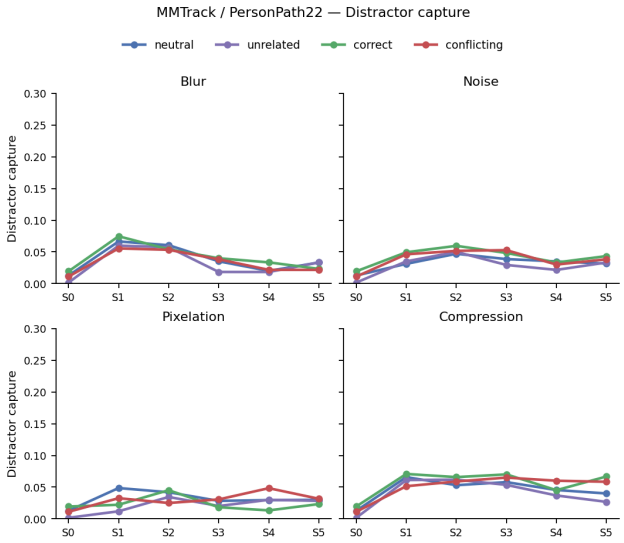

Figure 47. MMTrack on PersonPath22: distractor capture across all four perturbations, for all four text conditions.

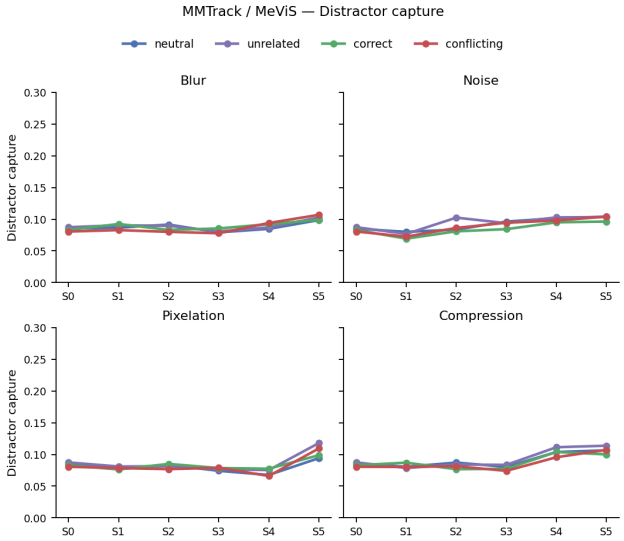

Figure 48. MMTrack on MeViS: distractor capture across all four perturbations, for all four text conditions.

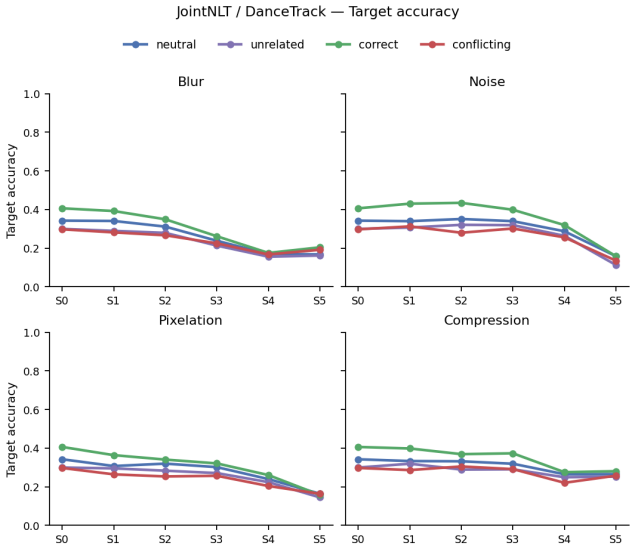

Figure 49. JointNLT on DanceTrack: target accuracy across all four perturbations, for all four text conditions.

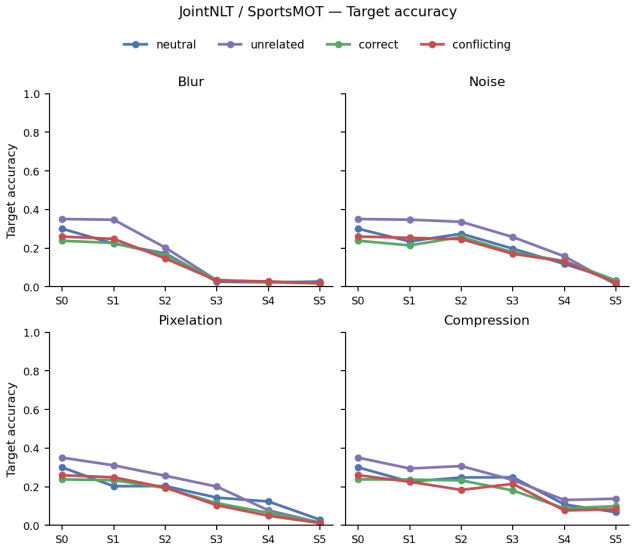

Figure 50. JointNLT on SportsMOT: target accuracy across all four perturbations, for all four text conditions.

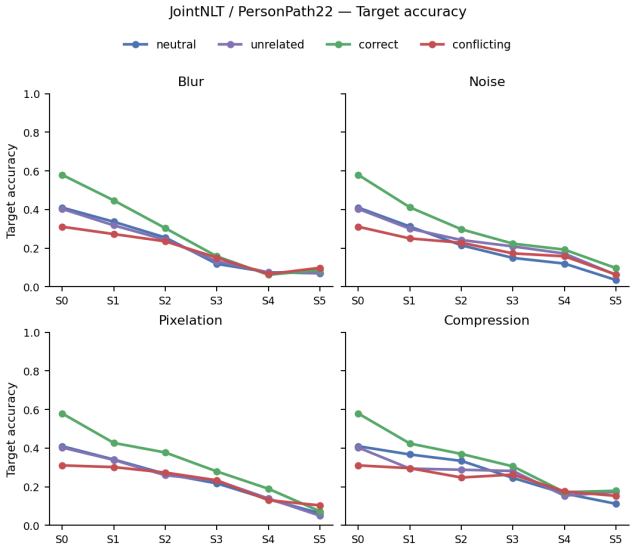

Figure 51. JointNLT on PersonPath22: target accuracy across all four perturbations, for all four text conditions.

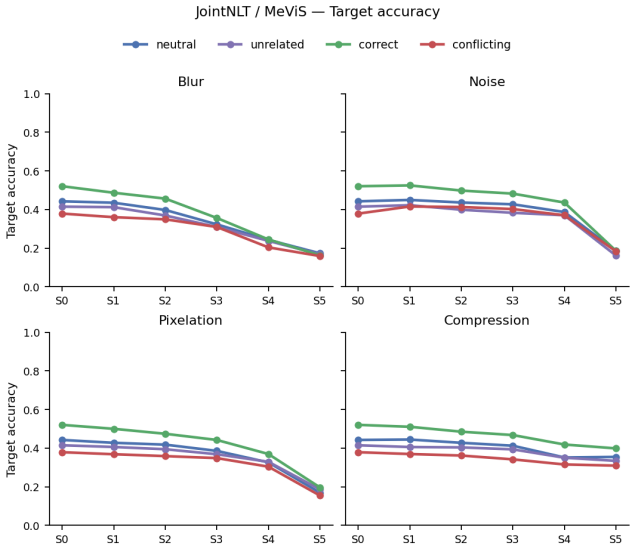

Figure 52. JointNLT on MeViS: target accuracy across all four perturbations, for all four text conditions.

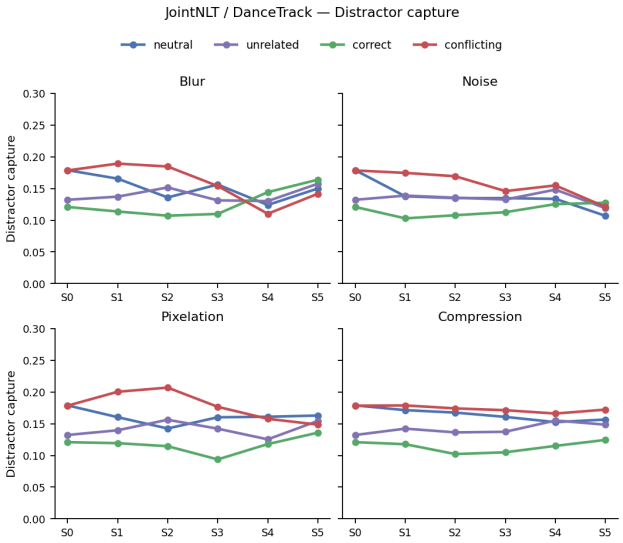

Figure 53. JointNLT on DanceTrack: distractor capture across all four perturbations, for all four text conditions.

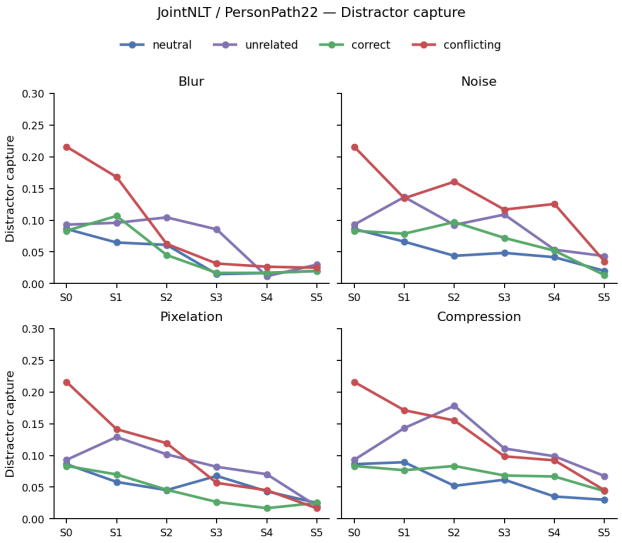

Figure 55. JointNLT on PersonPath22: distractor capture across all four perturbations, for all four text conditions.

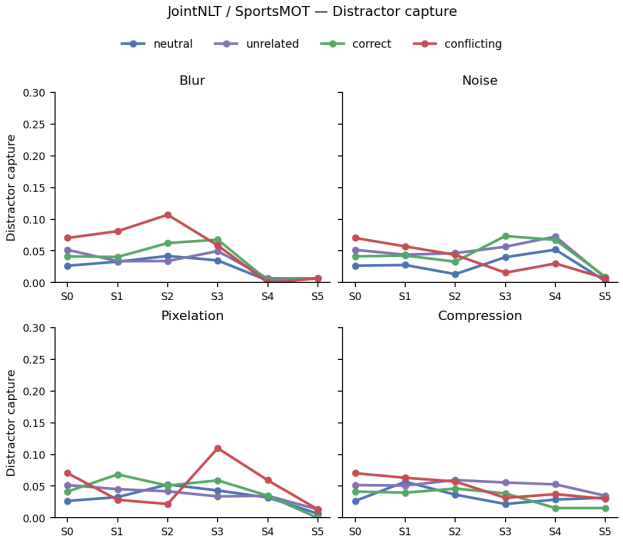

Figure 54. JointNLT on SportsMOT: distractor capture across all four perturbations, for all four text conditions.

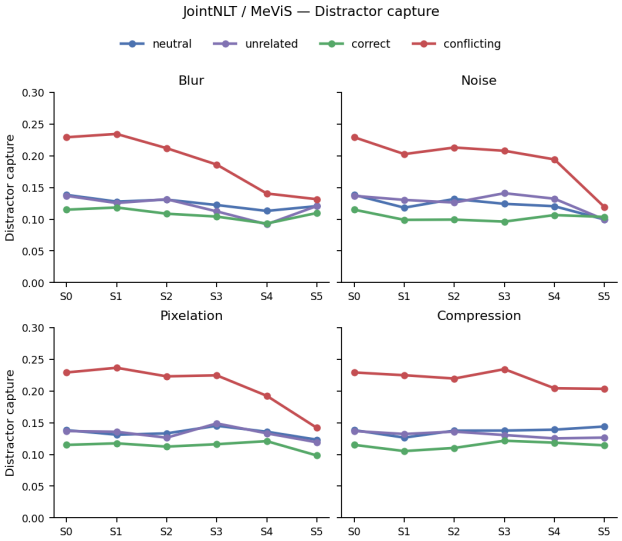

Figure 56. JointNLT on MeViS: distractor capture across all four perturbations, for all four text conditions.

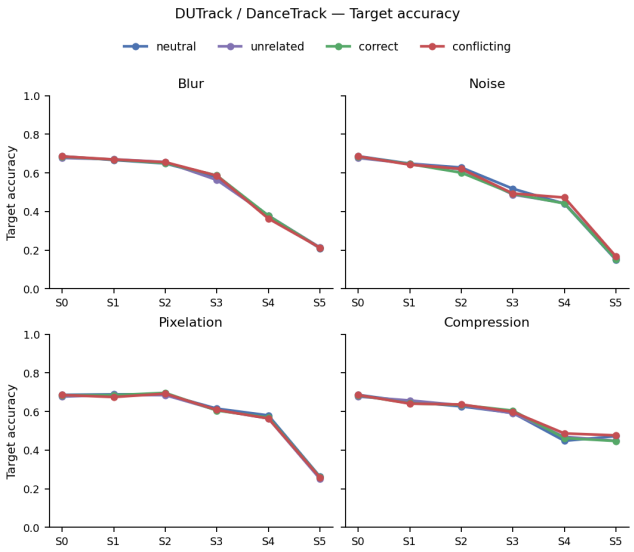

Figure 57. DUTrack on DanceTrack: target accuracy across all four perturbations, for all four text conditions.

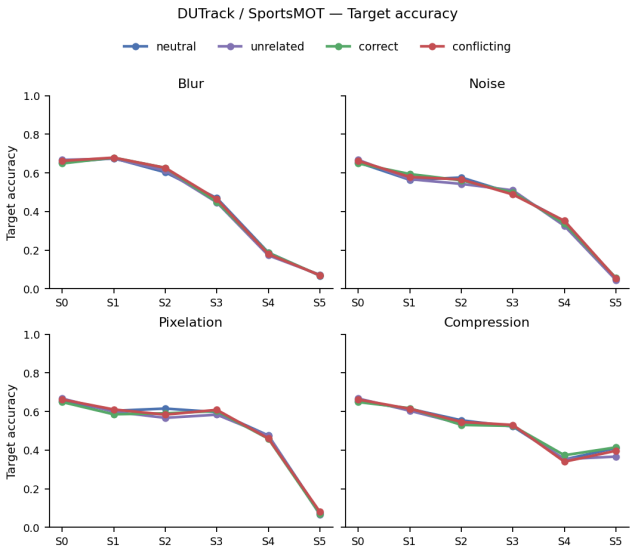

Figure 58. DUTrack on SportsMOT: target accuracy across all four perturbations, for all four text conditions.

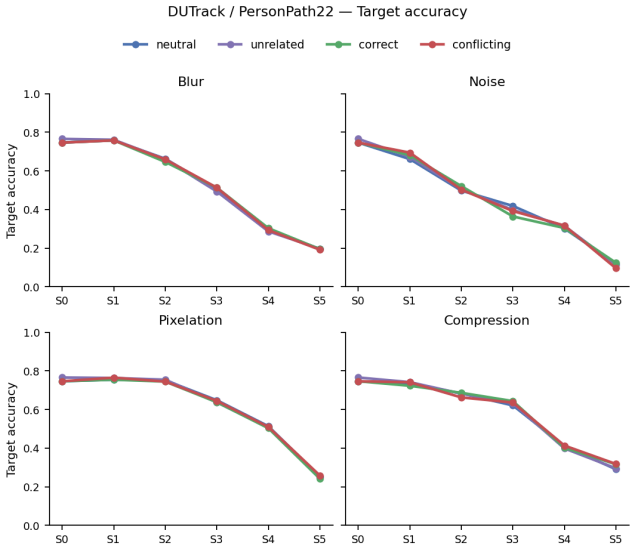

Figure 59. DUTrack on PersonPath22: target accuracy across all four perturbations, for all four text conditions.

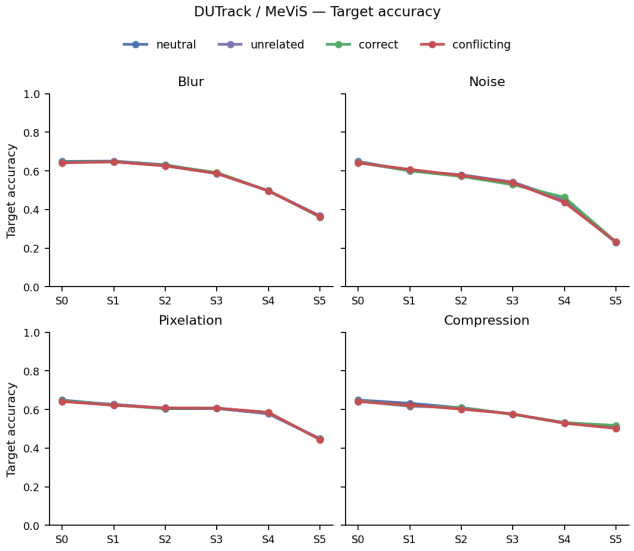

Figure 60. DUTrack on MeViS: target accuracy across all four perturbations, for all four text conditions.

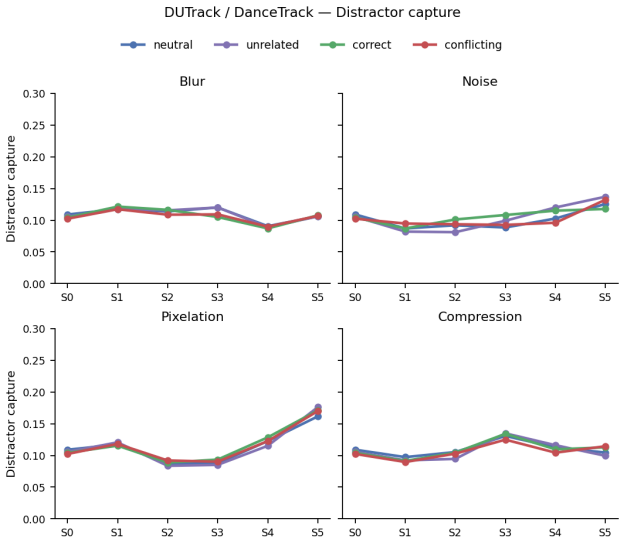

Figure 61. DUTrack on DanceTrack: distractor capture across all four perturbations, for all four text conditions.

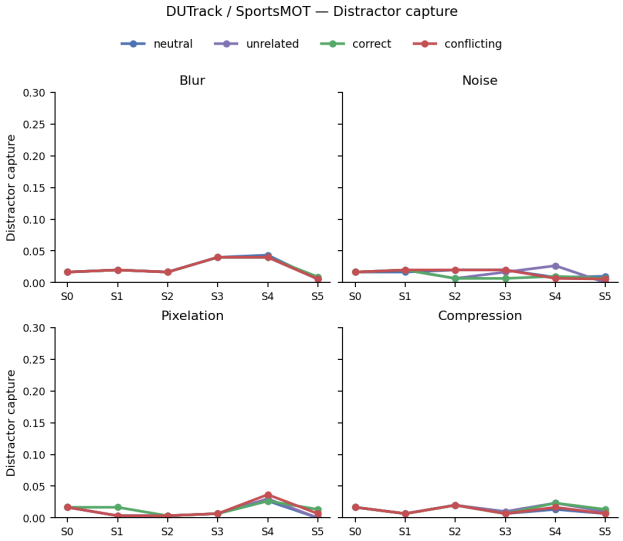

Figure 62. DUTrack on SportsMOT: distractor capture across all four perturbations, for all four text conditions.

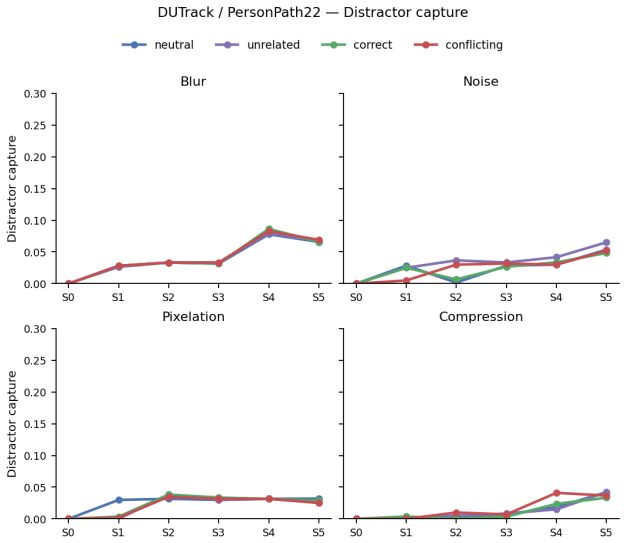

Figure 63. DUTrack on PersonPath22: distractor capture across all four perturbations, for all four text conditions.

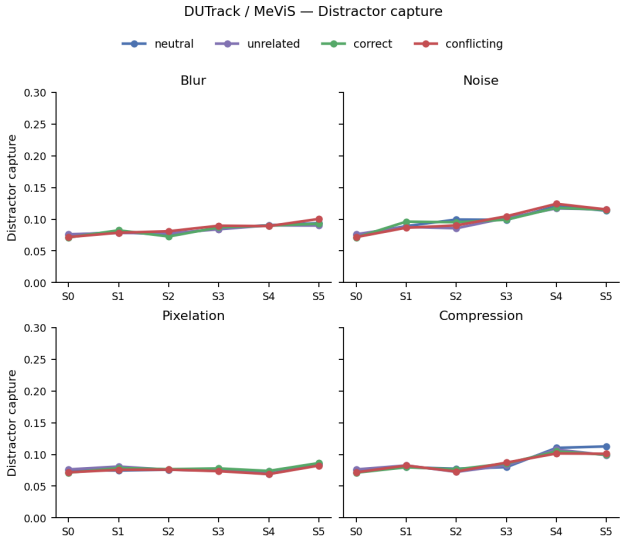

Figure 64. DUTrack on MeViS: distractor capture across all four perturbations, for all four text conditions.

B.5. Textual direction under degradation: full per-case statistics

These tables list the responsive-system effects with a CI-supported clean input value. S_0 is the clean signed effect; later entries are signed $S_5 - S_0$ changes in percentage points. Asterisks mark FDR-significant endpoint tests. Figure 65 visualizes the endpoint inventory, and Fig. 66 shows the corresponding full trajectories. Tables 9 and 10 provide the numerical target-accuracy and distractor-capture values.

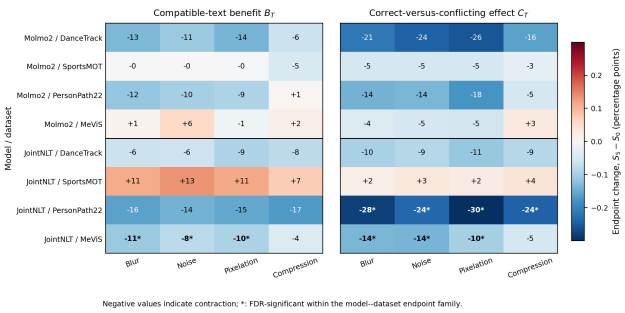

Figure 65. Endpoint change ($S_5 - S_0$) in responsive-system target-accuracy effects; asterisks mark FDR-significant endpoint tests.

Target-accuracy effects (B_T, C_T, or N_T).

Model	Dataset	Eff.	S_0	Blur	Noise	Pixel.	Compr.
JointNLT	MeViS	C_T	+14.2	-13.8^*	-13.9^*	-9.9^*	-5.2
JointNLT	PersonPath22	C_T	+26.9	-28.0^*	-23.5^*	-29.9^*	-24.2^*
Molmo2	DanceTrack	B_T	+12.8	-13.5	-11.0	-13.9	-6.4
JointNLT	MeViS	B_T	+10.6	-11.1^*	-8.1^*	-9.6^*	-4.1
Molmo2	DanceTrack	C_T	+25.3	-21.4	-24.3	-25.8	-16.0
JointNLT	MeViS	N_T	+7.8	-8.9^*	-7.8^*	-5.1	-3.4
JointNLT	DanceTrack	B_T	+10.7	-6.4	-6.1	-9.4	-7.8
JointNLT	DanceTrack	C_T	+10.9	-9.6	-8.6	-11.5	-8.5
JointNLT	PersonPath22	B_T	+17.6	-16.1	-14.1	-15.3	-17.0
JointNLT	PersonPath22	N_T	+17.0	-15.3	-10.8	-16.0	-10.2
JointNLT	SportsMOT	B_T	$-11.2^\dagger$	+11.2	+13.4	+11.3	+7.4
Molmo2	MeViS	B_T	+3.8	+1.1	+5.9	-1.4	+1.9
Molmo2	MeViS	C_T	+8.1	-4.4	-5.0	-5.4	+2.9
Molmo2	DanceTrack	N_T	+7.6	-9.8	-6.4	-8.6	-6.3

Table 9. Target-accuracy effects with a CI-supported clean-input value. Values after S_0 are signed endpoint changes; asterisks mark FDR-significant endpoint tests.

Distractor-capture effects (S_D).

Model	Dataset	S_0	Blur	Noise	Pixel.	Compr.
JointNLT	MeViS	+9.2	-8.2^*	-7.3^*	-7.0^*	-1.6
Molmo2	DanceTrack	+14.1	-7.6	-11.0	-12.8	-6.1
Molmo2	MeViS	+5.6	-4.4	-4.8	-5.6^*	+2.3
Molmo2	PersonPath22	+11.3	-12.6	-11.8	-11.9	-3.0

Table 10. Distractor-capture effects with a CI-supported clean-input value. Values after S_0 are signed endpoint changes; asterisks mark FDR-significant endpoint tests.

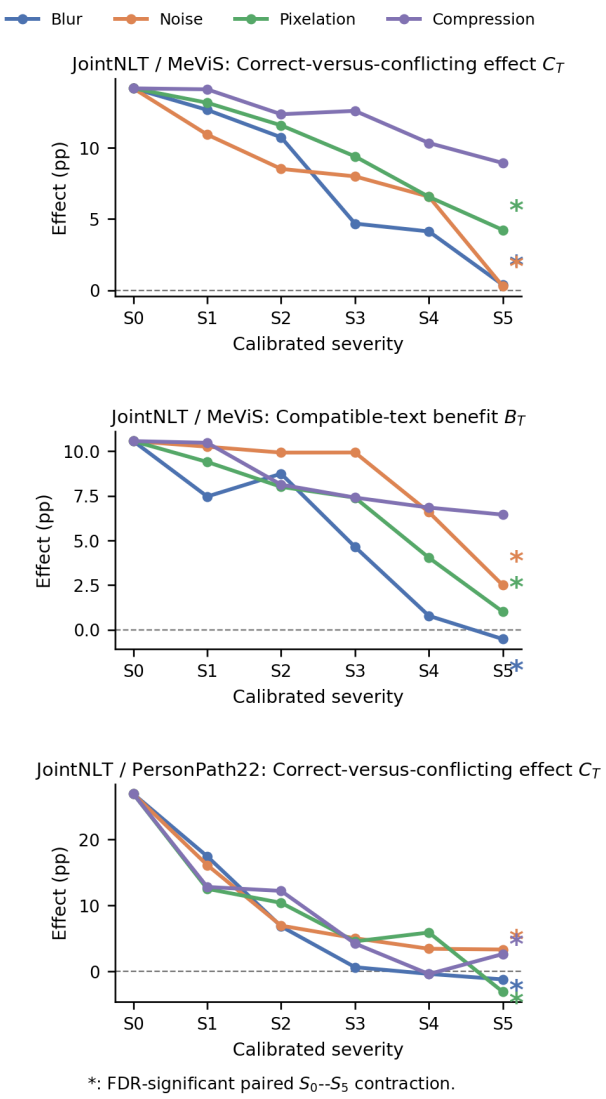

Figure 66. Full severity trajectories for the three model–dataset–effect cases with FDR-supported endpoint contraction.

B.6. Absolute performance and floor/ceiling diagnostics

Table 11 reports absolute target accuracy, distractor capture, and other/absent rate under correct text at S_0 and S_5.

	S_0			S_5 (pooled, 4 perturbations)		
Model	A_T	A_D	other/absent	A_T	A_D	other/absent
Molmo2	67.9%	2.4%	29.7%	37.4%	3.7%	58.8%
JointNLT	43.6%	9.0%	47.4%	19.3%	7.3%	73.4%
UVLTrack	64.4%	6.1%	29.5%	27.4%	7.6%	65.0%
MMTrack	62.9%	6.4%	30.7%	24.2%	7.0%	68.8%
DUTrack	68.1%	4.8%	27.1%	33.2%	6.6%	60.2%

Table 11. Correct-text outcome rates, pooled across datasets at S_0 and across datasets and perturbations at S_5. "other/absent" combines the two non-target, non-distractor labels.

B.7. Model internals

B.7.1. Saved states and hook semantics

Table 12 records the clean-MeViS hook map and pooled state semantics. Language states use [CLS] tokens or sentence vectors; visual states pool search tokens in target or distractor regions.

Model	Hooked modules/stages and stored pooled semantics
Molmo2	Decoder inputs 0–36 (post-connector input through final normalized decoder state): description-token state before coordinate generation and target/distractor visual states pooled over mask-covered image tokens.
JointNLT	BERT layers 0–11 under `language_backbone`; `text_proj`; six `vl_joint_trans` encoder blocks; target decoder; box-head input. We retain BERT/projected/fused [CLS], target/distractor search-region means, target query, and the spatially averaged pre-box-head feature.
UVLTrack	`backbone.bert` embeddings/layers 0–5; ViT blocks 0–11; backbone output. We retain text [CLS], target/distractor search-region means, and final text/search states before the box head.
MMTrack	Text encoder and `text_adj`; visual blocks; `bottleneck`; `vl_fusion`; coordinate decoder and predictor. We retain [CLS]/sentence states, target/distractor search-region means, and coordinate-query states through localization.
DUTrack	`backbone.descript_embedding`; first joint-block input; every text-reachable joint block; backbone output. We retain description/fusion-input [CLS], target/distractor search-region means, and the final visual state before temporal weighting/CENTER localization.

Table 12. Model-specific hook map and stored pooled states.

B.7.2. Full-depth Text BA sweeps

Tables 13–14 give the complete clean-input Text BA sweeps for JointNLT and Molmo2.

Stage	Text BA	q
text encoder L0	0.602	.002
text encoder L1	0.525	.098
text encoder L2	0.509	.316
text encoder L3	0.516	.202
text encoder L4	0.531	.040
text encoder L5	0.516	.196
text encoder L6	0.566	.002
text encoder L7	0.490	.760
text encoder L8	0.569	.002
text encoder L9	0.526	.067
text encoder L10	0.544	.006
text encoder L11	0.575	.002
text projection	0.563	.002
fusion L0 text	0.566	.002
fusion L0 $T-D$	0.518	.196
fusion L1 text	0.585	.002
fusion L1 $T-D$	0.535	.049
fusion L2 text	0.577	.002
fusion L2 $T-D$	0.690	.002
fusion L3 text	0.576	.002
fusion L3 $T-D$	0.699	.002
fusion L4 text	0.601	.002
fusion L4 $T-D$	0.673	.002
fusion L5 text	0.553	.006
fusion L5 $T-D$	0.685	.002
confirmatory late fusion $T-D$	**0.695**	.002
pre-loc. target query	0.519	.202
pre-loc. box feature	0.573	.002

Table 13. JointNLT correct-versus-conflicting Text BA at every hooked stage (0.5 is chance); q is FDR-corrected over the sweep. Bold denotes the decision-near stage used in the main paper.

Layer	BA	q
L0	0.499	1.000
L1	0.501	.567
L2	0.497	.688
L3	0.492	.702
L4	0.486	.797
L5	0.510	.441
L6	0.508	.459
L7	0.499	.688
L8	0.490	.709
L9	0.492	.701
L10	0.510	.441
L11	0.506	.441
L12	0.505	.554
L13	0.521	.275
L14	0.519	.370
L15	0.523	.252
L16	0.514	.441
L17	0.545	.053
L18	0.543	.057
L19	0.554	.033
L20	0.561	.033
L21	**0.634**	.015
L22	0.629	.015
L23	0.614	.015
L24	0.594	.015
L25	0.569	.015
L26	0.565	.015
L27	0.567	.015
L28	0.569	.015
L29	0.563	.015
L30	0.556	.025
L31	0.558	.015
L32	0.563	.015
L33	0.556	.025
L34	0.550	.057
L35	0.565	.015
L36	0.574	.015

Table 14. Molmo2 correct-versus-conflicting Text BA at all saved coordinate-query decoder depths (0.5 is chance); q is FDR-corrected over the 37-layer sweep. Bold denotes the peak.

B.7.3. JointNLT across visual degradation

JointNLT is evaluated on the clean MeViS manifest across all displayed severities. The query frame is clean; tracker state is reset for every run; and features are sampled at three fixed co-visible frames per source video. Figures 67, 68, 69, and 70 provide the four perturbation-specific trajectories, and Table 15 records their terminal late-fusion values.

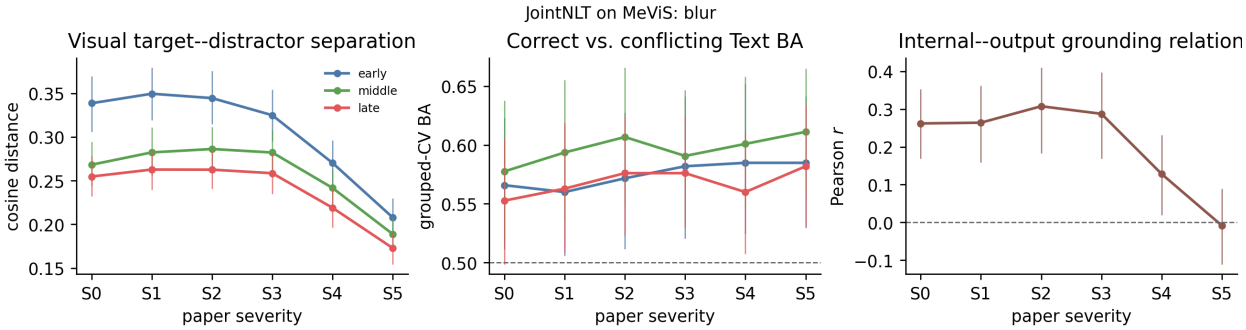

Figure 67. JointNLT MeViS internal trajectories under blur. Panels show visual separation, Text BA, and grounding-to-output association; error bars are 95% source-video bootstrap CIs.

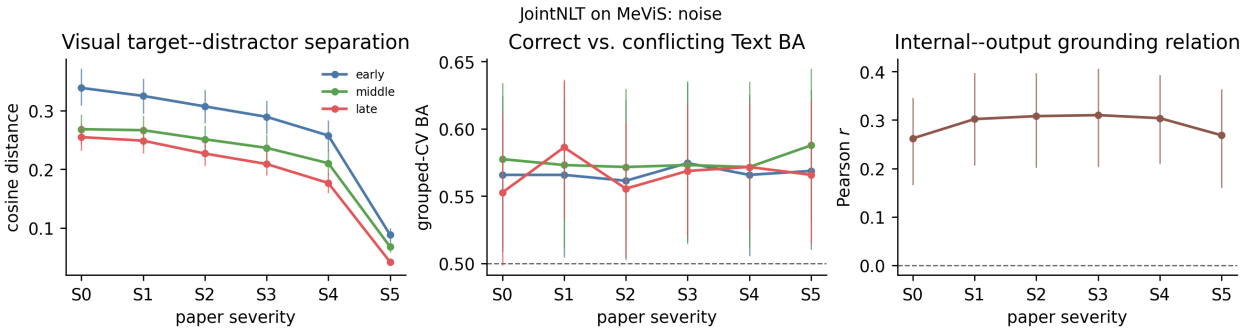

Figure 68. JointNLT MeViS internal trajectories under noise.

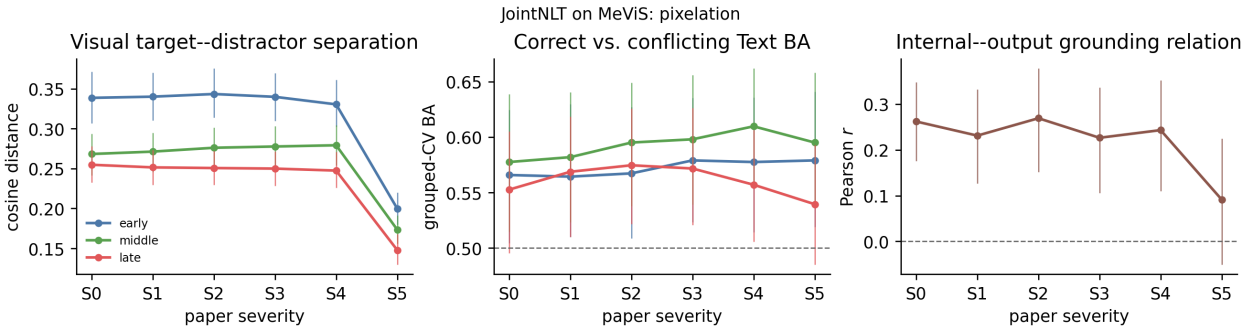

Figure 69. JointNLT MeViS internal trajectories under pixelation.

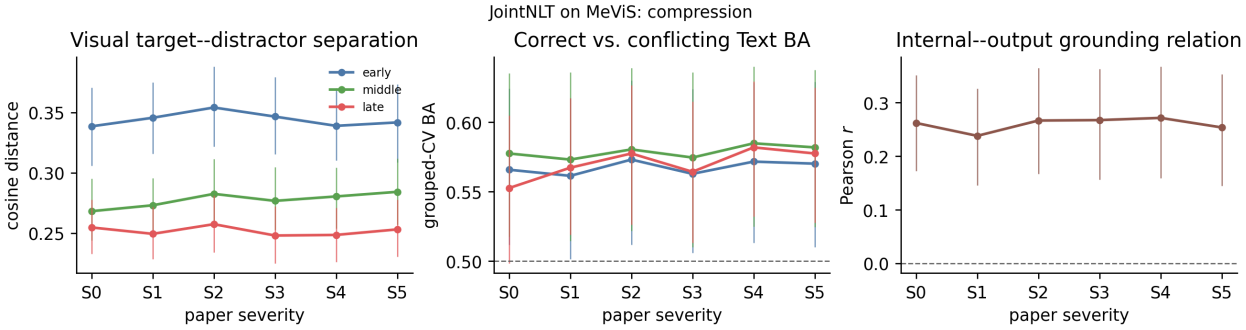

Figure 70. JointNLT MeViS internal trajectories under compression.

Perturbation	late visual sep.	late fused-text BA	output association r
blur	.173 [.154, .192]	.582 [.531, .633]	−.009 [−.111, .089]
noise	.042 [.036, .048]	.566 [.515, .621]	+.268 [.160, .363]
pixelation	.148 [.130, .168]	.539 [.485, .593]	+.091 [−.051, .225]
compression	.253 [.230, .278]	.577 [.528, .625]	+.254 [.145, .353]

Table 15. JointNLT late-fusion S_5 values by perturbation. Brackets are 95% source-video bootstrap CIs.

C. Ablations

This section reports protocol and intervention sensitivity checks.

C.1. Native-box-metrics sensitivity check

The four box trackers are also scored with native box IoU against tight mask-extent ground-truth boxes (IoU ≥ 0.5), using the same predictions and clustered inference as the point-in-mask metric. Table 16 reports the clean-input point and box-IoU effects side by side.

C.2. Causal donor-content controls

Each control replaces the state at the main-paper patch site while holding visual input, initialization, and sequential replay fixed.

JointNLT content controls. Table 17 lists the donor-state contrasts for the retention-control and persistent-capture strata.

Contrast (correct −)	Mean Δ	95% CI	q
Retention controls (n=67)			
neutral	+0.176	[0.090,0.274]	.0001
unrelated	+0.184	[0.101,0.280]	.0001
cross-video	+0.226	[0.140,0.326]	.0000
self (effect)	+0.049	[0.022,0.080]	.0006
Persistent captures (n=23)			
neutral	+0.359	[0.077,0.622]	.0251
unrelated	+0.193	[−0.045,0.413]	.1241
cross-video	+0.368	[0.110,0.634]	.0206
self (effect)	+1.193	[0.971,1.410]	.0000

Table 17. JointNLT donor-state contrasts by outcome stratum. Values are FDR-corrected within each layer–stratum family.

Molmo2 confirmation strata. Table 18 lists the L6/L12 donor-state contrasts by confirmation stratum.

Source state	Stratum	n	Layer	Mean Δm [95% CI]	p
correct	conflict capture	15	L6	+.868 [+.477,+1.266]	.0008
		15	L12	+.861 [+.470,+1.261]	.0012
correct	other/error	12	L6	+.242 [+.068,+.429]	.0332
		12	L12	+.232 [+.073,+.410]	.0604
correct	target retention	88	L6	+.013 [−.019,+.045]	.4253
		88	L12	+.021 [+.000,+.044]	.0618
unrelated	conflict capture	15	L6	+.770 [+.411,+1.136]	.0037
		15	L12	+.789 [+.401,+1.181]	.0027
unrelated	other/error	12	L6	+.029 [−.157,+.218]	.8219
		12	L12	+.133 [+.010,+.281]	.1282
unrelated	target retention	88	L6	+.007 [−.023,+.036]	.6450
		88	L12	+.005 [−.028,+.038]	.7749
correct−unrelated	conflict capture	15	L6	+.098 [+.020,+.218]	.0241
		15	L12	+.072 [−.024,+.204]	.3134
correct−unrelated	other/error	12	L6	+.213 [+.051,+.446]	.0308
		12	L12	+.099 [+.028,+.185]	.0635
correct−unrelated	target retention	88	L6	+.006 [−.037,+.051]	.7952
		88	L12	+.016 [−.016,+.052]	.3774

Model	Dataset	point	box-IoU	diff
B_T				
UVLTrack	DanceTrack	−0.8%	+0.4%	+1.2
UVLTrack	SportsMOT	+1.2%	+0.5%	−0.7
UVLTrack	PersonPath22	+5.3%	+4.1%	−1.2
UVLTrack	MeViS	+0.6%	+1.3%	+0.7
MMTrack	DanceTrack	−0.3%	+1.3%	+1.7
MMTrack	SportsMOT	−2.3%	−0.5%	+1.8
MMTrack	PersonPath22	+0.5%	−2.5%	−3.0
MMTrack	MeViS	+0.4%	+0.5%	+0.1
JointNLT	DanceTrack	+10.7%	+11.1%	+0.4
JointNLT	SportsMOT	−11.2%	−11.6%	−0.3
JointNLT	PersonPath22	+17.6%	+17.1%	−0.6
JointNLT	MeViS	+10.6%	+12.0%	+1.5
DUTrack	DanceTrack	+0.7%	+0.2%	−0.5
DUTrack	SportsMOT	−1.8%	−1.3%	+0.5
DUTrack	PersonPath22	−1.9%	−0.2%	+1.7
DUTrack	MeViS	+0.2%	−0.0%	−0.2
S_D				
UVLTrack	DanceTrack	−0.8%	+0.9%	+1.7
UVLTrack	SportsMOT	−3.0%	−3.3%	−0.3
UVLTrack	PersonPath22	−0.3%	−0.5%	−0.2
UVLTrack	MeViS	+0.4%	+0.3%	−0.1
MMTrack	DanceTrack	+0.9%	+1.3%	+0.5
MMTrack	SportsMOT	+0.0%	+0.0%	+0.0
MMTrack	PersonPath22	+0.9%	+0.2%	−0.8
MMTrack	MeViS	−0.7%	−0.6%	+0.1
JointNLT	DanceTrack	+4.6%	+6.0%	+1.3
JointNLT	SportsMOT	+1.9%	+2.6%	+0.8
JointNLT	PersonPath22	+12.3%	+11.4%	−0.9
JointNLT	MeViS	+9.2%	+7.4%	−1.9
DUTrack	DanceTrack	−0.3%	+0.0%	+0.3
DUTrack	SportsMOT	+0.0%	+0.0%	+0.0
DUTrack	PersonPath22	+0.0%	+0.0%	+0.0
DUTrack	MeViS	−0.5%	−0.1%	+0.3

Table 16. Clean-input B_T and S_D under point-in-mask and box-IoU scoring. "diff" is box-IoU minus point, in percentage points; underlining denotes a 95% source-video bootstrap CI excluding zero.

Table 18. Molmo2 L6/L12 donor-state contrasts. Δm is the trajectory-mean target-minus-distractor point-margin change from the conflicting baseline; CIs use source-video bootstrap and p is paired sign-flip. The bottom block compares correct and unrelated donors.

Causal audit of the behaviorally null systems. The three small-effect systems are tested on the 115-video manifest with matched-correct replacement and zero-text ablation at their declared text boundaries. Contrasts use the conflicting baseline; q is FDR-corrected within each system's two-contrast family. Table 19 reports the resulting semantic and presence contrasts.

Model	Contrast (vs. conflicting)	Mean Δ	q	Equiv.
UVLTrack	matched-correct (semantic)	+0.024	.365	No
	zero-ablation (presence)	+0.036	.007	No
MMTrack	matched-correct (semantic)	−0.012	.099	**Yes**
	zero-ablation (presence)	−0.022	.021	No
DUTrack	matched-correct (semantic)	+0.002	.573	**Yes**
	zero-ablation (presence)	−0.013	.379	**Yes**

Table 19. Output-margin change after matched-correct replacement (semantic) or zero-text ablation (presence), by system.

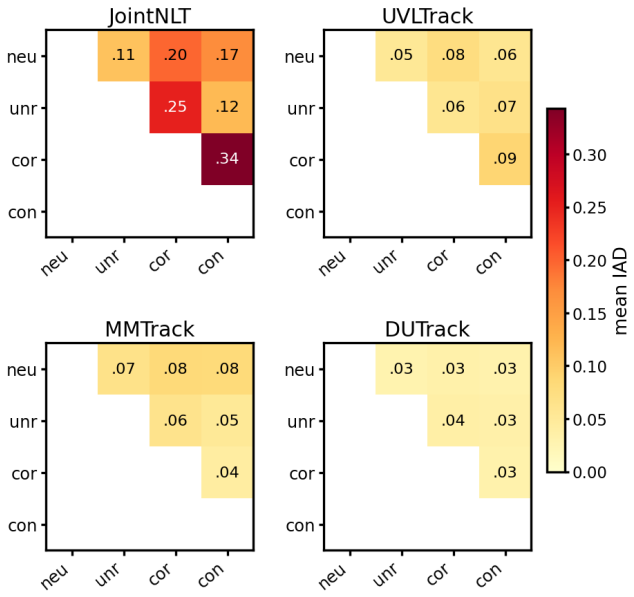

Figure 71. Outcome-independent pool: target-accuracy IAD, averaged equally across the three MOT sources and four perturbations over paper S_0–S_5. Labels abbreviate neutral/unrelated/correct/conflicting.

C.3. Outcome-independent sample confirmation

The main-paper MOT subsets intentionally retain rare Molmo2 conflicting-text takeover cases before seeded random filling. This confirmation instead retains the complete annotation-eligible DanceTrack, SportsMOT, and PersonPath22 pool: 404 pairs from 81 source videos (214/17, 110/31, and 80/33 pairs/videos, respectively). Pair inclusion requires at least five co-visible annotated frames, mean co-visible mask IoU below 0.3, and the required target, distractor, and deterministic category-matched unrelated expressions; it uses no model output, accuracy, or takeover label.

JointNLT, UVLTrack, MMTrack, and DUTrack (Molmo2 is out of scope) were run under all four text conditions at every displayed S_0–S_5 severity level of blur, noise, pixelation, and compression.

Figure 71 compactly summarizes each system's target-accuracy IAD matrices. Within a model, we equally average the three source-dataset matrices and, within a dataset, the four perturbation-specific IAD matrices over the paper's six-position severity ladder. Larger cells therefore indicate a larger average absolute condition gap, not a directional or inferential test.

Effect pattern. Figure 71 retains the main behavioral split under the outcome-independent distribution. JointNLT's correct–conflicting IAD is .50, .18, and .35 on DanceTrack, SportsMOT, and PersonPath22, respectively. The corresponding cells remain smaller for UVLTrack (.13, .06, .09), MMTrack (.03, .06, .04), and DUTrack (.02, .04, .03).

D. Reproducibility

We release the complete experimental pipeline as a self-contained, anonymized repository in the supplementary material, structured to reproduce every table, figure, and statistical test in this paper through a single configuration file and entry point rather than the roughly twenty scripts and five separate Python environments the pipeline spans internally. The repository provides the inference, validation, consolidation, statistics, and model-internals code exactly as used to generate the reported results; it introduces no reimplementation of the underlying methodology.

The archive does not redistribute raw third-party video or image frames. It provides code, frozen manifests and provenance metadata, annotations only where their upstream terms permit redistribution, and scripts that obtain and prepare source media from the original providers (Table 1, Table 2). The packaged operational smoke fixture is an author-created synthetic two-object sequence. It exercises the shared specialized-tracker input, four-condition text, perturbation, output-schema, and validation paths, but is not a scientific replication or an estimate of tracking or text-prior performance. Reproducing the full five-model study requires obtaining the source media, annotations, and checkpoints under their respective upstream terms. A CPU-only test suite covering adapter contracts, outcome-label invariants, and the statistical correction routines requires neither a GPU nor a model checkpoint.

Scaling from the synthetic operational fixture to the full paper-scale grid (every selected pair, all six severities) requires obtaining the complete datasets and model checkpoints per their documented licenses (Table 1, Table 2) and otherwise leaves the pipeline unchanged. The repository is anonymized for review; the complete, de-anonymized version will be made publicly available upon acceptance.

E. Ethics and Societal Impact

DISTRACT evaluates existing vision-language trackers; it does not train a tracker, collect new personal data, or release raw third-party media. Its purpose is to reveal failures under ordinary image degradation, not to certify a model as safe for deployment. Tracking can be used in privacy-sensitive or surveillance contexts, although the code adds no surveillance-specific capability. Users must obtain data lawfully, respect upstream terms and privacy obligations, and apply suitable oversight and safeguards before operational use.