Ensemble-Aware Structural Encoding (EASE): Conditioning Protein Encoders on Peptide Conformational Ensembles

Ruoyao Zhang[1] **Hongxuan Li**[2] **Jie Peng**[2] **Tianlong Chen**[2]
[1]University of Waterloo [2]University of North Carolina at Chapel Hill

arXiv:2610.27638v1 [cs.LG] 20 Oct 2026

Abstract

Protein foundation models provide transferable representations for downstream peptide prediction tasks. However, existing representations typically reduce a peptide to a single sequence or structure, which may miss conformational variability relevant to key peptide properties, including antibacterial activity, hemolysis, toxicity, stability, and binding. For example, protease resistance correlates more strongly with conformational dynamics measured by hydrogen exchange than with average helicity. Therefore, we introduce Ensemble-Aware Structural Encoding (EASE), a framework for effectively extracting structural information from conformational ensembles for peptide representation. To extract useful information from noisy conformers, EASE selects central conformers representing a peptide's major structural states and fuses their individual representations with a permutation-invariant operator to produce a peptide-level representation for downstream prediction. We then evaluate our framework on six peptide property-prediction tasks under leakage-controlled splits. Against matched single-structure baselines, EASE improves performance consistently on six tasks with a mean absolute gain of about 3.0% across two structural encoders. Gains keep growing as we add conformers through K=20. A single peptide conformer therefore provides a limited basis for structural representation, whereas our framework distilling and fusing multiple conformers yields a more informative representation.

1 Introduction

Peptides are short amino-acid chains, conventionally under fifty residues, that sit between small molecules and proteins in both size and pharmacology. Their ability to recognize biological targets, along with advances in peptide synthesis and chemical modification, made them increasingly important in therapeutics (Henninot et al., 2018; Muttenthaler et al., 2021). Accordingly, recent representation learning methods aim to predict peptide properties such as biological activity, toxicity, and stability (Zhang et al., 2025; Guntuboina et al., 2023; Zhang et al., 2026b). Many of these methods build on pretrained Protein Language Models (PLMs), which have become a widely used approach for learning transferable representations from amino-acid sequences. Structure aware models, such as ESM-3, extend sequence representations with structural information (Hayes et al., 2025), but typically process a single structure at a time. This presents a unique challenge for peptides. Peptides are conformationally heterogeneous, and their bioactive conformations are often low-probability states that a single static structure cannot fully capture (Zhang et al., 2026a). For example, protease resistance correlates more strongly with conformational dynamics measured by hydrogen exchange than with average helicity (Shi et al., 2013).

A conformational ensemble provides a more complete description of peptides' variability, but obtaining such ensembles traditionally required expensive molecular dynamics simulations (Karplus and McCammon, 2002; Hollingsworth and Dror, 2018). Recent generative models such as BioEmu, AlphaFlow, and Str2Str are trained on molecular dynamics to generate equilibrium-like structural ensembles directly, making the process affordable and scalable (Lewis et al., 2025; Jing et al., 2024; Lu et al., 2023). At the same time, recent studies suggest that conformational information can improve protein representations and downstream prediction (Shi and Oliver, 2026; Kalifa et al., 2026). However, existing peptide models primarily focus on sequence PLMs (Guntuboina et al., 2023; Chen et al., 2026), chemical representations (Feller and Wilke, 2025; Adamczyk et al., 2026), or a single predicted structure (Zhao et al., 2024; Aguilera-Puga and Plisson, 2024; Liu et al., 2026), leaving the value of conformational ensembles for peptide representation learning largely unexplored. Im-

portantly, simply generating more structures does not directly yield a better representation, as conformational ensembles can contain substantial redundancy, motivating the extraction of a compact subset of conformers that captures their structural diversity (Conev et al., 2023). This presents two related questions: how to effectively extract useful information from noisy conformational data, and to what extent such information can complement representations learned by PLMs.

To address these questions, we introduce Ensemble-Aware Structural Encoding (EASE), a general framework for extracting predictive structural information from conformational ensembles. From a large sampled conformers set, EASE performs structural clustering and selects a subset containing central conformers. Each representative conformer passes through a shared and adapted structural encoder, and their representations are fused through a permutation-invariant operator to obtain a single peptide representation. Our contributions include: ❶ **Ensemble-Aware Structural Encoding (EASE).** We introduce a framework that systematically exploits conformational diversity of peptides, producing representations that improve downstream prediction. ❷ **Consistent gains across structural encoders and peptide properties.** Across two distinct pretrained structural encoders, GearNet-Edge and ESM-3, and six downstream tasks, EASE outperforms a matched single-structure baseline by $+3.8\%$ on GearNet-Edge across all six tasks and $+2.2\%$ on ESM-3 across five of them. ❸ **Analysis of how ensemble information should enter the model.** We systematically compare EASE against six alternative mechanisms for injecting ensemble-derived summary statistics into the model. EASE provides the most consistent improvement, while the alternatives recover only a limited portion of the gain. ❹ **Analysis of ensemble information and protein language models.** We examine when conformational information provides additional predictive value beyond Protein Language Model (PLM) representations, through comparisons across PLM, structural, and PLM-conditioned structural encoders.

2 Related work

2.1 Peptide representation learning

Existing peptide models are mainly based on three approaches. Sequence-based approaches adapt protein language models to the peptide domain or train peptide-specific models on sequences, chemical representations use notations such as SMILES and HELM to non-standard amino-acids, and structure-aware approaches further incorporate geometric information from 3D structures (Fernández-Díaz et al., 2024; Guntuboina et al., 2023; Feller and Wilke, 2025; Lee et al., 2025; Zhang et al., 2025; Zhao et al., 2024; Liu et al., 2026; Yu et al., 2026). As a sequence-based example, PeptideBERT finetunes ProtBERT directly on downstream peptide properties prediction such as hemolysis and solubility (Guntuboina et al., 2023). For chemical representations, PeptideCLM learns from SMILES strings, enabling it to represent non-canonical residues and chemical modifications beyond standard amino-acid vocabularies (Feller and Wilke, 2025). For structure-aware learning, deepAMPNet combines sequence embeddings with AlphaFold2-predicted structures through a graph neural network to identify antimicrobial peptides (Zhao et al., 2024). Despite their different input modalities, most pretrained models consume a single sequence, chemical string, or structural conformation at a time. In contrast, our approach takes into account multiple conformers of the same peptide to learn an ensemble-informed representation.

2.2 Learning from conformational ensembles

Representation learning over conformational ensembles has been explored primarily on small molecules and later extended to peptides and proteins (Axelrod and Gómez-Bombarelli, 2023; Zhu et al., 2023; Nguyen et al., 2026; Feller et al., 2026; Wu et al., 2022; Rollins et al., 2024; Kalifa et al., 2026; Shi and Oliver, 2026). Axelrod and Gómez-Bombarelli encode multiple conformers for small molecules and apply attention over their representations (Axelrod and Gómez-Bombarelli, 2023). Additionally, MARCEL systematically evaluates conformer ensembles across molecular GNNs using aggregation methods such as mean pooling, DeepSets, and self-attention (Zhu et al., 2023). The work most closely related to our peptide setting is EnsembleEGNN, which encodes cyclic peptide conformers with shared EGNN layers, applies attention over them, and pretrains the resulting ensemble model on cyclic peptide conformations before downstream permeability task (Feller et al., 2026). As an example of protein representation models using ensembles, DynamicsPLM combines structural tokens from multiple conformers at each residue before processing them with a protein lan-

guage model (Kalifa et al., 2026). Alternatively, ENSEMBITS learns discrete tokens that represent the conformational ensemble per residue (Shi and Oliver, 2026). In contrast to these works, EASE additionally considers how a large sampled ensemble should be distilled before representation learning. Additionally, it addresses whether existing pretrained single-structure encoders can be adapted to learn peptide representations from conformational ensembles without ensemble-specific pretraining. Our inspiration came from multiple-instance learning, in which representations are learned from bags of instances through permutation-invariant aggregation (Zaheer et al., 2017; Ilse et al., 2018). We adopt a similar idea, providing an encoder-agnostic way to extract ensemble-informed peptide representations for downstream prediction.

3 Method

3.1 Pipeline Overview

Let a peptide's conformational ensemble $\mathcal{E}$ be an unordered, unweighted, and variable-size set of conformer structures, generated by an ensemble sampler. We perform structural clustering to select a fixed bank $\mathcal{B}_K \subset \mathcal{E}$ of K conformers from $\mathcal{E}$. Each conformer x_k passes through a shared adapted encoder trunk $f_{\theta,\phi}$, where pretrained weights θ are frozen and low-rank adapters ϕ are shared across conformers. Their representations are then fused to produce a single loss per peptide:

$$\begin{aligned} h_k &= f_{\theta,\phi}(x_k), \quad k = 1, \ldots, K; \\ z &= \Phi(h_1, \ldots, h_K); \\ \hat{y} &= g_\psi(\mathrm{LN}(z)), \end{aligned}$$

where Φ is a permutation-invariant fusion operator (§3.4) and g_ψ is the task head. Because the adapters are shared, gradients from all K conformers contribute to each update of ϕ.

To test whether EASE transfers across structural representations, we evaluate it with two pretrained structural encoders: ❶ **GearNet-Edge.** Each conformer is represented as a relational residue graph. Per-residue outputs are mean-pooled over valid positions to obtain $h_k \in \mathbb{R}^{512}$. LoRA adapters (Hu et al., 2021) are applied to the linear and self-loop maps of every node and edge layer, with one factor zero-initialized so the adapted model matches the pretrained encoder at initialization. ❷ **ESM-3.** Each conformer is tokenized by ESM-3's structure track and combined with the peptide's sequence track. The sequence track outputs remain the same across the K inputs so only the structure track varies. Per-residue outputs over valid positions are mean-pooled to obtain $h_k \in \mathbb{R}^{1536}$, which is then normalized before the fusion operator.

3.2 Ensemble generation

As the first step of the pipeline, we generate peptide conformational ensembles $\mathcal{E}$ using BioEmu (Lewis et al., 2025), sampling 150 conformers per peptide. After applying BioEmu's physical-validity filtering, we retain a median of 145 conformers (range: 72–150; all $\geq$ 50) per peptide. Since BioEmu generates backbone-frame structures only, we reconstruct side chains with H-Packer (McPartlon and Xu, 2023) and perform local energy minimization in OpenMM (Eastman et al., 2017) using Amber99SB (Hornak et al., 2006) and OBC2 implicit solvent (Onufriev et al., 2004).

3.3 Conformer banks and single-structure baseline selection method

The next step is selecting the K conformers from $\mathcal{E}$. For each peptide, we compute pairwise Cα RMSD between all sampled conformers after pairwise Kabsch alignment. We apply GROMOS clustering (Daura et al., 1999), where the conformer with the most neighbors within a cutoff forms a cluster with those neighbors; the selected cluster is then removed and the procedure repeats. We then take the medoid (the conformer with the lowest mean RMSD to all other conformers in the cluster) of each of the K largest clusters. To accommodate differences in peptide length and flexibility, the GROMOS clustering cutoff is defined relative to each peptide's pairwise-RMSD distribution. We progressively use quantiles $\{0.25, 0.15, 0.10, 0.05, 0.02, 0.01\}$ until at least K clusters are obtained. [1]

3.4 Fusion operators

After each conformer in $\mathcal{B}_K$ passes through the encoder, the output embeddings enter a fusion operator. We compare four permutation-invariant fusion operators Φ, mapping K d-dimensional embeddings to a single d-vector so that the downstream head is identically shaped across arms. ❶ **Mean.** Parameter-free mean pooling weights all conformers equally: $\Phi(h_1, \ldots, h_K) = \frac{1}{K}\sum_k h_k$. ❷ **Gated.** Each conformer receives a scalar gate,

[1] Across 900 peptides sampled from six tasks, 81% yield at least K clusters at the first cutoff and 100% at the second.

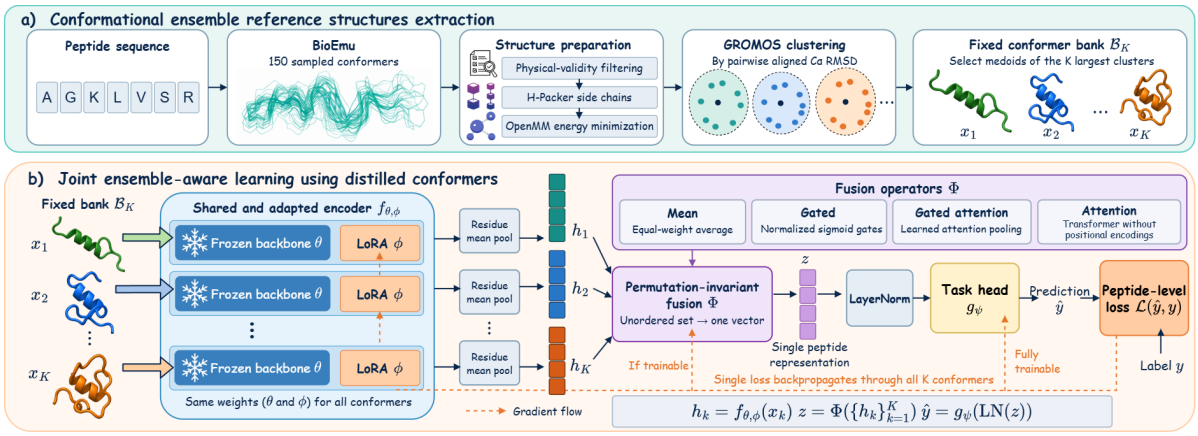

Figure 1: Ensemble-Aware Structural Encoding pipeline.

adding $d+1$ parameters: $s_k = \sigma(u^\top h_k + b)$, $w_k = s_k / \sum_j s_j$, $\Phi = \sum_k w_k h_k$. ❸ **Gated attention.** We use the multiple-instance-learning pooling of Ilse et al. (2018): $s_k = w^\top[\tanh(Vh_k) \odot \sigma(Uh_k)]$, $a_k = \exp(s_k)/\sum_j \exp(s_j)$, $\Phi = \sum_k a_k h_k$, where $V, U \in \mathbb{R}^{L_a \times d}$. We use the $L_a = 128$ from Ilse et al. (2018). ❹ **Attention.** We apply one Transformer encoder layer over the conformer axis without positional encodings and mean-pool its outputs. The layer is therefore permutation equivariant, and mean pooling makes Φ permutation invariant. Due to large added parameters, we reduce the width d before the Transformer layer, down-projecting to 64 for GearNet-Edge and 128 for ESM-3. Table 5 in the Appendix presents a summary of added parameters by each fusion operator.

4 Experimental setup

Tasks. We evaluate on six peptide property-prediction tasks (Table 1), comprising five peptide-intrinsic tasks and one peptide-protein interaction task. The tasks span key endpoints in peptide development, including safety (toxicity and hemolysis), efficacy (E. coli minimum inhibitory concentration), pharmacokinetics (blood stability), and target engagement (peptide-protein interaction). Five tasks are drawn from PepBenchmark (Zhang et al., 2026b); blood stability is drawn from PepMSND (Hu et al., 2025). For all tasks, we restrict peptides to canonical amino acids with lengths up to 50 residues. Two tasks include auxiliary inputs in addition to the peptide representation. For peptide-protein interaction, we represent the receptor using a frozen ESM-2-150M encoder, mean-pooled and fixed at 640 dimension (Lin et al., 2023). Receptor representation is concatenated with the peptide representation before the prediction head. For blood stability, assay conditions, including species and in vivo vs. in vitro, are encoded as a six-dimensional one-hot vector and concatenated before the prediction head. These auxiliary inputs are held fixed across all experiments, ensuring that only the peptide representation changes for each ablation.

Table 1: The six peptide property-prediction tasks.

Task	Peptides	Metric	Application
Toxicity	4,056	AUROC	Safety
Hemolytic	4,259	AUROC	Blood compatibility
Hemolytic HC50	1,926	RMSE	Blood compatibility
E. coli MIC	3,161	RMSE	AMP activity
Blood stability	406	AUROC	Proteolytic stability
PepPI	5,931	AUROC	Target binding

Splits. The five single peptide tasks use a homology-controlled five-fold cross validation. Sequences are clustered with MMseqs2 (Steinegger and Söding, 2017) at 50% identity, and entire sequence-similarity components are assigned to folds to minimize homology leakage. In each run, three folds are used for training, one for validation, and one for testing, with each fold serving as the test set once. For peptide–protein interaction, we use PepBenchmark's released peptide-cold splits, which ensure that test peptides do not occur in the corresponding training set.

Baselines. To evaluate our method, we compare EASE with several baselines. ❶ **Reference conformer.** We define the reference conformer as the medoid conformer from each peptide's $K = 10$ bank, as stated in §3.3. This baseline uses the same encoder, adapters, prediction head, and training protocol as the EASE method, but receives only the reference conformer instead of the full

K conformer bank. ❷ **PLM baseline.** We consider a sequence-only baseline as a conventional reference against protein language model representations for each encoder. For ESM-3, this baseline masks the structure track but keeps the sequence track with all other model and training settings unchanged. Since GearNet-Edge is fully structural, we use ESM-2 (Lin et al., 2023) with the same LoRA configuration and training protocol, selecting the 35M variant as the closest in size to the 20.2M parameter GearNet-Edge model. ❸ **Chemical representation baseline.** To offer a comparison with SMILES-based chemical models, we additionally evaluate the downstream tasks on ChemBERTa-77M (Chithrananda et al., 2020), a pretrained molecular language model.

Ablation studies on the ensemble size. To study the effect of the ensemble size on downstream performance, we conduct ablation studies on K, the size of the selected conformer bank. We evaluate $K \in \{5, 10, 20\}$ while holding the reference conformer fixed. For each K, the bank consists of the reference conformer and $K - 1$ cluster medoids selected by running the same GROMOS clustering procedure (§3.3) on the remaining conformer pool.

Training protocol. Hyperparameters are selected on validation performance separately for the two encoders. A table of hyperparameters are provided in the Appendix. For GearNet-Edge, we tune LoRA rank and prediction-head dropout using a grid shared across comparison arms and select one configuration per task, preventing arm-specific retuning. For ESM-3, we tune the trunk learning rate and batch size separately for each task and comparison arm. In all experiments, the checkpoint with the best validation performance is used for test evaluation. GearNet-Edge experiments use five folds and four training seeds per fold, and ESM-3 experiments use five folds and three seeds per fold. Seeds vary parameter initialization and minibatch ordering within each fixed split.

Statistics reporting. We report test performance as mean $\pm$ standard deviation over fold averages, with training seeds averaged within each fold. For comparisons with the baselines, we first average over seeds within each fold and compute paired differences between methods. We report the mean paired difference and the number of folds in which the EASE method improves over the baselines.

5 Results

5.1 EASE improves over a single structure

Table 2 reports absolute test performance for both encoders with all baselines specified in §4. Table 3 reports paired differences between EASE with mean fusion over ten conformers and the matched reference conformer baseline defined in §4. Averaged across the six tasks, EASE improves over the baseline by $+0.0274$ with GearNet-Edge and $+0.0162$ with ESM-3. Specifically, on GearNet-Edge, EASE improves over all six tasks and 28 of 30 folds. On ESM-3, EASE improves on five of six tasks over 23 of 30 folds. The exception is blood stability, where EASE underperforms the reference conformer by -0.0108 on average and loses on all five folds. A plausible explanation is the task's much smaller dataset compared to the other tasks, making its fold-to-fold performance volatile for both encoders.

Excluding blood stability, the remaining five tasks show a similar average improvement for both encoders ($+0.0248$ for GearNet-Edge, $+0.0216$ for ESM-3). This similar improvement over most tasks despite the choice of encoders is notable because GearNet-Edge and ESM-3 represent structure differently. GearNet-Edge directly builds a relational graph over the residues, whereas ESM-3 uses a VQ-VAE and represents local structure through discrete structure tokens. The task-level gains are less consistent across encoders. For example, GearNet-Edge gains more on peptide-protein interaction ($+0.0267$ versus $+0.0191$) and HC50 regression ($+0.0321$ versus $+0.0238$), while ESM-3 gains more on E. coli MIC regression ($+0.0327$ versus $+0.0196$). Thus, the two encoders agree on the overall benefit of multi-conformer encoding, but not on which tasks benefit most.

5.2 Simple mean fusion is sufficient for EASE

In this section, four permutation-invariant fusion methods are compared as stated in §3.4. Table 3 reports their paired differences over the same folds against the reference conformer baseline, and the absolute performance can be found in the Appendix 9. Mean pooling performs as well as or better than the learned fusion operators overall. On ESM-3, its average improvement over the baseline is $+0.0162$, compared with $+0.0140$ for the scalar gate, $+0.0111$ for transformer attention, and $+0.0108$ for gated attention. On GearNet-Edge, the scalar gate is marginally higher on av-

Table 2: Absolute test performance. Mean ± standard deviation over folds × training seeds. AUROC for classification and RMSE for regression. Average rank is obtained by ranking all arms on each task and averaging their ranks across the six tasks.

Encoder	Arm	Toxicity ↑	Hemolytic ↑	HC50 ↓	*E. Coli* MIC ↓	Stability ↑	Interaction ↑	Avg. rank ↓
ESM-3	Sequence track only	0.7896 ± 0.0640	0.9010 ± 0.0245	0.5540 ± 0.0328	0.6667 ± 0.0294	0.6983 ± 0.0231	0.8308 ± 0.0121	3.33
	Reference conformer	0.7792 ± 0.0534	0.9013 ± 0.0209	0.5793 ± 0.0230	0.6842 ± 0.0225	**0.7141 ± 0.0266**	0.8329 ± 0.0118	3.08
	EASE, mean $K = 10$	**0.8058 ± 0.0387**	**0.9071 ± 0.0161**	0.5554 ± 0.0334	**0.6515 ± 0.0183**	0.7032 ± 0.0314	**0.8520 ± 0.0100**	**1.67**
GearNet-Edge	ESM-2 (35M)	0.7652 ± 0.0597	0.8936 ± 0.0184	**0.5505 ± 0.0254**	0.6657 ± 0.0205	0.6558 ± 0.0694	0.8061 ± 0.0073	4.00
	Reference conformer	0.6893 ± 0.1008	0.8617 ± 0.0221	0.6460 ± 0.0191	0.6847 ± 0.0202	0.6737 ± 0.0458	0.8059 ± 0.0159	6.33
	EASE, mean $K = 10$	0.7231 ± 0.0938	0.8733 ± 0.0225	0.6139 ± 0.0164	0.6651 ± 0.0275	**0.7141 ± 0.0525**	0.8327 ± 0.0101	3.92
	ChemBERTa	0.6473 ± 0.0922	0.8677 ± 0.0247	0.6003 ± 0.0251	0.7094 ± 0.0251	0.5773 ± 0.1116	0.8351 ± 0.0156	5.67

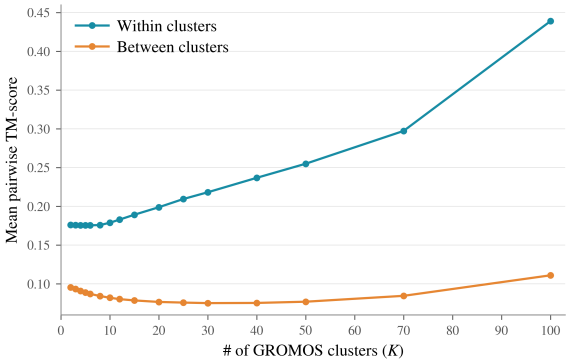

Figure 2: Structural separation over conformer bank sizes K. Mean within-cluster and between-cluster TM-scores.

erage (+0.0285 versus +0.0274), while gated attention and transformer attention reach +0.0268 and +0.0186, respectively. Across both encoders, no learned operator improves over mean pooling by more than 0.0011 on average. We note that mean pooling introduces no additional parameters. Therefore, the improvement of the mean fusion operator cannot be explained by additional fusion capacity, and the principal change is that the model receives multiple conformers instead of one.

Further, scalar gate and gated attention build on a strong premise that a few members of the bag carry the label while the rest are uninformative, and the fact that it fails to improve on a simple average suggests that identifying and upweighting a small subset of conformers is unnecessary in our setting. To offer potential explanation, we study additional characterizations of the sampled ensembles. Across 20,386 peptides, conformers from the same peptide are highly diverse: the median mean pairwise TM-score is 0.053, and 88% of peptides have a mean pairwise score below 0.2 (Table 4). This suggests that most peptides' sampled conformers exhibit a continuum instead of interconverting states. Diversity is particularly high among shorter peptides, which make up much of our benchmark.

We find similarly weak separation when applying the GROMOS clustering procedure used to construct the conformer banks. Across 1,000 peptides sampled across all length buckets, the mean within-cluster TM-score at $K = 10$ is only 0.179. Moreover, the difference between within- and between-cluster similarity continues to increase as K grows from 2 to 100, without showing an early plateau (Figure 2). These observations do not support a scenario where a small number of well-separated conformational groups dominate the sampled ensembles. Overall, uniformly aggregating representations across the sampled ensemble is at least as effective as learning conformer-specific weights while adding no trainable parameters. We therefore use mean pooling as the default fusion operator in the remaining experiments.

5.3 Performance improves with bank size

In this section, we report the ablation study on the conformer bank size K, varying K over 5, 10, 20. Figure 3 shows how performance changes with ensemble size relative to the reference conformer baseline (treated as K=1). Absolute and paired performance (Tables 10, 11) are given in the Appendix. Averaged across tasks, both encoders benefit increasingly from larger ensembles through $K = 20$, with mean improvements rising from +0.0213 at $K = 5$ to +0.0274 at $K = 10$ and +0.0313 at $K = 20$ for GearNet-Edge and from +0.0113 at $K = 5$ to +0.0162 at $K = 10$ and +0.0180 at $K = 20$ for ESM-3. The task-level trends are not uniformly monotonic, as some gains saturate or fluctuate and ESM-3 does not improve blood stability. Nevertheless, the overall trend shows that the benefits of EASE increase with diminishing returns as additional conformers are included.

This trend is consistent with the structural diversity measured in the sampled ensemble study in §5.2. As K increases, the selected conformers continue to cover structurally distinct regions of the sampled ensemble, while the marginal benefit of each additional conformer decreases. In this sense, K acts primarily as a sampling budget, as

Table 3: EASE with different fusion operators, paired difference over same folds against reference conformer (folds improved).

Encoder	Fusion	Toxicity	Hemolytic	HC50	*E. coli* MIC	Stability	Interaction	Mean
GearNet-Edge	Mean	**+0.0339 (5/5)**	+0.0116 (5/5)	+0.0321 (5/5)	+0.0196 (5/5)	**+0.0404 (3/5)**	+0.0267 (5/5)	+0.0274
	Gated	+0.0331 (5/5)	+0.0133 (5/5)	+0.0345 (5/5)	**+0.0208 (5/5)**	+0.0396 (3/5)	**+0.0294 (5/5)**	**+0.0285**
	Attention	+0.0318 (5/5)	+0.0093 (4/5)	**+0.0353 (5/5)**	+0.0091 (3/5)	+0.0109 (3/5)	+0.0151 (5/5)	+0.0186
	Gated attn	+0.0333 (5/5)	**+0.0136 (5/5)**	+0.0350 (5/5)	+0.0200 (4/5)	+0.0317 (3/5)	+0.0271 (5/5)	+0.0268
ESM-3	Mean	**+0.0266 (5/5)**	**+0.0058 (4/5)**	+0.0238 (4/5)	+0.0327 (5/5)	−0.0108 (0/5)	**+0.0191 (5/5)**	**+0.0162**
	Gated	+0.0133 (5/5)	+0.0046 (4/5)	+0.0150 (4/5)	**+0.0372 (5/5)**	**+0.0075 (4/5)**	+0.0062 (4/5)	+0.0140
	Attention	+0.0204 (4/5)	+0.0057 (4/5)	+0.0282 (4/5)	+0.0331 (5/5)	−0.0043 (1/5)	−0.0163 (0/5)	+0.0111
	Gated attn	+0.0116 (5/5)	+0.0036 (4/5)	**+0.0330 (5/5)**	+0.0355 (5/5)	−0.0217 (1/5)	+0.0026 (3/5)	+0.0108

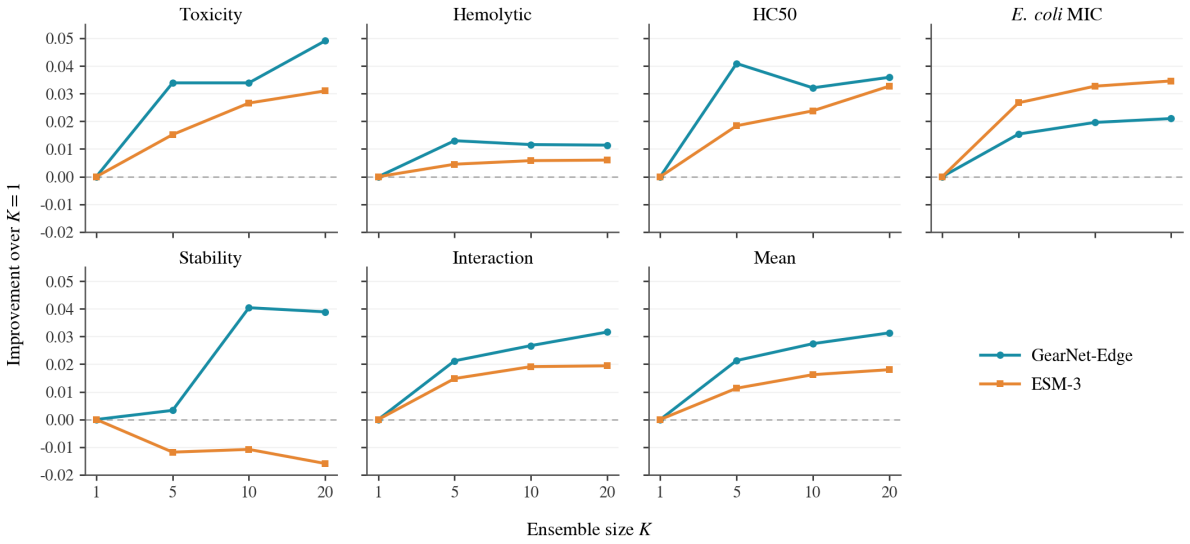

Figure 3: Effect of ensemble size on task performance. Each panel reports the paired difference with the reference conformer on the same folds. Positive values indicates improvement for both classification (AUROC) and regression (RMSE) tasks.

Table 4: Structural diversity of the sampled conformer sets. Length (Len.) and number of conformers (Conf.) are reported as medians across peptides. Pairwise TM-scores are computed between conformers of the same peptide and summarized over each task's sampled corpus.

Task	Len.	Conf.	Pairwise TM-score		
			Median [IQR]	$\% < 0.2$	$\% \geq 0.5$
Toxicity	36	142	0.095 [.055-.307]	67.0	12.8
Hemolytic	20	145	0.054 [.033-.113]	86.0	3.5
HC50	17	146	0.052 [.036-.083]	93.8	0.9
E. coli MIC	20	144	0.051 [.033-.086]	92.5	1.0
Stability	21	142	0.054 [.033-.087]	89.9	0.7
PepPI	11	145	0.041 [.027-.061]	99.1	0.1

larger K values provide broader ensemble coverage at proportionally greater computational cost. This trade-off motivates that we use $K = 10$ in the remaining experiments. It achieves approximately 88% of GearNet-Edge's average improvement observed at $K = 20$ and 90% of ESM-3's, while requiring half as many conformer encodings.

5.4 Alternative ways of using the ensemble

In addition to EASE, we compare six alternative mechanisms on GearNet-Edge, using the same conformational ensemble, data splits, and training protocol. The alternative methods incorporate ensemble summary statistics through graph features, edge geometry, topology, intermediate layers, the readout, or auxiliary supervision. Full definitions are provided in Appendix B.1.

Alternative mechanisms comparison. Table 12 in the Appendix reports the paired differences over same folds against each method's baseline. Across all methods of injecting ensemble-derived statistics into the model, average improvements over the base encoder span a fivefold range, from $+0.0028$ to $+0.0154$. The two mechanisms that make the ensemble statistics act throughout the encoder perform best: stochastic topology (3) improves by $+0.0154$, and layer-wise modulation (4) by $+0.0130$ with residue statistics and $+0.0117$ with residue and pair statistics. The mechanisms that inject a fixed summary at a single point tend to

perform worse: node and edge features (1) improve by $+0.0067$ and $+0.0064$, the pooled descriptor at the readout (5) by $+0.0057$, and auxiliary supervision (6) by $+0.0036$ and $+0.0028$. These results show that the injection point matters among the summary-statistics mechanisms.

Q: Does increasing the detail of the summary statistics help? A: No. Replacing the first two moments of the pairwise distance distribution with a 16-bin distributional representation (2) improves over the base encoder by $+0.0055$, slightly below the $+0.0064$ for the pairwise moments it replaces (1). Adding pairwise statistics on top of residue statistics similarly changes average performance by -0.0003. Within these tests, retaining additional pairwise or distributional detail does not recover the advantage of encoding the conformers directly.

Q: Direct conformer encoding vs. ensemble summary statistics? A: Direct conformer encoding. EASE improves over its reference-conformer baseline by $+0.0274$ on average, compared with $+0.0028$ to $+0.0154$ for the alternative summary-statistics approaches over their base encoder. These results favor preserving entire conformers as structural inputs rather than first compressing the ensemble to predefined statistics. One possible explanation is that summary statistics such as fluctuation, contact frequency, and distance moments marginalize over the ensemble before structural encoding and might abstract away spatial arrangement of residues within each structure. On the other hand, EASE preserves such joint geometry of each sampled structure for the pretrained encoder to process. We therefore conclude that the summary statistics tested here do not recover the information available from directly encoding the conformers.

5.5 When Does Ensemble Information Complement Protein Language Models?

Comparison with PLM baselines. We first investigate whether conformational ensembles provide additional predictive information beyond pretrained protein language model (PLM) representations. Table 2 compares EASE and its reference conformer baselines against sequence-only PLM representations. ESM-3 provides the most controlled comparison because its PLM baseline uses the same pretrained backbone, adapters, and training schedule, only with the structure track masked. The matched reference conformer underperforms sequence-only ESM-3 by 0.0058 on average and performs better on three of six tasks. EASE reverses this pattern, improving over sequence-only ESM-3 by $+0.0104$ on average and on five of six tasks. The only exception is HC50, where sequence-only ESM-3 is better by 0.0014. Thus, although adding a single structure provides no consistent advantage over the PLM track alone, conditioning the same model on a conformational ensemble using EASE can provide complementary predictive information across most tasks. On the other hand, since GearNet-Edge has no sequence channel on its own, we instead compare it with ESM-2-35M, a separately pretrained PLM of comparable size using the same downstream adaptation protocol. Unlike the ESM-3 comparison, differences here can reflect both input modality and pretrained encoder. The reference conformer on GearNet-Edge generally underperforms ESM-2, whereas EASE narrows this gap and outperforms ESM-2 on three of six tasks. The largest advantages occur on blood stability and peptide-protein interaction, where conformational information matters biologically: proteolytic stability depends in part on the accessibility and local geometry of potential cleavage sites, and protein binding depends on the availability of conformations compatible with the receptor protein surface. As a result, although EASE on GearNet-Edge does not consistently outperform the PLM baseline, conditioning the structure encoder on an ensemble provides competitive predictive information without relying on pretrained language representations.

Encoder-dependent ensemble gains and conformational variation. To further investigate how the integration of PLM representations influence the benefit of ensemble conditioning, we extend our evaluation to ESM-GearNet, which incorporates pretrained ESM-2 representations into a geometric graph encoder. EASE achieves an average paired improvement of +0.0077 across six tasks on ESM-GearNet, which is substantially smaller than +0.0162 on ESM-3 and +0.0274 on GearNet-Edge (Tables 13, 14). To examine whether these differences are related to how the encoders represent conformational variation, we measure the proportion of embedding variation attributable to differences between conformers of the same peptide. For each encoder and task, we compute $\eta^2 = SS_{\text{within}}/(SS_{\text{within}} + SS_{\text{between}})$, where SS_{within} is the variation across conformers of the same peptide and SS_{between} is the variation between different peptides. Figure 4 shows the relationship between η^2 and the paired improvement

of EASE over the matched reference conformer. GearNet-Edge exhibits the highest mean η^2 (0.562), followed by ESM-3 (0.330) and ESM-GearNet (0.164), and this ordering matches that of their average EASE gains. This suggests that encoders with greater conformer variation in their representations may benefit more from aggregating multiple conformations. One possible explanation for this trend is how PLM information enters each encoder. GearNet-Edge relies solely on geometric message passing without pretrained PLM components, which allows structural differences to directly influence its representations. ESM-3 jointly processes sequence and structure tokens through transformer layers, potentially reducing the relative contribution of conformer structure information. ESM-GearNet incorporates PLM information in two places: ESM-2 embeddings are both passed as node features for its graphs and concatenated directly to the final representation. This high PLM involvement may further reduce the influence of conformational differences, which is consistent with its lowest η^2 and smallest ensemble gains.

6 Conclusion

We introduced EASE, a framework for extracting predictive representations from peptide conformational ensembles. EASE improves over single-structure baselines on all six tasks on GearNet-Edge and five on ESM-3. Parameter-free mean pooling is competitive with learned fusion, and larger conformer banks provide diminishing gains through $K = 20$. Six alternative approaches that first reduce the ensemble to summary statistics yield smaller improvements over their respective baselines, supporting the value of encoding whole conformers directly instead of compressing ensemble information into dynamics descriptors. Comparisons with sequence-only PLMs further show that EASE provides predictive information beyond sequence representations. Across three encoders, greater conformer embedding variation is associated with larger ensemble gains, suggesting that how PLM information is integrated into structural representations influences the benefit of ensemble conditioning.

7 Limitations

There are several limitations to our current work. First, this study focuses on improving peptide structural representation learning but does not claim that structure is the strongest modality for peptide property prediction. Sequence-only baselines remain stronger on half of the six downstream tasks evaluated on GearNet-Edge (§5.5), and EASE only narrows this gap. Second, our experiments are restricted to linear peptides composed of canonical amino acids. This restriction excludes cyclic and non-canonical peptides, which constitute an important part of therapeutic peptide space but are not supported by BioEmu. Further, the ensembles generated by BioEmu approximate equilibrium conformational distributions instead of providing calibrated Boltzmann samples or molecular-dynamics trajectories. Our results therefore concern the use of multiple plausible conformations, without modeling how frequently each state occurs or how the peptide transitions between them. Finally, our work has potential ethical and environmental implications. Generating and encoding multiple conformers increases computational cost and associated energy consumption compared with single-structure approaches. Additionally, the predicted properties should not be treated as substitutes for experimental validation, particularly in safety-critical applications.

References

Jakub Adamczyk, Piotr Ludynia, and Wojciech Czech. 2026. Molecular fingerprints are strong models for peptide function prediction. *Bioinformatics*, 42(5).

Mariana D. C. Aguilera-Puga and Fabien Plisson. 2024. Structure-aware machine learning strategies for antimicrobial peptide discovery. *Scientific Reports*, 14(1).

Simon Axelrod and Rafael Gómez-Bombarelli. 2023. Molecular machine learning with conformer ensembles. *Machine Learning: Science and Technology*, 4(3):035025.

Leo Tianlai Chen, Zachary Quinn, Madeleine Dumas, Christina Peng, Lauren Hong, Moises Lopez-Gonzalez, Alexander Mestre, Rio Watson, Sophia Vincoff, Lin Zhao, Jianli Wu, Audrey Stavrand, Mayumi Schaepers-Cheu, Tian Zi Wang, Divya Srijay, Connor Monticello, Pranay Vure, Rishab Pulugurta, Sarah Pertsemlidis, and 7 others. 2026. Target sequence-conditioned design of peptide binders using masked language modeling. *Nature Biotechnology*, 44(6):1002–1010.

Seyone Chithrananda, Gabriel Grand, and Bharath Ramsundar. 2020. ChemBERTa: Large-scale self-supervised pretraining for molecular property prediction. *arXiv preprint arXiv:2010.09885*.

Anja Conev, Mauricio Menegatti Rigo, Didier Devaurs, André Faustino Fonseca, Hussain Kalavadwala, Martiela Vaz de Freitas, Cecilia Clementi, Geancarlo Zanatta, Dinler Amaral Antunes, and Lydia E. Kavraki. 2023. Engens: a computational framework for generation and analysis of representative protein conformational ensembles. *Briefings in Bioinformatics*, 24(4):bbad242.

Xavier Daura, Karl Gademann, Bernhard Jaun, Dieter Seebach, Wilfred F. van Gunsteren, and Alan E. Mark. 1999. Peptide folding: When simulation meets experiment. *Angewandte Chemie International Edition*, 38(1-2):236–240.

Peter Eastman, Jason Swails, John D. Chodera, Robert T. McGibbon, Yutong Zhao, Kyle A. Beauchamp, Lee-Ping Wang, Andrew C. Simmonett, Matthew P. Harrigan, Chaya D. Stern, Rafal P. Wiewiora, Bernard R. Brooks, and Vijay S. Pande. 2017. OpenMM 7: Rapid development of high performance algorithms for molecular dynamics. *PLOS Computational Biology*, 13(7):e1005659.

Aaron Feller, Kris Deibler, and Maxim Secor. 2026. Graph learning on ensembles of cyclic peptides: An investigation of molecular ensemble modeling. *Preprint*, arXiv:2607.21561.

Aaron L. Feller and Claus O. Wilke. 2025. Peptide-aware chemical language model successfully predicts membrane diffusion of cyclic peptides. *Journal of Chemical Information and Modeling*, 65(2):571–579.

Raúl Fernández-Díaz, Rodrigo Cossio-Pérez, Clement Agoni, Hoang Thanh Lam, Vanessa Lopez, and Denis C Shields. 2024. Autopeptideml: a study on how to build more trustworthy peptide bioactivity predictors. *Bioinformatics*, 40(9).

Chakradhar Guntuboina, Adrita Das, Parisa Mollaei, Seongwon Kim, and Amir Barati Farimani. 2023. Peptidebert: A language model based on transformers for peptide property prediction. *The Journal of Physical Chemistry Letters*, 14(46):10427–10434.

Thomas Hayes, Roshan Rao, Halil Akin, et al. 2025. Simulating 500 million years of evolution with a language model. *Science*, 387(6736):850–858.

Antoine Henninot, James C. Collins, and John M. Nuss. 2018. The current state of peptide drug discovery: Back to the future? *Journal of Medicinal Chemistry*, 61(4):1382–1414.

Scott A. Hollingsworth and Ron O. Dror. 2018. Molecular dynamics simulation for all. *Neuron*, 99(6):1129–1143.

Viktor Hornak, Robert Abel, Asim Okur, Bentley Strockbine, Adrian Roitberg, and Carlos Simmerling. 2006. Comparison of multiple Amber force fields and development of improved protein backbone parameters. *Proteins: Structure, Function, and Bioinformatics*, 65(3):712–725.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. Lora: Low-rank adaptation of large language models. *Preprint*, arXiv:2106.09685.

Haomeng Hu, Chengyun Zhang, Zhenyu Xu, Jingjing Guo, An Su, and Chengxi Li. 2025. PepMSND: integrating multi-level feature engineering and comprehensive databases to enhance in vivo/in vitro peptide blood stability prediction. *Digital Discovery*, 4(9):2478–2490.

Maximilian Ilse, Jakub M. Tomczak, and Max Welling. 2018. Attention-based deep multiple instance learning. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*.

Bowen Jing, Bonnie Berger, and Tommi Jaakkola. 2024. Alphafold meets flow matching for generating protein ensembles. *Preprint*, arXiv:2402.04845.

Dan Kalifa, Eric Horvitz, and Kira Radinsky. 2026. Learning protein representations with conformational dynamics. *Bioinformatics*, 42(5).

Martin Karplus and J. Andrew McCammon. 2002. Molecular dynamics simulations of biomolecules. *Nature Structural Biology*, 9(9):646–652.

Seungeon Lee, Takuto Koyama, Itsuki Maeda, Shigeyuki Matsumoto, and Yasushi Okuno. 2025. Helm-bert: A transformer for medium-sized peptide property prediction. *Preprint*, arXiv:2512.23175.

Sarah Lewis, Tim Hempel, José Jiménez-Luna, Michael Gastegger, Yu Xie, Andrew Y. K. Foong, Victor García Satorras, Osama Abdin, Bastiaan S. Veeling, Iryna Zaporozhets, Yaoyi Chen, Soojung Yang, Adam E. Foster, Arne Schneuing, Jigyasa Nigam, Federico Barbero, Vincent Stimper, Andrew Campbell, Jason Yim, and 9 others. 2025. Scalable emulation of protein equilibrium ensembles with generative deep learning. *Science*, 389(6761).

Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaniv Shmueli, Allan dos Santos Costa, Maryam Fazel-Zarandi, Tom Sercu, Salvatore Candido, and Alexander Rives. 2023. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130.

Enyan Liu, Yueming Hu, Liya Liu, Yifan Chen, Shilong Zhang, Sida Li, Haoyu Chao, Luyao Xie, Yi Shen, Liangwei Wu, Julio Raúl Fernández Massó, and Ming Chen. 2026. Pepanno: A structure-aware deep learning framework for bioactive peptide prediction, structural visualization, and physicochemical profiling. *PLOS Computational Biology*, 22(6):e1014369.

Jiarui Lu, Bozitao Zhong, Zuobai Zhang, and Jian Tang. 2023. Str2str: A score-based framework for zero-shot protein conformation sampling. *Preprint*, arXiv:2306.03117.

Matthew McPartlon and Jinbo Xu. 2023. An end-to-end deep learning method for protein side-chain packing and inverse folding. *Proceedings of the National Academy of Sciences*, 120(23):e2216438120.

Markus Muttenthaler, Glenn F. King, David J. Adams, and Paul F. Alewood. 2021. Trends in peptide drug discovery. *Nature Reviews Drug Discovery*, 20(4):309–325.

Duy M. H. Nguyen, Trung Q. Nguyen, Ha T. H. Le, Mai Thanh Nhat Truong, TrungTin Nguyen, Nhat Ho, Khoa D. Doan, Duy Duong-Tran, Li Shen, Daniel Sonntag, James Zou, Mathias Niepert, Hyojin Kim, and Jonathan E. Allen. 2026. Facet: A fragment-aware conformer ensemble transformer. In *International Conference on Learning Representations (ICLR)*.

Alexey Onufriev, Donald Bashford, and David A. Case. 2004. Exploring protein native states and large-scale conformational changes with a modified generalized Born model. *Proteins: Structure, Function, and Bioinformatics*, 55(2):383–394.

Ethan Perez, Florian Strub, Harm de Vries, Vincent Dumoulin, and Aaron Courville. 2018. Film: Visual reasoning with a general conditioning layer. In *Proceedings of the 32nd AAAI Conference on Artificial Intelligence (AAAI)*.

Zachary A Rollins, Talal Widatalla, Andrew Waight, Alan C Cheng, and Essam Metwally. 2024. Ablef: antibody language ensemble fusion for thermodynamically empowered property predictions. *Bioinformatics*, 40(5).

Kaiwen Shi and Carlos Oliver. 2026. Ensembits: an alphabet of protein conformational ensembles. *Preprint*, arXiv:2605.13789.

Xiangguo Eric Shi, Thomas E. Wales, Carl Elkin, Noriyuki Kawahata, John R. Engen, and D. Allen Annis. 2013. Hydrogen exchange–mass spectrometry measures stapled peptide conformational dynamics and predicts pharmacokinetic properties. *Analytical Chemistry*, 85(23):11185–11188.

Martin Steinegger and Johannes Söding. 2017. Mmseqs2 enables sensitive protein sequence searching for the analysis of massive data sets. *Nature Biotechnology*, 35(11):1026–1028.

Fang Wu, Shuting Jin, Yinghui Jiang, Xurui Jin, Bowen Tang, Zhangming Niu, Xiangrong Liu, Qiang Zhang, Xiangxiang Zeng, and Stan Z. Li. 2022. Pre-training of equivariant graph matching networks with conformation flexibility for drug binding. *Advanced Science*, 9(33).

Wenping Yu, Zhewen Li, Wei Xu, Yu Zhao, and Nan Sun. 2026. Sa-mtp: a structure-aware framework for multifunctional therapeutic peptide annotation. *Briefings in Bioinformatics*, 27(4).

Manzil Zaheer, Satwik Kottur, Siamak Ravanbakhsh, Barnabas Poczos, Ruslan Salakhutdinov, and Alexander J. Smola. 2017. Deep sets. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Buke Zhang, Junjie Zhu, Hai-Feng Chen, and Lu-Ning Liu. 2026a. Alphafold-based peptide structure prediction: Opportunities, limitations, and future directions. *Biotechnology Advances*, 92:108977.

Jiahui Zhang, Rouyi Wang, Kuangqi Zhou, Tianshu Xiao, Lingyan Zhu, Yaosen Min, and Yang Wang. 2026b. PepBenchmark: A standardized benchmark for peptide machine learning. In *International Conference on Learning Representations*.

Ruochi Zhang, Haoran Wu, Chang Liu, Qian Yang, Yuting Xiu, Kewei Li, Ningning Chen, Yu Wang, Yan Wang, Xin Gao, and Fengfeng Zhou. 2025. Pepland: a large-scale pre-trained peptide representation model for a comprehensive landscape of both canonical and non-canonical amino acids. *Briefings in Bioinformatics*, 26(4):bbaf367.

Fei Zhao, Junhui Qiu, Dongyou Xiang, Pengrui Jiao, Yu Cao, Qingrui Xu, Dairong Qiao, Hui Xu, and Yi Cao. 2024. deepampnet: a novel antimicrobial peptide predictor employing alphafold2 predicted structures and a bi-directional long short-term memory protein language model. *PeerJ*, 12:e17729.

Yanqiao Zhu, Jeehyun Hwang, Keir Adams, Zhen Liu, Bozhao Nan, Bo Stenfors, Yuanqi Du, Jatin Chauhan, Olaf Wiest, Olexandr Isayev, Connor W. Coley, Yizhou Sun, and Wei Wang. 2023. Learning over molecular conformer ensembles: Datasets and benchmarks. *Preprint*, arXiv:2310.00115.

Table 5: Parameters added by each fusion operator.

Encoder	Fusion	Internal width	Params added
GearNet-Edge ($d = 512$)	mean	512	0
	gated	512	513
	gated attn	512	131,457
	attention	64	99,584
ESM-3 ($d = 1536$)	mean	1536	0
	gated	1536	1,537
	gated attn	1536	393,601
	attention	128	132,480

A Fusion operators details

Table 5 reports the number of parameters introduced by each fusion operator for both encoders.

B Alternative methods details

B.1 Alternative uses of conformational ensembles

EASE preserves each sampled conformer as a structural input and aggregates their encoded representations. To test whether the same ensemble information can be used more effectively elsewhere in the model, we construct a set of six alternative mechanisms that expose the encoder to summary statistics computed from the ensemble. All alternatives draw on the same BioEmu-generated conformational ensembles. Since multiple alternative mechanisms apply to geometric graph features, we evaluate these mechanisms on GearNet-Edge only.

Ensemble summary statistics. The six alternatives first compact the sampled conformers into residue-level, pair-level, or peptide-level statistics. At the residue level, we compute the RMS fluctuation of each Cα position around its ensemble mean, the mean and standard deviation of per-residue solvent accessibility, and the circular variance of the backbone ϕ and ψ dihedral angles. At the pair level, we compute the mean and standard deviation of Cα–Cα distance and the contact frequency, defined as the fraction of conformers in which the pair is separated by less than 10 Å. Some of these quantities are injected in both raw and within-peptide z-scored form, as the raw values preserve magnitude across peptides while the z-scored values retain only the within-peptide profile. A summary of all statistics is listed in Table 6.

(1) Adding graph features and edges. We first inject ensemble statistics directly into GearNet-Edge. Residue-level statistics are concatenated to the node features, while pairwise statistics are concatenated to edge features. The pairwise variant additionally introduces an ensemble-contact edge type, defined as residue pairs whose contact frequency is at least 0.3.

(2) Distributional edge geometry. The pairwise features above summarize the distance between each residue pair using its mean and standard deviation. These moments cannot distinguish distributions with similar mean and variance but different shapes. We therefore replace the distance moments with a 16-bin soft histogram of the full distance distribution, obtained by averaging Gaussian radial basis functions centered between 2 and 32 Å across conformers. This representation preserves absolute distance scale while retaining more information about the distribution of distances across the ensemble. Graph topology is unchanged and only the features carried by each edge are replaced.

(3) Stochastic ensemble topology. The ensemble-contact edge type above converts contact frequency into a binary edge type using a fixed threshold, so a residue pair either is or is not in contact. In this mechanism, we instead use contact frequency directly to define a distribution over graph topology. During training, each non-sequential edge is retained independently with probability equal to its contact frequency, while sequential backbone edges are always retained. For instance, a pair in contact in half the conformers contributes to half the training steps. At evaluation, we use the expectation by weighting each edge's message by its contact frequency. This introduces ensemble-dependent topology without additional trainable parameters or a contact-frequency threshold.

(4) Layer-wise feature modulation. In the mechanisms above, the statistics enter the encoder only once, at the input layer. Thus, these statistics survive only as what the hidden state retains through six rounds of relational convolution, normalization and residual addition. We instead make them available throughout the encoder using feature-wise linear modulation (FiLM) (Perez et al., 2018). At each layer, a linear map transforms the residue-level statistics into a feature-wise scale and shift, and updates the hidden state as $h_l \leftarrow \gamma_l \odot h_l + \beta_l$. Each layer receives the original ensemble statistics rather than the output of the preceding modulation. The maps are initialized so that $\gamma_l = 1$ and $\beta_l = 0$, preserving the pretrained function at initialization. Modulation is applied independently at each residue, and the hidden dimensionality is unchanged.

Table 6: Ensemble statistics available to the alternative mechanisms. Residue and pairwise quantities enter the encoder as node features, edge features or an added relation; the peptide-level descriptor enters at the readout. The form column indicates whether a quantity is injected in raw form, within-peptide z-scored form, or both.

Level	Quantity	Form	Definition
residue	fluctuation	both	RMS deviation of the Cα position around the ensemble mean
residue	accessibility, mean and s.d.	both	per-residue solvent-accessible surface area across conformers
residue	dihedral variability	raw	circular variance of backbone ϕ and ψ across conformers
pair	distance, mean and s.d.	both	Cα–Cα distance across conformers
pair	contact frequency	raw	fraction of conformers with Cα–Cα separation below 10 Å
pair	ensemble-contact relation	—	additional edge type for pairs with contact frequency ≥ 0.3
peptide	pooled descriptor	both	22-dimensional summary of fluctuation, accessibility, distance, radius of gyration, dihedral variability, contact statistics, and peptide length

(5) Peptide-level descriptor at the readout. The mechanisms above provide per-residue or pairwise quantities, so no quantity at the entire peptide scope is used in the model. In this mechanism, we pool the ensemble statistics into a 22-dimensional peptide descriptor and concatenate it to the pooled peptide representation at the readout. The descriptor summarizes fluctuation, solvent accessibility, pairwise distance, radius of gyration, dihedral variability, contact frequency, long-range contacts, contact order, transient contacts, and peptide length. It includes both the raw and within-peptide normalized summaries where applicable. Descriptor dimensions are standardized using statistics from the training fold. The encoder itself is unchanged.

(6) Dynamics as auxiliary supervision. We next use ensemble statistics as training targets rather than model inputs. Auxiliary heads on the final residue representations predict three per-residue quantities: RMS fluctuation and the mean and standard deviation of solvent accessibility. A second variant additionally predicts pairwise contact probability from a down-projected inner product of residue representations. The training objective is $\mathcal{L} = \mathcal{L}\mathrm{task} + \lambda\mathcal{L}\mathrm{aux}$, with a second weighted term for the pairwise head in the latter variant. The auxiliary heads are discarded at inference, so the ensemble is used only during training. Rather than tuning λ as a task-specific hyperparameter, we calibrate it so that the task and auxiliary objectives contribute gradients of equal norm to the shared trunk, measured over training batches.

Baseline for alternative methods. The alternative mechanisms are compared against the unmodified GearNet-Edge encoder given a single structure, chosen as the conformer whose Cα coordinates lie closest to the ensemble mean. We note that this is a different structure from the reference conformer of §4, which is drawn from the $K = 10$ conformer bank. The reference conformer baseline remains specific to EASE, where the contrast has to isolate fusing the bank from selecting a member of it; here the encoder input is a single conformer in every arm, and what varies is the additional ensemble information supplied alongside it.

Evaluation of alternative methods. Each of these alternative methods is evaluated at five folds and four training seeds on all six tasks.

C Training protocol details

Configurations and hyperparameters. Tables 7 and 8 show the specific training configurations, hyperparameter search grids, and the selected hyperparameters.

Computational resources. Every training run used a single GPU, and the GPU is randomly scheduled across NVIDIA RTX 6000 Ada (48 GB), A100 (80 GB), and RTX PRO 6000 Blackwell (96 GB). The total computational budget was approximately 2,200 GPU-hours and 2,000 CPU-hours. We spent approximately 100 GPU-hours for BioEmu conformational ensemble generation (around 31.6 k peptides and 3.4 M conformers in total), including AlphaFold2 embedding and diffusion sampling. An additional 2,000 CPU-hours was required for the all-atom side-chain reconstruction and energy minimization. On the model training and evaluation part, we spent approximately 2,070 GPU-hours on the reported experiments involving EASE and its ablations, along with all baselines.

Models and trainable parameters. ESM-3 (esm3_sm_open_v1) contains 1.40 B parameters and GearNet-Edge contains 20.2 M. The ESM-GearNet includes a frozen ESM-2-650M (649 M) with a GearNet trunk (15.7 M). In every experiment, only the LoRA adapters, the fusion layer, and the prediction MLP heads are trained, which is around

Table 7: Training configuration for each encoder.

Setting	GearNet-Edge	ESM-3
Optimizer	AdamW	AdamW
Trunk learning rate	3×10^{-4}	tuned, $\{10^{-4}, 10^{-5}\}$
Head learning rate	3×10^{-4}	10^{-3}
Weight decay	0.1	0.01
Batch size	16	tuned, $\{16, 32\}$
LoRA rank	tuned, $\{4, 8\}$	8
Head dropout	tuned, $\{0.1, 0.3, 0.5\}$	0.1
Max epochs	100	20
Early-stopping patience	15	6
PPI batch size	64	tuned, $\{16, 32\}$
PPI max epochs	120	50
PPI patience	15	12
Tuning granularity	one config per task, shared by all arms	one config per (task, arm)
Seeds / fold	4	3
Head hidden width	64	128
Trainable parameters	0.46–0.63M	3.77–4.27M

Table 8: Selected configurations. GearNet-Edge, the ESM-2-35M sequence-only baseline, and the ChemBERTa chemical baseline each tune one configuration per task over LoRA rank $\{4, 8\}$ and head dropout $\{0.1, 0.3, 0.5\}$. ESM-3 tunes one configuration per (task, arm) over trunk learning rate $\{10^{-4}, 10^{-5}\}$ and batch size $\{16, 32\}$.

	GearNet-Edge		ESM-2 (35M)		ChemBERTa		ESM-3 (trunk learning rate / batch size)					
task	LoRA rank	head dropout	LoRA rank	head dropout	LoRA rank	head dropout	sequence	reference conformer	mean	gated	gated attn.	attention
toxicity	8	0.5	8	0.1	8	0.5	10^{-5}/16	10^{-4}/32	10^{-4}/16	10^{-4}/32	10^{-4}/16	10^{-4}/32
hemolytic	8	0.1	8	0.1	8	0.3	10^{-4}/16	10^{-4}/16	10^{-4}/16	10^{-4}/16	10^{-4}/16	10^{-4}/16
hemolytic HC50	8	0.1	8	0.5	4	0.3	10^{-4}/16	10^{-4}/16	10^{-4}/32	10^{-4}/32	10^{-4}/16	10^{-4}/16
E. coli MIC	8	0.5	8	0.5	4	0.5	10^{-4}/16	10^{-5}/16	10^{-4}/16	10^{-4}/32	10^{-4}/16	10^{-4}/32
blood stability	8	0.5	4	0.1	8	0.5	10^{-5}/16	10^{-5}/16	10^{-5}/32	10^{-5}/16	10^{-4}/16	10^{-5}/16
interaction	8	0.3	4	0.1	4	0.1	10^{-5}/32	10^{-5}/32	10^{-5}/32	10^{-5}/32	10^{-5}/32	10^{-5}/32

3.8 M trainable parameters for ESM-3, 0.46 M for GearNet-Edge, and 0.88 M for ESM-GearNet.

D Results: additional comparison tables and figures

Table 9 shows the absolute test performance of EASE with each fusion operator.

Tables 10 and 11 displays the absolute test performance and paired difference per fold against the reference conformer of the EASE method as the ensemble size K increases through 5, 10, 20.

The following table 12 shows the paired performance improvements over same folds of each alternative method against its own single-structure baseline, as defined in §B.1.

Figure 4 shows how conformer embedding variations in different encoders influence EASE improvements over the reference conformer. Tables 13 and 14 present the absolute and paired performance of EASE (mean $K = 10$) and reference conformer on ESM-GearNet, compared to ESM-3 and GearNet-Edge.

Table 9: Test performance of EASE with different fusion operators. Mean ± standard deviation over folds with training seeds averaged within each fold. AUROC is reported for classification and RMSE for regression. Average rank is obtained by ranking the five arms on each task and averaging their ranks across the six tasks. Bold indicates the best mean performance for each task within each encoder.

Encoder	Fusion	Toxicity ↑	Hemolytic ↑	HC50 ↓	*E. coli* MIC ↓	Stability ↑	Interaction ↑	Avg. rank ↓
GearNet-Edge	Ref. conformer	0.6893 ± 0.1008	0.8617 ± 0.0221	0.6460 ± 0.0191	0.6847 ± 0.0202	0.6737 ± 0.0458	0.8059 ± 0.0159	5.00
	Mean	**0.7231 ± 0.0938**	0.8733 ± 0.0225	0.6139 ± 0.0164	0.6651 ± 0.0275	**0.7141 ± 0.0525**	0.8327 ± 0.0101	2.50
	Gated	0.7224 ± 0.0952	0.8750 ± 0.0218	0.6115 ± 0.0182	**0.6639 ± 0.0313**	0.7134 ± 0.0413	**0.8353 ± 0.0097**	**2.00**
	Attention	0.7210 ± 0.0939	0.8710 ± 0.0208	**0.6107 ± 0.0217**	0.6755 ± 0.0232	0.6846 ± 0.0723	0.8210 ± 0.0110	3.50
	Gated attn	0.7226 ± 0.0956	**0.8753 ± 0.0217**	0.6110 ± 0.0191	0.6647 ± 0.0313	0.7055 ± 0.0495	0.8331 ± 0.0088	**2.00**
ESM-3	Ref. conformer	0.7792 ± 0.0534	0.9013 ± 0.0209	0.5793 ± 0.0230	0.6842 ± 0.0225	0.7141 ± 0.0266	0.8329 ± 0.0118	4.33
	Mean	**0.8058 ± 0.0387**	**0.9071 ± 0.0161**	0.5554 ± 0.0334	0.6515 ± 0.0183	0.7032 ± 0.0314	**0.8520 ± 0.0100**	**2.33**
	Gated	0.7925 ± 0.0515	0.9059 ± 0.0192	0.5642 ± 0.0545	**0.6470 ± 0.0189**	**0.7216 ± 0.0286**	0.8391 ± 0.0098	**2.33**
	Attention	0.7996 ± 0.0258	0.9070 ± 0.0171	0.5511 ± 0.0423	0.6511 ± 0.0124	0.7098 ± 0.0133	0.8166 ± 0.0090	2.83
	Gated attn	0.7908 ± 0.0514	0.9049 ± 0.0196	**0.5462 ± 0.0318**	0.6488 ± 0.0164	0.6924 ± 0.0163	0.8355 ± 0.0079	3.17

Table 10: Ensemble sizes K, absolute test performance. Mean ± standard deviation over folds with training seeds averaged within each fold. Mean fusion is used for every ensemble size. AUROC reported for classification and RMSE for regression.

encoder	arm	toxicity ↑	hemolytic ↑	HC50 ↓	*E. coli* MIC ↓	stability ↑	interaction ↑
GearNet-Edge	reference conformer	0.6893 ± 0.1008	0.8617 ± 0.0221	0.6460 ± 0.0191	0.6847 ± 0.0202	0.6737 ± 0.0458	0.8059 ± 0.0159
GearNet-Edge	K=5	0.7231 ± 0.0953	**0.8747 ± 0.0234**	**0.6051 ± 0.0231**	0.6692 ± 0.0249	0.6770 ± 0.0496	0.8271 ± 0.0108
GearNet-Edge	K=10	0.7231 ± 0.0938	0.8733 ± 0.0225	0.6139 ± 0.0164	0.6651 ± 0.0275	**0.7141 ± 0.0525**	0.8327 ± 0.0101
GearNet-Edge	K=20	**0.7384 ± 0.0876**	0.8731 ± 0.0197	0.6101 ± 0.0256	**0.6636 ± 0.0248**	0.7126 ± 0.0604	**0.8376 ± 0.0137**
ESM-3	reference conformer	0.7792 ± 0.0534	0.9013 ± 0.0209	0.5793 ± 0.0230	0.6842 ± 0.0225	**0.7141 ± 0.0266**	0.8329 ± 0.0118
ESM-3	K=5	0.7944 ± 0.0498	0.9058 ± 0.0161	0.5608 ± 0.0400	0.6575 ± 0.0217	0.7022 ± 0.0525	0.8477 ± 0.0116
ESM-3	K=10	0.8058 ± 0.0387	0.9071 ± 0.0161	0.5554 ± 0.0334	0.6515 ± 0.0183	0.7032 ± 0.0314	0.8520 ± 0.0100
ESM-3	K=20	**0.8102 ± 0.0318**	**0.9073 ± 0.0193**	**0.5466 ± 0.0458**	**0.6496 ± 0.0186**	0.6982 ± 0.0262	**0.8523 ± 0.0102**

Table 11: Ensemble sizes K, paired differences over same folds against the fixed reference-conformer baseline (folds improved out of five). Positive values indicate improvement.

encoder	K	toxicity	hemolytic	HC50	*E. coli* MIC	stability	interaction	mean	tasks ↑
GearNet-Edge	5	+0.0339 (5/5)	+0.0130 (5/5)	+0.0409 (5/5)	+0.0154 (4/5)	+0.0033 (2/5)	+0.0212 (5/5)	+0.0213	6/6
GearNet-Edge	10	+0.0339 (5/5)	+0.0116 (5/5)	+0.0321 (5/5)	+0.0196 (5/5)	+0.0404 (3/5)	+0.0267 (5/5)	+0.0274	6/6
GearNet-Edge	20	+0.0491 (5/5)	+0.0114 (5/5)	+0.0359 (5/5)	+0.0210 (5/5)	+0.0389 (3/5)	+0.0316 (5/5)	**+0.0313**	6/6
ESM-3	5	+0.0152 (5/5)	+0.0045 (4/5)	+0.0184 (4/5)	+0.0267 (5/5)	−0.0118 (2/5)	+0.0148 (5/5)	+0.0113	5/6
ESM-3	10	+0.0266 (5/5)	+0.0058 (4/5)	+0.0238 (4/5)	+0.0327 (5/5)	−0.0108 (0/5)	+0.0191 (5/5)	+0.0162	5/6
ESM-3	20	+0.0310 (5/5)	+0.0060 (4/5)	+0.0327 (4/5)	+0.0346 (5/5)	−0.0159 (0/5)	+0.0194 (5/5)	**+0.0180**	5/6

Table 12: Alternative uses of the conformational ensemble. GearNet-Edge, mean paired difference against each mechanism's own single-structure baseline over the same folds (folds improved out of five). The alternative mechanisms are paired against the base encoder with the mean structure and EASE is paired against the reference conformer. Positive is better, with mechanisms numbered as in §B.1. †This row alone is an increment over the row above rather than a difference against a baseline.

#	Mechanism	Toxicity	Hemolytic	HC50	*E. coli* MIC	Stability	PepPI	Mean	Tasks ↑
1	Per-residue statistics	+0.0083 (4/5)	+0.0054 (4/5)	+0.0116 (3/5)	−0.0013 (3/5)	+0.0109 (5/5)	+0.0056 (4/5)	+0.0067	5/6
1	Pairwise statistics	+0.0127 (4/5)	+0.0068 (3/5)	+0.0044 (3/5)	−0.0047 (2/5)	+0.0155 (5/5)	+0.0037 (4/5)	+0.0064	5/6
1	↪ Extra edge type†	+0.0044 (2/5)	+0.0014 (3/5)	−0.0071 (2/5)	−0.0034 (1/5)	+0.0046 (2/5)	−0.0019 (2/5)	−0.0003	3/6
2	Edge distance histogram	+0.0095 (3/5)	+0.0059 (4/5)	+0.0029 (3/5)	−0.0039 (2/5)	+0.0159 (4/5)	+0.0027 (4/5)	+0.0055	5/6
3	Contact-frequency topology	+0.0055 (4/5)	+0.0032 (2/5)	+0.0260 (5/5)	−0.0025 (1/5)	**+0.0499 (4/5)**	+0.0102 (5/5)	+0.0154	5/6
4	Modulation, node stats	+0.0257 (5/5)	+0.0077 (4/5)	+0.0184 (4/5)	−0.0024 (3/5)	+0.0206 (4/5)	+0.0081 (5/5)	+0.0130	5/6
4	Modulation, node + edge stats	+0.0265 (5/5)	+0.0043 (3/5)	+0.0188 (4/5)	−0.0044 (1/5)	+0.0169 (3/5)	+0.0081 (5/5)	+0.0117	5/6
5	Pooled descriptor	+0.0079 (4/5)	+0.0011 (3/5)	+0.0071 (3/5)	+0.0041 (4/5)	+0.0090 (3/5)	+0.0050 (3/5)	+0.0057	6/6
6	Auxiliary supervision, node	+0.0019 (2/5)	+0.0008 (3/5)	−0.0049 (3/5)	+0.0010 (3/5)	+0.0066 (3/5)	+0.0163 (5/5)	+0.0036	5/6
6	Auxiliary supervision, node + edge	+0.0086 (3/5)	−0.0023 (3/5)	+0.0007 (3/5)	−0.0035 (2/5)	−0.0111 (2/5)	+0.0246 (5/5)	+0.0028	3/6
—	**EASE, mean**	**+0.0339 (5/5)**	**+0.0116 (5/5)**	**+0.0321 (5/5)**	**+0.0196 (5/5)**	+0.0404 (3/5)	**+0.0267 (5/5)**	**+0.0274**	6/6

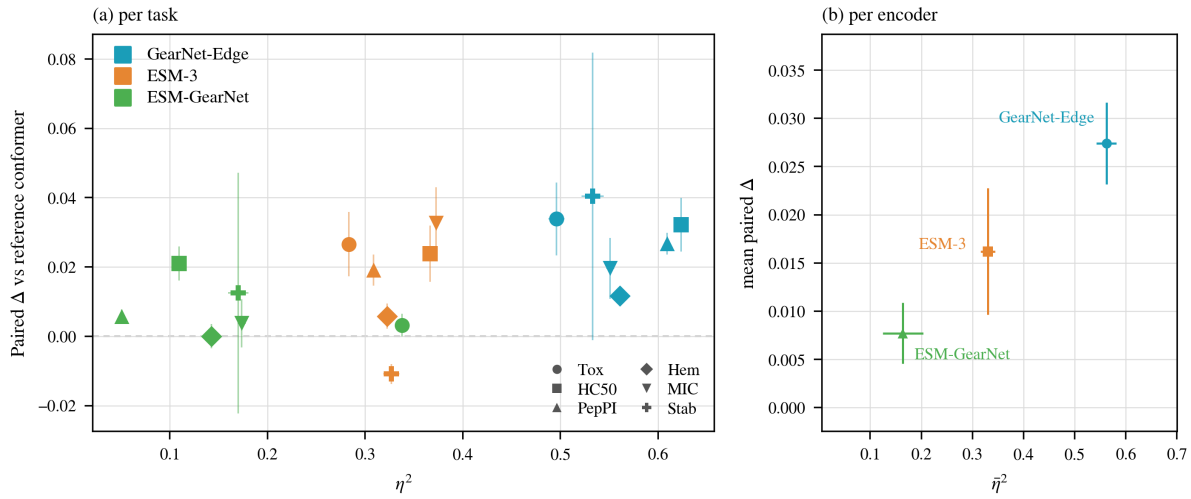

Figure 4: Relationship between conformer embedding variation and ensemble gains. **(a)** Task-level comparison across three pretrained encoders and six peptide prediction tasks. The vertical axis shows the paired performance improvement of EASE (mean fusion, $K = 10$) over the reference conformer baseline. **(b)** Encoder-level averages across the six tasks. Error bars indicate ± 1 standard error (bootstrap over peptides for x-axis in (a), across folds for y-axis in (a), and across tasks in (b)).

Table 13: Absolute test performance. Mean $\pm$ standard deviation over same folds. AUROC for classification, RMSE for regression. ESM-GearNet uses a frozen ESM-2-650M with LoRA on the GearNet trunk.

Encoder	Arm	Toxicity ↑	Hemolytic ↑	HC50 ↓	*E. coli* MIC ↓	Stability ↑	PepPI ↑
ESM-GearNet	Reference conformer (K=1)	0.7832 ± 0.0700	0.8895 ± 0.0252	0.5635 ± 0.0358	0.7059 ± 0.0168	0.6895 ± 0.0447	0.8352 ± 0.0097
	EASE, mean K=10	0.7864 ± 0.0722	0.8896 ± 0.0285	0.5425 ± 0.0418	0.7022 ± 0.0157	0.7020 ± 0.0474	0.8410 ± 0.0057
GearNet-Edge	Reference conformer (K=1)	0.6893 ± 0.1008	0.8617 ± 0.0221	0.6460 ± 0.0191	0.6847 ± 0.0202	0.6737 ± 0.0458	0.8059 ± 0.0159
	EASE, mean K=10	0.7231 ± 0.0938	0.8733 ± 0.0225	0.6139 ± 0.0164	0.6651 ± 0.0275	0.7141 ± 0.0525	0.8327 ± 0.0101
ESM-3	Reference conformer (K=1)	0.7792 ± 0.0534	0.9013 ± 0.0209	0.5793 ± 0.0230	0.6842 ± 0.0225	0.7141 ± 0.0266	0.8329 ± 0.0118
	EASE, mean K=10	0.8058 ± 0.0387	0.9071 ± 0.0161	0.5554 ± 0.0334	0.6515 ± 0.0183	0.7032 ± 0.0314	0.8520 ± 0.0100

Table 14: Paired difference of EASE with mean K=10 against the reference conformer over the same folds (folds improved). Training seeds are averaged within a fold before pairing, and the sign is set so positive always indicates improvement for both AUROC and RMSE. The Mean column averages the six tasks.

Encoder	Toxicity	Hemolytic	HC50	*E. coli* MIC	Stability	PepPI	Mean	
ESM-GearNet	+0.0032 (3/5)	+0.0000 (3/5)	+0.0210 (5/5)	+0.0037 (3/5)	+0.0125 (2/5)	+0.0058 (4/5)	+0.0077	
GearNet-Edge	+0.0339 (5/5)	+0.0116 (5/5)	+0.0321 (5/5)	+0.0196 (5/5)	+0.0404 (3/5)	+0.0267 (5/5)	+0.0274	
ESM-3		+0.0266 (5/5)	+0.0058 (4/5)	+0.0238 (4/5)	+0.0327 (5/5)	−0.0108 (0/5)	+0.0191 (5/5)	+0.0162