This work has been submitted to the IEEE for possible publication. Copyright may be transferred without notice, after which this version may no longer be accessible.

Preprint submitted to IEEE Access, October 2026

BRIDGE: An Acceptance Gate to Block False Improvements in Model Updates, with a Retrospective Clinical Evaluation

DWI SUNARYONO[1]

[1]Department of Informatics, Faculty of Intelligent Electrical and Informatics Technology, Institut Teknologi Sepuluh Nopember, Surabaya 60111, Indonesia

Corresponding author: Dwi Sunaryono (e-mail: dwi@its.ac.id).

This research is funded by the Indonesian Endowment Fund for Education (LPDP) on behalf of the Indonesian Ministry of Higher Education, Science, and Technology, and managed under the EQUITY Program (Contract No. 4299/B3/DT.03.08/2025 and No. 3029/PKS/ITS/2025), Directorate of Research and Community Service Institut Teknologi Sepuluh Nopember Ministry of Higher Education, Science, and Technology in accordance with the Implementation Agreement of the Strategic Research Grant (SRG) Assigned Research Program Institut Teknologi Sepuluh Nopember Fund for the Year 2026 Researcher Contract Number: 1669/PKS/ITS/2026, dated February 10, 2026.

arXiv:2610.29515v1 [cs.LG] 21 Oct 2026

ABSTRACT Models are updated repeatedly, and an update that looks better on validation data but not on new data is a false improvement that can misguide clinical decisions. The naive rule accepts every positive validation gain and is optimistic when the same data choose and judge the update. BRIDGE is a model-agnostic acceptance gate that decides each update from predictions and the available labels, with fixed settings. Each fold of the data judges the update chosen on the other folds (cross-fitting) to reduce selection bias; acceptance requires the gain's lower bootstrap bound to exceed a margin. The four settings are (A) 11 new datasets, (B) synthetic tasks, (C) 12 retrospective clinical tasks, and (D) 39 supporting datasets. The false-improvement rate falls from 19.1% (naive) to 1.1% (BRIDGE) in (A), 20.7% to 1.7% (B), 20.0% to 0% (C), and 12.5% to 0.5% (D); every 95% confidence interval of the reduction over datasets or tasks excludes zero (not over the clinical sources). An equal-size random draw from the naive acceptances contains more false improvements in (A)–(E) (post hoc, $p < 0.001$). On the OpenML collection (E) of 14 new datasets, the false-improvement rate falls from 25.6% to 1.0%. BRIDGE removes 97% to 100% of the false improvements that the naive rule accepts, and its mean gain per deployed update is 1.8 to 3.1 times the naive rule's. The 95% upper bound on the false-improvement rate is below 5.5%: formal in (B) for independent tasks (3.56%), and conditional on the real data tested (0.33% to 5.40%).

INDEX TERMS Acceptance gate, bootstrap, clinical prediction, cross-fitting, decision algorithm, false improvement, model selection, model updating, risk control

I. INTRODUCTION

Machine-learning models are typically selected and approved on the basis of training and validation scores, but performance on those scores is often not reproduced on previously unseen data [1], [2], [3], [4]. Deployed models therefore frequently have to be revised [5], [6], and the changes in question include a threshold shift [7], an ensemble [8], feature selection [9], [10], and hyperparameter tuning [11], [12].

The naive validation-improvement rule accepts such a change whenever the validation score improves, during development or after deployment. Formally, let b denote the baseline predictor, the model currently in use, and c a candidate update. Let $\mathcal{M}$ be a performance metric evaluated on a validation set $\mathcal{D}_v$ that is held out from fitting, but not from the selection of the candidate. The update is accepted whenever the observed difference

$$\hat{\Delta}(c, b) \;=\; \mathcal{M}(c; \mathcal{D}_v) \;-\; \mathcal{M}(b; \mathcal{D}_v) \qquad (1)$$

is positive. The same validation data serve two purposes: choosing the candidate and then estimating the gain of the chosen candidate against the baseline. The twofold use spans classical model comparison [13], [14], [15], [16] and modern automated pipelines [17], [18].

Because the validation data cannot judge the update chosen on that same data, in this study each accepted update is

scored on separate data set aside from the start, the locked reference. The outcome is one of two: the change improves, when the selected metric is higher than the baseline metric, or the change is a false improvement (a harmful update in the pre-registration documents), when the metric is lower or equal. A tie is counted as a false improvement, because a change that brings no improvement still carries the cost and the risk of replacing the model (formalized in Definition 1). The risk quantity of interest [19] is the proportion of false improvements among those accepted.

The false-improvement rate among accepted updates is difficult to bound, because two sources make the validation gain optimistic. The first source of optimism is selection bias: a candidate is chosen as the best of many on the validation data and then judged on those same data, so its apparent improvement is larger than the true one. Maximizing over many candidates is a multiple-comparisons problem, and the resulting optimism is known as selection bias or as the winner's curse [13], [14]. The bias grows with the number of candidates and the amount of tuning, as formalized in (2) (Section II-A; [10], [15]).

The second source of optimism is data leakage into the scores used for selection: Kapoor and Narayanan [20] found leakage in 294 machine-learning papers across 17 fields, including the use of test-set information in model or feature selection. The BRIDGE evaluation protocol closes that source: the locked reference is used only for judging, so leakage into the judging data is excluded by design and the optimism studied in this paper is that of selection.

Four directions of prior work address parts of this problem. First, estimating the performance of a selection procedure (nested cross-validation [13], bias correction of the selected error rate [15]) returns a performance estimate. Second, comparing two models or two training procedures [21], [22], [23] assumes the compared pair is fixed before the results are seen. Third, governing a released model or an evaluation resource (medical-device approval policies [19], [24], [25], no-regression certification [26], [27], adaptive data analysis [28], [29]) needs data beyond the labels already available. Fourth, controlling the risk of a chosen policy (Learn then Test [30]) needs a separate calibration set.

Despite these four directions, the naive rule remains in use in practice, and the danger is real. Across the five collections, (A) to (D) and the OpenML collection (E), 12.5% to 25.6% of the updates that the naive rule accepts are false improvements (Fig. 1, row "All"). The danger concentrates where the evidence is weakest: for a validation gain of $+0.5$ points or less, 38% to 64% of the accepted changes are false improvements, whereas above $+2$ points 0% to 4% are (Fig. 1). The pattern is descriptive; it was seen after the study was designed, and it was not used to set the gate parameters, which were fixed beforehand.

If left alone, the problem accumulates, because each false improvement replaces a model that is already running, with its cost and its risk. Repeated updates cause a slow decline in performance, the effect called bio-creep [19], [24], and add

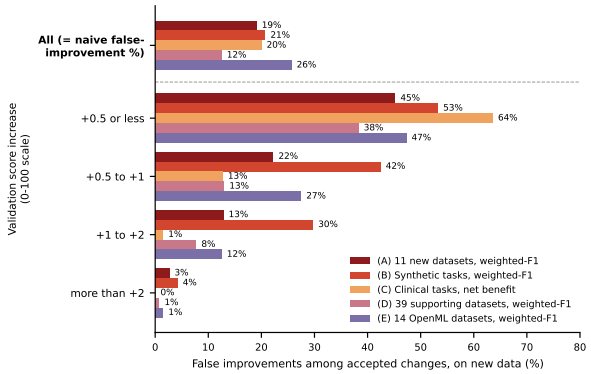

FIGURE 1. False improvements accepted by the naive rule, grouped by validation gain, in collections (A)–(E).

Note. The figure is descriptive. Current model: histogram gradient boosting (HGB) with the default candidate space (HGB, decision trees, and a support vector classifier); other learners, other candidate members, and a joint seven-family space (XGBoost, multilayer perceptron, logistic regression, naive Bayes, k-nearest neighbors) are tested separately (Tables H11 and H8). Data and comparison with BRIDGE: Sections IV-C–IV-E. Above $+5$ points none of the 726 accepted changes is a false improvement in any setting, so that band is merged into "more than $+2$".

technical debt to the machine-learning system [5], [6]. In the clinic, an update that makes a model look better in validation but not on new data is a false improvement, and a screening or diagnostic model updated on that basis can misguide decisions about patients [4], [2]. Regulators therefore ask for control over the changes made to a deployed model [25].

The four directions leave two gaps, a methodological gap and an empirical gap. For the methodological gap, a fixed acceptance gate for one selected update combines cross-fitted selection, a bootstrap decision score, and a margin; its false-improvement rate is evaluated across tasks (Table 1). Honest holdouts and tests with re-selection also separate choosing from judging and serve as comparators, whereas the closest governance and calibration methods use further data [19], [26] or bound the risk of a selected policy [30]. The empirical gap is the absence, in the literature reviewed in Section II, of a large-scale pre-registered evaluation across change types, metrics, and domains, with a factorial component analysis of the acceptance gate.

Pre-registration means that the design, the hypotheses, and the analysis are fixed before the runs. The gap matters because model comparisons often depend on the seed and the data split [31] and leakage and methodological failures are common [20], [2]. In this study, pre-registration is internal: each plan was written and hashed before its run, without a third-party registry, so the order of plan and run rests on the recorded times (Appendix C). Pre-registration limits the freedom to choose the analysis after the results are seen [32], whereas seed variation is handled by drawing a new seed for every trial.

The paper therefore proposes BRIDGE, an acceptance gate that decides one update using only the labels already available. BRIDGE separates the data that choose an update from the data that judge it (cross-fitting) and accepts the update only when the lower bootstrap bound on its gain exceeds

a minimum gain, the margin. Changes whose gain is small and uncertain—the group with the most false improvements in Fig. 1—are therefore rejected. On the collections tested (Section III-E), the false-improvement rate among accepted updates falls to at most 1.7% (Section IV-A). The contributions respond to the two gaps: C1 and C2 answer the methodological gap, C3 answers the empirical gap, and C4 makes both usable and checkable by others.

- **C1: BRIDGE, an acceptance gate with fixed settings.** BRIDGE blocks false improvements by accepting or rejecting one update after it has been chosen, using the current model's predictions, the candidate's predictions, and the labels already available, with no new data. What is new is the cross-fitting step, which separates the data that choose an update from the data that judge it, used together with the bootstrap lower bound and the margin. The settings are the seven frozen parameters (the fold count, the bootstrap count, the quantile, the margin, the bag count, the bag fraction, and the SEARCH size) and the candidate space. These settings are fixed once during development and used unchanged for every task, so no per-task tuning reintroduces selection bias. Binary and multiclass tasks use the same procedure (Section III).
- **C2: A finite-sample bound and a component analysis.** The finite-sample bound is Proposition 1: a 95% upper bound on the proportion of false improvements among those accepted, computed on the locked reference after each decision. The gate itself decides only from POOL, so the bound is an evaluation protocol. The bound applies to a finite number of tasks rather than only asymptotically, with the condition that the tasks are drawn independently from one population (Section III-D). The bound targets the false-improvement rate among accepted updates rather than a prediction error. The component analysis uses a factorial design over three components: the cross-fitted evidence, the bootstrap lower bound, and the margin. The design switches the three on and off in all eight combinations and also varies the size of the candidate space, to show which component lowers the false-improvement rate (Section IV-F).
- **C3: An internally pre-registered evaluation and a clinical application.** The evaluation is internally pre-registered, so the design, the pass criterion, the dataset list, and the seed rule are fixed before each run. Each file is recorded with a SHA-256 cryptographic hash [33] in a manifest that is kept in the code repository (chronology in Appendix C). A reader can therefore check that each reported result comes from the hashed code. All rules judge the same decisions in the four settings (A) to (D) (Section III-E), each with a new seed and ground truth on the locked reference; the OpenML collection (E) is run through the released program as well. BRIDGE is compared with acceptance rules built on classical test statistics, an honest holdout, and a fixed margin. The conditions of use are tested for base learners, tuning, and imbalanced data. The clinical application applies the gate retrospectively to 12 clinical tasks from 6 sources, with net benefit [34] as the decision metric (Section IV-E).
- **C4: A reproducible open-source implementation.** The open-source implementation is the package `bridge-gate` (`pip install bridge-gate`), which runs the paper's procedure as its default behavior and can be applied to a user's own data. The package reproduces every stored result row of 15 runs, and the repository and the release archive are listed (Appendix C; Data Availability Statement).

II. LITERATURE REVIEW

The related work is organized around three questions that match the three subsections: why validation-based model choice is biased, and how data leakage adds to that bias (Section II-A); what classifier comparison tests assess (Section II-B); and how prior work governs the approval of model updates (Section II-C). Table 1, at the end of Section II-C, compares the approaches closest to BRIDGE.

A. SELECTION BIAS AND DATA LEAKAGE IN VALIDATION-BASED MODEL CHOICE

Validation-based model choice is the common practice of selecting, from many candidates, the one with the highest validation score, and then reporting that score as evidence of improvement [14]. This practice causes selection bias, although for a single model that is already fixed the validation score is unbiased. Once the best of many candidates is chosen, the score of the selected candidate is the highest of many noisy scores. On average, that score therefore exceeds the gain that the candidate would show on new data.

In symbols, writing $\hat{\Delta}(c, b)$ for the validation gain of candidate c over the baseline b, with $\mathcal{C}$ the frozen candidate space (Section III-A),

$$\mathbb{E}\left[\max_{c \in \mathcal{C}} \hat{\Delta}(c, b)\right] \;\geq\; \max_{c \in \mathcal{C}} \mathbb{E}\left[\hat{\Delta}(c, b)\right], \tag{2}$$

where all expectations are over the random validation-data draw. The gap between the two sides of (2) can grow with the number of candidates and the amount of tuning, so hyperparameter search itself becomes a form of fitting [14]. Beyond this selection bias, a reported difference can also be lost in seed and partitioning variability, which replication studies find can be as large as the difference itself [31].

Three approaches reduce selection bias: (i) nested cross-validation [13] and bias correction of the minimum cross-validation error [15] estimate the performance of the selection procedure rather than of the model deployed [35]; (ii) post-selection inference accounts for the mechanism that selected the object [36], [37]; and (iii) cross-fitting separates the data used to choose from the data used to judge [38]. BRIDGE uses this principle within each round of the gate

(Section III-B) and returns an accept/reject decision for one update.

Besides selection bias, the second source of optimism noted in Section I is data leakage. Group leakage, where rows from the same group or subject cross the split, and temporal leakage, where information from a future time enters an earlier stage, both make a validation score too optimistic [39], [40], [20]. The paper closes that source with an explicit group and temporal audit before any split (Section III-E).

B. STATISTICAL TESTS FOR COMPARING CLASSIFIERS

Statistical tests are the standard way to compare classifiers, and the tests differ in the objects being compared: (i) McNemar's test [21] compares the predictions of two models that are already trained, on the same observations; (ii) two training procedures are compared by refitting on several data splits, using the 5×2 cross-validated (cv) paired t-test [22] or the corrected resampled t-test [23]; and (iii) two methods are compared across datasets by rank-based [41] or Bayesian tests [42]. The first two kinds assume that the compared pair is fixed before the results are inspected. Neither therefore assesses the decision that follows when the best configuration is chosen among many on the same data. The third kind answers which method is better, not whether a single update should be accepted.

C. GATES FOR APPROVING MODEL UPDATES

The work most directly related to this paper concerns gates, that is, rules that accept or reject a proposed model update before the update is used. Three approaches build such gates for model updates, and each approach has its own data needs. First, medical-device approval policies [19], [24] formalize bio-creep, the risk that approving too many changes slowly degrades performance, and test online on a single deployment stream. Second, slice-wise no-regression certification [26] evaluates a candidate and its baseline on independent certification data, with fallback to the current model. Third, disagreement labeling [27] uses unlabeled deployment data and requests labels only on the items where the candidate and baseline differ. All three need data beyond the labels already available, whereas BRIDGE uses only the labels already available.

Calibration methods also need a separate calibration set and work with policies, where a policy is a model together with its decision rule, for example a threshold. Learn then Test [30] uses the calibration set to certify that the risk of a chosen policy stays below a target level, with a multiple-testing procedure that controls the family-wise error across the candidate policies. Conformal risk control [43] uses it to calibrate a single threshold so that the expected loss of the resulting policy is bounded. BRIDGE instead answers a paired question, whether this update improves on the current model, from the cross-fit evidence (Section III-B).

Three further approaches answer other questions: (i) reusable-holdout and leaderboard mechanisms [28], [29] guard an evaluation resource reused across many queries rather than deciding one update; (ii) clinical model updating, such as recalibration and revision [44], [45], changes the model to fit new data without an accept/reject gate; and (iii) sharpness-aware minimization [46], invariant risk minimization [47], and group distributionally robust optimization [48] change how models are trained, not whether a post-training change is applied.

This paper asks whether a single update, chosen from many candidates, may be accepted using only the labels already available, and it evaluates how often the accepted updates are false improvements with an upper bound computed on a locked reference over many decisions. BRIDGE answers by separating the data that choose from the data that judge (Section III). Table 1 compares the closest approaches on what the method controls, whether it accounts for selection, what labels the decision needs, and what the bound covers.

Relative to the methods in Table 1, the closest comparison is with the calibration methods. Learn then Test and conformal risk control can take the calibration set from the same label budget, but these methods bound the risk of a selected policy, whereas the evaluation of BRIDGE bounds the share of false improvements over many decisions. The comparison with Learn then Test is conceptual, and no empirical advantage is claimed. BRIDGE chooses an update on one part of the data and judges it on another, and it uses the training labels that already exist. Its bound covers the acceptance procedure over a population of tasks rather than one deployed model.

III. THE BRIDGE PROCEDURE AND EXPERIMENTAL DESIGN

The BRIDGE gate (a label, not an acronym) turns a candidate update into an accept/reject decision from data that are already labeled. The gate operates within a complete procedure (Fig. 2) that runs from the split of the data to the deployment decision.

The proposed procedure (Fig. 2) has three parts: the fixed settings, the outer layer, and the inner layer. The fixed settings were chosen during development, before the confirmatory runs. These settings comprise the model type, the candidate space, and the gate parameters (Table 2). The outer layer splits the data of a task once into the locked reference and TRAIN. The locked reference is used solely to score the experiments, and the rows of TRAIN are re-drawn in every trial. TRAIN has $2{,}430$ rows, or $1{,}000$ on small datasets, and is split into SEARCH (630 rows) and POOL ($1{,}800$ rows, three folds of 600). The base models are fitted on SEARCH, and an update is selected and judged on POOL. Table D2 (Appendix D) gives the basis for these sizes.

In the inner layer, the cross-fit, POOL is split into three folds, and in each round a candidate is selected on two folds and judged on the third. The pooled judgments decide whether the update is accepted. An accepted update is re-selected on the whole of POOL and deployed; otherwise the baseline is retained. Every judgment therefore rests on data

TABLE 1. Positioning of BRIDGE against related gates and methods.

Method	What it controls	Accounts for selection	Labels needed at decision	Bound covers
Nested CV and bias corrections [13], [14], [35]	error-estimate optimism	yes (inside estimate)	existing labels only	selection procedure
McNemar / 5×2 cv / corrected t [21], [22], [23]	fixed-pair difference	no	existing labels only	prespecified comparison
Learn then Test [30]	policy risk	yes	reserved calibration set	selected policy
Update approval [19], [24]	bio-creep	yes (within policy)	labeled deployment stream	approval policy
No-regression certificates [26], [27]	per-slice regression	yes (within slices)	certification set (Zhang); disagreements only (Balachandran)	deployed model
BRIDGE	false-improvement rate among accepted updates	**yes** (per-fold separation)	**existing training labels only**	acceptance procedure, population-level

*Column meanings and reading: Section II-C; empirical comparison with an honest holdout: Table H3.

that were not used to fit or to select the candidate being judged.

Each component of the gate addresses one source of error: (i) selection bias, which makes the gain of the best candidate optimistic when the data that chose it also judge it (2), so cross-fitting judges each candidate on a fold that did not choose it; (ii) sampling noise, which can make a gain positive by chance, so the gate uses the bootstrap lower bound; and (iii) a gain too small to justify replacing the running model, so the margin sets the minimum improvement. The false-improvement rate is then evaluated with the finite-sample bound (Proposition 1) and the KL-ratio bound on the locked reference (Section III-D).

Sections III-A–III-D follow the figure: Section III-A presents the partition and the candidate space, and Section III-B presents the cross-fit and the bootstrap gate. Section III-C presents the metrics and multiclass tasks, and Section III-D presents the upper bound on the false-improvement rate. Section III-E then describes the experimental design, namely the data, the partition into trials, the comparators, and the evaluation metrics. Tables 3, 4, and 5 list the notation, and Algorithms 1–2 give the procedure step by step.

A. DATA PARTITION AND CANDIDATE SPACE

The data partition, the outer layer of Fig. 2, begins with the labeled examples of a task, $\{(x_i, y_i)\}_{i=1}^{N}$ with $y_i \in \{1, \ldots, K\}$ for K classes. A disjoint subset, the locked reference, is set aside once and used only to score the experiments; it supplies the ground truth at evaluation time. The rows that remain once the reference has been set aside are eligible for training. Each trial samples its training rows and splits the rows into a SEARCH set and a POOL set with a fresh seed (Section III-E),

$$\text{TRAIN} \;=\; \text{SEARCH} \sqcup \text{POOL}, \tag{3}$$

where SEARCH fits the base learners and POOL carries the evidence used both to select an update and to gate it (notation in Table 3). Base-learner parameters and preprocessing are fitted only on SEARCH. On POOL nothing is fitted: the gate selects one threshold-and-weight configuration from the prespecified candidate space.

TABLE 2. Frozen parameters of BRIDGE and of the evaluation protocol.

Parameter	Value
Cross-fit folds F	3
Bootstrap replicates N_{boot}	20,000
Lower quantile q	5%
Margin ε	0.005
Base learner	HGB, library defaults with at most 300 iterations and early stopping
Bagged base models	32 bags of 80% of SEARCH
Ensemble members	DT and SVC, library settings
Preprocessing	fitted on SEARCH only (WiDS exception: Table D1)
Candidate space	173 (binary); $14 + 8K$ for K classes
Training rows	2,430 (SEARCH 630, POOL 1,800); 1,000 when fewer rows
Locked reference	one pre-registered seed per task; size by a fixed rule
Trials	100 per dataset, 200 per clinical task, 60 per multiclass dataset, 1 per task of $\mathcal{Q}$; one fresh seed per trial

*Details of every value and its basis (where it was set and where it was tested): Table D2 in Appendix D.

a: Base learners.

The default base learner is a histogram gradient-boosted classifier (HGB) [49], [50], [51] with the settings fixed before any evaluation and listed in Table D2. The base learner uses a fixed project configuration: the library defaults, except at most 300 boosting iterations with early stopping after 10 iterations without improvement (Table D2). This fixed learner separates the effect of the gate from the effect of tuning. Tuning is treated separately as an update that can be proposed to the gate (Section IV-G).

All base models are bagged to reduce the variability of the predicted probabilities. Each bag is an 80% subsample of SEARCH, and 32 bags are drawn [52], [53] (the basis and the tests of 16 and 64 bags are in Table D2, Appendix D). A bag is retained provided that every class has enough examples for the internal validation split of the base model. The bagged score $\bar{p}_\ell(x)$ of base model ℓ is the average over the retained bags. Appendix E gives the minimum class count per bag and the few trials skipped under this rule.

b: Baseline and candidates.

The baseline b is the base learner alone at the fixed threshold 0.5, the usual default of a probabilistic classifier. A candidate c is a weighted, thresholded combination of L base models,

$$f_{w,t}(x) \;=\; \mathbb{1}\Big[\textstyle\sum_{\ell=1}^{L} w_\ell\, \bar{p}_\ell(x) \;\geq\; t\Big], \tag{4}$$

TABLE 3. Data units, partition, and predictors.

Symbol / term	Definition
Task	One dataset, or one draw from the synthetic population $\mathcal{Q}$.
Trial	One re-randomization of the training rows, the SEARCH/POOL split, and the bag draw under a fresh seed; produces one decision; trials per task: 100 (real), 200 (clinical), 60 (multiclass), 1 ($\mathcal{Q}$).
Decision; index j	Accept or reject in one trial; j runs over trials on real data and over tasks on $\mathcal{Q}$; J is the number of trials or tasks.
Source	Clinical: data provider (six sources; the four National Health and Nutrition Examination Survey (NHANES) tasks form one source, the four Behavioral Risk Factor Surveillance System (BRFSS) tasks another). Non-clinical: the dataset.
Locked reference R_j	Rows set aside once per task with a pre-registered seed; used only as the ground truth to judge the deployed update. Size: Appendix B.
TRAIN = SEARCH ⊔ POOL	SEARCH fits the base models; POOL holds the evidence for selection and gating. Size: Appendix B.
Development datasets	The 18 datasets on which $\mathcal{C}$ and the gate were frozen; a subset of the 34 confirmation datasets in (D).
New data (unseen during development)	Data held out from designing and freezing the gate (Section III-E).
b (baseline)	Default base learner (histogram gradient-boosted classifier, HGB) averaged over the kept bags at threshold 0.5 (Section III-A).
$\bar{p}_\ell(x)$	Mean predicted probability of base model ℓ over its retained bags (80% subsamples of SEARCH); bag-eligibility rule: Section III-A.
c (candidate), $\mathcal{C}$	A configuration $f_{w,t}$; $\mathcal{C}$ is the 173 configurations of the binary default space (threshold-only, ensembles with decision trees (DT), ensembles with DT and a support vector classifier (SVC)); other protocols use separate prespecified spaces; basis: Table D2.
Update	A proposed model change: a candidate from $\mathcal{C}$ selected on POOL that the gate may accept or reject; accepted update: the deployed configuration, re-selected on the whole of POOL.
Frozen	Gate parameters (F, N_{boot}, q, ε, bag count) are fixed before any evaluation run and never tuned per dataset; values: Table 2.

*Shading marks the four groups in order: *data units* (white), *data partition* (light gray), *models and updates* (gray), and *protocol* (darker gray).

where w lies on the simplex ($w_\ell \geq 0$, $\sum_\ell w_\ell = 1$) and $t \in (0, 1)$ is the decision threshold; the baseline is recovered at $L = 1, w_1 = 1, t = 0.5$.

In the binary default, $\mathcal{C}$ contains three families: other thresholds for the HGB base learner; mixtures of HGB with decision trees (DT) [54]; and mixtures of HGB, DT, and a support vector classifier (SVC) [55]. Each family is combined with a grid of thresholds, because moving the decision threshold is a standard way to trade errors between the classes [7]. Moving the threshold also handles class imbalance without reweighting the training data. The mixtures follow weighted ensemble averaging [8]. The grids are coarse by design to keep the space small, because the selection bias of (2) grows with the number of candidates.

The default space has 173 configurations. Appendix D gives the grids, the count, and the development history of the space, and Table D2 gives the basis of each setting. Section IV-F tests whether this coarse choice matters. In that section, the same grids are refined to 1,825 and 11,249 configurations (grids in Table H7). The multiclass space follows the same idea, with class scales in place of thresholds. Section III-C describes this space together with the multiclass metrics.

The three members, HGB, DT, and SVC, were fixed stepwise during development and represent three ways of learning: boosted trees, a single tree, and a kernel. During development, SVC replaced an earlier k-nearest neighbors (KNN) member [56]; Table C1 (Appendix C) lists the development runs behind this choice. The base model throughout is the HGB classifier; DT and SVC enter only as candidate members. Two other member pairs are evaluated as conditions of use in Section IV-G: logistic regression [57] with naive Bayes [58] and a multilayer perceptron (MLP) [59] with KNN. A joint space of seven member families (DT, SVC, logistic regression, naive Bayes, MLP, KNN, and XGBoost) is evaluated there as well. XGBoost [60] and an MLP are also evaluated there as alternative base learners. No member is tuned, and Appendix D gives the member settings.

With the space fixed in this way, the only adaptation per task on POOL is the choice of one configuration from $\mathcal{C}$, or no change. The space and the gate parameters were frozen during development on 18 datasets (dataset list in Table B1) before any evaluation. The confirmation datasets of (D), including those 18, are therefore re-evaluated with fresh per-trial seeds, whereas (A) contains task families not used during development (Section III-E).

B. CROSS-FITTED SELECTION AND BOOTSTRAP GATE

a: Cross-fitted selection.

POOL is partitioned into $F = 3$ stratified folds $Z_1, \ldots, Z_F$, with other values of F tested in Appendix D. In round k, the fold Z_k is the scoring fold and its complement POOL $\setminus Z_k$ is the selection data. A candidate c_k is chosen on POOL $\setminus Z_k$ by maximizing the metric of interest subject to strictly improving on the baseline, and the baseline and c_k are then evaluated on Z_k. When no candidate strictly improves on POOL $\setminus Z_k$, c_k is the baseline b (the fallback). Algorithm 1 gives the steps of the gate. Algorithm 2 gives the selection rule, including the tie rule.

Repeating over the F folds and pooling the held-out predictions yields the cross-fitted triple counts

$$n_{y,\tilde{b},\tilde{c}} \;=\; \big|\{\, i \in \text{POOL} : y_i = y,\ \hat{b}_i = \tilde{b},\ \hat{c}_i = \tilde{c}\,\}\big|, \quad (5)$$

where y, $\tilde{b}$, and $\tilde{c}$ range over the classes, and $\hat{b}_i$ and $\hat{c}_i$ are the predicted labels of the baseline and of the candidate. The label $\hat{c}_i$ records the prediction of the candidate c_k that was selected for the fold containing observation i. The

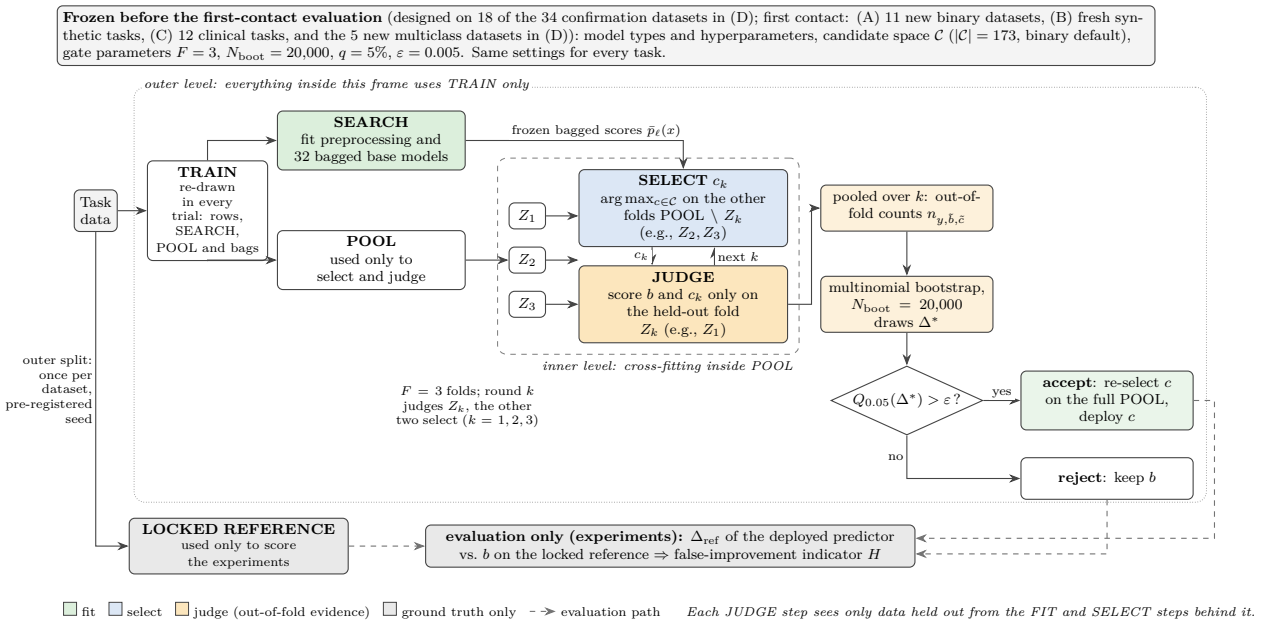

FIGURE 2. Two-level separation: outer train/reference split, inner three-fold cross-fit within POOL; colors mark subset roles.

pooled counts therefore combine the predictions of the per-fold candidates, which need not be the same configuration. Section III-C gives the coding of the labels for binary and multiclass tasks.

The baseline confusion table $n^{\text{base}}_{y,\hat{y}} = \sum_{\tilde{c}} n_{y,\hat{y},\tilde{c}}$ and the candidate confusion table $n^{\text{cand}}_{y,\hat{y}} = \sum_{\tilde{b}} n_{y,\tilde{b},\hat{y}}$ are both derived from these counts. A metric expressible in these counts is computed from the two tables, and the metric difference is defined as

$$\Delta \;=\; \mathcal{M}\big(\{n^{\text{cand}}_{y,\hat{y}}\}\big) \;-\; \mathcal{M}\big(\{n^{\text{base}}_{y,\hat{y}}\}\big)\,. \qquad (6)$$

Because the configuration evaluated on fold k was selected without fold k, the counts are out-of-sample with respect to selection. The separation is made explicit in Table D1 (Appendix D): within the cross-fit, the per-fold candidate is assessed on the fold that was not used to select it.

b: Bootstrap gate.

The gate resamples the pooled counts $N_{\text{boot}} = 20{,}000$ times from the multinomial distribution defined by the observed proportions and the observed total. The resamples form a bootstrap distribution $\{\Delta^*_{(r)}\}_{r=1}^{N_{\text{boot}}}$ [61]; Table 4 lists the quantities of the gate. The resampling does not repeat the candidate selection, so the lower bootstrap quantile $Q_{0.05}(\Delta^*)$, at the quantile level $q = 0.05$, is a decision statistic rather than a confidence bound. BRIDGE accepts the update exactly when that lower quantile exceeds a fixed margin,

$$\text{accept} \iff Q_{0.05}(\Delta^*) \;>\; \varepsilon, \qquad \varepsilon = 0.005, \qquad (7)$$

and otherwise retains the baseline.

The five gate parameters of (5)–(7) have been fixed since development and are the same for every dataset: $F = 3$, $N_{\text{boot}} = 20{,}000$, quantile level 5%, margin $\varepsilon = 0.005$, and bag count 32 (Table 2); the basis for each value is given in Table D2 (Appendix D). Together with the bag fraction and the SEARCH size, these values make up the seven frozen parameters (Section V-A). The sensitivity of the criterion to these values is reported in Appendix D. The margin $\varepsilon = 0.005$ is the operational minimum improvement fixed during development (0.5 percentage points of weighted F_1, 0.005 on the 0–1 scale), and its basis is also given in Table D2. For clinical decisions, benefit is assessed separately with net benefit in (C).

Once the gate accepts, the configuration to deploy is re-selected on the whole of POOL, so the deployed configuration may differ from the per-round candidates on which the decision was based. The bootstrap quantile on the cross-fit evidence is an empirical decision score, not an interval for the final configuration. Whether that score transfers to the final configuration is assessed empirically on the locked reference, through the false-improvement rate of the deployed configuration and Fig. 4. In that figure, the cross-fit estimate of the final configuration's gain is slightly pessimistic, by 0.0034 on average. Among the updates that BRIDGE accepts, the reference gain of the final configuration reaches the lower bound in 80% to 96% of cases. In the cases that fall short, the mean shortfall is 0.0026 to 0.0089, and most of these updates still improve on the reference (post hoc; Appendix H). The false-improvement rate and its bounds are then measured for that deployed configuration on the locked reference (Definition 1). The two layers of the design resemble nested cross-validation, and Remark 1 states how the two differ.

Remark 1 (Relation to nested cross-validation)**.** *Nested*

TABLE 4. Quantities of the bootstrap gate.

Symbol / term	Definition
$n_{y,\tilde{b},\tilde{c}}$	Cross-fitted triple counts: label, baseline prediction, and candidate prediction on each held-out fold, pooled over the $F = 3$ folds.
Δ, Δ^*	Metric difference (candidate minus baseline) from the pooled counts; Δ^* is one of $N_{\text{boot}} = 20{,}000$ multinomial bootstrap draws.
Accept rule	Accept if and only if $Q_{0.05}(\Delta^*) > \varepsilon$; otherwise retain b.
Margin ε	0.005 on the scale of the selected metric, not rescaled; for net benefit (10), five additional net true positives per 1,000 examples.
$\mathcal{M}$	Any higher-is-better metric computed from a confusion table (weighted F_1, balanced accuracy, macro-F_1, F_2, and, in the binary clinical protocol, net benefit); $\mathcal{M}$ is the metric, distinct from the accepted count $M = \sum_j A_j$ (Section III-D).

Algorithm 1 BRIDGE gate

Input: SEARCH and POOL with labels y; candidate space $\mathcal{C}$; metric $\mathcal{M}$; bag count, fold count F, replicate count N_{boot}, quantile level q, and margin ε (values in Table 2)
Output: deploy candidate c or retain baseline b

fit the bagged base models on SEARCH and store the bag predictions on POOL, including the baseline predictions $\hat{b}$
$\{Z_k\}_{k=1}^{F} \leftarrow$ stratified F-fold split of POOL; initialize counts $n_{y,\tilde{b},\tilde{c}} \leftarrow 0$
for $k = 1$ to F **do**
    $c_k \leftarrow \textsc{Select}(\text{POOL} \setminus Z_k,\ \mathcal{C},\ \mathcal{M},\ b)$ ▷ must improve on the baseline
    accumulate counts $n_{y,\tilde{b},\tilde{c}}$ on Z_k for $(\hat{b}, \hat{c}_k)$
end for
$\{\Delta^*_{(r)}\}_{r=1}^{N_{\text{boot}}} \leftarrow$ resample the pooled counts N_{boot} times from the multinomial distribution and compute (6) for each draw
if $Q_q(\Delta^*) > \varepsilon$ **then**
    $c \leftarrow \textsc{Select}(\text{POOL},\ \mathcal{C},\ \mathcal{M},\ b)$
    return deploy c
else
    return retain baseline b
end if

Algorithm 2 Candidate selection on a subset (used by Algorithm 1)

function $\textsc{Select}(Z, \mathcal{C}, \mathcal{M}, b)$
    score every $c \in \mathcal{C}$ on Z using the stored bag predictions
    return $\arg\max_c \mathcal{M}(c)$ subject to $\mathcal{M}(c) > \mathcal{M}(b) + 10^{-12}$; else b; candidates within 10^{-12} of the best take the first in list order
end function

cross-validation estimates the expected performance of a selection procedure (Section II-A), whereas the gate answers a different question. The outer layer of the design is the locked train/reference split, and the inner layer is the cross-fit within POOL (Fig. 2; data used for each choice in Table D1, Appendix D). The cross-fitted counts yield the decision statistic, whereas the population bound for the false-improvement rate among accepted updates comes from Proposition 1 (Section III-D); that bound holds on average over decisions from the task population, not for each decision. For this construction the cross-fit within POOL takes the role of the inner layer, so no additional nested layer is needed.

C. GENERALIZATION TO METRICS AND MULTICLASS TASKS

The same gate generalizes to any metric computed from the confusion table and to multiclass tasks with K classes, without changing the decision rule. The counts, the metric difference Δ, and the bootstrap are unchanged, so the gate applies to metrics defined on the confusion structure of (5) whose direction is higher-is-better; a lower-is-better metric would need a sign transform, which this implementation does not define. Labels and predictions are coded in $\{0, 1\}$ in the binary experiments, a relabeling of the two classes. For K classes the labels and predictions range over $[K] = \{1, \ldots, K\}$ as in Section III-A. In the binary case, applying (6) to the baseline and candidate confusion tables gives

$$\Delta \;=\; \mathcal{M}\big(\{n^{\text{cand}}_{y,\hat{y}}\}_{y,\hat{y}\in\{0,1\}}\big) \;-\; \mathcal{M}\big(\{n^{\text{base}}_{y,\hat{y}}\}_{y,\hat{y}\in\{0,1\}}\big)\,. \tag{8}$$

Multiclass tasks use a bagged HGB base learner with class-scale and simplex candidates. The class-scale family scales the score of one class at a time. The simplex family averages HGB, DT, and SVC, which are the same members as in the binary space. The alternative member pairs are also evaluated on the five multiclass datasets of (D) (Section IV-G): logistic regression with naive Bayes and an MLP with KNN. The joint space of the seven member families is evaluated there as well. The default multiclass space has $14 + 8K$ candidates for K classes, and the joint space has $42 + 32K$. Appendix D gives the composition, the member settings, and the tie rules, and Table D2 gives the basis of each setting. The baseline is the argmax of the bagged HGB, and the same gate procedure uses multiclass weighted-F_1. For K classes, the construction uses the confusion tables over $y, \hat{y} \in [K] = \{1, \ldots, K\}$,

$$\Delta \;=\; \mathcal{M}\big(\{n^{\text{cand}}_{y,\hat{y}}\}_{y,\hat{y}\in[K]}\big) \;-\; \mathcal{M}\big(\{n^{\text{base}}_{y,\hat{y}}\}_{y,\hat{y}\in[K]}\big)\,. \tag{9}$$

For binary and multiclass tasks alike, the evidence is stored as counts. The number of resampled categories therefore grows as K^3. The count-based representation limits the implementation to metrics computed from hard predicted labels.

Common metrics of this kind, namely weighted F_1, balanced accuracy, macro-F_1, F_2, sensitivity, and specificity [62], [63], [64], are computed from hard labels. Equations (8)

and (9) therefore cover these metrics. Net benefit at a threshold probability p_t is used in the binary clinical protocol,

$$\mathrm{NB}(p_t) \;=\; \frac{\mathrm{TP}}{n_{\mathrm{eval}}} \;-\; \frac{\mathrm{FP}}{n_{\mathrm{eval}}} \cdot \frac{p_t}{1-p_t}, \tag{10}$$

where TP and FP are the numbers of true and false positives on the data being assessed, and n_{eval} is the number of examples in those data. The clinical decision threshold p_t is distinct from the candidate prediction threshold t of (4). The definition of net benefit follows [34].

Because the deployed model decides at one threshold, the gate computes the decision metric at that threshold. By contrast, the area under the receiver operating characteristic curve (AUC) [65] scores the ranking of scores across all thresholds. AUC is therefore unchanged when an update changes the threshold alone, as many candidate updates do. The gate therefore uses a threshold metric for both settings. For clinical decisions the metric is net benefit, read as a decision curve [34], and for rare classes the metrics are the minority-sensitive ones, balanced accuracy and macro-F_1 (Section IV-G).

D. RISK AMONG ACCEPTED UPDATES: UPPER BOUND AND SCOPE

Whether an accepted update is a false improvement can be seen on the locked reference alone, so the risk of concern is a proportion over many decisions. This risk is the share of false improvements among the updates that the gate accepts (quantities in Table 5). The subsection first states when this share can be bounded (Assumption 1) and then defines the share (Definition 1). The subsection then gives an upper bound that holds for a finite number of tasks (Proposition 1). The scope of the bound is the average over tasks drawn from one population, not a single decision. Real data, where the assumption does not hold, are treated in Remark 2.

Assumption 1 (Task population)**.** *(i) Tasks $j = 1, \ldots, T$ are drawn i.i.d. from a fixed task population $\mathcal{Q}$; (ii) the batch construction and the training/reference split of each task are part of that draw; and (iii) the gate and candidate space are frozen before any evaluation data are seen. The proofs use only that the decision pairs (A_j, H_j) are i.i.d. across tasks.*

Definition 1 (Risk among accepted updates)**.** *For task j, let $A_j = \mathbb{1}\{\text{update accepted}\}$ and $H_j = \mathbb{1}\{\Delta_{\mathrm{ref},j} \le 0\}$, where $\Delta_{\mathrm{ref},j}$ is the metric difference between the deployed configuration and the baseline of task j measured on its locked reference (the update is a false improvement when it does not improve). The quantity of interest is the proportion of false improvements among those accepted, defined when the acceptance probability is positive ($\mathbb{E}[A_j] > 0$),*

$$r_{\mathcal{Q}} \;=\; \frac{\mathbb{E}[A_j H_j]}{\mathbb{E}[A_j]}. \tag{11}$$

The bound uses two counts over the T tasks: the number M of accepted updates and the number S of false improvements among those accepted. Given M, the count S behaves like the number of false improvements in M independent draws (Lemma 1). A binomial tail bound therefore applies (Proposition 1).

The gate makes each decision from POOL alone. The bound in this subsection belongs to the evaluation protocol: it is computed on the locked reference after the decision has been taken, so it measures risk rather than steering the gate. Proposition 1 applies to any frozen gate, and the contribution is the gate together with its evaluation, not a theorem for a single deployment.

Lemma 1 (Conditional binomial)**.** *Let $M = \sum_{j=1}^{T} A_j$ and $S = \sum_{j=1}^{T} A_j H_j$. Conditionally on $M = m$, $S \sim \mathrm{Binomial}(m, r_{\mathcal{Q}})$.*

Proposition 1 (Finite-sample upper bound)**.** *Under Assumption 1, with $\mathrm{kl}(x\|u) = x\log(x/u) + (1-x)\log((1-x)/(1-u))$, define*

$$U \;=\; \sup\Big\{\, u \in \big[\tfrac{S}{M}, 1\big] : M\,\mathrm{kl}\big(\tfrac{S}{M} \,\big\|\, u\big) \le \log\tfrac{1}{\delta} \,\Big\}, \tag{12}$$

and set $U = 1$ when $M = 0$. Then $\Pr(r_{\mathcal{Q}} \le U) \ge 1 - \delta$, where $\delta \in (0,1)$ is the error level (0.05 in the evaluation).

The bound U exceeds the true false-improvement rate with probability at least 95%, and it is computed from the M accepted updates and the S false improvements on the reference. The bound inverts an exact binomial tail inequality rather than a limit theorem, so the bound holds for any finite number of tasks, not merely asymptotically.

Proof sketch. Conditionally on $M = m$, Lemma 1 gives $S \sim \mathrm{Binomial}(m, r_{\mathcal{Q}})$; inverting the lower Chernoff–Kullback–Leibler tail of a binomial [67], [68] gives a one-sided upper bound of coverage $\ge 1 - \delta$ for every m, and averaging over the (stochastic) M preserves coverage. Full proofs of Lemma 1, Proposition 1 and Corollary 1 are in Appendix A. □

The bound counts an update as a false improvement exactly when the reference shows no improvement. A stricter reading also counts as a false improvement every accepted update whose reference result is uncertain. The bound remains valid under that reading (Corollary 1).

Corollary 1 (Conservative variant)**.** *Counting every accepted update of reference-uncertain status as a false improvement yields a valid but more conservative bound; this variant is reported alongside the Proposition 1 bound in Table C2.*

Remark 2 (Real data)**.** *Proposition 1 does not apply to real benchmark tasks. The trials of a task share one locked reference, whose error enters every trial of that task, and the tasks form a fixed collection rather than an i.i.d. sample. Instead, the false-improvement rate is reported through the KL-ratio bound, which is a ratio of one-sided Kullback–Leibler bounds [68],*

$$\hat{U} = \frac{U_{\mathrm{KL}}(\bar{W};\, J, \alpha/2)}{L_{\mathrm{KL}}(\bar{A};\, J, \alpha/2)}, \tag{13}$$

TABLE 5. Outcome and risk quantities.

Symbol / term	Definition
$\Delta_{\mathrm{ref},j}$	Metric difference between the deployed configuration c_j and b_j on the locked reference.
False improvement, H_j	$H_j = \mathbf{1}\{\Delta_{\mathrm{ref},j} \le 0\}$: the deployed configuration does not improve the selected metric on the locked reference; ties count as false improvements.
Clearly worse / reference-uncertain	From the 95% Bayesian-bootstrap interval of $\Delta_{\mathrm{ref},j}$ [66]: a Dirichlet posterior over the confusion cells with pseudocount 0.5 (the Jeffreys prior) on the binary and clinical protocols and over the positive cells on multiclass, 20,000 draws, quantiles 2.5% and 97.5%; entirely below 0 = clearly worse; entirely above 0 = clearly improving; otherwise reference-uncertain.
A_j; false-improvement rate $\hat{r}$	$A_j = 1$ if decision j accepts; $\hat{r} = \sum_j A_j H_j / \sum_j A_j$, defined when $\sum_j A_j > 0$.
Yield; total gain	Accepted updates per 100 decisions; $\sum_j A_j \Delta_{\mathrm{ref},j}$.
$r_{\mathcal{Q}}$	False-improvement rate among accepted updates, $\mathbb{E}[A_j H_j]/\mathbb{E}[A_j]$, on the i.i.d. task population $\mathcal{Q}$.
Proposition 1 bound U	KL upper confidence bound on $r_{\mathcal{Q}}$ at $\delta = 0.05$; valid on $\mathcal{Q}$ only.
KL-ratio bound	Real data: a 95% upper bound on the false-improvement rate for these datasets (not for new datasets), $U_{\mathrm{KL}}(\bar{W};\, J)/L_{\mathrm{KL}}(\bar{A};\, J)$, where U_{KL} and L_{KL} are the one-sided Kullback–Leibler (KL) upper and lower bounds at level $\alpha/2$ (Section III-D) and $W_j = A_j H_j$; conditional on the observed collection.
Leave-one-source-out (LOCO)	Maximum KL-ratio bound after leaving out one source (clinical) or one dataset (non-clinical), to check that no single source or dataset drives the result.
Conservative bound	The same bound (Proposition 1 on $\mathcal{Q}$, KL-ratio on the real data) with every accepted update that is not clearly improving counted as a false improvement.
Safer	A lower false-improvement rate (equivalently, a lower risk among accepted updates).

clamped to 1 *and set to* 1 *when the lower denominator is not positive. The quantities* $\bar{A}$ *and* $\bar{W}$ *are the sample means over the* J *trials of the acceptance indicator* A_j *and of* $W_j = A_j H_j$*. The bounds* U_{KL} *and* L_{KL} *are one-sided Kullback–Leibler bounds [68], [67], with* $\alpha = 0.05$ *split as* $\alpha/2$ *across the two bounds by the union bound. The bound* $\hat{U}$ *is conditional on the observed collection. This bound is reported for (A), (C), (D), and the OpenML collection (E) in Section IV. LOCO and a dataset-level bootstrap bound serve as sensitivity analyses [69], and the box in Section V-A collects the scope of each claim (the estimand* $r_{\mathcal{E}}$ *is defined in Appendix A). The scope of the method depends on the input that the gate requires as well as on the bound, as stated in Remark 3.*

Remark 3 (Interface and evidence scope)**.** *The gate consumes the baseline predictions, the candidate predictions, and the observed labels on POOL, together with a metric of the family defined in Section III-C. Because it uses only these predictions and labels, its interface does not depend on the model type, and the labels it needs are the POOL labels already available at the decision, not new labels. The empirical evidence covers tabular learners (trees, kernels, boosting, and a two-hidden-layer multilayer perceptron, Section IV-G).*

E. EXPERIMENTAL DESIGN: DATA, PARTITION, COMPARATORS, AND EVALUATION

a: Data.

The evidence rests on 76 real datasets or clinical tasks and on a set of synthetic tasks. The evidence is of three kinds: (i) a first evaluation on new data in (A), (B), and (E); (ii) a re-test on data seen before in (C) and (D), where 18 datasets of (D) were used during development and (C) repeats a failed run; and (iii) post hoc analyses (T3, Table H5, and the source-level checks). The data form five collections, and Section IV assigns a distinct role to each collection. (A) The 11 new real datasets supply the main comparison of acceptance rules (Section IV-C). No development run used these datasets. (B) The synthetic population $\mathcal{Q}$ supplies 1,400 independent tasks with reference-measured gain. These tasks use the task seeds of the generator that no pilot run had used (seed allocation and generator in Appendix B). These tasks test the formal bound. A second, fresh batch of 1,400 tasks repeats that test (Section IV-D). (C) The 12 retrospective clinical tasks, drawn from six sources, support the clinical application (Section IV-E). (D) The 39 supporting datasets are 34 binary confirmation datasets, 18 of which were used during development and are re-evaluated with fresh seeds, and 5 multiclass datasets. These datasets form the basis of the conditions-of-use tests. The gate is also run once on all 39 datasets together (Section IV-G). (E) The OpenML collection of 14 new datasets [70] passes through the released program under the same protocol and serves as supporting evidence (Section IV-A). In total, (A), (C), (D), and (E) contain $11 + 12 + 39 + 14 = 76$ real data units. Table B1 lists each unit with its domain, source, size, minority rate, and reference size (Appendix B). Table G1 gives the per-dataset decisions of (E) (Appendix G).

Datasets enter under the pre-specified inclusion and exclusion rules in Appendix B, and a group and temporal audit, de-duplication, and filtering are applied in the loaders before any split, so that rows from the same group or subject do not cross the split. Appendix G reports two further collections. The first is the pre-registered extension (E2), which adds 19 OpenML datasets, some of which carry curation flags (origin not verifiable, artificial data, or flattened images). The second is a later supplement with one dataset, Volkert.

b: Partition and trials.

A trial is one run of the procedure on one task under a fresh seed, and each seed re-draws the partition. The seed fixes the training rows, the SEARCH/POOL split, the bag draws, the folds, and the bootstrap draw. Each trial ends in one decision: accept or reject the update.

Each dataset receives 100 trials, the first count tested that passed the power check of Table 6 in development. That check gave LOCO_0 of 18.2% at 50 trials and 0.8% at 100 (development runs in Table C1). Each clinical task has 200 trials. This count was raised after the clinical run with 100 trials failed the per-source LOCO condition (run history in Appendix C). Each multiclass dataset has 60 trials, the count pre-registered for the multiclass confirmation (basis of each setting in Table D2). On (B), each task carries one trial, because every task is a new draw from $\mathcal{Q}$. These counts give 1,100 decisions in (A) (11×100), 1,400 in (B), 2,400 in (C) (12×200), and 3,700 in (D) ($34 \times 100 + 5 \times 60$). The four main settings therefore contain 8,600 decisions, and (E) adds 1,280 (11 binary datasets $\times 100 + 3$ multiclass datasets $\times 60$).

The locked reference and the random state of the learner stay fixed across trials. The reference is set aside once per task with a pre-registered seed, and the conclusion is unchanged under two further reference seeds (run-level numbers in Table C1). The reference, the outer split, and the SEARCH/POOL split are stratified by class. Because trials from the same dataset are correlated, the confidence intervals treat the dataset rather than the trial as the unit [69]; Appendix B lists the training rows and reference sizes per dataset. The experiments ran in Python 3.12 with scikit-learn 1.9.1 and XGBoost 3.4.1 on three CPU machines. A replay on a second machine reproduces the stored trials bit-identically, and Appendix C lists the exceptions and the full software versions.

c: Comparators.

All rules judge the same decisions, with the same candidates, the same base model, the same POOL counts, the same reference, and the same fixed learner configuration for every rule, so the only difference is the accept/reject rule. The test-based rules apply the statistics of the classical tests to fixed predictions on POOL splits. Because the base models are not refitted in each split, the 5×2 cv and corrected t statistics follow the original formulas [22], [23] but not the original retraining procedure. The comparator set comprises the naive rule, an honest holdout, a fixed validation margin, and a random control. The same set also includes three acceptance rules built on classical test statistics: the McNemar test, the 5×2 cv test, and the corrected resampled t-test, also with per-split re-selection. Table 6 defines each comparator and states the pre-registered pass criterion.

d: Evaluation metrics.

Each decision j is evaluated on the locked reference R_j, by comparing the deployed configuration c_j with the baseline b_j. This comparison produces the metric difference $\Delta_{\mathrm{ref},j}$ and the false-improvement indicator H_j of Definition 1. The main quantity is the false-improvement rate among accepted updates, $\hat{r} = \left(\sum_j A_j H_j\right) / \left(\sum_j A_j\right)$. Its upper bound is U (12) on $\mathcal{Q}$ and $\hat{U}$ (13) on the real data. Two further bounds check the real-data result: LOCO and a dataset-level bootstrap interval. LOCO is the maximum $\hat{U}$ obtained when one dataset (one source for the clinical tasks) is left out in turn. This bound indicates whether a single dataset drives the result. The dataset-level bootstrap interval takes datasets rather than trials as the sampling unit.

The main metric is weighted-F_1, fixed in advance as a choice of the study. That choice is not a claim that weighted-F_1 is better than balanced accuracy or macro-F_1. The comparison is repeated with those two metrics, and Table H6 reports the results by metric, size, and class count. For the clinical tasks the comparison is also repeated with two clinically oriented metrics (results in Section IV-E). The first is F_2 [62], which weighs sensitivity twice as heavily as precision. The second is net benefit (10), which credits true positives against false positives at the clinical threshold p_t. Further quantities are the yield (the number of accepted updates per 100 decisions) and the total gain $\sum_j A_j \Delta_{\mathrm{ref},j}$. For the clinical tasks, the reported quantities also include the clinical consequences of the accepted updates, namely the flagged patients and the extra true and false positives.

e: Statistical analysis.

The pre-registered criterion of Table 6 provides the primary decision of every run. This criterion is a decision rule on an upper bound, and it does not depend on the comparators. A run passes when the 95% upper bound on the false-improvement rate lies below 10%. The bound has error level $\delta = 0.05$. The bound component of the criterion is equivalent to rejecting $H_0 : r \geq 0.10$ in favor of $r < 0.10$ in a one-sided test at level 0.05. A pass requires every condition: the bound, LOCO, yield, the number of datasets, and LOCO_0. LOCO_0 is a check of precision sufficiency in the no-event scenario, not a general power proof. The quantity r is $r_{\mathcal{Q}}$ on (B) (Proposition 1) and the false-improvement rate of the observed collection on the real data (Remark 2). Three hypothesis tests support the criterion, and each test uses the significance level $\alpha = 0.05$. The first two tests resample datasets rather than trials. The third test is exact, and it is computed on all decisions and again within each dataset.

The first test, T1, compares the false-improvement rate of BRIDGE with that of a comparator rule k on the same decisions,

$$H_0 \colon r_{\mathrm{B}} - r_k = 0 \quad \text{versus} \quad H_1 \colon r_{\mathrm{B}} - r_k \neq 0, \tag{14}$$

where r_{B} and r_k denote the false-improvement rates among the updates accepted by BRIDGE and by rule k, respectively. The test is a paired cluster bootstrap [61], [69] with 10,000 resamples of the datasets (of the tasks in (B) and (C)). Both rules are evaluated on the same resample. The 95% confidence interval (CI) is the percentile interval of $\hat{r}_{\mathrm{B}} - \hat{r}_k$.

TABLE 6. Comparator rules and the pre-registered criterion.

Symbol / term	Definition
Naive rule	Accept whenever the configuration selected on the whole of POOL improves the validation metric, i.e., (1) with a positive gain.
McNemar	Exact one-sided binomial test on discordant pairs (candidate correct and baseline wrong, or the reverse) on POOL; tests per-example accuracy.
5×2 cv	Dietterich's paired t statistic, 5 d.f., on the fixed predictions of the final configuration over five 2-fold splits of POOL, without refitting.
Corrected t	Ten $2/3$–$1/3$ splits of POOL, Nadeau–Bengio variance correction, 9 d.f., final configuration; with re-selection: the candidate is re-chosen inside each split.
Fixed margin t	Accept when the validation gain on POOL is at least t; $t^* = 0.05$ pre-registered from development, other t post hoc.
Random control	Accept at random with BRIDGE's yield.
Dataset-level bootstrap	11 datasets resampled with replacement 10,000 times; 95th percentile of the pooled false-improvement rate.
Pre-registered criterion	Real data: pass if and only if (i) KL-ratio bound $< 10\%$, (ii) LOCO $< 10\%$, (iii) yield ≥ 5 per 100, (iv) at least 3 datasets (clinical: tasks) with yield ≥ 5 per 100, and (v) on the non-clinical runs, the power check $\mathrm{LOCO}_0 < 10\%$, the LOCO bound recomputed with zero false improvements. On $\mathcal{Q}$: Proposition 1 bound $< 10\%$ and yield ≥ 5 per 100.

The two-sided p is twice the smaller share of resampled differences on either side of zero. Within one collection, the p-values of all comparators are adjusted by the Holm procedure [71], which controls the family-wise error rate at α.

The second test, T2, asks whether a variant of the gate matches the default gate in false-improvement rate. Two one-sided tests (TOST) [72] are applied to the hypotheses

$$H_0 \colon |r_{\mathrm{B}} - r_k| \geq \eta \quad \text{versus} \quad H_1 \colon |r_{\mathrm{B}} - r_k| < \eta, \qquad (15)$$

using the equivalence margin $\eta = 2$ percentage points, the same cluster bootstrap, and the Holm adjustment over the variants compared in one run. The equivalence margin η was fixed in the pre-registration of the first equivalence test (run MINI-44) before any equivalence result was computed (Appendix C). Every run is labeled MINI-n and listed with its pre-registration in Table C2.

The third test, T3, asks whether BRIDGE attains a lower false-improvement rate by selecting updates rather than by accepting fewer updates. The null hypothesis H_0 of this test states that the n_{B} updates accepted by BRIDGE form a random draw of the same size from the n_{N} updates accepted by the naive rule. Of these n_{N} updates, h_{N} are false improvements. Under H_0 the number of false improvements h among the draw follows a hypergeometric distribution, and the one-sided p is the probability of at most the observed count h_{B},

$$p = \Pr\left(h \leq h_{\mathrm{B}}\right), \qquad h \sim \mathrm{Hypergeometric}\left(n_{\mathrm{N}},\, h_{\mathrm{N}},\, n_{\mathrm{B}}\right). \qquad (16)$$

A paired t-test on per-trial differences is a common choice for comparing classifiers [22], but it does not fit these comparisons for two reasons. First, the false-improvement rate is a ratio over accepted updates, and the number of accepted updates differs between rules. Second, the trials of one dataset are not independent. T1 was pre-registered for (C) and (D) and was applied afterwards to the stored runs of (A), (B), and (E). T2 was specified before its computation, and T3 is a post hoc analysis of the stored runs.

TABLE 7. Summary of main results: the naive rule against BRIDGE, (A)–(E).

Setting	Naive false % ↓	BRIDGE false % ↓	Reduction (pp) ↑	False removed % ↑	Gain per update ↑
(A) New datasets	19.1	1.1	18.0	99	2.5×
(B) Synthetic	20.7	1.7	19.0	97	2.1×
(C) Clinical	20.0	0.0	20.0	100	1.8×
(D) Supporting	12.5	0.5	12.0	99	2.1×
(E) OpenML†	25.6	1.0	24.6	99	3.1×

*False %: false improvements among the accepted updates, pooled over all decisions of a setting (one fresh seed per decision). Reduction: naive rule minus BRIDGE, in percentage points; all 95% intervals lie above zero. †(E): 14 additional OpenML datasets, run through the released program under the same protocol and reported as supporting evidence; the class-count subgroups, the extension (E2), and the dataset amendment appear in Appendix G. False removed: the false improvements of the naive rule that BRIDGE rejects (Table H4). Gain per update: BRIDGE relative to the naive rule. Definitions, intervals, and bounds: Table H1. Every setting: current model HGB with the fixed configuration; candidates from the default space built from HGB, DT and SVC (173 binary, $14 + 8K$ multiclass; Section III-A); other learners and ensemble members: Table H8. ↓/↑ = smaller/larger is better.

IV. RESULTS

Every confirmatory number reported in this section comes from a pre-registered run. The code was hashed before the run, and the seeds were drawn after pre-registration (chronology in Appendix C). A run passes when the pre-registered criterion of Section III-E is met (Table 6). This criterion is a 95% upper bound on the false-improvement rate below 10%. The 10% level was fixed in the first pre-registered run and stayed unchanged in every later run. The run list in Table C2 (Appendix C) records the criterion of each run. In every setting, the level lies below the false-improvement rate of the naive rule, which ranges from 12.5% to 25.6% (Fig. 1).

A. SUMMARY OF MAIN RESULTS

Table 7 summarizes the main results. In every collection, BRIDGE lowers the false-improvement rate among accepted updates by 12.0 to 24.6 percentage points relative to the naive rule, to at most 1.7%. Every 95% interval of the reduction over datasets or tasks lies above zero. The paired test (14) gives $p < 0.04$ in every collection, after the Holm correction wherever several rules are compared. BRIDGE removes 97% to 100% of the false improvements that the naive rule accepts (Table H4). The mean reference gain per update

that BRIDGE deploys is 1.8 to 3.1 times that of the naive rule (Table H4). All five collections meet the pre-registered criterion. Table H1 (Appendix H) lists the intervals, bounds, and per-collection counts.

The subsections that follow report the hypothesis tests (Section IV-B), the comparison of acceptance rules (Section IV-C), the formal bound (Section IV-D), the clinical application (Section IV-E), the component analysis (Section IV-F), and the conditions of use (Section IV-G). The OpenML collection (E) provides supporting evidence outside the four main settings, and Appendix G reports its class-count subgroups, the extension (E2), and the amendment of the (E2) dataset list.

B. HYPOTHESIS TESTS

The hypothesis tests T1 to T3 of Section III-E check whether the differences of Section IV-A could arise by chance. Questions 1 to 3 ask whether BRIDGE accepts a smaller share of false improvements than the naive rule, the test-statistic-based acceptance rules, and the rules that separate choosing from judging (the corrected t-test with per-split re-selection and the two holdouts). Question 4 asks whether the margin lowers the false-improvement rate, comparing BRIDGE with BRIDGE without the margin. These four questions use test T1, whose null hypothesis (14) states equal false-improvement rates.

Question 5 uses test T2, whose null hypothesis (15) states that the false-improvement rate at a margin between 0.0025 and 0.01 differs from that at the default margin 0.005 by at least 2 percentage points. Question 6 uses test T3, whose null hypothesis (16) states that the updates accepted by BRIDGE are a random draw from the acceptances of the naive rule. Table 8 gives the result and decision for each question, and Table H2 (Appendix H) gives the estimates, intervals, and p-values.

For questions 1 and 2, the naive rule and the test-statistic-based acceptance rules accept a larger share of false improvements than BRIDGE. The null hypothesis of equal rates is rejected against the naive rule in all five collections, against these test-based rules in (A), (B), and (D), and against the 5×2 cv test in (C). For question 3, the rules that also separate choosing from judging have a low false-improvement rate as well (0.0% to 3.0%), consistent with the source named in Section I: choosing and judging on the same data. Test T1 detects a difference in (D), where BRIDGE is 1.0 percentage point lower. Equal rates are a separate question from a difference test, so Section IV-C compares these rules by the pass criterion and by the number of accepted updates.

For question 4, the margin lowers the false-improvement rate by 1.1 to 1.8 percentage points in (A), (B), and (D), within the 2-point equivalence margin of question 5. On the binary datasets of (E), it lowers the rate by 7.7 points, from 8.9% without the margin to 1.2% with it (descriptive; Section IV-F). For question 5, a margin between 0.0025 and 0.01 changes the rate by less than 2 percentage points.

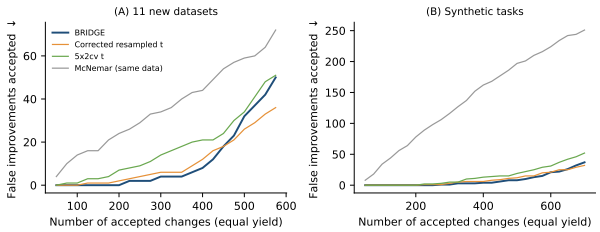

FIGURE 3. Acceptance rules compared at matched acceptance counts on (A) and (B).

Note. The curves rank the decisions by each rule's own statistic ($-p$ for the tests, the bootstrap lower bound for BRIDGE), keep the top N (ties broken by stored decision order), and count the false improvements among the selected trials; the metric is weighted-F_1. Runs on (A) and (B) store the decision statistic of every rule, so the curves use those collections. Runs on (C) and (D) store the accept/reject decision of each rule, and the program used for (E) records the decisions of the naive rule and BRIDGE. Current model: HGB with the default candidate space (HGB, DT, SVC). Other learners, other candidate members, and a joint seven-family space are tested separately (Tables H11 and H8).

For question 6, a random draw of the same size from the naive acceptances contains 49 to 240 false improvements, against 0 to 8 for BRIDGE, so H_0 is rejected in every collection. Within single datasets, H_0 is also rejected in (D) and (E), so BRIDGE selects among the trials of one dataset there as well. In (A) and (C), the selection is shown across datasets ($p = 0.15$ and 0.45 within datasets; (B) has one trial per task). This analysis is post hoc and uses the stored runs.

C. COMPARISON OF ACCEPTANCE RULES

This subsection compares BRIDGE with the naive rule, three acceptance rules built on classical test statistics, the corrected resampled t-test with per-split re-selection, and an honest holdout. The decisive comparators are the rules that also separate selection from evaluation, the same-POOL cell, and the fixed margin; the naive rule and the test-based rules are not. Table H5 gives the operating point of each rule, and a fixed margin chosen on six settings and applied to the seventh passes six of the seven (Appendix E). All rules are judged on the same decisions by the false-improvement rate among the updates that each rule accepts. The test-based rules and the holdout were run on the four main settings, (A) to (D). The OpenML collection (E) was run through the released program, which records the decisions of the naive rule and of BRIDGE. The collection (E) therefore enters this comparison through the naive rule and the random-selection control. Table 9 gives the false-improvement rate of every rule in every collection. Appendix H provides the detailed counts and bounds.

On the 11 new datasets of (A), the test-based rules fail the criterion, whereas BRIDGE meets it. The counts and bounds are in Table H3, and Table 9 gives the false-improvement rates. These tests assume that the compared pair is fixed before the data are seen. In update acceptance, however, the candidate is chosen among many configurations on the same data that then judge it (Section II-B). The failure is specific to (A). In (C) and (D), the 5×2 cv and corrected t-tests also meet the criterion. BRIDGE differs from the naive rule

TABLE 8. Hypothesis tests: results and decisions on questions 1–6.

No.	Test (H_0)	Result	Decision on H_0 ($\alpha = 0.05$) and meaning
1	T1 (14)	BRIDGE lower by 12.0 to 24.6 pp ($p < 0.04$)	Rejected in (A)–(E): BRIDGE accepts fewer false improvements than the naive rule in every collection
2	T1 (14)	BRIDGE lower by 2.0 to 33.1 pp wherever H_0 is rejected	Rejected in (A), (B), (D), and against the 5×2 cv test in (C): BRIDGE accepts fewer false improvements
3	T1 (14)	These rules have 0.0% to 3.0% false improvements; BRIDGE is 1.0 pp lower than the re-selection test in (D)	Not rejected (except that test in (D)): these rules share the principle of BRIDGE, separating the data that choose from the data that judge; BRIDGE accepts 1.1 to 1.9 times as many updates as the holdouts (Section IV-C)
4	T1 (14)	BRIDGE lower by 1.1 to 1.8 pp in (A), (B), (D); no difference in (C)	Rejected in (A), (B), (D): the margin lowers the false-improvement rate; in (C) both versions accept almost no false improvement
5	T2 (15)	All 12 comparisons within ± 2 pp ($p \leq 0.0498$)	Rejected in (A)–(D): rejecting this H_0 means that the margins are equivalent, so the false-improvement rate is stable across margins in this range
6	T3 (16), post hoc	Random draw 49 to 240 false improvements, BRIDGE 0 to 8 ($p < 10^{-23}$)	Rejected in (A)–(E): BRIDGE picks fewer false improvements than a random choice, also within single datasets in (D) and (E)

Questions 1–6: Section IV-B. T1–T3 and the null hypotheses: Section III-E. The test, not BRIDGE, decides on H_0. Rejecting H_0 means that the data show the effect named in the question (in question 5, that the margins give equivalent results). Not rejecting means that no effect is shown, which is not proof of equality. In (C), the pre-registered unit is the task (12 clusters, $p = 0.004$); over the six sources (6 clusters), the difference does not differ significantly ($p = 0.17$). pp = percentage points. p: Holm-adjusted wherever several rules are compared within one collection. Estimates, 95% intervals, and p-values of every comparison in every collection: Table H2 (Appendix H).

TABLE 9. Acceptance rules on collections (A)–(E): false-improvement rate of the accepted updates (%).

Rule	(A)	(B)	(C)	(D)	(E)
Naive (validation improves)	19.1* ×	20.7* ×	20.0* ×	12.5* ×	25.6* ×
McNemar, same POOL	12.0* ×	34.8* ×	0.0° ×	6.7* ×	–
McNemar, cross-fit	8.9 ×	26.0 ×	0.0° ×	2.5* ✓	–
5×2 cv paired t	6.3* ×	5.6* ✓	2.5* ✓	3.6* ✓	–
Corrected resampled t	6.3* ×	5.7* ✓	2.6° ✓	4.2* ✓	–
Corrected t, re-selection	2.3° ✓	3.0° ✓	0.4° ✓	1.5* ✓	–
BRIDGE without margin	2.8* ✓	3.5* ✓	0.6° ✓	1.6* ✓	–
Holdout 2/3 + 1/3	1.3	1.9 ✓	0.0° ✓	0.3° ✓	–
Holdout 1/2 + 1/2	3.2	2.1 ✓	0.0° ✓	0.8° ✓	–
BRIDGE	1.1 ✓	1.7 ✓	0.0 ✓	0.5 ✓	1.0 ✓

✓/×: pre-registered criterion met/not met (Table 6). *: BRIDGE has a significantly lower false-improvement rate than the rule; °: no significant difference (paired test (14), $\alpha = 0.05$, Holm over the comparators of one collection; in (E), a single comparison, $p < 0.0002$). No superscript: paired test not run. No mark: criterion not evaluated, because LOCO is not computed for the holdouts in (A). –: not run, because the program used for (E) records the decisions of the naive rule and BRIDGE. Accepted counts, bounds, LOCO, yield, and updates above the margin of each rule: Table H3; for (E): Table H1.

mainly on small validation gains. Of the 459 changes of (A) with a validation gain of at most 0.01, BRIDGE accepts none, whereas the naive rule accepts all 459 and is wrong 164 times (Table H3).

The corrected t-test with per-split re-selection is the closest comparator, because it also separates choosing from judging. On (A), this test meets the criterion with a false-improvement rate of 2.3%, against 1.1% for BRIDGE. In (A) and (B), the difference is not significant; in (A) it is -1.2 pp (95% CI $[-6.3; 1.3]$, Holm-adjusted $p = 0.32$; paired differences in Table E2). In (A), the same-POOL cell is equivalent at matched yield by TOST (post hoc), and the fixed margin chosen post hoc does not differ significantly either.

The rules accept different numbers of updates, so the comparison is also made at matched acceptance counts on (A) and (B) (Fig. 3). BRIDGE has an equal or lower false-improvement rate than the corrected t-test up to about 450 accepted updates, and a slightly higher proportion beyond that point. The matched-count claim therefore covers that range (matched-yield details in Appendix D).

A last comparator is an honest holdout on the same labeled data. This rule reserves part of POOL to test the chosen update, whereas BRIDGE reuses all of POOL through cross-fitting. On (A) to (D), BRIDGE accepts 1.1 to 1.9 times as many updates as the two holdout splits. BRIDGE deploys 1.1 to 1.9 times as many updates that gain more than the margin (post hoc, Table H5, Holm $p \leq 0.01$), and the share of accepted updates that gain at most the margin does not differ significantly (accepted counts in Table H3). The false-improvement rate of BRIDGE is close to that of the holdouts. In (A), the false-improvement rate is 1.1% for BRIDGE against 1.3% and 3.2% for the holdouts. In (B), it is 1.7% against 1.9% and 2.1% (descriptive, without a paired test). In (C), neither BRIDGE nor the holdouts accept a false improvement. In (D), the difference between BRIDGE and either holdout is not significant (Table 9).

a: Updates below the margin.

Two complementary measures use the margin $\varepsilon = 0.005$, which was fixed during development (Table 2), and both are post hoc (Table H5). The first measure is the share of accepted updates whose gain on the reference is at most ε, and the second is the number that gain more than ε per 100 decisions. Among the updates that the naive rule accepts, 33.9% to 57.4% gain at most ε, against 0.2% to 6.3% for BRIDGE, and the difference is significant in every collection (Holm $p \leq 0.005$). The two holdouts reach 13.4 to 23.3 in (A), (B), and (D) and 41.0 to 46.0 in (C), against 24.9, 32.8, 29.1, and 49.9 for BRIDGE, so BRIDGE deploys significantly more updates above the margin. No rule has a significantly smaller share below the margin than BRIDGE. In (A) and (B), the corrected t-test with per-split re-selection does not differ significantly from BRIDGE (6.0% and 5.3%

TABLE 10. Strongest comparators: false-improvement rate, yield, and updates above the margin, (A)–(D).

	(A) 11 new datasets			(B) Synthetic			(C) Clinical			(D) Supporting		
Rule	False % ↓	Yield ↑	$>\varepsilon$ ↑	False % ↓	Yield ↑	$>\varepsilon$ ↑	False % ↓	Yield ↑	$>\varepsilon$ ↑	False % ↓	Yield ↑	$>\varepsilon$ ↑
Corrected t, re-selection	2.3° ✓	31.7	29.8	3.0° ✓	37.9	35.9	0.4° ✓	58.4	56.4	1.5* ✓	39.6	35.2
Holdout 2/3 + 1/3	1.3	13.9	13.4	1.9 ✓	18.6	18.0	0.0° ✓	41.1	41.0	0.3° ✓	19.5	18.8
Holdout 1/2 + 1/2	3.2	17.3	16.2	2.1 ✓	24.2	23.3	0.0° ✓	46.0	46.0	0.8° ✓	22.3	21.2
Fixed margin, dataset held out†	5.8 ×	31.5	28.7	4.2 ✓	49.4	45.4	–	–	–	–	–	–
BRIDGE	1.1 ✓	25.8	24.9	1.7 ✓	33.9	32.8	0.0 ✓	50.0	49.9	0.5 ✓	30.1	29.1

False %: false improvements among the accepted updates. Yield: accepted updates per 100 decisions. $>\varepsilon$: accepted updates whose reference gain exceeds the margin $\varepsilon = 0.005$, per 100 decisions (post hoc; tests in Table H5). ✓/×: pre-registered criterion met/not met; no mark: LOCO not computed for the holdouts in (A). *: BRIDGE has a significantly lower false-improvement rate; °: no significant difference (paired cluster bootstrap, Holm over the comparators of a collection, Table 9). †Post hoc: for each dataset (in (B), each block of 140 tasks), the margin is chosen from the pre-registered grid 0.005–0.050 on all other data, by the rule of the leave-one-setting-out analysis, and applied to the held-out dataset; the margin 0.020 is chosen for 43 of 48 held-out units. No test is applied: the comparison is descriptive, because each held-out dataset's margin is chosen on data that include the other evaluated datasets, so the decisions are not independent of that choice. Not available for (C) and (D). Cost: BRIDGE re-scores stored predictions (about 0.03 s per run; Table D8).

TABLE 11. Finite-sample bound on the synthetic tasks (B) for each acceptance rule.

Rule	Accepted	False	Bound % ↓	Yield/100 ↑	Criterion
Naive (validation improves)	1,323	274	23.52	94.50	×
McNemar on same POOL	135	47	45.19	9.64	×
McNemar on cross-fit	50	13	42.65	3.57	×
5×2 cv paired t	610	34	8.14	43.57	✓
Corrected resampled t	774	44	7.95	55.29	✓
Corrected t, re-selection	531	16	5.19	37.93	✓
BRIDGE without margin	598	21	5.67	42.71	✓
BRIDGE	475	8	**3.56**	33.93	✓

*Setting (B): 1,400 synthetic tasks drawn independently from one population, one decision per task. Bound: Proposition 1. Conservative bounds: Table C2. ↓/↑ = smaller/larger is better; **bold** = best value per column, naive acceptance excluded; bound: one-sided 95% upper bound on the false-improvement rate, the uncertainty measure of each rule; criterion: Table 6.

against 3.5% and 3.4%), and neither do the two holdouts. BRIDGE keeps 34% to 80% of the above-margin updates that the naive rule accepts, from 24.9 against 59.5 per 100 decisions in (A) to 49.9 against 62.4 in (C). The rest are rejected because the POOL evidence is not yet sufficient.

b: Strongest comparators.

Table 10 lists the rules that also guard against selection. In (A) to (D), BRIDGE and the corrected t-test with re-selection meet the criterion, and the two holdouts deploy fewer updates above the margin than BRIDGE. A fixed margin selected without the evaluated dataset misses the criterion in (A), and its false-improvement rate in (B) is higher (4.2% against 1.7%); the comparison is post hoc and descriptive.

D. EMPIRICAL CONSISTENCY WITH THE FINITE-SAMPLE BOUND (SYNTHETIC POPULATION $\mathcal{Q}$)

The finite-sample bound of Proposition 1 is checked empirically on (B). This collection is a synthetic population of 1,400 binary tasks drawn independently from one generative population, with reference-measured gain. Each task has one trial and a 100,000-row reference, which is large enough that the error of the reference truth is about a quarter of the margin. The generator and the basis of these sizes are in Appendix B. The population is designed so that the naive rule is wrong about as often as on the real data: 20.7% of its accepted updates are false improvements, against 19.1% in (A) (Table 11). BRIDGE accepts 475 updates, of which 8 are false improvements, so $\hat{r} = 1.7\%$. The Proposition 1 bound on $r_{\mathcal{Q}}$ (11) is $U = 3.56\%$ (Table 11).

The bound is also computed on a fresh batch of 1,400 tasks that no earlier run used. With the same frozen gate, BRIDGE accepts 462 updates with 7 false improvements and gives $U = 3.35\%$ (base-learner results on $\mathcal{Q}$ in Table H10, row HGB).

A consistency check of the bound needs two batches, because within one batch $U \geq \hat{r} = S/M$ holds by construction (12). The bound of a batch therefore lies above the estimate from the same batch, and that comparison carries no evidence about the bound. The check is therefore made across the two independent batches. Proposition 1 gives $\Pr(r_{\mathcal{Q}} \leq U^{(b)}) \geq 1-\delta$ for each batch b, so the bound of one batch should lie above the false-improvement rate estimated on the other, $\hat{r}^{(2)} \leq U^{(1)}$ and $\hat{r}^{(1)} \leq U^{(2)}$. Both inequalities hold. The bound of the primary batch (3.56%) lies above the estimate of the fresh batch (1.5%). The bound of the fresh batch (3.35%) lies above the estimate of the primary batch (1.7%; Tables 11 and H10). Both bounds also lie above the upper end of the 95% Clopper–Pearson interval [73] for $r_{\mathcal{Q}}$ from the two batches together. That interval is 0.9% to 2.6% around an estimate of 1.6% (15 of 937). The 95% coverage itself rests on Proposition 1; the cross-batch comparison is a consistency check, not a validation of coverage.

E. CLINICAL APPLICATION (RETROSPECTIVE)

The clinical application is retrospective: the data come from existing records, not from a prospective trial, so the results show what would happen if the gate were used on that data (prospective validation: Section V-A). It is an illustration with one hypothetical action and one cost ratio at $p_t = 0.2$ for all tasks, not evidence of practical clinical benefit. Table B2 describes the 12 clinical tasks, which come from 6 independent sources. The sources are screening mammography from OpenML [70], an HIV-1 trial cohort from the UCI repository [74], and self-reported diagnoses from four NHANES [75] and four BRFSS [76] surveys. Two further sources are a national COVID-19 cohort from the open data of the Mexican

TABLE 12. Clinical application (C): results by metric.

Metric	Naive false impr. % ↓	BRIDGE accepted	BRIDGE false impr. % ↓	Bound % ↓	LOCO-source % ↓
Co-primary (pre-registered)					
Net benefit, $p_t = 0.2$	20.0	1,199	0.0	0.33	0.50
Balanced accuracy	0.3	2,273	0.4	0.82	1.25
Secondary					
Weighted F_1	23.6	435	2.1	5.12	7.84
F_2	0.0	2,379	0.0	0.16	0.24
Net benefit, $p_t = 0.1$	3.9	1,660	0.0	0.23	0.37

*Setting: 12 tasks from 6 independent sources, 200 trials per task (2,400 decisions per metric). Co-primary metrics, fixed in the pre-registration: net benefit at the clinical threshold $p_t = 0.2$ and balanced accuracy. Net benefit at the lower threshold $p_t = 0.1$ is a secondary sensitivity check. Bound and LOCO-source of BRIDGE. Clinical consequences, in extra true and false positives: Table H14. Criterion: Table 6. ↓/↑ = smaller/larger is better; bound: one-sided 95% upper bound on the false-improvement rate, the uncertainty measure of each rule; criterion: Table 6.

Ministry of Health [77] and an intensive care unit (ICU) admissions cohort from the WiDS Datathon 2020 [78].

Survey tasks are health-related cross-sectional classification tasks rather than prognosis tasks, and the predictors are restricted to information available at decision time. One hypothetical clinical action, such as a referral, is paired with each task, and the net-benefit calculation uses this action. The patient-identity audit of the sources (de-duplication and per-person filtering before any split) is reported in Table F1.

The confirmatory clinical run uses 200 trials per task and two co-primary metrics fixed in the pre-registration. The first is net benefit at $p_t = 0.2$ (10), under which one true positive is worth four false positives because the odds $0.2/0.8$ equal $1/4$ [34], and the second is balanced accuracy. Weighted-F_1, F_2, and net benefit at the lower threshold $p_t = 0.1$ are secondary metrics. Table 12 reports these metrics, and the pre-registration chronology is in Appendix C. In the clinical application and the binary collections, the baseline is the pre-registered HGB at threshold 0.5, and the candidate space contains thresholds from 0.05 to 0.95. A shift toward the clinical threshold is therefore a legitimate update, and part of the gain comes from correcting a baseline that was not set for the clinical objective.

Under net benefit at $p_t = 0.2$, 20.0% of the updates accepted by the naive rule are false improvements, whereas none of the 1,199 updates accepted by BRIDGE is. The reduction is 20.0 percentage points, with a 95% interval of 4.5 to 40.3 over tasks and 0.0 to 31.9 over the six sources. The task is the pre-registered unit of analysis; the four BRFSS tasks and the four NHANES tasks each share respondents, so the per-source analysis is reported as a sensitivity and is the conservative reading. Over the six sources, the reduction does not differ significantly ($p = 0.17$, interval up to 0.0). The false improvements of the naive rule lie only in BRFSS (282) and NHANES (171), whereas BRIDGE accepts none of its 1,199 updates as false improvements. The share of updates below the margin still differs significantly over sources ($p = 0.03$, post hoc, one unadjusted test). Three quantities answer different questions: the KL-ratio bound (0.33%) is conditional on the collection with the trial as unit, LOCO-source (0.50%) leaves one source out at a time, and the source-level test ($p = 0.17$) is the inference across sources. Under balanced accuracy, F_2, and net benefit at $p_t = 0.1$, the naive rule rarely accepts a false improvement (0.0% to 3.9%), so the gate has little effect on the result. Weighted-F_1 again separates the rules (23.6% against 2.1%; Table 12). On BRFSS kidney disease, all 133 updates accepted by the naive rule are false improvements, whereas BRIDGE accepts none, and the same pattern holds on BRFSS stroke (144 of 151) and NHANES coronary heart disease (97 of 182; Appendix F, Table F2).

Per 1,000 patients, the updates that BRIDGE accepts add 90.9 true positives and 197.0 false positives, so sensitivity rises and specificity falls, and the workload increases from 98 to 386 patients flagged per 1,000 (Table H14). The net-benefit run was pre-registered after an earlier weighted-F_1 run on the same tasks failed the LOCO criterion, so it is a re-test with fresh seeds, not an independent clinical confirmation (Appendix C).

F. COMPONENT ANALYSIS: ABLATION AND CANDIDATE-SPACE SIZE

a: Ablation.

The ablation switches each of the three gate components on and off while the deployed configuration stays fixed. The components are the evidence (cross-fitted counts or the same POOL), the decision statistic (bootstrap lower bound or point estimate), and the margin ($\varepsilon = 0.005$ or 0). Table 13 lists the eight resulting cells for (A) to (E). For (E), the cells cover the 11 binary datasets, which were rerun with every cell recorded (run MINI-52 in Table C2).

The table first shows which component lowers the false-improvement rate most when it is added to the naive rule alone, with the other two components switched off. The bootstrap lower bound alone gives the largest reduction in (A), (B), (D), and (E), at 7.6 to 14.1 percentage points below the naive rule. In (C), it ties with the margin alone (17.3 points each). The margin alone comes second in (A), (D), and (E), whereas cross-fitting alone comes second in (B). The bootstrap lower bound and the margin each lower the false-improvement rate at a cost in yield. Each of the two rejects updates that the naive rule would accept.

The binary datasets of (E), where every cell was recorded, follow the same ordering at a higher false-improvement rate. There, 27.7% of the updates accepted by the naive rule are false improvements. Of the eight cells, BRIDGE has the lowest false-improvement rate (1.2%, bound 6.1%) and comes closest to the criterion. Its LOCO bound is 10.3%, slightly above the 10% limit. BRIDGE without the margin reaches 8.9% (bound 15.6%), so no cell meets the criterion on these datasets (Table 13). On all 14 datasets of (E), BRIDGE meets the criterion, with a false-improvement rate of 1.0% and a bound of 5.4% (Table H1).

Cross-fitting is added on top of the bootstrap lower bound and the margin, and it has a small effect on the false-improvement rate in (A) and (B). In (A), BRIDGE and the

TABLE 13. Component analysis on (A)–(E): ablation of the three components of the gate.

Known as	Cross-fit	Bootstrap bound	Margin	(A)		(B)		(C)		(D)		(E)	
				False % ↓	Bound % ↓	False % ↓	Bound % ↓	False % ↓	Bound % ↓	False % ↓	Bound % ↓	False % ↓	Bound % ↓
Naive	–	–	–	19.08	22.90 ×	20.71	23.52 ×	19.99	22.68 ×	12.49	14.19 ×	27.65	32.13 ×
Fixed margin	–	–	✓	10.20	14.17 ×	16.42	19.19 ×	2.65	4.05 ✓	6.18	7.68 ✓	16.74	22.48 ×
Same-POOL bootstrap	–	✓	–	6.96	10.70 ×	6.61	8.95 ✓	2.67	4.06 ✓	4.87	6.33 ✓	13.73	19.16 ×
Same-POOL cell	–	✓	✓	2.65	5.85 ✓	3.18	5.14 ✓	0.00	0.31 ✓	1.08	2.06 ✓	3.37	8.17 ×‡
Cross-fit, point	✓	–	–	10.81	15.00 ×	12.07	14.72 ×	7.69	9.84 ×	7.28	8.89 ✓	18.98	24.82 ×
Cross-fit, point, margin	✓	–	✓	5.23	9.00 ×	9.49	12.06 ×	0.91	1.86 ✓	3.76	5.19 ✓	9.11	14.86 ×
BRIDGE without margin	✓	✓	–	2.78	6.16 ✓	3.51	5.67 ✓	0.56	1.34 ✓	1.64	2.75 ✓	8.86	15.63 ×
BRIDGE	✓	✓	✓	1.06	4.23 ✓	1.68	3.56 ✓	0.00	0.33 ✓	0.54	1.47 ✓	1.17	6.13 ×‡

*✓/× after the bound: pre-registered criterion met or not met. (E): the 11 binary datasets of the OpenML collection, with all eight cells from MINI-52 (run list in Table C2). ‡: bound below 10% but LOCO above it (same-POOL cell 11.95%, BRIDGE 10.30%). On the full 14 datasets of (E), the four cells computable from the released program give 1.04% for BRIDGE (bound 5.40%, met) and 7.89% for BRIDGE without the margin (bound 13.84%, not met). Each component switched on (✓) or off (–): cross-fitted evidence instead of the same POOL; the 5% bootstrap lower bound instead of the point estimate; the margin 0.005 instead of 0. Same-POOL cell: the rule that chooses and measures on the same POOL with the bootstrap bound and the margin. Every cell judges the same deployed configuration; current model HGB with the default candidate space. Bound: KL-ratio bound in (A), (C), and (D), and Proposition 1 bound in (B). Paired intervals for the same-POOL cell against BRIDGE: Table E2. Counts, LOCO, and yield of each cell are recorded in the code repository (runs MINI-32 and MINI-48, Table C2). Criterion: Table 6. An arrow ↓ or ↑ marks a column where a smaller or a larger value is better.

same-POOL cell with the bootstrap lower bound and the margin do not differ significantly. The difference is -1.6 pp (95% CI $[-5.1; 0.2]$). At matched yield, the two rules are equivalent by two one-sided tests (TOST) [72] of (15): -0.4 pp, 95% CI $[-1.9; 1.0]$, inside ± 2 pp (post hoc). In (B), BRIDGE is 1.5 pp lower than the same-POOL cell at the number of updates that each rule accepts (95% CI $[-2.7; -0.5]$). At matched yield, BRIDGE is 1.1 pp higher (95% CI $[0.2; 2.1]$; Table E2). The two comparisons disagree because the same-POOL cell accepts 660 updates and BRIDGE accepts 475 (paired differences in Table E2). The updates accepted by the same-POOL cell beyond that count carry most of its extra false improvements. In (E), by contrast, the same-POOL cell reaches a false-improvement rate of 3.4% against 1.2% for BRIDGE. The same-POOL cell also fails on LOCO there (12.0%; Table H7).

BRIDGE keeps cross-fitting because cross-fitting separates the data that select an update from the data that judge it. Through this separation, the decision statistic is out-of-sample with respect to selection (Section III-B). The naive estimate of the gain on the full POOL is optimistic by $+0.0034$ on average over the 1,064 decisions of (A) with a selected candidate. By contrast, the cross-fitted estimate is slightly pessimistic (-0.0034). Both errors amount to about two-thirds of the margin (Fig. 4; descriptive run MINI-27 in Table C1). Replacing cross-fitting with a cell found from these results would be selecting on the outcome, the practice this paper targets. The bound of Proposition 1 holds for any frozen gate, whether or not that gate uses cross-fitting. Two pre-registered predictions differ from the results. The bootstrap lower bound lowers the false-improvement rate more than cross-fitting does. On (A) to (D), the same-POOL cell also meets the criterion at the large candidate space (predictions in Appendix C).

b: Candidate-space size.

The second part of the analysis tests whether a larger candidate space increases the false-improvement rate, as the selection bias of (2) predicts. On (A) to (E), the same grids are refined from 173 to 1,825 and 11,249 binary configurations. HGB serves as the current model, with HGB, DT, and SVC as the members (results by candidate-space size in Table H7, Fig. 5). The multiclass datasets of (D) and (E) are left out of this analysis, because no equally dense multiclass grid is pre-registered. As the candidate space grows, the naive rule fails the criterion in every cell. Its false-improvement rate lies between 12.6% and 28.2% (Table H7). On (A) to (D), BRIDGE meets the criterion at every candidate-space size. Its false-improvement rate stays at or below 1.7% (Fig. 5). The estimates for BRIDGE do not fall steadily with size (in (A), 1.06%, 1.45%, and 1.37%; Table H7). On (A) to (D), therefore, a finer grid brings no material increase in false improvements, and the results do not show that a larger space lowers the false-improvement rate.

The binary datasets of (E) mark the range within which this finding holds. On those datasets, BRIDGE meets the criterion at 1,825 configurations. As reported for the ablation, BRIDGE narrowly misses the LOCO limit at 173. At 11,249 configurations the false-improvement rate rises to 3.4% (6

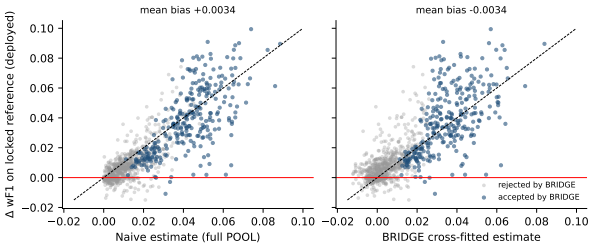

FIGURE 4. Component analysis: naive and cross-fitted gain estimates against the reference truth on (A).

Note. Each point is one of the 1,064 decisions of (A) with a selected candidate. The dashed line marks equality of estimate and reference truth. The naive estimate is measured on the same POOL that chose the update, and it is optimistic on average ($+0.0034$). The cross-fitted estimate, measured on folds not used to choose the update, is slightly pessimistic (-0.0034). Both errors amount to about two-thirds of the margin $\varepsilon = 0.005$. Blue marks updates accepted by BRIDGE, and gray marks rejected updates.

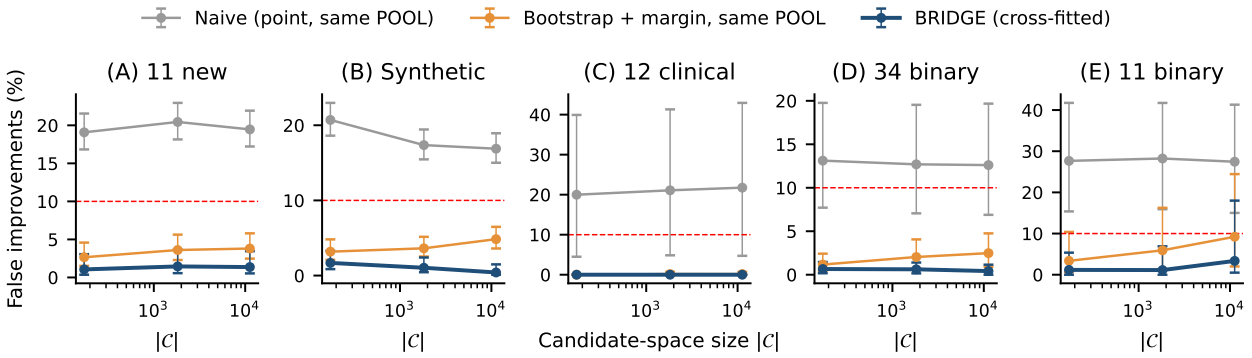

FIGURE 5. Component analysis: false-improvement rate against candidate-space size for (A)–(E).

Note. Error bars show 95% Wilson intervals [79]; the dashed line marks the 10% criterion. (E): the 11 binary datasets. Table H7 gives the grids, the counts, and the datasets of each panel (Appendix H).

of 179 accepted updates; bound 10.3%). This value lies 2.2 percentage points above that of the default space (95% CI 0.04 to 12.8; Table H7). On the reference, each of the six false improvements loses at most 0.006 in weighted-F_1. The same-POOL cell passes at every size on (A) to (D). Its false-improvement rate nevertheless grows in (D) from 1.2% to 2.5% and in (E) from 3.4% to 9.2%. This growth is consistent with the selection bias of (2), which grows with the number of candidates. The default space of 173 configurations therefore remains the tested condition of use, confirmed on (A) to (D) and on all 14 datasets of (E). A denser grid needs its own confirmation before use (Section V-A).

G. SCOPE AND CONDITIONS OF USE

a: Conditions of use.

A condition of use is a condition, such as the metric, the training size, the number of classes, the learner, the tuning, or the imbalance, that has been tested and meets the criterion. The claim holds for each condition tested, one at a time, as listed in Table H8; combinations of conditions that were not tested together lie outside the claim.

b: Metrics, training sizes, multiclass, and feature selection.

These four conditions are tested on a separate candidate space of thresholds for KNN and SVC. Each condition is tested with and without a mutual-information feature-selection filter [80] fitted on SEARCH alone (protocol in Appendix C). The gate meets the criterion for every metric tested (weighted-F_1, balanced accuracy, macro-F_1) and every training size (1,000, 2,430, and 6,000 rows). The gate also meets the criterion for the multiclass tasks (results by metric, size, and class count in Table H6). The gate also meets the criterion on the 39 supporting datasets of (D) taken together (detailed results in Table H1). Feature selection with SVC is the one exception, and Table H8 marks it. The OpenML collection (E) meets the criterion jointly with its extension (E2) (subgroup results in Appendix G).

c: Ensemble members.

Besides DT and SVC, two further pairs of ensemble members were tested. The first pairs logistic regression with naive Bayes, and the second pairs a multilayer perceptron (MLP) with KNN. With both pairs the gate meets the criterion in (A) to (D). A joint space of seven families, built with the same grids, meets the criterion in all five collections (711 binary configurations; composition in Appendix D; Table H11). On the OpenML collection (E), both pairs miss the criterion. Logistic regression with naive Bayes gives a bound of 7.35% and a LOCO of 10.16%, and the MLP with KNN gives a bound of 12.94% and a LOCO of 20.77% (Tables H11 and C1). A post hoc check identifies the cause. Each pair accepts 6 false improvements, against 2 for DT and SVC, and the MLP with KNN accepts fewer updates in total (147 against 192). Of the 12 false improvements, 11 lie on datasets with references of 1,125 to 2,112 rows, where the reference interval of the gain contains zero (Table C1). On (D) the joint space is practically equivalent to the default space, while on (A) it is not. The gate meets the criterion with every member set tested in (A) to (D), so the result does not rest on the HGB–DT–SVC members alone.

The default space is retained for three reasons. First, it was frozen during development, before the confirmatory runs; adopting the joint space after these results would be selecting on the outcome, the practice this paper targets. Second, a larger space increases the selection bias of (2): in (A), the false-improvement rate of the naive rule rises from 19.1% to 22.4%, while BRIDGE meets the criterion with the joint space (1.6% in (A), 0.7% in (D); Tables H11 and C1). Third, the joint space is a candidate for the next development cycle, to be developed first and confirmed on data not used before (Section V-A). DT and SVC were fixed during development, before this study.

d: Base learners.

The base learner was varied one factor at a time on fresh synthetic tasks (base-learner results on $\mathcal{Q}$ in Table H10)

and on the real data (base-learner results on real data in Table H11). XGBoost [60] and an MLP under early stopping [81] meet the criterion in all five collections (Table H11). The condition of use is therefore a base learner with iteration control, and Table H8 marks the variant without iteration control. The gate criterion and the ranking of absolute final quality are different questions: an MLP with early stopping has a lower final weighted-F_1 than HGB in (D) and yet meets the criterion (final model quality in Table H12).

The base-learner results can also be read through the skill of a base learner. Skill is 1 minus the ratio of the log-loss of the learner to the log-loss of a forecast based on the class prevalence [82]. Skill is computed on POOL, without reference data, so it is known before the gate is used. Across all 11 base-learner settings tested on $\mathcal{Q}$, higher skill is associated with a lower gate bound. The Spearman rank correlation [83], which runs from -1 to $+1$, is $\rho = -0.74$ ($p = 0.010$) over the full set and $\rho = -0.82$ for the nine HGB and XGBoost settings (run-level numbers in Table C1). The association has an exception and is not a decision rule: the MLP trained for 300 epochs without early stopping has a skill of -0.67, below every other setting, yet its bound of 4.1% meets the criterion (Tables H10 and E3). The association is descriptive rather than causal, because the settings differ in more than skill.

e: Tuning as an update.

Tuning is treated as an update proposed to the gate rather than as a way to set the baseline. During development, HGB was tuned with Optuna [18] using its tree-structured Parzen estimator sampler [84]. In the confirmatory runs, four learners (HGB, DT, SVC, KNN) were tuned in (A), in (B), and on three multiclass datasets of (D) (tuning results in Table H13). BRIDGE lowers the false-improvement rate in all 12 cells tested, and it meets the formal criterion in two of the 12 cells (tuning results in Table H13; status in Table H8). Tuning can therefore be proposed to the gate as an update. The default procedure nevertheless retains the fixed configuration.

Tuning the baseline in SEARCH raises its weighted-F_1 by about 0.3 points (Table D6), so keeping the fixed configuration has a small cost, and the final model quality after the gate is practically unchanged. The naive rule accepts false improvements more often (23.2% in (A), 15.3% in (D), against 19.1% and 12.5% without tuning), and BRIDGE again lowers the false-improvement rate by 14–21 pp (Table D6).

f: Imbalanced classes.

Five large binary datasets were re-sampled to minority prevalences of 20%, 10%, and 5% before the outer split. POOL then holds about 360, 180, and 90 minority cases, which are 20%, 10%, and 5% of the 1,800 POOL rows (imbalance results in Table H9). Under balanced accuracy and macro-F_1 the gate meets the criterion at every level down to 5% (imbalance results in Table H9). On imbalanced data, the condition of use is therefore a minority-sensitive metric, as Table H8 marks.

g: Parameter sensitivity and cost.

In the parameter sensitivity analysis, the criterion is met on $\mathcal{Q}$ when each of the fold count, the bootstrap count, and the bag count is varied in turn. The combined $q \times \varepsilon$ grid meets the criterion in all 24 real-data cells (3 quantile levels $\times$ 4 margins in (A) and in (D); quantile-margin grid in Table D5). Other margins give a false-improvement rate equivalent to the default margin 0.005 in (A) to (D) (margin comparison in Table D4). Variants of the bag fraction and of the SEARCH size pass in (D) and partly in (A), as Table H8 marks (details in Appendix D). The gate's own cost is about 0.03 s; the total cost is dominated by fitting the bagged base models, and the 16-bag setting cuts this fitting cost by about half (cost details in Appendix D).

V. CONCLUSION

This paper asked how often a model update that raises the validation score is a false improvement on data that did not choose it. In response, the paper proposed BRIDGE, an acceptance gate that decides each update from the labels already available, with settings fixed once. BRIDGE lets one part of the labeled data choose the update and another part judge it (cross-fitting). The gate accepts the update when the lower bootstrap bound of its gain exceeds a margin.

Across every collection tested, BRIDGE lowers the false-improvement rate among the updates it accepts. In the four main settings, the false-improvement rate falls from 19.1% to 1.1% in (A), from 20.7% to 1.7% in (B), from 20.0% to 0% in (C), and from 12.5% to 0.5% in (D) (Table 7). The OpenML collection (E) shows the same pattern, with the false-improvement rate falling from 25.6% to 1.0% (Table 7). The hypothesis tests reject equal false-improvement rates against the naive rule in all five collections, with (C) significant at the task level and not at the coarser source level. A random draw of the same size from the naive acceptances also contains more false improvements than the updates BRIDGE accepts (Section IV-B). On the synthetic population, the finite-sample bound gives a 95% upper bound of 3.56% on the false-improvement rate of BRIDGE. A second batch of 1,400 tasks gives 3.35% (Section IV-D).

BRIDGE removes 97% to 100% of the false improvements that the naive rule accepts, and its mean reference gain per deployed update is 1.8 to 3.1 times that of the naive rule (Table H4). Against an honest holdout on the same labeled data, BRIDGE deploys 1.1 to 1.9 times as many updates that gain more than the margin, and the share that gains at most the margin does not differ significantly (post hoc, Table H5; Table H3).

The finding also holds beyond the default learner, metric, and ensemble members. With XGBoost, with a multilayer perceptron under early stopping, and with the joint seven-family candidate space, the pre-registered criterion is met in all five collections. The criterion is also met in (A) to (D) with both alternative member pairs, and on imbalanced data with balanced accuracy and macro-F_1 down to a minority prevalence of 5% (Table H8). Among the three components

of the gate, the bootstrap lower bound reduces the false-improvement rate more than the other two components, and it ties with the margin in (C) (Section IV-F). For practice, these results support BRIDGE as an acceptance rule for tabular model updates within the tested conditions. In the retrospective clinical tasks, the gate kept updates that merely appeared better from reaching the deployed model. The open-source package `bridge-gate` (contribution C4) runs the procedure with the same settings and reproduces the stored results.

A. LIMITATIONS AND FUTURE WORK

The claims hold within the tested conditions listed in Table H8. Every collection as a whole meets the pre-registered criterion. The box states the population to which each bound applies. The paragraphs that follow give each limit with its reason and the next step.

What is bounded, and for which population?

(a) The formal bound (Proposition 1) on $r_{\mathcal{Q}}$ (11) applies to the specified i.i.d. synthetic task population $\mathcal{Q}$ and refers to that population, not to a single deployment. (b) On real data the primary bound is the KL-ratio bound, conditional on the observed collection and the frozen outer partition (the inner SEARCH/POOL split and the bag draw are re-drawn per trial). (c) The dataset-level bootstrap bound and LOCO are sensitivity analyses across datasets and sources, not the primary bound. (d) LOCO measures sensitivity across sources, not external validation in a new institution. (e) "Ground truth" means the locked reference, which is itself finite and imperfect; H_j is the false-improvement indicator measured on that finite reference of task j. (f) The Proposition 1 bound, like the empirical risk $\hat{r}$, concerns the false-improvement rate among the deployed configurations measured on the locked reference (Definition 1). The bootstrap quantile inside the gate is a decision rule, not a confidence interval, so no 95% coverage is claimed for it.

a: Comparators.

The best fixed validation margin was chosen after the results were seen. Its false-improvement rate does not differ significantly from that of BRIDGE, except at the smallest training size. At matched acceptance counts, BRIDGE retains its advantage between about 125 and 450 accepted updates (Appendix D). BRIDGE selects among updates across datasets in all five collections, and also within single datasets in (D) and (E) (Table H4). Learn then Test is compared with BRIDGE at the conceptual level (Section II-C). An empirical comparison at an equal label budget is the next step. In (A) and (B), the corrected t-test with per-split re-selection deploys more updates above the margin, with a share below the margin that does not differ significantly from that of BRIDGE (post hoc, Table H5).

b: Settings of the gate.

During development, the seven parameters of the gate and the partition were frozen (Table 2). In (D) every tested variant of these parameters meets the criterion. In (A) two variants fail, namely a bag drawn from all of SEARCH and a SEARCH of $1{,}000$ rows, so the frozen values are the tested condition of use (Appendix D). Likewise, the default candidate space of 173 configurations is the tested condition. A denser grid meets the criterion at every tested size in (A) to (D). On the binary datasets of (E) the bound exceeds 10% at $11{,}249$ configurations (Table H7). The next step is to develop a larger space, such as the joint seven-family space, through a new development cycle and a fresh confirmation.

c: Learners, metrics, and tuning.

The gate is confirmed for base learners that have iteration control. XGBoost with 300 boosting rounds and no early stopping fails in (A), (B), and (E). The failure lies in the gate statistic rather than in the final model, whose weighted-F_1 difference from HGB has an interval containing zero in (A) and (D) (Tables H10 and H12). On rare classes, the metrics of use are balanced accuracy and macro-F_1, because weighted-F_1 exceeds the 10% LOCO limit (Table H9). Two further conditions meet the criterion narrowly or in part. Feature selection with SVC misses the LOCO limit (11.4%), and tuning as an update meets the criterion in two of 12 cells because tuned updates are accepted rarely (Tables H6 and H13). The two alternative member pairs miss the criterion on the OpenML collection (E). Few updates are accepted there, and 11 of the 12 false improvements fall on small references (Tables H11 and C1). The next steps are a more stable feature-selection rule, a tuning path through development and confirmation, and larger references or more datasets for the member pairs in (E). The evidence covers tabular learners and label-based metrics. Deep and foundation models, raw images, and probability-level or ranking metrics are therefore directions for future work.

d: Data and clinical scope.

The clinical evaluation is retrospective and covers six sources, two of which are unweighted surveys. Net benefit therefore describes the sample rather than the population, and prospective validation in further institutions is the next step. The locked reference is finite and drawn once per task. On small datasets, near-zero changes are noisy in both directions, and two further reference seeds leave the conclusion unchanged (Table C1). The ACTG 175 task classifies the recorded event status (CD4 decline of at least 50%, AIDS, or death) without handling censoring, so it does not estimate a risk at a fixed horizon. Without ACTG 175, BRIDGE accepts $1{,}001$ updates with no false improvement (bound 0.39%; Table F2). The mammography source carries no patient identifier, and the main result is unchanged when that source is removed (Table F2). Two kinds of subgroup miss the criterion: the binary datasets of (E), narrowly on LOCO (10.3%), and multiclass subgroups of the OpenML collections, whose result varies with the number of classes (Appendix G); more multiclass datasets are the next step. The clinical baseline is the pre-registered threshold 0.5, not p_t; a baseline at $t = p_t$ and a full decision curve across thresholds are a further step [34]. Dataset shift and calibration drift after deployment lie outside the scope of the gate [85], [86].

e: Future work.

The evidence would be extended in five directions: (i) prospective validation, (ii) settings with explicit distribution shift, (iii) deep and foundation models, (iv) new datasets from the OpenML suite, and (v) an empirical comparison with Learn then Test at an equal label budget. Three further directions would extend the method itself rather than the evidence. The first direction is a series of updates on the same model, with a bound on the false-improvement rate that covers the whole series rather than each decision separately. The second direction is a count-based bootstrap for probability-level and ranking metrics. The third is the same-POOL cell as a separately pre-registered rule, together with per-task threshold selection. The quantile and the margin set the operating point between false improvements and yield, and choosing an operating point for a target yield, through a new development cycle on (A) and (B) and confirmation on new data, is a further step; such a result would be development, not a post hoc claim.

Appendices A to H follow: proofs, data, and pre-registration (A to C); data separation, frozen parameters, sensitivity, and cost (D); per-dataset results (E); the patient-identity audit (F); the OpenML collection (E) (G); and the detailed result tables (H).

APPENDIX A PROOFS

The proofs below cover Lemma 1, Proposition 1, and Corollary 1 of Section III-D. A final paragraph states the estimand used on the real data. The proofs rest on Assumption 1 and on a standard binomial tail bound alone.

Proof of Lemma 1. Under Assumption 1(i)–(iii), the tasks are i.i.d. and the gate is frozen before any data are seen (condition (ii) makes the batch construction and the training/reference split part of the task draw), so the pairs (A_j, H_j) are i.i.d. across tasks. Conditioning on $M = \sum_{j=1}^{T} A_j = m$ selects m accepted tasks, each with false-improvement indicator $H_j \sim \text{Bernoulli}(r_{\mathcal{Q}})$ independently; hence $S = \sum_{j=1}^{T} A_j H_j \sim \text{Binomial}(m, r_{\mathcal{Q}})$. □

Proof of Proposition 1. Fix $M = m \geq 1$ and write $r = r_{\mathcal{Q}}$. For $X \sim \text{Binomial}(m, r)$ the lower Chernoff–Kullback–Leibler bound gives, for a fixed threshold $x \leq r$, $\Pr(X \leq mx) \leq \exp(-m\,\text{kl}(x\|r))$. Let x^* be the fixed threshold with $m\,\text{kl}(x^*\|r) = \log(1/\delta)$ (it exists whenever $m\,\text{kl}(0\|r) \geq \log(1/\delta)$, and otherwise the event $r > U_m$ is empty so the bound holds trivially); since $\text{kl}(S/m\|r)$ decreases as S/m grows for $S/m \leq r$, the event $r > U_m$ is exactly $S/m < x^*$. Applying the bound with the fixed threshold x^* gives $\Pr(r > U_m) \leq \exp(-m\,\text{kl}(x^*\|r)) = \delta$, so the endpoint $U_m = \sup\{u \in [S/m, 1] : m\,\text{kl}(S/m\|u) \leq \log(1/\delta)\}$ (with $U_m = 1$ when $S = m$ or $m = 0$) is a one-sided upper bound on r with coverage $\geq 1 - \delta$ conditionally on $M = m$. Averaging over the distribution of M preserves coverage:

$$\Pr(r_{\mathcal{Q}} \leq U) = \sum_{m \geq 0} \Pr(M = m) \Pr(r_{\mathcal{Q}} \leq U_m \mid M = m) \geq \sum_{m \geq 0} \Pr(M = m)\,(1 - \delta) \;=\; 1 - \delta,$$

□

Proof of Corollary 1. Replace H_j by $H'_j \geq H_j$ for every task whose status is reference-uncertain, so that every accepted update is counted as a false improvement unless the reference indicates improvement, and let $S' = \sum_{j=1}^{T} A_j H'_j \geq S$. Bounds computed from (M, S') exceed the corresponding bounds computed from (M, S); since the bounds from (M, S) have coverage $\geq 1 - \delta$ by Proposition 1, so do the bounds from (M, S'). □

On the real data, let $\mathcal{E}$ denote the observed collection of tasks. The estimand is $r_{\mathcal{E}} = \mathbb{E}\left[\sum_j W_j \,\middle|\, \mathcal{E}\right] / \mathbb{E}\left[\sum_j A_j \,\middle|\, \mathcal{E}\right]$, the ratio of the expected sums of $W_j = A_j H_j$ and of A_j conditional on $\mathcal{E}$. The collection $\mathcal{E}$ includes the locked reference and the frozen outer partition, while the inner SEARCH/POOL split and the bag draw are re-drawn in every trial. The trials are independent given the fixed collection (independent seeds), with possibly different means across tasks, and the one-sided Chernoff–Kullback–Leibler bound applies to sums of independent Bernoulli variables with different means.

APPENDIX B DATA, AUDIT, AND GENERATOR

Tables of this appendix: p. 22.

Table B1 covers every dataset and task of (A) and (C)–(E), grouped by collection, with the domain, source, sample size, minority rate, and size of the locked reference. The table footnote defines the SEARCH/POOL split, the eligible training pool, and the number of trials.

a: Inclusion and exclusion rules.

A dataset qualified for the benchmark on three conditions: it was publicly available, it was tabular or image-derived in tabular form, and its labels were recorded. A dataset was excluded for any of the following reasons:

- The data were synthetic or simulated, or formed a time series.
- The binary target came from thresholding a continuous variable, or its provenance could not be verified.
- The class ratio was exactly 50%, which marks artificially balanced data, or more than 25% of the rows were duplicates.
- A single feature predicted the target almost without error (feature leakage, as for PhiUSIIL, AUC 0.996).
- The dataset family had already informed the design, or the task was a regression.
- The rows were grouped by subject or session, or identifiers indicated leakage that de-duplication could not remove.

TABLE B1. Data partition for each dataset or task in (A) and (C)–(E).

Dataset	Domain	Source	N	Min. %	Ref.
(A) 11 new datasets					
pc3	Software	OpenML 1050	1,563†	10.24	468
MIC	Medical	OpenML 45648	1,649†	15.59	494
Marketing_Campaign	Marketing	OpenML 46940	2,240†	14.91	672
Pumpkin_Seeds	Agriculture	OpenML 46951	2,500†	48.00	750
seismic-bumps	Mining	OpenML 46956	2,584†	6.58	775
Rice	Agriculture	UCI 545	3,810	42.78	1,143
Census-Income-KDD	Social	OpenML 42750	199,523	6.21	100,000
Santander2016	Banking	OpenML 46859	76,020	3.96	56,020
KDD98	Social	OpenML 42343	82,318	11.77	62,318
LICD	Employment	OpenML 46792	63,634	25.72	43,634
Homesite	Insurance	OpenML 46857	260,753	18.75	100,000
(C) 12 clinical tasks					
mammography	Radiology	OpenML 310	7,847	3.22	2,354
actg175	Clinical	UCI 890	2,139†	24.36	641
nhanes_hipertensi	Survey	NHANES 2017–18	5,559	38.12	1,667
nhanes_diabetes	Survey	NHANES 2017–18	5,393	16.24	1,617
nhanes_asma	Survey	NHANES 2017–18	5,564	15.03	1,669
nhanes_pjk	Survey	NHANES 2017–18	5,553	4.77	1,665
brfss_stroke	Survey	BRFSS 2015	440,166	4.15	100,000
brfss_asma	Survey	BRFSS 2015	439,963	13.50	100,000
brfss_ginjal	Survey	BRFSS 2015	439,879	3.56	100,000
brfss_depresi	Survey	BRFSS 2015	439,259	19.07	100,000
covid_mx	Public Health	Mexico COVID-19	218,784	12.29	100,000
wids_icu	ICU	WiDS 2020	91,713	8.63	71,713
(D) 34 binary supporting datasets					
miniboone	Physics	OpenML 41150	130,064	28.06	100,000
compas (dev)	Justice	OpenML 42192	5,278	47.04	1,583
credit-default-tw (dev)	Finance	OpenML 42477	30,000	22.12	8,000
noshow	Healthcare	OpenML 43439	110,527	20.19	90,527
cardio	Medical	OpenML 45547	70,000	49.97	50,000
heloc (dev)	Finance	OpenML 45553	9,871	47.97	2,961
lt-vehicle-loan	Finance	OpenML 46430	233,154	21.71	100,000
codexglue (dev)	Software	OpenML 46682	27,318	45.61	8,000
kickstarter	Crowdfund	OpenML 46810	86,502	31.96	66,502
bank-churn (dev)	Banking	OpenML 46911	10,000	20.37	3,000
ecom-shipping (dev)	E-commerce	OpenML 46924	10,999	40.33	3,299
employee (dev)	HR	OpenML 46925	4,653	34.39	1,395
amazon-access (dev)	HR	OpenML 4135	32,769	5.79	8,000
kick	Automotive	OpenML 41162	72,983	12.30	52,983
porto-seguro	Insurance	OpenML 42206	595,212	3.64	100,000
(D) 34 binary supporting datasets (cont.)					
relevant-images	Vision	OpenML 42680	134,731	4.14	100,000
newspaper-churn (dev)	Media	OpenML 44226	15,855	19.15	4,756
online-shoppers (dev)	E-commerce	OpenML 45060	12,330	15.47	3,699
give-some-credit	Finance	OpenML 45577	150,000	6.68	100,000
hmeq (dev)	Finance	OpenML 46431	5,960	19.95	1,788
avida-hil6	Biology	OpenML 46799	573,891	3.66	100,000
avida-cov2 (dev)	Biology	OpenML 46800	3,733	19.85	1,119
acs-drem	Social	OpenML 46802	1,138,289	7.90	100,000
fake-job (dev)	Employment	OpenML 46809	12,725	4.33	3,817
sepsis	Medical	OpenML 46853	110,341	7.37	90,341
dota2	Gaming	OpenML 45563	102,944	47.27	82,944
santander	Banking	OpenML 45566	200,000	10.05	100,000
home-credit	Finance	OpenML 45567	307,511	8.07	100,000
coil2000 (dev)	Insurance	OpenML 46281	9,822	5.97	2,946
bnp-paribas	Insurance	OpenML 46856	114,321	23.88	94,321
polish-bankruptcy (dev)	Finance	OpenML 46950	5,910	6.94	1,773
taiwan-bankruptcy (dev)	Finance	OpenML 46962	6,819	3.23	2,045
naticusdroid (dev)	Cyber	OpenML 46969	7,491	35.03	2,247
hr-analytics (dev)	HR	OpenML 46935	19,158	24.93	5,747
(D) 5 multiclass datasets					
Fashion-MNIST	Vision	OpenML 40996	70,000	10.00	50,000
DryBean	Agriculture	UCI 602	13,611	3.84	4,083
USPS	Vision	OpenML 41082	9,298	7.61	2,789
mnist_784	Vision	OpenML 554	70,000	9.02	50,000
Kuzushiji-MNIST	Vision	OpenML 41982	70,000	10.00	50,000
(E) 11 binary OpenML datasets					
airline-satisfaction	Marketing	OpenML 46920	129,880	45.27	100,000
KDDCup09-churn	Marketing	OpenML 1112	50,000	7.34	30,000
bank-marketing	Banking	OpenML 1461	45,211	11.70	25,211
pulsar-HTRU2	Physics	OpenML 45558	17,898	9.16	5,369
jm1	Software	OpenML 1053	10,885	19.35	3,265
support2	Medical	OpenML 46860	9,105	31.89	2,731
telco-customer-churn	Marketing	OpenML 42178	7,043	26.54	2,112
phoneme	Speech	OpenML 1489	5,404	29.35	1,621
spambase	Cyber	OpenML 44	4,601	39.40	1,380
sick	Medical	OpenML 38	3,772	6.12	1,131
bioresponse	Biology	OpenML 4134	3,751	45.77	1,125
(E) 3 multiclass OpenML datasets					
ipums-la-97-small	Survey	OpenML 382	7,019	3.68	2,105
otto-group	E-commerce	OpenML 45548	61,878	3.12	41,878
corporate-credit-ratings	Finance	OpenML 46372	5,000	10.00	1,500

*Locked reference: size (Ref.) and source of each real dataset or task. The rows that remain once the reference is set aside form the eligible training pool, from which every trial draws its training rows. Groups. (A): the 11 new binary datasets. (C): the 12 clinical tasks. (D): the 39 supporting datasets, namely the 34 binary confirmation datasets (including the 18 development datasets, marked (dev), on which the gate and the candidate space were frozen) and the 5 multiclass datasets. (E): the 14 datasets of the OpenML collection, with names shortened and identified by the OpenML ID. Setting (B), the synthetic population, is generated rather than drawn from a dataset and is described in this appendix. Min.: minority-class share; for a multiclass dataset, the smallest class. Split: 2,430 training rows by default (SEARCH 630, POOL 1,800); †: 1,000 training rows (SEARCH 259, POOL 741). Trials per task: 100 in (A), (D), and binary (E); 200 in (C); 60 in multiclass (D) and (E). Domains briefly characterize the public sources, and exact definitions live in the code-repository loaders. The 28 domain labels map to seven groups, all represented among the 64 real datasets: Health (Healthcare, Medical, Radiology, Clinical, Survey, Public Health, ICU); Finance, banking and insurance (Finance, Banking, Insurance, Crowdfund); Agriculture and mining (Agriculture, Mining); Cyber and software (Software, Cyber); Commerce, services and HR (E-commerce, Marketing, HR, Employment, Automotive); Social, legal and media (Social, Justice, Media, Gaming); Science and vision (Physics, Biology, Vision, Speech).

TABLE B2. Clinical tasks.

Task	Population	Outcome	Predictors / timing	Hypothetical action
mammography	Screening mammograms (OpenML 310)	Malignant calcification	Image-derived features at reading	Recall for biopsy
actg175	HIV-1 trial participants (UCI 890)	Recorded event: CD4 decline $\geq 50\%$, AIDS, or death	Baseline covariates only	Intensify monitoring
nhanes_*	US adults $\geq$20 y, NHANES 2017–18	Self-reported: hypertension, diabetes, asthma, CHD	Exam & questionnaire, same visit	Refer for testing
brfss_*	US adults, BRFSS 2015	Self-reported: stroke, asthma, kidney disease, depression	Survey responses, same interview	Refer for screening
covid_mx	Confirmed COVID-19 patients, Mexico	In-hospital/registered death	Admission data (comorbidities, triage)	Escalate care
wids_icu	ICU admissions, GOSSIS (WiDS 2020)	Hospital death	First-24 h ICU measurements	Escalate care

*Clinical tasks: population, outcome, timing of predictors and hypothetical action. The COVID-19 cohort is the open data released by the Mexican Secretaría de Salud. NHANES outcomes are physician diagnoses as reported by the respondent; the NHANES and BRFSS survey tasks are cross-sectional health-related classification (predictors from the same visit or interview), not prognosis.

The 25% duplicate threshold and the other rules were fixed in the pre-registration of run MINI-21 before any results were seen (Table C2). The benchmark expanded over successive selection rounds, so the rules above are the union of the rules of all rounds. The versioned list, with the reason for each dataset, is maintained in the code repository.

b: Synthetic population $\mathcal{Q}$: tasks and seeds.

The synthetic analysis draws 1,400 tasks i.i.d. from a fixed generative population $\mathcal{Q}$ and runs one trial per task; the population was pre-specified in run MINI-25 (Table C2). Each task carries the seed of the generator that creates it (generator in the code repository). The primary batch uses seeds 0 to 1,399, which no pilot run touched; the replication batch uses 1,500 to 2,899; and later runs take new seeds, such as 2,900 to 4,299 for the tuning runs.

c: Synthetic population $\mathcal{Q}$: label mechanism.

Each task is generated in its own batch of 103,000 rows, with 3,000 development rows and 100,000 reference rows. The development rows allow the same 2,430 training rows as on the real data, and the reference measures the true gain with a small error.

The label mechanism is drawn with equal probability from a tree mechanism and a linear mechanism. Under the tree mechanism, each task draws $p \in [5, 40]$ standard normal features, 4 to 12 pairwise-threshold interactions, weights $\sim \mathcal{N}(0, 2^2)$, thresholds $\sim \mathcal{N}(0, 0.7^2)$, and two feature indices uniformly. Under the linear mechanism, each task draws $p \in [5, 40]$ standard normal features with coefficients $\beta \sim \mathcal{N}(0, 1)$; a fraction of the coefficients is set to zero, and the rest are divided by the square root of the number of non-zero coefficients. Three parameters are drawn per task: the signal strength $s \in [0.05, 1]$, a class-balance parameter $m \in [0.05, 0.5]$, and the linear sparsity $\lambda \in [0, 0.7]$.

Three pilot rounds preceded the pre-registration of MINI-25 (Table C2), where the ranges and constants were fixed. With a strong signal, the naive rule accepted a false improvement on 1 of 59 pilot tasks; after the signal range was extended down to $s = 0.05$, it did so on 6 of 58.

Within each batch, the logit is standardized and scaled by s, and the labels are Bernoulli draws whose probability is the logistic transform of the scaled logit minus an intercept, the batch quantile of order $1 - m$ of the scaled logit. A stratified split sets the development rows apart from the reference rows. The parameter m shifts the intercept and does not fix the class proportion: the realized minority proportion averages 42% and lies between 29% and 50% for 90% of the tasks.

d: Synthetic population $\mathcal{Q}$: design and truth.

A pilot phase designed the population so that the naive false-improvement rate is realistic, at least 15% and close to the 19% seen on real data; the pilot tasks enter no reported result. A Gaussian mechanism was dropped because the naive rule was almost never wrong on it, and the results of BRIDGE were not used in these design choices. The true gain Δ is computed on each task's own reference of 100,000 rows, and the decision pairs are i.i.d. across tasks (Assumption 1). The standard error of the reference has a median of about 0.0014, roughly a quarter of $\varepsilon = 0.005$. The comparison uses the estimated rate $\hat{r}$, the bound U of Proposition 1 (12), and a conservative variant that counts updates of uncertain reference status as false improvements (Table 11).

e: Training rows and reference size.

A task has 2,430 training rows when at least 2,480 labeled examples remain after the reference is locked, and 1,000 training rows otherwise (Table B1). SEARCH takes 25.9% of the training rows and POOL the rest (3). The default tasks therefore have 630 SEARCH rows and 1,800 POOL rows, and Table D2 gives the basis of these sizes. When $N < 40{,}000$, the reference has $\lfloor 0.3N \rfloor$ rows, capped at 8,000. When $N \geq 40{,}000$, the reference has $\min(N - 20{,}000,\ 100{,}000)$ rows (Table D2).

f: Checks on the sizes.

Two checks examine the sizes chosen above (Table C1). First, the median standard error of the gain on the locked reference is 0.0003 to 0.0020 on the five new datasets whose reference holds at least 40,000 rows, and 0.0043 to 0.0152 on the six datasets with a shorter reference. Second, a SEARCH of 360, 630, or 1,000 rows on $\mathcal{Q}$ gives bounds of 1.65%, 3.56%, and 3.11%, at yields of 32.1 to 33.9 per 100, and all three sizes meet the criterion.

APPENDIX C PRE-REGISTRATION CHRONOLOGY

Tables of this appendix: pp. 24–25.

Nineteen development runs (MINI-01 to MINI-19) fixed the seven parameters and the candidate space, and the gate was then frozen (Table C2). Every later run tests that frozen gate on new data, in sensitivity analyses, and in the conditions of use; variants are reported beside the default, and no run changes the default.

The numbers quoted in the text from development, confirmatory, and sensitivity runs that appear in no other table are recomputed from the stored outputs and listed in Table C1, sorted by run number.

Table C2 summarizes the path from development to confirmation. For each confirmatory run, the table gives the primary endpoint and its outcome, including predictions that were not confirmed. The pre-registered predictions of the component analysis (Section IV-F) appear as P1, P2, and P3 in Table C2. The pre-registration documents record the labels H1, H2, and H35 for the same predictions. Run identifiers of the form MINI-n label the pre-registered runs [32] and match the scripts and CSV archives in the code repository.

a: Execution and documented bugs.

Every trial of the confirmatory runs MINI-20 to MINI-38 completed without an error (Table C2).

Four documented bugs appear in the chronology with the corrections (Table C2). A single-class bag error in the early-stopping split (MINI-18) was fixed by requiring at least 2 examples of each occurring class per bag. The first launch of MINI-29 stopped before producing data, because a helper function could not be passed to the parallel workers and the queue script reported completion despite the failure. The function was moved to module level, and the script now checks the exit status, with trial logic, data, and seeds

TABLE C1. Runs cited in the text, by run number.

Check	Acc./ false	Bound % ↓	LOCO % ↓	Yield
MINI-02 · 34 (D) × 100				
frozen gate, threshold-only HGB	586 / 1	1.24	1.45	17.2[a]
MINI-04 · 34 (D) × 100				
threshold-only KNN, cross-fit gate	150 / 7	14.04	19.95	4.4
MINI-09 · 18 dev. × 100				
third member KNN	218 / 3	5.71	8.70	12.1
third member SVC	269 / 4	5.27	8.41	14.9
MINI-12 · 18 dev. × 100				
tuned HGB as the update	193 / 5	8.61	12.32	10.7
tuned HGB, baseline and member	356 / 13	8.12	9.63	19.8
default HGB, baseline and member	455 / 10	5.20	6.41	25.3
MINI-14 · 27 (D) × 40 per size				
tuning on 630 rows	87 / 1	10.12	14.27	8.1
tuning on 2,000 rows	40 / 0	14.52	20.08	3.7
tuning on 5,000 rows	30 / 0	20.95	59.82	2.8
all three sizes (3,240 dec.)	157 / 1	5.21	6.67	4.8
MINI-16 · 18 dev. × 100				
accuracy	234 / 6	7.75	11.94	13.0
balanced accuracy	1,352 / 4	0.93	1.01	75.1
MINI-17 · size study, 3,240 trials				
skipped at 1,000 training rows: avida-hil6 2, porto-seguro 40				
MINI-19 · 8 wide dev. datasets				
union policy, first 50 trials (400 dec.)	61 / 1	15.10	32.37	15.2[b]
extended to 150 (MINI-19b; 1,200 dec.)	179 / 5	9.30	20.58	14.9[c]
MINI-20 · size study, 8,100 trials				
skipped at 1,000 training rows: avida-hil6 3, porto-seguro 100				
MINI-24 · 11 new, dataset-level bootstrap bound[d]				
naive		29.4		
5 × 2 cv		13.0		
corrected t		12.5		
BRIDGE		3.1		

Check	Acc./ false	Bound % ↓	LOCO % ↓	Yield
MINI-27 · 11 new, 1,064 dec.				
error of the gain estimate[e]: naive $+0.0034$; cross-fitted -0.0034				
MINI-28 · 12 clin. × 100, weighted-F_1 stage				
BRIDGE	222 / 6 (2.7)	8.15	12.80[f]	
naive	1,145 / 254 (22.2)	26.16	28.66[f]	
MINI-31 · $\mathcal{Q}$, 1,400				
XGBoost project default as base	138 / 16 (11.6)	19.33[g]		9.9
MINI-33b · $\mathcal{Q}$, 1,400				
regularized XGBoost below HGB in base weighted-F_1: 94.1% of tasks				
MINI-34 · $\mathcal{Q}$, 1,400, 9 configurations				
accepted, skill-positive base	2,161 / 1 (0.05)			
accepted, skill-negative base	1,142 / 77 (6.74)			
Spearman ρ, skill vs bound, 9 configurations: -0.82 ($p = 0.007$)				
Spearman ρ, 11 configurations incl. 2 MLPs: -0.74 ($p = 0.010$)				
MINI-37 · $\mathcal{Q}$, 1,400				
16 bags, SEARCH 630	459 / 9	3.98[g]		32.8
32 bags, SEARCH 630	475 / 8	3.56[g]		33.9
64 bags, SEARCH 630	465 / 5	2.69[g]		33.2
32 bags, SEARCH 360	464 / 2	1.65[g]		33.1
32 bags, SEARCH 1,000	450 / 6	3.11[g]		32.1
Monte Carlo error[h], $\mathcal{Q}$: SD 0.00014; 0.3%; 1.0% (14 of 1,400)				
Monte Carlo error[h], 11 new × 10: SD 0.00009; 0.7%; 1.8% (2 of 110)				
reference SE[i], 5 new datasets, reference $\geq$ 40,000: 0.0003–0.0020				
reference SE[i], 6 new datasets, reference $<$ 40,000: 0.0043–0.0152				
MINI-38 · 11 new × 100, three reference seeds				
original seed: naive	1,064 / 203 (19.1)	22.90	25.27	96.7
original seed: BRIDGE	284 / 3 (1.1)	4.23	5.95	25.8
seed 1: naive	1,078 / 244 (22.6)	26.60	29.29	98.0
seed 1: BRIDGE	340 / 8 (2.4)	6.09	8.12	30.9
seed 2: naive	1,067 / 236 (22.1)	26.15	28.83	97.0
seed 2: BRIDGE	289 / 2 (0.7)	3.42	5.05	26.3
MINI-50 · joint seven-family space				
(A): naive false impr. 22.4%; BRIDGE 1.6%				
(D): naive false impr. 13.2%; BRIDGE 0.7%				
MINI-53 · (E), 1,280 dec., post hoc check[j]				
default (DT and SVC)	192 / 2 (1.0)	5.40	8.43	15.0
members LR and NB	244 / 6 (2.5)	7.35	10.16	19.1
members MLP and KNN	147 / 6 (4.1)	12.94	20.77	11.5
false improvements of both pairs: 12; reference-uncertain 11, on references of 1,125–2,112 rows				
clearly worse: 1 (airline, reference 100,000 rows, gain -0.0017)				

*Runs MINI-02 to MINI-19: development runs (Table C2). The later runs are confirmatory, sensitivity, or descriptive. Data: datasets × trials per dataset, or tasks. (D): the 34 confirmation datasets of the supporting group, of which 27 are in MINI-14; dev.: the 18 development datasets. new: the 11 new datasets of (A). clin.: the 12 clinical tasks. dec.: decisions. Each run header gives the run number and the data of the run; a row that spans the columns states a result in words. Acc./false: accepted / false improvements, with the false-improvement share in parentheses (%). Bound: KL-ratio bound ($b = 1$), %. LOCO: leave-one-dataset-out bound, %. Yield: accepted updates per 100 decisions. Superscripts. Power check LOCO_0, the LOCO bound recomputed with zero false improvements: a, 0.81%; b, 18.22%; c, 5.54%. With 50 trials per dataset the bound exceeds 10% even when no decision is a false improvement, and 100 or more trials bring it below 10%. d: 95% upper bound of the false-improvement rate, from a bootstrap over the 11 datasets with 10,000 resamples, %. e: estimate minus reference truth for the deployed configuration, against the margin 0.005. f: LOCO by source. g: Proposition 1 bound on $\mathcal{Q}$. h: three Monte Carlo checks, namely the standard deviation of the 5% bootstrap quantile over 20 bootstrap seeds per trial, the share of re-runs whose decision one alternative seed changes, and the share of trials in which at least one of the 20 seeds changes the decision. i: range of the per-dataset median bootstrap standard error of the reference gain, separately for references of 43,634 to 100,000 rows ($\geq$ 40,000) and 468 to 1,143 rows ($<$ 40,000). MINI-12 and MINI-14 tune HGB with Optuna (tree-structured Parzen estimator, 25 search iterations) on the 630-row SEARCH split; MINI-14 repeats this at 630, 2,000 and 5,000 tuning rows. Mean reference weighted-F_1 in MINI-12: 0.8180 for the default HGB and 0.8206 for the tuned learner, a mean gain of $+0.0026$. MINI-14, the mean reference gain of the tuned learner at 630, 2,000, and 5,000 rows and over all sizes is $+0.0027$, $+0.0021$, $+0.0020$, and $+0.0023$. Share of trials whose 95% reference interval lies below zero: 12.3%, 9.0%, 6.7%, and 9.3%. MINI-17/MINI-20: trials skipped because a bag had fewer than 10 examples of a class. MINI-28: the clinical stage with weighted-F_1 as the primary endpoint, before the net-benefit stage (MINI-30). MINI-31: XGBoost with the project defaults (300 rounds, no early stopping) during development. MINI-33b: share of tasks on which the regularized XGBoost has a lower base weighted-F_1 than HGB on the locked reference. MINI-34: accepted updates pooled over the nine pre-registered HGB and XGBoost configurations, split by the sign of the base learner's log-loss skill on POOL; Spearman ρ between median skill and bound (Table H10). MINI-37: the bag count and the SEARCH size, varied one at a time. MINI-38: the outer split is repeated with two fresh reference seeds (original 20261125; seed 1: 1038000001; seed 2: 1038000002). j: post hoc, descriptive. reference-uncertain: the 95% Bayesian-bootstrap interval of the reference gain contains zero. clearly worse: the interval lies below zero.

unchanged. The first cross-machine verification of MINI-20 reported 83 differences, caused by a seed argument in the wrong position; after the fix, the second verification found none. A concatenation bug misreported the naive count, which was corrected in MINI-21.

b: Tuning search spaces.

The tuning runs (MINI-12, MINI-14, MINI-39, MINI-40) used the tree-structured Parzen estimator of Optuna with 25 iterations and three-fold cross-validation in SEARCH, scoring weighted-F_1 (Table C2). The budget was fixed in MINI-12 and used unchanged, and the search ranges were fixed in the pre-registration of each run.

MINI-12 and MINI-14 tuned HGB with the same sampler and space: learning rate 0.01 to 0.3 (log scale), 50 to 300 boosting iterations, 3 to 63 leaf nodes, 1 to 50 minimum samples per leaf, and L2 regularization 10^{-8} to 10 (log scale). MINI-39 and MINI-40 tuned four learners, with spaces in the scikit-learn parameter names. For DT: `max_depth` 2 to 30, `min_samples_split` 2 to 20, `min_samples_leaf` 1 to 20, and Gini or entropy. For SVC: C from 10^{-3} to 10^{3} (log scale), a radial basis function (RBF), linear, or polynomial

TABLE C2. Chronology of the pre-registered runs, ordered by run number.

Run	Fixed beforehand	End.	Outcome
01–19	mechanism studies; thresholded HGB, HGB–DT ensembles, the 173 space	Beh	validation–test independence hypothesis *not* supported
04	threshold-only KNN with the cross-fit gate	Crt	fails (bound 14.0%)
09	third ensemble member: KNN or SVC	Yld	SVC chosen: higher yield at a comparable false-improvement rate
12, 14	Optuna-tuned HGB on SEARCH	Dsc	small gains; clearly worse in up to 12% of trials, so the fixed configuration is frozen
17, 20	$1{,}000/2{,}430/6{,}000$ training rows (17 dev.); fresh seeds	Crt	pass (bound $3.04/1.05/1.37\%$)
21	11 new binary datasets, weighted-F_1	Crt	pass (bound 4.23%, LOCO 5.95%)
22	5 new multiclass datasets of (D)	Crt	pass (bound 2.11%, LOCO 3.16%)
23, 24	comparators on identical decisions	Bnd	test-statistic-based rules fail; BRIDGE passes (Table H3)
25	$1{,}400$ i.i.d. tasks of $\mathcal{Q}$, one trial each	Prp	pass (3.56%; conservative variant 4.70%)
26	honest holdout on the same labeled data	Yld	BRIDGE 1.40–$1.86\times$ as many updates; false improvements 1.1% against 1.3–3.2%
27	seeds and models of MINI-21; 173 candidates	Dsc	naive estimate $+0.0034$, cross-fitted -0.0034 (Table C1)
28	12 clinical tasks, weighted-F_1 primary; LOCO by source	Crt	LOCO fails at 100 trials per task
29	fixed margin t^* from development, at $6{,}000$ rows and on new data	Prd	failure confirmed: near $t^*{=}0.05$ yields 1.4, 0.1; bounds 16.1%, 100%
29	fixed margins chosen post hoc in sample	Dsc	$t \in \{0.020, \ldots, 0.035\}$ pass $7/7$ settings; leave-one-out: BRIDGE $7/7$, margin $6/7$
30	the same 12 tasks; net benefit $p_t{=}0.2$, balanced accuracy	Crt	pass (bound 0.33%)
31–36	XGBoost and MLP as base learners	Crt	XGBoost $\geq$ 100 rounds without early stopping fails; early-stopped learners pass
33–36	skill diagnostic and two MLP predictions	Prd	skill rule and both MLP predictions *not* confirmed
32, 35	ablation $2 \times 2 \times 2$; $\lvert\mathcal{C}\rvert$ up to $11{,}249$	Mec	same-POOL cell also meets the criterion; predictions **P1 and P3 not confirmed**
37	bag count, SEARCH size; bootstrap and reference precision	Crt	all variants pass (bags 2.7–4.0%, SEARCH 1.6–3.6%; Table C1)
38	11 new datasets; two fresh reference seeds	Crt	pass at both seeds (false improvements 0.7–2.4%, bounds 3.4–6.1%, worst LOCO 8.1%)
39, 40	HGB, DT, SVC, KNN tuned in SEARCH; (A), (B), multiclass (D)	Crt	false improvements lower in all 12 cells; criterion met in 2 (Table H13)
41	minority prevalence 20/10/5%; three metrics	Crt	balanced accuracy, macro-F_1 pass; weighted-F_1 exceeds LOCO (Table H9)
43	39 datasets of (D); fresh seeds	Crt	pass (bound 1.47%, LOCO 1.62%, LOCO_0 0.39%, yield 30.1)
44	stored seeds of (A), (D); tuned baseline; q–ε grid; $\eta = 2$ pp	Tst	tuning raises the baseline by about 0.3 points (Tables D6, D5)
45, 50	LR–NB and MLP–KNN pairs; joint seven-family space	Crt	met in (A) and (D) with every pair and the joint space (Table H11)
46, 49	XGBoost (300 rounds; early stopping), MLP; final quality	Crt	XGBoost with 300 rounds fails in (A); early-stopped learners pass (Table H11)
47	model seed, bag fraction, bag count, SEARCH size	Crt	(D) passes every variant; (A) two variants fail (Appendix D)
48	comparators, holdouts, ablation cells on (C), (D)	Tst	BRIDGE lower than or not different from every comparator (Table H3)
51, 52	grids of $1{,}825$ and $11{,}249$ on (C), (D), binary (E)	Crt	(C), (D) pass at every size; binary (E) at $1{,}825$ (Table H7)
53	members, base learners, joint space on (B), (C)	Crt	every design passes except XGBoost with 300 rounds on (B) (Table H11)
53	the same designs on (E)	Crt	early-stopped learners and joint space pass; member pairs and 300-round XGBoost miss
(E)	14 new OpenML datasets; frozen gate	Crt	pass jointly ($192/2$; bound 5.40%, LOCO 8.43%); subgroups do not meet every criterion
(E2)	19 new OpenML datasets; frozen gate	Crt	pass jointly ($452/7$; bound 4.16%, LOCO 5.44%)
(E2)	subgroups by class count	Crt	binary passes only with the flagged `churn`; multiclass fails LOCO
(E2)	amendment: `pendigits` replaces `shuttle` (60 trials stopped)	Dsc	replaced before any metric was read; deviates from the grouped-data rule (Appendix G)
post hoc	stored per-trial outputs, reanalyzed	Dsc	random-selection control and gain–loss; *not* pre-registered (Table H4)

*Run: MINI-n run number. Rows are ordered by the first run number, runs that answer the same question share a row, and a run with two findings has two rows. MINI-01 to MINI-19 are development runs (dev.). Endpoint codes: Crt, the pre-registered criterion (bound $< 10\%$, LOCO $< 10\%$, $\text{LOCO}_0 < 10\%$, yield ≥ 5; Table 6); Bnd, a bound below 10%; Prp, the bound of Proposition 1; Tst, paired differences or TOST; Prd, a pre-registered prediction; Beh, the behavior of the gate by metric, space, and size; Yld, the yield and the false-improvement rate; Mec, the mechanism and the stability in $\lvert\mathcal{C}\rvert$; Dsc, descriptive, with no endpoint. The OpenML collections (E) and (E2) and the post hoc reanalysis close the table, each with its own pre-registration or with none. The run number MINI-42 labels the one-pass analysis of the OpenML datasets (Appendix G). Code manifests (SHA-256) recorded before each confirmatory seed; the runs are distributed across three machines, per dataset or per block of tasks on the synthetic population, and the cross-machine checks are described in the reproducibility paragraph of this appendix; 0 trials stopped with an error on confirmatory runs. The pre-registration is internal: the times are recorded by the author. The pre-registration index in the code repository lists, per run, the pre-registration file, its recorded time, and the SHA-256 of the file and of the manifest, and the release tag at submission fixes the package contents; the hashes and the cross-machine checks verify the contents, not the times.

kernel, γ set to `scale` or `auto`, degree 2 to 4, and `coef0` 0 to 1. For KNN: `n_neighbors` 1 to 31, uniform or distance weights, and the Minkowski exponent $p \in \{1, 2\}$.

c: Development check.

On the development data, the gate met its bound on accuracy (7.75%) and exceeded the LOCO limit (11.9%). The default threshold is already close to optimal for accuracy, and candidates are seldom better; the yield is 13.0 per 100 decisions on accuracy, against 75.1 on balanced accuracy (MINI-16, Table C1). Rare acceptance widens the KL-ratio and LOCO bounds, as in the tuning limit of Section V-A.

d: Clinical runs.

The clinical evidence comes from two pre-registered runs on the same 12 tasks and 6 sources. The first fixed weighted-F_1 as its primary endpoint and failed the LOCO criterion at 100 trials per task (KL bound 8.15%, per-source LOCO 12.8%; MINI-28 in Table C1). The second was pre-registered after that failure, with net benefit at $p_t = 0.2$ and balanced accuracy as co-primary on the same tasks with fresh seeds.

e: Reproducibility.

Experiments used Python 3.12, scikit-learn 1.9.1, NumPy 2.5.3, pandas 3.0.6, SciPy 1.18.1, and XGBoost 3.4.1. On the tested Intel and AMD machines, all trials reproduce bit-identically, except the 17 trials that use Optuna tuning or mutual-information feature selection, which are identical on the same CPU family (Intel or AMD), and one MLP–KNN trial of MINI-53, whose reference gain differs between CPU families without changing the decision. The released software package (C4, Section I) runs the paper's procedure as its default behavior and reproduces every stored result row of the 15 development and confirmatory runs, including all 3,300 binary decisions of MINI-21 and the 300 multiclass decisions of MINI-22 (Table C2). The identity checks of every run are stored with its manifest in the code repository.

APPENDIX D DATA SEPARATION, FROZEN PARAMETERS, CANDIDATE SPACES, SENSITIVITY, FIXED MARGIN, AND COST

Tables of this appendix: pp. 26–28.

a: Data used for each choice.

For each choice in the procedure, Table D1 records the data with which the choice is made and the place where the choice is assessed. The table states the separation of Section III-B explicitly: within the cross-fit, each per-fold candidate is assessed on the fold that did not select it.

TABLE D1. What each choice is made with, and where it is assessed.

Choice	Made with	Assessed on
Model type and hyperparameters	fixed *a priori*, no tuning (rationale: Section III-A)	—
Candidate space (173 binary default) and gate parameters	18 development datasets, frozen	families of new data never seen in development (new binary, new multiclass, clinical, fresh synthetic); the 34 confirmation datasets of (D) with fresh seeds
Preprocessing (imputation, scaling, encoding)	fitted per bag, SEARCH only; WiDS category codes are listed from all rows before the split (covariates only, no labels)	POOL and reference
Base models (up to 32 bag draws, kept bags averaged); early stopping	SEARCH (internal bag validation)	POOL and reference
Candidate selection inside the gate	$2/3$ of POOL	the held-out $1/3$ of POOL
Accept / reject	lower bootstrap quantile of the cross-fit counts	—
Ground truth (false improvement or not)	—	locked reference, never used for any decision

*Within the cross-fit, every per-fold candidate is assessed on the fold that was not used to select it; the final configuration is re-selected on all of POOL (Section III-B).

b: Default binary space.

The default binary space combines three families of candidates on coarse grids. The thresholds run from 0.05 to 0.95 in steps of 0.05, and the mixture weights move in steps of 0.25. The three-model family uses the narrower threshold range 0.30 to 0.60. The grids are coarse by design, to keep the selection bias small (Section III-A). For HGB, the grids give 18 threshold variants, which are the 19 thresholds without the baseline value 0.5. The grids also give 57 HGB–DT mixtures, which are 3 weights $\times$ 19 thresholds. The grids also give 98 HGB–DT–SVC mixtures, which are 14 weight triples $\times$ 7 thresholds. The 14 triples are the 15 weight combinations in steps of 0.25 without pure HGB. The three families of the default space together give $18{+}57{+}98 = 173$ configurations. Run MINI-10, which introduced the three-model thresholds, preceded the pre-registration of MINI-11 that froze the space before the confirmatory runs (Table C2).

c: Refined and joint spaces.

The refined spaces of Table H7 keep the same three families with finer steps. The refined spaces also widen the three-model threshold range to 0.30 to 0.70. The space of $1{,}825$ configurations has 46 HGB thresholds (0.05 to 0.95, step 0.02) and 414 HGB–DT pairs (weights 0.1 to 0.9, step 0.1, $\times$ 46 thresholds). The space adds $1{,}365$ three-model points (65 weight vectors in steps of 0.1 without pure HGB, $\times$ 21 thresholds, step 0.02). The space of $11{,}249$ configurations has 90 HGB thresholds (step 0.01, without 0.50) and $1{,}729$ pairs (weights 0.05 to 0.95, step 0.05, $\times$ 91 thresholds). The space adds $9{,}430$ three-model points (230 weight vectors in steps of 0.05, $\times$ 41 thresholds, step 0.01). The joint space of the member study applies the same grids to every tested member. The joint space holds 18 HGB thresholds, 57 two-model mixtures for each of the seven members, and 98 three-model mixtures for each of the three tested pairs, for $18 + 7 \times 57 + 3 \times 98 = 711$ configurations (Table H11).

The members of the alternative pairs keep fixed settings and are not tuned. Logistic regression uses $C = 1$ and at most $1{,}000$ iterations. Gaussian naive Bayes uses the library defaults. The MLP has hidden layers of 64 and 32 units and at most 300 iterations. The early-stopping rule is that of the HGB configuration, with a 10% validation split and 10 iterations without improvement. These settings were fixed in the pre-registration of MINI-36 (Table C2) and not tuned. KNN uses $k = 5$. XGBoost in the joint space uses early stopping, as in the base-learner study (Table H10).

d: Multiclass spaces.

For K classes, the multiclass space has $14 + 8K$ candidates. The candidates are 14 simplex weights over HGB, DT, and SVC, plus 4 class scales applied to each of the K classes for two base scores, HGB and the HGB–DT average. The four scales, 0.5, 0.7, 1.4, and 2, are two down-weights and the approximate reciprocals of the two. Development run MINI-18 fixed these scales (Table C2). In the joint space, the multiclass candidates are the simplex weights of the three tested pairs (3×14) and the HGB class scales ($4K$). The space adds the equal average of HGB with each of the seven members under the class scales ($7 \times 4K$), for $42 + 32K$ candidates.

e: Parameter sweeps.

On the synthetic population $\mathcal{Q}$, ten parameter variants were swept. The variants are the nine combinations of the fold count F and the quantile level q at $N_{\text{boot}} = 20{,}000$, and $N_{\text{boot}} = 2{,}000$ at the standard configuration $F{=}3$, $q{=}5\%$. All ten pass the pre-registered criterion, with Proposition 1 bounds from 2.0% to 4.9% (Table D7). The quantile level governs the trade-off between the bound and the yield. With a Bonferroni correction [87] over the ten variants ($\delta = 0.05/10 = 0.005$ per variant), every variant stays below 10%: the highest bound is 5.73% ($F{=}3$, $q{=}10\%$) and the lowest 2.75%, computed from Table D7.

The margin sweep over 0.0025, 0.005, 0.0075, and 0.01 is reported in Table D3. Every setting that passes with the default margin also passes with the three other margins. Feature selection with SVC fails with every margin, as it does with the default. With 16, 32, and 64 bags on $\mathcal{Q}$, the bound is 3.98%, 3.56%, and 2.69%, at yields of 32.8, 33.9, and 33.2 per 100. All three bag counts meet the criterion (Table C1).

TABLE D2. Frozen parameters of BRIDGE and of the evaluation protocol: value and basis of each.

Parameter	Value	Set	Tested in
Cross-fit folds F	3	Dev	$F \in \{2, 3, 5\}$ on $\mathcal{Q}$ (Table D7)
Replicates N_{boot}	20,000, multinomial over the triple counts	Dev	MC error, 20 seeds; 2,000 (Table C1)
Lower quantile q	5%, decision statistic (not the error level $\delta = 0.05$)	Dev	2.5/5/10% (Tables D7, D5)
Margin ε	0.005	Dev	0.0025–0.01 (Tables D3, D4)
HGB base learner	lr 0.1, $\leq$ 300 iter., $\leq$ 31 leaves, $\geq$ 20/leaf, no L2; early stop: 10 iter., 10% of bag	Lib	tuned HGB (Tables H13, D6)
Bags	$\leq$ 32 draws of 80% of SEARCH; kept if every class $\geq$ 10 (multiclass $\geq$ 2); averaged	Dev	16, 32, 64 bags on $\mathcal{Q}$ (Table C1)
DT and SVC members	DT: library defaults; SVC: $C = 1$, RBF, $\gamma =$ scale, Platt probabilities	Dev	SVC over KNN (Table C1)
Preprocessing	median imputation, standardization; mode imputation, ordinal codes (-1 unseen)	Std	fitted per bag on SEARCH only, with the WiDS exception in Table D1
Candidate space	173 binary configurations; $14 + 8K$ for K classes (composition below)	Dev	1,825 and 11,249 configurations (Table H7)
Training rows	2,430 = SEARCH 630 + POOL 1,800; 1,000 below 2,480 rows	Dev	1,000/2,430/6,000 (Table H6); SEARCH 360–1,000
Locked reference	$0.3N \leq 8{,}000$ if $N < 40{,}000$, else $\min(N-20{,}000,\ 100{,}000)$; $\mathcal{Q}$: 10^5	Dev	reference SE; two more seeds (Table C1)
Trials	100 per dataset; 200 per clinical task; 60 per multiclass dataset; 1 per $\mathcal{Q}$ task	Pow, Pre	LOCO_0 18.2%/0.8% at 50/100 (Table C1)[a]
Seeds	fresh per trial (rows, split, bags, folds, bootstrap); learners 42; reference fixed	Pre	$\mathcal{Q}$ seeds unused by pilots (Appendix B)

*Set. Dev: set in development, on the 18 development datasets and the development runs (Table C1), and frozen before any confirmatory run (Table C2). Lib: library defaults, with the changes stated in the row, fixed before any evaluation. Std: standard practice. Pow: power check. Pre: pre-registration. Tested in: the run or table in which each value was varied or checked. The margin criterion is met at every tested margin wherever it is met at 0.005, and SEARCH sizes 360, 630, and 1,000 were tested on $\mathcal{Q}$. a: 200 trials per clinical task, adopted after the LOCO-source condition failed at 100 (Table C2); 60 per multiclass dataset as pre-registered, against 40 in development. MC: Monte Carlo. SE: standard error. iter.: boosting iterations. lr: learning rate.

TABLE D3. Parameter sensitivity: margin sweep.

Setting	$\varepsilon = 0.0025$	$\varepsilon = 0.005$ (default)	$\varepsilon = 0.0075$	$\varepsilon = 0.01$
(A), weighted-F_1	357/4: 3.83 / 5.10 ✓	284/3: 4.23 / 5.95 ✓	234/2: 4.32 / 6.36 ✓	209/1: 3.74 / 5.67 ✓
(A), balanced accuracy	856/1: 0.80 / 0.91 ✓	848/1: 0.81 / 0.92 ✓	836/0: 0.46 / 0.53 ✓	825/0: 0.47 / 0.54 ✓
(A), macro-F_1	852/0: 0.45 / 0.52 ✓	829/0: 0.47 / 0.53 ✓	813/0: 0.48 / 0.55 ✓	794/0: 0.49 / 0.56 ✓
(A), feature selection, KNN	263/5: 6.06 / 9.75 ✓	225/1: 3.45 / 6.00 ✓	193/0: 2.30 / 4.20 ✓	169/0: 2.66 / 5.32 ✓
(A), feature selection, SVC	296/10: 8.20 / 11.52 ×	253/7: 7.74 / 11.41 ×	211/5: 7.73 / 12.22 ×	171/3: 7.42 / 12.61 ×
(D) 26, 1,000 rows	648/12: 4.06 / 4.71 ✓	558/6: 3.04 / 3.59 ✓	468/4: 2.91 / 3.52 ✓	404/3: 2.93 / 3.64 ✓
(D) 27, 2,430 rows	843/7: 2.16 / 2.43 ✓	688/1: 1.05 / 1.20 ✓	554/1: 1.32 / 1.54 ✓	440/1: 1.69 / 2.01 ✓
(D) 27, 6,000 rows	1088/10: 2.08 / 2.30 ✓	825/3: 1.37 / 1.57 ✓	641/1: 1.13 / 1.35 ✓	498/1: 1.48 / 1.88 ✓
(D) 5 multiclass	209/0: 1.97 / 2.86 ✓	197/0: 2.11 / 3.16 ✓	180/0: 2.34 / 3.69 ✓	165/0: 2.59 / 4.23 ✓
(C), net benefit $p_t = 0.2$	1301/0: 0.30 / 0.44 ✓	1199/0: 0.33 / 0.50 ✓	1146/0: 0.34 / 0.53 ✓	1080/0: 0.36 / 0.59 ✓
(B) $\mathcal{Q}$	529/10: 3.72 / – ✓	475/8: 3.56 / – ✓	419/4: 2.62 / – ✓	373/3: 2.51 / – ✓

*Setting: collection, metric, or design. (D) 26/27: the datasets of (D) in the training-size study; rows = training rows. Cells: accepted/false: risk bound % / LOCO % and the pre-registered criterion of the setting, recomputed from the stored 5% bootstrap lower bound of every trial (the deployed configuration does not depend on the margin). Risk bound: KL-ratio bound (real data) or Proposition 1 (population Q); LOCO is not defined on Q. Trials, yield, the LOCO_0 power check, and the number of units with a yield of at least 5 per 100 are recorded in the code repository.

f: Monte Carlo error.

The Monte Carlo error of the bootstrap lower bound at $N_{\text{boot}} = 20{,}000$ is small (MINI-37, Table C2). The check re-runs the bootstrap with 20 seeds per trial on all 1,400 tasks of $\mathcal{Q}$ and on 110 real-data trials. The median standard deviation of the decision statistic is 0.00014 on $\mathcal{Q}$ and 0.00009 on the real data, about 3% and 2% of ε. A single alternative seed changes the decision in 0.3% and 0.7% of the re-runs, and at least one of the 20 seeds does so in 1.0% (14 of 1,400) and 1.8% (2 of 110). At $N_{\text{boot}} = 2{,}000$, BRIDGE accepts 481 updates with 7 false improvements, against 475 with 8 at 20,000.

g: Matched-yield comparison.

In the matched-yield comparison on the 11 new datasets, BRIDGE and the corrected t-test are tied at 0% below 125 accepted updates (Fig. 3). Between 125 and 450 accepted updates, BRIDGE is 0.4 to 1.1 percentage points lower; the curves cross near 450, both at 4.0%, and by 550 the corrected t-test is 1.6 points lower (7.6% against 6.0%). BRIDGE is below the 5×2 cv test at every counted yield, by up to 4.0 points.

TABLE D4. Parameter sensitivity: each margin tested against the default 0.005.

Margin	Acc./false	Δ, pp [95% CI]	TOST p	Acc. ratio
(A)				
0.0025	357 / 4	+0.06 [-0.29; +0.35]	0.011	1.26
0.0075	234 / 2	-0.20 [-0.88; +1.78]	0.0498	0.82
0.01	209 / 1	-0.58 [-1.40; +0.04]	0.0498	0.74
(B)				
0.0025	529 / 10	+0.21 [-0.22; +0.81]	$< 10^{-4}$	1.11
0.0075	419 / 4	-0.73 [-1.64; -0.02]	0.018	0.88
0.01	373 / 3	-0.88 [-1.89; -0.04]	0.0498	0.79
(C)				
0.0025	1,301 / 0	+0.00 [+0.00; +0.00]	$< 10^{-4}$	1.09
0.0075	1,146 / 0	+0.00 [+0.00; +0.00]	$< 10^{-4}$	0.96
0.01	1,080 / 0	+0.00 [+0.00; +0.00]	$< 10^{-4}$	0.90
(D)				
0.0025	1,375 / 13	+0.41 [+0.09; +0.86]	$< 10^{-4}$	1.24
0.0075	948 / 3	-0.22 [-0.63; +0.02]	$< 10^{-4}$	0.85
0.01	778 / 2	-0.28 [-0.69; +0.01]	$< 10^{-4}$	0.70

*Acc./false: accepted / false improvements. Δ: false-improvement rate minus the false-improvement rate at the default margin, in percentage points [95% CI]. TOST p: two one-sided tests of equivalence within ±2 percentage points, Holm-adjusted. Each of the 12 comparisons is equivalent. Acc. ratio: accepted updates, relative to the default margin. For margin ε, the decision compares the stored 5% lower bound with ε; all margins are therefore judged on the same trials. Default accepted / false: (A) 284 / 3, (B) 475 / 8, (C) 1,199 / 0, (D) 1,113 / 6. Brackets: cluster bootstrap over datasets, tasks or clinical tasks, 10,000 resamples; TOST with equivalence margin ±2 percentage points, Holm over the 12 comparisons; specification committed before the computation. In (C) no accepted update is a false improvement.

TABLE D5. Parameter sensitivity: quantile level q against margin ε on the real data.

q	$\varepsilon = 0.0025$	$\varepsilon = 0.005$	$\varepsilon = 0.0075$	$\varepsilon = 0.01$
(A)				
2.5%	299/3, 3.99 eq.	243/2, 4.15 n.e.	202/1, 3.89 n.e.	164/0, 2.76 n.e.
5%	357/4, 3.83 eq.	284/3, 4.23 def.	234/2, 4.32 n.e.	209/1, 3.74 n.e.
10%	422/11, 5.94 n.e.	357/4, 3.83 eq.	297/3, 4.02 eq.	251/2, 4.00 n.e.
(D)				
2.5%	1,180/7, 1.52 eq.	972/5, 1.52 eq.	799/2, 1.17 eq.	669/0, 0.61 eq.
5%	1,375/13, 1.96 eq.	1,113/6, 1.47 def.	948/3, 1.19 eq.	778/2, 1.21 eq.
10%	1,611/25, 2.68 eq.	1,326/12, 1.92 eq.	1,107/6, 1.48 eq.	934/3, 1.21 eq.

*Cells: accepted / false improvements, KL-ratio bound (%), and equivalence of the false-improvement rate with the default ($q = 5\%, \varepsilon = 0.005$) by a paired TOST with margin ± 2 percentage points and Holm correction over the 11 variants of a setting (eq.: equivalent; n.e.: not shown equivalent; def.: the default; which with 11 datasets reflects limited power, not a difference). All cells also meet LOCO, LOCO_0 and yield. Decisions at every (q, ε) are taken from the same cross-fit counts and bootstrap draws as the default. Pre-registration MINI-44 (run list: Table C2).

TABLE D6. Parameter sensitivity: tests of tuning on the real data.

Design	Naive	BRIDGE	Bound ↓	LOCO ↓	LOCO_0 ↓	Crit.
(A)						
Paper	1,064 / 203 (19.1)	284 / 3 (1.1)	4.23	5.95	2.11	✓
Tuned HGB	1,061 / 246 (23.2)	211 / 5 (2.4)	7.73	10.66	3.02	×
baseline change, F_1 points: +0.30 [+0.04; +0.61]						
Gate tuned	1,064 / 203 (19.1)	299 / 2 (0.7)	3.30	4.59	1.97	✓
(D)						
Paper	3,562 / 445 (12.5)	1,113 / 6 (0.5)	1.47	1.62	0.39	✓
Tuned HGB	3,523 / 540 (15.3)	906 / 10 (1.1)	2.53	2.81	0.49	✓
baseline change, F_1 points: +0.25 [+0.13; +0.39]						
Gate tuned	3,562 / 445 (12.5)	1,217 / 7 (0.6)	1.47	1.61	0.35	✓

*Naive, BRIDGE: accepted / false improvements (false %). Crit.: pre-registered criterion met (✓) or not met (×). Paper: the paper's design, fixed HGB and frozen gate. Same trials as the paper's runs (same seeds, splits and bags; the paper-design rows reproduce the stored decisions exactly in all 4,800 trials). Tuned HGB: tuned baseline, Optuna TPE, 25 iterations, 3-fold cross-validation inside SEARCH only, the HGB space and budget of the tuning runs; the tuned HGB replaces the fixed one as current model and member. Gate tuned: gate tuned per task, POOL split into two thirds and one third; the gate runs on the two thirds for every $q \in \{2.5, 5, 10\}\%$ and $\varepsilon \in \{0.0025, 0.005, 0.0075, 0.01\}$, the chosen candidate is scored on the third, and the best pair (ties: the default) is applied on the full POOL. Brackets: cluster bootstrap over datasets, 10,000 resamples. Pre-registered criterion as in Table 6. Pre-registration MINI-44 (run list: Table C2). Bound, LOCO, LOCO_0 in %.

TABLE D7. Parameter sensitivity on the synthetic tasks (B).

Variant	Acc.	False	U, % ↓	Cons., % ↓	Yield ↑	Gain ↑
F=2, q=2.5%	382	2	2.00	2.45	27.3	13.81
F=2, q=5%	445	4	2.47	4.41	31.8	15.09
F=2, q=10%	548	10	3.59	5.50	39.1	16.89
F=3, q=2.5%	406	4	2.70	3.08	29.0	14.34
F=3, q=5% **(standard)**	**475**	**8**	**3.56**	**4.70**	**33.9**	**15.61**
F=3, q=10%	563	16	4.90	6.24	40.2	17.12
F=3, q=5%, N_{boot}=2,000	481	7	3.22	4.37	34.4	15.73
F=5, q=2.5%	414	5	3.02	4.07	29.6	14.39
F=5, q=5%	483	10	4.07	5.70	34.5	15.62
F=5, q=10%	575	16	4.80	6.96	41.1	17.24

*Acc., False: accepted and false improvements; U: Proposition 1 bound at δ=0.05; Cons.: conservative variant (Corollary 1); yield: updates accepted per 100 decisions; gain: total reference gain. The standard configuration (bold) is mid-range.

h: Fixed validation margin.

A fixed validation margin offers a simple and strong alternative to the gate. The pre-registered margin t^*=0.05 was calibrated on the development size of 2,430 rows and fails under fresh seeds in the two settings that the pre-registration named: the tasks with 6,000 rows and the new datasets under weighted-F_1. A margin selected post hoc passes all settings. These comparisons are exploratory: the claim is that BRIDGE reaches a risk close to that of a well-chosen fixed margin without choosing a margin per task, not that BRIDGE dominates the margin. The fixed-margin paragraph of Appendix E gives the per-setting yields, the leave-one-setting-out analysis, and run MINI-29 (Table C2).

i: Parameter sensitivity on the real data.

Collection (D) passes every variant tested (Table H8). In (A), the 16- and 64-bag variants pass the criterion, but the TOST does not establish equivalence with the default within ± 2 percentage points. Two variants fail in (A): a bag drawn from all of SEARCH, a fraction of 100% (bound 13.95%, LOCO 21.33%), and a SEARCH of 1,000 rows, which finds no false improvement on the six datasets tested but fails the power check (LOCO_0 20.77%). The run labels and full numbers are in Table C1.

j: Computational cost.

The gate re-scores stored predictions instead of retraining, adding about 0.03 s per run, or 0.4–3.8% of the time to fit the bagged base models on SEARCH, and its inference cost equals that of the deployed predictor (Table D8). The 16-bag setting also meets the criterion (Table C1), so the fitting cost can be halved; the 16- and 64-bag variants are in Table H8.

TABLE D8. Wall-clock cost of BRIDGE.

	MIC	Mktg	Census	Homes.	Sant.
POOL rows	741	741	1,800	1,800	1,800
Features	104	25	41	297	369
Fit bags on SEARCH (s)	1.39	0.73	1.30	6.19	4.24
Naive selection (s) ↓	0.006	0.006	0.006	0.007	0.007
Cross-fitted selection (s) ↓	0.017	0.017	0.019	0.019	0.019
Bootstrap (s) ↓	0.011	0.011	0.010	0.009	0.006
BRIDGE gate total (s) ↓	0.027	0.028	0.029	0.027	0.024
Inference, baseline (ms/row) ↓	0.174	0.099	0.055	0.145	0.160
Inference, ensemble (ms/row) ↓	0.63	0.33	0.27	1.17	0.78
Gate / fit (%) ↓	1.96	3.84	2.20	0.44	0.58

*Columns: MIC; Mktg, Marketing_Campaign; Census, Census-Income-KDD; Homes., Homesite; Sant., Santander2016 (all datasets of (A)). Wall-clock cost (mean of 3 trials, one CPU core, Intel Core i9-13900K): fitting the bagged base models on SEARCH (32 bags each of HGB, DT and SVC), naive selection, the BRIDGE gate (cross-fitted selection and bootstrap) and inference latency of the deployed predictor (batch of 2,000 rows; baseline = 32 HGB bags, ensemble = 96 bagged models). Inference cost equals that of the deployed predictor, independent of BRIDGE. Gate total: cross-fitted selection plus bootstrap; percentages use unrounded times.

APPENDIX E PER-DATASET RESULTS AND PAIRED COMPARISONS

Tables of this appendix: p. 29.

For each rule, Table E1 gives the counts of accepted updates and false improvements on each of the 11 new datasets. The naive rule accepts the most false improvements on three datasets: KDD98 (83 accepted, 54 false improvements), Santander (90/39), and seismic (94/35). BRIDGE rejects every update on these three datasets.

a: Bags and skipped trials.

A drawn bag is retained when every class contributes at least 10 examples, so that each class is represented and the

TABLE E1. Per-dataset results on the 11 new datasets.

Dataset	Naive	5×2	Corr.	Re-sel.	No-m.	BRIDGE
Census-KDD	100/0	60/0	86/0	40/0	55/0	17/0
Homesite	99/3	49/0	71/0	41/0	53/0	30/0
KDD98	83/54	7/3	8/6	3/3	2/1	0/0
LICD	100/0	79/0	90/0	79/0	86/0	72/0
MIC	100/1	75/0	96/1	72/1	82/1	71/1
Marketing	100/0	78/0	98/0	81/0	87/0	76/0
Pumpkin	100/32	24/6	35/8	4/1	9/1	3/0
Rice	98/20	28/3	42/5	10/1	17/2	3/0
Santander	90/39	7/4	12/5	2/2	4/2	0/0
pc3	100/19	42/8	62/10	17/0	33/5	12/2
seismic	94/35	12/5	8/3	0/0	3/0	0/0
Total	1064/203	461/29	608/38	349/8	431/12	284/3

*Per-dataset results on the 11 new datasets, 100 trials each (1,100 in total). Columns: the naive rule; the 5×2 cv paired t-test; the corrected resampled t-test (Corr.); the corrected t with per-split re-selection (Re-sel.); BRIDGE without the margin (No-m.); BRIDGE. Dataset names shortened (Table B1). Cells: accepted / false improvements per rule; false improvement = gain ≤ 0 on the locked reference (Definition 1); 0/0 = no acceptance (false-improvement rate undefined). Totals reproduce Table H3.

TABLE E2. Paired comparisons: the difference BRIDGE minus each comparator.

Comparator	Δ, pp [95% CI]	p	Holm p
Real data (11 new datasets)			
naive	-18.0 $[-30.2; -7.1]$	$<0.0002^{\dagger}$	$<0.0012^{\dagger}$
McNemar, same POOL	-11.0 $[-19.2; -1.5]$	0.0026	0.0078
5×2 cv t	-5.2 $[-12.2; -1.4]$	0.0006	0.0024
corrected t	-5.2 $[-11.7; -1.6]$	0.0002	0.0010
BRIDGE without margin	-1.7 $[-4.8; -0.3]$	0.0046	0.0092
corrected t, re-selection	-1.2 $[-6.3; +1.3]$	0.32	0.32
same-POOL cell, own point	-1.6 $[-5.1; +0.2]$	0.22	–
same-POOL cell, matched yield	-0.4 $[-1.9; +1.0]$	1.00	–
margin 0.025, 11 new datasets	-0.2 $[-2.2; +0.6]$	0.68	–
margin 0.025, 1,000 rows	-2.2 $[-4.2; -0.6]$	0.0008	–
Synthetic population $\mathcal{Q}$			
naive	-19.0 $[-21.3; -16.8]$	$<0.0002^{\dagger}$	$<0.0012^{\dagger}$
McNemar, same POOL	-33.1 $[-41.2; -25.4]$	$<0.0002^{\dagger}$	$<0.0012^{\dagger}$
5×2 cv t	-3.9 $[-5.8; -2.1]$	$<0.0002^{\dagger}$	$<0.0012^{\dagger}$
corrected t	-4.0 $[-5.7; -2.4]$	$<0.0002^{\dagger}$	$<0.0012^{\dagger}$
BRIDGE without margin	-1.8 $[-3.1; -0.7]$	$<0.0002^{\dagger}$	$<0.0012^{\dagger}$
corrected t, re-selection	-1.3 $[-2.8; 0.0]$	0.051	0.051
same-POOL cell, own point	-1.5 $[-2.7; -0.5]$	0.0028	–
same-POOL cell, matched yield	$+1.1$ $[+0.2; +2.1]$	0.012	–
margin 0.025, 11 new datasets	-0.3 $[-1.4; +0.9]$	0.64	–
margin 0.02	-2.5 $[-4.1; -1.0]$	0.0018	–

*Δ: false-improvement rate of BRIDGE minus that of the comparator, in percentage points [95% CI]. p: two-sided. Holm p: adjusted over the six comparators. The fixed margins are selected post hoc. Paired cluster bootstrap at each rule's own operating point: in every resample both rules are evaluated on the same resampled decisions; clusters are the 11 datasets (real data) or the 1,400 tasks ($\mathcal{Q}$); 10,000 resamples, seed 20261012; percentile 95% intervals. Two-sided p = 2 $\times$ min(share of resampled differences ≥ 0, share ≤ 0), resolution 2/10,000 = 0.0002.† No resampled difference fell on the opposite side of zero (0 of 10,000): the empirical p is below its resolution, which is not an upper confidence bound on the true tail probability (likewise for the Holm-adjusted value); resamples in which a rule accepts nothing are left out for that rule. Holm adjustment over the six comparators; each test has the null hypothesis of equal false-improvement rates against the alternative of a difference. The same-POOL cell rows (at each rule's own operating point, and at matched yield, i.e., the same number of top-ranked accepted updates) and the fixed-margin rows are post-hoc, exploratory comparisons with nominal intervals, outside the Holm family.

internal validation split of the base model remains useful. The multiclass run uses a minimum of 2 examples instead (Appendix C). When a bag is discarded, the bagged score averages the remaining bags of the 32 drawn. A trial is skipped when the minority class is too small for any bag to meet the rule. All skipped trials belong to the training-size study with 1,000 rows. In that study porto-seguro is skipped in all trials, and avida-hil6 is skipped in some. The counts are 100 and 3 for the confirmatory run MINI-20 and 40 and 2 for the development run MINI-17 (runs in Table C2). Every other run has complete trials.

b: Paired comparisons.

Table E2 reports cluster-bootstrap confidence intervals for the difference in false-improvement rate between BRIDGE and each comparator. Each interval is computed at the number of accepted updates of that rule. The gain over BRIDGE without the margin is small and significant. This gain is -1.7 pp ($[-4.8; -0.3]$) on the real data and -1.8 pp ($[-3.1; -0.7]$) on $\mathcal{Q}$.

The corrected t-test with per-split re-selection and the best fixed margin chosen post hoc reach false-improvement rates that do not differ significantly from the false-improvement rate of BRIDGE. Two fixed margins are exceptions: $t = 0.025$ at 1,000 training rows and $t = 0.02$ on $\mathcal{Q}$. In these two settings BRIDGE is lower by 2.2 and 2.5 percentage points. The table note gives the test construction, the Holm correction over the six comparisons, and the resolution of the p-values; the fixed-margin rows are exploratory and have nominal intervals.

The difference between BRIDGE and the same-POOL cell with the bootstrap lower bound and the margin, at 173 candidates, is also in Table E2. The difference is computed twice, at the number of accepted updates of each rule and at matched numbers. Section IV-F discusses both results.

c: Fixed-margin detail.

The pre-registered grid of fixed margins spans 0.005 to 0.050 in steps of 0.005. Chosen post hoc in sample, the values 0.020 to 0.035 pass all 7 settings. The seven settings are the tasks with 1,000, 2,430, and 6,000 training rows, the new datasets under weighted-F_1, balanced accuracy, and macro-F_1, and the population $\mathcal{Q}$.

The yield of a fixed margin depends strongly on the chosen value. Across the weighted-F_1 settings the yield ranges from 5.4 to 49.4 per 100 decisions, against 21.5 to 33.9 for BRIDGE (computed from the stored runs of Table D3). On the tasks with 6,000 rows, the fixed margin reaches a yield of 5.4 to 21.0, below the 30.6 of BRIDGE. Near t=0.05 the fixed margin accepts few updates at the larger sizes. The yields are 1.4 and 0.1, and the bounds are 16.1% and 100% (MINI-29, Table C2). By contrast, the yield of BRIDGE grows with the sample size: 21, 25, and 31 at 1,000, 2,430, and 6,000 rows (Table H6).

Against the best fixed margin, the false-improvement rate of BRIDGE does not differ significantly, either on the 11 new datasets or on $\mathcal{Q}$. At the smallest training size, BRIDGE has a lower false-improvement rate than the best fixed margin (-2.2 pp, nominal 95% CI $[-4.2; -0.6]$). On large data BRIDGE accepts more improvements, with a yield of 31 against 14 per 100. The difference is $+16.8$ ($[+10.2; +23.7]$), computed with the paired bootstrap of Table E2. An exploratory leave-one-setting-out analysis, added

after the pre-registered results, chooses the best fixed margin from six settings and tests it on the seventh, in turn. The fixed margin passes six of the seven settings and fails on the new datasets under weighted-F_1 (LOCO 10.99%, bound 8.50%), whereas BRIDGE passes all seven, with a bound of 4.23% on that setting.

d: Base-learner audit.

The base-learner audit (Table E3) shows that AUC ranges from 0.566 to 0.579 across learners, while the number of boosting rounds, log-loss, extreme-probability rate (below 0.01 or above 0.99), and fit time differ. The differences relevant to the gate therefore lie in the predicted probabilities and in iteration control, not in the ranking of cases (Section IV-G).

TABLE E3. Audit of the base learners.

Learner	AUC ↑	Log-loss ↓	Extreme, % ↓	wF1 ↑	Rounds	Fit, s ↓
HGB	0.578	0.675	0.00	0.552	16.67	0.028
XGBoost, default	0.566	0.810	1.68	0.559	300.00	0.350
XGBoost, regul.	0.579	0.673	0.00	0.520	7.50	0.032

*Extreme: share of predicted probabilities below 0.01 or above 0.99; wF1: weighted-F_1 at threshold 0.5; rounds: boosting rounds in use; fit: fit time per bag. Audit of the base learners on the locked reference data (median over 1,400 tasks; mean for the share of probabilities below 0.01 or above 0.99; descriptive).

e: Dataset-level bootstrap.

A dataset-level bootstrap upper bound, which accounts for dataset clustering rather than conditioning on the observed collection, gives 3.1% for BRIDGE and 29.4% for the naive rule on the 11 new datasets of (A), against 13.0% and 12.5% for the 5×2 cv and corrected t tests (Table C1). On $\mathcal{Q}$ the 5×2 cv and corrected t tests meet the criterion (bounds 8.14% and 7.95%, Table 11) while failing on the real data. Under a stricter 5% threshold, BRIDGE passes on $\mathcal{Q}$ (Proposition 1 bound 3.56%) and on the real-data KL-ratio bound (4.23%) but not on the real-data LOCO (5.95%).

APPENDIX F PATIENT-IDENTITY AUDIT

Tables of this appendix: p. 30.

The patient-identity audit (Table F1) records, for each of the six clinical sources, the identity key, the row and unique-key counts, the excess duplicate rows and the rows actually removed; de-duplication and filtering are applied in the loaders before any split. On the co-primary net-benefit metric, the few accepted changes that are not better than both "treat none" and "treat all" are concentrated in the HIV-1 cohort (ACTG 175), where the prevalence exceeds p_t (Table H14). The mammography source carries no patient identifier, so its inclusion cannot be verified at the patient level, and the primary result is unchanged when that source is removed (Table F2).

The naive rule accepts the most false improvements on three clinical tasks: BRFSS kidney disease (133 false improvements of 133 accepted), BRFSS stroke (144 of 151), and NHANES coronary heart disease (97 of 182). BRIDGE accepts no update on any of the three tasks. On the real data the same pattern appears: BRIDGE accepts no update on KDD98, Santander, and seismic (Table E1).

TABLE F1. Patient-identity audit of the clinical sources.

Source	Identity key	Rows	Unique keys	Dup.
WiDS 2020 (GOSSIS)	patient_id	91,713	91,713	0
one row per patient				
unique encounter_id: 91,713				
hospitals: 147				
COVID-19 Mexico	id (record id)	566,602	563,201	3,401
duplicates are excess rows				
ids occurring more than once: 1,486				
4,887 rows with such ids removed before filtering				
confirmed cases kept: 218,784				
NHANES 2017-2018	SEQN	9,254	9,254	0
one row per respondent (survey design)				
the 4 outcomes share respondents and form one LOCO cluster				
ACTG 175 (UCI 890)	none in UCI copy	2,139		
randomized trial, one row per participant (documentation)				
BRFSS 2015	SEQNO (state-level)			
one record per interviewed respondent (survey design)				
the 4 outcomes form one LOCO cluster				
Mammography (OpenML 310)	none			
identical (X, y) rows kept once				
identical X with conflicting labels removed				
patient independence cannot be verified				

*Dup.: excess duplicate rows. Indented rows: notes per source. Patient-identity audit of the clinical sources (de-duplication and filtering applied in the loaders before any split). Participant identifier unavailable in the UCI copy of ACTG 175 and patient identifier unavailable in the OpenML mammography data; the BRFSS record number SEQNO is unique only within a state.

TABLE F2. Patient-identity audit: clinical results by source.

Source or subset	BRIDGE	Unc.	Bound ↓	Cons. ↓	LOCO ↓	Naive
Net benefit, $p_t = 0.2$ (co-primary)						
ACTG 175	198/0	2	–	–	–	200/0
BRFSS (4 tasks)	210/0	0	–	–	–	684/282
COVID-19 Mexico	200/0	0	–	–	–	200/0
Mammography	10/0	1	–	–	–	200/0
NHANES (4 tasks)	401/0	1	–	–	–	782/171
WiDS ICU	180/0	0	–	–	–	200/0
All sources	1,199/0	4	0.33	1.07	0.50	2,266/453
without ACTG 175	1,001/0	2	0.39	0.92	0.67	2,066/453
without BRFSS	989/0	4	0.39	1.29	0.67	1,582/171
without COVID-19 Mexico	999/0	4	0.39	1.29	0.67	2,066/453
without Mammography	1,189/0	3	0.33	0.92	0.50	2,066/453
without NHANES	798/0	3	0.50	1.40	0.67	1,484/282
without WiDS ICU	1,019/0	4	0.39	1.27	0.65	2,066/453
Balanced accuracy (co-primary)						
ACTG 175	177/0	23	–	–	–	200/0
BRFSS (4 tasks)	783/0	0	–	–	–	800/0
COVID-19 Mexico	200/0	0	–	–	–	200/0
Mammography	200/0	0	–	–	–	200/0
NHANES (4 tasks)	713/8	104	–	–	–	800/8
WiDS ICU	200/0	0	–	–	–	200/0
All sources	2,273/8	127	0.82	7.10	1.25	2,400/8
without ACTG 175	2,096/8	104	0.88	6.45	1.42	2,200/8
without BRFSS	1,490/8	127	1.25	10.85	1.45	1,600/8
without COVID-19 Mexico	2,073/8	127	0.90	7.79	1.45	2,200/8
without Mammography	2,073/8	127	0.90	7.79	1.45	2,200/8
without NHANES	1,560/0	23	0.24	2.49	0.48	1,600/0
without WiDS ICU	2,073/8	127	0.90	7.79	1.45	2,200/8

*Confirmatory clinical run with 200 trials per task (800 for the four BRFSS or NHANES tasks, 2,400 for all sources). BRIDGE and Naive: accepted / false improvements. Unc.: the accepted updates that are reference-uncertain. False improvement: gain ≤ 0 on the locked reference; reference-uncertain: the Bayesian-bootstrap interval of the reference gain contains 0. Bound: KL-ratio bound; conservative: every accepted update that is not clearly improving counted as a false improvement; LOCO: leave-one-source-out. Single-source rows give counts only. Bound, Cons., LOCO in %.

APPENDIX G OPENML COLLECTION (E): DATASETS, EXTENSION, AND SENSITIVITIES

Tables of this appendix: p. 32.

The datasets of the OpenML collection (E) and its extension (E2), with the decisions and the curation flags of each, are listed in Table G1. The OpenML collection (E) is the 14-dataset collection of Section III-E, evaluated with the frozen gate and reported in Table H1. BRIDGE accepts 192 of 1,280 decisions with 2 false improvements (1.0%), bound 5.40%, LOCO 8.43%, LOCO_0 3.62%, yield 15, and 8 datasets with a yield of at least 5. The pass is not uniform by class count. The binary subgroup of (E) (11 datasets) fails narrowly on LOCO (10.30%), a limitation found after the results. The multiclass subgroup (3 datasets) has no false improvement, but its bound is too wide to pass (32.59%, with LOCO and LOCO_0 at 100%). The row is retained as first reported, not replaced by the extension.

A second pre-registered collection, E2, adds 19 datasets and 1,700 decisions under the same protocol and the same frozen gate; the amendment of its dataset list deviates from the curation rule and is recorded below. E2 passes jointly: BRIDGE accepts 452 updates with 7 false improvements (1.55%), bound 4.16%, LOCO 5.44%, LOCO_0 1.19%, yield 26.6, and 14 datasets with a yield of at least 5. The naive rule accepts 1,633 updates with 303 false improvements (18.6%). The difference in false-improvement rate, BRIDGE minus naive, is -17.0 pp (95% CI $[-24.3, -9.6]$).

The joint result is read by class count, in addition to the pooled row. The binary subgroup of E2 (14 datasets, 1,400 decisions) meets every criterion: 315 accepted with 5 false improvements (1.59%), bound 5.00%, LOCO 7.60%, LOCO_0 2.04%, yield 22.5, and 9 datasets with a yield of at least 5. The BRIDGE-minus-naive difference is -19.6 pp (95% CI $[-28.1, -10.7]$). This pass depends on `churn`, a dataset flagged [b] as artificial, which contributes 100 accepted updates and no false improvement; the dependence was found after the results (a post hoc limitation). Without the flagged datasets, the combined binary subgroup of (E) and (E2) meets the LOCO condition (22 datasets; naive 24.9%, BRIDGE 1.59%; Table G1). The binary subgroup of (E2) alone, without the datasets flagged *a*–*c*, exceeds the LOCO limit (13.49%).

The multiclass subgroup of E2 (5 datasets, 300 decisions) meets the requirements on the false-improvement rate and yield. This subgroup has 137 accepted updates, 2 false improvements (1.46%), and a yield of 45.7. The LOCO bound of the subgroup is 14.10%, above the 10% condition, so the subgroup fails the criterion. The BRIDGE-minus-naive difference is -5.5 pp (95% CI $[-7.7, 0.0]$).

Pooling the two pre-registered collections, E+E2 covers 33 datasets and 2,980 decisions and meets every criterion: BRIDGE accepts 644 with 9 false improvements (1.40%), bound 3.38%, LOCO 4.04%, LOCO_0 0.75%, yield 21.6, and 22 datasets with a yield of at least 5. The BRIDGE-minus-naive difference is -20.2 pp (95% CI $[-26.7, -13.8]$). The binary subgroup (25 datasets, 2,500 decisions) also meets every criterion: 486 accepted with 7 false improvements (1.44%), bound 3.87%, LOCO 4.96%, LOCO_0 1.09%, yield 19.4, and 15 datasets with a yield of at least 5. The BRIDGE-minus-naive difference is -22.6 pp (95% CI $[-30.1, -15.3]$).

a: Deviation in the E2 amendment.

The frozen E2 list contains `shuttle`, whose 60 trials all stop with a structural error because a class is too small after the reference split. An amendment, frozen before any replacement metric was read, replaced `shuttle` with `pendigits` (OpenML 32). The post-run provenance audit then found a conflict with the frozen curation rule. The 10,992 rows of `pendigits` come from 44 writers, and the rule excludes data grouped by writer, subject, or session, whereas the runner splits rows because the OpenML table carries no writer identifier. `pendigits` is therefore not treated as passing the primary curation. The deviation is recorded, the run is retained as executed, and no third dataset is chosen after the results are seen.

b: Provenance flags.

The same provenance audit flags three further datasets of (E2). These are `churn`, whose original source is lost and whose data are described as artificial; `Fitness_Club`; and `Is-this-a-good-customer`, whose origin cannot be verified. All three remain reported as pre-registered, and no replacement dataset is introduced. Table G1 also reports the results without the datasets flagged *a*–*c*. With these datasets removed, the collection meets the criterion (29 datasets; naive 22.6% false, BRIDGE 1.67%).

c: Sensitivity without `pendigits`.

Because of this deviation, the pooled claim rests on the sensitivity analysis without `pendigits`. This analysis covers 32 datasets and 2,920 decisions and meets every criterion: 587 accepted, 9 false improvements (1.53%), bound 3.73%, LOCO 4.55%, LOCO_0 0.85%, yield 20.1, and 21 datasets with a yield of at least 5. The BRIDGE-minus-naive difference is -20.5 pp (95% CI $[-27.1, -14.2]$). The multiclass subgroup of this analysis (7 datasets, 420 decisions) meets the requirements on the false-improvement rate and yield. This subgroup has 101 accepted updates, 2 false improvements (1.98%), and a yield of 24.0. The subgroup fails all three checks: the bound (10.87%), the LOCO condition (30.27%), and the power check (LOCO_0 13.02%).

d: Image stress tests and the group structure of SVHN.

Two multiclass datasets of the extension, CIFAR-10 and the Street View House Numbers (SVHN), are flattened images, not tabular application data: each image contributes 3,072 pixel features (32×32 pixels in three colour channels), and both serve as high-dimensional stress tests. OpenML 42140 holds a 10% stratified subsample of SVHN without an identifier of the source photograph, so a group split cannot audit the independence of crops from the same photograph.

TABLE G1. Datasets of the OpenML collection (E) and its extension (E2), with per-dataset decisions and curation flags.

Dataset	Source	N	K	Feat.	Naive	BRIDGE
(E) 14 OpenML datasets						
airline-satisfaction	OML 46920	129,880	2	21	99/6	22/0
KDDCup09-churn	OML 1112	50,000	2	230	80/62	0/0
bank-marketing	OML 1461	45,211	2	16	100/0	53/0
pulsar-HTRU2	OML 45558	17,898	2	8	97/45	0/0
jm1	OML 1053	10,885	2	21	100/2	63/0
support2	OML 46860	9,105	2	46	93/26	5/0
telco-customer-churn	OML 42178	7,043	2	19	99/22	15/1
phoneme	OML 1489	5,404	2	5	96/26	3/0
spambase	OML 44	4,601	2	57	98/30	5/1
sick	OML 38	3,772	2	29	100/16	4/0
bioresponse	OML 4134	3,751	2	1,776	94/57	1/0
ipums-la-97-small	OML 382	7,019	8	60	60/14	2/0
otto-group	OML 45548	61,878	9	93	60/7	5/0
corporate-credit-ratings	OML 46372	5,000	10	7	60/4	14/0
(E2) 20 OpenML and UCI datasets						
in-vehicle-coupon	OML 46937	12,684	2	24	86/32	4/2
wilt	OML 40983	4,839	2	5	100/1	28/0
APSFailure	OML 41138	76,000	2	170	99/2	1/0
iranian-churn	UCI 563	3,150	2	13	100/0	92/0
SECOM	UCI 179	1,567	2	590	74/38	0/0
steel-plates-faults	UCI 198	1,941	2	27	100/2	31/0
myocardial-infarction	UCI 579	1,700	2	111	99/40	10/1
OVA-breast	OML 1128	1,545	2	10,935	100/33	12/1
lending-club	OML 43729	9,578	2	13	100/4	24/0
TVS-loan-default	OML 43743	119,528	2	30	83/47	0/0
personal-loan	OML 43826	5,000	2	13	97/31	5/0
fitness-club[c]	OML 46927	1,500	2	6	96/30	3/0
good-customer[c]	OML 46938	1,723	2	13	99/22	5/1
churn[b]	OML 46915	5,000	2	19	100/0	100/0
wine-white	OML 40498	4,898	7	11	60/8	5/1
shuttle[f]	OML 40685	58,000	7	9	–	–
covertype	OML 180	110,393	7	54	60/3	11/0
CIFAR-10-small[e]	OML 40926	20,000	10	3,072	60/10	6/1
SVHN-small[d]	OML 42140	9,927	10	3,072	60/0	58/0
pendigits (amend.)[a]	OML 32	10,992	10	16	60/0	57/0

*(E): pre-registered OpenML collection; (E2): its pre-registered extension; both run through the released program with the same protocol (HGB, default candidate space, weighted-F_1; 100 trials per binary and 60 per multiclass dataset). Naive, BRIDGE: accepted / false improvements. Decisions per dataset: 100 (binary) or 60 (multiclass). K: classes. Feat.: features. OML: OpenML ID. Names shortened. Flags: [a] rows grouped by 44 writers, against the grouped-data rule (amendment for `shuttle`); [b] original source lost, same data as OpenML 40701, described there as artificial; [c] origin of the data not verifiable (Kaggle practice or unattributed data); [d] flattened image whose source-photograph grouping cannot be audited; [e] flattened image (stress test); [f] all 60 trials stopped with a structural error (a class too small after the reference split), replaced by the amendment. Flagged datasets are kept as pre-registered and reported; without the datasets flagged *a*–*c*, (E)+(E2) covers 29 datasets and still meets the criterion (naive 22.6% false, BRIDGE 479 accepted / 8 false improvements, 1.67%; bound 4.28%, LOCO 5.38%), and so do its 22 binary datasets (naive 24.9%, BRIDGE 1.59%; bound 4.61%, LOCO 6.23%); the binary subgroup of (E2) alone does not (11 datasets, LOCO 13.49%).

This gap is reported as a limitation, not as a confirmed violation. Without `pendigits` and SVHN, the multiclass subgroup (6 datasets, 360 decisions) has 43 accepted updates, 2 false improvements (4.65%), and a yield of 11.9, and fails the bound (30.27%), the LOCO condition (49.85%), and the power check (LOCO_0 21.47%).

e: Volkert supplement.

A later supplement adds Volkert (OpenML 41166), an anonymous numerical tabular benchmark with 58,310 rows, 180 features, 10 classes, and 60 complete trials. The dataset was frozen after the confirmatory runs, so its analysis is exploratory. On Volkert alone, BRIDGE accepts 4 updates and no false improvement, at a yield of 6.7. The rule of at least three datasets with a yield of at least 5 needs more than one dataset, so this result is reported separately. Pooled with E+E2, the supplement gives 34 datasets and 3,040 decisions that meet every criterion: BRIDGE accepts 648 updates with 9 false improvements (1.39%), at a bound of 3.36%, LOCO 4.01%, and LOCO_0 0.75%. The yield is 21.3, 23 datasets reach a yield of at least 5, and the BRIDGE-minus-naive difference is -19.9 pp (95% CI $[-26.4, -13.6]$).

Because the supplement is exploratory, the pooled row leaves the confirmatory claim unchanged. The multiclass subgroup of the pooled row (9 datasets, 540 decisions) has a LOCO bound of 10.56%, slightly above the 10% condition.

f: Collection outside the validated class range.

Under the same program and protocol as (E), 11 further multiclass OpenML datasets were run as an exploratory collection with 60 trials each. These datasets have 3, 4, 6, 11, or 26 classes, outside the validated range of 7 to 10 classes. The collection meets the criterion jointly (175 accepted, 2 false improvements; bound 5.89%, LOCO 8.02%). Its class-count groups with $K = 3$ and $K = 6$ fail the criterion. Each of the remaining groups holds one dataset and is described without a criterion. The result therefore varies with the number of classes (Table H8).

g: One-pass analysis.

A one-pass analysis repeats the comparison with one decision per dataset instead of the 60 or 100 repeated seeds. On E2, BRIDGE accepts 7 of 19 decisions with no false improvement, against 2 false improvements among 17 naive acceptances (11.8%); pooled over E+E2, BRIDGE accepts 8 of 33 with no false improvement, against 7 among 31 naive acceptances (22.6%). These numbers are reported for completeness, because the main analysis rests on the full repetitions and a single decision per dataset carries no cluster interval.

APPENDIX H DETAILED RESULT TABLES

Tables of this appendix: pp. 33–36.

This appendix gives the detailed result tables behind Sections III and IV, numbered in print order. Each of the five groups below lists its tables and the main-text subsection that uses each table. In all appendix tables, ↓ and ↑ mark columns where a smaller or a larger value is better. A bound is the one-sided 95% upper bound on the false-improvement rate. The criterion is the pre-registered criterion of Table 6.

a: Main results and hypothesis tests.

For every collection, Table H1 gives the counts of accepted updates and false improvements, the 95% intervals, and the

TABLE H1. Main results in detail, (A)–(E).

Setting	Decisions	Naive: accepted	Naive % false [95% CI] ↓	BRIDGE: accepted	BRIDGE % false [95% CI] ↓	Reduction, pp [95% CI] ↑	Upper bound $\hat{U}$ or U (%) ↓	False removed % [95% CI] ↑
(A) 11 new datasets	1,100	1,064	19.1 [8.1; 31.6]	284	1.1 [0.0; 4.0]	18.0 [7.1; 30.2]	4.23	99 [95; 100]
(B) Synthetic tasks	1,400	1,323	20.7 [18.5; 22.9]	475	1.7 [0.6; 2.9]	19.0 [16.8; 21.3]	3.56	97 [95; 99]
(C) Clinical tasks, net benefit	2,400	2,266	20.0 [4.3; 40.4]	1,199	0.0	20.0 [4.5; 40.3]	0.33	100 [100; 100]
(D) 39 supporting datasets	3,700	3,562	12.5 [7.4; 18.6]	1,113	0.5 [0.1; 1.2]	12.0 [7.2; 17.9]	1.47	99 [97; 100]
(E) OpenML collection, 14 datasets	1,280	1,236	25.6 [15.0; 37.8]	192	1.0 [0.0; 4.0]	24.6 [14.0; 36.0]	5.40	99 [98; 100]

CI = confidence interval. Reduction: naive minus BRIDGE, in percentage points. Intervals: cluster bootstrap of the random-selection analysis (seed 20261014); the paired intervals of Table H3 come from the comparator runs (seed 20261022), so the interval limits can differ in the first decimal. Over the 6 clinical sources, the interval is 0.0 to 31.9 pp. Upper bound: the KL-ratio bound $\hat{U}$ (13) on the real data, and the Proposition 1 bound U (12) on (B). False removed: Table H4.

TABLE H2. Hypothesis tests: estimates and p-values, collection by collection.

Test	Comparison (T1: BRIDGE minus)	(A)	(B)	(C)	(D)	(E)
T1 (14)	Naive	-18.0 ($p < 0.0012^*$)	-19.0 ($p < 0.0012^*$)	-20.0 ($p = 0.034^*$)	-12.0 ($p = 0.0009^*$)	-24.6 ($p < 0.0002^*$)
	McNemar, same POOL	-11.0 ($p = 0.0078^*$)	-33.1 ($p < 0.0012^*$)	0.0 ($p = 1.00$)	-6.1 ($p = 0.0009^*$)	–
	McNemar, cross-fit	–	–	0.0 ($p = 1.00$)	-2.0 ($p = 0.018^*$)	–
	5×2 cv paired t	-5.2 ($p = 0.0024^*$)	-3.9 ($p < 0.0012^*$)	-2.5 ($p = 0.034^*$)	-3.1 ($p = 0.0009^*$)	–
	Corrected resampled t	-5.2 ($p = 0.0010^*$)	-4.0 ($p < 0.0012^*$)	-2.6 ($p = 0.10$)	-3.6 ($p = 0.0009^*$)	–
	Corrected t, re-selection	-1.2 ($p = 0.32$)	-1.3 ($p = 0.051$)	-0.4 ($p = 0.34$)	-1.0 ($p = 0.0009^*$)	–
	BRIDGE without margin	-1.7 ($p = 0.0092^*$)	-1.8 ($p < 0.0012^*$)	-0.6 ($p = 0.34$)	-1.1 ($p = 0.0016^*$)	–
	Holdout 2/3 + 1/3	–	–	0.0 ($p = 1.00$)	$+0.3$ ($p = 0.47$)	–
	Holdout 1/2 + 1/2	–	–	0.0 ($p = 1.00$)	-0.3 ($p = 0.47$)	–
T2 (15)	Margins vs 0.005	equivalent ($p \leq 0.0498^*$)	equivalent ($p \leq 0.0498^*$)	equivalent ($p < 10^{-4\,*}$)	equivalent ($p < 10^{-4\,*}$)	–
T3 (16)	Random draw, pooled	54.2 vs 3 ($p < 10^{-25\,*}$)	98.4 vs 8 ($p < 10^{-47\,*}$)	239.7 vs 0 ($p < 10^{-176\,*}$)	139.0 vs 6 ($p < 10^{-67\,*}$)	49.2 vs 2 ($p < 10^{-23\,*}$)
	Random draw, within dataset or task	5.5 vs 3 ($p = 0.15$)	–	0.7 vs 0 ($p = 0.45$)	27.4 vs 6 ($p < 10^{-9\,*}$)	12.9 vs 2 ($p < 10^{-5\,*}$)

T1: false-improvement rate of BRIDGE minus that of the comparator, in percentage points; a negative value means that BRIDGE accepts a smaller share of false improvements. Intervals: Table H3. T2: the largest Holm-adjusted p among the margin comparisons. T3: false improvements expected in a random draw, against the count of BRIDGE. *: H_0 rejected at $\alpha = 0.05$. –: comparison not run. Details: Tables E2, H3, D4, and H4.

TABLE H3. Acceptance rules on (A)–(E): counts, bounds, yield, gain, and paired differences, in one format.

Rule	Accepted	False (%) ↓	Bound % ↓	LOCO % ↓	Yield /100 ↑	Gain>ε /100 ↑	95% CI, pp
(A) 11 new datasets, 1,100 decisions							
Naive	1,064	203 (19.1)	22.90	25.27	96.7	59.5	$[-30.2;-7.1]^*$
McNemar, POOL	183	22 (12.0)	24.52	30.68	16.6	12.5	$[-19.2;-1.5]^*$
McNemar, cross-fit	45	4 (8.9)	41.14	54.35	4.1	3.2	–
5×2 cv t	461	29 (6.3)	10.99	13.43	41.9	36.2	$[-12.2;-1.4]^*$
Corrected t	608	38 (6.2)	10.09	12.15	55.3	46.8	$[-11.7;-1.6]^*$
Corr. t, re-selection	349	8 (2.3)	5.92	7.88	31.7	29.8	$[-6.3;+1.3]$
BRIDGE, $\varepsilon = 0$	431	12 (2.8)	6.16	7.85	39.2	36.2	$[-4.8;-0.3]^*$
Holdout 2/3 + 1/3	153	2 (1.3)	6.94	–	13.9	13.4	–
Holdout 1/2 + 1/2	190	6 (3.2)	9.67	–	17.3	16.2	–
BRIDGE	284	3 (1.1)	4.23	5.95	25.8	24.9	–
(B) synthetic population Q, 1,400 tasks							
Naive	1,323	274 (20.7)	23.52	–	94.5	59.2	$[-21.3;-16.8]^*$
McNemar, POOL	135	47 (34.8)	45.19	–	9.6	4.5	$[-41.2;-25.4]^*$
McNemar, cross-fit	50	13 (26.0)	42.65	–	3.6	2.2	–
5×2 cv t	610	34 (5.6)	8.14	–	43.6	37.6	$[-5.8;-2.1]^*$
Corrected t	774	44 (5.7)	7.95	–	55.3	47.6	$[-5.7;-2.4]^*$
Corr. t, re-selection	531	16 (3.0)	5.19	–	37.9	35.9	$[-2.8;+0.0]$
BRIDGE, $\varepsilon = 0$	598	21 (3.5)	5.67	–	42.7	39.9	$[-3.1;-0.7]^*$
Holdout 2/3 + 1/3	261	5 (1.9)	4.77	–	18.6	18.0	–
Holdout 1/2 + 1/2	339	7 (2.1)	4.55	–	24.2	23.3	–
BRIDGE	475	8 (1.7)	3.56	–	33.9	32.8	–

Rule	Accepted	False (%) ↓	Bound % ↓	LOCO % ↓	Yield /100 ↑	Gain>ε /100 ↑	95% CI, pp
(C) 12 clinical tasks, 2,400 decisions							
Naive	2,266	453 (20.0)	22.68	24.88	94.4	62.4	$[-39.4;-4.3]^*$
McNemar, POOL	17	0 (0.0)	45.35	100.00	0.7	0.4	$[+0.0;+0.0]$
McNemar, cross-fit	4	0 (0.0)	100.00	100.00	0.2	0.1	$[+0.0;+0.0]$
5×2 cv t	1,460	36 (2.5)	3.92	4.56	60.8	55.7	$[-6.9;-0.4]^*$
Corrected t	1,594	41 (2.6)	3.97	4.55	66.4	59.7	$[-7.2;-0.3]$
Corr. t, re-selection	1,402	5 (0.4)	1.03	1.21	58.4	56.4	$[-1.2;+0.0]$
BRIDGE, $\varepsilon = 0$	1,432	8 (0.6)	1.34	1.56	59.7	57.1	$[-1.9;+0.0]$
Holdout 2/3 + 1/3	987	0 (0.0)	0.40	0.67	41.1	41.0	$[+0.0;+0.0]$
Holdout 1/2 + 1/2	1,105	0 (0.0)	0.35	0.56	46.0	46.0	$[+0.0;+0.0]$
BRIDGE	1,199	0 (0.0)	0.33	0.50	50.0	49.9	–
(D) 39 supporting datasets, 3,700 decisions							
Naive	3,562	445 (12.5)	14.19	14.60	96.3	58.1	$[-17.7;-7.1]^*$
McNemar, POOL	1,142	76 (6.7)	9.55	10.48	30.9	22.8	$[-10.4;-3.2]^*$
McNemar, cross-fit	676	17 (2.5)	5.01	5.85	18.3	15.7	$[-4.6;-0.4]^*$
5×2 cv t	1,715	62 (3.6)	5.24	5.56	46.4	38.0	$[-5.2;-1.6]^*$
Corrected t	2,234	93 (4.2)	5.63	5.90	60.4	47.5	$[-5.9;-2.0]^*$
Corr. t, re-selection	1,465	22 (1.5)	2.69	2.89	39.6	35.2	$[-1.9;-0.3]^*$
BRIDGE, $\varepsilon = 0$	1,712	28 (1.6)	2.75	2.93	46.3	40.1	$[-2.0;-0.4]^*$
Holdout 2/3 + 1/3	721	2 (0.3)	1.31	1.48	19.5	18.8	$[-0.1;+0.8]$
Holdout 1/2 + 1/2	824	7 (0.8)	2.22	2.48	22.3	21.2	$[-1.2;+0.3]$
BRIDGE	1,113	6 (0.5)	1.47	1.62	30.1	29.1	–
(E) 14 OpenML datasets, 1,280 decisions							
Naive	1,236	317 (25.6)	29.62	32.21	96.6	41.2	$[-36.0;-14.0]^*$
BRIDGE	192	2 (1.0)	5.40	8.43	15.0	14.1	–

McNemar, POOL: McNemar applied on the same POOL. BRIDGE, $\varepsilon = 0$: BRIDGE without the margin. Holdout a + b: a share of POOL selects and b tests. Gain>ε/100: accepted updates whose reference gain exceeds the margin $\varepsilon = 0.005$, per 100 decisions (post hoc, Table H5). Bound: KL-ratio bound, and Proposition 1 bound on (B). 95% CI: paired difference, BRIDGE minus the rule, in percentage points. *: significant after the Holm correction. –: not computed. In (A), for small validation gains ($0 < \hat{\Delta} \leq 0.01$), the naive rule accepts 459 changes, 164 of which are false improvements, and BRIDGE accepts none. Point differences and p-values are given in Table H2, and the criterion marks in Table 9.

TABLE H4. Acceptance rules: random-selection control and gain per update on (A)–(E).

		Accepted (% false)		Random draw (T3): false, expected vs BRIDGE		False removed	Gain per update
Setting	Decisions	Naive	BRIDGE	Pooled	Within dataset	(%) ↑	BRIDGE / naive ↑
(A)	1,100	1,064 (19.1)	284 (1.1)	54.2 vs 3 ($p < 10^{-25}$)	5.5 vs 3 (p 0.15)	99 [95; 100]	0.0385 / 0.0156
(B)	1,400	1,323 (20.7)	475 (1.7)	98.4 vs 8 ($p < 10^{-47}$)	–	97 [95; 99]	0.0329 / 0.0157
(C)	2,400	2,266 (20.0)	1,199 (0.0)	239.7 vs 0 ($p < 10^{-176}$)	0.7 vs 0 (p 0.45)	100 [100; 100]	0.0416 / 0.0233
(D)	3,700	3,562 (12.5)	1,113 (0.5)	139.0 vs 6 ($p < 10^{-67}$)	27.4 vs 6 ($p < 10^{-9}$)	99 [97; 100]	0.0239 / 0.0115
(E)	1,280	1,236 (25.6)	192 (1.0)	49.2 vs 2 ($p < 10^{-23}$)	12.9 vs 2 ($p < 10^{-5}$)	99 [98; 100]	0.0197 / 0.0063

*False removed: the share of the false improvements of the naive rule that BRIDGE rejects. Gain per update: the mean reference gain per accepted update, for BRIDGE and for the naive rule. Brackets: 95% cluster-bootstrap intervals. Post hoc.

TABLE H5. Acceptance rules on (A)–(E): accepted updates below and above the margin on the reference (post hoc).

Rule	(A)	(B)	(C)	(D)	(E)
Accepted updates gaining at most ε (%) ↓					
Naive	38.4*	37.3*	33.9*	39.7*	57.4*
McNemar, same POOL	24.6*	53.3*	47.1°	26.1*	–
McNemar, cross-fit	22.2°	38.0*	25.0*	14.2*	–
5×2 cv paired t	13.7*	13.8*	8.5*	18.1*	–
Corrected resampled t	15.3*	14.0*	10.2*	21.4*	–
Corrected t, re-selection	6.0°	5.3°	3.4*	11.1*	–
BRIDGE without margin	7.7*	6.5*	4.3*	13.3*	–
Holdout 2/3 + 1/3	3.9°	3.4°	0.2°	3.7°	–
Holdout 1/2 + 1/2	6.3°	3.8°	0.0°	5.0°	–
BRIDGE	3.5	3.4	0.2	3.3	6.3
Accepted updates gaining more than ε, per 100 decisions ↑					
Naive	59.5†	59.2†	62.4†	58.1†	41.2†
McNemar, same POOL	12.5°	4.5*	0.4*	22.8°	–
McNemar, cross-fit	3.2*	2.2*	0.1*	15.7*	–
5×2 cv paired t	36.2†	37.6†	55.7†	38.0†	–
Corrected resampled t	46.8†	47.6†	59.7†	47.5†	–
Corrected t, re-selection	29.8†	35.9†	56.4†	35.2†	–
BRIDGE without margin	36.2†	39.9†	57.1†	40.1†	–
Holdout 2/3 + 1/3	13.4*	18.0*	41.0*	18.8*	–
Holdout 1/2 + 1/2	16.2*	23.3*	46.0*	21.2*	–
BRIDGE	24.9	32.8	49.9	29.1	14.1

Margin $\varepsilon = 0.005$: the minimum improvement fixed in development (Table 2); a gain at most ε includes every false improvement (gain ≤ 0). *: BRIDGE significantly better than the rule; °: no significant difference; †: the rule significantly better (paired cluster bootstrap of T1, 10,000 resamples, Holm over the rules of a collection, $\alpha = 0.05$). –: not run, because the program used for (E) records the decisions of the naive rule and BRIDGE.

TABLE H6. Conditions of use: the metrics, the training sizes, the multiclass tasks, and feature selection.

Setting	Accepted	False	Bound % ↓	LOCO % ↓	Yield /100 ↑	Crit.
(A) weighted F_1	284	3	4.23	5.95	25.82	✓
(A) balanced accuracy	848	1	0.81	0.92	77.09	✓
(A) macro-F_1	829	0	0.47	0.53	75.36	✓
(A) feat. selection + threshold (KNN)	225	1	3.45	6.00	20.45	✓
(A) feat. selection + threshold (SVC)	253	7	7.74	11.41	23.00	×
(D) 26 datasets, 1,000 rows	558	6	3.04	3.59	21.49	✓
(D) 27 datasets, 2,430 rows	688	1	1.05	1.20	25.48	✓
(D) 27 datasets, 6,000 rows	825	3	1.37	1.57	30.56	✓
(D) 5 multiclass datasets	197	0	2.11	3.16	65.67	✓

*Rows: training rows. Crit.: pre-registered criterion met (✓) or not met (×).

TABLE H7. Component analysis: candidate-space size on the collections (A)–(E).

		173 (default)		1,825		11,249	
Data	Rule	False % ↓	Bound % ↓	False % ↓	Bound % ↓	False % ↓	Bound % ↓
(A)	Naive	19.08	22.90 ×	20.44	24.23 ×	19.46	23.16 ×
	Same-POOL	2.65	5.85 ✓	3.61	7.02 ✓	3.80	7.17 ✓
	BRIDGE	1.06	4.23 ✓	1.45	5.07 ✓	1.37	4.75 ✓
(B)	Naive	20.71	23.52 ×	17.36	19.93 ×	16.88	19.43 ×
	Same-POOL	3.18	5.14 ✓	3.62	5.48 ✓	4.84	6.84 ✓
	BRIDGE	1.68	3.56 ✓	1.04	2.60 ✓	0.41	1.57 ✓
(C)	Naive	19.99	22.68 ×	21.09	23.74 ×	21.75	24.41 ×
	Same-POOL	0.00	0.31 ✓	0.16	0.71 ✓	0.15	0.70 ✓
	BRIDGE	0.00	0.33 ✓	0.00	0.33 ✓	0.00	0.33 ✓
(D)	Naive	13.12	14.94 ×	12.70	14.44 ×	12.61	14.34 ×
	Same-POOL	1.17	2.27 ✓	2.04	3.31 ✓	2.48	3.82 ✓
	BRIDGE	0.65	1.80 ✓	0.62	1.71 ✓	0.42	1.37 ✓
(E)	Naive	27.65	32.13 ×	28.23	32.35 ×	27.45	31.50 ×
	Same-POOL	3.37	8.17 ×	5.93	11.25 ×	9.22	15.39 ×
	BRIDGE	1.17	6.13 ×	1.13	5.89 ✓	3.35	10.34 ×

*Column groups: the candidate-space size $|\mathcal{C}|$. False: false-improvement rate among the accepted updates. Same-POOL: the bootstrap lower bound and the margin, without cross-fitting. Panels. (A): the 11 new datasets. (B): the synthetic tasks. (C): the 12 clinical tasks. (D): the 34 binary supporting datasets. (E): the 11 binary datasets of the OpenML collection. On (E), the LOCO bound of BRIDGE is 10.30%, 9.88%, and 17.37% at the three sizes, hence the × at 173. At 11,249 the cell has 6 false improvements among 179 accepted, a change of +2.2 pp from 173 (95% cluster-bootstrap CI [0.04; 12.8] over datasets). The same-POOL cell has LOCO 11.95% at 173.

bounds behind the summary of Section IV-A. Table H2 gives the estimate, the interval, and the p-value of every comparison in the hypothesis tests of Section IV-B. Every setting uses the current model HGB with the fixed configuration and the default candidate space (173 binary and $14 + 8K$ multiclass configurations; Section III-A). The false-improvement rate is pooled over all decisions of a setting, with one fresh seed per decision. The 95% intervals are cluster-bootstrap percentile intervals over datasets, tasks, or clinical tasks (10,000 resamples). In (C) BRIDGE accepts no false improvement, so the upper bound is the informative quantity. Collection (E) combines 11 binary and 3 multiclass datasets. In the tests, T1 resamples datasets ((A), (D), (E)) or tasks ((B), (C)). The Holm adjustment runs over the comparators of a collection: six in (A) and (B), nine in (C) and (D), and one comparison without adjustment in (E). The intervals are not adjusted, so in (C) the interval of the corrected t-test excludes zero while the Holm-adjusted p does not. The values $p < 0.0012$ and $p < 0.0002$ mark the resolution limit, where no resample fell on the other side of zero. T2 compares the margins 0.0025, 0.0075, and 0.01 with 0.005 within ±2 percentage points, with a Holm adjustment over the 12 comparisons. T3 is the one-sided hypergeometric test of a random draw, run post

TABLE H8. Conditions of use: setting, evidence, and status.

Row	Condition	Setting	Evidence	Status
1	Metrics	(A)	H6	met
2	Training sizes	(D)	H6	met
3	Multiclass tasks	(D)	H6	met
4	Feature selection	(A)	H6	conditional
5	Supporting datasets	(D)	H1	met
6	OpenML collections	(E), (E2)	H1, App. G	conditional
7	Volkert supplement	(E)+(E2)	App. G	exploratory
8	Base learners	(A)–(E)	H10, H11	conditional
9	Ensemble members	(A)–(E)	H11	conditional
10	Joint candidate space	(A)–(E)	H11	met
11	Denser candidate grids	(A)–(E)	H7	conditional
12	Gate parameters (q, ε)	(A)–(D)	D4, D5	met
13	Frozen parameters	(A), (D)	App. D	conditional
14	Tuning as an update	(A), (B), (D)	H13, D6	conditional
15	Imbalanced data	(A)	H9	conditional
16	Reproducibility	(A)–(E)	App. C	met

Status: met, the pre-registered criterion is met (Table 6); conditional, the criterion is met under the condition given below; exploratory, one dataset lies outside the pre-registered collections. Post hoc: the binary subgroup of (E) misses LOCO, and the pass of binary E2 depends on `churn` (Appendix G). Evidence: the table or appendix that reports the results; the conditions are discussed in Section V-A. Rows: (1) the metrics weighted-F_1, balanced accuracy, and macro-F_1; (2) the training sizes $1{,}000$, $2{,}430$, and $6{,}000$ rows; (3) the 7 to 10 classes; (4) mutual-information feature selection, where the criterion is met with KNN and missed with SVC (LOCO); (5) the 39 datasets of (D); (6) the criterion met by (E) and (E2) jointly; (7) one dataset; (8) the base learners HGB, XGBoost, and MLP, where the criterion is met with iteration control (early stopping); (9) two alternative member pairs, where the criterion is met in (A) to (D) and missed in (E); (10) seven families and 711 configurations; (11) the $1{,}825$ and $11{,}249$ configurations, where the criterion is met in (A)–(D) at both sizes and in (E) at $1{,}825$; (12) the quantile level q and the margin ε; (13) the bag fraction and the SEARCH size, where the criterion is met in (D) and for part of the values in (A); (14) the 12 tuning cells, where the criterion is met in 2; (15) the minority-sensitive metrics (balanced accuracy, macro-F_1), where the criterion is met; (16) the replay of the stored runs and the released package.

TABLE H9. Conditions of use: class imbalance on five datasets of (A).

Minority prev.	Naive: acc. / false (%)	BRIDGE: acc. / false (%)	Bound % ↓	LOCO % ↓	LOCO_0 % ↓	Yield ↑	Crit.
weighted-F_1							
20%	494 / 29 (5.9)	138 / 1 (0.7)	5.83	11.78	6.63	27.6	×
10%	474 / 95 (20.0)	100 / 0 (0.0)	4.76	14.77	14.77	20.0	×
5%	403 / 139 (34.5)	36 / 0 (0.0)	16.42	48.81	48.81	7.2	×
balanced accuracy							
20%	500 / 0 (0.0)	477 / 0 (0.0)	0.80	1.01	1.01	95.4	✓
10%	500 / 0 (0.0)	429 / 0 (0.0)	0.91	1.20	1.20	85.8	✓
5%	499 / 2 (0.4)	398 / 0 (0.0)	0.99	1.34	1.34	79.6	✓
macro-F_1							
20%	499 / 2 (0.4)	371 / 0 (0.0)	1.07	1.50	1.50	74.2	✓
10%	500 / 0 (0.0)	440 / 0 (0.0)	0.88	1.15	1.15	88.0	✓
5%	500 / 0 (0.0)	387 / 0 (0.0)	1.02	1.40	1.40	77.4	✓

*Rows grouped by metric. Acc.: accepted updates; false: false improvements; Crit.: pre-registered criterion met (✓) or not met (×).

TABLE H10. Conditions of use: all tested base-learner settings on (B).

Base learner setting	Acc./ false	U % ↓	Cons. % ↓	Yield ↑	Gain ↑	Naive % ↓	Skill ↑	Crit.
HGB, default (early stop.)	462/7	3.35	4.54	33.00	15.95	22.00	0.009	✓
XGBoost, regularized	893/5	1.41	1.74	63.79	61.28	9.23	0.012	✓
XGBoost, default (300 rounds)	146/15	17.48	23.21	10.43	1.96	31.87	-0.190	×
+ early stop.	963/1	0.59	1.15	68.79	71.04	6.75	0.011	✓
+ min_child_weight 5	138/18	21.08	26.99	9.86	1.77	31.59	-0.146	×
+ 31-leaf trees	125/9	14.21	21.24	8.93	1.69	32.34	-0.194	×
+ reg_lambda 0	122/10	15.62	21.74	8.71	1.68	33.65	-0.229	×
XGBoost, 30 rounds	287/3	3.25	6.33	20.50	7.26	26.52	-0.006	✓
XGBoost, 100 rounds	167/10	11.54	14.61	11.93	2.42	31.58	-0.076	×
MLP, early stop.	780/6	1.80	2.35	55.71	39.73	6.58	-0.010	✓
MLP, 300 epochs	373/7	4.14	6.67	26.64	9.18	18.06	-0.671	✓

*Acc./false: accepted and false improvements. U: Proposition 1 bound. Cons.: conservative variant (Corollary 1). Yield: accepted per 100. Gain: total reference gain. Naive %: false-improvement rate of the naive rule. Skill: median POOL log-loss skill. Crit.: criterion met (✓) or not met (×). +: one setting of the project default changed.

TABLE H11. Conditions of use: ensemble members and base learners across (A)–(E).

Design	(A)	(B)	(C)	(D)	(E)
Paper's design					
Default (DT, SVC)	1.1 ✓	1.7 ✓	0.0 ✓	0.5 ✓	1.0 ✓
Members replaced (base HGB)					
LR and NB	2.3 ✓	0.5 ✓	0.2 ✓	0.9 ✓	2.5 ×
MLP and KNN	1.2 ✓	1.5 ✓	0.0 ✓	1.9 ✓	4.1 ×
Base learner replaced (members DT, SVC)					
XGBoost, 300 rounds	4.1 ×	11.6 ×	0.1 ✓	1.7 ✓	4.2 ×
XGBoost, early stop.	0.8 ✓	0.1 ✓	0.0 ✓	0.8 ✓	0.7 ✓
MLP, early stop.	0.2 ✓	0.4 ✓	0.0 ✓	0.0 ✓	0.1 ✓
All seven member families (711 configurations)					
Seven families	1.6 ✓	0.9 ✓	0.2 ✓	0.7 ✓	1.4 ✓

*Cells: the false-improvement rate of BRIDGE among accepted updates (%), with the pre-registered criterion (✓ met, × not met). Metric: weighted-F_1; (C): net benefit, $p_t = 0.2$.

TABLE H12. Conditions of use: final model quality of four base learners.

Base learner	Baseline mean	Final mean	Final minus HGB, F_1 points [95% CI]	Better / equal / worse
(A) 11 new datasets				
HGB	0.855	0.864	+0.000 [+0.000; +0.000]	0 / 11 / 0
XGB0	0.863	0.864	-0.016 [-0.235; +0.210]	5 / 0 / 6
XGBES	0.849	0.866	+0.203 [-0.050; +0.521]	4 / 1 / 6
MLPES	0.823	0.862	-0.217 [-0.727; +0.367]	3 / 0 / 8
(D) 39 supporting datasets				
HGB	0.819	0.827	+0.000 [+0.000; +0.000]	0 / 39 / 0
XGB0	0.820	0.826	-0.072 [-0.259; +0.112]	16 / 0 / 23
XGBES	0.810	0.827	+0.012 [-0.093; +0.130]	15 / 2 / 22
MLPES	0.782	0.823	-0.415 [-0.587; -0.250]	9 / 0 / 30

Final mean: the mean weighted-F_1 on the locked reference after BRIDGE. better / equal / worse: datasets whose final weighted-F_1 is above, equal to, or below that of HGB. XGB0: XGBoost with 300 rounds, no early stopping. XGBES: XGBoost with early stopping. MLPES: a multilayer perceptron with early stopping.

hoc. Its within-dataset form keeps the count of BRIDGE in each dataset or task.

b: Comparison of acceptance rules.

In one format for (A) to (E), Table H3 gives the counts, bounds, LOCO, yield, and updates above the margin of every rule, including the two holdouts. The table also gives the 95% interval of the paired difference. Table H4 gives the random-selection control for question 6 (Section IV-B). For each rule, Table H5 gives the accepted updates that gain at most and more than the margin (post hoc). Among the updates that BRIDGE accepts, the reference gain of the final configuration reaches the cross-fit lower bound in 91%, 87%, 96%, 80%, and 82% of cases in (A) to (E), and the median excess $\Delta_{\text{ref}} - Q_{0.05}(\Delta^*)$ is 0.016, 0.013, 0.011, 0.007, and 0.008. The shortfall $\max\{Q_{0.05}(\Delta^*) - \Delta_{\text{ref}}, 0\}$ averages 0.0008, 0.0008, 0.0001, 0.0011, and 0.0007 over all accepted updates, and 0.0089, 0.0058, 0.0026, 0.0055, and 0.0040 over the updates that fall short; of those updates, 3 of 25, 8 of 64, 0 of 49, 6 of 219, and 2 of 34 are false improvements (post hoc, stored runs). The candidate selected in each fold was not stored, so the stability of that choice across folds

TABLE H13. Conditions of use: tuning treated as an update on (A), (B), and (D).

Learner	Naive acc./ false (%)	BRIDGE acc./ false (%)	Bound % ↓	LOCO % ↓	LOCO_0 % ↓	Yield ↑	Crit.
(A) 11 new datasets							
HGB	526 / 97 (18.4)	86 / 6 (7.0)	23.86	43.42	10.49	7.8	×
DT	677 / 120 (17.7)	171 / 6 (3.5)	10.88	20.29	5.46	15.5	×
SVC	503 / 92 (18.3)	163 / 6 (3.7)	11.48	24.84	6.00	14.8	×
KNN	571 / 169 (29.6)	39 / 4 (10.3)	49.17	59.92	33.67	3.5	×
(B) Synthetic tasks							
HGB	912 / 129 (14.1)	180 / 1 (0.6)	4.44	–	–	12.9	✓
DT	816 / 146 (17.9)	136 / 4 (2.9)	11.21	–	–	9.7	×
SVC	679 / 65 (9.6)	249 / 1 (0.4)	3.10	–	–	17.8	✓
KNN	732 / 232 (31.7)	96 / 4 (4.2)	16.69	–	–	6.9	×
(D) 3 multiclass datasets							
HGB	91 / 27 (29.7)	1 / 0 (0.0)	100.00	100.00	100.00	0.6	×
DT	55 / 13 (23.6)	7 / 0 (0.0)	100.00	100.00	100.00	3.9	×
SVC	147 / 11 (7.5)	73 / 0 (0.0)	6.55	22.63	22.63	40.6	×
KNN	136 / 11 (8.1)	38 / 1 (2.6)	26.25	52.91	39.69	21.1	×

*Acc./false: accepted and false improvements (false %). Crit.: criterion met (✓) or not met (×). Bound: KL-ratio bound on the real data, Proposition 1 bound on the synthetic population, where LOCO and LOCO_0 are not defined.

TABLE H14. Clinical application (C): consequences for each 1,000 observations.

		Per 1,000			Net benefit		
Task	Acc.	Flagged	TP	FP	Before	After	All
BRFSS depression	197	88→291	65.3	137.8	0.0495	0.0803	-0.0116
NHANES diabetes	200	37→390	100.3	253.4	0.0133	0.0502	-0.0467
NHANES asthma	1	0→85	22.2	62.9	0.0000	0.0064	-0.0620
BRFSS asthma	13	1→124	31.0	92.2	0.0002	0.0082	-0.0812
ICU (WiDS)	180	9→112	35.8	67.3	0.0059	0.0248	-0.1421
Mammography	10	10→36	13.1	13.4	0.0071	0.0168	-0.2096
COVID-19 Mexico	200	79→258	61.6	117.4	0.0334	0.0656	-0.0964
ACTG 175 (HIV-1)	198	40→614	161.5	412.3	0.0124	0.0709	0.0542
NHANES hypert.	200	339→661	123.7	198.8	0.1895	0.2635	0.2262
All tasks	1,199	98→386	90.9	197.0	0.0505	0.0921	-0.0031

*Acc.: accepted updates. Per 1,000 observations: observations flagged before→after the update (predicted positive), and extra true (TP) and false (FP) positives. Net benefit (10) before the update, after the update, and under treating all (every observation flagged).

cannot be assessed. Every rule judges the same decisions with the same deployed configuration, except the holdouts, which deploy the configuration selected by the holdout. In (E) the released program records the decisions of the naive rule and of BRIDGE, and of no other rule. LOCO leaves out one dataset, or one source in (C). LOCO is not defined on (B), where every task is independent, and is not computed for the holdouts in (A). The paired intervals use the cluster bootstrap of T1. The random-selection control is a post hoc, exploratory reanalysis of the stored runs. In every setting the updates BRIDGE accepts are a subset of those the naive rule accepts. The pooled draw takes the number of accepted updates of BRIDGE at random from the accepted updates of the naive rule. This draw gives the exact hypergeometric p of a count of false improvements no larger than the count of BRIDGE. It compares BRIDGE with a random choice of the same size, not a probability about future tasks. The within-dataset draw keeps the count of BRIDGE in each dataset or clinical task, so that the test concerns the choice among the trials of one dataset. In (E) the p-value comes from 100,000 draws.

c: Candidate-space size.

Table H7 gives the results of the three rules at 173, 1,825, and 11,249 configurations on (A) to (E). The table provides the data behind Fig. 5 (Section IV-F). The grids keep the HGB, DT, and SVC ensemble members (Appendix D). The panel of (E) covers the binary datasets of the OpenML collection in run MINI-52, pre-registered after the runs of (A) to (D). The multiclass datasets of (D) and (E) are left out, because no equally dense multiclass grid is pre-registered. At 173 configurations the rows of (E) reproduce the stored decisions of the released program in all 1,100 trials. For (A) to (D), the accepted counts, LOCO, and yield at 173 configurations are in Table H3.

d: Conditions of use.

Table H8 summarizes the status of every condition of use tested in Section IV-G. The tables that follow give the evidence. These tables cover the metrics, training sizes, and multiclass tasks (Table H6), imbalanced data (Table H9), and the base learners on $\mathcal{Q}$ (Table H10). The evidence also covers the ensemble members and base learners on all five collections (Table H11), the final model quality (Table H12), and the tuning study (Table H13). Across metrics, training sizes, and multiclass tasks, the criterion is met everywhere except with feature selection and SVC, where LOCO exceeds 10%. The imbalance study uses the five large binary datasets of (A) whose locked reference keeps at least 800 minority cases at every level (Census-Income-KDD, Homesite, KDD98, LICD, and Santander2016). Per dataset and level, the rows are subsampled once with a fixed seed to the stated prevalence before the locked outer split, so the reference has the same prevalence. Each cell has 100 trials per dataset (500 decisions), with the trial code, gate, and candidate space unchanged (pre-registration MINI-41). On $\mathcal{Q}$, the base-learner study changes one setting of the project-default XGBoost (300 rounds without early stopping) at a time on 1,400 fresh synthetic tasks. The study also adds a regularized variant with 4 settings changed and 30 or 100 rounds without early stopping. The same study adds an MLP with hidden layers $(64, 32)$, with and without early stopping, as a pre-registered prediction test. The POOL log-loss skill is 1 minus the log-loss over that of a prevalence-only forecast, which is computable without reference data. In the member and learner study every design runs on the same trials, seeds, splits, and bags as the design of the paper, and no learner is tuned. The runs are MINI-45, MINI-46, and MINI-50 for (A) and (D), and MINI-53 for (B), (C), and (E). The bounds of the cells that miss the criterion on (E) are in Table C1, and the bound, LOCO, LOCO_0, and yield of every cell are recorded in the code repository. The final model quality uses the first 20 trials of every dataset, with paired cluster-bootstrap intervals (MINI-49). In the tuning study each learner (HGB, DT, SVC, KNN) is tuned with Optuna on SEARCH, and the tuned model is the update against the

same learner with its fixed settings. The study covers the 11 new datasets (100 trials each), the synthetic population (task seeds 2,900 to 4,299, one trial each), and three multiclass datasets (Fashion-MNIST, USPS, DryBean; 60 trials each). The other two multiclass datasets were not run; this deviation was recorded before the results were seen (pre-registrations MINI-39, MINI-40).

e: Clinical consequences.

For each clinical task, Table H14 translates the accepted updates into flagged cases, extra true and false positives per 1,000, and net benefit before and after the update (Section IV-E). The values are exploratory means over the accepted updates of each task (MINI-30, 200 trials per task). Treat all is negative when the prevalence is below $p_t = 0.2$, because flagging every observation then costs more in false positives than it gains in true positives. In every task the mean net benefit after the update exceeds both treat none (0) and treat all. In NHANES asthma, ACTG 175, and NHANES hypertension the model first reaches this level after the update.

DATA AVAILABILITY STATEMENT

The datasets used in this study are publicly available from OpenML, the UCI Machine Learning Repository, scikit-learn, the CDC NHANES/BRFSS and the open Secretaría de Salud COVID-19 (Mexico) releases, and the OpenML mammography source; the GOSSIS (WiDS) cohort is available only through its original Kaggle competition, so the repository contains only the downloader and loader for that cohort, not the raw data. Code, per-trial outputs, and additional implementation details, packaged for direct installation with `pip install bridge-gate`, are available in the repository https://github.com/DwiSunaryono/BRIDGE-ML, and the release used in this paper is archived on Zenodo (doi:10.5281/zenodo.23523985). A pre-registration index lists, per run, the pre-registration file, its recorded time, and the SHA-256 of the file and of the manifest, and the plan files are available to the editor and reviewers on request; each pre-registration fixed the criteria, the dataset list and a deterministic seed formula before the run. The runs were split across three machines, and Appendix C describes the cross-machine checks. A single script rebuilds the reported tables from the stored per-trial outputs and does not train a model. The random-selection control and the gain-loss reanalysis behind Table H4 are post hoc, exploratory analyses of the stored runs and were not pre-registered.

DECLARATION OF GENERATIVE AI AND AI-ASSISTED TECHNOLOGIES

During the preparation of this work, the author used AI-assisted tools for language and formatting support and for compliance checking. After using these tools, the author reviewed and edited the content as needed and takes full responsibility for the content of the publication.

CONFLICT OF INTEREST

The author declares no competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

REFERENCES

[1] S. Kaufman, S. Rosset, C. Perlich, and O. Stitelman, "Leakage in data mining: Formulation, detection, and avoidance," ACM Transactions on Knowledge Discovery from Data, vol. 6, no. 4, pp. 1–21, 2012.

[2] G. Varoquaux and V. Cheplygina, "Machine learning for medical imaging: methodological failures and recommendations for the future," npj Digital Medicine, vol. 5, p. 48, 2022.

[3] B. Recht, R. Roelofs, L. Schmidt, and V. Shankar, "Do ImageNet classifiers generalize to ImageNet?" in Proceedings of the 36th International Conference on Machine Learning (ICML), vol. 97, 2019, pp. 5389–5400.

[4] M. Roberts, D. Driggs, M. Thorpe, J. Gilbey, M. Yeung, S. Ursprung, A. I. Aviles-Rivero, C. Etmann, C. McCague, L. Beer, J. R. Weir-McCall, Z. Teng, E. Gkrania-Klotsas, J. H. F. Rudd, E. Sala, and C.-B. Schönlieb, "Common pitfalls and recommendations for using machine learning to detect and prognosticate for COVID-19 using chest radiographs and CT scans," Nature Machine Intelligence, vol. 3, pp. 199–217, 2021.

[5] D. Sculley, G. Holt, D. Golovin, E. Davydov, T. Phillips, D. Ebner, V. Chaudhary, M. Young, J.-F. Crespo, and D. Dennison, "Hidden technical debt in machine learning systems," in Advances in Neural Information Processing Systems 28 (NIPS 2015), 2015, pp. 2503–2511.

[6] A. Paleyes, R.-G. Urma, and N. D. Lawrence, "Challenges in deploying machine learning: A survey of case studies," ACM Computing Surveys, vol. 55, no. 6, pp. 1–29, 2023, art. no. 114.

[7] Z. C. Lipton, C. Elkan, and B. Naryanaswamy, "Optimal thresholding of classifiers to maximize F1 measure," in Machine Learning and Knowledge Discovery in Databases (ECML PKDD 2014), ser. Lecture Notes in Computer Science, vol. 8725, 2014, pp. 225–239.

[8] R. Caruana, A. Niculescu-Mizil, G. Crew, and A. Ksikes, "Ensemble selection from libraries of models," in Proceedings of the 21st International Conference on Machine Learning (ICML), 2004, p. 18.

[9] I. Guyon and A. Elisseeff, "An introduction to variable and feature selection," Journal of Machine Learning Research, vol. 3, pp. 1157–1182, 2003.

[10] C. Ambroise and G. J. McLachlan, "Selection bias in gene extraction on the basis of microarray gene-expression data," Proceedings of the National Academy of Sciences, vol. 99, no. 10, pp. 6562–6566, 2002.

[11] J. Bergstra and Y. Bengio, "Random search for hyper-parameter optimization," Journal of Machine Learning Research, vol. 13, pp. 281–305, 2012.

[12] B. Bischl, M. Binder, M. Lang, T. Pielok, J. Richter, S. Coors, J. Thomas, T. Ullmann, M. Becker, A.-L. Boulesteix, D. Deng, and M. Lindauer, "Hyperparameter optimization: Foundations, algorithms, best practices, and open challenges," Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery, vol. 13, no. 2, 2023, art. no. e1484.

[13] S. Varma and R. Simon, "Bias in error estimation when using cross-validation for model selection," BMC Bioinformatics, vol. 7, p. 91, 2006.

[14] G. C. Cawley and N. L. C. Talbot, "On over-fitting in model selection and subsequent selection bias in performance evaluation," Journal of Machine Learning Research, vol. 11, pp. 2079–2107, 2010.

[15] R. J. Tibshirani and R. Tibshirani, "A bias correction for the minimum error rate in cross-validation," The Annals of Applied Statistics, vol. 3, no. 2, pp. 822–829, 2009.

[16] D. Krstajic, L. J. Buturovic, D. E. Leahy, and S. Thomas, "Cross-validation pitfalls when selecting and assessing regression and classification models," Journal of Cheminformatics, vol. 6, p. 10, 2014.

[17] M. Feurer, A. Klein, K. Eggensperger, J. T. Springenberg, M. Blum, and F. Hutter, "Efficient and robust automated machine learning," in Advances in Neural Information Processing Systems 28 (NIPS 2015), 2015, pp. 2962–2970.

[18] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, "Optuna: A next-generation hyperparameter optimization framework," in Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining (KDD), 2019, pp. 2623–2631.

[19] J. Feng, S. Emerson, and N. Simon, "Approval policies for modifications to machine learning-based software as a medical device: A study of bio-creep," Biometrics, vol. 77, no. 1, pp. 31–44, 2021.

[20] S. Kapoor and A. Narayanan, "Leakage and the reproducibility crisis in machine-learning-based science," Patterns, vol. 4, no. 9, p. 100804, 2023.

[21] Q. McNemar, "Note on the sampling error of the difference between correlated proportions or percentages," Psychometrika, vol. 12, no. 2, pp. 153–157, 1947.

[22] T. G. Dietterich, "Approximate statistical tests for comparing supervised classification learning algorithms," Neural Computation, vol. 10, no. 7, pp. 1895–1923, 1998.

[23] C. Nadeau and Y. Bengio, "Inference for the generalization error," Machine Learning, vol. 52, no. 3, pp. 239–281, 2003.

[24] J. Feng, "Learning to safely approve updates to machine learning algorithms," in Proceedings of the Conference on Health, Inference, and Learning (CHIL). ACM, 2021, pp. 164–173.

[25] U.S. Food and Drug Administration, "Proposed regulatory framework for modifications to artificial intelligence/machine learning (AI/ML)-based software as a medical device (SaMD): discussion paper and request for feedback," FDA, https://www.fda.gov/media/122535/download, 2019.

[26] S. Zhang, T. Wu, and F. Qian, "Certifying model upgrades with slice-wise non-regression and incumbent fallback," 2026, arXiv:2609.13714.

[27] V. B. Balachandran, "Pay only for disagreement: certified no-regression verdicts for model updates with matching label-complexity bounds," 2026, arXiv:2609.17560.

[28] C. Dwork, V. Feldman, M. Hardt, T. Pitassi, O. Reingold, and A. Roth, "The reusable holdout: preserving validity in adaptive data analysis," Science, vol. 349, no. 6248, pp. 636–638, 2015.

[29] A. Blum and M. Hardt, "The ladder: a reliable leaderboard for machine learning competitions," in Proceedings of the 32nd International Conference on Machine Learning (ICML), 2015, pp. 37–45.

[30] A. N. Angelopoulos, S. Bates, E. J. Candès, M. I. Jordan, and L. Lei, "Learn then test: calibrating predictive algorithms to achieve risk control," The Annals of Applied Statistics, vol. 19, no. 2, pp. 1641–1662, 2025.

[31] X. Bouthillier, P. Delaunay, M. Bronzi, A. Trofimov, B. Nichyporuk, J. Szeto, N. Mohammadi Sepahvand, E. Raff, K. Madan, V. Voleti, S. E. Kahou, V. Michalski, T. Arbel, C. Pal, G. Varoquaux, and P. Vincent, "Accounting for variance in machine learning benchmarks," in Proceedings of Machine Learning and Systems (MLSys), vol. 3, 2021. [Online]. Available: https://proceedings.mlsys.org/paper_files/paper/2021/hash/0184b0cd3cfb185989f858a1d9f5c1eb-Abstract.html

[32] B. A. Nosek, C. R. Ebersole, A. C. DeHaven, and D. T. Mellor, "The preregistration revolution," Proceedings of the National Academy of Sciences, vol. 115, no. 11, pp. 2600–2606, 2018.

[33] National Institute of Standards and Technology, "Secure hash standard (SHS)," Federal Information Processing Standards Publication 180-4, 2015.

[34] A. J. Vickers and E. B. Elkin, "Decision curve analysis: a novel method for evaluating prediction models," Medical Decision Making, vol. 26, no. 6, pp. 565–574, 2006.

[35] S. Bates, T. Hastie, and R. Tibshirani, "Cross-validation: what does it estimate and how well does it do it?" Journal of the American Statistical Association, vol. 119, no. 546, pp. 1434–1445, 2024.

[36] R. Berk, L. Brown, A. Buja, K. Zhang, and L. Zhao, "Valid post-selection inference," The Annals of Statistics, vol. 41, no. 2, pp. 802–837, 2013.

[37] J. Taylor and R. J. Tibshirani, "Statistical learning and selective inference," Proceedings of the National Academy of Sciences, vol. 112, no. 25, pp. 7629–7634, 2015.

[38] V. Chernozhukov, D. Chetverikov, M. Demirer, E. Duflo, C. Hansen, W. Newey, and J. Robins, "Double/debiased machine learning for treatment and structural parameters," The Econometrics Journal, vol. 21, no. 1, pp. C1–C68, 2018.

[39] S. Saeb, L. Lonini, A. Jayaraman, D. C. Mohr, and K. P. Kording, "The need to approximate the use-case in clinical machine learning," GigaScience, vol. 6, no. 5, pp. 1–9, 2017.

[40] D. Arp, E. Quiring, F. Pendlebury, A. Warnecke, F. Pierazzi, C. Wressnegger, L. Cavallaro, and K. Rieck, "Dos and don'ts of machine learning in computer security," in Proceedings of the 31st USENIX Security Symposium, 2022, pp. 3971–3988.

[41] J. Demšar, "Statistical comparisons of classifiers over multiple data sets," Journal of Machine Learning Research, vol. 7, pp. 1–30, 2006.

[42] A. Benavoli, G. Corani, J. Demšar, and M. Zaffalon, "Time for a change: a tutorial for comparing multiple classifiers through Bayesian analysis," Journal of Machine Learning Research, vol. 18, no. 77, pp. 1–36, 2017.

[43] A. N. Angelopoulos, S. Bates, A. Fisch, L. Lei, and T. Schuster, "Conformal risk control," in The Twelfth International Conference on Learning Representations (ICLR), 2024.

[44] K. J. M. Janssen, K. G. M. Moons, C. J. Kalkman, D. E. Grobbee, and Y. Vergouwe, "Updating methods improved the performance of a clinical prediction model in new patients," Journal of Clinical Epidemiology, vol. 61, pp. 76–86, 2008.

[45] E. W. Steyerberg, Clinical Prediction Models: A Practical Approach to Development, Validation, and Updating, 2nd ed. Springer, 2019.

[46] P. Foret, A. Kleiner, H. Mobahi, and B. Neyshabur, "Sharpness-aware minimization for efficiently improving generalization," in International Conference on Learning Representations (ICLR), 2021.

[47] M. Arjovsky, L. Bottou, I. Gulrajani, and D. Lopez-Paz, "Invariant risk minimization," 2019, arXiv:1907.02893.

[48] S. Sagawa, P. W. Koh, T. B. Hashimoto, and P. Liang, "Distributionally robust neural networks for group shifts: On the importance of regularization for worst-case generalization," in International Conference on Learning Representations (ICLR), 2020.

[49] J. H. Friedman, "Greedy function approximation: A gradient boosting machine," The Annals of Statistics, vol. 29, no. 5, pp. 1189–1232, 2001.

[50] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "LightGBM: A highly efficient gradient boosting decision tree," in Advances in Neural Information Processing Systems 30, 2017, pp. 3146–3154.

[51] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay, "Scikit-learn: Machine learning in Python," Journal of Machine Learning Research, vol. 12, pp. 2825–2830, 2011.

[52] L. Breiman, "Bagging predictors," Machine Learning, vol. 24, no. 2, pp. 123–140, 1996.

[53] P. Bühlmann and B. Yu, "Analyzing bagging," The Annals of Statistics, vol. 30, no. 4, pp. 927–961, 2002.

[54] L. Breiman, J. H. Friedman, R. A. Olshen, and C. J. Stone, Classification and Regression Trees. Belmont, CA, USA: Wadsworth, 1984.

[55] C. Cortes and V. Vapnik, "Support-vector networks," Machine Learning, vol. 20, no. 3, pp. 273–297, 1995.

[56] T. M. Cover and P. E. Hart, "Nearest neighbor pattern classification," IEEE Transactions on Information Theory, vol. 13, no. 1, pp. 21–27, 1967.

[57] D. R. Cox, "The regression analysis of binary sequences," Journal of the Royal Statistical Society: Series B, vol. 20, no. 2, pp. 215–232, 1958.

[58] D. J. Hand and K. Yu, "Idiot's Bayes—not so stupid after all?" International Statistical Review, vol. 69, no. 3, pp. 385–398, 2001.

[59] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning representations by back-propagating errors," Nature, vol. 323, pp. 533–536, 1986.

[60] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD), 2016, pp. 785–794.

[61] B. Efron and R. J. Tibshirani, An Introduction to the Bootstrap. Chapman & Hall, 1993.

[62] C. J. van Rijsbergen, Information Retrieval, 2nd ed. London, U.K.: Butterworths, 1979.

[63] K. H. Brodersen, C. S. Ong, K. E. Stephan, and J. M. Buhmann, "The balanced accuracy and its posterior distribution," in Proceedings of the 20th International Conference on Pattern Recognition, 2010, pp. 3121–3124.

[64] M. Sokolova and G. Lapalme, "A systematic analysis of performance measures for classification tasks," Information Processing & Management, vol. 45, no. 4, pp. 427–437, 2009.

[65] T. Fawcett, "An introduction to ROC analysis," Pattern Recognition Letters, vol. 27, no. 8, pp. 861–874, 2006.

[66] D. B. Rubin, "The Bayesian bootstrap," The Annals of Statistics, vol. 9, no. 1, pp. 130–134, 1981.

[67] S. Boucheron, G. Lugosi, and P. Massart, Concentration Inequalities: A Nonasymptotic Theory of Independence. Oxford University Press, 2013.

[68] J. Langford, "Tutorial on practical prediction theory for classification," Journal of Machine Learning Research, vol. 6, pp. 273–306, 2005.

[69] C. A. Field and A. H. Welsh, "Bootstrapping clustered data," Journal of the Royal Statistical Society: Series B, vol. 69, no. 3, pp. 369–390, 2007.

[70] J. Vanschoren, J. N. van Rijn, B. Bischl, and L. Torgo, "OpenML: Networked science in machine learning," SIGKDD Explorations, vol. 15, no. 2, pp. 49–60, 2014.

[71] S. Holm, "A simple sequentially rejective multiple test procedure," Scandinavian Journal of Statistics, vol. 6, no. 2, pp. 65–70, 1979.

[72] D. J. Schuirmann, "A comparison of the two one-sided tests procedure and the power approach for assessing the equivalence of average bioavailability," Journal of Pharmacokinetics and Biopharmaceutics, vol. 15, no. 6, pp. 657–680, 1987.

[73] C. J. Clopper and E. S. Pearson, “The use of confidence or fiducial limits illustrated in the case of the binomial,” Biometrika, vol. 26, no. 4, pp. 404–413, 1934.
[74] D. Dua and C. Graff, “UCI machine learning repository,” University of California, Irvine, School of Information and Computer Sciences, 2019.
[75] Centers for Disease Control and Prevention, National Center for Health Statistics, “National health and nutrition examination survey data, 2017–2018,” U.S. Department of Health and Human Services, Hyattsville, MD, USA, 2018.
[76] Centers for Disease Control and Prevention, “Behavioral risk factor surveillance system survey data, 2015,” U.S. Department of Health and Human Services, Atlanta, GA, USA, 2015.
[77] Secretaría de Salud, Dirección General de Epidemiología, “Datos abiertos de COVID-19 (Dirección General de Epidemiología),” Gobierno de México, https://www.gob.mx/salud/documentos/datos-abiertos-152127, 2020.
[78] M. Lee, J. Raffa, M. Ghassemi, T. Pollard, S. Kalanidhi, O. Badawi, K. Matthys, and L. A. Celi, “WiDS (Women in Data Science) Datathon 2020: ICU mortality prediction (version 1.0.0),” PhysioNet, 2020.
[79] E. B. Wilson, “Probable inference, the law of succession, and statistical inference,” Journal of the American Statistical Association, vol. 22, no. 158, pp. 209–212, 1927.
[80] A. Kraskov, H. Stögbauer, and P. Grassberger, “Estimating mutual information,” Physical Review E, vol. 69, no. 6, p. 066138, 2004.
[81] L. Prechelt, “Early stopping—but when?” in Neural Networks: Tricks of the Trade, ser. Lecture Notes in Computer Science. Springer, 1998, vol. 1524, pp. 55–69.
[82] T. Gneiting and A. E. Raftery, “Strictly proper scoring rules, prediction, and estimation,” Journal of the American Statistical Association, vol. 102, no. 477, pp. 359–378, 2007.
[83] C. Spearman, “The proof and measurement of association between two things,” The American Journal of Psychology, vol. 15, no. 1, pp. 72–101, 1904.
[84] J. Bergstra, R. Bardenet, Y. Bengio, and B. Kégl, “Algorithms for hyper-parameter optimization,” in Advances in Neural Information Processing Systems 24 (NIPS 2011), 2011, pp. 2546–2554.
[85] S. G. Finlayson, A. Subbaswamy, S. Singh, J. Bowers, A. Kupke, J. Zittrain, I. S. Kohane, and S. Saria, “The clinician and dataset shift in artificial intelligence,” New England Journal of Medicine, vol. 385, no. 3, pp. 283–286, 2021.
[86] S. E. Davis, T. A. Lasko, G. Chen, E. D. Siew, and M. E. Matheny, “Calibration drift in regression and machine learning models for acute kidney injury,” Journal of the American Medical Informatics Association, vol. 24, no. 6, pp. 1052–1061, 2017.
[87] O. J. Dunn, “Multiple comparisons among means,” Journal of the American Statistical Association, vol. 56, no. 293, pp. 52–64, 1961.

DWI SUNARYONO (Member, IEEE) was born in Surabaya, Indonesia, in 1972. He received the bachelor’s degree in computer engineering, the master’s degree in informatics engineering, and the Ph.D. degree in computer science from the Department of Informatics, Faculty of Intelligent Electrical and Informatics Technology, Institut Teknologi Sepuluh Nopember (ITS), Surabaya. He is currently an Associate Professor with the Department of Informatics, ITS, and the Head of the Algorithm and Programming Laboratory. His research interests include machine learning evaluation and model updating, intelligent systems, face recognition, signal processing, the Internet of Things, and management information systems.