DECODABILITY IS NOT CONTROLLABILITY: LOCAL AND ANISOTROPIC SPEAKER CONTROL IN CONTINUOUS-LATENT TTS

*Yuanzhong Zheng, Wei Huang, Ming Ma, Kailun Jin, Youhai Jiang**

VoxNexus, Singapore

arXiv:2610.29345v1 [cs.SD] 21 Oct 2026

ABSTRACT

High linear decodability does not entail causal controllability in continuous-latent text-to-speech (TTS). In VoxCPM2, position-matched probes and equal-norm interventions across 30 speakers and 60 directed dyads expose this separation: reference-audio states at the final residual acoustic language model layer (RALM L7) retain 0.970 probe accuracy, yet aligned edits yield 0/60 target-dominant dyads. Edits at early reference-audio positions produce the opposite pattern. At L0, pair-aligned updates increase target similarity by 0.326 and achieve target dominance in all 60 dyads, against at most 1/60 under any equal-norm direction or position control. Independent speaker encoders, retrieval against held-out speaker prototypes, and double-blind listening confirm that the shifted voice is heard as the intended target. Cross-architecture interventions in dots.tts distinguish backbone leverage from direct acoustic control: backbone edits displace the source without target attraction, while x-vector ablation reduces source similarity from 0.849 to 0.006. These findings identify two requirements for training-free speaker control: geometric specificity and a computational route that carries the intended speaker condition into acoustic synthesis.

***Index Terms*—** text-to-speech, speaker control, causal intervention, linear probes, representation analysis

1. INTRODUCTION

Prompt-conditioned text-to-speech (TTS) systems synthesize target voices from reference audio [1–5]. While linear probes measure the speaker information available in hidden states [6, 7], the critical distinction lies between reading that information and using it to control the generated voice [8–10]. We expose a direct counterexample to selecting control sites by probe accuracy: VoxCPM2's L7 reference-audio states retain 0.970 accuracy yet yield 0/60 target-dominant dyads under aligned intervention, whereas L0 aligned edits yield 60/60. Locating speaker control therefore requires attending to both the direction of an edit and its route into synthesis.

Activation patching, causal mediation, and representation steering test model behavior under direct intervention [11–15]. In speech synthesis, recent work uses activation steering to manipulate attributes such as emotion, style, and accent [16–19]. We couple position-matched probing with equal-norm geometric controls to localize speaker control and distinguish target attraction from acoustic disruption.

Target-specific control requires both increasing similarity to target speaker B and making the output closer to B than to source speaker A. An update can also suppress the source voice or disrupt speech without producing the target voice. We distinguish **decodability**, the ability to predict speaker identity from a state; **leverage**, the effect of changing that state on speech; and **specificity**, whether the change steers speech toward the intended target.

We study these properties in the residual acoustic language model (RALM) prefill states of VoxCPM2 [1]. Across 30 English LibriTTS-R speakers [20] organized into 60 role-balanced directed dyads (evaluated over two prompts and three generation seeds), we vary layer depth, token position, and perturbation direction in an equal-norm factorial. Independent speaker encoders, retrieval against held-out speaker prototypes, and double-blind listening measure the identity of the generated voice.

First, the experiments establish a **sharp dissociation between decodability and target control**: L0 reference-audio edits achieve target dominance in all 60 dyads, whereas L7 edits yield 0/60 despite a probe accuracy of 0.970. Second, control is **anisotropic and position-specific**: pair-aligned updates at early reference-audio positions produce target attraction, whereas equal-norm orthogonal updates reveal source-displacing leverage, and final-state updates primarily affect intelligibility. The comparison isolates the directions and positions that give leverage its target specificity. Finally, **readout reach and direct conditioning identify the architectural routes for speaker control**. In VoxCPM2, subsequent attention blocks carry early reference-state edits to the synthesis readout. In dots.tts, backbone edits and x-vector ablation distinguish behavioral leverage from a source-identity anchor supplied directly to the acoustic head. These experiments identify the editable states and conditioning inputs that a speaker-control interface must address.

2. METHOD

2.1. Model and probe

VoxCPM2 encodes reference waveforms into AudioVAE patches, maps them through a local audio encoder, interleaves the resulting audio embeddings with text embeddings, and processes the sequence through a TSLM and an eight-layer RALM; a flow-matching LocDiT and the AudioVAE decoder then synthesize the waveform [1]. Interventions modify post-block RALM hidden states at prefill step 0 and retain the intervened block's already-computed key/value (KV) cache. Subsequent autoregressive steps use the standard forward computation. Early-block updates propagate through later RALM blocks and their KV caches. The final RALM readout selects the last sequence position to condition LocDiT. We define **attention-mediated readout reach** as the capacity of a reference-audio intervention to influence the final conditioning state through subsequent RALM attention blocks. Early reference-audio edits have this route; post-block L7 reference-audio edits lie beyond the final mixing block. Edits to the final sequence position at L7 reach the acoustic generator directly.

The matched probe contains 300 held-out utterances from the same 30-speaker universe as the causal experiment. At layers L0,

*Corresponding author.

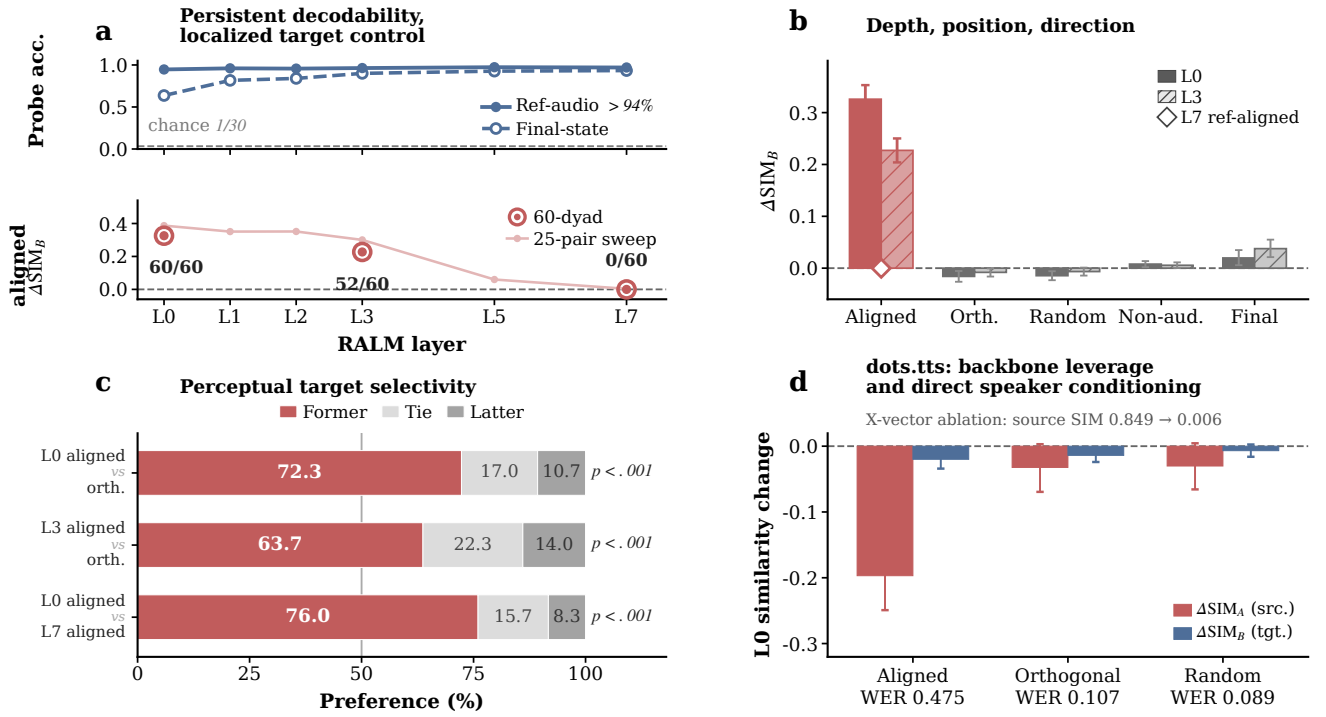

Fig. 1. (a) In VoxCPM2, speaker identity at reference-audio positions remains highly decodable (>94%, solid) while final-state decodability rises with depth (dashed). Bottom: 60-dyad target gain at L0/L3/L7 (markers) with target-dominant counts; a finer-grained six-layer sweep over 25 dyads (thin line) resolves a sharp decline in target control at L5. (b) Early reference-audio aligned updates produce the largest target gains among equal-norm direction and position conditions. (c) Held-out-reference listeners confirm perceptual direction and depth selectivity. (d) The dots.tts interventions separate backbone leverage from direct speaker conditioning. At L0, aligned updates produce the largest source displacement and highest WER; mean target gain is non-positive across directions. X-vector ablation reduces source similarity from 0.849 to 0.006. Uncertainty in (b,d) is evaluated by dyad-bootstrap (95% CIs; 60 directed dyads after averaging seeds and prompts within each dyad).

L1, L2, L3, L5, and L7, a linear classifier predicts speaker identity from the reference-audio mean, non-audio-token mean, final state, or whole-sequence mean. We use five-fold cross-validation grouped by transcript identity and 1,000 full-label permutations with a 24-condition max-statistic family-wise correction. The probe matches the causal experiment's speaker pool and reference/enrollment exclusions.

2.2. Equal-norm activation steering

For source A, target B, shared synthesis text x, and RALM layer ℓ, let h_A^ℓ and h_B^ℓ be post-block prefill states. Before subtraction, h_B^ℓ is length-aligned to the source sequence by linear interpolation when shorter and truncation when longer. The pair-specific direction is

$$d_{A\to B}^\ell = h_B^\ell - h_A^\ell. \tag{1}$$

Let M_p be a binary mask selecting position type p, broadcast across channels, and let $z^\ell \sim \mathcal{N}(0, I)$ have the same shape as h_A^ℓ. For direction condition $k \in \{\text{align}, \text{orth}, \text{rand}\}$, define $u_{p,\text{align}}^\ell = M_p d_{A\to B}^\ell$, $u_{p,\text{rand}}^\ell = M_p z^\ell$, and $u_{p,\text{orth}}^\ell = M_p z^\ell - \text{proj}_{u_{p,\text{align}}^\ell}(M_p z^\ell)$, where $\text{proj}_a(b) = \langle b, a\rangle_F a/\|a\|_F^2$. With reference-audio aligned norm $r^\ell = \|u_{\text{ref,align}}^\ell\|_F$, every condition uses

$$\widetilde{h}_A^\ell = h_A^\ell + r^\ell \frac{u_{p,k}^\ell}{\|u_{p,k}^\ell\|_F}. \tag{2}$$

At L0, L3, and L7, the factorial tests reference-audio aligned, reference-audio orthogonal, reference-audio random, non-audio aligned, and final-state aligned conditions. The non-audio mask contains BOS, separator, and text positions; the final-state mask contains only the last position. For each dyad and prompt at a given layer, all five conditions share r^ℓ, matching total Frobenius norm across directions and positions.

Probe accuracy measures decodability; changes in source similarity and WER measure behavioral leverage; target gain and target dominance measure specificity. For generated speech y, the primary metric is matched-baseline target gain

$$\Delta\text{SIM}_B = \text{SIM}(y_{\text{int}}, B) - \text{SIM}(y_{\text{base}}, B), \tag{3}$$

where SIM is cosine similarity between L2-normalized speaker embeddings. After averaging seeds within prompts and the two prompts within each directed dyad, a dyad satisfies the **target-dominant criterion** when both $\Delta\text{SIM}_B > 0$ and $\text{SIM}(y_{\text{int}}, B) > \text{SIM}(y_{\text{int}}, A)$. The criterion combines target attraction with target-over-source preference in the same dyad. The primary factorial uses ECAPA-TDNN [21]. We score the same generated speech with CAM++ [22], perform 30-way retrieval against five-utterance speaker prototypes disjoint from all intervention references, and use Whisper WER [23] as an intelligibility measure.

The directed dyad is the inferential unit. We report dyad-bootstrap 95% confidence intervals and two-way source/target

Table 1. Speaker identity remains decodable across depth, while target control concentrates at early reference-audio positions. Reference-audio probes use the token mean. Target gain is $\Delta\mathrm{SIM}_B$, with dyad-bootstrap 95% CIs; dominant counts are out of 60 directed dyads. Probe and intervention samples are disjoint and share the same 30-speaker pool.

	Reference-audio			Final-state		
Layer	**Probe acc.**	**Target gain [95% CI]**	**Dominant**	**Probe acc.**	**Target gain [95% CI]**	**Dominant**
L0	0.947	**+0.326 [0.298, 0.353]**	**60/60**	0.637	+0.020 [0.006, 0.035]	1/60
L3	0.963	**+0.227 [0.204, 0.250]**	**52/60**	0.900	+0.038 [0.021, 0.055]	4/60
L7	0.970	0.000 [−0.002, 0.001]	0/60	0.933	+0.042 [0.027, 0.057]	7/60

cluster-robust 95% confidence intervals accounting for repeated speakers in both roles [24].

3. RESULTS

3.1. Distributed speaker information, localized target control

The depth contrast separates speaker information from target control: L0 aligned reference-audio edits achieve target dominance in 60/60 dyads, while L7 edits yield 0/60 despite a probe accuracy of 0.970. Mean target gain falls from $\Delta\mathrm{SIM}_B = +0.326$ at L0 to 0.000 at L7; L3 retains a gain of $+0.227$ with 52/60 target-dominant dyads (Table 1). The final layer thus retains decodable speaker information after reference-audio edits have lost attention-mediated readout reach.

Speaker information is distributed more broadly than this control region. Reference-audio probe accuracy stays within 0.947–0.973 across layers, while non-audio accuracy rises from 0.223 at L0 to 0.903 at L7 and final-state accuracy from 0.637 to 0.933. All reported probe conditions satisfy $p_{\mathrm{FWE}} = 0.001$ under the 24-condition permutation correction. A finer-grained six-layer sweep over 25 dyads resolves the early control region spanning L0–L3, followed by a sharp decline at L5 (Fig. 1(a)); target dominance reaches 0/25 at L5 and L7.

3.2. Direction and position determine the steering response

Target specificity depends on both update direction and token position. At L0, pair-aligned reference-audio steering exceeds equal-norm orthogonal and random updates by $+0.342$ and $+0.340$ in target similarity, respectively (Fig. 1(b)). It also exceeds non-audio and final-state aligned interventions by $+0.318$ and $+0.306$. All four 95% confidence intervals lie above zero under both dyad-bootstrap and two-way source/target cluster-robust inference. At L3, aligned-minus-control gains range from $+0.189$ to $+0.235$; at L7, aligned and direction-control contrasts are centered near zero.

At L0, aligned reference-audio intervention both displaces the source and attracts the target ($\Delta\mathrm{SIM}_A = -0.467$; post-intervention $\mathrm{SIM}_B = 0.467 > \mathrm{SIM}_A = 0.173$). Orthogonal perturbations reduce source similarity ($\Delta\mathrm{SIM}_A = -0.195$) while also reducing target similarity ($\Delta\mathrm{SIM}_B = -0.016$; 95% CI excludes zero). Orthogonal edits exert behavioral leverage by moving speech away from both speakers; pair alignment channels that leverage toward the intended target. Final-state interventions concentrate their leverage in intelligibility, raising matched-baseline $\Delta\mathrm{WER}$ by $+0.113$ [$+0.058$, $+0.179$] at L0, $+0.127$ [$+0.043$, $+0.251$] at L3, and $+0.096$ [$+0.053$, $+0.147$] at L7, with low target-dominant counts (Table 1). Reference-audio aligned intervention achieves target gains with mean $\Delta\mathrm{WER}$ near zero (Table 2). Across the same 25 dyads used for the six-layer sweep, an L0 norm sweep at $0.5, 1, 2, 5\times$ the native norm yields $+0.186, +0.387, +0.302, +0.342$. The aligned direction supports target attraction at all four scales, spanning a ten-fold norm range, with the largest gain at the native norm.

3.3. Target specificity extends to held-out speaker references

The target-speaker match extends across evaluators and reference utterances. Independent CAM++ evaluation reproduces the depth ordering: aligned target gain is $+0.237, +0.190, +0.001$ at L0/L3/L7, with 55/60, 43/60, and 0/60 target-dominant dyads (Table 2). We also compare generated speech with five-utterance prototypes disjoint from the intervention references. At L0, ECAPA and CAM++ both retrieve the intended target as top-1 among 30 candidates for 46/60 dyads; at L3 this is 44/60 and 41/60, respectively, while at L7 it falls to 0/60 for both. Successful retrieval against these independent, multi-utterance prototypes establishes that target-speaker matching generalizes beyond the intervention reference to the speaker level.

Table 2. Cross-encoder and held-out target evaluation. CAM++ gain denotes aligned $\Delta\mathrm{SIM}_B$; dom. counts target-dominant dyads. Retrieval is top-1 over 30 held-out speaker prototypes. CAM++ and retrieval use the first prompt; Whisper $\Delta\mathrm{WER}$ uses both prompts. All counts are out of 60 dyads; brackets report dyad-bootstrap 95% CIs for $\Delta\mathrm{WER}$.

		Held-out target retrieval		
Layer	**CAM++ gain / dom.**	**ECAPA**	**CAM++**	**Aligned ΔWER**
L0	+0.237 / 55/60	**46/60**	**46/60**	-0.003 $[-0.018, 0.013]$
L3	+0.190 / 43/60	44/60	41/60	-0.003 $[-0.015, 0.007]$
L7	+0.001 / 0/60	0/60	0/60	-0.001 $[-0.003, 0.001]$

3.4. Direct acoustic conditioning anchors speaker identity in dots.tts

To trace control across distinct conditioning paths, we apply the direction-control protocol to dots.tts [25], a 2B-parameter continuous autoregressive TTS model with a 28-layer backbone. Its acoustic head receives a global speaker x-vector [26] alongside backbone states. Speakers, dyads, prompts, seeds, and evaluators are matched to the primary factorial. We probe and intervene at backbone layers L0, L14, and L27 with aligned, orthogonal, and random directions at reference positions.

X-vector ablation pinpoints a strong source-identity anchor at the direct acoustic input: disabling it reduces mean source similarity from 0.849 to 0.006. The x-vector and backbone states reach the acoustic head through separate routes; backbone interventions edit one route while keeping the source-speaker x-vector fixed.

The backbone itself carries both decodable identity and behavioral leverage. Final-state probe accuracy increases with depth, from 0.260 at L0 to 0.437 at L14 and 0.530 at L27, measured by

Table 3. Backbone leverage and target specificity in dots.tts under equal-norm direction controls across 60 directed dyads, two prompts, and three seeds. Similarity changes use ECAPA-TDNN and matched baselines; WER is absolute Whisper word-error rate. Every condition has 0/60 target-dominant dyads. Aligned updates show the strongest source-displacement and WER response at each layer; mean target gain is non-positive throughout.

Layer	Direction	Source change $\Delta\mathrm{SIM}_A$	Target gain $\Delta\mathrm{SIM}_B$	WER (absolute)
L0	Aligned	$\mathbf{-0.197}$	-0.020	$\mathbf{0.475}$
L0	Orthogonal	-0.032	-0.014	0.107
L0	Random	-0.030	-0.007	0.089
L14	Aligned	$\mathbf{-0.187}$	-0.004	$\mathbf{0.421}$
L14	Orthogonal	-0.049	-0.009	0.095
L14	Random	-0.022	-0.010	0.075
L27	Aligned	$\mathbf{-0.046}$	-0.001	$\mathbf{0.231}$
L27	Orthogonal	-0.014	-0.002	0.086
L27	Random	-0.014	-0.001	0.068

five-fold cross-validation. Early aligned intervention is direction-specific (Table 3): $\Delta\mathrm{SIM}_A = -0.197$ at L0 versus -0.032 and -0.030 for equal-norm orthogonal and random directions, with a weaker aligned effect at L27 (-0.046). CAM++ reproduces the aligned source-displacement profile ($-0.154, -0.155, -0.023$ at L0/L14/L27). L0 aligned Whisper WER is 0.475, versus 0.107 for orthogonal and 0.089 for random updates. Mean $\Delta\mathrm{SIM}_B$ is non-positive in every condition, and all conditions yield 0/60 target-dominant dyads. This contrast separates behavioral leverage from target specificity within the same backbone. The acoustic head receives edited backbone states alongside a source-speaker x-vector that remains fixed. Backbone intervention and x-vector ablation therefore act on different conditioning routes: one reveals effects on source similarity and intelligibility; the other identifies the direct input that anchors speaker identity.

Table 4. Double-blind listening over the 60 dyads, with five judgments per dyad per comparison. (a) Target-speaker preference using held-out references; ties remain in the denominator. Brackets show dyad-bootstrap 95% CIs; $^{***}p < 0.001$ (mixed-effects logistic regression). (b) Naturalness relative to baseline on a -3-to-$+3$ CMOS scale, with 95% CIs.

(a) Target-speaker ABX preference

Comparison (first vs. second)	First-condition preference [95% CI]
L0 aligned vs. L0 orthogonal	**72.3%** [67.2, 77.1]% ***
L3 aligned vs. L3 orthogonal	**63.7%** [58.4, 68.9]% ***
L0 aligned vs. L7 aligned	**76.0%** [71.1, 80.6]% ***

(b) Comparative naturalness

Comparison	ΔCMOS **[95% CI]**
L0 aligned vs. baseline	-0.12 [-0.21, -0.03]
L3 aligned vs. baseline	-0.07 [-0.15, $+0.01$]

3.5. Listeners hear the target-speaker shift

Human judgments evaluate five comparisons over the 60 dyads. Each comparison receives five double-blind judgments per dyad (300 per comparison; 1,500 total), with candidate order randomized and counterbalanced. ABX analysis uses all 300 trials per comparison in a binary mixed-effects logistic regression (tested-condition preference = 1; alternative or tie = 0), with random intercepts for dyads and listeners. Comparative naturalness uses a seven-point CMOS scale from -3 to $+3$ (intervention minus matched baseline).

With an unseen target-speaker utterance as reference, listeners prefer L0 aligned over its equal-norm orthogonal control in 72.3% of trials, L3 aligned over its orthogonal control in 63.7%, and L0 aligned over L7 aligned in 76.0%; all three contrasts are significant at $p < 0.001$ (Table 4). Listeners hear the target specificity of early aligned steering while baseline naturalness is largely preserved: mean ΔCMOS is -0.12 (95% CI $[-0.21, -0.03]$) at L0 and -0.07 ($[-0.15, +0.01]$) at L3.

4. DISCUSSION AND CONCLUSION

Speaker decodability and target control separate sharply under intervention: states that decode speaker identity almost perfectly can lie beyond the mixing path that reaches synthesis, while early reference-audio states admit precise, target-directed edits. Equal-norm controls identify the directions and positions responsible for this response. Held-out retrieval and listening judgments connect the resulting target specificity to speaker identity and perception.

The mechanism combines **geometric specificity with readout reach**. Subsequent attention blocks carry early reference-audio edits to VoxCPM2's synthesis readout; post-block L7 reference-audio edits lie beyond this mixing path. The dots.tts comparison reveals a different routing arrangement: a direct x-vector input anchors source identity alongside an editable backbone with behavioral leverage. These findings give a concrete design principle for training-free speaker control: expose editable states with downstream readout reach, and coordinate their target-directed updates with any direct acoustic speaker-conditioning inputs. Control-site selection must account for both intervention geometry and the paths that carry speaker identity into synthesis.

5. REFERENCES

[1] Y. Zhou et al., "VoxCPM2 technical report," arXiv preprint arXiv:2606.06928, 2026.

[2] S. Chen et al., "Neural codec language models are zero-shot text-to-speech synthesizers," *IEEE Trans. Audio, Speech, Lang. Process.*, vol. 33, pp. 705–718, 2025, doi: 10.1109/TASLPRO.2025.3530270.

[3] M. Le et al., "Voicebox: Text-guided multilingual universal speech generation at scale," in *Proc. NeurIPS*, 2023.

[4] E. Casanova, J. Weber, C. D. Shulby, A. Candido Junior, E. Gölge, and M. A. Ponti, "YourTTS: Towards zero-shot multi-speaker TTS and zero-shot voice conversion for everyone," in *Proc. ICML*, 2022, pp. 2709–2720.

[5] Z. Du et al., "CosyVoice: A scalable multilingual zero-shot text-to-speech synthesizer based on supervised semantic tokens," arXiv preprint arXiv:2407.05407, 2024.

[6] G. Alain and Y. Bengio, "Understanding intermediate layers using linear classifier probes," arXiv preprint arXiv:1610.01644, 2016.

[7] B. van Niekerk, L. Nortje, M. Baas, and H. Kamper, "Analyzing speaker information in self-supervised models to improve zero-resource speech processing," in *Proc. Interspeech*, 2021, pp. 1554–1558.

[8] J. Hewitt and P. Liang, "Designing and interpreting probes with control tasks," in *Proc. EMNLP-IJCNLP*, 2019, pp. 2733–2743.

[9] A. Ravichander, Y. Belinkov, and E. Hovy, "Probing the probing paradigm: Does probing accuracy entail task relevance?," in *Proc. EACL*, 2021, pp. 3363–3377.

[10] Y. Elazar, S. Ravfogel, A. Jacovi, and Y. Goldberg, "Amnesic probing: Behavioral explanation with amnesic counterfactuals," *Trans. Assoc. Comput. Linguistics*, vol. 9, pp. 160–175, 2021.

[11] K. Meng, D. Bau, A. Andonian, and Y. Belinkov, "Locating and editing factual associations in GPT," in *Proc. NeurIPS*, 2022.

[12] J. Vig, S. Gehrmann, Y. Belinkov, S. Qian, D. Nevo, Y. Singer, and S. Shieber, "Investigating gender bias in language models using causal mediation analysis," in *Proc. NeurIPS*, 2020.

[13] K. Li, A. K. Hopkins, D. Bau, F. Viégas, H. Pfister, and M. Wattenberg, "Emergent world representations: Exploring a sequence model trained on a synthetic task," in *Proc. ICLR*, 2023.

[14] A. M. Turner, L. Thiergart, G. Leech, D. Udell, J. J. Vazquez, U. Mini, and M. MacDiarmid, "Steering language models with activation engineering," arXiv preprint arXiv:2308.10248, 2023.

[15] N. Rimsky, N. Gabrieli, J. Schulz, M. Tong, E. Hubinger, and A. Turner, "Steering Llama 2 via contrastive activation addition," in *Proc. ACL*, 2024, pp. 15504–15522.

[16] T. Xie, S. Yang, C. Li, D. Yu, and L. Liu, "EmoSteer-TTS: Fine-grained and training-free emotion-controllable text-to-speech via activation steering," arXiv preprint arXiv:2508.03543, 2025.

[17] N. Koriagin, G. Aparin, N. Balagansky, and D. Gavrilov, "Interpreting and steering a text-to-speech language model with sparse autoencoders," arXiv preprint arXiv:2606.10029, 2026.

[18] S. Wang, J. Bailey, and T. Dang, "A geometric perspective on composable emotion steering in text-to-speech models," arXiv preprint arXiv:2607.00946, 2026.

[19] M. Yang and J. H. L. Hansen, "Activation steering for accent-neutralized zero-shot text-to-speech," arXiv preprint arXiv:2603.05977, 2026.

[20] Y. Koizumi et al., "LibriTTS-R: A restored multi-speaker text-to-speech corpus," in *Proc. Interspeech*, 2023, pp. 5496–5500.

[21] B. Desplanques, J. Thienpondt, and K. Demuynck, "ECAPA-TDNN: Emphasized channel attention, propagation and aggregation in TDNN based speaker verification," in *Proc. Interspeech*, 2020, pp. 3830–3834.

[22] H. Wang, S. Zheng, Y. Chen, L. Cheng, and Q. Chen, "CAM++: A fast and efficient network for speaker verification using context-aware masking," in *Proc. Interspeech*, 2023, pp. 5301–5305.

[23] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, "Robust speech recognition via large-scale weak supervision," in *Proc. ICML*, 2023, pp. 28492–28518.

[24] A. C. Cameron, J. B. Gelbach, and D. L. Miller, "Robust inference with multiway clustering," *J. Bus. Econ. Stat.*, vol. 29, no. 2, pp. 238–249, 2011, doi: 10.1198/jbes.2010.07136.

[25] S. Lian et al., "dots.tts technical report," arXiv preprint arXiv:2606.07080, 2026.

[26] D. Snyder, D. Garcia-Romero, G. Sell, D. Povey, and S. Khudanpur, "X-vectors: Robust DNN embeddings for speaker recognition," in *Proc. IEEE ICASSP*, 2018, pp. 5329–5333, doi: 10.1109/ICASSP.2018.8461375.