DISRL: Dual-Scale Incremental Semi-Supervised Reinforcement Learning for Dynamic Community Tracking

Tianming Zhang
College of Computer Science and Software Engineering
Zhejiang University of Technology
Hangzhou, China
tmzhang@zjut.edu.cn

Jinxu Wang
College of Computer Science and Software Engineering
Zhejiang University of Technology
Hangzhou, China
jinxuwang@zjut.edu.cn

Yixuan Zhang
College of Computer Science and Software Engineering
Zhejiang University of Technology
Hangzhou, China
zhangyixuan@zjut.edu.cn

Zhengyi Yang
School of Computer Science
The University of Sydney
Sydney, Australia
zhengyi.yang@sydney.edu.au

Chenyu Hou
College of Computer Science and Software Engineering
Zhejiang University of Technology
Hangzhou, China
houcy@zjut.edu.cn

Bin Cao
College of Computer Science and Software Engineering
Zhejiang University of Technology
Hangzhou, China
bincao@zjut.edu.cn

Jing Fan*
College of Computer Science and Software Engineering
Zhejiang University of Technology
Hangzhou, China
fanjing@zjut.edu.cn

arXiv:2610.29315v1 [cs.LG] 21 Oct 2026

Abstract

Dynamic local community detection tracks the evolving community of designated seed nodes across graph snapshots. However, it faces two critical challenges: *temporal supervision uncertainty*, where sparse and outdated labels fail to provide seed-specific guidance under topology changes, and *unbounded error cascades*, where sequential policy transfer accumulates exploration errors over long horizons. To address these challenges, we propose DISRL, a Dual-scale Incremental Semi-supervised Reinforcement Learning framework. DISRL introduces a *supervision construction module* that dynamically generates seed-specific guidance from sparse known communities, and a *dual-scale community exploration module* that stabilizes policy evolution by decoupling short-term adaptation within temporal windows from long-term parameter consolidation across windows. Furthermore, we construct a dynamic graph benchmark with diverse community evolution scenarios. Extensive experiments demonstrate that DISRL consistently outperforms state-of-the-art baselines under various dynamic settings.

CCS Concepts

• **Information systems** → **Social networks**; • **Theory of computation** → *Graph algorithms analysis.*

Keywords

Dynamic Graphs, Local Community Detection, Semi-supervised Learning, Reinforcement Learning

1 Introduction

Real-world networks, including social platforms [22], citation graphs [13], and fraud networks [32], exhibit prominent community structures. In these networks, partially observed communities [38] often provide critical structural guidance for discovering seed-centered communities around entities of interest. Semi-supervised local community detection aims to leverage such sparse supervisory information to identify the specific community surrounding a seed node without exhaustively partitioning the entire graph [37]. However, as interactions evolve across Web platforms and other real-world networks, communities surrounding queried users or entities may expand, shrink, merge, or split across temporal snapshots [28], requiring adaptive approaches that can track local structural changes while preserving temporal coherence.

Existing methods remain inadequate for dynamic local community tracking. Static approaches [3, 5, 12, 14, 16, 24–26, 33–35, 38] ignore temporal dependencies by processing snapshots independently, while traditional dynamic methods [2, 7, 19, 22, 36, 41] often rely on handcrafted transition rules. Although recent learning-based models [4, 8, 11, 15, 17, 20, 29, 31, 39] capture complex temporal patterns, they mainly process temporal streams passively via embedding aggregation rather than active exploration. Under non-stationary evolution, such continual transfer may accumulate tracking errors over long horizons.

To establish an active tracking paradigm, this paper investigates semi-supervised dynamic local community detection, aiming to accurately track the evolving community surrounding each seed node from graph snapshots with sparse supervision. This problem introduces two fundamental challenges.

*Corresponding author.

First challenge: Temporal supervision uncertainty. Rapid structural mutations (e.g., community splits or mergers) render historical predictions obsolete or misaligned with the current topology. Meanwhile, sparse supervision restricted to general known communities creates a structural mismatch with fine-grained, seed-specific local target boundaries.

Second challenge: Unbounded error cascade. In non-stationary temporal graphs, the environmental Markov property is highly volatile. While continuously passing policy parameters across sequential snapshots preserves short-term temporal dependencies, exploratory errors from noisy snapshots can compound over long horizons, causing policy drift and performance degradation.

To address these challenges, we propose **DISRL**, a **D**ual-scale **I**ncremental **S**emi-Supervised **R**einforcement **L**earning framework that formulates dynamic local community tracking as structured sequential decision-making. DISRL combats these challenges via two synergistic modules. First, a *semi-supervision construction module* resolves supervision uncertainty. By dynamically synthesizing seed-specific pseudo-reference communities from sparse annotations, it incrementally aligns historical predictions with current topologies, retrieving structurally relevant local guidance to counteract label sparsity. Second, a *dual-scale community exploration module* mitigates unbounded error cascades through a decoupled, asynchronous policy adaptation pipeline. Within each temporal window (intra-window level), a fast policy propagates and adapts across consecutive snapshots to capture fine-grained evolutionary trajectories, initialized from a stable window-level master parameter. Concurrently, the master parameter remains stationary throughout the window to serve as a reliable geometric anchor and is centrally updated at the inter-window level only at the window's end by aggregating all intra-window transition gradients.

Furthermore, to bridge the evaluation gap caused by synthetic or oversimplified datasets, we construct a dynamic graph benchmark based on empirical network topologies and authentic community structures. It generates temporally ordered snapshots with diverse community evolution patterns, including expansion, shrinkage, splitting, merging, and topological perturbations, enabling rigorous evaluation under non-stationary settings. The main contributions are summarized as follows:

- We pioneer the formulation of dynamic local community tracking as an incremental Reinforcement Learning (RL) sequence, shifting the traditional passive temporal embedding paradigms into a goal-directed, multi-step active exploration process around seed nodes.
- We develop DISRL, featuring a semi-supervised supervision construction module that synthesizes seed-specific adaptive guidance, and an asynchronous dual-scale exploration pipeline that decouples intra-window trajectory optimization from inter-window gradient aggregation, effectively stabilizing policy transfer and mitigating the risk of long-horizon error accumulation.
- We construct a non-stationary dynamic graph benchmark that explicitly simulates fine-grained community evolutions. Compared with the strongest baseline in each setting, DISRL achieves F1-score gains of 2.89 and up to 6.42 percentage points on the real-world dataset and the simulated-evolution datasets, respectively.

2 Related Work

Static Community Detection: Static community detection methods have been extensively studied based on optimization [5], matrix factorization [16], generative models [14], and graph representation learning [3, 9, 12, 30]. To exploit limited prior knowledge, Semi-supervised methods are proposed. SEAL [38] adopts a generator-discriminator framework, SLSS [24] and SLFM [35] exploit known communities to guide seed-centered local community detection. Recently, RL-based methods, including CLARE [34] and SLRL [26], formulate community expansion as sequential decision-making; prompt-based methods such as ProCom [33] and PPSL [25] improve label efficiency. However, these methods are designed for static topologies and cannot exploit temporal continuity, causing learned structural patterns to misalign with evolving communities.
Dynamic Community Detection: Traditional approaches model temporal evolution through evolutionary clustering [7], model-based approaches [19, 36], incremental updates [41], seed-centered tracking such as DevDynaP [2], and temporal community search methods [18, 40]. Recent learning-based methods leverage temporal graph neural networks [11, 17, 31, 39] and topology-regularized clustering [15, 20, 29] to capture complex evolutionary patterns. Nevertheless, these methods rarely address seed-centered local community detection under sparse supervision. Without explicit seed-specific supervision, sparse known communities may provide weak guidance for evolving local targets; meanwhile, continual parameter transfer across snapshots can propagate transient errors over long horizons. These limitations call for an adaptive and stable framework for dynamic local community detection.

3 Methodology

3.1 Problem Formulation and Overview

Let $\mathcal{G} = \{G^t = (\mathcal{V}^t, \mathcal{E}^t)\}_{t=1}^{T}$ denote a dynamic graph with T temporal snapshots, where $\mathcal{V}^t$ and $\mathcal{E}^t$ are the node and edge sets at timestamp t, respectively. At each snapshot, only a sparse set of known communities is available as incomplete supervision, denoted by $\mathcal{C}^t_{\text{known}} = \{C^t_i\}_{i=1}^{m}$, where $C^t_i \subseteq \mathcal{V}^t$ $(1 \le i \le m)$. Given a set of designated seed nodes $\mathcal{S}$, our goal is to track the evolving local community of each seed node over temporal snapshots. Specifically, for each seed $s \in \mathcal{S}$ at snapshot G^t, we aim to identify a community $C^t_{\text{pred}}(s) \subseteq \mathcal{V}^t$ satisfying $s \in C^t_{\text{pred}}(s)$. The objective is to sequentially infer the dynamic local communities $\mathcal{C}_{\text{pred}} = \{C^t_{\text{pred}}(s) \mid s \in \mathcal{S}, t = 1, \ldots, T\}$ by leveraging temporal graph evolution and sparse supervision.

Existing reinforcement learning-based community detection methods mainly target static graphs or short-term refinement, and cannot address *temporal supervision uncertainty* and *unbounded error cascade* in dynamic local tracking. To address these challenges, we propose DISRL, which is an incremental reinforcement learning framework. As shown in Fig. 1, DISRL consists of two synergistic modules. The *Supervision Construction Module* resolves temporal supervision uncertainty by dynamically refining sparse known communities into seed-specific adaptive guidance C^t_{rel}. The *Dual-scale*

Community Exploration Module performs reinforcement-learning-based boundary expansion under this guidance. To mitigate long-horizon error cascades, DISRL decouples policy optimization into intra-window fast adaptation across consecutive snapshots and inter-window master anchoring across temporal windows.

3.2 Supervision Construction Module

In dynamic graphs, coarse global annotations C^t_{known} capture macro-level partitions, often exhibiting spatial and structural mismatch with the micro-vicinity of a target seed s. To bridge this supervision gap without losing temporal consistency, the Supervision Construction module implements a two-stage process: reference community synthesis and relevant community retrieval.

3.2.1 Reference Community Synthesis. The objective is to construct a topology-congruent local structural anchor C^t_{ref} for the current snapshot G^t. Rather than serving as the final prediction, this reference acts as an inductive bias to filter out irrelevant global annotations. For a specified seed s, $C^t_{\text{ref}}(s)$ is defined as:

$$C^t_{\text{ref}}(s) = \text{CC}\left(G^t\left[C^{t-1}_{\text{pred}}(s) \cap \mathcal{V}^t\right]\right) \tag{1}$$

where $G^t[V']$ is the subgraph of G^t induced by V', and $\text{CC}(\cdot)$ extracts the connected component containing s.

Cold-start Bootstrapping ($t = 1$): At the temporal onset, historical predictions are inaccessible ($t = 1$). DISRL switches to the cold-start mode, executing the Community Explorer in the inference phase (as defined in Sec. 3.3.1) over G^1 initialized by a generic policy pre-trained on C^1_{known}. The generated initial seed-centric boundary is captured as C^1_{ref}.

Warm-start Alignment ($t > 1$): For subsequent snapshots, directly inheriting $C^{t-1}_{\text{pred}}(s)$ introduces geometric misalignment due to localized edge mutations in G^t. DISRL projects $C^{t-1}_{\text{pred}}(s)$ onto the current topology. As formalized in Eq. (1), valid active nodes are retained, while $\text{CC}(\cdot)$ prunes disconnected fragments caused by community shrinkage or splits, ensuring that C^t_{ref} remains seed-connected.

3.2.2 Relevant Community Retrieval. To eliminate macro-level topological noise, DISRL retrieves a compact set of known communities that are structurally similar to the synthesized reference community $C^t_{\text{ref}}(s)$. Let $\mathcal{G}^t_{\text{ref}}(s) = G^t[C^t_{\text{ref}}(s)]$ denote the reference subgraph. For each known community $C^t_i \in C^t_{\text{known}}$, let $\mathcal{G}^t_i = G^t[C^t_i]$ denote its induced subgraph. To capture fine-grained path-level structural patterns without resorting to NP-hard exact graph isomorphism, DISRL employs a normalized shortest-path graph kernel [6]. The structural similarity between the reference and the i-th known community is computed as:

$$\text{sim}(\mathcal{G}^t_{\text{ref}}(s), \mathcal{G}^t_i) = \frac{K_{\text{sp}}(\mathcal{G}^t_{\text{ref}}(s), \mathcal{G}^t_i)}{\sqrt{K_{\text{sp}}(\mathcal{G}^t_{\text{ref}}(s), \mathcal{G}^t_{\text{ref}}(s)) \cdot K_{\text{sp}}(\mathcal{G}^t_i, \mathcal{G}^t_i)}}, \tag{2}$$

where $K_{\text{sp}}(\cdot,\cdot)$ counts matching shortest-path lengths between graph node pairs, preserving geodesic path distributions that are sensitive to community boundary variations.

Subsequently, DISRL ranks all candidates in C^t_{known} in descending order of the similarity scores in Eq. (2) and selects the top-K communities to form the seed-specific relevant supervision set C^t_{rel}:

$$C^t_{\text{rel}}(s) = \text{Top-K}\left(\left\{C^t_i \in C^t_{\text{known}} \;\middle|\; \text{sim}(\mathcal{G}^t_{\text{ref}}(s), \mathcal{G}^t_i)\right\}\right). \tag{3}$$

3.3 Dual-scale Community Exploration Module

After obtaining $C^t_{\text{rel}}(s)$, DISRL formulates seed-centered community tracking on G^t as an MDP-based sequential expansion process. As shown in Fig. 1, the exploration module uses a dual-scale update pipeline to coordinate snapshot-level micro-exploration with temporal horizon-level macro-anchoring. At the intra-window scale, a fast policy $\Phi_{m,t}$ executes incremental transfer and rapid adaptation across adjacent snapshots to capture short-term topology changes. At the inter-window scale, a master parameter Θ_m serves as a long-term stable anchor, periodically consolidating transition meta-gradients at window boundaries to limit the consecutive propagation of historical exploration errors and mitigate the risk of unbounded error cascades across long horizons.

3.3.1 Intra-window Scale Exploration and Fast Adaptation. To efficiently approximate target boundaries without global graph traversal, DISRL employs a Community Explorer based on an active search policy. During training, the policy is fine-tuned within C^t_{rel}: for each $C^* \in C^t_{\text{rel}}$, the agent performs rollouts from a randomly sampled virtual seed, with rewards derived from C^*. During inference, the adapted policy is deployed on the target seed s to greedily explore the snapshot graph G^t and obtain the predicted community $C^t_{\text{pred}}(s)$. The exploration process is formulated as a snapshot-level MDP.

State Space. At exploration step k in snapshot G^t, the environment state $\mathbf{S}^t_k$ is composed of (i) the currently explored community C^t_k; (ii) its immediate candidate action boundary $\mathcal{N}(C^t_k)$, which is defined as:

$$\mathcal{N}(C^t_k) = \left\{u \in \mathcal{V}^t \setminus C^t_k \;\middle|\; \exists v \in C^t_k,\ (u,v) \in \mathcal{E}^t\right\}, \tag{4}$$

and (iii) the node representations computed on a subgraph centered around C^t_k.

To capture the local structural context efficiently, DISRL extracts an L-hop localized subgraph induced by the expanded community region:

$$\mathcal{V}^t_{\text{sub},k} = C^t_k \cup \mathcal{N}^L(C^t_k), \quad \mathcal{G}^t_{\text{sub},k} = G^t\left[\mathcal{V}^t_{\text{sub},k}\right], \tag{5}$$

where $\mathcal{N}^L(C^t_k)$ is the L-hop neighborhood set enclosing C^t_k.

Within the subgraph $\mathcal{G}^t_{\text{sub},k}$, we initialize a two-channel indicator matrix $\mathbf{X}^t_k \in \mathbb{R}^{|\mathcal{V}^t_{\text{sub},k}|\times 2}$ to explicitly parameterize the geometric orientation of the root seed s_{init} and the membership of the currently explored community C^t_k. For each node $u \in \mathcal{V}^t_{\text{sub},k}$, its corresponding row is defined as:

$$\mathbf{X}^t_k(u,:) = \left[\mathbb{I}(u = s_{\text{init}}), \mathbb{I}(u \in C^t_k)\right], \tag{6}$$

where $\mathbb{I}(\cdot)$ denotes the indicator function.

To effectively propagate the localized structural context and resolve the structural anonymity of partitions, an incremental Graph Pointer Network (iGPN) [38] is then deployed over $\mathcal{G}^t_{\text{sub},k}$. Setting the base hidden state $\mathbf{H}^{(0)} = \mathbf{X}^t_k$, the latent node embeddings

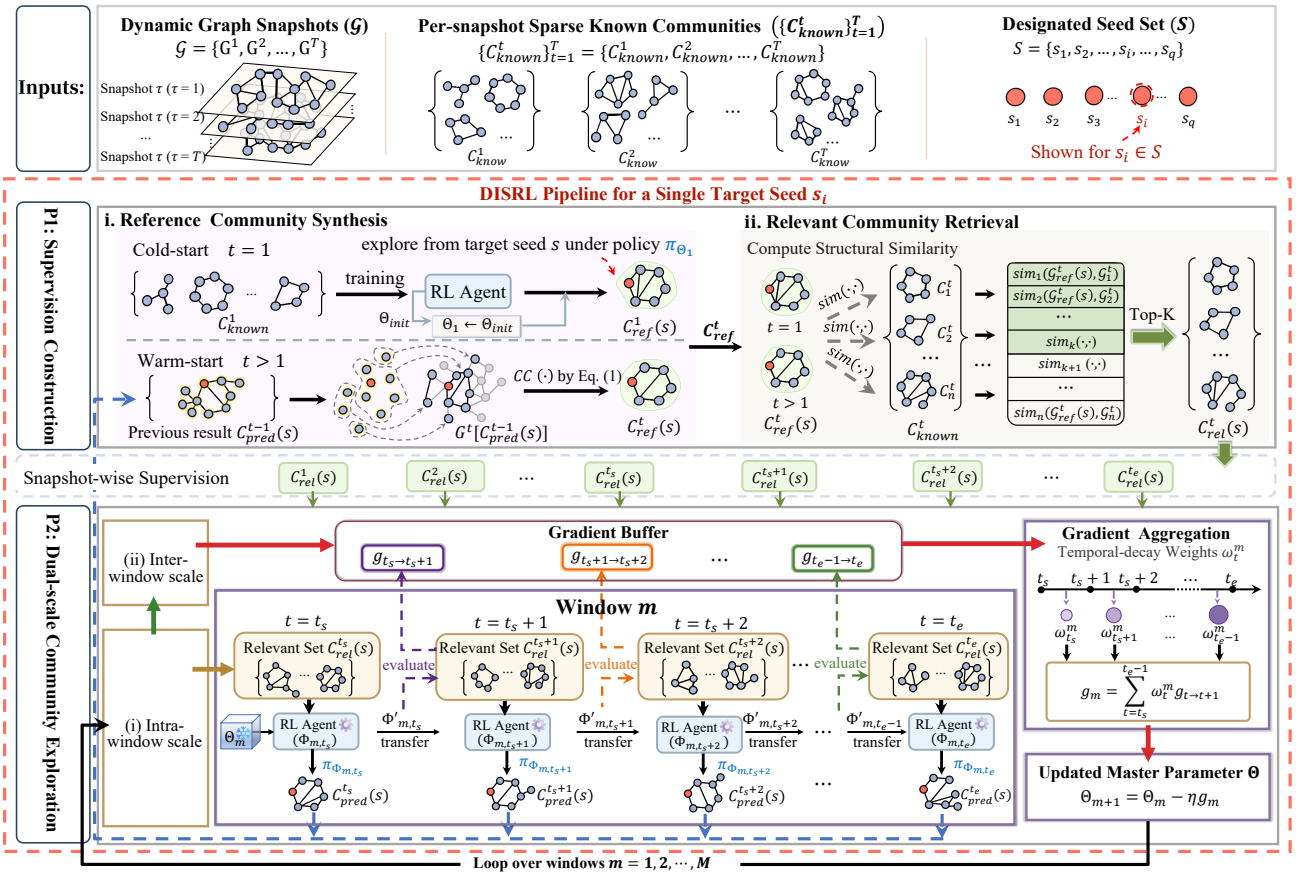

Figure 1: Overall framework of DISRL, consisting of supervision construction and dual-scale community exploration.

$\mathbf{H}^{(l)} \in \mathbb{R}^{|\mathcal{V}^t_{\text{sub},k}| \times 2}$ at the l-th propagation layer are iteratively computed as:

$$\mathbf{H}^{(l)} = \alpha \widetilde{\mathbf{A}}^t_{\text{sub},k} \mathbf{H}^{(l-1)} + (1-\alpha)\mathbf{X}^t_k, \quad l = 1, 2, \ldots, L \tag{7}$$

where $\widetilde{\mathbf{A}}^t_{\text{sub},k} \in \mathbb{R}^{|\mathcal{V}^t_{\text{sub},k}| \times |\mathcal{V}^t_{\text{sub},k}|}$ is the symmetrically normalized adjacency matrix of $\mathcal{G}^t_{\text{sub},k}$, and $\alpha \in (0, 1)$ is the propagation coefficient. The final output $\mathbf{H}^{(L)}$ yields continuous node representations that simultaneously encapsulate seed-aware and community-aware features.

To strictly penalize boundary leakage and shape distortions during sequential expansion, we further introduce a community-level structural deviation feature $\delta^t_{u,k} \in \mathbb{R}$ for each node $u \in \mathcal{V}^t_{\text{sub},k}$.

$$\delta^t_{u,k} = \left| \mathbf{H}^{(L)}(u, 2) - \frac{1}{|C^t_k|} \sum_{v \in C^t_k} \mathbf{H}^{(L)}(v, 2) \right|. \tag{8}$$

Subsequently, the propagated structural feature channels and localized deviation scalar are concatenated into a compact low-dimensional node feature vector $\mathbf{z}^t_{u,k} \in \mathbb{R}^3$:

$$\mathbf{z}^t_{u,k} = \left[\mathbf{H}^{(L)}(u, :) \parallel \delta^t_{u,k} \right]. \tag{9}$$

Finally, the complete environment state $\mathbf{S}^t_k$ is formalized as:

$$\mathbf{S}^t_k = \left(C^t_k, \mathcal{N}(C^t_k), \left\{ \mathbf{z}^t_{v,k} \,\middle|\, v \in C^t_k \cup \mathcal{N}(C^t_k) \right\} \right). \tag{10}$$

Action Space. At exploration step k, the action space $\mathcal{A}^t_k$ consists of boundary candidates and a termination token (EOS):

$$\mathcal{A}^t_k = \mathcal{N}(C^t_k) \cup \{\text{EOS}\}. \tag{11}$$

Selecting $a = u \in \mathcal{N}(C^t_k)$ expands the current community as $C^t_{k+1} = C^t_k \cup \{u\}$, while EOS terminates exploration and returns the final community. To bound exploration, we additionally terminate the rollout when the explored community reaches the maximum size among the available known communities.

The policy $\pi_{\Phi_{m,t}}(\cdot|\mathbf{S}^t_k)$ is parameterized by an MLP over node representations. During training, the fast policy is updated via policy gradient:

$$\Phi'_{m,t} = \Phi_{m,t} + \eta \sum_{k=0}^{T_e} \nabla_{\Phi_{m,t}} \log \pi_{\Phi_{m,t}}\left(a^t_k \mid \mathbf{S}^t_k\right) R^t_k, \tag{12}$$

where η is the learning rate, T_e is the terminal exploration step, and $R^t_k = \sum_{j=k}^{T_e} \gamma^{j-k} r^t_j$ is the discounted return with discount factor γ and immediate reward $r^t_j = \text{F1}(C^t_{j+1}, C^*) - \text{F1}(C^t_j, C^*)$. The updated policy is then transferred to initialize the explorer at the next

snapshot:

$$\Phi_{m,t+1} \leftarrow \Phi'_{m,t}, \quad \forall\, t \in \{t_s, \ldots, t_e - 1\}, \tag{13}$$

t_s and t_e are the start and end timestamps of the window, respectively. This incremental policy inheritance enables fast adaptation to local temporal changes while preserving evolutionary continuity within each temporal window.

3.3.2 Inter-window Scale Meta-Anchoring and Temporal-Decay Gradient Aggregation. To mitigate unbounded error cascades during long-term temporal adaptation, DISRL introduces an inter-window consolidation mechanism driven by first-order transition meta-gradients. The master parameter Θ_m provides a stable initialization anchor for fast policies across temporal windows. After intra-window adaptation, the transferred policy is evaluated upon the arrival of snapshot $t+1$, using the available training supervision C_{rel}^{t+1} to compute a transition meta-loss:

$$\mathcal{L}^{\text{meta}}_{t\to t+1}(\Phi_{m,t+1}) = -\sum_{i=1}^{B}\sum_{k=0}^{T_i-1} R_{i,k}^{t+1} \log \pi_{\Phi_{m,t+1}}\left(a_{i,k}^{t+1} \mid \mathrm{S}_{i,k}^{t+1}\right), \tag{14}$$

where B denotes the meta-batch size, and T_i denotes the exploration horizon of the i-th trajectory rollout.

Deriving the exact meta-gradient with respect to the master parameter Θ_m requires backpropagating through the complete unrolled intra-window adaptation path:

$$\nabla_{\Theta_m}\mathcal{L}^{\text{meta}}_{t\to t+1} = \left(\frac{\partial \Phi_{m,t+1}}{\partial \Theta_m}\right)^{\top} \nabla_{\Phi_{m,t+1}}\mathcal{L}^{\text{meta}}_{t\to t+1}. \tag{15}$$

Instead of computing the exact meta-gradient through the entire adaptation trajectory, DISRL adopts first-order meta-learning approximation [27] and directly estimates the transition correction:

$$\mathbf{g}_{t\to t+1} = \nabla_{\Phi_{m,t+1}}\mathcal{L}^{\text{meta}}_{t\to t+1}(\Phi_{m,t+1}). \tag{16}$$

These transition gradients capture the discrepancy between current policy adaptation and future topology evolution and are buffered without immediately updating the fast policy.

At the end of each temporal window, DISRL aggregates transition gradients with temporal decay weights ω_t^m controlled by a decay factor $\rho \in (0, 1]$:

$$\omega_t^m = \frac{\rho^{t_e-1-t}}{\sum_{q=t_s}^{t_e-1}\rho^{t_e-1-q}}, \quad t = t_s, \ldots, t_e - 1, \tag{17}$$

$$\mathbf{g}_m = \sum_{t=t_s}^{t_e-1} \omega_t^m \mathbf{g}_{t\to t+1}, \quad \Theta_{m+1} = \Theta_m - \eta \mathbf{g}_m, \tag{18}$$

where recent transitions receive larger weights. The updated master parameter initializes the next window, providing stable policy evolution and reducing long-term tracking drift.

3.4 Theoretical Analysis

We theoretically analyze why continuous fast-parameter transfer may suffer from long-horizon error accumulation and how the proposed dual-scale mechanism mitigates this issue. Detailed proofs are provided in Appendix A. Let $\Phi_{m,t}$ and $\Phi^{\star}_{m,t}$ denote the actual and ideal fast parameters at snapshot G^t, respectively. We define the tracking error as $d_{m,t} = \|\Phi_{m,t} - \Phi^{\star}_{m,t}\|_2$. We consider three standard assumptions: (1) bounded local adaptation perturbation $\|\boldsymbol{\xi}_{m,t}\|_2 \le \xi$, (2) bounded temporal drift $\|\Phi^{\star}_{m,t+1} - \Phi^{\star}_{m,t}\|_2 \le \Delta$, and (3) a Lipschitz adaptation operator with coefficient κ [10], which is expressed as

$$\|\mathcal{U}_t(\Phi_a) - \mathcal{U}_t(\Phi_b)\|_2 \le \kappa \|\Phi_a - \Phi_b\|_2. \tag{19}$$

The perturbed fast adaptation is

$$\Phi_{m,t+1} = \Phi'_{m,t} = \mathcal{U}_t(\Phi_{m,t}) + \boldsymbol{\xi}_{m,t}, \tag{20}$$

where the ideal parameter satisfies the fixed-point condition $\mathcal{U}_t(\Phi^{\star}_{m,t}) = \Phi^{\star}_{m,t}$. We establish the following proposition.

Proposition 1 (Single-step Error Propagation). For two consecutive snapshots within the same temporal window, the fast-parameter tracking error satisfies:

$$d_{m,t+1} \le \kappa d_{m,t} + \xi + \Delta. \tag{21}$$

The proposition indicates that each snapshot transition may introduce adaptation perturbations and temporal drift, potentially contributing to error accumulation during continuous policy transfer. By recursively applying Proposition 1 within a temporal window $\mathcal{W}_m$, the tracking error satisfies the following upper bound:

$$d_{m,t} \le \kappa^{t-t_s} d_{m,t_s} + (\xi + \Delta) \sum_{r=0}^{t-t_s-1} \kappa^r. \tag{22}$$

Therefore, without periodic policy anchoring, sequential fast-parameter transfer over long temporal horizons may allow historical errors to accumulate or amplify, increasing the risk of error cascades, particularly when $\kappa \ge 1$.

Corollary 1 (Window-confined Error Propagation). In DISRL, the fast parameter is reinitialized from the master parameter Θ_m at each temporal window boundary. Therefore, direct consecutive fast-parameter transfer is restricted to at most $w-1$ transitions within a window of size $w = |\mathcal{W}_m|$, yielding:

$$d_{m,t_e} \le d_{m,t_s} + (w-1)(\xi + \Delta), \quad \text{when } \kappa = 1. \tag{23}$$

Hence, the dual-scale design restricts direct error propagation within each temporal window, mitigating the risk of long-horizon error accumulation. This effect is further supported by the ablation study in Section 4.3 and the long-horizon evaluation in Section 4.4.

4 Experiments

4.1 Experimental Setup

Dynamic Benchmark Construction. Existing benchmarks provide limited support for dynamic local community tracking: real-world datasets are mostly static, while synthetic evolutionary benchmarks (e.g., CoDÆN [28]) mainly focus on global community partitioning and fail to simulate sparse, seed-centered local tracking. To address this limitation, we construct a non-stationary benchmark by synthesizing controlled, chronologically ordered evolutionary mutations onto four empirical networks while preserving their original topological complexity and community fidelity. Specifically, we generate temporal graph streams through four representative evolutionary operations, including community expansion, shrinkage, merging, and splitting, together with background topological perturbations (random node/edge insertions and deletions). The

Table 1: Dataset statistics. T denotes the number of snapshots; $\overline{|\mathcal{V}|}$, $\overline{|\mathcal{E}|}$, $\overline{|\mathcal{C}|}$, and $\overline{Q}$ denote the average numbers of nodes, edges, and communities, and the average modularity [23], respectively. All averages are taken over snapshots.

Dataset	T	$\overline{\lvert\mathcal{V}\rvert}$	$\overline{\lvert\mathcal{E}\rvert}$	$\overline{\lvert\mathcal{C}\rvert}$	$\overline{Q}$
Real-world dynamic dataset					
Football	24	137	653.9	19.8	0.5867
Expansion & shrinkage setting					
Amazon	25	18256.2	56573.5	1291.0	0.9301
DBLP	25	132172.3	540405.8	4520	0.9105
Twitter	13	99566.3	1334284.8	2651	0.6313
YouTube	13	227400.2	1426123.3	2821.9	0.5422
Merging & splitting setting					
Amazon	9	13178	38359.2	809.3	0.9837
DBLP	9	114095	485261.7	2829.7	0.9164

ground-truth communities are synchronously updated with graph evolution to provide reliable supervision.

Datasets and Evaluation Scenarios. We evaluate DISRL on five datasets covering simulated evolution and real-world dynamics. For real-world dynamics, we use Football, derived from the Century of College Football dataset [1], comprising 24 yearly snapshots from 2000–2024, with real game interactions and conference-based community labels. For simulated evolution, we use four SNAP datasets[1] with authentic community annotations: Amazon, DBLP, YouTube, and Twitter. We consider two scenarios: (1) expansion&shrinkage on all four datasets, which evaluates gradual community growth and contraction; and (2) merging&splitting on Amazon and DBLP, which tests structural reorganization among communities. Dataset statistics are summarized in Table 1. Detailed generation procedures and evaluation protocols are provided in Appendix B.

Baselines. We compare DISRL with representative baselines from two categories:

(i) Static Local Baselines: Applied independently to isolated snapshots using snapshot-specific topologies and annotations. **SLSS** [24] leverages graph-kernel similarity for local expansion; **SLRL** [26] uses an RL-based extractor-expander mechanism; **CLARE** [34] operates on a locator-rewriter framework; and **ProCom** [33] adopts a pre-training prompt paradigm for few-shot target localization.

(ii) Dynamic Tracking Baselines: Optimized across temporal streams. **DevDynaP** [2] executes incremental heuristic exploration; **MFC-TopoReg** [15] enforces topological regularization over dynamic matrix factorization partitions; **QTCS** [18] optimizes query-centered temporal searches; and **CS-TGN** [11] employs a query-driven temporal GNN.

Metrics and Implementation Details. Following [24, 26], we evaluate tracking accuracy across snapshots by comparing predictions against evolving ground-truth boundaries using four metrics: Precision, Recall, F1-score, and Jaccard. All metrics are computed per seed-snapshot pair and macro-averaged across all evaluated seeds and snapshots.

[1] https://snap.stanford.edu/data/

DISRL is implemented in PyTorch on an NVIDIA RTX 3090 GPU (24GB). For fair comparison, *MFC-TopoReg* is post-processed to extract the community containing the target seed; query-based methods (*CS-TGN*, *QTCS*, and *DevDynaP*) are evaluated under single-seed queries. Following [26], the number of generated candidate communities for CLARE and ProCom is set to 5000. Stochastic methods are evaluated over five random seeds, with results reported as mean ± standard deviation; deterministic methods are evaluated once. For each target seed, we exclude its ground-truth community from the available supervision. Specifically, Football uses all remaining annotated communities, while the other datasets randomly sample 100 communities per snapshot and remove the target community if present. Each target seed is tracked independently with its own policy parameters. Snapshots are processed chronologically, using only current and historical supervision, while target-community labels are reserved for evaluation. The detailed hyper-parameter configurations of DISRL are provided in Appendix Table 4. We highlight three key parameters corresponding to our proposed modules: the number of retrieved relevant communities $K = 5$ for Football and $K = 50$ for the other datasets, iGPN propagation layers $L = 2$ for Football and $L = 3$ for the other datasets, and temporal window size $w = 4$ for all datasets. The anonymized source code is available at https://anonymous.4open.science/r/DISRL-code-5D02.

4.2 Main Results

Comparison with Static Baselines. As shown in Table 2, DISRL achieves the best F1-score and Jaccard similarity on the real-world Football dataset and across all four benchmarks under expansion&shrinkage dynamics. Compared with the strongest static baseline, DISRL obtains absolute F1 gains of 8.07, 3.70, 10.01, 10.24, and 4.86 percentage points (pp) on Football, Amazon, DBLP, YouTube, and Twitter, respectively. Static baselines suffer from *temporal isolation*: processing snapshots independently prevents them from using historical trajectories to alleviate sparse supervision. Candidate-generation methods such as *CLARE* and *ProCom* further optimize global candidate pools rather than seed-centered local targets, often producing mismatched boundaries. In contrast, DISRL combines active seed-centric MDP rollouts with adaptive supervision construction, enabling temporally consistent tracking. Under merging&splitting dynamics, DISRL also maintains top F1 and Jaccard performance on both Amazon and DBLP, showing robustness to abrupt boundary mutations.

Comparison with Dynamic Baselines. Table 2 compares DISRL with dynamic tracking models. DISRL achieves the best F1-score and Jaccard similarity across all settings. On the real-world Football dataset, DISRL achieves an F1 gain of 2.89 pp over the strongest dynamic baseline. Under expansion&shrinkage and merging&splitting settings, its largest F1 gains reach 4.25 pp on Twitter and 6.42 pp on DBLP, respectively. Although dynamic baselines exploit temporal information, they remain weakly aligned with seed-centered tracking. Heuristic methods such as *DevDynaP* and *QTCS* rely on topological propagation or proximity thresholds, making them sensitive to edge noise. Matrix-factorization methods such as *MFC-TopoReg* prioritize global clustering consistency and may sacrifice local boundary fidelity. Temporal GNNs such as *CS-TGN* capture structural evolution but still rely on passive node representations

Table 2: Main comparison on simulated evolutionary settings and a real-world dynamic dataset. All metrics are reported in percentages (%). Best and second-best results are marked in bold and underlined within each comparison group, respectively.

Dataset	Static	Prec.	Rec.	F1	Jac.	Dynamic	Prec.	Rec.	F1	Jac.
Real-world dynamic dataset										
	DISRL	**97.00 ± 0.21**	**94.32 ± 0.63**	**95.24 ± 0.36**	**91.33 ± 0.62**	**DISRL**	**97.00 ± 0.21**	94.32 ± 0.63	**95.24 ± 0.36**	**91.33 ± 0.62**
	SLRL	91.99 ± 0.29	88.55 ± 0.90	87.17 ± 0.67	81.03 ± 0.88	CS-TGN	91.02 ± 4.27	**96.22 ± 0.85**	92.35 ± 1.63	87.53 ± 3.40
Football	CLARE	38.86 ± 1.10	47.49 ± 1.39	42.21 ± 1.19	32.02 ± 1.14	MFC-TopoReg	65.38 ± 3.40	62.18 ± 3.20	63.42 ± 3.28	52.92 ± 3.74
	ProCom	29.90 ± 1.58	54.08 ± 3.00	34.41 ± 1.63	23.58 ± 1.42	DevDynaP	88.11	96.14	91.80	86.17
	SLSS	82.09	90.00	85.51	77.42	QTCS	89.59	84.51	86.72	83.08
Expansion&shrinkage setting										
	DISRL	**86.78 ± 0.53**	95.15 ± 0.43	**89.86 ± 0.10**	**83.10 ± 0.16**	**DISRL**	86.78 ± 0.53	95.15 ± 0.43	**89.86 ± 0.10**	**83.10 ± 0.16**
	SLRL	79.77 ± 0.46	**97.34 ± 0.61**	86.16 ± 0.66	77.59 ± 0.99	CS-TGN	86.76 ± 0.69	92.87 ± 0.72	88.22 ± 0.02	80.68 ± 0.10
Amazon	CLARE	86.06 ± 0.03	75.36 ± 0.51	78.54 ± 0.33	68.45 ± 0.30	MFC-TopoReg	78.23 ± 0.08	93.75 ± 0.09	84.10 ± 0.08	74.81 ± 0.10
	ProCom	86.07 ± 0.28	75.79 ± 0.56	78.49 ± 0.45	69.48 ± 0.61	DevDynaP	82.73	**96.03**	87.85	80.01
	SLSS	81.13	93.60	85.36	77.38	QTCS	**93.44**	86.79	88.01	81.38
	DISRL	**76.10 ± 0.51**	83.83 ± 0.49	**78.48 ± 0.46**	**70.26 ± 0.38**	**DISRL**	76.10 ± 0.51	**83.83 ± 0.49**	**78.48 ± 0.46**	**70.26 ± 0.38**
	SLRL	56.92 ± 1.30	**90.90 ± 0.57**	68.47 ± 0.40	55.09 ± 0.51	CS-TGN	**76.31 ± 0.47**	82.97 ± 2.83	75.63 ± 0.65	66.12 ± 0.76
DBLP	CLARE	60.85 ± 0.79	61.83 ± 0.34	59.98 ± 0.49	49.88 ± 0.62	MFC-TopoReg	53.51 ± 0.24	72.89 ± 0.23	60.85 ± 0.24	49.24 ± 0.22
	ProCom	64.60 ± 0.73	72.02 ± 1.31	66.20 ± 1.03	58.16 ± 0.99	DevDynaP	61.96	83.77	69.57	58.46
	SLSS	56.71	81.62	65.38	54.27	QTCS	74.96	81.92	76.63	69.25
	DISRL	**72.02 ± 0.31**	70.93 ± 0.32	**67.76 ± 0.03**	**56.87 ± 0.08**	**DISRL**	**72.02 ± 0.31**	70.93 ± 0.32	**67.76 ± 0.03**	**56.87 ± 0.08**
	SLRL	57.16 ± 1.04	**74.76 ± 1.63**	57.52 ± 1.74	44.21 ± 1.50	CS-TGN	60.02 ± 1.00	**80.28 ± 1.06**	64.57 ± 0.49	51.77 ± 0.56
YouTube	CLARE	51.93 ± 0.22	42.74 ± 0.27	44.72 ± 0.25	33.94 ± 0.34	MFC-TopoReg	29.49 ± 0.55	38.85 ± 0.59	32.14 ± 0.54	22.23 ± 0.42
	ProCom	52.23 ± 0.39	49.94 ± 0.93	48.81 ± 0.78	38.97 ± 0.70	DevDynaP	50.74	61.99	53.67	41.72
	SLSS	53.24	63.66	56.01	45.48	QTCS	58.60	66.00	60.08	51.90
	DISRL	**65.32 ± 1.39**	64.11 ± 0.68	**61.30 ± 0.41**	**48.17 ± 0.53**	**DISRL**	**65.32 ± 1.39**	64.11 ± 0.68	**61.30 ± 0.41**	**48.17 ± 0.53**
	SLRL	48.35 ± 0.06	**77.34 ± 0.23**	56.44 ± 0.02	41.57 ± 0.07	CS-TGN	49.16 ± 2.84	**88.49 ± 1.28**	57.05 ± 1.81	43.56 ± 1.93
Twitter	CLARE	39.24 ± 0.59	35.15 ± 0.48	35.30 ± 0.52	23.66 ± 0.47	MFC-TopoReg	31.81 ± 0.81	40.57 ± 1.17	34.41 ± 0.92	23.92 ± 0.69
	ProCom	46.85 ± 0.84	50.32 ± 1.11	42.63 ± 0.61	31.55 ± 0.63	DevDynaP	33.87	56.80	40.87	28.85
	SLSS	47.44	59.72	50.88	38.34	QTCS	38.33	53.27	42.83	30.92
Merging&splitting setting										
	DISRL	77.71 ± 0.57	95.51 ± 0.28	**83.49 ± 0.50**	**73.97 ± 0.75**	**DISRL**	77.71 ± 0.57	95.51 ± 0.28	**83.49 ± 0.50**	**73.97 ± 0.75**
	SLRL	68.89 ± 1.08	**96.84 ± 0.88**	77.00 ± 1.35	66.36 ± 1.97	CS-TGN	84.32 ± 0.24	75.38 ± 1.35	76.32 ± 1.03	64.69 ± 1.18
Amazon	CLARE	**86.03 ± 0.24**	63.85 ± 0.42	69.12 ± 0.36	56.65 ± 0.38	MFC-TopoReg	71.24 ± 0.18	**96.43 ± 0.20**	80.00 ± 0.19	69.10 ± 0.21
	ProCom	85.72 ± 0.30	62.80 ± 0.66	68.05 ± 0.40	56.05 ± 0.54	DevDynaP	79.64	67.00	66.71	55.37
	SLSS	84.05	73.69	73.70	62.19	QTCS	**95.91**	57.69	66.95	55.65
	DISRL	53.34 ± 1.02	75.09 ± 1.40	**56.63 ± 1.27**	**44.56 ± 1.49**	**DISRL**	53.34 ± 1.02	**75.09 ± 1.40**	**56.63 ± 1.27**	**44.56 ± 1.49**
	SLRL	37.76 ± 1.05	**86.67 ± 0.19**	46.74 ± 0.14	34.63 ± 0.17	CS-TGN	48.24 ± 2.19	68.91 ± 2.14	50.21 ± 0.87	36.51 ± 0.86
DBLP	CLARE	53.28 ± 0.93	49.05 ± 0.53	45.39 ± 0.09	33.57 ± 0.12	MFC-TopoReg	25.64 ± 0.28	57.03 ± 0.31	31.74 ± 0.31	21.66 ± 0.25
	ProCom	**55.76 ± 1.00**	55.60 ± 0.76	49.92 ± 0.54	38.04 ± 0.47	DevDynaP	50.95	53.93	45.39	33.90
	SLSS	44.50	56.26	44.75	33.61	QTCS	**63.40**	49.80	49.61	39.39

Table 3: Ablation results. Each cell reports F1/Jaccard.

Variant	Real-world dynamics	Expansion&shrinkage		Merging&splitting	
	Football	Amazon	DBLP	Amazon	DBLP
w/o cc	93.42 ± 0.27/87.94 ± 0.44	88.02 ± 0.12/80.32 ± 0.16	75.89 ± 0.13/67.39 ± 0.23	81.48 ± 0.52/71.40 ± 0.87	52.54 ± 1.18/40.81 ± 1.40
w/o rcs	94.23 ± 0.30/90.18 ± 0.49	88.28 ± 0.15/80.63 ± 0.19	76.53 ± 0.24/67.67 ± 0.32	82.86 ± 0.44/73.35 ± 0.31	53.58 ± 1.08/41.65 ± 1.24
w/o sim	88.29 ± 1.86/81.95 ± 2.15	87.28 ± 1.00/79.27 ± 1.42	75.29 ± 1.71/66.80 ± 1.66	81.76 ± 1.94/71.46 ± 1.95	53.25 ± 2.24/41.72 ± 2.49
w/o win	92.97 ± 0.46/88.20 ± 0.62	87.14 ± 0.54/78.92 ± 0.64	74.81 ± 0.41/65.07 ± 0.76	79.96 ± 1.14/69.68 ± 1.17	49.07 ± 1.65/37.66 ± 1.53
w/o decay	94.86 ± 0.30/90.67 ± 0.49	88.53 ± 0.46/81.17 ± 0.80	77.08 ± 0.84/68.57 ± 1.12	82.46 ± 0.36/72.84 ± 0.54	54.98 ± 1.58/43.20 ± 1.40
DISRL	**95.24 ± 0.36/91.33 ± 0.62**	**89.86 ± 0.10/83.10 ± 0.16**	**78.48 ± 0.46/70.26 ± 0.38**	**83.49 ± 0.50/73.97 ± 0.75**	**56.63 ± 1.27/44.56 ± 1.49**

without explicit boundary expansion or stopping decisions. In contrast, DISRL dynamically constructs seed-specific supervision and employs dual-scale policy updating, effectively reducing supervision mismatch and error accumulation in evolving graphs.

4.3 Ablation Study

To investigate the contributions of the modules of DISRL, we conduct five ablation experiments, with results reported in Table 3.

Supervision Construction Module. To evaluate reference community synthesis, we consider two variants: *w/o cc*, which removes connected-component extraction and directly inherits previous predictions as reference communities; and *w/o rcs*, which replaces reference community synthesis with the two-hop neighborhood of the target seed. Their performance degradation highlights the importance of constructing topology-consistent reference communities for reliable seed-specific supervision. To assess relevant community retrieval, we further introduce *w/o sim*, which randomly

selects K known communities instead of using structural similarity. Its inferior performance suggests that structural similarity-based retrieval helps reduce supervision mismatch and provides more relevant guidance for community boundary expansion.

Dual-scale Community Exploration Module. To investigate the effect of window-level master anchoring, we consider *w/o win*, which removes window-level master anchoring and continuously transfers the fast policy across snapshots. Its performance degradation suggests that window-level anchoring helps mitigate the accumulation of localized adaptation errors across snapshots. We further evaluate temporal-decay gradient aggregation through *w/o decay*, which replaces temporal-decay weighting with uniform averaging. The inferior performance of this variant indicates that emphasizing more recent transition gradients benefits policy adaptation across windows. Overall, DISRL consistently outperforms all five variants under real-world dynamics, expansion&shrinkage, and merging&splitting settings.

4.4 Long-Horizon Error Mitigation Analysis

To empirically investigate the error-mitigation mechanism, we compare DISRL with its window-free variant (*w/o win*) across 24 consecutive snapshots of the real-world Football dataset, spanning six temporal windows ($w = 4$). As illustrated in Fig. 2, both variants achieve comparable performance in the early snapshots, whereas their performance diverges more noticeably in later stages. Specifically, the window-free variant exhibits more pronounced performance degradation and fluctuations, while DISRL maintains relatively high F1-score and Jaccard similarity throughout the sequence. As established by Corollary 1 in Section 3.4, periodic master-policy anchoring restricts direct fast-parameter error propagation to each temporal window. These experimental observations are consistent with this theoretical analysis, suggesting that the dual-scale mechanism helps mitigate the adverse effects of accumulated adaptation errors on long-horizon tracking performance.

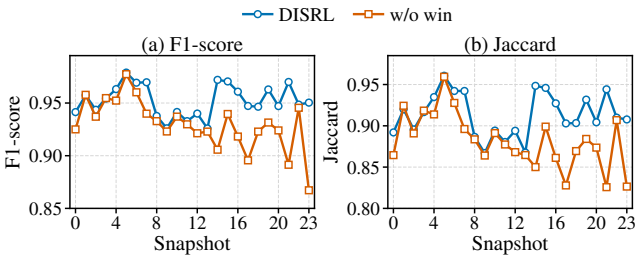

Figure 2: Long-horizon tracking performance on Football.

4.5 Hyper-parameter Sensitivity

Retrieved relevant communities. We report the supervision retrieval parameter K in Fig. 3. For the simulated-evolution datasets, DISRL achieves stable performance with $K = 50$, balancing supervision coverage and structural relevance. When $K = 100$, all known communities are included, leading to performance degradation due to potentially irrelevant supervision. On the Football dataset, we evaluate K from 1 to 10, with the best F1-score achieved at $K = 5$. Performance initially improves as K increases, but gradually declines beyond this point, highlighting the importance of selecting an appropriate number of relevant communities.

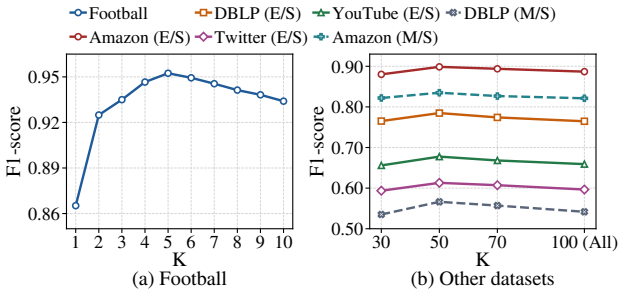

Figure 3: Effect of the retrieved community number K.

iGPN Propagation Layers. We then study the sensitivity of DISRL to the number of iGPN propagation layers L used by the localized state encoder on the real-world Football dataset and three expansion&shrinkage benchmarks (Amazon, DBLP, and Twitter). As shown in Fig. 4, $L = 2$ achieves the best performance on Football, whereas $L = 3$ achieves the best performance on Amazon, DBLP, and Twitter. As L controls the depth of localized structural propagation, too few layers may miss multi-hop boundary context, whereas too many layers may introduce remote structural noise into seed-centered representations. These results suggest that two-layer propagation is sufficient for Football, while three-layer localized propagation better balances structural context modeling and noise control on the other datasets.

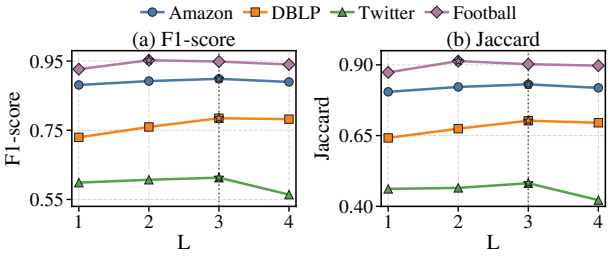

Figure 4: Sensitivity to iGPN propagation.

Temporal Window Horizon. We further evaluate the sensitivity of DISRL to the temporal window size $w \in \{2, 3, 4, 5, 6\}$ on Football, Amazon, and DBLP, with the latter two evaluated under the expansion&shrinkage setting. As illustrated in Fig. 5, $w = 4$ achieves the best overall performance across these three datasets, indicating a trade-off between policy stabilization and temporal drift:

- **Under-sized Windows ($w = 2$):** Restricting the horizon induces *temporal myopia*. The fast parameters ($\Phi_{m,t}$) overfit to transient single-frame fluctuations, starving the master anchor (Θ_m) of continuous, long-horizon meta-gradients.
- **Over-sized Windows ($w = 6$):** Expanding the horizon excessively introduces optimization lag. Under non-stationary graph streams, historical transition gradients from distant snapshots carry stale topological signals, causing *gradient smearing* that misaligns the window-level master update.

Consequently, setting the default horizon to $w = 4$ harmonizes short-term micro-evolution tracking with long-term macro-trend anchoring, mitigating the risk of long-horizon error accumulation.

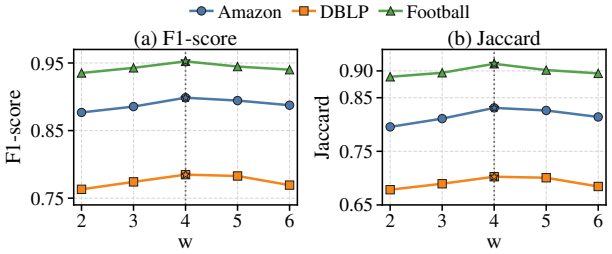

Figure 5: Sensitivity to temporal window size.

4.6 Additional Experimental Analyses

The qualitative case study and efficiency comparison are provided in Appendices C and D, respectively. DISRL achieves up to 10.35× speedup over CS-TGN and effectively tracks evolving community boundaries through adaptive supervision construction and dual-scale policy consolidation.

5 Conclusion

In this paper, we proposed DISRL, a dual-scale incremental semi-supervised reinforcement learning framework for dynamic local community tracking. DISRL addresses temporal supervision uncertainty by constructing seed-specific adaptive guidance from sparse communities, and mitigates long-horizon error cascades through a dual-scale policy optimization strategy. We further developed a dynamic graph benchmark with diverse community evolution patterns. Extensive experiments demonstrate that DISRL consistently outperforms state-of-the-art baselines with improved efficiency and tracking stability.

Acknowledgments

This work was supported by the National Natural Science Foundation of China under Grant Nos. 62302451 and 62276233.

References

[1] [n. d.]. Century of College Football Dataset. GitHub repository. https://github.com/JacobH140/century-of-college-football Accessed: 2026-10-08.

[2] S. Bakhtar and H. A. Harutyunyan. 2025. DevDynaP: A dynamic local community detection algorithm. *SN Computer Science* 6, 4 (2025), 299. doi:10.1007/s42979-025-03801-z

[3] A. Bhowmick, M. Kosan, Z. Huang, Ambuj Singh, and S. Medya. 2024. DG-CLUSTER: A neural framework for attributed graph clustering via modularity maximization. *Proceedings of the AAAI Conference on Artificial Intelligence* 38, 10 (2024), 11069–11077. doi:10.1609/aaai.v38i10.28983

[4] M. Bianco, L. Cagliero, and L. Vassio. 2025. Semi-supervised community detection in dynamic graphs. In *Proceedings of the Workshops of the EDBT/ICDT 2025 Joint Conference (CEUR Workshop Proceedings, Vol. 3946)*. https://ceur-ws.org/Vol-3946/DARLI-AP-1.pdf

[5] V. D. Blondel, J.-L. Guillaume, R. Lambiotte, and E. Lefebvre. 2008. Fast unfolding of communities in large networks. *Journal of Statistical Mechanics: Theory and Experiment* 2008, 10 (2008), P10008. doi:10.1088/1742-5468/2008/10/P10008

[6] K. M. Borgwardt and H.-P. Kriegel. 2005. Shortest-path kernels on graphs. In *Proceedings of the 5th IEEE International Conference on Data Mining*. 74–81. doi:10.1109/ICDM.2005.132

[7] D. Chakrabarti, R. Kumar, and A. Tomkins. 2006. Evolutionary clustering. In *Proceedings of the 12th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 554–560. doi:10.1145/1150402.1150467

[8] A. R. Costa and C. G. Ralha. 2023. AC2CD: An actor-critic architecture for community detection in dynamic social networks. *Knowledge-Based Systems* 261 (2023), 110202. doi:10.1016/j.knosys.2022.110202

[9] A. Grover and J. Leskovec. 2016. node2vec: Scalable feature learning for networks. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 855–864. doi:10.1145/2939672.2939754

[10] Moritz Hardt, Benjamin Recht, and Yoram Singer. 2016. Train faster, generalize better: Stability of stochastic gradient descent. In *Proceedings of the 33rd International Conference on Machine Learning*. PMLR, 1225–1234.

[11] F. Hashemi, A. Behrouz, and M. R. Hajidehi. 2023. CS-TGN: Community search via temporal graph neural networks. In *Companion Proceedings of the ACM Web Conference*. 1196–1203. doi:10.1145/3543873.3587654

[12] D. He, Y. Song, D. Jin, Z. Feng, B. Zhang, Z. Yu, and W. Zhang. 2020. Community-centric graph convolutional network for unsupervised community detection. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence*. 3515–3521. doi:10.24963/ijcai.2020/486

[13] G. He, Z. Xue, Z. Jiang, Y. Kang, S. Zhao, and W. Lu. 2023. H^2CGL: Modeling dynamics of citation network for impact prediction. *Information Processing & Management* 60, 6 (2023), 103512. doi:10.1016/j.ipm.2023.103512

[14] P. W. Holland, K. B. Laskey, and S. Leinhardt. 1983. Stochastic blockmodels: First steps. *Social Networks* 5, 2 (1983), 109–137. doi:10.1016/0378-8733(83)90021-7

[15] D. Kong, A. Zhang, and Y. Li. 2024. Learning persistent community structures in dynamic networks via topological data analysis. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 38. 8617–8626. doi:10.1609/aaai.v38i8.28706

[16] D. Kuang, C. Ding, and H. Park. 2012. Symmetric nonnegative matrix factorization for graph clustering. In *Proceedings of the SIAM International Conference on Data Mining*. 106–117. doi:10.1137/1.9781611972825.10

[17] X. Li, X. Zhen, X. Qi, H. Han, L. Zhang, and Z. Han. 2023. Dynamic community detection based on graph convolutional networks and contrastive learning. *Chaos, Solitons & Fractals* 176 (2023), 114157. doi:10.1016/j.chaos.2023.114157

[18] L. Lin, P. Yuan, R.-H. Li, C. Zhu, H. Qin, H. Jin, and T. Jia. 2024. QTCS: Efficient query-centered temporal community search. *Proceedings of the VLDB Endowment* 17, 6 (2024), 1187–1199. doi:10.14778/3648160.3648163

[19] Y.-R. Lin, Y. Chi, S. Zhu, H. Sundaram, and B. L. Tseng. 2008. FacetNet: A framework for analyzing communities and their evolutions in dynamic networks. In *Proceedings of the 17th International Conference on World Wide Web*. 685–694. doi:10.1145/1367497.1367590

[20] M. Liu, Y. Liu, K. Liang, W. Tu, S. Wang, S. Zhou, and X. Liu. 2024. Deep Temporal Graph Clustering. In *The Twelfth International Conference on Learning Representations*. https://openreview.net/forum?id=ViNe1fjGME

[21] Liam Madden, Stephen Becker, and Emiliano Dall'Anese. 2020. Bounds for the tracking error of first-order online optimization methods. *arXiv preprint arXiv:2003.02400* (2020).

[22] M. Mazza, G. Cola, and M. Tesconi. 2023. Modularity-based approach for tracking communities in dynamic social networks. *Knowledge-Based Systems* 281 (2023), 111067. doi:10.1016/j.knosys.2023.111067

[23] M. E. J. Newman. 2006. Modularity and community structure in networks. *Proceedings of the National Academy of Sciences* 103, 23 (2006), 8577–8582. doi:10.1073/pnas.0601602103

[24] L. Ni, J. Ge, Y. Zhang, W. Luo, and V. S. Sheng. 2024. Semi-supervised local community detection. *IEEE Transactions on Knowledge and Data Engineering* 36, 2 (2024), 823–839. doi:10.1109/TKDE.2023.3290095

[25] L. Ni, H. Xu, L. Mu, Y. Zhang, and W. Luo. 2025. Pre-trained prompt-driven semi-supervised local community detection. arXiv:2505.12304.

[26] L. Ni, R. Ye, W. Luo, Y. Zhang, L. Zhang, and V. S. Sheng. 2025. SLRL: Semi-supervised local community detection based on reinforcement learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 39. 631–639. doi:10.1609/aaai.v39i1.32044

[27] Alex Nichol, Joshua Achiam, and John Schulman. 2018. On First-Order Meta-Learning Algorithms. *arXiv preprint arXiv:1803.02999* (2018).

[28] G. Paoletti, L. Gioacchini, M. Mellia, L. Vassio, and J. Almeida. 2025. CoDÆN: Benchmarks and comparison of evolutionary community detection algorithms for dynamic networks. *ACM Transactions on the Web* 19, 3 (2025), 1–25. doi:10.1145/3718988

[29] N. Park, R. A. Rossi, E. Koh, I. A. Burhanuddin, S. Kim, F. Du, N. K. Ahmed, and C. Faloutsos. 2022. CGC: Contrastive graph clustering for community detection and tracking. In *Proceedings of the ACM Web Conference*. 1115–1126. doi:10.1145/3485447.3512160

[30] B. Perozzi, R. Al-Rfou, and S. Skiena. 2014. DeepWalk: Online learning of social representations. In *Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. 701–710. doi:10.1145/2623330.2623732

[31] E. Rossi, B. Chamberlain, F. Frasca, D. Eynard, F. Monti, and M. Bronstein. 2020. Temporal graph networks for deep learning on dynamic graphs. arXiv:2006.10637.

[32] X. Wang, J. Guo, X. Luo, and H. Yu. 2024. DyHDGE: Dynamic heterogeneous transaction graph embedding for safety-centric fraud detection in financial scenarios. *Journal of Safety Science and Resilience* 5, 4 (2024), 486–497. doi:10.1016/j.jnlssr.2024.05.005

[33] X. Wu, K. Xiong, Y. Xiong, X. He, Y. Zhang, Y. Jiao, and J. Zhang. 2024. ProCom: A few-shot targeted community detection algorithm. In *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. 3414–3424. doi:10.1145/3637528.3671749

[34] X. Wu, Y. Xiong, Y. Zhang, Y. Jiao, C. Shan, Y. Sun, Y. Zhu, and P. S. Yu. 2022. CLARE: A semi-supervised community detection algorithm. In *Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. 2059–2069. doi:10.1145/3534678.3539370

[35] H. Xiong, L. Sun, Y. Xiong, S. Zhou, H. Ren, and Y. Zhu. 2026. SLFM: Semi-supervised local community detection based on hyperbolic flow matching. In *Proceedings of the ACM Web Conference*. 1357–1367. doi:10.1145/3774904.3792548

[36] T. Yang, Y. Chi, S. Zhu, Y. Gong, and R. Jin. 2011. Detecting communities and their evolutions in dynamic social networks: A Bayesian approach. *Machine Learning* 82, 2 (2011), 157–189. doi:10.1007/s10994-010-5214-7

[37] M. Zafarmand, Y. Talebirad, E. Austin, C. Largeron, and O. R. Zaiane. 2023. Fast local community discovery relying on the strength of links. *Social Network Analysis and Mining* 13 (2023), 112. doi:10.1007/s13278-023-01115-7

[38] Y. Zhang, Y. Xiong, Y. Ye, T. Liu, W. Wang, Y. Zhu, and P. S. Yu. 2020. SEAL: Learning heuristics for community detection with generative adversarial networks. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*. 1103–1113. doi:10.1145/3394486.3403154

[39] Y. Zheng, X. Zhang, S. Chen, X. Zhang, X. Yang, and D. Wang. 2023. When convolutional network meets temporal heterogeneous graphs: An effective community detection method. *IEEE Transactions on Knowledge and Data Engineering* 35, 2 (2023), 2173–2178. doi:10.1109/TKDE.2021.3096122

[40] Y. Zhou, Y. Jiang, Y. Fang, W. Luo, Y. Hu, Y. Hu, and C. Chen. 2025. Effective durable community search in large temporal graph. *Proceedings of the VLDB Endowment* 19, 2 (2025), 127–140. doi:10.14778/3773749.3773753

[41] Di Zhuang, J. Morris Chang, and Mingchen Li. 2021. DynaMo: Dynamic Community Detection by Incrementally Maximizing Modularity. *IEEE Transactions on Knowledge and Data Engineering* 33, 5 (2021), 1934–1945. doi:10.1109/TKDE.2019.2951419

A Theoretical Analysis

We analyze how localized adaptation errors propagate and accumulate during sequential fast-parameter transfer within a temporal window $\mathcal{W}_m$ spanning from t_s to t_e. Let $\Phi_{m,t}$ and $\Phi^\star_{m,t}$ denote the actual fast parameter and the ideal error-free parameter at snapshot G^t, respectively.

Following time-varying optimization analyses [21], the snapshot-level tracking error is defined as $d_{m,t} = \left\|\Phi_{m,t} - \Phi^\star_{m,t}\right\|_2$, using the standard L2 norm. We establish the tracking dynamics under three standard stability assumptions.

Assumption 1 (Bounded Adaptation Perturbation). The local optimization perturbation $\boldsymbol{\xi}_{m,t}$ introduced during adaptation on G^t is uniformly bounded; that is, there exists a constant $\xi \geq 0$ such that

$$\left\|\boldsymbol{\xi}_{m,t}\right\|_2 \leq \xi. \tag{24}$$

Assumption 2 (Bounded Ideal-parameter Drift). The temporal variation of the ideal tracking target across adjacent snapshots is bounded; that is, there exists a constant $\Delta \geq 0$ such that

$$\left\|\Phi^\star_{m,t+1} - \Phi^\star_{m,t}\right\|_2 \leq \Delta. \tag{25}$$

Assumption 3 (Lipschitz Adaptation Operator). For theoretical analysis, let $\mathcal{U}_t(\cdot)$ denote the error-free counterpart of the snapshot-level adaptation on G^t. Corresponding to the policy-gradient update and fast-parameter transfer defined in the main paper, the adapted parameter and its subsequent transfer can be written as

$$\Phi'_{m,t} = \mathcal{U}_t(\Phi_{m,t}) + \boldsymbol{\xi}_{m,t}, \qquad \Phi_{m,t+1} \leftarrow \Phi'_{m,t}. \tag{26}$$

Following standard stability analyses of gradient-based optimization [10], we assume that $\mathcal{U}_t(\cdot)$ is uniformly Lipschitz continuous. Specifically, there exists a constant $\kappa \geq 0$ such that

$$\left\|\mathcal{U}_t(\Phi_a) - \mathcal{U}_t(\Phi_b)\right\|_2 \leq \kappa \left\|\Phi_a - \Phi_b\right\|_2. \tag{27}$$

The ideal parameter $\Phi^\star_{m,t}$ is assumed to be a fixed point of the error-free adaptation on G^t, i.e., $\mathcal{U}_t(\Phi^\star_{m,t}) = \Phi^\star_{m,t}$.

Proposition 1 (Single-step Error Propagation). Under Assumptions 1–3, for any adjacent snapshots G^t and G^{t+1} within the same temporal window $\mathcal{W}_m$, the fast-parameter error satisfies

$$d_{m,t+1} \leq \kappa d_{m,t} + \xi + \Delta. \tag{28}$$

Proof. By the definition of $d_{m,t+1}$ and the fast-parameter transfer relation in Eq. (26), we have

$$\begin{aligned} d_{m,t+1} &= \left\|\Phi_{m,t+1} - \Phi^\star_{m,t+1}\right\|_2 \\ &= \left\|\mathcal{U}_t(\Phi_{m,t}) + \boldsymbol{\xi}_{m,t} - \Phi^\star_{m,t+1}\right\|_2 \\ &\leq \left\|\mathcal{U}_t(\Phi_{m,t}) - \mathcal{U}_t(\Phi^\star_{m,t})\right\|_2 + \left\|\boldsymbol{\xi}_{m,t}\right\|_2 \\ &\quad + \left\|\mathcal{U}_t(\Phi^\star_{m,t}) - \Phi^\star_{m,t+1}\right\|_2. \end{aligned}$$

The first term is bounded by $\kappa d_{m,t}$ by Assumption 3, and the second term is bounded by ξ by Assumption 1. For the third term, using $\mathcal{U}_t(\Phi^\star_{m,t}) = \Phi^\star_{m,t}$ and Assumption 2, we obtain

$$\begin{aligned} &\left\|\mathcal{U}_t(\Phi^\star_{m,t}) - \Phi^\star_{m,t+1}\right\|_2 \\ &\quad = \left\|\Phi^\star_{m,t} - \Phi^\star_{m,t+1}\right\|_2 \leq \Delta. \end{aligned}$$

Combining the three bounds yields Eq. (28). □

Long-horizon Error Accumulation. By recursively applying Eq. (28) from the window start t_s to any snapshot $t \in \mathcal{W}_m$, we obtain

$$d_{m,t} \leq \kappa^{t-t_s} d_{m,t_s} + (\xi + \Delta) \sum_{r=0}^{t-t_s-1} \kappa^r. \tag{29}$$

When $0 \leq \kappa < 1$, the geometric sum is bounded by $1/(1-\kappa)$, so the tracking-error upper bound remains bounded for a fixed window-initial error. When $\kappa = 1$, this reduces to

$$d_{m,t} \leq d_{m,t_s} + (t - t_s)(\xi + \Delta), \tag{30}$$

showing that the tracking-error upper bound grows at most linearly with the length of continuous fast-parameter transfer. When $\kappa > 1$, Eq. (29) permits geometric growth of the worst-case bound. These upper bounds do not imply that the actual tracking error necessarily increases. Nevertheless, without periodic policy anchoring, continuous fast-parameter transfer may allow historical errors to accumulate or amplify, increasing the risk of error cascades, particularly when $\kappa \geq 1$.

Illustrative Example of Unbounded Error Accumulation. To illustrate that unbounded tracking-error accumulation is possible, consider a one-dimensional adaptation process over an arbitrarily long sequence without periodic anchoring. Let $\mathcal{U}_t(x) = x$, the ideal parameter be $x^\star_t = 0$, and the adaptation perturbation be a constant $\epsilon > 0$. Starting from $x_0 = 0$, the perturbed adaptation satisfies

$$x_{t+1} = x_t + \epsilon, \qquad x_t = t\epsilon.$$

Consequently, the tracking error is $d_t = |x_t - x^\star_t| = t\epsilon$, which grows without bound as $t \to \infty$. This construction satisfies Assumptions 1–3 with $\xi = \epsilon$, $\Delta = 0$, and $\kappa = 1$. It demonstrates that unbounded error accumulation is possible under the stated assumptions.

Corollary 1 (Window-confined Error Propagation). Let $w = |\mathcal{W}_m| = t_e - t_s + 1$ denote the size of temporal window $\mathcal{W}_m$. At the beginning of each window, DISRL initializes the snapshot-level fast parameter from the window-level master parameter, i.e., $\Phi_{m,t_s} \leftarrow \Theta_m$. Since the fast parameter is directly transferred only within the

current window, the recurrence in Eq. (29) can be expanded for at most $w-1$ consecutive transition steps.

When $\kappa = 1$, the window-end error is therefore bounded by

$$d_{m,t_e} \leq d_{m,t_s} + (w-1)(\xi + \Delta). \tag{31}$$

This bound highlights the role of periodic master-policy anchoring in limiting error propagation within each temporal window. Specifically, the additional error contribution from continuous fast-parameter transfer is bounded by $(w-1)(\xi+\Delta)$ when $\kappa = 1$, while the window-initial error $d_{m,t_s} = \|\Theta_m - \Phi^{\star}_{m,t_s}\|_2$ is determined by the evolving master parameter. This window-level bound does not imply a uniform error bound across windows, as the master-parameter initialization error remains unbounded by the analysis.

In contrast, without window-level anchoring, continuous fast-parameter transfer over the entire graph sequence would make the direct accumulation length grow with the total number of snapshots. Since the window size w is fixed and smaller than the full sequence length T, the dual-scale design restricts direct error propagation to each temporal window through periodic master-policy anchoring, thereby mitigating the risk of unbounded long-horizon drift.

B Evaluation Settings and Hyper-parameters

B.1 Dynamic Generation Procedures

We implement four core topological operators alongside continuous background noise (random node/edge insertions and deletions) to evolve graphs into temporal snapshot sequences:

- **Expansion & Shrinkage:** Expansion injects boundary nodes mimicking local homophily connectivity patterns, whereas shrinkage prunes peripheral members while preserving the structural core.
- **Merging & Splitting:** Merging introduces bridge edges between structurally related community pairs to unify their labels, whereas splitting fractures a cohesive group into disjoint, internally coherent sub-components.

B.2 Evaluation Protocols

We evaluate simulated evolution under two scenarios and additionally use a real-world dynamic dataset:

- **Expansion & Shrinkage Dynamic Setting:** This setting is evaluated on the four SNAP datasets. Starting from $t = 1$, Amazon and DBLP undergo six consecutive expansion steps followed by six shrinkage steps, and then twelve additional steps in which expansion or shrinkage is randomly selected at each step, producing $T = 25$ snapshots. YouTube and Twitter undergo six consecutive expansion steps followed by six shrinkage steps, producing $T = 13$ snapshots. We randomly sample 100 target seed nodes from each of the Amazon, DBLP, YouTube, and Twitter datasets.
- **Merge & Split Dynamic Setting:** This setting is evaluated on Amazon and DBLP due to their sharp modularity boundaries. Each sequence undergoes four merge operations followed by four split operations, producing $T = 9$ snapshots over 100 sampled target seed nodes per dataset.
- **Real-world Dynamic Setting:** Football is derived from the Century of College Football dataset [1]. We aggregate games from 2000–2024 into 24 yearly snapshots, excluding 2020 due to the COVID-19 pandemic, and use each team's annual conference affiliation as its community label. We track 10 fixed target teams present in all snapshots, without introducing any artificial community evolution events.

The query-specific supervision exclusion is performed independently at each snapshot by the evaluation procedure. The target community's ground-truth annotation is used to identify and remove it from the candidate supervision pool and to compute evaluation metrics, but is not supplied as training supervision for that query. For each independently tracked seed, the fast policy is adapted using the available annotations before target prediction. Upon the arrival of the next snapshot, its permitted supervision is used to compute transition meta-gradients, which are buffered rather than immediately applied. These gradients are aggregated to update the master policy only at window boundaries, without accessing unobserved future snapshots.

B.3 Hyper-parameter Configurations

Table 4 summarizes the key hyper-parameter configurations used in the experiments.

Table 4: Hyper-parameter configurations of DISRL.

Component	Configuration	Value
Supervision Retrieval	Known communities	100,all except target
	Top-K	50,5
Community Explorer	Batch size	16
	Epochs per snapshot	50
	Hidden dimension	64
	iGPN layers	3,2
	MLP layers	2
	Learning rate	10^{-2}
	Discount factor γ	0.99
Dual-scale Update	Window size	4
	Decay factor ρ	0.8
	Learning rate η	0.01

C Case Study

To visually dissect the boundary tracking behavior of DISRL under fine-grained topological mutations, Fig. 6 qualitatively orchestrates a temporal case study on DBLP under expansion&shrinkage dynamics. The tracking sequence tracks a specific target seed node s across four consecutive snapshots (t_0 to t_3), during which the local ground-truth community size dynamically swells from 7 to 16 nodes.

Empirically, DISRL demonstrates remarkable adaptive stabilization throughout this structural expansion, registering high F1-scores of 0.93, 0.95, 0.96, and 0.97 chronologically. In sharp contrast, both *ProCom* [33] and *CS-TGN* [11] suffer from severe synchronization delay at the initial onset (t_0), stagnating at an F1-score of 0.55. As the stream advances to t_1, *ProCom* fails to capture the newly incorporated members and its tracking accuracy drops to 0.47. Concurrently, while the dynamic GNN baseline *CS-TGN* slowly recaptures the community boundary through continuous representation updates, its accuracy consistently trails DISRL across the entire temporal window.

This qualitative divergence provides a clear microcosm of the underlying algorithmic mechanisms. The early-stage performance

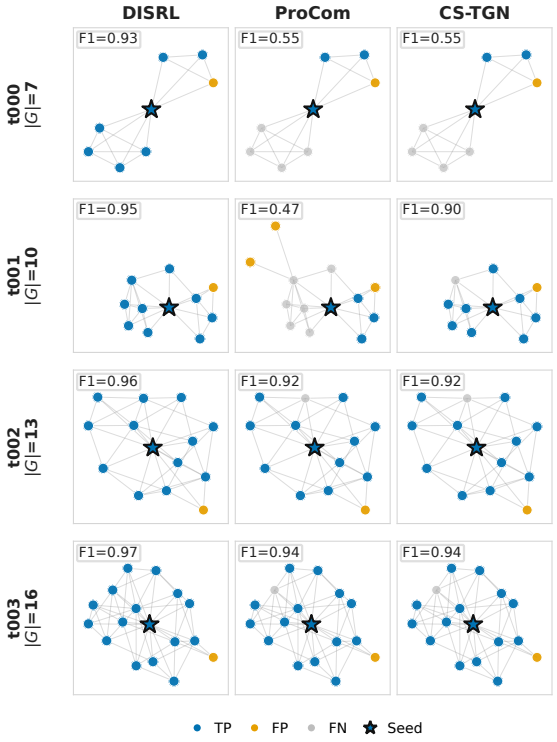

Figure 6: DBLP case study over four snapshots. Rows show snapshots and columns show methods; each panel reports F1. The target community grows from 7 to 16 nodes.

collapse of *ProCom* and *CS-TGN* uncovers their vulnerability to cold-start topological changes; without specialized localized historical anchors, they easily over-expand into adjacent dense graph fragments. Conversely, DISRL circumvents this tracking lag via its dual-stage coordination. At the initial frame t_0, the cold-start bootstrapping policy isolates the immediate vicinity of seed s. In subsequent snapshots ($t_1 \rightarrow t_3$), the reference community synthesis operator explicitly applies topological filtering through the reference-synthesis rule defined in the main paper to prune disconnected structural noise. This mechanism allows the explorer agent to seamlessly inherit previous predictions as historical context, accurately identifying boundary-expanding members while preserving temporal tracking continuity.

D Efficiency Study

To evaluate the computational efficiency of DISRL, we benchmark its wall-clock execution time against competitive baselines under the expansion&shrinkage setting. As illustrated in Fig. 7, DISRL consistently outperforms all baselines, achieving the lowest execution latency across all three datasets. Compared with CS-TGN, DISRL obtains 1.96×, 4.78×, and 10.35× speedups, respectively. The advantage is especially clear on Twitter, where the larger and denser topology increases the cost of repeatedly processing temporal graph structures. The efficiency difference comes from the execution patterns of the compared methods. CS-TGN learns temporal GNN

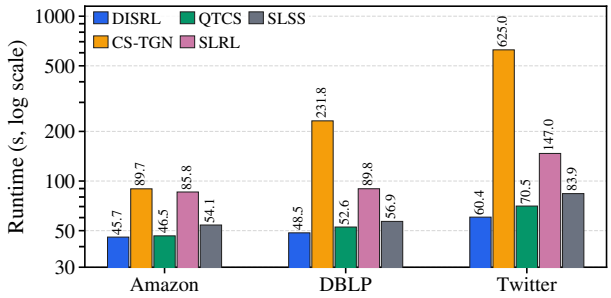

Figure 7: Efficiency comparison under the expansion&shrinkage setting. The reported runtime is normalized per target seed per snapshot over 50 independent queries.

representations over all snapshot graphs and repeatedly scans training queries across epochs. SLRL retrains a local detector for each seed-snapshot pair, where each detector performs two optimization rounds and repeatedly updates the expander across multiple epochs. SLSS incurs costly similar-community matching during community expansion, while QTCS is slowed by seed-wise TPPR-guided ALS search over temporal transition states. In contrast, DISRL performs policy adaptation on a compact Top-K seed-specific relevant supervision set and localized subgraphs, avoiding full-graph temporal representation learning.