Gate Initialisation versus Recurrent Dropout and Zoneout in Learning Long Dependencies: A Controlled Study of LSTM and GRU Networks on Synthetic Tasks

Oleksandr Kholodniak *o.o.kholodniak@khai.edu*
National Aerospace University "Kharkiv Aviation Institute", Kharkiv, Ukraine

Oleksandr Prokhorov *o.prokhorov@khai.edu*
National Aerospace University "Kharkiv Aviation Institute", Kharkiv, Ukraine

arXiv:2610.29034v1 [cs.LG] 21 Oct 2026

Abstract

Dropout and zoneout are routinely added to long short-term memory (LSTM) and gated recurrent unit (GRU) networks, but, to our knowledge, their placements have not been crossed with gate initialisation in a factorial design, and it is not known whether early gradient propagation tells which runs will learn long dependencies. We trained one-layer LSTMs and GRUs on the adding problem and on delayed recall in a factorial grid of 720 runs, crossing conventional and chrono gate initialisation with up to six recurrent regulariser placements at one rate each, with pre-specified follow-ups, and recorded how sensitive the final hidden state was to the input at every lag. Because both tasks supply fresh or exhaustive training data, the study measures the regularisers' effects on optimisation and does not assess generalisation from a finite training set. At lengths chosen to make conventional initialisation difficult, all 360 chrono-initialised runs reached the validation criterion within 5000 updates, as did 119 of 120 unregularised runs with a horizon about four times too short or too long, whereas conventionally initialised networks learned slowly or did not reach the criterion within the budget at the longer lengths. Under conventional initialisation, dropout generally delayed learning and zoneout helped only in some conditions. The early input sensitivity added little to predicting which conditions learn once the early validation loss was modelled per task, architecture and initialisation. Within conditions, a flatter early decay went with the seeds that learned, but in a pre-specified test the association was modest (concordance 0.61), condition-dependent and not robust to resampling seeds. The results support chrono initialisation, even with a rough horizon, on such tasks.

1 Introduction

In a recurrent network trained by backpropagation through time (BPTT), the computations that process an input seen T steps before the loss receive direct credit only through the gradient carried back over those steps (Werbos, 1990). When the Jacobians along this path contract, the gradient decays exponentially with the lag and long dependencies become hard to learn (Bengio et al., 1994; Pascanu et al., 2013). Gated architectures, the long short-term memory (LSTM) (Hochreiter & Schmidhuber, 1997; Gers et al., 2000) and the gated recurrent unit (GRU) (Cho et al., 2014), mitigate this with multiplicative gates whose biases set how long a unit keeps its state, and the initial value of these biases matters: a positive forget-gate bias is standard for LSTMs (Gers et al., 2000; Jozefowicz et al., 2015), and chrono initialisation draws the gate biases so that the units' memory timescales cover the horizon of the task (Tallec & Ollivier, 2018). Tallec & Ollivier (2018) showed that chrono initialisation improves the learning of long dependencies by LSTMs over a forget-gate bias of one, with faster convergence on the adding problem and a clearer advantage on copy tasks.

The same networks are usually regularised with stochastic noise in the recurrence: dropout on the recurrent inputs of the gates, with a new mask at every step or with one mask per sequence (Gal & Ghahramani, 2016),

dropout of the cell state with one mask per sequence (Moon et al., 2015), dropout on the candidate update (Semeniuta et al., 2016), or zoneout, which keeps a random subset of units at their previous value (Krueger et al., 2017). These placements differ in how they interact with the gradient path. Zoneout was introduced partly because identity connections should help the gradient propagate, and Krueger et al. (2017) showed, in one figure for one task, that after one epoch the gradient of an LSTM with zoneout reached earlier steps than with dropout of the cell state or without regularisation. What has not been measured is how large these effects are compared with gate initialisation, whether they hold for LSTMs and GRUs alike, and whether the reach of the gradient, measured early in training, tells which runs will eventually learn.

This last question matters in practice because training recurrent networks on long dependencies is slow and its outcome varies between seeds. If early gradient reach predicted success, it could be used to stop unpromising runs or to choose between configurations before training has gone far. Comparable signals are already used for this: early validation performance for fine-tuning runs of pretrained language models (Dodge et al., 2020), extrapolated learning curves and successive halving for hyperparameter search (Domhan et al., 2015; Li et al., 2018), and zero-cost proxies computed from gradients at initialisation for architectures (Abdelfattah et al., 2021). Theory already links propagation timescales at initialisation to trainability (Chen et al., 2018; Gilboa et al., 2019), and the decay of gradient norms has been estimated during training for other purposes (Aicher et al., 2019; Chernodub & Nowicki, 2016). Whether such a measurement adds information beyond what the training curve itself shows is an empirical question.

We answer these questions with a factorial grid of 720 training runs of one-layer LSTM and GRU networks on two long-dependency tasks, crossing gate initialisation with six regulariser placements for the LSTM and four for the GRU, at one rate each, and five seeds, followed by runs with more seeds, a longer budget and mismatched chrono horizons (Table 1). Both tasks supply fresh or exhaustive training data, so the study measures how the regularisers affect optimisation; their protection against overfitting to a finite training set is outside its scope. These networks are trained with a library whose BPTT passes finite-difference checks and agrees with PyTorch to about 10^{-16} in double precision (Kholodniak, 2026); 120 further runs repeat the unregularised conditions with PyTorch's built-in cells. The contributions are:

1. We find that gate initialisation produced a larger contrast in whether and how fast the networks learned than any regulariser at the rates tested, although the regularisers still changed attainment and learning time in some conditions (Figure 1). For the LSTM on the adding problem this replicates Tallec & Ollivier (2018); it extends their comparison to the GRU, to delayed recall, to the crossing with regularisers and to mismatched horizons, under which chrono initialisation kept its advantage with a horizon about four times too short or too long.

2. We measure how each regulariser placement changes the slope of the input sensitivity with lag after 100 updates (Table 3), and find that a flattening by zoneout of similar size is already present in evaluation mode before training.

3. We find that this slope adds little to predicting which *conditions* learn once the early validation loss is modelled per architecture, initialisation and task (Table 4). Within a condition it is associated with which *seeds* learn, but a test specified before those runs supported this only when resampling conditions, not seeds (Figure 3).

4. We release the code, the per-run results and the gradient curves of every run.

Figure 1 summarises the main grid.

2 Related work

Regularising the recurrence. Zaremba et al. (2014) applied dropout only to non-recurrent connections because dropping the recurrent state at every step erased memory. Moon et al. (2015) kept one mask per sequence and applied it to the cell state, and Gal & Ghahramani (2016) derived the same per-sequence (variational) mask from approximate Bayesian inference and applied it to the input and to the recurrent state where they enter the gates. Semeniuta et al. (2016) argued that dropout should act on the candidate update

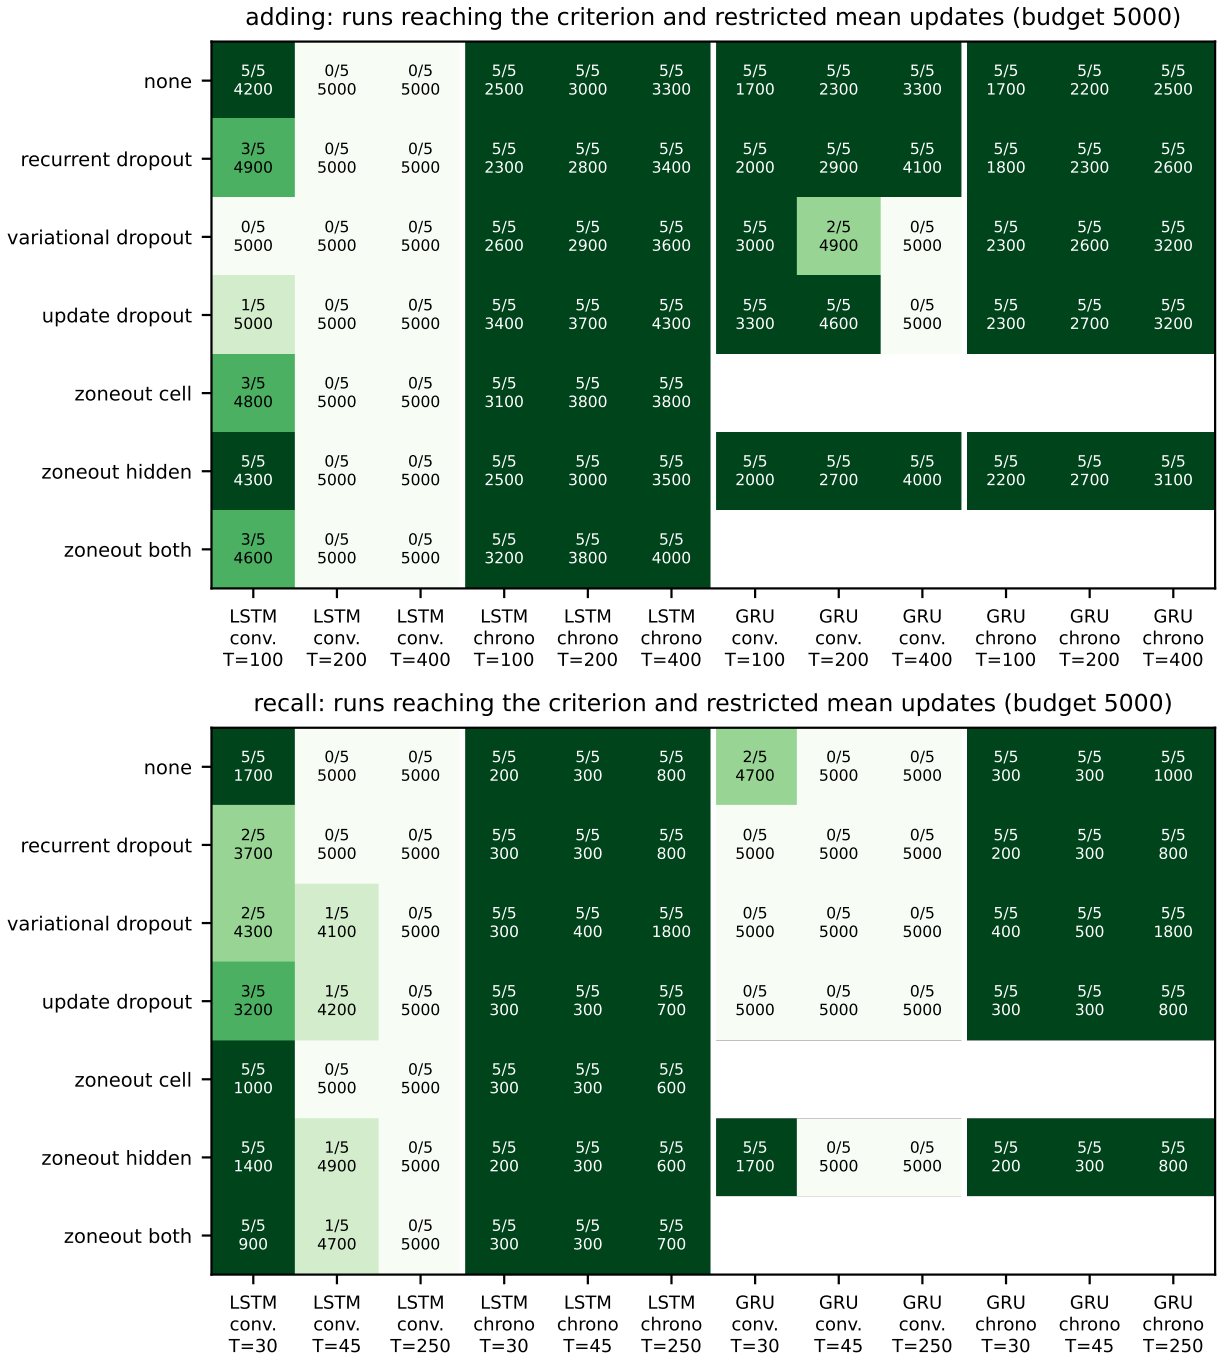

Figure 1: Chrono initialisation reached the criterion in every main-grid run; conventional initialisation often did not at the longer lengths, where the regulariser could change the outcome. Each cell gives the runs that reached the validation criterion within 5000 updates (of five seeds) and the restricted mean number of updates to the criterion (runs that did not reach it count as 5000); darker shading means more runs reached it. Blank cells: cell-state and both-state zoneout, not defined for the GRU. Columns combine architecture, initialisation (conv. = conventional) and length.

rather than on the state, so that what is already stored is never erased. Zoneout replaces dropped units by their previous values instead of zeros (Krueger et al., 2017). These works compare placements by task performance; only Krueger et al. (2017) report a gradient measurement, a normalised snapshot of cell-state gradient norms on one task.

Gating and initialisation. A forget-gate bias of one is a standard default for LSTMs (Gers et al., 2000; Jozefowicz et al., 2015), and large studies of LSTM variants found the forget gate to be among their most important components (Greff et al., 2017). Chrono initialisation links the gate biases to the timescales a task requires (Tallec & Ollivier, 2018). Gu et al. (2020b) found chrono initialisation sensitive to the choice of $T_{\max}$ and proposed a uniform gate initialisation that needs no horizon, together with a refine gate against saturation, and Chandar et al. (2019) proposed non-saturating recurrent units. Other approaches to long memory constrain or initialise the recurrent matrix itself (Le et al., 2015; Arjovsky et al., 2016; Henaff et al., 2016), and recent recurrences designed for long sequences set the timescales of the state through optimal polynomial projections (Gu et al., 2020a) or through linear recurrences with eigenvalues initialised near the unit circle and then learned (Orvieto et al., 2023). Our study concerns the classic non-linear gated cells, in which gate biases play this role.

Gradient propagation and trainability. Exponential decay of gradients in recurrent networks was analysed by Bengio et al. (1994) and Pascanu et al. (2013). Mean-field theory predicts propagation timescales from initialisation and relates them to which networks can be trained (Chen et al., 2018; Gilboa et al., 2019). Zucchet & Orvieto (2024) show that long memory makes optimisation difficult even when gradients neither vanish nor explode, and h-detach modifies the gradient flowing through the LSTM hidden path to aid the cell path (Arpit et al., 2019). The decay rate of gradient norms has been estimated during training to choose truncation lengths (Aicher et al., 2019), and per-example estimates of the long-term part of the gradient have been used to assemble mini-batches (Chernodub & Nowicki, 2016). We use such measurements not to change training but as an observable, and test whether they predict its outcome.

3 Methods

3.1 Networks

All networks have one recurrent layer of $H = 64$ units and a linear output layer applied to the last hidden state $\boldsymbol{h}_T$. The LSTM computes, for input $\boldsymbol{x}_t$,

$$\boldsymbol{c}_t = \boldsymbol{f}_t \odot \boldsymbol{c}_{t-1} + \boldsymbol{i}_t \odot \boldsymbol{g}_t, \qquad \boldsymbol{h}_t = \boldsymbol{o}_t \odot \tanh(\boldsymbol{c}_t), \tag{1}$$

where the input, forget and output gates $\boldsymbol{i}_t, \boldsymbol{f}_t, \boldsymbol{o}_t$ are logistic and the candidate $\boldsymbol{g}_t$ is a tanh of affine functions of $\boldsymbol{x}_t$ and $\boldsymbol{h}_{t-1}$, and $\odot$ is the elementwise product. We write the GRU in the convention of Chung et al. (2014), with the reset gate applied before the recurrent matrix as in Cho et al. (2014):

$$\boldsymbol{n}_t = \tanh\big(W_n \boldsymbol{x}_t + U_n(\boldsymbol{r}_t \odot \boldsymbol{h}_{t-1}) + \boldsymbol{b}_n\big), \qquad \boldsymbol{h}_t = (1 - \boldsymbol{z}_t) \odot \boldsymbol{h}_{t-1} + \boldsymbol{z}_t \odot \boldsymbol{n}_t, \tag{2}$$

with logistic reset and update gates $\boldsymbol{r}_t$ and $\boldsymbol{z}_t$. Weight matrices are drawn from $U(-0.1, 0.1)$, output weights from $U(-a, a)$ with $a = \sqrt{2/(H+O)}$ for O outputs, and biases are zero except as stated next. The networks were implemented and trained with rust-lstm 0.10.0 (Kholodniak, 2026), whose test suite checks every backward pass, including those of dropout and zoneout in training mode, against central finite differences, and in which LSTM and GRU outputs and gradients agree with PyTorch to about 10^{-16} in double precision on reference cases. The initialisation contrast was also rerun with PyTorch's own modules (Section 4.1).

Gate initialisation. *Conventional* initialisation sets the LSTM forget-gate bias to 1 and all other biases to 0; the GRU update gate then starts at $\sigma(0) = 0.5$, so, unlike the LSTM, the conventional GRU has no bias towards keeping its state. *Chrono* initialisation (Tallec & Ollivier, 2018) draws a timescale $u_j \sim U[1, T_{\max} - 1]$ for every unit j and sets the LSTM biases to $b_{f,j} = \log u_j$ and $b_{i,j} = -\log u_j$. For the GRU equation above, the corresponding choice is $b_{z,j} = -\log u_j$, so that $z_j \approx 1/(1 + u_j)$ and unit j retains about a fraction $1 - 1/(1 + u_j)$ of its state per step. In the main grid $T_{\max}$ equals the sequence length T; a follow-up uses $T_{\max} = \max(3, \lfloor T/4 \rfloor)$ and $4T$.

Regularisers. Six placements for the LSTM and four for the GRU are compared with an unregularised baseline:

- *recurrent dropout*: dropout on the recurrent input $\boldsymbol{h}_{t-1}$ of the gates with a new mask at every step;
- *variational dropout*: the same with one mask per sequence (Gal & Ghahramani, 2016); Keras's `recurrent_dropout` is of this kind;
- *update dropout*: dropout on the candidate $\boldsymbol{g}_t$ of the LSTM or $\boldsymbol{n}_t$ of the GRU (Semeniuta et al., 2016);
- *zoneout* (Krueger et al., 2017) of the LSTM cell state, of the hidden state, or of both (GRU: hidden state only), which replaces a new state $\tilde{\boldsymbol{s}}_t$ by $\boldsymbol{d}_t \odot \boldsymbol{s}_{t-1} + (1 - \boldsymbol{d}_t) \odot \tilde{\boldsymbol{s}}_t$ with Bernoulli masks $\boldsymbol{d}_t$.

Dropout masks are inverted (scaled by $1/(1-p)$) and never touch the carried state. Rates are $p = 0.25$ for dropout and $q = 0.15$ for zoneout; at evaluation, dropout is off and zoneout uses its expectation $\boldsymbol{d}_t = q$, which adds a fixed leak to the carried state.

3.2 Tasks

In the *adding problem* (Hochreiter & Schmidhuber, 1997), in the form of Le et al. (2015), each input is a pair (v_t, m_t) with $v_t \sim U[0,1]$ and a marker m_t that equals one at one position in each half of the sequence; the target is the sum of the two marked values, predicted at the last step and scored by squared error. Predicting the mean, 1, gives a mean squared error of $1/6$. In *delayed recall* one of eight symbols is presented as a one-hot vector at the first step, the inputs are zero for the next $T-2$ steps, and a query flag at the last step asks for the symbol, scored by softmax cross-entropy. For a given length recall has only eight distinct sequences. Because adding draws fresh sequences for every batch and recall validates on all of its eight sequences, neither task assesses the usual effect of regularisation on generalisation from a finite training set.

3.3 Training and endpoint

Networks are trained with Adam (Kingma & Ba, 2015) at learning rate 10^{-3}, batches of 32 fresh sequences, and the recurrent and output gradients clipped jointly to norm 1. Every 100 updates the network is evaluated on fixed validation data: 1024 adding sequences, or all eight recall sequences. A run *reaches the criterion* at the first evaluation at which the validation mean squared error falls below 0.01 (a 94 % reduction from the constant prediction) or all eight symbols are recalled, and training then stops; otherwise it stops at the budget of 5000 updates (20 000 in one follow-up). Reaching the criterion is the endpoint throughout. In the main grid a final score on 4096 fresh adding sequences confirms it for 243 of the 262 adding runs that reached it (169 of the 180 chrono-initialised ones); the remaining 19 have test errors just above 0.01.

3.4 Gradient probes

At updates 0, 100, 250, 500, 1000 and 2000, unless the run has already stopped, the gradients of a fixed probe batch of 64 sequences are recorded at every lag $k = 0, \ldots, T-1$, where lag k denotes the input k steps before the last one. The main quantity is the input sensitivity of projections of the final hidden state,

$$G(k) = \frac{1}{4} \sum_{j=1}^{4} \frac{1}{64} \sum_{b=1}^{64} \log \left\| \frac{\partial\, \boldsymbol{u}_j^{\top} \boldsymbol{h}_{T,b}}{\partial \boldsymbol{x}_{T-k,b}} \right\|, \tag{3}$$

where b indexes the sequences of the probe batch and $\boldsymbol{u}_1, \ldots, \boldsymbol{u}_4$ are fixed random unit vectors. $G(k)$ is independent of the loss and of the output layer. Its trend with lag is summarised by the slope s of a least-squares line of $G(k)$ on k over $1 \le k \le T/2$, in natural-log units per step. This window was fixed in the first analysis code; it excludes the lag of the recall symbol and all but one possible lag of the first adding marker, so s describes the trend over the more recent half of the sequence, and slopes over $1 \le k \le T/4$ and $1 \le k \le T-1$ are computed as sensitivity analyses. We use s as a descriptive summary of input sensitivity: a value near zero means little relative attenuation over the window and says nothing about the earliest inputs,

Table 1: Experimental plans. Conditions combine task, length, architecture, initialisation and regulariser; the GRU has five regulariser settings (no cell-state zoneout), the LSTM seven.

Plan	Content	Seeds	Runs	Budget
Calibration	unregularised; adding $T \in \{50, \ldots, 400\}$, recall $T \in \{10, \ldots, 500\}$	0–1	104	5000
Main grid	2 architectures × 2 initialisations × 5 or 7 regularisers × 6 lengths	0–4	720	5000
More seeds	the 13 main-grid conditions with mixed outcomes	5–19	195	5000
Long budget	32 conventional conditions below 5/5 at shorter lengths	0–4	160	20 000
Horizon mismatch	unregularised, chrono with $T_{\max} \approx T/4$ or $4T$	0–4	120	5000
PyTorch replication	the 24 unregularised main-grid conditions, PyTorch `nn.LSTM`/`nn.GRU`; outcomes only	0–4	120	5000

which lie outside it. Only for a negative slope and a curve close to exponential is $-1/s$ the lag over which the sensitivity shrinks by a factor of e. We also record the gradient of each sequence's own loss and, as task-specific checks, the derivative of the adding prediction with respect to the two marked values and of the true-class logit with respect to the recall symbol. Probes are taken in evaluation mode and in training mode, the latter averaged over eight draws of the dropout and zoneout masks from a random stream separate from training. In evaluation mode dropout is off, so its probes show only the dynamics it induced during training, whereas the expected zoneout mask enters the probe directly.

3.5 Experimental design

Table 1 summarises the plans. Two calibration plans with unregularised networks and seeds 0 and 1 selected the lengths of the main grid so that each task has a length that most conditions solve, one in transition, and one that most fail without chrono initialisation. The main grid crosses architecture, initialisation, regulariser, task and length with five seeds; of its 144 conditions, 53, all conventionally initialised, had fewer than five of five runs reaching the criterion. Three follow-ups were specified, with their analyses, before they were run (repository commit `a54e1cb`): more seeds for the 13 conditions in which some but not all seeds reached the criterion, a fourfold budget for the 32 of the 53 at the two shorter lengths of each task, and chrono initialisation with mismatched horizons. The 48 unregularised calibration runs with seeds 0 and 1 at the main-grid lengths coincide with main-grid runs and gave identical results. The long-budget runs repeat main-grid runs for their first 5000 updates and continue them to 20 000; the 25 that reached the criterion within 5000 updates are identical to their main-grid counterparts. All random streams are seeded for each fixed run configuration and implementation. Runs used rust-lstm 0.10.0 and rustc 1.86 on one Apple M1 Pro.

Main-grid lengths are $T \in \{100, 200, 400\}$ for adding and $T \in \{30, 45, 250\}$ for recall.

3.6 Statistical analysis

Landmark prediction across conditions. Following the landmark approach to dynamic prediction (van Houwelingen, 2007), prediction is made from a fixed point in training, the landmark at update 100: every main-grid run not yet at the criterion there (all 720) is followed to the end of its budget. A gradient-free baseline predicts reaching the criterion from task, length, architecture, initialisation, regulariser and the log validation loss at update 100; the full model adds the slope s at update 100 in evaluation mode. The *additive* baseline enters all factors as main effects. It was committed with the first analysis code while the main grid was running.[1] A *stratified* baseline, added after a review of our first analysis, gives each task × architecture × initialisation stratum its own intercept and its own effects of log length and log landmark loss. Several conditions separate the outcomes completely, so unpenalised maximum-likelihood estimates would not exist (Heinze & Schemper, 2002); the models are therefore L_2-penalised logistic regressions on standardised features with inverse penalty strength $C = 1$, which keeps coefficients finite. Predictions are made out of sample, holding out all runs of a condition together in ten grouped folds (primary) or holding

[1] Commit `f9cf778`, three minutes after the main grid started.

Table 2: Long budget: conventionally initialised conditions below 5/5 at 5000 updates, rerun with 20 000. Runs (of five) that reached the criterion by 5000 → by 20 000 updates, and median updates among those that reached it. Selected conditions; all 32 are tabulated in the repository. rec. = recurrent, var. = variational, upd. = update dropout; zo. = zoneout of the cell, hidden or both states.

		none	rec.	var.	upd.	zo. cell	zo. hid.	zo. both
LSTM adding $T = 200$	by 5000 → 20 000	0 → 5	0 → 5	0 → 5	0 → 5	0 → 5	0 → 5	0 → 5
	median	8100	9000	11 200	10 100	8500	8700	9100
GRU recall $T = 30$	by 5000 → 20 000	2 → 5	0 → 5	0 → 5	0 → 2	–	5 → 5*	–
GRU recall $T = 45$	by 5000 → 20 000	0 → 5	0 → 0	0 → 5	0 → 0	–	0 → 2	–
LSTM recall $T = 45$	by 5000 → 20 000	0 → 3	0 → 1	1 → 1	1 → 1	0 → 2	1 → 3	1 → 3

*Already 5 of 5 within 5000 updates in the main grid; not rerun.

out whole seeds, regulariser families or task lengths (Roberts et al., 2017). We report the reduction in log loss and its 95 % interval from resampling conditions, and for the primary comparison also from resampling conditions and seeds crossed; these intervals hold the out-of-sample predictions fixed.

Ranking within conditions. Within each condition with both outcomes, every pair of a run that reached the criterion and one that did not is scored 1 if the run that reached it has the larger predictor, $\frac{1}{2}$ for ties and 0 otherwise; the mean over pairs is a concordance (an area under the curve within conditions). For the loss the predictor is $-\log$ of the validation loss at the landmark, so that a lower loss counts as predicting success. The exploratory analysis used the main grid. The confirmatory test, written into the project README before the follow-up runs (commit `a54e1cb`), used only seeds 5–19, predicted a concordance above 0.5 for the slope and not for the landmark loss, and specified a 95 % interval from resampling conditions. Because a seed fixes the initial weights and the training data for every condition of a given architecture, task and length, runs with the same seed are not independent across conditions; we therefore also report intervals that resample seeds (a pair counts once per draw of each of its two seeds) and that resample conditions and seeds crossed.

Time to the criterion. A discrete-time hazard model (Singer & Willett, 1993) uses one record per run and evaluation interval from update 200 until the criterion is reached or the budget ends, and models the probability of reaching the criterion in an interval by a penalised logistic regression on log time, the additive baseline covariates and the slope, standardised across runs.

4 Results

4.1 In this grid, initialisation decided whether and how fast the tasks were learned

Figure 1 shows the main grid. With chrono initialisation all 360 runs in 72 conditions reached the criterion within 5000 updates; with conventional initialisation 120 of 360 did, and 53 of its 72 conditions had fewer than five of five. Without regularisation, the conventionally initialised LSTM reached it in none of ten adding runs at $T = 200$ and 400 and in none of ten recall runs at $T = 45$ and 250, and the conventionally initialised GRU in 2 of 5 recall runs at $T = 30$ and none at the two longer recall lengths.

A fourfold budget shows that many of these failures are slow learning rather than inability (Table 2). Of the 160 runs in the 32 conventional conditions below 5/5 at the two shorter lengths of each task, 25 reached the criterion within 5000 updates and 121 within 20 000. All 35 LSTM adding runs at $T = 200$ reached it, with medians between 8100 updates without regularisation and 11 200 with variational dropout. Recall at $T = 45$ remained hard: without regularisation 3 of 5 LSTM and 5 of 5 GRU runs reached it.

Chrono initialisation does not owe its advantage to knowing the task length exactly. With $T_{\max} = \max(3, \lfloor T/4 \rfloor)$ or $4T$, 119 of 120 unregularised runs in 24 conditions reached the criterion within 5000 updates; the exception was one LSTM recall run at $T = 250$ with the short horizon. On recall the mismatched horizons were somewhat slower than the matched one at short lengths (for example 420 against 240 updates for the LSTM at $T = 30$ with $T/4$), but at $T = 30$ and 45 their restricted mean numbers of updates were 4 to 17 times lower than with conventional initialisation (at $T = 30$: 1700 updates for the LSTM and 4680 for

the GRU, counting runs that did not reach the criterion as 5000, against 260–540 with mismatched horizons). On adding the short horizon was the fastest GRU setting at every length. With $T_{\max} = T/4$ and gates held at their initial biases, the slowest possible unit would retain $(1 - 1/T_{\max})^T$, about 1–2 % of its state over a whole sequence at these lengths. The experiment shows that chrono initialisation tolerates this mismatch on tasks with one long lag.

Replication in PyTorch. After the main analysis, the 24 unregularised conditions of the main grid (two architectures, two initialisations, six task lengths, seeds 0–4) were rerun with the LSTM and GRU modules of PyTorch 2.9.1 (Paszke et al., 2019), with PyTorch's default weight initialisation, single precision and independently drawn training data, but the same gate biases, tasks, optimiser, batch size, clipping, budget, validation schedule and criteria. PyTorch's GRU differs from the cell of Section 3.1 in two ways: it keeps the old state with weight z, so its chrono biases are $b_{z,j} = \log u_j$, and it applies the reset gate after the recurrent matrix, $\boldsymbol{r}_t \odot (U\boldsymbol{h}_{t-1} + \boldsymbol{b})$, rather than before it. All 60 chrono-initialised runs reached the criterion, as in the main grid, and 27 of 60 conventionally initialised runs did in each implementation. The number of seeds that reached the criterion was the same in 21 of the 24 conditions and differed by at most two (LSTM recall at $T = 30$: 5 against 3 of 5); across the 18 conditions in which runs reached the criterion in both implementations, the median numbers of updates had a Spearman correlation of 0.97. The advantage of chrono initialisation in unregularised learning outcomes thus also appeared with PyTorch's built-in cells, which is evidence against an artefact of the implementation used for the main grid. Only these outcomes were replicated: the gradient measurements were not repeated, and the regulariser comparisons rest on the main implementation alone, because PyTorch's recurrent modules do not provide the recurrent dropout and zoneout placements.

4.2 Regularisers changed learning speed, and conventional initialisation exposed the difference

Under conventional initialisation dropout in every placement generally delayed learning, and in some conditions it prevented it within the budget (Figure 1, Table 2). The conventional GRU reached the criterion on adding at $T = 400$ in 5 of 5 runs without regularisation but in none with variational or update dropout. On recall at $T = 30$ it reached it in 2 of 5 runs without regularisation and in none with any dropout placement within 5000 updates; at $T = 45$ it reached it in 5 of 5 runs within 20 000 updates without regularisation but in none with recurrent or update dropout. Zoneout of the hidden state let the same GRU reach recall at $T = 30$ in 5 of 5 runs, and cell or both-state zoneout shortened the conventional LSTM's restricted mean time on recall at $T = 30$ from 1700 to 960 and 940 updates. Zoneout was not uniformly helpful, however: it slowed the conventional GRU on adding at $T = 400$ (4000 against 3260 updates), lowered the LSTM on adding at $T = 100$ from 5 to 3 of 5 runs within 5000 updates in the cell and both-state variants, and with 20 000 updates the GRU reached recall at $T = 45$ in only 2 of 5 runs with hidden zoneout against 5 of 5 without regularisation. Dropout was not uniformly harmful either: on recall at $T = 45$ the conventional LSTM reached the criterion within 5000 updates in one of five runs with variational or update dropout and in none without regularisation.

Under chrono initialisation every regulariser reached the criterion in every run, but learning times still differed: variational dropout roughly doubled the time on recall at $T = 250$ (1800 against 820 updates for the LSTM, 1760 against 960 for the GRU), whereas update dropout shortened it slightly (740 and 760), as did recurrent dropout for the GRU (820). These are contrasts at one rate per regulariser family and do not rank the families in general.

4.3 Initialisation set the decay of input sensitivity more than the regularisers did

Without regularisation and before training, the slope of $G(k)$ per step, averaged over the six task lengths, was -0.514 for the conventional GRU, -0.185 for the conventional LSTM, -0.044 for the chrono GRU and -0.042 for the chrono LSTM. The GRU value is of the order of $\log 0.5 = -0.69$ expected from an update gate at 0.5 in the direct carry alone. Conventional curves were close to exponential (median R^2 of the line at least 0.95 in every task, length and architecture at update 100), whereas chrono curves were not (median R^2 between 0.61 and 0.94) and in short recall were sometimes increasing; for chrono networks the slope is a descriptive summary rather than a decay rate. Figure 2 shows the resulting contrast for recall at $T = 45$:

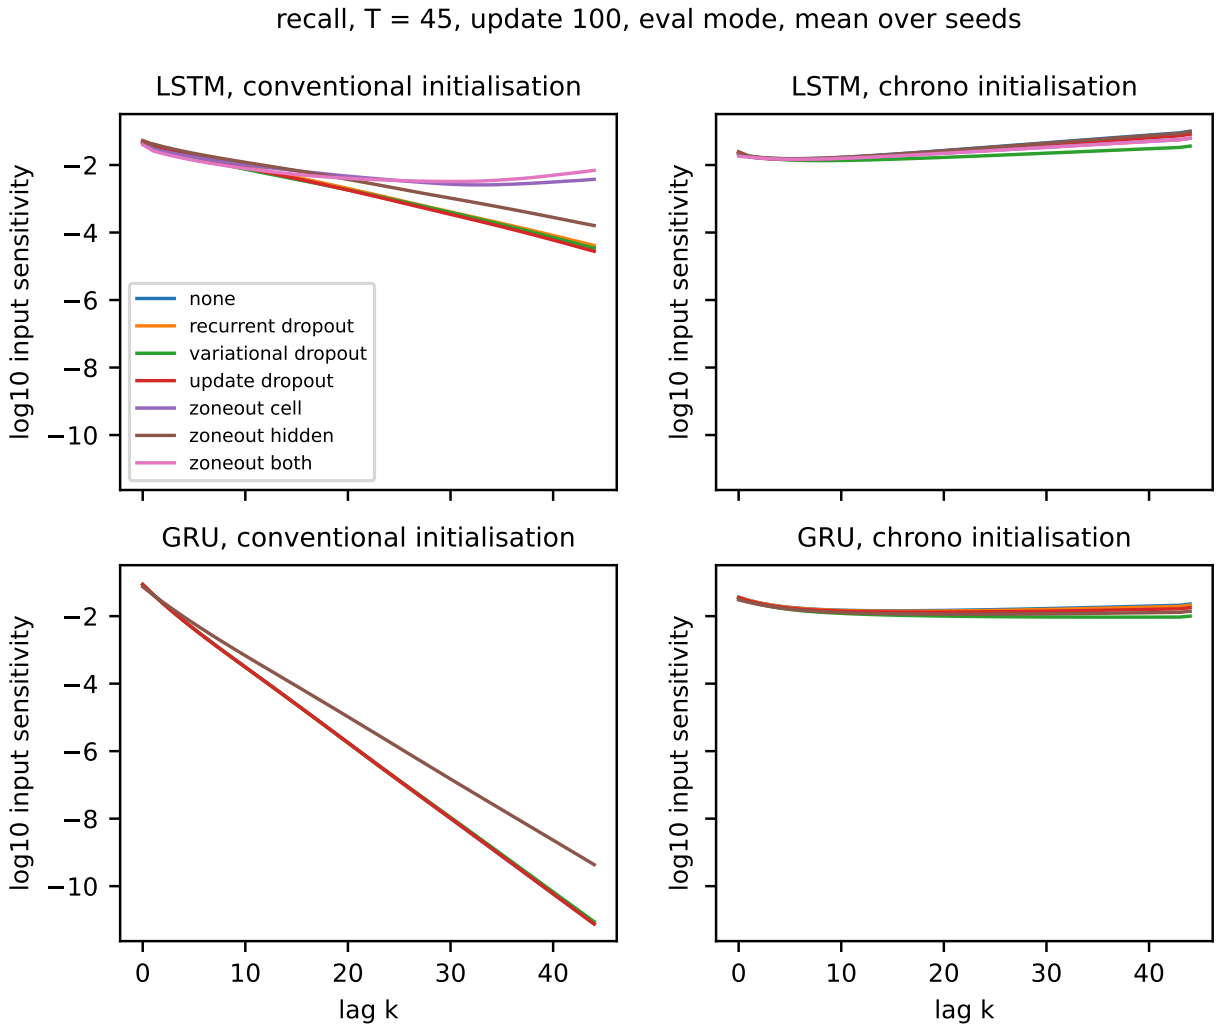

Figure 2: Under conventional initialisation the input sensitivity decays roughly exponentially with lag; chrono networks attenuate far less. Input sensitivity of the final hidden state, $G(k)/\ln 10$, against lag k for recall at $T = 45$ after 100 updates, evaluation mode, mean over five seeds.

after 100 updates, the sensitivity of the unregularised conventional GRU to the first input is about ten orders of magnitude below its sensitivity to the last one (about eight with hidden zoneout), that of the unregularised conventional LSTM about three (about one with cell or both-state zoneout), and chrono networks are nearly flat. The task-specific derivatives, averaged over the three lengths and five seeds of each task in evaluation mode after 100 updates, agree: the derivative of the unregularised adding prediction with respect to the first marked value was 1.6×10^{-8} for the conventional GRU and 5.7×10^{-6} for the conventional LSTM, against 2.8×10^{-3} and 5.3×10^{-3} with chrono initialisation, and the derivative of the recall logit with respect to the symbol was 2×10^{-9} for the conventional GRU against 0.07 for the conventional LSTM and 0.23 and 0.41 for the chrono networks. These derivatives show that early inputs reach the output, not that the networks already use them selectively: for the chrono LSTM the mean absolute derivative with respect to unmarked values (5.3×10^{-3}) was as large as that with respect to the marked ones, against target derivatives of 1 and 0.

Table 3 gives the change of the slope caused by each regulariser after 100 updates. Under conventional initialisation variational dropout steepened the decay most for both architectures, update and recurrent dropout less, and zoneout flattened it where it acts on the path that carries memory: the cell state of the

Table 3: Change of the slope s at update 100 (evaluation mode) against the unregularised network with the same seed, task and length, averaged over lengths within each seed; mean and pointwise 95 % t interval over five seeds (exploratory). Positive values mean slower decay; rec., var., upd. = recurrent, variational, update. Zoneout values include the direct effect of its expected mask in evaluation mode, present before training (see text).

	Conventional		Chrono	
Regulariser	LSTM	GRU	LSTM	GRU
rec. dropout	−0.006 [−0.023, 0.011]	−0.017 [−0.022, −0.012]	−0.001 [−0.002, 0.001]	−0.001 [−0.002, 0.000]
var. dropout	−0.038 [−0.057, −0.019]	−0.060 [−0.068, −0.052]	−0.006 [−0.012, −0.001]	−0.007 [−0.013, −0.002]
upd. dropout	−0.008 [−0.032, 0.016]	−0.034 [−0.043, −0.024]	−0.001 [−0.004, 0.001]	−0.002 [−0.007, 0.002]
zoneout, cell	0.033 [0.015, 0.050]	–	0.000 [−0.006, 0.006]	–
zoneout, hidden	0.003 [−0.009, 0.016]	0.066 [0.062, 0.070]	0.000 [−0.001, 0.001]	−0.003 [−0.011, 0.004]
zoneout, both	0.045 [0.032, 0.057]	–	−0.001 [−0.006, 0.005]	–

Table 4: Landmark prediction of reaching the criterion from update 100 (720 runs, 480 reached it; seeds 2–4 only: 432 runs, 288 reached it). Log loss of the gradient-free baseline and of the baseline plus the slope s, out of sample, and the reduction with a 95 % interval from resampling conditions; for held-out conditions also from resampling conditions and seeds crossed. Intervals hold the out-of-sample predictions fixed; reductions are computed before rounding.

Baseline	Held out	Baseline	$+\,s$	Reduction	Interval
additive	conditions (10 folds)	0.239	0.165	0.074	[0.044, 0.104]
	crossed interval				[0.022, 0.120]
stratified	conditions (10 folds)	0.162	0.154	0.008	[−0.002, 0.018]
	crossed interval				[−0.009, 0.025]
stratified	seeds	0.141	0.130	0.011	[0.002, 0.022]
stratified	regulariser families	0.301	0.270	0.030	[0.006, 0.059]
stratified	task lengths	0.316	0.309	0.006	[−0.008, 0.021]
stratified	conditions, seeds 2–4 only	0.179	0.174	0.006	[−0.004, 0.016]

LSTM and the hidden state of the GRU. A flattening of similar size is present before any training: at update 0, when the networks differ only in the evaluation-mode zoneout mask, the slope changed by +0.092 for the conventional GRU with hidden zoneout and by +0.028 and +0.030 for the conventional LSTM with cell and both-state zoneout. After training, the contrast combines this effect of the expected mask with learned dynamics, and the two cannot be separated here. Under chrono initialisation the mean changes were at most about 0.007 in absolute value, the largest for variational dropout. Measured in training mode before any training, the picture differs: variational dropout then flattened the curve (+0.027 for the GRU and +0.020 for the LSTM under conventional initialisation), although it delayed learning in several of the conditions above. Input sensitivity at initialisation therefore does not rank the dropout placements by how well they train.

4.4 The early slope added little to predicting which conditions learned

Against the additive baseline of the first analysis, the slope lowered the out-of-sample log loss from 0.239 to 0.165 (Table 4), and the reduction was positive when pooled over held-out seeds, regulariser families and task lengths, with the training-mode slope, with other fitting windows, and on seeds 2–4 alone, which were not used for calibration. Most of this gain, however, came from structure the additive baseline could not express. The stratified baseline, which lets the effects of length and of the early loss differ between task, architecture and initialisation, reached a log loss of 0.162 without the slope, and adding the slope lowered it only to 0.154, a reduction of 0.008 whose interval includes zero. Small reductions remained when whole seeds (0.011) or whole regulariser families (0.030) were held out, but not for held-out lengths or for seeds 2–4. In these models, the early validation loss together with the condition therefore predicted nearly as well without the slope as with it; the models add one linear slope term and do not test non-linear or stratum-specific slope effects, and the slope may help when predicting for a regulariser the model has not seen. Pooled reductions also hide groups in which the slope made predictions worse: with the stratified baseline, one of the five held-out

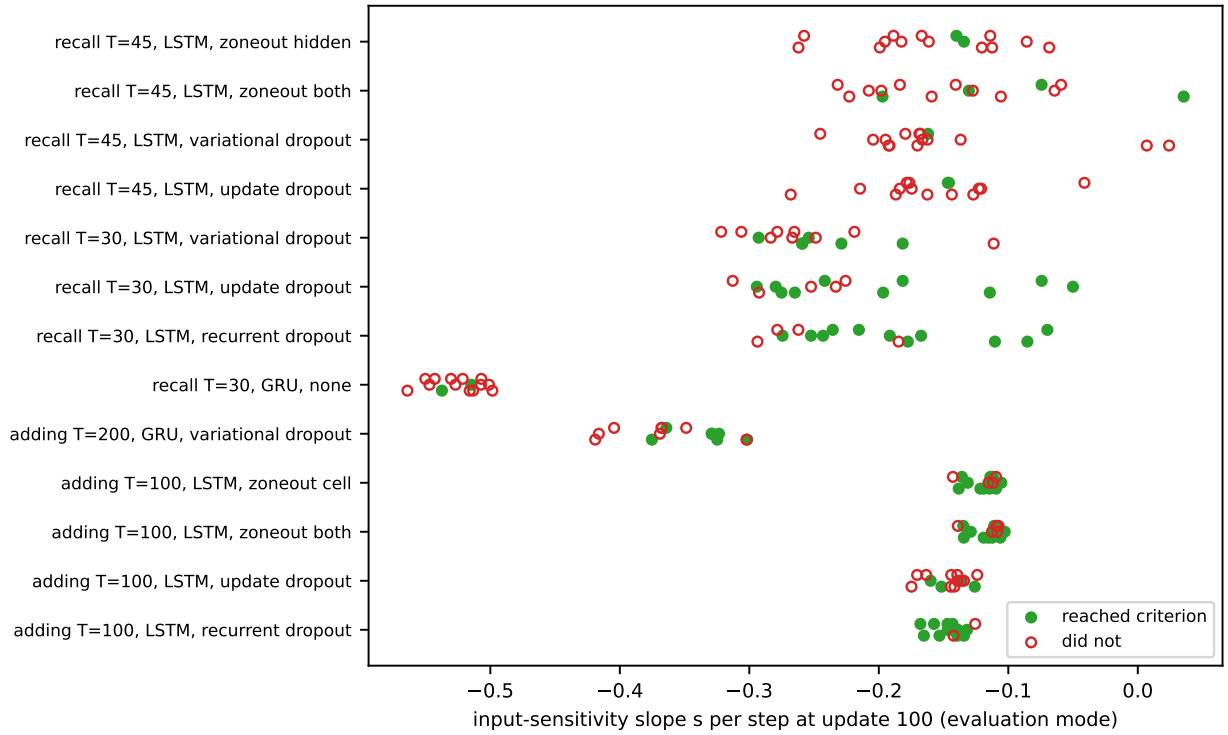

Figure 3: Seeds that learned tended to have flatter early slopes in the LSTM recall conditions, but not in LSTM adding at $T = 100$. Confirmatory runs: slope s at update 100 for every run with seeds 5–19 in the 13 conditions with mixed outcomes in the main grid. Filled markers reached the criterion within 5000 updates.

seeds and two of the six held-out lengths. With reaching the criterion on the test set as the outcome, the additive-baseline reduction shrank from 0.074 to 0.029 (interval −0.003 to 0.057). The slope does not order success consistently across architectures either: the unregularised conventional GRU reached the criterion on adding at $T = 200$ and 400 with a mean slope of about −0.26 after 100 updates, while the unregularised conventional LSTM did not reach it there within 5000 updates with −0.13.

4.5 Within some conditions, flatter early decay went with the seeds that learned

Within a condition, runs differ only in their seed, and the question becomes whether the slope tells which seeds will learn. In the main grid 13 conditions, all conventionally initialised, had both outcomes. Among their 65 runs the slope ranked runs that reached the criterion above runs that did not with a concordance of 0.735 (95 % interval 0.618–0.853 from resampling conditions, 0.556–0.925 from resampling seeds), whereas a lower validation loss at update 100 did not (0.426, 0.297–0.567). Because this was found after looking at the data, we tested it on 15 new seeds per condition, as specified beforehand. On these 195 runs (504 pairs of a run that reached the criterion and one that did not) the concordance was 0.613 for the slope and 0.462 for the early loss (Figure 3). With the prespecified interval, which resamples conditions, the slope passed the test (0.515–0.698) and the loss did not (0.340–0.586). Seeds, however, are shared across conditions, and resampling them widens the slope interval to 0.482–0.744, and resampling conditions and seeds crossed to 0.409–0.779; both include 0.5. The paired difference between the two predictors on the same pairs was 0.151 (0.014–0.285 resampling conditions, −0.015 to 0.305 resampling seeds). The association was not uniform: the concordance exceeded 0.6 in all seven conventional LSTM recall conditions and in the GRU adding condition, but was 0.19–0.55 in the four LSTM adding conditions at $T = 100$ and 0.46 for the GRU on recall. The evidence therefore supports a modest, condition-dependent association rather than a general ranking of seeds; the absence of an association for the early loss is a failure to detect one, not evidence that none exists. In the hazard model, a slope flatter by one standard deviation across runs was associated with higher odds

of reaching the criterion in each interval (0.93 on the log-odds scale, interval 0.59–1.54 from refitting on resampled conditions; 720 runs, 20 954 intervals); because this model uses the additive covariates, which Section 4.4 found insufficient across conditions, it is descriptive only.

5 Discussion

The clearest result is that gate initialisation produced larger contrasts than any regulariser we tested. The comparison is between particular levels: a conventional default suited to short dependencies, a horizon-matched chrono initialisation and one rate of each regulariser, at lengths chosen so that most conventionally initialised conditions did not reach the criterion at the longest length within 5000 updates. Other rates, tasks or lengths could change the ranking. With matched chrono initialisation the input sensitivity attenuated far less with lag than under conventional initialisation, every run reached the criterion, and the regulariser changed only how long it took; with conventional initialisation the sensitivity decayed by orders of magnitude over the sequence and the choice of regulariser could decide whether a run learned within the budget. For the LSTM on the adding problem this replicates Tallec & Ollivier (2018), and it agrees with mean-field accounts in which gate biases set the propagation timescales of LSTMs and GRUs (Gilboa et al., 2019). That a horizon off by about a factor of four still worked in 119 of 120 unregularised runs shows that chrono initialisation tolerates this mismatch on tasks with one long lag; Gu et al. (2020b) found it sensitive to $T_{\max}$ in other settings, and whether a spread of timescales or a suitable common bias is what matters cannot be separated without fixed-bias baselines.

Consistent with the gradient-propagation motivation of zoneout (Krueger et al., 2017), applying it to the state that carries memory flattened the input-sensitivity curve. The flattening was similar in size to the direct effect of the expected mask in evaluation mode before training. Zoneout helped short recall, but its learning benefits were inconsistent, and it often slowed learning on adding. Dropout generally slowed learning in all three placements and, once training had begun, steepened the decay, clearly so for the GRU and, in the LSTM, for variational dropout; variational dropout, which keeps one mask per sequence and was derived from approximate variational inference (Gal & Ghahramani, 2016), steepened it most. At initialisation, in training mode, the same placement flattened the decay, so measurements before training would have ranked it above the other dropout placements. One plausible reading is that a fixed mask removes the same recurrent inputs at every step, which the network compensates for by learning faster-decaying dynamics; we did not test this mechanism.

Across conditions, the validation loss after 100 updates, modelled per architecture, initialisation and task, predicted learning nearly as well without the gradient slope as with it, so these models give little support for using gradient probes to predict which conditions will learn. Within conditions, we detected an association between the slope and the outcome in the pooled analysis and in the prespecified test, but did not detect one for the early validation loss. The slope association was modest, occurred only in some conditions, and was not distinguishable from chance after accounting for the dependence between runs that share a seed. Whether a gradient probe after a few hundred updates can help select the seeds to continue remains to be tested directly. The within-condition findings concern conventional initialisation; all chrono-initialised main-grid runs reached the criterion, leaving no outcome variation to predict.

Practical implications. The results support chrono initialisation, even with a rough horizon, for one-layer gated networks that must carry information over many steps. They do not say whether dropout or zoneout improve generalisation, because the design did not assess that benefit with a fixed, finite training set, and they do not cover non-recurrent dropout on the inputs and outputs. We did not evaluate an early-stopping policy based on either predictor.

Limitations. The study uses one-layer networks with 64 units, one optimiser and learning rate, one rate per regulariser family and two synthetic tasks with fresh or exhaustive training data, so the study measures the regularisers' effects on optimisation. The task lengths were chosen on unregularised networks so that most conventionally initialised conditions do not reach the criterion at the longest length within 5000 updates, which shapes the size of the initialisation contrast. The conventional GRU, unlike the LSTM, has no bias towards keeping its state, and the zoneout rate of 0.15 is lower for the cell state and higher for the hidden

state than the values Krueger et al. (2017) found best on character-level Penn Treebank (they used 0.15 for both states on permuted MNIST); recall has only eight distinct sequences per length and zero inputs between symbol and query, so derivatives with respect to those inputs measure sensitivity to perturbations the task never presents. The endpoint is the first validation success; for 19 of 262 adding runs it was not confirmed on the test set. Seeds 0 and 1 were also used to choose the lengths, and the stratified baseline was added after a review of the first analysis; the seeds 2–4 analysis and the confirmatory runs address these points only partly. The measured quantity is the input sensitivity of projections of the final hidden state, not a direct measurement of the LSTM cell-state Jacobian, and a single exponential describes chrono curves poorly. The prediction intervals resample conditions and, for the primary comparison, conditions and seeds, but hold the fitted out-of-sample predictions fixed. The many sensitivity analyses are exploratory; only the within-condition comparison on seeds 5–19 was confirmatory, and its conclusion depends on how the shared seeds are treated.

6 Conclusion

We compared gate initialisation with recurrent dropout and zoneout in one-layer LSTM and GRU networks on two synthetic long-dependency tasks with fresh or exhaustive training data, at one rate per regulariser. Initialisation produced the largest differences, both in how far the input sensitivity reached and in whether and how fast the networks learned. Dropout and zoneout mattered mainly when initialisation left the sensitivity decaying fast: dropout generally slowed learning, and zoneout helped in some settings and not in others. Chrono initialisation kept its advantage with a horizon about four times too short or too long. The decay measured after 100 updates added little to predicting which conditions learn once the early validation loss was modelled per task, architecture and initialisation. Within conditions a flatter early decay was associated with the seeds that went on to learn, but in the test specified beforehand this association was modest, condition-dependent, and robust only to resampling conditions, not seeds.

Code and data availability

The experiment code, the analysis and the per-run results are available at `https://github.com/SyntaxSpirits/rnn-gradient-reach`; the gradient curves of every run are archived at `https://doi.org/10.5281/zenodo.23285386`. Networks were trained with rust-lstm 0.10.0 (Kholodniak, 2026).

References

Mohamed S. Abdelfattah, Abhinav Mehrotra, Łukasz Dudziak, and Nicholas D. Lane. Zero-cost proxies for lightweight NAS. In *International Conference on Learning Representations*, 2021.

Christopher Aicher, Nicholas J. Foti, and Emily B. Fox. Adaptively truncating backpropagation through time to control gradient bias. In *Proceedings of the 35th Conference on Uncertainty in Artificial Intelligence*, 2019. doi: 10.48550/arXiv.1905.07473.

Martin Arjovsky, Amar Shah, and Yoshua Bengio. Unitary evolution recurrent neural networks. In *Proceedings of the 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, pp. 1120–1128, 2016. doi: 10.48550/arXiv.1511.06464.

Devansh Arpit, Bhargav Kanuparthi, Giancarlo Kerg, Nan Rosemary Ke, Ioannis Mitliagkas, and Yoshua Bengio. h-detach: Modifying the LSTM gradient towards better optimization. In *International Conference on Learning Representations*, 2019. doi: 10.48550/arXiv.1810.03023.

Yoshua Bengio, Patrice Simard, and Paolo Frasconi. Learning long-term dependencies with gradient descent is difficult. *IEEE Transactions on Neural Networks*, 5(2):157–166, 1994. doi: 10.1109/72.279181.

Sarath Chandar, Chinnadhurai Sankar, Eugene Vorontsov, Samira Ebrahimi Kahou, and Yoshua Bengio. Towards non-saturating recurrent units for modelling long-term dependencies. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pp. 3280–3287, 2019. doi: 10.1609/aaai.v33i01.33013280.

Minmin Chen, Jeffrey Pennington, and Samuel S. Schoenholz. Dynamical isometry and a mean field theory of RNNs: Gating enables signal propagation in recurrent neural networks. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pp. 873–882, 2018. doi: 10.48550/arXiv.1806.05394.

Artem Chernodub and Dimitri Nowicki. Sampling-based gradient regularization for capturing long-term dependencies in recurrent neural networks. In *Neural Information Processing (ICONIP 2016)*, volume 9948 of *Lecture Notes in Computer Science*, pp. 90–97. Springer, 2016. doi: 10.1007/978-3-319-46672-9_11.

Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using RNN encoder–decoder for statistical machine translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing*, pp. 1724–1734, 2014. doi: 10.3115/v1/D14-1179.

Junyoung Chung, Caglar Gulcehre, KyungHyun Cho, and Yoshua Bengio. Empirical evaluation of gated recurrent neural networks on sequence modeling, 2014.

Jesse Dodge, Gabriel Ilharco, Roy Schwartz, Ali Farhadi, Hannaneh Hajishirzi, and Noah Smith. Fine-tuning pretrained language models: Weight initializations, data orders, and early stopping. *arXiv preprint arXiv:2002.06305*, 2020. doi: 10.48550/arXiv.2002.06305.

Tobias Domhan, Jost Tobias Springenberg, and Frank Hutter. Speeding up automatic hyperparameter optimization of deep neural networks by extrapolation of learning curves. In *Proceedings of the 24th International Joint Conference on Artificial Intelligence*, pp. 3460–3468, 2015.

Yarin Gal and Zoubin Ghahramani. A theoretically grounded application of dropout in recurrent neural networks. In *Advances in Neural Information Processing Systems*, volume 29, pp. 1019–1027, 2016. doi: 10.48550/arXiv.1512.05287.

Felix A. Gers, Jürgen Schmidhuber, and Fred Cummins. Learning to forget: Continual prediction with LSTM. *Neural Computation*, 12(10):2451–2471, 2000. doi: 10.1162/089976600300015015.

Dar Gilboa, Bo Chang, Minmin Chen, Greg Yang, Samuel S. Schoenholz, Ed H. Chi, and Jeffrey Pennington. Dynamical isometry and a mean field theory of LSTMs and GRUs. *arXiv preprint arXiv:1901.08987*, 2019. doi: 10.48550/arXiv.1901.08987.

Klaus Greff, Rupesh K. Srivastava, Jan Koutník, Bas R. Steunebrink, and Jürgen Schmidhuber. LSTM: A search space odyssey. *IEEE Transactions on Neural Networks and Learning Systems*, 28(10):2222–2232, 2017. doi: 10.1109/TNNLS.2016.2582924.

Albert Gu, Tri Dao, Stefano Ermon, Atri Rudra, and Christopher Ré. HiPPO: Recurrent memory with optimal polynomial projections. In *Advances in Neural Information Processing Systems*, volume 33, pp. 1474–1487, 2020a. doi: 10.48550/arXiv.2008.07669.

Albert Gu, Caglar Gulcehre, Thomas Paine, Matt Hoffman, and Razvan Pascanu. Improving the gating mechanism of recurrent neural networks. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pp. 3800–3809, 2020b. doi: 10.48550/arXiv.1910.09890.

Georg Heinze and Michael Schemper. A solution to the problem of separation in logistic regression. *Statistics in Medicine*, 21(16):2409–2419, 2002. doi: 10.1002/sim.1047.

Mikael Henaff, Arthur Szlam, and Yann LeCun. Recurrent orthogonal networks and long-memory tasks. In *Proceedings of the 33rd International Conference on Machine Learning*, volume 48 of *Proceedings of Machine Learning Research*, pp. 2034–2042, 2016. doi: 10.48550/arXiv.1602.06662.

Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997. doi: 10.1162/neco.1997.9.8.1735.

Rafal Jozefowicz, Wojciech Zaremba, and Ilya Sutskever. An empirical exploration of recurrent network architectures. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pp. 2342–2350, 2015.

Oleksandr Kholodniak. rust-lstm: Recurrent neural networks with verified backpropagation through time in Rust, version 0.10.0. Software, Zenodo, https://doi.org/10.5281/zenodo.23200142, 2026.

Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations*, 2015. doi: 10.48550/arXiv.1412.6980.

David Krueger, Tegan Maharaj, János Kramár, Mohammad Pezeshki, Nicolas Ballas, Nan Rosemary Ke, Anirudh Goyal, Yoshua Bengio, Aaron Courville, and Chris Pal. Zoneout: Regularizing RNNs by randomly preserving hidden activations. In *International Conference on Learning Representations*, 2017. doi: 10.48550/arXiv.1606.01305.

Quoc V. Le, Navdeep Jaitly, and Geoffrey E. Hinton. A simple way to initialize recurrent networks of rectified linear units. *arXiv preprint arXiv:1504.00941*, 2015. doi: 10.48550/arXiv.1504.00941.

Lisha Li, Kevin Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. Hyperband: A novel bandit-based approach to hyperparameter optimization. *Journal of Machine Learning Research*, 18(185): 1–52, 2018.

Taesup Moon, Heeyoul Choi, Hoshik Lee, and Inchul Song. RNNDROP: A novel dropout for RNNs in ASR. In *2015 IEEE Workshop on Automatic Speech Recognition and Understanding*, pp. 65–70, 2015. doi: 10.1109/ASRU.2015.7404775.

Antonio Orvieto, Samuel L. Smith, Albert Gu, Anushan Fernando, Caglar Gulcehre, Razvan Pascanu, and Soham De. Resurrecting recurrent neural networks for long sequences. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 26670–26698, 2023. doi: 10.48550/arXiv.2303.06349.

Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. On the difficulty of training recurrent neural networks. In *Proceedings of the 30th International Conference on Machine Learning*, volume 28 of *Proceedings of Machine Learning Research*, pp. 1310–1318, 2013.

Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems*, volume 32, pp. 8024–8035, 2019. doi: 10.48550/arXiv.1912.01703.

David R. Roberts, Volker Bahn, Simone Ciuti, Mark S. Boyce, Jane Elith, Gurutzeta Guillera-Arroita, Severin Hauenstein, José J. Lahoz-Monfort, Boris Schröder, Wilfried Thuiller, David I. Warton, Brendan A. Wintle, Florian Hartig, and Carsten F. Dormann. Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure. *Ecography*, 40(8):913–929, 2017. doi: 10.1111/ecog.02881.

Stanislau Semeniuta, Aliaksei Severyn, and Erhardt Barth. Recurrent dropout without memory loss. In *Proceedings of the 26th International Conference on Computational Linguistics*, pp. 1757–1766, 2016. doi: 10.48550/arXiv.1603.05118.

Judith D. Singer and John B. Willett. It's about time: Using discrete-time survival analysis to study duration and the timing of events. *Journal of Educational Statistics*, 18(2):155–195, 1993. doi: 10.3102/10769986018002155.

Corentin Tallec and Yann Ollivier. Can recurrent neural networks warp time? In *International Conference on Learning Representations*, 2018. doi: 10.48550/arXiv.1804.11188.

Hans C. van Houwelingen. Dynamic prediction by landmarking in event history analysis. *Scandinavian Journal of Statistics*, 34(1):70–85, 2007. doi: 10.1111/j.1467-9469.2006.00529.x.

Paul J. Werbos. Backpropagation through time: What it does and how to do it. *Proceedings of the IEEE*, 78 (10):1550–1560, 1990. doi: 10.1109/5.58337.

Wojciech Zaremba, Ilya Sutskever, and Oriol Vinyals. Recurrent neural network regularization. *arXiv preprint arXiv:1409.2329*, 2014. doi: 10.48550/arXiv.1409.2329.

Nicolas Zucchet and Antonio Orvieto. Recurrent neural networks: Vanishing and exploding gradients are not the end of the story. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.48550/arXiv.2405.21064.