Interpretable Cutting Plane Selection for Mixed-Integer Linear Programming Using Genetic Programming

Gwen Maudet[1], **Mathis Da Cunha**[1], **Grégoire Danoy**[2]

[1]SnT, University of Luxembourg, Esch-sur-Alzette, Luxembourg
[2]FSTM/DCS, SnT, University of Luxembourg, Esch-sur-Alzette, Luxembourg
{gwen.maudet,gregoire.danoy}@uni.lu

arXiv:2610.29311v1 [cs.LG] 21 Oct 2026

Abstract

Cutting planes (cuts) remove fractional solutions from mixed-integer linear programs (MILPs) without excluding any feasible integer solutions, making them a key ingredient in modern MILP solvers. Current state-of-the-art solvers rely on fine-tuned handcrafted cut strategies, which, despite their effectiveness, remain inherently static and can be suboptimal. More recently, learning-based approaches have been proposed to design adaptive cut strategies ; however, most of these methods are based on neural networks, offering limited interpretability and providing no guarantees of convergence. In this paper, we propose *Genetic Programming and Statistics-based Cut Selection* (GPSCS), the first learning-based framework that explicitly generalizes handcrafted cut selection heuristics into a fully interpretable policy. First, a genetic programming model evolves a symbolic scoring function that ranks candidate cuts using features derived from the cuts and the solving context. Second, a lightweight statistics-based mechanism adaptively determines the number of top-ranked cuts to add at each separation round by detecting significant score dropouts. We evaluate GPSCS within the SCIP solver on seven classes of synthetic and real-world benchmark instances, with a strong emphasis on generalization : evaluation instances are on average 41 times harder than those used for training. Experimental results show that GPSCS improves average solving time by 6% compared to SCIP's default cut selection strategy and achieves a two-fold speedup over both a recent neural-network-based approach and a no-cut baseline, demonstrating strong robustness and generalization across heterogeneous MILP families.

1 Introduction

Mixed-Integer Linear Programming (MILP) plays a central role in modeling and solving a wide range of real-world decision-making problems (Archetti, Peirano, and Speranza 2022; Rais and Viana 2011). Modern MILP solvers like SCIP (Bolusani et al. 2024), gurobi (gurobi 2024), CPLEX (CPLEX 2022) combine multiple complementary components to achieve high solving efficiency. At their core is the branch-and-bound (B&B) algorithm, which recursively decomposes the original problem into smaller subproblems. Each subproblem is evaluated through its linear programming (LP) relaxation, leading either to a feasible solution of the original MILP if integrality constraints are satisfied, or to a lower bound that can be used to prune the corresponding subtree when a better feasible solution is already known. A key mechanism for enhancing the effectiveness of the B&B process is the use of cutting planes (cuts), i.e., valid inequalities that remove fractional LP solutions without excluding any feasible integer solutions. By strengthening the LP relaxation, cuts improve bound quality and, consequently, the solver's pruning capability. However, the design of an effective cut strategy involves an inherent trade-off : while adding cuts can significantly tighten the LP relaxation, each additional constraint also increases the computational burden of solving the LP, potentially slowing down the overall optimization process. An effective strategy must therefore carefully balance the number of generated cuts and their strengthening power, aiming to select a small subset of highly effective cuts that maximally improve the LP bound.

In practice, cut strategies in modern MILP solvers are largely based on handcrafted heuristics, which use scoring functions defined as linear combinations of cut features to rank candidate cuts and select a fixed number of them depending on the node type (root or non-root) (Achterberg 2007; Suhl and Wesselmann 2012). While effective in many settings, such heuristics may exhibit highly variable performance depending on the problem structure as the strategy is fixed for all instances. Recent advances in machine learning (ML) have opened new perspectives for improving decision-making within MILP solvers by enabling the automatic design of adaptive cut strategies (Lodi and Zarpellon 2017; Bengio, Lodi, and Prouvost 2021; Cappart et al. 2021; Scavuzzo et al. 2024). Most existing approaches in this direction rely on neural networks (NN) (Paulus et al. 2022; Tang, Agrawal, and Faenza 2020; Turner et al. 2023; Wang et al. 2023), and have demonstrated significant improvements over static strategies. In particular, the work of (Wang et al. 2023) proposes *HEM*, a comprehensive cutting strategy in which, at each node of the B&B tree, cuts are ranked and a decision is made on the number of cuts to apply. Despite their potential effectiveness, these methods suffer from limited interpretability and have no proof of convergence. Related work is discussed in more detail in section 3.

In this work, we pursue the same objective of learning a complete cutting strategy that operates at each node of the B&B tree and is tailored to a specific class of MILP instances.

Copyright © 2027, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

Motivated by recent evidence showing that simple and interpretable learning approaches grounded in existing handcrafted methods can outperform general-purpose NN approaches in combinatorial optimization (Maudet and Danoy 2025), we adopt a similar perspective. We introduce *Genetic Programming and Statistics-based Cut Selection* (GPSCS), an interpretable framework that jointly learns how to rank candidate cuts according to their impact on bound tightening and how many cuts should be added to the formulation. Specifically, GP is used to evolve explicit scoring functions that evaluate cut relevance based on problem-, node-, and cut-level features, thereby generalizing classical handcrafted scoring rules over a much richer and more flexible search space. Moreover, we propose a lightweight statistics-based dropout detection mechanism that dynamically determines the number of cuts to retain at each node by identifying score breakpoints among the non-parallel cuts. Together, these components define a fully interpretable and adaptive cut strategy, described in detail in section 4.

We implement our approach within the SCIP framework and evaluate it on seven benchmark classes covering both synthetic and real-world MILP instances. The synthetic benchmarks consist of primal-hard instances, where evaluation problems are significantly larger than those used for training, allowing us to assess generalization within homogeneous problem families. The remaining four classes are drawn from established benchmark libraries and represent heterogeneous real-world instances. Here again, training is performed on simpler instances than those used for evaluation. Overall, evaluation instances require, on average, 41 times more solving time than the corresponding training instances. We compare GPSCS against several alternatives that combine GP with different strategies for determining the number of retained cuts, including selecting all cuts, only non-parallel cuts, applying statistical dropout detection to all cuts, and a reinforcement-learning (RL)-based cut quantity policy adapted from the HEM method (Wang et al. 2023). We also evaluate the transferability of the learned GP scoring functions by applying functions trained on one problem class to instances from different classes. In addition, we compare GPSCS against standard baselines, including the default SCIP cut selection strategy, a configuration without cutting planes (NoCuts), and the neural-network-based cut selection method HEM. Among all evaluated methods, GPSCS achieves the best overall performance : it outperforms HEM and *NoCuts* with an average two-fold speedup and improves upon SCIP's default cut selection strategy by 6% on average. Detailed experimental results are presented in section 5.

The main contributions of this paper are summarized as follows :

- We define a general, feature-based cut scoring space that encompasses and extends traditional handcrafted scoring rules.
- We propose a GP-based method to automatically explore this space to generate interpretable cut scoring functions trained to be efficient on training instances.
- We introduce a simple statistics-based procedure to adaptively determine the number of cuts to apply by detecting breakpoints in the score distribution.
- We conduct an extensive empirical evaluation on seven benchmark classes, including both synthetic and real-world instances.
- We experimentally demonstrate the effectiveness of our adaptive cut number selection by comparing it against four alternative strategies for choosing the number of cuts.
- GPSCS achieves significant performance improvements over the default solver strategy, a NoCut configuration, and the state-of-the-art NN-based method HEM, while also highlighting potential convergence issues in HEM on certain instance classes.

2 Background

This section presents the key concepts underlying our approach. We first review the B&B algorithm, which forms the foundation of modern MILP solvers. Next, we discuss cutting planes and the strategies for selecting them. Finally, we introduce GP, a method that evolves a population of programs over successive generations.

B&B A MILP P is defined by decision variables $x = (x_i)_{i\in\{1,\dots,n\}}$, where a subset indexed by $\mathcal{J} \subseteq \{1,\dots,n\}$ is required to take integer values. The objective is to minimize a linear function $c^\top x$ subject to m linear constraints. MILPs are typically solved using the B&B algorithm, which recursively partitions the feasible region into subproblems organized as a search tree. At each node N, integrality constraints are relaxed, yielding an LP relaxation that can be solved efficiently. Let x^{LP} denote an optimal LP solution with objective value $z^{LP} = c^\top x^{LP}$. If x^{LP} satisfies all integrality constraints, it is optimal for the subproblem and updates the best known feasible objective value z^*. Otherwise, z^{LP} provides a valid lower bound on the subproblem. If this bound is no better than the incumbent solution (i.e., $z^{LP} \geq z^*$), the entire subtree rooted at N can be *pruned*. If not, the subproblem is branched into two or more child subproblems whose union covers the feasible region of the parent. This process continues until all nodes are either solved or pruned. The efficiency of B&B critically depends on its ability to prune large portions of the search tree early.

Cut Selection Strategy As the B&B's efficiency largely derives from its pruning capabilities, it depends on the quality of the bounds provided by LP relaxations. Cutting planes are introduced to strengthen these bounds. A *cut* is a valid linear inequality of the form $r : \alpha^\top x \leq \beta$ that eliminates fractional LP solutions without excluding any feasible integer solution. By tightening the LP relaxation, cuts increase the likelihood of early pruning during the search. Cuts are incorporated through an iterative procedure where for each B&B node, candidate cuts are generated by dedicated cut generators (such as Gomory cut generation (Gilmore and Gomory 1961) and others), a subset is selected, and the LP is reoptimized with the additional constraints. This process defines the *cut strategy*, which can be characterized by two main components (Dey and Molinaro 2018) : (i) a ranking or scoring mechanism used to assess the relevance of candidate cuts, and (ii) a rule determining the number of cuts ultimately added to the formulation. While adding cuts generally strengthens the LP relaxation, including too many cuts can

substantially increase the size of the formulation and degrade LP solving performance. Further background on MILP solving, B&B, and cutting planes can be found in (Achterberg 2007).

GP Evolutionary algorithms are population-based metaheuristics inspired by the principles of natural evolution. They operate by maintaining a population of candidate solutions, referred to as *individuals*, which is iteratively evolved over successive *generations* according to a fitness criterion that quantifies solution quality with respect to the objective of interest. A distinctive characteristic of GP is that each individual encodes a computer program, most commonly represented as a tree structure. Leaf nodes, called *terminals*, correspond to the basic elements of the program, while internal nodes represent *operators* that combine these elements. The evolutionary process of an evolutionary algorithm starts from an initialized population and, at each generation, produces a new population through the application of three main operators : (i) *selection*, which favors individuals with higher fitness for reproduction ; (ii) *crossover*, which combines two or more individuals by exchanging parts ; and (iii) *mutation*, which introduces random modifications to individuals. This process continues until a predefined stopping condition is met, such as a maximum number of generations. The population-based and stochastic nature of GP allows all individuals within a generation to be evaluated independently, enabling embarrassingly parallel execution and making GP particularly well suited for large and unstructured search spaces. When this space corresponds to a space of heuristics, GP can be viewed as a generative hyper-heuristic (Dokeroglu, Kucukyilmaz, and Talbi 2024). A comprehensive overview of evolutionary algorithms is provided in (Talbi 2009).

3 Related Work

This section reviews prior work on cut selection strategies. We first discuss handcrafted heuristics that are still widely used as default strategies in modern solvers. We then survey recent ML-based approaches that aim to learn adaptive cut selection policies. Finally, we position our work as a lightweight and interpretable alternative that leverages GP to generalize handcrafted scoring rules and combines it with a simple drop detection for determining the number of cuts.

Handcrafted Heuristics Modern MILP solvers such as SCIP, Gurobi, and CPLEX (Bolusani et al. 2024; gurobi 2024; CPLEX 2022) typically rely on handcrafted scoring functions to evaluate and rank candidate cuts. For instance, SCIP defines its scoring function as a weighted combination of four primary cut quality metrics : (i) *efficacy* (eff), measuring the distance between the cut and the LP solution, defined as $(\alpha^\top x^{LP} - \beta)/\|\alpha\|$; (ii) *directed cutoff distance* (dcd), which incorporates information from the best known feasible solution, defined as $(\alpha^\top x^{LP} - \beta)/|\alpha^\top y|$, with $y = (x^* - x^{LP})/\|x^* - x^{LP}\|$; (iii) *objective parallelism* (obp), capturing the alignment between the cut and the objective function vector, defined as $|\alpha^\top c/(\|\alpha\|\|c\|)|$; and (iv) *integer support* (isp), measuring the involvement of integer-constrained variables, defined as $N_{\text{int}}/N_{\text{nonz}}$, where $N_{\text{nonz}} = \sum_{i=1}^{n} \mathbb{1}(\alpha_i \neq 0)$ and $N_{\text{int}} = \sum_{i \in \mathcal{J}} \mathbb{1}(\alpha_i \neq 0)$, with $\mathbb{1}$ denoting the indicator function. Another important metric is the *normalized violation* (nvi) (Wang et al. 2023; Huang et al. 2021), defined as $(\alpha^\top x^{LP} - \beta)/|\beta|$. Empirical and theoretical studies indicate that using fixed global weights for these metrics is suboptimal on heterogeneous instance distributions (Turner et al. 2023), motivating the need for adaptive, instance-aware cut scoring strategies.

Similarly, the number of cuts added per separation round is typically governed by fixed limits ; in SCIP, for instance, larger limits are applied at the root node and smaller limits elsewhere, often retaining most cuts. Procedures are also used to discard nearly parallel cuts to preserve effectiveness. However, as highlighted in (Dey and Molinaro 2018; Wang et al. 2023), determining the appropriate number of cuts is critical, and fixed limits are generally suboptimal, motivating adaptive strategies that balance LP strengthening with computational overhead.

ML Approaches Recent advances in ML have led to the formulation of several solver decisions as learning problems, including node selection, branching, and cut selection, as summarized in several surveys (Lodi and Zarpellon 2017; Bengio, Lodi, and Prouvost 2021; Cappart et al. 2021; Scavuzzo et al. 2024). These approaches can be broadly categorized into imitation learning and self-learning. For cut selection, (Paulus et al. 2022) propose an imitation learning approach in which a NN is trained to mimic an oracle that selects the best scoring metric from a predefined set of handcrafted measures, thereby reducing the computational cost of oracle-based decisions at inference time. However, this method only selects a single cut per separation round, whereas modern solvers typically add multiple cuts simultaneously to achieve stronger LP tightening. Moreover, such an oracle can only be heuristic in nature, which inherently bounds the performance of the learned model by the quality of the oracle itself and may therefore lead to suboptimal strategies. For these reasons, most recent work has focused on self-learning approaches based on RL. One of the earliest applications in this direction is proposed in (Tang, Agrawal, and Faenza 2020), where deep RL is used to select Gomory cuts, again restricting the selection to a single cut per round, which limits practical applicability. The approach proposed in (Turner et al. 2023) extends this line of work by enabling the selection of multiple cuts through the dynamic adjustment of the weights used in the cut scoring mechanism of SCIP. However, the number of cuts added remains fixed, and the learned policy is evaluated solely based on bound improvements at the root node. A more comprehensive proposition of cut selection is given with *HEM* in (Wang et al. 2023), which jointly addresses cut ranking and cut quantity decisions that are done along all the solving process. The authors introduce a hierarchical NN based RL framework composed of a high-level policy that determines the number of cuts to add and a lower-level sequence model that ranks candidate cuts. Although this approach improves upon static solver configurations, the resulting policies are black-box models, limiting interpretability. More importantly, the absence of prior structural or heuristic knowledge requires exploring a greatly enlarged decision space, combining both the number

of cuts and a total ordering over all candidates. Because the reward is the solving time after fully exploring the B&B tree, this large space coupled with sparse, delayed feedback makes training highly demanding, potentially unstable, and without guaranteed convergence.

Positioning While NN-based approaches—particularly those relying on self-learning—offer a general and powerful framework for modeling solver decisions, their all-purpose nature often entails expensive training times and provides no guarantee of stable or rapid convergence toward high-quality policies. In contrast, leveraging decades of domain knowledge embedded in handcrafted heuristics to guide automated method generation enables faster convergence, stronger inductive bias, and the production of interpretable strategies.

GP provides a lightweight and interpretable alternative to NN-based methods by directly evolving executable and human-readable heuristics. Recent work (Maudet and Danoy 2025) demonstrates that GP can be used as a hyper-heuristic to learn effective B&B search strategies by generalizing existing handcrafted heuristics through symbolic expressions, achieving improvements over default solver configurations and NN methods, while requiring significantly less computational effort. Our work builds on this idea by applying GP to the cut selection problem, which has not previously been explored using GP. In contrast to NN-based RL approaches, our framework produces lightweight scoring functions that are evaluated at each B&B node. These functions combine features derived from the cut, the current node, and the problem instance, and extend classical handcrafted scoring rules over a substantially richer search space, enabling the discovery of novel and effective cut ranking policies that generalize across instance classes.

Beyond cut ranking, the number of cuts added at each separation round and node has a significant impact on solver performance (Dey and Molinaro 2018; Wang et al. 2023). Rather than learning this decision through an additional learning procedure that would imply building a looping procedure to ensure convergences with the GP, we propose a simple deterministic cutoff method based on statistical analysis of score distributions. By assuming that the scoring function reflects cut quality, sharp drops in the distribution signal lower-ranked cuts unlikely to add value ; our method then retains only the highest-scoring cuts. In our framework, the GP component ensures this accurate scoring, thereby guaranteeing the effectiveness of the statistics-based cutoff.

4 GP-based Cut Selection Strategy

We introduce *genetic programming and statistics-based cut selection* (GPSCS), a learning-based cut strategy built around two complementary components. First, we propose a selection procedure that filters highly parallel cuts and adaptively determines how many cuts should be added at each separation round. Second, we leverage GP to automatically learn a cut scoring function that ranks candidate cuts based on their expected impact. In contrast to traditional solver heuristics, which rely on manually designed and fixed scoring rules, GPSCS learns an explicit and fully interpretable strategy that adapts both across problem instances and dynamically across nodes of the B&B tree during the solving process.

Selection Procedure According to a Cut Scoring Function

We first describe the procedure used to determine the final subset of cuts to be applied at each separation round and at each node, given a cut scoring function. Let $f(r, N, P)$ denote a scoring function that assigns a numerical value to a cut r generated at node N of the B&B tree for a problem instance P. Evaluating f over the set of generated cuts induces an ordering $\{(r_i)_{1\leq i\leq l}\}$ such that $i \leq j$ implies $f(r_i) \geq f(r_j)$. The objective is to select a subset $\mathcal{R} \subseteq \{r_1, \ldots, r_l\}$ of cuts to be added to the formulation. The selection procedure consists of two successive stages. First, a filtering step removes highly parallel cuts to promote geometric diversity. Second, a dropout detection heuristic is applied to the filtered list in order to identify a natural cutoff point in the ranking, retaining only the most promising cuts.

Prioritization of Orthogonal Cuts Starting from the initial ranking induced by the learned scoring function, we apply a filtering procedure commonly used both in solvers and in recent research (Bolusani et al. 2024; Turner et al. 2023). The objective is to prioritize cuts that introduce genuinely new geometric information, rather than reinforcing directions already well represented by previously selected cuts. Such a criterion has been consistently shown to improve the effectiveness of cutting planes by reducing redundancy and strengthening LP relaxations more efficiently (Achterberg 2007; Dey and Molinaro 2018; Turner et al. 2023; Wang et al. 2023). Since cut parallelism can be computed explicitly, handling this filtering step separately simplifies the learning task faced by the scoring function.

Concretely, cuts are processed iteratively from highest to lowest score. For each selected cut $r : \alpha^\top x \leq \beta$, the absolute cosine similarity with every lower-ranked cut $r' : {\alpha'}^\top x \leq \beta'$ is computed as $s(r, r') = |\alpha^\top \alpha'|/(\|\alpha\|\|\alpha'\|)$, where $s(r, r') = 1$ denotes perfect parallelism and $s(r, r') = 0$ orthogonality. Any lower-ranked cut whose similarity exceeds a predefined threshold (set to 0.9 in our experiments) is discarded.

Heuristic for Determining the Number of Cuts After filtering highly parallel cuts from the ranked list, the remaining question is how many of the remaining candidates should be retained at the current separation round. We address this question through a score dropout detection heuristic, based on the assumption that the learned scoring function induces a meaningful ordering of cut quality : a pronounced drop between consecutive scores indicates diminishing marginal benefits from adding further cuts. Since adding cuts is generally beneficial, the goal is not aggressive filtering but rather removing only the lowest-quality cuts at the bottom of the ranking.

Let $\{(r_i)_{1\leq i\leq l}\}$ denote the filtered and ordered list of cuts, and let $\Delta = \{|f(r_i) - f(r_{i+1})|\}_{1\leq i\leq l-1}$ be the sequence of absolute differences between consecutive scores. Since scoring functions may operate on different numerical

scales, these differences are normalized using a z-score (Norman and PhD 2003), by computing for each Δ_i the value $z_i = (\Delta_i - \text{mean}(\Delta))/\text{std}(\Delta)$. We aim to detect a significant gap near the end of the ranked list. To this end, we introduce a threshold $z_{\max}$ and scan the sequence $\{z_i\}$ for large gaps. If one or more indices satisfy $z_i > z_{\max}$, we select the largest such index ; otherwise, we select the index of the largest observed gap. Formally, the stopping index is $i_{\text{threshold}} = \max(\{i \mid z_i > z_{\max}\})$ if non-empty, and $i_{\text{threshold}} = \arg\max_i z_i$ otherwise. Only the first $i_{\text{threshold}}$ cuts in the ranked list are added to the formulation.

GP-based Cut Scoring Method

We now present the GP framework for learning cut scoring functions, first defining the policy space and then describing how it is explored.

Definition of the Cut Scoring Policy Space We define the cut scoring policy space as the set of symbolic expressions generated by combining cut- and context-related variables with arithmetic operators. Each expression assigns a score to every candidate cut, thereby inducing a ranking from the most to the least promising. Our objective is to generalize the scoring strategies used in modern MILP solvers by defining a richer policy space that subsumes existing handcrafted heuristics while exploiting additional information about the current B&B node and the problem instance.

As terminal symbols, we include the classical cut quality features eff, dcd, obp, N_{int}, N_{nonz}, and nvi. To capture the local solving context, we additionally include the node depth d_i, while the number of variables n and constraints m encode instance-level characteristics. Finally, a large positive constant M enables strong prioritization when required. The operator set consists of the arithmetic operators $+, -, \times, \div$ together with the unary operators $\exp(\cdot)$ and $\sqrt{\cdot}$, allowing both linear and non-linear scoring functions. For instance, a policy prioritizing efficacy while favoring sparse cuts is given by $f(r) = eff(r) \times M - N_{\text{nonz}}(r)$.

GP Framework We employ GP to explore the space of cut scoring heuristics introduced above, thereby framing our approach as a generative hyper-heuristic (Dokeroglu, Kucukyilmaz, and Talbi 2024). We adopt a standard GP framework (Talbi 2009), in which a population of p individuals—each encoding a candidate scoring function—is evolved over a fixed number of generations N. The initial population is generated randomly, with tree depths bounded between $D_{\text{init-min}}$ and $D_{\text{init-max}}$.

The population evolves across generations according to a fitness-driven selection mechanism. In our setting, the objective is to learn a scoring function that efficiently solves a set of training set of instances belonging to the same problem class. Since all training instances are solved to optimality, the most reliable and least biased fitness measure is the solving time. We therefore define the fitness of an individual as the average solving time over the training set, with a time limit of 500 seconds assigned as a penalty when the limit is reached.

For selection, we employ a double tournament strategy (Luke and Panait 2002) designed to limit excessive tree growth and mitigate overfitting. In the first tournament, k individuals are sampled uniformly at random and the one with the best fitness is selected. In the second tournament, two winners from the first stage are compared based on tree size, and the shallower tree is selected with probability $p_{\text{pars}}/2$, where $p_{\text{pars}} \in [1, 2]$.

After selection, crossover is applied with probability P_{mate} (typically above 0.8). We use one-point subtree crossover : a node is randomly selected in each parent tree, and the corresponding subtrees are exchanged to produce two offspring. Mutation is then applied independently to each individual with probability P_{mut} (typically below 0.15). During mutation, a randomly selected subtree is replaced by a newly generated subtree whose depth lies between $D_{\text{mut-min}}$ and $D_{\text{mut-max}}$.

Training and Deployment During training, each GP-derived scoring function is evaluated by solving all training instances to optimality or until the time limit is reached. For each instance and at each B&B node, candidate cuts are ranked using the GP-based scoring function, filtered to promote orthogonality, and selected using the dropout detection heuristic described earlier. At deployment time, the learned scoring function is applied unchanged, following the exact same ranking, filtering, and selection procedure, ensuring full consistency between training and inference.

5 Experiments

This section presents an experimental evaluation of our approach. We first assess several variants of GPSCS by considering different strategies for determining the number of retained cuts and by evaluating GP scoring functions trained on classes different from those used for testing. We then compare GPSCS against established cut selection strategies, including a recent learning-based method. Experiments are conducted on seven classes of instances. The first three classes are synthetically generated, with evaluation instances significantly larger than those used for training, in order to assess the generalization capability of the proposed method. The remaining four classes are drawn from state-of-the-art benchmarks and consist of heterogeneous, real-world instances, with training performed on simpler instances than those in evaluation, allowing us to test performance on realistic problem distributions.

Experimental Setup All methods are implemented within the open-source solver SCIP version 9.2.2, with all default solver components left unchanged except for the cut selection strategy. The source code is publicly available at https://.....git. For both training and evaluation, each run is executed on a dedicated node of an HPC cluster. CPU-only runs are carried out on nodes equipped with 128 cores and 256 GB of RAM, while runs involving GPU resources use nodes with 28 CPU cores, 768 GB of RAM, and four NVIDIA V100 GPUs with 16 GB each. Since SCIP itself solves each instance using a single CPU core, allocating an entire node primarily provides access to large memory capacities ; in particular, GPU-equipped nodes offer a hardware advantage in terms of available RAM and additional GPUs. These

multi-core and multi-GPU resources are exploited through parallelism during the training phase, whereas evaluation instances are solved sequentially, one at a time. To prevent excessive computation times, we focus on instances that can be solved within a reasonable duration and impose a time limit of 500 seconds per instance.

Experiments are conducted on seven instance classes, each split into 50 instances for training and 50 for evaluation. The same training sets are used for all methods requiring training (all GP variants and the NN-based HEM). The first three classes are synthetically generated to evaluate generalization to larger instance from homogeneous training set : (i) Generalized Independent Set (GISP) (Chmiela et al. 2021), with training graphs of size $n_g \in [60, 70]$ and evaluation graphs of size $n_g \in [120, 140]$; (ii) Maximum Satisfiability (MAXSAT) (Béjar et al. 2009), using the same size scaling ; and (iii) Fixed Charge Multicommodity Network Flow (FCMCNF) (Hewitt, Nemhauser, and Savelsbergh 2009), with $n_g = 15$ nodes for training and $n_g = 20$ for evaluation, and $m = 1.5n_g$ commodities (Chmiela et al. 2021). Graphs are generated via the Erdos-Renyi model (Erdos 1961) with edge probabilities 0.3 for GISP and MAXSAT, and 0.6 for FCMCNF. The remaining four classes are drawn from established benchmarks to evaluate performance on heterogeneous, real-world instances, also using easier instances for training and harder ones for evaluation. For selection, we first choose 100 instances per class that can be solved within 300 seconds, then use the 50 fastest for training and the remaining 50 for evaluation. These include the Mixed-Integer Knapsack Polyhedron (MIK) class (Atamturk 2003) and the three StrIPlib (Bastubbe et al. 2025) classes with the largest number of instances—binpacking, cuttingstock, and capacitated lot-sizing—to ensure sufficient heterogeneity.

We compare our approach against three baselines : a configuration without cutting planes (*NoCuts*), the default SCIP cut selection strategy (*SCIP*), and a recent neural-network-based method that jointly learns cut ranking and cut quantity decisions (Wang et al. 2023) (*HEM*). For the GP-based methods, we investigate several alternatives by considering the following configurations : (i) *GPSCS*, our proposed method applied to the list of non-parallel cuts only, where the statistics-based decision rule determines the number of retained cuts for non parallel cuts ; (ii) *GPSAll*, which applies the same decision rule to the complete set of generated cuts. To favor geometric diversity, non-parallel cuts are ranked first, followed by the filtered parallel cuts, each group remaining internally ordered according to the GP score. The score-dropout detection is applied independently within each group, preventing the transition between the two groups from influencing the cutoff decision ; (iii) *GPAll*, which retains all generated cuts ; (iv) *GPNP*, which retains all non-parallel cuts ; and (v) *GPRL*, a hybrid variant that uses the GP scoring function learned by *GPSCS* while replacing the statistical decision rule with the RL-based cut quantity module of *HEM*. All these variants follow the same pipeline : the specific cut ranking function is incorporated during GP training, producing a distinct GP-derived scoring function for each alternative. Finally, to evaluate the transferability of GP scoring functions across problem classes, we consider *WGPSCS*. For each evaluation instance, *WGPSCS* randomly selects a GP scoring function trained on a different problem class and combines it with the same statistics-based decision rule as *GPSCS*. For *GPSCS*, *WGPSCS* and *GPSAll*, the dropout-based cut quantity mechanism uses a maximum z-score threshold of $z_{\max} = 1$. All GP trainings employ a population of $p = 64$ individuals evolved over $N = 20$ generations, with fitness evaluations performed in parallel on a CPU node. Selection is based on tournament selection with size $k = 5$ and a parsimony coefficient of $p_{\text{pars}} = 1.2$, while crossover and mutation are applied with probabilities $P_{\text{mate}} = 0.9$ and $P_{\text{mutate}} = 0.1$, respectively. Initial tree depths are bounded between 1 and 17, and mutation-generated subtrees are restricted to depths between 1 and 5. Since GP is inherently stochastic, for all GP-based variants, we perform 15 independent training runs with different seeds and retain the scoring function achieving the best performance on the training set. For the NN baseline, the *HEM* approach needs training two neural networks in parallel over 100 epochs on four GPUs of a single computing node.

Problem	**GP Function**	**Train perf**
FCMCNF	$-\left(m + \frac{\exp(\exp(-N_{\text{int}}))}{eff}\right)$	3.53 ± 0.10
MAXSAT	$nvi \times (dcd - d_i)$	4.25 ± 0.06
GISP	$\frac{d_i - \sqrt{eff}}{nvi}$	2.53 ± 0.02
MIK	$\exp(dcd \times m)$	16.65 ± 0.48
Cuttingstock	$\left(dcd - \left(\frac{m}{dcd} + obp\right)\right) \times n$	7.12 ± 1.40
Binpacking	$-\frac{nvi}{dcd \times eff}$	6.83 ± 0.14
Cap lot-sizing	$-\frac{N_{\text{nonz}}}{dcd \times eff}$	4.19 ± 0.02

TABLE 1 – Best GPSCS GP functions selected from 15 runs based on training solving time ; values report the mean and standard deviation of the selected functions across runs.

Performance Evaluation The retained scoring functions learned by *GPSCS*, namely the candidates achieving the lowest average solving time on the training set, are reported in table 1. We also summarize the overall GP training behavior by reporting, across the 15 independent runs, the mean solving time and standard deviation of the selected best-performing candidates. Although not shown explicitly, the GP runs do not converge to identical symbolic expressions. Nevertheless, their performance is highly consistent : for the synthetically generated instances, the standard deviation remains below 0.06 seconds for mean solving times of only a few seconds, while slightly larger variations are observed on benchmark classes due to their higher instance heterogeneity. Convergence plots leading to the final *GPSCS* scoring functions are provided in Appendix C. While the learned functions differ substantially across problem classes, highlighting the need for problem-specific cut selection strategies, some cut quality features consistently appear across the learned policies. Appendix D provides a frequency-based analysis

Method	FCMCNF	MAXSAT	GISP	MIK	Binpacking	Cuttingstock	Cap lot-sizing	Mean
NoCuts	500.0 ± 0	262.2 ± 202	71.3 ± 85	234.7 ± 215	293.2 ± 189	452.3 ± 124	500.0 ± 0	330.6 ± 207
SCIP	25.7 ± 23	281.6 ± 170	334.2 ± 107	14.5 ± 10	169.1 ± 54	181.0 ± 43	147.9 ± 53	172.1 ± 142
HEM	500.0 ± 0	447.1 ± 111	204.3 ± 77	295.0 ± 221	400.2 ± 175	399.9 ± 192	500.0 ± 0	392.4 ± 111
GPRL	27.1 ± 22	283.2 ± 174	399.7 ± 145	33.4 ± 34	225.6 ± 160	176.4 ± 184	255.1 ± 138	200.1 ± 122
GPAll	28.8 ± 23	309.4 ± 170	292.6 ± 107	26.5 ± 22	199.8 ± 170	182.8 ± 158	201.1 ± 105	183.7 ± 163
GPNP	24.1 ± 18	280.2 ± 168	321.6 ± 123	20.3 ± 19	188.8 ± 160	128.1 ± 159	178.7 ± 117	170.2 ± 166
GPSAll	23.8 ± 17	276.3 ± 165	301.8 ± 119	19.1 ± 15	211.9 ± 155	144.7 ± 165	190.8 ± 107	173.2 ± 162
GPSCS	23.1 ± 17	245.2 ± 150	306.3 ± 100	18.7 ± 14	191.4 ± 159	137.4 ± 139	172.4 ± 117	162.4 ± 153
WGPSCS	500.0 ± 0	423.4 ± 124	269.1 ± 103	258.7 ± 218	296.6 ± 190	450.3 ± 130	500.0 ± 0	383.6 ± 166

TABLE 2 – Comparative results across seven benchmark classes for three state-of-the-art cut selection strategies and five GP-based variants. Reported values are the mean (standard deviation) of solving times. Green highlights the best baseline per class, red indicates solving times at least two times longer than the best, with a gradient used for intermediate values.

of feature occurrences and their contribution to score variations, distinguishing whether increasing a feature value tends to increase or decrease the resulting cut score.

table 2 reports the mean solving times of the three baselines and the five GP-based variants across the seven problem classes. Evaluation instances are on average 41 times harder than training instances in terms of solving time, highlighting the significant generalization challenge addressed in this study. Among the GP-based variants, the RL-based strategy for determining the number of retained cuts (*GPRL*) consistently underperforms the other approaches, showing that using RL solely for cut quantity decisions remains ineffective, even when combined with a learned GP ranking function. Selecting only non-parallel cuts (*GPNP*), which is close to the strategy employed by SCIP, is generally outperformed by *GPSCS*, although the difference remains limited, with *GPNP* requiring only 5% more solving time on average. As reported in Appendix E, the statistics-based cutoff mechanism reduces the number of selected cuts by approximately 10% compared to retaining all non-parallel cuts, suggesting that a fraction of these additional non-parallel cuts provides limited benefit for the solving process. Applying the statistics-based cutoff to the complete ranked cut list (*GPSAll*) also achieves competitive results but remains slightly below *GPSCS* and *GPNP*. Overall, *GPSCS* provides the best performance among the GP-based variants. Compared with the other approaches, *GPSCS* also achieves the best overall performance, being the fastest method on three out of seven problem classes and obtaining the lowest average solving time across all classes. It improves over *SCIP* by 6% on average, despite *SCIP* being a highly optimized solver that jointly exploits many handcrafted components, while the GP functions are trained on instances that are on average 41 times smaller than those used for evaluation. The transferability experiment further confirms the importance of learning class-specific scoring functions. The *WGPSCS* variant, which randomly applies a GP function trained on a different problem class, consistently achieves the worst performance, showing that the learned scoring functions are effective precisely because they capture characteristics specific to their target problem classes. Finally, *HEM* also exhibits limited performance. Except on MAXSAT, its solving times are generally at least twice as large as those of the best-performing methods and are often comparable to the *NoCuts* baseline (with the exception of GISP). Moreover, whenever *NoCuts* reaches the time limit of 500 seconds, *HEM* also fails to solve the instances, even in cases where *GPSCS* and *SCIP* solve them within 25 seconds. This suggests potential convergence limitations of the RL-based framework in these settings.

6 Conclusions

In this work, we investigated the problem of cut selection, a key component governing the efficiency of MILP solvers. We introduced GPSCS, a fully explicit and human-readable cut strategy that jointly addresses cut ranking and cut quantity decisions. Cut relevance is modeled through nonlinear scoring functions combining cut-specific and contextual features, while selection is further refined by prioritizing non-parallel cuts and dynamically determining the number of cuts to add via a statistical dropout detection mechanism applied to score distributions.

Extensive experimental results demonstrate that GPSCS outperforms the default SCIP configuration and the most practically applicable NN-based baseline for cut selection. Notably, our experiments also reveal potential convergence issues in the NN based RL framework, highlighting limitations induced by delayed and sparse reward signals in large decision spaces.

From the perspective of practical deployment, our contributions can be integrated at two distinct levels. First, the proposed statistical rule for dynamically determining the number of cuts is solver-agnostic and independent of the problem class; it can therefore be directly incorporated into existing MILP solvers, with possible refinements. Second, the GP-based scoring mechanism learns instance-class-specific cut ranking functions, making it particularly suitable as a tuned cutting strategy for researchers or practitioners who repeatedly solve large collections of instances from the same problem family.

Overall, this paper demonstrates that embedding decades of expert knowledge on cutting planes into the design space of learning algorithms enables the discovery of high-performing and interpretable strategies with faster and more stable convergence. Our results highlight population-based methods, such as GP, as a compelling and complementary alternative to prevailing NN-based approaches for improving solver performance.

Références

Achterberg, T. 2007. *Constraint Integer Programming*. Ph.D. thesis, ZIB Berlin.

Archetti, C. ; Peirano, L. ; and Speranza, M. G. 2022. Optimization in multimodal freight transportation problems : A Survey. *European Journal of Operational Research*, 299(1) : 1–20.

Atamturk, A. 2003. On the facets of the mixed–integer knapsack polyhedron. *Mathematical Programming*, 98(1) : 145–175.

Bastubbe, M. ; Helber, A. ; Kirchhart, L. ; L{\"u}bbecke, M. ; Rieken, N. ; and Witt, J. 2025. strIPlib : Structured Integer Programming Library.

Bengio, Y. ; Lodi, A. ; and Prouvost, A. 2021. Machine Learning for Combinatorial Optimization : a Methodological Tour d'Horizon. *European Journal of Operational Research*. ArXiv :1811.06128 [cs, stat].

Bolusani, S. ; Besançon, M. ; Bestuzheva, K. ; Chmiela, A. ; Dionisio, J. ; Donkiewicz, T. ; Doornmalen, J. v. ; Eifler, L. ; Ghannam, M. ; Gleixner, A. ; Graczyk, C. ; Halbig, K. ; Hedtke, I. ; Hoen, A. ; Hojny, C. ; Hulst, R. v. d. ; Kamp, D. ; Koch, T. ; Kofler, K. ; Lentz, J. ; Manns, J. ; Mexi, G. ; Mühmer, E. ; Pfetsch, M. E. ; Schlösser, F. ; Serrano, F. ; Shinano, Y. ; Turner, M. ; Vigerske, S. ; Weninger, D. ; and Xu, L. 2024. The SCIP Optimization Suite 9.0.

Béjar, R. ; Cabiscol, A. ; Manyà, F. ; and Planes, J. 2009. Generating Hard Instances for MaxSAT. In *2009 39th International Symposium on Multiple-Valued Logic*, 191–195.

Cappart, Q. ; Chételat, D. ; Khalil, E. B. ; Lodi, A. ; Morris, C. ; and Veličković, P. 2021. Combinatorial Optimization and Reasoning with Graph Neural Networks. In *Proceedings of the Thirtieth International Joint Conference on Artificial Intelligence*, 4348–4355. Montreal, Canada : International Joint Conferences on Artificial Intelligence Organization. ISBN 978-0-9992411-9-6.

Chmiela, A. ; Khalil, E. B. ; Gleixner, A. ; Lodi, A. ; and Pokutta, S. 2021. Learning to Schedule Heuristics in Branch-and-Bound. ArXiv :2103.10294 [cs, math].

CPLEX, I. 2022. IBM ILOG CPLEX Optimization Studio.

Dey, S. S. ; and Molinaro, M. 2018. Theoretical challenges towards cutting-plane selection. *Mathematical Programming*, 170(1) : 237–266.

Dokeroglu, T. ; Kucukyilmaz, T. ; and Talbi, E.-G. 2024. Hyper-heuristics : A survey and taxonomy. *Computers & Industrial Engineering*, 187 : 109815.

Erdos, P. 1961. On the evolution of random graphs. *Bulletin of the Institute of International Statistics*, 38 : 343–347.

Gilmore, P. C. ; and Gomory, R. E. 1961. A Linear Programming Approach to the Cutting-Stock Problem. *Operations Research*.

gurobi. 2024. Gurobi Optimizer Reference Manual.

Hewitt, M. ; Nemhauser, G. L. ; and Savelsbergh, M. W. P. 2009. Combining Exact and Heuristic Approaches for the Capacitated Fixed-Charge Network Flow Problem. *INFORMS Journal on Computing*.

Huang, Z. ; Wang, K. ; Liu, F. ; Zhen, H.-l. ; Zhang, W. ; Yuan, M. ; Hao, J. ; Yu, Y. ; and Wang, J. 2021. Learning to Select Cuts for Efficient Mixed-Integer Programming. ArXiv :2105.13645 [math].

Lodi, A. ; and Zarpellon, G. 2017. On learning and branching : a survey. *TOP*, 25(2) : 207–236.

Luke, S. ; and Panait, L. 2002. Fighting Bloat with Nonparametric Parsimony Pressure. In Guervós, J. J. M. ; Adamidis, P. ; Beyer, H.-G. ; Schwefel, H.-P. ; and Fernández-Villacañas, J.-L., eds., *Parallel Problem Solving from Nature — PPSN VII*, 411–421. Berlin, Heidelberg : Springer. ISBN 978-3-540-45712-1.

Maudet, G. ; and Danoy, G. 2025. Search Strategy Generation for Branch and Bound Using Genetic Programming. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(11) : 11299–11308.

Norman, G. R. ; and PhD, D. L. S. 2003. *PDQ Statistics*. PMPH USA. ISBN 978-1-60795-077-6. Google-Books-ID : Lw_2DwAAQBAJ.

Paulus, M. B. ; Zarpellon, G. ; Krause, A. ; Charlin, L. ; and Maddison, C. 2022. Learning to Cut by Looking Ahead : Cutting Plane Selection via Imitation Learning. In *Proceedings of the 39th International Conference on Machine Learning*, 17584–17600. PMLR.

Rais, A. ; and Viana, A. 2011. Operations Research in Healthcare : a survey. *International Transactions in Operational Research*, 18(1) : 1–31.

Scavuzzo, L. ; Aardal, K. ; Lodi, A. ; and Yorke-Smith, N. 2024. Machine learning augmented branch and bound for mixed integer linear programming. *Mathematical Programming*.

Suhl, U. H. ; and Wesselmann, F. 2012. Implementing cutting plane management and selection techniques.

Talbi, E.-G. 2009. *Metaheuristics : From Design to Implementation*. John Wiley & Sons. ISBN 978-0-470-49690-9. Google-Books-ID : SIsa6zi5XV8C.

Tang, Y. ; Agrawal, S. ; and Faenza, Y. 2020. Reinforcement Learning for Integer Programming : Learning to Cut. In *Proceedings of the 37th International Conference on Machine Learning*, 9367–9376. PMLR.

Turner, M. ; Koch, T. ; Serrano, F. ; and Winkler, M. 2023. Adaptive Cut Selection in Mixed-Integer Linear Programming. *Open Journal of Mathematical Optimization*, 4 : 1–28.

Wang, Z. ; Li, X. ; Wang, J. ; Kuang, Y. ; Yuan, M. ; Zeng, J. ; Zhang, Y. ; and Wu, F. 2023. Learning Cut Selection for Mixed-Integer Linear Programming via Hierarchical Sequence Model. In *The Eleventh International Conference on Learning Representations*. The Eleventh International Conference on Learning Representations.