arXiv:2610.31436v1 [cs.AI] 22 Oct 2026

Conservative Meta-Agent Search: Architecture Adaptation under Non-Stationarity

Ali Atiah Alzahrani[†]

Abstract

Automated design of agentic systems is typically evaluated on static task suites with immediate scalar rewards. Deployed systems face persistent shift, delayed dependent feedback, correlated component failures, and costly architecture changes. We formalize meta-agent architecture search as *sequential incumbent replacement* with switching costs and hard risk constraints, and introduce Conservative Meta-Agent Search (CMAS). Throughout, *meta-agent* denotes the entire outer loop: proposer, compiler, evaluator, gate, and rollback controller. CMAS retrieves archived architectures and proposes typed graph edits, then scores each against the incumbent by paired shadow replay under common random numbers. Component credit is interventional. Promotion requires an *empirically calibrated, variance-penalized simultaneous score* above a pre-specified margin, and is followed by a simulated canary stage and pre-specified rollback. Backtested return cannot certify meta-level skill, so we introduce METAINVEST-BENCH, a regime-switching simulator whose specified worker competence makes a surrogate-optimal architecture computable, with leakage-controlled replay, adversarial governance tracks, and five reusable negative controls. Financial streams give an instrumented, risk-sensitive test of meta-level generalization; we claim no alpha. Across 96 streams, CMAS cuts normalized mean surrogate dynamic-oracle regret from 0.116 to 0.104 and raises held-out utility from 0.44 to 0.48 over the strongest adaptive common-budget reimplementation; when extra workers add no value, it contracts to near-single-agent teams. Its *deployment* inference cost is $\sim$6% lower, though its search-inclusive token cost is the highest of the methods compared. The observed per-update false-promotion rate is 3.4% against a nominal 5% target, and 4.8% in a sealed rerun. The statistical ingredients are standard. The calibration is empirical: it concerns an update-level replay statistic on an author-designed benchmark, and guarantees neither full-utility improvement, nor absence of simulated canary-stage harm, nor sequence-level error control.

1 Introduction

Meta-agents act on agentic systems. They design harnesses, optimize compound workflows, allocate resources, and supervise workers. An LLM system's performance depends on weights, but also on prompts, tools, memory, communication, verification, and control flow [1, 3, 5], and a growing literature searches code-defined workflows [19], modular components [14], communication graphs [23, 18], and query-dependent architectures [9, 20]. Yet most evaluations remain *static*: candidates receive immediate scores on exchangeable tasks, and the best design is deployed without modeling persistent change or the risk of a bad update.

Architecture is therefore a *decision variable* that should track the regime (Figure 1). Investment systems expose non-stationarity, delayed and dependent feedback, correlated component errors,

[†]Corresponding author: aliatiah100@gmail.com. The views expressed are those of the author and do not reflect the views of any other individual or entity. This material is for research purposes only and does not constitute investment advice.

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026).

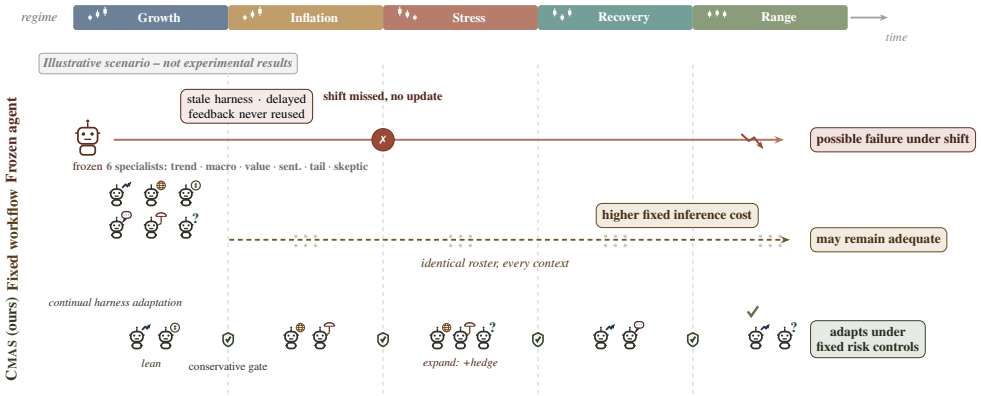

Figure 1: **Illustrative schematic of architecture as a sequential decision variable; annotations are conceptual and do not report experimental results.** A frozen workflow may miss persistent shifts; a fixed committee stays expensive in every regime; CMAS expands, rewires, contracts, retains, or rolls back under fixed external controls, after reversible calibrated promotion.

switching costs, and tail-risk constraints. Finance is useful here as an instrumented, risk-sensitive test of meta-level generalization and conservative update evaluation. Backtested return is not intrinsically persuasive [6, 7, 13]. What is missing is (i) a controlled benchmark whose surrogate-optimal worker architecture is computable and changes over time, and (ii) a method that *screens* candidate updates under noisy, delayed feedback [17, 15, 12, 2]. An update whose simulator-evaluated mean improvement is at most the promotion margin is a *false promotion*, the meta-level analogue of a false discovery and a near-neutral change; Table 2 separates it from the outcomes that do make the system worse.

Generated multi-agent systems can also be costlier yet weaker than strong single-agent baselines [4]; a credible meta-agent must learn both when to assemble specialists and when to collapse to a simpler system. We ask whether a meta-agent can adapt the composition, communication graph, and compute allocation of a compound agentic system under regime shifts while empirically limiting false promotion, architecture bloat, and operational risk.

We answer with CMAS (Figure 2). The *inner loop* executes a bounded agent architecture behind a deterministic constraint and authorization layer. The *outer loop* observes the pre-decision regime descriptor z_k and a trace archive, proposes typed edits, evaluates candidates against the incumbent by paired shadow replay, assigns counterfactual credit, and promotes only when the empirically calibrated simultaneous score clears the margin. Promotion enters a simulated canary stage and is reversible. Every method sees the same pre-decision descriptor z_k, available before the decisions it informs, while the latent simulator regime s_t is evaluator-only. Throughout, *meta-agent* denotes this entire outer loop (proposer, compiler, evaluator, credit estimator, gate, and rollback controller), not the proposal model alone. It is a *hybrid*: generation is LLM-based and archive-conditioned, while the compiler, risk projection, authorization layer, gate, and rollback rules are deterministic and outside anything the LLM can edit (Figure 2, locked box).

Three contributions follow. (1) We formulate agent-architecture search as sequential *incumbent replacement* over typed architectures, with a structural distance (Eq. (1)) as switching cost, delayed dependent utility (Eq. (2)), hard eligibility constraints, an explicit false-promotion target, and meta-level generalization (Definition 1) as the scored quantity. (2) CMAS couples archive-conditioned typed edits, paired candidate–incumbent replay under common random numbers, interventional Shapley-style [34, 35] credit fed back to the proposer, and a conservative promotion-and-rollback gate. The statistical ingredients are standard; the contribution is their integration. Proposition 1 states only the *idealized* conditions under which such a gate would control family-wise false promotion, and Section 4 scopes the departures. (3) METAINVEST-BENCH's controlled track makes surrogate-oracle regret, selection accuracy, and credit fidelity computable, and five reusable controls (clone-collapse, shuffled-regime, gate-calibration, proposer-value, Goodhart-drift) measure whether adaptation happens for the intended reason, exposing reward gaming under a return-only evaluator (RQ7).

The main empirical finding is a *promotion-calibration separation*. Holding the evaluator and gate fixed, false promotion stays at 3.2–3.5% across proposers while utility tracks proposer quality (Table 7).

2 Related work and positioning

Automated design and optimization of agent systems. ADAS, AgentSquare/AFlow, MaAS/AutoMaAS, and GEPA search executable programs, workflows, architectures, or prompts [3, 14, 19, 20, 9, 1]. GPTSwarm and G-Designer optimize communication graphs [23, 18]. Adaptive Auto-Harness directly studies sustained harness construction and routing on open-ended, heterogeneous, non-stationary streams [43]. HarnessAudit evaluates trajectory-level boundary and information-flow compliance, while SafeHarness integrates verification, privilege control, degradation, and rollback across the agent lifecycle [44, 45]. What CMAS makes first-class is repeated candidate–incumbent replacement, paired replay on shared randomness, update-level promotion calibration, switching cost, and a reversible simulated canary stage. The closest systems route at solve time and leave the deployed incumbent in place [43, 20], gate on candidate quality without a paired incumbent comparison or rollback [39, 40], or audit and secure execution without an update-level false-promotion estimand [44, 45]; full comparisons in Appendix I.

Safe improvement, verification, and evaluation validity. Four literatures supply ingredients (Appendix I): safe policy improvement and conservative bandits give the gate its form [26–28]; non-stationary switching-cost learning the regret metric [31, 32]; safety and verification work the threat model [11, 16, 8]; off-policy evaluation [38] motivates paired replay. New here are a typed action space with a distance (Eq. (1)), adaptively LLM-generated pools (so the guarantee is pool-conditional), and interventional trace-level credit. Because financial evaluation is vulnerable to leakage [6, 7, 13, 22, 17, 15, 12, 10, 2], the simulator is primary and market replay provides external validity only.

3 Problem: sequential meta-agent design under non-stationarity

Inner decision system. At time t an information packet x_t is $\mathcal{F}_t$-measurable; architecture h produces an unconstrained proposal $\tilde{w}_t(h)$ and trace $\tau_t(h)$. A deterministic projection maps it to the feasible set, $w_t(h) = \Pi_{\mathcal{C}_t}(\tilde{w}_t(h))$, with $\mathcal{C}_t$ encoding leverage, concentration, turnover, volatility, and restricted-asset constraints. The LLM system cannot modify $\Pi_{\mathcal{C}_t}$, accounting, authorization, or data cutoffs (Figure 2, locked box). Net return is $R^{\text{net}}_{t+1}(h) = w_t(h)^\top r_{t+1} - \kappa_t \|w_t(h) - w_{t-1}(h)\|_1$; during paired replay $w_{t-1}(h)$ is the candidate's *own* counterfactual weights, so each pays its own path's turnover. Because an active projection can mask proposal differences, we fix its form throughout: Euclidean projection onto the constraint polytope via a QP solver.

Architectures, distance, and meta-level generalization. An executable architecture is a typed object $h = (V, E, \rho, \mu, b, g, v) \in \mathcal{H}$: worker set V, directed acyclic communication graph E, typed roles and prompts ρ, model assignments μ, compute/tool budgets b, aggregation rule g, verification nodes v. We equip $\mathcal{H}$ with a distance

$$d(h, h') = \omega_V d_{\text{Jac}}(V, V') + \omega_E d_{\text{GED}}(E, E') + \omega_b d_{\text{budget}}(b, b') + \omega_\rho d_{\text{typed}}\big((\rho, \mu, g, v), (\rho', \mu', g', v')\big). \tag{1}$$

used both as a switching cost and as a drift diagnostic. All four terms are normalized to $[0, 1]$, so no architecture field changes at zero switching or novelty cost. With $\leq$10 edges (Appendix B) d_{GED} is exact; weights are fixed in advance (Table 17) and act through λ_{switch}. At outer-loop time k, the meta-agent sees the observable pre-decision regime descriptor z_k, the incumbent h_k, and an archive $\mathcal{D}_k$ of traces, credits, and outcomes; it proposes a candidate pool $\mathcal{H}_k \subset \mathcal{H}$.

Definition 1 (Meta-level generalization)**.** *A meta-policy π $(\varepsilon_{\text{gen}}, \varepsilon_{\text{safe}})$-*generalizes *over a held-out factor family F (tasks, regime mixtures, worker libraries, model backbones) if expected held-out utility is within ε_{gen} of in-distribution utility and each held-out violation or false-promotion rate is no more than $\varepsilon_{\text{safe}}$ worse; the strict version sets $\varepsilon_{\text{safe}} = 0$. We report achieved deltas rather than asserting a threshold (Table 13): the strict clause is* not *met, since held-out false promotion rises 0.7 pp (RQ3).*

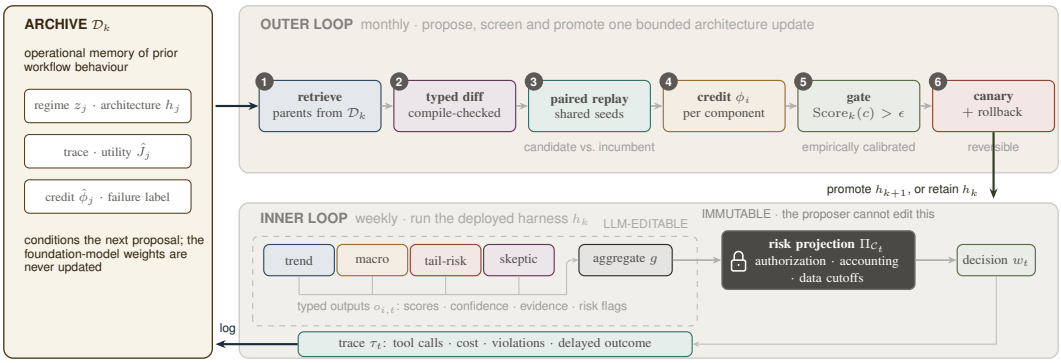

Figure 2: **CMAS as a governed adaptation loop for a multi-agent investment workflow.** The risk projection, authorization, accounting and data cutoffs sit outside the LLM-editable boundary. *Left:* the archive $\mathcal{D}_k$ is operational memory of prior workflow behavior, storing regimes, architectures, traces, utility and credit estimates, and failure labels; the foundation-model weights are never updated. *Bottom (inner loop, weekly):* incumbent h_k runs typed specialists whose outputs $o_{i,t}$ are aggregated by g, after which the immutable projection $\Pi_{\mathcal{C}_t}$ enforces the financial control boundary and the trace τ_t is logged back to the archive. *Top (outer loop, monthly):* ❶ retrieve parents by regime similarity and novelty; ❷ propose a compile-checked, auditable typed diff with hypothesis; ❸ paired shadow replay of candidate against incumbent under common random numbers, plus red-teaming; ❹ interventional credit ϕ_i fed back; ❺ promote only if the empirically calibrated simultaneous score exceeds ϵ; ❻ simulated canary stage with rollback, after which the promoted h_{k+1} (or the retained h_k) returns to the inner loop.

Constrained utility over shadow window W_k scalarizes soft preferences; hard constraints stay eligibility filters. A candidate failing authorization, evidence, leakage, or risk checks is ineligible regardless of utility. For candidate h and the sequence's previous architecture h^-,

$$\begin{aligned} J_k(h; h^-) \;=\; & \mathrm{CEQ}_\gamma\big(R^{\mathrm{net}}(h)\big) \;-\; \lambda_{\mathrm{cvar}}\,\mathrm{CVaR}_{.95}\big(-R^{\mathrm{net}}(h)\big) \;-\; \lambda_{\mathrm{dd}}\,\mathrm{MDD}(h) \\ & -\; \lambda_{\mathrm{cost}}\,\mathrm{Cost}(h) \;-\; \lambda_{\mathrm{turn}}\,\mathrm{Turnover}(h) \;-\; \lambda_{\mathrm{switch}}\,d(h, h^-) \;-\; \lambda_{\mathrm{viol}}\,\mathrm{Viol}(h). \end{aligned} \tag{2}$$

$\mathrm{CEQ}_\gamma(R) = \mathbb{E}[R] - \frac{\gamma}{2}\mathrm{Var}(R)$ is the certainty equivalent, $\mathrm{CVaR}_{.95}$ the mean worst 5%, MDD maximum drawdown; penalties are nonnegative and normalized by frozen training statistics (Appendix G). Viol counts only *soft* violations (hard authorization and risk violations are eligibility filters), and "OOD utility" is held-out J_k under the same frozen constants. The gate tests only the additive return/cost/turnover/switch statistic $\tilde{J}$, which omits the path-risk terms, so a candidate can clear the gate yet reduce full J_k within eligibility. A *false promotion* is distinct from both a *negative-$\tilde{J}$ promotion* and a *full-J reduction* (Table 2). Let $h^\star_{1:K}$ denote the computable surrogate dynamic-oracle sequence that maximizes $\sum_k J_k(h^\star_k; h^\star_{k-1})$ over the enumerated feasible surrogate space. The reported normalized mean surrogate dynamic-oracle regret is $K^{-1}\sum_k \big[J_k(h^\star_k; h^\star_{k-1}) - J_k(h_k; h_{k-1})\big]$; each sequence pays switching cost against its own predecessor.

The promotion gate uses the following explicitly additive block statistic, with every term normalized by the same frozen training constants used in Eq. (2):

$$\begin{aligned} \tilde{J}_{k,b}(h; h^-) = \overline{R}^{\mathrm{net}}_{k,b}(h) - \lambda_{\mathrm{cost}}\mathrm{Cost}_{k,b}(h) - \lambda_{\mathrm{turn}}\mathrm{Turnover}_{k,b}(h) \\ - \lambda_{\mathrm{switch}} d(h, h^-) - \lambda_{\mathrm{viol}}\mathrm{Viol}_{k,b}(h). \end{aligned} \tag{3}$$

The CEQ variance term, CVaR, and MDD are excluded from $\tilde{J}$ as non-additive path functionals; they act through full-J reporting and hard eligibility filters. A candidate pays $d(c, h_k)$, while the incumbent's own comparison $\tilde{J}_{k,b}(h_k; h_k)$ pays zero incremental switching cost. False-promotion labels concern the held-out mean of this paired statistic. They say nothing about full utility or downstream harm.

4 CMAS: conservative meta-agent search

Typed worker and architecture interfaces. Workers are bounded specialists, each returning a typed tuple $o_{i,t} = (s_{i,t}, c_{i,t}, u_{i,t}, e_{i,t}, q_{i,t})$: asset scores, calibrated confidence, uncertainty, evidence identifiers, risk flags. The library holds trend, macro, value/carry, sentiment, liquidity/tail-risk, and skeptic agents. Before execution, an architecture *compiler* validates schemas, acyclicity, permissions, budgets, and evidence contracts, rejecting invalid diffs with machine-readable errors the proposer can learn from.

Archive-conditioned candidate generation. The archive stores tuples $(z_j, h_j, \tau_j, \hat{J}_j, \hat{\phi}_j, f_j)$, where $\hat{\phi}_j$ are component credits and f_j failure labels. Parents are retrieved by

$$S_j \;=\; \beta_z \,\mathrm{sim}(z_k, z_j) \;+\; \beta_J \,\hat{J}_j \;-\; \beta_u \,\hat{\sigma}_j \;+\; \beta_d \,\mathrm{novelty}(h_j), \tag{4}$$

with Gaussian regime kernel $\mathrm{sim}(z_k, z_j) = \exp(-\|z_k - z_j\|^2/2\ell^2)$ and $\mathrm{novelty}(h_j) = \min_{h' \in \text{deployed}} d(h_j, h')$ from Eq. (1). Records dominated on (utility, cost) are dropped first, so the archive acts as stepping stones [21]. The proposer sees only retrieved records and may apply one *auditable typed edit* per candidate, drawn from a closed set (workers, edges, one role field, budget, aggregation, verification; Table 3). Each proposal is a machine-readable diff with a hypothesis, mechanism, and anticipated failure mode (Figure 2, ❷). We exclude unrestricted code generation: free-form self-rewriting [21, 37] would change the candidate distribution and evaluation surface at once, confounding calibration attribution.

Paired evaluation and interventional credit assignment. Candidate and incumbent run on the *same* packets, seeds, market paths, and tool responses; pairing removes path luck and sharply reduces the variance of the improvement estimate. For component i (worker, edge, verification node, budget increment) we estimate Shapley-style interventional credit

$$\phi_i(h) \;=\; \mathbb{E}_\pi\Big[\, J\big(\mathrm{Pre}_\pi(i) \cup \{i\}\big) - J\big(\mathrm{Pre}_\pi(i)\big) \,\Big], \tag{5}$$

where π is a uniform random permutation of $C(h)$ and $\mathrm{Pre}_\pi(i)$ the components preceding i, estimated over a fixed permutation budget (Table 17) under common random numbers, counted in search compute. The removal operator is *compiler-valid*: dropping i re-routes orphaned edges to the aggregator, substitutes a neutral pass-through for a removed verifier, and holds the token budget fixed. Infeasible coalitions take the incumbent's projected utility, so $J(S)$ is always defined. Because the imputation is a design choice shrinking the credit of components whose removal breaks compilation, these are *compiler-defined Shapley-style credits* rather than canonical Shapley values (alternative rule: Appendix C). We do not report the Monte Carlo standard error of $\hat{\phi}$, so the evaluated outcome is the *rank* correlation; individual magnitudes should not be read quantitatively. Credits from round k enter the archive only at round $k{+}1$, never within the round that generated them. The pool $\mathcal{H}_k$ is therefore fixed before the shadow evidence that evaluates it, the pre-generation condition of Proposition 1. Attribution is *interventional*, so a persuasive worker that adds no counterfactual utility receives no credit. Appendix I positions this against observational textual backpropagation [25, 1].

Adversarial and negative-control evaluation. The evaluator applies five intervention families (dropout, corrupted workers, stale data, cost shocks, descriptor corruption; Appendix C). Two negative controls are central: under **clone-collapse** all workers share identical information and competence, so a sound meta-agent should contract [4]; permuting z_k (**shuffled-regime**) should make any conditional advantage vanish (Table 10).

Conservative promotion with calibrated simultaneous correction. For candidate c, let $D_{c,b} = \tilde{J}_{k,b}(c; h_k) - \tilde{J}_{k,b}(h_k; h_k)$ be the paired improvement of Eq. (3) on effective block b, where B_{eff} counts approximately independent blocks (circular block bootstrap [33], pre-selected length). At update k the archive and pool are frozen at cutoff t_k; the pre-designated window W_k is scored without revealing outcomes to the proposer, and labels and credits enter $\mathcal{D}_{k+1}$ only afterwards. Successive 26-week windows overlap at monthly cadence, so later rounds may reuse earlier observations, but no round adapts its pool to its own held-out W_k (Appendix G). The *implemented* gate applies an empirically calibrated, variance-penalized simultaneous score,

$$\mathrm{Score}_k(c) \;=\; \bar{D}_c \;-\; r_{k,c}, \qquad r_{k,c} \;=\; \sqrt{\frac{2\,\widehat{V}_{k,c}\,\log\!\big(2|\mathcal{H}_k|/\delta_k\big)}{B_{\text{eff}}}}, \tag{6}$$

treating the tightly concentrated paired deltas (empirical range ≈ 0.04 under common random numbers) as sub-Gaussian; the idealized Bernstein analysis appears in Appendix A. Here $\widehat{V}_{k,c}$ is the sample variance of the $D_{c,b}$. Because both $\widehat{V}_{k,c}$ and B_{eff} are estimated, Score_k is *not a confidence bound*; it is calibrated empirically (Figure 4; Remark 1, Appendix A). Among feasible candidates, the gate promotes the highest-scoring one (ties broken toward fewer workers) only if $\text{Score}_k(c) > \epsilon$; otherwise it retains the incumbent. Promotion first enters a simulated, limited-authority canary stage alongside the incumbent, with pre-specified rollback triggers.

The result below assumes a *valid* confidence radius, which Eq. (6) does not provide. We state it because it forces the design conditions.

Proposition 1 (Idealized conditional family-wise false-promotion result; *not* a property of Eq. (6))**.** *Fix the candidate pool $\mathcal{H}_k$ before the shadow sample is drawn ($\mathcal{H}_k$ measurable w.r.t. information prior to W_k), and let $r_{k,c}^{\text{valid}}$ be any valid $1-\delta_k/|\mathcal{H}_k|$ confidence radius for candidate c's mean block improvement. Define the* idealized *lower confidence bound* $\text{LCB}_k^{\text{ideal}}(c) = \bar{D}_c - r_{k,c}^{\text{valid}}$*. Then, with probability at least $1-\delta_k$, every candidate promoted under $\text{LCB}_k^{\text{ideal}}(c) > \epsilon$ has expected* replay *improvement exceeding ϵ; choosing $\delta_k = \delta/K$ (or $\delta_k = 6\delta/\pi^2 k^2$) bounds the probability of any false promotion across updates by δ.*

Departures. The proof is a standard simultaneous-concentration argument [26–28]. Eq. (6) breaks three of its premises (Appendix A): overlapping windows break cross-update disjointness; the variance-only penalty drops the Bernstein range correction and plugs in an *estimated* block count, so no finite-sample theorem applies; and the deltas are winsorized. We therefore call the gate *empirically calibrated* (Figure 4). A *sealed-protocol* rerun removes cross-update overlap and realizes 4.8% at the 5% target. It does not validate Proposition 1, which concerns *replay* improvement and says nothing about a misspecified objective (RQ7). A real deployment would inflate the margin to $\epsilon + \xi_k$; every number here uses $\xi_k = 0$, and nothing was run live.

Scoping. *Within* an update, δ is split across the $|\mathcal{H}_k|$ candidates and budgeted at $\delta = 0.05$ *per update*, which is the denominator that the 3.4% rate and Figure 4 report against. Sequence-level control would require $\delta_k = \delta/K$ or the anytime schedule, neither activated for Track A. The implied per-stream bound is *vacuous*, so RQ4's per-stream rate is covered by no 5% guarantee (Appendix A).

5 METAINVEST-BENCH: benchmarking meta-level adaptation

No single backtest can establish meta-agent quality, so METAINVEST-BENCH triangulates controlled capability (Track A), external validity (Track B), and adversarial governance (Track C) over a shared worker library, risk layer, and budget accounting (Figure 6).

Track A: controlled regime environment. The primary controlled benchmark is a hidden-state factor process, a Markov regime-switching model after [36]. The simulator specifies regime-dependent worker competence, so a surrogate-optimal architecture is computable. The regime evolves, a signal about next-period factors is emitted *before* the decision, and the return realizes afterward:

$$\begin{aligned} &s_{t+1} \sim P(\cdot \mid s_t), \qquad q_t,\ \nu_{i,t},\ \eta_{i,t} \in \mathcal{F}_t, \qquad \nu_{t+1},\ \epsilon_{t+1} \perp \mathcal{F}_t, \\ &u_{t+1} = A_{s_t} q_t + \nu_{t+1}, \qquad v_{i,t} = C_{i,s_t} q_t + \nu_{i,t}, \qquad f_{t+1} = \mu_{s_t} + u_{t+1}, \\ &y_{i,t} = H_{i,s_t} \mu_{s_t} + G_{i,s_t} v_{i,t} + \eta_{i,t}, \qquad r_{t+1} = B_{s_t} f_{t+1} + \epsilon_{t+1}. \end{aligned} \tag{7}$$

Here q_t is the pre-decision predictable component; worker-observation noises $\nu_{i,t}, \eta_{i,t}$ are $\mathcal{F}_t$-measurable at decision time, while the time-$(t{+}1)$ innovations $\nu_{t+1}, \epsilon_{t+1}$ are independent of $\mathcal{F}_t$ (all sampled independently given the regime). A_{s_t} and C_{i,s_t} induce the worker–future-shock correlation without exposing a future-indexed quantity; G_{i,s_t} sets signal strength; workers never see f_{t+1}, r_{t+1}, or oracle labels. The tensors H, G define hidden regime-dependent competence (some specialists useful only in some regimes; one persuasive but anti-correlated). Knowing them, the evaluator computes a *surrogate dynamic-oracle sequence* on the same schedule and budget, charging switching cost rather than taking the per-regime best. That makes regret, selection accuracy, and credit fidelity computable.

LLM workers in Track A. A worker is an LLM specialist: it receives a rendered *evidence packet* built from its noisy signal $y_{i,t}$, the context z_k, and role instructions, and returns $o_{i,t}$. The simulator sets signal quality via H, G, while the LLM supplies calibration, disagreement, and failure modes (over-confidence, susceptibility to the "persuasive" anti-correlated worker), so Track A tests

orchestration (worked packet in Appendix C). The evaluator observes s_t only to score the surrogate oracle; every method sees the same z_k (no-leakage conditions: Eq. (7)).

Track B: market replay; Track C: worker-fleet governance. Track B is a secondary, underpowered external-validity check on historical data, using date-controlled and counterfactual structural replay against documented LLM-backtest leakage [6, 7, 22] (Appendix C). Track C is a market-noise-free suite run under a budgeted meta-adversary, following the compromised-insider threat model of [11]. It tests whether the meta-agent detects a corrupted worker, isolates the affected path, picks a minimal repair, avoids leaking restricted context, and rolls back after a simulated governance incident.

6 Experimental protocol

Baselines. Seven Track A methods: a strong single agent with self-consistency under the same per-decision *deployment* budget cap; a fixed full committee; a rule-based regime router; a LinUCB bandit [29] over a 15-architecture menu (partially logged; Table 6); AFlow-style workflow search [14, 19]; MaAS/AutoMaAS-style adaptive selection [20, 9], the strongest adaptive competitor and what "best adaptive baseline" denotes throughout; and GEPA-style prompt-only optimization [1]. All are scored against the same surrogate dynamic oracle and share data packets, z_k, worker library, risk layer, model pool, and equal *search-inclusive budget caps* (Shapley and paired-replay compute counted), with no information advantage. These are our reimplementations, not official releases, so every superiority claim is scoped to *common-budget reimplementations under a shared risk layer*. Each method's defining search mechanism is retained; only what the shared environment forces is changed (retained/changed detail in Appendix C). Routing the best baseline's proposals through our gate still reaches 0.45 OOD utility against CMAS's 0.48 (Table 7), bounding the risk that a weak reimplementation drives the result.

Research questions. RQ1 regret and cost-adjusted utility; **RQ2** interventional worker competence vs. evaluator noise; **RQ3** transfer to unseen regimes, libraries, and backbones (Definition 1); **RQ4** gate calibration under serial dependence; **RQ5** sparsity when extra agents add no value; **RQ6** recovery from compromised workers; **RQ7** drift under a misspecified (return-only) evaluator.

Metrics. Track A: normalized mean surrogate dynamic-oracle regret; selection accuracy; *adaptation delay* (updates until post-shift regret returns within ϵ_{delay}, distinct from the margin ϵ); credit rank correlation. Track B: CEQ, Sharpe, MDD, CVaR, turnover, factor α. All tracks: cost, size, churn, violations, rollback precision/recall, and *false-promotion rate*. That rate is the fraction of per-update *decisions* promoting a candidate whose simulator-evaluated mean improvement is $\leq \epsilon$ (undefined where a method never switches).

Statistics. Nested walk-forward with an untouched final test; search is confined to training/validation intervals. The unit of replication is a *stream* ($96 = 12$ seeds $\times$ 8 schedules), with paired stream-level block-bootstrap CIs as the primary inference and schedule-clustered hierarchical intervals as a sensitivity analysis (Tables 1, 7–10). With only eight top-level schedules, clustered intervals are read cautiously and significance rests on the *paired* stream-level difference under common random numbers; clustering widens intervals without overturning the ablations (e.g. no-paired-evaluation OOD utility $[0.42, 0.48]$; Table 10). Paired CMAS-vs-best-baseline differences in regret and OOD utility exclude zero after Holm correction. The full pairwise matrix is omitted for space (seed separation in Appendix G).

This *simulator-evaluated mean improvement* is not directly observed: the evaluator estimates it from 50 Monte Carlo paths per candidate (SE < 0.004; ties within ± 0.005 of ϵ labeled false conservatively). Score form, block length, winsorization, ϵ, and pool size were frozen at validation before the untouched final test (Tables 15–16; configuration Table 17).

7 Results and diagnostic analyses

RQ1: Adaptation and cost-adjusted utility. Among adaptive methods, CMAS attains the lowest regret, adaptation delay, and false promotion, and the highest OOD utility and credit fidelity (Tables 1, 8). Versus the MaAS/AutoMaAS-style reimplementation, regret fell $0.116 \to 0.104$ (10.3%), OOD utility rose $0.44 \to 0.48$, credit correlation $0.32 \to 0.41$, deployment inference cost was 6.0% lower, and recovery took 2.7 updates vs. 3.1 (Figure 3). Paired differences exclude zero after Holm correction: regret -0.012 ($[-0.020, -0.004]$), OOD utility $+0.038$ ($[+0.010, +0.066]$). CMAS is nevertheless

Table 1: **Primary meta-level results** (Track A; 96 streams; ±95% stream-level block-bootstrap CIs). Search rows are our common-budget "-style" reimplementations. Cost is *deployment* inference only and is *not* a net saving: CMAS is heaviest on search-inclusive tokens (Table 9). "n/a" = never updates, so false promotion is undefined, not zero.

Method	Regret ↓	OOD utility ↑	False promo. ↓	Team size	Cost ↓
Single agent + SC	0.226±.022	0.28±.05	n/a	1.0	1.00
Fixed committee	0.183±.017	0.32±.03	n/a	6.0	2.40
Rule router	0.158±.016	0.38±.03	0.071	3.5	1.45
AFlow-style	0.137±.014	0.41±.05	0.093	4.7	2.05
MaAS/AutoMaAS-style	0.116±.013	0.44±.03	0.082	4.0	1.83
GEPA prompt-only	0.166±.019	0.36±.03	0.063	4.0	1.55
CMAS (ours)	**0.104±.011**	**0.48±.05**	**0.034**	3.4	1.72

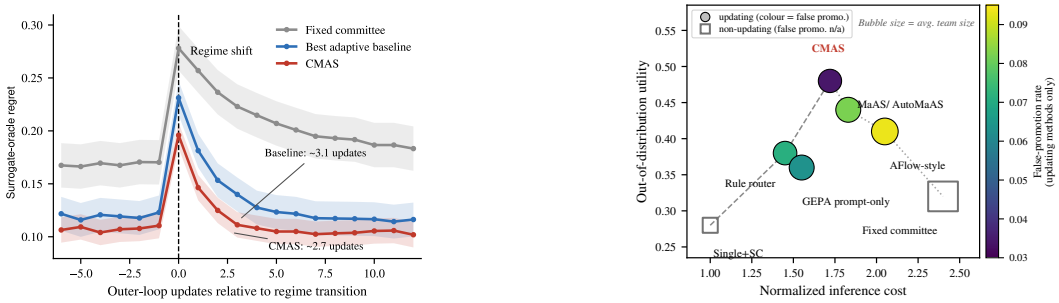

Figure 3: **Headline diagnostics.** *Left:* regret around evaluator-known regime transitions, 95% bands; CMAS recovers in ∼2.7 steps. *Right:* utility–cost frontier on *deployment-only* cost (CMAS is heaviest on search-inclusive tokens; Table 9). Bubble size is average team size. Color shows false-promotion rate for updating methods; static methods make no promotion decisions, so they are drawn as hollow grey squares outside the color scale.

heaviest in search-inclusive tokens (Table 9); the frontier claim is deployment-only. Track B is directionally consistent but underpowered (N=18; Table 12); the LinUCB menu also trails (Table 6).

RQ2: Interventional fidelity. Exact surrogate-oracle selection is 63.8% vs. 59.1%, credit correlation $\rho = 0.41$ vs. 0.32 (per-stream median 0.41, CI $[0.35, 0.47]$); selection rotates with regime, and counterfactually useless persuasive workers get low credit (Figure 5).

RQ3: Transfer. Utility retention is 83.5% on unseen regime mixtures (78.5% baseline) and 81.2% with two new specialists after four updates; backbone-switch loss is 7.8% vs. AFlow-style's 11.8% (Table 13). Held-out false promotion rises 3.4%→4.1%. Strict safety generalization is not established. Shuffled descriptors remove the advantage.

RQ4: Promotion calibration. False promotions are 3.4% of 2,112 updates and 5.86% of actual promotions; the incumbent is retained on 42%. These are *near-neutral* updates: 1.2% of decisions are full-J reductions and 0.4% are negative-$\tilde{J}$ promotions (Table 2); 11.4% of streams contain at least one false promotion. The sealed rerun yields 4.8%; removing simultaneous correction or pairing raises the rate to 8.4% or 6.1% (Table 10). No full-utility, canary-harm, or sequence-level guarantee is claimed.

RQ5 (sparsity). Under clone-collapse CMAS contracts 3.4→1.6 (4.6 without parsimony), a negative control rather than a general guarantee. **RQ6 (governance).** *Simulated* recovery from a budgeted adversary is 78%, rollback precision/recall 0.85/0.80 (Table 14). **RQ7 (Goodhart).** A return-only objective develops a growing gap between simulated canary and deployed-configuration utility, unlike the constrained evaluator (Figure 4); one misspecification is tested. Holding evaluator and gate fixed, proposal quality raises OOD utility 0.38→0.48 while false promotion stays 3.2–3.5% (Table 7), separating search quality from calibration.

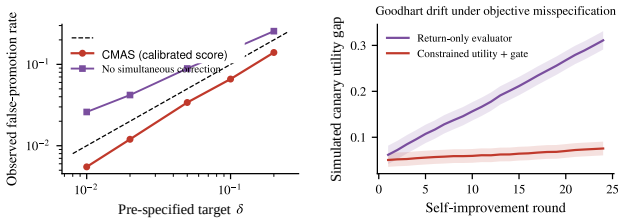

Figure 4: **Promotion calibration and Goodhart drift.** Held-out final test (2,112 decisions, 96 streams); all plotted parameters frozen at validation (Table 16), so neither panel is post-hoc fitted. *Left:* realized false promotion vs. pre-specified δ, observed below the nominal target on the tested distributions. This is empirical calibration; the score is not a valid confidence bound (Remark 1). *Right:* one objective-misspecification probe illustrating a failure mode.

8 Discussion, limitations, and responsible use

Evidence and limitations, ordered by how much each bounds the conclusions. *Statistical:* the implemented score falls outside Proposition 1 and certifies only $\tilde{J}$. *Benchmark:* the simulator is author-designed, so Track A shows internal validity only, and Track B is underpowered. *Baseline:* comparisons are common-budget "-style" reimplementations. *Scope:* one domain, an assumed informative z_k, and exact-architecture matching that can mark a near-optimal graph wrong, so regret is primary. *Deployment:* the canary stage and rollback are simulated; nothing was run live. Claims rest on surrogate-oracle regret, negative controls, calibrated promotion, and held-out transfer, not backtested return [6, 7].

Responsible-use statement. (Extended version in Appendix H.) *Intended use:* offline or shadow evaluation of *bounded* architecture changes, as decision support for authorized operators, not autonomous authorization of high-impact actions, and over a closed search space (Table 3). *Stakeholders and harms:* updates affect operators, users of the managed system, owners of the data and tools its workers touch, downstream decision recipients, and here investors. A wrong selection can install an overconfident worker, a correlated team, an evaluator-tuned architecture, or a verifier that blocks useful actions or certifies unsafe ones. *Controls and authority:* risk projection, authorization, accounting, and data cutoffs sit outside the LLM boundary (Figure 2); hard checks reject ineligible candidates before the gate; promotion is reversible. The owner sets objective, ϵ, limits, and triggers, while sign-off, halt, and audit ownership should rest with a reviewer independent of the proposer's authors. We implement the technical controls and specify the human ones. Those we did not evaluate. *Misuse and residual risk:* the loop is objective-agnostic and could optimize fleets for surveillance or manipulation, or search around a misspecified evaluator. A candidate can pass the gate and still reduce full J (1.2% here); sequence-level error is uncontrolled; simulated canary-stage harm preceded rollback in one stream (Appendix F); sign-off can decay into rubber-stamping.

9 Conclusion

CMAS treats meta-agent search as reversible sequential deployment, and METAINVEST-BENCH makes "did it adapt for the right reason?" measurable. Together they give a bounded, auditable framework for deciding when to replace an incumbent. That is not proof that meta-agent updates are safe.

References

[1] L. A. Agrawal, S. Tan, D. Soylu, et al. GEPA: Reflective prompt evolution can outperform reinforcement learning. *arXiv:2507.19457*, 2025.

[2] A. Ang, N. Azimbayev, and A. Kim. The self-driving portfolio: Agentic architecture for institutional asset management. *arXiv:2604.02279*, 2026.

[3] S. Hu, C. Lu, and J. Clune. Automated design of agentic systems. *arXiv:2408.08435*, 2024.

[4] P. Jwalapuram, H. Lin, C. Li, et al. The illusion of multi-agent advantage. *arXiv:2606.13003*, 2026.

[5] Y. Lee, R. Nair, Q. Zhang, K. Lee, O. Khattab, and C. Finn. Meta-Harness: End-to-end optimization of model harnesses. *arXiv:2603.28052*, 2026.

[6] X. Li, Y. Zeng, X. Xing, J. Xu, and X. Xu. Profit Mirage: Revisiting information leakage in LLM-based financial agents. *arXiv:2510.07920*, 2025.

[7] W. W. Li, M. Wang, and T. Ma. Summoning the oracle to slay it: Mitigating look-ahead bias in financial backtesting with large language models. *arXiv:2605.24564*, 2026.

[8] X. Lu et al. The Meta-Agent Challenge: Are current agents capable of autonomous agent development? *arXiv:2606.04455*, 2026.

[9] B. Ma, H. Li, Z. Hu, X. Gui, L. Liu, and S. Liu. AutoMaAS: Self-evolving multi-agent architecture search for large language models. *arXiv:2510.02669*, 2025.

[10] K. Miyazaki, T. Kawahara, S. Roberts, and S. Zohren. Toward expert investment teams: A multi-agent LLM system with fine-grained trading tasks. *arXiv:2602.23330*, 2026.

[11] J. Nöther, A. Singla, and G. Radanovic. MaMa: A game-theoretic approach for designing safe agentic systems. *arXiv:2602.04431*, 2026.

[12] C. Papadakis, A. Dimitriou, G. Filandrianos, et al. ATLAS: Adaptive trading with LLM agents through dynamic prompt optimization and multi-agent coordination. *arXiv:2510.15949*, 2025.

[13] L. Qian, X. Peng, Y. Wang, et al. When agents trade: Live multi-market trading benchmark for LLM agents. *arXiv:2510.11695*, 2025.

[14] Y. Shang, Y. Li, K. Zhao, L. Ma, J. Liu, F. Xu, and Y. Li. AgentSquare: Automatic LLM agent search in modular design space. In *ICLR*, 2025.

[15] Y. Xiao, E. Sun, D. Luo, and W. Wang. TradingAgents: Multi-agents LLM financial trading framework. *arXiv:2412.20138*, 2024.

[16] A. Xu and Y.-W. Tai. Meta-Agent: From task descriptions to verified multi-agent systems. *arXiv:2605.25233*, 2026.

[17] Y. Yu, H. Li, Z. Chen, et al. FinMem: A performance-enhanced LLM trading agent with layered memory and character design. *arXiv:2311.13743*, 2023.

[18] G. Zhang, Y. Yue, X. Sun, et al. G-Designer: Architecting multi-agent communication topologies via graph neural networks. *arXiv:2410.11782*, 2024.

[19] J. Zhang, J. Xiang, Z. Yu, et al. AFlow: Automating agentic workflow generation. In *ICLR*, 2025.

[20] G. Zhang, L. Niu, J. Fang, K. Wang, L. Bai, and X. Wang. Multi-agent architecture search via agentic supernet. *arXiv:2502.04180*, 2025.

[21] J. Zhang, S. Hu, C. Lu, R. Lange, and J. Clune. Darwin Gödel Machine: Open-ended evolution of self-improving agents. *arXiv:2505.22954*, 2025.

[22] T. Zhu, W. Zhao, R. Sun, et al. From Knowing to Doing: A Memory-Controlled Benchmark for LLM Trading Agents on Stock Markets. *arXiv:2605.28359*, 2026.

[23] M. Zhuge, W. Wang, L. Kirsch, F. Faccio, D. Khizbullin, and J. Schmidhuber. Language agents as optimizable graphs. In *ICML*, 2024.

[24] O. Khattab, A. Singhvi, P. Maheshwari, et al. DSPy: Compiling declarative language model calls into self-improving pipelines. *arXiv:2310.03714*, 2023.

[25] C.-A. Cheng, A. Nie, and A. Swaminathan. Trace is the next AutoDiff: Generative optimization with rich feedback, execution traces, and LLMs. In *NeurIPS*, 2024.

[26] P. S. Thomas, G. Theocharous, and M. Ghavamzadeh. High-confidence policy improvement. In *ICML*, 2015.

[27] R. Laroche, P. Trichelair, and R. Tachet des Combes. Safe policy improvement with baseline bootstrapping. In *ICML*, 2019.

[28] Y. Wu, R. Shariff, T. Lattimore, and Cs. Szepesvári. Conservative bandits. In *ICML*, 2016.

[29] L. Li, W. Chu, J. Langford, and R. E. Schapire. A contextual-bandit approach to personalized news article recommendation. In *WWW*, 2010.

[30] S. R. Howard, A. Ramdas, J. McAuliffe, and J. Sekhon. Time-uniform, nonparametric, nonasymptotic confidence sequences. *Annals of Statistics*, 49(2):1055–1080, 2021.

[31] O. Besbes, Y. Gur, and A. Zeevi. Stochastic multi-armed-bandit problem with non-stationary rewards. In *NeurIPS*, 2014.

[32] M. Herbster and M. K. Warmuth. Tracking the best expert. *Machine Learning*, 32(2):151–178, 1998.

[33] H. R. Künsch. The jackknife and the bootstrap for general stationary observations. *Annals of Statistics*, 17(3):1217–1241, 1989.

[34] L. S. Shapley. A value for n-person games. In *Contributions to the Theory of Games II*, 1953.

[35] S. M. Lundberg and S.-I. Lee. A unified approach to interpreting model predictions. In *NeurIPS*, 2017.

[36] J. D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989.

[37] Google DeepMind. AlphaEvolve: A coding agent for scientific and algorithmic discovery. Technical report, 2025.

[38] D. Precup, R. S. Sutton, and S. Singh. Eligibility traces for off-policy policy evaluation. In *ICML*, 2000.

[39] X. Luo, F. Wang, C. Hu, et al. Self-evolving agent harnesses via gated semantic quality-diversity. *arXiv:2607.13683*, 2026.

[40] B. Sengupta and J. Wang. HARBOR: Automated harness optimization. *arXiv:2604.20938*, 2026.

[41] S. Karten et al. Continual Harness: Online adaptation for self-improving foundation agents. *arXiv:2605.09998*, 2026.

[42] Y. Yue, X. Zhu, Y. Ma, et al. AutoRAS: Learning robust agentic systems with primitive representations. *arXiv:2606.21445*, 2026.

[43] Z. Liu, Z. Shi, Y. Sang, et al. Adaptive Auto-Harness: Sustained Self-Improvement for Agentic System Deployment on Open-Ended Task Streams. *arXiv:2606.01770*, 2026.

[44] C. Liu, Y. Guo, Y. Liu, et al. Auditing Agent Harness Safety. *arXiv:2605.14271*, 2026.

[45] X. Lin, Y. Liu, Y. Chen, et al. SafeHarness: Lifecycle-Integrated Security Architecture for LLM-based Agent Deployment. *arXiv:2604.13630*, 2026.

A Proof of Proposition 1 (false-promotion control)

For each candidate $c \in \mathcal{H}_k$, let $\Delta_c = \mathbb{E}[D_{c,b} \mid \mathcal{H}_k]$. By the assumed concentration bound,

$$\Pr\big(\Delta_c < \bar{D}_c - r_{k,c} \,\big|\, \mathcal{H}_k\big) \;\le\; \frac{\delta_k}{|\mathcal{H}_k|},$$

where, for this idealized statement, $r_{k,c}$ is the empirical-Bernstein radius (variance term of Eq. (6) plus the range correction $3M\log(2|\mathcal{H}_k|/\delta_k)/B_{\text{eff}}$), valid for improvements bounded in $[-M, M]$ with sample variance $\widehat{V}_{k,c}$ over B_{eff} approximately independent blocks [30]. The *implemented* score is the variance-only form of Eq. (6), which is not itself a valid radius (Remark 1). Writing $\mathrm{LCB}_k^{\text{ideal}}(c) = \bar{D}_c - r_{k,c}^{\text{valid}}$, a union bound over the finite candidate pool implies that, with probability at least $1-\delta_k$, $\Delta_c \ge \mathrm{LCB}_k^{\text{ideal}}(c)$ simultaneously for all candidates; hence any candidate promoted under $\mathrm{LCB}_k^{\text{ideal}}(c) > \epsilon$ satisfies $\Delta_c > \epsilon$. A second union bound over K updates with $\delta_k = \delta/K$ yields the global statement.

For serially dependent replay data, B_{eff} must come from a valid block construction. If the mixing or block assumptions fail, the guarantee does not apply. This condition is stress-tested empirically in Figure 4 (left); a heavy-tail/autocorrelation calibration suite is left to future work.

Table 2: **Terminology used for promotion outcomes.**

Term	Meaning
Certified statistic $\tilde{J}$	additive block statistic in Eq. (3); excludes variance/CVaR/MDD
Full utility J	complete scalarization incl. CVaR/MDD penalties; risk terms act as eligibility filters
Promotion	candidate replaces incumbent after $\mathrm{Score}_k(c) > \epsilon$ and eligibility pass
Canary stage	simulated post-promotion, limited-authority run alongside the incumbent with pre-specified rollback triggers
Rollback	reversion to the incumbent when a trigger fires
False promotion	promoted candidate whose simulator-evaluated mean improvement in $\tilde{J}$ is $\leq \epsilon$; near-neutral, and not necessarily a negative-$\tilde{J}$ promotion
Negative-$\tilde{J}$ promotion	promoted candidate whose simulator-evaluated mean improvement in $\tilde{J}$ is below zero (0.4% of all decisions)
Full-J reduction	promotion improving $\tilde{J}$ but reducing full J within eligibility (1.2% of decisions)
Simulated canary-stage harm	adverse outcome observed during the simulated canary run, before a rollback trigger fires; distinct from all three statistics above and not measured by the gate
Per-update rate	false promotions / all update decisions (3.4% overlapping protocol; 4.8% sealed)
Conditional rate	false promotions / actual promotions (5.86%)
Per-stream family-wise rate	streams with $\geq$1 false promotion / all streams (11.4%)

The configuration constant $M{=}0.15$ bounds the *raw, unpaired* utility scale; under common random numbers candidate and incumbent share the same shocks, so the paired deltas $D_{c,b}$ concentrate far more tightly, with empirical range ≈ 0.04. Empirically, the median B_{eff} is 5.8 and $M = 0.15$, and implemented penalties fall in $r_{k,c} \in [0.014, 0.032]$. Promoted candidates have $\bar{D}_c \in [0.028, 0.048]$, while rejected candidates peak near $\bar{D}_c \approx 0.015$.

Implemented score versus idealized radius. At these constants the Bernstein range correction alone evaluates to ≈ 0.45 and is therefore not the implemented quantity. Worked implemented-score example: $\widehat{V}_{k,c} = 1.8{\times}10^{-4}$, $B_{\text{eff}} = 5.8$, $\delta_k = 0.05$, $|\mathcal{H}_k| = 8$ ($\log 320 \approx 5.768$) give $r_{k,c} \approx 0.0189$; with $\bar{D}_c = 0.045$, $\mathrm{Score}_k(c) = 0.0261 > \epsilon = 0.02$, so the candidate is promoted. □

Remark 1 (idealized sub-Gaussian radius). If the block improvements were σ_c-sub-Gaussian with *known* proxy σ_c, the radius would reduce to $\sigma_c\sqrt{2\log(2|\mathcal{H}_k|/\delta_k)/B_{\text{eff}}}$. A plug-in variance in that form is *not* a valid radius; the empirical-Bernstein (or a self-normalized) bound is what licenses estimated variance. The implemented Score_k of Eq. (6) is exactly such a plug-in construction, which is why it is treated as empirically calibrated (Section 4).

Anytime-valid variant. Replacing $\delta_k = \delta/K$ by $\delta_k = 6\delta/\pi^2k^2$ gives $\sum_{k\geq 1}\delta_k = \delta$, so the family-wise guarantee holds for an *unbounded* sequence of updates without pre-committing to K; a time-uniform confidence sequence [30] yields the same conclusion under martingale dependence. The sequence of pools $\mathcal{H}_1, \ldots$ is itself adaptive (the proposer sees cross-round archive feedback). Formally, let $\mathcal{F}_{k-1} \supseteq \sigma(\mathcal{H}_1, \ldots, \mathcal{H}_k, W_1, \ldots, W_{k-1})$ be the information available before shadow window W_k is drawn; because each $\mathcal{H}_k$ is $\mathcal{F}_{k-1}$-measurable, each round's radius is valid *conditionally on* $\mathcal{F}_{k-1}$, and the union over rounds applies to the resulting conditional bounds. As stated in Section 4 (Scoping), the reported Track A runs budget δ per update and do not activate this schedule.

Idealized versus implemented. Three premises of Proposition 1 are not established for the overlapping evaluation protocol: (i) cross-update disjointness fails because 26-week windows at monthly cadence overlap, so proposals condition on outcomes contained in earlier evaluation windows; (ii) the implemented score drops the Bernstein range correction and plugs an estimated $\widehat{V}_{k,c}$ and effective block count B_{eff} into a variance-only form, which carries no finite-sample concentration theorem under dependence (Remark 1); and (iii) the paired deltas are winsorized at their validation-calibrated 99% quantile (binding in 1.2% of blocks, mainly volatility shocks), while the configuration constant $M{=}0.15$ is a raw-utility-scale bound, loose by roughly $4\times$ on the paired-delta scale.

Table 3: **Pre-specified search space.**

Dimension	Pre-specified values
Workers	any subset of six roles; per-role backbone is a separate model-assignment field (two open-weight options; twelve role–backbone instances)
Graph	DAG; ≤ 10 edges; no self-loops; permissions constrain edge types
Aggregation	parallel merge, confidence-weighted, sequential, debate, critic-then-aggregate
Models	one small open model, one stronger open model, one hosted reasoning model
Budget	$1\times$, $2\times$, $4\times$ token/tool budgets with a hard per-episode ceiling
Prompt mutation	one typed field per candidate; no hidden instruction changes
Verification	none, evidence validator, numerical recomputation, counterargument, dual-model critic
Promotion	calibrated simultaneous score, hard policy checks, parsimony tie-break, limited-authority canary stage, rollback snapshot

Gate sensitivity: block length 2/4/8 weeks gives realized FP 3.8/3.4/2.9%; replacing the circular-block B_{eff} estimator with autocorrelation-based ESS gives 3.5%. Removing the winsorization inflates $\widehat{V}_{k,c}$, widening the radius, lowering the promotion rate by $\sim$4%, and raising realized false promotion to 4.1%, still under target. The reported false-promotion rates under the overlapping evaluation protocol are empirical calibration.

The direct test of the idealized guarantee is a sealed-sample calibration study with fresh, non-overlapping evaluation paths per update. It realizes 4.8% false promotion at the nominal 5% target (vs. 3.4% under the overlapping protocol), which shows that calibration is not an artifact of cross-window overlap. Because the implemented score was used, this is not a direct finite-sample validation of Proposition 1. Ground-truth labels use 50 independent Monte Carlo paths per candidate (standard error < 0.004 normalized units); candidates within ± 0.005 of ϵ are conservatively classified as false promotions; propagated label uncertainty contributes < 0.2 pp to the 3.4% rate, and the sealed evaluation paths are disjoint from W_k and independent across updates. Under that reported per-update budget, the implied worst-case per-stream bound over $K \approx 22$ updates is vacuous ($K\delta > 1$; $\approx$0.68 even under cross-update independence). The observed 11.4% per-stream family-wise rate (RQ4) is far below what the configuration nominally permits, and above the per-update 5% target, which applies to a different denominator.

B Pre-specified search space

Table 3 fixes the search space; every candidate edit selects values from these dimensions, and nothing outside them was searched.

Chance baselines for selection accuracy (for the RQ2 figures of Section 7). Uniform guessing over the surrogate oracle's realized 6-architecture support gives 16.7%, the majority-architecture baseline 18.2% (full-space chance far lower, Appendix B).

C Simulator specification and required checks

Primary suite. *Five recurring macro-regime types* (Growth, Inflation, Stress, Recovery, Range; Figure 1) are arranged into *eight distinct latent-regime schedules* (transition sequences), so the unit of replication is a *stream* and 96 streams = 12 simulation seeds $\times$ 8 schedules. *Six worker roles* (trend, macro, value/carry, sentiment, liquidity/tail-risk, skeptic), each instantiated on two open-weight backbones drawn from the three-model pool (Appendix B), so the default library holds twelve specialists; the third, hosted model is reserved for the backbone-transfer test (RQ3). Each stream spans $\sim$122 weeks, including a 26-week *unscored burn-in* that fills the first shadow window; the remaining 96 scored weeks yield $K \approx 22$ monthly gate decisions, and streams include recurrent, abrupt, and gradual transitions. The Track C governance suite injects six incident scenarios per stream: the five single-attack families of Section 5 and Table 14 (instruction corruption, anti-correlated worker, stale/missing data, inference-cost shock, worker dropout), plus one multi-attack combination sampled per stream. That gives $96 \times 6 = 576$ governance episodes. The 5/10/25 bps cost levels apply to Track B replay, not to governance. The competence tensor includes specialists useful only jointly,

redundant specialists, a cheap weak agent, an expensive strong agent, and a persuasive anti-correlated agent.

Oracle utility is computed by exhaustive search for small graphs and mixed-integer optimization for the larger space. The oracle is a *surrogate*: it optimizes over the enumerated feasible subset of the pre-specified space (valid DAGs within budget) using the simulator's specified competence model. In that model, worker outputs enter through their analytic expected utility under the regime's factor distribution, marginalizing the simulator's noise terms; no actual LLM executions are involved. Prompt fields and verification nodes are fixed categorical choices in the MIP. Its optimality is exact for this surrogate model, not for arbitrary LLM behavior.

Track B detail. Historical evaluation uses two safeguards against documented LLM-backtest leakage [6, 7, 22]: *date-controlled replay* restricted to information released after each model's documented training cutoff (for Mixtral the cutoff is undisclosed and January 2024 is used, so genuine post-training-cutoff status cannot be guaranteed for that model), and *counterfactual structural replay* that fits a factor model, scrambles identities and calendars, and resamples factor/residual paths preserving cross-asset dependence. Agents never receive real tickers, dates, future documents, or web access. Performance is decomposed into market, style, and residual components so passive beta is not counted as skill, and publication lags and vintage data are frozen. The universe spans liquid equity, duration, credit, commodity, gold, real-estate, and international ETFs; decisions weekly, updates monthly (a quarterly variant is not evaluated here), costs 5/10/25 bps. Transfer is evaluated on a held-out universe and a different model family. Track B CEQ and factor α are reported in annualized percent throughout (e.g. CMAS CEQ $= 7.8\%$); the cross-window standard deviation over the 18 replay windows is ± 1.1 pp.

Pre-specified analysis protocol. We separate analyses *reported in this paper* from the benchmark's standing evaluation suite. Reported here: strong single agent under the same per-decision deployment budget cap (Table 1); the random/evolutionary proposer inside the CMAS evaluator+gate (Table 7); no/shuffled regime context and the component ablations (Table 10); identical-worker clones and different evaluator/worker family at test time; objective misspecification (return-only evaluator, RQ7); and search-inclusive cost (below). Not evaluated here, and therefore not claimed: Bayesian/change-point proposers, an oracle-context upper bound, candidate-pool scaling curves, and heavy-tail/autocorrelation gate calibration.

Baselines: full Track A roster (summarized in Section 6). We compare against seven methods in Track A: a strong single agent with self-consistency, evaluated under the same per-decision *deployment* budget cap (no search budget; Table 9); a fixed full committee; a rule-based regime router (rules in Appendix C); a LinUCB bandit [29] on a fixed 15-architecture menu over z_k; AFlow-style workflow search [14, 19]; MaAS/AutoMaAS-style adaptive selection [20, 9], which is the strongest adaptive competitor and is what "best adaptive baseline" denotes in every figure and table; and GEPA-style prompt-only optimization [1] on the fixed manager–worker graph. Track B adds risk parity. All Track A methods are scored against the same surrogate dynamic oracle and share data packets, the context z_k, worker library, risk layer, model pool, and equal *search-inclusive budget caps* (Shapley and paired-replay compute counted), with no information advantage. Realized usage differs under the cap: CMAS is highest in tokens though not wall-clock (Table 9, Appendix C). Because the environment mandates a shared risk layer and a common upper budget cap, published systems are reimplemented as "-style" variants (Appendix C); Table 7 bounds the resulting risk for the gate comparison, Section 8 for the search-quality comparison. Not run here: equal-weight/trend strategies, a human-designed manager–worker graph, and GPTSwarm/G-Designer-style topology search [23, 18]; random and evolutionary proposers *are* run inside the CMAS evaluator+gate (Table 7).

Baseline fidelity: retained vs. changed. Every comparison method is a common-interface reimplementation, not an official release, and all comparative claims in the paper are scoped accordingly. *Retained* for each method is its defining search mechanism: AFlow-style keeps graph-edit search with lookahead; MaAS/AutoMaAS-style keeps supernet optimization over discrete architectural choices; GEPA-style keeps reflective prompt mutation with tournament selection (operators, depths and population sizes below). *Changed*, and only where the shared environment forces it: (i) the identical 0.85M-token/update search-inclusive budget cap; (ii) the common typed worker interface $o_{i,t}$, so that every method emits the same schema; (iii) the immutable risk projection $\Pi_{\mathcal{C}_t}$ and authorization layer, which no method may bypass; (iv) the common three-model pool and worker library; and (v) the identical per-decision deployment budget cap. *Information available* is the same for all methods: the

Table 4: **Closest search and adaptation methods.** Entries summarize first-class mechanisms in the cited descriptions.

Method	Search object / proposal	Adaptation scope	Primary distinction from CMAS
ADAS / AFlow / AgentSquare [3, 19, 14]	executable programs or workflows	offline/static or per task	no repeated paired incumbent replacement
MaAS / AutoMaAS [20, 9]	architecture and resources	query-level	query routing rather than deployment-time promotion calibration
GEPA / Darwin Gödel Machine [1, 21]	prompts or executable code	iterative/open-ended	population fitness rather than paired replay and rollback
Adaptive Auto-Harness [43]	harness tree and stateful evolver	open-ended non-stationary streams	solve-time routing/human steering; no explicit update-level false-promotion target
Continual Harness / AutoRAS [41, 42]	online or robustness-aware harnesses	continual/robust	no explicit paired incumbent gate with measured false-promotion error
CMAS	typed architecture and archive-conditioned edits	sequential deployment	paired replay, calibrated score, switching cost, canary stage, rollback

pre-decision descriptor z_k, the same data packets, seeds, market paths and tool responses, and no access to the latent regime s_t or to evaluator oracle labels. Two measurements bound the residual risk that a weak reimplementation drives the headline result: routing the strongest baseline's proposals through the CMAS evaluator and gate still reaches 0.45 OOD utility against CMAS's 0.48 (Table 7), and the random/evolutionary proposer control in the same table separates proposal quality from gate calibration.

Baseline specifications. The rule-based regime router maps z_k to a hand-specified team via three conditional branches: *Calm* (volatility descriptor ≤ 0.15): Trend + Sentiment (2 specialists); *Volatile* (> 0.15): Tail-Risk + Skeptic + Macro + Value/carry (4); *Crisis* (> 0.30 or drawdown trigger): those four plus Liquidity (5). Weighted by realized regime frequencies across the 8 schedules, the expected team size is 3.48, the 3.5 of Table 1. For the search baselines: AFlow-style runs DFS/BFS over 12 graph-edit operators with lookahead depth 3; MaAS/AutoMaAS-style optimizes a supernet with 15 discrete architectural choices via REINFORCE with a moving-average baseline; GEPA-style runs reflective prompt mutation with population 20 and tournament selection. All are capped at the identical 0.85M-token/update search-inclusive budget as CMAS.

The infeasible-coalition imputation used in credit estimation shrinks the estimated credit of components whose removal breaks compilation toward the incumbent's level rather than measuring their true marginal value. The direction and size of the resulting bias depend on the repair rule, which we treat as a design choice (an alternative repair rule is evaluated in Appendix C). **Alternative credit repair.** A drop-and-penalize rule (infeasible coalitions scored 0) raises credit variance 14% without materially changing selection accuracy (63.1% vs. 63.8%). **Proposer-component ablations** (OOD utility; full proposer 0.48): removing archive retrieval drops it to 0.41, removing credit feedback to 0.44, and removing regime conditioning to 0.44, so retrieval carries the largest share of the proposer's edge.

The LinUCB menu bandit runs LinUCB [29] on context z_k over a fixed menu of 15 architectures, enumerated by taking the 5 most frequently Pareto-optimal architectures from historical (pre-test) runs and crossing them with the 3 pre-specified budget levels ($1\times, 2\times, 4\times$; Appendix B), with exploration parameter $\alpha = 1.0$.

Search-inclusive cost. The "Cost" column in Table 1 is *deployment-only*; Table 9 reports realized search-inclusive usage per update. CMAS is the heaviest token consumer, while its wall-clock is below the sequential graph-search baselines because paired replay parallelizes across candidates; the 6.0% cost edge is therefore a deployment-only figure.

Table 5: **Closest safety and deployment-control methods.** Related ingredients are acknowledged; CMAS contributes their deployment-level combination.

Method	First-class safety mechanism	Relation to CMAS
Gated semantic QD [39]	deterministic candidate-quality gate	no paired incumbent replay, rollback, or update-level calibration
HARBOR [40]	posterior chance-constrained safety check	no repeated incumbent replacement or rollback
HarnessAudit [44]	full-trajectory boundary and information-flow audit	complementary audit surface; does not optimize sequential promotion
SafeHarness [45]	lifecycle verification, privilege control, adaptive degradation, rollback	complementary execution security; no false-promotion estimand
CMAS	paired calibrated promotion, hard eligibility, canary stage, rollback	focuses on whether and when to replace the incumbent under drift

Table 6: **Partially logged LinUCB menu baseline.** Only the two metrics below were retained; it is therefore excluded from the fully populated primary comparison in Table 1.

Method	Normalized mean surrogate dynamic-oracle regret ↓	OOD utility ↑
LinUCB menu bandit	0.142	0.38

Table 7: **Proposer-value and gate decomposition.** Proposer varied; evaluator, credit, and gate fixed. "+ CMAS gate" routes that baseline's proposals through our gate. Hierarchical (schedule-clustered) 95% CIs on OOD utility, top to bottom: $[.46, .50]$, $[.43, .47]$, $[.40, .44]$, $[.36, .40]$. Proposer-component ablations: Appendix C.

Proposer / config. (fixed evaluator+gate)	OOD util. ↑	Dyn. regret ↓	False promo. ↓
LLM proposer (full CMAS)	**0.48**	**0.104**	0.034
Best adaptive baseline + CMAS gate	0.45	0.113	0.035
Evolutionary candidates	0.42	0.129	0.033
Random candidates	0.38	0.151	0.032

Table 8: **Secondary Track A metrics**, held out of the headline table for readability. Adaptation delay is right-censored at the eight-update event-study horizon; exact architecture matching is used for credit/selection diagnostics, while regret remains the primary oracle-relative metric.

Method	Adaptation delay ↓	Credit ρ ↑	Notes
Single agent + SC	7.3±.8	–	static
Fixed committee	6.2±.6	0.11±.06	static
Rule router	4.3±.5	0.18±.05	hand-routed
AFlow-style	3.5±.4	0.25±.05	search
MaAS/AutoMaAS-style	3.1±.3	0.32±.04	strongest adaptive baseline
GEPA prompt-only	4.6±.5	0.18±.05	fixed graph
CMAS	**2.7±.3**	**0.41±.05**	paired calibrated promotion

Table 9: **Realized search-inclusive cost per update** (mean over streams; budget cap 0.85M tokens/update for all search methods).

Method	Tokens/update	A100-hours/update
CMAS	850k	4.2
AFlow-style	820k	6.5
MaAS-style	780k	5.1
GEPA-style	450k	1.8
Single agent	50k	0.2

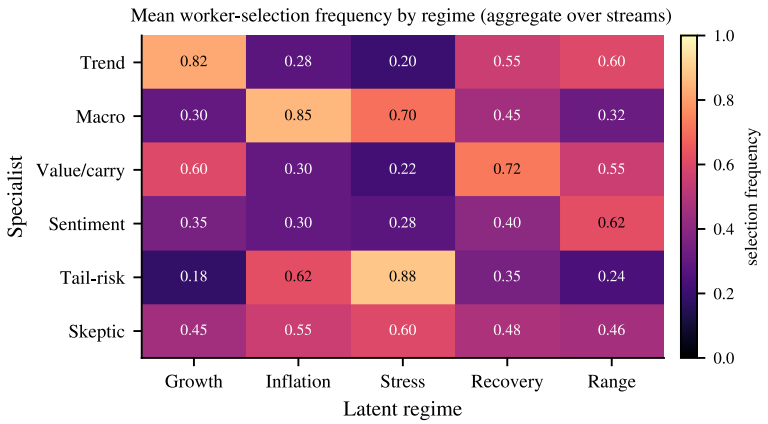

Figure 5: **Selection vs. regime (RQ2).** Mean worker-selection frequency by latent regime over the 96 streams; rotation with the hidden regime (macro triples from Growth to Inflation, tail-risk peaks in Stress) is evidence consistent with regime-conditioned selection, and the shuffled-regime control (Table 10) removes it.

Table 10: **Key ablations and controls** (Track A held-out streams). The rows above the rule are one-component-removed ablations; the two rows below the rule are *negative controls* run on full CMAS. *Clone team size* is each variant's mean team size under the clone-collapse control; for the clone-control row itself this is by construction the same quantity as Full CMAS's entry (1.6), and its OOD utility of 0.28 matches the single agent, as expected under identical workers. 95% hierarchical block-bootstrap CIs (clustered by regime schedule, seeds within) widen slightly relative to stream-level intervals (e.g., no-paired-evaluation OOD utility $[0.42, 0.48]$), but every utility and false-promotion degradation relative to Full CMAS remains assessed by paired stream-level differences with Holm correction (full matrix omitted for space); the marginal CIs shown are descriptive, and the "No parsimony" and "No adversarial tests" false-promotion deltas are near-neutral.

Variant	OOD utility ↑	False promo. ↓	Cost ↓	Clone team size
Full CMAS	**0.48**	**0.034**	1.72	**1.6**
No regime conditioning	0.44	0.040	1.70	3.2
No paired evaluation	0.45	0.061	1.77	3.4
No interventional credit	0.46	0.046	1.75	3.5
No parsimony penalty	0.47	0.037	2.02	4.6
No simultaneous correction	0.48	0.084	1.73	3.4
No adversarial tests	0.47	0.036	**1.69**	3.3
Shuffled regimes (control)	0.40	0.038	1.71	3.3
Identical-worker clones (control)	0.28	0.033	1.46	1.6

Leakage and reproducibility protocol. Freeze raw snapshots, release times, model versions, retrieval timestamps; lag features by actual publication availability with vintage macro data; date-controlled or counterfactual replay (never prompt-based memory suppression); strip real dates, event names, tickers in the controlled benchmark; disable web access and uncontrolled memory during replay; common random numbers for candidate/incumbent pairs; report factor exposures, residuals, costs, turnover, failed episodes; store diffs, traces, tool calls, approvals, rollbacks; pre-specify primary metrics, confidence methods, attack budgets, and final test intervals; release simulator, schema, prompts, and harness where licensing permits.

D Additional tables, figures, and ablations

Table 11 tabulates the RQ2 interventional-fidelity numbers quoted in the main text, and Table 13 the RQ3 transfer numbers as stream-level means. The RQ6 governance suite is decomposed by attack family in Table 14; Table 7 adds the proposer-value control (random/evolutionary candidates inside the CMAS evaluator+gate), and Figure 7 visualizes the Table 10 degradations.

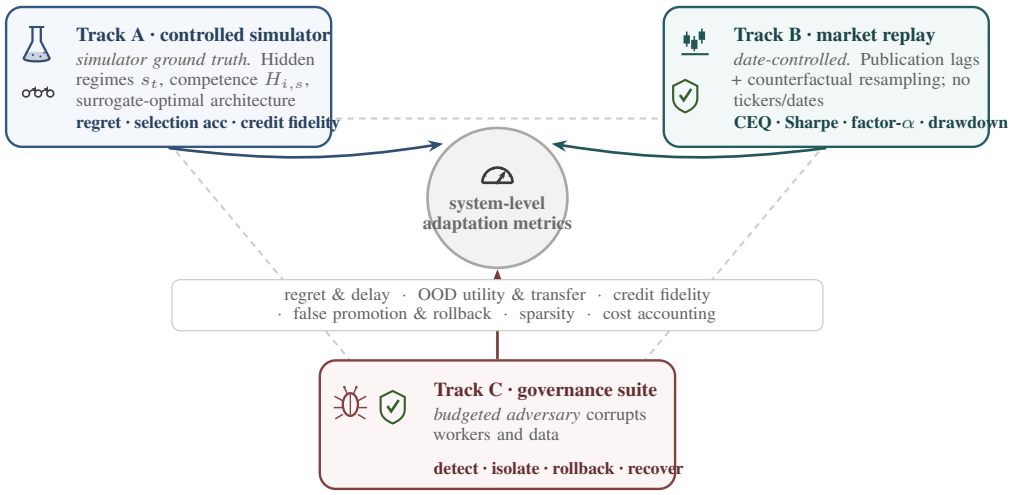

Figure 6: **MetaInvest-Bench triangulation.** Track A measures controlled meta-level capability against a computable surrogate dynamic oracle; Track B provides leakage-controlled external validity on markets (18 replay windows); Track C measures governance under a budgeted meta-adversary (576 episodes). All tracks feed a shared meta-level metric suite.

Table 11: **RQ2 interventional fidelity** (Track A). Architecture-selection accuracy is the fraction of outer-loop updates matching the architecture selected by the surrogate dynamic oracle at that update; credit ρ is the rank correlation between estimated interventional credit and surrogate-model component value.

Method	Selection accuracy ↑	Credit ρ ↑
Best adaptive baseline	59.1%	0.32
CMAS (ours)	**63.8%**	**0.41**

Table 12: **Track B external validity** ($N = 18$ date-controlled replay windows; net of 10 bps costs; "best adaptive baseline" is MaAS/AutoMaAS-style selection). This track tests whether Track A conclusions survive contact with market data under leakage controls; with $N = 18$ it is *underpowered to separate adaptive methods*, differences are not tested for significance, and no row should be read as evidence of deployable alpha. Factor α is the residual after market/style attribution. Worker-compromise recovery is a governance quantity and is reported only in Table 14 (Track C) to avoid conflating the two tracks. Because $N = 18$, we report the sign pattern only: CMAS CEQ dispersion is ± 1.1 pp (across-window standard deviation; mean 7.8%, min 6.8%, max 11.2%; the window distribution is right-skewed), so its worst windows overlap the baseline mean, and no cell should be read as a significance claim. Units: CEQ and factor α in annualized %; Sharpe dimensionless; MDD and CVaR %; turnover annualized fraction; violations fraction of decisions.

Method	CEQ % ↑	Sharpe ↑	MDD ↓	CVaR ↓	α % ↑	Turnover ↓	Violations ↓
Risk parity	4.2	0.68	19.2	2.24	0.2	0.18	0.000
Single agent + SC	5.1	0.84	17.8	2.05	0.9	0.31	0.034
Fixed committee	5.8	0.95	18.4	2.12	1.2	0.44	0.049
Rule router	6.7	1.04	15.6	1.88	1.6	0.32	0.026
Best adaptive basel.	7.4	1.12	14.5	1.79	2.0	0.41	0.031
CMAS (ours)	7.8	1.15	14.1	1.68	2.2	0.35	0.021

Table 13: **Meta-level transfer** (Track A held-out streams; retention = held-out / in-distribution utility). "Best baseline" is MaAS/AutoMaAS-style selection except in the backbone-switch row, where the relevant competitor is the search-based AFlow-style method (noted inline). The held-out-library baseline cell is "–" because MaAS-style selection cannot ingest new worker types without re-running search.

Held-out factor	CMAS	Best baseline	Relative
Unseen regime mixtures	83.5%	78.5%	+5.0 pp
Held-out worker library (2 new)	81.2%	–	recovers in 4 updates
Backbone switch (OOD util. cost)	7.8%	11.8% (AFlow-style)	1.5× smaller
Held-out violation rate (vs. 2.1% ID)	1.8%	–	safety clause met (no increase)
Held-out false promotion (vs. 3.4% ID)	4.1%	–	strict clause not met (+0.7 pp)

Table 14: **Governance suite by attack family** (Track C; 576 episodes nested in 96 streams; recovery within three updates). Aggregate Wilson 95% interval $[0.75, 0.81]$; per-family Wilson and cluster-by-stream intervals are omitted for space. Five single-attack families plus one multi-attack combination; the aggregate spans all six scenarios. Aggregate rollback precision/recall $0.85/0.80$.

Attack family	CMAS recovery	(k/n)	Best baseline	Committee
Instruction-corrupted worker	0.81	78/96	0.74	0.54
Anti-correlated worker	0.79	76/96	0.72	0.53
Stale/missing data	0.77	74/96	0.70	0.51
Inference-cost shock	0.76	73/96	0.69	0.50
Worker dropout	0.80	77/96	0.73	0.52
Multi-attack combination	0.75	72/96	0.68	0.49
Aggregate	0.78	450/576	0.71	0.52

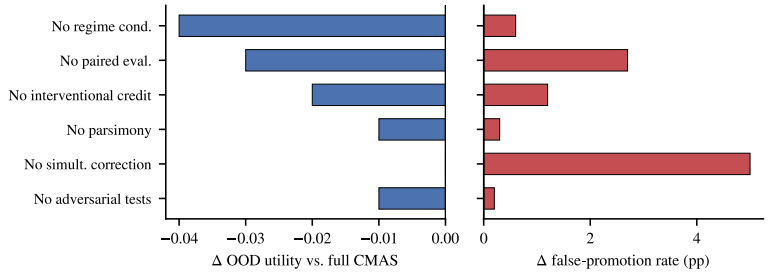

Figure 7: **Ablation sensitivity** (visualizes Table 10). Regime conditioning, paired evaluation, and interventional credit dominate utility; paired evaluation and the simultaneous correction dominate promotion safety; the parsimony penalty prevents architecture bloat under clone-collapse.

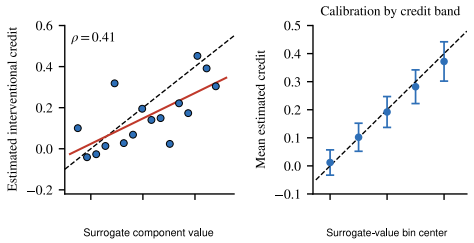

Figure 8: **Credit fidelity (RQ2).** Estimated interventional credit shows moderate positive rank agreement with surrogate-model component value (ρ=0.41, n=16 component types pooled across streams; regression line shown), strongest among high-credit components. The pooled scatter is illustrative only; the statistic we rely on is the *per-stream* distribution of ρ over the 96 streams: median $\rho = 0.41$, 95% CI $[0.35, 0.47]$, excluding zero.

Algorithm 1 CMAS at outer-loop update k

Require: incumbent h_k, context z_k, archive $\mathcal{D}_k$, shadow window W_k, risk policy $\mathcal{C}$

retrieve diverse parent architectures from $\mathcal{D}_k$ via Eq. (4)
meta-agent proposes typed candidates $\mathcal{H}_k$ with hypotheses and anticipated failure modes
compile candidates; reject invalid permissions, schemas, or graphs
for $c \in \mathcal{H}_k$ **do**
    run paired shadow replay against h_k (common random numbers)
    estimate the certified statistic $\tilde{J}$ and full utility J, component credits ϕ (Eq. (5)), uncertainty, and cost
    execute failure interventions (deployment-time eligibility checks; the benchmark-level negative controls of Section 4 are run by the evaluator, not per candidate)
end for
compute calibrated simultaneous scores of $\tilde{J}$ (Eq. (6)) and the Pareto-feasible set
if some feasible c has $\mathrm{Score}_k(c) > \epsilon$ **then**
    canary-deploy the highest-score feasible candidate (ties broken toward fewer workers); retain rollback snapshot
else
    retain h_k
end if
append traces, credits, failures, and deployment outcome to $\mathcal{D}_{k+1}$

E Example update trace

One end-to-end update from the promotion logs: the proposer submits the diff {add Tail-Risk, remove Sentiment} with hypothesis and expected mechanism; the compiler validates; a 26-week paired shadow replay yields $\mathrm{Score}_k = 0.028 > \epsilon = 0.02$ with all eligibility filters passing; the candidate is promoted and enters canary deployment; no rollback trigger fires over the subsequent four weeks. The full trace (prompts, per-block deltas, credit vector, trigger telemetry) is omitted for space.

F Failure modes and negative results

CMAS is not uniformly best. It underperforms the static rule router in 14 of 96 streams, predominantly during prolonged *Range* regimes, where switching costs and exploration penalties outweigh the benefit of adaptation and a fixed cheap policy is hard to beat. The gate is not infallible in either direction. In Stream 42, a tail-risk specialist was promoted immediately before a sudden volatility crush. The canary-deployment rollback trigger fired, but only after a 1.2% simulated canary-path drawdown was incurred, which is the concrete instance behind the pre-rollback-harm limitation in Section 8. Conversely, in Stream 11 a sparse architecture that the oracle ranks highly was rejected because its calibrated score missed the $\epsilon = 0.02$ margin by 0.002 under high shadow-window variance, an example of the conservatism cost quantified in RQ4 (42% retention). Per-stream loss maps and the full promotion/rejection logs for these cases are omitted for space.

G Reproducibility

This appendix is the source of record for the configuration that produced the reported results. A reproducibility package (simulator, environment interface, baseline adapters, evaluation scripts, seed manifests, frozen configurations, one full worker/aggregation prompt, and ten fully traced streams) is in preparation for public release. Licensed datasets and hosted-model credentials cannot be redistributed, and scripts and manifests document the access procedure. Failed worker calls are retried up to three times (retries counted in the search-inclusive budget); after three failures the worker returns a neutral no-opinion tuple, and wasted tokens (2.1% of search tokens on average) are logged. Settings are mirrored in `config_frozen.json`.

Denominator transparency. The main text reports both false promotions/all updates (3.4%) and false promotions/actual promotions (5.86%), plus full-J reduction (1.2%), negative-$\tilde{J}$ promotion (0.4%), and streams with at least one false promotion (11.4%). The manuscript source contains

Table 15: **Information chronology at update k.** This distinguishes the overlapping protocol from the sealed rerun and makes the proposal/evaluation boundary explicit.

Step	Time / object	Information boundary
1	archive cutoff t_k	build $\mathcal{D}_k$ only from outcomes and credits released before t_k
2	candidate generation	freeze $\mathcal{H}_k$; no access to outcomes in the designated W_k
3	paired replay W_k	run candidate and incumbent on identical packets, shocks, seeds, and tool responses
4	gate decision	apply eligibility filters and Score_k; retain incumbent or enter simulated canary stage
5	canary interval	candidate has limited simulated authority alongside the incumbent until the next update; rollback restores the incumbent architecture and state snapshot
6	archive release	replay labels, canary outcomes, and credits become available only to $\mathcal{D}_{k+1}$
Overlap	overlapping protocol	successive 26-week W_k overlap at monthly cadence; a later pool can use earlier released outcomes, but no pool sees its own W_k before freezing
Sealed check	calibration rerun	fresh, non-overlapping paths per update; same frozen score and margin, no retuning
Final test	reported results	untouched during search/calibration; only final frozen configuration is scored

Table 16: **Calibration and evaluation stages.** The source records the roles and frozen choices below; it does not contain a public preregistration timestamp or anonymous commit hash, so none is claimed.

Stage	Used for	Frozen before subsequent stage
Development/training	proposer/archive development; utility normalizers	search space, typed edit interface, normalizers
Validation/calibration	score form and sensitivity analysis	block length 4, 99th-percentile winsorization, $\epsilon = 0.02$, candidate pool 8, rollback thresholds
Final Track A test	96 streams; 2,112 update decisions	no tuning or candidate-pool adaptation to final outcomes
Sealed rerun	non-overlapping empirical calibration check	reuses the same score, margin, and eligibility rules

rounded rates rather than raw integer event logs; rounded-rate equivalents are approximately 72 false promotions among 2,112 updates, approximately 1,225 promotions, approximately 25 full-J reductions, approximately 8 negative-$\tilde{J}$ promotions, and approximately 11 affected streams. These are explicitly approximate; exact integer logs are part of the planned reproducibility release.

Metric details. Adaptation delay is right-censored at eight post-transition updates; a static method can therefore receive a finite capped delay. Selection accuracy uses exact equality with the surrogate-optimal architecture, so functionally equivalent or near-optimal graphs count as mismatches; normalized regret is the primary metric. OOD normalization uses frozen training minima/maxima without clipping test values.

Per-decision Wilson intervals (secondary). Treating the $2{,}112$ update-level promotion decisions as independent, the Wilson interval for observed false promotion is $[2.7, 4.3]\%$ at the 5% target vs. $[7.3, 9.6]\%$ without the simultaneous correction; these intervals ignore within-stream and within-schedule clustering and are secondary to the schedule-clustered estimates reported in the main text.

Seed separation and planned reproducibility release (referenced from Section 6). Seeds for sampling, simulation, and proposal are separated; an open-weight track reduces API dependence; failed runs, costs, prompts, versions, and traces are logged in that package.

Configuration (stated in-paper). Architecture space and promotion rule are pre-specified in Appendix B: DAG with ≤ 10 edges; any subset of six worker roles (per-role backbone a separate field);

Table 17: **Configuration used for the reported results.** These values are the source of record for the reported results. Proposer and red-team/critic model: Llama-3-70B-Instruct (cutoff Dec 2023; temperature 0.2, top-p 0.9); candidate utilities are computed by deterministic replay code. Workers: Mixtral-8x7B-Instruct-v0.1 (vendor cutoff undisclosed; replay restricted to post–January 2024 data; temperature 0.4) and Llama-3-8B-Instruct (cutoff Mar 2023; temperature 0.3). Hosted reasoning model (backbone-transfer test): Claude 3 Opus via API (cutoff Aug 2023; temperature 0.1). Seeds: 42 (sampling), 1337 (simulation), 2024 (proposal).

Symbol	Meaning	Value
$\lvert\mathcal{H}_k\rvert$	candidate-pool size per update	8
K	outer-loop updates per stream	≈22 (monthly)
ϵ	promotion margin (normalized utility)	0.02
δ	family-wise false-promotion budget (per update)	0.05
W_k	paired replay-window length	26 weeks
block length	circular block bootstrap	4 weeks
Shapley budget	adaptive (screen all; full on top-3)	8 screen; 24×3
γ	CEQ risk aversion	2.0
$\lambda_{\text{cvar}}, \lambda_{\text{dd}}, \lambda_{\text{cost}}$	risk/cost weights	0.5, 0.3, 0.2
$\lambda_{\text{switch}}, \lambda_{\text{turn}}, \lambda_{\text{viol}}$	switch/turnover/violation weights	0.15, 0.1, 1.0
retrieved parents	per proposal	4
β	retrieval weights (sim, utility, unc., novelty); Pareto prefilter first	0.4, 0.3, 0.2, 0.1
ℓ	regime-kernel bandwidth (standardized z_k)	0.35
$(\omega_V, \omega_E, \omega_b, \omega_\rho)$	worker/graph/budget/typed-field distance weights, Eq. (1)	0.3, 0.4, 0.1, 0.2
ϵ_{delay}	adaptation-delay tolerance (Section 6)	0.05
rollback thresholds	canary MDD / CVaR violation / utility drift	>8%; > 2× baseline; >0.03
LinUCB	menu size; exploration α	15; 1.0
M	raw-utility-scale range bound (paired deltas concentrate to ≈0.04)	0.15
B_{eff}	effective independent blocks (median)	5.8
streams	seeds × schedules	96 (12 × 8)
gate decisions	total across streams	2,112
tokens/update	staged Shapley + 8B routing	0.85M
total compute	tokens; aggregate A100-hours (4.2/update, Table 9)	1.79B; ≈8,870
ξ_k	replay-to-deployment margin (Track A)	0
worker model pool	small / strong open-weight; hosted transfer model	3 models (proposer/critic separate)

budgets $\{1\times, 2\times, 4\times\}$; five aggregation and five verification options; one typed field edited per candidate.

Statistics: 96 streams (12 seeds × 8 regime schedules), ∼122-week streams (26-week unscored burn-in; 96 scored weekly decisions) with monthly updates ($K \approx 22$), 95% circular block-bootstrap CIs with pre-selected block length, Holm correction *across the seven inferential Table 1 comparisons* (average team size is descriptive and excluded) for CMAS vs. best baseline, common random numbers for candidate/incumbent pairs, separated seeds for sampling/simulation/proposal. We apply the budget *per update*: $\delta = 0.05$ is the family-wise target across the $|\mathcal{H}_k| = 8$ candidates evaluated in one update (per-candidate level $\delta/|\mathcal{H}_k| \approx 0.006$). The anytime-valid schedule of Appendix A would additionally control error across the open-ended update sequence, but, as scoped in Section 4, it was *not activated* for the reported Track A runs; the per-stream family-wise rate of RQ4 should be read accordingly. The reported 3.4% is the per-update false-promotion rate over 2,112 decisions ($96 \times K$), directly comparable to the 5% per-update target (Figure 4).

Normalization of Eq. (2). Every penalty term is mapped to a common scale before the λ weights are applied: per-term min–max normalization over the training interval, with the normalizers frozen at deployment (no test-time re-normalization). As a local scale illustration only, holding all other normalized penalty terms fixed, OOD utility $0.44 \to 0.48$ maps to approximately $4.12\% \to 4.58\%$ annualized CEQ; this is not an independent economic-performance claim.

Worked packet (Track A). A "macro" worker at week t receives a packet of the form {`context`: z_k descriptor; `signal`: rendered $y_{i,t}$ as bucketed factor tilts with a stated noise level; `history`: last-4-week own scores and realized outcomes; `role`: "estimate 4-week factor direction and size, flag regime risk"} and returns $o_{i,t}$ = {`scores`: per-asset tilts; `confidence`: 0–1 calibrated; `uncertainty`: interval; `evidence_ids`; `risk_flags`}. The simulator sets signal quality via H, G; the LLM contributes calibration, aggregation-relevant disagreement, and failure modes.

Cost accounting. Per-method realized usage is summarized in Table 9. Structurally, CMAS's search (paired replay + Shapley) is its dominant cost, and its total exceeds every compared method's, so break-even against the Table 1 deployment-only edge depends on the per-update deployment horizon.

H Extended responsible use and governance

This appendix expands the statement in Section 8. Nothing here is an evaluated result: it specifies the governance envelope the method is designed to sit inside, and separates what we implemented from what a real operator would still have to supply.

Implemented in this work. The deterministic risk projection $\Pi_{\mathcal{C}_t}$, authorization layer, accounting, and data cutoffs are outside the LLM boundary and cannot be edited by any proposal (Figure 2). Hard authorization, evidence, leakage, and risk checks are *eligibility filters* applied before the statistical gate, so a candidate failing them is rejected regardless of measured utility. The candidate space is closed (Table 3) and every promotion is logged as a typed diff with hypothesis, replay evidence, credit vector, and outcome.

Specified but not evaluated. Human sign-off, emergency halt, incident review, and audit-log ownership are design requirements, not measured quantities. We recommend that the party setting the objective, margin ϵ, risk limits, and rollback triggers be distinct from the party building the proposer, and that promotion authority and halt authority rest with a reviewer independent of both. The canary stage in this paper is *simulated*: a real limited-authority deployment would need separate approval, position and blast-radius limits, and an independent monitoring path.

Failure modes of the oversight itself. Oversight can fail in four ways. Automation bias: a reviewer presented with a passing score and a plausible hypothesis may approve without inspecting the diff. Rubber-stamping turns routine sign-off into a formality, and is likelier the higher the promotion rate, which is one reason the 42% incumbent-retention rate matters. Evaluator misspecification: if the objective is wrong, the gate faithfully certifies the wrong thing, which is exactly what the return-only probe of RQ7 demonstrates. Specification error in the immutable layer completes the list; constraints protect only if correctly written, and nothing in CMAS detects a wrong constraint.

Misuse. A meta-agent search loop is objective-agnostic. The same machinery could optimize agent fleets for surveillance, manipulation, or exploitative coordination; scale a harmful objective more cheaply; search around an imperfect safety evaluator, since a gate that screens on a measured statistic will also find architectures that satisfy the statistic without satisfying its intent; or produce architectures whose complexity obscures accountability. The parsimony tie-break and the bounded edit set mitigate the last of these, but do not address the first three, which are governance rather than technical problems.

Residual risk. A candidate can pass the gate and reduce full J (1.2% of decisions here). Sequence-level error is uncontrolled and the per-stream bound implied by the reported configuration is vacuous (Appendix A). Canary-stage harm can occur before a trigger fires, as in Stream 42 (Appendix F). The simulator may omit correlated failure modes that matter in deployment. Track B is underpowered and is not evidence of deployable alpha. Brokerage connectivity and live-trading validation are excluded entirely, and a real deployment would require validation by parties independent of the developers.

I Extended related work and positioning

On the optimization side, DSPy compiles declarative pipelines and GEPA evolves prompts reflectively, while Trace generalizes backpropagation-style credit to computational graphs [24, 1, 25]; Meta-Harness searches full harness code with prior traces [5].

Concurrent 2026 work studies gated harness evolution separating LLM proposals from deterministic crediting [39], automated harness optimization [40], online adaptation for self-improving agents [41], robust agentic-system learning [42], and sustained construction/routing over open-ended non-stationary streams [43]; CMAS differs in targeting repeated candidate–incumbent *replacement* under non-stationarity, with paired incumbent comparison, rollback, and false-promotion measurement first-class (Tables 4–5).

Safe improvement, conservative exploration, and anytime-valid inference. Our gate belongs to the tradition of high-confidence policy improvement [26], safe policy improvement with baseline bootstrapping [27], and conservative bandits/exploration bounding performance below a baseline while learning [28]; the surrogate dynamic-oracle-regret metric follows non-stationary and switching-cost online learning [31, 32]; our paired shadow replay is related to off-policy evaluation with confidence intervals [38]; it pairs candidate and incumbent under common random numbers rather than reweighting a logged policy. Paired replay is available because the simulator and counterfactual replay permit *re-execution* of candidate and incumbent on the same shocks, the step that importance-weighted, doubly robust, and fitted-Q estimators are designed to make unnecessary when only logged data are available. Three properties are new relative to this literature: (a) the action space is *typed architectures* with structure and a distance (Eq. (1)), not a fixed policy class; (b) candidates are generated *adaptively by an LLM proposer*, breaking i.i.d. candidate assumptions, which is why the guarantee is conditional on the pool; and (c) evaluation uses LLM execution traces with interventional credit. For open-ended deployment, δ/K can be replaced by an anytime-valid schedule or confidence sequence [30] (Appendix A).

Safety, verification, and meta-level failure. MaMa casts safe agent-system design as a game against a worker-compromising meta-adversary [11]; HarnessAudit audits full trajectories for boundary and information-flow violations [44], SafeHarness integrates lifecycle verification and rollback [45], and other work verifies agent graphs and attributes failures structurally [16]. The Meta-Agent Challenge argues that task benchmarks do not measure the ability to *develop* agent systems [8]. CMAS operationalizes these via false-promotion measurement, worker-compromise tests, trace-level credit, canary deployment, and rollback.

Financial agents and evaluation validity. Financial systems include memory-based agents [17], trading-firm committees [15], prompt adaptation from market feedback [12], fine-grained expert teams [10], and self-driving portfolios [2]. Leakage-controlled studies show LLM backtests are inflated by parametric memory and passive factor exposure [6, 7, 13, 22]; we therefore make the simulator the primary controlled ground-truth test and treat market replay as *external validity* only.