NN-Based Transistor Modeling for Ultra-Fast SPICE Simulation and Standard Cell Library Characterization

Rodion Novkin, and Hussam Amrouch, *Member, IEEE*

arXiv:2610.31024v1 [cs.ET] 22 Oct 2026

***Abstract*—Neural network (NN)-based transistor models have proven to be a promising solution to accelerate device modeling. Although these models demonstrated remarkable speedup in circuit simulations, they were not rigorously tested in complex tasks within standard EDA tool flows for chip design, which require very high transistor model precision to provide accurate results and converge properly. To investigate their full potential and promise, we develop NN-based transistor models that accurately predict the drain current and various charges for a $3\,\mathrm{nm}$ nanosheet, and employ the models for a standard cell library characterization. This enables us to evaluate the accuracy and speedup under a very wide range of SPICE simulations, benchmarking our models against an industry-standard implementation. Since standard cell library characterization involves thousands of SPICE simulations, it serves as a great benchmark for transistor model quality and consistency. To further evaluate model accuracy, we also perform analog- and digital-level simulations for various complex circuits. Our experiments reveal that NN-based models are a valid alternative to existing transistor compact models, capable of producing accurate delay and power estimates with a much higher speed and convergence time compared to the conventional Verilog-A-based industry-standard (BSIM-CMG) transistor compact model. Our NNs achieve sub-0.11% and 0.76% errors for delay and total power evaluation on a wide range of circuits, while taking up to 8 times less time during the standard cell library characterization. Additionally, we provide comprehensive information and guidance on NN accuracy, unveiling the relationship between NN size and its effect on circuit simulation precision.**

***Index Terms*—Neural network, transistor model, circuit simulation, library characterization, static timing analysis**

I. Introduction

ACCURATE transistor models are essential for the semiconductor industry since they are required for a wide variety of complex tasks, such as design and simulation of various circuits. Berkeley short-channel IGFET (BSIM) [1], [2] is the industry-standard model that is commonly used to perform these tasks. BSIM relies on physics-based equations to accurately reproduce underlying transistor behavior, enabling the model to achieve excellent performance in precise modeling of device characteristics. However, developing an analytical solution to accurately describe the underlying physics of device operation is often very time-consuming and requires substantial domain expertise since transistor behavior becomes more and more complex in advanced sub-5nm nodes and emerging technologies, where multiple new effects and challenges attributed to advanced transistor structures start to take place. Furthermore, all new effects must be seamlessly integrated into an existing BSIM code and built upon it without altering the model's prior behavior, which unavoidably introduces additional fitting parameters alongside the hundreds of already existing ones and greatly increases device simulation time. Finally, if the underlying physics were not properly captured, an analytical-based model may not even provide the correct output for the necessary transistor characteristics, leading to faulty circuit behavior.

NN-based compact transistor modeling received significant attention over the recent years due to NN's promising ability to learn complex data patterns. NNs can approximate any function [3] without the need to learn underlying physical processes and can effectively create a completely standalone, compact model for specific working conditions dictated by the range of data and input parameters used for NN training. A simplistic nature of matrix multiplication, which is prevalent in most NN architectures, provides compact models with the potential to greatly reduce the complexity of the model, thus lowering the required time for circuit simulation, which can be reduced even further with GPU-acceleration of NN-based transistor models and SPICE solvers.

An outstanding performance of a standalone NN-based compact modeling approach was successfully demonstrated in multiple prior works [4]–[9] on a wide range of transistor types and technology. Even though multi-layer perceptron (MLP) was a common choice for device characteristics modeling, other architectures, e.g, Kolmogorov-Arnold network (KAN) [10], graph neural network (GNN) [11], and a semantic segmentation network [12], were also explored. In addition to pure NN-based compact modeling, multiple hybrid approaches that utilize NN to correct the output of a low-fidelity equation-based model were also presented [13]–[15]. Finally, knowledge-based neural networks (KNBB), which allow NNs to learn quantities through empirical physics-based equations and even enable some post-training calibration, were introduced and expanded upon very recently [16]–[18].

Most of the NN-based transistor models demonstrated high accuracy and great fitting capabilities for important transistor characteristics like drain and gate currents (I), charges (Q), and pairing capacitances (C), which resulted in accurate circuit simulations of different complexity [4], [5], [8], [9], [15], [18]. However, most circuit simulations using NN-based transistor models demonstrated just a visual match with an

Rodion Novkin and Hussam Amrouch are with the Technical University of Munich; TUM School of Computation, Information and Technology; Chair of AI Processor Design; Munich Institute of Robotics and Machine Intelligence, Munich, Germany (e-mail: {rodion.novkin,amrouch}@tum.de).

industry-standard reference model on a very limited set of circuits. Additionally, most research works were focused exclusively on nMOS while keeping pMOS model source and implementation undisclosed. Therefore, a comprehensive and statistically supported evaluation of NN-based transistor model performance for circuit simulation was never performed. Since standard cell library characterization involves thousands of SPICE simulations across a wide variety of circuits, capacitive loads, and transition slews, it provides a strong benchmark for transistor model robustness, accuracy, and convergence. Standard cell library characterization and static timing analysis (STA) represent vital and fundamental SPICE-based tasks in the EDA flow, which require accurate transistor models for precise timing and power evaluation. Machine learning was previously used for aging analysis, standard cell library characterization, and power/delay estimation [19]–[25]. However, NNs were never employed as a model on a transistor level. Thus, it is not clear whether they are capable of producing precise delay and power values for complex circuit simulations and how NN accuracy and size affect the precision of the delay and power estimation later.

Our main contributions within this paper are as follows:
(1) We demonstrate, for the first time, the speed and accuracy of NN-based transistor models for standard cell library characterization, as well as power and timing signoff estimation, benchmarking nMOS and pMOS NN-based transistor models against industry-standard implementation of the 3nm nanosheet using standard EDA tools and a variety of circuits.
(2) We perform standard cell library characterization employing 71 combinational and sequential cells in 1701 SPICE simulations containing 7×7 slew-capacitive load combinations each, achieving 0.06% and 0.33% average relative errors for delay and cell power estimation among all simulations.
(3) Using the generated libraries, we evaluate timing and power signoff for 12 industrial circuits with the number of gates ranging from 70 to 104286, where NN-based transistor models achieve sub-0.11% and 0.76% relative errors for delay and total power estimation, respectively.
(4) We perform a SPICE circuit simulation for a full 16-bit divider and 1001-stage NAND ring oscillator (RO) consisting of 5270 and 4004 transistors, demonstrating a simulation speedup of up to 10x and 43x, respectively.

II. Methodology for NN-based Transistor Modeling

In the following, we describe the selected device and dataset generation process. NN architecture, used for transistor modeling, and the training algorithm (including conversion, loss, and error functions) are presented in Section II-B. Simulation setup for the standard cell library characterization and circuit analysis is detailed in Section II-C. A complete experimental workflow is outlined in Fig. 1.

A. Dataset Generation

For systematic evaluation of the NN-based transistor models, a $3\,\mathrm{nm}$ nanosheet with a regular threshold voltage was selected from an open-source PDK [26]. For a single device,

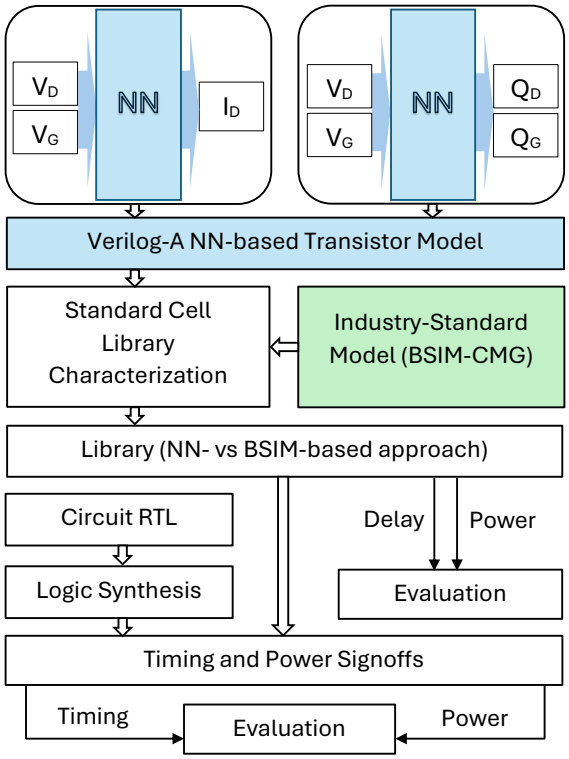

Fig. 1: NN-based transistor model is implemented in Verilog-A out of pre-trained networks. Then, all transistor models are employed for the standard cell library characterization, and the output libraries are used for timing and power signoff estimation and compared to the industry-standard BSIM-CMG.

we model the drain current I_D, drain charge Q_D, and gate charge Q_G characteristics by varying the drain voltage V_D and gate voltage V_G. The bulk charge Q_B is not considered since it is extremely small compared to other charges and does not contribute any noticeable pairing capacitances for the selected transistor. The source charge is not modeled and calculated as $Q_S = -(Q_D + Q_G)$ to maintain charge conservation. Training data is generated with a voltage step of $5\,\mathrm{mV}$ on a range from $-0.2\,\mathrm{V}$ to $1\,\mathrm{V}$ for nMOS and from $-1\,\mathrm{V}$ to $0.2\,\mathrm{V}$ for pMOS. Opposite voltage regions (negative for nMOS and positive for pMOS) are mainly generated to account for overshoots that are common in circuit simulations. Additionally, to learn the behavior of derivatives correctly with NN, it is necessary to extend the boundaries of the data so that they slightly exceed the intended operating range. In the scope of this work, the intended operating range, including overshoots, is from $-0.1\,\mathrm{V}$ to $0.9\,\mathrm{V}$ for nMOS and from $-0.9\,\mathrm{V}$ to $0.1\,\mathrm{V}$ for pMOS. Test data was generated with a slight shift in a range from $-0.097\,\mathrm{V}$ to $0.893\,\mathrm{V}$ for nMOS and in a range from $-0.893\,\mathrm{V}$ to $0.097\,\mathrm{V}$ for pMOS with the same $5\,\mathrm{mV}$ step. In addition to drain current and charge values, their first and second order derivatives with respect to V_D and V_G are also calculated since they are required for training. All data was obtained with a commercial SPICE simulator using BSIM-CMG for gate-all-around transistors (version 111.2.1) [27].

B. Neural Network Architecture and Training Algorithm

Since our goal is to explore the relation between NN model size and its precision for circuit simulation, we employ two different MLPs with the same architecture, composed of input, output, and two hidden layers connected with the $Tahn$ activation function. A smaller NN features a hidden dimension of 8 neurons, while a bigger NN has 16. In total, we employ two NNs for each size: the drain current NN output y_I produces a magnitude of the drain current using drain voltage V_D and gate voltage V_G as inputs; and the output of the charge NN y_Q produces scaled values for Q_D and Q_G with the same inputs. The final values for I_D and both charges are calculated with the following conversion functions:

$$I_D = V_D \cdot e^{y_I} \tag{1}$$

$$Q_{D,G} = y_Q \cdot \lambda \tag{2}$$

Similar to other works [4], [5], [7], the sign of the drain current is dictated by the sign of V_D, while the signs of the charges have to be learned since they cannot be easily determined. The charge scaling factor λ in Equation 2 helps to keep NN output values within a reasonable range and is based on the maximum charge value of a transistor.

All networks are trained using the PyTorch framework [28] and weighted loss functions with various coefficients that adjust the influence of different loss function terms:

$$\begin{aligned} L_I = a \cdot Er(I_D) + Er(y_I) + b \cdot Er(g_m) + \\ + c \cdot Er(g_{DS}) + d \cdot Er(g'_m) + e \cdot Er(g'_{DS}) \end{aligned} \tag{3}$$

$$\begin{aligned} L_Q = a \cdot Er(Q) + b \cdot Er(\frac{\partial Q}{\partial V_G}) + \\ + c \cdot Er(\frac{\partial Q}{\partial V_D}) + d \cdot Er(\frac{\partial^2 Q}{\partial V_D^2}) + e \cdot Er(\frac{\partial^2 Q}{\partial V_G^2}) \end{aligned} \tag{4}$$

$$Er(Y) = \frac{1}{p} \sum_{i=1}^{p} |y_i - \hat{y_i}| \tag{5}$$

We employ mean absolute error (MAE) as an error function to avoid error spikes when training NNs with scaled charges and to keep drain current errors reasonably high since they are not scaled. Training is conducted in an iterative manner, meaning that different loss terms are enabled sequentially.

All NNs are trained with an Adam optimizer [29] and an initial learning rate of 0.001, which is adjusted with the learning rate scheduler during training. For the drain current I_D, **1)** the training starts with just fitting the margin error $L_I = Er(y_I)$. **2)** In the next step, coefficients a and b are enabled to ensure better fitting in the saturation region and to add the transconductance component. **3)** Then, an output conductance loss term c is introduced. Finally, **4)** both second-order derivatives are enabled with coefficients d, e, and after that **5)** NNs are additionally trained with a complete loss function and reduced initial learning rate.

Since the precision of capacitances is more relevant than the precision of charges themselves, **1)** charge NN training starts with the loss $L_Q = Er(\frac{\partial Q}{\partial V_G})$. **2)** In the next step, another first-order partial derivative is enabled with c followed by **3)** an introduction of the second-order derivatives with coefficients d and e. Then, **4)** an actual charge component is added with coefficient a. Finally, **5)** NNs are additionally trained with a complete loss function and reduced initial learning rate, similar to drain current networks.

Outside of conversion functions shown in Equation 1 and 2, neither input nor output data is normalized in any other way to avoid additional computation during gradient calculation, which is necessary for the loss functions. For the same reason, no normalization layers are used in the NN architecture. The values of the loss term coefficients are selected to ensure a reasonably equal contribution from different loss terms. However, this rule is not very strict, and the coefficients are rather flexible when used with the state-of-the-art NN optimizers. Depending on the task's focus, more important terms can be given a stronger influence. However, to remain helpful, the loss contribution of any term should not be hundreds or thousands of times smaller than that from other terms.

TABLE I: Circuits used for timing and power signoff estimation.

Abbreviation	Description	Cells
Norm	Normalization for Fractional Input	70
CRC	Cyclic Redundancy Check	150
Adder128	128-bit Adder	194
InvSqrtFP	Floating-Point Reciprocal of Square Root	1271
ExpFP	Floating-Point Exponential	4214
SinCosFP	Floating-Point Sine or Cosine	5017
MAC32	32-bit Multiply-Accumulate	6512
Mult64	64-bit Multiplier	19898
Exp64	64-bit Exponential	65914
Div64	64-bit Divider	72304
RISC-V	64-bit RISC-V	6529
MAC-array	16x16 8-bit MAC-Array for AI acceleration	104286

C. Simulation Setup for Standard Cell Library Characterization and Circuit Analysis

In our library, we employ 71 combinatorial and sequential standard cells, with each cell characterized at rising and falling edges. The function and behavior of each cell are described by the measurement information about timing, power, leakage, and area. During the standard cell library characterization process, each cell is simulated and analyzed for 7 signal slews and 7 load capacitances, resulting in 49 total combinations. Full standard cell library characterization employs 1701 post-layout SPICE netlists, including parasitics. The library is essentially a database containing all data from these simulations. It has a well-defined format and serves as an input for further circuit analysis, synthesis, and timing/power estimation.

In order to evaluate how model accuracy scales from the individual cell level to the circuit level, we perform STA for 12 different circuits, described in Table I. To cover a wide range of applications and data types, we selected 10 industrial arithmetic circuits, RISC-V (Rocket) [30], and an in-house AI-accelerator, which consists of 16x16 systolic MAC-arrays. All circuits were synthesized with a supply voltage V_{dd} = 0.8 V, resulting in a range of 70 to 104286 standard cells.

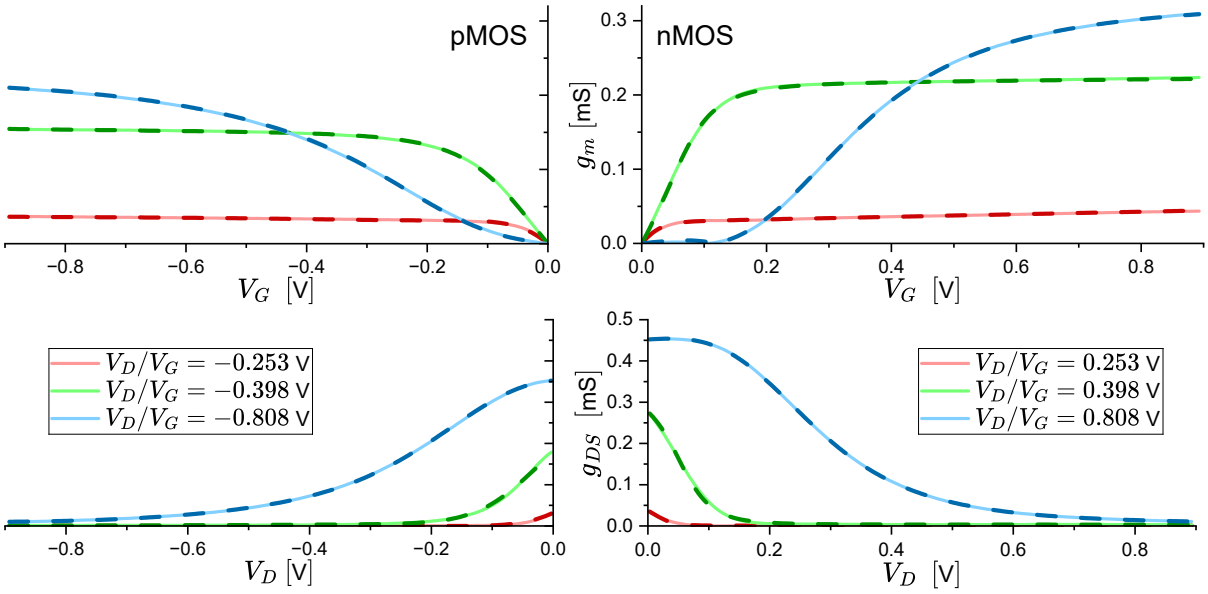

Fig. 2: NN8 model exhibits an excellent fit for nMOS and pMOS transconductance and output conductance. Lines are NN8 outputs, similar-colored dashed lines - BSIM-CMG industry-standard compact model. Either V_{G} or V_{D} is fixed, depending on the plot. Y-scale is the same for nMOS (right) and pMOS (left) plots.

III. Accuracy Evaluation at the Transistor Level

Before conducting any experiments and performing measurements, all trained neural networks are coded into a standalone transistor model written in Verilog-A. Similar to other works (e.g. [8]), we unroll all the loops within matrix multiplications. All data, shown in this section, is obtained from commercial EDA tools using different transistor models: industry-standard BSIM-CMG and our two NN-based transistor models (NN16, NN8), where the number indicates the number of neurons in the hidden layers.

We perform a comprehensive evaluation of the accuracy for all NN-based transistor models using a wide range of characteristics necessary for precise circuit simulations. All visual plots for characteristics are shown for both nMOS and pMOS on the same y-scale to highlight their asymmetry. To avoid clutter, NN16 was omitted from the visual evaluation, as it produces even more precise plots than NN8. Fig. 2 demonstrates an excellent fit for transconductance g_{m} and output conductance g_{DS} produced by the 8-neuron NN-based model. The conductance behavior is captured very well throughout the entire testing voltage range. Similar quality can be observed in all capacitance plots shown in Fig. 3, where NN8 is able to capture the behavior of various capacitances with great accuracy in the most important operating regions, i.e., near-zero voltage and near the intended supply voltage equal to $0.8\,\mathrm{V}$. However, in the region near $0.4\,\mathrm{V}$, NN8 struggles to capture charge derivative behavior properly, and this observation will be important later.

Although visual estimation is extremely beneficial for a general evaluation of NN-based transistor model performance, it does not provide a full picture of how accurate every NN is. Thus, we employ two quantitative metrics to evaluate absolute and relative accuracy of all important characteristics - mean

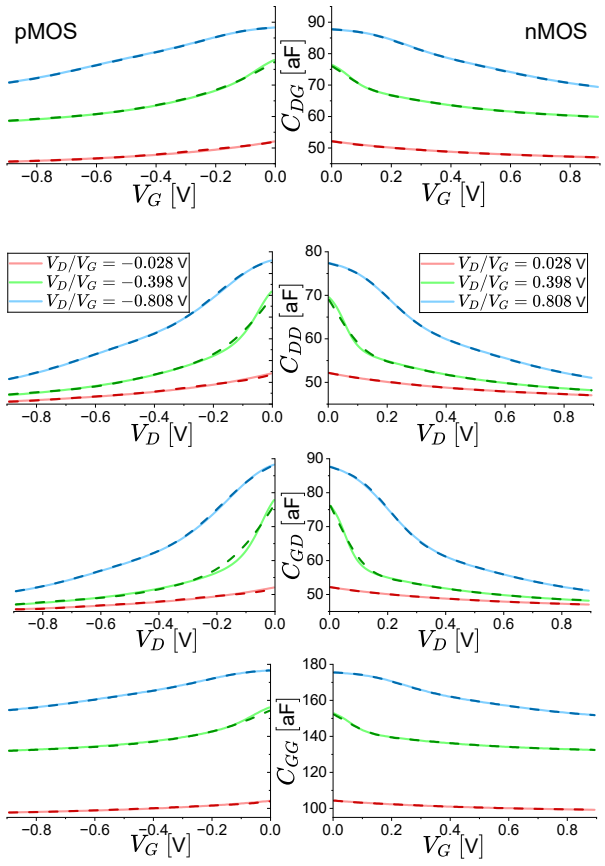

Fig. 3: NN8 model demonstrates a good fit for capacitances in necessary regions. Lines are NN8 outputs, similar-colored dashed lines - BSIM-CMG industry-standard compact model. Either V_{G} or V_{D} is fixed, depending on the plot. Y-scale is the same for nMOS (right) and pMOS (left) plots.

TABLE II: Prediction error of trained neural networks.

Variable	nMOS				pMOS			
	NN16		NN8		NN16		NN8	
	RMSE	MAPE	RMSE	MAPE	RMSE	MAPE	RMSE	MAPE
I_{D}	0.026 µA	0.32 %	0.056 µA	0.73 %	0.037 µA	0.51 %	0.072 µA	0.59 %
g_m	0.290 µS	2.54 %	0.914 µS	8.84 %	0.369 µS	0.89 %	0.553 µS	1.03 %
g_{DS}	0.677 µS	4.28 %	1.256 µS	8.54 %	0.737 µS	4.42 %	1.077 µS	7.54 %
Q_{D}	0.007 aC	0.22 %	0.011 aC	0.29 %	0.041 aC	3.84 %	0.036 aC	1.93 %
Q_{G}	0.012 aC	0.24 %	0.021 aC	0.24 %	0.091 aC	0.69 %	0.070 aC	0.42 %
Q_{S}	0.008 aC	0.09 %	0.014 aC	0.12 %	0.051 aC	0.19 %	0.043 aC	0.53 %
C_{DG}	0.052 aF	0.04 %	0.111 aF	0.10 %	0.102 aF	0.13 %	0.297 aF	0.31 %
C_{DD}	0.098 aF	0.06 %	0.170 aF	0.12 %	0.080 aF	0.08 %	0.312 aF	0.27 %
C_{GG}	0.089 aF	0.03 %	0.177 aF	0.07 %	0.214 aF	0.13 %	0.571 aF	0.27 %
C_{GD}	0.137 aF	0.08 %	0.211 aF	0.15 %	0.132 aF	0.12 %	0.523 aF	0.43 %
C_{GS}	0.138 aF	0.09 %	0.224 aF	0.17 %	0.256 aF	0.27 %	0.606 aF	0.52 %
C_{SG}	0.050 aF	0.04 %	0.106 aF	0.10 %	0.122 aF	0.13 %	0.373 aF	0.39 %
C_{SS}	0.093 aF	0.07 %	0.181 aF	0.16 %	0.141 aF	0.16 %	0.439 aF	0.45 %

absolute percentage error (MAPE) and root mean square error (RMSE), calculated as:

$$\mathrm{MAPE} = \frac{\sum_{i=1}^{n} |y_i - \hat{y}_i|}{\sum_{i=1}^{n} |y_i|} \tag{6}$$

$$\mathrm{RMSE} = \sqrt{\frac{\sum_{i=1}^{n} (y_i - \hat{y}_i)^2}{n}} \tag{7}$$

Table II provides an overview of the NN-based model test accuracy compared to the BSIM-CMG, which was used to generate train and test data. All capacitances not listed in the table are several magnitudes lower than the listed ones and do not significantly affect circuit simulations. As expected, bigger NNs demonstrate better results than their smaller counterparts for all important characteristics. Since the pMOS data pattern is different from the nMOS, same-size NNs did not achieve the same accuracy for different transistor types. While RMSE values are easily interpretable, MAPE ones have more depth and may require additional context. The majority of data are on an NN-friendly scale, meaning the ground truth and predicted values are neither extremely low or high on a linear scale. In contrast, g_{m} and g_{DS} values can get really close to zero, resulting in rather high average relative errors (e.g., a 100% relative error from $\hat{y}_i = 0.002$ and $y_i = 0.001$ significantly contributes to the average value, especially when the average is very low). Thus, MAPE for these characteristics is noticeably higher than in other cases, which skews the perception of the results but barely contributes to an actual performance of the NN-based models, since RMSE for these cases is still low.

Going from 16 neurons to 8, nMOS NNs lose roughly half of the absolute precision. pMOS NNs experience a similar accuracy drop for the drain current, but capacitance absolute errors increase by almost three times. The comparison between nMOS and pMOS NNs shows that pMOS NN16 achieves a similar accuracy as nMOS NN8 which indicates that the data pattern for pMOS capacitances is harder to learn. Therefore, for the same NN capacity (i.e., the number of weights), the fitting capabilities will differ for different devices. When NN capacity is further reduced by lowering the number of neurons in hidden layers, the pMOS charge model is more impacted and experiences a greater drop in capacitance accuracy, as it was already struggling to fit the data on the same level as nMOS. In contrast to a regular function regression, NNs in transistor models not only have to learn the function output, but its derivatives with respect to multiple variables as well. Therefore, if NN does not have a sufficient number of parameters to learn everything, the optimizer can select an easier-to-learn loss term to optimize during the training process, which may result in a situation demonstrated by pMOS charge NN - for NN8 charge accuracy goes slightly up compared to NN16, but capacitance accuracy decreases significantly. Although it is possible to balance loss terms with coefficients to mitigate this effect and potentially enhance capacitance accuracy, we will keep the charge NN8 for pMOS unchanged to assess how larger errors impact circuit simulation and further analysis.

IV. Analog Level Circuit Simulation

In this section, we perform analog level analysis using SPICE simulations for two circuits, comparing precision and speedup of NN-based models with respect to BSIM-CMG.

A. Accuracy of Analog Circuits

We perform a comprehensive evaluation of the transistor model performance for a 16-bit divider circuit designed with 5270 nMOS and pMOS transistors in total. The simulation is conducted with the same netlist and various transistor models in SPICE using two setups representing the precision of the circuit simulation. The default simulation setup involves low tolerances, as well as a predefined optimization profile that trades off the simulation speed for a much higher accuracy. This setup is commonly used when a very precise simulation is required for an accurate delay estimation. The second simulation setup is a speed-oriented one, featuring higher tolerances and a predefined optimization profile that trades off simulation accuracy for higher speed and faster convergence. The latter setup will be referred to as a low-precision simulation (LPS) in all further experiments. LPS is primarily used to demonstrate

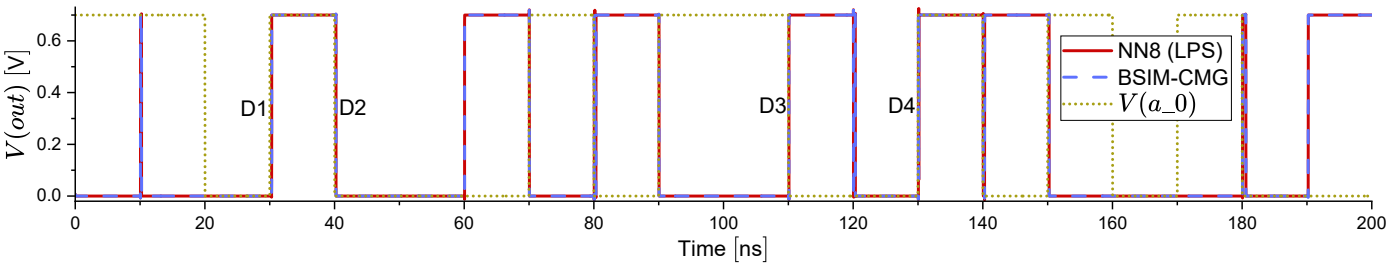

Fig. 4: 16-bit divider circuit containing 5270 transistors. Simulation with NN8 model and a low-precision simulation (LPS) setup shows an excellent visual match with BSIM-CMG simulated with the default precision setup. D1-D4 indicate the points for delay measurement for the output voltage at *quotient_0* node with respect to the input *a_0*. V_{dd} = 0.7 V.

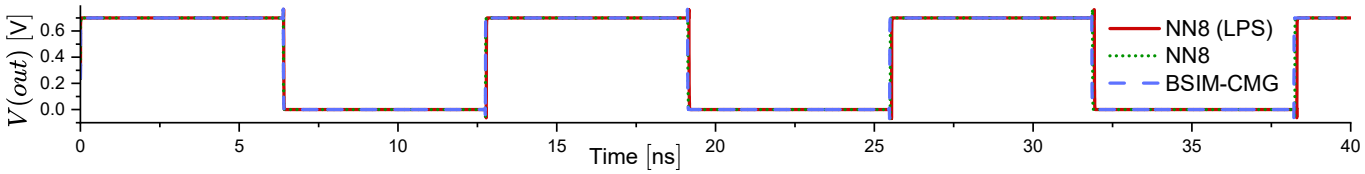

Fig. 5: 1001-stage NAND ring oscillator (RO) circuit containing 4004 transistors. Simulation with NN8 model and a low-precision simulation (LPS) setup shows minor visual deviations with BSIM-CMG and NN8 simulated with the default precision setup. V_{dd} = 0.7 V.

TABLE III: Delay measurements for the 16-bit divider for delay points shown in Fig. 4 and 1001-stage NAND ring oscillator (RO) period measurements. Relative error (RE) is calculated with respect to BSIM-CMG.

Divider delay point	BSIM-CMG Delay [ps]	NN16 Delay [ps]	NN16 RE [%]	NN8 Delay [ps]	NN8 RE [%]	NN8 (LPS) Delay [ps]	NN8 (LPS) RE [%]
D1	296.00	295.93	-0.024	295.66	-0.115	295.37	-0.213
D2	216.00	216.07	+0.032	216.06	+0.028	215.93	-0.032
D3	91.40	91.41	+0.011	91.29	-0.120	91.19	-0.230
D4	51.10	51.07	-0.059	50.99	-0.215	50.77	-0.646
RO period [ps]	12720	12732	+0.094	12734	+0.110	12752	+0.252

its immense benefit for NN-based modeling speedup, as well as unavoidable limitations for precise circuit simulation, which come with the low precision.

Fig. 4 demonstrates that even with the smallest and the least accurate NN8 model, combined with the speed-oriented simulation setup, the output of the circuit shows a perfect visual match with the default high precision simulation using an industry-standard BSIM-CMG model. NN-based model accurately represents all fronts and overshoots for an entire 200 ns time-frame and can provide all necessary visual information on circuit behavior. However, accurate delay estimation cannot be performed on the shown in Fig. 4 output since the difference is way beyond the time scale of the plot. Thus, we measure delay values in several points at the $V_{dd}/2$ level, indicated by $D1-D4$ in Fig. 4. Table III demonstrates that the delay produced by NN16 is extremely close to the BSIM-CMG model and exhibits only some negligible deviation due to numerical differences and minor fitting errors. NN8 demonstrates a slightly higher error on average, and the output becomes less stable with relative errors ranging from 0.028% to 0.215% for the divider circuit. Finally, NN8 (LPS) exhibits significantly higher errors compared to other setups, despite providing an excellent visual match for the circuit output. However, based on the comparison with NN8, the low-precision setup itself is a major source of these errors.

The impact of the LPS becomes more noticeable in the 1001-stage NAND ring oscillator (RO) - a circuit with 4004 total number of transistors, which essentially multiplies an error from a single NAND cell. Fig. 5 demonstrates that NN8 (LPS) starts to deviate slightly from the reference BSIM-CMG output after a certain period of time, while NN8 with the default simulation setup is more precise. Table III shows a much higher oscillation period difference for NN8 (LPS) compared to the default setup, and the visual deviation is primarily driven by this accumulating over time difference. Nevertheless, all NN-based transistor models eventually deviate from the reference output for the 1001-stage NAND RO, as even the error caused by small numerical differences is multiplied and accumulated over time. Results of the analog level circuit simulations demonstrate that even very small NN-based transistor models can provide high-quality visual output for circuit behavior. However, the impact of NN size and simulation precision on accurate delay estimation is noticeable and needs to be considered.

TABLE IV: Time and number of iterations required for the 16-bit divider and 1001-stage NAND ring oscillator (RO) simulations.

Model	16-bit divider				1001-stage NAND RO			
	Time [s]	Iterations Total	Iterations Converged	Computational throughput [iter/s]	Time [s]	Iterations Total	Iterations Converged	Computational throughput [iter/s]
BSIM-CMG	1573	29609	16198	18.82	5165	139191	70313	27
NN16	657	33635	20087	51.19	146	156789	88929	1074
NN8	272	33518	20016	123.23	130	149391	86996	1149
BSIM-CMG (LPS)	794	20094	6383	25.31	2615	89584	28380	34.3
NN16 (LPS)	76	20619	6399	271.30	61	84830	26386	1390
NN8 (LPS)	44	28667	6345	651.52	53	85563	26386	1614

B. Circuit Simulation Speedup of NN-based Models

Another key aspect of the precision-speed balance is the simulation speed. Table IV shows the time that was spent on the simulation of the 16-bit divider circuit with various transistor models and simulation setups. Verilog-A implementation of BSIM-CMG, selected for a fair comparison with Verilog-A NN implementation, has the highest simulation time of $1573\,\mathrm{s}$ due to having thousands of lines of code and hundreds of complex functions. NNs need only $657\,\mathrm{s}$ and $272\,\mathrm{s}$ for 16- and 8-neuron models, respectively. Notably, the number of iterations for both NN simulations is very similar, while Verilog-A BSIM-CMG did roughly 4000 fewer iterations, resulting in an extremely low computational throughput compared to NN models. Computational throughput represents the number of total iterations performed in a fixed amount of time. This relation changes even more when a low-precision simulation setup is used, and NN speed becomes more than $25\times$ and $10\times$ higher than Verilog-A BSIM-CMG for NN8 and NN16, respectively. Further analysis reveals that NN8 (LPS) was able to perform approximately 40% more iterations over the same time, resulting in this extremely high computational speed. However, the number of required iterations is selected by an optimizer and is based on many factors, including model stability, consistency, and derivative smoothness, which smaller NNs can struggle with. Since NN16 is more stable than NN8, it performs similarly to BSIM-CMG in terms of the number of operations in the low-precision case but takes more time due to its bigger size.

In contrast to the 16-bit divider, 1001-stage NAND RO demonstrates even larger speedups compared to BSIM-CMG. Results in Table IV show that computational throughput for NN16 and NN8 is almost $40\times$ and $43\times$ higher than BSIM-CMG, respectively. When a low-precision simulation setup is introduced, the computational throughput ratio remains mostly the same, with a slight increase for both NNs.

Simulation time reduction in Fig. 6 follows similar pattern as computational throughput. Low-precision simulation setup significantly reduces simulation time with NNs for the 16-bit divider, while 1001-stage NAND RO introduces only a minor change in relative time reduction. Based on consistent results across multiple experiments, it is evident that NN16 achieves similar accuracy to BSIM-CMG while maintaining superior computational efficiency in Verilog-A.

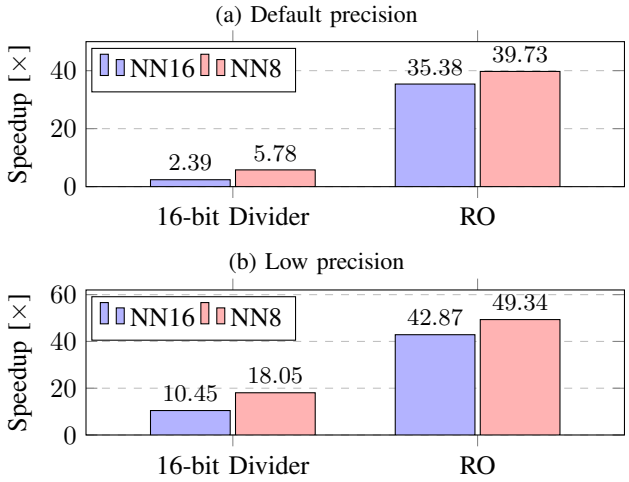

Fig. 6: Relative simulation time reduction for NN-based transistor models compared to BSIM-CMG. All relative values are calculated using data from Table IV.

V. Accuracy Evaluation at the Digital Level

In the following section we perform standard cell library characterization followed by digital circuit synthesis and static timing analysis. In addition to cell and circuit delays, we evaluate power and the speed of characterization, comparing NNs to the industry-standard BSIM-CMG.

A. Standard Cell Library Characterization

Standard cell library characterization is performed with commercial software using the setup described in Section II-C. First, we investigate how NN-based transistor models perform with different supply voltages executing standard cell library characterization with default simulation setup, synthesis, and STA for $V_{dd} \in [0.5\,\mathrm{V},\ 0.6\,\mathrm{V},\ 0.7\,\mathrm{V}]$, excluding low-precision setups. Fig. 7 demonstrates that all models still maintain good accuracy in delay estimation. However, a noticeable trend is observed in a slight performance reduction for all NNs when the supply voltage is significantly lowered. This happens because capacitance behavior in the region near $0.4\,\mathrm{V}$ was not learned as precisely as in the region near $0.8\,\mathrm{V}$, which was demonstrated previously in Fig. 3. However, this error still remains negligible, especially for NN16, which experiences a significantly lower delay error increase than NN8 due to its

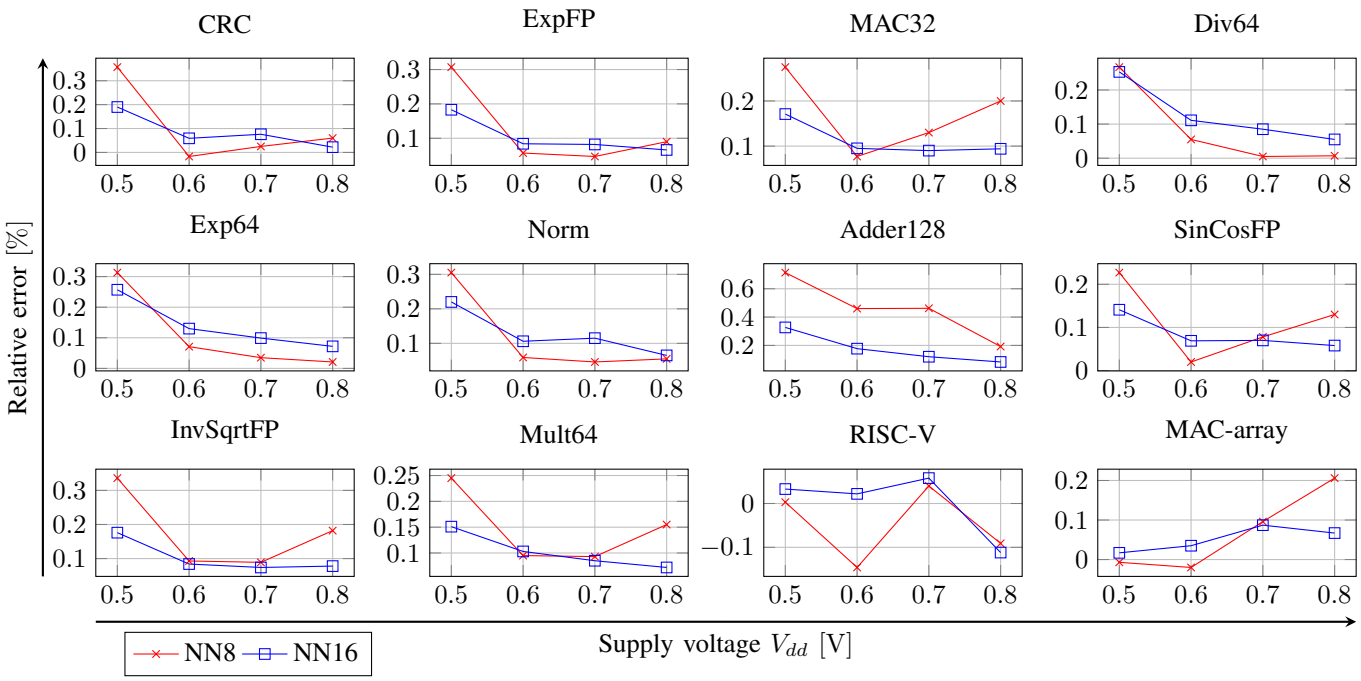

Fig. 7: Relative error for circuit delay estimation calculated with respect to BSIM-CMG for different supply voltages.

larger size and model capacity, allowing it to cover a wider voltage region better. Additionally, NN16 demonstrates less chaotic behavior than NN8, since NN16 error follows a smooth and rather stable curve while NN8 fluctuates a lot.

Since both NNs demonstrated remarkable delay prediction accuracy for the 16-bit divider circuit and V_{dd} = 0.7 V, the supply voltage for detailed analysis during characterization is set to 0.8 V. Each cell type has several variations and incorporates multiple paths for delay measurement. All delay data from 1701 simulations is combined into a database (library). Fig. 8 illustrates relative error distribution for 18620 delay entries (190 paths with 49 slew/capacitance combinations for rise and fall) with 0% representing a perfect match with BSIM-CMG up to 0.1 ps. NN16, shown in Fig. 8a, demonstrates the error distribution with almost 8000 perfect matches and an average relative error of 0.06%, while NN8 in Fig. 8b has fewer perfect matches with 0.08% average error, which is still relatively low. Errors for the low-precision simulation setup, shown in Fig. 8c, fluctuate a lot and are relatively high compared to the reference. Multiple errors for NN8 (LPS) are also not shown, as for some x10 and x12 circuits, they exceeded even 100%, making them unsuitable for any reasonable plot scale. The comparison of two simulation setups for NN8 indicates that the low-precision simulation is the main source of errors and should be employed with caution for standard cell library characterization, although it is perfectly capable of producing a good visual match for a circuit output, which was demonstrated earlier in Section IV. Nevertheless, we will further investigate the performance of NN8 (LPS) in static timing analysis to determine how errors in standard cell delays are propagated into circuit delays.

B. Standard Cell Library Characterization Speedup

The runtime for the standard cell library characterization, as shown in Fig. 9, follows a similar pattern, which was observed for circuits in Section IV-B. NN8 completes all simulations in the shortest amount of time of 4.1 minutes, while Verilog-A BSIM-CMG takes 55 minutes, which is the longest for the same task. Since computational throughput plays a minor role in the standard cell library characterization, NN16 is relatively close to NN8, with a runtime of 7 minutes. In contrast to the 16-bit divider circuit simulation, where the number of calculations is the major factor, characterization speed greatly depends on the amount of code that needs to be compiled. For a single large circuit, Verilog-A code is compiled only once, and then a significant number of calculations is executed. The complexity of standard cells is much lower than that of the circuits used in Section IV. Consequently, a simulation of a single cell will be performed almost instantly and will not be significantly affected by a low-precision simulation setup. However, all Verilog-A implementations require compiling the transistor model for each netlist, meaning that standard cell library characterization runtime greatly depends on the size of the code. With the Verilog-A BSIM-CMG implementation having the largest code base, compiling it multiple times during characterization takes a considerable amount of time. Standard cell library characterization speed with NN-based models is significantly lower than that of the Verilog-A BSIM-CMG due to the simplistic nature of the NN code, which primarily contains matrix multiplications. Notably, compilation time for the Verilog-A code is not affected by the low-precision simulation setup and remains constant for the same implementation, keeping the benefit really low compared to analog level circuit simulations in Section IV-B.

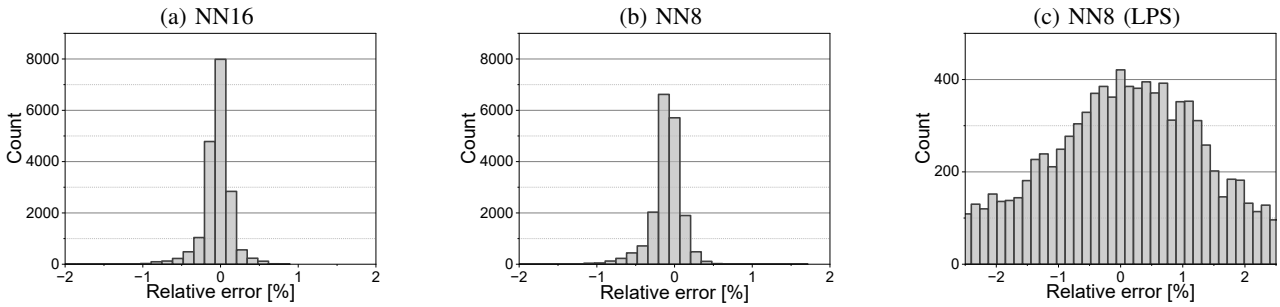

Fig. 8: Distribution of relative errors for standard cell delays measured with respect to the BSIM-CMG. V_{dd} = 0.8 V. NN16 model demonstrates marginally better results over NN8 model. Standard cell library characterization performed with low-precision simulations shows significant deviations from the reference.

TABLE V: Static timing analysis. Relative error (RE) is calculated with respect to BSIM-CMG. V_{dd} = 0.8 V.

Circuit	BSIM-CMG	NN16		NN8		NN8 (LPS)	
	Delay $[ns]$	Delay $[ns]$	RE [%]	Delay $[ns]$	RE [%]	Delay $[ns]$	RE [%]
CRC	0.343551	0.343626	+0.022	0.343756	+0.060	0.341851	-0.495
ExpFP	3.177209	3.179308	+0.066	3.180068	+0.090	3.172469	-0.149
MAC32	2.179060	2.181101	+0.094	2.183413	+0.200	2.176965	-0.096
Div64	8.563543	8.568236	+0.055	8.564155	+0.007	8.540630	-0.268
Exp64	8.669079	8.675332	+0.072	8.670883	+0.021	8.739084	+0.808
Norm	0.437502	0.437785	+0.065	0.437742	+0.055	0.438675	+0.268
Adder128	3.355594	3.358387	+0.083	3.362043	+0.192	3.386376	+0.917
SinCosFP	4.819123	4.821896	+0.058	4.825410	+0.130	4.808432	-0.222
InvSqrtFP	4.086938	4.090143	+0.078	4.094364	+0.182	4.067212	-0.483
Mult64	4.752072	4.755509	+0.072	4.759443	+0.155	4.754639	+0.054
RISC-V	1.482343	1.480690	-0.112	1.480992	-0.091	1.486655	+0.292
MAC-array	0.949363	0.949995	+0.067	0.951319	+0.206	0.951861	+0.263

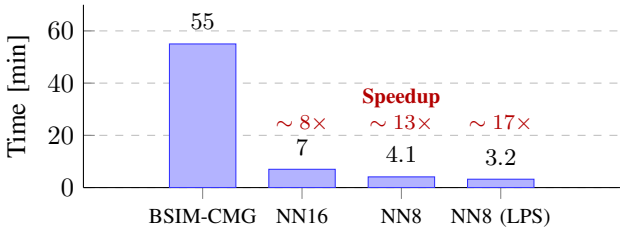

Fig. 9: NN-based transistor models show a significant reduction in time required for standard cell library characterization, with NN16 being approximately 8× faster than BSIM-CMG.

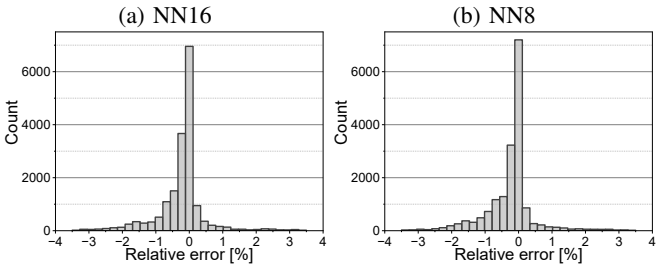

Fig. 10: Distribution of relative errors for cell internal power measured with respect to the BSIM-CMG. V_{dd} = 0.8 V. NN8 error distribution looks visually similar to the NN16 one. However, the average error for NN16 is lower than that for NN8 due to higher accuracy.

C. Static Timing Analysis

Static timing analysis (STA) is performed with a commercial software using libraries (databases) generated during the standard cell library characterization described in Section V-A. Circuit synthesis is done once and is based on the BSIM-CMG library. Then, all 12 synthesized circuits are used to calculate the respective delays for each transistor model. Table V shows that NN16, on average, outperforms NN8 and produces delay measurements closer to the reference ones from BSIM-CMG. Additionally, the NN16 model demonstrates rather stable results, with the relative error ranging from -0.112% to 0.094%, while the NN8 model fluctuates between -0.091% and 0.206%. Nevertheless, NN8, despite having much higher errors than NN16 for various transistor characteristics (see Table II), still shows remarkable delay prediction accuracy. Given the high precision across all 12 circuits, the results indicate that all NN-based transistor models perform with an accuracy level very close to the reference BSIM-CMG. In contrast to other setups, NN8 (LPS) exhibits significant instability and higher delay errors, reaching 0.917%, which is primarily caused by the low simulation precision itself.

TABLE VI: Circuit power analysis with NN-based transistor model. Absolute values are given for NN16, and relative errors (RE) are calculated with respect to BSIM-CMG.

Circuit	Net switching power		Cell internal power		Cell leakage power		Total power	
	Value [W]	RE [%]	Value [W]	RE [%]	Value [W]	RE [%]	Value [W]	RE [%]
CRC	7.032e-05	-0.043	4.395e-05	+0.023	1.251e-07	+0.160	1.144e-04	+0.000
ExpFP	1.664e-04	+0.000	1.570e-04	+0.512	3.811e-06	+0.184	3.271e-04	+0.214
MAC32	5.252e-04	+0.019	7.313e-04	+0.647	9.299e-06	+0.216	1.266e-03	+0.397
Div64	6.088e-04	-0.033	4.116e-04	+0.195	3.661e-05	+0.192	1.057e-03	+0.095
Exp64	2.079e-03	+0.000	1.406e-03	+0.214	3.653e-05	+0.192	3.521e-03	+0.057
Norm	1.539e-05	-0.065	7.716e-06	+0.065	4.667e-08	+0.215	2.316e-05	+0.000
Adder128	2.500e-05	-0.080	1.959e-05	-0.659	6.283e-07	+0.207	4.522e-05	-0.309
SinCosFP	4.627e-05	+0.000	4.602e-05	+0.590	5.200e-06	+0.212	9.749e-05	+0.298
InvSqrtFP	6.698e-05	-0.030	6.722e-05	+0.164	1.221e-06	+0.246	1.354e-04	+0.074
Mult64	8.945e-04	-0.011	1.015e-03	+0.396	2.074e-05	+0.193	1.930e-03	+0.208
RISC-V	9.750e-05	+0.062	7.231e-04	+0.865	6.665e-06	+0.226	8.272e-04	+0.755
MAC-array	4.010e-03	+0.025	1.650e-02	+1.227	1.618e-04	+0.186	2.060e-02	+0.488

D. Circuit Power Signoff Analysis

In addition to timing data, standard cell library characterization provides information on the internal cell power. Fig. 10 shows the relative error distribution for the same 18620 entries employed in timing analysis. Even though NN8 seems to have a very similar error distribution as NN16, the average error of 0.3275% for NN16 is lower than 0.3324%, demonstrated by NN8. Since the cell power library is used to estimate the power of complex circuits, it is more beneficial to employ a model with stable and consistent behavior, which is capable of producing smooth transitions between look-up table (LUT) elements for accurate circuit power estimation. The effect of model stability can be easily noticed in the next experiment.

Circuit analysis provides multiple power characteristics: net switching power, cell internal power, and cell leakage power. Table VI illustrates the power report for a 16-neuron NN-based transistor model and compares its power estimation with BSIM-CMG. Net switching power estimation is the most accurate among the three power components and shows an ideal match for three circuits, with an error below 0.08% for the other ones, indicating a good gate capacitance and ON-current match between two transistor models. Cell leakage power, being relatively low compared to other powers, does not contribute significantly to the total circuit power. Nevertheless, the NN-based transistor model was able to estimate it with high accuracy and relative errors below 0.25% compared to the industry-standard BSIM-CMG.

In contrast to other power components, cell internal power demonstrates the largest relative error, ranging from -0.659% to 1.227%, since it is extremely sensitive to any numerical differences between models. Internal power is primarily influenced by CMOS short circuit conditions and slew propagation, which lead to highly nonlinear behavior, instability, and overall difficulty during estimation. Additionally, circuit power estimation is based on the LUTs from individual cells, which introduces another level of computational errors. All these factors result in error accumulation during different stages of calculation, leading to a more noticeable deviation in the cell internal power. Total power, as a sum of all three powers, also exhibits higher error than switching power on average, but this deviation is primarily driven by the high error in the cell internal power. Notably, power estimation performed with an 8-bit NN-based transistor model exhibits only slightly higher errors for net switching power and similar ones for cell leakage power, with the maximum errors of 0.27% and 0.25%, respectively. However, the relative error for cell internal power, estimated with the NN8 model, ranges from 0.43% to 4.8% and is spread uniformly across this range for all 12 circuits, highlighting the instability of smaller NNs and their struggle with modeling highly nonlinear transistor behavior.

VI. Conclusion

In this paper, we demonstrated that NN-based transistor models can provide high-quality results for standard cell library characterization and further circuit analysis, while achieving significantly higher execution and convergence speeds compared to the industry-standard Verilog-A implementation. Even though circuit-level analysis involves multiple levels of simulations, NN-based transistor models were able to maintain remarkable performance when employed as a foundation in complex multi-stage tasks. Given the right architecture and size, NN-based transistor models open up the possibility of performing precise circuit simulations for any emerging technology without the need to describe underlying physical processes, significantly accelerating the research and electronic design automation flow, especially when employed with GPU acceleration.

Acknowledgement

Authors would like to thank Anirban Kar, Shivendra Singh Parihar, and Girish Pahwa for insightful discussions regarding transistor compact modeling in Verilog-A; Mahdi Benkhelifa and Albi Mema for valuable support during standard cell library characterization and circuit analysis.

REFERENCES

[1] Y. S. Chauhan, S. Venugopalan, M.-A. Chalkiadaki, M. A. U. Karim, H. Agarwal, S. Khandelwal, N. Paydavosi, J. P. Duarte, C. C. Enz, A. M. Niknejad, and C. Hu, "Bsim6: Analog and rf compact model for bulk mosfet," *IEEE Transactions on Electron Devices*, vol. 61, no. 2, pp. 234–244, 2014.

[2] J. P. Duarte, S. Khandelwal, A. Medury, C. Hu, P. Kushwaha, H. Agarwal, A. Dasgupta, and Y. S. Chauhan, "Bsim-cmg: Standard finfet compact model for advanced circuit design," in *ESSCIRC Conference 2015 - 41st European Solid-State Circuits Conference (ESSCIRC)*, 2015, pp. 196–201.

[3] K. Hornik, M. Stinchcombe, and H. White, "Multilayer feedforward networks are universal approximators," *Neural Networks*, vol. 2, no. 5, pp. 359–366, 1989. [Online]. Available: https://www.sciencedirect.com/science/article/pii/0893608089900208

[4] C.-T. Tung, M.-Y. Kao, and C. Hu, "Neural network-based and modeling with high accuracy and potential model speed," *IEEE Transactions on Electron Devices*, vol. 69, no. 11, pp. 6476–6479, 2022.

[5] C.-T. Tung and C. Hu, "Neural network-based bsim transistor model framework: Currents, charges, variability, and circuit simulation," *IEEE Transactions on Electron Devices*, vol. 70, no. 4, pp. 2157–2160, 2023.

[6] J. Wang, Y.-H. Kim, J. Ryu, C. Jeong, W. Choi, and D. Kim, "Artificial neural network-based compact modeling methodology for advanced transistors," *IEEE Transactions on Electron Devices*, vol. 68, no. 3, pp. 1318–1325, 2021.

[7] S. Eom, H. Yun, H. Jang, K. Cho, S. Lee, J. Jeong, and R. Baek, "Neural compact modeling framework for flexible model parameter selection with high accuracy and fast spice simulation," *Advanced Intelligent Systems*, vol. 6, 01 2024.

[8] C.-T. Tung, S. Salahuddin, and C. hu, "A spice-compatible neural network compact model for efficient ic simulations," 09 2024, pp. 01–04.

[9] R. Novkin and H. Amrouch, "Dual neural network framework with spice integration for fast and accurate transistor modeling," *Advanced Intelligent Systems*, vol. 7, 04 2025.

[10] ——, "Kolmogorov–arnold network for transistor compact modeling," *Advanced Intelligent Systems*, 08 2025.

[11] W. Jang, S. Myung, J. M. Choe, Y.-G. Kim, and D. S. Kim, "Tcad device simulation with graph neural network," *IEEE Electron Device Letters*, vol. 44, no. 8, pp. 1368–1371, 2023.

[12] S. Thomann, R. Novkin, J. Li, Y. Hu, J. Xiong, and H. Amrouch, "Ml-tcad: Accelerating fefet reliability analysis using machine learning," *IEEE Transactions on Electron Devices*, vol. 71, no. 1, pp. 213–222, 2024.

[13] K. Sheelvardhan, S. Guglani, M. Ehteshamuddin, S. Roy, and A. Dasgupta, "Machine learning augmented compact modeling for simultaneous improvement in computational speed and accuracy," *IEEE Transactions on Electron Devices*, vol. 71, no. 1, pp. 239–245, 2024.

[14] M.-Y. Kao, H. Kam, and C. Hu, "Deep-learning-assisted physics-driven mosfet current-voltage modeling," *IEEE Electron Device Letters*, vol. 43, no. 6, pp. 974–977, 2022.

[15] G. Choe, J. Kwak, and S. Yu, "Machine learning-assisted compact modeling of w-doped indium oxide channel transistor for back-end-of-line applications," *IEEE Transactions on Electron Devices*, vol. 71, no. 1, pp. 231–238, 2024.

[16] Y. Kim, S. Myung, J. Ryu, C. Jeong, and D. S. Kim, "Physics-augmented neural compact model for emerging device technologies," in *2020 International Conference on Simulation of Semiconductor Processes and Devices (SISPAD)*, 2020, pp. 257–260.

[17] M. Li, Q.-J. Zhang, F. Feng, Z. Li, T. He, S. Zhong, and K. Ma, "A scalable knowledge-based neural network model for gan hemts with accurate trapping and self-heating effects characterization," *IEEE Transactions on Microwave Theory and Techniques*, vol. 71, no. 9, pp. 3747–3760, 2023.

[18] C.-T. Tung, J.-H. Chen, S. Salahuddin, and C. Hu, "Physics-enhanced neural network compact model for fast and flexible transistor modeling," *IEEE Transactions on Electron Devices*, pp. 1–7, 2025.

[19] F. Klemme, Y. Chauhan, J. Henkel, and H. Amrouch, "Cell library characterization using machine learning for design technology co-optimization," in *Proceedings of the 39th International Conference on Computer-Aided Design*, ser. ICCAD '20. New York, NY, USA: Association for Computing Machinery, 2020. [Online]. Available: https://doi.org/10.1145/3400302.3415713

[20] F. Klemme and H. Amrouch, "Machine learning for on-the-fly reliability-aware cell library characterization," *IEEE Transactions on Circuits and Systems I: Regular Papers*, vol. 68, no. 6, pp. 2569–2579, 2021.

[21] ——, "Efficient learning strategies for machine learning-based characterization of aging-aware cell libraries," *IEEE Transactions on Circuits and Systems I: Regular Papers*, vol. 69, no. 12, pp. 5233–5246, 2022.

[22] ——, "Scalable machine learning to estimate the impact of aging on circuits under workload dependency," *IEEE Transactions on Circuits and Systems I: Regular Papers*, vol. 69, no. 5, pp. 2142–2155, 2022.

[23] P. R. Genssler, H. E. Barkam, K. Pandaram, M. Imani, and H. Amrouch, "Modeling and predicting transistor aging under workload dependency using machine learning," *IEEE Transactions on Circuits and Systems I: Regular Papers*, vol. 70, no. 9, pp. 3699–3711, 2023.

[24] M. B. Rakesh, P. Das, K. R. S. Pranav, and A. Acharyya, "Inductive gnn-based methodology for accurate and fast average power estimation of synthesized asic designs from rtl simulation bypassing gate-level simulation," *IEEE Transactions on Circuits and Systems I: Regular Papers*, vol. 72, no. 6, pp. 2818–2831, 2025.

[25] P. Cao, Z. Deng, Y. Dong, X. Cheng, Y. Ye, and J. Yang, "Statchar: Statistical timing characterization framework via heterogeneous graph attention network and active learning with parasitic rc reduction," *IEEE Transactions on Circuits and Systems I: Regular Papers*, pp. 1–11, 2025.

[26] D. Eun Shim, P. Kumar, A. Ashok Kini, M. Mallikarjuna, M. Nahid Haque Shazon, and A. Naeemi, "Gt3: An open-source 3-nm gaafet pdk and platform for end-to-end evaluation of emerging technologies," *IEEE Transactions on Electron Devices*, vol. 72, no. 4, pp. 1582–1588, 2025.

[27] *BSIM-CMG Model*, 2025. [Online]. Available: https://www.bsim.berkeley.edu/models/bsimcmg/

[28] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Kopf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, "Pytorch: An imperative style, high-performance deep learning library," in *Advances in Neural Information Processing Systems 32*. Curran Associates, Inc., 2019, pp. 8024–8035. [Online]. Available: https://pytorch.org/

[29] D. Kingma and J. Ba, "Adam: A method for stochastic optimization," *International Conference on Learning Representations*, 12 2014.

[30] *Chipyard Framework*, 2025. [Online]. Available: https://github.com/ucb-bar/chipyard

Rodion Novkin is a Doctoral Researcher at the Chair of AI Processor Design within the Technical University of Munich (TUM). He is also with the Munich Institute of Robotics and Machine Intelligence in Germany. He received the B.Sc. in Robotics and Mechatronics from the Bauman Moscow State Technical University, Russian Federation, in 2013 and the M.Sc. in Automation and Robotics from the Technical University Dortmund, Germany, in 2021. He is currently working towards the Ph.D. degree at the Chair of AI Processor Design within TUM. His research interests include low-power deep neural networks, machine learning for embedded systems and computer-aided design.

Hussam Amrouch (S'11-M'15) Professor heading the Chair of AI Processor Design within the Technical University of Munich. He is the Founding Director of the Munich Advanced-Technology Center for AI Chips. He is the head of the Brain-inspired Computing within the Munich Institute of Robotics and Machine Intelligence. He is the head of the Semiconductor Test and Reliability at the Uni. of Stuttgart. He received his Ph.D. degree with the highest distinction (Summa cum laude) from KIT in 2015. He holds 10x HiPEAC Paper Awards. He has more than 320 publications in multidisciplinary research areas across the entire computing stack, starting from semiconductor physics to circuit design and EDA.