arXiv:2610.36524v1 [cs.CL] 27 Oct 2026

Causal Localization of Evidence Effects in Transformers

Jipeng Li Pramod Khargonekar Yanning Shen

Abstract

Can controlled symbolic inputs reveal attention heads that influence a language model's decisions from natural language evidence? We introduce controlled evidence aggregation (CEA), which varies the counts and order of two symbols in otherwise fixed prompts and measures changes in a frozen Transformer's preference between two answer tokens. The preference is measured by their logit difference; the prompts prescribe no correct answer. Balanced utility causal selection (BUC) identifies compact head sets by jointly evaluating how much ablating their outputs removes the count effect and how closely patching mean outputs between count conditions recovers the effect in both directions. We derive conditional bounds linking head attribution to intervention effects and specifying when the attribution shortlist can be recovered from finite samples. In Pythia-410M, ablating sets of 1–12 heads selected separately for four ordering regimes removes 65.5–91.6% of the count effect. The sets retain their influence on count comparisons not used for selection. To test relevance beyond symbolic inputs, we evaluate a natural language task asking which state is supported by more reports. Ablating eight heads selected only on symbolic `yes/no` prompts lowers mean accuracy by 8.2 percentage points, versus 0.8 for random sets matched by size and layer. Our results demonstrate that controlled symbolic probes can identify compact attention head sets whose causal influence extends to natural language decisions.

1 Introduction

Understanding how language models use evidence requires identifying internal components that connect input changes to answer changes (Elhage et al., 2021; Bereska and Gavves, 2024; Rai et al., 2024). Natural language reports pose a control problem: changing the balance of support can change wording and invoke different learned associations. Symbolic inputs allow more restricted changes, but the identified components' relevance to language behavior remains an empirical question. Prior studies demonstrate component reuse across tasks and prompt variants (Merullo et al., 2024; Nainani et al., 2024). We ask: *Can controlled symbolic inputs identify attention heads whose causal influence extends to decisions from natural language evidence?*

We introduce controlled evidence aggregation (CEA), which varies the counts and order of B and C in prompts to a frozen Transformer while fixing the template, symbol identities, sequence length, and local answer context. The FRONT ordering, for example, compares

`50 B | 150 C | 50 A` $\longleftrightarrow$ `150 B | 50 C | 50 A`.

Numbers denote repetition counts; both sequences share a tail of 50 A symbols. Both prompts request either `various` or `successful`. Their logit difference measures answer preference; its change between count conditions is the *count effect*. Neither answer is designated correct. CEA therefore supplies a controlled response to localize without training a symbolic rule or assuming pretrained associations are absent.

To localize the count effect, balanced utility causal selection (BUC) evaluates compact head sets through two complementary interventions at the final prompt position. Ablation zeros selected outputs and measures how much the count effect decreases. Mean patching replaces them with averages from the other count condition and measures how closely replacement reproduces the count effect in both directions. Removal and recovery are distinct requirements, consistent with prior work on intervention and evaluation choices (Miller et al., 2024; Shi et al., 2024). BUC uses both scores in greedy search, evaluating sets jointly rather than summing individual head effects. Reuse tests keep selected heads fixed; reselection for each target would test a different claim.

Our theory separates attribution accuracy from selection uncertainty. Under a bound on attribution mismatch, we bound errors in predicting removal and recovery scores from head attribution. We also derive finite sample conditions for recovering the population attribution shortlist. A separate greedy bound applies under monotonicity and weak submodularity to a weighted sum of unnormalized intervention gains, whereas BUC optimizes normalized removal and recovery scores.

To test symbolic localization on labeled behavior, we introduce a natural language (NL) task with sixteen reports supporting two mutually exclusive states. Affirmative and negated reports balance total state word frequencies: `the item is open` and `the item is not closed`, for example, both support `open`. The model answers whether one state has more supporting reports than the other; reversing the comparison reverses the correct `yes/no` answer. At equal head budget, we compare sets selected only on symbolic `yes/no` prompts, sets selected on NL prompts, and random sets matched by layer. Patching in the expanded study uses NL donor outputs, so transfer concerns head identities rather than unchanged activation vectors.

In Pythia-410M, ablating sets of 1–12 heads selected separately for four ordering regimes removes 65.5–91.6% of the count effect. The same sets affect count comparisons not used for selection (Section 5.3). In the original NL evaluation, ablating eight heads selected only on symbolic `yes/no` prompts lowers mean accuracy by 8.2 percentage points, versus 0.8 for matched random sets (Section 5.5). Further tests assess reuse across orderings, answer tokens, and replacement operators. The results show that controlled symbolic probes can identify compact attention head sets whose causal influence extends beyond selection prompts to predictions in a labeled language task.

2 Related work

Behavioral targets. Studies of induction, factual recall, indirect-object identification, and task vectors localize components in pretrained models (Olsson et al., 2022; Geva et al., 2021; Meng et al., 2022; Wang et al., 2022; Hendel et al., 2023; Todd et al., 2024). Task-trained symbolic models permit comparison with a designed algorithm (Brinkmann et al., 2024); structured tasks also reveal compositional failures (Dziri et al., 2023). CEA controls count and order without training a symbolic rule, then tests fixed-set effects on labeled evidence.

Circuit reuse and evaluation. Cross-task reuse (Merullo et al., 2024), prompt-variant generalization (Nainani et al., 2024), ablation dependence (Miller et al., 2024), and tests of preservation, localization, and minimality (Shi et al., 2024) establish distinct aspects of circuit validity. CEA examines which controlled changes preserve the effect of the same head set selected for a count-dependent response, separating this question from successful reselection.

Selection methods. Activation and path patching test a supplied behavioral target (Goldowsky-Dill et al., 2023; Zhang and Nanda, 2024). Automated circuit discovery (ACDC) uses intervention-based search; edge attribution patching (EAP) estimates relevant edges (Conmy et al., 2023; Syed et al., 2024). BUC scores set-level removal and recovery in both directions; equal head budgets do not equate search cost or intervention granularity. Appendix G provides extended comparisons of behavioral targets, selection methods, circuit reuse, and in-context learning interpretations.

3 Controlled target and causal tests

3.1 Symbolic CEA target

Let f_θ be a frozen autoregressive Transformer and $x = (x_1, \ldots, x_T)$ a tokenized prompt. The logit $\ell_T(w \mid x)$ at the final prompt position predicts answer token w. For two audited single-token candidates w_+ and w_-, define the readout

$$y(x) = \ell_T(w_+ \mid x) - \ell_T(w_- \mid x). \tag{3.1}$$

The logit difference $y(x)$ is also the log-odds after restricting the softmax to these two tokens. The candidate order specifies a measurement direction, not correct labels.

The canonical prompt has a fixed prefix, $L_E = 200$ evidence symbols, a tail of $L_A = 50$ copies of A, and a fixed answer cue. The evidence region contains b copies of B and $L_E - b$ copies of C. Its ordering regime belongs to $\mathcal{R} = \{\textsc{Front}, \textsc{Back}, \textsc{Interleave}, \textsc{Random}\}$.

FRONT puts all B symbols first; BACK puts all C symbols first. INTERLEAVE starts with B, alternates while both symbols remain, and appends the excess. RANDOM samples a permutation with the specified counts. Let $D_{r,b}$ be the prompt distribution for regime r and count b. The first three distributions are point masses at each fixed count. For counts $b_H > b_L$ and an integrable scalar function g, define the contrast operator and clean output contrast

$$\mathcal{C}_r^{b_H, b_L}[g] = \mathbb{E}_{D_{r,b_H}}[g(x)] - \mathbb{E}_{D_{r,b_L}}[g(x)], \tag{3.2}$$

$$\Delta_r = \mathcal{C}_r^{b_H, b_L}[y]. \tag{3.3}$$

A nonzero Δ_r is a response to this count manipulation under ordering rule r. Count changes also change symbol positions; the design does not identify count independently of every positional feature. Separate comparisons hold the counts fixed and change ordering.

3.2 Head attribution and intervention unit

Let d_{model} be the residual width and $\mathcal{H}$ the set of attention heads. Let $v_i(x) \in \mathbb{R}^{d_{\text{model}}}$ be head i's write to the residual stream at the final prompt position, after its output projection. For token unembedding vectors $W_U^{(w_+)}$ and $W_U^{(w_-)}$, define

$$d = W_U^{(w_+)} - W_U^{(w_-)}, \tag{3.4}$$

$$s_i(x) = \langle v_i(x), d \rangle, \qquad \Delta s_i(r) = \mathcal{C}_r^{b_H, b_L}[s_i], \tag{3.5}$$

$$F(H, r) = \sum_{i \in H} \Delta s_i(r), \qquad H \subseteq \mathcal{H}. \tag{3.6}$$

The attribution contrast $\Delta s_i(r)$ ranks individual heads; cumulative attribution $F(H, r)$ sums those contrasts. Neither is assumed to decompose the true intervention effect: final normalization and downstream interactions can be nonlinear.

Ablation sets the selected writes to zero. **Mean patching** replaces each selected write by its mean over clean donors from the opposite count. All selected heads are intervened on in one forward pass, including sets spanning layers. Unselected components and downstream computation remain active. Thus recovery is an effect of replacement within the full model, not proof that the selected heads form a sufficient subnetwork. Removal alone can reflect disruption without showing a controlled change toward the opposite count. Conversely, patching can alter the response even when other components preserve the contrast after ablation. BUC therefore evaluates both tests.

Let $\Delta_r^{\text{abl}}(H)$ be the output contrast after ablating set H. High-to-low (HL) patching acts on low-count targets; its oriented gain $\Gamma_r^{\text{HL}}(H)$ is patched minus clean. Low-to-high (LH) patching acts on high-count targets; its oriented gain $\Gamma_r^{\text{LH}}(H)$ is clean minus patched. For $\Delta_r \neq 0$, define

$$\text{shrink}(H, r) = 1 - \frac{|\Delta_r^{\text{abl}}(H)|}{|\Delta_r|}, \tag{3.7}$$

$$\text{match}^{\tau}(H, r) = 1 - \frac{|\Delta_r - \Gamma_r^{\tau}(H)|}{|\Delta_r|}, \quad \tau \in \{\text{HL}, \text{LH}\}, \tag{3.8}$$

$$\text{match}_{\text{avg}}(H, r) = \tfrac{1}{2}\big(\text{match}^{\text{HL}}(H, r) + \text{match}^{\text{LH}}(H, r)\big). \tag{3.9}$$

Shrink measures removal; match measures recovery. A score of one means zero mismatch; zero means no improvement over the clean mismatch. Match penalizes overshoot. Both scores can be negative. Mean match can hide a weak direction, so directional values are reported separately. "Necessity" and "sufficiency" refer only to these operational tests, not to an isolated circuit.

3.3 BUC selection and evaluation

BUC ranks all heads by $|\widehat{\Delta s}_i(r)|$, screens the attribution shortlist with single-head interventions, and runs forward greedy search in the resulting candidate pool $\mathcal{P}_r$. For a weight $\lambda \in [0, 1]$, the set objective is

$$Q_\lambda(H, r) = \lambda \, \text{shrink}(H, r) + (1 - \lambda) \, \text{match}_{\text{avg}}(H, r). \tag{3.10}$$

Each candidate is evaluated jointly with the current set, so its value can change as heads are added. The search accounts for measured interactions that a fixed ranking of individual effects omits. The main configuration uses $\lambda = 1/2$ and at most 12 heads. The FP32 rule takes the first prefix with shrink ≥ 0.85 and mean match ≥ 0.70, or returns the full budget with a failed-threshold flag. Search also records all 12 prefixes for fixed-budget comparisons. Layer-matched controls preserve set size and the number of heads per layer.

Selection and evaluation use separate batches, but deterministic regimes repeat inputs. Exceeding controls, passing both thresholds, and predicting NL labels correctly are distinct criteria. Interpretation depends on the model, inputs, readout, intervention position, and donor rule.

4 Conditional bounds for attribution and selection

The results separate attribution error, shortlist estimation, and optimization over a screened pool. Each requires its own assumptions; success at one step does not establish the conditions for the next.

Fix regime r and counts $b_H > b_L$. The clean contrast Δ_r and cumulative attribution $F(H, r)$ are defined in Section 3. For a head set H, define the signed ablation gain, mean oriented patch gain, and raw joint utility

$$\begin{aligned} N(H,r) &= \Delta_r - \Delta_r^{\mathrm{abl}}(H), \qquad S(H,r) = \tfrac{1}{2} \sum_{\tau \in \{\mathrm{HL},\mathrm{LH}\}} \Gamma_r^\tau(H), \\ J_\lambda(H,r) &= \lambda N(H,r) + (1-\lambda) S(H,r). \end{aligned} \tag{4.1}$$

Unlike the normalized objective Q_λ, raw utility J_λ has no absolute-value penalty for overshooting the target contrast.

Assumption 4.1 (Bounded attribution mismatch)**.** For $a \in \{b_H, b_L\}$, let $\mu_i^a(r) = \mathbb{E}_{D_{r,a}}[s_i(x)]$ be a donor's mean attribution for donor count a. Write the intervention residuals as

$$y^{\mathrm{abl}(H)}(x) = y(x) - \sum_{i \in H} s_i(x) + \xi_H^{\mathrm{abl}}(x), \tag{4.2}$$

$$y^{\mathrm{pat}(H;a)}(x) = y(x) + \sum_{i \in H} (\mu_i^a(r) - s_i(x)) + \xi_H^a(x). \tag{4.3}$$

Assume the absolute ablation residual contrast and the absolute mean patch residual on each opposite-count target are at most $\varepsilon_{\mathrm{int}}(H, r)$. The residual identities are definitional; a small error bound is an additional condition, not a consequence of residual-stream additivity.

Theorem 4.2 (T1: attribution error and intervention scores)**.** *Under Assumption 4.1,*

$$\max\left\{ |N(H,r) - F(H,r)|, \max_{\tau \in \{\mathrm{HL},\mathrm{LH}\}} |\Gamma_r^\tau(H) - F(H,r)| \right\} \le \varepsilon_{\mathrm{int}}(H,r). \tag{4.4}$$

The same bound holds for $|S - F|$ and $|J_\lambda - F|$. If $\Delta_r \neq 0$, then

$$\left| Q_\lambda(H,r) - \left(1 - \frac{|\Delta_r - F(H,r)|}{|\Delta_r|}\right) \right| \le \frac{\varepsilon_{\mathrm{int}}(H,r)}{|\Delta_r|}. \tag{4.5}$$

Even with small mismatch, strong removal and recovery require cumulative attribution $F(H, r)$ to be close to the clean contrast Δ_r, not merely large in magnitude. Weak contrasts amplify normalized error. Appendix A.2 proves T1; Appendix A.3 gives the separate shrink and match bounds. A conditional position-profile approximation and its transfer bound are in Appendix A.4. Similar attribution across regimes alone does not guarantee similar normalized scores.

Theorem 4.3 (T2: recovery of the attribution shortlist)**.** *Let $(x_j^H, x_j^L)_{j=1}^n$ be $n \ge 1$ independent matched pairs with marginals D_{r,b_H} and D_{r,b_L}. For each head i, suppose the centered paired score $Z_{i,j} = s_i(x_j^H) - s_i(x_j^L) - \Delta s_i(r)$ satisfies $\mathbb{E}e^{tZ_{i,j}} \le e^{t^2\nu_r^2/2}$ for every real t. Define $\widehat{\Delta s}_i(r) = n^{-1}\sum_j [s_i(x_j^H) - s_i(x_j^L)]$. For $\delta \in (0,1)$, with probability at least $1-\delta$,*

$$\max_{i \in \mathcal{H}} |\widehat{\Delta s}_i(r) - \Delta s_i(r)| \le \eta_{n,\delta}(r) := \nu_r \sqrt{\frac{2\log(2|\mathcal{H}|/\delta)}{n}}. \tag{4.6}$$

For $1 \le m < |\mathcal{H}|$, an absolute population-score gap greater than $2\eta_{n,\delta}(r)$ between ranks m and $m+1$ guarantees exact recovery of the top-m attribution shortlist on the same event.

T2 permits dependence within pairs and across heads. It addresses the attribution ranking, not the subsequent intervention screening. Repeated deterministic inputs do not estimate uncertainty over other prompts. Appendix A.5 gives the proof and the independent-sample special case.

Assumption 4.4 (Monotonicity and weak submodularity)**.** On a fixed screened pool $\mathcal{P}_r$, raw utility J_λ is monotone and has submodularity ratio $\gamma_r \in (0, 1]$ as defined in Equation A.10, including unions of greedy prefixes and budget-k comparators.

Theorem 4.5 (T3: conditional greedy bound)**.** *For $1 \le k \le |\mathcal{P}_r|$, let G_k be the exact forward-greedy set for J_λ, and let O_k maximize J_λ over pool subsets of size at most k. Under Assumption 4.4, with $\alpha_r = 1 - e^{-\gamma_r}$,*

$$J_\lambda(G_k, r) \ge \alpha_r J_\lambda(O_k, r). \tag{4.7}$$

If $|J_\lambda(H, r) - F(H, r)| \le \epsilon_{J,k}(r)$ for every pool subset of size at most k, then

$$F(G_k, r) \ge \alpha_r F(O_k, r) - (1 + \alpha_r)\epsilon_{J,k}(r). \tag{4.8}$$

The comparator O_k maximizes raw utility J_λ, not attribution F. T3 does not guarantee BUC's empirical maximization of Q_λ: the objectives differ, and the FP32 audits find negative marginal gains for observed $Q_{1/2}$. Small regret within a tested eight-head pool supplies no guarantee for the full candidate pool. Appendix A.6 gives the proof.

5 Experiments

The original BF16 study (Appendices B–D) and native FP32 audits (Appendices E–F) use separate configurations and are not pooled. Brackets give two-sided 95% confidence intervals (CIs) for eligible sampling units, not across-condition spans. Deterministic single-pair conditions and finite audit summaries are reported as point values (Appendix H).

5.1 Setup and main symbolic result

The frozen `EleutherAI/pythia-410m-deduped` has 24 layers and 16 heads per layer (Biderman et al., 2023). The BF16 study uses five execution seeds and final prompt position interventions, without training. At $(b_L, b_H) = (50, 150)$, selection/evaluation use separate 48/96-row batches per count; deterministic inputs repeat. BUC shortlists 192 heads, screens to 48, and searches up to 12; controls use 120 layer-matched draws per seed.

Table 1 compares simultaneous set-level interventions with layer-matched random controls; the scores include non-additive interactions.

Table 1: Main BF16 symbolic result at $(b_L, b_H) = (50, 150)$. BUC mean match averages both directions. The last two columns compare the same HL direction with random controls; their values are reported in Appendix B.3. Uncertainty is the standard deviation over five executions.

Regime	Δ_r	Size	BUC shrink / mean match	BUC HL match	Random shrink / HL match
FRONT	0.239	1	**0.906 / 0.974**	0.998	0.020 / 0.048
BACK	1.965	12	**0.655 / 0.759**	0.765	0.043 / 0.075
INTERLEAVE	1.659	3	**0.916 / 0.914**	0.937	0.061 / 0.051
RANDOM	3.184 ± 0.097	7 ± 1	$\mathbf{0.866 \pm 0.021 / 0.823 \pm 0.084}$	0.813	0.069 / 0.061

Compactness is regime-dependent: FRONT uses one head; BACK uses all 12 and its shrink remains below 0.85. Deterministic sets repeat across seeds 0–4 but do not test stability over new prompts. The three documented RANDOM sets (seeds 0–2) have minimum Jaccard overlap $5/9$ (Appendix B.5).

5.2 Budget-matched baseline comparison

Table 2 shows a BF16 removal–recovery trade-off. In BACK, INTERLEAVE, and RANDOM, BUC increases bidirectional match by 0.083–0.111 and reduces shrink by 0.018–0.049 relative to sensitivity-only; FRONT ties in the summed score at displayed precision. Both methods satisfy the joint 0.85/0.70 thresholds in the FRONT, INTERLEAVE, and RANDOM means; both fail in BACK. The improvement is therefore in recovery match rather than joint-threshold coverage or both component metrics. Appendix C.4 gives the arithmetic decomposition and ActPatch comparison.

FP32 fixed-budget comparison. At $k = 8$, BUC-minus-sensitivity differences in normalized joint score $Q_{1/2}$ are $+0.006$, $+0.030$, $+0.031$, and $+0.039$ for FRONT, BACK, INTERLEAVE, and RANDOM, respectively. Only RANDOM has distinct matched permutations for a prompt-bootstrap CI: $[+0.015, +0.065]$. The first three differences describe fixed deterministic inputs and have no prompt-sampling CI. Head count, not search cost, is matched (Appendix E.6).

Table 2: BF16 comparison with matched head counts within each regime (Table 1), not a common $k = 8$. ActPatch ranks 192 sensitivity-shortlisted heads by single-head HL recovery, then reports the jointly evaluated set score $2Q_{1/2}$ (Appendix B.10).

Regime	BUC shrink / match	Sens. shrink / match	ActPatch sum
FRONT	0.906 / 0.974	0.920 / 0.960	1.865
BACK	0.655 / 0.759	0.704 / 0.648	1.285
INTERLEAVE	0.916 / 0.914	0.934 / 0.831	1.695
RANDOM	0.866 / 0.823	0.895 / 0.716	1.520

5.3 Count generalization and ordering dependence

Heads selected only at $(50, 150)$ are held fixed over $b \in \{0, 10, \ldots, 200\}$ and new count pairs $(0, 150)$, $(25, 175)$, and $(75, 125)$. Table 3 shows slope reductions of 0.646–0.896 after ablation and positive effects on the new count pairs. The pair $(0, 150)$ shares one selection endpoint, whereas the other pairs use new endpoints; the effects are not confined to the selection pair. The clean sweep describes the count response; ablation tests whether the selected writes contribute across that range. The order-only tests below instead ask whether equal counts can produce different responses.

Table 3: Fixed-set count generalization. The final column averages evaluation on three unseen gaps without head reselection.

Regime	Spearman $\rho(b, y)$	Linear R^2	Slope reduction	Unseen-gap shrink / match
FRONT	0.986	0.948	0.880	0.832 / 0.874
BACK	0.972	0.931	0.646	0.612 / 0.689
INTERLEAVE	0.979	0.944	0.896	0.839 / 0.855
RANDOM	0.958	0.906	0.845	0.786 / 0.811

Four pre-specified order-only contrasts have identical B/C counts, yet absolute output effects of 0.321–1.846. BUC shrink/match ranges are 0.603–0.842/0.651–0.884, versus random shrink/HL match of 0.049–0.081/0.056–0.088. The effects cannot be explained by count differences (Appendix C.3).

Regime transfer and cue wording. BF16 BACK→FRONT match is 0.958 versus 0.129 in reverse, but source budgets are 12 versus 1 (Appendix B.8). At equal FP32 budget $k = 8$, off-diagonal bidirectional match ranges from -0.84 to $+0.89$. Negative match means greater average residual mismatch than no patch, not necessarily an opposite-signed gain. Fixed-set cue changes retain shrink 0.64–0.93 and match 0.70–0.96: reuse tolerates these wording changes but is not regime-invariant. These are descriptive spans over tested conditions, not confidence intervals.

5.4 Readout-pair audit

We repeat the full protocol with seven audited single-token readout pairs: `various/successful`, `yes/no`, `true/false`, `1/0`, `A/B`, `cat/dog`, and `red/blue`. Across pairs, mean $|\Delta_r|$ is 0.721–2.083, mean selected-set size is 4.50–7.50, mean BUC shrink is 0.642–0.861, and mean match is 0.701–0.884. Layer-matched random shrink remains 0.048–0.076 and random high-to-low match remains 0.059–0.083.

Table 4 reports target-specific reselection and fixed-set effects on representative pairs. These selected-set summaries do not establish an equal-budget reselection advantage. The `A/B` row tests lexical overlap with prompt symbols.

Table 4: Representative readout-pair results, averaged over matched regime–seed sets. The fixed set was selected with `various/successful`.

Target pair	Reselected shrink / match	Fixed-set shrink / match	Jaccard
`yes/no`	0.812 / 0.846	0.772 / 0.810	0.610
`A/B`	0.861 / 0.884	0.704 / 0.741	0.390
`cat/dog`	0.642 / 0.701	0.382 / 0.431	0.240

Registered-readout scan (FP32). Of $4 \times 256 = 1024$ fixed-prompt, fixed-set regime–pair conditions, 960 have $|\widehat{\Delta}_r| \geq 0.01$. Fractions passing shrink ≥ 0.85 and mean match ≥ 0.70 are 0.060 of all conditions and 0.064 of evaluable

conditions. Strong fixed-set reuse is uncommon in this finite registry. Failing a joint threshold need not mean zero effect: the scan measures the prevalence of strong reuse under the stated thresholds, not of any detectable influence. The seven targeted pairs and this registry are different samples, not a paired fixed-versus-reselected comparison. Weak conditions remain registered (Appendix E.8).

5.5 Lexically balanced natural-language bridge

The labeled NL bridge uses 16 reports about mutually exclusive states. Balanced direct and negative forms keep eight occurrences of each state word and eight negative reports. The same word totals can therefore accompany opposite majority labels; state word frequencies alone cannot solve the task. Selection uses three state pairs and Bank A; evaluation uses three unseen pairs and separately written Bank B paraphrases. Each sequence is queried in both directions. A constant answer fails one query; paired metrics also check a rule that uses query orientation alone.

Here x denotes the reports, b counts reports supporting the first state, and query polarity $q \in \{+1, -1\}$ selects which state is queried first. The model input combines those reports and the specified query. The correct restricted label (1 for `yes`) and aligned readout are

$$a^\star(x, q) = \mathbf{1}[q(2b - 16) > 0], \qquad y_{\mathrm{NL}}(x, q) = q\big[\ell_T(\texttt{yes} \mid x, q) - \ell_T(\texttt{no} \mid x, q)\big]. \tag{5.1}$$

The main NL contrast uses $(b_L, b_H) = (4, 12)$. Table 5 reports the original BF16 bridge. The mean restricted accuracy is 0.772 on held-out state pairs and paraphrases. PairAcc requires both reversed queries to be correct; MinPairAcc takes the weaker positive- and negative-margin paired score and detects a shortcut that succeeds on only one margin sign. The graded response has Spearman correlation 0.901 and linear $R^2 = 0.812$.

Table 5: Original BF16 natural-language bridge, five-seed mean. The behavior row uses held-out domains and paraphrases. The intervention rows use equal budget $k = 8$. Accuracy drop is in percentage points.

Condition	Accuracy	PairAcc	MinPairAcc	Shrink / match	Slope reduction	Acc. drop
Base behavior	0.772	0.604	0.583	–	–	–
NL set → NL	–	–	–	0.734 / 0.766	0.721	15.2
CEA set → NL	–	–	–	0.428 / 0.463	0.411	8.2
Layer-matched random	–	–	–	0.059 / 0.072	0.051	0.8

At $k = 8$, the fixed symbolic `yes/no` set's NL effect lies between the NL-selected and random effects. Ablating the symbolic set lowers accuracy by 8.2 percentage points, versus 0.8 for random controls. The symbolic set sees no NL prompt during selection. The transfer ratio (TR) divides the transferred set's normalized joint score $Q_{1/2}$ above random by the NL-selected set's score above random. The BF16 value, 0.555, compares intervention effects; it is not a fraction of shared heads. Shrink measures one count contrast, slope reduction covers the count grid, and accuracy drop measures effects on discrete predictions. Together, these metrics support partial intervention transfer, but do not distinguish shared evidence processing from shared answer readout. The BF16 summaries lack the bootstrap intervals required by the full declared bridge criterion (Appendix D).

Expanded bridge (FP32). Table 6 covers eight test domains and Banks B–D. Behavioral evaluation uses three seeds; causal selection uses one. At $k = 8$, fixed sets with query-stratified target calibration yield mean TRs of 0.438–0.557 and accuracy decreases of 4.9–8.6 percentage points. These model summaries are descriptive, not model-level significance tests (Appendix E.10).

Table 6: Expanded FP32 NL bridge, separate from Table 5: model-level point estimates without model-level CIs. Behavior uses gap $(4, 12)$; transfer uses fixed symbolic `yes/no` sets. n_{est} counts estimable TR conditions out of 12; n_{suppr} counts those with a suppressed *condition-level* interval. Accuracy drop is in percentage points on the causal cohort.

Model	Accuracy	PairAcc	MinPairAcc	CEA → NL shrink	CEA → NL match	TR	$n_{\mathrm{est}}/12$	n_{suppr}	Acc. drop (pp)
Pythia-410M	0.752	0.566	0.541	0.401	0.442	0.520	11/12	1	7.4
Pythia-1B	0.805	0.651	0.611	0.462	0.501	0.557	12/12	0	8.6
GPT-2	0.681	0.458	0.414	0.306	0.348	0.438	9/12	3	4.9

The FP32 ratios use separate controls matched to source and target layer patterns (Equation E.14). Model summaries average all finite ratios passing the target-contrast (at least 0.01 in magnitude) and target-excess (more than 0.05 above

matched random) gates. Unavailable ratios are not zero-filled; interval suppression does not remove a finite point. A condition-level interval is suppressed above 5% unstable bootstrap replicates (Appendix E.10). Estimability is not a positive-transfer success rate.

5.6 Scope checks

The LAST-versus-SEQ comparison tests intervention location. Applying the selected sets only at symbolic-sequence positions yields shrink 0.001–0.059 and match -0.006 to 0.006. For the reported scope diagnostic, these heads have much larger effects at the final prompt position. The RANDOM diagnostic is not documented as the same five-execution cohort as the main table (Appendix B.7).

Post-selection split patches distinguish query–key (QK) attention routing from output–value (OV) writing. Representative FRONT, INTERLEAVE, and RANDOM heads have OV-only match 0.88/0.92/0.85 versus QK-only 0.35/0.12/0.18. BACK has QK/OV match 0.72/0.61, indicating a more mixed contribution (Appendix B.6).

Word-symbolic, signal-token, and short-sentence inputs retain mean shrink/match of 0.812/0.845, 0.820/0.852, and 0.765/0.790. Across Pythia-70M through Pythia-1B, GPT-2, and Qwen3-4B-Base (Radford et al., 2019; Yang et al., 2025), four-regime mean BUC shrink/match ranges from 0.650/0.700 to 0.895/0.925 (Appendix B.19). Positive effects do not establish shared head identity across models.

Input-only Greedy Occlusion gives held-out shrink/recovery of 0.795–0.915/0.812–0.898, versus random-subset recovery of 0.248–0.262. The search uses token scores and 8,775 candidate-subset evaluations per regime before batching; generated text alone is insufficient. Input recovery does not penalize overshoot and is not the head-patching match score. BUC requires activation access.

5.7 Attribution and selection audits

Many-set T1 diagnostic (BF16). Across 500 sampled head sets per regime, cumulative attribution $F(H, r)$ has Pearson correlation 0.841–0.923 with the ablation gain $N(H, r)$, and 0.755–0.884 with the two oriented patch gains $\Gamma_r^{\mathrm{HL}}(H)$ and $\Gamma_r^{\mathrm{LH}}(H)$ (Appendix B.9). The sampled sets show an association, not a bound on attribution error or a guarantee of correct ranking.

Operator and T3 audits (FP32). Observed operator minima in Table 7 remain positive but do not uniformly retain the joint thresholds. Negative marginal fractions are 17–31%, and minimum empirical ratios are -5.8 to -2.6. These witnesses in the finite pools establish nonmonotonicity of observed $Q_{1/2}$ only on the tested data and candidate pools. Maximum selection regret is 0.021–0.043 within the registered local pools; small local regret certifies neither full-pool optimality nor the raw-$J_{1/2}$ bound (Appendix E.5).

Table 7: FP32 operator and local selection audits: point estimates. Shrink and match minima across eight operators need not share a minimizer. Evaluation marginals/ratios use both eight-head pools and tolerance 10^{-3}. Maximum regret compares exact and greedy selection utilities over both pools and budgets $k = 1, \ldots, 8$, not held-out performance. These finite summaries have no population CIs (Appendix H).

Regime	Minimum shrink	Minimum match	Negative marginals (%)	Minimum $\widehat{\gamma}$	Maximum selection regret
FRONT	0.81	0.84	31	-5.8	0.021
BACK	0.52	0.57	17	-2.6	0.043
INTERLEAVE	0.80	0.81	28	-4.9	0.027
RANDOM	0.71	0.68	22	-3.9	0.035

Selection-matched null (FP32). For RANDOM, 99 completed paired-label trials rerun the adaptive pipeline. Raw magnitude-improvement utility is 2.69 $[2.55,\ 2.80]$; Holm-adjusted $p = 0.01$ is at Monte Carlo resolution. The interval for the null exceedance probability is given in Appendix E.9. Only RANDOM meets the unique-pair requirement. The result supports rejection of the specified exchangeability null, not mechanism uniqueness (Appendices E.9, H).

6 Discussion and limitations

Scientific and practical scope. Count generalization and limited transfer are compatible. A selected set can influence several count contrasts without being independent of ordering or readout. Comparing layer-matched sets controls head count and layer allocation, so the difference concerns which heads are selected rather than simply how many or how deep. The labeled bridge adds evidence that a fixed symbolic set affects predictions, but does not identify whether the shared computation processes reports or writes the `yes/no` response. Phrase-based strategies remain possible; general NL reasoning is not established. Interventions identify writes at the final prompt position, not complete head–MLP paths; compactness is neither uniqueness nor global minimality.

Selection and statistical scope. Head-budget comparisons do not control search cost or isolate greedy search from ranking by the same joint objective in the same pool (Appendix C.4). Bootstrap CIs condition on fixed sets and donors; repeated deterministic inputs do not support prompt-sampling intervals. The BF16 bridge lacks its required bootstrap intervals. Expanded FP32 model summaries have no CIs, use one seed for causal selection, and average TR only over estimable conditions.

Use and open questions. Further tests should distinguish evidence aggregation from answer-readout effects, examine computations at earlier tokens and along head–MLP paths, and extend to longer heterogeneous reports and decisions involving multiple tokens. Localized components could enable targeted manipulation; the present audits do not certify deployment safety.

7 Conclusion

CEA tests whether controlled symbolic inputs can identify attention heads relevant to natural language decisions. BUC selects compact head sets by jointly evaluating ablation and bidirectional mean patching. Conditional theory links attribution to intervention effects and specifies assumptions for shortlist recovery and greedy bounds on raw utility. In Pythia-410M, selected heads affect count contrasts not used for selection. Ablating a fixed symbolic `yes/no` set reduces labeled NL accuracy more than random controls matched by size and layer. Audits characterize how reuse depends on ordering, readout, and replacement operators. The results show that controlled symbolic probes can identify attention heads whose causal influence extends to predictions in a labeled language task.

AI use statement

Generative AI tools assisted with language editing, document structure, LaTeX formatting and checking of mathematical claims and proofs. The authors are responsible for the final text, mathematical claims, code, experimental measurements, and statistical procedures.

Ethics statement

The work analyzes open pretrained language models with synthetic and controlled natural-language prompts. It does not use personal data or human subjects. A localized component could be used for targeted model manipulation, and controlled results could be overstated as evidence of general reasoning or safety. Section 6 states these risks and restricts the supported conclusions.

Reproducibility statement

Sections 3–5 define the target, interventions, selection objective, and experimental comparisons. Appendix A contains the proofs; Appendices B–D specify the BF16 experiments. Appendices E–F describe FP32 sampling, controls, code, and numerical checks, and Appendix H specifies which statistics have confidence intervals. The accompanying source and configurations are under `anc/code/` and `anc/protocol/`; aggregate measurements are under `anc/results/`. Appendix F.5 describes the package contents, runtime dependencies, and experiments outside the supplied implementation.

References

Kwangjun Ahn, Xiang Cheng, Hadi Daneshmand, and Suvrit Sra. Transformers learn to implement preconditioned gradient descent for in-context learning. *Advances in Neural Information Processing Systems*, 36:45614–45650, 2023.

Ekin Akyürek, Dale Schuurmans, Jacob Andreas, Tengyu Ma, and Denny Zhou. What learning algorithm is in-context learning? investigations with linear models. *arXiv preprint arXiv:2211.15661*, 2022.

Yu Bai, Fan Chen, Huan Wang, Caiming Xiong, and Song Mei. Transformers as statisticians: Provable in-context learning with in-context algorithm selection. *Advances in neural information processing systems*, 36:57125–57211, 2023.

Leonard Bereska and Efstratios Gavves. Mechanistic interpretability for ai safety–a review. *arXiv preprint arXiv:2404.14082*, 2024.

Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O'Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, et al. Pythia: A suite for analyzing large language models across training and scaling. In *International conference on machine learning*, pages 2397–2430. PMLR, 2023.

Jannik Brinkmann, Abhay Sheshadri, Victor Levoso, Paul Swoboda, and Christian Bartelt. A mechanistic analysis of a transformer trained on a symbolic multi-step reasoning task. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 4082–4102, 2024.

Arthur Conmy, Augustine Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. *Advances in Neural Information Processing Systems*, 36: 16318–16352, 2023.

Guy Dar, Mor Geva, Ankit Gupta, and Jonathan Berant. Analyzing transformers in embedding space. In *Proceedings of the 61st annual meeting of the Association for Computational Linguistics (volume 1: Long papers)*, pages 16124–16170, 2023.

Abhimanyu Das and David Kempe. Approximate submodularity and its applications: Subset selection, sparse approximation and dictionary selection. *Journal of Machine Learning Research*, 19(3):1–34, 2018.

Andrea Gregor de Varda, Sana Pandey, Pengrui Han, Jacob Andreas, and Evelina Fedorenko. Shared circuits predict whether llms generalize across formats in arithmetic reasoning. *arXiv preprint arXiv:2609.04463*, 2026.

Nouha Dziri, Ximing Lu, Melanie Sclar, Xiang Lorraine Li, Liwei Jiang, Bill Yuchen Lin, Sean Welleck, Peter West, Chandra Bhagavatula, Ronan Le Bras, et al. Faith and fate: Limits of transformers on compositionality. *Advances in neural information processing systems*, 36:70293–70332, 2023.

Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. https://transformer-circuits.pub/2021/framework/index.html.

Javier Ferrando and Elena Voita. Information flow routes: Automatically interpreting language models at scale. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 17432–17445, 2024.

Shivam Garg, Dimitris Tsipras, Percy S Liang, and Gregory Valiant. What can transformers learn in-context? a case study of simple function classes. *Advances in neural information processing systems*, 35:30583–30598, 2022.

Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy. Transformer feed-forward layers are key-value memories. In *Proceedings of the 2021 conference on empirical methods in natural language processing*, pages 5484–5495, 2021.

Mor Geva, Jasmijn Bastings, Katja Filippova, and Amir Globerson. Dissecting recall of factual associations in auto-regressive language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12216–12235, 2023.

Nicholas Goldowsky-Dill, Chris MacLeod, Lucas Sato, and Aryaman Arora. Localizing model behavior with path patching. *arXiv preprint arXiv:2304.05969*, 2023.

Roee Hendel, Mor Geva, and Amir Globerson. In-context learning creates task vectors. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 9318–9333, 2023.

Yingcong Li, Muhammed Emrullah Ildiz, Dimitris Papailiopoulos, and Samet Oymak. Transformers as algorithms: Generalization and stability in in-context learning. In *International conference on machine learning*, pages 19565–19594. PMLR, 2023.

Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in gpt. *Advances in neural information processing systems*, 35:17359–17372, 2022.

Jack Merullo, Carsten Eickhoff, and Ellie Pavlick. Circuit component reuse across tasks in transformer language models. In *International Conference on Learning Representations*, volume 2024, pages 18349–18377, 2024.

Joseph Miller, Bilal Chughtai, and William Saunders. Transformer circuit faithfulness metrics are not robust. *arXiv preprint arXiv:2407.08734*, 2024.

Jatin Nainani, Sankaran Vaidyanathan, AJ Yeung, Kartik Gupta, and David Jensen. Adaptive circuit behavior and generalization in mechanistic interpretability. *arXiv preprint arXiv:2411.16105*, 2024.

Catherine Olsson, Nelson Elhage, Neel Nanda, Nicholas Joseph, Nova DasSarma, Tom Henighan, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, et al. In-context learning and induction heads. *arXiv preprint arXiv:2209.11895*, 2022.

Madhur Panwar, Kabir Ahuja, and Navin Goyal. In-context learning through the bayesian prism. In *International Conference on Learning Representations*, volume 2024, pages 49789–49843, 2024.

Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.

Daking Rai, Yilun Zhou, Shi Feng, Abulhair Saparov, and Ziyu Yao. A practical review of mechanistic interpretability for transformer-based language models. *arXiv preprint arXiv:2407.02646*, 2024.

Claudia Shi, Nicolas Beltran-Velez, Achille Nazaret, Carolina Zheng, Adrià Garriga-Alonso, Andrew Jesson, Maggie Makar, and David M Blei. Hypothesis testing the circuit hypothesis in llms. *Advances in neural information processing systems*, 37:94539–94567, 2024.

Aaquib Syed, Can Rager, and Arthur Conmy. Attribution patching outperforms automated circuit discovery. In *Proceedings of the 7th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 407–416, 2024.

Eric Todd, Millicent Li, Arnab Sen Sharma, Aaron Mueller, Byron Wallace, and David Bau. Function vectors in large language models. In *International conference on learning representations*, volume 2024, pages 17282–17333, 2024.

Johannes Von Oswald, Eyvind Niklasson, Ettore Randazzo, João Sacramento, Alexander Mordvintsev, Andrey Zhmoginov, and Max Vladymyrov. Transformers learn in-context by gradient descent. In *International Conference on Machine Learning*, pages 35151–35174. PMLR, 2023.

Kevin Wang, Alexandre Variengien, Arthur Conmy, Buck Shlegeris, and Jacob Steinhardt. Interpretability in the wild: a circuit for indirect object identification in gpt-2 small. *arXiv preprint arXiv:2211.00593*, 2022.

Sang Michael Xie, Aditi Raghunathan, Percy Liang, and Tengyu Ma. An explanation of in-context learning as implicit bayesian inference. *arXiv preprint arXiv:2111.02080*, 2021.

An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods. In *International Conference on Learning Representations*, volume 2024, pages 1651–1678, 2024.

A Proofs and scope of the theoretical claims

All expectations below are finite. Fix an ordering regime r, counts $b_H > b_L$, and a head set $H \subseteq \mathcal{H}$, where $\mathcal{H}$ is the set of model heads. Write

$$\mathcal{C}_r g = \mathbb{E}_{D_{r,b_H}}[g] - \mathbb{E}_{D_{r,b_L}}[g]$$

for the count contrast of an integrable scalar function g. Thus the clean logit contrast is $\Delta_r = \mathcal{C}_r y$, a head's attribution contrast is $\Delta s_i(r) = \mathcal{C}_r s_i$, and cumulative attribution is $F(H, r) = \sum_{i \in H} \Delta s_i(r)$. Interventions act simultaneously on the specified heads at the final prompt position. Theorems concern population quantities and the stated operators; finite donor estimates, other positions, and other operators require their own error control.

A.1 Readout and intervention conventions

For distinct answer tokens w_+ and w_-, the readout is $y(x) = \ell_T(w_+ \mid x) - \ell_T(w_- \mid x)$. If $p_+(x)$ denotes the softmax probability restricted to these two tokens, then

$$p_+(x) = \frac{e^{\ell_T(w_+|x)}}{e^{\ell_T(w_+|x)} + e^{\ell_T(w_-|x)}}, \qquad \log \frac{p_+(x)}{1 - p_+(x)} = y(x).$$

The identity follows by cancelling the common denominator. It does not make an arbitrary answer-token pair a task with correct labels.

Let $v_i(x)$ be head i's residual write and $d = W_U^{w_+} - W_U^{w_-}$ the difference of unembedding directions. Direct attribution is $s_i(x) = \langle v_i(x), d \rangle$. Subsequent normalization and nonlinear computation can make s_i differ from the head's intervention effect. For a donor count $a \in \{b_H, b_L\}$, let $\mu_i^a(r) = \mathbb{E}_{D_{r,a}} s_i$. The residual terms in Assumption 4.1 account for that difference; writing the identities alone imposes no restriction on a model.

For clarity, the error bound assumed by T1 is

$$\max\left\{ |\mathcal{C}_r \xi_H^{\mathrm{abl}}|, |\mathbb{E}_{D_{r,b_L}} \xi_H^{b_H}|, |\mathbb{E}_{D_{r,b_H}} \xi_H^{b_L}| \right\} \le \varepsilon_{\mathrm{int}}(H, r). \tag{A.1}$$

A mean donor estimated from a finite calibration bank is not the population mean in this definition. Its discrepancy must be included in the error term or bounded separately. The bound also does not follow from a large correlation between attribution and intervention scores.

A.2 Proof of T1

Taking count contrasts in Equation 4.2 gives

$$\Delta_r^{\mathrm{abl}}(H) = \Delta_r - F(H,r) + \mathcal{C}_r\xi_H^{\mathrm{abl}}, \qquad N(H,r) = F(H,r) - \mathcal{C}_r\xi_H^{\mathrm{abl}}.$$

For a high-count donor patched into a low-count target, Equation 4.3 gives

$$\Gamma_r^{\mathrm{HL}}(H) = F(H,r) + \mathbb{E}_{D_{r,b_L}}\xi_H^{b_H}.$$

For a low-count donor patched into a high-count target, the orientation in the definition of Γ_r^{LH} reverses the residual sign:

$$\Gamma_r^{\mathrm{LH}}(H) = F(H,r) - \mathbb{E}_{D_{r,b_H}}\xi_H^{b_L}.$$

Equation A.1 bounds all three differences from cumulative attribution $F(H,r)$. The mean patch gain $S(H,r)$ and raw utility $J_\lambda(H,r)$, for $\lambda \in [0,1]$, are convex combinations of these quantities, so their errors are bounded by the same $\varepsilon_{\mathrm{int}}(H,r)$. This proves Equation 4.4 and its stated consequences. □

A.3 Normalized scores and error amplification

For a nonzero clean contrast Δ, define $g_\Delta(u) = 1 - |\Delta - u|/|\Delta|$. The reverse triangle inequality gives

$$|g_\Delta(u) - g_\Delta(v)| \le \frac{|u-v|}{|\Delta|}. \tag{A.2}$$

Because $\Delta_r^{\mathrm{abl}} = \Delta_r - N$, shrink equals $g_{\Delta_r}(N)$. Each directional match equals $g_{\Delta_r}(\Gamma_r^\tau)$. Applying Equation A.2 to T1 yields

$$|\mathrm{shrink}(H,r) - g_{\Delta_r}(F(H,r))| \le \frac{\varepsilon_{\mathrm{int}}(H,r)}{|\Delta_r|}, \tag{A.3}$$

$$|\mathrm{match}^\tau(H,r) - g_{\Delta_r}(F(H,r))| \le \frac{\varepsilon_{\mathrm{int}}(H,r)}{|\Delta_r|}, \quad \tau \in \{\mathrm{HL},\mathrm{LH}\}. \tag{A.4}$$

Taking the weighted average with weights $\lambda, (1-\lambda)/2, (1-\lambda)/2$ proves Equation 4.5 for Q_λ. In particular, each normalized score is at least

$$1 - \frac{|\Delta_r - F(H,r)| + \varepsilon_{\mathrm{int}}(H,r)}{|\Delta_r|}.$$

Good recovery therefore needs both attribution alignment with the clean contrast and small intervention error. A large absolute attribution can instead overshoot the contrast. Shrink and match are at most one but have no finite lower bound; a negative match can reflect reversal or excessive movement in the intended direction. The average of two matches can pass a threshold while one direction fails it. None of these scores is defined at $\Delta_r = 0$; denominator floors used in computation define different, regularized statistics.

A.4 A conditional approximation by evidence position

Let $e_j(x) \in \{\mathtt{B},\mathtt{C}\}$ be the evidence token at position $j \in \{1,\dots,L_E\}$. Define the difference in positional token probabilities

$$\delta p_r(j) = \Pr_{D_{r,b_H}}(e_j = \mathtt{B}) - \Pr_{D_{r,b_L}}(e_j = \mathtt{B}).$$

Suppose an integrable scalar observable g admits the representation

$$g(x) = \beta_g(r) + \sum_{j=1}^{L_E}\phi_g(e_j(x), j; r) + \zeta_g(x), \tag{A.5}$$

where $\beta_g(r)$ does not depend on count, ϕ_g is a contribution assigned to one token and position, and ζ_g is the approximation residual. Define its position kernel $\kappa_g(r,j) = \phi_g(\mathtt{B}, j; r) - \phi_g(\mathtt{C}, j; r)$. Since the two token probabilities at each position sum to one,

$$\mathcal{C}_r g = \sum_{j=1}^{L_E}\kappa_g(r,j)\,\delta p_r(j) + \mathcal{C}_r\zeta_g. \tag{A.6}$$

The constant contribution cancels under the contrast, and the change in the probability of C is the negative change in the probability of B; these two facts prove the identity. Arbitrary g can be written in this form by placing all error in ζ_g. The useful assumption is a small residual contrast, not the existence of the representation.

For the full readout, define $K_\star(r,j) = \kappa_y(r,j)$ and $\Pi_\star(r) = \sum_j K_\star(r,j)\delta p_r(j)$. If $|\mathcal{C}_r\zeta_y| \le \varepsilon_\star(r)$, then

$$|\Delta_r - \Pi_\star(r)| \le \varepsilon_\star(r).$$

For head attribution, define $K_H(r,j) = \sum_{i\in H}\kappa_{s_i}(r,j)$ and $\Pi_H(r) = \sum_j K_H(r,j)\delta p_r(j)$. If $\sum_{i\in H}|\mathcal{C}_r\zeta_{s_i}| \le \varepsilon_{\mathrm{ker}}(H,r)$, then

$$|F(H,r) - \Pi_H(r)| \le \varepsilon_{\mathrm{ker}}(H,r). \tag{A.7}$$

Together with T1, Equation A.7 bounds the difference between each signed intervention gain and $\Pi_H(r)$ by $\varepsilon_{\mathrm{int}}(H,r) + \varepsilon_{\mathrm{ker}}(H,r)$. These are conditional approximations; the experiments do not estimate uniformly small values of either error. Higher-order token interactions can enter the residuals, and an additive representation need not be unique.

Proposition A.1 (Attribution change across regimes)**.** *Fix a head set H and two regimes s,t with the same counts and evidence length. Suppose Equation A.7 holds in both regimes. Define*

$$\begin{aligned} B_H(s,t) = &\sum_j |K_H(s,j)|\,|\delta p_t(j) - \delta p_s(j)| \\ &+ \sum_j |K_H(t,j) - K_H(s,j)|\,|\delta p_t(j)| \\ &+ \varepsilon_{\mathrm{ker}}(H,s) + \varepsilon_{\mathrm{ker}}(H,t). \end{aligned} \tag{A.8}$$

Then $|F(H,t) - F(H,s)| \le B_H(s,t)$.

Proof. Subtract the two profile approximations and use

$$\Pi_H(t) - \Pi_H(s) = \sum_j K_H(s,j)(\delta p_t(j) - \delta p_s(j)) + \sum_j (K_H(t,j) - K_H(s,j))\delta p_t(j).$$

The triangle inequality, followed by the two approximation bounds, gives the result. □

An attribution bound alone does not control normalized transfer: the clean contrast and intervention error can also change. Under T1 in both regimes, with $\Delta_s \neq 0$ and $\Delta_t \neq 0$, the same fixed set and λ satisfy

$$\begin{aligned} |Q_\lambda(H,t) - Q_\lambda(H,s)| \le &\frac{\varepsilon_{\mathrm{int}}(H,t)}{|\Delta_t|} + \frac{\varepsilon_{\mathrm{int}}(H,s)}{|\Delta_s|} + \frac{B_H(s,t)}{|\Delta_t|} \\ &+ \frac{|F(H,s)|\,|\Delta_t - \Delta_s|}{|\Delta_t|\,|\Delta_s|}. \end{aligned} \tag{A.9}$$

To verify the bound, write $g_\Delta(F) = 1 - |1 - F/\Delta|$, apply the reverse triangle inequality, and bound

$$\left|\frac{F(H,t)}{\Delta_t} - \frac{F(H,s)}{\Delta_s}\right| \le \frac{|F(H,t) - F(H,s)|}{|\Delta_t|} + \frac{|F(H,s)|\,|\Delta_t - \Delta_s|}{|\Delta_t|\,|\Delta_s|}.$$

Equation 4.5 supplies the two intervention-error terms. The bound concerns a fixed set evaluated with each regime's own donor distribution, not a claim that activations themselves transfer unchanged.

A.5 Proof of T2 and sampling conditions

Let $M = |\mathcal{H}|$, let $n \ge 1$ be the number of independent matched pairs, and let ν_r^2 be the sub-Gaussian variance proxy for each centered paired score in T2. Independence across pairs gives

$$\mathbb{E}\exp\Big(t\,[\widehat{\Delta s}_i(r) - \Delta s_i(r)]\Big) \le \exp\left(\frac{t^2\nu_r^2}{2n}\right).$$

For $\nu_r > 0$, the Chernoff bound applied to both signs yields

$$\Pr\Big(|\widehat{\Delta s}_i(r) - \Delta s_i(r)| > u\Big) \le 2\exp\left(-\frac{nu^2}{2\nu_r^2}\right).$$

A union bound over M heads and the choice $u = \nu_r\sqrt{2\log(2M/\delta)/n}$ prove Equation 4.6. If $\nu_r = 0$, every centered paired score is zero almost surely and the claim holds with zero error. No independence between heads is used.

On the uniform event, reverse triangle inequality bounds the error of each absolute score by $\eta_{n,\delta}(r)$. If the population absolute-score gap between ranks m and $m+1$ exceeds twice this error, every population top-m head still has a larger empirical absolute score than every excluded head. The shortlist is therefore recovered exactly. A strict boundary gap removes ambiguity from ties at that boundary. □

Independent samples as a special case. Suppose each centered high-count and low-count head score is sub-Gaussian with variance proxy σ_r^2. Independent high and low samples give a paired difference with proxy $\nu_r^2 = 2\sigma_r^2$, so

$$\eta_{n,\delta}(r) = 2\sigma_r\sqrt{\frac{\log(2M/\delta)}{n}}.$$

For dependent observations within a pair, Cauchy–Schwarz instead gives, for centered scores X, Y,

$$\mathbb{E}e^{t(X-Y)} \le (\mathbb{E}e^{2tX})^{1/2}(\mathbb{E}e^{-2tY})^{1/2} \le e^{2t^2\sigma_r^2}.$$

Thus $\nu_r = 2\sigma_r$ is a valid general bound; a smaller proxy needs information about the pairing. In either case, different pairs must remain independent for the stated concentration argument. Duplicate deterministic prompts contain no variation over prompts, regardless of how many row identifiers they receive. T2 neither validates a richer target distribution nor controls the later intervention-based screening and greedy search.

A.6 Proof and limits of T3

Fix the screened pool $\mathcal{P}_r$, budget $1 \le k \le |\mathcal{P}_r|$, and $\lambda \in [0,1]$. Write $J(H) = J_\lambda(H, r)$ within this proof. Empty interventions leave the model unchanged, so $J(\varnothing) = 0$. Monotonicity means $J(A) \le J(B)$ for $A \subseteq B$. The required weak-submodularity condition is

$$\sum_{i\in B}\big[J(A\cup\{i\}) - J(A)\big] \ge \gamma_r\big[J(A\cup B) - J(A)\big], \tag{A.10}$$

for disjoint $A, B \subseteq \mathcal{P}_r$ with $|A| \le k-1$, $|B| \le k$, and $\gamma_r \in (0,1]$. In particular, the condition includes unions with as many as $2k-1$ heads, not only sets within budget k. Monotonicity is required on these sets as well. The proof is the standard greedy argument for the submodularity ratio (Das and Kempe, 2018); its application here is conditional on the raw utility satisfying the assumptions.

Let $G_0 = \varnothing$, and let exact greedy search add a head with maximum marginal raw utility to form G_{t+1} from G_t. Let O_k maximize J among pool subsets of size at most k. For $t < k$, Equation A.10 and monotonicity give

$$\sum_{i\in O_k\setminus G_t}\big[J(G_t\cup\{i\}) - J(G_t)\big] \ge \gamma_r[J(G_t\cup O_k) - J(G_t)] \ge \gamma_r[J(O_k) - J(G_t)].$$

If $O_k \subseteq G_t$, the remaining optimality gap is zero. Otherwise at most k summands appear, so greedy search obtains

$$J(G_{t+1}) - J(G_t) \ge \frac{\gamma_r}{k}[J(O_k) - J(G_t)].$$

Starting from $J(G_0) = 0$ and iterating the gap recurrence yields

$$J(G_k) \ge \left[1 - \left(1 - \frac{\gamma_r}{k}\right)^k\right] J(O_k) \ge (1 - e^{-\gamma_r})J(O_k).$$

The last step uses $(1 - a/k)^k \le e^{-a}$ and $J(O_k) \ge 0$, which follows from monotonicity and normalization. This proves Equation 4.7.

For the attribution comparison, set $\alpha_r = 1 - e^{-\gamma_r}$ and let $\epsilon_{J,k}(r)$ be the uniform difference between raw utility and cumulative attribution assumed in T3. Apply the error bound at the greedy set and at the same raw-utility optimum:

$$\begin{aligned} F(G_k, r) &\geq J_\lambda(G_k, r) - \epsilon_{J,k}(r) \\ &\geq \alpha_r J_\lambda(O_k, r) - \epsilon_{J,k}(r) \\ &\geq \alpha_r F(O_k, r) - (1 + \alpha_r)\epsilon_{J,k}(r). \end{aligned}$$

This proves Equation 4.8; O_k is not replaced by an attribution optimum. □

What the bound does not cover. BUC searches empirical normalized utility Q_λ, rather than exact raw utility J_λ. Absolute-value penalties can make the objectives rank sets differently even without sampling error. The theorem also conditions on the screened pool: it says nothing about heads excluded earlier, a minimum circuit size, or uniqueness. Noisy greedy evaluations require additional bounds, and the observed nonmonotonicity of $Q_{1/2}$ neither proves nor disproves monotonicity of the different population utility $J_{1/2}$. The finite audits in Appendix E.5 describe the tested pools and data, not a uniform guarantee.

B Additional experimental details and results

This appendix gives the BF16 configurations and results supporting Section 5: tokenization audits, data splits, head identities, transfer matrices, attribution diagnostics, baselines, input-level interventions, and count, order, tail-length, surface-form, and model comparisons. It also describes compute requirements and reproduction coverage.

Protocol scope. The configurations and tables in this appendix belong to the original BF16 study. The separately recorded native FP32 configurations, calibration banks, and controls are specified in Appendix E; the two configurations are analyzed separately.

B.1 Prompt, model, and tokenization details

All experiments use the controlled evidence aggregation prompt family from Section 3. Each prompt has the fixed form

$$\begin{aligned} \text{Prompt} = &\texttt{Tokens:\textbackslash n} + \text{(space-separated symbolic sequence)} \\ &+ \texttt{\textbackslash nChoose one word: various or successful\textbackslash nAnswer:}. \end{aligned}$$

The symbolic sequence length is $L_{\text{seq}} = 250$, with evidence length $L_E = 200$ and A-only tail length $L_A = 50$. The two answer strings are `various` and `successful`. For each model, we audit whether the strings `various` and `successful`, including the leading space, are single-token candidates under the model tokenizer.

B.2 Data splits, selection pipeline, and control construction

For each regime $r \in \{\text{FRONT}, \text{BACK}, \text{INTERLEAVE}, \text{RANDOM}\}$, head selection and final evaluation use disjoint prompt batches. The selection split contains 48 prompts from D_{r,b_H} and 48 prompts from D_{r,b_L}. The final evaluation split contains 96 prompts from D_{r,b_H} and 96 prompts from D_{r,b_L}. A full-layer audit uses 16 prompts from each evidence level to decide whether middle-layer heads should be included. In the Pythia-410M runs, the audit expands the candidate set to all 24 layers.

The selected head set $H_{\text{BUC}}(r)$ is obtained in three stages. First, heads are ranked by empirical absolute sensitivity $|\widehat{\Delta s}_i(r)|$, and the top 192 heads are kept. Second, individual heads are scored by the joint removal and recovery score, and the top 48 heads form the greedy candidate pool. Third, forward greedy selection is run with budget 12. The selected set is evaluated on a held-out evaluation split.

Layer-matched random controls. For a selected set $H_{\text{BUC}}(r)$, the layer-matched random control samples the same number of heads from each selected layer. For example, if $H_{\text{BUC}}(r)$ contains two heads from layer 10, one head from layer 15, and one head from layer 19, the matched control samples two random heads from layer 10, one from layer 15, and one from layer 19. The layer-matched control keeps both head count and layer pattern fixed.

Table 8: Prompt and tokenization audit for the main Pythia-410M experiment. The SEQ token slice refers to symbolic-sequence positions and excludes the final prompt position.

	Layers	Heads per layer	Hook tensor
Value	24	16	`attn.hook_z`
Notes	no fine-tuning	total 384 heads	projected through W_O before logit-direction projection
	Evaluation dtype	**Prefix token length**	**Symbolic sequence token length**
Value	`bfloat16`	5	250
Notes	inference only	tokenizer audit	one token per symbol in the tested tokenizer
	Suffix token length	**SEQ token slice**	**Answer candidates**
Value	11	$[5 : 255)$	`various`, `successful`
Notes	tokenizer audit	symbolic positions only; excludes the final prompt position	both single-token under the tested tokenizer

Table 9: Tokenizer audit for the model sweep. The answer pair is usable only when both answer strings are single-token candidates under the model tokenizer.

Model	Single token	w_+	w_-	Notes
`EleutherAI/pythia-70m-deduped`	yes	`various`	`successful`	audited
`EleutherAI/pythia-160m-deduped`	yes	`various`	`successful`	audited
`EleutherAI/pythia-410m-deduped`	yes	`various`	`successful`	main model
`EleutherAI/pythia-1b-deduped`	yes	`various`	`successful`	audited
`gpt2`	yes	`various`	`successful`	out-of-family baseline
`Qwen/Qwen3-4B-Base`	yes	`various`	`successful`	modern open-weight baseline

Table 10: Selection and evaluation configuration for the main Pythia-410M experiment.

Configuration item	Value
Evidence setting	MIXED $(b_L, b_H) = (50, 150)$
Ordering regimes	FRONT, BACK, INTERLEAVE, RANDOM
Seeds	$\{0, 1, 2, 3, 4\}$
Selection prompts per evidence level	48
Evaluation prompts per evidence level	96
Audit prompts per evidence level	16
Sensitivity shortlist size	192
Single-head candidate pool size	48
Greedy budget	12
Layer-matched random draws	120 per seed
Primary intervention scope	LAST
Additional intervention scope	SEQ on symbolic positions $[5 : 255)$

B.3 Base evidence effects and directional patching

Table 11 reports the base evidence effects used in the main Pythia-410M experiment. The quantity Δ_r is the average high-minus-low restricted two-answer log-odds shift.

Table 12 reports directional sufficiency. The main text reports the average of high-to-low and low-to-high recovery.

B.4 Full white-box result with matched random controls

Table 13 reports set sizes, intervention scores, and layer-matched controls. BUC match averages both patch directions, whereas the random control uses HL match.

Table 11: Base evidence effects in the MIXED setting $(b_L, b_H) = (50, 150)$ on `pythia-410m-deduped`.

Regime r	Δ_r	Notes
FRONT	0.239 ± 0.000	fixed deterministic input
BACK	1.965 ± 0.000	large shift
INTERLEAVE	1.659 ± 0.000	large shift
RANDOM	3.184 ± 0.097	largest and seed-dependent

Table 12: Directional mean-patching recovery in the MIXED setting. The average of the two directional scores is the BUC match score reported in the main text.

Regime r	match^{HL}	match^{LH}	$\text{match}_{\text{avg}}$
FRONT	0.998 ± 0.000	0.950 ± 0.000	0.974 ± 0.000
BACK	0.765 ± 0.000	0.753 ± 0.000	0.759 ± 0.000
INTERLEAVE	0.937 ± 0.000	0.891 ± 0.000	0.914 ± 0.000
RANDOM	0.813 ± 0.092	0.833 ± 0.080	0.823 ± 0.084

Table 13: Full white-box result with matched random controls in the MIXED setting. BUC values are mean $\pm$ standard deviation across five seeds. Random match is measured in the high-to-low direction.

Regime r	$\lvert H_{\text{BUC}}(r)\rvert$	BUC shrink	BUC match	Random shrink	Random HL match
FRONT	1.00 ± 0.00	0.906 ± 0.000	0.974 ± 0.000	0.020	0.048
BACK	12.00 ± 0.00	0.655 ± 0.000	0.759 ± 0.000	0.043	0.075
INTERLEAVE	3.00 ± 0.00	0.916 ± 0.000	0.914 ± 0.000	0.061	0.051
RANDOM	7.00 ± 1.00	0.866 ± 0.021	0.823 ± 0.084	0.069	0.061

B.5 Selected head identities and stability

Table 14 reports selected head identities. A head is written as (ℓ, h), with zero-indexed layer and head indices. The deterministic regimes use the same set across all five seeds. For RANDOM, the explicitly listed sets cover seeds 0–2; Table 15 summarizes only those displayed sets.

Table 14: Selected BUC head sets in the MIXED setting. Heads are written as (ℓ, h). For FRONT, BACK, and INTERLEAVE, the listed set is identical across five seeds.

Regime r	Seed(s)	$H_{\text{BUC}}(r)$
FRONT	0–4	$(10, 2)$
BACK	0–4	$(19, 6), (20, 2), (13, 5), (15, 3), (19, 9), (17, 3), (15, 10), (15, 5), (19, 14), (14, 13), (18, 1), (11, 8)$
INTERLEAVE	0–4	$(17, 3), (10, 2), (15, 10)$
RANDOM	0	$(17, 3), (14, 7), (10, 2), (19, 14), (16, 1), (15, 10), (19, 6), (16, 14)$
RANDOM	1	$(17, 3), (14, 7), (12, 14), (19, 14), (10, 2), (15, 10), (19, 6)$
RANDOM	2	$(17, 3), (14, 7), (12, 14), (19, 14), (10, 2), (15, 10)$

Table 15: Stability of the head sets explicitly listed in Table 14. Deterministic-regime sets are identical across seeds 0–4. The RANDOM value is the minimum over listed seeds 0–2 only; seeds 3 and 4 are not included in this calculation.

Regime r	Minimum pairwise Jaccard	Interpretation
FRONT	1.000	identical selected set
BACK	1.000	identical selected set
INTERLEAVE	1.000	identical selected set
RANDOM	0.556	listed seeds 0–2 only

Pairwise overlap. For the three RANDOM sets listed in Table 14, the pairwise overlaps for seeds $(0, 1)$, $(0, 2)$, and $(1, 2)$ are $6/9$, $5/9$, and $6/7$. Thus

$$\min\left\{\frac{6}{9}, \frac{5}{9}, \frac{6}{7}\right\} = \frac{5}{9} \approx 0.556.$$

The minimum is attained by seeds 0 and 2. Table 15 excludes RANDOM seeds 3 and 4. The overlap statistic covers three selected sets, whereas the intervention-effect summaries cover five executions.

B.6 Mechanistic diagnostics of selected heads

The main white-box tests measure the removal and recovery effects of selected head sets. The following post-selection diagnostics examine how the selected heads act at the final prompt position. The diagnostics separate two components of an attention head: QK routing, which determines where the final prompt position attends, and OV writing, which determines what vector is written into the residual stream and projected onto the logit-difference direction d.

Let $R_B(x)$, $R_C(x)$, and $R_E(x) = R_B(x) \cup R_C(x)$ denote the B-positions, C-positions, and all evidence-region positions in prompt x. For a selected head or head group, $\Delta\mathsf{Attn}(R, r)$ denotes the high-minus-low change in attention at the final prompt position mass to region R, and $\Delta\mathsf{OV}(R, r)$ denotes the high-minus-low change in source-wise OV contribution along d. QK / OV match reports pattern-only and OV-only patch recovery, using the same recovery-match convention as the main text.

Representative heads. Table 16 reports post-selection diagnostics for the heads in Table 14. Multi-head rows are group-level measurements.

Table 16: Mechanistic diagnostics for representative selected heads at the answer position. $\Delta\mathsf{Attn}$ is a high-minus-low attention-mass contrast, and $\Delta\mathsf{OV}$ is a high-minus-low source-wise OV contribution contrast along d. The attention and OV difference columns report $\Delta\mathsf{Attn}(R_B, r) - \Delta\mathsf{Attn}(R_C, r)$ and $\Delta\mathsf{OV}(R_B, r) - \Delta\mathsf{OV}(R_C, r)$, respectively. QK / OV match reports pattern-only and OV-only patch recovery.

Regime r	Head(s)	Main routed region	Patch comparison
FRONT	$(10, 2)$	$R_B \cup R_C$	Larger OV patch match
BACK	$(19, 6), (20, 2), (13, 5)$	R_C-side routing shift	Positive QK and OV matches
INTERLEAVE	$(17, 3), (10, 2), (15, 10)$	Broad over R_E	Larger OV patch match
RANDOM	$(17, 3), (14, 7), (19, 14)$	Broad over R_E	Larger OV patch match

Regime r	Attention difference	OV difference	QK / OV match
FRONT	$+0.42$	$+1.28$	0.35 / 0.88
BACK	-0.65	$+2.15$	0.72 / 0.61
INTERLEAVE	$+0.05$	$+1.85$	0.12 / 0.92
RANDOM	$+0.08$	$+2.30$	0.18 / 0.85

Interpretation. OV-only patching has larger match than QK-only patching in the listed FRONT, INTERLEAVE, and RANDOM examples; both split interventions have substantial effects in BACK. These measurements distinguish routing and writing interventions for the tested heads. They do not provide an additive decomposition of the complete effect or identify all earlier computations.

B.7 Final-position intervention scope

Table 17 reports the reported LAST-versus-SEQ scope diagnostic. The same head set $H_{\mathrm{BUC}}(r)$ is used under both intervention scopes. The LAST intervention modifies only the answer position. The SEQ intervention modifies symbolic-sequence positions and excludes the final prompt position.

The RANDOM row is retained as the separately reported scope diagnostic. Its size range 7–8 and LAST shrink range 0.835–0.917 are not documented as the same five-execution cohort summarized in Table 13; the listed head identities also include a six-head set. The supplied aggregates do not resolve that cohort difference. The range must not be treated as the five-seed uncertainty or as an additional stability result.

Table 17: LAST versus SEQ using the same selected head set. The SEQ effect is near zero in every regime.

Regime r	$k^\star$	LAST shrink	SEQ shrink	SEQ match
FRONT	1	0.906	0.059	-0.006
BACK	12	0.655	0.022	$+0.003$
INTERLEAVE	3	0.916	0.001	$+0.006$
RANDOM	7–8	0.835–0.917	$\approx$ 0.008–0.010	≈ -0.001

B.8 Full cross-regime transfer

The transfer experiment selects $H_{\mathrm{BUC}}(r_s)$ on source regime r_s, then evaluates mean-patching recovery on target regime r_t. A negative recovery-match score indicates severe mismatch, not necessarily a sign reversal.

Table 18: Representative off-diagonal cross-regime transfer matrix under LAST interventions (seed 0). Rows are source regimes; columns are target regimes. Entries are recovery-match values. Diagonal effects are reported separately by direction in Table 12.

Source \ Target	FRONT	BACK	INTERLEAVE	RANDOM
FRONT	–	$+0.129$	$+0.173$	$+0.066$
BACK	$+0.958$	–	$+0.961$	$+0.850$
INTERLEAVE	-0.467	$+0.230$	–	$+0.611$
RANDOM	-0.925	$+0.413$	$+0.687$	–

Table 19: Off-diagonal cross-regime transfer matrix averaged over seeds. Entries are mean $\pm$ standard deviation of the recovery match; within-regime directional effects are in Table 12.

Source \ Target	FRONT	BACK	INTERLEAVE	RANDOM
FRONT	–	0.129 ± 0.000	0.173 ± 0.000	0.066 ± 0.001
BACK	0.958 ± 0.000	–	0.961 ± 0.000	0.851 ± 0.002
INTERLEAVE	-0.467 ± 0.000	0.230 ± 0.000	–	0.616 ± 0.006
RANDOM	-0.912 ± 0.021	0.428 ± 0.015	0.695 ± 0.011	–

B.9 Many-set T1 diagnostic

Under its error assumption, Theorem 4.2 bounds how closely cumulative attribution $F(H, r)$ tracks the gains produced by ablation and mean patching. The diagnostic instead measures correlation, which is not an error bound. It samples multiple head sets H, computes $F(H, r)$, and compares $F(H, r)$ to $N(H, r)$, $\Gamma_r^{\mathrm{HL}}(H)$, and $\Gamma_r^{\mathrm{LH}}(H)$.

Table 20: Many-set T1 diagnostic. Pearson and Spearman correlations are computed over 500 sampled head sets H on held-out batches.

Regime r	Sets	Pearson $\rho(F, N)$	Pearson $\rho(F, \Gamma^{\mathrm{HL}})$	Pearson $\rho(F, \Gamma^{\mathrm{LH}})$	Spearman $\rho(F, N)$	Spearman $\rho(F, S)$
FRONT	500	0.923	0.851	0.884	0.941	0.892
BACK	500	0.875	0.782	0.816	0.899	0.835
INTERLEAVE	500	0.908	0.814	0.842	0.927	0.861
RANDOM	500	0.841	0.755	0.793	0.865	0.804

B.10 Selection baselines

Table 21 compares BUC with budget-matched baselines. Sensitivity-only selects the top-k heads by $|\widehat{\Delta s_i}(r)|$, where $k = |H_{\mathrm{BUC}}(r)|$.

ActPatch is the BF16 single-head activation-patching baseline. Among the 192 sensitivity-shortlisted heads, ActPatch ranks heads by single-head high-to-low (HL) patch recovery on the selection split, takes the top $k = |H_{\mathrm{BUC}}(r)|$, and evaluates that set jointly on the held-out split. The reported ActPatch value is the set-level $\mathrm{shrink} + \mathrm{match}_{\mathrm{avg}}$ score, including both patch directions in the evaluation match. Single-head HL scores determine ranking; the reported bidirectional score is evaluated jointly on the selected set, rather than obtained by adding single-head effects.

Same-layer top norm selects the largest write-vector norms in the selected layers; layer-random preserves the selected layer counts. No method-specific quantitative comparison with EAP or ACDC is claimed. The separate FP32 single-head ranking rules are defined in Appendix E.6.

Table 21: BF16 selection baselines under matched head budgets. BUC, sensitivity-only, ActPatch, and top-norm use $\mathrm{shrink} + \mathrm{match}_{\mathrm{avg}}$. Layer-random uses shrink plus *HL-only* match and is a directional control, not an identical-objective comparison.

Regime r	BUC	Sens.-only	ActPatch	Top-norm	Layer-rand (HL)
FRONT	1.880	1.880	1.865	0.451	0.068
BACK	1.414	1.352	1.285	0.523	0.118
INTERLEAVE	1.830	1.765	1.695	0.614	0.112
RANDOM	1.689	1.611	1.520	0.488	0.130

B.11 Sensitivity to the joint-utility trade-off

The main experiments use the balanced joint utility with $\lambda = 1/2$. This subsection checks whether the selected head sets and held-out intervention scores are sensitive to the choice of λ. For each value of λ, we rerun the greedy selection stage on the same candidate pool and evaluate the selected set on the same held-out split as in the main experiment. The normalized empirical selection score is

$$Q_\lambda(H, r) := \lambda\,\mathrm{shrink}(H, r) + (1 - \lambda)\,\mathrm{match}_{\mathrm{avg}}(H, r), \tag{B.1}$$

where $\lambda = 0$ emphasizes recovery, $\lambda = 1$ emphasizes removal, and $\lambda = 1/2$ is the balanced setting used in the main text. For each regime r, let $H_\lambda(r)$ denote the selected head set under $Q_\lambda(H, r)$. We also report the Jaccard overlap with the balanced set $H_{1/2}(r)$:

$$\mathrm{Jac}_\lambda(r) := \frac{|H_\lambda(r) \cap H_{1/2}(r)|}{|H_\lambda(r) \cup H_{1/2}(r)|}. \tag{B.2}$$

Table 22: Sensitivity to the joint-utility trade-off λ. Values are averaged over FRONT, BACK, INTERLEAVE, and RANDOM in the MIXED setting. The balanced setting ($\lambda = 0.50$) matches the primary causal sets used in the main text.

λ	Selection emphasis	Mean $\lvert H_\lambda \rvert$	Mean shrink	Mean match	Mean Jaccard to $H_{1/2}$
0.00	recovery only	6.25	0.725	0.895	0.650
0.25	recovery weighted	5.75	0.810	0.880	0.885
0.50	balanced	5.75	0.836	0.868	1.000
0.75	removal weighted	5.50	0.865	0.812	0.820
1.00	removal only	5.00	0.890	0.645	0.580

Table 22 reports the tested weight grid. Relative to $\lambda = 1/2$, the values $\lambda = 0.25$ and $\lambda = 0.75$ retain mean Jaccard overlaps of 0.885 and 0.820. Increasing the removal weight trades patching match for shrink. The comparison is limited to the tested objective weights.

B.12 Greedy-prefix curves

The greedy-prefix diagnostic evaluates the prefix sequence of heads selected by the greedy algorithm. Let $H_k(r)$ denote the first k heads selected under regime r. The diagnostic reports $\mathrm{shrink}(H_k, r)$ and $\mathrm{match}_{\mathrm{avg}}(H_k, r)$ for each k. FRONT and INTERLEAVE have small increments after the first few heads; BACK continues improving through the maximum tested budget of 12. The measurements do not establish a plateau beyond that budget.

Table 23: Greedy-prefix scores at the evaluated head budgets.

Regime r	Metric	$k=1$	$k=2$	$k=3$	$k=4$	$k=6$	$k=8$	$k=12$
FRONT	shrink	0.906	0.908	0.910	0.911	0.915	0.918	0.920
FRONT	match	0.974	0.975	0.976	0.976	0.978	0.979	0.981
BACK	shrink	0.142	0.285	0.371	0.463	0.541	0.598	0.655
BACK	match	0.165	0.312	0.404	0.518	0.627	0.693	0.759
INTERLEAVE	shrink	0.465	0.781	0.916	0.919	0.924	0.927	0.931
INTERLEAVE	match	0.482	0.815	0.914	0.917	0.920	0.922	0.926
RANDOM	shrink	0.254	0.448	0.612	0.745	0.841	0.875	0.892
RANDOM	match	0.281	0.472	0.635	0.721	0.803	0.835	0.851

B.13 Black-box input-level interventions

The black-box experiments use only input-level changes. Let $S \subseteq \{1, \ldots, L_E\}$ denote a subset of evidence positions. In masking, the positions in S are replaced by A. In add-back, the positions in S are copied from a high-evidence donor prompt into a low-evidence target prompt. Let $x^{\mathrm{mask}(S)}$ be a masked input and let $x_L \leftarrow_S x_H$ denote the low-count input with positions S copied from its high-count donor. Then

$$\Delta_r^{\mathrm{mask}}(S) = \mathcal{C}_r[y(x^{\mathrm{mask}(S)})], \qquad \Delta_r^{\mathrm{add}}(S) = \mathbb{E}[y(x_L \leftarrow_S x_H) - y(x_L)],$$

where the second expectation uses the donor–target pairs of the input experiment. The supplied BF16 aggregates do not specify the complete donor pairing registry.

The input removal score is

$$\mathrm{shrink}_{\mathrm{bb}}(S, r) := 1 - \frac{|\Delta_r^{\mathrm{mask}}(S)|}{|\Delta_r|}. \tag{B.3}$$

The input recovery score is

$$\mathrm{rec}_{\mathrm{bb}}(S, r) := \frac{\Delta_r^{\mathrm{add}}(S)}{\Delta_r}. \tag{B.4}$$

Input recovery is a signed ratio and does not penalize overshooting the clean contrast. It is not the head-patching match metric and does not establish that the selected heads explain the input intervention. All normalized input scores require $\Delta_r \neq 0$.

B.14 Greedy Occlusion algorithm

This subsection gives the search procedure used for the input-level black-box validation in Section 5. Greedy Occlusion uses only input-output queries to the scalar output $y(x)$. The procedure does not use gradients, attention weights, head activations, or any other internal model states.

Algorithm 1: Greedy Occlusion for input-level subset selection

Input: Regime r, evidence length $L_E = 200$, budget $K_{\mathrm{bb}} = 50$, selection split, scalar output $y(x)$.
Output: Selected evidence-position subset $S^\star(r)$.
Initialize $S_0 \leftarrow \emptyset$;
Compute the empirical base effect $\widehat{\Delta}_r$ on the selection split;
for $m = 1, \ldots, K_{\mathrm{bb}}$ **do**
  foreach $j \in \{1, \ldots, L_E\} \setminus S_{m-1}$ **do**
    Construct masked prompts by replacing positions in $S_{m-1} \cup \{j\}$ with A;
    Compute $\widehat{\mathrm{shrink}}_{\mathrm{bb}}(S_{m-1} \cup \{j\}, r)$ from input-output queries to $y(x)$;
  Select

$$j_m \in \arg\max_{j \notin S_{m-1}} \widehat{\mathrm{shrink}}_{\mathrm{bb}}(S_{m-1} \cup \{j\}, r)$$

  ;
  Set $S_m \leftarrow S_{m-1} \cup \{j_m\}$;
Return $S^\star(r) \leftarrow S_{K_{\mathrm{bb}}}$;

For a regime $r \in \mathcal{R}$, let $S \subseteq \{1, \ldots, L_E\}$ denote a subset of evidence positions. Let $\widehat{\Delta}_r$ denote the empirical evidence effect on the black-box selection split. Let $\widehat{\Delta}_r^{\mathrm{mask}}(S)$ denote the empirical evidence effect after replacing positions in S by A. The masking objective used by Greedy Occlusion is

$$\widehat{\mathrm{shrink}}_{\mathrm{bb}}(S, r) := 1 - \frac{\left|\widehat{\Delta}_r^{\mathrm{mask}}(S)\right|}{\left|\widehat{\Delta}_r\right|}. \tag{B.5}$$

The selected subset is built by forward greedy addition. Starting from the empty set, Greedy Occlusion adds one evidence position at a time:

$$j_m \in \arg\max_{j \in \{1,\ldots,L_E\} \setminus S_{m-1}} \widehat{\mathrm{shrink}}_{\mathrm{bb}}(S_{m-1} \cup \{j\}, r), \qquad S_m := S_{m-1} \cup \{j_m\}. \tag{B.6}$$

The algorithm stops after $K_{\mathrm{bb}} = 50$ steps and returns

$$S^\star(r) := S_{K_{\mathrm{bb}}}.$$

The add-back recovery score is not used to select $S^\star(r)$. Add-back is evaluated only after selection, on a held-out evaluation split.

The naive number of candidate subset evaluations for one regime is

$$\sum_{m=0}^{K_{\mathrm{bb}}-1} (L_E - m) = \frac{K_{\mathrm{bb}}(2L_E - K_{\mathrm{bb}} + 1)}{2}. \tag{B.7}$$

With $L_E = 200$ and $K_{\mathrm{bb}} = 50$, the number of candidate subset evaluations is 8775 per regime. Each candidate evaluation uses a prompt batch, so 8,775 is not a count of individual model queries. Batching changes execution cost but not the number of candidate subsets.

After $S^\star(r)$ is selected, the held-out masking and recovery scores are

$$\mathrm{shrink}_{\mathrm{bb}}(S^\star, r) := 1 - \frac{\left|\Delta_r^{\mathrm{mask}}(S^\star)\right|}{|\Delta_r|}, \qquad \mathrm{rec}_{\mathrm{bb}}(S^\star, r) := \frac{\Delta_r^{\mathrm{add}}(S^\star)}{\Delta_r}. \tag{B.8}$$

Matched random controls use random subsets $S_{\mathrm{rand}} \subseteq \{1, \ldots, L_E\}$ with $|S_{\mathrm{rand}}| = |S^\star(r)| = 50$, and are evaluated with the same held-out masking and add-back metrics.

Table 24: Black-box masking and add-back results using Greedy Occlusion position selection.

Regime r	Subset family	$\lvert S^\star(r)\rvert$	BB shrink	BB recovery	Random shrink	Random recovery
FRONT	Greedy Occlusion	50	0.852	0.871	0.245	0.251
BACK	Greedy Occlusion	50	0.915	0.898	0.258	0.262
INTERLEAVE	Greedy Occlusion	50	0.840	0.855	0.249	0.255
RANDOM	Greedy Occlusion	50	0.795	0.812	0.252	0.248

Table 25: Candidate subset families for masking and add-back. White-box-informed positions are diagnostic only and are not used for the strict black-box claim.

Subset family	Description	Status
Contiguous spans	$S = \{a, a+1, \ldots, a+\ell-1\}$	Evaluated
Union of spans	S is a union of several disjoint contiguous spans	Evaluated
Prefix/suffix evidence spans	Spans anchored at the beginning or end of the evidence region	Low performance
White-box-informed positions	Positions scored from internal signals; diagnostic only	Not strict black-box
Matched random subsets	Same size or same span length as $S^\star(r)$	Baseline

B.15 Order effects at fixed evidence count

The order-effect diagnostic separates evidence count from evidence order. For two regimes $r_1, r_2 \in \mathcal{R}$ and a fixed count b, define

$$\Delta^{\text{order}}(r_1 \to r_2; b) := \mathbb{E}_{x \sim D_{r_2,b}}[y(x)] - \mathbb{E}_{x \sim D_{r_1,b}}[y(x)]. \tag{B.9}$$

Table 26: Order effects at fixed evidence count. Positive shifts indicate that the target regime produces a larger output log-odds response than the source regime.

Regime change	$\Delta^{\text{order}}(\cdot; b_L = 50)$	$\Delta^{\text{order}}(\cdot; b_H = 150)$
FRONT $\to$ BACK	-0.258	$+1.468$
FRONT $\to$ INTERLEAVE	-0.194	$+1.226$
BACK $\to$ INTERLEAVE	$+0.064$	-0.242
INTERLEAVE $\to$ RANDOM	$+0.321$	$+1.846$

B.16 Evidence-gap robustness

The main experiment uses the MIXED setting $(b_L, b_H) = (50, 150)$. Table 27 reports robustness across alternate evidence gaps.

Table 27: Evidence-gap robustness. Values are averaged over the four ordering regimes.

Setting	(b_L, b_H)	Mean $\lvert H\rvert$	Mean shrink	Mean match	Notes
MIXED	$(50, 150)$	5.75	0.836	0.868	main configuration
EXTREME	$(0, 150)$	6.25	0.892	0.915	larger contrast
SMALL-GAP	$(75, 125)$	4.50	0.712	0.745	lower signal-to-noise
COUNT-SWEEP	multiple	5.50	0.810	0.835	increments of 10

B.17 A-only tail length ablation

The main setting uses an A-only tail of length $L_A = 50$. Table 28 evaluates a tail-length ablation. The $L_A = 50$ row repeats the main MIXED result.

Table 28: Tail-length ablation. Intervention scores remain positive over the tested tail lengths; the mean output contrast decreases as the tail grows.

Tail length L_A	Mean Δ_τ	Mean $\lvert H\rvert$	Mean shrink	Mean match
0	1.950	5.25	0.850	0.880
25	1.850	5.50	0.842	0.875
50	1.762	5.75	0.836	0.868
100	1.520	6.50	0.815	0.840

B.18 Naturalized surface validation

The main task uses the alphabet $\Sigma = \{A, B, C\}$. Table 29 reports alternative token surfaces with the same count/order structure, testing the dependence of the results on single-letter inputs.

Table 29: Surface-form validation. The tested word and sentence surfaces yield larger selected sets and positive intervention effects; set size does not measure all processing in the model.

Surface	Example evidence tokens	Model	Mean $\lvert H\rvert$	Mean shrink	Mean match
ABC symbolic	B, C, A	`pythia-410m-deduped`	5.75	0.836	0.868
Word-symbolic	`blue`, `red`, `gray`	`pythia-410m-deduped`	6.25	0.812	0.845
Signal-token	`signal-x`, `signal-y`, `neutral`	`pythia-410m-deduped`	6.00	0.820	0.852
Sentence-template	short evidence sentences	`pythia-410m-deduped`	8.50	0.765	0.790

B.19 Model sweep

Table 30 gives the aggregate model-sweep results, including GPT-2 and Qwen3-4B-Base. Table 31 expands the Pythia sweep to per-regime rows. Table 32 expands per-regime rows for GPT-2 and Qwen3.

Aggregation scope. Each numerical aggregate in Table 30 is the unweighted arithmetic mean of the four displayed per-regime values in Tables 31–32, rounded only after averaging. The aggregates describe the tested regimes and are not used for cross-model significance tests. BUC match is bidirectional; random match is HL only. The Pythia-410M primary-configuration averages repeated in Tables 22, 27, 28, and 29 use the same displayed-row calculation.

Table 30: Aggregate model sweep. Each numerical aggregate is the unweighted arithmetic mean of the four corresponding displayed per-regime values in Tables 31–32. Random match uses HL only.

Model	Token pair	Mean $\lvert H\rvert$	BUC shrink	BUC match	Random shrink	Random HL match	LAST effect
`EleutherAI/pythia-70m-deduped`	yes	3.50	0.650	0.700	0.030	0.040	yes
`EleutherAI/pythia-160m-deduped`	yes	4.25	0.720	0.750	0.035	0.045	yes
`EleutherAI/pythia-410m-deduped`	yes	5.75	0.836	0.868	0.048	0.059	yes
`EleutherAI/pythia-1b-deduped`	yes	6.50	0.880	0.910	0.040	0.050	yes
`gpt2`	yes	4.00	0.750	0.780	0.035	0.045	yes
`Qwen/Qwen3-4B-Base`	yes	7.50	0.895	0.925	0.035	0.042	yes

Table 31: Per-regime Pythia model sweep.

Model	Regime r	Token pair	$\lvert H\rvert$	BUC shrink	BUC match	Random shrink	Random HL match
pythia-70m	FRONT	yes	1.00	0.751	0.812	0.012	0.024
pythia-70m	BACK	yes	5.00	0.552	0.614	0.045	0.051
pythia-70m	INTERLEAVE	yes	4.00	0.684	0.722	0.038	0.046
pythia-70m	RANDOM	yes	4.00	0.613	0.652	0.025	0.039
pythia-160m	FRONT	yes	1.00	0.824	0.861	0.015	0.028
pythia-160m	BACK	yes	7.00	0.635	0.681	0.048	0.059
pythia-160m	INTERLEAVE	yes	3.00	0.748	0.772	0.031	0.042
pythia-160m	RANDOM	yes	6.00	0.673	0.686	0.046	0.051
pythia-410m	FRONT	yes	1.00	0.906	0.974	0.020	0.048
pythia-410m	BACK	yes	12.00	0.655	0.759	0.043	0.075
pythia-410m	INTERLEAVE	yes	3.00	0.916	0.914	0.061	0.051
pythia-410m	RANDOM	yes	7.00	0.866	0.823	0.069	0.061
pythia-1b	FRONT	yes	1.00	0.952	0.985	0.018	0.031
pythia-1b	BACK	yes	11.00	0.784	0.812	0.055	0.062
pythia-1b	INTERLEAVE	yes	4.00	0.918	0.931	0.034	0.048
pythia-1b	RANDOM	yes	10.00	0.866	0.912	0.053	0.059

Table 32: Per-regime GPT-2 and Qwen3 model sweep.

Model	Regime r	Token pair	$\lvert H\rvert$	BUC shrink	BUC match	Random shrink	Random HL match
gpt2	FRONT	yes	1.00	0.820	0.860	0.015	0.025
gpt2	BACK	yes	6.00	0.680	0.710	0.045	0.055
gpt2	INTERLEAVE	yes	4.00	0.780	0.810	0.035	0.045
gpt2	RANDOM	yes	5.00	0.720	0.740	0.045	0.055
Qwen/Qwen3-4B-Base	FRONT	yes	2.00	0.945	0.970	0.012	0.015
Qwen/Qwen3-4B-Base	BACK	yes	12.00	0.850	0.885	0.045	0.055
Qwen/Qwen3-4B-Base	INTERLEAVE	yes	6.00	0.920	0.945	0.035	0.042
Qwen/Qwen3-4B-Base	RANDOM	yes	10.00	0.865	0.900	0.048	0.056

B.20 Compute and reproduction information

All reported experiments are inference-only: no model fine-tuning is performed. The main compute cost comes from batched forward passes under activation interventions, model sweeps, and the input-level Greedy Occlusion search. The FP32 code and configurations are described in Appendix F.5; the BF16 experiments require their corresponding prompt and intervention records.

Table 33: Compute and reproduction summary for the reported experiments.

Item	Value
Hardware	NVIDIA A100 GPU
Software environment	Python / PyTorch / Transformers / TransformerLens; the BF16 environment records are separate from the supplied FP32 package (Appendix F.5)
Primary model	EleutherAI/pythia-410m-deduped
Evaluation precision	bfloat16 inference
Training or fine-tuning	none
Main seeds	$\{0, 1, 2, 3, 4\}$
Selection / evaluation prompts	$48/96$ prompts per evidence level
Layer-matched random controls	120 draws per seed
Greedy Occlusion budget	$K_{\mathrm{bb}} = 50$ positions from $L_E = 200$ evidence positions
Greedy Occlusion candidate evaluations	8775 subset evaluations per regime before batching
Model sweep checkpoints	Pythia-70M, Pythia-160M, Pythia-410M, Pythia-1B, GPT-2, Qwen3-4B-Base
Implementation coverage	The supplied package contains FP32 code and aggregate results; separate BF16 scripts and records are not included (Appendix F.5)

C Targeted readout, count, order, and baseline experiments

This appendix specifies the original targeted readout, count, order, and baseline experiments. The additional randomly registered readout pairs and native FP32 controls have their own specification in Appendix E. All reported values come from frozen Pythia-410M and use five seeds unless stated otherwise.

C.1 Seven readout pairs

The audited single-token pairs are

`various/successful`, `yes/no`, `true/false`, `1/0`, `A/B`, `cat/dog`, `red/blue`.

For every pair, the two tokens define the readout direction in Equation 3.1; neither token is a symbolic CEA label. The `A/B` pair is an overlap stress test because the symbols also occur in the prompt. Averaged over four regimes and five seeds, mean $|\Delta_r|$ ranges from 0.721 to 2.083, mean selected-set size from 4.50 to 7.50, BUC shrink from 0.642 to 0.861, and BUC match from 0.701 to 0.884. Layer-matched random shrink ranges from 0.048 to 0.076, and random high-to-low match ranges from 0.059 to 0.083.

Table 34: Representative readout-pair results. Reselected BUC tests procedure robustness. Fixed reference-set transfer tests reuse of the set selected with `various/successful`.

Target pair	Reselected shrink / match	Fixed-set shrink / match	Mean Jaccard
`yes/no`	0.812 / 0.846	0.772 / 0.810	0.610
`A/B`	0.861 / 0.884	0.704 / 0.741	0.390
`cat/dog`	0.642 / 0.701	0.382 / 0.431	0.240

The targeted pairs show positive local intervention effects and partial fixed-set reuse. The selected source and target sets are not constrained to a common budget. In the separate FP32 registry (Appendix E.8), 6.0% of fixed-set conditions pass the joint thresholds. The different target samples and budgets preclude a paired estimate of the benefit of reselection. Failure of a fixed source set leaves open whether a different compact set explains the target.

C.2 Fixed-set count sweep and unseen gaps

For each regime, heads are selected only at $(b_L, b_H) = (50, 150)$. The same set is then evaluated over $b \in \{0, 10, \ldots, 200\}$ and on the unseen gaps $(0, 150)$, $(25, 175)$, and $(75, 125)$. Let $\widehat{\beta}_{\mathrm{base}}$ and $\widehat{\beta}_{\mathrm{abl}}$ be least-squares slopes of y on b before and after ablation. The slope reduction is

$$1 - \frac{|\widehat{\beta}_{\mathrm{abl}}|}{|\widehat{\beta}_{\mathrm{base}}|}.$$

Table 35: Fixed-set count-sweep results. The final column averages the three unseen gaps without reselection.

Regime	Spearman ρ	Linear R^2	Slope reduction	Gap shrink / match
FRONT	0.986	0.948	0.880	0.832 / 0.874
BACK	0.972	0.931	0.646	0.612 / 0.689
INTERLEAVE	0.979	0.944	0.896	0.839 / 0.855
RANDOM	0.958	0.906	0.845	0.786 / 0.811

C.3 Causal localization of order-only contrasts

The original order diagnostic in Appendix B.15 measures the output difference between two regimes at the same count. We additionally apply BUC to four pre-specified order-only contrasts whose two sides have identical B/C counts. Across

the contrasts, $|\Delta_{\mathrm{ord}}|$ ranges from 0.321 to 1.846, BUC shrink from 0.603 to 0.842, and BUC match from 0.651 to 0.884. Layer-matched random shrink is 0.049–0.081, and random high-to-low match is 0.056–0.088. Because the counts are equal by construction, the localized effects cannot be attributed to count changes.

C.4 Balanced-objective baseline decomposition

Table 36 decomposes the BF16 comparison in Table 2. Shrink measures removal of the original contrast, and $\mathrm{match}_{\mathrm{avg}}$ measures bidirectional patch recovery. The equal-weight score $Q_{1/2}$, defined in Equation 3.10, averages those two quantities. Each table entry below is BUC minus sensitivity-only, computed from the displayed three-decimal regime means. These descriptive differences do not provide a significance test. Head counts are matched within regimes as in Table 1; they are not a common $k = 8$.

Table 36: BF16 removal–recovery trade-off, derived from Table 2. Every difference is BUC minus sensitivity-only. Values are dimensionless metric differences, not accuracy percentage points.

Regime	Shrink difference	Mean-match difference	$Q_{1/2}$ difference
FRONT	-0.014	$+0.014$	0.0000
BACK	-0.049	$+0.111$	$+0.0310$
INTERLEAVE	-0.018	$+0.083$	$+0.0325$
RANDOM	-0.029	$+0.107$	$+0.0390$

Removal–recovery trade-off. For BACK, INTERLEAVE, and RANDOM, BUC exchanges 0.018–0.049 of shrink for 0.083–0.111 of mean match. In FRONT, the losses and gains cancel at displayed precision. Relative to ActPatch, the BUC-minus-ActPatch differences in the summed score $2Q_{1/2}$ are 0.015, 0.129, 0.135, and 0.169 in regime order. These comparisons support a recovery-oriented trade-off, not an improvement in both component metrics. The FP32 common-budget differences in Section 5 use a separate configuration.

Joint-threshold performance. Applying shrink ≥ 0.85 and mean match ≥ 0.70 to the displayed BF16 regime means gives the same three passing rows for BUC and sensitivity-only: FRONT, INTERLEAVE, and RANDOM. Both fail in BACK. BUC raises BACK's mean match from 0.648 to 0.759, but its shrink is only 0.655. Thus Table 2 does not show more regimes passing the joint criterion. Passage of three aggregate rows is not a per-seed success rate.

Selection objective versus search procedure. Sensitivity-only ranks direct attributions; BF16 ActPatch ranks single-head HL recovery within a 192-head sensitivity shortlist (Appendix B.10). BUC instead evaluates joint removal and bidirectional recovery for candidate sets. These comparisons change the objective as well as the search procedure. They do not isolate the benefit of greedy augmentation over selecting the top k single heads by $Q_{1/2}(\{i\}, r)$ in the same candidate pool and evaluating their union jointly. The FP32 single-bidirectional-patch ranking also omits the shrink term. The benefit of greedy search under an identical objective and candidate pool therefore remains untested. The timings in Appendix E.6 exclude shared selection costs, so they do not establish an end-to-end cost advantage.

D Full natural-language bridge specification

Question and configuration. The BF16 bridge separates model accuracy, localization on labeled NL prompts, and the effect of a fixed symbolic CEA set on the NL target. The frozen `EleutherAI/pythia-410m-deduped` has 24 layers and 16 heads per layer and uses the audited single-token `yes/no` readout. Interventions modify head writes only at the final prompt position. Every selected or transferred set has budget $k = 8$. Random controls preserve head count and the selected layer pattern, with 120 draws per seed. Tables report means and standard deviations over five seeds; these standard deviations are not bootstrap confidence intervals.

Lexically balanced reports. For domain c, let u_c, v_c be mutually exclusive and exhaustive states, as explicitly stipulated in the prompt. Every prompt has $L_{\mathrm{NL}} = 16$ reports. Bank A uses four forms, indexed by report number j:

$$\begin{aligned} E^{+}_{\mathrm{dir}}(j,c) &= \text{“Report } j\text{: the item is } u_c\text{.”},\\ E^{+}_{\mathrm{neg}}(j,c) &= \text{“Report } j\text{: the item is not } v_c\text{.”},\\ E^{-}_{\mathrm{dir}}(j,c) &= \text{“Report } j\text{: the item is } v_c\text{.”},\\ E^{-}_{\mathrm{neg}}(j,c) &= \text{“Report } j\text{: the item is not } u_c\text{.”}. \end{aligned}$$

Here $+$ means support for the first state u_c, and $-$ means support for the second state v_c. At even count b, the numbers of these four forms are, respectively,

$$b/2, \qquad b/2, \qquad (L_{\mathrm{NL}} - b)/2, \qquad (L_{\mathrm{NL}} - b)/2.$$

Consequently, the evidence contains eight occurrences of each state word, eight direct reports, and eight negative reports at every count. Bank A also contains eight occurrences of `not`. Bank B replaces the forms with separately written alternatives, such as “Report j records the item as u_c” and “Report j rules out the item being v_c,” with corresponding forms supporting v_c. Each bank preserves its own lexical multiset and audited token length across counts; the claim is not that different banks have identical tokenization.

Selection uses Bank A and state pairs `open/closed`, `active/inactive`, and `present/absent`. Evaluation uses Bank B and `valid/invalid`, `available/unavailable`, and `correct/incorrect`. Neither evaluation state pairs nor report templates appear in NL selection. Symbolic CEA selection sees no NL prompt. Lexical balance excludes a strategy based only on the total frequencies of the state words or `not`; it does not exclude local phrase processing or reliance on a subset of reports.

Order, counts, and query reversal. The four regimes order reports by semantic support: FRONT puts all reports supporting u_c first, BACK reverses the blocks, INTERLEAVE alternates support as far as possible, and RANDOM uses matched random permutations. Selection uses $(b_L, b_H) = (4, 12)$. The fixed-set sweep uses $b \in \{0, 2, 4, 6, 10, 12, 14, 16\}$, excluding the tie at eight; new count pairs are $(0, 16), (2, 14), (6, 10)$.

Each report sequence x receives two questions. Query polarity $q = +1$ asks “Do more reports indicate u_c rather than v_c?”; $q = -1$ reverses the states in that question. The evidence is identical for the two questions. With $m(x) = 2b - L_{\mathrm{NL}}$ denoting the semantic margin, the correct label and aligned readout are

$$a^{\star}(x,q) = \mathbf{1}[qm(x) > 0], \qquad y_{\mathrm{NL}}(x,q) = q[\ell_T(\texttt{yes} \mid x,q) - \ell_T(\texttt{no} \mid x,q)],$$

where label one means `yes`. A larger aligned readout indicates stronger model preference for u_c, regardless of query orientation. Let $D^{\mathrm{NL}}_{r,b}$ denote the report distribution at regime r and count b. Averaging equally over both query polarities, define

$$\Delta^{\mathrm{NL}}_r = \mathbb{E}_{x \sim D^{\mathrm{NL}}_{r,b_H}, q}[y_{\mathrm{NL}}(x,q)] - \mathbb{E}_{x \sim D^{\mathrm{NL}}_{r,b_L}, q}[y_{\mathrm{NL}}(x,q)].$$

Behavioral metrics. The restricted prediction is $\widehat{a}(x,q) = \mathbf{1}[\ell_T(\texttt{yes} \mid x,q) > \ell_T(\texttt{no} \mid x,q)]$, so a tied logit pair predicts `no`. Restricted accuracy is $\mathrm{Acc} = \Pr(\widehat{a} = a^{\star})$. Define success on both queries separately for positive and negative semantic margins:

$$\begin{aligned} P_+ &= \Pr(\widehat{a}(x,+1) = 1,\ \widehat{a}(x,-1) = 0 \mid m(x) > 0),\\ P_- &= \Pr(\widehat{a}(x,+1) = 0,\ \widehat{a}(x,-1) = 1 \mid m(x) < 0),\\ \mathrm{PairAcc} &= (P_+ + P_-)/2, \qquad \mathrm{MinPairAcc} = \min(P_+, P_-). \end{aligned}$$

A rule that ignores the evidence and answers `yes` for $q = +1$, `no` for $q = -1$, obtains PairAcc 0.5 but MinPairAcc zero. Independent fair guesses have population paired accuracy 0.25 on each margin sign. The count sweep reports Spearman correlation between semantic margin $m(x)$ and aligned readout $y_{\mathrm{NL}}(x, q)$, and the linear R^2 from regressing the readout on the margin. These statistics do not prove that every report contributes to the decision.

Sampling and head selection. Per seed, selection contains $3 \times 4 \times 2 \times 2 \times 24 = 1152$ prompts: three state pairs, four regimes, two counts, two query polarities, and 24 instances. Evaluation contains $3 \times 4 \times 8 \times 2 \times 32 = 6144$ prompts across three new state pairs and the eight counts. Splits are disjoint. Even in deterministic semantic orders, instances vary the assignment of direct and negative forms; RANDOM also varies the semantic order.

Write $H^{\mathrm{yn}}_{\mathrm{CEA}}(r)$ for the first eight heads of the symbolic `yes/no` BUC ordering and $H_{\mathrm{NL}}(r)$ for the first eight selected on NL prompts. Let $d_{\mathrm{yn}} = W_U^{(\mathtt{yes})} - W_U^{(\mathtt{no})}$. NL attribution is $s_i^{\mathrm{NL}}(x, q) = q\langle v_i(x, q), d_{\mathrm{yn}}\rangle$, where the head write depends on the full input, including the query. The pipeline ranks absolute attribution contrasts, keeps 192 heads, screens 48 with single-head joint scores, and greedily selects eight. Each set is evaluated on both its own target and the other target, without changing head identities after transfer.

Intervention scores and transfer ratio. For target behavior B, with nonzero clean contrast Δ_B, define shrink and bidirectional match as in Section 3:

$$\mathrm{shrink}_B(H) = 1 - \frac{|\Delta_B^{\mathrm{abl}}(H)|}{|\Delta_B|}, \qquad \mathrm{match}_B(H) = \frac{1}{2} \sum_{\tau \in \{\mathrm{HL},\mathrm{LH}\}} \left(1 - \frac{|\Delta_B - \Gamma_B^{\tau}(H)|}{|\Delta_B|}\right).$$

The normalized joint score is $Q_B(H) = [\mathrm{shrink}_B(H) + \mathrm{match}_B(H)]/2$, the target-specific version of $Q_{1/2}$; it is distinct from the raw mean patch gain $S(H, r)$ in the theory. If $\widehat{\beta}_{\mathrm{base}} \neq 0$ and $\widehat{\beta}_{\mathrm{abl}}$ are the clean and ablated regression slopes over the count sweep, slope reduction is $1 - |\widehat{\beta}_{\mathrm{abl}}|/|\widehat{\beta}_{\mathrm{base}}|$. For NL, accuracy decrease is $\mathrm{Acc}_{\mathrm{base}} - \mathrm{Acc}_{\mathrm{abl}}$; tables express that decrease in percentage points.

For source A and target B, let H_A, H_B be source-selected and target-selected sets of eight heads. The BF16 transfer ratio uses the reported matched-random reference score $Q_B(H_{\mathrm{rand}})$:

$$\mathrm{TR}_{A \to B} = \frac{Q_B(H_A) - Q_B(H_{\mathrm{rand}})}{Q_B(H_B) - Q_B(H_{\mathrm{rand}})}.$$

A positive denominator is required to interpret the ratio relative to a target-specific effect above random. A ratio between zero and one places transferred utility between the random reference and target-selected utility. It is not a fraction of shared heads or a measure of how much of an algorithm is reused. The BF16 summary uses one random reference per target; the expanded FP32 protocol instead matches source and target layer patterns separately (Appendix E.10).

D.1 Results

Table 37: Behavioral validation on held-out domains and Bank B paraphrases. Accuracy, paired-query accuracy, and $|\Delta_r^{\mathrm{NL}}|$ use the main gap $(4, 12)$. Spearman correlation and linear R^2 use the full count sweep. Values are mean and standard deviation over five seeds.

Regime	Accuracy	PairAcc	MinPairAcc	Spearman $\rho(m, y_{\mathrm{NL}})$	Linear R^2	Mean $\lvert\Delta_r^{\mathrm{NL}}\rvert$
FRONT	0.742 ± 0.019	0.556 ± 0.028	0.533 ± 0.031	0.879 ± 0.014	0.774 ± 0.022	0.93 ± 0.07
BACK	0.759 ± 0.018	0.581 ± 0.026	0.560 ± 0.029	0.891 ± 0.013	0.798 ± 0.021	1.07 ± 0.08
INTERLEAVE	0.812 ± 0.015	0.674 ± 0.024	0.652 ± 0.026	0.931 ± 0.009	0.863 ± 0.017	1.28 ± 0.06
RANDOM	0.773 ± 0.021	0.605 ± 0.031	0.585 ± 0.034	0.901 ± 0.015	0.814 ± 0.025	1.12 ± 0.09
Mean	0.772 ± 0.014	0.604 ± 0.021	0.583 ± 0.023	0.901 ± 0.010	0.812 ± 0.016	1.10 ± 0.06

Behavioral interpretation. At gap $(4, 12)$, mean restricted accuracy is 0.772, PairAcc is 0.604, and MinPairAcc is 0.583. The means indicate success on both signs of the semantic margin rather than only one query orientation. Across the count sweep, Spearman correlation is 0.901 and linear $R^2 = 0.812$; neither statistic establishes exact monotonicity

Table 38: Held-out gap generalization with the same model and no head reselection. Values are averaged over four ordering regimes and five seeds. The semantic gap is $|b_H - b_L|$.

Gap	Semantic gap	Accuracy	PairAcc	Mean $\lvert\Delta_r^{\mathrm{NL}}\rvert$
$(0, 16)$	16	0.861 ± 0.014	0.748 ± 0.022	2.05 ± 0.11
$(2, 14)$	12	0.829 ± 0.016	0.690 ± 0.025	1.62 ± 0.09
$(4, 12)$	8	0.772 ± 0.014	0.604 ± 0.021	1.10 ± 0.06
$(6, 10)$	4	0.646 ± 0.023	0.397 ± 0.030	0.46 ± 0.06

Table 39: Per-regime natural-language causal localization and fixed symbolic-to-natural-language transfer. All sets have budget $k = 8$. The final two columns report the slope reduction and restricted-accuracy decrease caused by ablating the fixed symbolic CEA set on the natural-language target. Accuracy decrease is in percentage points.

Regime	$H_{\mathrm{NL}} \to$ NL shrink/match	$H_{\mathrm{CEA}}^{\mathrm{yn}} \to$ NL shrink/match	Random $\to$ NL shrink/match	CEA-transfer SlopeRed	CEA-transfer AccDrop
FRONT	0.694/0.731	0.391/0.425	0.057/0.070	0.373	6.9
BACK	0.721/0.758	0.427/0.462	0.054/0.069	0.411	8.2
INTERLEAVE	0.789/0.814	0.472/0.509	0.060/0.073	0.455	9.4
RANDOM	0.733/0.761	0.422/0.454	0.065/0.076	0.406	8.3
Mean	0.734/0.766	0.428/0.463	0.059/0.072	0.411	8.2

Table 40: Aggregate fixed-set transfer between symbolic CEA and the natural-language bridge. All values are averaged over four ordering regimes and five seeds. Base and post-ablation accuracy are shown only for the labeled natural-language target. Accuracy decrease is in percentage points.

Source set $\to$ target	k	Target shrink	Target match	SlopeRed
$H_{\mathrm{NL}} \to$ NL	8	0.734 ± 0.024	0.766 ± 0.026	0.721 ± 0.028
$H_{\mathrm{CEA}}^{\mathrm{yn}} \to$ NL	8	0.428 ± 0.031	0.463 ± 0.034	0.411 ± 0.033
Layer-matched random $\to$ NL	8	0.059 ± 0.008	0.072 ± 0.009	0.051 ± 0.010
$H_{\mathrm{CEA}}^{\mathrm{yn}} \to$ CEA	8	0.812 ± 0.021	0.846 ± 0.023	0.834 ± 0.025
$H_{\mathrm{NL}} \to$ CEA	8	0.499 ± 0.029	0.539 ± 0.031	0.486 ± 0.033
Layer-matched random $\to$ CEA	8	0.050 ± 0.007	0.060 ± 0.008	0.043 ± 0.009

Source set $\to$ target	Base Acc.	Ablated Acc.	AccDrop	Transfer ratio
$H_{\mathrm{NL}} \to$ NL	0.772	0.620	15.2	1.000
$H_{\mathrm{CEA}}^{\mathrm{yn}} \to$ NL	0.772	0.690	8.2	0.555
Layer-matched random $\to$ NL	0.772	0.764	0.8	0.000
$H_{\mathrm{CEA}}^{\mathrm{yn}} \to$ CEA	–	–	–	1.000
$H_{\mathrm{NL}} \to$ CEA	–	–	–	0.599
Layer-matched random $\to$ CEA	–	–	–	0.000

or identifies the source of deviations from linearity. As the gap decreases from 16 to four, accuracy falls from 0.861 to 0.646, PairAcc from 0.748 to 0.397, and the absolute contrast from 2.05 to 0.46. INTERLEAVE has the highest accuracy point estimate and FRONT the lowest, but the comparison does not isolate recency as the cause.

The task label is a majority over report meanings under the stated exclusivity assumption. Constant word counts rule out a total-frequency explanation, not phrase-based or partial-report strategies. The observed performance therefore does not establish that the model computes all 16 report meanings and then applies a majority rule.

Intervention interpretation. On the NL target, the NL-selected set gives mean shrink/match $0.734/0.766$, slope reduction 0.721, and a 15.2-point accuracy decrease. The fixed symbolic set gives $0.428/0.463$, 0.411, and 8.2 points, respectively; random controls give $0.059/0.072$, 0.051, and 0.8 points. Shrink concerns one count pair, slope reduction

concerns the full grid, and accuracy concerns discrete predictions. The aggregate measurements do not identify which examples change predictions. The selected sets have mean Jaccard overlap 0.296 ± 0.041; overlap alone cannot distinguish shared evidence processing from a shared answer readout.

Transfer-ratio arithmetic. From the rounded aggregate scores in Table 40,

$$\mathrm{TR}_{\mathrm{CEA}\to\mathrm{NL}} = \frac{0.4455 - 0.0655}{0.750 - 0.0655} \simeq 0.555, \qquad \mathrm{TR}_{\mathrm{NL}\to\mathrm{CEA}} = \frac{0.519 - 0.055}{0.829 - 0.055} \simeq 0.599.$$

Each input to these ratios is the average of the corresponding shrink and match. The two directions have different targets and denominators; their difference cannot identify additional heads specific to semantic processing. Unlike the expanded FP32 model summaries, these calculations are ratios of the displayed aggregate scores, not means of per-condition ratios.

Declared decision criteria and available evidence. The design's full criterion requires held-out Acc > 0.5, PairAcc > 0.5, and MinPairAcc > 0.25, with prompt-pair bootstrap intervals above those references. It also requires positive count correlation, transferred shrink and match above random, a positive transfer ratio, positive slope reduction, and positive accuracy decrease. The displayed BF16 means satisfy the point-value requirements at the main gap. However, the supplied BF16 results contain means and standard deviations, not the required prompt-pair bootstrap intervals. The complete statistical criterion is therefore not established by the reported summaries. The smaller gap $(6, 10)$ also has PairAcc 0.397, below the 0.5 reference.

Scope. The BF16 results describe intervention transfer of fixed head writes at the final prompt position to a labeled task. They do not establish a shared algorithm for computing a majority: both tasks use `yes/no`, and the rest of the target network remains active. The expanded FP32 study reports separate descriptive model means without model-level intervals, so it does not establish the interval-based criterion either. GPT-2's PairAcc point estimate is 0.458, below 0.5. Positive transfer scores and accuracy decreases do not replace the missing behavioral evidence. BF16 and FP32 measurements are not pooled.

E Native FP32 supplemental experimental protocol

Study design. The FP32 audits evaluate intervention operators, selection conditions, fixed-budget baselines, fixed-set transfer, readout dependence, selection-matched nulls, and an expanded NL bridge (Table 41). Their sampling, donors, and controls differ from the BF16 experiments in Appendices B–D; the two studies are not pooled. The expanded NL bridge uses three models, distinct from the six-model symbolic sweep.

Table 41: FP32 experimental comparisons and their relation to the BF16 study.

Module	Comparison	Relation to the BF16 study
Operators	Fixed selected set under eight ablation/patch combinations	Operator-dependence audit
T3	Exhaustive local subsets, sampled larger subsets, greedy versus exact selection	Tests objective conditions, separately from the many-set T1 correlation
Baselines	Five methods and three random-control families at five common budgets	Extends the original budget-matched comparison
Generalization	Count, unseen gaps, cue wording, equal-budget regime transfer, order, and seeds	Tests count, order, cue, and seed dependence
Readouts	Random token-pair registry, fixed-prompt scan, and two reselection protocols	Extends the seven targeted readout pairs
Selection null	Full adaptive search under paired label swaps	Controls for adaptive selection
Expanded NL	Three models, shared domain split, four template banks, shortcuts, and fixed transfer	Extends the design in Appendix D

Reported outcomes. Section 5 reports FP32 point estimates, observed condition spans, and the eligible 95% bootstrap intervals. Appendix H distinguishes their statistical units. TR summaries are conditional on estimability; neither an all-condition success rate nor simultaneous interval coverage is inferred.

E.1 Shared configuration, pairing, and randomization

The primary checkpoint is `EleutherAI/pythia-410m-deduped`. The expanded NL module additionally uses `EleutherAI/pythia-1b-deduped` and `gpt2`. The models are frozen and loaded sequentially, without training, fine-tuning, quantization, or TransformerLens weight conversion. The supplemental model computations use FP32 with TensorFloat-32 (TF32) disabled. Runtime files record the resolved model revision, tokenizer, software versions, GPU, and actual execution policy. Appendix F describes the numerical checks; FP32 and BF16 measurements are not pooled into one seed average.

Exact symbolic construction. The canonical prompt is the concatenation

```
Tokens:
space-separated symbolic sequence
Choose one word:  various or successful
Answer:
```

The readout strings include their leading spaces: ` various` and ` successful`; both must encode to one token. The symbolic sequence has b copies of B, $200 - b$ copies of C, and a final 50 copies of A. FRONT and BACK place the two evidence symbols in opposite block orders. INTERLEAVE starts with B, alternates while both symbols remain, and appends the excess symbol. For a RANDOM replicate, a permutation of the 200 evidence positions is drawn once. Its first b positions receive B and the remainder receive C. Reusing that permutation across low and high counts produces a matched, nested count contrast. Prefix, suffix, readout order, and tail are fixed within a comparison.

Seeds and split identities. A task seed is obtained by joining the master seed and task identifiers with `||`, taking SHA-256 of the UTF-8 string, converting its first 16 hexadecimal digits to an integer, and reducing modulo $2^{31} - 1$. NumPy's `default_rng` uses the resulting seed. Python, NumPy, and PyTorch seeds are also recorded. Selection, evaluation, and calibration use distinct task keys. The source files specify the exact key strings, avoiding dependence

Table 42: FP32 experimental settings. Execution microbatches are specified separately in Appendix F.

Setting	Value
Master seed	1729
Symbolic evidence and tail lengths	$L_E = 200$, $L_A = 50$
Primary low/high counts	$(b_L, b_H) = (50, 150)$
Selection / evaluation / calibration rows	48 / 96 / 48 per count and regime
Attribution shortlist / single-head pool	192 / 48 heads
Greedy maximum budget / weight	$K_{\max} = 12$, $\lambda = 1/2$
First-crossing thresholds	shrink ≥ 0.85 and mean match ≥ 0.70
Continue after first crossing	Yes; retain the full greedy ordering
Common comparison budgets	$k \in \{1, 2, 4, 8, 12\}$
Random-control draws	30 per family and comparison
Bootstrap replicates	1,000
Weak raw-contrast flag	$\lvert\widehat{\Delta}\rvert < 0.01$
Random readout scan / full-reselection pairs	256 / first 8 registered pairs
Paired-label null trials	99 per estimable target
Stochastic operator repetitions	12 per combination
T3 local pool size / pools per regime	8 / 2 (ranked and uniformly sampled)
T3 larger-pool comparisons	128 per regime
T3 tolerances	10^{-5}, 10^{-3}, 10^{-2}
Symbolic stability seeds	1729, 1730, 1731

on Python's process-randomized built-in hash. A replicate identifier pairs the low and high rows. In NL, both query directions remain associated with the same evidence sequence; query polarity is also a donor stratum.

Independent sampling units. At a fixed count, FRONT, BACK, and INTERLEAVE each define one deterministic symbolic sequence. Their 48 or 96 rows are repeated evaluations, not 48 or 96 independent prompts. RANDOM varies the permutation, with exact duplicate inputs assigned to the same sampling unit. Input hashes and pairing metadata identify these units. Statistical sample sizes and the independence across pairs in Theorem 4.3 refer to independent pairs rather than repeated deterministic executions.

Donor separation. Candidate selection uses donor means from the opposite side of its selection pair. For final primary symbolic evaluation, donor means instead come from the separate 48-row-per-count calibration pair. Independent generation only creates content separation where the prompt generator has random variation. Target-domain calibration is allowed for a fixed transferred set: fixed transfer prohibits *head reselection*, not the use of labeled target-domain donors. The matched-example operator and the NL validation diagnostic use evaluation-partner donors, as specified below, rather than independent calibration donors.

E.2 Intervention unit, empirical scores, and BUC selection

For a head $i = (\ell, h)$ at layer ℓ, let $z_i(x) \in \mathbb{R}^{d_{\text{head}}}$ be its head vector before the output projection at the last prompt position, and let $W_{O,i}$ be the corresponding block of the attention output projection. The write from Section 3 is $v_i(x) = z_i(x)W_{O,i}$. The native implementation intervenes on the selected channels of the concatenated head vector immediately *before* this projection: on `attention.dense` for GPT-NeoX/Pythia and `attn.c_proj` for GPT-2. Zeroing $z_i(x)$ removes its linear write; mean replacement commutes with this linear projection. The shared projection bias is not attributed to a head and is not removed. Unselected heads, model weights, and all downstream model operations remain in the original model. Every selected head is changed in the same forward pass, including sets spanning multiple layers.

For an empirical high/low pair, write $\bar{y}_H$ and $\bar{y}_L$ for the mean native clean readouts. With a selected head set H, write $\bar{y}_H^{\text{abl}}$ and $\bar{y}_L^{\text{abl}}$ for ablated means; $\bar{y}_L^{\text{HL}}$ and $\bar{y}_H^{\text{LH}}$ denote means after the two donor-direction patches. The empirical

counterparts of Section 3 are

$$\widehat{\Delta} = \bar{y}_H - \bar{y}_L, \qquad \widehat{\Delta}^{\mathrm{abl}}(H) = \bar{y}_H^{\mathrm{abl}} - \bar{y}_L^{\mathrm{abl}}, \tag{E.1}$$

$$\widehat{\Gamma}^{\mathrm{HL}}(H) = \bar{y}_L^{\mathrm{HL}} - \bar{y}_L, \qquad \widehat{\Gamma}^{\mathrm{LH}}(H) = \bar{y}_H - \bar{y}_H^{\mathrm{LH}}. \tag{E.2}$$

The LH patch gain is oriented as clean high minus patched high. For a denominator floor $a \geq 0$, define $D_a = \max\{|\widehat{\Delta}|, a, 10^{-12}\}$ and

$$\widehat{\mathrm{shrink}}_a(H) = \frac{|\widehat{\Delta}| - |\widehat{\Delta}^{\mathrm{abl}}(H)|}{D_a}, \tag{E.3}$$

$$\widehat{\mathrm{match}}_a^{\tau}(H) = \frac{|\widehat{\Delta}| - |\widehat{\Delta} - \widehat{\Gamma}^{\tau}(H)|}{D_a}, \qquad \tau \in \{\mathrm{HL}, \mathrm{LH}\}, \tag{E.4}$$

$$\widehat{Q}_{\lambda,a}(H) = \lambda \widehat{\mathrm{shrink}}_a(H) + (1-\lambda)\frac{\widehat{\mathrm{match}}_a^{\mathrm{HL}}(H) + \widehat{\mathrm{match}}_a^{\mathrm{LH}}(H)}{2}. \tag{E.5}$$

For nonzero contrasts and inactive flooring, these are exactly the normalized scores in Equations 3.7–3.10. Scores are not clipped; negative values remain visible. Stabilizing the denominator does not replace the numerator $|\widehat{\Delta}|$ by D_a, so the empty intervention has score zero. The 10^{-12} safeguard prevents numerical division by zero; normalized scores at zero contrast have no effect-size interpretation. Raw contrasts are retained, and weak contrasts are flagged at 0.01. The nonzero floor is specific to the null protocol that repeats selection in Appendix E.9.

Selection, stopping, and fixed budgets. All 384 Pythia-410M heads are ranked by the absolute selection-split contrast of $s_i(x) = \langle v_i(x), d \rangle$. The largest 192 enter single-head intervention screening. The largest 48 single-head $\widehat{Q}_{1/2,0}$ values form the greedy pool. Coordinate order resolves exact ties. Forward greedy search adds the head that maximizes the score of the entire augmented set on selection data. Evaluation scores are not used for head selection. The stopping set is the first prefix with shrink at least 0.85 and mean bidirectional match at least 0.70; if no prefix meets both conditions, the full 12-head budget is returned with a failed-threshold flag. The supplemental configuration continues to record the ordering to 12 heads even when an earlier prefix crosses both thresholds. Fixed-k comparisons use the first k heads of the recorded greedy ordering.

Normalized objective versus raw theory. The executable selection objective is Q_λ in Equation 3.10. The raw signed utility J_λ in Equation 4.1 is a different set function. Theorem 4.5, as stated for J_λ, is not automatically an approximation guarantee for Q_λ. The same greedy argument would require monotonicity and an appropriate submodularity-ratio condition for Q_λ itself. The finite audit below tests the empirical objective, including negative marginal gains and violations of its required conditions. Neither first threshold crossing nor an empirical greedy curve proves a unique or globally smallest causal set. Numerical candidate refinement is specified in Appendix F.3.

E.3 Statistical units and uncertainty

The basic sampling unit for intervention uncertainty is a unique matched high/low input pair, not an execution row. Symbolic pairs are identified by unpadded token hashes, while NL pairs use paired evidence identifiers; reversed query orientations remain in the same cluster. We use 1,000 cluster-bootstrap replicates and report two-sided 95% percentile intervals, with endpoints given by the 2.5th and 97.5th percentiles. Within each replicate, the same resampled cluster indices are used for the clean, ablated, HL-patched, and LH-patched arrays. The intervals are conditional on the selected head set and fixed donor bank; head selection is not rerun inside the bootstrap.

For normalized intervention metrics, only bootstrap replicates satisfying $|\widehat{\Delta}| \geq 0.01$ are retained. These intervals are therefore conditional on a non-weak resampled contrast. Raw contrasts and raw utilities retain their resampled values without this conditioning. Baseline differences use matched bootstrap indices for the two methods.

At a fixed symbolic count, FRONT, BACK, and INTERLEAVE each contain only one unique symbolic sequence. Repeated execution of that same sequence is not treated as independent prompt sampling. Consequently, a prompt-sampling bootstrap interval is not reported for a condition containing only one unique matched pair. RANDOM conditions use distinct matched permutations as sampling clusters.

Execution-seed variation and variation across stochastic operator draws are summarized separately from prompt-sampling uncertainty. For stochastic operators, the 12 intervention repetitions are first averaged at the readout level for the

primary score; the across-repetition standard deviation and standard error of the separately computed scores are Monte Carlo summaries and are not bootstrap confidence intervals. Repeated deterministic executions likewise do not increase the number of independent prompt units.

Cross-regime or cross-condition point summaries are arithmetic summaries of the corresponding condition-level point estimates. We do not assign them a cross-regime confidence interval by treating regimes as independent samples. Any aggregate interval reported for an experiment with genuine sampling units is obtained by recomputing the same aggregate within every bootstrap replicate before taking the 2.5th and 97.5th percentiles. Comparisons in the paper do not claim simultaneous confidence coverage.

The condition-level NL procedures and the point-only model aggregation are specified separately in Appendix E.10; Table 6 has no model-level CI. Counts and observed condition spans are descriptive. Table 7 reports finite operator/T3 summaries without population confidence intervals. Appendix H specifies which statistics receive intervals.

E.4 Fixed-set intervention-operator audit

The primary stopping set for each regime is fixed across the operator audit; the same heads are used for every operator. The eight combinations are listed in Table 43. Pooled ablation donors concatenate the high and low calibration banks. Directional patch donors use the corresponding high or low calibration bank. In NL-capable replacement functions, donor matching is stratified by query orientation, rather than mixing opposite orientations.

Table 43: Operator combinations. Stochastic conditions use 12 draws; deterministic mean and exact matched-example conditions use one evaluation.

Ablation	Directional patch	Donor dependence
Zero	Donor mean	Marginal means of calibration donors
Pooled mean	Donor mean	Pooled calibration ablation mean
Zero	Matched example	Corresponding opposite-side evaluation partner
Zero	Joint resampling	One donor row for all selected heads
Zero	Independent-head resampling	Separate donor row for each head
Zero	Empirical joint Gaussian	Joint mean and full empirical covariance
Pooled mean	Joint resampling	Shared patch donor row
Pooled resampling	Joint resampling	Joint donor rows for both replacements

Example matching and resampling. Matched-example patching copies the opposite high/low evaluation partner with the same pairing key and donor stratum. Missing matches terminate the comparison. The donors are evaluation partners rather than independent calibration samples. Joint resampling draws one donor example for each target example and copies all selected layer/head vectors from that donor. Independent-head resampling draws separately for each selected layer/head, thereby removing their empirical cross-head dependence. The two resampling procedures use the same source distribution but different joint replacements.

Full empirical covariance without a dense covariance matrix. Let $X \in \mathbb{R}^{n_D \times D_H}$ concatenate the preprojection vectors of all heads in H across $n_D \geq 2$ donors in one stratum, where $D_H = |H| d_{\text{head}}$. Let $\mu \in \mathbb{R}^{D_H}$ be the row mean of X. For n_T target rows, draw $E \in \mathbb{R}^{n_T \times n_D}$ with independent standard-normal entries and set

$$\widetilde{X} = \mathbf{1}_{n_T} \mu^\top + \frac{E(X - \mathbf{1}_{n_D} \mu^\top)}{\sqrt{n_D - 1}}. \tag{E.6}$$

Writing $X_c = X - \mathbf{1}_{n_D} \mu^\top$, a centered replacement row is $X_c^\top e / \sqrt{n_D - 1}$, with $\mathbb{E}e = 0$ and $\mathrm{Cov}(e) = I$. Its covariance is therefore $X_c^\top X_c / (n_D - 1)$, the full empirical covariance, and its mean is μ. There is no principal-component truncation or ridge perturbation in this operator. A zero-variance donor bank gives its mean; fewer than two donors cannot estimate a nontrivial covariance. Matching mean and covariance is not a guarantee that a generated activation is a natural model state.

Monte Carlo aggregation. For a stochastic operator, the 12 intervention readouts are averaged for each example before the primary shrink, directional match, and $Q_{1/2}$ are computed. For the 12 separately computed $Q_{1/2}$ scores, the output records their mean, sample standard deviation, and Monte Carlo standard error. The standard error equals the sample standard deviation divided by $\sqrt{12}$. The score of averaged intervention readouts need not equal the average of scores. This Monte Carlo uncertainty is distinguished from the matched-input bootstrap.

E.5 Finite T3 audit and exact local comparisons

Let $\mathcal{P}_r$ be the 48-head single-intervention candidate pool for regime r, not the 192-head attribution shortlist. Two 8-head pools are registered for each regime: the first eight ranked candidates and eight candidates sampled uniformly without replacement from $\mathcal{P}_r$. The registry is fixed before examining evaluation scores. For each pool, the protocol evaluates all $2^8 = 256$ subsets, including the empty set, on both selection and evaluation data. The design specifies $4 \times 2 \times 2 \times 256 = 4096$ local set scores; cached duplicate sets may reduce actual forward calls.

For a pool P, subset $A \subseteq P$, and head $i \in P \setminus A$, define the empirical marginal

$$m_i(A) = \widehat{Q}_{1/2,0}(A \cup \{i\}) - \widehat{Q}_{1/2,0}(A). \tag{E.7}$$

For each tolerance $\epsilon \in \{10^{-5}, 10^{-3}, 10^{-2}\}$, the audit records the fraction of the $8 \times 2^7 = 1024$ pool edges with $m_i(A) < -\epsilon$, the minimum marginal, and marginal quantiles. For every nonempty $S \subseteq P \setminus A$ with $\widehat{Q}_{1/2,0}(A \cup S) - \widehat{Q}_{1/2,0}(A) > \epsilon$, the empirical ratio is

$$\widehat{\gamma}(A, S) = \frac{\sum_{i \in S} m_i(A)}{\widehat{Q}_{1/2,0}(A \cup S) - \widehat{Q}_{1/2,0}(A)}. \tag{E.8}$$

The minimum, lower quantiles, median, and nonpositive fraction are reported without clipping negative values to zero or values above one to one. Ratios are omitted when the joint gain does not exceed the tolerance.

Selection versus an evaluation oracle. For each local budget $k = 1, \ldots, 8$, the audit constructs a greedy prefix using only selection scores and independently enumerates the best selection-score subset with at most k heads. Both fixed choices are then evaluated on evaluation data. The best evaluation-score subset with at most k heads is recorded separately as an oracle diagnostic, not as an implementable held-out selection rule. The greedy/oracle score ratio is reported only for an oracle score above the diagnostic tolerance. Selection regret is the selection optimum minus the greedy selection score.

Larger-pool sampling. For each regime, the registry contains 128 pairs (A, S) of disjoint subsets drawn from $\mathcal{P}_r$. A random permutation supplies the heads; $|A|$ is uniform on $0, \ldots, 12$ and $|S|$ is uniform on $1, \ldots, \min\{12, 48 - |A|\}$. The audit evaluates A, $A \cup S$, and each $A \cup \{i\}$ for $i \in S$ on evaluation data. The resulting 512 comparisons record the actual sizes, joint gain, marginal sum, minimum marginal, and the ratio when joint gain exceeds 10^{-5}. The union may contain up to 24 heads, so the diagnostic extends beyond the 12-head search budget. An exhaustive minimum on an 8-head pool, or a sampled minimum on the 48-head pool, is not a global lower bound for all model heads. Empirical violations limit the applicability of T3 to the tested objective.

Finite-audit results. Table 7 reports negative evaluation marginal fractions of 31%, 17%, 28%, and 22% for FRONT, BACK, INTERLEAVE, and RANDOM, respectively, at tolerance 10^{-3}. The corresponding minimum empirical ratios are $-5.8, -2.6, -4.9, -3.9$. These are witnesses in the finite pools of nonmonotonicity and negative ratios for the observed normalized objective, not estimates establishing a negative ratio throughout a prompt population. Maximum selection regrets are $0.021, 0.043, 0.027, 0.035$. They compare selection-data optima with greedy selection within the two registered eight-head pools; they are neither held-out nor guarantees for the full pool. No confidence intervals are used for these finite-audit summaries. The observations do not contradict the conditional Theorem 4.5 for raw utility J_λ.

E.6 Fixed-budget method and random-control comparisons

The protocol compares five rules at each $k \in \{1, 2, 4, 8, 12\}$ and each regime. BUC takes the first k greedy heads. Sensitivity-only takes the largest absolute attribution contrasts across all heads. Single-HL patch ranking uses the single-head HL match values measured during 192-head screening. Single-bidirectional patch ranking uses the corresponding average match values. Write-norm ranking uses all heads' mean $\|z_i(x)W_{O,i}\|_2$ over the concatenated high/low selection rows. All five sets receive the same set-level evaluation metrics. The design contains $4 \times 5 \times 5 = 100$ method comparisons.

These rules do not have identical candidate universes or equal search costs. The comparison matches final head count, not end-to-end compute. In particular, the single-head rankings are separate baselines, not implementations of EAP or ACDC. Reported evaluation forward calls and elapsed times exclude shared selection and precomputation costs.

For every regime/budget pair, the protocol draws 30 sets from each of three control families. Layer-matched sampling preserves the exact count of heads in each selected layer and samples without replacement within a set. Shortlist-uniform sampling draws k heads from the 48-head single-intervention pool. Layer/norm-nearest sampling processes each reference head, ranks still-available heads in the same layer by absolute difference in log write norm (with a 10^{-12} norm floor), and samples uniformly from the four nearest candidates, or all remaining candidates if fewer than four remain. The last control is approximate norm matching, not equality of norms. Selected heads remain eligible for sampling in all three families. There are $4 \times 5 \times 3 \times 30 = 1800$ random-control conditions.

BUC-minus-alternative differences are computed for the four alternative methods at each regime and budget, giving 80 comparison rows. Their matched-pair bootstrap uses the same clusters for both methods, with the limitations in Appendix E.3. The paired-label null in Appendix E.9 separately controls for adaptive BUC selection.

E.7 Count, prompt, ordering, and seed generalization

Fixed-set count sweep. The primary stopping set selected at $(50, 150)$ is held fixed. For each regime, the protocol evaluates $b \in \{0, 10, \ldots, 200\}$ using 24 rows per count and the same matched-permutation seed across count levels. Per-example clean and ablated readouts are retained and averaged within each count. The analysis computes Spearman association between count and the clean count means, together with ordinary least-squares fits to the clean and ablated count means. Let $\widehat{\beta}_{\mathrm{base}}$ and $\widehat{\beta}_{\mathrm{abl}}$ be their slopes. The count-sweep summary reports the clean fit's R^2 and $1-|\widehat{\beta}_{\mathrm{abl}}|/|\widehat{\beta}_{\mathrm{base}}|$ when the clean slope is numerically nonzero. These are descriptive count-curve statistics, not symbolic accuracy or a new selection criterion.

Unseen contrasts. The same set is evaluated at $(0, 150)$, $(25, 175)$, and $(75, 125)$, using 96 rows per side and a separate 48-row-per-side calibration bank. There is no reselection. The first contrast shares the count 150 with selection, so "unseen gap" denotes an unseen count pair, not necessarily two unseen endpoint values. Each of the 12 regime/gap combinations retains its own raw contrast and intervention scores.

Prompt-cue controls. The readout and head set remain fixed while the prompt is changed by either replacing `Tokens:` by `Symbol sequence:`, or replacing `Choose one word:` by `Continue with one word:` and `Answer:` by `Response:`. The two controls give eight regime/cue comparisons with 96 evaluation and 48 calibration rows per count. The symbolic contents are paired across cue variants; tokenized lengths may change between variants. These controls are different from the original alphabet/word/sentence surface-form sweep.

Equal-budget cross-regime transfer. For every source regime, target regime, and common budget $k \in \{1, 2, 4, 8, 12\}$, the first k source greedy heads are transferred to the target without reselection. The $5 \times 4 \times 4 = 80$ comparisons include the diagonal where source and target coincide. Ablation, both directional patch gains, and both matches are evaluated on the target's primary evaluation/calibration data. Equal head count controls for differences in source-set size, but target-specific donor calibration and different layer patterns remain part of the intervention design.

Order-only localization. Four ordered comparisons are fixed in the registry: FRONT→BACK at $b = 50$ and at $b = 150$, FRONT→INTERLEAVE at $b = 150$, and INTERLEAVE→RANDOM at $b = 150$. Both sides have identical symbol counts. The difference between the target and source mean readouts replaces the high-count minus low-count contrast. Each comparison performs its own BUC selection with 48 rows per side, evaluation with 96, calibration with 48, and 30 layer-matched control sets. These are causal order tests, not just raw output differences.

Permutation-seed stability. For RANDOM, selection is repeated at seeds 1729, 1730, and 1731 with a common primary evaluation and calibration pair. The output records selected sets, threshold status, scores, and the three pairwise Jaccard overlaps. The seeds index prompt generation and search execution, not independently pretrained models. Exact head identity is not assumed across seeds, readouts, models, or budgets.

E.8 Random readout registry and separate prompt/readout changes

Eligible tokens and registration. The eligible vocabulary contains token IDs whose decoded string is exactly one leading space followed by 3–14 ASCII letters, whose string encodes back to the same single ID, and whose ID is present in the model output vocabulary. This is an alphabetic-token filter, not an English dictionary or a uniform distribution over all residual-space directions. The registry samples 256 ordered pairs of distinct IDs uniformly without duplicate ordered pairs, using a task-derived seed before model responses are evaluated. The reverse ordered pair is a different registry entry. The registry, decoded strings, and IDs are saved; weak and negative-effect pairs remain in the registry.

Broad fixed-prompt, fixed-set scan. The canonical prompt still lists `various/successful`, and the primary head set is held fixed. The scan caches vocabulary-logit contrast vectors for the clean, ablated, HL-patched, and LH-patched model and evaluates the difference between the two registered token coordinates for each pair. The scan has $4 \times 256 = 1024$ rows. It checks the identity obtained by swapping the readout order: the raw contrast and oriented gains reverse sign, while the absolute-value normalized scores do not change where defined. The scan evaluates reuse of a fixed set on new output coordinates; it does not claim that the model follows an instruction to generate those sampled words.

Two protocols that repeat selection. Full BUC and 30 layer-matched controls are applied to the first eight pairs in registry order in each regime, independently of their measured effects. In `fixed_prompt_reselected`, only the measured token IDs change; the original prompt and symbolic contents remain fixed. In `prompt_changed_reselected`, the candidate words in the prompt also change, while the underlying symbolic permutations and split keys stay matched. The $8 \times 2 \times 4 = 64$ tasks separately measure sensitivity to the output direction and joint prompt/readout sensitivity. Each records the selected size, native final intervention scores, controls, and Jaccard overlap with the primary set.

The broad scan summary uses raw-effect thresholds 0.001, 0.01, 0.05, and 0.1. At each threshold, the summary records the numbers of registered pairs, evaluable pairs, and pairs passing shrink ≥ 0.85 and average match ≥ 0.70, with pass fractions relative to both all registered pairs and the evaluable subset. Nonfinite and weak rows remain in the underlying scan. These targeted and randomly registered readout studies do not establish head identity that is invariant to the readout. At threshold 0.01, the fixed-set scan has 960 evaluable conditions out of 1024. The rounded pooled passing fractions are 0.060 overall and 0.064 among evaluable conditions, not averages of regime-specific rates. Counts and fractions describe this registered finite scan; no interval over resampled readout registries is reported. They measure fixed-set reuse under the canonical prompt, not the success rate of full BUC reselection.

E.9 Paired-label selection-matched null

The null tests joint high/low label exchangeability within unique matched pairs. It is only applied when both selection and evaluation have at least eight unique pair clusters. Otherwise the target is recorded as not testable and its permutation probability is unavailable. Replicating one deterministic contrast cannot satisfy this requirement. Thus the four symbolic regimes are all listed, but their eligibility is determined from unique pairs, not batch-row count.

Let s_{pool} be the standard deviation of the concatenated clean high/low selection readouts, computed with `ddof=0`. Set

$$a = \max\{10^{-6},\ 0.05 s_{\mathrm{pool}}\}. \tag{E.9}$$

Pooled variation is invariant to label swaps. The observed run and all null runs use the same a, $\widehat{Q}_{1/2,a}$, shortlist sizes, greedy budget, thresholds, and continue-after-crossing rule. The floor depends on pooled standard deviation rather than the observed clean contrast. In each of 99 trials, the two sides of every matched cluster are independently swapped with probability $1/2$. All duplicate rows or query orientations in a cluster receive the same swap. Selection, evaluation, and calibration banks each receive their own cluster swaps. Each trial recomputes attribution, the 192-head shortlist, 48-head screening pool, greedy ordering, stopping set, and donor patches.

The test statistic is the unnormalized *magnitude-improvement* utility on the corresponding evaluation labels:

$$T_\lambda(H) = \lambda\big(|\widehat{\Delta}| - |\widehat{\Delta}^{\mathrm{abl}}(H)|\big) + (1-\lambda)\left(|\widehat{\Delta}| - \frac{1}{2}\sum_{\tau\in\{\mathrm{HL},\mathrm{LH}\}} |\widehat{\Delta} - \widehat{\Gamma}^{\tau}(H)|\right). \tag{E.10}$$

This statistic avoids dividing by a nearly zero permuted evaluation contrast. It differs from the signed raw utility J_λ used in the theory. Let T_{obs} be the observed statistic, T_t the statistic after null trial t, and B the integer number of completed

trials. With $c = \sum_{t=1}^{B} \mathbf{1}[T_t \geq T_{\text{obs}}]$, the Monte Carlo upper-tail probability is

$$\widehat{p} = \frac{1+c}{B+1}, \qquad \text{resolution} = \frac{1}{B+1}. \tag{E.11}$$

The completed-trial count has no confidence interval. A separate two-sided 95% Clopper–Pearson interval describes the underlying null exceedance probability, using beta quantiles at 0.025 and 0.975 with zero/one endpoints when $c = 0$ or $c = B$. The interval describes Monte Carlo uncertainty in the exceedance rate, not effect-size uncertainty, an interval for the adjusted p-value, or uncertainty in the number of completed trials. Holm adjustment applies to the finite estimable permutation probabilities; non-testable targets remain unavailable.

Only RANDOM is estimable in this fixed-count symbolic design, so the Holm family contains one finite probability. The 99 completed trials give $T_{\text{obs}} = 2.69$, with effect-size CI $[2.55,\ 2.80]$, and adjusted $p = 0.01$. The probability is at the $1/100$ resolution limit, not below it. With one estimable hypothesis, $p = 0.01$ and $B = 99$ imply $c = 0$; the corresponding Clopper–Pearson interval for the null exceedance probability is

$$[0,\ 1 - 0.025^{1/99}] \approx [0,\ 0.0366].$$

This binomial interval does not replace the calibrated permutation-test decision or describe the effect size.

The null-selected set is also evaluated on the unchanged clean target, but its clean-score rank is explicitly a diagnostic, not the permutation probability in Equation E.11. Interpretation depends on the paired exchangeability null and all stated split/selection rules; rejection does not prove mechanism uniqueness, universal causal validity, or semantic correctness.

E.10 Expanded natural-language bridge

Models, seeds, and domain registry. Behavior and causal transfer use Pythia-410M, Pythia-1B, and GPT-2. The `paper` profile uses behavioral seeds 1729, 1730, and 1731, while causal selection uses only seed 1729. The behavioral and causal seed counts describe different repetitions. The single-model, five-seed bridge in Appendix D uses a separate configuration.

The source lists 30 ordered state pairs. A shared task seed derived from `1729||shared-state-registry-v4` permutes this list once. Its first four pairs form selection domains, the next four validation domains, and the next eight test domains (Table 44). The same split is used for every model and behavioral seed. It is not selected from model accuracy or intervention effects. The complete 30-pair list and the deterministic registry are included in `anc/protocol/nl_state_registry.json` and the source.

Table 44: Shared state-pair split derived from the supplemental master seed. The prompt declares the two named states mutually exclusive and exhaustive within each task; this is a task assumption, not a claim about unrestricted word meaning.

Split	Ordered state pairs
Selection	local/remote; legal/illegal; stable/unstable; known/unknown
Validation	inside/outside; positive/negative; enabled/disabled; above/below
Test	valid/invalid; safe/unsafe; alive/dead; connected/disconnected; hot/cold; approved/rejected; ready/unready; secure/insecure

Balanced report construction. For one domain c, let (u_c, v_c) be its ordered state pair and let $L_{\text{NL}} = 16$ be the number of reports. For each even count b of reports supporting u_c, the prompt contains $b/2$ direct reports of u_c, $b/2$ reports negating v_c, $(L_{\text{NL}} - b)/2$ direct reports of v_c, and $(L_{\text{NL}} - b)/2$ reports negating u_c. Each state word occurs eight times *in the evidence region*, with eight direct and eight negative report forms. The prefix and question also mention the states and are not counted in that evidence-only statement. The exact sentence forms are in Table 45; negative evidence for one state supports the other under the task's two-state assumption.

Banks A, C, and D contain eight occurrences of the literal word `not` in the evidence region. Bank B instead has eight `rules out` constructions; it does not have eight literal `not` tokens. Lexical balancing is *within* a bank, not equality of tokenization between banks. State words may be multiple tokens; only the `yes/no` answer strings must be single tokens. For every registered domain, bank, count, and query polarity, the tokenizer audit compares actual token multisets using a

Table 45: Exact supplemental report templates. Here j is the report number, u_c is the supported state, and v_c is the other state. Exchange the two states for negative-polarity reports.

Bank	Direct support for u_c	Support for u_c by negating v_c
A	Report j: the item is u_c.	Report j: the item is not v_c.
B	Record j describes the item as u_c.	Record j rules out the item being v_c.
C	Witness j says that the item remains u_c.	Witness j says that the item is not v_c.
D	Statement j: u_c is the reported condition.	Statement j: v_c is not the reported condition.

fixed audit generator. Any failed audit raises an error rather than dropping that domain. The paired causal batches also check equal real token length between the matched count conditions. Length checks use unpadded inputs.

The canonical expanded prompt is

> Exactly one state holds: the item is u_c or the item is v_c.
> Read the reports and answer the final question with yes or no.
> *16 reports in the specified order*
> Question: Do more reports indicate u_c rather than v_c?
> Answer:

For reversed query polarity, only the two state positions in the question are exchanged. The reports are identical for the two queries of a given example.

Report units are shuffled before being ordered by semantic polarity. The four ordering regimes then use positive-first blocks, negative-first blocks, positive-first alternating polarity with any excess appended, or a random report permutation. Assignment and order of direct versus negative forms can vary even within a non-random semantic-order regime. The count grid is $\{0, 2, 4, 6, 10, 12, 14, 16\}$; the tied case $b = 8$ is excluded. Selection and primary accuracy summaries use the $(4, 12)$ contrast; the full count grid is analyzed separately. Low/high NL examples are paired by domain, bank, regime, and replicate. Because the generation seed includes the count, low and high report permutations need not be identical; only reversed queries share exactly the same reports. This differs from the nested symbolic count pairing.

Labels, aligned scores, and paired correctness. Let $q \in \{+1, -1\}$ be the query orientation and $m(x) = 2b - L_{\mathrm{NL}}$ the semantic margin. The correct label and prediction are

$$a^\star(x, q) = \mathbf{1}[qm(x) > 0], \qquad \widehat{a}(x, q) = \mathbf{1}[\ell_T(\texttt{yes} \mid x, q) > \ell_T(\texttt{no} \mid x, q)]. \tag{E.12}$$

A logit tie is predicted as `no`. The semantic readout $y_{\mathrm{NL}}(x, q) = q[\ell_T(\texttt{yes} \mid x, q) - \ell_T(\texttt{no} \mid x, q)]$ aligns both query orientations toward u_c. Attribution uses $q\langle v_i(x, q), d_{\mathrm{yn}}\rangle$, where d_{yn} is the `yes`-minus-`no` unembedding direction. Donor means are stratified by q and pooled over the specified donor domains within each orientation. Donors from opposite query orientations are kept separate.

Restricted accuracy averages $\mathbf{1}[\widehat{a} = a^\star]$ over query rows. Let P_+ be the probability that both queries are correct conditional on $m > 0$, and P_- the corresponding probability conditional on $m < 0$. The paired metrics are $\mathrm{PairAcc} = (P_+ + P_-)/2$ and $\mathrm{MinPairAcc} = \min\{P_+, P_-\}$. Always answering according to query orientation while ignoring the reports can achieve PairAcc $= 0.5$ but MinPairAcc $= 0$. Independent fair random answers instead have expected paired correctness 0.25. These are different reference behaviors. The report includes both side-specific scores rather than interpreting all paired metrics as binary accuracy.

The full-count summaries separately report Spearman association between $m(x)$ and $y_{\mathrm{NL}}(x, q)$ and the ordinary least-squares slope and R^2. In addition, let $p_\theta(w \mid x, q)$ be the model's full-vocabulary probability. The output records the mean and median of

$$M_{\mathrm{yn}}(x, q) = p_\theta(\texttt{yes} \mid x, q) + p_\theta(\texttt{no} \mid x, q). \tag{E.13}$$

This candidate mass is distinct from the sigmoid-normalized two-answer probability. Good forced-choice accuracy alone does not establish that an unrestricted generation would choose either answer token.

Sample accounting and donor use. Bank A and the four selection domains provide three replicates per count and query orientation. Thus a single regime has $4 \times 3 \times 2 = 24$ selection rows per count, or 48 across $(4, 12)$, and the four regimes have 192 selection rows per model and causal seed. Four held-out validation domains use Bank A with the same

replicate count. Validation scores are excluded from head selection. This validation pair uses its own opposite side for mean donors; it is a held-out-domain diagnostic, not an independent-calibration evaluation.

Final behavioral evaluation uses eight test domains, Banks B–D, six replicates, four regimes, eight count levels, and both queries. Per behavioral seed and model, there are $8 \times 3 \times 6 \times 4 \times 8 \times 2 = 9216$ query rows, or 27,648 across the three behavioral seeds. These are design counts, not counts of independent observations. Final NL causal evaluation uses the same test domains and Banks B–D at $(4, 12)$: 96 target rows per count and regime/bank. Its separate calibration bank uses three replicates, giving 48 donor rows per count and regime/bank. The target and calibration generators use distinct keys for the same domain, bank, and semantic count. This is target-calibrated transfer, not a donor-free zero-shot intervention.

Table 46: Expanded NL sampling units. Rows count questions, including both query orientations; intervals retain the paired evidence units.

Use	Domains / banks	Rows and seed convention
NL selection	4 / A	24 per count/regime; causal seed 1729
NL validation	4 / A	24 per count/regime; not used to choose heads
Behavior, original prompts	8 / B–D	9,216 per model/behavioral seed; three seeds
Behavior, four shortcut controls	8 / B–D	9,216 additional rows per model; main gap, first seed only
Causal target	8 / one of B–D	96 per count/regime/bank; causal seed 1729
Causal calibration	8 / matching bank	48 per count/regime/bank; separate generation key
Symbolic yes/no source	Canonical symbolic	48 / 96 / 48 selection/evaluation/calibration per count/regime

Shortcut controls. For the main counts 4 and 12 and behavioral seed 1729, four transformations are applied separately to every test domain/regime/bank condition. The query-only control removes the reports but keeps the two prefix lines and the question. Reverse-report order reverses the report lines without renumbering them. Shuffled-report numbers replace the number in each line by a seeded permutation of $1, \ldots, 16$ without reordering the lines. Reversed-prefix states swaps u_c and v_c only in the opening two-state declaration, leaving reports and question unchanged. The correct labels are retained. These controls test different shortcuts; they are not four independent replications of the main behavior. They add $8 \times 3 \times 4 \times 2 \times 2 \times 6 \times 4 = 9216$ query rows per model.

Condition-level behavioral uncertainty. The behavior summary preserves model, regime, bank, control type, and semantic gap. Headline accuracy and paired scores use $(4, 12)$; full-count association is summarized separately. For 1,000 bootstrap replicates, tested state domains are sampled with replacement, followed by paired-report records within each sampled domain and each semantic-margin sign. Reversed query orientations remain in the same resampled record. The 2.5th and 97.5th percentiles define the two-sided 95% condition-level interval, conditional on the tested domain/template design. Query-only controls can create repeated identical final prompts; distinct originating report IDs do not make those model responses independent.

Model-level point summaries. Table 6 aggregates the condition-level point estimates over ordering regimes and Banks B–D. Model-level CIs are not reported, and no model-level significance claim is made from these summaries. Condition-level interval endpoints are not averaged to construct an interval for a model mean; regimes, banks, and repeated execution rows are not treated as independent bootstrap replicates. The condition-level transfer-ratio procedure and its estimability counts are specified below, separately from behavioral uncertainty.

Fixed sets and bidirectional task transfer. Within each model/regime, $H^{\mathrm{yn}}_{\mathrm{CEA}}(r)$ contains the first eight greedy heads selected on symbolic CEA with the `yes/no` readout, and $H_{\mathrm{NL}}(r)$ contains the first eight from Bank-A NL selection. Both use up to 192 attribution-ranked heads (all heads when fewer are available) and up to 48 screened candidates; the budget for causal transfer is fixed at eight regardless of earlier threshold crossing. The four tested roles are

$$H^{\mathrm{yn}}_{\mathrm{CEA}} \to \mathrm{CEA}, \quad H_{\mathrm{NL}} \to \mathrm{NL}, \quad H^{\mathrm{yn}}_{\mathrm{CEA}} \to \mathrm{NL}, \quad H_{\mathrm{NL}} \to \mathrm{CEA}.$$

The reverse NL-to-symbolic target is evaluated once per model/regime; the forward symbolic-to-NL target is evaluated separately on Banks B, C, and D. No target-bank reselection is used. There are $3 \times 4 \times (1 + 3) = 48$ configured target

comparisons. Each comparison evaluates the source set and target-specific set, plus two separate 30-draw layer-matched random families, for 62 set-level role/control rows. The source and target layer-count patterns need not match even though both sets have eight heads.

For a target task B, let $Q_B(H)$ denote the equal-weight normalized score of set H on that target, and let $\overline{Q}_B(R_A)$ and $\overline{Q}_B(R_B)$ denote mean scores of controls matched to the source and target-specific sets, respectively. The expanded transfer ratio is

$$\mathrm{TR}_{A\to B} = \frac{Q_B(H_A) - \overline{Q}_B(R_A)}{Q_B(H_B) - \overline{Q}_B(R_B)}. \tag{E.14}$$

Using separate control families distinguishes this protocol from the single-control arithmetic of the original bridge. The numerator and denominator are recorded separately. A point ratio is estimable only when both gates hold:

$$|\widehat{\Delta}_B| \geq 0.01, \qquad Q_B(H_B) - \overline{Q}_B(R_B) > 0.05.$$

Here $\widehat{\Delta}_B$ is the empirical clean contrast on target task B, and $Q_B(H_B) - \overline{Q}_B(R_B)$ is the target-selected set's score in excess of its own layer-matched random control. Otherwise the point ratio is unavailable, not clipped or set to zero. Ratios outside $[0, 1]$ are possible; a ratio is not a fraction of a uniquely identified circuit.

Condition-level transfer intervals. For each estimable comparison, the interval procedure uses 1,000 replicates. Matched target clusters are resampled, and the draws within each of the two random-control families are resampled independently. The same target indices are used for the source, target-selected, and arrays from random sets. Replicates must pass the same gates for the raw contrast and excess target score. A condition-level interval is suppressed if more than 5% of replicates fail either gate; otherwise the 2.5th and 97.5th percentiles of the stable replicates define the interval. These intervals condition on fixed selected sets and donor banks, not on repeated causal selection. A finite point ratio remains in the output when its interval is suppressed.

Transfer coverage and aggregation. For each model, forward symbolic-to-NL transfer has $4 \times 3 = 12$ regime–bank target conditions. Let n_{est} be the number of finite ratios satisfying both gates for estimating a point ratio, and let n_{suppr} be the number of these estimable conditions whose interval is suppressed. The model-level TR in Table 6 is the unweighted arithmetic mean of all n_{est} finite condition-level ratios, including those with suppressed intervals. It is not the ratio of separately averaged numerators and denominators. Conditions without an estimable point ratio are excluded from the TR mean, not assigned zero.

In Pythia-410M, Pythia-1B, and GPT-2 order, $n_{\mathrm{est}} = 11, 12, 9$, and $n_{\mathrm{suppr}} = 1, 0, 3$. Thus $12 - n_{\mathrm{est}}$, the number of unavailable point ratios, is 1, 0, and 3, respectively. The number of retained *condition-level* intervals, $n_{\mathrm{est}} - n_{\mathrm{suppr}}$, is 10, 12, and 6, respectively. These counts do not provide a model-level interval or the number of conditions with positive transfer.

The labeled NL analysis records clean and ablated accuracy and their paired difference. CSV differences are fractions and are multiplied by 100 for reporting in percentage points. The expanded causal module evaluates the main gap only, so it does not estimate the ablated full-count slope reduction reported for the BF16 bridge. Fixed-set Jaccard overlap is reported separately from behavioral and intervention transfer.

Interpretation of the expanded aggregate results. Table 6 reports accuracy/PairAcc/MinPairAcc of 0.752/0.566/0.541 for Pythia-410M, 0.805/0.651/0.611 for Pythia-1B, and 0.681/0.458/0.414 for GPT-2. Accuracy point estimates exceed 0.5 in all three models; GPT-2's PairAcc point estimate is below the 0.5 reference. No model-level CI is reported, so these point comparisons do not establish the interval-based behavioral criteria in Appendix D or statistical differences between models.

The transferred sets have shrink/match 0.401/0.442, 0.462/0.501, and 0.306/0.348, respectively. None of these model means meets the original joint symbolic thresholds. Mean TR values 0.520, 0.557, and 0.438 are based on the 11, 12, and 9 estimable target conditions, respectively; unavailable ratios are excluded. The means are consistent with partial intervention transfer relative to the target-selected reference, but do not establish positive transfer for every condition, full circuit recovery, or population significance. A TR is not a percentage of an identified reasoning circuit.

Accuracy decreases of 7.4, 8.6, and 4.9 percentage points are paired differences on the causal cohort. Subtracting those decreases from the accuracy means over three behavioral seeds does not produce matched-cohort ablated accuracies. These descriptive FP32 results remain separate from the original BF16 bridge and do not replace its full behavioral validation criteria.

F Execution and numerical validation

F.1 Native model path and reference computation

The implementation is divided into the experiment definitions in `anc/code/cea_native_suite.py`, the prefix-cache engine in `anc/code/cea_fast_v9.py`, and numerical routing in `anc/code/cea_numeric_v9.py`. Source checksums in the accompanying protocol identify the implementation.

Models use the native Hugging Face GPT-NeoX/Pythia or GPT-2 implementation, in evaluation mode without gradients. The output embedding is obtained through `get_output_embeddings()`, and model-specific projection layouts are handled explicitly. FP32 inference is retained without autocast, TF32, lower-bit weight conversion, or model fine-tuning. Attention uses the native eager path. The reference clean readout for a cached prompt is obtained from a full native forward pass. Storage can use -1 padding to hold unequal-length inputs, but the model never receives those padding values: inputs are grouped by unpadded length, and outputs are restored to the original row order. Native reference calls use unpadded singleton inputs where required for the established reference path.

The numerical protocol distinguishes exact mathematical identities from finite floating-point comparisons. Replacement of head channels before a linear output projection implements the specified write intervention, but reordering batch computations or reusing attention caches can change floating-point results. Native and cached intervention outputs are therefore compared numerically.

F.2 Causal-prefix caching and independent intervention branches

All supplemental white-box interventions act at the final prompt position. Causal attention prevents changes at that position from modifying earlier positions. For a prompt of length T, the engine selects a suffix length s from $\{1, 8, 32, 64\}$, limited to the available context length. A full native prefill stores the key/value tensors for positions before the suffix. For each head-set evaluation, the engine reuses only the unmodified prefix and recomputes the suffix, including the final token, every affected attention operation, all downstream layers, layer normalizations, and MLP computations. The intervention is still applied only at the final position. Earlier suffix tokens are recomputed without head replacement.

The prefix tensors are read-only. Each intervention branch receives an independent cache object, and the selected head mask and donor row are aligned with the original prompt/branch indices. Multiple head sets for the *same* prompt can be evaluated in parallel; each distinct prompt has its own prefix and donors. There is no addition of isolated single-head scores to approximate a set effect. This optimization does not apply unchanged to an intervention that alters prefix tokens or earlier positions, such as the original SEQ scope experiment.

Table 47: Hardware-specific execution defaults. Sample budgets are unchanged. Run metadata and engine configurations record the actual settings; imported measurements retain their original hardware provenance.

Setting	A100	L4	T4
Initial intervention branches per prompt	64	32	8
Logits-only NL behavior microbatch	16	8	2
Model precision	FP32	FP32	FP32
Shared execution and validation settings			
Metric set block	256 sets		
Prefix CPU cache budget	768 MiB		
Suffix candidates	1, 8, 32, 64 tokens		
Actual-pair probe sets	Up to 3 per registered numerical condition		
Replay raw-readout tolerance	0.002 absolute		
Replay normalized-score tolerance	0.002 absolute		
Full-logit guard	0.01 absolute		
Branch-order/isolation tolerance	0.0004 absolute		
Empirical refinement safety factor	2		
Periodic replay recheck	Every 5,000 replay calls		
Minimum warm-prefix speedup gate	2, unless explicitly overridden and recorded		

Execution microbatches control memory use, not the number of evaluated prompts, candidate sets, or null trials. Hardware-specific defaults are listed in Table 47; run metadata records the actual settings.

For output-only NL behavior, the execution path batches unpadded equal-length prompts and stores only required final logits, instead of caching every head's activation. Causal selection and donor construction retain the head activations they require. A cache of pure zero-ablation outputs can reuse an identical model/prompt/head-set/readout calculation across label permutations. The raw cache does not apply query-orientation signs or high/low labels; those are applied later. Patching caches additionally depend on donor and replacement identities.

F.3 Validation stages and native candidate refinement

Software checks. Before the pretrained-model experiments, the notebook runs synthetic causal-attention regressions, model-interface checks, and a full small-budget software workflow. These exercise both GPT-NeoX and GPT-2 layouts, cache isolation, interventions, migration, and resume behavior. Synthetic tests check software behavior and are distinct from pretrained-model measurements.

Pretrained replay checks. Each pretrained checkpoint is compared with its full native forward path on symbolic and NL probe prompts, original and `yes/no` readouts, and nonempty head interventions. Candidate suffix lengths and branch sizes are tested before the main computation. Raw readout differences, full-logit maximum errors, and reversed branch-order errors use the thresholds in Table 47. The initial global diagnostic also records normalized score errors for sufficiently large probe contrasts ($|\widehat{\Delta}| \geq 0.05$). The global diagnostic is followed by checks for each condition on the actual paired data. Every record of raw values includes the compared readout IDs, baseline, intervention values, denominator, and errors.

Common native baseline. Both the replay and reference score calculations use the same native clean $\bar{y}_H$ and $\bar{y}_L$. Only intervention outputs differ in the comparison. Scores computed with a clean baseline computed with replay are separate diagnostics; reported effect estimates use the native clean baseline. Small raw differences can be amplified when $|\widehat{\Delta}|$ is small; passing a full-vocabulary logit check alone is therefore insufficient for normalized shrink and match.

Condition-specific paired-data checks. Each numerical condition is keyed by model, target pair, donor pair, replacement identity, weight, and denominator floor. The validation routine selects up to three evenly spaced nonempty candidate sets from the requested list and compares native and replay per-example readouts, raw contrasts and gains, shrink, both directional matches, their average, and Q_λ. The compared arrays and errors are stored. The maximum observed raw readout error is denoted e_y and the maximum score error e_Q. Conditions with $e_y > 0.002$ route their set-score evaluation to native inference. Non-selection conditions with $e_Q > 0.002$ also use native set scores. The same tolerances apply to small-contrast conditions.

Selection boundaries and winners. Large candidate lists can use replay to screen, followed by native refinement. Let D_a be the denominator in Equation E.5. The empirical readout envelope is $\eta_y = \max\{10^{-7}, 2e_y\}$, and the initial score radius for an unrefined candidate is $2\eta_y/D_a$. Native checks of candidates near the top-k boundary or the winner update that radius to at least twice any observed replay/native Q_λ discrepancy. Candidate intervals that overlap the membership or winner boundary trigger more native measurements. Every selected greedy candidate is re-evaluated with the native forward pass. The single-head screening boundary uses $k = 48$; greedy augmentation uses $k = 1$. Native comparisons resolve exact ties by the same coordinate rule.

This refinement rule uses a finite empirical error envelope, not a proved uniform bound over unmeasured candidates. The error envelope does not establish a uniform ranking-equivalence guarantee beyond the tested conditions. The number of natively rechecked candidates, actual errors, final radius, and selected winner are logged. Raw-check failures route the full affected selection evaluation to native computation instead.

Final arrays and intermediate diagnostics. A single requested final set and the exported selected-set effect arrays use full native reference calculations. Multi-set random controls use replay only if their condition passes the raw and score checks; otherwise their arrays are native. The scan over all vocabulary coordinates with fixed head sets uses native mean logits, because certifying many weak readout ratios with one unrelated output axis would be insufficient. Count-curve intervention calls can use replay with periodic checks of raw readouts. Output records distinguish imported measurements, native reference arrays, replay screening values, and measurements evaluated through other validated execution paths.

F.4 Compute accounting, task continuation, and numerical cohorts

The engine benchmarks full-forward and warm-prefix intervention times on the current GPU before the main computation. The default engine rejects a replay policy with less than a twofold warm-prefix speedup unless the slow-execution override is explicit. This benchmark excludes prefix construction, native refinement, and final native arrays. The benchmark therefore measures local intervention speedup, not end-to-end experiment runtime. Forward counts, input tokens, native-reference time, task duration, and completed null-trial latencies are saved separately. These task-level records are distinct from total experiment runtime.

Checkpoint identity. A checkpoint is indexed by scientific configuration, source hashes, model and tokenizer revisions, numerical policy, input hashes, readout IDs, selected heads, and donor/replacement identity. Restarts preserve these identities and do not count as new experimental seeds. Hardware or critical package changes require a new numerical comparison before reuse.

Importing a previous numerical cohort. A numerical cohort means measurements made with a specified model revision, precision, hardware, and execution/validation policy. Imported tables and raw arrays retain their original metadata. Before direct reuse, the compatibility check reruns 12 stored baseline cases: four regimes at $k \in \{1, 4, 12\}$. It compares saved and current per-example clean, ablated, and patched readouts, together with shrink, directional matches, and Q. The check covers these baseline cases only.

Direct reuse requires every comparison to have per-example readout error at most 0.002 and score error at most 0.002. If all 12 comparisons agree, compatible checkpoints can be used for continuation. If score errors exceed the tolerance but errors in raw readouts remain within the migration limit, completed measurements retain their original numerical cohort, while incompatible unfinished tasks are recomputed under the current configuration. A source/current native readout difference greater than 0.01 stops migration. Measurements from different hardware or precision cohorts are not pooled as repetitions of one configuration. The runtime outputs `NUMERICAL_COHORTS.json` and `table_provenance.csv` record table origins and the corresponding numerical configuration. Synthetic software-test outputs remain separate from pretrained-model measurements.

F.5 Reproduction files and implementation coverage

The accompanying package provides the FP32 implementation and software tests in `anc/code/`, configuration records and source hashes in `anc/protocol/`, and the reported aggregate results in `anc/results/`. The continuation notebook requires a compatible experimental snapshot, which is not included. The supplied FP32 implementation does not cover the separate BF16 T1, QK/OV, LAST/SEQ, objective-weight, tail-length, surface-form, input-occlusion, or six-model experiments.

The environment pins `transformers==5.16.1` and `tokenizers==0.23.2`. The recorded environment is Python 3.13.15, PyTorch 2.11.0+cu128, CUDA 12.8, NumPy 2.5.3, pandas 2.3.3, SciPy 1.18.1, and Matplotlib 3.11.2. Source hashes and configuration snapshots are in `anc/protocol/`. The runtime records the environment and numerical checks for each new execution.

Table 48 lists the condition-level files generated by experiment runs. These output families are distinct from the aggregate files supplied in `anc/results/`. The package does not include a complete experiment-output archive: the per-example effect arrays, run-specific head identities and readout registries, numerical-validation and bootstrap outputs, condition-level FP32 tables, individual permutation trials, original BF16 archives, and postprocessor for intervals on model means are not supplied. The shared NL domain registry is included in `anc/protocol/nl_state_registry.json`. The supplied `anc/results/transfer_availability.json` contains TR coverage counts. Condition-level interval procedures are specified in Appendices E.3 and E.10; Table 6 reports model-level point estimates only. Appendix H summarizes the statistical reporting scope.

Table 48: Condition-level files generated by the FP32 implementation. Rows represent experimental units or conditions. The package coverage is specified in the accompanying text.

Files in the run's `tables/` directory	Recorded content
`primary_buc.csv`	Current-cohort reference sets and primary scores
`operator_audit_summary.csv`, `operator_audit_draws.csv`	Operator summaries and Monte Carlo draws
`t3_local_summary.csv`, `t3_local_subsets.csv`	Local conditions and complete subset values
`t3_greedy_vs_exact.csv`, `t3_full_shortlist_sampled.csv`	Selection/evaluation comparisons and larger-pool samples
`fixed_budget_baselines.csv`, `fixed_budget_random_controls.csv`	Five methods and three control families
`paired_baseline_differences.csv`	Matched method differences and uncertainty
`fixed_set_count_examples.csv`, `fixed_set_count_summary.csv`	Per-example count values and count-mean fits
`unseen_gap_fixed_set.csv`, `fixed_set_cue_transfer.csv`	Fixed-set gap and wording controls
`fixed_budget_cross_regime_transfer.csv`, `order_only_localization.csv`	Equal-k transfer and fixed-count order localization
`readout_random_scan.csv`, `readout_random_summary.csv`, `readout_random_full_buc.csv`	Registered-pair scan and two reselection protocols
`selection_matched_null_trials.csv`, `selection_matched_null_summary.csv`	Trial statistics, eligibility, and permutation inference
`nl_behavior_all_models.csv`, `nl_causal_all_models.csv`, `nl_transfer_all_models.csv`	Model/domain behavior, causal roles, and transfer ratios

G Extended related work

This appendix expands the comparisons in Section 2. Table 49 distinguishes the behavioral target from the intervention and selection rule. A method for finding relevant components and an experimental design for testing their reuse address different questions.

Table 49: Boundary between nearby lines of work. "Same set" means that one selected set is evaluated by ablation and both mean-patching directions.

Work line	Model setting	Behavioral target	Causal criterion
Indirect-object identification and factual recall	Frozen pretrained models	Semantic task behavior	Study-specific interventions
Synthetic algorithm studies	Often task-trained models	Designed algorithm	Study-specific mechanism tests
Patching, ACDC, EAP	Model- and target-dependent	Researcher-supplied target	General localization machinery
CEA/BUC	Frozen pretrained models	Controlled count/order readout; controlled NL bridge	Same set tested by ablation and two-way patching

G.1 Behavioral targets and controlled evidence

Studies of induction, factual recall, and indirect-object identification connect model outputs to particular internal components (Olsson et al., 2022; Geva et al., 2021; Meng et al., 2022; Wang et al., 2022). In those settings, the target behavior supplies a task-specific criterion for an explanation. Task-trained symbolic models offer a different starting point: the training task specifies an algorithmic problem, allowing a discovered mechanism to be compared with the computation required to solve it (Brinkmann et al., 2024). Behavioral studies of structured tasks also examine failures of compositional generalization (Dziri et al., 2023).

CEA does not train a model to map symbols to answers. Its symbolic contrast changes the count or order of fixed tokens and measures the resulting output response. Count sweeps test extension beyond the selection endpoints; fixed-count order comparisons test ordering dependence. Keeping a head set fixed distinguishes persistence of its effect from the ability to select another set. The symbolic target has no external correctness label, so the labeled NL task provides a separate test of behavioral relevance. Lexical balance and query reversal constrain simple shortcuts but do not eliminate all phrase-based strategies or pretrained associations.

G.2 Interventions and component selection

Activation patching replaces internal activations to measure their effect on a supplied output metric; path patching restricts the routes through which a replacement can influence that metric (Goldowsky-Dill et al., 2023; Zhang and Nanda, 2024). ACDC searches a computational graph using intervention tests, while EAP ranks candidate edges using attribution estimates (Conmy et al., 2023; Syed et al., 2024). Their targets, graph units, and search budgets need not match a head-set comparison.

BUC uses head writes at the final prompt position as its intervention unit and evaluates removal plus recovery in both patching directions on the same set. Its scores include set interactions because interventions are evaluated jointly. The BF16 ActPatch baseline instead ranks single-head high-to-low recovery within a sensitivity shortlist before jointly evaluating the selected union. These comparisons change the ranking objective as well as the search procedure. They therefore demonstrate a measured removal–recovery trade-off, not the independent benefit of greedy search over ranking single heads with the same objective or an end-to-end cost advantage (Appendix C.4). No quantitative result in the paper is attributed to a separate reproduction of EAP or ACDC.

G.3 Component reuse across tasks and prompts

Merullo et al. (2024) examine shared components in indirect-object identification and the Colored Objects task. Their intervention analysis provides evidence of common computational steps across these tasks. Nainani et al. (2024) test

the indirect-object-identification circuit under prompt variants and find reuse with additional input edges. These studies establish that circuit reuse can occur; neither reuse nor failure should be presumed for another experimental setting.

More recently, de Varda et al. (2026) study whether circuit overlap and loading predict arithmetic accuracy across numeric and verbal formats. Their predictive question differs from CEA's direct measurement of a fixed source set's intervention effects on specified target conditions. CEA makes no claim to predict correctness on arbitrary new tasks.

Our comparisons distinguish changes in count, ordering, wording, output coordinates, and replacement operators. In the labeled bridge, the fixed symbolic set and an NL-selected set are evaluated on the same NL targets at equal head budget. Source selection does not use NL prompts; the expanded FP32 transfer uses target-specific donor calibration. Positive effects support partial reuse under that intervention protocol. Because both tasks use the `yes/no` readout, the comparison does not isolate shared evidence aggregation from shared answer processing. Likewise, the seven targeted readout pairs and the larger registered scan are different samples, not a paired estimate of a reselection advantage.

G.4 Faithfulness, localization, and statistical tests

Miller et al. (2024) show that circuit faithfulness scores can change substantially with ablation methodology. Our operator audit measures that dependence for fixed CEA head sets under removal and two patch directions. It tests sensitivity of the reported effect to the replacement distribution rather than assuming a unique operator-independent circuit.

Shi et al. (2024) develop statistical tests for behavior preservation, localization, and minimality. Our selection-matched null asks a different question: whether the observed utility is unusual under the specified within-pair label-exchangeability null when the adaptive selection pipeline is repeated. The trial design and eligibility rules are in Appendix E.9. Rejection of that null neither establishes minimality or uniqueness nor assigns semantic correctness to the symbolic response. Causal effects of small sets, hypothesis rejection, and cross-setting reuse remain separate claims.

G.5 Representations and in-context learning interpretations

Analyses of information routing and residual representations provide context for interpreting writes at the final prompt position (Geva et al., 2023; Dar et al., 2023; Ferrando and Voita, 2024). Task-vector and function-vector work studies internal representations of demonstrated functions and their effects in other contexts (Hendel et al., 2023; Todd et al., 2024). CEA retains the identity of selected heads while allowing their activations to depend on the target inputs and donor distribution. A fixed head set is therefore not a fixed activation vector or a complete description of the earlier computation producing its writes.

Bayesian and gradient-based accounts of in-context learning relate model behavior to learning rules under specified tasks and assumptions (Xie et al., 2021; Garg et al., 2022; Akyürek et al., 2022; Von Oswald et al., 2023; Panwar et al., 2024; Bai et al., 2023; Ahn et al., 2023; Li et al., 2023). The exact two-token log-odds identity in CEA supplies a scalar measurement; it does not identify a Bayesian update or optimization algorithm. Our theory instead states conditions for attribution fidelity, shortlist recovery, and greedy utility bounds. The finite empirical audits do not establish all those conditions for the implemented normalized objective.

Remaining mechanism questions. The evidence leaves open whether NL transfer reflects shared count/order processing, shared answer readout, or both; which computations at earlier tokens and along head–MLP paths supply the writes at the final prompt position; and whether the same components remain effective across pretraining checkpoints or longer heterogeneous evidence. These questions require interventions beyond the reported tests with fixed head sets.

H FP32 confidence intervals and statistical scope

Intervals and their sampling units. Table 6 gives descriptive model means and transfer-coverage counts, not model-level CIs. The condition-level procedures in Appendix E.10 do not supply intervals for those means. The two bootstrap intervals in Table 50 use matched clusters as specified in Appendix E.3, with head sets and donor banks fixed. Behavioral and causal-selection seeds remain separate; no simultaneous coverage or inference across independently pretrained models is claimed.

Table 50: FP32 outcomes outside Tables 6 and 7. A dash indicates that no CI is reported; it does not denote zero uncertainty. Only the RANDOM baseline difference and raw null-test utility receive bootstrap intervals here.

Quantity	Point estimate or span	95% bootstrap CI
BUC − sensitivity, FRONT	$+0.006$	–
BUC − sensitivity, BACK	$+0.030$	–
BUC − sensitivity, INTERLEAVE	$+0.031$	–
BUC − sensitivity, RANDOM	$+0.039$	$[+0.015, +0.065]$
Off-diagonal match: observed span	-0.84 to $+0.89$	–
Cue-change shrink: observed span	0.64–0.93	–
Cue-change match: observed span	0.70–0.96	–
Evaluable conditions, of 1024	960	–
Pass fraction: all registered	0.060	–
Pass fraction: evaluable	0.064	–
Completed permutation trials	99	–
Observed raw null-test utility	2.69	$[2.55, 2.80]$
Holm-adjusted permutation probability	0.01	–

Deterministic conditions. At each fixed count, FRONT, BACK, and INTERLEAVE contain one unique symbolic input. Repeated copies of one matched pair provide no independent prompt variation. Consequently, these conditions receive no prompt-bootstrap CI. Their point estimates describe the specified inputs, not uncertainty over a richer prompt population. The RANDOM difference uses distinct matched permutations, and its interval in Table 50 excludes zero.

Finite-design summaries. Off-diagonal and cue spans describe 12 source–target and eight regime–wording conditions, respectively. The observed counts of 960 evaluable conditions and 99 completed null trials are fixed; their values determine the denominators. The passing fractions summarize the registered readout scan without a population CI over alternative registries. Table 7 likewise gives finite operator minima, negative-marginal fractions, ratio minima, and local selection regrets without population intervals. The summaries describe the tested conditions, not a population guarantee or a bound for the full pool.

Permutation uncertainty. With $B = 99$ completed trials and one estimable hypothesis, $\widehat{p} = 0.01$ implies zero null exceedances. The separate 95% Clopper–Pearson interval for the null exceedance probability is $[0, 1 - 0.025^{1/99}] \approx [0, 0.0366]$. It is not a bootstrap interval, an effect-size interval, an adjusted-p interval, or a trial-count interval (Appendix E.9).

NL interpretation. The three model-level mean TR values are between zero and one, and the observed accuracy decreases are positive. These are point estimates, not model-level significance tests. The TR means use 11, 12, and 9 estimable conditions in Pythia-410M, Pythia-1B, and GPT-2, respectively; their 1, 0, and 3 suppressed condition-level intervals do not remove the corresponding finite ratios from the means. Retained condition-level intervals are not model-level CIs. Neither the positive means nor the estimability counts establish every paired-behavior criterion, positive transfer in all target conditions, complete circuit recovery, or differences between models.

BF16 bridge. The BF16 tables supply execution means and standard deviations, not the prompt-pair bootstrap intervals required by the declared criterion. Neither those summaries nor the separate FP32 procedures establish the complete BF16 criterion (Appendix D).