Distilling Large-Scale Machine Learning Interatomic Potentials to Predict Long-Timescale Electrolyte Transport Properties

Ishan Amin,[1,*] Yue Jian,[2,*] Nitesh Kumar,[3] Christine Zhang,[1] Jianwei Lai,[4] Casey S. Mezerkor,[5] Aaron M. Melemed,[6] Jonathan R. Thurston,[6] Neil P. Dasgupta,[6] J. David Bazak,[5] Ethan J. Crumlin,[4] Samuel Blau,[4,†] and Aditi S. Krishnapriyan[1,2,‡]

[1] *Department of Computer Science, University of California, Berkeley, Berkeley, California 94720, USA*
[2] *Department of Chemical and Biomolecular Engineering, University of California, Berkeley, Berkeley, California 94720, USA*
[3] *Materials Sciences Division, Lawrence Berkeley National Laboratory, Berkeley, California 94720, USA*
[4] *Lawrence Berkeley National Laboratory, Berkeley, California 94720, USA*
[5] *Physical and Computational Sciences Directorate, Pacific Northwest National Laboratory, Richland, Washington 99354, USA*
[6] *Department of Mechanical Engineering, University of Michigan, Ann Arbor, Michigan 48109, USA*

arXiv:2610.39513v1 [physics.chem-ph] 29 Oct 2026

Accurately predicting the transport properties of liquid electrolytes is critical for next-generation battery design and requires force fields that are both quantum-mechanically accurate and fast enough for long-timescale simulation. Recent large-scale pretrained machine learning interatomic potentials (MLIPs) approach density-functional-theory accuracy, but remain too slow to access the necessary timescales to converge transport properties. We provide a knowledge distillation procedure to transfer the accuracy of a large-scale "teacher" MLIP into fast, system-specific "students", without any additional ab initio data. Our distilled models run ∼40× faster, reproduce teacher-predicted structural observables, and produce stable 20 ns trajectories in days, rather than the months the teacher would require. Across a diverse set of sodium- and lithium-ion electrolytes, student-predicted self-diffusivities, ionic conductivities, and shear viscosities track experiment and outperform classical force-field predictions, establishing distillation as a practical route to accurate, long-timescale simulations necessary for high-throughput electrolyte screening.

Ion transport is an essential factor in battery electrolyte performance. Transport coefficients such as ionic conductivity, the per-species self-diffusivities, and the shear viscosity determine an electrochemical cell's rate capability, power density, and low-temperature effectiveness [1]. Since the number of viable electrolytes is combinatorially large, thorough experimental exploration of the space is costly and time-consuming. Atomistic simulations offer an alternative for rapid screening without requiring manual synthesis and testing, provided the simulation method can both accurately and quickly predict the transport properties of the electrolyte [2–4].

However, calculating converged transport coefficients through simulation requires not only quantum-mechanically accurate forces and stresses, but also multiple nanoseconds of well-sampled molecular dynamics (MD) simulations averaged over multiple replicas [5–7]. Evaluating self-diffusivity and ionic conductivity requires calculating the mean-squared displacement across extended timescales, while viscosity relies on time-integrating autocorrelation functions until they fully decay [5, 8]. As a result, these calculations depend on extended lag times, decreasing the number of independent sampling windows, which in turn increases the statistical noise [9]. Ab initio methods are too expensive for the millions of steps needed to converge transport properties, while classical force fields, though computationally efficient, often misrepresent solvation structures and ion–ion correlations [10]. Per-system fitting to ab initio or experimental data can improve classical force-field accuracy, but is impractical for high-throughput screening.

Machine learning interatomic potentials (MLIPs), which approximate ab initio accuracy at a fraction of the cost [11–18], have been increasingly deployed to compute electrolyte transport properties, though existing efforts remain limited in chemical scope [19–23]. Prior works mostly rely on bespoke datasets that would need substantial expansion to cover the entire space of potential battery candidates. BAMBOO and PhyNEO-Electrolyte are trained on data covering at most a few dozen solvents and salts, and validate transport only for lithium electrolytes [19, 20]. Ju et al. applied SevenNet-0, pretrained on inorganic crystals, zero-shot to lithium electrolytes, but had to fix the simulation box to experimental densities in order to reproduce experimental diffusivities [21]. ByteFF-Pol trains a graph neural network to predict the parameters of a polarizable classical force field; while its dataset includes hundreds of solvent molecules, a significant advancement, it contains only three anions and lithium as the only cation [22]. Wang et al. go furthest toward a universal electrolyte potential, creating a dataset covering multiple cations, but their dataset is still only 160,000 structures and the model is again only validated on lithium systems [23]. In each case, moving beyond the demonstrated scope would require costly large-scale ab initio data generation.

Recently, large-scale MLIPs pretrained on chemically

* These authors contributed equally to this work.
† smblau@lbl.gov
‡ aditik1@berkeley.edu

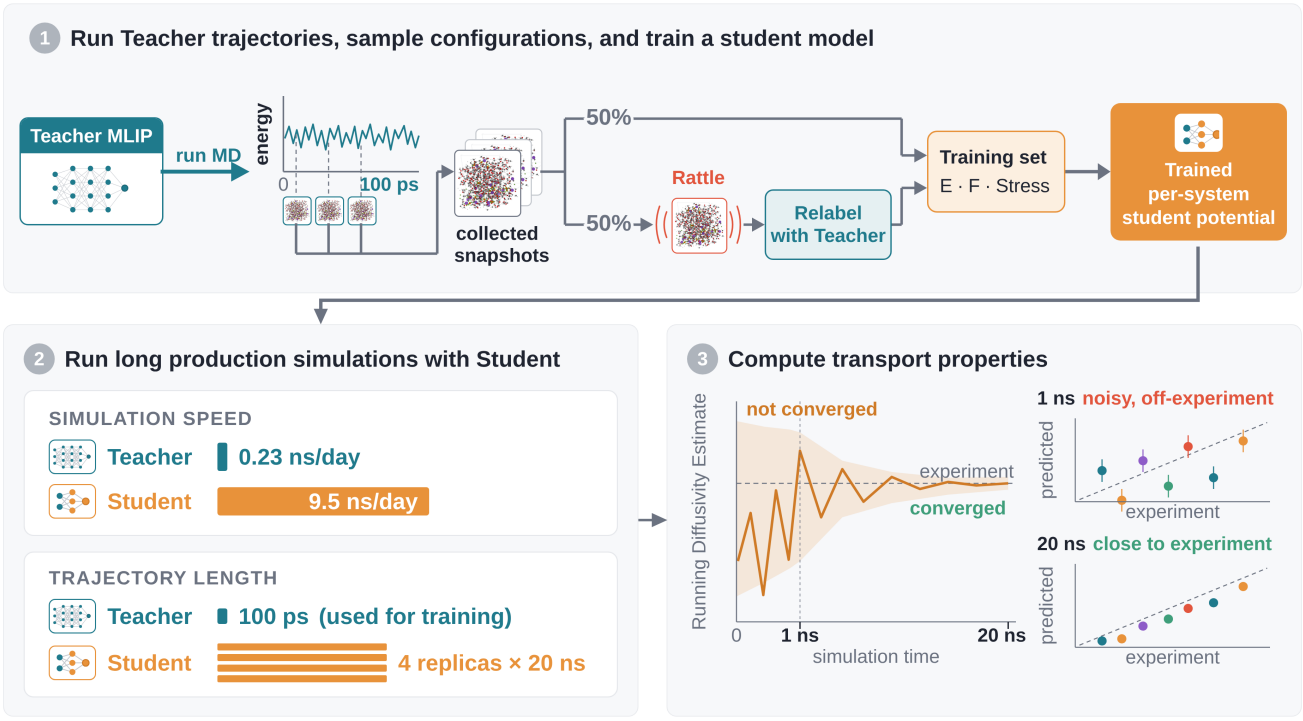

FIG. 1: Teacher–student distillation workflow for converged electrolyte transport properties. The workflow distills an accurate but slow teacher potential into a fast, system-specific student that can be run long enough to converge transport coefficients. **(1) Data generation and distillation.** The teacher potential first runs a short reference trajectory, from which training configurations are sampled. To broaden coverage beyond near-equilibrium structures, half of the sampled configurations are rattled with Gaussian noise and relabeled with the teacher to provide consistent energy, force, and stress targets. The student potential is then trained on this mixed dataset of both perturbed and unperturbed samples. **(2) Production simulations.** Once trained, the student replaces the teacher for production molecular dynamics, running ∼40× faster and reaching trajectory lengths of 20 ns. **(3) Convergence and validation.** The 20 ns trajectories allow transport observables to converge. Transport coefficients start out noisy at short timescales (1 ns), but converge near experimental values once the trajectory is sufficiently long.

diverse datasets have demonstrated the ability to extrapolate to systems not seen during training, reducing the need for system-specific fitting [24–28]. In particular, the Open Molecules 2025 (OMol25) dataset [29] thoroughly samples electrolyte configurations across the periodic table, unlocking the potential of models pretrained on the data to accurately simulate diverse electrolyte compositions. Prior work has shown that OMol25-trained MLIPs, such as the Universal Model for Atoms (UMA) [30], can closely reproduce short timescale structural properties of sodium-based electrolytes [31]. However, the size and speed of the models needed to fit such datasets make them too slow to reach the timescales required for converged transport coefficients.

Knowledge distillation of MLIPs offers a way to use OMol25-pretrained MLIPs to predict electrolyte transport across chemistries, by accessing the speedups granted by smaller models without sacrificing accuracy. A large pretrained MLIP (the "teacher") labels configurations of a specific system or region of chemical space through a systematic procedure. Those configurations are used to train a smaller, faster model with fewer parameters (the "student") [32–36]. The student can then simulate in the teacher's place, ideally with the same accuracy as the teacher but orders of magnitude faster, without any new expensive ab initio labels.

The ability of distilled models to accurately predict long-timescale transport behavior, however, has not been systematically tested. Prior distillation work has primarily focused on static observables, such as energies, forces, and radial distribution functions, mostly computed from trajectories of a nanosecond or less [37–39]. To date, transport properties reported for distilled potentials have largely been limited to a single transport coefficient for a handful of systems per study [34, 35, 40–45]. Excellent performance on static observables over short trajectories does not necessarily imply correct dynamical properties, or student fidelity across long timescales [46–49]. Accurate transport properties therefore require a distilled model to both remain stable and track ion–ion correlations across long, thoroughly sampled simulations.

In this study, we provide a procedure to distill large-

scale MLIPs into students capable of running multi-nanosecond simulations, which are necessary to converge transport coefficients. Our simulations both preserve teacher-model accuracy and track experimentally measured transport properties. To our knowledge, this is the first application of knowledge distillation to transport coefficients in liquid electrolytes, validated against experiment across more than a dozen systems. Our distilled models run ∼40× faster than their teacher, and enable simulations two hundred times the length of the trajectories used to generate their training data. This allows us to access dynamical regimes otherwise inaccessible to the original teacher model. To verify that our students are faithful to the teacher, we periodically spot-check the student's energy, force, and stress errors against the teacher with single-point calculations throughout the simulation, at intervals that make it relatively cheap to do so. We apply our approach to sixteen next-generation candidates for battery electrolytes, focusing on sodium-ion systems [50, 51], with various solutes, solvents, and concentrations (detailed in the Results).

For each salt–solvent pair, we distill the small variant of UMA evaluated with its OMol task head (checkpoint `uma-s-1p1`; hereafter referred to as UMA-OMol) into a system-specific student MLIP built on PaiNN, an equivariant graph neural network [52]. The student is trained on a set of configurations sampled from a 100 ps teacher trajectory, with some of the structures noised and relabeled to ensure student robustness. We then run four independent 20 ns replica simulations per system, and periodically check the student model's errors against single-point predictions from UMA-OMol along each trajectory to confirm that they remain low and stable. The student trajectories are then used to compute the self-diffusivities, ionic conductivities, and shear viscosities for every electrolyte system, which are compared against experiment. An overview of the workflow is shown in Fig. 1.

Across all sixteen systems, the distilled potentials capture experimental trends and agree more closely with experiment than the widely used OPLS-AA force field [53], reducing the median deviation from experiment for every transport property studied. By transferring the accuracy of pretrained MLIPs into compact models, distillation enables high-throughput electrolyte transport simulations that achieve a balance of accuracy and efficiency unavailable from classical force fields, direct ab initio MD, or large-scale models alone. More broadly, the approach can be used to evaluate the behavior of the large-scale pretrained MLIP on timescales it cannot directly access.

Results

The distilled PaiNN potentials enable simulations long enough to converge transport coefficients, tracking experiment across sixteen electrolyte systems. We first evaluate the student's faithfulness at the 1 ns timescale accessible to the teacher, by comparing both models' radial distribution functions and mass densities. We additionally check the student's per-frame force and stress errors periodically across 20 ns of student simulation. The students' ∼40× speed advantage enables 20 ns simulations, from which we compute ionic conductivities, shear viscosities, and self-diffusivities. Student-predicted transport properties, benchmarked against experiment, outperform the OPLS-AA classical force field [53]. The sixteen systems are sodium electrolytes combining $NaPF_6$ or NaOTf with monoglyme (DME), diglyme (DEGDME), tetraglyme (TEGDME), or propylene carbonate (PC) at 0.1 and 1 M ($NaPF_6$/TEGDME at 1 M only), plus a lithium baseline, $LiPF_6$ in DME at 0.5 M. Full compositions are listed in Table S1; details of the box setup, MD ensemble settings, and student training are given in the Methods and Secs. S1–S2 of the Supplementary Information (SI).

Reproducing the behavior of the teacher model

At the 1 ns timescales accessible to the teacher, the student reproduces both its radial distribution functions (RDFs) and its mass density. The RDF curves of the student and teacher overlap across all pair types, and the agreement holds across all 1 M systems (full set in Sec. S6), indicating that the student accurately captures the teacher's first-shell coordination as well as its longer-range features (Fig. 2a). The distilled model also matches UMA-OMol's mass densities, agreeing with UMA-OMol to within 1.1% on average across all sixteen electrolytes (Fig. 2b). Both models overestimate the experimental densities by 6–15% (about 10% on average); however, this is a known tendency of UMA-OMol [31], and distillation adds little error to the existing bias.

The student additionally maintains low force and stress errors relative to the teacher throughout long simulations, as a result of augmenting the training data with rattled structures. For a representative system (1 M NaOTf in TEGDME), we perform single-point UMA-OMol force and stress evaluations at 20 ps intervals along a student 20 ns NPT trajectory. This serves as a low-cost spot check on the student's accuracy that avoids running the teacher itself for the full 20 ns. The force mean absolute error (MAE) against UMA-OMol is around 7 meV/Å and the stress MAE near 30 bar. Both errors are low with no upward drift (Fig. 2c). The force errors stay similarly low along the NVT production trajectories used for the transport coefficients, for every system (Figs. S3 and S4). The student trajectory's stability is a direct result of adding rattled structures to the training set. We isolate the effect of including noised data during training in Sec. S3. Trained without noised data, the student quickly wanders out of the narrow configurational region its training covers. The student's force error grows as the simulation progresses, until the system destabilizes within a nanosecond. Exposing the student to off-equilibrium configurations teaches it a restoring force that corrects small

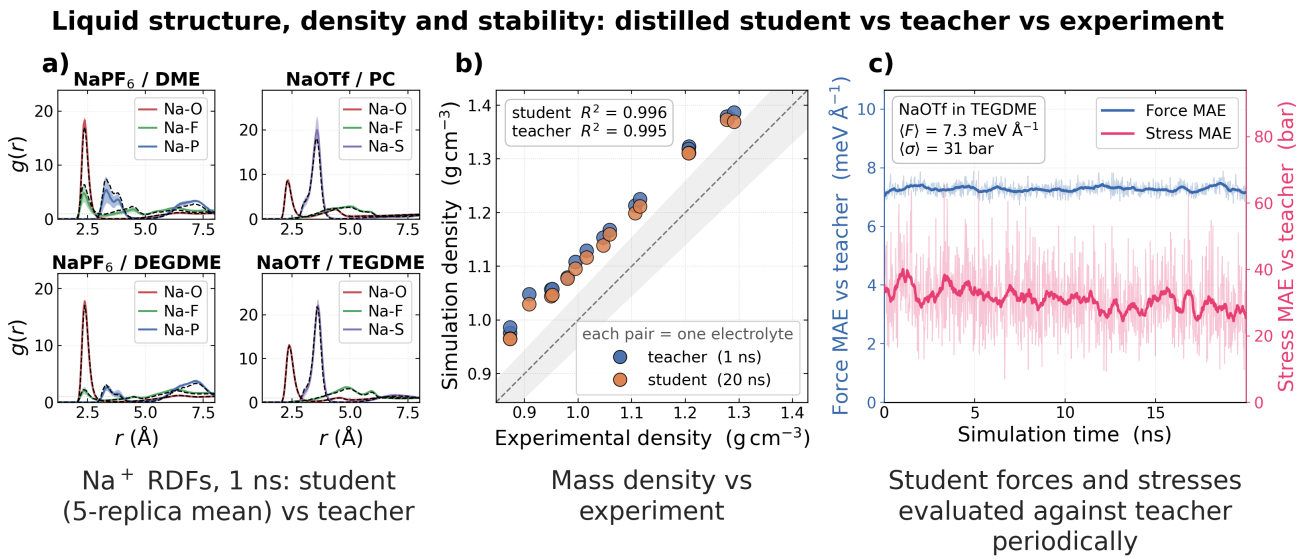

FIG. 2: Distilled PaiNN potentials reproduce the static observables of the UMA-OMol teacher at short timescales. **(a)** Radial distribution functions of four electrolyte systems across multiple correlation pairs, obtained from 1 ns trajectories at 293 K, with both the teacher model (UMA-OMol) and the student (PaiNN) displayed. Each PaiNN line is the mean over five independent 1 ns replicas, where the shading indicates the ±1 s.d. replica spread. The teacher curve falls within the shaded band for almost every pair and system, indicating that the teacher lies within the student's per-replica noise and that the student accurately reproduces teacher behavior. **(b)** Parity plot of simulated versus experimentally measured mass density, where each dot denotes a different electrolyte system. Both UMA-OMol and PaiNN show a systematic upward bias against experiment, but UMA-OMol and PaiNN coincide for almost every system, so distillation adds little error of its own. The individual densities for each electrolyte are listed in Sec. S7. **(c)** For a 20 ns *NPT* 1 M NaOTf in TEGDME simulation, we evaluate force mean absolute error (MAE) and stress MAE relative to UMA-OMol at 20 ps intervals by using UMA-OMol to compute single-point estimates. Light traces are per-frame errors, while bold curves show the ∼500 ps running average. Both errors stay low, with no upward drift over 20 ns: a linear fit changes the force MAE by <0.2%, and the stress MAE decreases slightly over time. Similar plots for the production runs are shown in Sec. S5.

deviations and prevents instability from compounding.

Computing transport properties with the student model

The distilled models' ∼40× speedup over the teacher enables converged transport properties that track experimental trends and consistently outperform OPLS-AA, the classical force-field baseline. Using the distilled potentials, we run 20 ns production simulations and extract ionic conductivities, shear viscosities, and self-diffusivities, which are then compared to experiment (Figs. 3 and 4). Every reported coefficient is the mean over four 20 ns NVT replicas launched from a common configuration with independently seeded velocities. Along each trajectory, we confirm that the energy, temperature, and pressure are stable, and that the force error against the teacher remains low for the full 20 ns (Sec. S5). For each property, we evaluate model performance against experiment using three metrics. A log–log Pearson correlation coefficient (r) near unity indicates that predictions tightly track experiment; a low median absolute percentage deviation (MdAPD) indicates small typical errors; and a log–log linear regression slope near one indicates that the model reproduces the experimental spread across systems.

Computed student conductivities track experiment across over an order of magnitude, and possess a lower MdAPD than OPLS-AA. Results are reported for the nine concentrated (≥0.5 M) systems; at 0.1 M the box contains too few charge carriers for the collective estimator to converge. Seven of the nine systems sit within a factor of two, and the model reaches an MdAPD of 36% (Fig. 3a). OPLS-AA is notably less accurate, as it underpredicts seven of the nine conductivities, places only three within the factor-of-two band, and has an MdAPD of 67%. Its log–log correlation ($r = 0.72$) is comparable to the students' (0.66) (Fig. 3b). The student does have a large error for 1 M NaOTf in DME, overpredicting the conductivity by a factor of five, unlike OPLS-AA. This system has the lowest ionicity of the concentrated systems, and in Sec. S12 we use ion-pair scans to show that this error likely arises from the teacher's finite 6 Å graph cutoff, which causes ions to pair less than they should.

Student viscosities likewise track the experimental ranking closely; their consistent upward bias is attributable to the teacher's systematic density overshoot. Because the students inherit the teacher's overestimate of the

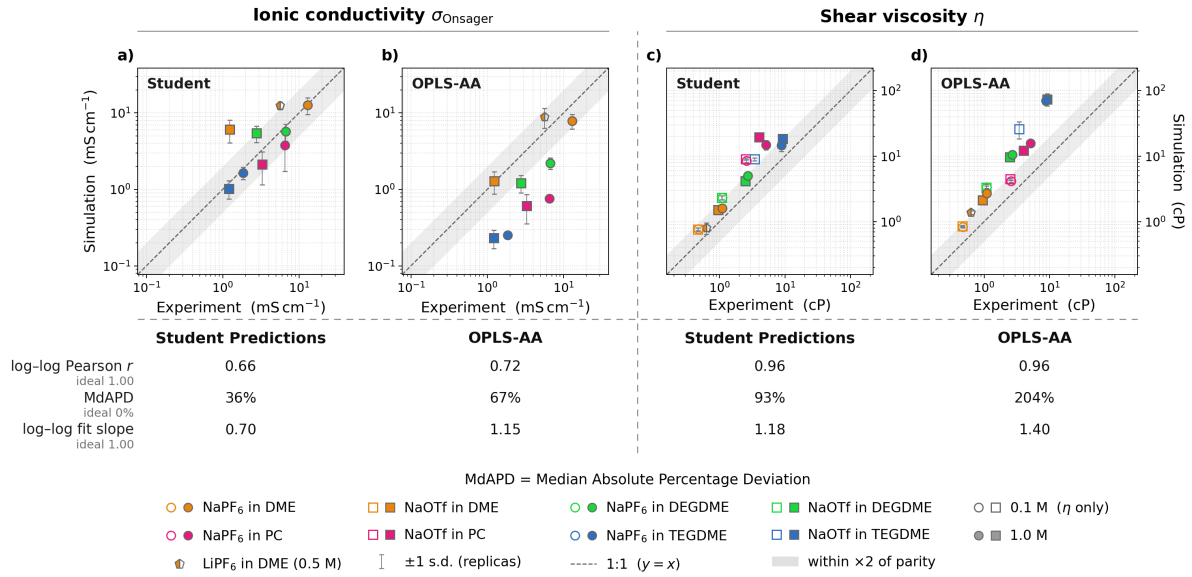

FIG. 3: Ionic conductivity and shear viscosity vs. experiment at 298 K, for the distilled student models and the OPLS-AA force field. Each point is the mean over four independent 20 ns replicas, and error bars show ±1 s.d. across the replicas. **(a,b)** Ionic conductivity σ for the PaiNN student and OPLS-AA. Conductivity is only computed for the concentrated ($\geq$0.5 M) systems, since at 0.1 M the box holds too few charge carriers to converge the collective estimator. The student roughly halves the median absolute percentage deviation of OPLS-AA (MdAPD; 36% vs. 67%), with a comparable log–log correlation. **(c,d)** Viscosity η against experiment for the two force fields. Again the student tracks experiment much more closely than OPLS-AA, with an MdAPD less than half that of OPLS-AA and a log–log fit slope closer to one.

liquid densities (6–13% for the students; Fig. 2b), the simulated liquids are all over-compressed compared to experiment. Most of the reduction in space comes out of the free volume in the box (the volume not occupied by the molecules) [55, 56]. In Sec. S11.2, using high-pressure viscometry data of comparable liquids from the literature, we estimate that a density overshoot of this size should inflate our predicted viscosity by roughly three- to four-fold. Consistent with this, every student viscosity lies above experiment, by 2.1× on geometric average (Fig. 3c). Student predictions are tightly correlated with experiment ($r = 0.96$) and mostly follow a linear trend, with a log–log slope of 1.18. OPLS-AA likewise overpredicts the viscosity, but much more strongly, with an MdAPD of 204% against the students' 93% (Fig. 3d). OPLS-AA's overprediction is also much less uniform, growing as the viscosity itself grows (log–log fit slope 1.40), from under 2× for the least viscous systems to about 8× the experimental value for the two 1 M TEGDME systems.

Self-diffusivities follow the same pattern, with the distilled models reproducing experimental trends more faithfully than the classical baseline (Fig. 4). The student models achieve $r = 0.92$–0.98 and MdAPDs of 55%, 41%, and 56% for D^+, D^-, and D^0. The students underpredict the diffusivities by roughly 2×, which mirrors the overshoot we saw for viscosity: by the approximate Stokes–Einstein relation $D \propto k_B T/\eta$, a 2× increase in viscosities should result in a decrease in diffusivities by roughly the same factor, which is what we observe. Both the viscosity overprediction and diffusivity underprediction are likely caused by the density bias inherited from the teacher, not by shortcomings introduced by distillation. OPLS-AA performs notably worse, underpredicting the diffusivities by roughly a factor of three with MdAPDs of 63–71% (Fig. 4d–f). Both models achieve comparable log–log correlations, but the distilled models follow a much more linear relationship against experiment, with log–log slopes of 1.09–1.21 compared with 1.30–1.57 for OPLS-AA.

Beyond the absolute values, the students recover expected chemical trends. Self-diffusivities decrease with both salt concentration and glyme chain length (DME > DEGDME > TEGDME). The solvent moves fastest in each electrolyte, and the predicted ordering of the cation and anion diffusivities agrees with experiment, except in cases where the two are nearly equal.

The transport properties reported here required thoroughly sampled MD simulations; trajectories limited to 1–5 ns were insufficient to converge the coefficients (Sec. S13). Critically, using the teacher to simulate these trajectories would be impractical; at UMA-OMol throughput (~0.23 ns/day per system, Sec. S4), each 20 ns replica would take roughly three months to finish. In contrast, the PaiNN student models complete 20 ns in about two days.

Discussion

The student's combination of speed and accuracy makes large-scale pretrained MLIPs practical for trans-

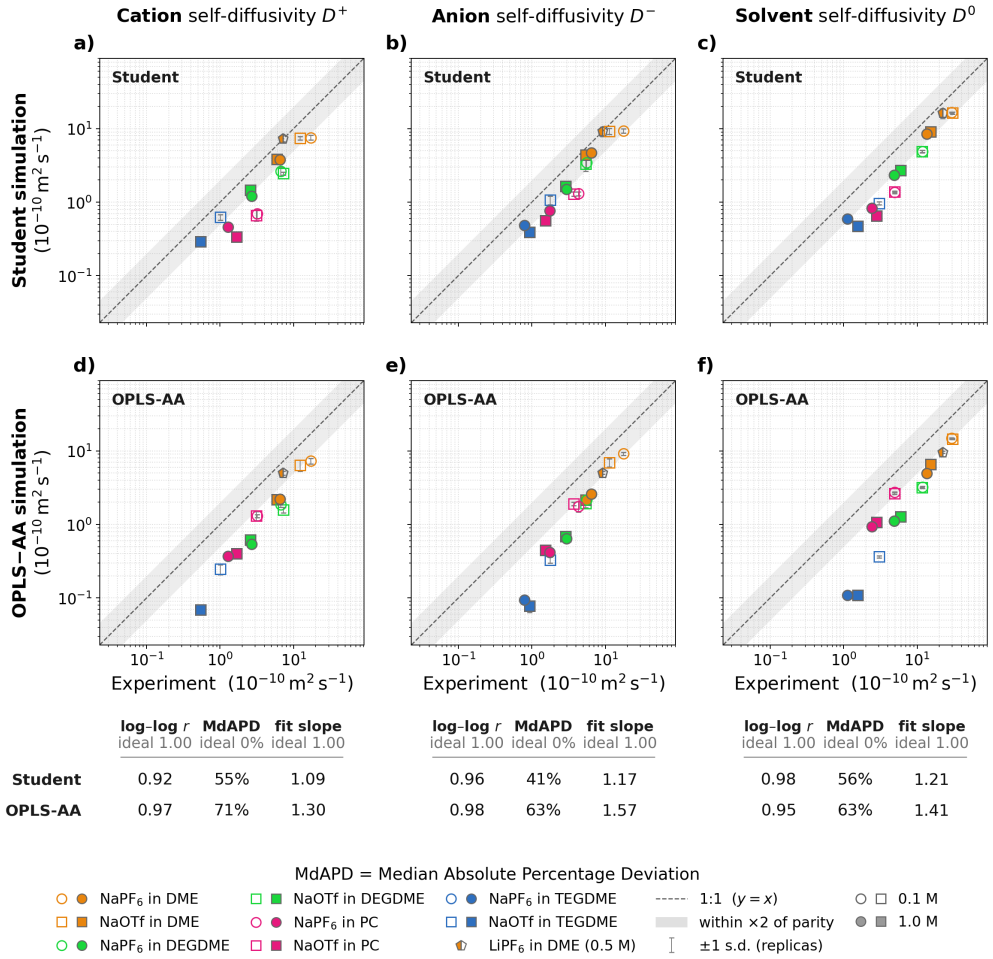

	Cation D^+			Anion D^-			Solvent D^0		
	log–log r (ideal 1.00)	MdAPD (ideal 0%)	fit slope (ideal 1.00)	log–log r (ideal 1.00)	MdAPD (ideal 0%)	fit slope (ideal 1.00)	log–log r (ideal 1.00)	MdAPD (ideal 0%)	fit slope (ideal 1.00)
Student	0.92	55%	1.09	0.96	41%	1.17	0.98	56%	1.21
OPLS-AA	0.97	71%	1.30	0.98	63%	1.57	0.95	63%	1.41

FIG. 4: Self-diffusivities vs. experiment at 298 K, for the distilled students and the OPLS-AA force field. Self-diffusivities of the cation D^+ **(a,d)**, anion D^- **(b,e)**, and solvent D^0 **(c,f)**, for the PaiNN student (top row) and OPLS-AA (bottom row), where each point is the mean over four independent 20 ns replicas (error bars, ±1 s.d. across the replicas), with the Yeh–Hummer finite-size correction applied [54]. The student outperforms OPLS-AA for every species, both in MdAPD and in log–log fit slope.

port screening. A new electrolyte requires only a short teacher trajectory, roughly ten GPU-hours of student training, and about two days of production MD to obtain transport coefficients. Although each student is specific to a single salt–solvent pair, the cost of training is small compared to the time saved by using the student for inference. Because OMol25 covers a broad range of molecular chemistry, the teacher can be applied to new electrolytes beyond the systems studied here, without additional ab initio data.

The student's fidelity to the teacher bounds the accuracy of its predictions, but it also provides a new way to evaluate the long-timescale behavior of large-scale pretrained MLIPs. UMA-OMol overestimates the densities of our electrolytes by about 10% on average, and distillation reveals that this translates into a systematic ∼2× viscosity overshoot and a corresponding ∼2× diffusivity undershoot. Without access to the long trajectories produced by the student, we would not know whether this error was due to statistical fluctuations or an actual shortcoming of the teacher model. The student's inheritance of teacher shortcomings diminishes accuracy against experiment; however, this is a transient limitation. Large-scale pretrained models are quickly improving, with training datasets growing in diversity and new architectures incorporating explicit long-range electrostatics [26] that directly remove the cutoff artifact identified in Sec. S12. Our pipeline consumes only the energies, forces, and stresses the teacher assigns, and is agnostic to the choice of architecture or the data the teacher was trained on. Every improvement to the large-scale model therefore propagates to the distilled transport predictions, without any change to the method.

Large-scale pretrained models have been shown to reproduce the structure of liquid electrolytes; we show distillation extends them to electrolyte dynamics at scale. Converged transport coefficients at the accuracy of large-scale models have so far required a months-long computa-

tion per system. Knowledge distillation reduces this cost enough that such calculations could be incorporated into the screening loops needed for next-generation electrolyte design.

Methods

Electrolyte systems and simulation cells

We study sixteen systems in this work: $NaPF_6$ and NaOTf in monoglyme (DME), diglyme (DEGDME), tetraglyme (TEGDME), and propylene carbonate (PC) at 0.1 and 1 M ($NaPF_6$ in TEGDME at 1 M only), plus $LiPF_6$ in DME at 0.5 M; per-system compositions and box sizes are listed in Table S1. The final boxes for production molecular dynamics (MD) runs contain 2700–3700 atoms. For MD simulations, we prepare initial configurations using an earlier structural study's protocol [31]. Specifically, we use PACKMOL to place molecules into cubic cells that are 30 Å on each side [57], perform geometric relaxation, then equilibrate each box in the isothermal–isobaric (NPT) ensemble with the OPLS-AA force field [53] for 1 ns to relax the initial packing. Because the boxes are packed at a non-equilibrium volume, the molarity drifts significantly from its target when the barostat relaxes. To mitigate this, we use each system's simulated equilibrium density to recompute the correct solute-to-solvent ratio, and repack the box. The corrected boxes equilibrate to within 1.4% of the target molarity on average. We provide the derivation of our correction, as well as per-system shifts, in Sec. S1.

Teacher reference data

Reference data for training the distilled models is generated with UMA-OMol (checkpoint `uma-s-1p1`) using its OMol task head [29, 30]. For each system and concentration we run a 1 ns trajectory in the NPT ensemble at 293 K and 1 bar using the Atomic Simulation Environment (ASE) [58], with a Nosé–Hoover thermostat (coupling time 100 fs) [59, 60], a Parrinello–Rahman-type barostat [61], and a 1 fs timestep. We then sample the first 100 ps of each trajectory to make the training configurations; the last 70% of the trajectory is used as the reference for the radial distribution function comparison between student and teacher.

Unperturbed snapshots are sampled every 50 fs from the teacher trajectory. A second set of snapshots, sampled every 50 fs but offset by 20 fs from the first grid, is collected and then perturbed by adding Gaussian noise with a standard deviation of 0.1 Å to every atom. The simulation cell is not perturbed. Previous work has found that augmenting training with Gaussian-noised structures improves the robustness of MLIPs [62]. The noised points allow the model to see a wider region of configuration space, and help stabilize the distilled potential during simulations. We isolate the effect of the noised data in Sec. S3. Every configuration is labeled with UMA-OMol energies, forces, and the full stress tensor. We use a 90:10 train–validation split, giving 3600 training and 400 validation configurations per concentration. When two concentrations exist for a salt–solvent pair, we combine them into one dataset so that a single potential is trained across both loadings.

Student potential and training

We use the PaiNN equivariant message passing graph neural network [52] for our student model, as implemented in the FeNNol repository in JAX [63]. The model has three message passing layers, an embedding dimension of 128, and a radial neighbor cutoff of 5.1 Å. We also add a trainable Ziegler–Biersack–Littmark (ZBL) term [64] with an inner cutoff of 1.1 Å for explicit short-range nuclear repulsion. To train the model, we minimize the loss

$$\mathcal{L} = \frac{|E - \hat{E}|}{N} + \frac{1}{N}\sum_{i=1}^{N}\|\mathbf{F}_i - \hat{\mathbf{F}}_i\| + \frac{1}{9}\sum_{\alpha\beta}|\sigma_{\alpha\beta} - \hat{\sigma}_{\alpha\beta}|, \quad (1)$$

where N is the number of atoms and hats denote UMA-OMol labels. We use AdamW [65] as the optimizer and train for 275 epochs with a one-cycle cosine learning-rate schedule (peak learning rate 8×10^{-4}). Training a single model takes 9–22 hours on one NVIDIA GH200 GPU, depending on the number of pooled concentrations (Table S3). The full architecture and optimization settings are given in Table S2.

Structural validation against the teacher

Each distilled potential is validated against the teacher using radial distribution functions and mass densities, which are short timescale observables. Between species A and B, the radial distribution function is defined as

$$g_{AB}(r) = \frac{1}{N_A\,\rho_B}\left\langle \sum_{i\in A}\sum_{j\in B}\frac{\delta(r - r_{ij})}{4\pi r^2}\right\rangle, \quad (2)$$

where r_{ij} is the minimum-image distance between atoms i and j, N_A is the number of A atoms, and $\rho_B = N_B/V$ is the bulk number density of B. For each system, we run five independent 1 ns student NPT simulations at 293 K and compare them with a matched 1 ns UMA-OMol run. Each RDF is computed identically, using 3000 configurations sampled from the last 70% of the trajectory, with a 0.05 Å bin width out to 12 Å (Sec. S6). The 0.1 M cells only contain two ion pairs, which is not enough to converge the RDFs in 1 ns of simulation. We calculate the students' mass densities from a 20 ns NPT trajectory at 298 K, averaged after discarding the first 0.5 ns, and obtain the corresponding UMA-OMol densities from our

own 1 ns NPT simulations at 298 K with the teacher checkpoint.

Production simulations

We run production simulations with the student model using the JAX-based MD engine implemented in FeNNol [63, 66]. Each molarity-corrected box (Sec. S1) is first equilibrated for 1 ns in an NPT ensemble at 298 K and 1 bar, using a Langevin thermostat (friction 1 ps^{-1}) and an isotropic Langevin-piston barostat (piston relaxation time 200 fs). We compute the mean density over the last 500 ps of this trajectory and select the frame from the last 300 ps whose instantaneous density lies closest to the mean. This frame is then used as the starting configuration of the production simulation (Sec. S1.3).

The four 20 ns NVT production replicas run at 298 K, each starting with different initial Maxwell–Boltzmann velocities and thermostat noise. We use the global Bussi (CSVR) velocity-rescaling thermostat [67] with a 10 ps relaxation time. Holding the volume fixed removes barostat fluctuations from the stress autocorrelations, and the slow global thermostat leaves sub-picosecond correlations unbiased [5, 68]. We use a 1 fs timestep and save frames every 1 ps and the stress tensor used to calculate the shear viscosity every 2 fs. All MD settings are reported in Table S5.

To verify that the distilled potentials remain consistent with UMA-OMol over these long trajectories, we perform single-point UMA-OMol evaluations on 500 frames per student trajectory, spaced 40 ps apart across the full 20 ns, and monitor the mean and maximum force errors throughout the simulation, together with other observables (potential energy per atom, temperature, and pressure). We observe low and flat force errors for every system, and the other diagnostic observables are very stable throughout each trajectory (Sec. S5).

Computational throughput

We test the throughput of both student and teacher models on a single NVIDIA GH200. The student reaches a median of 9.5 ns/day across the benchmarked boxes versus 0.23 ns/day for UMA-OMol, a 41× speedup. Per-system throughputs are reported in Sec. S4.

Classical force-field baseline

We use GROMACS 2025.3 [69] to run simulations with the OPLS-AA classical force field [53]. Solvent parameters (DME, DEGDME, TEGDME, PC) are OPLS-AA with 1.14*CM1A charges from LigParGen [70]. Both anions use the CL&P parameter set, which is built on the OPLS-AA functional form: PF_6^- from Canongia Lopes et al. [71] with the bonded terms of Kaminski and Jorgensen [72], and OTf^- from Canongia Lopes and Pádua [73] with the triflate oxygen charge of Gouveia et al. [74]. Na^+ and Li^+ use the Åqvist parameters [75] as distributed with GROMACS. Lennard-Jones cross terms follow geometric combination rules and 1–4 interactions are scaled by 0.5, as in OPLS-AA. All ions carry their full formal charge. The parameter sources are summarized in Table S6, and the full per-species parameters are provided with the released code. We correct the molarity of the initial boxes with the same procedure as for the student model, but with OPLS-AA's own equilibrium density. Steepest-descent energy minimization is used to relax the packed boxes, and the boxes are then equilibrated for 1 ns of NPT at 298 K and 1 bar (using a velocity-rescaling (CSVR) thermostat with a 0.1 ps relaxation time and an isotropic stochastic cell-rescaling (C-rescale) barostat [76] with a 2 ps relaxation time). Finally, we run four independent 20 ns NVT production replicas at the equilibrated volume; the full settings are tabulated in Table S6. During the production run, we use a 1 fs timestep with no bond constraints (fully flexible molecules), and each replica is initialized with Maxwell–Boltzmann velocities from an independent seed. Particle-mesh Ewald [77] is used for the electrostatics with a 12 Å real-space cutoff, and the Lennard-Jones interactions are cut at 12 Å with an analytic long-range dispersion correction. We save frames every 1 ps and the pressure tensor every 2 fs.

Experimental measurements

All electrolytes were prepared and handled inside an argon-filled glovebox (O_2 below 0.2 ppm, H_2O below 0.1 ppm). The experimental densities, self-diffusivities, ionic conductivities, and shear viscosities were measured in-house; the values are tabulated in Sec. S9. The densities were also used in an earlier structural study [31]. Self-diffusivities were measured by pulsed-field-gradient nuclear magnetic resonance (PFG-NMR) on 11.7 T Bruker Avance NEO spectrometers, following ^{1}H for the solvent, ^{7}Li or ^{23}Na for the cation, and ^{19}F for the anion; the most viscous ^{23}Na samples were measured on a 14.1 T Avance III HD system with a Diff50 probe. Ionic conductivities were measured in the glovebox with a Mettler Toledo SevenDirect SD30 meter and an InLab 710 four-electrode platinum probe, with three independent readings per electrolyte. Density and viscosity were obtained from a single sample filling at 298 K with an Anton Paar DMA 4500 M oscillating U-tube density meter coupled to a Lovis ME rolling-ball viscometer, both calibrated with certified reference water. Full protocols for all four measurements are given in Sec. S8.

Transport-property analysis

All transport coefficients, simulated and experimental, are reported at 298 K. Simulated coefficients are

computed independently for each production replica. To alleviate the noise introduced by early-stage equilibration, we discard the first 2 ns of each trajectory and use the remainder for the computation. Results are reported as the mean and standard deviation over the four replicas.

Self-diffusivity

We calculate self-diffusivity by fitting the slope of the mean-squared displacement (MSD) over time lag τ based on the Einstein relation [9, 78, 79],

$$D = \frac{1}{6} \lim_{\tau\to\infty} \frac{d}{d\tau} \big\langle |\mathbf{r}(t_0+\tau) - \mathbf{r}(t_0)|^2 \big\rangle, \tag{3}$$

where t_0 is the time origin, and $\langle \cdots \rangle$ denotes an ensemble average evaluated over all available time origins as well as over all particles of the species of interest (specifically, cation, anion, and solvent).

We first compute the unwrapped coordinates of each trajectory using the minimum-image convention, then subtract the center-of-mass drift of the entire system. We use the molecular centers of mass for the per-species displacements. MSDs across different time lags are calculated using a fast Fourier transform (FFT) and fitted following the best practices of Maginn *et al.* [5]. Instead of using a fixed lag window, for each species we evaluate the logarithmic slope $\nu(\tau) = \mathrm{d}\ln \mathrm{MSD}/\mathrm{d}\ln\tau$ on logarithmically spaced lags between 5 ps and 36% of the maximum lag. We then fit the MSD over the longest contiguous run satisfying $|\nu - 1| \le 0.15$, requiring at least twelve points. If no such window exists, we perform the fit over the twelve-point stretch with the smallest $|\nu - 1|$ value.

After calculating the periodic-box diffusivity D_{PBC} from the fitted slope, we apply the Yeh–Hummer correction to remove finite box size effects [54], $D = D_{\mathrm{PBC}} + \Delta D$, with

$$\Delta D = \frac{k_{\mathrm{B}} T\, \xi}{6\pi\, \eta\, L}, \tag{4}$$

where $\xi = 2.837297$ for a cubic cell, L is the box edge length, and $k_{\mathrm{B}}T$ is the thermal energy. For η we use the Green–Kubo viscosity of the same system, averaged over its replicas, so the correction is fully self-consistent and involves no experimental input. Both the uncorrected diffusivities and the corrections are listed in Tables S8 and S9.

Ionic conductivity

The ionic conductivity is calculated from the Onsager transport coefficients L_{ij} as [8, 80],

$$\sigma = e^2 \sum_{ij} z_i z_j L_{ij}, \tag{5}$$

where e is the elementary charge and z_i is the signed integer charge of species i. Each ion is treated as a point charge at its center of mass, with $z = +1$ for the Na^+ and Li^+ cations, $z = -1$ for the PF_6^- and OTf^- anions, and $z = 0$ for the solvent molecules. We first define the cumulative displacement of all species-i particles over a lag τ as

$$\Delta\mathbf{R}_i(t_0,\tau) = \sum_{k=1}^{N_i} \left[\mathbf{r}_{ik}(t_0+\tau) - \mathbf{r}_{ik}(t_0)\right], \tag{6}$$

where $\mathbf{r}_{ik}$ is the center-of-mass position of the k-th molecule of species i, measured relative to the system's center of mass, and N_i is the total number of species-i particles in the simulation box. We next define the cross-correlation of species i and j as

$$C_{ij}(\tau) = \langle \Delta\mathbf{R}_i(t_0,\tau) \cdot \Delta\mathbf{R}_j(t_0,\tau) \rangle . \tag{7}$$

The angle brackets $\langle \cdots \rangle$ denote an ensemble average, which is evaluated as an average over all available time origins t_0. The transport coefficient is then given by the long-time slope of the cross-correlation $C_{ij}(\tau)$ as

$$L_{ij} = \frac{\beta}{6V} \lim_{\tau\to\infty} \frac{\partial}{\partial\tau} C_{ij}(\tau), \tag{8}$$

where $\beta = 1/(k_{\mathrm{B}}T)$ and V is the volume of the simulation box. The slope is extracted by a least-squares fit over lag times of 0.5–2 ns. We use a shorter window than for the self-diffusivities, because the displacement for conductivity is collective and does not include particle averaging, making noise a significant issue at longer lags. Unlike the Nernst–Einstein estimate, this correlated estimator retains the ion–ion cross-correlations [8]; the derivation and the comparison between the two estimators are given in Sec. S12. Conductivities are reported for the nine concentrated ($\ge$0.5 M) systems only: at 0.1 M the cell contains only two ion pairs, which is not enough to accurately measure distinct-pair (cross) correlations in L_{ij}.

Shear viscosity

The shear viscosity η is computed from the equilibrium Green–Kubo stress-autocorrelation function [9, 79, 81–83],

$$\eta \;=\; \frac{V}{k_{\mathrm{B}}T} \int_0^\infty \big\langle\, \sigma_{\alpha\beta}(0)\, \sigma_{\alpha\beta}(t) \,\big\rangle \, \mathrm{d}t, \tag{9}$$

where $\sigma_{\alpha\beta}$ is an off-diagonal component, such as σ_{xy}, of the instantaneous stress tensor, logged along the trajectory. To reduce statistical noise, we average the autocorrelation in Eq. (9) over five shear components of the traceless symmetric stress tensor, each of which integrates to the same value η [5, 84]:

$$\sigma_{xy}, \quad \sigma_{xz}, \quad \sigma_{yz}, \quad \tfrac{1}{2}(\sigma_{xx} - \sigma_{yy}), \quad \tfrac{1}{2}(\sigma_{yy} - \sigma_{zz}). \tag{10}$$

The autocorrelation is evaluated by FFT and integrated by the trapezoidal rule to give the running estimate $\eta(t)$, accumulated over lags up to 1 ns. We use the time-decomposition method described by Zhang *et al.* [5, 6] to obtain the reported value. The integral is fit with a double-exponential form over a window that ends at the lag t_{cut}, defined as the first point at which the standard deviation across five equal time blocks of the trajectory exceeds 40% of the running value. Each fit point is weighted inversely by the standard deviation. We extrapolate our fit to infinite time to get η_∞, and this value is reported. This procedure prevents the noisy tail of the integral from biasing the estimate. Convergence checks and example fits are given in Sec. S11.

Use of large language models

Large language models (Claude, Anthropic) were used to assist with writing code. All code was thoroughly reviewed by the authors, who take full responsibility for the content.

Data availability

The trained student models, the initial and molarity-corrected simulation boxes, and the per-replica transport analysis curves are available at `https://github.com/ASK-Berkeley/distillation-electrolytes`. The molecular dynamics trajectories and distillation datasets can be regenerated with the released code, models, and starting configurations.

Code availability

All code used in this work is available at `https://github.com/ASK-Berkeley/distillation-electrolytes`. The student models are trained and simulated with a fork of FeNNol [63], included in the repository as a submodule.

Acknowledgements

We would like to thank Kranthi K. Mandadapu for the guidance, help, and discussion on the theory behind transport property calculation and error quantification in this paper. We would also like to thank Bryan D. McCloskey for allowing us to use his experimental equipment to make measurements for the work.

We would also like to thank Sakib Matin, Sanjeev Raja, and Eric Qu for insightful discussions that benefited the work.

This work used the DeltaAI GPU resource at the National Center for Supercomputing Applications through allocation BIO250235 from the Advanced Cyberinfrastructure Coordination Ecosystem: Services & Support (ACCESS) program [85], which is supported by U.S. National Science Foundation grants #2138259, #2138286, #2138307, #2137603, and #2138296. This research used the DeltaAI advanced computing and data resource, which is supported by the National Science Foundation (award OAC 2320345) and the State of Illinois. DeltaAI is a joint effort of the University of Illinois Urbana-Champaign and its National Center for Supercomputing Applications.

Funding

This work was partially funded by the Office of Science, Office of Basic Energy Sciences in the Chemical Sciences, Geosciences, and Biosciences subprogram at Lawrence Berkeley National Laboratory under contract #DE-AC02-05CH11231. We also acknowledge funding from the Energy Storage Research Alliance (ESRA) (DE-AC02-06CH11357), an Energy Innovation Hub funded by the U.S. Department of Energy, Office of Science, Basic Energy Sciences.

Author contributions

I.A. and Y.J. formulated the distillation procedure and ran all model training, simulations, and transport calculations. N.K. helped supply the initial boxes for the work, as well as code for static observables. C.Z. helped train and simulate initial versions of distilled models. J.L. prepared the electrolytes and provided the salts and solvents. C.S.M. and J.D.B. conducted the density, viscosity, and self-diffusivity measurements. A.M.M. and J.R.T. performed ionic conductivity measurements. E.J.C., J.D.B., and N.P.D. supervised the experimental work. N.K., I.A., Y.J., S.B., and A.S.K. conceptualized and designed the study. A.S.K. and S.B. supervised the work. I.A., Y.J., N.K., S.B., and A.S.K. all helped write up the final manuscript.

Competing interests

The authors declare no competing interests.

[1] K. Xu, Nonaqueous liquid electrolytes for lithium-based rechargeable batteries, Chemical Reviews **104**, 4303 (2004).

[2] S. Zhu, B. Ramsundar, E. Annevelink, H. Lin, A. Dave, P.-W. Guan, K. Gering, and V. Viswanathan, Differentiable modeling and optimization of non-aqueous li-based battery electrolyte solutions using geometric deep learning, Nature Communications **15**, 8649 (2024).

[3] J. Lai, Y. Guo, H.-E. Lai, F. A. Ospina-Acevedo, W. Tian, D. Kuai, D. Chen, P. B. Balbuena, and F. Shi, Linking solvation equilibrium thermodynamics to electrolyte transport kinetics for lithium batteries, Journal of the American Chemical Society **147**, 14348 (2025).

[4] J. Hu, Y. Ji, G. Zheng, W. Huang, Y. Lin, L. Yang, and F. Pan, Influence of electrolyte structural evolution on battery applications: cationic aggregation from dilute to high concentration, Aggregate **3**, e153 (2022).

[5] E. J. Maginn, R. A. Messerly, D. J. Carlson, D. R. Roe, and J. R. Elliott, Best practices for computing transport properties 1. self-diffusivity and viscosity from equilibrium molecular dynamics [article v1.0], Living Journal of Computational Molecular Science **1**, 6324 (2019).

[6] Y. Zhang, A. Otani, and E. J. Maginn, Reliable viscosity calculation from equilibrium molecular dynamics simulations: A time decomposition method, Journal of Chemical Theory and Computation **11**, 3537 (2015).

[7] A. J. Ringsby, K. D. Fong, J. Self, H. K. Bergstrom, B. D. McCloskey, and K. A. Persson, Transport phenomena in low temperature lithium-ion battery electrolytes, Journal of The Electrochemical Society **168**, 080501 (2021).

[8] K. D. Fong, H. K. Bergstrom, B. D. McCloskey, and K. K. Mandadapu, Transport phenomena in electrolyte solutions: Non-equilibrium thermodynamics and statistical mechanics, AIChE Journal **66**, e17091 (2020).

[9] R. E. Jones and K. K. Mandadapu, Adaptive Green–Kubo estimates of transport coefficients from molecular dynamics based on robust error analysis, The Journal of Chemical Physics **136**, 154102 (2012).

[10] O. Borodin, Polarizable force field development and molecular dynamics simulations of ionic liquids, The Journal of Physical Chemistry B **113**, 11463 (2009).

[11] J. Behler and M. Parrinello, Generalized neural-network representation of high-dimensional potential-energy surfaces, Physical Review Letters **98**, 146401 (2007).

[12] C. Schran, F. L. Thiemann, P. Rowe, E. A. Müller, O. Marsalek, and A. Michaelides, Machine learning potentials for complex aqueous systems made simple, Proceedings of the National Academy of Sciences **118**, e2110077118 (2021).

[13] I. Batatia, D. P. Kovács, G. N. C. Simm, C. Ortner, and G. Csányi, MACE: Higher order equivariant message passing neural networks for fast and accurate force fields, in *Advances in Neural Information Processing Systems*, Vol. 35 (2022) pp. 11423–11436.

[14] S. Batzner, A. Musaelian, L. Sun, M. Geiger, J. P. Mailoa, M. Kornbluth, N. Molinari, T. E. Smidt, and B. Kozinsky, E(3)-equivariant graph neural networks for data-efficient and accurate interatomic potentials, Nature Communications **13**, 2453 (2022).

[15] E. Qu, B. M. Wood, A. S. Krishnapriyan, and Z. W. Ulissi, A recipe for scalable attention-based ML potentials: unlocking long-range accuracy with all-to-all node attention, in *Forty-third International Conference on Machine Learning* (2026).

[16] R. Jacobs, D. Morgan, S. Attarian, J. Meng, C. Shen, Z. Wu, C. Y. Xie, J. H. Yang, N. Artrith, B. Blaiszik, G. Ceder, K. Choudhary, G. Csanyi, E. D. Cubuk, B. Deng, R. Drautz, X. Fu, J. Godwin, V. Honavar, O. Isayev, A. Johansson, S. Martiniani, S. P. Ong, I. Poltavsky, K. J. Schmidt, S. Takamoto, A. P. Thompson, J. Westermayr, B. M. Wood, and B. Kozinsky, A practical guide to machine learning interatomic potentials: Status and future, Current Opinion in Solid State and Materials Science **35**, 101214 (2025).

[17] E. Qu and A. Krishnapriyan, The importance of being scalable: Improving the speed and accuracy of neural network interatomic potentials across chemical domains, Advances in Neural Information Processing Systems **37**, 139030 (2024).

[18] J. Behler, M. Ceriotti, C. Clementi, G. Csányi, A.-M. Elena, A. Krishnapriyan, J. W. Abbott, F. Affinito, A. P. Bartók, I. Batatia, *et al.*, A strategic roadmap for an atomistic machine-learning ecosystem, arXiv preprint arXiv:2609.39090 (2026).

[19] S. Gong, Y. Zhang, Z. Mu, Z. Pu, H. Wang, X. Han, Z. Yu, M. Chen, T. Zheng, Z. Wang, L. Chen, Z. Yang, X. Wu, S. Shi, W. Gao, W. Yan, and L. Xiang, A predictive machine learning force-field framework for liquid electrolyte development, Nature Machine Intelligence **7**, 543 (2025).

[20] J. Chen, Q. Gao, Y. Lin, M. Huang, Z. Cheng, W. Feng, J. Huang, B. Wang, and K. Yu, A hybrid physics-driven neural network force field for liquid electrolytes, Journal of Chemical Theory and Computation **22**, 3011 (2026), phyNEO-Electrolyte.

[21] S. Ju, J. You, G. Kim, Y. Park, H. An, and S. Han, Application of pretrained universal machine-learning interatomic potential for physicochemical simulation of liquid electrolytes in li-ion batteries, Digital Discovery **4**, 1544 (2025).

[22] T. Zheng, X. Xu, Z. Wang, Z. Yang, Y. Wang, X. Han, L. Chen, Z. Mu, Z. Zhang, S. Liu, S. Gong, K. Yu, and W. Yan, Bridging quantum mechanics to liquid properties via a universal organic force field, Nature Communications **17**, 8276 (2026).

[23] F. Wang, Y.-H. Tang, Z.-B. Ma, Y.-C. Jin, and J. Cheng, Domain oriented universal machine learning potential enables fast exploration of chemical space of battery electrolytes, Nature Communications **17**, 1226 (2026).

[24] I. Batatia, P. Benner, Y. Chiang, A. M. Elena, D. P. Kovács, J. Riebesell, X. R. Advincula, M. Asta, M. Avaylon, W. J. Baldwin, *et al.*, A foundation model for atomistic materials chemistry, The Journal of Chemical Physics **163**, 184110 (2025).

[25] M. Neumann, J. Gin, B. Rhodes, S. Bennett, Z. Li, H. Choubisa, A. Hussey, and J. Godwin, Orb: A fast, scalable neural network potential (2024), arXiv:2410.22570 [cond-mat.mtrl-sci].

[26] I. Batatia, W. J. Baldwin, D. Kuryla, J. Hart, E. Kasoar, A. M. Elena, H. Moore, M. J. Gawkowski, B. X. Shi, V. Kapil, P. Kourtis, I.-B. Magdău, and G. Csányi, MACE-POLAR-1: A polarisable electrostatic foundation

model for molecular chemistry (2026), arXiv:2602.19411 [physics.chem-ph].

[27] D. S. Levine, N. Liesen, L. Chua, J. Diffenderfer, H. Ingolfsson, M. P. Kroonblawd, N. Kumar, A. Maiti, S. S. Mohottalalage, M. Shuaibi, B. Van Essen, B. M. Wood, C. L. Zitnick, S. M. Blau, and E. R. Antoniuk, The open polymers 2026 (OPoly26) dataset and evaluations (2025), arXiv:2512.23117 [physics.chem-ph].

[28] E. C.-Y. Yuan, Y. Liu, J. Chen, P. Zhong, S. Raja, T. Kreiman, S. Vargas, W. Xu, M. Head-Gordon, C. Yang, S. M. Blau, B. Cheng, A. Krishnapriyan, and T. Head-Gordon, Foundation models for atomistic simulation of chemistry and materials, Nature Reviews Chemistry **10**, 212 (2026).

[29] D. S. Levine, M. Shuaibi, E. W. C. Spotte-Smith, M. G. Taylor, M. R. Hasyim, K. Michel, I. Batatia, G. Csányi, M. Dzamba, P. Eastman, N. C. Frey, X. Fu, V. Gharakhanyan, A. S. Krishnapriyan, N. Kumar, J. A. Rackers, S. Raja, A. Rizvi, A. S. Rosen, Z. Ulissi, S. Vargas, C. L. Zitnick, S. M. Blau, and B. M. Wood, The open molecules 2025 (OMol25) dataset, evaluations, and models (2025), arXiv:2505.08762 [physics.chem-ph].

[30] B. M. Wood, M. Dzamba, X. Fu, M. Gao, M. Shuaibi, L. Barroso-Luque, K. Abdelmaqsoud, V. Gharakhanyan, J. R. Kitchin, D. S. Levine, K. Michel, A. Sriram, T. Cohen, A. Das, A. Rizvi, S. J. Sahoo, Z. W. Ulissi, and C. L. Zitnick, UMA: A family of universal models for atoms, in *Advances in Neural Information Processing Systems*, Vol. 38 (2025) pp. 143528–143564.

[31] N. Kumar, J. Lai, C. S. Mezerkor, J. Wang, K. M. Wiaderek, J. D. Bazak, S. M. Blau, and E. J. Crumlin, Prediction and experimental verification of electrolyte solvation structure from an OMol25-trained interatomic potential, The Journal of Physical Chemistry Letters **17**, 9073 (2026).

[32] G. Hinton, O. Vinyals, and J. Dean, Distilling the knowledge in a neural network, in *NIPS Deep Learning and Representation Learning Workshop* (2015) arXiv:1503.02531 [stat.ML].

[33] I. Amin, S. Raja, and A. S. Krishnapriyan, Towards fast, specialized machine learning force fields: Distilling foundation models via energy hessians, in *The Thirteenth International Conference on Learning Representations* (2025).

[34] W. Li, N. Charoenphakdee, Y.-B. Zhuang, R. Okuno, Y. Tsuboi, S. Takamoto, J. Ishida, and J. Li, Lightpfp: A lightweight route to ab initio accuracy at scale, arXiv preprint arXiv:2510.23064 (2025), arXiv:2510.23064 [cond-mat.mtrl-sci].

[35] S. Oh, J. You, J. Kim, J. Lee, H. An, S. Han, and Y. Kang, A lightweight universal machine-learning interatomic potential via knowledge distillation for scalable atomistic simulations, Journal of Chemical Information and Modeling **66**, 8875 (2026).

[36] X. Wang and B. Cheng, Distilling latent electrostatics from foundation machine learning interatomic potentials (2026), arXiv:2606.15001 [physics.comp-ph].

[37] M. Kim, S. Kim, J.-Y. Jung, M. Y. Ha, and W. B. Lee, Density diversity in training data governs thermodynamic transferability of machine learning interatomic potentials, arXiv preprint arXiv:2605.05733 (2026).

[38] J. L. A. Gardner, D. F. T. du Toit, C. B. Mahmoud, Z. F. Beaulieu, V. Juraskova, L.-B. Paşca, L. A. M. Rosset, F. Duarte, F. Martelli, C. J. Pickard, and V. L. Deringer, Distillation of atomistic foundation models across architectures and chemical domains (2025), arXiv:2506.10956 [physics.comp-ph].

[39] S. Matin, A. E. A. Allen, E. Shinkle, A. Pachalieva, G. T. Craven, B. Nebgen, J. S. Smith, R. Messerly, Y. W. Li, S. Tretiak, K. Barros, and N. Lubbers, Teacher-student training improves accuracy and efficiency of machine learning interatomic potentials (2025), arXiv:2502.05379 [physics.chem-ph].

[40] D. Zhang, X. Liu, X. Zhang, C. Zhang, C. Cai, H. Bi, Y. Du, X. Qin, A. Peng, J. Huang, B. Li, Y. Shan, J. Zeng, Y. Zhang, S. Liu, Y. Li, J. Chang, X. Wang, S. Zhou, J. Liu, X. Luo, Z. Wang, W. Jiang, J. Wu, Y. Yang, J. Yang, M. Yang, F.-Q. Gong, L. Zhang, M. Shi, F.-Z. Dai, D. M. York, S. Liu, T. Zhu, Z. Zhong, J. Lv, J. Cheng, W. Jia, M. Chen, G. Ke, W. E, L. Zhang, and H. Wang, DPA-2: a large atomic model as a multi-task learner, npj Computational Materials **10**, 293 (2024).

[41] R. Wang, Y. Gao, H. Wu, and Z. Zhong, Pre-training, fine-tuning, and distillation (PFD): Automatically generating machine learning force fields from universal models, Physical Review Materials **9**, 113802 (2025).

[42] W. Zhang, X. Wu, C. Wang, S. Hu, Y. Liu, and L.-W. Wang, Constructing machine learning interatomic potentials with minimum amount of ab initio data, npj Computational Materials **12**, 174 (2026).

[43] N. Matsumura, Y. Yoshimoto, Y. Iwasaki, M. Yamazaki, and Y. Sakai, Knowledge distillation framework for accelerating high-accuracy neural network-based molecular dynamics simulations, in *Machine Learning and Principles and Practice of Knowledge Discovery in Databases (ECML PKDD 2025)*, Communications in Computer and Information Science (Springer, 2026) pp. 41–56.

[44] S. Chen, N. O'Neill, B. X. Shi, and V. Kapil, Distilling first-principles accuracy into compact machine learning potentials for condensed-phase chemistry (2026), arXiv:2606.06848 [physics.chem-ph].

[45] X. Lian and A. Pasquarello, Data-efficient and fast machine learning molecular dynamics through integrated active learning and knowledge distillation, Journal of Chemical Theory and Computation **22**, 8730 (2026).

[46] X. Fu, Z. Wu, W. Wang, T. Xie, S. Keten, R. Gomez-Bombarelli, and T. Jaakkola, Forces are not enough: Benchmark and critical evaluation for machine learning force fields with molecular simulations, Transactions on Machine Learning Research (2023), arXiv:2210.07237; OpenReview `https://openreview.net/forum?id=A8pqQipwkt`.

[47] S. Stocker, J. Gasteiger, F. Becker, S. Günnemann, and J. T. Margraf, How robust are modern graph neural network potentials in long and hot molecular dynamics simulations?, Machine Learning: Science and Technology **3**, 045010 (2022).

[48] S. Raja, I. Amin, F. Pedregosa, and A. S. Krishnapriyan, Stability-aware training of machine learning force fields with differentiable boltzmann estimators, Transactions on Machine Learning Research (2025).

[49] S. Miret, K. L. K. Lee, C. Gonzales, S. Mannan, and N. M. A. Krishnan, Energy & force regression on dft trajectories is not enough for universal machine learning interatomic potentials (2025), arXiv:2502.03660 [cond-mat.mtrl-sci].

[50] J.-Y. Hwang, S.-T. Myung, and Y.-K. Sun, Sodium-ion batteries: Present and future, Chemical Society Reviews

46, 3529 (2017).
[51] N. Yabuuchi, K. Kubota, M. Dahbi, and S. Komaba, Research development on sodium-ion batteries, Chemical Reviews **114**, 11636 (2014).
[52] K. T. Schütt, O. T. Unke, and M. Gastegger, Equivariant message passing for the prediction of tensorial properties and molecular spectra, in *Proceedings of the 38th International Conference on Machine Learning (ICML)*, PMLR, Vol. 139 (2021) pp. 9377–9388.
[53] W. L. Jorgensen, D. S. Maxwell, and J. Tirado-Rives, Development and testing of the opls all-atom force field on conformational energetics and properties of organic liquids, Journal of the American Chemical Society **118**, 11225 (1996).
[54] I.-C. Yeh and G. Hummer, System-size dependence of diffusion coefficients and viscosities from molecular dynamics simulations with periodic boundary conditions, The Journal of Physical Chemistry B **108**, 15873 (2004).
[55] A. K. Doolittle, Studies in Newtonian flow. II. the dependence of the viscosity of liquids on free-space, Journal of Applied Physics **22**, 1471 (1951).
[56] M. H. Cohen and D. Turnbull, Molecular transport in liquids and glasses, The Journal of Chemical Physics **31**, 1164 (1959).
[57] L. Martínez, R. Andrade, E. G. Birgin, and J. M. Martínez, PACKMOL: A package for building initial configurations for molecular dynamics simulations, Journal of Computational Chemistry **30**, 2157 (2009).
[58] A. H. Larsen *et al.*, The atomic simulation environment—a python library for working with atoms, Journal of Physics: Condensed Matter **29**, 273002 (2017).
[59] S. Nosé, A unified formulation of the constant temperature molecular dynamics methods, The Journal of Chemical Physics **81**, 511 (1984).
[60] W. G. Hoover, Canonical dynamics: Equilibrium phase-space distributions, Physical Review A **31**, 1695 (1985).
[61] M. Parrinello and A. Rahman, Polymorphic transitions in single crystals: A new molecular dynamics method, Journal of Applied Physics **52**, 7182 (1981).
[62] E. C. Y. Yuan and T. Head-Gordon, Teachers that teach the irrelevant: Pre-training machine learned interaction potentials with classical force fields for robust molecular dynamics simulations (2026), arXiv:2509.14205 [physics.chem-ph].
[63] T. Plé, O. Adjoua, L. Lagardère, and J.-P. Piquemal, FeNNol: An efficient and flexible library for building force-field-enhanced neural network potentials, The Journal of Chemical Physics **161**, 042502 (2024).
[64] J. F. Ziegler, J. P. Biersack, and U. Littmark, *The Stopping and Range of Ions in Matter* (Pergamon, New York, 1985).
[65] I. Loshchilov and F. Hutter, Decoupled weight decay regularization, in *International Conference on Learning Representations (ICLR)* (2019).
[66] J. Bradbury, R. Frostig, P. Hawkins, M. J. Johnson, C. Leary, D. Maclaurin, G. Necula, A. Paszke, J. VanderPlas, S. Wanderman-Milne, and Q. Zhang, JAX: Composable transformations of Python+NumPy programs, `http://github.com/google/jax` (2018).
[67] G. Bussi, D. Donadio, and M. Parrinello, Canonical sampling through velocity rescaling, The Journal of Chemical Physics **126**, 014101 (2007).
[68] J. E. Basconi and M. R. Shirts, Effects of temperature control algorithms on transport properties and kinetics in molecular dynamics simulations, Journal of Chemical Theory and Computation **9**, 2887 (2013).
[69] M. J. Abraham, T. Murtola, R. Schulz, S. Páll, J. C. Smith, B. Hess, and E. Lindahl, Gromacs: High performance molecular simulations through multi-level parallelism from laptops to supercomputers, SoftwareX **1–2**, 19 (2015).
[70] L. S. Dodda, I. Cabeza de Vaca, J. Tirado-Rives, and W. L. Jorgensen, Ligpargen web server: an automatic opls-aa parameter generator for organic ligands, Nucleic Acids Research **45**, W331 (2017).
[71] J. N. Canongia Lopes, J. Deschamps, and A. A. H. Pádua, Modeling ionic liquids using a systematic all-atom force field, The Journal of Physical Chemistry B **108**, 2038 (2004).
[72] G. A. Kaminski and W. L. Jorgensen, Host–guest chemistry of rotaxanes and catenanes: application of a polarizable all-atom force field to cyclobis(paraquat-p-phenylene) complexes with disubstituted benzenes and biphenyls, Journal of the Chemical Society, Perkin Transactions 2 , 2365 (1999).
[73] J. N. Canongia Lopes and A. A. H. Pádua, Molecular force field for ionic liquids composed of triflate or bistriflylimide anions, The Journal of Physical Chemistry B **108**, 16893 (2004).
[74] A. S. L. Gouveia, C. E. S. Bernardes, L. C. Tomé, E. I. Lozinskaya, Y. S. Vygodskii, A. S. Shaplov, J. N. Canongia Lopes, and I. M. Marrucho, Ionic liquids with anions based on fluorosulfonyl derivatives: from asymmetrical substitutions to a consistent force field model, Physical Chemistry Chemical Physics **19**, 29617 (2017).
[75] J. Åqvist, Ion-water interaction potentials derived from free energy perturbation simulations, The Journal of Physical Chemistry **94**, 8021 (1990).
[76] M. Bernetti and G. Bussi, Pressure control using stochastic cell rescaling, The Journal of Chemical Physics **153**, 114107 (2020).
[77] U. Essmann, L. Perera, M. L. Berkowitz, T. Darden, H. Lee, and L. G. Pedersen, A smooth particle mesh ewald method, The Journal of Chemical Physics **103**, 8577 (1995).
[78] A. Einstein, Über die von der molekularkinetischen Theorie der Wärme geforderte Bewegung von in ruhenden Flüssigkeiten suspendierten Teilchen, Annalen der Physik **322**, 549 (1905).
[79] J.-P. Hansen and I. R. McDonald, *Theory of Simple Liquids: with Applications to Soft Matter*, 4th ed. (Academic Press, Oxford, 2013).
[80] K. D. Fong, J. Self, B. D. McCloskey, and K. A. Persson, Onsager transport coefficients and transference numbers in polyelectrolyte solutions and polymerized ionic liquids, Macromolecules **53**, 9503 (2020).
[81] M. S. Green, Markoff random processes and the statistical mechanics of time-dependent phenomena. II. irreversible processes in fluids, The Journal of Chemical Physics **22**, 398 (1954).
[82] R. Kubo, Statistical-mechanical theory of irreversible processes. I. general theory and simple applications to magnetic and conduction problems, Journal of the Physical Society of Japan **12**, 570 (1957).
[83] D. J Evans and G. P Morriss, *Statistical mechanics of nonequilbrium liquids* (ANU Press, 2007).
[84] P. J. Daivis and D. J. Evans, Comparison of constant pressure and constant volume nonequilibrium simulations

of sheared model decane, The Journal of Chemical Physics **100**, 541 (1994).

[85] T. J. Boerner, S. Deems, T. R. Furlani, S. L. Knuth, and J. Towns, ACCESS: Advancing innovation: NSF's advanced cyberinfrastructure coordination ecosystem: Services & support, in *Practice and Experience in Advanced Research Computing (PEARC '23)* (ACM, Portland, OR, USA, 2023).

Supporting Information

Distilling Large-Scale Machine Learning Interatomic Potentials to Predict Long-Timescale Electrolyte Transport Properties

Ishan Amin[1,*], Yue Jian[2,*], Nitesh Kumar[3], Christine Zhang[1], Jianwei Lai[3], Casey S. Mezerkor[4], Aaron M. Melemed[5], Jonathan R. Thurston[5], Neil P. Dasgupta[5], J. David Bazak[4], Ethan J. Crumlin[3], Samuel Blau[3,†], Aditi S. Krishnapriyan[1,2,‡]

[1]Department of Computer Science, University of California, Berkeley, Berkeley, California 94720, USA
[2]Department of Chemical and Biomolecular Engineering, University of California, Berkeley, Berkeley, California 94720, USA
[3]Lawrence Berkeley National Laboratory, Berkeley, California 94720, USA
[4]Physical and Computational Sciences Directorate, Pacific Northwest National Laboratory, Richland, Washington 99354, USA
[5]Department of Mechanical Engineering, University of Michigan, Ann Arbor, Michigan 48109, USA

*These authors contributed equally to this work.
†smblau@lbl.gov ‡aditik1@berkeley.edu

Contents

S1 Initial box setup and system compositions **2**
S1.1 Initial packing 2
S1.2 Correcting the molarity 2
S1.3 Selecting the NVT seed boxes 4
S1.4 System compositions 4

S2 Student architecture and training **5**

S3 Results of noising training configurations **7**

S4 Teacher vs. student speed benchmark **7**

S5 Simulation settings and production stability metrics **8**
S5.1 Production stability and teacher fidelity 11

S6 Radial distribution functions for 1 M systems **12**

S7 Per-system mass densities **14**

S8 Details of experimental measurements **14**
S8.1 Electrolyte preparation 14
S8.2 Self-diffusivity 15
S8.3 Ionic conductivity 15
S8.4 Density and viscosity 16

S9 Experimental and simulated transport data **16**

S10 Diffusivity calculation details **19**
S10.1 Example mean-squared displacement curves 19

S11 Viscosity calculation details **19**
S11.1 Example Green–Kubo viscosity integrals 19
S11.2 Expected viscosity overshoot from the density bias 20

S12 Conductivity **23**
S12.1 Justifying the choice of conductivity calculation 23
S12.2 Derivation of the Onsager conductivity 23
S12.3 The conductivity outlier is a low-ionicity, correlation-dominated system 26
S12.4 Origin of the teacher's conductivity error from rigid-scan interaction profiles . . . 27

S13 Convergence analysis **28**
S13.1 Expanding-window convergence 28
S13.2 Error analysis 28

S14 The smaller multi-layer-perceptron student **31**
S14.1 MLP student architecture and training 31
S14.2 Transport results 32

S1 Initial box setup and system compositions

S1.1 Initial packing

To create each system, we place the number of ion pairs corresponding to the target molarity in a $30 \times 30 \times 30$ Å box, and use PACKMOL to fill the rest of the space with solvent. The boxes are then equilibrated with OPLS-AA in the NPT ensemble at 298 K for 1 ns. The resulting configurations are then used as the starting point for the UMA-OMol (checkpoint `uma-s-1p1` with the OMol task head) teacher runs of Table S5, whose first 100 ps supply the configurations used to train the student. Once the student is trained, we use the same starting configurations for the 1 ns NPT student trajectories used to compare against the teacher in the main text.

S1.2 Correcting the molarity

It is important to ensure that the production boxes used to compute the observables compared against experiment are close to the actual target molarity. Placing the correct number of ion pairs in a 30 Å cube will only yield the correct molarity if the volume is held fixed. Once the barostat relaxes the cell to the force field's equilibrium density, the volume will change and the molarity will drift from the target. We empirically found that during equilibration, the volume would decrease from its original $30 \times 30 \times 30$ Å size, causing the molarity to overshoot the target. Table S1 shows the concentrations of the equilibrated boxes before the correction described below: the initial nominal 1 M boxes equilibrate at 1.19–1.32 M and the nominal 0.1 M boxes at 0.15–0.17 M.

To create boxes that equilibrate to the correct target molarity, we fix the solute-to-solvent ratio instead of the ratio between solute and volume. For each system, we take the student's equilibrium

density ρ from its 298 K NPT run on the initial box (Section S5), and use it to solve a mass balance equation for the number of solvent molecules.

Let

$$\begin{aligned} N_{\mathrm{su}} &= \text{number of ion pairs (solute)} \\ N_{\mathrm{sv}} &= \text{number of solvent molecules} \\ m_{\mathrm{su}},\ m_{\mathrm{sv}} &= \text{mass of one ion pair, one solvent molecule (amu)} \\ M &= \text{target molarity} \\ \rho &= \text{target (equilibrium) density (amu/Å}^3) \\ V &= \text{box volume} \end{aligned} \tag{S1}$$

A density in $\mathrm{g\,cm^{-3}}$ is converted to amu/Å^3 by multiplying by $N_A \times 10^{-24} = 0.6022$. Then, we aim to find the number of solvent molecules N_{sv} that satisfies

$$N_{\mathrm{su}}\, m_{\mathrm{su}} + N_{\mathrm{sv}}\, m_{\mathrm{sv}} = \rho V. \tag{S2}$$

The volume of our system needs to correspond to the correct target molarity, so that the molarity satisfies

$$M = \frac{N_{\mathrm{su}} \text{ particles}}{V \text{ [L]}} \times \frac{1 \text{ mol}}{N_A \text{ particles}}. \tag{S3}$$

Solving for V and converting liters to Å^3 (using 1 L $= 10^{27}$ Å^3) gives

$$V\ [\text{Å}^3] = \frac{N_{\mathrm{su}}}{N_A\, M} \times \frac{10^{27}\ \text{Å}^3}{1\ \text{L}} = \frac{N_{\mathrm{su}}}{\mathcal{N}\, M}, \qquad \mathcal{N} \equiv \frac{N_A}{10^{27}} = 6.022 \times 10^{-4}\ \frac{\text{particles/Å}^3}{\text{mol/L}}, \tag{S4}$$

where the constant $\mathcal{N}$ collects Avogadro's number and the liter-to-Å^3 conversion, so a solution at molarity M holds $\mathcal{N}M$ particles per Å^3. Substituting the expression for the volume into the mass balance equation, and then solving for the solvent count, we get

$$N_{\mathrm{sv}} = \mathrm{round}\left[\frac{N_{\mathrm{su}}}{m_{\mathrm{sv}}}\left(\frac{\rho}{\mathcal{N}\, M} - m_{\mathrm{su}}\right)\right]. \tag{S5}$$

After obtaining the solvent count for each system, we create boxes with the original solute count and the newly specified solvent count using PACKMOL. We set the initial box to a density of 80% of ρ, which is loose enough that packing does not have to resolve overlaps. The packed box is then relaxed with UMA-OMol to remove the close contacts left by packing, using a FIRE optimizer at fixed cell until the maximum force falls below 0.5 eV/Å or 100 steps have elapsed. Finally, the relaxed box is run for a 1 ns NPT simulation (298 K, 1 bar) with the student model, to equilibrate the density. A single frame from this run is then selected to seed the NVT production runs (see Section S1.3 for details), from which transport observables are derived. We found that the above procedure results in boxes that closely match the target molarity: the corrected boxes equilibrate at 0.96–0.99 M and 0.099–0.100 M against nominal 1 M and 0.1 M targets. Figure S1 shows this shift for every system, from the original boxes of Section S1.1 to the ratio-corrected ones.

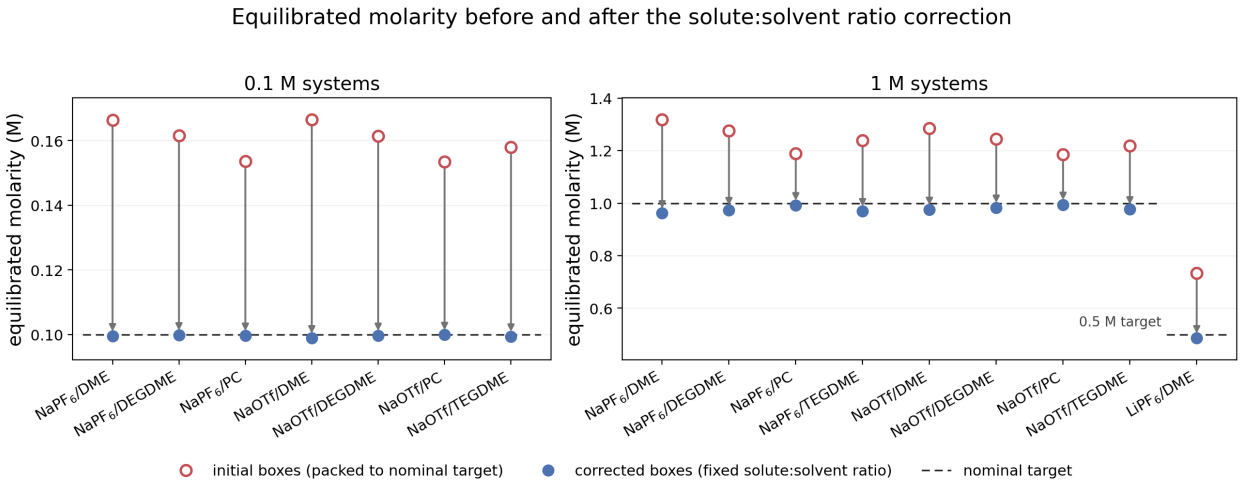

Figure S1: Equilibrated molarity before and after the solute:solvent ratio correction. Each arrow runs from the molarity the initial box equilibrates to (open circle) to that of the corrected box (filled circle), against the nominal target (dashed). The initial boxes are always over-concentrated, by 50–66% at 0.1 M and 19–32% at 1 M, because equilibration reduces the original volume. $LiPF_6$/DME is grouped with the concentrated systems but targets 0.5 M, so it carries its own reference segment.

S1.3 Selecting the NVT seed boxes

To choose a frame from the NPT simulations to start our fixed-volume NVT simulations, we first average the mass density over the last 500 ps of the NPT run, yielding a target density. We then select the frame from the last 300 ps of the run that is closest to our target density, and use that frame for the following NVT simulations used to compute transport coefficients.

S1.4 System compositions

Table S1 lists every simulated electrolyte cell. The "Ion pairs" column counts salt formula units (equal numbers of cations and anions), which the correction leaves unchanged; the remaining columns are given as *initial / production* where the two differ.

Table S1: Composition of every simulated electrolyte cell, as *initial / production* where the molarity correction changed the value. Ion-pair counts are not affected by the correction.

Salt	Solvent	Nominal (M)	Ion pairs	Solvent mol.	Atoms	Equilibrated (M)
$NaPF_6$	DME	0.1	2	125 / 210	2016 / 3376	0.166 / 0.100
$NaPF_6$	DME	1.0	17	125 / 175	2136 / 2936	1.318 / 0.962
$NaPF_6$	PC	0.1	2	164 / 254	2148 / 3318	0.154 / 0.100
$NaPF_6$	PC	1.0	17	164 / 200	2268 / 2736	1.189 / 0.992
$NaPF_6$	DEGDME	0.1	2	94 / 153	2178 / 3535	0.162 / 0.100
$NaPF_6$	DEGDME	1.0	17	94 / 126	2298 / 3034	1.275 / 0.974
$NaPF_6$	TEGDME	1.0	17	62 / 80	2430 / 3096	1.240 / 0.969
NaOTf	DME	0.1	2	125 / 211	2018 / 3394	0.166 / 0.099
NaOTf	DME	1.0	17	125 / 170	2153 / 2873	1.285 / 0.975
NaOTf	PC	0.1	2	164 / 253	2150 / 3307	0.153 / 0.100
NaOTf	PC	1.0	17	164 / 200	2285 / 2753	1.185 / 0.994
NaOTf	DEGDME	0.1	2	94 / 153	2180 / 3537	0.161 / 0.100
NaOTf	DEGDME	1.0	17	94 / 122	2315 / 2959	1.245 / 0.983
NaOTf	TEGDME	0.1	2	62 / 99	2312 / 3681	0.158 / 0.099
NaOTf	TEGDME	1.0	17	62 / 79	2447 / 3076	1.220 / 0.977
$LiPF_6$	DME	0.5	9	125 / 190	2072 / 3112	0.733 / 0.486

S2 Student architecture and training

Table S2 reports the full PaiNN architecture and optimization settings used for every distilled student. We report per-system training set sizes, which depend on the number of concentrations for each system, in Table S3.

We train a distilled student for every distinct salt/solvent pair, where datapoints from the same pairing but different concentrations (Table S1) are combined into one dataset to train a single model.

We obtain training configurations from a 100 ps UMA-OMol trajectory, using two interleaved 50 fs grids: an unperturbed MD grid and a grid offset by 20 fs from the first. The positions of the offset grid's configurations are rattled with isotropic Gaussian noise (standard deviation 0.1 Å), and relabeled with the teacher. Both the noised and unperturbed configurations are split 90/10 into train/val, then pooled together. This results in 1800 MD + 1800 noised = 3600 training and 200 MD + 200 noised = 400 validation configurations per concentration. Table S3 reports the resulting per-system dataset totals.

Table S2: Distilled-potential architecture and training settings.

Setting	Value
Architecture	
Backbone	PaiNN (FeNNol implementation)
Feature dimension (`dim`)	128
Message-passing layers	3
Channels (`nchannels`)	128
Message / update MLP hidden	[128]
Nonlinearity	SiLU
Radial cutoff	5.1 Å
Radial basis	128 trainable Gaussians
Cutoff switching	trainable polynomial ($p = 3$)
Species encoding	random, dim 32 ($Z_{\max} = 86$)
Readout head	eSEN energy head (hidden 128)
Short-range repulsion	trainable ZBL, inner cutoff 1.1 Å
Loss	
Energy loss	per-atom MAE (weight 1)
Force loss	Euclidean-norm (L_2) error (weight 1)
Stress loss	MAE (weight 1)
Loss normalization	dataset force RMS (all three terms)
Optimization	
Optimizer	AdamW
Weight decay	10^{-3} (kernel weights)
Gradient clipping	100
LR schedule	one-cycle cosine, 1% warmup
Learning rate (init/peak/final)	1.6×10^{-4} / 8×10^{-4} / 8×10^{-6}
Batch size	4
Epochs	275
Validation interval	every 5 epochs
EMA	none

Table S3: Per-system training dataset statistics, as well as the cost of training on each dataset. Each concentration contributes 1800 MD + 1800 noised = 3600 training and 200 MD + 200 noised = 400 validation configurations. The two-concentration systems have twice as many datapoints, and therefore cost roughly twice the wall time to train on. Training times are for 275 epochs on a single NVIDIA GH200.

Salt	Solvent	Conc. (M)	Train configs	Val configs	Training time (h)
$NaPF_6$	DME	0.1, 1.0	7200	800	17.4
$NaPF_6$	PC	0.1, 1.0	7200	800	17.7
$NaPF_6$	DEGDME	0.1, 1.0	7200	800	20.2
$NaPF_6$	TEGDME	1.0	3600	400	10.9
NaOTf	DME	0.1, 1.0	7200	800	17.5
NaOTf	PC	0.1, 1.0	7200	800	16.9
NaOTf	DEGDME	0.1, 1.0	7200	800	19.3
NaOTf	TEGDME	0.1, 1.0	7200	800	21.8
$LiPF_6$	DME	0.5	3600	400	8.7

S3 Results of noising training configurations

As discussed in Section S2, half of every training set consists of structures whose positions are rattled with isotropic Gaussian noise (standard deviation 0.1 Å). The noised structures give the student information about geometries that molecular dynamics at 293 K would rarely visit.

This section compares the performance of models trained with and without noise. We retrain two students, $NaPF_6$/TEGDME at 1 M and $LiPF_6$/DME at 0.5 M, on datasets built from identical time grids, with the rattling step disabled so that the second grid's samples come from pure molecular dynamics. Everything else about training, including frame count (4000 per system, split 3600/400), architecture, loss, optimizer, schedule, and epoch budget, remains the same.

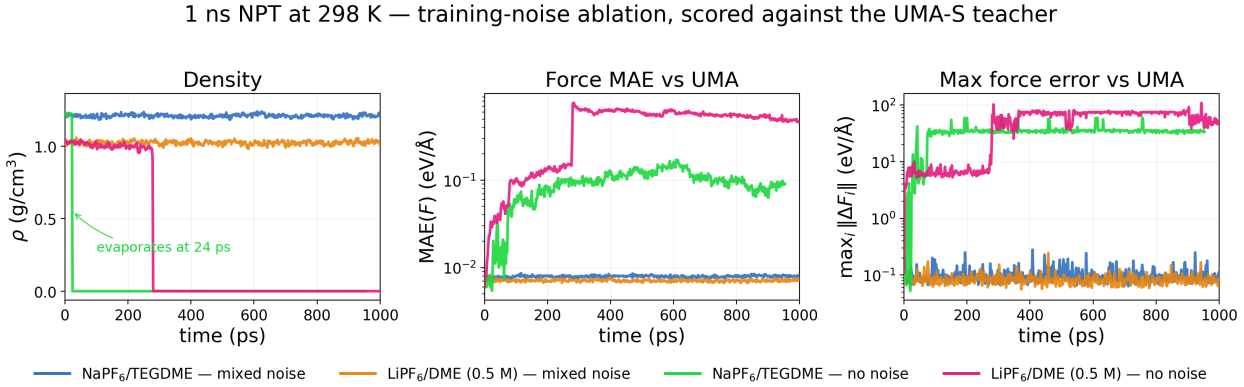

Figure S2: Training-set noising is required for stable dynamics. Stability and error statistics for 1 ns NPT simulations at 298 K. Two systems are displayed, with two models trained on each: one model is trained with noised structures, and the other without. *Left:* density. *Center and right:* mean and maximum per-atom force error against the UMA-OMol teacher, evaluated every 2 ps on the configurations in the student simulation (note the logarithmic scale on the two force panels). Both noised students have stable densities throughout, as well as low mean and maximum force errors. In contrast, the no-noise students evaporate at 24 ps ($NaPF_6$/TEGDME) and 280 ps ($LiPF_6$/DME). In both no-noise models' simulations, the error against the teacher grows large before the density collapses.

After training, both models exhibit low force errors on a static held-out validation set, at 7.0–7.4 meV/Å for the noised student and 5.6–5.8 meV/Å for the no-noise ablation. However, their behavior differs drastically during simulation. Over 1 ns of NPT at 298 K (settings as in Table S5, seeded from equilibrated teacher frames), the noised students have stable densities and low force errors. In contrast, the simulations with the no-noise students start exhibiting high force errors, and eventually their boxes expand and their simulations break apart (Figure S2).

For both systems, it is clear that noising some of the data pays off massively during simulation. The reason is intuitive: a student fit only to molecular-dynamics frames has seen only the thin manifold of thermally accessible configurations. During simulation, any deviation from that manifold produces incorrect forces that drive the system even farther away, compounding the error. Giving the model noised data allows it to learn the forces of off-manifold configurations, preventing small deviations from accumulating.

S4 Teacher vs. student speed benchmark

We benchmark the simulation throughput of all models used in this work on identical hardware: a single NVIDIA GH200 Grace–Hopper superchip with fp32 precision. Both the PaiNN and UMA-

OMol models support tf32 precision; however, we observed occasional instabilities when simulating with both models at tf32 at multi-ns timescales, so we used fp32 throughout. We also include the throughput of a lightweight multi-layer perceptron (MLP) baseline not discussed in the main text (see Section S14 for its transport results), as an example of a faster but less accurate model. Throughput was measured by running 100 ps of simulation with each model on every system, discarding the first 50 ps and reporting the steady-state rate over the remainder. Table S4 lists the throughput (ns/day) of the UMA-OMol teacher, the PaiNN student, and the MLP student across all 16 electrolyte boxes, sorted by descending PaiNN throughput. Both students run more than an order of magnitude faster than the teacher on every system: at the median, the PaiNN student reaches 9.53 ns/day, against 0.23 ns/day for UMA-OMol, a speedup of roughly 41×.

Table S4: Per-system simulation throughput (ns/day) for the UMA-OMol teacher, PaiNN student, and MLP student.

System (salt / solvent, conc., temp.)	UMA-OMol (ref., fp32)	PaiNN	MLP
NaOTf / PC, 1 M, 298 K	0.28	11.40 (40.7×)	27.84 (99.4×)
$NaPF_6$ / PC, 1 M, 298 K	0.28	11.14 (39.8×)	27.27 (97.4×)
NaOTf / DME, 1 M, 298 K	0.26	10.59 (40.7×)	26.13 (100.5×)
$LiPF_6$ / DME, 0.5 M, 298 K	0.25	10.20 (40.8×)	24.44 (97.8×)
$NaPF_6$ / DME, 1 M, 298 K	0.25	9.90 (39.6×)	25.70 (102.8×)
NaOTf / DEGDME, 1 M, 298 K	0.24	9.75 (40.6×)	23.09 (96.2×)
$NaPF_6$ / TEGDME, 1 M, 298 K	0.23	9.66 (42.0×)	23.79 (103.4×)
$NaPF_6$ / DEGDME, 1 M, 298 K	0.23	9.60 (41.7×)	24.05 (104.6×)
NaOTf / TEGDME, 1 M, 298 K	0.23	9.45 (41.1×)	23.60 (102.6×)
$NaPF_6$ / PC, 0.1 M, 298 K	0.22	9.41 (42.8×)	22.78 (103.5×)
NaOTf / PC, 0.1 M, 298 K	0.23	9.36 (40.7×)	22.09 (96.0×)
$NaPF_6$ / DME, 0.1 M, 298 K	0.22	9.31 (42.3×)	22.87 (104.0×)
NaOTf / DME, 0.1 M, 298 K	0.22	8.98 (40.8×)	22.12 (100.5×)
NaOTf / DEGDME, 0.1 M, 298 K	0.20	8.50 (42.5×)	21.10 (105.5×)
$NaPF_6$ / DEGDME, 0.1 M, 298 K	0.20	8.41 (42.1×)	20.64 (103.2×)
NaOTf / TEGDME, 0.1 M, 298 K	0.18	7.86 (43.7×)	20.33 (112.9×)
Median	**0.23**	**9.53** (41.4×)	**23.35** (101.5×)

S5 Simulation settings and production stability metrics

Table S5 reports the simulation settings for every UMA-OMol and PaiNN simulation in this work. Table S6 reports the simulation settings for the OPLS-AA baseline. The first 100 ps of the 293 K UMA-OMol simulations supply the training configurations for the student, while the rest of the trajectory up to 1 ns is used to compute the radial distribution functions that the student is compared against (Section S6). The 298 K NPT UMA-OMol simulations are used for density comparisons against the student. Likewise, the student 293 K runs are used to extract radial distribution functions, while the 298 K simulations are used to compute the density (the density runs are extended to 20 ns). Once the student's density is computed, we create new boxes with the molarity procedure described in Section S1, and use the same 298 K NPT settings for 1 ns to equilibrate with the student. We then choose a frame from this NPT equilibration closest to the average volume (see Section S1.3), and hold that equilibrated cell fixed for the NVT simulation.

We use constant volume (NVT) simulations for production, as this removes volume-related fluctuations from the stress autocorrelations that the Green–Kubo estimator for viscosity integrates.

The global Bussi thermostat used ($\gamma = 0.1$ THz, $\tau = 10$ ps) is chosen to be slow enough to leave sub-picosecond correlations unbiased.

Table S5: Molecular dynamics settings for teacher sampling, student NPT, and student NVT production. The UMA-OMol column covers both the 293 K run (training and structural reference) and the 298 K run (density reference); the student NPT column covers both the teacher comparison (293 K) and the production equilibration (298 K). Within each column the two runs are otherwise identical.

Setting	UMA-OMol teacher	Student NPT	Student NVT production
Purpose	training (first 100 ps) + structural reference (293 K); density reference (298 K)	teacher comparison (293 K); production equilibration (298 K)	transport coefficients
Engine	ASE	FeNNol/JAX	FeNNol/JAX
Ensemble	NPT	NPT	NVT
Thermostat	Nosé–Hoover (100 fs)	Langevin (γ=1 THz)	Bussi/CSVR (γ=0.1 THz)
Barostat	Parrinello–Rahman-type (p_f=0.1)	Langevin piston (iso., τ_p=200 fs, γ_p=20 THz)	— (cell fixed)
Temperature	293 K or 298 K	293 K or 298 K	298 K
Pressure	1 bar	1 bar	— (V from equilibration)
Timestep	1 fs	1 fs	1 fs
Duration	1 ns	1 ns / 20 ns (density)	20 ns
Precision	fp32	fp32	fp32
Stress log	—	—	every 2 fs

Table S6: Molecular dynamics NPT and NVT settings for the OPLS-AA baseline.

Setting	OPLS-AA NPT equilibration	OPLS-AA NVT production
Purpose	density/box equilibration	transport coefficients
Engine	GROMACS 2025.3 (mixed precision)	GROMACS 2025.3 (mixed precision)
Solvent parameters	OPLS-AA, 1.14*CM1A charges (LigParGen) [1]	
Anion parameters	CL&P: PF_6^- [2, 3], OTf^- [4, 5]	
Cation parameters	Åqvist Na^+, Li^+ [6], full +1 charge	
Combination rule / 1–4 scaling	geometric / 0.5 (OPLS-AA)	
Pair-list buffer tolerance	10^{-3} kJ mol^{-1} ps^{-1}	10^{-3} kJ mol^{-1} ps^{-1}
Ensemble	NPT	NVT
Thermostat	CSVR/Bussi (τ_T=0.1 ps)	CSVR/Bussi (τ_T=10 ps)
Barostat	C-rescale (iso., τ_p=2 ps)	— (cell fixed)
Temperature	298 K	298 K
Pressure	1 bar	— (V from equilibration)
Timestep	1 fs	1 fs
Duration	1 ns	20 ns
Constraints	none (fully flexible)	none (fully flexible)
Electrostatics	PME, 12 Å real-space cutoff	PME, 12 Å real-space cutoff
van der Waals	12 Å cutoff, long-range dispersion correction	12 Å cutoff, long-range dispersion correction
Initial velocities	fresh Maxwell–Boltzmann	fresh Maxwell–Boltzmann (independent seed per replica)
Pressure-tensor log	—	every 2 fs
Coordinate log	—	every 1 ps

S5.1 Production stability and teacher fidelity

The production runs used to calculate transport coefficients should be thermodynamically stable, and the student should have low errors relative to the teacher throughout the trajectory. To check this, for each 20 ns NVT simulation we monitor the energy, temperature, and pressure to confirm that they are stable throughout the simulation (top panels of Figures S3 and S4). We also re-evaluate 500 frames per run, spaced 40 ps apart across the full trajectory, with the UMA-OMol teacher, giving both models' energies, forces, and stresses on identical configurations (bottom panels of Figures S3 and S4).

Overall, all 64 runs are stable. Potential energy per atom, temperature, and pressure have no drift across 20 ns. Additionally, the force mean absolute error (MAE) against UMA-OMol is small, lying between 6 and 10 meV/Å for all runs, with no drift in any trajectory. The maximum per-atom force error is likewise low and stable.

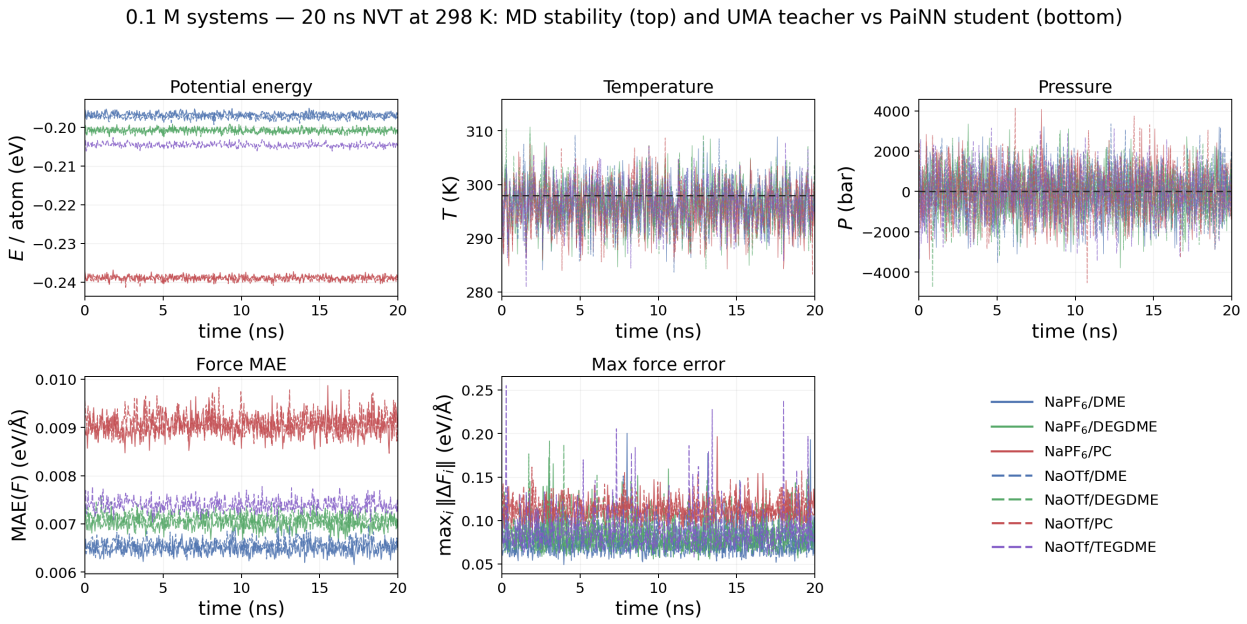

Figure S3: Production stability and teacher fidelity for the 0.1 M systems. A random replica for each system was selected to be shown here. Top row: the run's own potential energy per atom, temperature, and pressure over the full 20 ns NVT trajectory; the horizontal black reference line in each of the last two panels marks the 298 K setpoint and zero pressure, respectively. Bottom row: mean and maximum per-atom force error between the UMA-OMol teacher and the student on the same 500 frames, where frames are picked every 40 ps. All traces are flat, with no thermodynamic drift or growth in teacher–student disagreement.

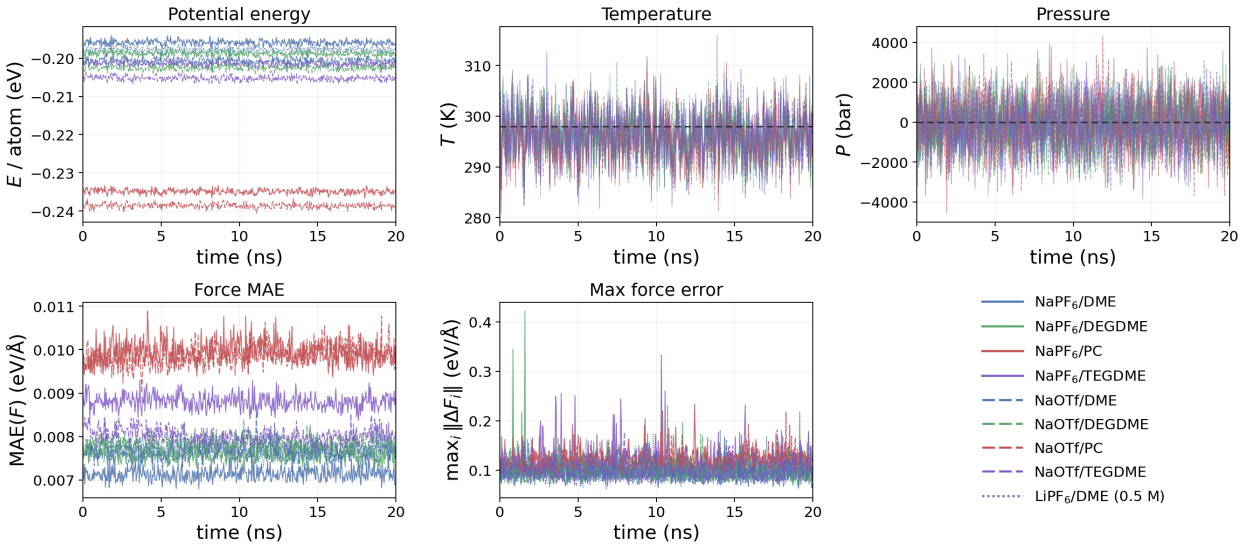

Figure S4: Production stability and teacher fidelity for the 1 M systems. Panels and sampling as in Figure S3; $LiPF_6$/DME, which exists only at 0.5 M, is grouped here.

S6 Radial distribution functions for 1 M systems

Fig. 2 of the main text shows student-vs-teacher radial distribution functions (RDFs) for four representative 1 M electrolytes. We extend this comparison to all of the 1 M salt/solvent pairs in Figure S5, and confirm that the student reproduces the RDFs of the teacher across all systems.

Each student curve is the average over five independent replicas. Replicas are differentiated by their starting velocities and the random seed of the integrator. The shaded band indicates one standard deviation across the five replicas. In nearly every panel the teacher $g(r)$ lies within this band, indicating that the teacher–student deviations are consistent with ordinary sampling fluctuations, rather than a systematic difference between the two potentials.

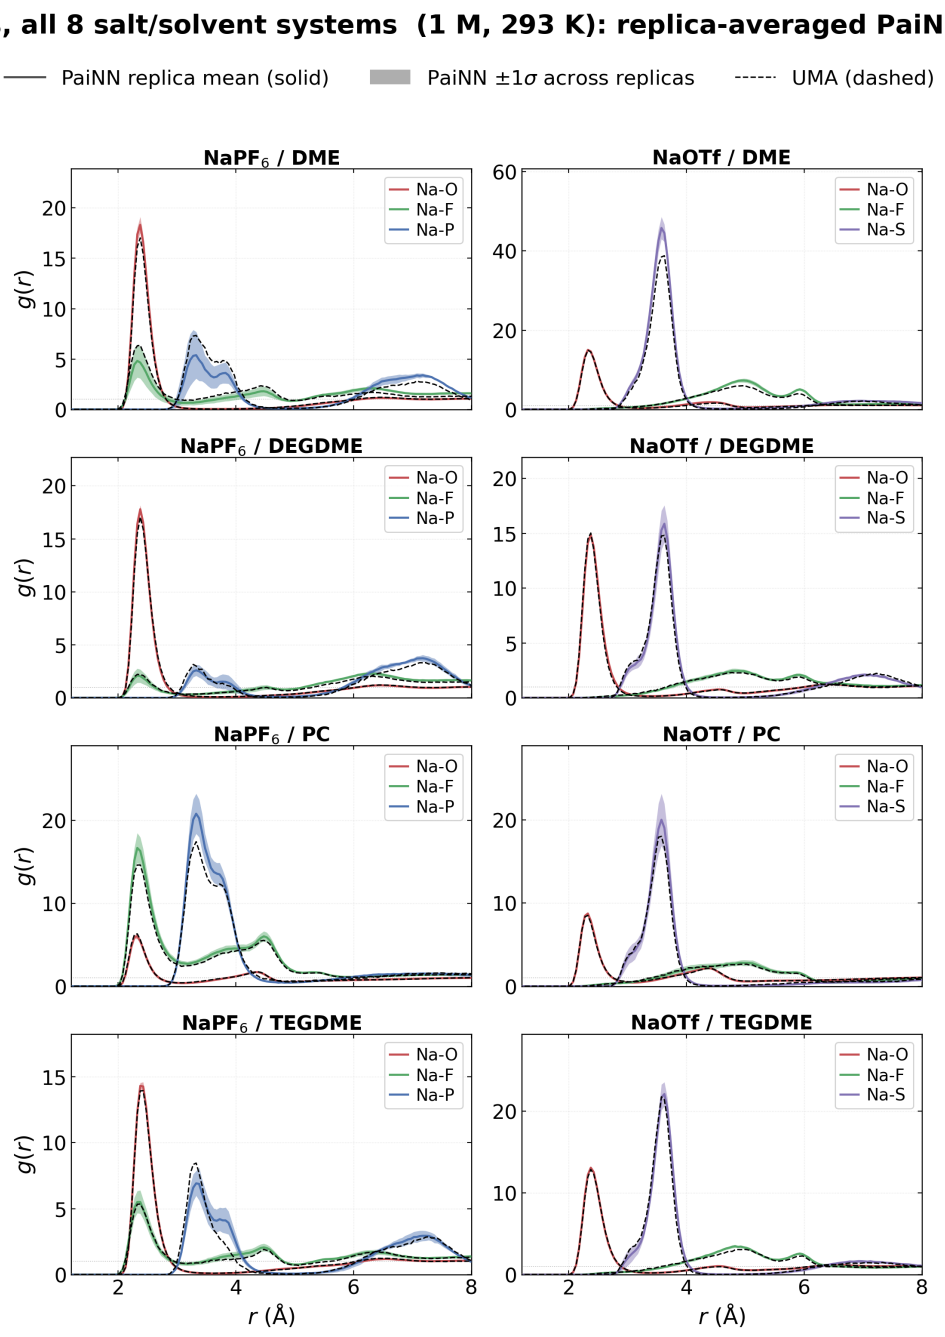

Figure S5: Radial distribution functions for all eight 1 M electrolytes. Student $g(r)$ (solid, colored by atom pair) is the mean over five 1 ns, 293 K replica trajectories, with the shaded band indicating one standard deviation across replicas; the UMA-OMol teacher (black dashed) is a single 1 ns, 293 K trajectory.

S7 Per-system mass densities

Panel (b) of Fig. 2 in the main text combines all sixteen electrolytes into a single density-parity plot, without identifying individual systems. We list the per-system density numbers in Table S7. The distilled student densities are obtained from a single 20 ns NPT trajectory (first 0.5 ns discarded), while the UMA-OMol teacher densities are obtained from a single ∼1 ns trajectory. We also include each system's experimentally measured density.

Table S7: Mass densities ($g\,cm^{-3}$, 298 K) for every simulated electrolyte: the distilled PaiNN student (single 20 ns NPT trajectory), the UMA-OMol teacher (single ∼1 ns trajectory, this work), and experiment.

Salt	Solvent	Conc. (M)	Student (20 ns)	Teacher UMA-OMol (1 ns)	Experiment
$NaPF_6$	DME	0.1	0.965	0.987	0.873
$NaPF_6$	DME	1.0	1.096	1.109	0.995
$NaPF_6$	PC	0.1	1.311	1.318	1.206
$NaPF_6$	PC	1.0	1.369	1.386	1.291
$NaPF_6$	DEGDME	0.1	1.046	1.058	0.952
$NaPF_6$	DEGDME	1.0	1.160	1.168	1.059
$NaPF_6$	TEGDME	1.0	1.212	1.225	1.115
NaOTf	DME	0.1	0.966	0.976	0.873
NaOTf	DME	1.0	1.077	1.080	0.980
NaOTf	PC	0.1	1.311	1.323	1.207
NaOTf	PC	1.0	1.373	1.379	1.277
NaOTf	DEGDME	0.1	1.045	1.057	0.950
NaOTf	DEGDME	1.0	1.139	1.154	1.047
NaOTf	TEGDME	0.1	1.116	1.129	1.016
NaOTf	TEGDME	1.0	1.199	1.213	1.106
$LiPF_6$	DME	0.5	1.029	1.048	0.909

S8 Details of experimental measurements

The self-diffusivities, ionic conductivities, densities, and viscosities used as the experimental reference throughout this work were measured in-house for the same salt/solvent/concentration combinations as the simulations. This section details the protocols followed to obtain the transport coefficients in the main text; the resulting experimental values are listed in Sec. S9.

S8.1 Electrolyte preparation

All electrolyte preparation was conducted inside an argon glovebox (O_2 content <0.2 ppm, H_2O content <0.1 ppm, VAC-ATM). A standalone solvent trap with an activated carbon filter (VAC-ATM) was utilized. 1,2-dimethoxyethane (DME; 99.5%, anhydrous, Sigma-Aldrich), bis(2-methoxyethyl) ether (DEGDME; 99+%, Thermo Scientific), tetraethylene glycol dimethyl ether (TEGDME; 99%, Thermo Scientific), and propylene carbonate (PC; 99.5%, anhydrous, Thermo Scientific) were dried over activated 4 Å molecular sieves for at least 72 hours. Water content for each solvent was determined to be <10 ppm by KF titration (Mettler Toledo C10SD). Sodium hexafluorophosphate ($NaPF_6$; 99.9%, Sigma-Aldrich), sodium trifluoromethanesulfonate (NaOTf; 99.5%, Solvionic), and lithium hexafluorophosphate ($LiPF_6$, 99.99%, Sigma-Aldrich) were each vacuum-dried at 100 °C overnight in a Buchi glass oven apparatus.

S8.2 Self-diffusivity

Species-resolved self-diffusion coefficients were obtained by pulsed-field-gradient nuclear magnetic resonance (PFG-NMR), with each species followed through a distinct nucleus so that the cation, anion, and solvent diffusivities are measured independently on the same sample: ^{1}H for the solvent, ^{7}Li or ^{23}Na for the cation, and ^{19}F for the PF_6^-/OTf^- anion. Measurements were performed on 11.7 T magnets equipped with Avance NEO consoles and Bruker iProbe BBFO auto-tuning HX solution probes, with temperature regulated by a BCU/2 unit and automatically corrected. For each sample, $\pi/2$ pulse calibrations were carried out at both 293 K and 298 K to maximize accuracy, with recycle delays set from reference samples and confirmed sufficient by the observation of sinusoidal nutation behavior. The reference samples were 1% $CHCl_3$ in acetone-d_6 (^{1}H at 8.02 ppm), 0.1 M LiCl(aq) (^{7}Li at 0 ppm), and 0.1 M NaCl(aq) (^{23}Na at 0 ppm), all at 298 K. Longitudinal relaxation times T_1 were estimated from the null point of an inversion-recovery experiment once the pulse width was determined, and used to set the recycle delays for the 1D spectra (as $5T_1$) and for the PFG-NMR measurements (as $3T_1$).

For all samples and all nuclei except ^{23}Na, convection-compensated double-stimulated echoes with bipolar pulse-pair gradients were used, together with a 5 ms longitudinal eddy-current delay. The gradient duration was 2 ms (2.4 ms for ^{19}F) and the diffusion time 100–300 ms, chosen per sample according to viscosity and relaxation parameters, with the gradient strength scaled to each sample's diffusion coefficient using two-step, first-and-last trial experiments for each nucleus. A total of 16 linearly spaced gradient steps and 16 scans per slice of signal averaging were used for all full PFG-NMR experiments. The diffusion coefficient is extracted from the gradient-dependent signal attenuation via the Stejskal–Tanner relation, $I/I_0 = \exp[-D\,\gamma^2 g^2 \delta^2 (\Delta - \delta/3)]$, where γ is the gyromagnetic ratio of the observed nucleus, g and δ the gradient-pulse amplitude and duration, and Δ the diffusion time. Because PFG-NMR tracks the same molecular centers of mass as the Einstein/MSD analysis of the trajectories, the experimental and simulated diffusivities are directly comparable.

For the most viscous ^{23}Na compositions, where relaxation and gradient-strength constraints limited the feasibility of measurements on regular solution-state probes, a 14.1 T Avance III HD system equipped with a Diff50 probe was used instead. Here both inversion-recovery and Carr–Purcell–Meiboom–Gill (CPMG) measurements were performed to determine T_1 and T_2 quantitatively for parameter setup. Owing to the higher viscosity and the correspondingly reduced propensity for convection (and rapid relaxation times), these measurements used a regular stimulated-echo PFG pulse sequence with single gradients. Both the gradient duration (0.8–1.1 ms) and the diffusion time (7–12 ms) were scaled to the sample's ^{23}Na relaxation parameters, with the gradient strength adjusted to achieve reasonable attenuation while minimizing lineshape and baseline distortion. Between 128 and 2048 transients per gradient step were acquired, depending on the sample's relaxation parameters.

S8.3 Ionic conductivity

Ionic conductivity was measured inside an argon-filled glovebox using a Mettler Toledo SevenDirect SD30 benchtop conductivity meter fitted with a Mettler Toledo InLab 710 four-electrode platinum probe with a chemically resistant glass body. The probe was immersed in 4 mL of each electrolyte and a reading was taken once it had stabilized; three independent measurements were recorded per electrolyte. For temperature-dependent measurements, the probe was inserted through a hole drilled in the vial cap, the joint was sealed with Parafilm and electrical tape to prevent solvent evaporation, and the vial was heated in an oil bath and allowed to equilibrate for 10 minutes at

each temperature before measurement; the sample temperature was recorded with each conductivity reading. The conductivities reported in this work are at 298 K.

This four-electrode conductivity-cell measurement is compared in the main text against the Onsager (collective) conductivity from simulation, which is the correct simulated counterpart because it retains the ion–ion cross-correlations the cell measurement inherently includes.

S8.4 Density and viscosity

Density and viscosity were measured with an Anton Paar DMA 4500 M oscillating U-tube density meter coupled to an Anton Paar Lovis ME rolling-ball viscometer, so that both quantities were obtained from a single filling of sample. Both instruments were calibrated with certified reference water supplied by an ISO 17034-accredited producer. The filled U-tube was inspected with the integrated camera before every measurement to confirm that no bubbles were present. Measurements were carried out at 298 K, held by the built-in Peltier thermostat. A stream of dry N_2 was directed through the measuring cell to suppress condensation on the outer surface of the U-tube. After each sample, the cell was flushed three times with absolute ethanol ($\geq$99.8%) and then purged with N_2 for 600 s. Instrument performance was checked at regular intervals during data collection using the built-in air and water density check, and the ethanol flush and N_2 purge were repeated after each check before the next sample was loaded.

The rolling-ball viscometer measures the shear viscosity η directly, which is compared in the main text.

S9 Experimental and simulated transport data

In this section, we list the raw transport data that underlie the parity plots in the main text. All entries, simulated and experimental, are at 298 K. We report simulated entries as a mean $\pm$ std over the replica trajectories, for both the PaiNN student and OPLS-AA.

In Tables S8 and S9, we report the self-diffusion coefficients of the cation (D^+), anion (D^-), and solvent (D^0), with separate tables for PaiNN and OPLS-AA. We report the raw simulated diffusivities *without* any finite-size correction, as well as the additive Yeh–Hummer correction [7] $\Delta D_{\mathrm{YH}} = \xi\, k_{\mathrm{B}}T/(6\pi\eta L)$. Two versions of the Yeh–Hummer correction are included: one uses the simulated viscosity, and the other uses the experimental viscosity. The corrected diffusivity is obtained by adding the Yeh–Hummer correction to the corresponding raw simulated diffusivity. We list the experimental self-diffusivities in each table for reference. The simulated diffusivities use the Fickian fit window described in Methods. For 7 of the 192 PaiNN fits (16 systems, four replicas, three species) and 14 of the 192 OPLS-AA fits, no window satisfied $|\nu - 1| \leq 0.15$, and the fallback twelve-point fit was used; every mean includes all four replicas.

Table S10 lists the ionic conductivity σ and the shear viscosity η, for PaiNN, OPLS-AA, and experiment. The conductivity is the Onsager conductivity, which accounts for ion–ion cross-correlations, and the viscosity is obtained from the Green–Kubo integral of the stress autocorrelation function.

Table S8: Raw (uncorrected) self-diffusivities and additive Yeh–Hummer finite-size corrections ΔD_{YH} for PaiNN, with experimental self-diffusivities for reference. Simulated entries are mean ± std over replicas; ΔD_{YH} is species-independent ("sim"/"exp" denote the viscosity used to evaluate it, and the finite-size corrected diffusivity is calculated by adding it to the raw value). All values in 10^{-10} m^2/s; "–" denotes values not available.

			D_{raw} (sim.)			ΔD_{YH}		Experimental		
Salt	Solvent	Conc.	D^+	D^-	D^0	sim	exp	D^+	D^-	D^0
$LiPF_6$	DME	0.5	4.91 ± 0.66	6.72 ± 1.35	13.84 ± 2.48	2.43	3.12	7.14	9.09	22.10
$NaPF_6$	DME	0.1	5.06 ± 0.54	6.82 ± 0.65	14.15 ± 0.25	2.49	4.08	17.20	17.60	29.50
$NaPF_6$	DME	1	2.53 ± 0.19	3.41 ± 0.45	7.18 ± 0.16	1.26	1.83	6.49	6.42	13.60
$NaPF_6$	PC	0.1	0.47 ± 0.10	1.08 ± 0.08	1.14 ± 0.05	0.23	0.75	3.21	4.29	4.97
$NaPF_6$	PC	1	0.32 ± 0.03	0.63 ± 0.11	0.69 ± 0.04	0.14	0.39	1.29	1.76	2.43
$NaPF_6$	DEGDME	0.1	1.81 ± 0.36	2.63 ± 0.53	4.09 ± 0.14	0.83	1.74	6.67	5.55	11.80
$NaPF_6$	DEGDME	1	0.80 ± 0.09	1.09 ± 0.16	1.93 ± 0.11	0.40	0.75	2.67	2.96	4.89
$NaPF_6$	TEGDME	1	0.23 ± 0.02	0.35 ± 0.05	0.45 ± 0.03	0.14	0.23	–	0.80	1.12
NaOTf	DME	0.1	4.89 ± 0.40	6.67 ± 0.85	13.74 ± 0.59	2.52	4.10	12.30	11.40	30.10
NaOTf	DME	1	2.52 ± 0.13	3.03 ± 0.16	7.67 ± 0.42	1.33	2.12	5.90	5.44	15.30
NaOTf	PC	0.1	0.45 ± 0.11	1.08 ± 0.09	1.17 ± 0.02	0.21	0.76	3.11	3.70	4.92
NaOTf	PC	1	0.23 ± 0.02	0.46 ± 0.05	0.54 ± 0.02	0.10	0.51	1.68	1.53	2.82
NaOTf	DEGDME	0.1	1.62 ± 0.10	2.47 ± 0.67	4.06 ± 0.22	0.83	1.78	7.25	5.44	11.60
NaOTf	DEGDME	1	0.97 ± 0.09	1.17 ± 0.04	2.23 ± 0.14	0.49	0.82	2.58	2.87	5.94
NaOTf	TEGDME	0.1	0.41 ± 0.06	0.85 ± 0.14	0.75 ± 0.04	0.21	0.56	1.01	1.77	3.05
NaOTf	TEGDME	1	0.18 ± 0.02	0.28 ± 0.03	0.36 ± 0.01	0.11	0.22	0.54	0.94	1.55

Table S9: Raw (uncorrected) self-diffusivities and additive Yeh–Hummer finite-size corrections ΔD_{YH} for OPLS-AA, with experimental self-diffusivities for reference. Simulated entries are mean ± std over replicas; ΔD_{YH} is species-independent ("sim"/"exp" denote the viscosity used to evaluate it, and the finite-size corrected diffusivity is calculated by adding it to the raw value). All values in 10^{-10} m^2/s; "–" denotes values not available.

			D_{raw} (sim.)			ΔD_{YH}		Experimental		
Salt	Solvent	Conc.	D^+	D^-	D^0	sim	exp	D^+	D^-	D^0
$LiPF_6$	DME	0.5	3.55 ± 0.27	3.61 ± 0.26	8.16 ± 0.33	1.44	3.18	7.14	9.09	22.10
$NaPF_6$	DME	0.1	4.96 ± 0.61	6.86 ± 0.44	12.48 ± 0.31	2.33	4.09	17.20	17.60	29.50
$NaPF_6$	DME	1	1.46 ± 0.09	1.85 ± 0.14	4.20 ± 0.12	0.76	1.86	6.49	6.42	13.60
$NaPF_6$	PC	0.1	0.83 ± 0.07	1.29 ± 0.27	2.27 ± 0.04	0.47	0.75	3.21	4.29	4.97
$NaPF_6$	PC	1	0.24 ± 0.02	0.29 ± 0.04	0.80 ± 0.03	0.13	0.39	1.29	1.76	2.43
$NaPF_6$	DEGDME	0.1	1.28 ± 0.13	1.59 ± 0.33	2.63 ± 0.10	0.57	1.75	6.67	5.55	11.80
$NaPF_6$	DEGDME	1	0.34 ± 0.03	0.45 ± 0.04	0.91 ± 0.03	0.19	0.76	2.67	2.96	4.89
$NaPF_6$	TEGDME	1	0.04 ± 0.00	0.07 ± 0.01	0.08 ± 0.01	0.03	0.23	–	0.80	1.12
NaOTf	DME	0.1	4.14 ± 1.08	4.68 ± 0.87	12.32 ± 0.27	2.24	4.12	12.30	11.40	30.10
NaOTf	DME	1	1.22 ± 0.09	1.20 ± 0.30	5.64 ± 0.36	0.96	2.14	5.90	5.44	15.30
NaOTf	PC	0.1	0.88 ± 0.18	1.47 ± 0.11	2.21 ± 0.07	0.43	0.76	3.11	3.70	4.92
NaOTf	PC	1	0.23 ± 0.02	0.28 ± 0.01	0.91 ± 0.05	0.17	0.51	1.68	1.53	2.82
NaOTf	DEGDME	0.1	1.00 ± 0.14	1.35 ± 0.25	2.59 ± 0.07	0.58	1.78	7.25	5.44	11.60
NaOTf	DEGDME	1	0.41 ± 0.05	0.48 ± 0.06	1.06 ± 0.06	0.21	0.82	2.58	2.87	5.94
NaOTf	TEGDME	0.1	0.17 ± 0.04	0.25 ± 0.03	0.29 ± 0.01	0.07	0.56	1.01	1.77	3.05
NaOTf	TEGDME	1	0.04 ± 0.01	0.05 ± 0.01	0.08 ± 0.00	0.03	0.22	0.54	0.94	1.55

Table S10: Simulated and experimental ionic conductivity σ (Onsager, mS/cm) and shear viscosity η (cP). Simulated entries are mean $\pm$ std over replicas for the distilled PaiNN model and OPLS-AA; "–" denotes values not available.

			σ (mS/cm)			η (cP)		
Salt	Solvent	Conc.	PaiNN	OPLS-AA	Exp.	PaiNN	OPLS-AA	Exp.
$LiPF_6$	DME	0.5	12.55 ± 1.49	8.89 ± 2.57	5.68	0.80 ± 0.16	1.38 ± 0.08	0.63
$NaPF_6$	DME	0.1	2.18 ± 1.24	2.52 ± 1.36	0.61	0.77 ± 0.04	0.83 ± 0.02	0.47
$NaPF_6$	DME	1	12.74 ± 3.15	7.84 ± 1.71	12.96	1.59 ± 0.08	2.69 ± 0.13	1.10
$NaPF_6$	PC	0.1	0.59 ± 0.13	0.40 ± 0.11	1.95	8.49 ± 0.44	4.14 ± 0.15	2.58
$NaPF_6$	PC	1	3.77 ± 2.04	0.76 ± 0.06	6.51	14.94 ± 2.42	15.63 ± 0.33	5.16
$NaPF_6$	DEGDME	0.1	0.73 ± 0.18	0.77 ± 0.24	0.52	2.32 ± 0.17	3.37 ± 0.15	1.10
$NaPF_6$	DEGDME	1	5.71 ± 1.43	2.21 ± 0.37	6.69	5.01 ± 0.15	10.53 ± 0.88	2.68
$NaPF_6$	TEGDME	1	1.65 ± 0.30	0.25 ± 0.03	1.86	14.60 ± 2.72	70.56 ± 10.44	8.81
NaOTf	DME	0.1	0.84 ± 0.20	0.09 ± 0.15	0.05	0.76 ± 0.05	0.86 ± 0.02	0.47
NaOTf	DME	1	6.08 ± 1.99	1.28 ± 0.41	1.23	1.52 ± 0.13	2.12 ± 0.12	0.95
NaOTf	PC	0.1	0.50 ± 0.21	0.87 ± 0.42	1.52	9.05 ± 0.31	4.48 ± 0.12	2.52
NaOTf	PC	1	2.13 ± 0.98	0.61 ± 0.25	3.31	19.38 ± 2.14	12.18 ± 0.83	4.01
NaOTf	DEGDME	0.1	0.47 ± 0.12	0.18 ± 0.15	0.11	2.32 ± 0.18	3.33 ± 0.38	1.08
NaOTf	DEGDME	1	5.43 ± 1.32	1.21 ± 0.31	2.78	4.15 ± 0.11	9.69 ± 0.95	2.48
NaOTf	TEGDME	0.1	0.24 ± 0.10	0.13 ± 0.01	0.09	8.94 ± 0.48	25.83 ± 7.57	3.43
NaOTf	TEGDME	1	1.02 ± 0.27	0.23 ± 0.06	1.22	18.30 ± 1.79	73.92 ± 15.76	9.20

S10 Diffusivity calculation details

S10.1 Example mean-squared displacement curves

Figure S6 illustrates how the self-diffusivities tabulated in Section S9 are extracted from the trajectories. For eight different systems, we show the time-origin-averaged mean-squared displacement against lag time of the cation, anion, and solvent, from a single 20 ns production replica of the PaiNN student model. The dashed segments denote the fitted line over the automatically selected Fickian window, which is the longest contiguous stretch with logarithmic slope $|\nu - 1| \leq 0.15$. The long-lag tails of the ion MSDs are clearly noisier for the 0.1 M systems than for the 1.0 M ones, as a result of averaging over around ten times fewer particles.

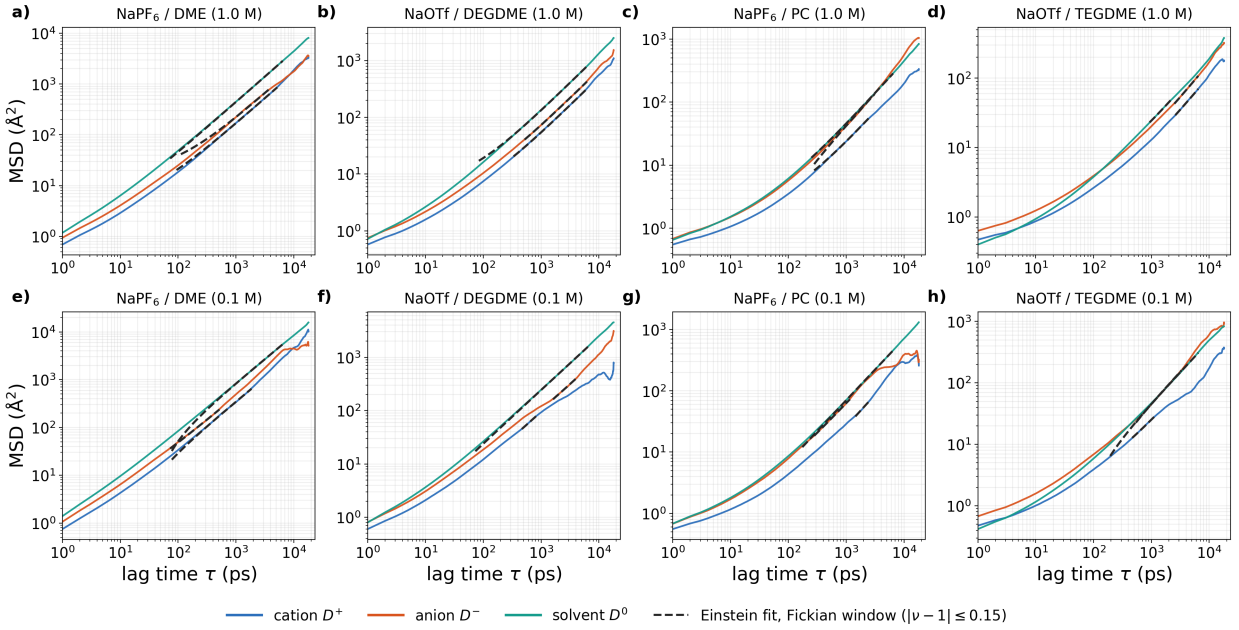

Figure S6: Example MSD curves and Fickian diffusivity fits. Time-origin-averaged mean-squared displacement of the cation, anion, and solvent for the PaiNN production NVT simulations of eight selected systems. Dashed black lines show the linear fit over the longest contiguous window whose logarithmic slope deviates at most 15% from 1.

S11 Viscosity calculation details

S11.1 Example Green–Kubo viscosity integrals

Figure S7 illustrates how the shear viscosities in Section S9 are calculated, using the same eight systems as Figure S6. In each panel, we show the running Green–Kubo integral $\eta(\tau)$ of the stress autocorrelation function. The integral is averaged over the five independent traceless stress components, and evaluated up to a maximum lag time of 1 ns. We take stress samples every 2 fs, and discard the first 2 ns of each NVT run so that only equilibrium samples are used. The integral rises sharply at short timescales, since the stress is highly correlated, then levels off into noise as the correlations go to zero.

We extract the plateau value of η in the following way. First, we split the equilibrated 18 ns of stress data into five equal blocks, and compute the Green–Kubo integral for each. The shaded

band is the standard deviation across the five blocks at each lag. The final fit then uses lags from 1 ps up to the first point where the standard deviation is greater than 40% of the running value of the integral. Everything beyond that point is discarded. Both the band and the running value are smoothed with a 5 ps moving average. The cut takes effect in panels **a**–**c**, **f**, and **g**; in the remaining panels, the noise stays below the threshold for the full 1 ns. The part of the integral that is kept is then fit with the double exponential given in Zhang *et al.* [8],

$$\eta(\tau) \;=\; b_1\big(1 - e^{-\tau/\tau_1}\big) \;+\; b_2\big(1 - e^{-\tau/\tau_2}\big), \qquad \eta_\infty = b_1 + b_2, \tag{S6}$$

where the fit is evaluated on ∼400 logarithmically spaced lags within the window. Each point is weighted by the inverse of the block band, so noisier points contribute less to the fit. The infinite-time limit η_∞ (dashed line) is the value reported for that replica. We finally average over our four replicas to get the reported viscosity for each system.

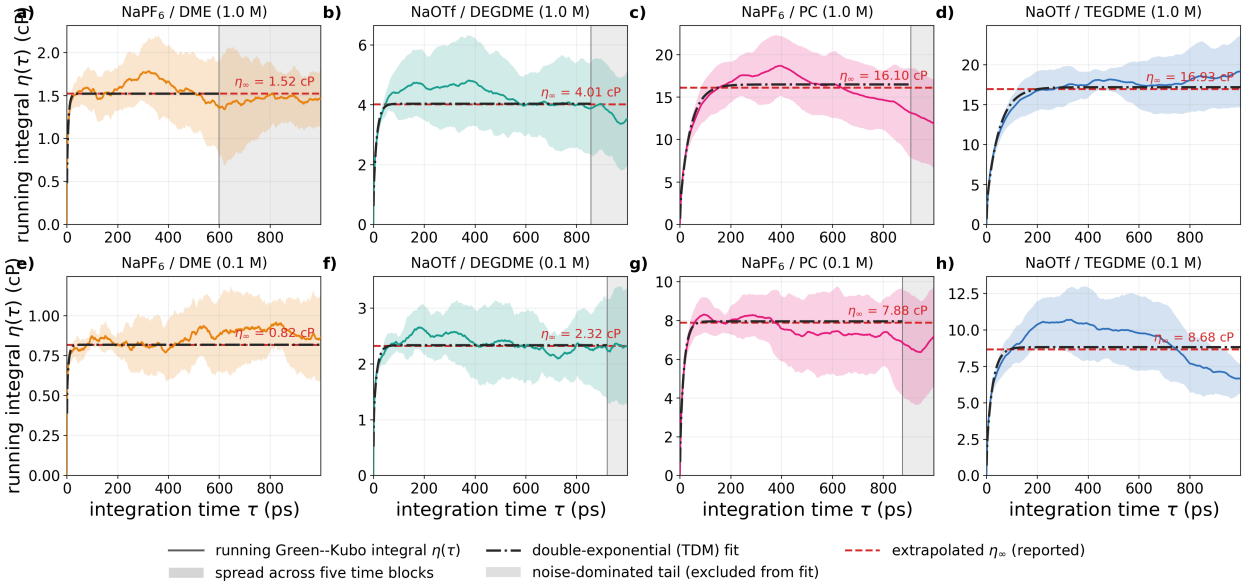

Figure S7: Example running Green–Kubo viscosity integrals. Running Green–Kubo integral $\eta(\tau)$ of the stress autocorrelation function for eight 20 ns PaiNN replicas. The shaded band denotes the standard deviation across five equal time blocks. The dash-dotted black curve is the double-exponential time-decomposition fit (Equation (S6)). The gray region is excluded from the fit due to noise. The dashed red line denotes the extrapolated infinite-time limit η_∞, which is the final viscosity value reported for the replica.

S11.2 Expected viscosity overshoot from the density bias

Since both the UMA-OMol teacher model and the student overshoot the experimental density, we ask how this bias impacts the simulated transport coefficients downstream. Each NVT production box is held at the density predicted by its own model, and both UMA-OMol and the distilled students routinely overpredict the density by $\delta = 6$–15% compared to experiment (Table S7). Writing the viscosity of the liquid as a function of density and temperature, $\eta(\rho, T)$, and the simulated density as $\rho_{\mathrm{sim}} = (1+\delta)\,\rho_{\mathrm{exp}}$, we are interested in how this change in density at constant

temperature impacts the viscosity downstream; in other words,

$$\frac{\eta_{\mathrm{sim}}}{\eta_{\mathrm{exp}}} = \frac{\eta\big((1+\delta)\rho_{\mathrm{exp}}, T\big)}{\eta(\rho_{\mathrm{exp}}, T)}. \tag{S7}$$

In the rest of this section, we estimate this ratio for the liquids considered here, given a δ of 6–15%.

Viscosity is hypersensitive to density. Liquids at room temperature consist of densely packed molecules, and are nearly incompressible on the scale of a few hundred MPa. When compression does happen, the reduction in space is absorbed almost entirely by the free volume between molecules; the bonds in the actual molecules are compressed very little. Flow, or the rearrangement of molecules, requires a molecule to find a free space of at least its size, so the viscosity depends directly on the free volume in the liquid, rather than the total volume. The classic statement of this idea is Doolittle's equation [9], which Cohen and Turnbull later derived from a statistical model of void redistribution [10],

$$\eta = A \exp\!\left(B\, \frac{v_0}{v - v_0} \right), \tag{S8}$$

where v is the specific volume, v_0 the "occupied" volume of the molecules, $v_{\mathrm{f}} = v - v_0$ the free volume, and A, B are constants with B of order one. To measure the sensitivity of η with respect to the density, we take the logarithm of the above, $\ln\eta = \ln A + B v_0/(v - v_0)$, and differentiate with respect to $\ln\rho$ holding the occupied volume v_0 fixed. Holding v_0 fixed is the central assumption of the free-volume picture, since the compression is absorbed entirely by v_{f}. Since $\rho = 1/v$, $\partial/\partial\ln\rho = -v\,\partial/\partial v$, and $\partial(v - v_0)^{-1}/\partial v = -(v - v_0)^{-2}$, so

$$n \equiv \left(\frac{\partial \ln\eta}{\partial \ln\rho} \right)_T = -v\, \frac{\partial}{\partial v} \left[B\, \frac{v_0}{v - v_0} \right] = B\, \frac{v_0\, v}{(v - v_0)^2} = B\, \frac{1 - f}{f^2}, \qquad f \equiv \frac{v_{\mathrm{f}}}{v}, \tag{S9}$$

where the last step substitutes $v - v_0 = fv$ and $v_0 = (1 - f)v$. The $1/f^2$ factor is what makes n large. For a small density overshoot, we get $\Delta\ln\eta \approx n\,\Delta\ln\rho = n\ln(1 + \delta)$, i.e.

$$\frac{\eta_{\mathrm{sim}}}{\eta_{\mathrm{exp}}} \approx (1 + \delta)^n, \tag{S10}$$

treating n as constant over the compression (according to Equation (S9) it actually grows as f shrinks, so this slightly underestimates the effect). The exponent n is system-dependent. As a representative example, take $B = 1$ and a free-volume fraction $f = 0.25$, which are values close to Doolittle's fits for the normal alkanes at room temperature [9]. Then Equation (S9) gives $n = 0.75/0.0625 = 12$. Concretely, making the liquid 10% denser reduces its volume by 9%, and because all of that loss comes out of the free volume, v_{f} falls from $0.25v$ to $0.16v$, a 36% reduction; Equation (S10) then gives $\eta_{\mathrm{sim}}/\eta_{\mathrm{exp}} \approx 1.1^{12} \approx 3$.

The free-volume model, while illustrative, is a crude approximation. First, it assumes that v_0 is fixed, which is only valid for small compressions. For high-pressure data, the occupied volume of the molecules can actually shrink with pressure (for example, the pressure-dependent V_0 that Comuñas *et al.* report for TEGDME in Table 4 of Ref. [11] falls from 211.6 to 205.7 $\mathrm{cm^3\,mol^{-1}}$ between 20 and 100 MPa), and a shrinking v_0 makes η rise more slowly than Equation (S9) predicts. Furthermore, fixing v_0 causes $\eta \to \infty$ once the volume has been reduced by the fraction f, because at that point the model predicts zero free volume; in real liquids, the occupied volume also starts to decrease. The model does, however, capture why viscosity is very sensitive to small changes in the density: compressing a liquid disproportionately affects the free volume, the quantity that controls flow. We turn to measurements for realistic values of n for our systems.

The viscosity rise in real liquids. High-pressure viscometry can be used to measure n directly, by recording η along an isotherm while compressing the liquid under high pressures. If the density is measured simultaneously, $n = \Delta \ln \eta / \Delta \ln \rho$ follows directly. We focus our analysis on TEGDME below, where measured η–density curves exist in the literature.

Comuñas *et al.* [12] measured the density of TEGDME up to 60 MPa, and the viscosity up to 100 MPa using a falling-ball viscometer, along isotherms ranging from 283 to 353 K. Along the 293.15 K isotherm (Tables I and IV of that paper), the density rises from 1.0112 to 1.0435 $\mathrm{g\,cm^{-3}}$ between 0.1 and 60 MPa, an increase of only 3.2%, while the viscosity rises from 3.82 to 6.03 mPa s, a factor of 1.58. The measured exponent over 0.1–60 MPa is then

$$n = \frac{\ln(6.03/3.82)}{\ln(1.0435/1.0112)} = 14.5. \tag{S11}$$

Inspecting the table, n grows with compression as the free-volume picture predicts (13 over 0.1–20 MPa, 16 over 40–60 MPa). At 100 MPa the viscosity reaches 7.99 mPa s, but the density was not measured above 60 MPa. To obtain a state point at 100 MPa, we use the modified Tait equation shown in the paper's Eq. 2:

$$\rho(p) = \frac{\rho(p_{\mathrm{atm}})}{1 - C_{\mathrm{T}} \ln\left[\dfrac{B_{\mathrm{T}} + p}{B_{\mathrm{T}} + p_{\mathrm{atm}}}\right]}, \tag{S12}$$

where $p_{\mathrm{atm}} = 0.1$ MPa, and B_{T} and C_{T} are parameters to be fitted. From these two parameters we can obtain the compressibility at atmospheric pressure, $C_{\mathrm{T}}/(B_{\mathrm{T}} + p_{\mathrm{atm}})$. The authors fit the form of Equation (S12) to their full (T, p) data, but do not list their parameters. To fill the gap, we fit B_{T} and C_{T} to the thirteen densities they tabulate at $T = 293.15$ K, and obtain $C_{\mathrm{T}} = 0.083$ and $B_{\mathrm{T}} = 133$ MPa. The fit reproduces the measured densities to 1.5×10^{-4} $\mathrm{g\,cm^{-3}}$ and gives an atmospheric-pressure compressibility of 6.2×10^{-4} $\mathrm{MPa^{-1}}$, matching the 6.0×10^{-4} the authors list at 5 MPa. When we evaluate Equation (S12) at 100 MPa, we get a density rise of 4.9%. This, alongside the 2.09× viscosity rise at that pressure, corresponds to an exponent of $n = 15.5$, close to the $n = 14.5$ value we got at lower pressures. Thus, $n \approx 14.5$: a 1% increase in the density of TEGDME at room temperature raises its viscosity by about 15%.

To expand our analysis beyond one solvent, we also utilize Bridgman's survey of forty-three liquids [13, 14], which lists, for each liquid, the logarithm of the ratio of its viscosity at a given pressure to its viscosity at atmospheric pressure. Using the 30 °C entries at 2000 $\mathrm{kg\,cm^{-2}}$ (196 MPa), the ratio $\eta(196\ \mathrm{MPa})/\eta(0.1\ \mathrm{MPa})$ for the unassociated organic liquids in the table is 2–4.4 (CS_2 2.0, acetone 2.4, ethyl acetate 2.9, toluene 3.1, *n*-pentane 3.3, ethyl ether 3.3, *n*-hexane 3.6, *n*-octane 4.4). Bridgman notes that the increase is larger the more complicated the molecule. Bridgman does not report densities, but 196 MPa typically compresses liquids of this kind by around 10% (for TEGDME, the Tait extrapolation above gives 8.1%), so these 2–4.4× increases in the viscosity, combined with a ∼10% increase in the density, correspond to $n \approx 7$–18, reinforcing the TEGDME exponent from Comuñas *et al.*

Now that we have an estimate of $n \approx 14.5$ for liquids like ours, we can calculate what the viscosity overshoot would be given our average density overshoot of $\delta = 9.4\%$ (Table S7). Plugging $n = 14.5$ into Equation (S10), we get $1.094^{14.5} = 3.7\times$ for the average density overshoot; a 3.1–4.0× increase in the viscosity for the 8–10% overshoot of the TEGDME boxes; and a 2.3–5.9× viscosity overshoot across the full 6–13% range of student density errors. We conclude that an isothermal compression of ∼10% should be expected to raise the viscosity of liquids like ours by a factor of 3–4, so the observed geometric-mean viscosity overshoot of 2.1× (Fig. 3 of the main text; Table S10) is in the right ballpark. The remaining gap in our predictions is expected: our liquids also contain ions,

while Comuñas *et al.* and Bridgman only measure pure liquids; the exponent is also solvent-specific, and Equation (S10) is only an approximation. The above analysis is only a plausibility estimate, not proof of direct causation. The mapping from density error to viscosity error additionally assumes that the potential reproduces $\eta(\rho, T)$ of the real liquid, which is not necessarily guaranteed: a potential can also incorrectly estimate the viscosity, even given the correct density.

S12 Conductivity

S12.1 Justifying the choice of conductivity calculation

The ionic conductivity is often calculated using the Nernst–Einstein (NE) relation [15, 16], which uses only the diffusion coefficients of the charged species to obtain the conductivity:

$$\sigma_{\mathrm{NE}} = \frac{\beta e^2}{V} \sum_i N_i z_i^2 D_i, \tag{S13}$$

where $\beta = 1/(k_B T)$ is the inverse temperature, e the elementary charge, V the simulation-cell volume, and the sum runs over the ionic species i, each contributing N_i ions of charge number z_i and self-diffusion coefficient D_i.

Because the Nernst–Einstein equation only requires self-diffusivities, it is easy to compute. However, the equation assumes that all charged species move independently of each other, which is not necessarily the case in real electrolytes. Cations and anions tend to pair together, and ions with the same charge push each other apart, so the cross terms in the double sum of Equation (S29) typically make a negative contribution to the current [15, 17]. The NE estimate therefore systematically overpredicts the conductivity, especially at high concentration [15, 16].

To capture cross-correlations, we instead calculate the conductivity with the Onsager transport framework [17, 18], which retains the cross-correlation terms discarded by Nernst–Einstein. The ratio $\sigma_{\mathrm{Onsager}}/\sigma_{\mathrm{NE}}$, called the ionicity [15], helps quantify the impact of cross-correlations on the conductivity. In the following subsection, we derive the Onsager estimator used.

S12.2 Derivation of the Onsager conductivity

We start with the definition of the macroscopic total electrical current $\mathbf{J}(t)$ in a volume V, where the sum runs over all N charged particles with charges q_k and velocities $\mathbf{v}_k(t)$ [15]:

$$\mathbf{J}(t) = \sum_{k=1}^{N} q_k \mathbf{v}_k(t). \tag{S14}$$

According to the Green–Kubo relations from linear response theory [19, 20, 15, 17], the conductivity σ is defined as the time integral of the current auto-correlation function. Using the inverse temperature $\beta = 1/(k_B T)$:

$$\sigma = \frac{\beta}{3V} \int_0^\infty \langle \mathbf{J}(0) \cdot \mathbf{J}(t) \rangle \, dt. \tag{S15}$$

We now relate the conductivity to collective displacements. We define the total dipole moment of the cell $\mathbf{M}(t)$ [15] as:

$$\mathbf{M}(t) = \sum_{k=1}^{N} q_k \mathbf{r}_k(t). \tag{S16}$$

Taking the time derivative gives the current:

$$\dot{\mathbf{M}}(t) = \mathbf{J}(t). \tag{S17}$$

Substituting this back into the Green–Kubo expression gives:

$$\sigma = \frac{\beta}{3V} \int_0^\infty \langle \dot{\mathbf{M}}(0) \cdot \dot{\mathbf{M}}(t) \rangle \, dt. \tag{S18}$$

Next, we analyze the mean squared change in collective displacement over a time interval t, $\Delta\mathbf{M}(t) = \mathbf{M}(t) - \mathbf{M}(0) = \int_0^t \dot{\mathbf{M}}(s)\, ds$. Expanding the square:

$$\begin{aligned} \langle |\Delta\mathbf{M}(t)|^2 \rangle &= \left\langle \left(\int_0^t \dot{\mathbf{M}}(t_1)\, dt_1 \right) \cdot \left(\int_0^t \dot{\mathbf{M}}(t_2)\, dt_2 \right) \right\rangle \\ &= \int_0^t \int_0^t \langle \dot{\mathbf{M}}(t_1) \cdot \dot{\mathbf{M}}(t_2) \rangle \, dt_1 \, dt_2. \end{aligned} \tag{S19}$$

By the time-translation invariance of equilibrium systems, this double integral simplifies to [15]:

$$\langle |\Delta\mathbf{M}(t)|^2 \rangle = 2 \int_0^t (t - \tau) \langle \dot{\mathbf{M}}(\tau) \cdot \dot{\mathbf{M}}(0) \rangle \, d\tau. \tag{S20}$$

We evaluate the long-time limit of the time derivative of this quantity. Utilizing Leibniz's Integral Rule:

$$\lim_{t\to\infty} \frac{\partial}{\partial t} \langle |\Delta\mathbf{M}(t)|^2 \rangle = \lim_{t\to\infty} 2 \int_0^t \langle \dot{\mathbf{M}}(\tau) \cdot \dot{\mathbf{M}}(0) \rangle \, d\tau = 2 \int_0^\infty \langle \dot{\mathbf{M}}(\tau) \cdot \dot{\mathbf{M}}(0) \rangle \, d\tau. \tag{S21}$$

Using the asymptotic relation $\lim_{t\to\infty} \frac{\partial}{\partial t} f(t) = \lim_{t\to\infty} \frac{f(t)}{t}$ for these bounds, we write:

$$\lim_{t\to\infty} \frac{\langle |\Delta\mathbf{M}(t)|^2 \rangle}{t} = 2 \int_0^\infty \langle \dot{\mathbf{M}}(\tau) \cdot \dot{\mathbf{M}}(0) \rangle \, d\tau. \tag{S22}$$

Substituting this result back into our Green–Kubo conductivity expression gives the Einstein relation for the collective displacement [15, 16]:

$$\sigma = \lim_{t\to\infty} \frac{\beta}{6Vt} \langle |\Delta\mathbf{M}(t)|^2 \rangle. \tag{S23}$$

To evaluate this, we expand the change in the total dipole moment, $\Delta\mathbf{M}(t) = \sum_{k=1}^{N} q_k[\mathbf{r}_k(t) - \mathbf{r}_k(0)]$. We group the system by ionic species. For a species i with N_i particles, we define the per-species displacement correlation function $C_{ij}(t)$:

$$C_{ij}(t) = \left\langle \left(\sum_{k=1}^{N_i} [\mathbf{r}_{ik}(t) - \mathbf{r}_{ik}(0)] \right) \cdot \left(\sum_{l=1}^{N_j} [\mathbf{r}_{jl}(t) - \mathbf{r}_{jl}(0)] \right) \right\rangle. \tag{S24}$$

This allows us to define the Onsager phenomenological coefficients L_{ij} [17, 18]:

$$L_{ij} = \frac{\beta}{6V} \lim_{t\to\infty} \frac{1}{t} C_{ij}(t). \tag{S25}$$

Given ionic valencies $q_i = z_i e$, the macroscopic conductivity is represented as a weighted sum of these coefficients [17, 18]:

$$\sigma = e^2 \sum_{i,j} z_i z_j L_{ij}. \tag{S26}$$

Finally, we decompose the total sum into self-diffusion and cross-correlation terms:

$$\sigma_{\text{Onsager}} = \lim_{t\to\infty} \frac{\beta}{6Vt} \sum_{i,j} \sum_{k,l} q_i q_j \langle \Delta\mathbf{r}_{ik} \cdot \Delta\mathbf{r}_{jl} \rangle. \tag{S27}$$

Separating the self terms ($i = j$ and $k = l$) from the distinct-particle cross-terms ($i \neq j$, or $i = j$ with $k \neq l$) yields [15, 18]:

$$\sigma_{\text{Onsager}} = \frac{\beta}{V} \sum_i N_i q_i^2 D_i \; + \; \lim_{t\to\infty} \frac{\beta}{6Vt} \sum_{i,j} \sum_{k\neq l} q_i q_j \langle \Delta\mathbf{r}_{ik} \cdot \Delta\mathbf{r}_{jl} \rangle, \tag{S28}$$

where $D_i = \lim_{t\to\infty} \frac{\langle |\Delta\mathbf{r}_i|^2 \rangle}{6t}$ is the self-diffusion coefficient of species i [15]. The first term depends only on the single-ion self-diffusivities and, using $q_i = z_i e$, is identical to the Nernst–Einstein conductivity of Eq. (S13). We can therefore write

$$\sigma_{\text{Onsager}} = \sigma_{\text{NE}} \; + \; \lim_{t\to\infty} \frac{\beta}{6Vt} \sum_{i,j} \sum_{k\neq l} q_i q_j \langle \Delta\mathbf{r}_{ik} \cdot \Delta\mathbf{r}_{jl} \rangle, \tag{S29}$$

which makes explicit that the Onsager conductivity is the Nernst–Einstein estimate plus the distinct-particle cross-correlations that Nernst–Einstein neglects [15, 16, 18].

S12.3 The conductivity outlier is a low-ionicity, correlation-dominated system

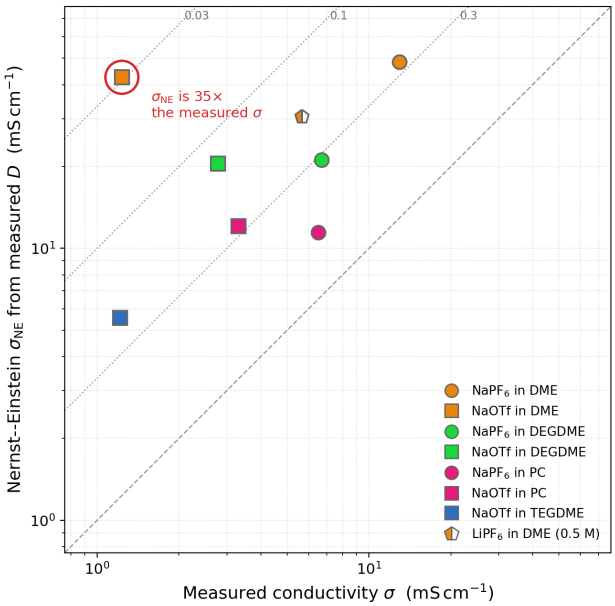

Figure S8: Experimental Nernst–Einstein vs. measured conductivity for the concentrated systems. We plot the experimental Nernst–Einstein conductivity σ_{NE}, obtained only from the experimentally measured diffusivities, against the experimentally measured ionic conductivity σ, for the concentrated ($\geq$0.5 M) systems. The farther a point lies above the $y = x$ line, the lower its experimental ionicity; dotted guides mark ionicities of 0.3, 0.1, and 0.03. NaOTf in DME is a clear outlier, with ion–ion cross-correlations reducing the conductivity from its Nernst–Einstein value by $\sim$97%. 1 M $NaPF_6$ in TEGDME is omitted, as it has no measured cation diffusivity.

In Fig. 3 of the main text, 1 M NaOTf in DME has a much higher error in conductivity compared to the other 1 M systems. This system is also the 1 M system with the lowest ionicity: its ion–ion cross-correlations decrease the conductivity the most of any system. Figure S8 plots the experimental Nernst–Einstein conductivity against the experimentally measured conductivity for the concentrated systems. The Nernst–Einstein conductivity was obtained by inserting the experimentally measured self-diffusivities into Eq. (S13). A system's height above the parity line grows as its experimental ionicity, $\sigma/\sigma_{\mathrm{NE}}$, falls. NaOTf in DME's ionicity is only $\approx$0.03, meaning that the Nernst–Einstein estimate overpredicts the true conductivity by roughly 35×. Every other concentrated electrolyte has an ionicity between roughly 0.1 and 0.6: for these systems, the cross-correlation term cancels a smaller share of σ_{NE}, so an error in the cross-correlations has a much smaller effect on the predicted σ than it does for NaOTf/DME.

Reproducing the system's conductivity depends on correctly capturing the ion–ion interactions that influence the cross-correlations, especially the interaction between cation and anion that binds oppositely charged pairs. We hypothesize that the student overpredicts the NaOTf/DME conductivity because the teacher it is distilled from, UMA-OMol, has a finite radial graph cutoff that prevents it from fully capturing the true attractive force between cation–anion pairs. The student inherits this error, causing it to miss the correlation between oppositely charged ion pairs. In the next subsection, we inspect the teacher's ion–ion interaction profiles with rigid vacuum scans, and find that the teacher under-binds the long-range part of the cation–anion attraction.

S12.4 Origin of the teacher's conductivity error from rigid-scan interaction profiles

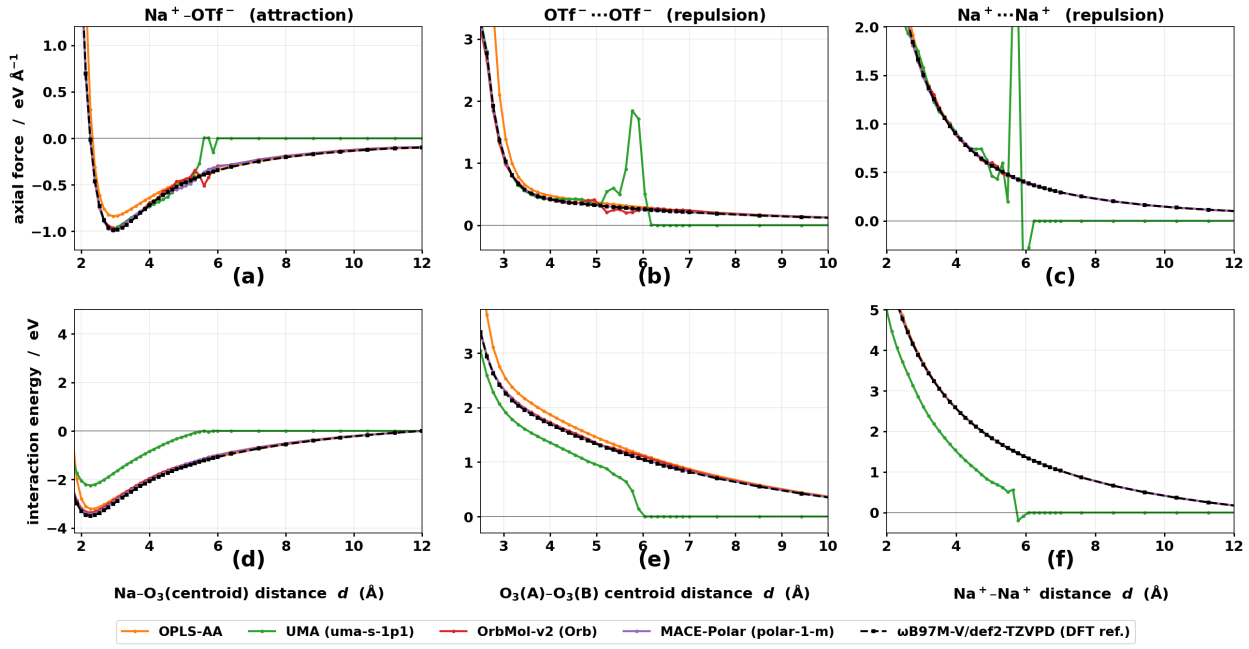

Figure S9: Rigid-scan interaction profiles of the UMA-OMol teacher model. (**a**–**c**) Axial force and (**d**–**f**) interaction energy along rigid separation scans of three representative ion–ion pairs: the (**a**,**d**) Na^+–OTf^- attraction as a function of the Na–O_3 centroid distance, the (**b**,**e**) $OTf^- \cdots OTf^-$ repulsion as a function of the O_3(A)–O_3(B) centroid distance, and the (**c**,**f**) $Na^+ \cdots Na^+$ repulsion as a function of the Na^+–Na^+ distance. We display several different potentials: OPLS-AA (orange), UMA-OMol (`uma-s-1p1`, green), OrbMol-v2 (red), and MACE-Polar (checkpoint `polar-1-m`, purple). As a reference, we show DFT at the ωB97M-V/def2-TZVPD level of theory (black dashed). UMA-OMol's finite radial graph cutoff causes the interaction energies and forces to contain kinks, and to abruptly truncate to zero between 5 and 6 Å.

To test whether the teacher's interactions explain the student's underestimated ion–ion correlations and overpredicted conductivity for NaOTf/DME, we inspect NaOTf/DME's pairwise ion–ion interactions with rigid vacuum scans. For each ion pair, we hold the internal geometry of each ion fixed, and vary the distance between the ions, recording the interaction energy and axial force along the scan. We do this for several potentials, including three different MLIPs and the OPLS-AA baseline. We compare all potentials against the ωB97M-V/def2-TZVPD DFT level of theory, as shown in Figure S9.

The figure shows two main results. First, at short range, the forces of the MLIPs track the DFT reference more closely than those of OPLS-AA, as expected, since the MLIPs were all trained on DFT labels at the ωB97M-V/def2-TZVPD level of theory. Second, UMA-OMol, unlike the other two MLIPs, develops non-physical kinks in its force profiles. Its forces and energies also go to zero by 6 Å, as a result of UMA-OMol's 6 Å radial graph cutoff. This force truncation suppresses the long-range tail of the ion–ion interactions, which likely prevents the full effect of ion–ion correlations from materializing during simulations. This could explain why the student, which inherits these interactions from the teacher, overpredicts the NaOTf/DME conductivity.

S13 Convergence analysis

S13.1 Expanding-window convergence

To verify that the transport coefficients are converged at the 20 ns production simulation length, as well as to illustrate that long timescales are needed for convergence, we recompute estimates of our self-diffusivities, viscosities, and conductivities using varying window lengths of simulation. For each observable, we fix the equilibration cut t_{start}, then sweep the window end t_{end} and recompute the observable on each window. We plot the values of the transport coefficients in Figure S10, normalized to the value at the end of the full 20 ns trajectory. Every system starts off noisy and fluctuating, indicating that timescales greater than a few ns are needed to converge transport coefficients. However, the coefficients settle inside a $\pm 5\%$ band of their final values well before the end of the 20 ns trajectory, indicating convergence.

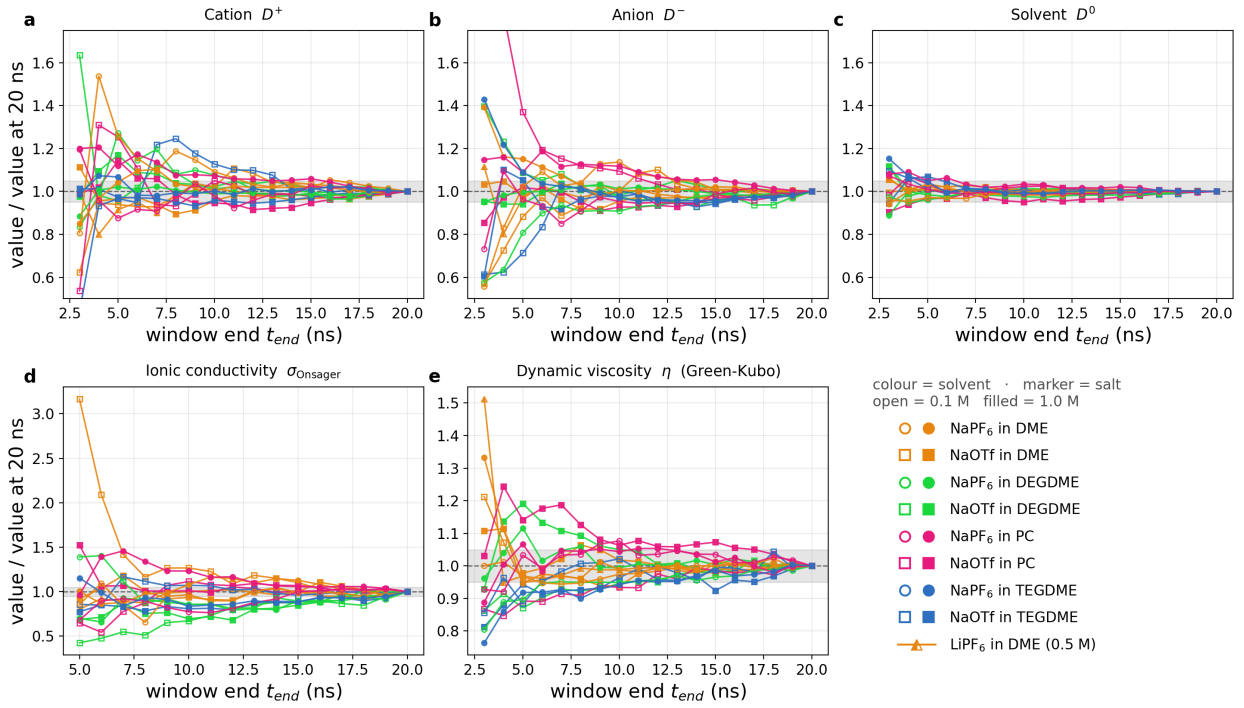

Figure S10: Expanding-window convergence of the transport observables. Each transport coefficient is recomputed over windows of increasing length. We plot each value against the window end t_{end} (the total simulation time used), normalized to the final value at the full 20 ns window: (**a**–**c**) the cation D^+, anion D^-, and solvent D^0 self-diffusivities; (**d**) the Onsager ionic conductivity σ; and (**e**) the Green–Kubo shear viscosity η. The dashed line marks the converged value ($y = 1$) and the shaded band $\pm 5\%$. All observables are unconverged within the first few ns, but settle within the $\pm 5\%$ band before 20 ns has elapsed. To compare different-length trajectories on equal terms, we fit D and σ over a fixed 250–1000 ps lag range rather than the per-species Fickian window (D) and the 0.5–2 ns range (σ) used for the reported values; η uses the production estimator.

S13.2 Error analysis

Previous works have derived theoretical bounds on the relative error of transport coefficients estimated from Green–Kubo integration methods, due to finite simulation time [21]. The "ground

truth" here is the ensemble value of the transport coefficient under our force field—that is, the value we would obtain if we simulated for infinitely long. Experimental values do not enter the analysis anywhere. The bound thus gives us an error estimate between the calculated value from a finite trajectory and the true simulated value, before we compare anything to the experimental value.

We use Jones and Mandadapu's bound for the relative error of a Green–Kubo estimate from a single trajectory [21]:

$$\frac{\mathrm{std}[D(\tau_c)]}{D} \leq 2\sqrt{\frac{\tau_c}{\tau}}, \tag{S30}$$

where $D(\tau_c) = \int_0^{\tau_c} C(t)\,\mathrm{d}t$ is the Green–Kubo estimate from a single trajectory of length τ, $C(t)$ is the relevant flux autocorrelation function, τ_c is the upper limit of the Green–Kubo integral, D is the ensemble value of the transport coefficient, and $\mathrm{std}[D(\tau_c)]$ is the standard deviation of the estimator over independent realizations. The left-hand side of the equation is therefore the relative error: the standard error of the simulated estimator divided by the mean. The bound is subject to the following assumptions:

- $\tau_c \gg \tau_C$, where τ_C is the decorrelation time of the flux autocorrelation function;
- the simulation trajectory used is in equilibrium;
- for multi-replica simulations, the replicas are independent and identically distributed.

The bound can also be extended to the case of multiple replicas:

$$\frac{\mathrm{std}[\bar{D}(\tau_c)]}{D} \leq 2\sqrt{\frac{\tau_c}{\tau}}\,\frac{1}{\sqrt{N_r}}, \tag{S31}$$

where N_r is the number of independent replicas and $\bar{D}(\tau_c)$ is the mean of the single-trajectory estimates. τ_c is a property of the system being studied; the Methods of the main text and Section S11 describe how the integration cutoff is chosen for each simulation.

We can calculate this bound to get an upper limit on the uncertainty of our viscosity estimates, which we compute via Green–Kubo integration of the stress autocorrelation function. Plugging in $N_r = 4$ replicas, $\tau = 18$ ns of equilibrated stress data per replica, and $\tau_c = 1$ ns (the maximum lag of the running viscosity integral),

$$\frac{\mathrm{std}[\bar{\eta}]}{\eta} \leq \frac{2}{\sqrt{4}}\sqrt{\frac{1\ \mathrm{ns}}{18\ \mathrm{ns}}} \approx 24\%. \tag{S32}$$

This is a conservative worst-case bound. In practice, the cutoff is chosen per system and usually falls below 1 ns, the fit weights points by their inverse variance rather than weighting all lags up to the cutoff equally, the fit points are logarithmically spaced so later lag times contribute less to the fit, and the autocorrelation is averaged over five approximately independent stress components [22]. Consistent with this, the actual standard error is much smaller than the bound above. The replica-to-replica spread of the viscosities reported in Table S10 corresponds to a standard error of the four-replica mean of around 2–5%, and at most ∼10%.

The bound can also be applied to self-diffusivity, if calculated via Green–Kubo integration of the velocity autocorrelation function. We use the 1 M NaOTf/DME system as a demonstration. We run 32 replicas of 20 ns with our MLP student (Section S14), discard the first 1 ns of each for equilibration, and use the remaining 19 ns to calculate the self-diffusivities with the Green–Kubo method. We choose 20 ps as our integration cutoff to ensure that $\tau_c \gg \tau_C$, as verified in Figure S12.

Results for the relative error, as well as the underlying theoretical bound, are shown in Figure S11. The figure shows that the relative error of the estimated diffusivity can reach as low as 1%, and is much lower than the bound.

As a sanity check, we compare the Green–Kubo self-diffusivity with the self-diffusivity from the mean-squared displacement (MSD) in Table S11. Neither value includes the finite-size correction. The two estimates agree within 16% for all three species.

Table S11: Self-diffusivities of the 1 M NaOTf/DME system simulated with the MLP student at 298 K, from the MSD and from the Green–Kubo integral with $\tau_c = 20$ ps, both without the finite-size correction. Entries are mean $\pm$ std over replicas (4 replicas for the MSD and 32 replicas for Green–Kubo). All values in 10^{-10} m^2/s.

Species	D_{MSD}	D_{GK}
Na^+	3.69 $\pm$ 1.18	4.02 $\pm$ 0.08
OTf^-	4.49 $\pm$ 1.01	5.18 $\pm$ 0.13
DME	10.22 $\pm$ 0.25	11.68 $\pm$ 0.09

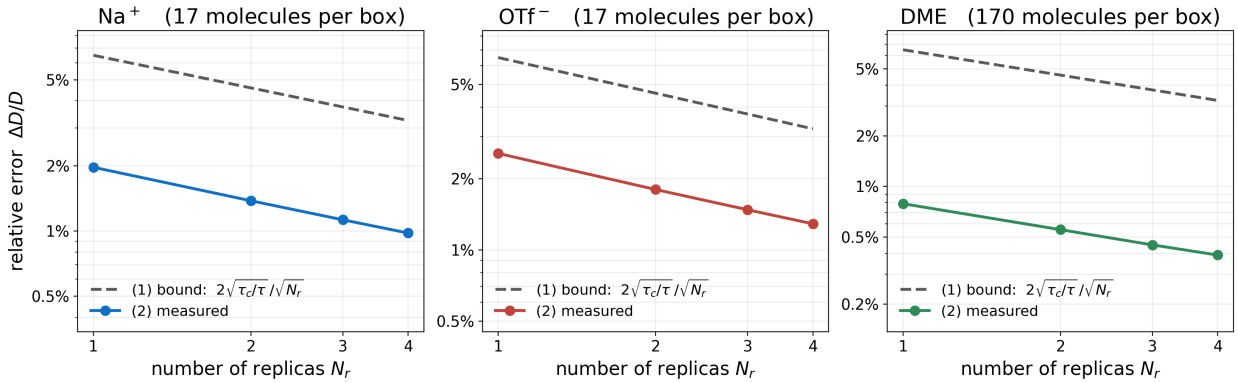

Figure S11: Relative error of the self-diffusivities against the number of replicas. Relative error of the Green–Kubo self-diffusivities of Na^+, OTf^-, and DME in the 1 M NaOTf/DME system, as a function of the number of replicas N_r averaged in the estimate. The dashed line is the theoretical upper bound on the relative error, $2\sqrt{\tau_c/\tau}/\sqrt{N_r}$, with $\tau_c = 20$ ps and $\tau = 19$ ns. The solid line is the measured relative error: for each N_r, we form eight independent N_r-replica means from $8N_r$ of the 32 replicas (e.g. 8 replicas for $N_r = 1$, 16 for $N_r = 2$) and report the relative standard deviation across these eight means. The measured error decreases with N_r and stays well below the bound for all three species.

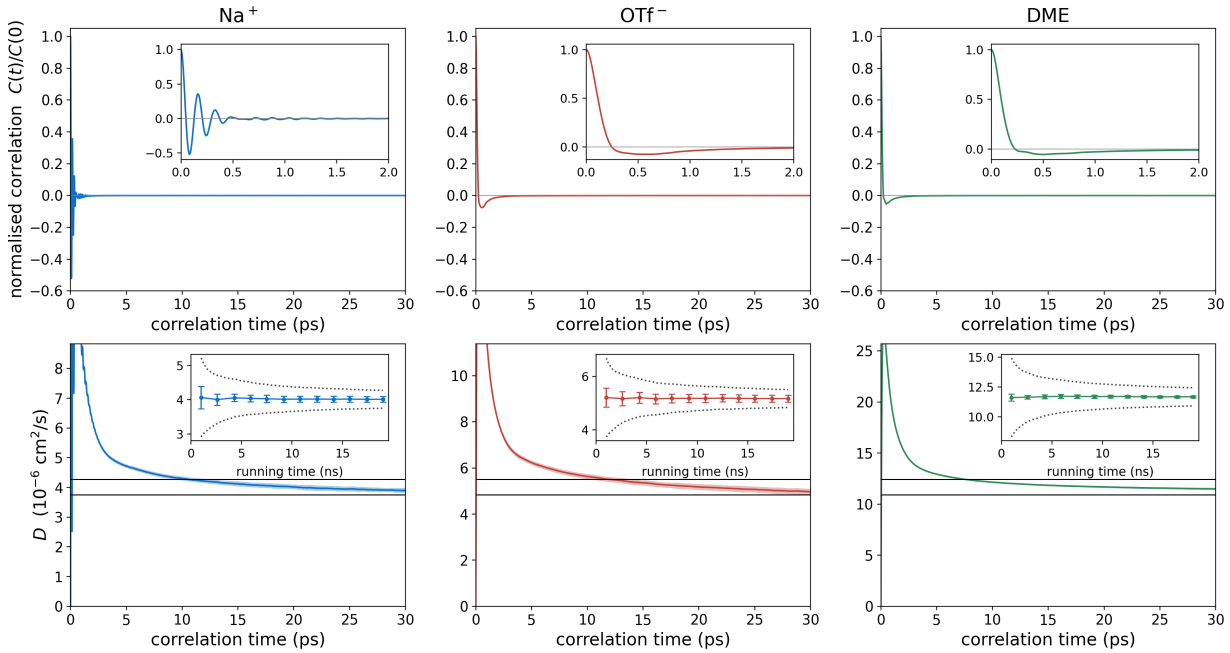

Figure S12: Green–Kubo self-diffusivity convergence for 1 M NaOTf/DME. Estimates of the (top) velocity autocorrelation and (bottom) self-diffusivity with error bars for Na^+, OTf^-, and DME. The inset in the top row is a close-up of the short-time correlation decay. The inset in the bottom row is the evolution of the estimate and error with simulation time. The shaded bands and error bars are the standard deviation over 32 replicas. The two horizontal lines in the bottom row mark the theoretical error bound for a single 19 ns trajectory, drawn around the diffusivity at the $\tau_c = 20$ ps integration cutoff: they sit at $\pm 2\sqrt{\tau_c/\tau} \approx \pm 6.5\%$ of $D(\tau_c)$, with $\tau = 19$ ns. The running integral is shown out to 30 ps, past the cutoff. The dotted lines in the bottom-row insets are the same bound evaluated at each running time τ, so they narrow as the trajectory grows.

S14 The smaller multi-layer-perceptron student

As an ablation to test the speed–accuracy frontier of distilled models, we distill an even lighter student model, which runs around 100× faster than the UMA-OMol teacher it was distilled from (Section S4). The student's architecture is a simple multi-layer perceptron (MLP) applied on top of a Filtered Overlap of Atomic Moments (FOAM) descriptor. We detail the model's architecture and training in Section S14.1, and examine its transport properties compared to experiment in Section S14.2.

S14.1 MLP student architecture and training

The faster student model employed here uses a fixed-form local FOAM descriptor, as implemented in FeNNol [23]. FOAM is a linearly filtered Smooth Overlap of Atomic Positions (SOAP)-type embedding [24]. Because FOAM is purely local, all cross-atom information enters through a single neighbor sum, making the model extremely fast. The FOAM descriptor is then converted into a scalar energy by the same type of energy head used for the PaiNN student. We also include a trainable short-range ZBL repulsion to prevent close nuclear contacts. Table S12 reports the full architecture and optimization settings. The distillation protocol, datasets, loss, and optimizer are the same as those of the PaiNN student (Section S2).

Table S12: Smaller MLP-student architecture and training settings. Identical across every system; only the per-system dataset and its loss normalizers differ. The training scheme (loss, optimizer, schedule) is identical to the PaiNN student in Table S2.

Setting	Value
Architecture	
Backbone	FOAM descriptor + MLP head (FeNNol implementation)
Message passing	none (single-shot local descriptor)
Descriptor	FOAM (filtered SOAP-type invariants)
Angular order	$l_{\max} = 2$
Channels (`nchannels`)	32
Radial cutoff	5.1 Å
Radial basis	8 trainable Bessel functions
Cutoff switching	trainable polynomial ($p = 3$)
Species encoding	random, dim 32 ($Z_{\max} = 86$)
Central-atom species	concatenated into descriptor
Readout head	eSEN energy head, 3-layer MLP (hidden 256)
Nonlinearity	SiLU
Short-range repulsion	trainable ZBL, inner cutoff 1.1 Å
Loss	
Energy loss	per-atom MAE (weight 1)
Force loss	Euclidean-norm (L_2) error (weight 1)
Stress loss	MAE (weight 1)
Loss normalization	dataset force RMS (all three terms)
Optimization	
Optimizer	AdamW
Weight decay	10^{-3} (kernel weights)
Gradient clipping	100
LR schedule	one-cycle cosine, 1% warmup
Learning rate (init/peak/final)	1.6×10^{-4} / 8×10^{-4} / 8×10^{-6}
Batch size	4
Epochs	275
Validation interval	every 5 epochs
EMA	none

S14.2 Transport results

After training the MLP student, we use it to compute the transport coefficients of our sixteen electrolyte systems, with the same protocol as for PaiNN (four 20 ns NVT replicas, analyzed as in Section S9). Figure S13 compares its transport observables against experiment. The MLP student recovers the qualitatively correct transport behavior across different systems, but has systematically looser correlations against experiment than the PaiNN student. The model routinely misrepresents the PC systems, which appear far from parity for every transport coefficient shown. Ionic conductivity is degraded the most compared to PaiNN, with a log–log Pearson correlation of $r = 0.45$, against the PaiNN student's 0.66. Overall, the MLP student preserves some of the teacher's accuracy at a further $\sim$2.5$\times$ speedup over PaiNN; however, its reduced accuracy limits its practical use, and PaiNN, while slower, remains a much more viable model for distillation.

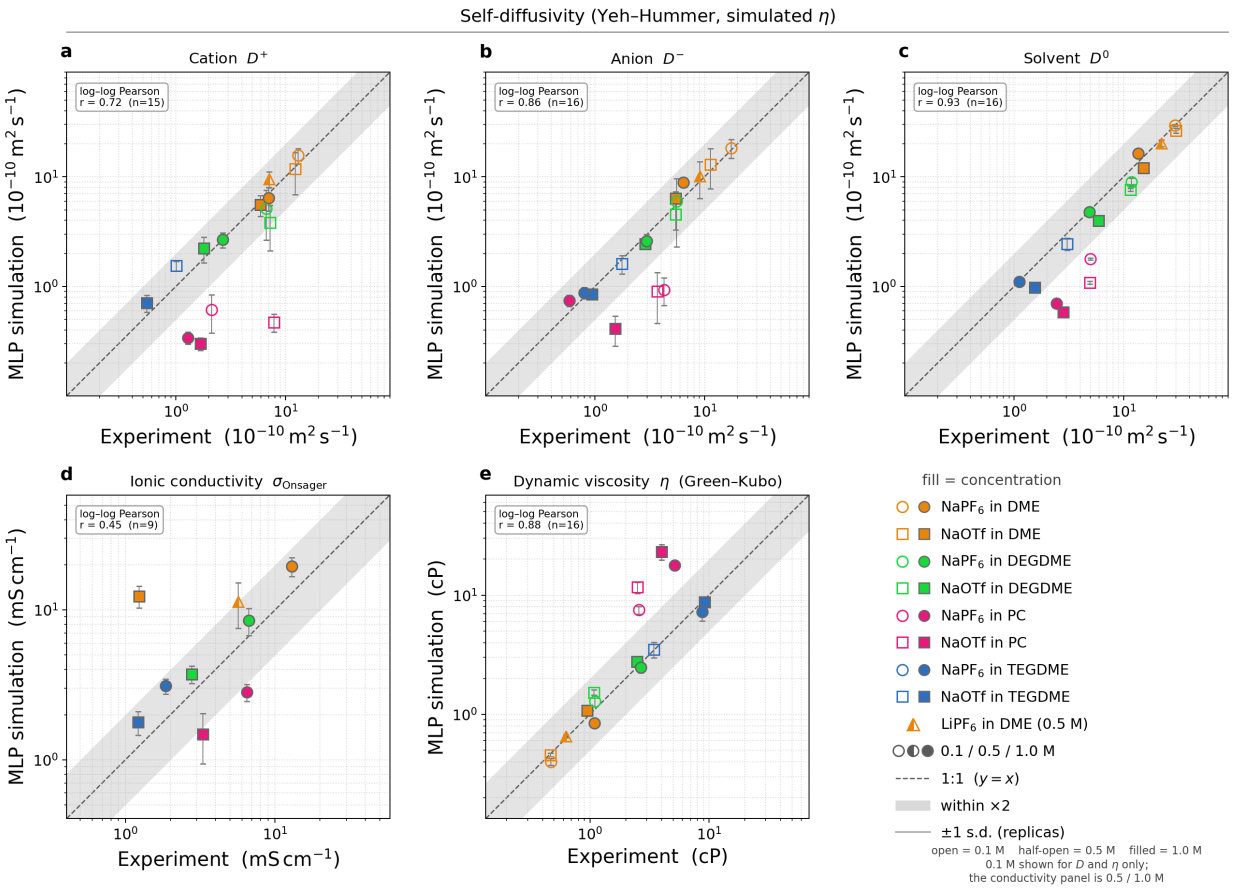

Figure S13: The smaller MLP student trades transport accuracy for speed. Parity of MLP-student simulation against experiment at 298 K for (**a**–**c**) the finite-size-corrected (Yeh–Hummer, simulated η) self-diffusivities of the cation D^+, anion D^-, and solvent D^0; (**d**) the Onsager ionic conductivity σ; and (**e**) the Green–Kubo shear viscosity η. Each point is the mean over four independent replicas and error bars show ± 1 s.d. across replicas. The conductivity panel shows the ≥ 0.5 M systems only.

References

[1] Leela S. Dodda, Israel Cabeza de Vaca, Julian Tirado-Rives, and William L. Jorgensen. Ligpargen web server: an automatic opls-aa parameter generator for organic ligands. *Nucleic Acids Research*, 45(W1):W331–W336, 2017.

[2] José N. Canongia Lopes, Johnny Deschamps, and Agílio A. H. Pádua. Modeling ionic liquids using a systematic all-atom force field. *The Journal of Physical Chemistry B*, 108(6):2038–2047, 2004.

[3] George A. Kaminski and William L. Jorgensen. Host–guest chemistry of rotaxanes and catenanes: application of a polarizable all-atom force field to cyclobis(paraquat-p-phenylene) complexes with disubstituted benzenes and biphenyls. *Journal of the Chemical Society, Perkin Transactions 2*, (11):2365–2375, 1999.

[4] José N. Canongia Lopes and Agílio A. H. Pádua. Molecular force field for ionic liquids composed of triflate or bistriflylimide anions. *The Journal of Physical Chemistry B*, 108(43):16893–16898, 2004.

[5] Andreia S. L. Gouveia, Carlos E. S. Bernardes, Liliana C. Tomé, Elena I. Lozinskaya, Yakov S. Vygodskii, Alexander S. Shaplov, José N. Canongia Lopes, and Isabel M. Marrucho. Ionic liquids with anions based on fluorosulfonyl derivatives: from asymmetrical substitutions to a consistent force field model. *Physical Chemistry Chemical Physics*, 19(43):29617–29624, 2017.

[6] Johan Åqvist. Ion-water interaction potentials derived from free energy perturbation simulations. *The Journal of Physical Chemistry*, 94(21):8021–8024, 1990.

[7] In-Chul Yeh and Gerhard Hummer. System-size dependence of diffusion coefficients and viscosities from molecular dynamics simulations with periodic boundary conditions. *The Journal of Physical Chemistry B*, 108(40):15873–15879, 2004.

[8] Yong Zhang, Akihito Otani, and Edward J. Maginn. Reliable viscosity calculation from equilibrium molecular dynamics simulations: A time decomposition method. *Journal of Chemical Theory and Computation*, 11(8):3537–3546, 2015.

[9] Arthur K. Doolittle. Studies in Newtonian flow. II. the dependence of the viscosity of liquids on free-space. *Journal of Applied Physics*, 22(12):1471–1475, 1951.

[10] Morrel H. Cohen and David Turnbull. Molecular transport in liquids and glasses. *The Journal of Chemical Physics*, 31(5):1164–1169, 1959.

[11] M. J. P. Comuñas, A. Baylaucq, C. Boned, and J. Fernández. Dynamic viscosity for HFC-134a + polyether mixtures up to 373.15 K and 140 MPa at low polyether concentration. measurements and modeling. *Industrial & Engineering Chemistry Research*, 43(3):804–814, 2004.

[12] M. J. P. Comuñas, A. Baylaucq, C. Boned, and J. Fernández. High-pressure measurements of the viscosity and density of two polyethers and two dialkyl carbonates. *International Journal of Thermophysics*, 22(3):749–768, 2001.

[13] P. W. Bridgman. The viscosity of liquids under pressure. *Proceedings of the National Academy of Sciences*, 11(10):603–606, 1925.

[14] P. W. Bridgman. The effect of pressure on the viscosity of forty-three pure liquids. *Proceedings of the American Academy of Arts and Sciences*, 61(3):57–99, 1926.

[15] Jean-Pierre Hansen and Ian R. McDonald. *Theory of Simple Liquids: with Applications to Soft Matter*. Academic Press, Oxford, 4 edition, 2013.

[16] Arthur France-Lanord and Jeffrey C. Grossman. Correlations from ion pairing and the nernst-einstein equation. *Physical Review Letters*, 122(13):136001, 2019.

[17] Kara D. Fong, Helen K. Bergstrom, Bryan D. McCloskey, and Kranthi K. Mandadapu. Transport phenomena in electrolyte solutions: Non-equilibrium thermodynamics and statistical mechanics. *AIChE Journal*, 66(12):e17091, 2020.

[18] Kara D. Fong, Julian Self, Bryan D. McCloskey, and Kristin A. Persson. Onsager transport coefficients and transference numbers in polyelectrolyte solutions and polymerized ionic liquids. *Macromolecules*, 53(21):9503–9512, 2020.

[19] Melville S. Green. Markoff random processes and the statistical mechanics of time-dependent phenomena. II. irreversible processes in fluids. *The Journal of Chemical Physics*, 22(3):398–413, 1954.

[20] Ryogo Kubo. Statistical-mechanical theory of irreversible processes. I. general theory and simple applications to magnetic and conduction problems. *Journal of the Physical Society of Japan*, 12(6):570–586, 1957.

[21] Reese E. Jones and Kranthi K. Mandadapu. Adaptive Green–Kubo estimates of transport coefficients from molecular dynamics based on robust error analysis. *The Journal of Chemical Physics*, 136(15):154102, 2012.

[22] Peter J. Daivis and Denis J. Evans. Comparison of constant pressure and constant volume nonequilibrium simulations of sheared model decane. *The Journal of Chemical Physics*, 100(1):541–547, 1994.

[23] Thomas Plé, Olivier Adjoua, Louis Lagardère, and Jean-Philip Piquemal. FeNNol: An efficient and flexible library for building force-field-enhanced neural network potentials. *The Journal of Chemical Physics*, 161(4):042502, 2024.

[24] Albert P. Bartók, Risi Kondor, and Gábor Csányi. On representing chemical environments. *Physical Review B*, 87(18):184115, 2013.