arXiv:2611.04760v1 [cs.LG] 3 Nov 2026

Koopman Generative Operators for Efficient Probabilistic Time-Series Forecasting

Raz Marshanski[1] **Liran Nochumsohn**[1] **Mayank Jauhari**[2*]

Boris N. Oreshkin[2*] **Omri Azencot**[1]

[1]Ben-Gurion University of the Negev [2]Amazon Science

Abstract

Probabilistic time-series forecasting requires models that simultaneously capture structured temporal dynamics, expressive uncertainty, and efficient inference, yet existing approaches fall short of this goal: latent dynamical models impose structure but rely on restrictive generation mechanisms, while modern generative methods such as diffusion and flow matching achieve flexibility at the cost of iterative and computationally expensive sampling. We introduce the Koopman Generative Operator (KGO), a probabilistic forecasting method that conceptualizes prediction as the evolution of structured uncertainty. KGO integrates three core components: (1) Koopman Patch Embedding (KoPE) for temporally consistent latent trajectory extrapolation; (2) Koopman Flow Matching (KoFM), which enables fast, single-step generation through a closed-form matrix exponential in a Koopman latent space, bypassing iterative sampling; and (3) an Adaptive Uncertainty Gate (AUG), which provides probabilistic forecasts by adapting uncertainty per-variable and per-horizon, corresponding to a learned decomposition of aleatoric uncertainty without manual tuning. By unifying these components, KGO achieves state-of-the-art accuracy on 13/17 ProbTS benchmarks in CRPS and 16/17 in NMAE. Further, by eliminating iterative sampling, KGO delivers at least a 25$\times$ reduction in inference time compared to iterative generative models. These results position KGO as a practical and principled framework that advances the accuracy–efficiency trade-off in probabilistic time-series forecasting, combining leading accuracy on the evaluated benchmarks with fast, single-step generation.

https://github.com/azencot-group/KGO

1 Introduction

Probabilistic time-series forecasting is essential for decision-making under uncertainty, requiring models to predict not only expected outcomes but full predictive distributions over trajectories, rather than single-point estimates [40, 43]. In applications such as energy systems, finance, and healthcare, miscalibrated uncertainty, i.e., inaccurate estimates of the range and likelihood of possible outcomes, can have severe consequences: underestimation leads to critical failures, while overestimation results in overly conservative and inefficient decisions [14, 19]. Recent advances in deep learning have significantly improved the modeling of complex temporal dependencies and multimodal predictive distributions, with approaches based on flows, diffusion, and VAEs demonstrating strong empirical performance [42, 41, 50, 35, 34, 16, 30]. However, practical deployment remains challenging, as modern probabilistic forecasting systems must achieve high predictive accuracy, reliable uncertainty estimates, and efficient inference in real-world settings [43, 41]. Bridging this gap between modeling capability and deployment requirements remains a central challenge in probabilistic forecasting.

Correspondence to: razmar@post.bgu.ac.il. *This work is not related to the authors' positions at Amazon.

40th Conference on Neural Information Processing Systems (NeurIPS 2026).

Despite these advances, existing probabilistic forecasting approaches address different aspects of the problem in a fragmented manner, and no single framework simultaneously captures expressive uncertainty, structured temporal dynamics, and efficient inference. First, modern generative approaches based on diffusion and flow matching provide flexible modeling of complex, multimodal predictive distributions, but rely on *iterative sampling* procedures whose computational cost scales with the number of steps, limiting their applicability in large-scale or real-time settings [41, 46, 28, 7]. Second, many methods compress the historical context into a single representation, either as an explicit latent state, as in latent-variable and Koopman-based models, or through local conditioning mechanisms such as last-patch representations. While latent-variable approaches enable structured temporal evolution, they typically tie uncertainty to the initial state, inducing variability through its propagation rather than modeling how uncertainty evolves over time. This results in a tight coupling between temporal dynamics and uncertainty modeling, limiting flexibility in how predictive distributions are formed, while local conditioning strategies can obscure long-range dependencies and *limit extrapolation* across extended horizons [40, 50, 30]. Finally, existing approaches often adopt fixed or heuristic strategies for uncertainty modeling, treating uncertainty in a *global or implicit manner*, despite the fact that predictability, and thus the appropriate level of uncertainty, *varies across variables and forecast horizons* [25, 26]. Together, these limitations highlight the need for a unified approach that jointly supports efficient generation, models temporal dynamics, and adapts uncertainty to the underlying predictability of the data.

To address these limitations, we argue that probabilistic forecasting should be viewed as modeling the *evolution of structured uncertainty over time*, rather than predicting distributions in a static manner or generating samples through iterative simulation. Future outcomes arise from underlying temporal dynamics, implying that effective forecasting requires representations that preserve and extrapolate this structure across the prediction horizon. At the same time, uncertainty is inherently heterogeneous: the predictability of future observations varies across variables and time, suggesting that uncertainty should be adapted to the underlying signal rather than treated uniformly. This perspective reframes forecasting as a transformation of uncertainty governed by temporal structure, rather than a static prediction or an iterative generative process. It suggests decoupling temporal evolution from uncertainty modeling, allowing each to be handled in a principled and efficient manner.

To instantiate this perspective, we propose *KGO*, a unified framework for probabilistic forecasting. In contrast to prior Koopman-based approaches that evolve latent states and rely on sampling-based decoding, KGO directly models the transport from noise to future trajectories. At its core is *Koopman Flow Matching (KoFM)*, a formulation that replaces iterative generative processes with a closed-form transformation in a learned latent space [13, 4, 5]. Specifically, rather than simulating trajectories through multiple steps as in diffusion or flow-based models, KoFM learns a structured transport that maps noise to future samples via linear evolution, enabling expressive yet tractable dynamics and efficient single-step generation. To support this mechanism, we introduce *Koopman Patch Embedding (KoPE)*, which constructs latent representations that preserve temporal structure and admit extrapolation across the forecast horizon. By lifting observed time series into a space where evolution can be modeled in a structured manner, KoPE provides the state representation on which KoFM operates. Finally, to account for heterogeneous uncertainty, we propose an *Adaptive Uncertainty Gate (AUG)* that modulates the contribution of deterministic and stochastic components based on the predictability of each variable and horizon, allowing the model to allocate uncertainty dynamically. Together, these components realize forecasting as a transformation of structured uncertainty, yielding accurate probabilistic forecasts with efficient inference.

We evaluate KGO on the ProbTS benchmark [52] and observe consistent improvements across a wide range of probabilistic forecasting tasks. KGO achieves the lowest CRPS on **13/17** evaluated settings and the lowest NMAE on **16/17**, with strong results in both regimes: it ranks first on **7/9** long-term datasets in CRPS and **9/9** in NMAE, and on **6/8** short-term datasets in CRPS and **7/8** in NMAE. Notably, KGO enables single-step inference, yielding over **25**$\times$ faster inference than iterative methods. Ablation studies show that replacing KoPE with a last-patch representation, the sLSTM encoder with a Transformer, or KoFM with a normalizing-flow head increases average CRPS by 16.9%, 4.7%, and 9.4%, respectively. These results demonstrate that modeling forecasting as the evolution of structured uncertainty yields both strong predictive accuracy and substantial gains in inference efficiency.

Contributions. We introduce the *Koopman Generative Operator (KGO)*, a framework for probabilistic time-series forecasting that models prediction as the *evolution of structured uncertainty* rather than iterative generative sampling.

1. **A unifying perspective.** We formulate probabilistic forecasting as the evolution of structured uncertainty, providing an alternative to iterative diffusion and flow-based generation and motivating a transport-based view of prediction.
2. **A unified methodology.** We develop KGO, which integrates (i) *Koopman Patch Embedding (KoPE)* for structured latent evolution, (ii) *Koopman Flow Matching (KoFM)* for closed-form, single-step probabilistic generation, and (iii) an *Adaptive Uncertainty Gate (AUG)* for per-variable, per-horizon uncertainty modeling. Together, these components bridge the gap between temporal structure, expressive uncertainty, and efficient inference.
3. **Strong empirical performance.** On the ProbTS benchmark, KGO achieves state-of-the-art results on **29/34** configurations, improving CRPS by up to **27%**, and reducing iterative inference cost by at least **25**×.

2 Related Work

2.1 Deterministic Time-Series Forecasting

Classical time-series forecasting has long relied on statistical models such as ARIMA [8] and vector autoregression [44], which provide strong theoretical foundations but struggle with complex non-linearities and high-dimensional dependencies. The advent of deep learning introduced more expressive models, including recurrent architectures such as LSTMs and GRUs [12, 18], which improved flexibility but remain limited in capturing long-range temporal dependencies. Recent approaches have explored a wide range of architectures, including convolutional, Transformer-based, and MLP-based models [49, 36, 51, 48, 38, 37, 31]. Among these, Transformer and MLP-based methods have emerged as dominant paradigms due to their ability to model long-range dependencies and scale to high-dimensional data. More recently, there has been renewed interest in recurrent architectures through the development of xLSTM and its variants [6, 24, 23, 1], which improve memory capacity and training stability, making them competitive with modern sequence models. Despite these advances, these methods primarily produce point forecasts and do not model predictive uncertainty, limiting their applicability in decision-making settings where uncertainty is critical.

2.2 Probabilistic Time-Series Forecasting

Probabilistic time-series forecasting has been addressed through a range of approaches that differ in how predictive distributions are modeled and generated. Early methods such as DeepAR [43] and DeepState [40] parameterize simple distributions conditioned on historical context, offering efficient inference but limited expressivity. To capture more complex, multimodal distributions, subsequent works have introduced generative models based on normalizing flows [42], variational autoencoders [50], and diffusion and flow matching-based approaches [41, 46, 22, 34, 30]. From the perspective of generation, diffusion and flow-based methods excel in predicting distributions via flexible generative transport processes, but rely on *iterative sampling* procedures whose computational cost scales with the number of steps, limiting their use in real-time settings [30]. This has led to the perception of a trade-off between expressivity and efficiency, which we address by enabling flexible probabilistic modeling with single-step inference.

In terms of *temporal modeling*, Koopman-based approaches impose structured, approximately linear dynamics in a learned latent space for sequence prediction and forecasting [5, 29, 47, 33]. This principle has also been extended to probabilistic and generative sequence modeling [35, 50, 53], where latent uncertainty is propagated through structured dynamics. In contrast, many forecasting architectures capture temporal dependencies implicitly through learned conditioning mechanisms, without explicitly constraining the evolution of the latent state [36]. K^2VAE [50] models uncertainty through a latent initial state sampled from a learned posterior, whose deterministic evolution induces a distribution over future trajectories. In contrast, KGO decouples temporal evolution from generative modeling: Koopman dynamics structure latent trajectories, while a generative operator directly maps source noise to future observations, allowing uncertainty to be modeled at the level of the forecasted sequence. Beyond modeling predictive distributions, KGO explicitly learns how to balance deterministic and stochastic predictions across variables and forecast horizons.

3 Method

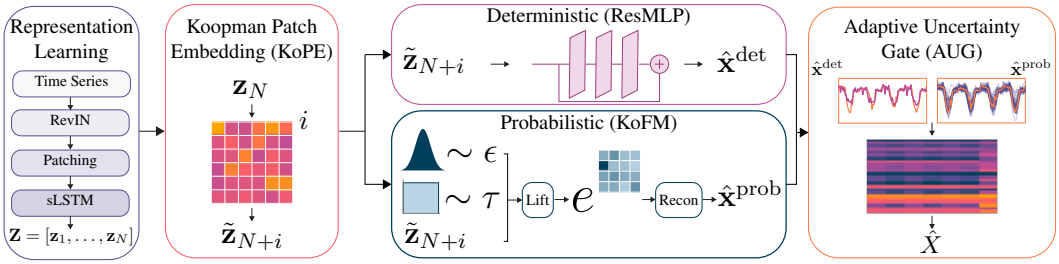

Figure 1: KGO consists of four main parts. (Left) A representation learning stage transforms the input time series via RevIN normalization, patching, and an sLSTM encoder into a latent state. (Middle-left) Koopman Patch Embedding (KoPE) extrapolates latent states by applying powers of a learned linear operator in latent space. (Middle-right, top) A residual MLP (ResMLP) decodes the KoPE latents into a deterministic point forecast. (Middle-right, bottom) Koopman Flow Matching (KoFM) generates probabilistic forecasts by interpolating between noise and data and transporting samples through a matrix exponential conditioned on the KoPE latents. (Right) An Adaptive Uncertainty Gate (AUG) combines deterministic and probabilistic predictions for the final forecast. Extended details in App. A.

3.1 Modeling the Evolution of Structured Uncertainty

We denote by $\mathbf{x}_{1:L}$ the observed historical time series of length L, and by $\mathbf{x}_{L+1:L+H}$ the future horizon of length H. We formalize probabilistic time-series forecasting as the problem of modeling the predictive distribution of future observations given historical context:

$$p(\mathbf{x}_{L+1:L+H} \mid \mathbf{x}_{1:L}) \,. \tag{1}$$

This requires capturing not only the expected trajectory, but also the variability and structure of possible future outcomes. We propose to view probabilistic forecasting as modeling the *evolution of structured uncertainty over time*. Under this perspective, future trajectories arise from underlying temporal dynamics that govern both the signal and its uncertainty, rather than being generated independently or through repeated sampling procedures. This view imposes three key requirements on an effective forecasting model: (i) *efficient generation*, enabling direct sampling of future trajectories without iterative simulation; (ii) *structured temporal evolution*, allowing the model to encode and extrapolate dynamics over long horizons; and (iii) *adaptive uncertainty*, capturing the fact that predictability varies across variables and forecast horizons. These considerations motivate a unified approach that explicitly models forecasting as the evolution of structured uncertainty, jointly satisfying the above requirements within a single framework.

3.2 Overview of KGO

We realize this perspective through *KGO*, a unified framework for probabilistic forecasting that models prediction as the evolution of structured uncertainty. KGO consists of three main components: (i) a *Koopman Flow Matching (KoFM)* module that enables efficient, single-step generation of future trajectories, (ii) a *Koopman Patch Embedding (KoPE)* module that encodes and extrapolates temporal dynamics in a structured latent space, and (iii) an *Adaptive Uncertainty Gate (AUG)* that dynamically balances deterministic and stochastic predictions according to the predictability of the data.

Although these components are conceptually introduced in terms of their roles, their implementation follows a different dependency structure. In particular, KGO first constructs a latent representation of the historical context, which is then extrapolated into the future by KoPE. The resulting sequence of latent states conditions KoFM, which generates probabilistic forecasts in a single step. Finally, AUG fuses the outputs of deterministic and probabilistic decoding streams to produce the final prediction. Formally, given an input time series $\mathbf{x}_{1:L}$, KGO first encodes it into a sequence of latent states $\mathbf{Z} = [\mathbf{z}_1, \ldots, \mathbf{z}_N]$. KoPE then extrapolates the final latent $\mathbf{z}_N$ to produce future latent states $\tilde{\mathbf{Z}}_{\text{KoPE}} = [\tilde{\mathbf{z}}_{N+1}, \ldots, \tilde{\mathbf{z}}_{N+M}]$. These latents serve as conditioning variables for KoFM, which maps noise samples to future trajectories, $\hat{\mathbf{x}}^{\text{prob}}_{L+1:L+H}$. In parallel, a deterministic decoder produces a point forecast, $\hat{\mathbf{x}}^{\text{det}}_{L+1:L+H}$, from the same latent trajectory via a residual MLP. The two outputs are combined by AUG to yield the final prediction. An overview of the architecture is illustrated in Fig. 1.

3.3 Representation Learning

KGO models the evolution of structured uncertainty by first constructing a latent representation of the historical context. We apply reversible instance normalization (RevIN) [21] to mitigate non-stationarity, and partition the input $\mathbf{x}_{1:L}$ into $N = \lfloor L/P \rfloor$ non-overlapping patches of length P, which are projected into D-dimensional embeddings $\mathbf{X}^P \in \mathbb{R}^{N \times D}$ [36]. These embeddings are processed by an sLSTM encoder [6], which captures temporal dependencies while maintaining efficient recurrent computation. Compared to Transformer-based encoders, sLSTM provides lightweight state tracking with linear complexity, making it suitable for long-context forecasting. The encoder outputs latent states $\mathbf{Z} = [\mathbf{z}_1, \ldots, \mathbf{z}_N]$, where the final state $\mathbf{z}_N$ serves as a compact representation of the system state and initializes subsequent Koopman-based temporal evolution.

3.4 Koopman Patch Embedding (KoPE)

A key requirement for modeling the evolution of structured uncertainty is the ability to represent and extrapolate temporal dynamics in a form that supports long-range forecasting. Standard conditioning strategies often rely on compressed representations of the past [36, 30], which can limit the preservation and propagation of temporal structure across the prediction horizon. To address this, we introduce *Koopman Patch Embedding (KoPE)*, which constructs a structured latent trajectory by modeling temporal evolution via a learned linear operator in latent space. This formulation is motivated by the Koopman perspective, where nonlinear dynamics can be represented through linear evolution in a lifted space, an approach that has been successfully exploited for sequence forecasting and generative modeling [5, 35], yielding future states by repeatedly applying the operator.

Given the final latent state $\mathbf{z}_N \in \mathbb{R}^D$ produced by the encoder (see Sec. 3.3), KoPE generates a sequence of future latent states by modeling temporal evolution via a learned linear operator. Specifically, we learn an operator $K \in \mathbb{R}^{D \times D}$ that approximates Koopman evolution, and obtain future states through repeated application:

$$\tilde{\mathbf{z}}_{N+i} = K^{i-1}\mathbf{z}_N \,, \qquad i = 1, \ldots, M \,, \tag{2}$$

where $M = \lceil H/P \rceil$ is the number of patches covering the forecast horizon. To keep the rollout stable, K is normalized by its spectral norm whenever this norm exceeds 1, so that $\|K\|_2 \leq 1$. This rollout propagates the latent state forward in time, preserving temporal structure across the horizon and yielding a structured latent trajectory $\tilde{\mathbf{Z}}_{\text{KoPE}} = [\tilde{\mathbf{z}}_{N+1}, \ldots, \tilde{\mathbf{z}}_{N+M}]$.

Importantly, unlike prior Koopman-based approaches where latent evolution directly defines the generative process [50], KoPE is not used to generate forecasts. Instead, it constructs a latent trajectory $\tilde{\mathbf{Z}}_{\text{KoPE}}$ over the forecast horizon that encodes temporal dynamics and serves as conditioning for the subsequent generative module. This decoupling separates temporal evolution from uncertainty modeling, allowing each to be handled independently and more flexibly.

3.5 Koopman Flow Matching (KoFM)

While KoPE provides a structured latent representation of future dynamics, a second requirement is to efficiently model the predictive distribution over future observations. Existing generative approaches, such as diffusion and flow matching [45, 17, 20, 28], construct flexible distributions but rely on *iterative sampling*, requiring multiple function evaluations at inference time. To address this limitation, we propose *Koopman Flow Matching (KoFM)*, a generative framework that replaces iterative transport with a *closed-form transformation*. Building on the use of Koopman modeling for efficient generative transport [7], KoFM directly parameterizes the transport using a linear dynamical system in a learned latent space, rather than learning a time-dependent velocity field that incrementally transports noise to data, enabling efficient one-step inference.

The flow matching formulation [28] defines a linear interpolation between noise and data:

$$\mathbf{x}_{N+i}^{(\tau)} = (1-\tau)\boldsymbol{\epsilon} + \tau \mathbf{x}_{N+i} \,, \qquad \tau \sim \mathcal{U}[0,1] \,, \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \,, \tag{3}$$

where τ is the transport time. Sampling is performed by drawing noise $\boldsymbol{\epsilon}$ and transporting it along this path; multiple trajectories are obtained by sampling $\boldsymbol{\epsilon}$ independently. In practice, flow matching learns a time-dependent velocity field $\mathbf{v}_\theta(\mathbf{x}^{(\tau)}, \tau)$ that transports samples along this path from noise ($\tau = 0$) to data ($\tau = 1$) by matching the conditional displacement between endpoints. This induces a

nonlinear flow defined by the ordinary differential equation $d\mathbf{x}/d\tau = \mathbf{v}_\theta(\mathbf{x}^{(\tau)}, \tau)$. In contrast, we restrict this transport to a *time-independent* vector field in a learned latent space, corresponding to a finite-dimensional approximation of the Koopman generator [13]. This replaces the general nonlinear flow with a linear dynamical system whose evolution admits a closed-form solution.

KoFM generates future patches through lifting, latent transport, and decoding. Here, i indexes patches along the forecast horizon, whereas $\tau \in [0, 1]$ denotes transport time from noise to data. The interpolated input $\mathbf{x}_{N+i}^{(\tau)}$ is first lifted through a learned transformation conditioned on the KoPE latent $\tilde{\mathbf{z}}_{N+i}$, producing $\tilde{\mathbf{z}}_{N+i}^{(\tau)}$. KoFM then applies a closed-form affine transport:

$$\tilde{\mathbf{z}}_{N+i}^{(1)} = e^{(1-\tau)\mathbf{A}} \tilde{\mathbf{z}}_{N+i}^{(\tau)} + (1-\tau)\mathbf{b} \, . \tag{4}$$

Here, $\mathbf{A}$ is a learned generator and $\mathbf{b}$ is a learned offset. The matrix exponential represents the exact homogeneous linear evolution, while the offset provides a time-scaled translation. A reconstruction head decodes the transported latents into forecast patches, which are concatenated to produce $\hat{\mathbf{x}}_{L+1:L+H}^{\text{prob}}$. At inference, $\tau = 0$, so the cached matrix $e^{\mathbf{A}}$ enables single-step transport through a matrix product and bias addition. During training, $\tau \sim \mathcal{U}[0, 1]$, and to improve optimization we mask τ to 0 for each sample with probability 0.5. App. B.2 details the affine correction and training-time computation.

3.6 Adaptive Uncertainty Gate (AUG)

While KoFM enables efficient generation of probabilistic forecasts, a final requirement is to account for the heterogeneous nature of uncertainty. In time-series forecasting, predictability varies across variables and forecast horizons, implying that the appropriate balance between deterministic and stochastic components should be adaptive rather than fixed [25, 26]. To this end, we combine two complementary prediction streams. A deterministic decoder maps the KoPE latent trajectory $\tilde{\mathbf{Z}}_{\text{KoPE}}$ to a point forecast $\hat{\mathbf{x}}_{L+1:L+H}^{\text{det}}$, capturing the predictable component of the signal, while KoFM produces a stochastic prediction $\hat{\mathbf{x}}_{L+1:L+H}^{\text{prob}}$ that models the full predictive distribution.

We introduce an *Adaptive Uncertainty Gate (AUG)* that dynamically combines these two components. Specifically, for each horizon step h, we compute a gating coefficient $g^h \in [0, 1]$ based on the historical context, and define the final prediction as:

$$\hat{x}_{L+h} = (1 - g^h)\, \hat{x}_{L+h}^{\text{det}} + g^h\, \hat{x}_{L+h}^{\text{prob}} \, . \tag{5}$$

The gating coefficients are produced by a lightweight encoder–router network that maps the input context to per-(variate, horizon) weights, allowing the model to allocate uncertainty dynamically. This enables the model to rely more on deterministic predictions when the signal is predictable, and increase the contribution of the probabilistic component in uncertain regimes. AUG learns to modulate predictive dispersion across variables and forecast horizons, supporting adaptive modeling of aleatoric uncertainty. Appendix B.1 characterizes this effect under the prediction rule in Eq. (5).

3.7 Training Objective

KGO is trained end-to-end on the gate-fused output. For each input, we draw J samples $\{\hat{\mathbf{x}}_{L+1:L+H}^{(j)}\}_{j=1}^{J}$ and optimize a CRPS-based objective composed of two complementary terms:

$$\mathcal{L} = \frac{1}{2} \underbrace{\left(\mathbb{E}\|\hat{X} - X^\star\|_1 - \frac{1}{2}\mathbb{E}\|\hat{X} - \hat{X}'\|_1 \right)}_{\text{Energy Score}} + \frac{1}{2|Q|} \sum_{q \in Q} \rho_q\big(X^\star - \hat{X}_q\big). \tag{6}$$

The first term corresponds to the *Energy Score* [15, 10] formulation of CRPS, which combines an L_1 accuracy term with a repulsion term that discourages sample collapse. Here, $\hat{X}$ and $\hat{X}'$ denote independent samples from the model; in practice, $\hat{X}'$ is created by randomly shuffling the samples in $\hat{X}$. Optimizing this objective for a conditional generative model enables learning forecasts that avoid systematic bias and dispersion errors. An ablation study demonstrating the effectiveness of this component is provided in App. E.5. The second term is the *CRPS quantile loss* [43, 41], computed over a set of quantile levels Q, where $\hat{X}_q$ denotes the empirical q-quantile and $\rho_q(u) = u(q - \mathbb{I}[u < 0])$ is the pinball loss. This term directly supervises predictive quantiles across probability levels, complementing the sample-based term in learning the location and spread of the predictive distribution.

4 Experiments

4.1 Experimental Setup

We evaluate KGO on the ProbTS benchmark [52] across short- and long-horizon probabilistic forecasting. Experiments cover 17 data configurations from standard benchmarks (ETT, Electricity, Exchange, Traffic, Weather, Solar, ILI) under the official protocol. In the short-term regime, context length equals the horizon (24–30), while in the long-term regime, the context length is fixed at 96 and performance is averaged over horizons $\{96, 192, 336, 720\}$; for ILI, the context length is 36 and the horizons are $\{24, 36, 48, 60\}$. We compare against eight baselines spanning major probabilistic forecasting paradigms: flow-based models (GRU MAF, Trans MAF, GRU NVP) [42], diffusion models (TimeGrad, TSDiff, CSDI) [41, 22, 46], and variational autoencoder-based models (K^2VAE, D^3VAE) [50, 27]. These cover autoregressive, flow, diffusion, and latent-dynamics-based designs, providing a comprehensive comparison across model families. We follow the ProbTS evaluation pipeline using CRPS for probabilistic accuracy and NMAE for point forecasting (lower is better), reporting averages over 5 independent runs. Additional details on datasets, metrics, and implementation are provided in App. C.1–C.3. Extended results, ablation studies, and interpretability analyses are provided in Appendices D–F.

Table 1: Long-term forecasting. Results across nine datasets, averaged over horizons and five independent runs (Sec. 4.1). Lower CRPS/NMAE is better. **Red bold**: best; blue underline: second best; "–": unavailable. Per-horizon results and standard deviations appear in Tables 5–6.

Dataset	TSDiff		CSDI		TimeGrad		GRU MAF		Trans MAF		GRU NVP		K^2VAE		**KGO (Ours)**	
Metric	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTm1-L	0.407	0.519	0.324	0.420	0.587	0.736	0.412	0.528	0.476	0.605	0.453	0.585	0.262	0.328	**0.244**	**0.308**
ETTm2-L	0.275	0.338	0.173	0.222	0.447	0.545	0.309	0.377	0.272	0.339	0.414	0.542	0.165	0.194	**0.138**	**0.173**
ETTh1-L	0.457	0.586	0.479	0.603	0.502	0.651	0.353	0.440	0.354	0.453	0.441	0.559	0.294	0.374	**0.270**	**0.347**
ETTh2-L	0.339	0.417	0.241	0.311	0.435	0.545	0.480	0.553	0.282	0.358	0.597	0.736	0.221	0.236	**0.160**	**0.204**
Electricity-L	0.407	0.519	–	–	0.101	0.126	0.094	0.121	–	–	0.101	0.127	**0.066**	0.105	0.080	**0.095**
Exchange-L	0.084	0.111	0.069	0.086	0.082	0.095	0.070	0.083	0.067	0.086	0.073	0.093	0.045	0.053	**0.039**	**0.048**
Traffic-L	0.327	0.392	–	–	0.211	0.245	–	–	–	–	0.198	0.245	**0.142**	0.238	0.193	**0.223**
Weather-L	0.132	0.134	0.077	0.093	0.125	0.155	0.133	0.165	0.123	0.154	0.119	0.147	0.081	0.090	**0.060**	**0.074**
ILI-L	0.248	0.259	0.276	0.290	0.284	0.310	0.129	0.161	0.124	0.151	0.283	0.309	0.116	0.144	**0.097**	**0.118**
1^{st} **Count**	0	0	0	0	0	0	0	0	0	0	0	0	2	0	**7**	**9**
2^{nd} **Count**	0	0	1	0	0	0	0	0	0	0	0	0	6	9	2	0

Table 2: Short-term forecasting. Results across eight datasets, averaged over five independent runs. Lower CRPS/NMAE is better. **Red bold**: best; blue underline: second best. Full results and standard deviations appear in Tables 7–8.

Dataset	TSDiff		D^3VAE		CSDI		TimeGrad		GRU MAF		Trans MAF		K^2VAE		**KGO (Ours)**	
Metric	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTm1-S	0.209	0.276	0.198	0.250	0.140	0.169	0.186	0.246	0.151	0.198	0.139	0.162	0.135	0.152	**0.104**	**0.133**
ETTm2-S	0.124	0.162	0.303	0.378	0.144	0.181	0.148	0.189	0.146	0.193	0.180	0.224	0.122	0.146	**0.092**	**0.118**
ETTh1-S	0.304	0.400	0.324	0.410	0.262	0.339	0.273	0.356	0.258	0.326	0.309	0.400	0.256	0.312	**0.228**	**0.280**
ETTh2-S	0.204	0.272	0.216	0.267	0.133	0.161	0.184	0.224	0.160	0.208	0.200	0.256	0.128	0.140	**0.104**	**0.130**
Electricity-S	0.111	0.115	0.071	0.092	**0.051**	0.066	0.057	0.073	0.056	0.073	0.054	0.071	0.053	0.068	0.052	**0.064**
Exchange-S	0.077	0.096	0.011	0.012	0.009	0.013	0.009	0.012	0.012	0.016	0.012	0.016	0.009	**0.009**	**0.007**	**0.009**
Traffic-S	0.189	0.206	0.143	0.178	0.147	0.175	0.130	0.155	0.144	0.182	0.133	0.160	0.129	0.157	**0.127**	**0.152**
Solar-S	0.568	0.635	0.769	0.998	0.392	0.533	0.465	0.609	0.486	0.603	0.442	0.577	**0.367**	**0.480**	0.456	0.569
1^{st} **Count**	0	0	0	0	1	0	0	0	0	0	0	0	1	2	**6**	**7**
2^{nd} **Count**	0	0	0	1	2	2	1	2	0	0	0	0	6	4	1	0

4.2 Main Results

Tables 1 and 2 demonstrate that KGO achieves the lowest CRPS on 13/17 evaluated settings and the lowest NMAE on 16/17, highlighting its ability to produce accurate predictive distributions (CRPS) while maintaining superior point-forecasting accuracy (NMAE). In the long-term forecasting regime, KGO is particularly dominant, ranking first on 7/9 datasets in CRPS and 9/9 in NMAE. By comparison, the closest competitor, K^2VAE, attains only 2/9 first-place results in CRPS and none in NMAE. Notable gains are observed on ETTh2-L, where KGO reduces CRPS by more than 27% and NMAE by more than 13% relative to K^2VAE. Similarly, on Weather-L, KGO reduces CRPS by approximately 26% and NMAE by approximately 18% relative to K^2VAE. In the short-term forecasting regime, KGO also delivers the best overall performance, obtaining 13 of the 16 best scores

across CRPS and NMAE, compared with 3 for K^2VAE, the next strongest baseline. The largest gains are on the ETT datasets, where KGO reduces CRPS relative to K^2VAE by approximately 23% on ETTm1-S, 25% on ETTm2-S, and 19% on ETTh2-S. Overall, these findings position KGO as a leading method for probabilistic time-series forecasting across the evaluated benchmarks. In particular, KGO substantially outperforms prior flow-based approaches such as Trans MAF, while also surpassing stronger generative baselines such as K^2VAE.

4.3 Efficiency

To assess the trade-off between accuracy and efficiency, we benchmark KGO against the same ProbTS baselines on forecast quality (CRPS, NMAE), inference latency, and peak GPU memory. All models are evaluated with matched context/horizon lengths, batch size 1, FP32 precision, and 100 forecast trajectories per window on the same hardware. Figure 2 shows the (CRPS, NMAE) plane with bubble size encoding inference latency (left) and peak memory (right); full numbers and the measurement protocol are in App. D.4. KGO occupies the lower-left corner of the (CRPS, NMAE) plane and has the lowest average latency in the panel: 11 ms/window versus 27 ms/window for K^2VAE, a 2.5× speedup, averaged over 36 dataset–horizon pairs. Its average peak GPU memory is 0.11 GB, above K^2VAE (0.07 GB) and the CSDI and MAF baselines (0.03–0.05 GB), but below TimeGrad and GRU NVP (0.12 GB) and TSDiff (0.34 GB).

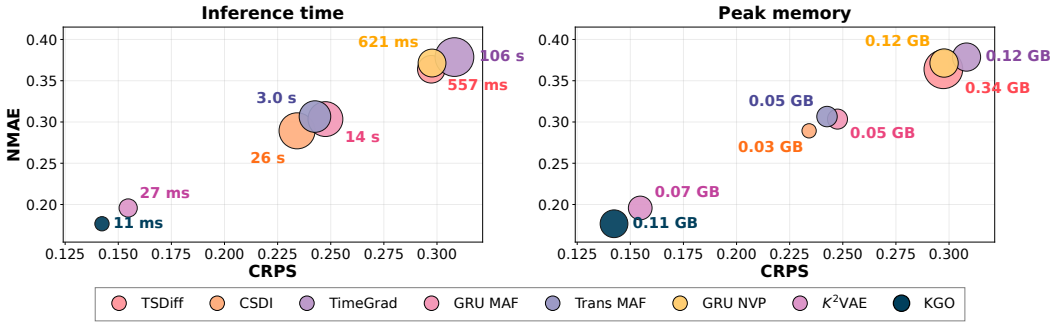

Figure 2: Accuracy–efficiency trade-off on the (CRPS, NMAE) plane. *Left:* bubble size encodes inference latency per window (log scale), generating 100 forecast trajectories. *Right:* bubble size encodes peak GPU memory. Each method is averaged over the datasets for which its accuracy is available in Table 1.

KGO is channel-independent, processing variates in parallel, while future patches and generated samples are batched within the decoder. For fixed patch size and latent dimensions, inference computation and decoder activation memory scale linearly in each of these axes when all samples are processed together. The sLSTM encoder likewise has linear complexity in the number of context patches, while KoFM generates samples in one shot through a matrix-exponential transport, avoiding iterative sampling. Batching exploits available GPU parallelism and can limit latency growth until hardware saturation, after which runtime increases with the computational workload. This parallelism comes with a memory trade-off: processing more trajectories simultaneously increases peak activation memory.

4.4 Foundation Models

We compare full-shot KGO against three off-the-shelf zero-shot foundation models, TiRex [3], Chronos-2 [2], and Sundial [30], under the same ProbTS evaluator; full protocol and the extended table are in App. D.3. The results in Tab. 3 show a clear and consistent advantage for KGO over the foundation-model baselines. Across most datasets, KGO achieves the lowest CRPS and NMAE, ranking first on 5/6 datasets for both metrics, while the baselines rarely take the top spot. The gains are particularly pronounced on Weather and the ETT variants, reflecting improvements in both probabilistic and point-forecasting accuracy. A minor exception is Exchange, where Chronos-2 and TiRex slightly outperform KGO. Overall, despite the growing popularity of these foundation models for zero-shot forecasting, the results demonstrate that a full-shot approach, training directly on the target data, remains highly effective, especially in probabilistic forecasting where precise distributional modeling is critical.

Table 3: Foundation-model comparison. Per-dataset CRPS/NMAE for KGO (trained from scratch) versus three zero-shot time-series foundation models, averaged across horizons $\{96, 192, 336, 720\}$ at context L=96, seed 0. Lower is better. **Red bold**: best; blue underline: second best; bottom rows count per-dataset rankings.

Dataset	TiRex		Chronos-2		Sundial		**KGO (Ours)**	
Metric	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTh1	0.291	0.375	0.274	0.356	0.296	0.347	**0.269**	**0.345**
ETTh2	0.169	0.216	0.168	0.216	0.186	0.216	**0.159**	**0.203**
ETTm1	0.313	0.381	0.291	0.366	0.285	0.344	**0.244**	**0.307**
ETTm2	0.151	0.187	0.147	0.181	0.158	0.186	**0.138**	**0.174**
Weather	0.109	0.131	0.095	0.115	0.087	0.102	**0.060**	**0.075**
Exchange	0.039	0.046	**0.037**	**0.045**	0.045	0.049	0.040	0.048
1^{st} **Count**	0	0	1	1	0	0	**5**	**5**
2^{nd} **Count**	1	2	3	2	2	4	0	0

5 Ablation Study

5.1 Representation and Architecture

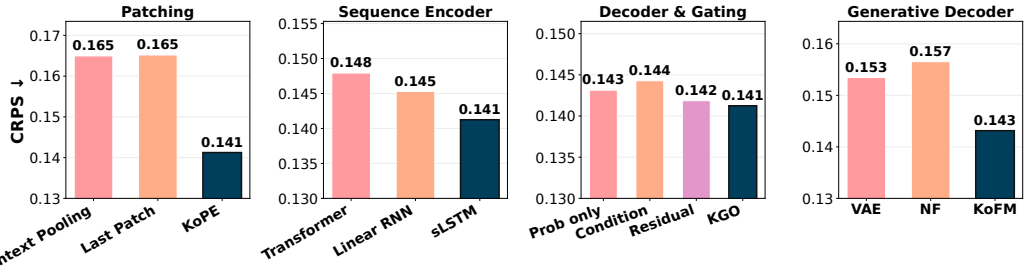

Figure 3: Architecture ablations averaged across datasets and forecasting horizons (lower CRPS is better). Panels evaluate the four core design choices underlying KGO: latent trajectory representation (KoPE), sequence encoder, deterministic–probabilistic coupling (AUG), and generative decoder (KoFM; each head used alone, without the gate or deterministic branch). Koopman-based latent rollout, sLSTM encoding, adaptive gating, and closed-form Koopman transport yield the lowest average CRPS. Full results appear in Appendices E.1–E.4.

In this section, we ablate the three main components, KoPE, KoFM, and AUG, as well as the backbone choice (sLSTM), reporting CRPS averaged over datasets and forecasting horizons (Fig. 3). Across 36 dataset–horizon pairs, KoPE achieves lower average CRPS (0.141) than both the *Last Patch* baseline (0.165) and the *Context Pooling* head (0.165). The two non-Koopman alternatives perform nearly identically, indicating that the gain arises from the Koopman-based representation propagated to the decoder. Consistent with this, the dominant eigenvalues of the learned KoPE operator recover the daily and weekly cycles of the data (App. F.2.1). On the encoder side, holding depth and width fixed, replacing the *sLSTM* backbone with either a *Transformer* or a *Linear RNN* degrades average CRPS (0.141 vs. 0.148 and 0.145, respectively), suggesting that recurrent backbones are better suited than attention-based alternatives in this setting. Combining the deterministic and probabilistic streams through AUG gives the lowest average CRPS among the coupling variants (0.141, vs. 0.144 for conditioning KoFM on the deterministic prediction and 0.142 for fixed residual addition), although the margins are small. AUG thus matches fixed coupling rules in accuracy while adaptively modeling aleatoric uncertainty, learning the contribution of the stochastic branch per variable and forecast horizon (App. F.3). Finally, with each generative head used alone (no gate or deterministic branch) and trained with its native loss, replacing KoFM with either a single-step *VAE* or a *normalizing flow* head increases average CRPS by 7.2% and 9.4%, respectively, demonstrating that *KoFM*, the closed-form Koopman push-forward, produces a stronger predictive distribution than standard one-shot generative heads.

5.2 Flow Matching

To investigate KoFM's ability to deliver accurate predictions while keeping inference efficient, we compare KGO against Sundial-base-128m [30], a strong pretrained flow-matching foundation model, by sweeping its Euler-ODE step budget N and tracing the resulting accuracy–efficiency trade-off in

CRPS, NMAE, and inference latency per window (Fig. 4). KGO occupies the lower-left of the trade-off plane: at the budget $N{=}50$ adopted in the original Sundial paper, KGO achieves $\sim$13% lower CRPS and $\sim$7% lower NMAE, with average latency of 7.8 ms/window versus 387 ms/window for Sundial, a $\sim$50$\times$ speedup. These averages cover 28 dataset–horizon pairs, with each call generating 100 forecast trajectories (per-dataset results in Appendix D.2).

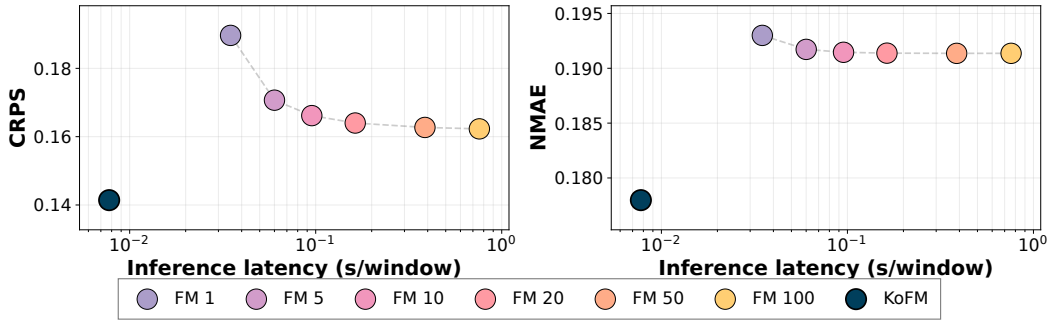

Figure 4: KoFM vs Sundial-base-128m (N-step Euler ODE): aggregate CRPS (left) and NMAE (right) vs inference latency (s/window, logarithmic scale). Each forecast call generates 100 trajectories per multivariate window.

6 Conclusion

We presented the Koopman Generative Operator (KGO), a probabilistic forecasting framework that reinterprets prediction as the evolution of structured uncertainty in a Koopman latent space. By combining Koopman Patch Embedding (KoPE) for consistent latent trajectory extrapolation, Koopman Flow Matching (KoFM) for closed-form, single-step generation, and an Adaptive Uncertainty Gate (AUG) for adaptive, horizon- and variable-aware uncertainty modeling, KGO unifies structure, expressivity, and efficiency within a single model. Empirically, this design yields state-of-the-art accuracy on 29/34 ProbTS configurations, in both short- and long-term forecasting, while reducing inference cost by over an order of magnitude compared to iterative generative methods. These results highlight the advantage of embedding generative modeling within linearizable latent dynamics, enabling both principled uncertainty propagation and scalable inference. More broadly, KGO suggests a new perspective on probabilistic forecasting: rather than treating generation and dynamics as separate processes, they can be jointly governed through operator-theoretic evolution. An important direction for future work is to develop a unified Koopman framework that simultaneously models evolution in physical time and transport time, further bridging dynamical systems and generative modeling.

Acknowledgements

This research was partially supported by the Lynn and William Frankel Center of The Stein Faculty of Computer and Information Science, Ben-Gurion University of the Negev, and by an ISF grant 1299/25.

References

[1] M. Alharthi and A. Mahmood. xLSTMTime: Long-term time series forecasting with xLSTM. *Ai*, 5(3):1482–1495, 2024.

[2] A. F. Ansari, O. Shchur, J. Küken, A. Auer, B. Han, P. Mercado, S. S. Rangapuram, H. Shen, L. Stella, X. Zhang, et al. Chronos-2: From univariate to universal forecasting. *arXiv preprint arXiv:2510.15821*, 2025.

[3] A. Auer, P. Podest, D. Klotz, S. Böck, G. Klambauer, and S. Hochreiter. TiRex: Zero-shot forecasting across long and short horizons with enhanced in-context learning. *arXiv preprint arXiv:2505.23719*, 2025.

[4] O. Azencot, M. Ben-Chen, F. Chazal, and M. Ovsjanikov. An operator approach to tangent vector field processing. In *Computer Graphics Forum*, volume 32, pages 73–82. Wiley Online Library, 2013.

[5] O. Azencot, N. B. Erichson, V. Lin, and M. Mahoney. Forecasting sequential data using consistent Koopman autoencoders. In *International conference on machine learning*, pages 475–485. PMLR, 2020.

[6] M. Beck, K. Pöppel, M. Spanring, A. Auer, O. Prudnikova, M. Kopp, G. Klambauer, J. Brandstetter, and S. Hochreiter. xLSTM: Extended long short-term memory. In *Advances in Neural Information Processing Systems 38*, 2024.

[7] N. Berman, I. Naiman, M. Eliasof, H. Zisling, and O. Azencot. One-step offline distillation of diffusion-based models via Koopman modeling. *Advances in Neural Information Processing Systems*, 38:14776–14811, 2026.

[8] G. E. P. Box, G. M. Jenkins, G. C. Reinsel, and G. M. Ljung. *Time Series Analysis: Forecasting and Control*. Wiley, 5 edition, 2015.

[9] P. M. Catt. Forecastability as an information-theoretic limit on prediction. *arXiv preprint arXiv:2603.27074*, 2026.

[10] J. Chen, T. Janke, F. Steinke, and S. Lerch. Generative machine learning methods for multivariate ensemble postprocessing. *The Annals of Applied Statistics*, 18(1):159–183, 2024.

[11] T. Q. Chen, Y. Rubanova, J. Bettencourt, and D. Duvenaud. Neural ordinary differential equations. In *Advances in Neural Information Processing Systems 31*, 2018.

[12] J. L. Elman. Finding structure in time. *Cognitive Science*, 14(2):179–211, 1990.

[13] K.-J. Engel and R. Nagel. *One-parameter semigroups for linear evolution equations*. Springer, 2000.

[14] T. Gneiting and M. Katzfuss. Probabilistic forecasting. *Annual Review of Statistics and Its Application*, 1(1):125–151, 2014.

[15] T. Gneiting and A. E. Raftery. Strictly proper scoring rules, prediction, and estimation. *Journal of the American statistical Association*, 102(477):359–378, 2007.

[16] T. Gonen, I. Pemper, I. Naiman, N. Berman, and O. Azencot. Time series generation under data scarcity: A unified generative modeling approach. *Advances in Neural Information Processing Systems*, 38:172751–172790, 2026.

[17] J. Ho, A. Jain, and P. Abbeel. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851, 2020.

[18] S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997.

[19] T. Hong and S. Fan. Probabilistic electric load forecasting: A tutorial review. *International Journal of Forecasting*, 32(3):914–938, 2016.

[20] T. Karras, M. Aittala, T. Aila, and S. Laine. Elucidating the design space of diffusion-based generative models. *Advances in neural information processing systems*, 35:26565–26577, 2022.

[21] T. Kim, J. Kim, Y. Tae, C. Park, J. Choi, and J. Choo. Reversible instance normalization for accurate time-series forecasting against distribution shift. In *The Tenth International Conference on Learning Representations, ICLR*, 2022.

[22] M. Kollovieh, A. F. Ansari, M. Bohlke-Schneider, J. Zschiegner, H. Wang, and Y. Wang. Predict, refine, synthesize: Self-guiding diffusion models for probabilistic time series forecasting. In *Advances in Neural Information Processing Systems 36*, 2023.

[23] Y. Kong, Z. Wang, Y. Nie, T. Zhou, S. Zohren, Y. Liang, P. Sun, and Q. Wen. Unlocking the power of LSTM for long term time series forecasting. In *Thirty-Ninth AAAI Conference on Artificial Intelligence*, 2025.

[24] M. Kraus, F. Divo, D. S. Dhami, and K. Kersting. xLSTM-mixer: Multivariate time series forecasting by mixing via scalar memories. *arXiv preprint arXiv:2410.16928*, 2024.

[25] Q. Li, Z. Zhang, L. Yao, Z. Li, T. Zhong, and Y. Zhang. Diffusion-based decoupled deterministic and uncertain framework for probabilistic multivariate time series forecasting. In *The Thirteenth International Conference on Learning Representations, ICLR*, 2025.

[26] Y. Li, W. Chen, X. Hu, B. Chen, B. Sun, and M. Zhou. Transformer-modulated diffusion models for probabilistic multivariate time series forecasting. In *The Twelfth International Conference on Learning Representations, ICLR*, 2024.

[27] Y. Li, X. Lu, Y. Wang, and D. Dou. Generative time series forecasting with diffusion, denoise, and disentanglement. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

[28] Y. Lipman, R. T. Chen, H. Ben-Hamu, M. Nickel, and M. Le. Flow matching for generative modeling. In *The Eleventh International Conference on Learning Representations (ICLR)*, 2023.

[29] Y. Liu, C. Li, J. Wang, and M. Long. Koopa: Learning non-stationary time series dynamics with Koopman predictors. In *Advances in Neural Information Processing Systems 36*, 2023.

[30] Y. Liu, G. Qin, Z. Shi, Z. Chen, C. Yang, X. Huang, J. Wang, and M. Long. Sundial: A family of highly capable time series foundation models. In *Forty-second International Conference on Machine Learning, ICML*, Proceedings of Machine Learning Research, 2025.

[31] Z. Liu. FreqMoE: Enhancing time series forecasting through frequency decomposition mixture of experts. *arXiv preprint arXiv:2501.15125*, 2025.

[32] J. E. Matheson and R. L. Winkler. Scoring rules for continuous probability distributions. *Management Science*, 22(10):1087–1096, 1976.

[33] I. Naiman and O. Azencot. An operator theoretic approach for analyzing sequence neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, volume 37, pages 9268–9276, 2023.

[34] I. Naiman, N. Berman, I. Pemper, I. Arbiv, G. Fadlon, and O. Azencot. Utilizing image transforms and diffusion models for generative modeling of short and long time series. *Advances in Neural Information Processing Systems*, 37:121699–121730, 2024.

[35] I. Naiman, N. B. Erichson, P. Ren, M. W. Mahoney, and O. Azencot. Generative modeling of regular and irregular time series data via Koopman VAEs. In *International conference on learning representations*, volume 2024, pages 50547–50573, 2024.

[36] Y. Nie, N. H. Nguyen, P. Sinthong, and J. Kalagnanam. A time series is worth 64 words: Long-term forecasting with transformers. In *The Eleventh International Conference on Learning Representations, ICLR*, 2023.

[37] L. Nochumsohn, R. Marshanski, H. Zisling, and O. Azencot. Super-linear: A lightweight pretrained mixture of linear experts for time series forecasting. *arXiv preprint arXiv:2509.15105*, 2025.

[38] L. Nochumsohn, H. Zisling, and O. Azencot. A multi-task learning approach to linear multivariate forecasting. *arXiv preprint arXiv:2502.03571*, 2025.

[39] L. Perko. *Differential equations and dynamical systems*, volume 7. Springer Science & Business Media, 2001.

[40] S. S. Rangapuram, M. W. Seeger, J. Gasthaus, L. Stella, Y. Wang, and T. Januschowski. Deep state space models for time series forecasting. In *Advances in Neural Information Processing Systems 31*, 2018.

[41] K. Rasul, C. Seward, I. Schuster, and R. Vollgraf. Autoregressive denoising diffusion models for multivariate probabilistic time series forecasting. In *Proceedings of the 38th International Conference on Machine Learning, ICML*, 2021.

[42] K. Rasul, A. Sheikh, I. Schuster, U. M. Bergmann, and R. Vollgraf. Multivariate probabilistic time series forecasting via conditioned normalizing flows. In *9th International Conference on Learning Representations, ICLR*, 2021.

[43] D. Salinas, V. Flunkert, J. Gasthaus, and T. Januschowski. DeepAR: Probabilistic forecasting with autoregressive recurrent networks. *International journal of forecasting*, 36(3):1181–1191, 2020.

[44] C. A. Sims. Macroeconomics and reality. *Econometrica*, 48(1):1–48, 1980.

[45] J. Sohl-Dickstein, E. Weiss, N. Maheswaranathan, and S. Ganguli. Deep unsupervised learning using nonequilibrium thermodynamics. In *International conference on machine learning*, pages 2256–2265. pmlr, 2015.

[46] Y. Tashiro, J. Song, Y. Song, and S. Ermon. CSDI: conditional score-based diffusion models for probabilistic time series imputation. In *Advances in Neural Information Processing Systems 34*, 2021.

[47] R. Wang, Y. Dong, S. Ö. Arik, and R. Yu. Koopman neural operator forecaster for time-series with temporal distributional shifts. In *The Eleventh International Conference on Learning Representations, ICLR*, 2023.

[48] S. Wang, H. Wu, X. Shi, T. Hu, H. Luo, L. Ma, J. Y. Zhang, and J. Zhou. TimeMixer: Decomposable multiscale mixing for time series forecasting. In *The Twelfth International Conference on Learning Representations, ICLR*, 2024.

[49] H. Wu, J. Xu, J. Wang, and M. Long. Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting. In *Advances in Neural Information Processing Systems 34*, 2021.

[50] X. Wu, X. Qiu, H. Gao, J. Hu, B. Yang, and C. Guo. K2VAE: A koopman-kalman enhanced variational autoencoder for probabilistic time series forecasting. In *Forty-second International Conference on Machine Learning, ICML*, 2025.

[51] A. Zeng, M. Chen, L. Zhang, and Q. Xu. Are transformers effective for time series forecasting? In *Thirty-Seventh AAAI Conference on Artificial Intelligence, AAAI*, 2023.

[52] J. Zhang, X. Wen, Z. Zhang, S. Zheng, J. Li, and J. Bian. ProbTS: Benchmarking point and distributional forecasting across diverse prediction horizons. In *Advances in Neural Information Processing Systems 38*, 2024.

[53] Y. Zhang, L. Ma, A. Valkanas, B. N. Oreshkin, and M. Coates. SKOLR: structured Koopman operator linear RNN for time-series forecasting. In *Forty-second International Conference on Machine Learning, ICML*, Proceedings of Machine Learning Research, 2025.

Appendix

Table of Contents

A Architecture **15**

A.1 Representation Learning 15

A.2 Koopman Patch Embedding (KoPE) 15

A.3 Koopman Flow Matching (KoFM) 15

A.4 Deterministic Decoder (ResMLP) 16

A.5 Adaptive Uncertainty Gate (AUG) 16

B Theory **17**

B.1 Theoretical Analysis of the Uncertainty Gate 17

B.2 Matrix-Exponential Transport and Affine Correction 17

C Experiment Details **20**

C.1 Training Details 20

C.2 Datasets 20

C.3 Metrics 20

D Extended Results **21**

D.1 Full Results Tables 21

D.2 Flow Matching 24

D.3 Foundation Model Comparison 26

D.4 Efficiency 28

E Ablation Studies **29**

E.1 Patching (KoPE) 29

E.2 Sequence Encoder 30

E.3 Decoder & Gating 31

E.4 Generative Decoder 32

E.5 Loss Function 33

E.6 Ablation Summary 34

F Interpretability & Analysis **35**

F.1 Sampling Convergence 35

F.2 Spectral and Stability Behavior 36

F.3 Uncertainty Gate 38

F.4 Qualitative Forecasts 39

A Architecture

The main paper focuses on the conceptual design and motivation underlying KGO. In this appendix, we provide additional implementation details for each module, following the structure of Sec. 3 and Fig. 1: representation learning, Koopman Patch Embedding (KoPE), Koopman Flow Matching (KoFM), the deterministic decoder (ResMLP), and the Adaptive Uncertainty Gate (AUG).

A.1 Representation Learning

Given an input sequence $\mathbf{x}_{1:L} \in \mathbb{R}^L$, we first apply RevIN normalization independently per variable. The normalized sequence is partitioned into non-overlapping patches of length P, producing $N = \lfloor L/P \rfloor$ patches. Each patch is flattened and projected through a learned linear embedding layer into a D-dimensional latent representation.

The sequence encoder is implemented using stacked sLSTM blocks [6]. Each block combines recurrent state updates with multiplicative gating and normalization, allowing efficient long-context processing while maintaining stable hidden-state dynamics. Compared to Transformer-based alternatives, the recurrent formulation reduces memory complexity from quadratic to linear in sequence length and empirically improves forecasting performance in our setting (Sec. 5). The encoder outputs a latent sequence $\mathbf{Z} = [\mathbf{z}_1, \ldots, \mathbf{z}_N]$, from which the final latent state $\mathbf{z}_N$ initializes the KoPE rollout.

A.2 Koopman Patch Embedding (KoPE)

The approximate Koopman operator $K \in \mathbb{R}^{D\times D}$ is initialized orthogonally and normalized by its spectral norm whenever this norm exceeds 1, so that $\|K\|_2 \leq 1$. LayerNorm is applied before the operator rollout and to each resulting conditioning state. Future latent states are generated recursively:

$$\tilde{\mathbf{z}}_{N+i} = K\tilde{\mathbf{z}}_{N+i-1}\ , \quad i = 2, \ldots, M\ , \qquad \tilde{\mathbf{z}}_{N+1} = \mathbf{z}_N\ .$$

In practice, this recursive rollout is fully differentiable and efficiently computed in parallel over the horizon dimension. The resulting latent trajectory $\tilde{\mathbf{Z}}_{\text{KoPE}} = [\tilde{\mathbf{z}}_{N+1}, \tilde{\mathbf{z}}_{N+2}, \ldots, \tilde{\mathbf{z}}_{N+M}]$, $M = \lceil H/P \rceil$, serves as a conditioning signal shared by both deterministic and probabilistic decoding streams.

A.3 Koopman Flow Matching (KoFM)

Lift. Within KoFM, the interpolated inputs $\mathbf{x}^{(\tau)}_{N+i}$ given by

$$\mathbf{x}^{(\tau)}_{N+i} = (1-\tau)\boldsymbol{\epsilon} + \tau\mathbf{x}_{N+i}\ , \qquad \tau \in [0,1], \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})\ ,$$

are first mapped into the latent space through a learned linear projection. Conditioning on the KoPE trajectory is introduced through adaptive Layer Normalization (AdaLN): each KoPE latent $\tilde{\mathbf{z}}_{N+i}$ determines scale, shift, and residual modulation parameters. The normalized and modulated activations pass through an MLP, whose output is scaled and added to the projected input. The resulting conditioned latent $\tilde{\mathbf{z}}^{(\tau)}_{N+i}$ is then transported toward $\tau = 1$.

Transport. The conditioned latent representation $\tilde{\mathbf{z}}^{(\tau)}_{N+i}$ is transported by a learned matrix $\mathbf{A} \in \mathbb{R}^{D\times D}$ and offset $\mathbf{b} \in \mathbb{R}^D$ through the closed-form affine map

$$\tilde{\mathbf{z}}^{(1)}_{N+i} = e^{(1-\tau)\mathbf{A}}\tilde{\mathbf{z}}^{(\tau)}_{N+i} + (1-\tau)\mathbf{b}\ ,$$

where the matrix exponential is computed using the differentiable operator provided by PyTorch. The exponential represents the exact homogeneous linear evolution, while the additive offset retains the first-order term of the corresponding affine ODE's offset integral (Appendix B.2). During training, $\tau \sim \mathcal{U}[0,1]$ is masked to 0 for each sample with probability 0.5 (Sec. 3.5), and the remaining transport time $1-\tau$ is rounded to the nearest of 32 uniformly spaced values in $[0,1]$. At inference, $\tau = 0$, and the cached matrix exponential and learned bias enable single-step latent transport. Multiple trajectories are generated by sampling independent Gaussian noise vectors $\boldsymbol{\epsilon}$ in parallel.

Reconstruction (Recon). A reconstruction head decodes the transported latents $\tilde{\mathbf{z}}^{(1)}_{N+1:N+M}$ into forecast patches. Each transported latent $\tilde{\mathbf{z}}^{(1)}_{N+i}$ is layer-normalized, modulated by AdaLN scale and shift parameters computed from the same KoPE latent $\tilde{\mathbf{z}}_{N+i}$ that conditions the Lift, and linearly projected to a patch of length P. The decoded patches are concatenated to produce $\hat{\mathbf{x}}^{\text{prob}}_{L+1:L+H}$.

A.4 Deterministic Decoder (ResMLP)

In parallel to KoFM, KGO includes a deterministic decoder designed to capture the predictable component of the signal. Each KoPE latent $\tilde{\mathbf{z}}_{N+i}$ is processed by a residual MLP consisting of linear layers with GELU activations, Layer Normalization, and dropout regularization. A linear skip connection projects the latent directly into the output space:

$$\hat{\mathbf{x}}_{N+i}^{\mathrm{det}} = \mathrm{MLP}(\tilde{\mathbf{z}}_{N+i}) + W_{\mathrm{res}}\tilde{\mathbf{z}}_{N+i}.$$

Predicted patches are concatenated to form the deterministic forecast $\hat{\mathbf{x}}_{L+1:L+H}^{\mathrm{det}}$.

A.5 Adaptive Uncertainty Gate (AUG)

AUG applies a lightweight encoder–router network independently to each variable's historical context. The router produces two logits per forecast step, corresponding to the stochastic and deterministic branches. During training, differentiable Gumbel-Softmax with a learned temperature produces the routing weights; at inference, ordinary softmax produces deterministic weights conditional on the context. The stochastic weight is denoted by g^h, with complementary deterministic weight $1 - g^h$, following Eq. (5). This fine-grained routing mechanism enables KGO to interpolate adaptively between deterministic and probabilistic predictions depending on local predictability. The learned gates exhibit horizon- and variate-dependent weighting, with increasing stochastic contributions at later forecast steps on several datasets (Appendix F.3), consistent with AUG's motivation to accommodate heterogeneous aleatoric uncertainty.

B Theory

B.1 Theoretical Analysis of the Uncertainty Gate

Let y denote a future value and $\mathbf{x}_{1:L}$ the historical context. Following Catt [9], the maximum achievable reduction in expected log loss relative to the marginal baseline equals the mutual information between the future and the context:

$$\max_{\theta} \Big(H(y) + \mathbb{E}_{p(\mathbf{x}_{1:L}, y)} \big[\log p_\theta(y \mid \mathbf{x}_{1:L}) \big] \Big) = I(y; \mathbf{x}_{1:L}). \tag{7}$$

This forecastability differs across variables and forecast steps, and the uncertainty that remains given the context is aleatoric. AUG is motivated by this heterogeneity: the gate is trained end-to-end, so the gradient of the CRPS objective determines, per variable and forecast step, how much the forecast relies on the stochastic branch. Moreover, by the data processing inequality, any compression f of the context cannot increase the available forecastability, $I(y; f(\mathbf{x}_{1:L})) \leq I(y; \mathbf{x}_{1:L})$; AUG therefore conditions the gate on the raw context rather than on an encoded representation.

AUG adapts the contribution of stochastic predictions to account for differences in aleatoric uncertainty across variables and forecast horizons. We characterize its effect using the prediction rule in Eq. (5). Let $\mathbf{x} = \mathbf{x}_{1:L}$ denote the historical context. We analyze the predictive distribution at inference. Conditional on $\mathbf{x}$, the deterministic prediction $\hat{x}^{\text{det}}_{L+h}$ and gate $g^h \in [0, 1]$ are fixed, while $\hat{x}^{\text{prob}}_{L+h}$ varies across generated samples.

Proposition 1 (Gate-Controlled Predictive Dispersion)**.** *Assuming the stochastic prediction has finite conditional variance, the gated forecast*

$$\hat{x}_{L+h} = (1 - g^h)\hat{x}^{\text{det}}_{L+h} + g^h \hat{x}^{\text{prob}}_{L+h}$$

satisfies

$$\mathbb{E}[\hat{x}_{L+h} \mid \mathbf{x}] = (1 - g^h)\hat{x}^{\text{det}}_{L+h} + g^h \mathbb{E}[\hat{x}^{\text{prob}}_{L+h} \mid \mathbf{x}], \tag{8}$$

$$\text{Var}(\hat{x}_{L+h} \mid \mathbf{x}) = (g^h)^2 \,\text{Var}(\hat{x}^{\text{prob}}_{L+h} \mid \mathbf{x}). \tag{9}$$

Proof. Conditional on $\mathbf{x}$, the deterministic prediction and gate are constants. The identities follow from linearity of expectation and the variance of an affine transformation. □

Thus, AUG jointly controls the location and dispersion of the predictive distribution. For a fixed stochastic branch, smaller gate values contract generated samples toward the deterministic prediction, while larger values retain more of the stochastic branch's variability. KGO learns this balance through its CRPS-based training objective, allowing the predictive spread to adapt across variables and horizons without a manually specified uncertainty schedule.

This provides a direct interpretation of AUG as a mechanism for modeling aleatoric uncertainty. The gate itself is not an identifiable estimate of that uncertainty, since predictive dispersion also depends on the learned stochastic branch. Appendix F.3 presents empirical analyses of the resulting gate behavior.

B.2 Matrix-Exponential Transport and Affine Correction

In standard Flow Matching [28] and Continuous Normalizing Flows (CNFs) [11], generating a sample requires integrating a learned, time-dependent neural vector field $v_\theta(z, \tau)$ over the continuous flow time $\tau \in [0, 1]$. Standard implementations rely on numerical ODE solvers (e.g., Runge–Kutta), which introduce truncation errors and require multiple sequential function evaluations.

In Koopman Flow Matching (KoFM), we project the data into a latent space. By Koopman semigroup theory [13], lifted observables of any nonlinear dynamical system evolve linearly under the Koopman generator $\mathcal{L}$ (i.e., $\dot{\psi} = \mathcal{L}\psi$). We approximate this infinite-dimensional operator with a finite-dimensional, linear, time-invariant (LTI) generator matrix A. Below, we establish the exact flow map for the homogeneous latent dynamics and its density transformation, assuming the latent distribution admits a density. We then describe the affine correction and time discretization used in the implementation.

B.2.1 The Linear Vector Field and the Exact Flow Map

Let $z_\tau \in \mathbb{R}^d$ be the latent state at flow-matching time $\tau \in [0, 1]$. We parameterize the vector field as a linear autonomous ordinary differential equation:

$$\frac{dz}{d\tau} = v(z) = Az \tag{10}$$

where $A \in \mathbb{R}^{d\times d}$ is the learned generator matrix.

Theorem 1 (Exact Flow Map)**.** *The unique solution to the initial value problem $\frac{dz}{d\tau} = Az$ is given by the matrix exponential flow map $\Phi_s : \mathbb{R}^d \to \mathbb{R}^d$:*

$$\Phi_s(z) = e^{As}z \tag{11}$$

where $e^{As} = \sum_{k=0}^{\infty} \frac{(As)^k}{k!}$.

Proof. Using the Taylor series expansion of the matrix exponential, we differentiate term-by-term (which is valid due to uniform convergence):

$$\frac{d}{ds}e^{As} = \sum_{k=1}^{\infty} \frac{A^k s^{k-1}}{(k-1)!} = A\sum_{k=0}^{\infty} \frac{A^k s^k}{k!} = Ae^{As} \tag{12}$$

Applying this to a state z, we obtain $\frac{d}{ds}(e^{As}z) = A(e^{As}z)$, verifying that the mapping satisfies the differential equation. Furthermore, because the vector field $v(z) = Az$ is globally Lipschitz continuous, the Picard–Lindelöf theorem guarantees that this is the unique solution of the initial value problem [39]. □

For the homogeneous dynamics, a latent state z_τ at transport time τ is advanced to transport time 1. The remaining integration time is $\Delta\tau = 1 - \tau$, giving:

$$z_1 = \Phi_{1-\tau}(z_\tau) = e^{A(1-\tau)}z_\tau \tag{13}$$

This mapping exactly parameterizes the homogeneous latent trajectory and motivates the matrix-exponential component of KoFM.

B.2.2 Exact Log-Density Evaluation

To demonstrate how tracking the single trajectory $\Phi_{1-\tau}$ correctly evolves the probability density, we analyze the density pointwise along the continuous flow.

Lemma 1. *For a linear vector field $v(z) = Az$, the divergence is the trace of the generator matrix: $\nabla \cdot v(z) = Tr(A)$.*

Theorem 2 (Log-Density Evaluation)**.** *Under the linear vector field $v(z) = Az$, the log-density evaluates pointwise along the trajectory as:*

$$\log p_1(z_1) = \log p_\tau(\Phi_{-(1-\tau)}(z_1)) - Tr(A)(1-\tau) \tag{14}$$

Proof. From the Continuous Normalizing Flow framework [11], the instantaneous change of variables formula tracks the log-density $p_\tau(z(\tau))$ along a deterministic trajectory:

$$\frac{d}{d\tau}\log p_\tau(z(\tau)) = -\nabla \cdot v(z(\tau)) \tag{15}$$

Substituting Lemma 1, the divergence is the constant scalar $\mathrm{Tr}(A)$. Integrating this from current time τ to target time 1 yields:

$$\log p_1(z(1)) - \log p_\tau(z(\tau)) = -\int_\tau^1 \mathrm{Tr}(A)\, ds = -\mathrm{Tr}(A)(1-\tau) \tag{16}$$

Let $z_1 = z(1)$ and $z_\tau = z(\tau)$. Because the flow is invertible, the pre-image of the state is $z_\tau = \Phi_{-(1-\tau)}(z_1) = e^{-A(1-\tau)}z_1$. Substituting this into the integrated equation yields the final log-density identity. □

This identity describes the density transformation under the homogeneous latent flow. It does not establish an exact likelihood for the complete forecasting model, which also includes learned lifting and decoding maps.

B.2.3 Efficient Computation and Affine Extension

Evaluating `torch.matrix_exp` directly on a batch of samples with unique, uniformly distributed flow times $\tau_i \sim \mathcal{U}(0,1)$ requires computing B distinct matrix exponentials, which scales poorly during training. Furthermore, dynamically computing the matrix exponential at inference time is suboptimal.

Training via Time-Binning: During training, we discretize the time domain $[0,1]$ into n_b uniformly spaced bins. Rather than computing B distinct matrix exponentials, we compute exactly n_b exponentials:

$$\mathcal{E} = \left\{ e^{A \cdot \Delta\tau_k} \right\}_{k=1}^{n_b} \tag{17}$$

We then assign each sample's integration time $\Delta\tau_i = (1 - \tau_i)$ to its nearest bin. This reduces the forward pass complexity to n_b exponentiations shared across the batch, heavily accelerating training while maintaining the continuous-time parameterization.

Cached Single-Step Inference: At inference, standard Gaussian noise is first lifted into the latent space at $\tau = 0$. The remaining transport time is fixed at $\Delta\tau = 1$, so the implemented affine transport becomes:

$$z_1 = e^A z_0 + b \tag{18}$$

Because A is static after training, we precompute and cache the single operator matrix $C = e^A$. Latent transport then reduces to a single $\mathcal{O}(d^2)$ matrix multiplication and bias addition per sample.

Affine Implementation Note: For the affine dynamics $\frac{d\mathbf{z}}{d\tau} = \mathbf{A}\mathbf{z} + \mathbf{b}$, the exact solution over the remaining transport time $\Delta\tau = 1 - \tau$ is

$$\mathbf{z}_1 = e^{\Delta\tau \mathbf{A}} \mathbf{z}_\tau + \int_0^{\Delta\tau} e^{s\mathbf{A}} \mathbf{b}\, ds.$$

Equation (4) uses $\Delta\tau \mathbf{b}$ in place of this integral, retaining its first-order term. Thus, the matrix exponential exactly represents the homogeneous linear evolution, while the additive offset generally approximates the corresponding affine ODE solution.

C Experiment Details

C.1 Training Details

KGO is implemented in PyTorch and trained with PyTorch Lightning under the ProbTS pipeline [52], using a single GPU. Larger datasets (Electricity, Traffic) are trained on an NVIDIA RTX 6000 Ada, while the remaining smaller datasets are trained on an RTX 4090. We optimize with AdamW ($\beta = (0.9, 0.95)$, weight decay 0.1) under a cosine schedule with 10% linear warmup, placing the Koopman generator parameters in a separate group at $0.1\times$ the base learning rate.

C.2 Datasets

Table 4 summarizes the public benchmarks used throughout the paper. Long-term datasets (*-L*) and short-term datasets (*-S*) follow the ProbTS protocol of [52]; splits, normalization, and context/horizon settings are inherited without modification.

Table 4: Statistical information of the datasets.

Horizon	Dataset	#var.	range	freq.	timesteps	Description
Long-term	ETTh1/h2-L	7	$\mathbb{R}^+$	H	17,420	Electricity transformer temperature per hour
	ETTm1/m2-L	7	$\mathbb{R}^+$	15min	69,680	Electricity transformer temperature every 15 min
	Electricity-L	321	$\mathbb{R}^+$	H	26,304	Electricity consumption (kWh)
	Traffic-L	862	$(0, 1)$	H	17,544	Road occupancy rates
	Exchange-L	8	$\mathbb{R}^+$	Busi. Day	7,588	Daily exchange rates of 8 countries
	ILI-L	7	$(0, 1)$	W	966	Ratio of patients seen with influenza-like illness
	Weather-L	21	$\mathbb{R}^+$	10min	52,696	Local climatological data
Short-term	ETTh1/h2-S	7	$\mathbb{R}^+$	H	17,420	Electricity transformer temperature per hour
	ETTm1/m2-S	7	$\mathbb{R}^+$	15min	69,680	Electricity transformer temperature every 15 min
	Exchange-S	8	$\mathbb{R}^+$	Busi. Day	6,071	Daily exchange rates of 8 countries
	Solar-S	137	$\mathbb{R}^+$	H	7,009	Solar power production records
	Electricity-S	370	$\mathbb{R}^+$	H	5,833	Electricity consumption
	Traffic-S	963	$(0, 1)$	H	4,001	Road occupancy rates

C.3 Metrics

Normalized Mean Absolute Error (NMAE). The Normalized Mean Absolute Error is a normalized version of the MAE, dimensionless and therefore comparable across datasets of different magnitude. It is defined as

$$\mathrm{NMAE} \;=\; \frac{\sum_{k=1}^{K}\sum_{t=1}^{T}\left|x_t^k - \hat{x}_t^k\right|}{\sum_{k=1}^{K}\sum_{t=1}^{T}\left|x_t^k\right|}. \tag{19}$$

Continuous Ranked Probability Score (CRPS). The CRPS [32] quantifies the agreement between a cumulative distribution function (CDF) F and an observation x:

$$\mathrm{CRPS} \;=\; \int_{\mathbb{R}} \big(F(z) - \mathbb{I}\{x \le z\}\big)^2 \, dz, \tag{20}$$

where $\mathbb{I}\{x \le z\}$ denotes the indicator function, equal to one if $x \le z$ and zero otherwise. As a strictly proper scoring rule, CRPS is minimized in expectation when the predictive distribution matches the true conditional distribution. When F is replaced by the empirical CDF $\hat{F}(z) = \frac{1}{n}\sum_{i=1}^{n}\mathbb{I}\{X_i \le z\}$ with samples $X_i \sim F$, the CRPS can be computed directly from samples of the conditional distribution $p_\theta(\mathbf{x}_t \mid \mathbf{h}_t)$. We use $n = 100$ samples to estimate the empirical CDF.

D Extended Results

D.1 Full Results Tables

Table 5: Full long-term probabilistic forecasting results across all baselines and datasets; protocol as in Table 1. Lower CRPS values indicate better predictions. Each cell reports the mean ± standard deviation of the CRPS across 5 independent runs of retraining and evaluation. **Red bold**: best; blue underline: second best. Due to the excessive time and memory consumption, some results are unavailable and denoted as "–".

Dataset	Horizon	TSDiff	CSDI	TimeGrad	GRU MAF	Trans MAF	GRU NVP	K^2VAE	**KGO (Ours)**
ETTm1-L	96	0.344±0.050	0.236±0.006	0.522±0.105	0.295±0.055	0.313±0.045	0.383±0.053	0.232±0.010	**0.215±0.001**
	192	0.345±0.035	0.291±0.025	0.603±0.092	0.389±0.033	0.424±0.029	0.396±0.030	0.259±0.013	**0.236±0.001**
	336	0.462±0.043	0.322±0.033	0.601±0.028	0.429±0.021	0.481±0.019	0.486±0.032	0.262±0.030	**0.249±0.000**
	720	0.478±0.027	0.448±0.038	0.621±0.037	0.536±0.033	0.688±0.043	0.546±0.036	0.294±0.026	**0.276±0.001**
ETTm2-L	96	0.175±0.019	0.115±0.009	0.427±0.042	0.177±0.024	0.227±0.013	0.319±0.044	0.126±0.007	**0.106±0.001**
	192	0.255±0.029	0.147±0.008	0.424±0.061	0.411±0.026	0.253±0.037	0.326±0.025	0.148±0.009	**0.130±0.001**
	336	0.328±0.047	0.190±0.018	0.469±0.049	0.377±0.023	0.253±0.013	0.449±0.145	0.164±0.010	**0.147±0.004**
	720	0.344±0.046	0.239±0.035	0.470±0.054	0.272±0.029	0.355±0.043	0.561±0.273	0.221±0.023	**0.170±0.001**
ETTh1-L	96	0.395±0.052	0.437±0.018	0.455±0.046	0.293±0.034	0.333±0.045	0.379±0.030	0.264±0.020	**0.240±0.001**
	192	0.467±0.044	0.496±0.051	0.516±0.038	0.348±0.075	0.351±0.063	0.425±0.019	0.290±0.016	**0.265±0.001**
	336	0.450±0.027	0.454±0.025	0.512±0.026	0.377±0.026	0.371±0.031	0.458±0.054	0.308±0.021	**0.287±0.004**
	720	0.516±0.027	0.528±0.012	0.523±0.027	0.393±0.043	0.363±0.053	0.502±0.039	0.314±0.011	**0.289±0.005**
ETTh2-L	96	0.336±0.021	0.164±0.013	0.358±0.026	0.239±0.019	0.263±0.020	0.432±0.141	0.162±0.009	**0.134±0.001**
	192	0.265±0.043	0.226±0.018	0.457±0.081	0.313±0.034	0.273±0.024	0.625±0.170	0.186±0.018	**0.155±0.002**
	336	0.350±0.031	0.274±0.022	0.481±0.078	0.376±0.024	0.265±0.042	0.793±0.319	0.257±0.023	**0.173±0.002**
	720	0.406±0.056	0.302±0.040	0.445±0.016	0.990±0.023	0.327±0.033	0.539±0.090	0.280±0.014	**0.177±0.002**
Elec.-L	96	0.344±0.006	0.153±0.137	0.096±0.002	0.083±0.009	0.088±0.014	0.094±0.003	0.073±0.002	**0.066±0.000**
	192	0.345±0.006	0.200±0.094	0.100±0.004	0.093±0.024	0.097±0.009	0.097±0.002	0.080±0.004	**0.074±0.000**
	336	0.462±0.054	–	0.102±0.007	0.095±0.001	–	0.099±0.001	**0.054±0.001**	0.081±0.000
	720	0.478±0.005	–	0.108±0.003	0.106±0.007	–	0.114±0.013	**0.057±0.005**	0.099±0.000
Exch.-L	96	0.079±0.007	0.028±0.003	0.068±0.003	0.026±0.010	0.028±0.002	0.071±0.006	0.031±0.002	**0.020±0.001**
	192	0.093±0.011	0.045±0.003	0.087±0.013	0.034±0.009	0.046±0.017	0.068±0.004	0.032±0.010	**0.029±0.001**
	336	0.081±0.010	0.060±0.004	0.074±0.009	0.058±0.023	0.045±0.010	0.072±0.002	0.048±0.004	**0.040±0.002**
	720	0.082±0.010	0.143±0.020	0.099±0.015	0.160±0.019	0.148±0.017	0.079±0.009	0.069±0.005	**0.068±0.002**
Traffic-L	96	0.294±0.003	–	0.202±0.004	0.215±0.003	0.208±0.004	0.187±0.002	**0.086±0.001**	0.172±0.001
	192	0.306±0.004	–	0.208±0.003	–	–	0.192±0.001	**0.088±0.002**	0.178±0.000
	336	0.317±0.006	–	0.213±0.003	–	–	0.201±0.004	0.195±0.003	**0.192±0.000**
	720	0.391±0.002	–	0.220±0.002	–	–	0.211±0.004	**0.200±0.001**	0.228±0.003
Weather-L	96	0.104±0.020	0.068±0.008	0.130±0.017	0.139±0.008	0.105±0.011	0.116±0.013	0.080±0.007	**0.053±0.001**
	192	0.134±0.012	0.068±0.006	0.127±0.019	0.143±0.020	0.142±0.022	0.122±0.021	0.079±0.009	**0.058±0.001**
	336	0.137±0.010	0.083±0.002	0.130±0.006	0.129±0.012	0.133±0.014	0.128±0.011	0.082±0.010	**0.062±0.000**
	720	0.152±0.003	0.087±0.003	0.113±0.011	0.122±0.006	0.113±0.004	0.110±0.011	0.084±0.005	**0.068±0.001**
ILI-L	24	0.228±0.024	0.250±0.013	0.275±0.047	0.097±0.010	0.092±0.019	0.257±0.003	0.087±0.003	**0.067±0.010**
	36	0.235±0.010	0.285±0.010	0.272±0.057	0.117±0.017	0.115±0.011	0.281±0.004	0.113±0.005	**0.088±0.003**
	48	0.265±0.039	0.285±0.036	0.295±0.033	0.128±0.019	0.133±0.022	0.288±0.008	0.124±0.010	**0.108±0.014**
	60	0.263±0.022	0.283±0.013	0.295±0.083	0.172±0.034	0.155±0.018	0.307±0.005	0.142±0.008	**0.124±0.003**

Table 6: Full long-term probabilistic forecasting results across all baselines and datasets; protocol as in Table 1. Lower NMAE values indicate better predictions. Results are reported as mean ± standard deviation over 5 independent runs of retraining and evaluation. **Red bold**: best; blue underline: second best. Due to the excessive time and memory consumption, some results are unavailable and denoted as "–".

Dataset	Horizon	TSDiff	CSDI	TimeGrad	GRU MAF	Trans MAF	GRU NVP	K^2VAE	**KGO (Ours)**
ETTm1-L	96	0.441±0.021	0.308±0.005	0.645±0.129	0.402±0.087	0.456±0.042	0.488±0.058	0.284±0.011	**0.270±0.002**
	192	0.441±0.019	0.377±0.026	0.748±0.084	0.476±0.046	0.553±0.012	0.514±0.042	0.323±0.020	**0.297±0.002**
	336	0.571±0.033	0.419±0.042	0.759±0.015	0.522±0.019	0.590±0.047	0.630±0.029	0.330±0.014	**0.316±0.000**
	720	0.622±0.045	0.578±0.051	0.793±0.034	0.711±0.081	0.822±0.034	0.707±0.050	0.373±0.032	**0.349±0.001**
ETTm2-L	96	0.224±0.033	0.146±0.012	0.525±0.047	0.212±0.082	0.279±0.031	0.413±0.059	0.144±0.011	**0.133±0.001**
	192	0.316±0.040	0.189±0.012	0.530±0.060	0.535±0.029	0.292±0.041	0.427±0.033	0.170±0.009	**0.162±0.001**
	336	0.397±0.051	0.248±0.024	0.566±0.047	0.407±0.043	0.309±0.032	0.580±0.169	0.187±0.021	**0.184±0.005**
	720	0.416±0.065	0.306±0.040	0.561±0.044	0.355±0.048	0.475±0.029	0.749±0.385	0.275±0.035	**0.213±0.002**
ETTh1-L	96	0.510±0.029	0.557±0.022	0.585±0.058	0.371±0.034	0.423±0.047	0.481±0.037	0.336±0.041	**0.309±0.002**
	192	0.596±0.056	0.625±0.065	0.680±0.058	0.430±0.022	0.451±0.012	0.531±0.018	0.372±0.023	**0.341±0.001**
	336	0.581±0.035	0.574±0.026	0.666±0.047	0.462±0.049	0.481±0.041	0.580±0.064	0.394±0.022	**0.365±0.004**
	720	0.657±0.017	0.657±0.014	0.672±0.015	0.496±0.019	0.455±0.025	0.643±0.046	0.396±0.012	**0.375±0.006**
ETTh2-L	96	0.421±0.033	0.214±0.018	0.448±0.031	0.292±0.012	0.345±0.042	0.548±0.158	0.189±0.010	**0.170±0.002**
	192	0.339±0.033	0.294±0.027	0.575±0.089	0.376±0.112	0.343±0.044	0.766±0.223	0.213±0.021	**0.198±0.003**
	336	0.427±0.041	0.353±0.028	0.606±0.095	0.454±0.057	0.333±0.078	0.942±0.408	0.263±0.039	**0.220±0.003**
	720	0.482±0.022	0.382±0.030	0.550±0.018	1.092±0.019	0.412±0.020	0.688±0.161	0.278±0.020	**0.228±0.002**
Elec.-L	96	0.441±0.013	0.203±0.189	0.119±0.003	0.108±0.009	0.114±0.010	0.118±0.003	0.093±0.002	**0.079±0.000**
	192	0.441±0.005	0.264±0.129	0.124±0.005	0.120±0.033	0.131±0.008	0.121±0.003	0.102±0.010	**0.087±0.000**
	336	0.571±0.022	–	0.126±0.008	0.122±0.018	–	0.123±0.001	0.107±0.002	**0.096±0.000**
	720	0.622±0.142	–	0.134±0.004	0.136±0.098	–	0.144±0.017	**0.117±0.019**	0.118±0.001
Exch.-L	96	0.090±0.010	0.036±0.005	0.079±0.002	0.033±0.003	0.036±0.009	0.091±0.009	0.032±0.002	**0.025±0.001**
	192	0.106±0.010	0.058±0.005	0.100±0.019	0.044±0.004	0.058±0.007	0.087±0.005	0.040±0.005	**0.036±0.001**
	336	0.106±0.010	0.076±0.006	0.086±0.008	0.074±0.017	0.058±0.009	0.091±0.002	0.054±0.001	**0.050±0.002**
	720	0.142±0.009	0.173±0.020	0.113±0.016	0.182±0.010	0.191±0.006	0.103±0.009	0.084±0.017	**0.083±0.002**
Traffic-L	96	0.342±0.042	–	0.234±0.006	0.274±0.012	0.265±0.007	0.231±0.003	0.230±0.010	**0.198±0.001**
	192	0.354±0.012	–	0.239±0.004	–	–	0.236±0.002	0.234±0.003	**0.205±0.000**
	336	0.392±0.006	–	0.246±0.003	–	–	0.248±0.006	0.242±0.007	**0.222±0.000**
	720	0.478±0.006	–	0.263±0.001	–	–	0.264±0.006	**0.248±0.010**	0.269±0.003
Weather-L	96	0.113±0.022	0.087±0.012	0.164±0.023	0.176±0.011	0.139±0.010	0.145±0.017	0.086±0.011	**0.065±0.001**
	192	0.144±0.020	0.086±0.007	0.158±0.024	0.166±0.022	0.160±0.037	0.147±0.025	0.083±0.011	**0.072±0.001**
	336	0.138±0.033	0.098±0.002	0.162±0.006	0.168±0.014	0.170±0.027	0.160±0.012	0.093±0.010	**0.076±0.001**
	720	0.141±0.026	0.102±0.005	0.136±0.020	0.149±0.034	0.148±0.040	0.135±0.008	0.099±0.009	**0.083±0.001**
ILI-L	24	0.242±0.086	0.263±0.012	0.296±0.044	0.124±0.019	0.118±0.033	0.283±0.001	0.116±0.011	**0.089±0.011**
	36	0.246±0.117	0.298±0.011	0.298±0.048	0.144±0.011	0.143±0.089	0.307±0.007	0.142±0.008	**0.109±0.004**
	48	0.275±0.044	0.301±0.034	0.320±0.025	0.159±0.020	0.160±0.039	0.314±0.009	0.152±0.017	**0.130±0.014**
	60	0.272±0.020	0.299±0.013	0.325±0.068	0.216±0.014	0.183±0.019	0.333±0.005	0.167±0.007	**0.146±0.004**

Table 7: Full short-term probabilistic forecasting results across all baselines and datasets; protocol as in Table 2. Lower CRPS values indicate better predictions. Each cell reports the mean ± standard deviation of the CRPS across 5 independent runs of retraining and evaluation. **Red bold**: best; blue underline: second best.

Dataset	TSDiff	D^3VAE	CSDI	TimeGrad	GRU MAF	Trans MAF	GRU NVP	K^2VAE	**KGO (Ours)**
ETTm1-S	0.209±0.013	0.198±0.015	0.140±0.012	0.186±0.003	0.151±0.009	0.139±0.005	0.455±0.029	0.135±0.008	**0.104±0.002**
ETTm2-S	0.124±0.008	0.303±0.024	0.144±0.018	0.148±0.004	0.146±0.011	0.180±0.010	0.276±0.014	0.122±0.008	**0.092±0.002**
ETTh1-S	0.304±0.016	0.324±0.019	0.262±0.012	0.273±0.007	0.258±0.013	0.309±0.009	0.398±0.034	0.256±0.008	**0.228±0.001**
ETTh2-S	0.204±0.006	0.216±0.015	0.133±0.006	0.184±0.006	0.160±0.008	0.200±0.012	0.309±0.023	0.128±0.006	**0.104±0.001**
Electricity-S	0.111±0.013	0.071±0.009	**0.051±0.001**	0.057±0.002	0.056±0.002	0.054±0.002	0.062±0.003	0.053±0.002	0.052±0.000
Exchange-S	0.077±0.019	0.011±0.002	0.009±0.001	0.009±0.001	0.012±0.003	0.012±0.001	0.019±0.006	0.009±0.001	**0.007±0.000**
Traffic-S	0.189±0.009	0.143±0.008	0.147±0.014	0.130±0.005	0.144±0.022	0.133±0.004	0.168±0.008	0.129±0.004	**0.127±0.000**
Solar-S	0.568±0.015	0.769±0.029	0.392±0.006	0.465±0.016	0.486±0.007	0.442±0.011	0.530±0.008	**0.367±0.005**	0.456±0.010

Table 8: Full short-term probabilistic forecasting results across all baselines and datasets; protocol as in Table 2. Lower NMAE values indicate better predictions. Results are reported as mean ± standard deviation over 5 independent runs of retraining and evaluation. **Red bold**: best; blue underline: second best.

Dataset	TSDiff	D^3VAE	CSDI	TimeGrad	GRU MAF	Trans MAF	GRU NVP	K^2VAE	**KGO (Ours)**
ETTm1-S	0.276±0.008	0.250±0.018	0.169±0.021	0.246±0.007	0.198±0.004	0.162±0.006	0.584±0.047	0.152±0.007	**0.133±0.003**
ETTm2-S	0.162±0.008	0.378±0.031	0.181±0.024	0.189±0.006	0.193±0.008	0.224±0.009	0.349±0.028	0.146±0.009	**0.118±0.003**
ETTh1-S	0.400±0.025	0.410±0.016	0.339±0.009	0.356±0.013	0.326±0.016	0.400±0.011	0.477±0.040	0.312±0.008	**0.280±0.002**
ETTh2-S	0.272±0.015	0.267±0.018	0.161±0.013	0.224±0.014	0.208±0.003	0.256±0.009	0.375±0.024	0.140±0.007	**0.130±0.001**
Electricity-S	0.115±0.018	0.092±0.013	0.066±0.001	0.073±0.004	0.073±0.004	0.071±0.003	0.081±0.006	0.068±0.002	**0.064±0.001**
Exchange-S	0.096±0.024	0.012±0.002	0.013±0.001	0.012±0.002	0.016±0.002	0.016±0.001	0.024±0.007	**0.009±0.001**	**0.009±0.000**
Traffic-S	0.206±0.011	0.178±0.013	0.175±0.013	0.155±0.007	0.182±0.029	0.160±0.006	0.209±0.013	0.157±0.007	**0.152±0.001**
Solar-S	0.635±0.012	0.998±0.049	0.533±0.007	0.609±0.015	0.603±0.009	0.577±0.014	0.670±0.011	**0.480±0.008**	0.569±0.012

D.2 Flow Matching

This subsection details the comparison in Sec. 5.2 between KGO, using single-step closed-form KoFM transport, and pretrained Sundial-base-128m. We sweep Sundial's Euler integration budget over $N \in \{1, 5, 10, 20, 50, 100\}$ to examine its accuracy–latency trade-off relative to KGO's matrix-exponential transport. This compares complete forecasting models; their architectures and training procedures differ in addition to their sampling mechanisms.

Datasets and forecast accuracy. We evaluate seven LTSF datasets: ETTh1/2, ETTm1/2, Electricity, Weather, and Exchange. Traffic and ILI are excluded because they appear in Sundial's pretraining corpus. Both models receive a context of length 96 and generate 100 forecast trajectories for each prediction horizon $H \in \{96, 192, 336, 720\}$. CRPS and NMAE are computed on real test windows using the standard ProbTS Evaluator, with KGO results from seed 0. The latency benchmark is run separately and does not recompute these accuracy results.

Inference-time measurement. We measure latency in FP32 on a single NVIDIA RTX 6000 Ada GPU, using batch size 1 and 100 forecast trajectories per window. We report median latency after warm-up, with CUDA synchronization around each timed call. Latency is reported in seconds per window, excluding model loading and metric computation.

Table 9 reports CRPS and inference latency for all 28 dataset–horizon pairs.

Table 9: CRPS and inference latency (s/window) for KGO (KoFM) and Sundial (FM N, where N is the number of Euler steps); lower is better. See Sec. D.2 for the evaluation and timing protocol. **Red bold**: best; blue underline: second best.

		KoFM		FM 1		FM 5		FM 10		FM 20		FM 50		FM 100	
Dataset	Horizon	CRPS	Time	CRPS	Time	CRPS	Time	CRPS	Time	CRPS	Time	CRPS	Time	CRPS	Time
ETTm1-L	96	**0.2145**	**0.0030**	0.3027	0.0124	0.2677	0.0173	0.2586	0.0229	0.2546	0.0343	0.2524	0.0696	0.2517	0.1269
	192	**0.2363**	**0.0030**	0.3336	0.0127	0.2954	0.0174	0.2853	0.0235	0.2809	0.0353	0.2784	0.0704	0.2777	0.1309
	336	**0.2491**	**0.0030**	0.3480	0.0133	0.3084	0.0186	0.2979	0.0245	0.2934	0.0353	0.2909	0.0697	0.2902	0.1283
	720	**0.2747**	**0.0030**	0.3754	0.0140	0.3360	0.0185	0.3252	0.0248	0.3205	0.0361	0.3180	0.0733	0.3173	0.1315
ETTm2-L	96	**0.1064**	**0.0054**	0.1511	0.0109	0.1359	0.0157	0.1316	0.0196	0.1295	0.0335	0.1282	0.0668	0.1277	0.1167
	192	**0.1309**	**0.0054**	0.1741	0.0112	0.1584	0.0149	0.1538	0.0210	0.1515	0.0315	0.1501	0.0658	0.1496	0.1111
	336	**0.1450**	**0.0054**	0.1920	0.0112	0.1759	0.0154	0.1708	0.0230	0.1683	0.0342	0.1667	0.0743	0.1662	0.1119
	720	**0.1711**	**0.0054**	0.2145	0.0118	0.1982	0.0152	0.1930	0.0231	0.1904	0.0315	0.1887	0.0628	0.1882	0.1174
ETTh1-L	96	**0.2405**	**0.0031**	0.3031	0.0105	0.2707	0.0148	0.2653	0.0197	0.2625	0.0300	0.2609	0.0599	0.2603	0.1082
	192	**0.2649**	**0.0031**	0.3414	0.0107	0.3084	0.0145	0.3026	0.0195	0.2995	0.0294	0.2977	0.0608	0.2970	0.1080
	336	**0.2856**	**0.0031**	0.3589	0.0110	0.3250	0.0149	0.3189	0.0205	0.3157	0.0294	0.3137	0.0607	0.3130	0.1084
	720	**0.2841**	**0.0031**	0.3608	0.0114	0.3240	0.0151	0.3177	0.0205	0.3146	0.0314	0.3126	0.0610	0.3119	0.1113
ETTh2-L	96	**0.1338**	**0.0032**	0.1836	0.0124	0.1652	0.0176	0.1608	0.0231	0.1585	0.0354	0.1571	0.0741	0.1566	0.1300
	192	**0.1549**	**0.0031**	0.2094	0.0125	0.1912	0.0173	0.1867	0.0232	0.1843	0.0352	0.1827	0.0710	0.1822	0.1368
	336	**0.1721**	**0.0032**	0.2295	0.0126	0.2115	0.0175	0.2070	0.0236	0.2045	0.0349	0.2029	0.0710	0.2024	0.1320
	720	**0.1764**	**0.0032**	0.2299	0.0133	0.2113	0.0188	0.2067	0.0245	0.2041	0.0359	0.2025	0.0711	0.2019	0.1316
Elec.-L	96	**0.0666**	**0.0166**	0.0811	0.1242	0.0708	0.2728	0.0693	0.4767	0.0686	0.8964	0.0682	2.2059	0.0681	4.3803
	192	**0.0736**	**0.0168**	0.0882	0.1288	0.0775	0.2794	0.0759	0.4887	0.0751	0.9044	0.0747	2.2060	0.0745	4.3846
	336	**0.0814**	**0.0225**	0.0976	0.1427	0.0861	0.2874	0.0843	0.4912	0.0834	0.9066	0.0829	2.2133	0.0827	4.3850
	720	0.0982	**0.0494**	0.1150	0.2614	0.1021	0.3969	0.0998	0.5996	0.0987	0.9124	0.0980	2.2090	**0.0978**	4.3862
Exch.-L	96	**0.0193**	**0.0065**	0.0272	0.0129	0.0256	0.0175	0.0250	0.0225	0.0246	0.0340	0.0244	0.0677	0.0243	0.1244
	192	**0.0294**	**0.0065**	0.0374	0.0129	0.0357	0.0175	0.0350	0.0233	0.0346	0.0346	0.0343	0.0708	0.0342	0.1256
	336	**0.0407**	**0.0064**	0.0494	0.0134	0.0477	0.0178	0.0469	0.0238	0.0465	0.0350	0.0462	0.0684	0.0461	0.1267
	720	**0.0689**	**0.0065**	0.0792	0.0153	0.0776	0.0189	0.0769	0.0245	0.0765	0.0361	0.0762	0.0698	0.0761	0.1281
Weather-L	96	**0.0527**	**0.0067**	0.0975	0.0174	0.0837	0.0258	0.0800	0.0399	0.0785	0.0663	0.0778	0.1531	0.0776	0.3085
	192	**0.0583**	**0.0068**	0.1017	0.0177	0.0892	0.0266	0.0857	0.0411	0.0842	0.0690	0.0833	0.1589	0.0831	0.3138
	336	**0.0620**	**0.0068**	0.1103	0.0181	0.0970	0.0289	0.0931	0.0399	0.0913	0.0687	0.0902	0.1601	0.0899	0.3146
	720	**0.0682**	**0.0070**	0.1175	0.0197	0.1029	0.0280	0.0988	0.0400	0.0970	0.0700	0.0959	0.1593	0.0956	0.3144

D.3 Foundation Model Comparison

We compare KGO against three zero-shot time-series foundation models on the standard long-term probabilistic benchmark. KGO is trained from scratch per dataset; the foundation models (35M–128M parameters) are pre-trained on much larger corpora and used off the shelf.

Baselines. We evaluate three publicly released zero-shot models off the shelf, with no fine-tuning on the target datasets:

- **Sundial** (`thuml/sundial-base-128m`, 128M params) – univariate; the wrapper loops over the variable axis.
- **TiRex** (`NX-AI/TiRex`, 35M params) – per-series; the wrapper flattens the variable axis into the batch and decodes each series independently.
- **Chronos-2** (`amazon/chronos-2`, 120M params) – natively multivariate via long-format input with covariates.

All three are run with their published checkpoints and no fine-tuning.

Protocol and context length. Every CRPS / NMAE number in the table, including the foundation-model rows, is computed by the same ProbTS `Evaluator`. Context length is fixed to 96 for every model so the comparison sits inside the standard long-term forecasting protocol used elsewhere in the paper. The foundation models natively support much longer contexts (2k–8k), so the table reflects their behavior at the benchmark's input length, not their best zero-shot setting.

Sampling and quantile handling. All three models are scored by the same ProbTS evaluator at the same 19 quantile levels in $[0.05, 0.95]$, so any gap between them reflects model behavior rather than evaluation drift. Sundial's 100 stochastic samples are passed in directly; TiRex and Chronos-2 emit quantile grids (9 and 19 levels respectively), which we transport to the evaluator's sample-based interface by piecewise linear inverse-CDF interpolation. This is lossless for Chronos-2 at every scored level. TiRex emits only $[0.1, 0.9]$, so the evaluator's $q_{0.05}$ and $q_{0.95}$ queries fall back to TiRex's nearest emitted quantile; of the other 17 levels, 9 lie on TiRex's grid and 8 are linearly interpolated between grid levels. We keep all 19 in order to make a fair comparison across models.

Datasets, horizons, and seed. We evaluate on ETTh1, ETTh2, ETTm1, ETTm2, Exchange, Weather, and (Sundial only) Electricity at horizons $\{96, 192, 336, 720\}$ with a single seed (0); Traffic, ILI, and TiRex/Chronos-2 on Electricity are excluded as pretraining leakage.

Table 10: Multivariate probabilistic time series forecasting results. **Red bold**: best; blue underline: second best. "–" = unavailable or omitted because the foundation model was pretrained on that dataset (data-leakage avoidance).

		KGO		Sundial		TiRex		Chronos-2	
Dataset	Horizon	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTm1-L	96	**0.2145**	**0.2690**	0.2532	0.3080	0.2876	0.3467	0.2726	0.3449
	192	**0.2363**	**0.2964**	0.2788	0.3368	0.2998	0.3630	0.2876	0.3641
	336	**0.2491**	**0.3158**	0.2913	0.3503	0.3115	0.3786	0.2925	0.3684
	720	**0.2747**	**0.3469**	0.3183	0.3799	0.3531	0.4375	0.3113	0.3873
ETTm2-L	96	**0.1064**	**0.1343**	0.1284	0.1535	0.1264	0.1540	0.1147	0.1453
	192	**0.1309**	**0.1638**	0.1500	0.1769	0.1440	0.1784	0.1357	0.1690
	336	**0.1450**	**0.1824**	0.1667	0.1949	0.1569	0.1956	0.1558	0.1924
	720	**0.1711**	**0.2142**	0.1887	0.2174	0.1784	0.2215	0.1805	0.2172
ETTh1-L	96	**0.2405**	**0.3077**	0.2609	0.3087	0.2477	0.3185	0.2421	0.3149
	192	**0.2649**	**0.3396**	0.2977	0.3470	0.2803	0.3609	0.2720	0.3546
	336	**0.2856**	**0.3629**	0.3137	0.3647	0.3041	0.3891	0.2891	0.3750
	720	**0.2841**	0.3694	0.3126	**0.3667**	0.3302	0.4331	0.2938	0.3788
ETTh2-L	96	**0.1338**	**0.1703**	0.1571	0.1862	0.1427	0.1838	0.1411	0.1824
	192	**0.1549**	**0.1968**	0.1826	0.2126	0.1638	0.2096	0.1621	0.2093
	336	**0.1721**	**0.2179**	0.2027	0.2325	0.1819	0.2316	0.1815	0.2320
	720	**0.1764**	**0.2265**	0.2025	0.2326	0.1864	0.2383	0.1891	0.2397
Elec.-L	96	**0.0666**	**0.0783**	0.0681	0.0831	–	–	–	–
	192	**0.0736**	**0.0864**	0.0745	0.0905	–	–	–	–
	336	**0.0814**	**0.0961**	0.0828	0.1001	–	–	–	–
	720	0.0982	**0.1168**	**0.0978**	0.1177	–	–	–	–
Exch.-L	96	0.0193	0.0242	0.0243	0.0274	0.0195	0.0247	**0.0184**	**0.0236**
	192	0.0294	0.0367	0.0344	0.0377	0.0282	0.0348	**0.0268**	**0.0338**
	336	0.0407	0.0498	0.0462	0.0497	0.0392	0.0469	**0.0372**	**0.0456**
	720	0.0689	0.0821	0.0763	0.0794	0.0684	0.0761	**0.0640**	**0.0751**
Weather-L	96	**0.0527**	**0.0655**	0.0783	0.0913	0.1176	0.1457	0.0977	0.1226
	192	**0.0583**	**0.0725**	0.0834	0.0973	0.1149	0.1386	0.0962	0.1170
	336	**0.0620**	**0.0769**	0.0902	0.1059	0.1040	0.1235	0.0895	0.1071
	720	**0.0682**	**0.0839**	0.0960	0.1127	0.1003	0.1168	0.0972	0.1125

D.4 Efficiency

Measurement protocol. We benchmark all methods on a single NVIDIA RTX 6000 Ada GPU in FP32 in eager mode (without compilation), using batch size 1 and generating 100 forecast trajectories per input window. All latencies are measured after warm-up, with CUDA synchronization around each timed forward pass; we report the mean latency for baselines and the median for KGO. Peak GPU memory is the maximum allocated memory, reported in GB. Each baseline uses its native sampler; KGO reuses sample-independent conditioning computations and generates trajectories in chunks of five on Electricity and Traffic to limit peak memory.

Table 11: Inference latency (s/window) and peak GPU memory (GB) across nine datasets; lower is better. See Sec. D.4 for the measurement protocol. **Red bold**: best; blue underline: second best.

		TSDiff		CSDI		TimeGrad		GRU MAF		Trans MAF		GRU NVP		K^2VAE		**KGO (Ours)**	
Dataset	Horizon	Time	Mem	Time	Mem	Time	Mem	Time	Mem	Time	Mem	Time	Mem	Time	Mem	**Time**	**Mem**
ETTm1-L	96	0.454	0.092	24.969	0.018	31.406	0.015	0.685	0.022	0.734	0.079	0.222	**0.013**	0.018	0.015	**0.0030**	0.019
	192	0.454	0.132	24.920	0.022	63.597	0.016	1.265	0.022	1.553	0.080	0.456	**0.014**	0.022	0.017	**0.0030**	0.028
	336	0.451	0.198	25.193	0.027	111.064	0.018	2.266	0.024	2.616	0.082	0.795	**0.016**	0.027	0.021	**0.0030**	0.039
	720	0.484	0.357	24.062	0.040	254.576	**0.021**	4.950	0.029	5.804	0.087	1.776	**0.021**	0.041	0.031	**0.0030**	0.070
ETTm2-L	96	0.462	0.092	25.103	0.018	34.132	**0.012**	0.684	0.014	0.731	0.040	0.153	0.014	0.018	0.015	**0.0054**	0.019
	192	0.451	0.132	24.991	0.022	67.229	**0.013**	1.403	0.015	1.586	0.041	0.300	0.015	0.022	0.017	**0.0054**	0.028
	336	0.484	0.198	25.976	0.027	119.362	**0.015**	2.228	0.017	2.828	0.043	0.495	0.017	0.027	0.021	**0.0054**	0.038
	720	0.485	0.357	24.905	0.040	249.600	**0.020**	5.065	0.022	5.920	0.048	1.152	0.022	0.044	0.031	**0.0054**	0.069
ETTh1-L	96	0.482	0.092	25.048	0.018	33.937	0.015	0.635	0.015	0.735	0.041	0.225	**0.013**	0.018	0.015	**0.0031**	0.019
	192	0.453	0.132	24.012	0.022	66.838	0.017	1.384	0.016	1.600	0.042	0.481	**0.015**	0.022	0.017	**0.0031**	0.028
	336	0.486	0.198	24.674	0.027	118.148	0.018	2.363	0.018	2.819	0.043	0.793	**0.016**	0.027	0.021	**0.0031**	0.039
	720	0.483	0.357	25.638	0.040	256.945	**0.023**	5.227	**0.023**	5.835	0.047	1.741	**0.023**	0.043	0.031	**0.0031**	0.070
ETTh2-L	96	0.453	0.092	25.116	0.018	34.045	**0.013**	0.692	0.015	0.728	0.041	0.160	0.014	0.018	0.015	**0.0032**	0.022
	192	0.452	0.132	25.605	0.022	67.710	**0.014**	1.278	0.016	1.465	0.042	0.321	0.015	0.022	0.017	**0.0031**	0.031
	336	0.451	0.198	25.886	0.027	118.779	**0.016**	2.429	0.018	2.723	0.043	0.565	0.017	0.027	0.021	**0.0032**	0.042
	720	0.485	0.357	25.893	0.040	256.475	**0.022**	5.154	0.023	6.004	0.047	1.224	0.023	0.041	0.031	**0.0032**	0.073
Elec.-L	96	0.489	0.238	96.645	0.296	33.504	0.132	25.746	0.134	26.389	0.165	0.125	0.132	0.019	**0.067**	**0.0166**	0.117
	192	0.537	0.352	165.237	0.435	66.456	0.182	53.662	0.184	54.411	0.215	0.251	0.182	0.022	**0.099**	**0.0168**	0.176
	336	0.643	0.513	286.608	0.647	119.140	0.258	94.409	0.261	94.014	0.292	0.477	0.259	0.029	**0.145**	**0.0225**	0.271
	720	0.986	0.960	712.445	1.212	248.970	0.455	202.201	0.458	196.062	0.491	0.942	0.456	**0.040**	**0.268**	0.0494	0.523
Exch.-L	96	0.443	0.092	24.805	0.019	33.105	**0.013**	0.776	0.014	0.805	0.040	0.134	0.015	0.018	0.015	**0.0065**	0.020
	192	0.445	0.133	25.398	0.023	66.895	**0.014**	1.506	0.015	1.721	0.041	0.269	0.016	0.021	0.017	**0.0065**	0.025
	336	0.474	0.199	25.458	0.030	116.148	**0.017**	2.702	0.018	3.015	0.043	0.473	0.018	0.027	0.021	**0.0064**	0.034
	720	0.480	0.358	25.551	0.043	250.412	**0.021**	5.723	0.022	6.258	0.048	1.087	0.023	0.041	0.031	**0.0065**	0.056
Traffic-L	96	0.653	0.619	386.904	0.772	33.289	0.336	67.177	0.336	67.180	0.342	0.273	0.343	**0.019**	**0.159**	0.0225	0.223
	192	0.824	0.910	650.234	1.150	66.347	0.469	134.161	0.470	133.921	0.476	0.551	0.477	**0.022**	**0.241**	0.0229	0.264
	336	1.093	1.359	1057.751	1.717	114.615	0.670	234.524	0.670	234.948	0.677	1.055	0.677	**0.027**	**0.361**	0.0329	0.429
	720	1.852	2.558	2371.636	3.236	250.309	1.204	507.890	1.203	508.772	1.212	2.300	1.211	**0.042**	**0.688**	0.0731	0.852
Weather-L	96	0.479	0.097	25.588	0.032	33.777	**0.016**	1.741	0.025	1.960	0.081	0.226	0.018	0.018	0.017	**0.0067**	0.037
	192	0.446	0.139	24.682	0.041	67.780	**0.020**	3.710	0.027	3.935	0.083	0.456	0.022	0.023	**0.020**	**0.0068**	0.052
	336	0.446	0.208	24.927	0.056	117.750	**0.025**	6.603	0.031	6.934	0.086	0.829	0.028	0.027	**0.025**	**0.0068**	0.072
	720	0.488	0.376	36.440	0.096	254.377	0.039	13.993	0.046	13.877	0.096	1.713	0.041	0.041	**0.037**	**0.0070**	0.126
ILI-L	24	0.447	0.036	25.964	0.013	7.779	**0.011**	0.156	0.014	0.180	0.024	0.050	0.014	0.018	0.017	**0.0043**	0.016
	36	0.441	0.041	24.946	0.014	12.765	**0.011**	0.234	0.014	0.276	0.024	0.076	0.014	0.020	0.018	**0.0043**	0.018
	48	0.467	0.047	25.527	0.014	16.660	**0.011**	0.314	0.014	0.362	0.024	0.100	0.014	0.022	0.018	**0.0043**	0.020
	60	0.443	0.051	24.585	0.015	21.114	**0.012**	0.431	0.014	0.470	0.024	0.126	0.014	0.023	0.020	**0.0043**	0.022

E Ablation Studies

E.1 Patching (KoPE)

In this ablation we examine the KoPE component. To check its importance, we compare this component against alternative techniques such as using the last patch representation and pooling the full context. To isolate the representation alone, we keep all other components the same. We evaluate on a grid of 36 dataset–horizon settings using seed 0, reporting CRPS and NMAE with the standard ProbTS evaluator.

Table 12: Patch representation ablation. Multivariate probabilistic forecasting results. **Red bold**: best; blue underline: second best.

		KoPE		Last Patch		Context Pooling	
Dataset	Horizon	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTm1-L	96	**0.2145**	**0.2690**	0.3127	0.3925	0.3119	0.3938
	192	**0.2363**	**0.2964**	0.3263	0.4107	0.3255	0.4093
	336	**0.2491**	**0.3158**	0.3371	0.4236	0.3351	0.4173
	720	**0.2747**	**0.3469**	0.3563	0.4412	0.3558	0.4400
ETTm2-L	96	**0.1064**	**0.1343**	0.1242	0.1602	0.1250	0.1611
	192	**0.1309**	**0.1638**	0.1454	0.1849	0.1438	0.1836
	336	**0.1450**	**0.1824**	0.1565	0.1990	0.1554	0.1972
	720	**0.1711**	**0.2142**	0.1783	0.2234	0.1780	0.2239
ETTh1-L	96	**0.2405**	**0.3077**	0.2413	0.3105	0.2442	0.3167
	192	**0.2649**	0.3396	0.2652	**0.3389**	0.2712	0.3478
	336	0.2856	0.3629	**0.2842**	**0.3595**	0.2932	0.3706
	720	**0.2841**	**0.3694**	0.2992	0.3887	0.2959	0.3884
ETTh2-L	96	**0.1338**	**0.1703**	0.1369	0.1743	0.1376	0.1752
	192	0.1549	0.1968	0.1538	0.1947	**0.1529**	**0.1932**
	336	**0.1721**	**0.2179**	0.1724	0.2184	0.1748	0.2240
	720	0.1764	0.2265	**0.1740**	**0.2228**	0.1771	0.2280
Elec.-L	96	**0.0666**	**0.0783**	0.0687	0.0817	0.0681	0.0816
	192	**0.0736**	**0.0864**	0.0764	0.0899	0.0758	0.0895
	336	**0.0814**	**0.0961**	0.0840	0.0992	0.0830	0.0991
	720	**0.0982**	0.1168	0.0989	**0.1165**	0.1013	0.1167
Exch.-L	96	0.0193	0.0242	**0.0191**	0.0238	**0.0191**	**0.0237**
	192	0.0294	0.0367	**0.0286**	**0.0350**	0.0292	0.0359
	336	0.0407	0.0498	**0.0393**	**0.0482**	0.0410	0.0498
	720	0.0689	**0.0821**	**0.0680**	0.0822	0.0683	**0.0821**
Traffic-L	96	**0.1732**	**0.1988**	0.2062	0.2464	0.2050	0.2438
	192	**0.1777**	**0.2048**	0.2283	0.2797	0.2268	0.2759
	336	**0.1921**	**0.2217**	0.2350	0.2860	0.2318	0.2811
	720	**0.2274**	**0.2671**	0.2490	0.3025	0.2461	0.3003
Weather-L	96	**0.0527**	**0.0655**	0.0679	0.0828	0.0680	0.0834
	192	**0.0583**	**0.0725**	0.0791	0.0966	0.0795	0.0975
	336	**0.0620**	**0.0769**	0.0830	0.1010	0.0831	0.1008
	720	**0.0682**	**0.0839**	0.0881	0.1073	0.0878	0.1054
ILI-L	24	**0.0531**	**0.0750**	0.1252	0.1609	0.1222	0.1600
	36	**0.0821**	**0.1024**	0.1492	0.1933	0.1490	0.1925
	48	**0.0991**	**0.1191**	0.1453	0.2012	0.1422	0.1978
	60	**0.1206**	**0.1423**	0.1435	0.1919	0.1336	0.1790

E.2 Sequence Encoder

In this subsection, we investigate the Sequence Encoder that produces the patch representation. This representation is propagated through KoPE; beyond memory and time efficiency, it is important to model the temporal dynamics well. We therefore compare three cornerstone sequential encoders: Transformer, Linear RNN, and sLSTM. To isolate the representation alone, we keep all other components the same. We evaluate on a grid of 36 dataset–horizon settings using seed 0, reporting CRPS and NMAE with the standard ProbTS evaluator.

Table 13: Sequence encoder ablation. KGO with different sequential encoders (sLSTM is the KGO baseline). **Red bold**: best; blue underline: second best.

		sLSTM		Transformer		Linear RNN	
Dataset	Horizon	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTm1-L	96	**0.214**	**0.269**	0.237	0.297	0.220	0.275
	192	**0.236**	**0.296**	0.265	0.333	0.238	0.299
	336	**0.249**	**0.316**	0.276	0.349	0.253	0.319
	720	**0.275**	**0.347**	0.301	0.381	0.278	0.349
ETTm2-L	96	0.106	0.134	0.106	0.134	**0.105**	**0.132**
	192	0.131	0.164	0.129	0.161	**0.127**	**0.158**
	336	0.145	0.182	0.144	0.181	**0.142**	**0.177**
	720	0.171	0.214	0.168	0.210	**0.166**	**0.207**
ETTh1-L	96	0.240	0.308	**0.238**	**0.307**	0.241	0.310
	192	**0.265**	**0.340**	0.267	0.341	0.267	0.341
	336	0.286	0.363	0.288	0.366	**0.281**	**0.358**
	720	**0.284**	**0.369**	0.297	0.382	0.287	0.372
ETTh2-L	96	**0.134**	**0.170**	**0.134**	0.171	0.135	0.172
	192	0.155	0.197	**0.154**	**0.196**	0.155	0.198
	336	0.172	0.218	**0.171**	**0.217**	0.174	0.221
	720	0.176	0.226	**0.175**	**0.225**	0.177	0.227
Elec.-L	96	**0.067**	**0.078**	**0.067**	0.079	**0.067**	0.080
	192	**0.074**	**0.086**	**0.074**	0.087	0.075	0.088
	336	**0.081**	**0.096**	0.082	0.097	0.082	0.097
	720	0.098	0.117	**0.097**	**0.116**	0.099	0.118
Exch.-L	96	**0.019**	**0.024**	**0.019**	**0.024**	0.020	**0.024**
	192	0.029	0.037	**0.027**	**0.034**	0.029	0.036
	336	0.041	0.050	0.040	0.050	**0.039**	**0.047**
	720	0.069	0.082	0.065	**0.078**	**0.064**	0.079
Traffic-L	96	**0.173**	0.199	**0.173**	**0.197**	0.178	0.205
	192	**0.178**	0.205	**0.178**	**0.204**	0.182	0.210
	336	**0.192**	**0.222**	0.200	0.231	0.199	0.230
	720	**0.227**	**0.267**	0.313	0.366	0.262	0.307
Weather-L	96	**0.053**	**0.066**	0.054	0.068	**0.053**	**0.066**
	192	**0.058**	**0.073**	0.060	0.075	**0.058**	**0.073**
	336	**0.062**	**0.077**	0.063	0.078	**0.062**	0.078
	720	**0.068**	**0.084**	0.069	0.085	0.069	0.085
ILI-L	24	**0.053**	**0.075**	0.072	0.093	0.099	0.124
	36	0.082	0.102	**0.075**	**0.098**	0.100	0.125
	48	**0.099**	**0.119**	0.118	0.143	0.117	0.139
	60	**0.121**	**0.142**	0.129	0.151	0.132	0.153

E.3 Decoder & Gating

In this ablation we examine the coupling between the deterministic and probabilistic decoders. The Adaptive Uncertainty Gate (AUG) modulates their contributions per (variate, horizon) based on the predictability of the signal, allocating uncertainty dynamically rather than with a fixed split. To check the importance of this learned gating, we compare AUG against three alternatives: removing the deterministic stream entirely, conditioning KoFM on the deterministic prediction, and a fixed residual split. To isolate the coupling rule alone, we keep all other components the same. We evaluate on a grid of 36 dataset–horizon settings using seed 0, reporting CRPS and NMAE with the standard ProbTS evaluator.

Table 14: Decoder and gating ablation. KGO is compared with a stochastic-only decoder (Prob Only), KoFM conditioned on the deterministic prediction (Condition), and fixed residual addition (Residual). Lower CRPS/NMAE is better. **Red bold**: best; blue underline: second best.

		KGO		Prob Only		Condition		Residual	
Dataset	Horizon	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTm1-L	96	0.2145	0.2690	0.2150	0.2695	**0.2129**	0.2675	0.2158	**0.2642**
	192	0.2363	0.2964	**0.2342**	0.2955	0.2349	0.2945	0.2382	**0.2922**
	336	**0.2491**	**0.3158**	**0.2491**	0.3159	0.2511	0.3170	0.2568	0.3166
	720	**0.2747**	**0.3469**	0.2766	0.3490	0.2765	0.3493	0.2831	0.3481
ETTm2-L	96	0.1064	0.1343	**0.1047**	0.1321	0.1060	0.1333	0.1058	**0.1320**
	192	0.1309	0.1638	0.1306	0.1637	**0.1292**	0.1615	0.1298	**0.1610**
	336	0.1450	0.1824	0.1449	0.1823	0.1520	0.1902	**0.1440**	**0.1789**
	720	0.1711	0.2142	**0.1679**	0.2092	0.1691	0.2106	0.1685	**0.2082**
ETTh1-L	96	**0.2405**	0.3077	0.2434	0.3112	0.2460	0.3139	0.2454	**0.3069**
	192	**0.2649**	**0.3396**	0.2654	0.3410	0.2682	0.3414	0.2784	0.3447
	336	0.2856	0.3629	**0.2821**	0.3596	0.2857	0.3657	0.2849	**0.3572**
	720	**0.2841**	**0.3694**	0.2910	0.3781	0.2870	0.3744	0.2917	**0.3694**
ETTh2-L	96	0.1338	**0.1703**	0.1345	0.1716	0.1373	0.1746	**0.1334**	0.1713
	192	0.1549	0.1968	0.1552	0.1970	**0.1539**	0.1961	0.1563	**0.1940**
	336	**0.1721**	**0.2179**	0.1765	0.2250	0.1741	0.2210	0.1745	0.2190
	720	**0.1764**	**0.2265**	0.1804	0.2326	0.1802	0.2344	0.1819	0.2343
Elec.-L	96	0.0666	**0.0783**	0.0664	0.0786	0.0670	0.0787	**0.0658**	0.0784
	192	0.0736	0.0864	0.0739	0.0865	0.0741	0.0869	**0.0733**	**0.0862**
	336	0.0814	**0.0961**	0.0816	0.0969	0.0825	0.0968	**0.0806**	0.0962
	720	0.0982	0.1168	0.0998	0.1192	0.1025	0.1221	**0.0963**	**0.1153**
Exch.-L	96	0.0193	0.0242	0.0194	**0.0241**	0.0201	0.0250	**0.0190**	0.0242
	192	0.0294	0.0367	0.0290	0.0352	0.0285	0.0360	**0.0271**	**0.0344**
	336	0.0407	0.0498	0.0406	0.0514	**0.0378**	**0.0471**	0.0407	0.0515
	720	0.0689	0.0821	0.0658	0.0797	0.0676	0.0824	**0.0617**	**0.0778**
Traffic-L	96	0.1732	0.1988	0.1721	0.1979	**0.1717**	0.1970	0.1719	**0.1957**
	192	0.1777	0.2048	0.1782	0.2044	0.1790	0.2042	**0.1768**	**0.2006**
	336	0.1921	0.2217	0.1934	0.2225	0.1964	0.2235	**0.1898**	**0.2160**
	720	0.2274	0.2671	0.2302	0.2694	0.2617	0.3039	**0.2202**	**0.2515**
Weather-L	96	0.0527	0.0655	0.0533	0.0661	**0.0518**	**0.0641**	0.0531	0.0643
	192	0.0583	0.0725	0.0582	0.0724	**0.0579**	**0.0720**	0.0592	**0.0720**
	336	0.0620	0.0769	0.0617	0.0764	**0.0606**	**0.0749**	0.0623	**0.0749**
	720	0.0682	0.0839	0.0677	0.0835	**0.0640**	**0.0793**	0.0702	0.0844
ILI-L	24	**0.0531**	**0.0750**	0.0768	0.1026	0.0942	0.1191	0.0680	0.0964
	36	0.0821	0.1024	0.0959	0.1233	0.0932	0.1236	**0.0742**	**0.0973**
	48	0.0991	**0.1191**	0.1146	0.1410	0.1032	0.1293	**0.0958**	0.1215
	60	0.1206	0.1423	0.1227	0.1479	0.1161	0.1412	**0.1122**	**0.1408**

E.4 Generative Decoder

In this ablation we examine the generative head that produces the probabilistic forecasts. KoFM models the predictive distribution via a closed-form Koopman push-forward in a learned latent space, replacing the iterative sampling of standard generative paradigms. To check its importance, we compare KoFM against two alternative stochastic heads: a conditional VAE and a conditional normalizing flow (RealNVP). To isolate the generative paradigm alone, each head is used on its own, without AUG or the deterministic decoder, and all three consume the same KoPE latent representation $\tilde{\mathbf{z}}_{N+i}$ as their conditioning input; the encoder, KoPE, and training budget are kept identical. KoFM is trained with the KGO objective, while the VAE and NF heads are trained only with the loss native to their paradigm (ELBO for the VAE, exact log-likelihood for the NF), since enforcing a shared loss would not be well-defined across these stochastic decoders. We evaluate on a grid of 36 dataset–horizon settings using seed 0, reporting CRPS and NMAE with the standard ProbTS evaluator.

Table 15: Generative decoder ablation. Multivariate probabilistic forecasting results. **Red bold**: best; blue underline: second best.

		KoFM		VAE		NF	
Dataset	Horizon	CRPS	NMAE	CRPS	NMAE	CRPS	NMAE
ETTm1-L	96	**0.2150**	**0.2695**	0.2205	0.2818	0.2206	0.2752
	192	**0.2342**	**0.2955**	0.2412	0.3081	0.2370	0.2997
	336	**0.2491**	**0.3159**	0.2632	0.3324	0.2689	0.3329
	720	**0.2766**	**0.3490**	0.3008	0.3798	0.3223	0.3979
ETTm2-L	96	**0.1047**	**0.1321**	0.1098	0.1377	0.1077	0.1357
	192	0.1306	0.1637	**0.1292**	**0.1634**	0.1306	0.1642
	336	**0.1449**	**0.1823**	0.1500	0.1870	0.1491	0.1864
	720	**0.1679**	**0.2092**	0.1724	0.2135	0.1792	0.2245
ETTh1-L	96	0.2434	**0.3112**	0.2445	0.3139	**0.2395**	0.3124
	192	0.2654	0.3410	0.2712	0.3463	**0.2647**	**0.3392**
	336	**0.2821**	**0.3596**	0.2890	0.3670	0.2880	0.3706
	720	**0.2910**	**0.3781**	0.3170	0.3950	0.2954	0.3833
ETTh2-L	96	**0.1345**	**0.1716**	0.1390	0.1761	0.1431	0.1766
	192	**0.1552**	0.1970	0.1555	**0.1967**	0.1572	0.2007
	336	**0.1765**	**0.2250**	0.1789	0.2259	0.1835	0.2254
	720	**0.1804**	**0.2326**	0.1993	0.2370	0.1857	0.2378
Elec.-L	96	0.0664	**0.0786**	0.0678	0.0870	**0.0656**	0.0835
	192	**0.0739**	**0.0865**	0.0751	0.0965	0.0747	0.0951
	336	**0.0816**	**0.0969**	0.0854	0.1104	0.0841	0.1076
	720	**0.0998**	**0.1192**	0.1074	0.1398	0.1037	0.1332
Exch.-L	96	0.0194	0.0241	**0.0191**	**0.0240**	0.0194	0.0246
	192	0.0290	**0.0352**	0.0291	0.0359	**0.0276**	0.0354
	336	**0.0406**	**0.0514**	0.0422	0.0535	0.0518	0.0592
	720	**0.0658**	**0.0797**	0.0719	0.0891	0.0755	0.0856
Traffic-L	96	**0.1721**	**0.1979**	0.1966	0.2479	0.2024	0.2528
	192	**0.1782**	**0.2044**	0.2078	0.2628	0.2065	0.2579
	336	**0.1934**	**0.2225**	0.2361	0.3017	0.2699	0.3413
	720	**0.2302**	**0.2694**	0.3364	0.4361	0.3325	0.4181
Weather-L	96	**0.0533**	**0.0661**	0.0558	0.0680	0.0601	0.0723
	192	**0.0582**	**0.0724**	0.0628	0.0754	0.0991	0.1158
	336	**0.0617**	**0.0764**	0.0712	0.0845	0.0787	0.0970
	720	**0.0677**	**0.0835**	0.0787	0.0974	0.0993	0.1121
ILI-L	24	0.0768	0.1026	**0.0735**	**0.0985**	0.0956	0.1333
	36	0.0959	0.1233	0.0964	0.1289	**0.0699**	**0.0992**
	48	0.1146	0.1410	**0.1075**	**0.1327**	0.1207	0.1568
	60	0.1227	**0.1479**	**0.1221**	0.1499	0.1256	0.1632

E.5 Loss Function

The L_1 accuracy term and the CRPS quantile loss $\rho_q(u) = u(q - \mathbb{I}[u < 0])$ over levels $q \in Q$ are by now standard ingredients of forecasting objectives: the L_1 term penalizes deviations of forecast samples from the observed target, while the quantile (pinball) term directly supervises the empirical quantiles $\hat{X}_q$ across probability levels, helping learn the location and spread of the predictive distribution. In this ablation we isolate the contribution of the *Energy Score* repulsion term $\mathbb{E}\,\|\hat{X} - \hat{X}'\|$, where $\hat{X}$ and $\hat{X}'$ are independent samples from the model (in practice $\hat{X}'$ is obtained by shuffling the samples in $\hat{X}$). This term is meaningful at training time because KoFM produces per-sample trajectories: each instance is replicated along a sample dimension and each replica draws its own flow time and noise, so the pairwise repulsion expectation $\mathbb{E}\|\hat{X} - \hat{X}'\|$ is a well-defined training signal that discourages sample collapse to a single point prediction. We measure the impact of this term by comparing the full loss against the loss with the $\mathbb{E}\|\hat{X} - \hat{X}'\|$ term removed, reporting CRPS and NMAE per (dataset, horizon). Across eight datasets (excluding ILI) and four horizons per dataset, using seed 0, including the repulsion term reduces mean CRPS by 5.5%, with the largest reduction of 10.4% on ETTh1 at horizon 720. The effect is consistent across horizons, with mean CRPS reductions of 5.3%–5.7% at each horizon. Removing the term slightly improves NMAE, indicating a modest trade-off between probabilistic and point-forecasting performance in this ablation. Overall, this cost is small relative to the CRPS gain.

Table 16: Loss-function ablation. KGO trained with the Repulsion term removed (Full is the KGO baseline). **Red bold**: best; blue underline: second best.

		Full		w/o Repulsion	
Dataset	Horizon	CRPS	NMAE	CRPS	NMAE
ETTm1-L	96	**0.214**	0.269	0.226	**0.262**
	192	**0.236**	0.296	0.254	**0.293**
	336	**0.249**	0.316	0.267	**0.312**
	720	**0.275**	0.347	0.293	**0.343**
ETTm2-L	96	**0.106**	0.134	0.110	**0.131**
	192	**0.131**	0.164	0.134	**0.158**
	336	**0.145**	0.182	0.152	**0.181**
	720	**0.171**	0.214	0.174	**0.207**
ETTh1-L	96	**0.240**	**0.308**	0.258	**0.308**
	192	**0.265**	**0.340**	0.286	0.346
	336	**0.286**	**0.363**	0.308	0.368
	720	**0.284**	**0.369**	0.317	0.380
ETTh2-L	96	**0.134**	**0.170**	0.139	0.173
	192	**0.155**	0.197	0.162	**0.196**
	336	**0.172**	**0.218**	0.182	**0.218**
	720	**0.176**	**0.226**	0.184	0.228
Elec.-L	96	**0.067**	**0.078**	0.071	**0.078**
	192	**0.074**	**0.086**	0.078	**0.086**
	336	**0.081**	0.096	0.087	**0.095**
	720	**0.098**	0.117	0.104	**0.115**
Exch.-L	96	**0.019**	**0.024**	0.021	0.025
	192	**0.029**	0.037	0.030	**0.035**
	336	**0.041**	0.050	0.043	**0.049**
	720	**0.069**	0.082	0.070	**0.079**
Traffic-L	96	**0.173**	0.199	0.180	**0.196**
	192	**0.178**	0.205	0.186	**0.202**
	336	**0.192**	0.222	0.200	**0.219**
	720	**0.227**	0.267	0.235	**0.262**
Weather-L	96	**0.053**	0.066	0.057	**0.065**
	192	**0.058**	0.073	0.062	**0.072**
	336	**0.062**	0.077	0.065	**0.075**
	720	**0.068**	0.084	0.071	**0.082**

E.6 Ablation Summary

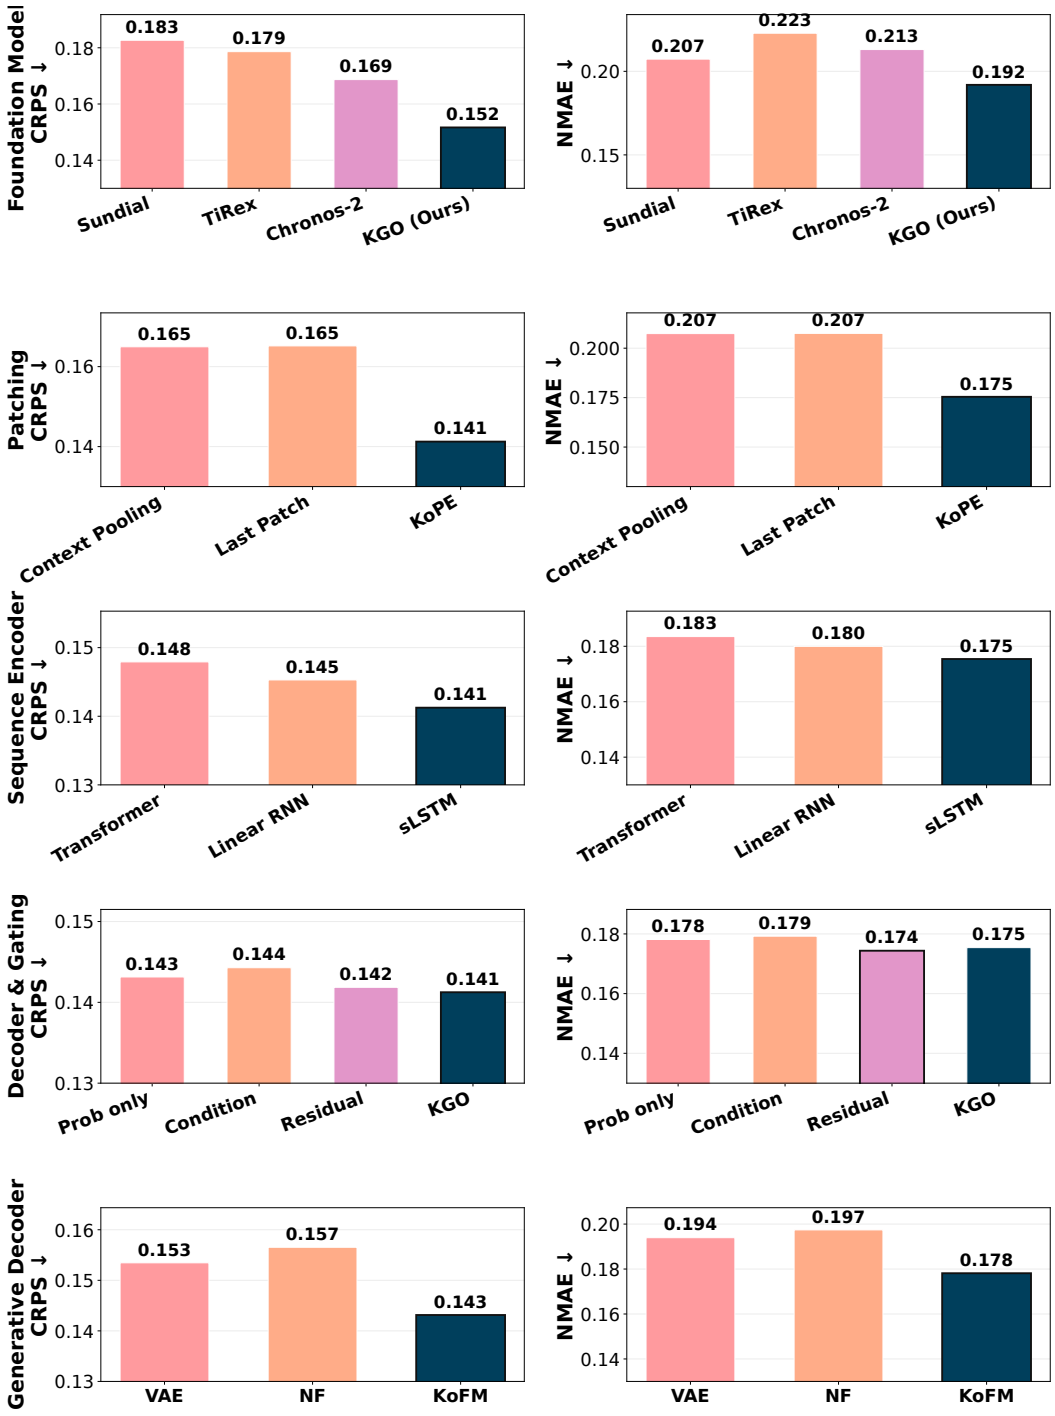

Figure 5: Combined ablation summary across all five categories. Rows correspond to ablations (Foundation Models, Patching, Sequence Encoder, Decoder & Gating, Generative Decoder); columns show CRPS (left) and NMAE (right). Each subplot averages over all 36 (dataset, horizon) cells with seed 0; the Foundation Models row averages over the 24 cells where all four models have results (ILI, Traffic, and Electricity excluded). The best value in each subplot is highlighted with a darker edge.

F Interpretability & Analysis

F.1 Sampling Convergence

In this subsection, we examine the number of Monte Carlo samples required for the metrics to converge. To quantify this, we take the checkpoints of the trained models and run inference on the ProbTS benchmark while varying the number of samples. We sweep $S \in \{1, 5, 10, 20, 50, 100\}$ on the four ETT datasets at horizon $H = 96$ (seed 0) and report CRPS and NMAE.

Both curves drop sharply from $S = 1$ to $S \approx 10$ and then plateau: by $S = 20$, CRPS is within 0.005 of its $S = 100$ value, and NMAE is also within 0.005. On these datasets, using 20 samples retains near-100-sample accuracy while generating five times fewer trajectories, offering a practical option under memory constraints.

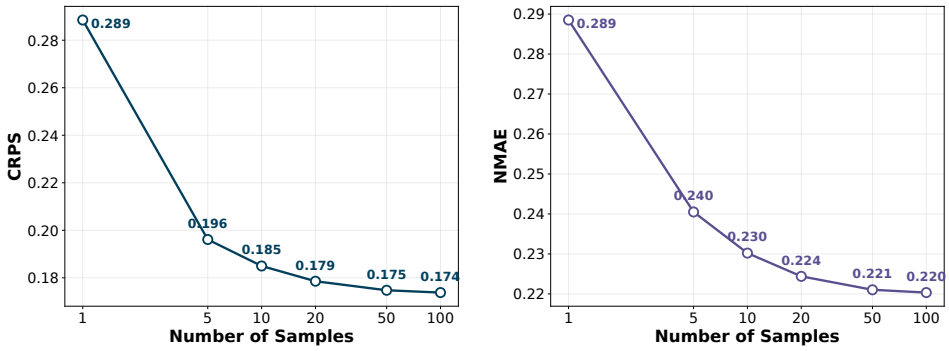

Figure 6: Effect of sample count S on CRPS (left) and NMAE (right), averaged over the four ETT datasets at $H = 96$.

F.2 Spectral and Stability Behavior

F.2.1 Spectral Analysis

We examine the eigenvalue spectrum of the learned KoPE operator K across nine datasets and four forecast horizons per dataset. Complementing the spectrum visualization in Fig. 7, Table 17 reports the mean eigenvalue magnitude and the spectral radius $\rho(K) = \max_j |\lambda_j|$. Each statistic is averaged over the evaluated horizons.

Table 17: Spectral statistics of the learned KoPE operator, averaged over horizons $\{96, 192, 336, 720\}$, or $\{24, 36, 48, 60\}$ for ILI, using seed 0.

Dataset	Mean $\lvert\lambda\rvert$	$\rho(K)$
ETTh1	0.607	0.87
ETTh2	0.659	0.87
ETTm1	0.622	0.79
ETTm2	0.667	0.84
Electricity	0.490	0.78
Traffic	0.392	0.80
Exchange	0.465	0.83
Weather	0.570	0.74
ILI	0.655	0.77

All eigenvalues lie inside the unit circle: the dataset-level averages of eigenvalue magnitude range from 0.392 to 0.667, and $\rho(K) < 1$ in all 36 dataset–horizon configurations, ranging from 0.59 to 0.96. The horizon-averaged spectral radius ranges from 0.74 to 0.87. The bound $\rho(K) \leq \|K\|_2 \leq 1$ follows from the spectral-norm normalization of K (Sec. 3.4); in all evaluated configurations the learned operators are strictly contracting in spectrum and therefore asymptotically stable.

The largest-magnitude eigenvalues (the top 2%, pooled over horizons) are not scattered at random: their phase angles match the seasonal periods of the data. A mode with phase θ repeats every $360^\circ/\theta$ patches, so a season of T time steps appears at $\theta = 360^\circ P/T$. With $P = 24$, the daily cycle of the 15-minute ETTm datasets ($T = 96$) appears at 90.1° (expected 90°), with its half-day harmonic at 180.0° on ETTm1; the daily cycle of the 10-minute Weather dataset ($T = 144$) at 59.9° (expected 60°), with a second harmonic at 120.3°; and the weekly cycle of the hourly Electricity and Traffic datasets ($T = 168$) at 51.5° and 51.6° (expected 51.4°), with second harmonics at 102.9° and 102.8°. On the hourly ETTh datasets a patch spans exactly one day, so the daily cycle folds onto the positive real axis, where all of their largest eigenvalues lie; Exchange, which has no calendar seasonality, shows no consistent angle. These modes appear as the isolated high-magnitude points outside the bulk in Fig. 7. KoPE thus recovers the dominant daily and weekly cycles of each dataset without being given them.

F.2.2 Stability

KoPE is initialized orthogonally, with LayerNorm applied before the operator rollout and to each resulting conditioning state. In addition, K is normalized by its spectral norm whenever this norm exceeds 1 (Sec. 3.4), so $\|K^i \mathbf{z}\|_2 \leq \|\mathbf{z}\|_2$ for every power and the rollout cannot expand. KoPE also uses a finite rollout determined by the prediction horizon and patch size: $M = \lceil H/P \rceil$, giving 30 conditioning states for $H = 720$ and $P = 24$.

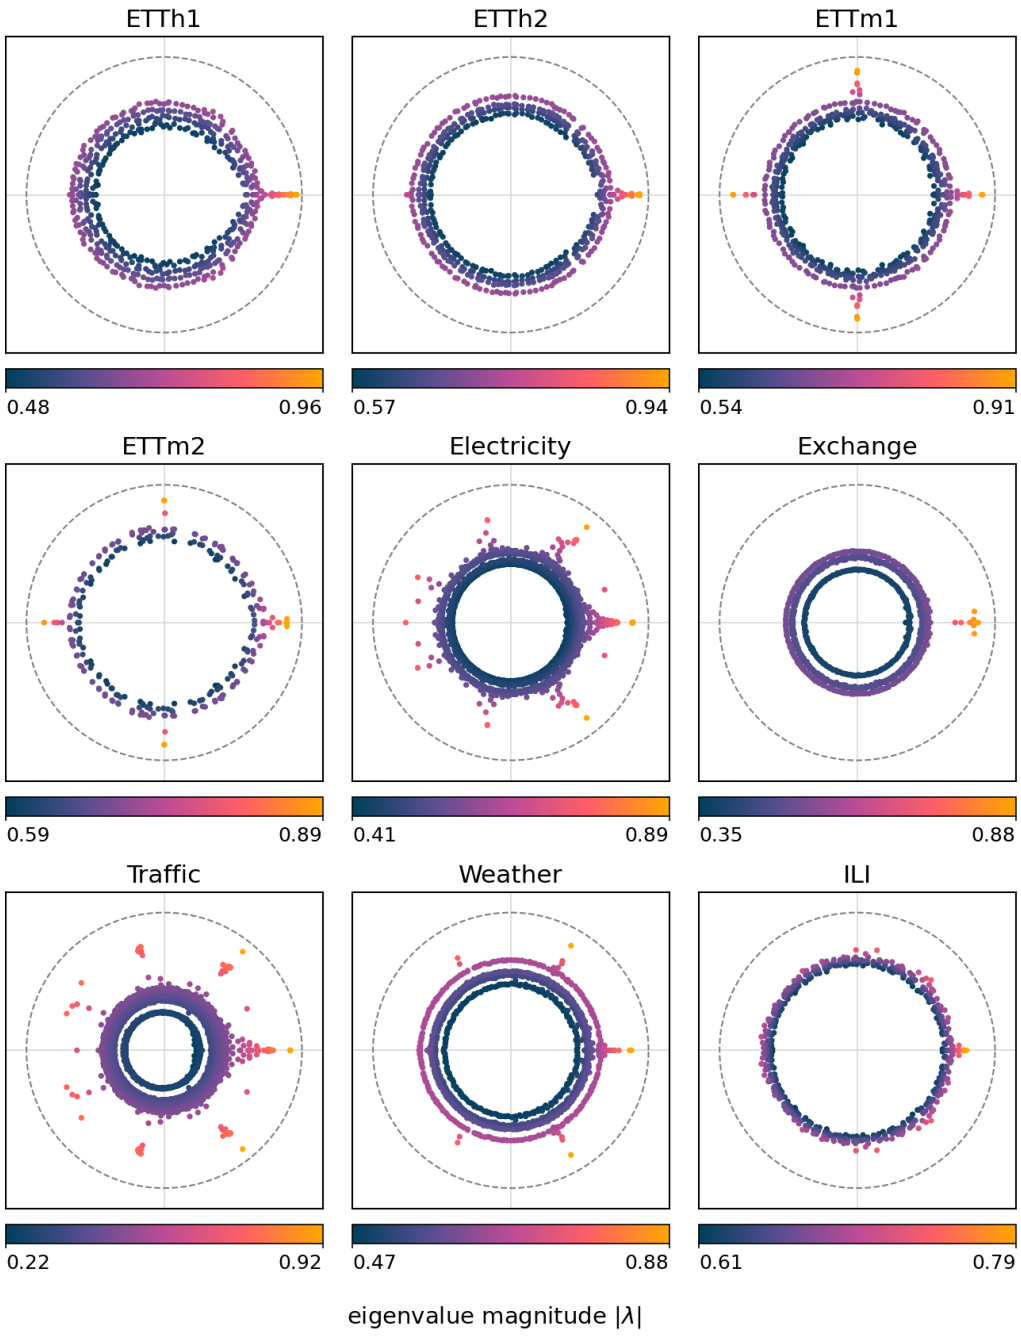

Figure 7: Eigenvalues of the learned KoPE operator $\mathbf{K}$ on the nine long-term datasets, pooled over the four horizons (seed 0); the dashed circle is the unit circle. Color encodes $|\lambda|$ on each panel's own range, given by the bar below the panel. The isolated high-magnitude modes outside the bulk include the seasonal modes discussed in App. F.2.1.

F.3 Uncertainty Gate

Predictability in time series varies across both forecast horizons and variates: longer-range forecasts generally involve greater uncertainty, while different underlying processes and measurement conditions introduce different levels of noise. AUG accommodates this heterogeneity by learning the relative contributions of deterministic and stochastic predictions through the forecasting objective. Its weights describe the contributions of the two branches and should not be interpreted as direct estimates of aleatoric uncertainty.

Figure 8 shows the stochastic weight across forecast steps (vertical axis) and variates (horizontal axis) for eight representative dataset–horizon pairs. Three structures stand out. First, except on Traffic, the first forecast steps are the most deterministic and the stochastic weight rises with lead time (on Exchange-L from 0.27 to around 0.8): the immediate future is well predicted from the recent context, while aleatoric uncertainty accumulates further ahead. Second, the short-term panels keep a substantial deterministic share (mean weight 0.64 on ETTm1-S), whereas on most long-term datasets the gate is nearly fully stochastic beyond the first steps. Third, the gate distinguishes variates: on ILI-L the noisy influenza-like-illness rates are almost fully stochastic (0.98), while the slowly changing number of reporting providers relies more on the deterministic branch (0.78). Together, these views show that AUG places more weight on the stochastic branch where aleatoric uncertainty is expected to be larger.

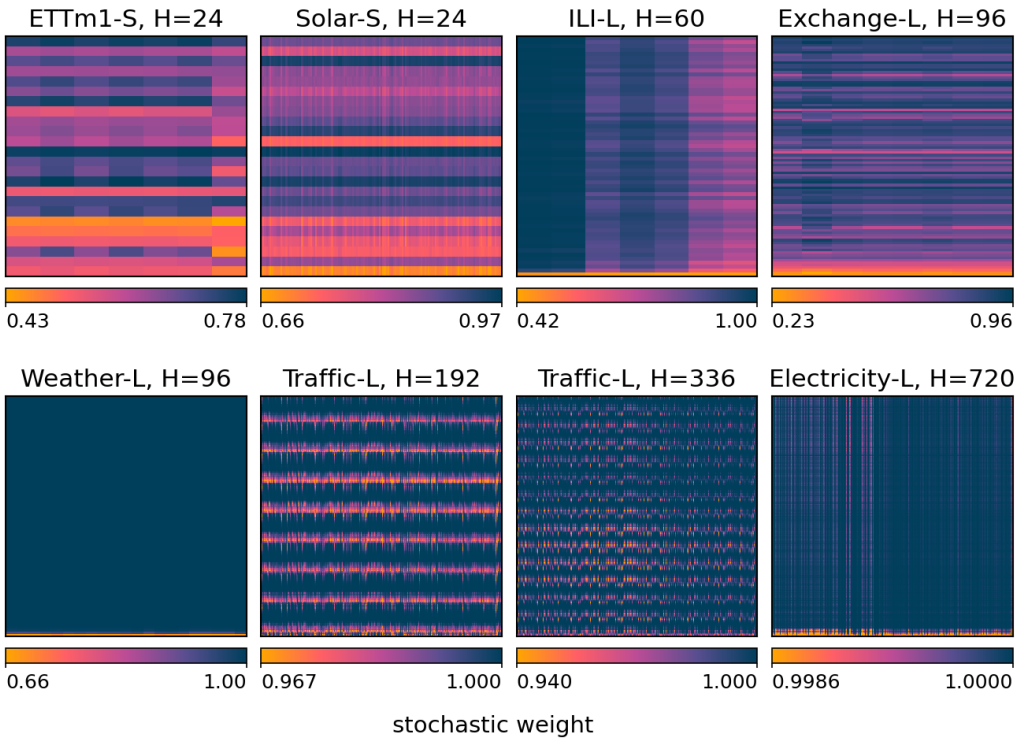

Figure 8: Gate weights averaged over the test windows (seed 0) for eight dataset–horizon pairs, named in the panel titles. Within each panel, the horizontal axis indexes variates and the vertical axis indexes forecast steps, increasing from bottom to top. Color encodes the stochastic weight on the panel's own range, from its lowest value (amber; the 1st percentile for Traffic-L and Electricity-L, whose few lowest cells would otherwise wash out the rest) to its highest (navy), given by the bar below the panel.

F.4 Qualitative Forecasts

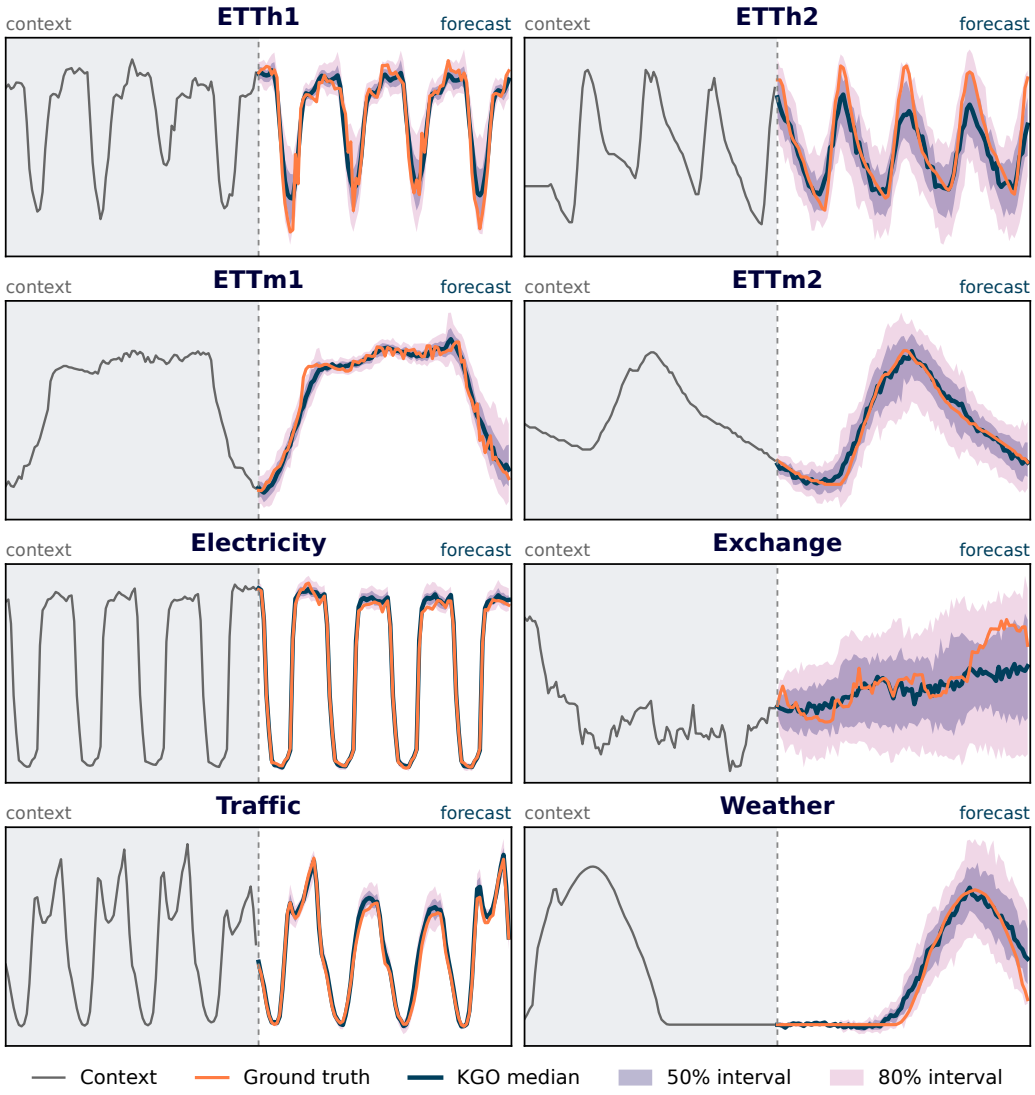

Figure 9: Qualitative forecasts at horizon $H{=}96$ (predictive interval view): context (gray), ground truth (orange), and KGO median (navy) with 50% (purple) and 80% (light magenta) predictive intervals.

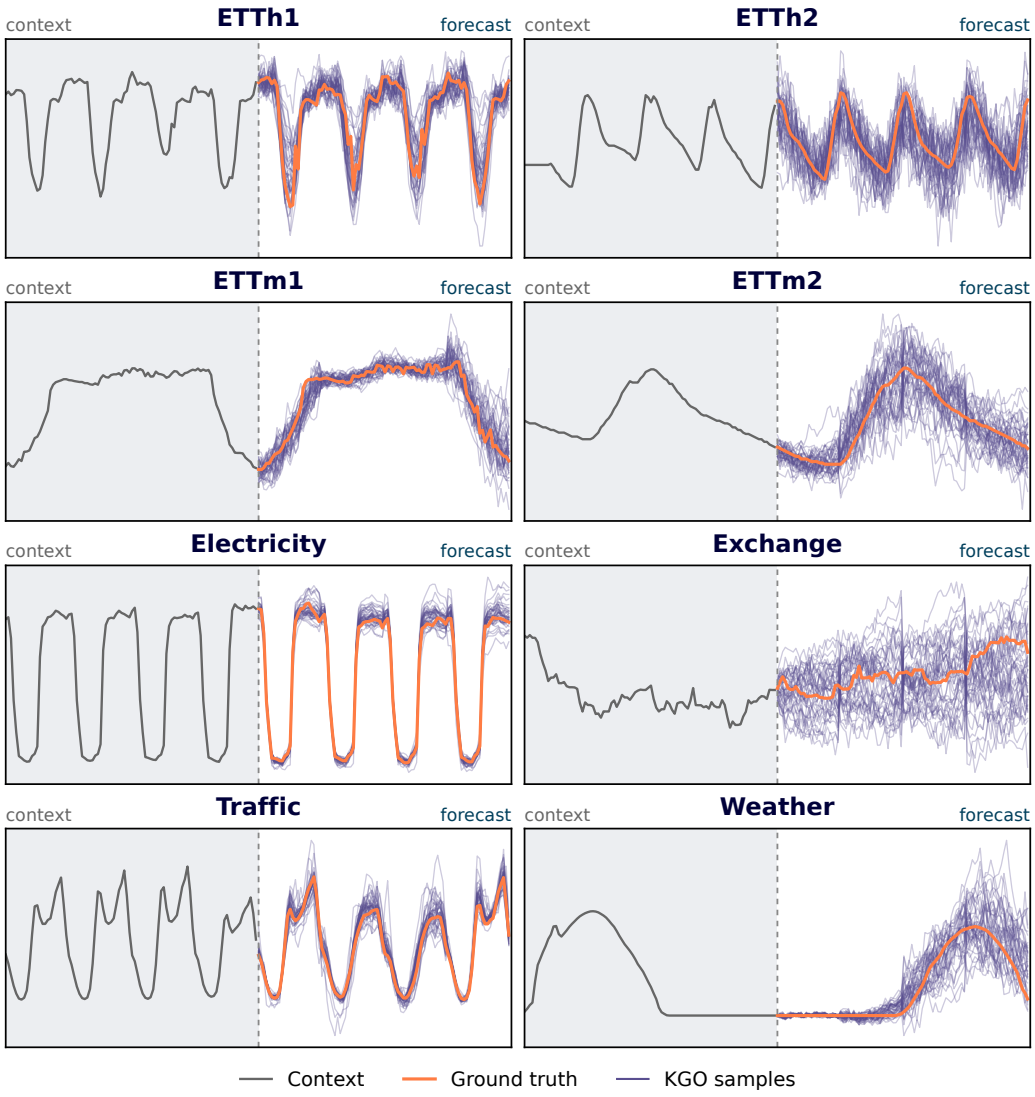

Figure 10: Qualitative forecasts at horizon $H{=}96$ (generative view): same windows as Figure 9, but rendered as 40 individual forecast trajectories sampled from KGO (purple) overlaid with the ground truth (orange). The density of trajectories visualizes the predictive distribution directly.