arXiv:2611.06052v1 [cs.LG] 4 Nov 2026

Alleviating Limited Effectiveness and Multi-Class Collapse in Deep Graph-level Clustering

Xiaoyang Lin
Hong Kong Baptist University
Hong Kong SAR, China
csxylin@comp.hkbu.edu.hk

Renchi Yang*
Hong Kong Baptist University
Hong Kong SAR, China
renchi@hkbu.edu.hk

Yuqi Chen
Hong Kong Baptist University
Hong Kong SAR, China
csyqchen@comp.hkbu.edu.hk

Xin Huang
Hong Kong Baptist University
Hong Kong SAR, China
xinhuang@comp.hkbu.edu.hk

Abstract

Given a collection of N graphs, *graph-level clustering* (GLC) seeks to partition the N graphs into semantically meaningful groups, which serves as a fundamental task in analyzing graph-structured data. While *Graph Isomorphism Networks* (GINs) have become the de facto encoders for GLC, they still fail to distinguish graphs due to their neglect of global topology. On top of that, existing deep clustering frameworks suffer from severe "clustering collapse" issues, particularly in imbalanced multi-class scenarios.

To address these challenges, we propose S^{3}GLC, a novel framework that integrates global topological characteristics with local neighborhood features for balanced GLC. Theoretically, we leverage the Hoffman-Wielandt inequality to construct Spectral Signature as a rigorous lower bound for *graph edit distance* (GED), providing a computationally efficient proxy for global topological features. We additionally propose a *joint manifold-preserving sparse regression* module to fuse Spectral Signature with Spatial Representation, eliminating noise while preserving the complementary information. Furthermore, to mitigate clustering collapse, we introduce a balanced clustering objective based on *optimal transport*, harnessing the *Sinkhorn-Knopp* normalization to enforce equipartition constraints on cluster assignments. Extensive experiments on 10 real-world datasets demonstrate that S^{3}GLC significantly outperforms 15 state-of-the-art baselines, achieving remarkable improvements on benchmark datasets such as *REDDIT-12K* and *OGBG-PPA*.

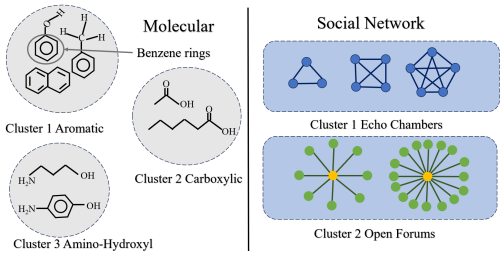

Figure 1: Example GLC applications.

1 Introduction

Graph-level clustering (GLC) [38, 72, 76] involves partitioning a collection of graphs into disjoint groups where graphs within the same cluster (resp. distinct clusters) exhibit high (resp. low) semantic similarity in topological structure or attributes. The unsupervised paradigm of GLC bypasses the need for expert priors [72, 76] and is critical in domains where label acquisition is cost-prohibitive or impractical [26, 61], such as screening molecular compounds for drug discovery [1, 7, 56], typological analysis of social communities [29, 44, 67], and analyzing states in brain networks [32, 33].

As illustrated in Fig. 1, molecular clustering primarily depends on functional groups (e.g., Benzene rings). Molecules that contain similar functional groups are regarded as belonging to the same compound (e.g., Aromatic). This highlights the importance of preserving key subgraph structures when learning graph representations. Therefore, effective GLC relies on *graph representation learning* (GRL) to encode such structural characteristics into informative feature vectors. However, relying solely on spectral methods [2, 75] to capture global topology overlooks the importance of subgraph structure, causing graphs that share the same subgraphs to fail to obtain similar graph-level representations. Meanwhile, most subgraph-matching-based approaches are inflexible and have high time complexity [27]. Moreover, as shown in the social network clustering in Fig. 1, graphs with similar structures often differ in scale. For example, in the cluster *Open Forums*, the two graphs are structurally similar but have different numbers of nodes (9 vs. 17). Since many methods based on GINs commonly use sum pooling [11, 12, 60], differences in graph size tend to overshadow structural differences, making the graph-level embeddings less discriminative. In summary, obtaining graph embeddings solely with spectral methods or with GINs has non-negligible drawbacks, as a subtask of unsupervised GRL, GLC inevitably inherits these challenges.

Prior work on graph-level clustering have evolved from early graph-kernel methods [8, 9, 20, 28, 43, 50, 50, 51] to recent GNN-based contrastive learning-based methods [16, 23, 31, 42, 58, 73] and end-to-end clustering frameworks [11, 12, 14]. Although these approaches have substantially improved graph representation learning and clustering efficiency, two fundamental challenges remain: *limited expressivity* and *multi-class collapse*. First, as discussed in §4.1, GINs cannot surpass the 1-WL test and thus fail to distinguish certain non-isomorphic graphs (see Fig. 2). Unfortunately,

*Corresponding Author

this limitation stems from the inability of GNNs to perceive the global structure and cannot be remedied by contrastive learning or hierarchical pooling [37]. *Graph edit distance* (GED) [48] can precisely measure the similarity between two graphs and distinguish graphs that the 1-WL test fails to tell apart. However, its computation is NP-hard. Second, to enable end-to-end clustering, both SOTA methods minimize the *Kullback–Leibler* (KL) divergence between cluster assignments and pseudo-labels, but this strategy lacks cluster-size constraints, leading to cluster collapse. In §4.2, we quantify the imbalance between ground-truth labels and the predicted labels of these SOTA methods on 10 real-world datasets and find that collapse is prevalent (see Fig. 3), especially pronounced on multi-class datasets such as *ENZYMES* [45].

To address the two fundamental challenges above, we present S^{3}GLC, a novel Spectral Self-Supervised Graph-Level Clustering framework that jointly tackles limited expressivity and multi-class collapse in a unified, end-to-end pipeline. We introduce *Spectral Signature* based on the Hoffman-Wielandt inequality, proving that adjacency eigenvalue distances constitute a rigorous lower bound on GED and enable efficient global topology encoding beyond the 1-WL limit, complementing GIN-based *Spatial Representation* (Fig. 2). We propose *joint manifold-preserving sparse regression* (JMPSR) to fuse spectral and spatial views through row-sparse projections and manifold preservation, together with *GED-guided contrastive learning* that aligns spatial embeddings via topology-aware positive pairs. Finally, we formulate balanced pseudo-label generation as an *optimal transport* problem with equipartition constraints, solved efficiently by the *Sinkhorn-Knopp* (SK) algorithm to ensure balanced cluster assignments. In summary, our contributions are threefold:

- We derive Spectral Signature from the Hoffman-Wielandt inequality, establishing it as a computationally tractable yet rigorous lower bound for GED that breaks the 1-WL expressivity ceiling without incurring NP-hard pairwise matching.
- We propose JMPSR together with GED-guided contrastive learning to integrate global spectral and local spatial views, preserving critical complementary semantics through manifold structure while suppressing redundancy via row-sparse projections.
- We introduce an optimal-transport-based balanced clustering objective with SK normalization, which enforces equipartition constraints on cluster assignments and mitigates the multi-class collapse prevalent in existing SOTA methods.

Extensive experiments on 10 real-world datasets against 15 competitive baselines demonstrate that S^{3}GLC establishes new state-of-the-art results, with particularly pronounced gains on challenging multi-class benchmarks (e.g., +14.6% ACC on *REDDIT-12K* and +35.2% NMI on *DD*), validating the effectiveness of topology-enhanced, balanced GLC.

2 Related Work

2.1 Graph-level Representation Learning

Current frameworks for unsupervised GRL can be broadly categorized into two streams: mutual information-based methods [40, 46, 68] and contrastive learning-based approaches [16, 23, 58]. Regarding the former, InfoGraph [53] learns graph embeddings by maximizing the mutual information between the global graph representation and multi-scale substructures. Similarly, HeGCL [52] enhances the mutual information between similar views through a cross-view contrastive mechanism. The second category, contrastive learning, has gained significant traction. GraphCL [73] employs diverse graph augmentations with a contrastive loss to yield highly transferable and robust graph embeddings. GraphACL [42] introduces an adversarial branch to generate hard negative samples. Beyond these, several approaches emphasize capturing global graph semantics [39, 41, 54, 71]. GraphLoG [66] enhances instance-level contrast with hierarchical prototypes representing global semantic clusters. These frameworks provide robust and transferable graph embeddings without supervision.

2.2 Graph-level Clustering

GLC aims to partition a set of unlabeled graphs into distinct groups based on semantic similarity. Traditional approaches typically employ graph kernels [8, 28, 43, 50, 51] to measure pairwise similarities via subgraph structures, subsequently applying spectral clustering algorithms to perform the partition. In recent years, GNNs have become the dominant approach for graph encoding [30, 64]. Inspired by advancements in deep learning [14, 21, 70], end-to-end clustering frameworks have been widely adopted. For instance, DGLC [11] and DCGLC [12] utilize pseudo-labels to supervise the clustering process within a unified framework. However, this self-training paradigm can inadvertently exacerbate the risk of clustering collapse. Furthermore, to effectively integrate global structural features with local node attributes, multi-view graph clustering methods [22, 36] have been proposed. However, existing approaches often fail to fully address the challenge of effectively aligning and exploiting the complementary information in different views.

3 Preliminaries

3.1 Graph Nomenclature and Terminology

Symbols and Problem Statement. Given a set of N graphs $\mathcal{G} = \{\mathcal{G}^{(1)}, \mathcal{G}^{(2)}, \ldots, \mathcal{G}^{(N)}\}$, where the i-th graph $\mathcal{G}^{(i)}$ is associated with a node set $\mathcal{V}^{(i)}$, an edge set $\mathcal{E}^{(i)}$, and a node attribute matrix $\mathbf{X}^{(i)}$, denoted by $\mathcal{G}^{(i)} = (\mathcal{V}^{(i)}, \mathcal{E}^{(i)}, \mathbf{X}^{(i)})$. For each node $u \in \mathcal{G}^{(i)}$, we use $\mathcal{N}(u)$ to represent the set of neighboring nodes of node u. The adjacency matrix of graph $\mathcal{G}^{(i)}$ is symbolized by $\mathbf{A}^{(i)}$. We denote by $\mathbf{L}^{(i)}$ and $\tilde{\mathbf{L}}^{(i)}$ the Laplacian and normalized Laplacian of graph $\mathcal{G}^{(i)}$, respectively. Throughout this paper, sets are denoted by calligraphic letters, e.g., $\mathcal{V}$. Matrices (resp. vectors) are written in bold uppercase (resp. lowercase) letters, e.g., $\mathbf{E}$ (resp. $\mathbf{e}$). We use $\mathbf{E}[i]$ and $\mathbf{E}[\cdot, i]$ to represent the i^{th} row and column of $\mathbf{E}$, respectively. $\|\mathbf{E}\|_F$ denotes the Frobenius norm of matrix $\mathbf{E}$, as for the computation of $\|\mathbf{E}\|_{2,1}$, is defined as the sum of the L_2-norms of its rows. Table 1 summarizes frequently used symbols.

The goal of *graph-level clustering* (GLC) is to partition the graph set $\mathcal{G}$ into K disjoint groups: $\{C_1, C_2, \ldots, C_K\}$ and $C_i \cap C_j = \emptyset$ for any $i \neq j$, such that graphs within the same cluster C_i are similar to each other in terms of topology and attribute semantics, while those across clusters are dissimilar to each other.

Manifold Structure. Given a data matrix $\mathbf{X}$, the manifold structure underlying it can be characterized by its graph Laplacian [6], $\mathbf{L}_{\mathbf{X}}$. Particularly, a lower-dimensional matrix $\mathbf{Y}$ is said to *preserve the*

Table 1: Frequently used notations.

Symbol	Description
$\mathcal{G}^{(i)}, \mathcal{V}^{(i)}, \mathcal{E}^{(i)}$	The i-th graph sample and its node and edge sets.
N, K, B	The total number of graphs, clusters and batch size.
$\mathbf{X}^{(i)}, \mathbf{A}^{(i)}$	The node attribute matrix and adjacency matrix of $\mathcal{G}^{(i)}$.
$\boldsymbol{\lambda}(\mathbf{A}^{(i)})$	The vector of sorted eigenvalues of $\mathbf{A}^{(i)}$.
$\mathbf{H}, \mathbf{F}$	The Spatial Representation and Spectral Signature of graphs.
$\mathbf{Y}, \mathbf{Z}$	The consensus matrix and fused embeddings of graphs.
h, r, d	The dimensions of $\mathbf{H}, \mathbf{F}, \mathbf{Z}$, respectively
$\mathbf{W}_{\mathrm{H}}, \mathbf{W}_{\mathrm{F}}$	The sparse projection matrices for $\mathbf{H}$ and $\mathbf{F}$.
$\mathbf{Q}, \mathbf{T}$	The soft cluster assignment and the target distribution.

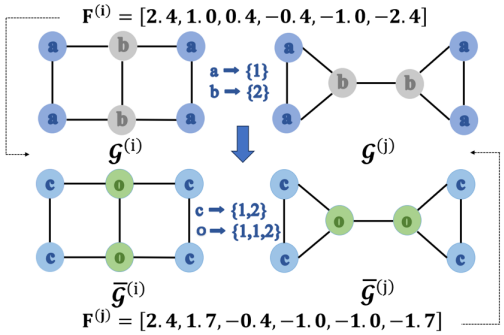

Figure 2: A toy example for 1-WL test.

manifold structure of $\mathbf{X}$ if for any node pair $i \neq j$ with the lower value in $\mathbf{L}_{\mathrm{X}}[i, j]$, $\|\mathbf{Y}[i] - \mathbf{Y}[j]\|_2^2$ is minimized.

Sparse Regression. The *sparse regression* [25] eliminates feature redundancy by jointly evaluating feature subsets via sparse multi-output regression against the target matrix $\mathbf{Y}$, formulated as:

$$\min_{\mathbf{W}} \|\mathbf{Y} - \mathbf{X}\mathbf{W}\|_{\mathrm{F}}^2 + \lambda \|\mathbf{W}\|_{2,1}. \tag{1}$$

Intuitively, the $L_{2,1}$-norm induces *row-sparsity* on the projection matrix $\mathbf{W}$. Since each row of $\mathbf{W}$ corresponds to a feature dimension, driving a row's L_2-norm to zero effectively eliminates that feature.

3.2 Graph Isomorphism, WL test, and GED

Given two graphs $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$, we say they are isomorphic if there exists a bijective mapping function $f: \mathcal{G}^{(i)} \rightarrow \mathcal{G}^{(j)}$ such that an edge $(u, v) \in \mathcal{G}^{(i)}$ iff the edge $(f(u), f(v))$ exists in $\mathcal{G}^{(j)}$. The graph isomorphism problem involves determining whether two graphs are isomorphic, i.e., topologically identical, which has no known polynomial-time algorithms for it yet [4].

The Weisfeiler-Lehman (WL) test of graph isomorphism [35] is an effective and fast (polynomial-time) tool for distinguishing a broad class of graphs. Its 1-dimensional form (1-WL test) iteratively (i) aggregates the labels of nodes and their neighbors, and (ii) hashes the aggregated labels into unique new labels. The test determines that two graphs are non-isomorphic if, at some iteration, the labels of the nodes between the two graphs are different.

In the literature, *graph edit distance* (GED) [48] generalizes isomorphism and acts as a standard continuous measure for quantifying the topological similarity of a graph pair. More concretely, the GED $\mathrm{ged}(\mathcal{G}^{(i)}, \mathcal{G}^{(j)})$ between two graphs $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$ is defined as the minimum number of edit operations to make $\mathcal{G}^{(i)}$ isomorphic to $\mathcal{G}^{(j)}$:

$$\mathrm{ged}(\mathcal{G}^{(i)}, \mathcal{G}^{(j)}) = \min_{\Upsilon \in \mathcal{P}(\mathcal{G}^{(i)}, \mathcal{G}^{(j)})} \sum_{e_i \in \Upsilon} c(e_i), \tag{2}$$

where $\mathcal{P}(\mathcal{G}^{(i)}, \mathcal{G}^{(j)})$ denotes the set of all valid edit paths transforming $\mathcal{G}^{(i)}$ to $\mathcal{G}^{(j)}$. An edit operation[1] e_i can be either one of the following node-wise/edge-wise operations of cost $c(e_i)$: (i) insert or delete an isolated vertex; and (ii) insert or delete an edge.

[1]We consider the case where nodes/edges are unlabeled.

3.3 Graph Isomorphism Networks

Compared to generic GNNs, *graph isomorphism networks* (GINs) [65] achieve higher discriminative power through an *injective* neighborhood aggregation scheme. That is, the rule for updating node representations therein is formulated as:

$$\mathbf{h}_v^{(\ell)} = \mathsf{MLP}^{(\ell)}\big((1 + \gamma^{(\ell)}) \cdot \mathbf{h}_v^{(\ell-1)} + \sum_{u \in \mathcal{N}(v)} \mathbf{h}_u^{(\ell-1)}\big), \tag{3}$$

where $\mathbf{h}_v^{(\ell)}$ represents the latent embedding for node v at the ℓ-th layer ($1 \leq \ell \leq L$), and $\gamma^{(\ell)}$ is a weighting coefficient. Distinct from generic GNNs that employ mean operators, GINs aggregates the features from neighborhood via a summation (i.e. $\sum_{u \in \mathcal{N}(v)} \mathbf{h}_u^{(\ell-1)}$), which guarantees the injectivity of the aggregation function over multisets, thereby allowing the model to achieve discriminative power equivalent to the 1-WL test. Subsequently, GINs derive the Spatial Representation for each graph $\mathcal{G}^{(i)}$ by applying global summation pooling over all final-layer node embeddings in $\mathcal{G}^{(i)}$: $\mathbf{H}_i = \mathsf{READOUT}(\{\mathbf{h}_v^{(L)} : v \in \mathcal{V}^{(i)}\}) = \sum_{v \in \mathcal{V}^{(i)}} \mathbf{h}_v^{(L)}$.

4 Design Motivation

In this section, we elaborate on our motivation for proposing $\mathsf{S}^3\mathsf{GLC}$ by analyzing the inherent deficiencies of existing solutions.

4.1 Limited Effectiveness of GINs

While GINs match the theoretical limit of the 1-WL test, and thus, surpass generic GNNs in discriminability, the nature of their aggregation mechanism inherently makes them blind to global topological properties. Particularly, some graphs cannot be distinguished by the 1-WL test, e.g., distinct graphs with similar neighborhoods but different global topologies are possibly mapped to identical embeddings by GINs (Theorem 4.1).

Theorem 4.1. *If $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$ cannot be distinguished by the 1-WL test, then their Spatial Representations by GINs are indistinguishable.*

We use a toy example in Fig. 2 for illustration. It can be observed that $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$ present distinct structures, i.e., are non-isomorphic, given that $\mathcal{G}^{(i)}$ contains 4-cycles and is bipartite, whilst $\mathcal{G}^{(j)}$ contains 3-cycles and is non-bipartite. Initially, nodes in graphs $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$ are labeled with 1 or 2. After neighborhood aggregations, each node in the updated graph $\overline{\mathcal{G}}^{(i)}$ and $\overline{\mathcal{G}}^{(j)}$ possesses a

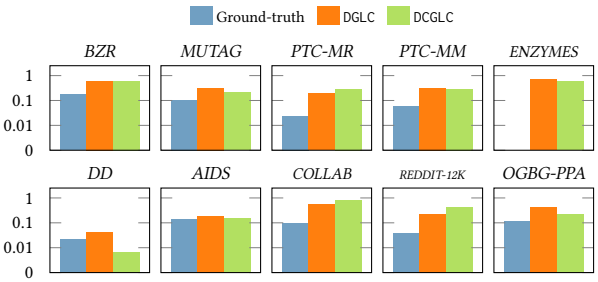

Figure 3: The PCR of existing models on real datasets.

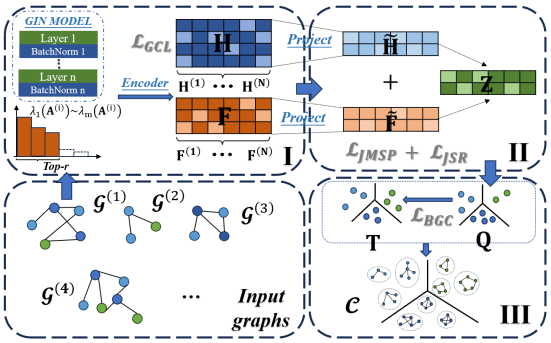

Figure 4: The overall framework of S³GLC.

new set of labels (c or o). At the graph level, the label multi-sets of $\overline{\mathcal{G}}^{(i)}$ and $\overline{\mathcal{G}}^{(j)}$ are identical, i.e., $\{c, c, o, o, c, c\}$, which corroborates our above claim and implies that the Spatial Representations produced by GINs cannot solely represent the graphs effectively due to its inability to capture the global topology.

Lemma 4.2 ([48]). *$\forall \mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$, $\mathsf{ged}(\mathcal{G}^{(i)}, \mathcal{G}^{(j)}) = 0$ iff $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$ are isomorphic ($\mathcal{G}^{(i)} \cong \mathcal{G}^{(j)}$).*

According to Lemma 4.2, the GED $\mathsf{ged}(\mathcal{G}^{(i)}, \mathcal{G}^{(j)}) = 4$ can be used as an effective indicator for graph isomorphism and quantitative metric for measuring the topological differences between $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$, particularly where GINs fails. However, integrating the GED as augmented features into the representation learning of graphs presents a prohibitive computational barrier. The exact computation of GED is known to be NP-hard [10, 74], and even the approximation of a single pair of graphs with dozens of nodes is intractable [19], not to mention for the $N \times N$ graph pairs.

4.2 Imbalanced Multi-Class Clustering

The SOTA GLC models DCGLC [12] and DGLC [11] enable end-to-end training for effective self-supervised clustering by minimizing the Kullback-Leibler (KL) divergence between a soft cluster assignment $\mathbf{Q}$ and a sharpened target distribution $\mathbf{T}$ [62].

Despite sharpening the soft cluster assignments, this mechanism suffers from a critical flaw: *it lacks explicit global constraints on the cluster sizes and tends to produce imbalanced clusters*, which is particularly severe in multi-class scenarios. To be specific, unlike binary clustering (widely adopted in previous GLC works [11, 12]), where assignments are naturally antagonistic, in multi-class settings, existing models can easily collapse the majority of samples into a few dominant groups while leaving others empty [3, 13].

As empirical evidence, we quantify the discrepancy in distributional imbalance between the ground-truth class distribution and the predicted clustering results by DCGLC and DGLC. To this end, we employ a metric derived from the *Normalized Shannon Entropy* [49], referred to here as the *Partition Collapse Ratio* (PCR), defined as:

$$\mathrm{PCR}(C_1, \ldots, C_K) = 1 - \frac{-\sum_{k=1}^{K} p_k \log p_k}{\log K}. \tag{4}$$

Here, $p_k = |C_k|/N$ denotes the proportion of samples assigned to the k-th cluster, and the denominator $\log K$ represents the maximum possible entropy. The PCR of clusters $\{C_1, C_2, \ldots, C_K\}$ ranges from 0 to 1, and the larger (resp. smaller) value indicates a higher (resp. lower) degree of cluster collapse or imbalance. Fig. 3 reports the PCR values calculated upon ground-truth clusters and those output by baselines DCGLC and DGLC over 10 real datasets with multiple classes (see Table 2). It can be observed that the PCRs achieved by DCGLC and DGLC are consistently and conspicuously higher than the ground-truth PCR values. For example, on *ENZYMES* dataset, the PCR values by DCGLC and DGLC are over 0.55, whereas the ground-truth PCR is almost 0. This significant gap demonstrates that the baseline models suffer from severe clustering collapse problems on multi-class graph datasets, which, in turn, hinders them from gaining further inroads on clustering performance.

5 Methodology

In this section, we present our proposed S³GLC model to cope with the foregoing limitations of existing works in GLC. We begin with a synoptic overview of S³GLC in §5.1, followed by elaborating on the algorithmic details for constructing Spectral Signature in §5.2, the fusion of graph feature vectors in §5.3, and training and balanced clustering loss functions in §5.4. Lastly, §5.5 provides the complexity analysis of S³GLC.

5.1 Synoptic Overview

As illustrated in Fig. 4, the major stages S³GLC include learning the final graph representations $\mathbf{Z} \in \mathbb{R}^{N \times d}$ and deriving clusters $\{C_1, C_2, \ldots, C_K\}$ therefrom. More concretely, the representation learning involves three main components. The first is to construct the Spectral Signature $\mathbf{F} \in \mathbb{R}^{N \times r}$ for N graphs, whose high-level idea is to derive the spectral signatures of each graph's adjacency matrix to approximate the GED, and thus, capture the global structural signatures. In the meantime, Spatial Representation $\mathbf{H} \in \mathbb{R}^{N \times h}$ that encodes both node attributes and local neighborhoods in each graph $\mathcal{G}^{(i)}$ is generated by classic GINs models:

$$\mathbf{H}^{(i)} = \mathsf{GIN}(\mathcal{G}^{(i)}, \mathbf{X}^{(i)}) \in \mathbb{R}^{h}. \tag{5}$$

Afterwards, S³GLC seeks to fuse $\mathbf{F}$ and $\mathbf{H}$ into the final representations $\mathbf{Z}$ so as to retain complementarity while eliminating redundancy. Towards this end, S³GLC learns two projection matrices $\mathbf{W}_{\mathrm{H}}$ and $\mathbf{W}_{\mathrm{F}}$ through JMPSR module, which projects $\mathbf{F}$ and $\mathbf{H}$ into a unified lower-dimensional space, i.e.,

$$\tilde{\mathbf{H}} = \mathbf{H}\mathbf{W}_{\mathrm{H}} \in \mathbb{R}^{d}, \ \tilde{\mathbf{F}} = \mathbf{F}\mathbf{W}_{\mathrm{F}} \in \mathbb{R}^{d}. \tag{6}$$

Based thereon, the unified representation $\mathbf{Z}$ is obtained by

$$\mathbf{Z} = \tilde{\mathbf{H}} + \tilde{\mathbf{F}} = \mathbf{H}\mathbf{W}_{\mathrm{H}} + \mathbf{F}\mathbf{W}_{\mathrm{F}}. \tag{7}$$

In the first stage, S^{3}GLC facilitates the learning of $\mathbf{H}$ via a *GED-guided contrastive learning objective*, and includes a self-supervised clustering loss that explicitly harnesses the Sinkhorn-Knopp normalization to avert cluster collapse.

5.2 Spectral Signature for Graphs

Instead of calculating the GED of graph pairs for the construction of embedding vectors [5], which entails prohibitive computational and training overhead, S^{3}GLC resorts to directly deriving signatures from the spectrum of each graph independently as the GED-approximate embedding for the graph.

More concretely, we reformulate the GED computation as an equivalent problem of finding the optimal permutation matrix $\mathbf{P}$ that optimizes the following objective [77]:

$$\mathrm{ged}(\mathcal{G}^{(i)}, \mathcal{G}^{(j)}) = \min_{\mathbf{P}\in\mathbb{\Pi}} \frac{1}{2} \left\| \mathbf{A}^{(i)} - \mathbf{P}^\top \mathbf{A}^{(j)} \mathbf{P} \right\|_{\mathrm{F}}^2, \tag{8}$$

where $\mathbb{\Pi}$ denotes the set of all valid permutation matrices and $\mathbf{P}^\top\mathbf{A}^{(j)}\mathbf{P}$ represents the permutations of the rows and columns in $\mathbf{A}^{(j)}$ for node matching between $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$. Note that when $|\mathcal{V}^{(i)}| \neq |\mathcal{V}^{(j)}|$, we can simply pad zero rows and columns to the smaller adjacency matrices to ensure the dimension consistency.

Lemma 5.1 (Spectral Permutation Invariance). *Let $\mathbf{A}$ be the adjacency matrix of a graph. For any permutation matrix $\mathbf{P} \in \mathbb{\Pi}$, the spectrum of the permuted matrix $\mathbf{P}^\top\mathbf{A}\mathbf{P}$ is identical to that of $\mathbf{A}$.*

Theorem 5.2 (Hoffman-Wielandt Inequality [24]). *Let $\lambda_1(\mathbf{A}) \geq \ldots \geq \lambda_n(\mathbf{A})$ denote the eigenvalues of a real symmetric matrix $\mathbf{A}$ sorted in descending order. For any two such matrices $\mathbf{A}$ and $\mathbf{B}$, the following inequality holds: $\sum_{k=1}^{n} (\lambda_k(\mathbf{A}) - \lambda_k(\mathbf{B}))^2 \leq \|\mathbf{A} - \mathbf{B}\|_F^2$.*

The direct optimization of Eq. (8) is still NP-hard. Our core idea is to leverage our Lemma 5.1 and the Hoffman-Wielandt theorem [24] (Theorem 5.2) for an approximate solution. More precisely, from Lemma 5.1, the eigenvalue sequences correspond exactly under any permutation:

$$\lambda_k(\mathbf{A}) \equiv \lambda_k(\mathbf{P}^\top\mathbf{A}\mathbf{P}), \quad \forall k = 1, \ldots, n.$$

Using Theorem 5.2, our optimization objective in Eq. (8) can be lower bounded as follows:

$$\begin{aligned} \min_{\mathbf{P}\in\mathbb{\Pi}} \left\| \mathbf{A}^{(i)} - \mathbf{P}^\top \mathbf{A}^{(j)} \mathbf{P} \right\|_{\mathrm{F}}^2 &\geq \min_{\mathbf{P}\in\mathbb{\Pi}} \sum_{k=1}^{n} \left(\lambda_k(\mathbf{A}^{(i)}) - \lambda_k(\mathbf{P}^\top\mathbf{A}^{(j)}\mathbf{P}) \right)^2 \\ &= \sum_{k=1}^{n} \left(\lambda_k(\mathbf{A}^{(i)}) - \lambda_k(\mathbf{A}^{(j)}) \right)^2 = \|\boldsymbol{\lambda}(\mathbf{A}^{(i)}) - \boldsymbol{\lambda}(\mathbf{A}^{(j)})\|_2^2, \end{aligned} \tag{9}$$

where $\boldsymbol{\lambda}(\mathbf{A}^{(i)})$ and $\boldsymbol{\lambda}(\mathbf{A}^{(j)})$ denote the vectors containing the respective sorted eigenvalues of $\mathbf{A}^{(i)}$ and $\mathbf{A}^{(j)}$, referred to as *Spectral Signature* hereafter.

Eq. (9) reveals that the Euclidean distance between Spectral Signature serves as a rigorous *lower bound* for the GED of graph pair $(\mathcal{G}^{(i)}, \mathcal{G}^{(j)})$, which accurately captures the topological similarity between graphs. On top of that, the Spectral Signature of each graph can be very efficiently computed by merely applying the high-performance eigenvalue solvers [34] in common numeric computation libraries over its small and sparse adjacency matrix (often with only dozens of nodes and edges) in GLC tasks. Since the small eigenvalue part is negligible, we pick the top-r largest eigenvalues in $\boldsymbol{\lambda}(\mathbf{A}^{(i)})$ of each graph $\mathcal{G}^{(i)}$ to act as its Spectral Signature $\mathbf{F}^{(i)}$:

$$\mathbf{F}^{(i)} = [\lambda_1(\mathbf{A}^{(i)}), \lambda_2(\mathbf{A}^{(i)}), \ldots, \lambda_r(\mathbf{A}^{(i)})]^\top. \tag{10}$$

We revisit the example in Fig. 2. With the Spectral Signature $[2.4, 1, 0.4, -0.4, -1, -2.4]$ and $[2.4, 1.7, -0.4, -1, -1, -1.7]$ for $\mathcal{G}^{(i)}$ and $\mathcal{G}^{(j)}$, respectively, these two graphs can be easily distinguished now.

5.3 Joint Manifold-Preserving Sparse Regression Feature Fusion

We fuse two complementary graph views, i.e., Spatial Representation $\mathbf{H}$ and Spectral Signature $\mathbf{F}$, whose discrepant information is critical for discriminating graphs. Accordingly, our fusion preserves view-specific information while suppressing noisy information.

5.3.1 **Fusion Objectives.** To suppress noisy information during fusion, we introduce a latent consensus space $\mathbf{Y}$, which can guide the learning of the projection matrices $\mathbf{W}_{\mathrm{H}}$ and $\mathbf{W}_{\mathrm{F}}$, specifically, we propose the *joint sparse regression* scheme:

$$\min_{\mathbf{W}_{\mathrm{H}},\mathbf{W}_{\mathrm{F}}} \mathcal{L}_{\mathrm{JSR}} = \underbrace{\|\mathbf{Y} - \tilde{\mathbf{H}}\|_{\mathrm{F}}^2 + \|\mathbf{Y} - \tilde{\mathbf{F}}\|_{\mathrm{F}}^2}_{\text{Regression}} + \underbrace{(\|\mathbf{W}_{\mathrm{H}}\|_{2,1} + \|\mathbf{W}_{\mathrm{F}}\|_{2,1})}_{\text{Sparse Regularization}}, \tag{11}$$

where the regression term aligns both views to the same consensus space, while the sparse regularization term promotes row-sparsity in $\mathbf{W}_{\mathrm{H}}$ and $\mathbf{W}_{\mathrm{F}}$, which is formally justified in Appendix A.5. This row-wise sparsification can act as a denoising function during fusion.

In multi-view fusion, discrepant information is often treated as noise and removed [59], so we expect the consensus representation $\mathbf{Y}$ to guide the Spatial Representation $\mathbf{H}$ and Spectral Signature $\mathbf{F}$ to preserve view-specific complementary information during projection.

Theorem 5.3 ([47]). *For $\mathbf{X} \in \mathbb{R}^{n\times n}$, minimizing $\sum_{i=1}^{n} \|\mathbf{X}[i] - \sum_{j=1}^{n} \mathbf{W}[i,j]\mathbf{X}[j]\|$ can force $\mathbf{X}$ preserve the manifold structure of $\mathbf{W}$.*

Inspired by Theorem 5.3, we further require $\mathbf{Y}$ capturing both the shared and view-specific manifold structures of the Spatial Representation $\mathbf{H}$ and Spectral Signature $\mathbf{F}$, which minimizes the weighted pairwise distances in both feature spaces [25]. The corresponding *joint manifold structure preservation* objective can be formulated as:

$$\min_{\mathbf{Y}} \mathcal{L}_{\mathrm{JMSP}} = \frac{1}{2} \sum_{i=1}^{B} \|\mathbf{Y}[i] - \sum_{j=1}^{B} \mathbf{S}_{\mathrm{CA}}[i,j] \cdot \mathbf{Y}[j]\|^2 = \mathrm{Tr}\left(\mathbf{Y}^\top \mathbf{L}_{\mathrm{CA}} \mathbf{Y}\right), \tag{12}$$

where $\mathbf{L}_{\mathrm{CA}} = (\mathbf{I} - \mathbf{S}_{\mathrm{CA}})^\top(\mathbf{I} - \mathbf{S}_{\mathrm{CA}})$ and the $\mathbf{S}_{\mathrm{CA}}$ represents the consensus affinity matrix of projected features $\tilde{\mathbf{H}}$ and $\tilde{\mathbf{F}}$ defined in Eq. (13).

$$\mathbf{S}_{\mathrm{CA}} = (\tilde{\mathbf{H}} + \tilde{\mathbf{F}})(\tilde{\mathbf{H}} + \tilde{\mathbf{F}})^\top = \underbrace{\tilde{\mathbf{H}}\tilde{\mathbf{H}}^\top}_{\text{self affinity}} + \underbrace{\tilde{\mathbf{F}}\tilde{\mathbf{F}}^\top}_{\text{self affinity}} + \underbrace{(\tilde{\mathbf{H}}\tilde{\mathbf{F}}^\top + \tilde{\mathbf{F}}\tilde{\mathbf{H}}^\top)}_{\text{cross affinity}}. \tag{13}$$

Finally, we keep $\mathbf{S}_{\mathrm{CA}}[i,j] > 0$ and set the remaining entries to 0 via element-wise thresholding. The cross-affinity terms preserve the shared manifold structure, whereas the self-affinity terms preserve the view-specific manifold structure.

In a nutshell, the overall fusion objective is formulated as:

$$\min_{\mathbf{Y},\mathbf{W}_{\mathrm{H}},\mathbf{W}_{\mathrm{F}}} \quad \mathcal{L}_{\mathrm{JMSP}} + \mathcal{L}_{\mathrm{JSR}}. \tag{14}$$

5.3.2 **Closed-form Solutions.** Notice that our optimization objective in Eq. (14) is not jointly convex with respect to all variables. However, it is convex with respect to each variable when the others are fixed. Therefore, we resort to an *alternating optimization* strategy to iteratively update the latent feature space $\mathbf{Y}$, and projection matrices $\mathbf{W}_\mathrm{H}$ and $\mathbf{W}_\mathrm{F}$.

Update $\mathbf{W}_\mathbf{H}$ and $\mathbf{W}_\mathbf{F}$. When fixing $\mathbf{Y}$, the optimization for $\mathbf{W}_\mathrm{H}$ (resp. $\mathbf{W}_\mathrm{F}$) becomes a standard sparse regression problem. Taking the derivative of the objective function with respect to $\mathbf{W}_\mathrm{H}$ and setting it to zero, we obtain $-2\mathbf{H}^\top(\mathbf{Y}-\mathbf{H}\mathbf{W}_\mathrm{H})+2\Delta_\mathrm{H}\mathbf{W}_\mathrm{H}=0$, where Δ_H is a diagonal matrix whose the k-th diagonal element is $\frac{1}{2\|\mathbf{W}_\mathrm{H}[k]\|_2}$,which essentially re-weights the features during. Accordingly, the closed-form solution is:

$$\mathbf{W}_\mathrm{H}=(\mathbf{H}^\top\mathbf{H}+\Delta_\mathrm{H})^{-1}\mathbf{H}^\top\mathbf{Y}. \tag{15}$$

Analogously, we can derive $\mathbf{W}_\mathrm{F}=(\mathbf{F}^\top\mathbf{F}+\Delta_\mathrm{F})^{-1}\mathbf{F}^\top\mathbf{Y}$, where Δ_F is a diagonal matrix whose the k-th diagonal element is $\frac{1}{2\|\mathbf{W}_\mathrm{F}[k]\|_2}$.

Update $\mathbf{Y}$. When the projection matrices $\mathbf{W}_\mathrm{H}$ and $\mathbf{W}_\mathrm{F}$ are fixed, the optimization problem w.r.t. $\mathbf{Y}$ can be rewritten by substituting the optimal $\mathbf{W}$ back into the original objective, yielding

$$\min_{\mathbf{Y}^\top\mathbf{Y}=\mathbf{I}}\ \mathrm{Tr}\left(\mathbf{Y}^\top\mathbf{M}\mathbf{Y}\right), \tag{16}$$

where the auxiliary matrix $\mathbf{M}$ is defined as:

$$\mathbf{M}=\mathbf{L}_\mathrm{CA}+2\mathbf{I}-\left(\mathbf{H}(\mathbf{H}^\top\mathbf{H}+\Delta_\mathrm{H})^{-1}\mathbf{H}^\top+\mathbf{F}(\mathbf{F}^\top\mathbf{F}+\Delta_\mathrm{F})^{-1}\mathbf{F}^\top\right). \tag{17}$$

The detailed derivative steps are deferred to Appendix A.7 for the interest of space. As per Ky Fan's trace maximization principle [18], the trace minimization problem in Eq. (16) can be optimized by computing the d smallest eigenvectors of $\mathbf{M}$ as $\mathbf{Y}$, and the shape of $\mathbf{M}$ depends on the batch size B, i.e. $\mathbf{M}\in\mathbb{R}^{B\times B}$.

5.4 Self-Supervised Training and Clustering

To facilitate the learning of Spatial Representation $\mathbf{H}$ by GINs and its resulting representations $\mathbf{Z}$ (Eq. (7)), as well as the clustering of $\mathbf{Z}$, we include two self-supervised training objectives: *GED-guided contrastive learning* loss $\mathcal{L}_\mathrm{GCL}$ and *balanced graph-level clustering* loss $\mathcal{L}_\mathrm{BGC}$, leading to our overall training loss: $\mathcal{L}_\mathrm{GCL}+\mathcal{L}_\mathrm{BGC}$.

5.4.1 **GED-guided Contrastive Learning.** To create discriminative Spatial Representation $\mathbf{H}$ from GINs, we propose to introduce a new graph contrastive learning objective [63] to incorporate the topological relevance (i.e., GED) of graph samples, which is shown to be preserved in their Spectral Signature. Specifically, given the Spectral Signature $\mathbf{F}$, for each anchor graph sample $\mathcal{G}^{(i)}$, we pick its β-nearest neighbors in terms of their embedding similarities in $\mathbf{F}$, constituting its positive neighbor set $P(\mathcal{G}^{(i)})$ as follows:

$$P(\mathcal{G}^{(i)})=\underset{\mathcal{G}^{(j)}\in\mathcal{G},j\neq i}{\arg\,\text{top-}\beta}\left(\mathbf{F}[i]\cdot\mathbf{F}[j]^\top\right), \tag{18}$$

where $\arg\,\text{top-}\beta(\cdot)$ is an operator that retrieves the top-β elements based on the similarity. For simplicity, we fix $\beta=10$ in all experiments.

The contrastive loss is then formulated to maximize the similarity between positive pairs in the neighbor set:

$$\mathcal{L}_\mathrm{GCL}=-\frac{1}{B\beta}\sum_{i=1}^{B}\sum_{\mathcal{G}^{(j)}\in P(\mathcal{G}^{(i)})}\log\frac{\exp\left(\frac{\cos(\mathbf{H}[i],\mathbf{H}[j])}{\tau}\right)}{\sum_{l=1}^{B}\exp\left(\frac{\cos(\mathbf{H}[i],\mathbf{H}[l])}{\tau}\right)-\exp\left(\frac{1}{\tau}\right)}, \tag{19}$$

where $\cos(\cdot,\cdot)$ denotes the cosine similarity, and τ is a temperature parameter controlling the smoothness of the distribution. Notably, $e^{1/\tau}$ is subtracted in the denominator to eliminate the self-similarity effect of the anchor sample itself [15].

5.4.2 **Balanced Graph-level Clustering.** We also adopt the following loss [62] for deep clustering:

$$\mathcal{L}_\mathrm{BGC}=\mathrm{KL}(\mathbf{T}\|\mathbf{Q})=\sum_{i=1}^{N}\sum_{k=1}^{K}\mathbf{T}[i,k]\cdot\log\frac{\mathbf{T}[i,k]}{\mathbf{Q}[i,k]}, \tag{20}$$

which is to minimize the KL divergence between the soft assignment $\mathbf{Q}$ and the target distribution $\mathbf{T}$. In particular, we follow [62] to construct the soft assignment $\mathbf{Q}$, but propose to achieve cluster balance by explicitly constraining the target distribution $\mathbf{T}$.

Specifically, after obtaining the final representations $\mathbf{Z}$, we partition the N graphs into K groups by applying K-`Means` over $\mathbf{Z}$, yielding K initial cluster cluster centroids: $\{\mathbf{m}_1,\mathbf{m}_2,\ldots,\mathbf{m}_K\}$. Following the standard practice in [3, 62], we utilize the Student's t-distribution kernel to quantify the probability of assigning the i-th graph to the k-th cluster:

$$\mathbf{Q}[i,k]=\frac{(1+\|\mathbf{Z}[i]-\mathbf{m}_k\|^2)^{-1}}{\sum_{j=1}^{K}(1+\|\mathbf{Z}[i]-\mathbf{m}_j\|^2)^{-1}}, \tag{21}$$

This formulation ensures that $\sum_{k=1}^{K}\mathbf{Q}[i,k]=1$, forming a valid probability distribution for each sample.

To cope with the multi-class imbalance issue stated in §4.2, we take inspiration from the *optimal transport* [17]. That is, we first impose the following constraint over the target distribution $\mathbf{T}$, i.e., $\mathbf{T}\mathbf{1}_K=\frac{1}{N}\mathbf{1}_N,\mathbf{T}^\top\mathbf{1}_N=\frac{1}{K}\mathbf{1}_K$, which forces the probability distribution in $\mathbf{T}$ to be doubly normalized. Most importantly, the target distribution $\mathbf{T}$ should optimize the following objective:

$$\min_{\mathbf{T}\in\mathcal{U}(\mathbf{g},\mathbf{c})}\ \langle\mathbf{T},\mathbf{C}\rangle_\mathrm{F}+\varepsilon\sum_{i=1}^{N}\sum_{k=1}^{K}\mathbf{T}[i,k]\log\mathbf{T}[i,k], \tag{22}$$

where $\mathbf{C}=-\log\mathbf{Q}$ and $\langle\cdot,\cdot\rangle_\mathrm{F}$ denotes the Frobenius dot-product. The first term corresponds to the cross-entropy loss, encouraging $\mathbf{T}$ to align with high-confidence predictions, whereas the second term acts as an entropic regularizer with coefficient ε, which ensures the problem is strictly convex and solvable. The solution space is restricted to the *transport polytope* $\mathcal{U}(\mathbf{g},\mathbf{c})=\{\mathbf{T}\in\mathbb{R}_+^{N\times K}\mid\mathbf{T}\mathbf{1}_K=\mathbf{g},\mathbf{T}^\top\mathbf{1}_N=\mathbf{c}\}$. By setting the marginals to $\mathbf{g}=\frac{1}{N}\mathbf{1}_N$ and $\mathbf{c}=\frac{1}{K}\mathbf{1}_K$, we strictly enforce the equipartition constraint, theoretically guaranteeing that no cluster collapses.

Lemma 5.4. *Let* $\mathbf{u}=\mathbf{g}\oslash(\mathbf{K}\mathbf{v})$ *and* $\mathbf{v}=\mathbf{c}\oslash(\mathbf{K}^\top\mathbf{u})$, *where* $\mathbf{K}=\exp\left(\frac{\log\mathbf{Q}}{\varepsilon}\right)$ *is the kernel matrix and* $\oslash$ *denotes the element-wise division. The optimal solution to Eq. (22) is* $\mathbf{T}=\mathrm{diag}(\mathbf{u})\cdot\mathbf{K}\cdot\mathrm{diag}(\mathbf{v})$.

According to Lemma 5.4, the target distribution $\mathbf{T}$ can be constructed through a dual scaling of the kernel matrix $\mathbf{K}$, which is efficiently solved using the celebrated *Sinkhorn-Knopp algorithm* [17] as outlined in Algorithm 1.

5.5 Complete Algorithm and Analysis

Algorithm 2 summarizes the training pipeline of S^{3}GLC. Given a graph dataset $\mathcal{G}$ and hyperparameters (K,d,r,h,T), the method first initializes the GINs parameters θ and the projection matrices $(\mathbf{W}_\mathrm{H},\mathbf{W}_\mathrm{F})$, and then computes Spectral Signatures $\mathbf{g}^{(i)}$ for each graph. For each epoch, it extracts Spatial Representations $\mathbf{H}^{(i)}$ using

Algorithm 1: Sinkhorn-Knopp Normalization [17]

Input: $\mathbf{Q} \in \mathbb{R}^{N\times K}$, ε, E.
Output: $\mathbf{T} \in \mathbb{R}^{N\times K}$.

Calculate $\mathbf{K} = \exp(\frac{1}{\varepsilon}\log \mathbf{Q})$;
Initialize $\mathbf{v} \leftarrow \mathbf{1}_K$;
Define $\mathbf{g} = \frac{1}{N}\mathbf{1}_N$ and $\mathbf{c} = \frac{1}{K}\mathbf{1}_K$;
for *step* $= 1$ **to** E **do**
  $\mathbf{u} \leftarrow \mathbf{g} \oslash (\mathbf{K}\mathbf{v})$;
  $\mathbf{v} \leftarrow \mathbf{c} \oslash (\mathbf{K}^\top \mathbf{u})$;
end
Calculate $\mathbf{T} = \mathrm{diag}(\mathbf{u})\mathbf{K}\mathrm{diag}(\mathbf{v})$;
return $\mathbf{T}$

GINs with sum pooling. Next, it performs alternating optimization to fuse views by iteratively updating $(\mathbf{W}_\mathrm{H}, \mathbf{W}_\mathrm{F})$ and the consensus matrix $\mathbf{Y}$ until convergence, and then obtains the fused embedding $\mathbf{Z}$. Based on $\mathbf{Z}$, it computes the topology-aware contrastive loss $\mathcal{L}_{\mathrm{GCL}}$, derives soft cluster assignments $\mathbf{Q}$, and computes an OT-based transport plan to enforce balanced clustering, yielding $\mathcal{L}_{\mathrm{BGC}}$. The parameters θ are updated by backpropagation with the joint loss. After training, hard clustering labels C are obtained from $\mathbf{T}$ by row-wise argmax (one-hot).

Algorithm 2: The pipeline of S^{3}GLC

Input: Dataset $\mathcal{G}$, Clusters K, d, r, h, E.
Output: Clustering results C.
`/* Initialization & Preprocessing */`
Initialize θ; Initialize $\mathbf{W}_H, \mathbf{W}_F$ randomly;
Compute $\mathbf{F}^{(i)}$ via Eq. (10) $\forall \mathcal{G}_i \in \mathcal{G}$;
`/* Joint Training Loop */`
for *step* $= 1$ **to** E **do**
  $\mathbf{H}[i] \leftarrow \mathtt{SUMPOOL}\,(\mathrm{GIN}(\mathcal{G};\theta))\ \forall \mathcal{G}_i \in \mathcal{G}$;
  repeat
    Update $\mathbf{W}_\mathrm{H}, \mathbf{W}_\mathrm{F}$ via Eq. (15);
    Update $\mathbf{Y}$ via Eq. (17);
  until *Convergence*;
  Calculate $\mathbf{Z}$ via Eq. (7);
  Compute $\mathcal{L}_{GCL}$ via Eq. (19);
  Obtain $\mathbf{Q}$ via Eq (21);
  Obtain $\mathbf{T}$ via **Algorithm 1**;
  Compute $\mathcal{L}_{BGC}$ via Eq. (20) ;
  Update θ via Backpropagation;
end
Invoke $\mathbf{T}$ to C;
return C

Complexity Analysis. We analyze the complexity of S^{3}GLC based on the complete algorithm above. Following [65], and assuming a as the average number of nodes per graph, the complexity of GINs is $O(Nah)$. The updates of $(\mathbf{W}_\mathrm{H}, \mathbf{W}_\mathrm{F})$ and $\mathbf{Y}$ are performed within mini-batches of size B. Specifically, constructing the affinity matrix and updating $\mathbf{Y}$ involve operations on $B \times B$ matrices, costing $B^2 d$ per batch, the total complexity for the fusion module is $O(NBd)$. Similarly, the cost for $\mathcal{L}_{\mathrm{GCL}}$ is reduced to linear dependence on N via batching. The number of Sinkhorn-Knopp iterations is set to a constant, each of which operate on an $N \times K$ matrix and thus costs $O(NK)$. Overall, the time complexity of S^{3}GLC is $O(NBd)$, which is linear with respect to N.

Table 2: The statistics of tested datasets

Dataset	#Graphs	#Avg. Nodes	#Avg. Edges	#Clusters
BZR	405	35.8	38.4	2
MUTAG	188	17.9	19.8	2
PTC-MR	344	14.3	14.7	2
PTC-MM	336	13.97	14.32	2
ENZYMES	600	32.6	62.1	6
DD	1178	284.3	715.7	2
AIDS	2000	15.7	16.2	2
COLLAB	5, 000	74.5	2457.8	3
REDDIT-12K	11, 929	391.4	456.9	11
OGBG-PPA	20, 000	238.7	2266.1	13

6 Experiments

This section experimentally evaluates S^{3}GLC against 15 competitors regarding clustering quality and efficiency on 10 real GLDs of varied volumes. All experiments are conducted on a Linux machine with an NVIDIA Ampere A100 GPU (80 GB memory), AMD EPYC 7513 CPUs (2.6 GHz), and 1TB RAM. The codes of all algorithms are collected from their respective authors.

6.1 Experimental Setup

Datasets. We evaluate across three dataset families [45, 55, 69]: molecular datasets (*BZR*, *MUTAG*, *PTC-MR*, *PTC-MM*, *AIDS*), biological datasets (*ENZYMES*, *DD*, *OGBG-PPA*), and social networks datasets (*COLLAB*, *REDDIT-12K*). All datasets include node attributes except *COLLAB*. The class labels range from binary to multi-class (e.g., *ENZYMES*: 6, *REDDIT-12K*: 11, *COLLAB*: 3). Relatively large collections by graph count are *REDDIT-12K* and *OGBG-PPA* (>10k graphs). For detailed statistics of these datasets, please refer to Table 2.

Baselines. We select methods with state-of-the-art results and implement them using their recommended hyperparameter settings. The baselines fall into three categories: (1) Graph Kernels (GK): `SP` [8], `RW` [43], `WL` [50],`LT` [28], `GLK` [51]; (2) Unsupervised Graph Representation Learning (UGRL): `InfoGraph` [53], `GraphCL` [73], `CuCo` [16], `MVGRL` [23], `GZL` [41], and (3) Graph-Level Clustering (GLC): `GWF` [64], `GLCC` [30], `DGLC` [11], `DCGLC` [12], `GPC` [14]. For GK and UGRL, clustering is performed with KMeans on the eigenvectors of kernel-based similarity matrix and learned graph embeddings, respectively.

6.2 Clustering Quality Evaluation

We conduct a comprehensive evaluation of clustering quality on 10 graph datasets against 15 baselines. When a method encounters out-of-memory (OOM) or its runtime exceeds 48 h on a dataset, we mark the entry as "−". We categorize datasets by the number of graphs, datasets with more than (resp. less than) 1000 graphs are categorized as Large datasets (resp. Small datasets). For each metric column, the top-1 result is highlighted in dark blue and the top-2 result in

Table 3: Clustering quality on small datasets (best is highlighted in blue and best baseline underlined).

Method	*PTC-MM*			*BZR*			*PTC-MR*			*ENZYMES*			*MUTAG*		
	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑
SP [8]	$61.0_{\pm 0.0}$	$0.8_{\pm 0.0}$	$2.6_{\pm 0.0}$	$52.5_{\pm 0.0}$	$1.6_{\pm 0.0}$	$2.7_{\pm 0.0}$	$60.4_{\pm 0.0}$	$1.3_{\pm 0.0}$	$1.5_{\pm 0.0}$	$23.9_{\pm 0.0}$	$5.0_{\pm 0.0}$	$2.0_{\pm 0.0}$	$68.2_{\pm 0.0}$	$28.7_{\pm 0.0}$	$11.8_{\pm 0.0}$
RW [43]	$65.3_{\pm 0.0}$	$0.2_{\pm 0.0}$	$0.8_{\pm 0.0}$	$73.6_{\pm 0.0}$	$0.3_{\pm 0.0}$	$1.2_{\pm 0.0}$	$59.2_{\pm 0.0}$	$0.4_{\pm 0.0}$	$0.1_{\pm 0.0}$	$17.7_{\pm 0.0}$	$1.9_{\pm 0.0}$	$0.0_{\pm 0.0}$	$68.9_{\pm 0.0}$	$0.3_{\pm 0.0}$	$0.9_{\pm 0.0}$
WL [50]	$62.3_{\pm 0.0}$	$0.2_{\pm 0.0}$	$1.8_{\pm 0.0}$	$50.7_{\pm 0.0}$	$1.4_{\pm 0.0}$	$1.2_{\pm 0.0}$	$55.3_{\pm 0.0}$	$0.1_{\pm 0.0}$	$-0.1_{\pm 0.0}$	$23.5_{\pm 0.0}$	$3.7_{\pm 0.0}$	$1.7_{\pm 0.0}$	$71.9_{\pm 0.0}$	$25.4_{\pm 0.0}$	$18.4_{\pm 0.0}$
LT [28]	$64.6_{\pm 0.3}$	$2.2_{\pm 0.5}$	$5.0_{\pm 1.1}$	$78.3_{\pm 3.8}$	$12.1_{\pm 0.2}$	$18.2_{\pm 1.4}$	$59.2_{\pm 1.2}$	$0.9_{\pm 0.0}$	$2.2_{\pm 0.1}$	$24.1_{\pm 0.04}$	$4.0_{\pm 0.1}$	$1.2_{\pm 0.1}$	$55.6_{\pm 2.8}$	$5.6_{\pm 0.8}$	$0.6_{\pm 0.2}$
GLK [51]	$63.3_{\pm 0.0}$	$0.1_{\pm 0.0}$	$0.9_{\pm 0.0}$	$69.8_{\pm 0.0}$	$8.1_{\pm 0.0}$	$7.7_{\pm 0.0}$	$57.9_{\pm 0.0}$	$0.1_{\pm 0.0}$	$0.0_{\pm 0.0}$	$26.1_{\pm 0.0}$	$5.2_{\pm 0.0}$	$2.8_{\pm 0.0}$	$67.4_{\pm 0.0}$	$23.8_{\pm 0.0}$	$11.8_{\pm 0.0}$
InfoGraph [53]	$61.8_{\pm 1.1}$	$2.5_{\pm 0.9}$	$3.8_{\pm 1.6}$	$62.8_{\pm 1.8}$	$1.2_{\pm 0.3}$	$2.4_{\pm 0.2}$	$57.5_{\pm 0.8}$	$0.6_{\pm 0.2}$	$1.4_{\pm 0.3}$	$23.9_{\pm 0.8}$	$4.4_{\pm 0.2}$	$2.3_{\pm 0.1}$	$77.8_{\pm 1.6}$	$34.2_{\pm 2.8}$	$30.5_{\pm 2.4}$
GraphCL [73]	$58.9_{\pm 0.6}$	$0.4_{\pm 0.2}$	$0.8_{\pm 0.3}$	$73.2_{\pm 1.3}$	$8.5_{\pm 0.7}$	$10.3_{\pm 1.4}$	$60.0_{\pm 0.2}$	$1.0_{\pm 0.0}$	$2.4_{\pm 0.1}$	$21.2_{\pm 0.7}$	$3.1_{\pm 0.3}$	$0.9_{\pm 0.1}$	$74.7_{\pm 1.8}$	$23.7_{\pm 2.3}$	$22.7_{\pm 0.8}$
CuCo [16]	$60.2_{\pm 0.3}$	$0.3_{\pm 0.0}$	$2.2_{\pm 0.9}$	$68.3_{\pm 0.1}$	$1.1_{\pm 0.1}$	$0.6_{\pm 0.0}$	$53.6_{\pm 0.1}$	$0.0_{\pm 0.0}$	$0.0_{\pm 0.0}$	$23.9_{\pm 0.3}$	$4.8_{\pm 0.2}$	$2.1_{\pm 0.1}$	$73.3_{\pm 0.4}$	$29.8_{\pm 1.2}$	$21.1_{\pm 0.4}$
MVGRL [23]	$61.6_{\pm 0.7}$	$0.2_{\pm 0.0}$	$0.7_{\pm 0.2}$	$62.2_{\pm 0.3}$	$1.7_{\pm 0.0}$	$2.7_{\pm 0.1}$	$55.3_{\pm 0.1}$	$0.0_{\pm 0.0}$	$0.0_{\pm 0.0}$	$24.5_{\pm 0.3}$	$5.0_{\pm 0.2}$	$2.0_{\pm 0.0}$	$68.9_{\pm 0.7}$	$22.5_{\pm 2.4}$	$13.3_{\pm 1.2}$
GZL [41]	$54.8_{\pm 1.1}$	$2.0_{\pm 0.1}$	$0.1_{\pm 0.0}$	$71.3_{\pm 0.1}$	$13.1_{\pm 0.8}$	$4.1_{\pm 0.2}$	$59.0_{\pm 0.0}$	$4.9_{\pm 0.0}$	$2.1_{\pm 0.0}$	$25.2_{\pm 0.6}$	$5.5_{\pm 0.2}$	$1.6_{\pm 0.2}$	$69.7_{\pm 0.1}$	$6.9_{\pm 0.0}$	$7.0_{\pm 0.0}$
GWF [64]	$56.3_{\pm 2.7}$	$1.4_{\pm 0.7}$	$2.1_{\pm 0.7}$	$58.8_{\pm 1.2}$	$0.9_{\pm 0.2}$	$2.0_{\pm 0.1}$	$58.1_{\pm 0.3}$	$1.5_{\pm 0.3}$	$2.3_{\pm 0.4}$	$21.5_{\pm 1.3}$	$1.5_{\pm 0.2}$	$0.2_{\pm 0.0}$	$77.7_{\pm 5.2}$	$29.5_{\pm 3.4}$	$30.3_{\pm 6.8}$
GLCC [30]	$61.8_{\pm 3.1}$	$1.2_{\pm 0.8}$	$1.8_{\pm 1.1}$	$63.3_{\pm 6.4}$	$3.3_{\pm 0.2}$	$2.4_{\pm 0.1}$	$55.6_{\pm 0.2}$	$1.3_{\pm 0.1}$	$0.9_{\pm 0.0}$	$23.6_{\pm 0.7}$	$2.2_{\pm 0.2}$	$0.6_{\pm 0.1}$	$63.1_{\pm 1.2}$	$6.8_{\pm 1.3}$	$4.5_{\pm 0.6}$
DGLC [11]	$63.8_{\pm 0.2}$	$2.7_{\pm 0.5}$	$5.5_{\pm 0.6}$	$81.2_{\pm 1.1}$	$11.4_{\pm 0.3}$	$20.4_{\pm 0.3}$	$59.5_{\pm 2.7}$	$2.4_{\pm 0.6}$	$3.3_{\pm 0.8}$	$24.1_{\pm 1.2}$	$5.4_{\pm 0.3}$	$1.6_{\pm 0.2}$	$84.4_{\pm 2.1}$	$33.1_{\pm 1.3}$	$46.1_{\pm 1.8}$
DCGLC [12]	$64.9_{\pm 0.7}$	$3.8_{\pm 1.2}$	$6.0_{\pm 0.6}$	$81.7_{\pm 1.7}$	$15.2_{\pm 1.1}$	$20.6_{\pm 1.3}$	$58.2_{\pm 1.4}$	$2.1_{\pm 0.1}$	$2.3_{\pm 0.3}$	$26.6_{\pm 1.1}$	$5.1_{\pm 0.6}$	$2.4_{\pm 0.7}$	$86.7_{\pm 0.8}$	$42.4_{\pm 3.2}$	$53.2_{\pm 2.3}$
GPC [14]	$61.5_{\pm 1.2}$	$4.6_{\pm 0.3}$	$4.9_{\pm 0.4}$	$81.0_{\pm 0.0}$	$9.3_{\pm 0.0}$	$18.6_{\pm 0.0}$	$51.8_{\pm 0.1}$	$2.3_{\pm 0.0}$	$2.1_{\pm 0.0}$	$26.8_{\pm 0.3}$	$6.9_{\pm 0.1}$	$3.4_{\pm 0.1}$	$75.5_{\pm 0.5}$	$31.4_{\pm 0.2}$	$25.6_{\pm 0.1}$
S^3GLC	$67.9_{\pm 0.6}$	$3.3_{\pm 0.9}$	$6.2_{\pm 0.5}$	$79.7_{\pm 0.0}$	$16.1_{\pm 0.0}$	$23.5_{\pm 0.2}$	$64.0_{\pm 1.5}$	$6.2_{\pm 2.8}$	$6.0_{\pm 1.6}$	$36.5_{\pm 2.3}$	$26.0_{\pm 2.3}$	$11.3_{\pm 2.2}$	$87.4_{\pm 0.8}$	$41.4_{\pm 2.0}$	$55.0_{\pm 2.4}$
Improv.	+2.6	-1.3	+0.2	-2.0	+0.9	+2.9	+3.6	+1.3	+2.7	+9.7	+19.1	+7.9	+0.7	-1.0	+1.8

Table 4: Clustering quality on large datasets (best is highlighted in blue and best baseline underlined).

Method	*AIDS*			*DD*			*COLLAB*			*REDDIT-12K*			*OGBG-PPA*		
	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑	ACC ↑	NMI ↑	ARI ↑
SP [8]	$83.6_{\pm 0.0}$	$35.4_{\pm 0.0}$	$38.0_{\pm 0.0}$	$58.5_{\pm 0.0}$	$0.2_{\pm 0.0}$	$0.1_{\pm 0.0}$	$55.7_{\pm 0.0}$	$8.1_{\pm 0.0}$	$2.9_{\pm 0.0}$	$19.2_{\pm 0.0}$	$7.8_{\pm 0.0}$	$1.3_{\pm 0.0}$	$15.6_{\pm 0.0}$	$6.7_{\pm 0.0}$	$1.8_{\pm 0.0}$
RW [43]	$72.3_{\pm 0.0}$	$7.6_{\pm 0.0}$	$5.7_{\pm 0.0}$	$58.7_{\pm 0.0}$	$0.1_{\pm 0.0}$	$0.2_{\pm 0.0}$	$52.1_{\pm 0.0}$	$1.1_{\pm 0.0}$	$0.1_{\pm 0.0}$	–	–	–	–	–	–
WL [50]	$87.2_{\pm 0.0}$	$45.7_{\pm 0.0}$	$50.7_{\pm 0.0}$	$58.4_{\pm 0.0}$	$0.4_{\pm 0.0}$	$-0.1_{\pm 0.0}$	$55.4_{\pm 0.0}$	$6.7_{\pm 0.0}$	$2.8_{\pm 0.0}$	$20.6_{\pm 0.0}$	$12.7_{\pm 0.0}$	$5.1_{\pm 0.0}$	$15.9_{\pm 0.0}$	$7.2_{\pm 0.0}$	$2.5_{\pm 0.0}$
LT [28]	$53.2_{\pm 1.3}$	$5.6_{\pm 0.6}$	$3.4_{\pm 0.1}$	–	–	–	–	–	–	–	–	–	–	–	–
GLK [51]	$83.5_{\pm 0.0}$	$35.2_{\pm 0.0}$	$37.1_{\pm 0.0}$	$57.9_{\pm 0.0}$	$1.5_{\pm 0.0}$	$1.1_{\pm 0.0}$	–	–	–	$18.3_{\pm 0.0}$	$7.3_{\pm 0.0}$	$0.0_{\pm 0.0}$	–	–	–
InfoGraph [53]	$82.8_{\pm 2.7}$	$26.0_{\pm 2.4}$	$37.2_{\pm 3.8}$	$58.4_{\pm 0.0}$	$0.3_{\pm 0.0}$	$0.0_{\pm 0.0}$	$56.7_{\pm 1.1}$	$16.3_{\pm 1.4}$	$7.3_{\pm 0.9}$	$22.3_{\pm 1.1}$	$4.2_{\pm 0.3}$	$0.5_{\pm 0.2}$	$13.2_{\pm 0.2}$	$1.7_{\pm 0.4}$	$0.1_{\pm 0.0}$
GraphCL [73]	$84.8_{\pm 0.5}$	$38.7_{\pm 0.6}$	$42.7_{\pm 1.1}$	$57.6_{\pm 1.3}$	$1.5_{\pm 0.4}$	$0.0_{\pm 0.0}$	$55.2_{\pm 0.3}$	$9.1_{\pm 0.1}$	$3.1_{\pm 0.1}$	$22.1_{\pm 0.6}$	$3.3_{\pm 0.0}$	$0.3_{\pm 0.0}$	$16.2_{\pm 0.8}$	$4.3_{\pm 0.3}$	$1.1_{\pm 0.2}$
CuCo [16]	$84.1_{\pm 0.6}$	$36.7_{\pm 0.2}$	$39.6_{\pm 0.3}$	$58.4_{\pm 0.3}$	$0.4_{\pm 0.1}$	$0.0_{\pm 0.0}$	$56.7_{\pm 0.5}$	$9.5_{\pm 1.1}$	$3.4_{\pm 0.7}$	$22.4_{\pm 0.1}$	$4.2_{\pm 0.0}$	$0.5_{\pm 0.0}$	$14.1_{\pm 0.5}$	$3.0_{\pm 0.4}$	$0.8_{\pm 0.2}$
MVGRL [23]	$87.4_{\pm 2.1}$	$46.3_{\pm 1.4}$	$51.3_{\pm 1.4}$	–	–	–	–	–	–	–	–	–	–	–	–
GZL [41]	$78.4_{\pm 1.2}$	$26.5_{\pm 0.7}$	$15.3_{\pm 1.2}$	$62.3_{\pm 0.0}$	$3.3_{\pm 0.0}$	$4.2_{\pm 0.0}$	$59.3_{\pm 0.7}$	$18.6_{\pm 0.6}$	$15.9_{\pm 0.7}$	$19.2_{\pm 1.1}$	$11.2_{\pm 0.5}$	$5.2_{\pm 0.3}$	$17.4_{\pm 0.1}$	$7.9_{\pm 0.0}$	$2.8_{\pm 0.0}$
GWF [64]	$82.1_{\pm 5.4}$	$31.8_{\pm 4.9}$	$39.1_{\pm 6.4}$	$50.8_{\pm 0.3}$	$0.3_{\pm 0.0}$	$0.6_{\pm 0.1}$	$48.5_{\pm 3.2}$	$9.2_{\pm 2.1}$	$2.4_{\pm 0.3}$	–	–	–	–	–	–
GLCC [30]	$70.7_{\pm 3.4}$	$6.1_{\pm 1.3}$	$0.1_{\pm 0.0}$	$64.8_{\pm 3.3}$	$3.2_{\pm 0.4}$	$3.5_{\pm 0.2}$	$52.8_{\pm 3.3}$	$13.5_{\pm 2.2}$	$9.6_{\pm 1.8}$	$22.6_{\pm 0.2}$	$10.8_{\pm 0.4}$	$5.6_{\pm 0.0}$	$16.1_{\pm 0.4}$	$4.3_{\pm 0.3}$	$2.4_{\pm 0.1}$
DGLC [11]	$88.5_{\pm 3.4}$	$39.8_{\pm 2.3}$	$47.6_{\pm 2.3}$	$58.7_{\pm 1.4}$	$1.7_{\pm 0.3}$	$2.1_{\pm 0.2}$	$61.2_{\pm 2.1}$	$23.5_{\pm 1.8}$	$16.0_{\pm 0.9}$	$23.2_{\pm 0.6}$	$11.5_{\pm 1.4}$	$7.2_{\pm 0.7}$	$14.9_{\pm 0.8}$	$3.6_{\pm 0.3}$	$0.0_{\pm 0.0}$
DCGLC [12]	$95.5_{\pm 2.3}$	$64.8_{\pm 3.4}$	$78.4_{\pm 3.3}$	$74.2_{\pm 1.3}$	$12.8_{\pm 0.1}$	$20.3_{\pm 0.2}$	$55.2_{\pm 0.0}$	$9.1_{\pm 0.0}$	$3.1_{\pm 0.0}$	$22.2_{\pm 0.1}$	$3.2_{\pm 0.0}$	$0.2_{\pm 0.0}$	$17.6_{\pm 0.3}$	$8.2_{\pm 1.4}$	$6.7_{\pm 1.4}$
GPC [14]	$84.2_{\pm 0.3}$	$20.6_{\pm 0.4}$	$23.8_{\pm 0.3}$	$60.9_{\pm 0.2}$	$3.2_{\pm 0.0}$	$2.1_{\pm 0.0}$	$60.4_{\pm 0.3}$	$19.9_{\pm 0.1}$	$11.5_{\pm 0.1}$	$21.9_{\pm 0.1}$	$9.3_{\pm 0.3}$	$1.7_{\pm 0.0}$	$16.1_{\pm 0.2}$	$4.0_{\pm 0.1}$	$1.7_{\pm 0.1}$
S^3GLC	$99.3_{\pm 0.2}$	$93.6_{\pm 1.6}$	$96.9_{\pm 0.9}$	$84.1_{\pm 0.7}$	$48.0_{\pm 3.6}$	$46.4_{\pm 1.8}$	$69.3_{\pm 0.5}$	$29.5_{\pm 2.9}$	$25.5_{\pm 1.1}$	$37.8_{\pm 0.6}$	$39.5_{\pm 1.2}$	$18.9_{\pm 0.5}$	$25.8_{\pm 2.1}$	$24.8_{\pm 2.8}$	$10.9_{\pm 1.7}$
Improv.	+3.8	+28.8	+18.5	+9.9	+35.2	+26.1	+8.1	+6.0	+9.5	+14.6	+26.8	+11.7	+8.2	+16.6	+4.2

light blue; the best-performing baseline is underlined. The last row ("Improv.") reports the absolute improvement (in percentage points) of our method over the strongest baseline in that column.

S^3GLC establishes a new state-of-the-art by delivering substantial improvements across datasets of varying sizes and complexities. Specifically, it achieves remarkable boosts on widely used binary datasets in UGRL and GLC. For instance, on *AIDS* and *DD*, it increases NMI by 28.8% and 35.2%, and ARI by 18.5% and 26.1%, respectively, culminating in a near-perfect 99.3% ACC on *AIDS*. Unlike many GLC methods that yield limited gains on complex multi-class tasks, our approach proves highly effective. On *ENZYMES*, a representative multi-class benchmark, S^3GLC markedly outperforms the runner-up method GPC by improving ACC, NMI, and ARI by 9.7%, 19.1%, and 7.9%, respectively.

Most notably, S^3GLC demonstrates robust scalability on large graphs where SOTA baselines frequently fail or degrade drastically. On massive datasets with over ten classes, such as *REDDIT-12K* and *OGBG-PPA*, our method consistently maintains leading capabilities. Specifically, it improves ACC by 14.6% and 8.2%, NMI by 26.8% and 16.6%, and ARI by 11.7% and 4.2% on *REDDIT-12K* and *OGBG-PPA*, respectively, validating its superiority in handling large-scale graph clustering.

6.3 Ablation Study

To verify the effectiveness of each critical component within the S^3GLC framework, we conducted a comprehensive ablation analysis as presented in Table 5 and Table 6.

As discussed in §4.2, multi-class datasets are prone to clustering collapse. Empirically, removing $\mathcal{L}_{\text{BGC}}$ causes substantial degradation: ACC drops by 5.3% on *ENZYMES* (6) and 2.5% on *REDDIT-12K* (11), whereas fluctuations are negligible on binary datasets like *MUTAG*. This confirms that balanced constraints via Sinkhorn-Knopp algorithm successfully prevent trivial solutions. An intuitive manifestation of the cluster-balancing effect induced by $\mathcal{L}_{\text{BGC}}$ can be found in the visualization results reported in Appendix B.4.

To evaluate the JMPSR, we replaced it with standard MLP reduction. Results reveal significant declines on complex datasets like *DD* and *REDDIT-12K*, where ACC decreases by 10.7% and 13.2%, respectively. As designed in §5.3, our mechanism filters noise while preserving complementary semantics between Spatial Representation and Spectral Signature. Furthermore, removing $\mathcal{L}_{GCL}$ degrades clustering consistency; on *ENZYMES*, the NMI of w/o $\mathcal{L}_{GCL}$ drops by 7.6%.

To demonstrate the effectiveness of Spectral Signature, we replaced the global structural features with common GK methods, and

Table 5: Ablation study.

Dataset		$\mathcal{L}_{\text{BGC}}$	w/o $\mathcal{L}_{\text{GCL}}$	JMPSR	S^{3}GLC
ENZYMES	ACC ↑	$31.2_{\pm 2.7}$	$34.3_{\pm 3.2}$	$26.9_{\pm 0.8}$	$36.5_{\pm 2.3}$
	NMI ↑	$18.4_{\pm 6.3}$	$18.4_{\pm 6.2}$	$5.8_{\pm 0.7}$	$26.0_{\pm 2.3}$
	ARI ↑	$7.3_{\pm 2.8}$	$9.9_{\pm 3.8}$	$3.1_{\pm 0.4}$	$11.3_{\pm 2.2}$
MUTAG	ACC ↑	$86.5_{\pm 0.9}$	$86.2_{\pm 0.4}$	$86.5_{\pm 0.5}$	$87.4_{\pm 0.8}$
	NMI ↑	$39.4_{\pm 2.2}$	$38.7_{\pm 0.9}$	$38.8_{\pm 1.7}$	$41.4_{\pm 2.0}$
	ARI ↑	$51.8_{\pm 3.0}$	$51.3_{\pm 1.4}$	$52.2_{\pm 1.7}$	$55.0_{\pm 2.4}$
DD	ACC ↑	$82.9_{\pm 1.2}$	$82.8_{\pm 1.1}$	$73.4_{\pm 2.3}$	$84.1_{\pm 0.7}$
	NMI ↑	$38.0_{\pm 6.3}$	$43.4_{\pm 5.0}$	$26.4_{\pm 4.3}$	$48.0_{\pm 3.6}$
	ARI ↑	$43.2_{\pm 3.2}$	$43.1_{\pm 2.9}$	$22.3_{\pm 3.8}$	$46.4_{\pm 1.8}$
REDDIT-12K	ACC ↑	$35.3_{\pm 3.0}$	$37.7_{\pm 2.5}$	$24.6_{\pm 1.5}$	$37.8_{\pm 0.6}$
	NMI ↑	$39.2_{\pm 1.7}$	$39.2_{\pm 1.7}$	$15.8_{\pm 0.6}$	$39.5_{\pm 1.2}$
	ARI ↑	$17.5_{\pm 0.9}$	$18.2_{\pm 1.3}$	$6.4_{\pm 0.7}$	$18.9_{\pm 0.5}$

Table 6: The effectiveness of Spectral Signature.

Dataset	Structural Features	ACC ↑	NMI ↑	ARI ↑
ENZYMES	SP [8]	$33.2_{\pm 2.2}$	$16.3_{\pm 2.7}$	$10.1_{\pm 2.5}$
	RW [43]	$26.4_{\pm 2.3}$	$6.7_{\pm 2.3}$	$4.2_{\pm 2.2}$
	WL [50]	$32.9_{\pm 0.9}$	$13.4_{\pm 0.6}$	$9.8_{\pm 1.0}$
	GLK [51]	$33.9_{\pm 2.5}$	$15.3_{\pm 4.0}$	$9.7_{\pm 2.4}$
	ours	$\mathbf{36.5}_{\pm 2.3}$	$\mathbf{26.0}_{\pm 2.3}$	$\mathbf{11.3}_{\pm 2.2}$
MUTAG	SP [8]	$82.2_{\pm 3.8}$	$31.7_{\pm 7.0}$	$25.2_{\pm 7.5}$
	RW [43]	$82.7_{\pm 3.9}$	$24.8_{\pm 9.7}$	$30.7_{\pm 11.0}$
	WL [50]	$84.5_{\pm 1.9}$	$34.2_{\pm 2.4}$	$41.6_{\pm 3.8}$
	GLK [51]	$81.0_{\pm 3.1}$	$29.4_{\pm 11.6}$	$35.1_{\pm 6.5}$
	ours	$\mathbf{87.4}_{\pm 0.8}$	$\mathbf{41.4}_{\pm 2.0}$	$\mathbf{55.0}_{\pm 2.4}$
DD	SP [8]	$71.0_{\pm 5.3}$	$21.4_{\pm 8.3}$	$18.5_{\pm 8.8}$
	RW [43]	$72.3_{\pm 5.6}$	$20.2_{\pm 12.7}$	$20.5_{\pm 10.8}$
	WL [50]	$68.7_{\pm 3.7}$	$12.8_{\pm 6.0}$	$14.4_{\pm 6.1}$
	GLK [51]	$69.5_{\pm 6.3}$	$20.2_{\pm 8.0}$	$16.4_{\pm 9.8}$
	ours	$\mathbf{84.1}_{\pm 0.7}$	$\mathbf{48.0}_{\pm 3.6}$	$\mathbf{46.4}_{\pm 1.8}$
REDDIT-12K	SP [8]	$23.2_{\pm 1.1}$	$16.5_{\pm 2.1}$	$4.4_{\pm 1.3}$
	RW [43]	-	-	-
	WL [50]	$22.4_{\pm 1.2}$	$11.2_{\pm 3.4}$	$2.1_{\pm 0.8}$
	GLK [51]	$20.7_{\pm 0.8}$	$8.8_{\pm 2.3}$	$1.3_{\pm 0.5}$
	ours	$\mathbf{37.8}_{\pm 0.6}$	$\mathbf{39.5}_{\pm 1.2}$	$\mathbf{18.9}_{\pm 0.5}$

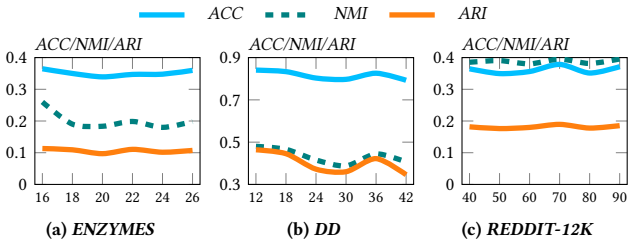

(a) *ENZYMES* (b) *DD* (c) *REDDIT-12K*

Figure 5: Varying r

the results are presented in Table 6. It can be seen that our method comprehensively outperforms traditional GK methods. Notably, on *DD*, compared to the best-performing SP, our method improves NMI by 26.6%, demonstrating that our method effectively captures global topological differences ignored by Spatial Representation and traditional GK methods.

6.4 Hyperparameter Analysis

We investigate the sensitivity of two pivotal hyperparameters: the Spectral Signature dimension r and the fused feature dimension d. As illustrated in Fig. (5), the model demonstrates strong robustness to variations in r, particularly on large-scale datasets like *REDDIT-12K*. This corroborates our theoretical basis relying on the Hoffman-Wielandt inequality: the top-r eigenvalues suffice as

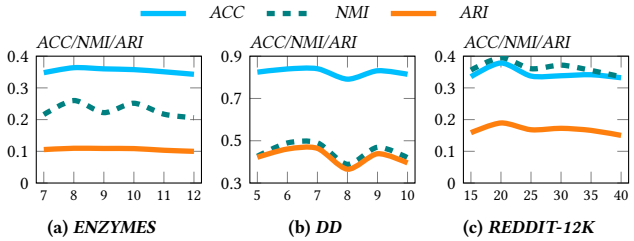

(a) *ENZYMES* (b) *DD* (c) *REDDIT-12K*

Figure 6: Varying d

a rigorous lower bound for GED to capture core global topology without introducing excessive noise. Regarding d in Fig. (6), the performance exhibits minimal fluctuation. This stability is attributed to the JMPSR, which effectively filters multi-view redundancy, ensuring the learning of a compact and consistent manifold representation Y regardless of dimensional shifts.

6.5 Efficiency Analysis

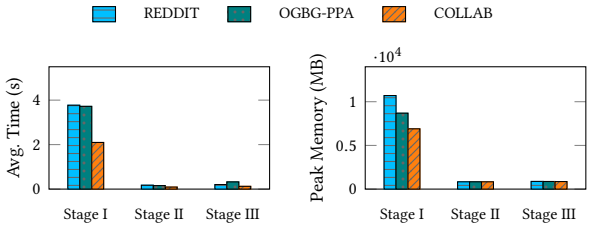

(a) ***Runtime Analysis*** (b) ***Memory Usage Analysis***

Figure 7: Efficiency analysis across training stages. (a) Average runtime per epoch. (b) Peak memory usage.

Figure 7 presents a detailed efficiency analysis of S^{3}GLC. Specifically, Fig. 7 (a) reports the average runtime per mini-batch in each epoch, while Fig. 7 (b) records the global peak memory usage during training for each epoch. As illustrated, both JMPSR (Stage II) and Balanced clustering (Stage III) incur substantially lower time and memory overheads than GIN encoding (Stage I), demonstrating that our core methodological enhancements introduce negligible complexity.

7 Conclusion

In this paper, to address the limited expressivity of GINs in GLC and the clustering collapse in multi-class settings, we propose S^{3}GLC. The method leverages Spectral Signature to capture global topology, performs JMPSR with feature fusion, and introduces a constraint to achieve balanced clustering. Nevertheless, the current design is mainly tailored to homogeneous graphs and is limited to static graph structures. In future work, we will extend the algorithm to the analysis of heterogeneous or dynamic graph data.

References

[1] José María Amigó, Jorge Gálvez, and Vincent M Villar. 2009. A review on molecular topology: applying graph theory to drug discovery and design. *Naturwissenschaften* 96, 7 (2009), 749–761.

[2] Tejaswini Ananthanarayana, Lipisha Chaudhary, and Ifeoma Nwogu. 2022. SignNet: Single Channel Sign Generation using Metric Embedded Learning. *2023 IEEE 17th International Conference on Automatic Face and Gesture Recognition (FG)* (2022), 1–7. https://api.semanticscholar.org/CorpusID:254274967

[3] Yuki M Asano, Christian Rupprecht, and Andrea Vedaldi. 2020. Self-labelling via simultaneous clustering and representation learning. In *International Conference on Learning Representations (ICLR)*. https://openreview.net/forum?id=B1esx6EYvr

[4] László Babai. 2016. Graph isomorphism in quasipolynomial time. In *Proceedings of the forty-eighth annual ACM symposium on Theory of Computing*. 684–697.

[5] Yunsheng Bai, Hao Ding, Song Bian, Ting Chen, Yizhou Sun, and Wei Wang. 2019. Simgnn: A neural network approach to fast graph similarity computation. In *Proceedings of the twelfth ACM international conference on web search and data mining*. 384–392.

[6] Mikhail Belkin and Partha Niyogi. 2003. Laplacian eigenmaps for dimensionality reduction and data representation. *Neural computation* 15, 6 (2003), 1373–1396.

[7] Pietro Bongini, Monica Bianchini, and Franco Scarselli. 2021. Molecular generative graph neural networks for drug discovery. *Neurocomputing* 450 (2021), 242–252.

[8] Karsten M. Borgwardt and Hans-Peter Kriegel. 2005. Shortest-path kernels on graphs. *Fifth IEEE International Conference on Data Mining (ICDM'05)* (2005), 8 pp.–. https://api.semanticscholar.org/CorpusID:1550330

[9] Karsten M. Borgwardt, Cheng Soon Ong, Stefan Schönauer, S. V. N. Vishwanathan, Alex Smola, and Hans-Peter Kriegel. 2005. Protein function prediction via graph kernels. *Bioinformatics* 21 Suppl 1 (2005), i47–56. https://api.semanticscholar.org/CorpusID:8174592

[10] Horst Bunke. 1997. On a relation between graph edit distance and maximum common subgraph. *Pattern Recognition Letters* 18, 8 (1997), 689–694.

[11] Jinyu Cai, Yi Han, Wenzhong Guo, and Jicong Fan. 2023. Deep graph-level clustering using pseudo-label-guided mutual information maximization network. *Neural Computing and Applications* 36 (2023), 9551 – 9566. https://api.semanticscholar.org/CorpusID:256615801

[12] Jinyu Cai, Yunhe Zhang, Jicong Fan, Yali Du, and Wenzhong Guo. 2024. Dual Contrastive Graph-Level Clustering with Multiple Cluster Perspectives Alignment. In *International Joint Conference on Artificial Intelligence*. https://api.semanticscholar.org/CorpusID:271508056

[13] Mathilde Caron, Piotr Bojanowski, Armand Joulin, and Matthijs Douze. 2018. Deep clustering for unsupervised learning of visual features. In *Proceedings of the European Conference on Computer Vision (ECCV)*. 132–149.

[14] Mansheng Chen, Pei-Yuan Lai, De-Zhang Liao, Changdong Wang, and Jian-Huang Lai. 2025. Graph Prompt Clustering. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 47 (2025), 5794–5805. https://api.semanticscholar.org/CorpusID:277189865

[15] Ting Chen, Simon Kornblith, Mohammad Norouzi, and Geoffrey Hinton. 2020. A simple framework for contrastive learning of visual representations. In *International conference on machine learning*. PmLR, 1597–1607.

[16] Guanyi Chu, Xiao Wang, Chuan Shi, and Xunqiang Jiang. 2021. CuCo: Graph Representation with Curriculum Contrastive Learning. In *International Joint Conference on Artificial Intelligence*. https://api.semanticscholar.org/CorpusID:235671564

[17] Marco Cuturi. 2013. Sinkhorn distances: Lightspeed computation of optimal transport. In *Advances in Neural Information Processing Systems (NIPS)*, Vol. 26. 2292–2300.

[18] Ky Fan. 1949. On a theorem of Weyl concerning eigenvalues of linear transformations I. *Proceedings of the National Academy of Sciences* 35, 11 (1949), 652–655.

[19] Stefan Fankhauser, Kaspar Riesen, and Horst Bunke. 2011. Speeding up graph edit distance computation through fast bipartite matching. In *International Workshop on Graph-Based Representations in Pattern Recognition*. Springer, 102–111.

[20] Thomas Gärtner, Peter Flach, and Stefan Wrobel. 2003. On graph kernels: Hardness results and efficient alternatives. In *Learning theory and kernel machines: 16th annual conference on learning theory and 7th kernel workshop, COLT/kernel 2003, washington, DC, USA, August 24-27, 2003. proceedings*. Springer, 129–143.

[21] Renda Han, Wenxuan Tu, Junlong Wu, Wenxin Zhang, Jingxin Liu, Haotian Wang, Jieren Cheng, Huajie Lei, Guangzhen Yao, Lingren Wang, Mengfei Li, and Yu Li. 2025. Dual Boost-Driven Graph-Level Clustering Network. *2025 International Joint Conference on Neural Networks (IJCNN)* (2025), 1–8. https://doi.org/10.1109/ijcnn64981.2025.11228488

[22] Renda Han, Guangzhen Yao, Wenxin Zhang, Yu Li, Wen Xin, Huajie Lei, Mengfei Li, Zeyu Zhang, and Yahe Tian. 2025. Multi-Relation Graph-Kernel Strengthen Network for Graph-Level Clustering. *2025 International Joint Conference on Neural Networks (IJCNN)* (2025), 1–8. https://doi.org/10.1109/ijcnn64981.2025.11227596

[23] Kaveh Hassani and Amir Hosein Khas Ahmadi. 2020. Contrastive Multi-View Representation Learning on Graphs. In *International Conference on Machine Learning*. https://api.semanticscholar.org/CorpusID:219558740

[24] Alan J Hoffman, Helmut W Wielandt, et al. 1953. The variation of the spectrum of a normal matrix. *Duke Math. J* 20, 1 (1953), 37–39.

[25] Chenping Hou, Feiping Nie, Xuelong Li, Dong yun Yi, and Yi Wu. 2014. Joint Embedding Learning and Sparse Regression: A Framework for Unsupervised Feature Selection. *IEEE Transactions on Cybernetics* 44 (2014), 793–804. https://api.semanticscholar.org/CorpusID:15474528

[26] Weihua Hu, Bowen Liu, Joseph Gomes, Marinka Zitnik, Percy Liang, Vijay S. Pande, and Jure Leskovec. 2019. Strategies for Pre-training Graph Neural Networks. *arXiv: Learning* (2019). https://api.semanticscholar.org/CorpusID:213085920

[27] Akihiro Inokuchi, Takashi Washio, and Hiroshi Motoda. 2000. An Apriori-Based Algorithm for Mining Frequent Substructures from Graph Data. In *European Conference on Principles of Data Mining and Knowledge Discovery*. https://api.semanticscholar.org/CorpusID:7847517

[28] Fredrik D. Johansson, Vinay Jethava, Devdatt P. Dubhashi, and Chiranjib Bhattacharyya. 2014. Global graph kernels using geometric embeddings. In *International Conference on Machine Learning*. https://api.semanticscholar.org/CorpusID:6718308

[29] Simon Jones and Eamonn O'Neill. 2010. Feasibility of structural network clustering for group-based privacy control in social networks. In *proceedings of the sixth symposium on usable privacy and security*. 1–13.

[30] Wei Ju, Yiyang Gu, Binqi Chen, Gongbo Sun, Yifang Qin, Xing Liu, Xiao Luo, and Ming Zhang. 2022. GLCC: A General Framework for Graph-level Clustering. In *AAAI Conference on Artificial Intelligence*. https://api.semanticscholar.org/CorpusID:253080427

[31] Wei Ju, Yiyang Gu, Xiao Luo, Yifan Wang, Haochen Yuan, Huasong Zhong, and Ming Zhang. 2023. Unsupervised graph-level representation learning with hierarchical contrasts. *Neural networks : the official journal of the International Neural Network Society* 158 (2023), 359–368. https://doi.org/10.1016/j.neunet.2022.11.019

[32] Sofia Ira Ktena, Sarah Parisot, Enzo Ferrante, Martin Rajchl, Matthew Lee, Ben Glocker, and Daniel Rueckert. 2017. Distance metric learning using graph convolutional networks: Application to functional brain networks. In *International conference on medical image computing and computer-assisted intervention*. Springer, 469–477.

[33] Sofia Ira Ktena, Sarah Parisot, Enzo Ferrante, Martin Rajchl, Matthew Lee, Ben Glocker, and Daniel Rueckert. 2018. Metric learning with spectral graph convolutions on brain connectivity networks. *NeuroImage* 169 (2018), 431–442.

[34] Richard B Lehoucq and Danny C Sorensen. 1996. Deflation techniques for an implicitly restarted Arnoldi iteration. *SIAM J. Matrix Anal. Appl.* 17, 4 (1996), 789–821.

[35] Andrei Leman and Boris Weisfeiler. 1968. A reduction of a graph to a canonical form and an algebra arising during this reduction. *Nauchno-Technicheskaya Informatsiya* 2, 9 (1968), 12–16.

[36] Zeyi Li, Renda Han, Tianyu Hu, Mengfei Li, and Caimao Li. 2025. Multi-View Enhancement Graph-Level Clustering Network. *IEEE Access* 13 (2025), 47996–48006. https://doi.org/10.1109/access.2025.3550185

[37] Zhipeng Li, Si-Guo Wang, Qinhu Zhang, Yi-Jie Pan, Naian Xiao, Jia-Yang Guo, Chang-An Yuan, Wen-Jian Liu, and De-Shuang Huang. 2024. Graph pooling for graph-level representation learning: a survey. *Artificial Intelligence Review* 58 (2024). https://api.semanticscholar.org/CorpusID:274944366

[38] Zhi-Peng Li, Si-Guo Wang, Qin-Hu Zhang, Yi-Jie Pan, Nai-An Xiao, Jia-Yang Guo, Chang-An Yuan, Wen-Jian Liu, and De-Shuang Huang. 2024. Graph pooling for graph-level representation learning: a survey. *Artificial Intelligence Review* 58, 2 (2024), 45.

[39] Yixin Liu, Yu Zheng, Daokun Zhang, Hongxu Chen, Hao Peng, and Shirui Pan. 2022. Towards Unsupervised Deep Graph Structure Learning. *Proceedings of the ACM Web Conference 2022* (2022). https://doi.org/10.1145/3485447.3512186

[40] Ziwen Liu, Chenguang Wang, Congying Han, and Tiande Guo. 2021. Learning graph representation by aggregating subgraphs via mutual information maximization. *Neurocomputing* 548 (2021), 126392. https://doi.org/10.1016/j.neucom.2023.126392

[41] Ziyang Liu, Chaokun Wang, Yunkai Lou, and Hao Feng. 2023. Fast Unsupervised Graph Embedding via Graph Zoom Learning. In *2023 IEEE 39th International Conference on Data Engineering (ICDE)*. 2551–2564. https://doi.org/10.1109/ICDE55515.2023.00196

[42] Xiao Luo, Wei Ju, Yiyang Gu, Zhengyan Mao, Luchen Liu, Yuhui Yuan, and Ming Zhang. 2023. Self-supervised Graph-level Representation Learning with Adversarial Contrastive Learning. *ACM Transactions on Knowledge Discovery from Data* 18 (2023), 1 – 23. https://doi.org/10.1145/3624018

[43] Pierre Mahé, Nobuhisa Ueda, Tatsuya Akutsu, Jean-Luc Perret, and Jean-Philippe Vert. 2004. Extensions of marginalized graph kernels. *Proceedings of the twenty-first international conference on Machine learning* (2004). https://api.semanticscholar.org/CorpusID:18350669

[44] Giovanna Menardi and Domenico De Stefano. 2022. Density-based clustering of social networks. *Journal of the Royal Statistical Society Series A: Statistics in Society* 185, 3 (2022), 1004–1029.

[45] Christopher Morris, Nils M. Kriege, Franka Bause, Kristian Kersting, Petra Mutzel, and Marion Neumann. 2020. TUDataset: A collection of benchmark datasets for learning with graphs. *ArXiv* abs/2007.08663 (2020). https://api.semanticscholar.org/CorpusID:220633407

[46] Zhen Peng, Wenbing Huang, Minnan Luo, Q. Zheng, Yu Rong, Tingyang Xu, and Junzhou Huang. 2020. Graph Representation Learning via Graphical Mutual Information Maximization. *Proceedings of The Web Conference 2020* (2020). https://doi.org/10.1145/3366423.3380112

[47] Patrick Rubin-Delanchy. 2020. Manifold structure in graph embeddings. *ArXiv* abs/2006.05168 (2020). https://api.semanticscholar.org/CorpusID:219558619

[48] Alberto Sanfeliu and King-Sun Fu. 1983. A distance measure between attributed relational graphs for pattern recognition. *IEEE transactions on systems, man, and cybernetics* 3 (1983), 353–362.

[49] Claude E Shannon. 1948. A mathematical theory of communication. *The Bell System Technical Journal* 27, 3 (1948), 379–423.

[50] Nino Shervashidze, Pascal Schweitzer, Erik Jan van Leeuwen, Kurt Mehlhorn, and Karsten M. Borgwardt. 2011. Weisfeiler-Lehman Graph Kernels. *J. Mach. Learn. Res.* 12 (2011), 2539–2561. https://api.semanticscholar.org/CorpusID:1797579

[51] Nino Shervashidze, S. V. N. Vishwanathan, Tobias Petri, Kurt Mehlhorn, and Karsten M. Borgwardt. 2009. Efficient graphlet kernels for large graph comparison. In *International Conference on Artificial Intelligence and Statistics*. https://api.semanticscholar.org/CorpusID:17557614

[52] Gen Shi, Yifan Zhu, Jian K. Liu, and Xia Li. 2023. HeGCL: Advance Self-Supervised Learning in Heterogeneous Graph-Level Representation. *IEEE Transactions on Neural Networks and Learning Systems* 35 (2023), 13914–13925. https://doi.org/10.1109/tnnls.2023.3273255

[53] Fan-Yun Sun, Jordan Hoffmann, and Jian Tang. 2019. InfoGraph: Unsupervised and Semi-supervised Graph-Level Representation Learning via Mutual Information Maximization. *ArXiv* abs/1908.01000 (2019). https://api.semanticscholar.org/CorpusID:199441876

[54] Xuecheng Sun, Zonghui Wang, Zheming Lu, and Ziqian Lu. 2024. Self-supervised graph representations with generative adversarial learning. *Neurocomputing* 592 (2024), 127786. https://doi.org/10.1016/j.neucom.2024.127786

[55] Damian Szklarczyk, Annika L. Gable, David Lyon, Alexander Junge, Stefan Wyder, Jaime Huerta-Cepas, Milan Simonovic, Nadezhda T. Doncheva, John H. Morris, Peer Bork, Lars Juhl Jensen, and Christian von Mering. 2018. STRING v11: protein–protein association networks with increased coverage, supporting functional discovery in genome-wide experimental datasets. *Nucleic Acids Research* 47 (2018), D607 – D613. https://api.semanticscholar.org/CorpusID:53751796

[56] Ichigaku Takigawa and Hiroshi Mamitsuka. 2013. Graph mining: procedure, application to drug discovery and recent advances. *Drug discovery today* 18, 1-2 (2013), 50–57.

[57] Rita Vázquez, Avenilde Romo, Rebeca Romo-Vázquez, and María Trigueros. 2016. La separación ciega de fuentes: un puente entre el álgebra lineal y el análisis de señales. *Educación matemática* 28, 2 (2016), 31–57.

[58] Liangtian Wan, Zhenqiang Fu, Lu Sun, Xianpeng Wang, Gang Xu, Xiaoran Yan, and Feng Xia. 2023. Self-Supervised Teaching and Learning of Representations on Graphs. *Proceedings of the ACM Web Conference 2023* (2023). https://doi.org/10.1145/3543507.3583441

[59] Ren Wang, Haoliang Sun, Yuxiu Lin, Chuanhui Zuo, Yongshun Gong, Yilong Yin, and Wenjia Meng. 2025. SeqMvRL: A Sequential Fusion Framework for Multi-view Representation Learning. In *Proceedings of the Computer Vision and Pattern Recognition Conference*. 25822–25831.

[60] Shangyou Wang, Zezhong Ding, and Xike Xie. 2025. SamGoG: A Sampling-Based Graph-of-Graphs Framework for Imbalanced Graph Classification. arXiv:2507.13741 [cs.LG] https://arxiv.org/abs/2507.13741

[61] Zonghan Wu, Shirui Pan, Fengwen Chen, Guodong Long, Chengqi Zhang, and Philip S. Yu. 2019. A Comprehensive Survey on Graph Neural Networks. *IEEE Transactions on Neural Networks and Learning Systems* 32 (2019), 4–24. https://api.semanticscholar.org/CorpusID:57375753

[62] Junyuan Xie, Ross Girshick, and Ali Farhadi. 2016. Unsupervised deep embedding for clustering analysis. In *International conference on machine learning*. PMLR, 478–487.

[63] Binbin Xu, Jun Yin, and Nan Zhang. 2024. Graph based Consistency Learning for Contrastive Multi-View Clustering. *Proceedings of the 32nd ACM International Conference on Multimedia* (2024). https://api.semanticscholar.org/CorpusID:273643520

[64] Hongteng Xu, Jiachang Liu, Dixin Luo, and Lawrence Carin. 2022. Representing Graphs via Gromov-Wasserstein Factorization. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 45 (2022), 999–1016. https://api.semanticscholar.org/CorpusID:247083430

[65] Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. 2018. How Powerful are Graph Neural Networks? *ArXiv* abs/1810.00826 (2018). https://api.semanticscholar.org/CorpusID:52895589

[66] Minghao Xu, Hang Wang, Bingbing Ni, Hongyu Guo, and Jian Tang. 2021. Self-supervised Graph-level Representation Learning with Local and Global Structure. *ArXiv* abs/2106.04113 (2021).

[67] Xiaowei Xu, Nurcan Yuruk, Zhidan Feng, and Thomas AJ Schweiger. 2007. Scan: a structural clustering algorithm for networks. In *Proceedings of the 13th ACM SIGKDD international conference on Knowledge discovery and data mining*. 824–833.

[68] Yuhua Xu, Junli Wang, Mingjian Guang, Chungang Yan, and Changjun Jiang. 2024. Graph contrastive learning with min-max mutual information. *Inf. Sci.* 665 (2024), 120378. https://doi.org/10.1016/j.ins.2024.120378

[69] Pinar Yanardag and S. V. N. Vishwanathan. 2015. Deep Graph Kernels. *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (2015). https://api.semanticscholar.org/CorpusID:207227372

[70] Jinbin Yang, Jinyu Cai, Yunhe Zhang, Sujia Huang, and Shiping Wang. 2025. Towards Adaptive Masked Structural Learning for Graph-Level Clustering. *IEEE Transactions on Network Science and Engineering* 12 (2025), 2021–2032. https://doi.org/10.1109/tnse.2025.3543194

[71] Yining Yang, Ke Xu, and Ying Tang. 2024. GAN-based self-supervised message passing graph representation learning. *Expert Syst. Appl.* 251 (2024), 124012. https://doi.org/10.1016/j.eswa.2024.124012

[72] Zhenyu Yang, Ge Zhang, Jia Wu, Jian Yang, Quan Z Sheng, Shan Xue, Chuan Zhou, Charu Aggarwal, Hao Peng, Wenbin Hu, et al. 2024. State of the art and potentialities of graph-level learning. *Comput. Surveys* 57, 2 (2024), 1–40.

[73] Yuning You, Tianlong Chen, Yongduo Sui, Ting Chen, Zhangyang Wang, and Yang Shen. 2020. Graph Contrastive Learning with Augmentations. *ArXiv* abs/2010.13902 (2020). https://api.semanticscholar.org/CorpusID:225076220

[74] Zhiping Zeng, Anthony KH Tung, Jianyong Wang, Jianhua Feng, and Lizhu Zhou. 2009. Comparing stars: On approximating graph edit distance. *Proceedings of the VLDB Endowment* 2, 1 (2009), 25–36.

[75] Mingda Zhang, Chun-Te Chu, Andrey Zhmoginov, Andrew G. Howard, Brendan Jou, Yukun Zhu, Li Zhang, Rebecca Hwa, and Adriana Kovashka. 2021. BasisNet: Two-stage Model Synthesis for Efficient Inference. *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)* (2021), 3075–3084. https://api.semanticscholar.org/CorpusID:234094277

[76] Sheng Zhou, Hongjia Xu, Zhuonan Zheng, Jiawei Chen, Zhao Li, Jiajun Bu, Jia Wu, Xin Wang, Wenwu Zhu, and Martin Ester. 2024. A comprehensive survey on deep clustering: Taxonomy, challenges, and future directions. *Comput. Surveys* 57, 3 (2024), 1–38.

[77] Wei Zhuo and Guang Tan. 2022. Efficient graph similarity computation with alignment regularization. *Advances in Neural Information Processing Systems* 35 (2022), 30181–30193.

A Theoretical Proofs

A.1 Proof of Theorem 4.1

Proof. Our goal is to prove that for any iteration b, if two nodes v and u have the same WL label, i.e., $l_u^{(b)} = l_v^{(b)}$, then they must also have the same GINs node representation, i.e., $\mathbf{h}_u^{(b)} = \mathbf{h}_v^{(b)}$.

- (b=0) We initialize $l_u^{(0)} = l_v^{(0)}$ and $\mathbf{h}_u^{(0)} = \mathbf{h}_v^{(0)}$.
- (b-1) In the b-1 step in 1-WL test, if $l_u^{(b-1)} = l_v^{(b-1)}$ holds, then $\mathbf{h}_u^{(b-1)} = \mathbf{h}_v^{(b-1)}$ also holds.
- (b) At iteration b, if two nodes v and u have the same label ($l_u^{(b)} = l_v^{(b)}$), it implies that their labels at the previous iteration were identical i.e. $l_u^{(b-1)} = l_v^{(b-1)}$, and the multisets of their neighbors' labels from the previous iteration, i.e. $l_x^{(b-1)}, x \in N(u)$ and $l_y^{(b-1)}, y \in N(v)$, were also identical. In that case, we can deduce that $\sum_{x\in N(u)} \mathbf{h}_x^{(b-1)} = \sum_{y\in N(v)} \mathbf{h}_y^{(b-1)}$, according to Eq. 3, we have $\mathbf{h}_u^{(b)} = \mathbf{h}_v^{(b)}$.

When all node representations in two graphs are identical, the Spatial Representation is also identical. In conclusion, if a pair of graphs cannot be identical by 1-WL test, it also cannot be identical by the GINs model. □

A.2 Proof of Lemma 4.2

Proof. $c(e)$ was defined in §3.2.
(Sufficiency): Assume $\mathcal{G}_i \cong \mathcal{G}_j$ with isomorphism ϕ. Let $\Upsilon_\phi = \{u \to \phi(u) \mid u \in \mathcal{V}_i\}$.

$$\mathcal{G}_i \cong \mathcal{G}_j \implies C(\Upsilon_\phi) = \sum_{e\in\Upsilon_\phi} c(e) = 0. \tag{23}$$

$$\therefore 0 \le \text{ged}(\mathcal{G}_i, \mathcal{G}_j) \le C(\Upsilon_\phi) = 0 \implies \text{ged}(\mathcal{G}_i, \mathcal{G}_j) = 0.$$

(Necessity): Assume $\text{ged}(\mathcal{G}_i, \mathcal{G}_j) = 0$ with optimal path Υ^*. Let $\mathcal{E}_{str}$ be the set of insertion/deletion operations.

$$\sum_{e\in\Upsilon^*} c(e) = 0 \wedge c(\mathcal{E}_{str}) > 0 \implies \Upsilon^* \cap \mathcal{E}_{str} = \emptyset \implies |\mathcal{V}_i| = |\mathcal{V}_j|.$$

Thus, $\phi : \mathcal{V}_i \to \mathcal{V}_j \implies \mathcal{G}_i \cong \mathcal{G}_j$. □

A.3 Proof of Lemma 5.1

Proof. A permutation matrix $\mathbf{P}$ is by definition an orthogonal matrix, satisfying $\mathbf{P}^\top = \mathbf{P}^{-1}$. consequently, the transformation of the adjacency matrix under node permutation can be rewritten as a similarity transformation:

$$\tilde{\mathbf{A}} = \mathbf{P}\mathbf{A}\mathbf{P}^\top = \mathbf{P}\mathbf{A}\mathbf{P}^{-1}. \tag{24}$$

In linear algebra, similar matrices share the same characteristic polynomial . We derive the characteristic polynomial of $\tilde{\mathbf{A}}$ as follows:

$$\begin{aligned}\det(\lambda\mathbf{I} - \tilde{\mathbf{A}}) &= \det(\lambda\mathbf{I} - \mathbf{P}\mathbf{A}\mathbf{P}^{-1}) \\ &= \det(\mathbf{P}(\lambda\mathbf{I} - \mathbf{A})\mathbf{P}^{-1}) \\ &= \det(\mathbf{P})\det(\lambda\mathbf{I} - \mathbf{A})\det(\mathbf{P}^{-1}) \\ &= \det(\lambda\mathbf{I} - \mathbf{A}).\end{aligned} \tag{25}$$

Since $\mathbf{A}$ and $\tilde{\mathbf{A}}$ have identical characteristic polynomials, they yield the exact same set of eigenvalues. Thus, $\boldsymbol{\lambda}(\mathbf{A}) = \boldsymbol{\lambda}(\mathbf{P}\mathbf{A}\mathbf{P}^\top)$ for any $\mathbf{P} \in \Pi$. □

A.4 Proof of Theorem 5.2

Proof. Let $\mathbf{A}$ and $\mathbf{B}$ be two real symmetric matrices with spectral decompositions $\mathbf{A} = \mathbf{U}\boldsymbol{\Lambda}_A\mathbf{U}^\top$ and $\mathbf{B} = \mathbf{V}\boldsymbol{\Lambda}_B\mathbf{V}^\top$, where $\mathbf{U}, \mathbf{V}$ are orthogonal matrices, and $\boldsymbol{\Lambda}_A = \text{diag}(\lambda_1(\mathbf{A}), \ldots, \lambda_n(\mathbf{A}))$, $\boldsymbol{\Lambda}_B = \text{diag}(\lambda_1(\mathbf{B}), \ldots, \lambda_n(\mathbf{B}))$ contain eigenvalues sorted in descending order.

First, we expand the squared Frobenius norm of the difference:

$$\begin{aligned}\|\mathbf{A} - \mathbf{B}\|_\text{F}^2 &= \text{Tr}((\mathbf{A} - \mathbf{B})^\top(\mathbf{A} - \mathbf{B})) \\ &= \text{Tr}(\mathbf{A}^2 + \mathbf{B}^2 - 2\mathbf{A}\mathbf{B}) \\ &= \text{Tr}(\mathbf{A}^2) + \text{Tr}(\mathbf{B}^2) - 2\text{Tr}(\mathbf{A}\mathbf{B}) \\ &= \sum_{i=1}^n \lambda_i^2(\mathbf{A}) + \sum_{i=1}^n \lambda_i^2(\mathbf{B}) - 2\text{Tr}(\mathbf{A}\mathbf{B}).\end{aligned} \tag{26}$$

Similarly, expanding the summation term in the inequality gives:

$$\sum_{i=1}^n (\lambda_i(\mathbf{A}) - \lambda_i(\mathbf{B}))^2 = \sum_{i=1}^n \lambda_i^2(\mathbf{A}) + \sum_{i=1}^n \lambda_i^2(\mathbf{B}) - 2\sum_{i=1}^n \lambda_i(\mathbf{A})\lambda_i(\mathbf{B}). \tag{27}$$

Comparing the two expansions, proving the inequality is equivalent to proving:

$$\text{Tr}(\mathbf{A}\mathbf{B}) \le \sum_{i=1}^n \lambda_i(\mathbf{A})\lambda_i(\mathbf{B}). \tag{28}$$

Now, we analyze the term $\text{Tr}(\mathbf{A}\mathbf{B})$:

$$\begin{aligned}\text{Tr}(\mathbf{A}\mathbf{B}) &= \text{Tr}(\mathbf{U}\boldsymbol{\Lambda}_A\mathbf{U}^\top\mathbf{V}\boldsymbol{\Lambda}_B\mathbf{V}^\top) \\ &= \text{Tr}(\boldsymbol{\Lambda}_A(\mathbf{U}^\top\mathbf{V})\boldsymbol{\Lambda}_B(\mathbf{V}^\top\mathbf{U})) \\ &= \text{Tr}(\boldsymbol{\Lambda}_A\mathbf{P}\boldsymbol{\Lambda}_B\mathbf{P}^\top),\end{aligned} \tag{29}$$

where $\mathbf{P} = \mathbf{U}^\top\mathbf{V}$. Since $\mathbf{U}$ and $\mathbf{V}$ are orthogonal, $\mathbf{P}$ is also an orthogonal matrix. Let $\mathbf{S}[i,j] = (\mathbf{P}[i,j])^2$. Since $\mathbf{P}$ is orthogonal, we have $\sum_i \mathbf{S}[i,j] = \sum_i \mathbf{P}[i,j]^2 = 1$ and $\sum_j \mathbf{S}[i,j] = \sum_j \mathbf{P}[i,j]^2 = 1$. Thus, $\mathbf{S}$ is a *doubly stochastic matrix*.

Expanding the trace term element-wise:

$$\text{Tr}(\boldsymbol{\Lambda}_A\mathbf{P}\boldsymbol{\Lambda}_B\mathbf{P}^\top) = \sum_{i=1}^n\sum_{j=1}^n \lambda_i(\mathbf{A})\mathbf{S}_{ij}\lambda_j(\mathbf{B}). \tag{30}$$

According to the Hardy-Littlewood-Pólya rearrangement inequality [57] generalized to doubly stochastic matrices (Birkhoff's theorem), the weighted sum $\sum_{i,j}\mathbf{S}_{ij}\lambda_i(\mathbf{A})\lambda_j(\mathbf{B})$ is maximized when the weights $\mathbf{S}_{ij}$ align the largest $\lambda_i(\mathbf{A})$ with the largest $\lambda_j(\mathbf{B})$. Specifically, since both sets of eigenvalues are sorted in descending order, the maximum is attained when $\mathbf{S}$ is the identity matrix $\mathbf{I}$:

$$\sum_{i=1}^n\sum_{j=1}^n \mathbf{S}_{ij}\lambda_i(\mathbf{A})\lambda_j(\mathbf{B}) \le \sum_{i=1}^n \lambda_i(\mathbf{A})\lambda_i(\mathbf{B}). \tag{31}$$

Substituting this back into Eq. (28), we confirm that $\text{Tr}(\mathbf{A}\mathbf{B}) \le \sum_{i=1}^n \lambda_i(\mathbf{A})\lambda_i(\mathbf{B})$, which implies:

$$-2\text{Tr}(\mathbf{A}\mathbf{B}) \ge -2\sum_{i=1}^n \lambda_i(\mathbf{A})\lambda_i(\mathbf{B}). \tag{32}$$

Finally, adding $\sum\lambda_i^2(\mathbf{A}) + \sum\lambda_i^2(\mathbf{B})$ to both sides completes the proof:

$$\|\mathbf{A} - \mathbf{B}\|_F^2 \ge \sum_{i=1}^n (\lambda_i(\mathbf{A}) - \lambda_i(\mathbf{B}))^2. \tag{33}$$

□

A.5 Proof of row-sparsity in Eq. (11)

Lemma A.1 ($\ell_{2,1}$ induces row-sparsity). *In the optimization problem*

$$\min_{\mathbf{W}\in\mathbb{R}^{d\times k}} f(\mathbf{W}) + \lambda\|\mathbf{W}\|_{2,1}, \qquad \|\mathbf{W}\|_{2,1} = \sum_{r=1}^{d} \|\mathbf{W}[i]\|_2,\ \lambda > 0,$$

where $f(\mathbf{W})$ is convex and differentiable (e.g., a least-squares regression loss), the $\ell_{2,1}$ regularizer drives many entire rows *of the optimal $\mathbf{W}^*$ to be exactly zero (i.e., row-sparsity / group sparsity).*

Proof. Since f is convex differentiable and $\|\cdot\|_{2,1}$ is convex, a matrix $\mathbf{W}^*$ is optimal iff it satisfies the first-order (subgradient) optimality condition

$$\mathbf{0} \in \nabla f(\mathbf{W}^*) + \lambda\,\partial\|\mathbf{W}^*\|_{2,1}. \tag{34}$$

Because $\|\mathbf{W}\|_{2,1}$ is separable across rows, its subdifferential is also row-separable:

$$\partial\|\mathbf{W}\|_{2,1} = \Big\{\mathbf{G}\in\mathbb{R}^{d\times k} : \mathbf{G}[i] \in \partial\|\mathbf{W}[i]\|_2,\ \forall r\Big\},$$

and for each row i,

$$\partial\|\mathbf{W}[i]\|_2 = \begin{cases} \left\{\dfrac{\mathbf{W}[i]}{\|\mathbf{W}[i]\|_2}\right\}, & \mathbf{W}[i]\neq\mathbf{0}, \\ \left\{\mathbf{g}\in\mathbb{R}^{1\times k} : \|\mathbf{g}\|_2 \le 1\right\}, & \mathbf{W}[i]=\mathbf{0}. \end{cases}$$

Plugging the above into (34) and examining the i-th row gives two cases:

(i) Nonzero row. If $\mathbf{W}[i]^* \neq \mathbf{0}$, then $\partial\|\mathbf{W}[i]^*\|_2$ is a singleton and

$$\nabla f(\mathbf{W}^*)[i] + \lambda\frac{\mathbf{W}[i]^*}{\|\mathbf{W}[i]^*\|_2} = \mathbf{0},$$

which implies $\left\|\nabla f(\mathbf{W}^*)[i]\right\|_2 = \lambda$.

(ii) Zero row. If $\mathbf{W}[i]^* = \mathbf{0}$, then there exists some $\mathbf{g}$ with $\|\mathbf{g}\|_2 \le 1$ such that

$$\nabla f(\mathbf{W}^*)[i] + \lambda\mathbf{g} = \mathbf{0} \quad\Longrightarrow\quad \left\|\nabla f(\mathbf{W}^*)[i]\right\|_2 = \lambda\|\mathbf{g}\|_2 \le \lambda.$$

Equivalently, whenever $\left\|\nabla f(\mathbf{W}^*)[i]\right\|_2 < \lambda$, the KKT condition can only be satisfied by $\mathbf{W}[i]^* = \mathbf{0}$.

Therefore, $\ell_{2,1}$ imposes a row-wise thresholding effect: rows whose contribution to decreasing $f(\mathbf{W})$ is insufficient (i.e., with small row-gradient norm) are set exactly to zero, yielding row-sparsity (group-wise feature selection) in $\mathbf{W}^*$. □

A.6 Proof of Theorem 5.3

Proof. Let

$$\begin{aligned} E(\mathbf{X}) &= \sum_{i=1}^{n}\left\|\mathbf{X}[i] - \sum_{j=1}^{n}\mathbf{W}[i,j]\,\mathbf{X}[j]\right\|_2^2 \\ &= \left\|\mathbf{X} - \mathbf{W}\,\mathbf{X}\right\|_F^2 \\ &= \mathrm{Tr}\left((\mathbf{X}-\mathbf{W}\,\mathbf{X})^\top(\mathbf{X}-\mathbf{W}\,\mathbf{X})\right). \end{aligned}$$

Set

$$\mathbf{M} \;=\; (\mathbf{I}-\mathbf{W})^\top(\mathbf{I}-\mathbf{W}) \;\in\; \mathbb{R}^{n\times n},$$

so that

$$E(\mathbf{X}) \;=\; \mathrm{Tr}\left(\mathbf{X}^\top\mathbf{M}\,\mathbf{X}\right).$$

Since $\mathbf{M}$ is symmetric p.s.d. and

$$\mathbf{M}\,\mathbf{1} = (\mathbf{I}-\mathbf{W})^\top(\mathbf{I}-\mathbf{W})\,\mathbf{1} = (\mathbf{I}-\mathbf{W})^\top\,\mathbf{0} = \mathbf{0}.$$

we have $\lambda_0 = 0$ as an eigenvalue of $\mathbf{M}$. Order the eigenpairs $\{(\lambda_k, u_k)\}_{k=0}^{n-1}$ so that

$$0 = \lambda_0 \le \lambda_1 \le \cdots \le \lambda_{n-1}, \quad U = [\,u_0, u_1, \ldots, u_{n-1}\,] \in \mathbb{R}^{n\times n}$$

orthonormal.

Consider the constrained problem

$$\min_{\mathbf{X}\in\mathbb{R}^{n\times d}} \mathrm{Tr}\left(\mathbf{X}^\top\mathbf{M}\,\mathbf{X}\right) \quad \text{s.t.} \quad \mathbf{X}^\top\mathbf{1} = \mathbf{0}, \quad \mathbf{X}^\top\mathbf{X} = \mathbf{I}_d.$$

By the Courant–Fischer theorem, the minimizer $\mathbf{X}^*$ is given (up to an orthogonal transform) by the eigenvectors $\{u_1, \ldots, u_d\}$ corresponding to the d smallest nonzero eigenvalues of $\mathbf{M}$:

$$\mathbf{X}^* \;=\; [\,u_1, \ldots, u_d\,].$$

On the other hand, let $\phi : M \to \mathbb{R}^d$ be an isometric coordinate chart of the underlying d-manifold, and define

$$\Phi = (\phi(x_1), \ldots, \phi(x_n))^\top \;\in\; \mathbb{R}^{n\times d}.$$

By construction of $\mathbf{W}$ via local linear reconstruction on sufficiently small neighborhoods one shows $\mathbf{M}\,\Phi = \mathbf{0}$ and $\Phi^\top\mathbf{1} = \mathbf{0},\ \Phi^\top\Phi = \mathbf{I}_d$ (up to a global scale). Uniqueness of the above minimizer then implies

$$\mathbf{X}^* \;=\; \alpha\,\Phi\,R, \quad \alpha > 0,\ R \in O(d),$$

i.e. $\mathbf{X}^*$ recovers the manifold coordinates $\phi(x_i)$ up to scale and rotation. This completes the proof. □

A.7 The Derivation of Eq (17)

To optimize $\mathbf{Y}$, we start with the overall objective function defined in Eq. (14). By fixing the projection matrices $\mathbf{W}_\mathrm{H}$ and $\mathbf{W}_\mathrm{F}$, the objective function related to $\mathbf{Y}$ can be expanded as:

$$\begin{aligned} \mathcal{J}(\mathbf{Y}) =& \mathrm{Tr}(\mathbf{Y}^\top\mathbf{L}_{\mathrm{CA}}\mathbf{Y}) + \|\mathbf{Y}-\mathbf{H}\mathbf{W}_\mathrm{H}\|_F^2 + \|\mathbf{Y}-\mathbf{F}\mathbf{W}_\mathrm{F}\|_F^2 \\ &+ \mathrm{Tr}(\mathbf{W}_\mathrm{H}^\top\Delta_\mathrm{H}\mathbf{W}_\mathrm{H}) + \mathrm{Tr}(\mathbf{W}_\mathrm{F}^\top\Delta_\mathrm{F}\mathbf{W}_\mathrm{F}), \end{aligned} \tag{35}$$

where the $L_{2,1}$-norm terms are reformulated using the auxiliary diagonal matrices Δ_H and Δ_F (e.g., $\lambda\|\mathbf{W}\|_{2,1} = \mathrm{Tr}(\mathbf{W}^\top\Delta\mathbf{W})$).

We first focus on the terms associated with the view $\mathbf{H}$. Substituting the optimal solution $\mathbf{W}_\mathrm{H} = (\mathbf{H}^\top\mathbf{H}+\Delta_\mathrm{H})^{-1}\mathbf{W}_\mathrm{H}^\top\mathbf{Y}$ derived in Eq. (15) into the regression and regularization terms:

$$\begin{aligned} \mathcal{J}_\mathrm{H}(\mathbf{Y}) &= \|\mathbf{Y}-\mathbf{H}\mathbf{W}_\mathrm{H}\|_F^2 + \mathrm{Tr}(\mathbf{W}_\mathrm{H}^\top\Delta_\mathrm{H}\mathbf{W}_\mathrm{H}) \\ &= \mathrm{Tr}\left((\mathbf{Y}-\mathbf{H}\mathbf{W}_\mathrm{H})^\top(\mathbf{Y}-\mathbf{H}\mathbf{W}_\mathrm{H}) + \mathbf{W}_\mathrm{H}^\top\Delta_\mathrm{H}\mathbf{W}_\mathrm{H}\right) \\ &= \mathrm{Tr}\left(\mathbf{Y}^\top\mathbf{Y} - 2\mathbf{Y}^\top\mathbf{H}\mathbf{W}_\mathrm{H} + \mathbf{W}_\mathrm{H}^\top(\mathbf{H}^\top\mathbf{H}+\Delta_\mathrm{H})\mathbf{W}_\mathrm{H}\right). \end{aligned} \tag{36}$$

Let $\mathbf{A}_\mathrm{H} = \mathbf{H}^\top\mathbf{H}+\Delta_\mathrm{H}$. Then $\mathbf{W}_\mathrm{H} = \mathbf{A}_\mathrm{H}^{-1}\mathbf{H}^\top\mathbf{Y}$. Substituting this back:

$$\begin{aligned} \mathcal{J}_\mathrm{H}(\mathbf{Y}) &= \mathrm{Tr}\left(\mathbf{Y}^\top\mathbf{Y} - 2\mathbf{Y}^\top\mathbf{H}\mathbf{A}_\mathrm{H}^{-1}\mathbf{H}^\top\mathbf{Y} + \mathbf{Y}^\top\mathbf{H}\mathbf{A}_\mathrm{H}^{-1}\mathbf{A}_\mathrm{H}\mathbf{A}_\mathrm{H}^{-1}\mathbf{H}^\top\mathbf{Y}\right) \\ &= \mathrm{Tr}\left(\mathbf{Y}^\top\mathbf{Y} - 2\mathbf{Y}^\top\mathbf{H}\mathbf{A}_\mathrm{H}^{-1}\mathbf{H}^\top\mathbf{Y} + \mathbf{Y}^\top\mathbf{H}\mathbf{A}_\mathrm{H}^{-1}\mathbf{H}^\top\mathbf{Y}\right) \\ &= \mathrm{Tr}\left(\mathbf{Y}^\top\mathbf{Y} - \mathbf{Y}^\top\mathbf{H}(\mathbf{H}^\top H+\Delta_\mathrm{H})^{-1}\mathbf{H}^\top\mathbf{Y}\right) \\ &= \mathrm{Tr}\left(\mathbf{Y}^\top\left(\mathbf{I}-\mathbf{H}(\mathbf{H}^\top\mathbf{H}+\Delta_\mathrm{H})^{-1}\mathbf{H}^\top\right)\mathbf{Y}\right). \end{aligned} \tag{37}$$

Analogously, for the spectral view F, we obtain:

$$\mathcal{J}_\mathrm{F}(\mathbf{Y}) = \mathrm{Tr}\left(\mathbf{Y}^\top\left(\mathbf{I}-\mathbf{F}(\mathbf{F}^\top\mathbf{F}+\Delta_\mathrm{F})^{-1}\mathbf{F}^\top\right)\mathbf{Y}\right). \tag{38}$$

Combining $\mathcal{J}_\mathrm{H}(\mathbf{Y})$, $\mathcal{J}_\mathrm{F}(\mathbf{Y})$, and the manifold preservation term $\mathrm{Tr}(\mathbf{Y}^\top\mathbf{L}_{\mathrm{CA}}\mathbf{Y})$, the total objective becomes:

$$\min_{\mathbf{Y}^\top\mathbf{Y}=I} \mathrm{Tr}\left(\mathbf{Y}^\top\left[\mathbf{L}_{\mathrm{CA}} + 2\mathbf{I} - \mathbf{H}(\mathbf{H}^\top\mathbf{H}+\Delta_\mathrm{H})^{-1}\mathbf{H}^\top - \mathbf{F}(\mathbf{F}^\top\mathbf{F}+\Delta_\mathrm{F})^{-1}\mathbf{F}^\top\right]\mathbf{Y}\right). \tag{39}$$

So we define the auxiliary matrix $\mathbf{M}$ as:

$$\mathbf{M} = \mathbf{L}_{\mathrm{CA}} + 2\mathbf{I} - \left(\mathbf{H}(\mathbf{H}^\top\mathbf{H} + \Delta_{\mathrm{H}})^{-1}\mathbf{H}^\top + \mathbf{F}(\mathbf{F}^\top\mathbf{F} + \Delta_{\mathrm{F}})^{-1}\mathbf{F}^\top\right). \tag{40}$$

Thus, the optimization problem is reduced to Eq. (16), which can be solved by performing eigendecomposition on $\mathbf{M}$.

A.8 Proof of Lemma 5.4

Proof. To derive the optimal solution for the objective function in Eq. (22), we employ the method of Lagrange multipliers. The problem seeks to minimize the entropy-regularized transport cost subject to marginal constraints. Recall that the cost matrix is defined as $\mathbf{C} = -\log \mathbf{Q}$.

Taking the partial derivative with respect to $\mathbf{T}[i,k]$ and setting it to zero yields the first-order optimality condition:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{T}[i,k]} = \mathbf{C}[i,k] + \varepsilon(1 + \log \mathbf{T}[i,k]) = 0. \tag{41}$$

Rearranging terms to solve for $\mathbf{T}[i,k]$:

$$\log \mathbf{T}[i,k] = \frac{-\mathbf{C}[i,k]}{\varepsilon} - 1. \tag{42}$$

Substituting the cost definition $\mathbf{C}[i,k] = -\log \mathbf{Q}[i,k]$, we have:

$$\begin{aligned} \mathbf{T}[i,k] &\propto \exp\left(\frac{-\mathbf{C}[i,k]}{\varepsilon}\right) \\ &= \exp\left(\frac{\log \mathbf{Q}[i,k]}{\varepsilon}\right). \end{aligned} \tag{43}$$

To strictly enforce the marginal constraints, we introduce the scaling vectors $\mathbf{u}$ and $\mathbf{v}$. By defining the kernel matrix $\mathbf{K}$ with entries $\mathbf{K}[i,k] = \exp(\frac{1}{\varepsilon}\log \mathbf{Q}[i,k])$, the optimal solution can be factorized into the matrix form:

$$\mathbf{T}[i,k] = \mathbf{u}_i\mathbf{K}[i,k]\mathbf{v}_k \iff \mathbf{T} = \mathrm{diag}(\mathbf{u})\mathbf{K}\mathrm{diag}(\mathbf{v}). \tag{44}$$

Finally, substituting the optimal $\mathbf{T}$ back into the marginal constraints allows us to solve for $\mathbf{u}$ and $\mathbf{v}$:

$$\begin{cases} \mathbf{T}\mathbf{1}_K = \mathbf{g} \implies \mathrm{diag}(\mathbf{u})(\mathbf{K}\mathbf{v}) = \mathbf{g} \implies \mathbf{u} = \mathbf{g} \oslash (K\mathbf{v}), \\ \mathbf{T}^\top\mathbf{1}_N = \mathbf{c} \implies \mathrm{diag}(\mathbf{v})(K^\top\mathbf{u}) = \mathbf{c} \implies \mathbf{v} = \mathbf{c} \oslash (K^\top\mathbf{u}), \end{cases} \tag{45}$$

where $\oslash$ denotes element-wise division. These equations correspond to the iterative updates in the Sinkhorn-Knopp algorithm provided in Algorithm 1. □

B Additional Experimental Details

B.1 Experimental Settings

B.1.1 Datasets.

- ***BZR***: Graphs represent benzodiazepine receptor ligands. Nodes are atoms, and edges are chemical bonds. The binary classification task determines whether the compound is active or inactive against the receptor.
- ***MUTAG***: Graphs represent mutagenic aromatic and heterochromatic nitro compounds. Nodes denote atoms and edges represent chemical bonds. The binary classification task distinguishes between mutagenic and non-mutagenic effects on bacteria.
- ***PTC-MR***: Graphs depict chemical compounds tested for carcinogenicity. Nodes represent atoms and edges represent bonds. The binary task classifies compounds based on their carcinogenicity in male rats.
- ***AIDS***: Graphs represent chemical compounds from an antiviral screen. Nodes are atoms and edges are bonds. The binary classification task predicts whether a compound is active or inactive against HIV.
- ***ENZYMES***: Graphs represent protein tertiary structures. Nodes are secondary structure elements (SSEs) and edges denote spatial or sequential neighbors. The task classifies proteins into one of six top-level EC enzyme types.
- ***DD***: Graphs represent protein structures where nodes are amino acids and edges indicate spatial proximity. The binary classification task distinguishes between enzymes and non-enzymes.
- ***OGBG-PPA***: Graphs represent protein-protein association networks. Nodes are proteins and edges denote biologically meaningful associations. The multi-class task predicts the taxonomic group from which the network originates.
- ***COLLAB***: Graphs are researcher ego-networks derived from bibliographies. Nodes represent researchers and edges denote collaboration. The multi-class task identifies the field of study: High Energy, Condensed Matter, or Astrophysics.
- ***REDDIT-12K***:Graphs represent discussion threads where nodes are users and edges denote replies. The classification task identifies which of the 11 subreddits a thread belongs to based on interaction structure.

B.1.2 Baselines.

- `SP` [8]: To balance expressiveness and efficiency in graph kernels, the authors propose path-based kernels instantiated as shortest-path kernels, computable in polynomial time and positive definite. On protein graph classification, these kernels outperform walk-based kernels, achieving higher accuracy while preserving strong representational power.
- `RW` [43]: Positive definite graph kernels enable SVM-based analysis of labeled graphs by marginalizing path kernels over random walks. The authors enhance them via Morgan-enriched vertex labels and a no-backtracking random walk to reduce computation and improve similarity, yielding promising gains on chemical compound classification benchmarks.
- `WL` [50]: A graph kernel using the Weisfeiler–Lehman test to generate sequences of relabeled graphs that capture topology and labels. A subtree-like WL kernel runs in linear time in edges and sequence length, and experiments show superior accuracy and runtime, enabling large applications.
- `LT` [28]: This method introduces two graph kernels for unlabeled graphs that capture global properties via the Lovász number and its orthonormal representation, addressing limitations of local subgraph-based kernels. Providing a margin-based theory guiding kernel choice and showing on synthetic and benchmark datasets that these kernels achieve better or competitive accuracy.
- `GLK` [51]: To scale graph kernels to large graphs, the paper compares graphs via counts of graphlets. Since exhaustive enumeration is costly, they introduce theoretically grounded speedups: sampling and a bounded-degree method, and demonstrate efficient comparison of large graphs beyond the reach of existing kernels.
- `InfoGraph` [53]: The paper proposes a method learns graph-level representations by maximizing mutual information between

whole-graph embeddings and multi-scale substructure embeddings (nodes, edges, triangles), improving generalization over kernels and graph2vec.

- GraphCL [73]: This is a graph contrastive learning framework that learns unsupervised graph representations via four principled graph augmentations. Across semi-supervised, unsupervised, transfer, and adversarial settings, GraphCL achieves comparable or superior generalization, transferability, and robustness without heavy tuning, with parameterized augmentations yielding further gains.
- CuCo [16]: They study the role of negative samples in contrastive graph-level representation learning, using four augmentations, GNN encoders, and a scoring and pacing strategy to progressively select negatives.
- MVGRL [23]: Proposing a self-supervised framework that learns node and graph representations by contrasting structural views. The method contrasts encodings from first-order neighbors with a graph diffusion, avoiding multi-view (>2) or multi-scale contrast that offer no gains.
- GZL [41]: The work proposes a self-supervised, unsupervised graph embedding framework orthogonal to random-walk and contrastive paradigms, which comprises graph zoom-out to create a lightweight miniature graph for efficient training, and miniature-scale learning that injects community structure and employs negative-free point-to-point contrast to learn node/graph representations.
- GWF [64]: Reconstructing each graph as a weighted combination of shared factors under GW discrepancy, producing permutation-invariant representations via factor weights. Jointly learn factors and weights by minimizing reconstruction error with unconstrained reparametrization, and compute GW via PPA or BADMM; extensions integrate GNN auto-encoding, node attributes, and semi-supervised regularization.
- GLCC [30]: Performing graph-level clustering across multiple graphs by building an adaptive affinity graph and applying contrastive learning at instance and cluster levels. It employs a graph-Laplacian contrastive loss, neighbor-informed cluster representations, and neighbor-aware pseudo-labels in an alternating optimization.
- DGLC [11]: Learning graph-level representations with a graph isomorphism network by maximizing mutual information between whole-graph embeddings and substructure representations. A clustering module supplies pseudo labels to regularize and enforce discriminative embeddings, enabling end-to-end joint representation learning and graph clustering.
- DCGLC [12]: Clustering graphs via contrastive learning with two heads—Euclidean-based and subspace-based—to capture complementary cluster structures. A contrastive alignment mechanism fuses their cluster information, yielding end-to-end mutual enhancement of representation learning and clustering.
- GPC [14]: Pretraining a graph model on a source dataset via mutual information maximization with self-supervised clustering regularization. It then freezes the model and learns dataset-specific prompt vectors to adapt graph-level representations to diverse target distributions.

Table 7: Hyper-parameters on different datasets.

DS	lr	h	B	epochs	iter	seed	r	d
ENZYMES	$2e^{-3}$	64	512	20	5	42	16	8
AIDS	$1e^{-3}$	64	256	20	5	42	10	5
BZR	$8e^{-4}$	64	512	20	5	42	6	4
MUTAG	$1e^{-4}$	64	512	20	5	42	8	4
PTC-MR	$1e^{-3}$	64	512	20	5	42	5	3
DD	$2e^{-3}$	64	512	20	5	42	12	4
PTC-MM	$2e^{-3}$	64	256	20	5	42	5	3
COLLAB	$1e^{-3}$	128	4096	20	5	42	40	20
REDDIT-12K	$1e^{-3}$	128	4096	20	5	42	70	20
OGBG-PPA	$1e^{-3}$	128	4096	20	5	42	60	20

B.1.3 Hyper-parameters. Table 7 reports the hyper-parameters used for each dataset. For the GINs encoder, the learning rate (*lr*) controls the optimizer step size, h denotes the dimensionality of the Spatial Representation $\mathbf{H}$, and *batch-size* is the number of training samples in each mini-batch. We run training for multiple repetitions (*iter*); all results in both the main text and the appendix are averaged over these *iter* runs. Within each run, *epochs* is the number of training epochs. *seed* fixes the random initialization and data shuffling for reproducibility. r is the dimensionality of $\mathbf{F}$, and d is the dimensionality of the fused feature $\mathbf{Z}$.

B.2 Ablation Study

Table 8 reports the ablation study of S^3GLC by removing one key component at a time. Consistent with the motivation in § 4.2, the balanced clustering loss $\mathcal{L}_{\text{BGC}}$ is crucial for preventing clustering collapse, especially on multi-class datasets. For example, on *COLLAB* and *OGBG-PPA*, removing $\mathcal{L}_{\text{BGC}}$ leads to clear drops in ACC/NMI/ARI, showing that the Sinkhorn-Knopp equipartition constraint helps avoid degenerate assignments. Moreover, removing the GED-guided contrastive loss $\mathcal{L}_{\text{GCL}}$ generally hurts performance, indicating that topology-aware positive sampling improves representation consistency. Finally, replacing the proposed Joint Manifold-Preserving Sparse Regression and manifold preservation module ($\mathcal{L}_{\text{JMSP}} + \mathcal{L}_{\text{JSR}}$) results in the largest degradation on several datasets (e.g., *BZR* and *COLLAB*), validating that our fusion mechanism is necessary to remove redundancy while retaining complementary semantics from the Spatial Representation $\mathbf{H}$ and Spectral Signature $\mathbf{F}$, as discussed in § 5.3.

Table 9 further verifies the effectiveness of Spectral Signature. Across all shown datasets, replacing our spectral features with classical GKs yields consistently worse clustering results. This supports our claim in § 4.1 that Spatial Representation alone may miss global topology, while the proposed Spectral Signature provides an efficient, permutation-invariant proxy of GED that captures global structural differences beyond 1-WL.

B.3 Parameter Analysis

Fig. 8 reports the sensitivity S^3GLC to the Spectral Signature size r and the fused feature dimension d on *AIDS* . In Fig. 8 (a), ACC/NMI/ARI improve as r increases from 6 to about 10, indicating that the top-r eigenvalues already capture sufficient global topology as a GED proxy. When r becomes larger (e.g., $r \geq 14$), performance drops, suggesting that extra eigenvalues contribute limited signal and

Table 8: Ablation study.

Dataset		w/o $\mathcal{L}_{\text{BGC}}$	w/o $\mathcal{L}_{\text{GCL}}$	w/o $\mathcal{L}_{\text{JMSP}}$ + $\mathcal{L}_{\text{JSR}}$	S^{3}GLC
BZR	ACC ↑	$79.1_{\pm 0.4}$	$79.1_{\pm 0.9}$	$77.0_{\pm 1.8}$	$79.7_{\pm 0.0}$
	NMI ↑	$14.3_{\pm 1.3}$	$14.5_{\pm 2.1}$	$11.1_{\pm 3.7}$	$16.1_{\pm 0.0}$
	ARI ↑	$25.1_{\pm 0.8}$	$24.6_{\pm 1.5}$	$20.3_{\pm 4.0}$	$23.5_{\pm 0.2}$
PTC-MR	ACC ↑	$63.8_{\pm 0.3}$	$63.3_{\pm 1.8}$	$63.2_{\pm 2.1}$	$64.0_{\pm 1.5}$
	NMI ↑	$3.9_{\pm 0.3}$	$3.8_{\pm 1.5}$	$3.6_{\pm 1.4}$	$6.2_{\pm 2.8}$
	ARI ↑	$6.1_{\pm 0.3}$	$5.7_{\pm 1.6}$	$5.7_{\pm 1.9}$	$6.0_{\pm 1.6}$
PTC-MM	ACC ↑	$67.5_{\pm 0.6}$	$67.5_{\pm 0.9}$	$66.1_{\pm 0.2}$	$67.9_{\pm 0.6}$
	NMI ↑	$2.4_{\pm 0.7}$	$2.4_{\pm 1.2}$	$1.3_{\pm 0.3}$	$3.3_{\pm 0.9}$
	ARI ↑	$5.5_{\pm 0.3}$	$5.3_{\pm 1.0}$	$4.1_{\pm 0.6}$	$6.2_{\pm 0.5}$
AIDS	ACC ↑	$98.7_{\pm 0.4}$	$98.7_{\pm 0.5}$	$99.3_{\pm 0.4}$	$99.3_{\pm 0.2}$
	NMI ↑	$89.9_{\pm 2.9}$	$90.1_{\pm 3.0}$	$94.0_{\pm 2.8}$	$93.6_{\pm 1.6}$
	ARI ↑	$94.7_{\pm 1.6}$	$94.7_{\pm 1.9}$	$97.3_{\pm 1.4}$	$96.9_{\pm 0.9}$
COLLAB	ACC ↑	$65.5_{\pm 4.5}$	$69.2_{\pm 0.7}$	$55.3_{\pm 1.2}$	$69.3_{\pm 0.5}$
	NMI ↑	$25.7_{\pm 4.5}$	$29.3_{\pm 3.2}$	$18.7_{\pm 1.2}$	$29.5_{\pm 2.9}$
	ARI ↑	$21.7_{\pm 3.0}$	$25.0_{\pm 1.0}$	$10.6_{\pm 2.2}$	$25.5_{\pm 1.1}$
OGBG-PPA	ACC ↑	$21.5_{\pm 0.2}$	$25.4_{\pm 2.6}$	$15.4_{\pm 0.2}$	$25.8_{\pm 2.1}$
	NMI ↑	$27.9_{\pm 1.8}$	$24.2_{\pm 4.2}$	$4.0_{\pm 0.1}$	$24.8_{\pm 2.8}$
	ARI ↑	$8.2_{\pm 1.0}$	$10.8_{\pm 3.8}$	$2.7_{\pm 0.1}$	$10.9_{\pm 1.7}$

Table 9: The effectiveness of Spectral Signature.

Dataset	Structural Features	ACC ↑	NMI ↑	ARI ↑
BZR	SP [8]	$71.0_{\pm 3.6}$	$4.3_{\pm 0.7}$	$8.0_{\pm 2.4}$
	RW [43]	$73.9_{\pm 0.0}$	$0.7_{\pm 0.0}$	$0.1_{\pm 0.2}$
	WL [50]	$72.9_{\pm 2.0}$	$3.9_{\pm 1.5}$	$7.9_{\pm 3.6}$
	GLK [51]	$72.5_{\pm 2.1}$	$3.1_{\pm 0.7}$	$5.0_{\pm 3.8}$
	ours	**$79.7_{\pm 0.0}$**	**$16.1_{\pm 0.0}$**	**$23.5_{\pm 0.2}$**
PTC-MM	SP [8]	$62.4_{\pm 4.6}$	$1.7_{\pm 0.7}$	$3.1_{\pm 2.3}$
	RW [43]	$67.4_{\pm 0.6}$	**$3.5_{\pm 1.7}$**	$3.5_{\pm 1.2}$
	WL [50]	$59.9_{\pm 3.3}$	$0.9_{\pm 0.3}$	$1.3_{\pm 0.5}$
	GLK [51]	$62.3_{\pm 2.0}$	$1.1_{\pm 0.4}$	$3.2_{\pm 1.3}$
	ours	**$67.9_{\pm 0.6}$**	$3.3_{\pm 0.9}$	**$6.2_{\pm 0.5}$**
PTC-MR	SP [8]	$58.9_{\pm 0.6}$	$1.5_{\pm 0.3}$	$2.2_{\pm 0.7}$
	RW [43]	$59.9_{\pm 1.0}$	$2.5_{\pm 1.3}$	$3.4_{\pm 0.9}$
	WL [50]	$58.0_{\pm 1.7}$	$1.7_{\pm 0.3}$	$1.5_{\pm 0.5}$
	GLK [51]	$59.2_{\pm 0.2}$	$1.8_{\pm 0.4}$	$2.1_{\pm 0.7}$
	ours	**$64.0_{\pm 1.5}$**	**$6.2_{\pm 2.8}$**	**$6.0_{\pm 1.6}$**
AIDS	SP [8]	$91.4_{\pm 5.6}$	$46.5_{\pm 6.1}$	$51.4_{\pm 9.7}$
	RW [43]	$78.4_{\pm 0.5}$	$13.3_{\pm 1.1}$	$17.9_{\pm 2.4}$
	WL [50]	$86.4_{\pm 6.7}$	$37.3_{\pm 4.2}$	$41.8_{\pm 6.0}$
	GLK [51]	$91.9_{\pm 4.9}$	$45.9_{\pm 7.4}$	$55.1_{\pm 15.7}$
	ours	**$99.3_{\pm 0.2}$**	**$93.6_{\pm 1.6}$**	**$96.9_{\pm 0.9}$**

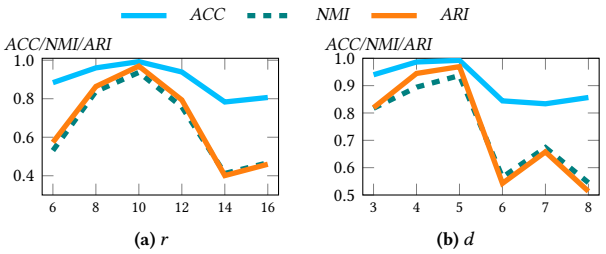

Figure 8: Clustering quality on *AIDS*

may introduce noise. In Fig. 8 (b), the best results are achieved around $d \approx 5$, while larger d degrades metrics, consistent with §5.3: the Joint Manifold-Preserving Sparse Regression favors a compact fused space that removes redundancy between $\mathbf{H}$ and $\mathbf{F}$.

B.4 Cluster Visualizations

Fig. 9 reports t-SNE visualizations of the learned graph embeddings on four datasets. Compared with DGLC and DCGLC, S^{3}GLC

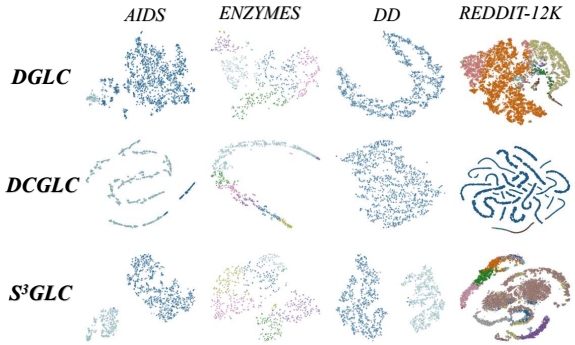

Figure 9: t-SNE visualization of graph embeddings learned by DGLC, DCGLC, and S^{3}GLC

yields clearer inter-cluster margins, notably in *AIDS* and *DD*. On *ENZYMES* and *REDDIT-12K*, S^{3}GLC maintains balanced, distinct cluster regions. Conversely, DGLC and DCGLC exhibit severe clustering collapse, with samples concentrated in a single dominant cluster, failing to preserve the intrinsic multi-cluster structure.