Attraction–Ranking Mismatch in Out-of-Distribution Detection: Theory and Evidence from Hamiltonian Feature Dynamics

Zilin Wang[a], Lianghai Xiao[a,*]

[a]School of Information Science and Technology, Jinan University, Guangzhou, 510632, China

arXiv:2611.06021v1 [cs.LG] 4 Nov 2026

Abstract

Attracting features toward learned in-distribution (ID) regions is an intuitive strategy for out-of-distribution (OOD) detection, but does such attraction improve detection? We investigate this question using class-conditioned Hamiltonian dynamics as a controlled testbed and identify an *attraction–ranking mismatch*: improving individual class-affinity scores need not improve the relative ordering of ID and OOD samples. Our analysis separates misalignment between the driving field and the detection score from ranking failure that persists even when they are aligned. An exact AUROC decomposition and a local boundary analysis identify relative score changes at the ID–OOD ranking boundary as the source of first-order AUROC change under regularity conditions. Experiments demonstrate that a score-aligned rule can improve performance on a held-out source validation split yet consistently degrade it after transfer, even without a single winning-class change. A separate perturbation study finds that near-boundary relative score changes track the direction of small AUROC changes in the tested settings. Finally, centroid approximation and shared scales connect the zero-step score to a static centroid family that includes the existing CTM detector. Together, these results explain why class attraction is insufficient to justify feature refinement and motivate evaluating its effect on the ranking boundary.

Keywords: Out-of-Distribution Detection, Attraction–Ranking Mismatch, Feature Dynamics, Hamiltonian Framework, Pairwise Ranking

1. Introduction

Post-hoc out-of-distribution (OOD) detection asks whether a test sample belongs to the distribution represented by a fixed classifier (Hendrycks and Gimpel, 2017; Liang et al., 2017). Existing detectors extract evidence from confidence, logits, or feature geometry (Liu et al., 2020; Wang et al., 2022; Sun et al., 2022). A recurring intuition is that an in-distribution (ID) feature should lie near, or move toward, a learned class region. The unresolved question is whether this pointwise geometric attraction improves the *relative ordering* of ID and OOD samples.

We turn that intuition into a minimal dynamical problem. Normalized training features define class-conditioned wells on the unit sphere; a query starts from rest; and the strongest class affinity visited along its trajectory is used as ID evidence. Hamiltonian dynamics provide a transparent testbed because the force, trajectory, and observed score are all explicit. We use this model to examine whether attraction implies better detection.

The implication fails in two logically distinct ways. First, the field that moves a feature can be misaligned with the instantaneous class evidence used by the detector: summing class forces may reduce the affinity of the currently winning class. Second, even a score-aligned flow can raise OOD scores more than ID scores near the baseline ranking boundary. We call these failures *attraction–ranking mismatch*. The distinction matters because AUROC is a functional of ID–OOD pairs, not an average of samplewise score changes.

The central ranking result identifies the relative score change at the ID–OOD ranking boundary that determines the first-order AUROC change under regularity conditions. An exact decomposition into corrected and spoiled pairs explains finite changes. A first-step analysis links the driving field to the detection score, while an ordering-preservation result states the limits of universal improvement by deterministic refinement.

Our experiments show that local score alignment does not ensure better OOD ranking. A rule that improves AUROC on a held-out source split reduces it on the target benchmarks when the refinement rule and step count are kept fixed. This failure occurs without observed winning-class changes within the fixed candidate sets, showing that class switching is not required to explain the loss in ranking quality. In a separate perturbation study, near-boundary relative score changes track the direction of small AUROC changes more consistently than their global average. The experiments thus illustrate ranking failure despite local alignment; they do not establish the destructive class interference allowed by the theory.

The $T = 0$ reference separates the static geometric information from the effect of motion. After removing candidate truncation, we bound the error of replacing each class field by a radial centroid score. Under shared scales, this approximate score is exactly rank-equivalent to centroid cosine similarity. CTM (Nguyen et al., 2023) is an existing full-bank, feature-norm-weighted member of the resulting family. This relation

*Corresponding author.
Email address: xiaolh@jnu.edu.cn (Lianghai Xiao)

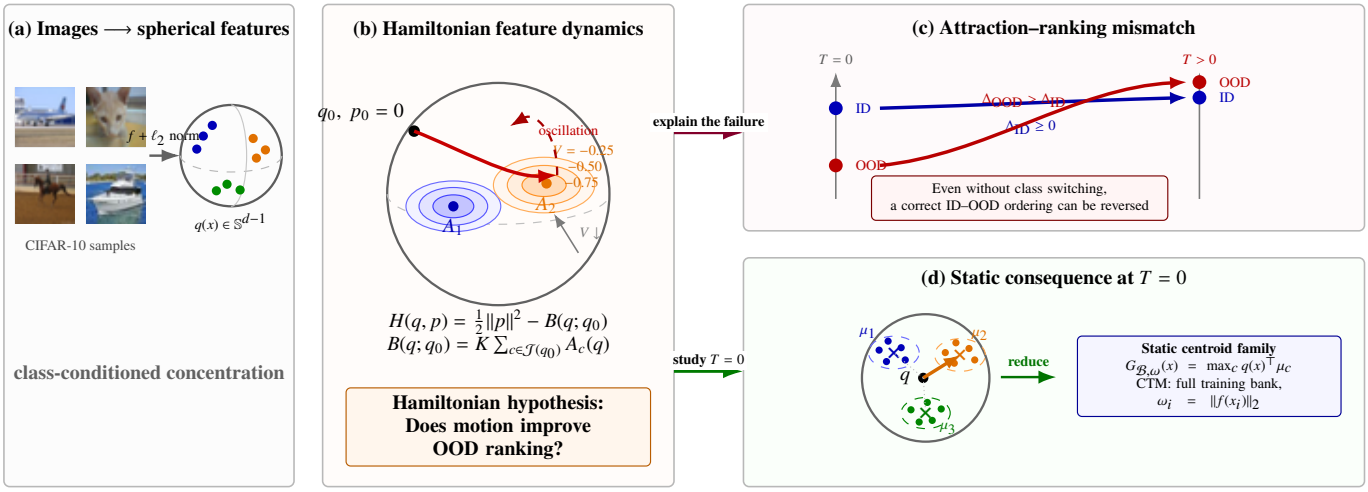

Figure 1: Research map. Class-conditioned concentration on the feature sphere defines potential wells and a testable Hamiltonian hypothesis. Positive-time attraction need not preserve ID–OOD ordering; the locked score-aligned audit shows that the mismatch can persist without a winning-class switch. The exact $T = 0$ reference exposes a static geometry that reduces, after explicit candidate expansion, centroid approximation, and scale sharing, to a centroid family containing CTM as a full-bank member.

provides a static baseline for studying the contribution of motion.

The contribution of this study is threefold:

- first, we analyze the attraction–ranking mismatch by separating field–score alignment from ID–OOD ranking and identifying the boundary condition that governs first-order AUROC change under regularity conditions;
- second, we show through controlled experiments that local score alignment does not guarantee better ranking, even without winning-class changes, and examine the boundary-based explanation in a separate perturbation study;
- third, we connect zero-step anchor scores to a static centroid family through approximation bounds and shared-scale ranking equivalence. This analysis also shows that CTM is an existing member of this family.

The rest of the paper is organized as follows. Section 2 positions the problem. Section 3 states the fundamental model, Section 4 develops the theory, and Section 5 derives its static consequence. Section 6 tests the theoretical chain, followed by discussion and limitations.

2. Related Work

2.1. Post-hoc OOD scores and feature shaping

Post-hoc methods leave the classifier fixed. Output-based scores include maximum softmax probability (MSP), ODIN, Energy, MaxLogit, generalized entropy, and LogitGap (Hendrycks and Gimpel, 2017; Liang et al., 2017; Liu et al., 2020; Hendrycks et al., 2022; Liu et al., 2023; Liang et al., 2025). Feature shaping instead modifies the representation used for scoring: ReAct clips activations, DICE retains high-contribution connections, ASH prunes and rescales activations, Scale rescales features, and RankFeat removes a dominant rank-one component (Sun et al., 2021; Sun and Li, 2022; Djurisic et al., 2023; Xu et al., 2024; Song et al., 2022). Gradient Short-Circuit studies post-hoc feature intervention, whereas FEVER-OOD uses dimensionality reduction and singular-value regularization to mitigate free-energy vulnerabilities (Gu et al., 2025; Isaac-Medina et al., 2025). These distinct training-time and inference-time approaches motivate our central question: when can a representation change be expected to improve ranking?

2.2. Geometric and prototype-based detection

Geometric detectors use class statistics, neighborhoods, or residual subspaces. Representative examples include Mahalanobis distance, RMDS, KNN, ViM, and SHE (Lee et al., 2018; Ren et al., 2021; Sun et al., 2022; Wang et al., 2022; Zhang et al., 2023a). NNGuide and SIRC combine geometric evidence with classifier confidence (Park et al., 2023; Xia and Bouganis, 2024). Recent prototype methods refine the class representative or enrich its comparison rule: SPROD suppresses spurious features, POT uses prototype transport with virtual outliers, GROOD introduces an artificial OOD prototype, and uncertainty-aware refinement couples prototype geometry with energy (Zohrabi et al., 2025; Ke et al., 2025; ElAraby et al., 2025; Wei and Liu, 2026). CADRef retains class-aware relative information beyond a single centroid (Ling et al., 2025).

The CTM method is the closest static construction to our zero-step family (Nguyen et al., 2023). It averages unnormalized training features class-wise, normalizes the resulting means, and uses maximum cosine similarity. Section 5 bounds the centroid approximation error and shows that the resulting radial centroid score is rank-equivalent to cosine similarity under shared scales. CTM is recovered as a full-bank special case.

2.3. From geometric attraction to ranking

Manifold and class-concentration arguments motivate localized feature geometry (Cayton, 2008; Fefferman et al., 2016), while spectral and information-geometric analyses relate learned representations to gradient or predictive geometry (Seleznova et al., 2026; Nguyen et al., 2026). These works explain why geometric structure can be informative. They do not imply that motion generated by this structure improves a distributional ranking.

Hamiltonian mechanics supplies a clean counterexample framework: a scalar potential generates coupled position–momentum motion (Arnold et al., 1989), whereas an OOD score evaluates a separate observable. We analyze why attraction need not improve ranking by separating (i) alignment between the dynamical driver and the observed score from (ii) alignment between samplewise score changes and ID–OOD pair ordering. The resulting pairwise decomposition, boundary derivative, and limitation on universal improvement apply to score refinement beyond this Hamiltonian implementation when their stated assumptions hold.

3. A Hamiltonian Model of Feature Refinement

To make the distinction between attraction and ranking explicit, we construct a model in which both the feature dynamics and the detection score are defined by class affinities. Normalized ID training features serve as anchors, and their affinities define a potential that drives query features toward the learned class regions. We then specify the Hamiltonian dynamics and the score evaluated along each trajectory, providing a concrete setting for analyzing how feature motion affects ID–OOD ranking. Further geometric motivation and details of the numerical update schemes are provided in Appendix A.1.

3.1. A class-induced landscape

Let $f : \mathcal{X} \to \mathbb{R}^d$ be the penultimate representation of a fixed classifier, with $f(x) \neq 0$ on the populations under consideration, and retain only feature direction,

$$q(x) = \frac{f(x)}{\|f(x)\|_2} \in \mathbb{S}^{d-1}. \tag{1}$$

Let C be the number of classes and $K \geq 1$ the number of anchors per class. For $c \in \{1, \ldots, C\}$ and $j \in \{1, \ldots, K\}$, let $a_{cj} \in \mathbb{S}^{d-1}$ be normalized training features used as anchors, $w_{cj} \geq 0$ their weights, and $\sigma_{cj} > 0$ bandwidths. A decreasing radial kernel $\phi : [0, \infty) \to \mathbb{R}$ defines the class affinity

$$A_c(q) = \frac{1}{K} \sum_{j=1}^{K} w_{cj} \phi\left(\frac{\|q - a_{cj}\|_2}{\sigma_{cj}}\right). \tag{2}$$

The full-class reference potential is $V(q) = -\sum_{c=1}^{C} A_c(q)$. In computation, a nonempty candidate set $\mathcal{J}(q_0) \subseteq \{1, \ldots, C\}$ is selected at the initial feature q_0 and held fixed. The scalar driving the motion is

$$B(q; q_0) = K \sum_{c \in \mathcal{J}(q_0)} A_c(q). \tag{3}$$

The dynamics then use the conditional potential $-B(q; q_0)$.

3.2. The Hamiltonian question

Let $P_q = I - qq^\top$. On the tangent bundle $T\mathbb{S}^{d-1} = \{(q, p) : \|q\|_2 = 1,\ q^\top p = 0\}$, augment q with momentum p and start the query from rest under

$$\begin{aligned} H(q, p) &= \tfrac{1}{2}\|p\|_2^2 - B(q; q_0), \\ \dot{q} &= p, \\ \dot{p} &= P_q \nabla_q B(q; q_0) - \|p\|_2^2 q, \qquad (q_0, p_0) = (q(x), 0). \end{aligned} \tag{4}$$

The normal term preserves tangency. Let $\Delta t > 0$ be the numerical step size, $T \in \{0, 1, \ldots\}$ the number of steps, and q_n the feature after step n. Define the instantaneous score and the maximum score along the trajectory by

$$m_{\mathcal{J}}(q; q_0) = \max_{c \in \mathcal{J}(q_0)} A_c(q), \qquad S_T(q_0) = \max_{0 \leq n \leq T} m_{\mathcal{J}}(q_n; q_0). \tag{5}$$

Thus $S_T \geq S_0$ by construction even if instantaneous evidence decreases. The implementation stores KS_T because it sums rather than averages the K anchors; this common positive factor does not change rankings. Figure 2 shows the complete construction: spherical class wells generate motion, while the detector observes a maximum class affinity along the resulting trajectory.

Hypothesis 3.1 (Hamiltonian attraction hypothesis)**.** *Within a specified field, step-size, and trajectory-length grid, nonzero attraction toward the learned class regions improves the ID–OOD ranking induced by the zero-step score S_0, and this advantage persists across benchmark scale.*

By construction, the trajectory score satisfies $S_T(q_0) \geq S_0(q_0)$ for every query. This samplewise guarantee does not establish the hypothesis, because OOD scores may increase more than ID scores. Moreover, the aggregate field driving the motion differs from the maximum class affinity used for scoring. These distinctions motivate the analysis in the next section.

4. Attraction–Ranking Mismatch Theory

We examine two links between feature motion and detection performance. First, does motion driven by the aggregate field increase the instantaneous detection score? Second, do the resulting score changes improve ID–OOD ranking? We begin with a step from rest to characterize the local effect of the driving field, then turn to the relative ordering of ID and OOD scores. Proofs and the connection between a single trajectory step and the continuous score family used in the boundary analysis are provided in Appendix A.2.

4.1. Field–score mismatch

Let $P_q = I - qq^\top$ be the tangent projector of the unit sphere. Both the constrained geodesic kick–drift–kick scheme and the original position-normalized update have the same first-step expansion from rest; Appendix A.1 treats both schemes. Gradients are taken with respect to the feature argument; the initial candidate set is held fixed. Smoothness near a point on the sphere refers to a smooth extension to an ambient neighborhood.

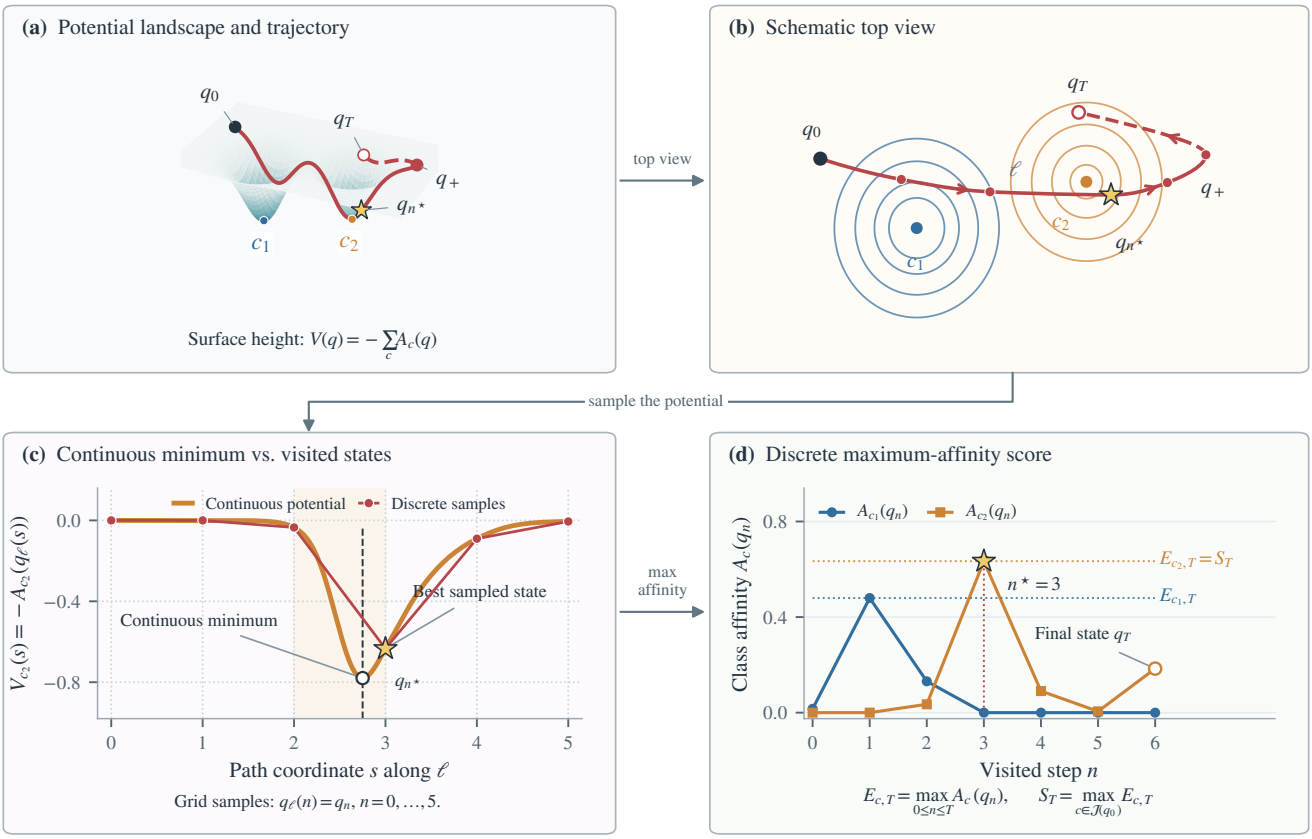

Figure 2: Schematic construction of the trajectory score. (a) The surface height is $V(q) = -\sum_c A_c(q)$, with wells labeled c_1, c_2. (b) The same field and path in top view. Both views identify the initial state q_0, best sampled state $q_{n^\star}$, illustrative turning state q_+, and final state q_T; solid and dashed red curves show outward and returning motion. (c) Class potential along the outward branch ℓ, parameterized by $q_\ell(n) = q_n$ for $n = 0, \ldots, 5$. Grid lines locate samples, and the red polygon joins their potential values. The continuous minimum (open circle) lies between samples and is lower than the best sampled value (star). (d) Each class contributes its maximum visited affinity $E_{c,T}$; their largest value defines S_T, which need not equal the final-state affinity. Stars mark the same state in all panels. All panels share affinities and sampled states; the planar motion is schematic, not a Hamiltonian simulation.

Theorem 4.1 (First-step field–score mismatch)**.** *Fix $\mathcal{J}(q_0)$ and suppose B and A_c are twice continuously differentiable near q_0. Write $b_0 = P_{q_0}\nabla B(q_0; q_0)$ and $g_c = P_{q_0}\nabla A_c(q_0)$. Starting from $p_0 = 0$,*

$$q_1 = q_0 + \frac{\Delta t^2}{2} b_0 + O(\Delta t^4). \tag{6}$$

Consequently,

$$B(q_1; q_0) - B(q_0; q_0) = \frac{\Delta t^2}{2}\|b_0\|_2^2 + O(\Delta t^4), \tag{7}$$

$$A_c(q_1) - A_c(q_0) = \frac{\Delta t^2}{2}\langle g_c, b_0\rangle + O(\Delta t^4). \tag{8}$$

Hence the aggregate driver's leading change is nonnegative, and is positive when $b_0 \neq 0$, whereas an individual class affinity can either increase or decrease.

The theorem makes the first mismatch explicit: intuitive attraction follows the sum B, but the detector observes the largest class affinity. Other classes can therefore pull against the class that currently provides the score.

Definition 4.2 (Class-interference index)**.** *Let $c^\star$ be the unique maximizer of $A_c(q_0)$ over $\mathcal{J}(q_0)$ and set $g_c = P_{q_0}\nabla A_c(q_0)$. For $g_{c^\star} \neq 0$, define*

$$I(q_0) = \frac{\sum_{c \in \mathcal{J}(q_0)\setminus\{c^\star\}} \langle g_{c^\star}, g_c\rangle}{\|g_{c^\star}\|_2^2}. \tag{9}$$

Corollary 4.3 (Destructive class interference)**.** *Under Theorem 4.1 and Definition 4.2,*

$$A_{c^\star}(q_1) - A_{c^\star}(q_0) = \frac{K\Delta t^2}{2}\|g_{c^\star}\|_2^2[1 + I(q_0)] + O(\Delta t^4). \tag{10}$$

If $I(q_0) < -1$, then for every sufficiently small nonzero Δt, the aggregate field increases while the instantaneous maximum class affinity decreases.

For these sufficiently small steps, $S_1(q_0) = S_0(q_0)$: the maximum along the trajectory retains the initial score. The instantaneous score decreases, but the trajectory score does not.

Definition 4.4 (Driver–score alignment)**.** *For a differentiable driver F and instantaneous score s on the sphere, the driver is* locally score-aligned *at q when*

$$\mathcal{A}_{F,s}(q) = \langle P_q\nabla s(q), P_q\nabla F(q)\rangle \geq 0, \tag{11}$$

and strictly aligned when the inequality is strict.

For twice continuously differentiable F and s, a step from rest changes s by $\frac{1}{2}\Delta t^2 \mathcal{A}_{F,s}(q) + O(\Delta t^4)$. Choosing $F = s$ therefore removes the local field–score conflict away from ties. This is a statement about a step from rest. At a general point on the continuous trajectory,

$$\frac{d}{dt} s(q(t)) = \langle P_{q(t)} \nabla s(q(t)), p(t) \rangle,$$

so alignment of the driver alone does not imply a monotone instantaneous score. The trajectory score S_T is nondecreasing in T because it retains the maximum. Neither property determines the relative changes of ID and OOD scores. Figure 3 contrasts these two levels.

4.2. From fixed winners to pairwise ranking

For score-aligned dynamics, let $c_0 = \arg\max_{c \in \mathcal{J}(q_0)} A_c(q_0)$, using deterministic tie-breaking. The initial-winner driver is $F_{\text{init}} = KA_{c_0}$, while the current-winner driver is $F_{\text{cur}} = Km_{\mathcal{J}}(\cdot; q_0)$ and reselects its active class at each force evaluation. The initial-winner control retains only $\max_{n \le T} A_{c_0}(q_n)$; the current-winner control uses S_T.

Proposition 4.5 (Fixed-winner equivalence)**.** *Run the two systems from the same initial state with the same deterministic kick–drift–kick scheme and candidate set. If c_0 is the unique maximizer at every retained state of the initial-winner trajectory, then their states, active classes, driver values, and retained scores are identical at every step. For an integrator with additional internal force evaluations, the same condition is required at those evaluation points.*

When the winner stays fixed, these controls cannot differ through class selection. Their common score can still change ID–OOD ordering. Define the score increase $\Delta_T(q) = S_T(q) - S_0(q) \ge 0$.

Let $X_{\text{in}} \sim P_{\text{in}}$ and $X_{\text{out}} \sim P_{\text{out}}$ be independent input samples, and set $Q_{\text{in}} = q(X_{\text{in}})$ and $Q_{\text{out}} = q(X_{\text{out}})$. For a measurable input score $s : \mathcal{X} \to \mathbb{R}$, larger values indicate stronger ID evidence. We use

$$\begin{aligned} \text{AUROC}(s) = \Pr\{s(X_{\text{in}}) > s(X_{\text{out}})\} \\ + \tfrac{1}{2} \Pr\{s(X_{\text{in}}) = s(X_{\text{out}})\}. \end{aligned} \tag{12}$$

For feature scores, write $\text{AUROC}_T = \text{AUROC}(S_T \circ q)$.

Definition 4.6 (Corrected and spoiled pairs)**.** *For the independent features defined above, let*

$$M_0 = S_0(Q_{\text{in}}) - S_0(Q_{\text{out}}), \qquad D_T = \Delta_T(Q_{\text{in}}) - \Delta_T(Q_{\text{out}}).$$

The sets

$$\begin{aligned} C_T &= \{M_0 \le 0,\ M_0 + D_T > 0\}, \\ \mathcal{S}_T &= \{M_0 > 0,\ M_0 + D_T \le 0\}. \end{aligned} \tag{13}$$

contain baseline errors corrected by the trajectory and baseline-correct pairs spoiled by it, respectively.

Theorem 4.7 (Exact AUROC decomposition)**.** *Assume score ties have probability zero for S_0 and S_T, and write $\Delta\,\text{AUROC} = \text{AUROC}_T - \text{AUROC}_0$. Then*

$$\begin{aligned} \text{AUROC}_T - \text{AUROC}_0 = \Pr(C_T) - \Pr(\mathcal{S}_T), \\ |\Delta\,\text{AUROC}| \le \Pr(|M_0| \le |D_T|). \end{aligned} \tag{14}$$

The identity shows why nonnegative score increases are insufficient: ranking improves only when corrected-pair mass exceeds spoiled-pair mass.

This finite-change identity applies to the trajectory score for any fixed T. To identify which pairs control a small change, we next consider a score family indexed by a continuous parameter η. Its expansion and distributional assumptions must be checked separately; integer step counts and maxima along a trajectory do not automatically satisfy them.

Theorem 4.8 (Boundary-conditioned AUROC derivative)**.** *For the independent inputs $X_{\text{in}}, X_{\text{out}}$ defined above, let $s_\eta : \mathcal{X} \to \mathbb{R}$ be a measurable score family with*

$$\begin{aligned} s_\eta(x) &= s_0(x) + \eta r(x) + R_\eta(x), \\ M_0 &= s_0(X_{\text{in}}) - s_0(X_{\text{out}}), \\ U &= r(X_{\text{in}}) - r(X_{\text{out}}). \end{aligned}$$

Suppose:

1. *$\varepsilon_\eta = \|R_\eta\|_{L^\infty(P_{\text{in}}+P_{\text{out}})} = o(|\eta|)$;*
2. *$\mathbb{E}|U| < \infty$, and (M_0, U) has a joint density $p_{M_0,U}(m,u)$ continuous in m at zero for almost every u;*
3. *for some $\delta > 0$ and measurable $g \ge 0$, $p_{M_0,U}(m,u) \le g(u)$ whenever $|m| \le \delta$, with $\int (1+|u|) g(u)\, du < \infty$, and $p_{M_0}(0) = \int p_{M_0,U}(0,u)\, du > 0$.*

Then

$$\frac{d}{d\eta} \text{AUROC}(s_\eta)\Big|_{\eta=0} = p_{M_0}(0)\, \mathbb{E}[U \mid M_0 = 0]. \tag{15}$$

Here the conditional mean is defined by the density:

$$\mathbb{E}[U \mid M_0 = 0] = \frac{\int u\, p_{M_0,U}(0,u)\, du}{p_{M_0}(0)}.$$

If the score family is defined only for $\eta \ge 0$, the derivative is right-sided. The ranking boundary $M_0 = 0$ consists of ID–OOD pairs with equal baseline scores; it is not a classifier's class decision boundary. Relative score change at this boundary controls the first-order AUROC change, whereas the global mean $\mathbb{E}[U]$ need not have the same sign.

Theorem 4.9 (No universal AUROC gain from deterministic refinement)**.** *Let $s, s_{\text{ref}} : \mathcal{X} \to \mathbb{R}$ be measurable baseline and deterministic refined scores. If $s(x) > s(y)$ but $s_{\text{ref}}(x) \le s_{\text{ref}}(y)$ for some x, y, then there exist ID and OOD distributions for which $\text{AUROC}(s_{\text{ref}}) < \text{AUROC}(s)$. Universal non-decrease over all ID/OOD distributions without baseline ties therefore requires strict preservation of every ordering induced by s; under that condition, the two AUROCs are equal for all such distributions.*

Proof. Choose the point-mass distributions $P_{\text{in}} = \delta_x$ and $P_{\text{out}} = \delta_y$. Equation (12) gives $\text{AUROC}(s) = 1$, while

$$\text{AUROC}(s_{\text{ref}}) = \begin{cases} 0, & s_{\text{ref}}(x) < s_{\text{ref}}(y), \\ 1/2, & s_{\text{ref}}(x) = s_{\text{ref}}(y). \end{cases}$$

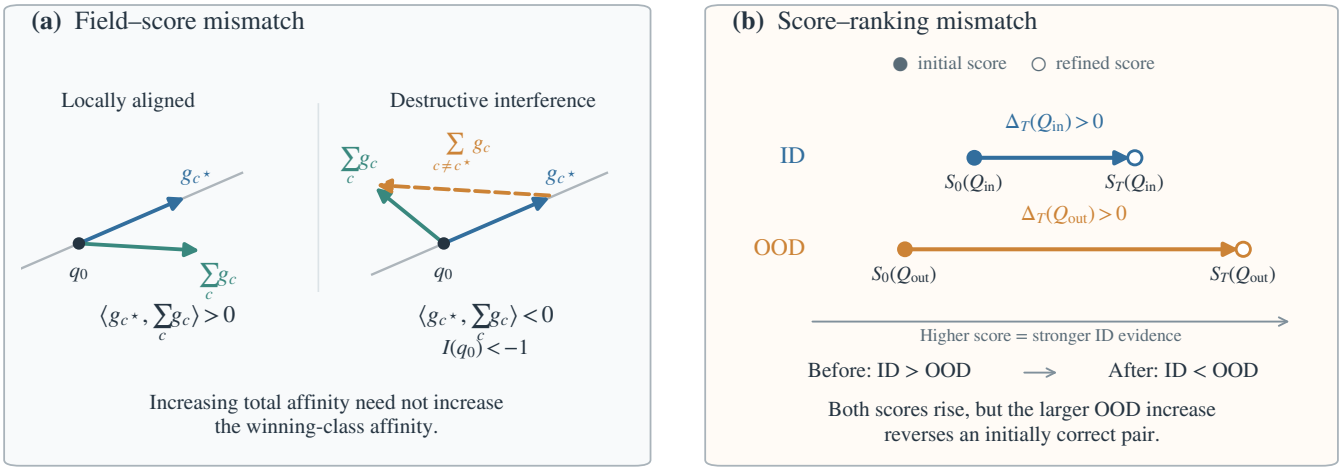

Figure 3: Two levels of the attraction–ranking mismatch. (a) A projected aggregate force can align with the winning-class score gradient (left), whereas nonwinning classes can reverse that alignment and produce the destructive condition $I(q_0) < -1$ (right). (b) Even when each sample's retained trajectory score is nondecreasing, unequal ID and OOD increases can reverse an initially correct pair. All class sums are over $\mathcal{J}(q_0)$. Panel (a) is a theoretical schematic, not an observed audit sample; the geometry and values are illustrative.

Thus universal non-decrease requires $s(x) > s(y) \Rightarrow s_{\text{ref}}(x) > s_{\text{ref}}(y)$ for every pair. Applying this condition also with x, y exchanged preserves the sign of every nonzero baseline difference. If baseline ties have probability zero, then

$$\mathbf{1}_{\{s(X_{\text{in}}) > s(X_{\text{out}})\}} = \mathbf{1}_{\{s_{\text{ref}}(X_{\text{in}}) > s_{\text{ref}}(X_{\text{out}})\}} \quad \text{almost surely,}$$

and refined ties also have probability zero. Taking expectations proves equality of the two AUROCs. □

The theorem limits guarantees over all ID/OOD distributions; it does not rule out improvement on a particular task. Preserving every strict baseline ordering leaves AUROC unchanged when baseline ties have probability zero, whereas changing such an ordering can reduce AUROC for some distributions.

5. Static Scores at Zero Steps

The preceding analysis studies how feature motion changes the ranking relative to the initial score S_0. We now examine the structure of this baseline: what geometric information does the detector use before any motion takes place? Setting $T = 0$ in Eq. (5) gives $S_0(q) = \max_{c \in \mathcal{J}(q)} A_c(q)$.

To study its relation to centroid scoring, we consider the full-class score $S_0^{\text{all}}(q) = \max_{1 \le c \le C} A_c(q)$. Removing candidate truncation is an additional modeling choice and need not preserve the original score. We first bound the error of replacing each class field by a centroid score, then examine when the resulting scores induce the same ranking as cosine similarity. Proofs are provided in Appendix A.3.

Write $W_c = K^{-1} \sum_j w_{cj}$, $\pi_{cj} = w_{cj} / \sum_\ell w_{c\ell}$, $\bar{a}_c = \sum_j \pi_{cj} a_{cj}$, and $\mu_c = \bar{a}_c / \|\bar{a}_c\|_2$. Assume $\sum_j w_{cj} > 0$ and $\bar{a}_c \neq 0$ for every class considered.

Proposition 5.1 (Centroid approximation). *Suppose class c uses a common bandwidth σ_c and ϕ is L_ϕ-Lipschitz. Define $\widetilde{A}_c(q) = W_c \phi(\|q - \mu_c\|_2 / \sigma_c)$. Then*

$$|A_c(q) - \widetilde{A}_c(q)| \le \frac{W_c L_\phi}{\sigma_c} \left(\sum_j \pi_{cj} \|a_{cj} - \bar{a}_c\|_2 + |1 - \|\bar{a}_c\|_2| \right). \quad (16)$$

Consequently, the maximum-field error is bounded by the largest class-wise right-hand side.

Let ϵ be the maximum, over classes, of the right-hand side of Eq. (16), and write $\widetilde{S}_0(q) = \max_c \widetilde{A}_c(q)$. Then $\sup_q |S_0^{\text{all}}(q) - \widetilde{S}_0(q)| \le \epsilon$. In particular,

$$S_0^{\text{all}}(q) - S_0^{\text{all}}(q') > 2\epsilon \implies \widetilde{S}_0(q) > \widetilde{S}_0(q'). \quad (17)$$

Smaller margins need not be preserved. The following exact equivalence concerns the approximated centroid score, not the original anchor score.

Proposition 5.2 (Common-scale ranking equivalence). *If $W_c = W > 0$ and $\sigma_c = \sigma > 0$ are shared across classes and ϕ is strictly decreasing, then for unit q and μ_c,*

$$\max_c \widetilde{A}_c(q) = W \phi\left(\frac{\sqrt{2 - 2G(q)}}{\sigma} \right), \qquad G(q) = \max_c q^\top \mu_c. \quad (18)$$

The two scores select the same classes and rank all samples identically.

Definition 5.3 (Static centroid family). *Let $\mathcal{B} = (\mathcal{B}_c)_{c=1}^C$, where $\mathcal{B}_c$ is a nonempty set of training-example indices for class c. For positive directional weights ω_i such that $\sum_{i \in \mathcal{B}_c} \omega_i q(x_i) \neq 0$, define*

$$\mu_c(\mathcal{B}, \omega) = \frac{\sum_{i \in \mathcal{B}_c} \omega_i q(x_i)}{\|\sum_{i \in \mathcal{B}_c} \omega_i q(x_i)\|_2}, \qquad G_{\mathcal{B},\omega}(x) = \max_c q(x)^\top \mu_c(\mathcal{B}, \omega). \quad (19)$$

The collection of all such $G_{\mathcal{B},\omega}$ is the static centroid family induced by the zero-step reduction.

Unequal bank sizes are also compatible with a common class multiplier. For $K_c = |\mathcal{B}_c|$, replace $K^{-1}\sum_j$ in the class affinity by $K_c^{-1}\sum_{i\in\mathcal{B}_c}$ and choose

$$w_{ci} = \frac{K_c W \omega_i}{\sum_{j\in\mathcal{B}_c}\omega_j}.$$

Then $K_c^{-1}\sum_i w_{ci} = W$, while the normalized anchor weights are $\pi_{ci} = \omega_i / \sum_j \omega_j$. Thus bank selection and directional weighting identify different members of the same family under shared bandwidths.

Proposition 5.4 (CTM is a special case)**.** *Let $\mathcal{B}_{\text{all}}$ contain all training indices for each class. Write each training feature as $f(x_i) = r_i q(x_i)$ with $r_i = \|f(x_i)\|_2$. CTM's normalized class mean is $\mu_c(\mathcal{B}_{\text{all}}, r)$, and its maximum-cosine score is exactly $G_{\mathcal{B}_{\text{all}},r}$. Hence CTM is the full-bank, feature-norm-weighted member of Definition 5.3.*

The reduction separates three operations: removing candidate truncation, approximating each class field by a centroid, and imposing shared scales. The approximation is controlled by an error bound; the resulting radial centroid and cosine scores are exactly rank-equivalent under the shared-scale conditions. CTM is an existing member of this static family. This connection provides static controls for separating feature motion from score construction; the ranking equivalence does not extend to the original anchor field without accounting for the approximation error.

6. Experiments

The experiments address three questions: whether the tested attraction improves OOD ranking, whether local score alignment prevents ranking failure, and whether a boundary statistic explains the direction of small AUROC changes. A final comparison separates the steps of the static reduction. Destructive first-step interference is a theoretical possibility; the real-data studies below do not establish its occurrence.

6.1. Protocol

We use OpenOOD v1.5 splits (Zhang et al., 2023b). CIFAR-10, CIFAR-100, and ImageNet-200 use three official ResNet-18 checkpoints; ImageNet-1K uses a pinned torchvision ResNet-50 V1, with DenseNet-121 as an additional backbone. Tables A.19–A.20 report the individual Near-/Far-OOD datasets and results.

The original position-normalized dynamics use Gaussian anchor fields, $\Delta t = 0.05$, and $T \in \{0, 1, 3, 10\}$. CIFAR uses 80 anchors per class and all-class forces; ImageNet uses five anchors from 12 setup features per class and a fixed initial top-20 candidate set. The separate score-aligned audit uses five anchors from 12 setup features per class on every benchmark, all 10 classes on CIFAR-10, fixed top-20 candidates elsewhere, exact geodesic drift, $\Delta t = 0.05$, and $T \in \{0, 1, 3, 5, 10\}$. Current-winner, $T = 5$ is selected on an ImageNet-1K tuning split, checked on its disjoint holdout, and transferred without retuning to CIFAR-10, CIFAR-100, and ImageNet-200. The score-aligned rule and step count are fixed across benchmarks; each benchmark uses its own classifier and anchor field. The score-aligned study and the separate stability probe use distinct validation partitions, with parameter selection and holdout evaluation kept separate. The probe step size is selected on the source and fixed across targets. Neither study uses Near-/Far-OOD test loaders. The original trajectory comparisons and static baselines report Near-/Far-OOD test results. Proxy OOD is Tiny ImageNet for CIFAR and OpenImage-O for ImageNet.

Table 1: Gaussian trajectory ablation. The best nonzero T is selected retrospectively from $\{1, 3, 10\}$. Entries are AUROC/FPR@95; Δ is the AUROC change from $T = 0$.

ID	Group	$T = 0$	best $T > 0$	T	Δ
CIFAR-10	Near	88.59/42.36	88.95/40.91	10	+.36
	Far	91.07/31.09	91.56/29.95	10	+.49
CIFAR-100	Near	79.46/63.16	79.72/61.88	3	+.26
	Far	83.41/51.10	84.01/49.61	3	+.60
IN-200	Near	79.44/64.42	80.22/61.89	3	+.78
	Far	91.18/31.66	91.29/31.36	1	+.11
IN-1K	Near	**66.14/77.84**	64.51/79.27	1	−1.63
	Far	**77.44/65.16**	74.37/69.75	3	−3.07

The static comparisons use the same backbones and loaders. CTM uses the full training bank, while the small-bank control uses 80 directions per class on CIFAR and 12 on ImageNet. Its maximum cosine similarity is combined with MSP at the previously selected weight $\alpha = 0.8$ after ID-only standardization. We call this the centroid–MSP baseline and use it to examine the static construction in Section 5.

We report AUROC and FPR@95; higher AUROC and lower FPR@95 are better. Following the evaluation code, FPR@95 uses OOD as positive: it is the ID rejection rate at 95% OOD recall. AUROC assigns half credit to ties. CIFAR and ImageNet-200 results average three independently trained model seeds. ImageNet-1K variability for anchor-based methods is over three sampled banks on a fixed backbone. Selection and calibration settings are described in Appendix A.14.

6.2. Testing the attraction hypothesis

Table 1 compares the zero-step Gaussian field with the best nonzero step count in the tested grid. Small gains occur on CIFAR and ImageNet-200, but every tested nonzero Gaussian trajectory is worse on ImageNet-1K: even the retrospective best loses 1.63 Near-OOD and 3.07 Far-OOD AUROC points. The complete T grid, IMQ control, five-kernel study, and anchor-weighting ablation appear in Tables A.9, A.10, A.11, and A.12. Nonzero attraction is therefore not consistently better than $T = 0$ across these benchmarks. This comparison does not exclude improvements at other step sizes or trajectory lengths.

6.3. Ranking failure without class switching

The current-winner flow raises AUROC from 71.79 to 72.93 on the ImageNet-1K holdout, but this improvement does not transfer. Table 2 shows negative AUROC change on every seed of the three target benchmarks. Spoiled pairs exceed corrected pairs in all nine runs, as accounted for exactly by Theorem 4.7. Moreover, none of the 20,205 audited sample–seed trajectories changes its initial winning class. Thus ranking can deteriorate without a change of winning class.

Away from ties, the current-winner flow satisfies the local alignment condition in Definition 4.4. Numerical checks support the local expansion and confirm that the geodesic update preserves the sphere and tangent momentum to numerical precision. The audit finds no destructive first-step interference. Appendix A.4 reports the numerical checks and the comparison with the initial-winner flow; the exact equivalence in Proposition 4.5 additionally requires a unique winner at every force evaluation.

6.4. Evaluating the boundary-based diagnostic

To examine the local diagnostic suggested by Theorem 4.8, we use a separate perturbation with an explicit score derivative. A small repulsive step from a fixed class centroid provides this control: its score decrease can be calculated for each sample and compared across ID–OOD pairs. The probe uses the score at the perturbed feature, not the maximum along an attractive trajectory. Let $m(x)$ be MSP, $\hat{c}(x) = \arg\max_{1\le c\le C} q(x)^\top \mu_c$, $G(x) = q(x)^\top \mu_{\hat{c}(x)}$, with a deterministic rule for ties, and $D(x) = (1 - m(x))(1 - G(x)^2)$. Freeze m, $\hat{c}$, and the ID-fit calibration statistics, and set

$$\begin{aligned} q_\eta(x) &= \mathrm{Exp}_{q(x)}\Big[-\eta(1-m(x))P_{q(x)}\mu_{\hat{c}(x)}\Big], \\ s_\eta^{\mathrm{probe}}(x) &= (1-\alpha)\frac{m(x)-\bar{m}}{\sigma_m} + \alpha\frac{q_\eta(x)^\top\mu_{\hat{c}(x)} - \bar{G}}{\sigma_G}, \qquad \alpha = 0.8. \end{aligned} \tag{20}$$

Here Exp is the spherical exponential map. The means $\bar{m}, \bar{G}$ and positive standard deviations σ_m, σ_G are fitted on the ID-fit split and held fixed. The probe keeps the initial centroid and uses no maximum over the trajectory. Its score derivative is

$$\left.\frac{\partial s_\eta^{\mathrm{probe}}(x)}{\partial \eta}\right|_{\eta=0} = -\frac{\alpha}{\sigma_G} D(x).$$

Thus D measures the decrease up to the common positive factor α/σ_G. Define

$$M_\eta^{\mathrm{probe}} = s_\eta^{\mathrm{probe}}(X_{\mathrm{in}}) - s_\eta^{\mathrm{probe}}(X_{\mathrm{out}}).$$

The repulsive step gives

$$M_\eta^{\mathrm{probe}} = M_0^{\mathrm{probe}} + \frac{\alpha\eta}{\sigma_G}\{D(X_{\mathrm{out}}) - D(X_{\mathrm{in}})\} + O(\eta^2). \tag{21}$$

Under the theorem's density conditions, the first-order AUROC change is positive exactly when $\mathbb{E}[D(X_{\mathrm{out}}) - D(X_{\mathrm{in}}) \mid M_0^{\mathrm{probe}} = 0] > 0$. The uniform score remainder holds here because $m \in [0,1]$, $G \in [-1,1]$, and the calibration scales are fixed. To summarize pairs near this boundary, let $\mathcal{I}_\rho$ contain the empirical fraction ρ of ID–OOD pairs with the smallest absolute baseline margins, and define

$$B_\rho = \frac{1}{|\mathcal{I}_\rho|}\sum_{(i,j)\in\mathcal{I}_\rho}[D(x_j^{\mathrm{out}}) - D(x_i^{\mathrm{in}})]. \tag{22}$$

Here $x_i^{\mathrm{in}}, x_j^{\mathrm{out}}$ are validation inputs. We report $B_{5\%}$ and use $B_{100\%}$ for the mean over all pairs. Writing $p_{M_0^{\mathrm{probe}}}$ for the baseline margin density, the theorem gives

$$\begin{aligned} \Delta\,\mathrm{AUROC} = \eta\, p_{M_0^{\mathrm{probe}}}(0)\frac{\alpha}{\sigma_G} \\ \times \mathbb{E}[D(X_{\mathrm{out}}) - D(X_{\mathrm{in}}) \mid M_0^{\mathrm{probe}} = 0] + o(\eta). \end{aligned} \tag{23}$$

Both the boundary density and the calibration scale can vary across settings. We therefore use $B_{5\%}$ as a diagnostic of direction, rather than as a predictor of AUROC-change magnitude across settings.

The primary $\eta = 0.04$ is selected on ImageNet-1K/ResNet-50 and then fixed before all reported transfers. Table 3 shows small positive changes on both ImageNet-1K backbones but no positive target mean. The 5% boundary statistic has the same sign as AUROC change in 14 of 15 runs across five dataset–backbone settings. In contrast, the unconditional statistic is positive on all three transfer targets and has the opposite sign to their observed changes. Figure 4 visualizes this separation. This finite-window agreement is descriptive, not a direct estimate of the infinitesimal derivative or a significance test. The probe tests this local mechanism, not a performance-improvement claim; detailed boundary fractions and descriptive correlations are in Appendix A.5.

6.5. Auditing the zero-step reduction

The static comparisons distinguish the effect of feature motion from the choices used to construct the zero-step score. The six-stage comparison in Table 4 separates field, scale, centroid, bank, and fusion effects on fixed ImageNet-1K validation features. Candidate truncation is removed before S1, which already uses the full class set; S2 then shares the class scales. These are additional modeling choices beyond setting $T = 0$. Sharing class scale produces the largest single increase; radial centroid affinity and cosine then give identical rankings, exactly as Proposition 5.2 predicts. The empirical gain from centroid replacement is not certified by the global approximation bound, which preserves no audited pair by itself. Enlarging the bank and adding MSP are empirical controls rather than consequences of the reduction.

The static controls are competitive but not uniformly dominant. CTM is slightly stronger than the centroid–MSP baseline on four of eight Near-/Far-OOD AUROC tasks, CADRef leads on several ImageNet groups, and Scale is strongest on ImageNet-1K. Complete baselines, component ablations, class-count analysis, bank/calibration sensitivity, and batch-one efficiency are therefore placed in Appendix A.9. These comparisons show how training-bank estimation and confidence fusion affect the static baselines.

Table 2: Validation-only score-alignment audit. ImageNet-1K is the holdout after selection; other rows transfer the selected current-winner geodesic $T = 5$ rule without retuning. Entries are mean±standard deviation AUROC/FPR@95; Δ_0 is aligned $T = 5$ AUROC minus $T = 0$ AUROC.

Benchmark	$T = 0$	Aggregate $T = 5$	Current winner $T = 5$	Δ_0
ImageNet-1K	71.79 ± .82/72.63 ± 3.17	65.99 ± .23/74.59 ± 1.40	72.93 ± .57/72.03 ± 2.66	+1.14
CIFAR-10	89.42 ± .37/39.60 ± 3.22	90.13 ± .19/36.27 ± 7.10	87.98 ± .39/54.67 ± 7.63	−1.43
CIFAR-100	84.26 ± .88/47.07 ± 6.88	85.38 ± 1.02/45.27 ± 4.97	83.84 ± .89/48.93 ± 4.06	−.42
ImageNet-200	88.09 ± .69/38.78 ± 2.67	86.73 ± 1.48/38.61 ± 3.15	87.53 ± .93/40.17 ± 4.09	−.56

Table 3: Stability probe with a fixed step size. Static and probe columns report AUROC; $\overline{B}_{5\%}$ is $B_{5\%}$ from Eq. (22), averaged over three seeds.

Target	Static	Probe	Δ	$10^3\overline{B}_{5\%}$	Positive seeds
IN-1K/ResNet-50	91.67 ± .45	91.73 ± .44	+.056	+13.329	3/3
IN-1K/DenseNet-121	89.82 ± .40	90.04 ± .38	+.219	+34.395	3/3
CIFAR-10	89.46 ± .81	89.45 ± .81	−.006	−.551	0/3
CIFAR-100	84.96 ± .73	84.95 ± .74	−.016	−.589	1/3
ImageNet-200	92.11 ± .32	92.08 ± .33	−.027	−1.871	0/3

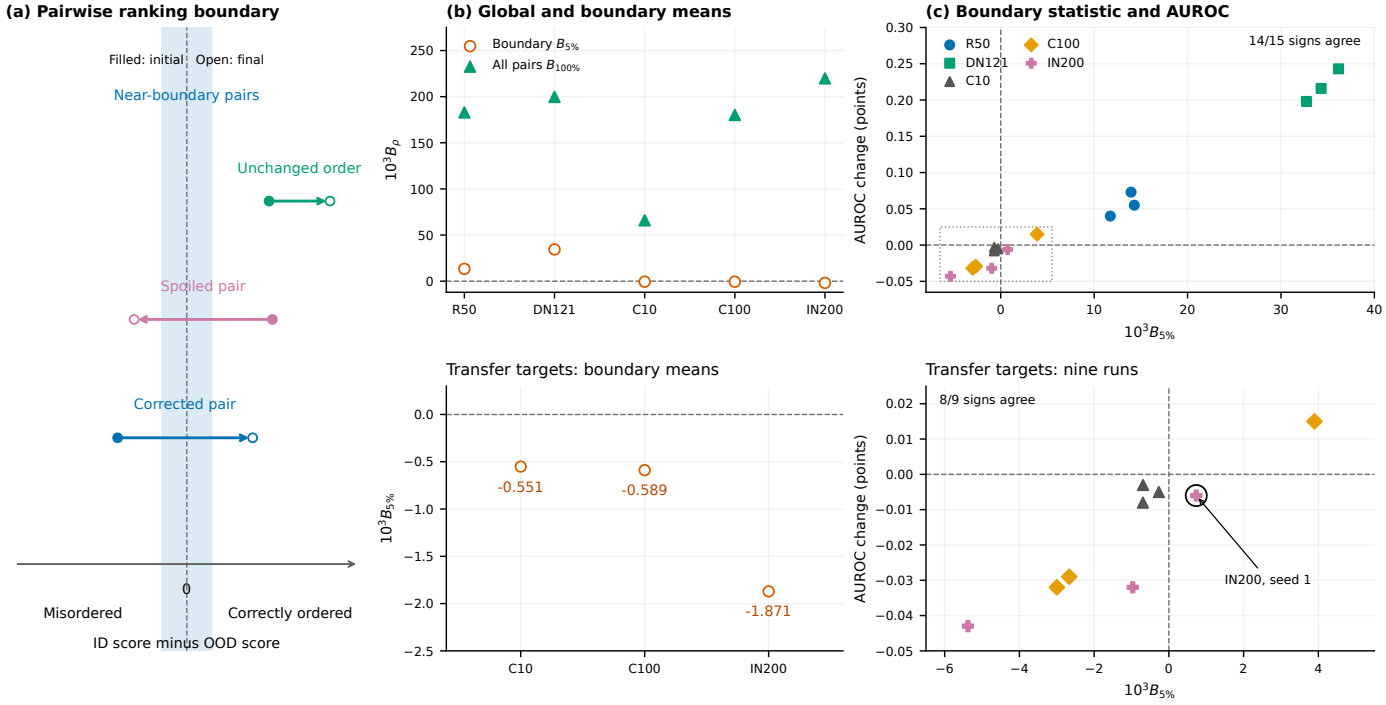

Figure 4: Boundary-conditioned mechanism and validation evidence. (a) AUROC changes when an ID–OOD pair crosses the ranking boundary; filled and open markers denote initial and final margins. (b) Mean $D_{\text{out}} - D_{\text{in}}$ over all pairs and over the closest 5% of baseline margins. The lower view enlarges the negative boundary means on the three transfer targets. (c) Boundary statistic $B_{5\%}$ versus AUROC change at $\eta = .04$, with the nine cross-dataset target–seed audits enlarged below. Signs agree in 14 of 15 audits overall and 8 of 9 cross-dataset audits; the circled point is the sole disagreement. AUROC changes are percentage points. R50 and DN121 denote the ImageNet-1K ResNet-50 and DenseNet-121 settings; C10, C100, and IN200 denote CIFAR-10, CIFAR-100, and ImageNet-200. Both statistics are computed on the same validation holdout.

7. Discussion and Limitations

7.1. Interpretation of the mismatch theory

The theory separates the driving field, the instantaneous detection score, and the relative ordering of ID and OOD scores. The first-step theorem connects the field to the instantaneous score for a step from rest. Keeping the maximum score along a trajectory guarantees a nonnegative score increase, but the exact pair decomposition shows why this need not improve AUROC. For score families satisfying the additional expansion and density conditions, the boundary theorem shows that the mean relative score derivative at equal baseline scores, together with the margin density there, determines the first-order AUROC change.

The transfer experiment shows that local score alignment and an unchanged winning class are insufficient to ensure better ranking. This failure does not require the destructive class interference allowed by the first-step analysis. The separate probe

Table 4: Step-by-step static comparison on ImageNet-1K validation data. Values are mean±standard deviation over three detector banks; deltas are relative to MSP.

Stage	Interpretation	AUROC	FPR@95	ΔAUROC	ΔFPR@95
MSP	output reference	84.77 ± .00	49.26 ± .00	0.00	0.00
S1	exact full-class field	73.16 ± .86	72.34 ± 2.61	−11.61	+23.08
S2	common class scale	87.06 ± .69	46.44 ± .13	+2.30	−2.82
S3	radial anchor centroid	90.24 ± .61	38.61 ± .75	+5.47	−10.65
S4	cosine implementation	90.24 ± .61	38.61 ± .75	+5.47	−10.65
S5	setup-bank centroid	91.19 ± .67	35.36 ± .34	+6.43	−13.90
S6	fixed centroid–MSP control	**92.02 ± .52**	**31.58 ± .71**	**+7.26**	**−17.68**

illustrates why averaging score changes over all pairs can miss the direction of AUROC change: the pairs near equal baseline scores can move differently from the population as a whole.

The zero-step analysis identifies the geometric information present before motion. Removing candidate truncation is an additional modeling choice. The approximation bound then controls the replacement of full-class anchor fields by radial centroid scores. Under shared scales, the resulting centroid score is rank-equivalent to maximum cosine similarity. CTM is an existing full-bank, feature-norm-weighted member of this centroid family. The step-by-step experiment examines these scoring choices separately from the additional effects of bank size and confidence fusion. These comparisons help attribute static performance to scoring choices; they do not establish that centroid replacement improves AUROC in general.

7.2. A practical diagnostic and its scope

Theorem 4.8 suggests a validation diagnostic for score refinements satisfying its regularity conditions. On ID and proxy-OOD validation data allowed by the evaluation protocol, estimate each sample's score derivative, form ID–OOD pairs, retain a fixed fraction with the smallest absolute baseline margins, and average the ID-minus-OOD relative derivative. Its sign provides a local prediction to compare with the observed validation AUROC change. Stability across seeds and boundary fractions should be checked in the intended validation setting. The mean relative derivative over all pairs provides a useful comparison, because its sign can differ from that near the ranking boundary.

The empirical conclusions are limited to the tested settings. The dynamics use fixed initial candidate sets, and the no-switch audit covers $T \leq 10$; the ImageNet audit excludes classes outside the retained top-20 set. The score-aligned transfers and the separate probe use validation data. The probe's 15 runs cover three seeds in each of five dataset–backbone settings, with shared validation images within each setting. For each ImageNet-1K backbone, variation is over sampled banks. These runs provide descriptive evidence, not 15 independent datasets or a significance test.

The diagnostic also has a local scope. The fixed 5% window summarizes pairs near the ranking boundary; it does not directly estimate the conditional mean at exactly equal scores. The observed AUROC changes at $\eta = 0.04$ are finite-step changes, several of which are very small. Agreement between their signs and the boundary statistic therefore supports a local interpretation without guaranteeing gains at other step sizes or on other distributions. Theorem 4.9 further shows that changing a strict baseline ordering can reduce AUROC for some ID/OOD distributions.

The numerical model and static approximation impose further limits. The original position-normalized scheme has substantial energy and tangency errors, so the geodesic control provides the evidence for constrained feature motion. The centroid approximation bound is too loose to certify the observed static rankings. Whether the diagnostic extends to transformers and foundation-model features, or can be used with ID-only validation objectives, remains open.

8. Conclusion

Class attraction and local score alignment do not guarantee improved OOD ranking. The transfer study demonstrates this failure even without observed winning-class changes within the fixed candidate sets. Under regularity conditions, the mean relative score derivative at the ID–OOD ranking boundary determines the direction of first-order AUROC change. A separate perturbation study finds that near-boundary relative score changes generally track the direction of small AUROC changes. Through centroid approximation and shared-scale ranking equivalence, the static analysis relates zero-step anchor scores to a centroid family that includes CTM. Together, these results show why samplewise score increases are insufficient and why feature refinement should be evaluated by its effect on ID–OOD ordering.

Code availability

Implementation and evaluation scripts are available in our public GitHub repository. Software versions, pinned dependencies, and checkpoint identifiers are specified in the appendix.

Appendix A. Mathematical Details and Supplementary Results

Appendix A.1. Fundamental-model properties and numerical update

Lemma A.1 (Transport of local class concentration)**.** *Let $\mathcal{M}_c \subset \mathcal{X}$ be the class-c image manifold and let $q : \mathcal{X} \to \mathbb{S}^{d-1}$ be L_c-Lipschitz on $U_c = \{x : d_{\mathcal{X}}(x, \mathcal{M}_c) \le r_c\}$. If a class-c random image X_c satisfies*

$$\Pr[d_{\mathcal{X}}(X_c, \mathcal{M}_c) \le r_c] \ge 1 - \delta_c, \tag{A.1}$$

then its normalized feature satisfies

$$\Pr[d_{\mathbb{S}}(q(X_c), q(\mathcal{M}_c)) \le L_c r_c] \ge 1 - \delta_c, \tag{A.2}$$

where $d_{\mathbb{S}}(u, B) = \inf_{v \in B} \|u - v\|_2$.

Proof. For every $x \in U_c$ and $m \in \mathcal{M}_c$, Lipschitz continuity gives

$$d_{\mathbb{S}}(q(x), q(\mathcal{M}_c)) \le \|q(x) - q(m)\|_2 \le L_c d_{\mathcal{X}}(x, m).$$

Taking the infimum over $m \in \mathcal{M}_c$ yields $d_{\mathbb{S}}(q(x), q(\mathcal{M}_c)) \le L_c d_{\mathcal{X}}(x, \mathcal{M}_c)$. Applying this deterministic inequality on the event in Eq. (A.1) gives Eq. (A.2). □

Original position-normalized update.. The original trajectory experiments use ambient kicks with position normalization. With the initial candidate set fixed, one update is

$$\begin{aligned} p_{n+\frac{1}{2}} &= p_n + \frac{\Delta t}{2} \nabla_q B(q_n; q_0), \\ \widetilde{q}_{n+1} &= q_n + \Delta t\, p_{n+\frac{1}{2}}, \\ q_{n+1} &= \widetilde{q}_{n+1} / \|\widetilde{q}_{n+1}\|_2, \\ p_{n+1} &= p_{n+\frac{1}{2}} + \frac{\Delta t}{2} \nabla_q B(q_{n+1}; q_0). \end{aligned} \tag{A.3}$$

The normalization keeps positions on $\mathbb{S}^{d-1}$, but momentum is not constrained to the tangent space. This is a distinct heuristic discretization, not a symplectic integrator for Eq. (4); no long-time energy or tangency guarantee is claimed for it.

Constrained geodesic update.. For the score-aligned audit, let F denote the aggregate or single-class driver in Section 4.2. Starting with tangent p_n, set

$$\widetilde{v}_n = p_n + \tfrac{1}{2} \Delta t P_{q_n} \nabla F(q_n), \qquad \nu_n = |\widetilde{v}_n\|_2.$$

Exact kinetic drift followed by a tangent kick gives

$$\begin{aligned} q_{n+1} &= \cos(\Delta t \nu_n) q_n + \frac{\sin(\Delta t \nu_n)}{\nu_n} \widetilde{v}_n, \\ \widetilde{p}_{n+1} &= \cos(\Delta t \nu_n) \widetilde{v}_n - \nu_n \sin(\Delta t \nu_n) q_n, \\ p_{n+1} &= \widetilde{p}_{n+1} + \tfrac{1}{2} \Delta t P_{q_{n+1}} \nabla F(q_{n+1}). \end{aligned} \tag{A.4}$$

The quotient is interpreted as Δt when $\nu_n = 0$. For a tangent vector v, the spherical exponential map is

$$\mathrm{Exp}_q(v) = \cos(\|v\|_2) q + \frac{\sin(\|v\|_2)}{\|v\|_2} v, \qquad \mathrm{Exp}_q(0) = q.$$

The two kicks and the drift preserve the tangent bundle in exact arithmetic. A deterministic active class is selected at a current-winner tie, and smooth-flow claims are made only away from ties. Both schemes keep $\mathcal{J}(q_0)$ fixed; only the active class within that set may change.

Appendix A.2. Proofs for the attraction–ranking analysis

Proof of Theorem 4.1. The first half-step and position update give

$$p_{1/2} = \frac{\Delta t}{2} \nabla B(q_0; q_0), \qquad \widetilde{q}_1 = q_0 + \frac{\Delta t^2}{2} \nabla B(q_0; q_0).$$

For a unit vector q and small h,

$$\frac{q + hv}{\|q + hv\|_2} = q + h(I - qq^\top)v + O(h^2).$$

Substitution of $h = \Delta t^2/2$ proves Eq. (6) for the original update. For the geodesic scheme, $\nu_0 = \frac{1}{2}\Delta t\|b_0\|_2$ and $q_1 = \mathrm{Exp}_{q_0}(\frac{1}{2}\Delta t^2 b_0)$, which has the same expansion with remainder $O(\Delta t^4)$. Taylor expansion at q_0 gives

$$B(q_1; q_0) - B(q_0; q_0) = \frac{\Delta t^2}{2}\nabla B(q_0; q_0)^\top P_{q_0}\nabla B(q_0; q_0) + O(\Delta t^4),$$

which is Eq. (7) because P_{q_0} is an orthogonal projector. The same expansion for A_c yields Eq. (8). □

Proof of Proposition 4.5. Write $z_n^{\mathrm{init}} = (q_n^{\mathrm{init}}, p_n^{\mathrm{init}})$ and $z_n^{\mathrm{cur}} = (q_n^{\mathrm{cur}}, p_n^{\mathrm{cur}})$. The initial states agree. Suppose $z_n^{\mathrm{init}} = z_n^{\mathrm{cur}} = z_n$. Since c_0 is the unique maximizer at q_n, continuity gives a neighborhood on which $F_{\mathrm{cur}} = F_{\mathrm{init}} = KA_{c_0}$. Hence

$$P_{q_n}\nabla F_{\mathrm{cur}}(q_n) = P_{q_n}\nabla F_{\mathrm{init}}(q_n).$$

The first kicks produce identical intermediate momenta, and the deterministic drift produces the same q_{n+1} and intermediate momentum. The unique-winner assumption at q_{n+1} gives the same equality of projected gradients there, so the second kicks imply $z_{n+1}^{\mathrm{init}} = z_{n+1}^{\mathrm{cur}}$. Induction proves equality of all states and active classes. The same argument applies at any additional internal force evaluation points covered by the assumption. Finally, at every retained state,

$$m_{\mathcal{J}}(q_n; q_0) = A_{c_0}(q_n), \qquad F_{\mathrm{cur}}(q_n) = F_{\mathrm{init}}(q_n).$$

Taking maxima over $0 \le n \le T$ proves equality of trajectory scores. □

Proof of Corollary 4.3. Linearity of P_{q_0} and Eq. (3) imply

$$P_{q_0}\nabla B(q_0; q_0) = K \sum_{c \in \mathcal{J}(q_0)} g_c.$$

Hence

$$\left\langle g_{c^\star}, P_{q_0}\nabla B(q_0; q_0)\right\rangle = K\|g_{c^\star}\|_2^2[1 + I(q_0)],$$

which proves Eq. (10). If $I(q_0) < -1$, its leading coefficient is negative. It also follows that $P_{q_0}\nabla B(q_0; q_0) \neq 0$, so Eq. (7) has a positive leading coefficient. The condition $I(q_0) < -1$ requires at least two candidate classes. Since $c^\star$ is the unique maximizer, define

$$\delta_* = A_{c^\star}(q_0) - \max_{c \in \mathcal{J}(q_0)\setminus\{c^\star\}} A_c(q_0) > 0.$$

Continuity and finiteness of the candidate set give a neighborhood U of q_0 on which $|A_c(q) - A_c(q_0)| < \delta_*/3$ for every candidate class. Thus $A_{c^\star}(q) - A_c(q) > \delta_*/3$ for all $q \in U$ and $c \neq c^\star$. Since $q_1 \to q_0$ as $\Delta t \to 0$, the maximizer is unchanged for all sufficiently small nonzero steps. The instantaneous maximum therefore decreases while B increases. □

Proof of Theorem 4.7. Under the no-tie assumption,

$$\mathrm{AUROC}_T = \Pr(M_0 + D_T > 0), \qquad \mathrm{AUROC}_0 = \Pr(M_0 > 0).$$

The pointwise indicator identity is

$$\mathbf{1}_{\{M_0+D_T>0\}} - \mathbf{1}_{\{M_0>0\}} = \mathbf{1}_{C_T} - \mathbf{1}_{\mathcal{S}_T}.$$

Taking expectations proves the first statement. A changed ordering requires $|D_T| \ge |M_0|$, so

$$|\Delta\,\mathrm{AUROC}\,| \le \mathbb{E}|\mathbf{1}_{C_T} - \mathbf{1}_{\mathcal{S}_T}| \le \Pr(|M_0| \le |D_T|).$$

If ties are present, Eq. (12) instead gives the exact identity

$$\begin{aligned}\Delta\,\mathrm{AUROC} &= \Pr(C_T) - \Pr(\mathcal{S}_T) \\ &\quad + \tfrac{1}{2}[\Pr(M_0 + D_T = 0) - \Pr(M_0 = 0)].\end{aligned} \tag{A.5}$$

Thus empirical pair counts must include the change in tie mass when compared with AUROC computed using half credit for ties. □

Proof of Theorem 4.8. First control the tail of the relative score derivative. Omit the remainder and write $Z_\eta = M_0 + \eta U$. Set $a_\eta = \delta/(2|\eta|)$ for $\eta \neq 0$. Since $\mathbb{E}|U| < \infty$,

$$\frac{\Pr(|U| > a_\eta)}{|\eta|} \leq \frac{2}{\delta}\mathbb{E}\Big[|U|\mathbf{1}_{\{|U|>a_\eta\}}\Big] \longrightarrow 0.$$

Next evaluate the derivative on the remaining region. For $|u| \leq a_\eta$, the oriented integral identity

$$\frac{1}{\eta}\int_{-\eta u}^{0} p_{M_0,U}(m,u)\,dm = u\int_0^1 p_{M_0,U}(-t\eta u,u)\,dt$$

is bounded in absolute value by $|u|g(u)$. Continuity in m, dominated convergence, and the negligible tail therefore give

$$\lim_{\eta\to 0}\frac{\Pr(Z_\eta > 0) - \Pr(M_0 > 0)}{\eta} = \int u\, p_{M_0,U}(0,u)\,du.$$

Finally restore the remainder. Since

$$|R_\eta(X_{\text{in}}) - R_\eta(X_{\text{out}})| \leq 2\varepsilon_\eta \quad \text{almost surely,}$$

a change in ordering, including the creation of a tie, requires $|Z_\eta| \leq 2\varepsilon_\eta$. For sufficiently small $|\eta|$, the corresponding intervals in m lie in $[-\delta,\delta]$ whenever $|u| \leq a_\eta$, so

$$\Pr(|Z_\eta| \leq 2\varepsilon_\eta) \leq 4\varepsilon_\eta\int g(u)\,du + \Pr(|U| > a_\eta) = o(|\eta|).$$

Thus the remainder does not change the derivative, also under the usual half-credit convention for ties. Finally, $\int u\, p_{M_0,U}(0,u)\,du = p_{M_0}(0)\mathbb{E}[U \mid M_0 = 0]$ by the stated density definition. □

Connection to a single trajectory step.. For a unique initial winner, Theorem 4.1 and the positive class margin imply, with $\eta = \Delta t^2/2 \geq 0$,

$$m_{\mathcal{J}}(q_1;q_0) - m_{\mathcal{J}}(q_0;q_0) = \eta\langle g_{c^\star}, b_0\rangle + O(\eta^2).$$

Since $u \mapsto \max\{u,0\}$ is 1-Lipschitz, retaining the initial score gives

$$S_1(q_0) - S_0(q_0) = \eta\max\{\langle g_{c^\star}, b_0\rangle, 0\} + O(\eta^2).$$

This is a pointwise expansion. Applying the boundary theorem to this family would additionally require a uniform remainder over the two populations and the stated joint-density conditions. The separate probe in Section 6.4 provides a simpler score family with a uniform remainder.

Appendix A.3. Proofs for the static reduction

Proof of Proposition 5.1. The definitions give

$$A_c(q) = W_c\sum_j \pi_{cj}\phi\left(\frac{\|q - a_{cj}\|_2}{\sigma_c}\right).$$

The reverse triangle inequality and Lipschitz continuity of ϕ imply

$$\begin{aligned}
&\left|A_c(q) - W_c\phi\left(\frac{\|q-\bar a_c\|_2}{\sigma_c}\right)\right| \\
&\quad\leq \frac{W_c L_\phi}{\sigma_c}\sum_j \pi_{cj}\left|\|q-a_{cj}\|_2 - \|q-\bar a_c\|_2\right| \\
&\quad\leq \frac{W_c L_\phi}{\sigma_c}\sum_j \pi_{cj}\|a_{cj}-\bar a_c\|_2.
\end{aligned}$$

Because $\bar a_c = \|\bar a_c\|_2\mu_c$ and the anchors lie on the unit sphere,

$$\|\bar a_c - \mu_c\|_2 = |1 - \|\bar a_c\|_2|\,.$$

A second application of the Lipschitz inequality bounds the replacement of $\bar{a}_c$ by μ_c by $W_c L_\phi \left|1 - \|\bar{a}_c\|_2\right| / \sigma_c$. The triangle inequality combines the two displays into Eq. (16). To obtain the score bound, use

$$\left|\max_c A_c(q) - \max_c \widetilde{A}_c(q)\right| \le \max_c \left|A_c(q) - \widetilde{A}_c(q)\right|$$

and apply the class-wise bound. For two queries q, q', it follows that

$$\widetilde{S}_0(q) - \widetilde{S}_0(q') \ge S_0^{\text{all}}(q) - S_0^{\text{all}}(q') - 2\epsilon.$$

This proves the margin condition in Eq. (17). □

Proof of Proposition 5.2. For unit vectors,

$$\|q - \mu_c\|_2^2 = 2 - 2q^\top \mu_c.$$

Thus maximizing cosine similarity over classes minimizes radial distance. Because ϕ is strictly decreasing and $W > 0$, substituting the minimizing distance gives Eq. (18). The map $g \mapsto W\phi(\sqrt{2 - 2g}/\sigma)$ is strictly increasing on $[-1, 1]$, which proves both the sample-ranking and class-selection statements. □

Proof of Proposition 5.4. CTM forms an unnormalized class mean proportional to

$$\sum_{i \in \mathcal{B}_c} f(x_i) = \sum_{i \in \mathcal{B}_c} r_i q(x_i).$$

Normalizing this vector removes the irrelevant factor $|\mathcal{B}_c|^{-1}$ and yields $\mu_c(\mathcal{B}_{\text{all}}, r)$ in Eq. (19). Maximizing its cosine with $q(x)$ is therefore exactly $G_{\mathcal{B}_{\text{all}}, r}(x)$. □

Appendix A.4. Score-aligned transfer and winner-switch audit

Table A.5 gives the per-seed validation results under the locked geodesic $T = 5$ rule. Corrected and spoiled columns are the probability masses from Theorem 4.7, expressed as percentages. AUROC uses half credit for ties; when tie mass changes, the exact relation includes Eq. (A.5). Displayed pair masses and AUROC changes are rounded separately.

Table A.5: Per-seed score-aligned mechanism audit. Static, aggregate, and aligned columns report AUROC; C/S gives corrected/spoiled pair mass in percent for the aligned flow relative to $T = 0$.

Validation benchmark	Seed	$T = 0$	Aggregate	Aligned	C/S
ImageNet-1K	0	72.27	65.77	73.07	2.74/1.94
	1	70.84	66.24	72.30	3.02/1.57
	2	72.27	65.95	73.41	3.02/1.88
CIFAR-10	0	89.23	89.93	87.63	0.57/2.17
	1	89.18	90.18	87.93	0.77/2.01
	2	89.85	90.29	88.39	0.46/1.92
CIFAR-100	0	85.15	86.44	84.73	0.50/0.91
	1	83.40	84.40	82.95	0.53/0.97
	2	84.24	85.31	83.84	0.48/0.88
ImageNet-200	0	87.62	85.79	87.16	0.78/1.25
	1	88.88	88.44	88.59	0.75/1.04
	2	87.77	85.98	86.85	0.59/1.51

The post-hoc switch audit uses all retained states $T = 0, \ldots, 10$ rather than only the selected $T = 5$ result. Table A.6 aggregates the three seeds. "Candidate classes" denotes the set fixed at initialization; the ImageNet-200 and CIFAR-100 statement does not claim that classes outside the top-20 set were re-evaluated during motion. The two driver implementations also have zero maximum absolute difference in trajectory score, instantaneous score, driver value, and Hamiltonian energy.

Table A.6: Winner-switch audit over 20,205 sample–seed trajectories and 11 retained states. Switch counts are summed over seeds and populations.

Benchmark	Candidates	ID	OOD	Current switches	Initial switches
CIFAR-10	all 10	3000	3000	0	0
CIFAR-100	top 20	3000	2916	0	0
ImageNet-200	top 20	3000	5289	0	0

The mechanism and numerical checks average over ID/OOD populations and three anchor banks. The aggregate and current-winner alignment cosines are 0.634 and 1.000; neither produces a first-step winner decrease, and the audited fraction satisfying $I(q_0) < -1$ is zero. Thus these data do not demonstrate destructive class interference, although they test its local alignment expansion. Both attain 100% Taylor sign agreement. Winner and driver Taylor residuals have log–log slopes 4.138–4.162 and 3.833–4.179, consistent with the $O(\Delta t^4)$ remainder underlying Definition 4.4. The geodesic controls keep maximum sphere error below 2.38×10^{-7} and maximum tangency error below 1.49×10^{-7}; their mean maximum relative energy drifts are 0.0352 (aggregate) and 0.00265 (current winner). For comparison, the original retraction-based aggregate integrator has drift 1.327 and tangency error 8.604. The stable geodesic experiments therefore provide the basis for the aligned-flow mechanism claim; their failed locked transfer is not explained by leaving the sphere or tangent bundle.

Appendix A.5. Boundary-conditioned stability-probe audit

The primary stability-probe rule fixes $\eta = 0.04$ from the ImageNet-1K ResNet-50 selection split. No transfer target re-selects the step size. A secondary scale control sets

$$\eta_{\mathrm{ID}} = 0.0025/\mathbb{E}_{\mathrm{ID,fit}}[D]$$

using only the ID fit split; its results are retained in the accompanying CSV archive as an additional control rather than in the primary tables. Both rules retain the same MSP component and the same ID-fit standardization as the static score.

The probe in Eq. (20) is generated by the conditional potential $V_x(q) = \gamma(1 - m(x))q^\top \mu_{\hat{c}(x)}$, where $\gamma > 0$ and the initial confidence and centroid are held fixed. One geodesic kick–drift step from rest has $\eta = \gamma\Delta t^2/2$; the final kick does not change its position. Thus this is a one-step repulsive perturbation, not an exact finite-time solution of the full nonlinear flow or a maximum score along a trajectory. For completeness, with $\theta = \eta(1 - m)\sqrt{1 - G^2}$, the initial-winning geodesic has the exact endpoint

$$G_\eta = G\cos\theta - \sqrt{1 - G^2}\sin\theta = G - \eta(1 - m)(1 - G^2) + O(\eta^2).$$

Substitution into the fixed standardized fusion yields Eq. (21) in the main text. Table A.7 reports the primary rule. Corrected and spoiled masses are percentages of all ID–OOD pairs. Their difference gives the AUROC change in percentage points, with the tie adjustment in Eq. (A.5) when needed.

Table A.7: Per-seed stability-probe results. $B_{5\%}$ is the mean $D_{\mathrm{out}} - D_{\mathrm{in}}$ among the nearest 5% of pairs to the static ranking boundary. C/S denotes corrected/spoiled pair mass in percent.

Validation target	Seed	Static	Probe	Δ	$10^3 B_{5\%}$	C/S
IN-1K / R50	0	92.165	92.205	+0.040	+11.734	0.296/0.256
	1	91.269	91.342	+0.073	+13.950	0.336/0.263
	2	91.583	91.638	+0.055	+14.302	0.325/0.269
IN-1K / DN121	0	90.065	90.281	+0.216	+34.302	0.553/0.337
	1	89.357	89.600	+0.243	+36.159	0.581/0.338
	2	90.049	90.246	+0.198	+32.726	0.537/0.339
CIFAR-10	0	88.859	88.851	−0.008	−0.694	0.008/0.016
	1	89.135	89.131	−0.003	−0.691	0.006/0.010
	2	90.377	90.372	−0.005	−0.268	0.001/0.006
CIFAR-100	0	85.375	85.343	−0.032	−2.998	0.108/0.140
	1	84.127	84.098	−0.029	−2.666	0.122/0.152
	2	85.392	85.407	+0.015	+3.897	0.143/0.128
IN-200	0	92.069	92.026	−0.043	−5.378	0.106/0.149
	1	92.448	92.441	−0.006	+0.733	0.113/0.120
	2	91.819	91.788	−0.032	−0.966	0.109/0.141

The single sign disagreement occurs for ImageNet-200 seed 1, where both the boundary statistic and the AUROC change are close to zero. The pair decomposition accounts for AUROC change by construction. The additional finding is the sign agreement with $B_{5\%}$, which is computed from baseline margins and score derivatives on the same holdout. Across the 15 runs in five dataset–backbone settings, the descriptive Pearson and Spearman correlations between $B_{5\%}$ and AUROC change are 0.991 and 0.979. These pooled correlations do not establish a common quantitative relation across settings: the boundary density and calibration scale also enter the first-order AUROC change. Runs within a setting share validation images, so these correlations are not significance tests based on 15 independent datasets.

Table A.8 varies the width of the empirical boundary layer. ImageNet-1K has a positive boundary statistic throughout. For every transfer target the full-pair mean is positive, yet the narrowest layers are negative. This is the observed global-to-boundary sign reversal behind the failed transfer.

Table A.8: Mean boundary statistic B_ρ, multiplied by 10^3, over three seeds. Columns retain the stated fraction of ID–OOD pairs with the smallest absolute static margin.

Validation target	1%	2.5%	5%	10%	25%	100%
IN-1K / R50	11.128	12.013	13.329	18.480	43.652	182.907
IN-1K / DN121	32.548	33.590	34.395	38.253	59.390	199.719
CIFAR-10	−0.527	−0.567	−0.551	−0.457	0.277	65.996
CIFAR-100	−0.934	−1.616	−0.589	1.831	16.693	180.182
IN-200	−4.551	−3.871	−1.871	4.976	40.295	219.853

The predefined transfer criteria for AUROC and boundary consistency are not met, although the FPR@95 criterion is met. The probe was therefore not evaluated on the Near-/Far-OOD test sets; no corresponding test loader or score file was accessed in this study.

Appendix A.6. Full ImageNet trajectory ablation

Table A.9 reports the aggregated ImageNet trajectory grid. Values are Near-/Far-OOD macro averages; each cell is AUROC/FPR@95. The Gaussian and IMQ fields share the same anchors, optimizer schedule, candidate count, and step size.

Table A.9: ImageNet trajectory ablation. ImageNet-200 varies three official model seeds; ImageNet-1K varies three centroid banks on one fixed backbone.

ID	Potential	$T = 0$	$T = 1$	$T = 3$	$T = 10$
IN-200 Near	Gaussian	79.44/64.42	79.64/63.99	80.22/61.89	79.63/62.10
IN-200 Far	Gaussian	91.18/31.66	91.29/31.36	90.68/33.10	88.82/34.61
IN-200 Near	IMQ	78.23/65.27	78.41/64.95	79.34/62.99	78.16/63.37
IN-200 Far	IMQ	89.79/34.33	89.90/34.06	90.17/33.50	87.09/37.39
IN-1K Near	Gaussian	66.14/77.84	64.51/79.27	61.42/79.63	60.94/79.63
IN-1K Far	Gaussian	77.44/65.16	72.71/71.29	74.37/69.75	68.51/70.71
IN-1K Near	IMQ	63.46/80.45	60.91/82.52	61.94/81.33	58.03/82.91
IN-1K Far	IMQ	73.76/69.48	65.68/78.21	75.58/66.79	63.44/75.60

Table A.10 gives the uncertainty and precision–recall statistics for the ImageNet-1K $T = 10$ column.

Table A.10: ImageNet-1K $T = 10$ full-evaluation results over three independently sampled centroid banks on the fixed ResNet-50. Values are mean±standard deviation; ID accuracy is 76.1867% for every run.

Potential	Group	AUROC	FPR@95	AUPR-IN	AUPR-OUT
Gaussian	Near	60.94 ± .07	79.63 ± 1.29	75.14 ± .17	34.70 ± .24
Gaussian	Far	68.51 ± .33	70.71 ± 1.75	90.86 ± .26	26.15 ± .04
IMQ	Near	58.03 ± .83	82.91 ± 1.79	73.52 ± .22	33.27 ± .14
IMQ	Far	63.44 ± 1.05	75.60 ± 1.52	88.96 ± .53	23.21 ± .15

Table A.11: Five potential families on CIFAR. Bandwidths are fitted from zero-step class affinities and scores use $T = 120$. Values are macro averages over all six OOD datasets and three seeds.

ID	Potential	AUROC	FPR@95
CIFAR-10	Gaussian	89.77 ± 0.13	35.18 ± 1.24
	Laplacian	**90.97 ± 0.36**	32.82 ± 2.19
	Cauchy	89.94 ± 0.24	34.55 ± 1.11
	IMQ	90.66 ± 0.05	**31.73 ± 1.27**
	Matérn-3/2	89.97 ± 0.44	37.46 ± 3.18
CIFAR-100	Gaussian	**82.58 ± 0.61**	**53.70 ± 0.91**
	Laplacian	81.60 ± 0.46	55.16 ± 0.94
	Cauchy	81.19 ± 0.36	54.99 ± 0.81
	IMQ	82.35 ± 0.53	54.44 ± 0.73
	Matérn-3/2	81.07 ± 0.74	54.46 ± 0.92

Appendix A.7. Effective-rank anchor-weighting ablation

Table A.12 reports the fitted anchor-weighting variants, and Fig. A.5 displays changes from uniform weights. Image weights are the effective rank of mean-centered luminance arrays, computed as the exponential entropy of normalized singular values. Bandwidths are refitted for each variant; global weight normalization alone would leave a fixed static field's ranking unchanged.

Table A.12: Static Gaussian anchor-weighting ablation. Each Near/Far entry is AUROC/FPR@95 at $T = 0$ over three seeds. "Class" and "global" divide raw effective-rank weights by their class-wise or global mean.

ID	Anchor weighting	Near-OOD	Far-OOD
CIFAR-10	Uniform	88.59/42.36	91.07/31.09
	Effective rank/raw	76.99/77.43	81.43/55.91
	Effective rank/class	88.53/42.85	91.04/31.12
	Effective rank/global	79.35/71.87	83.17/55.49
CIFAR-100	Uniform	79.46/63.16	83.41/51.10
	Effective rank/raw	72.55/74.41	84.33/50.11
	Effective rank/class	79.41/63.35	83.66/50.63
	Effective rank/global	75.69/71.72	85.27/48.29
IN-200	Uniform	79.44/64.42	91.18/31.66
	Effective rank/raw	77.17/69.34	89.14/41.41
	Effective rank/class	78.67/65.05	90.11/33.86
	Effective rank/global	75.03/69.75	85.22/48.57

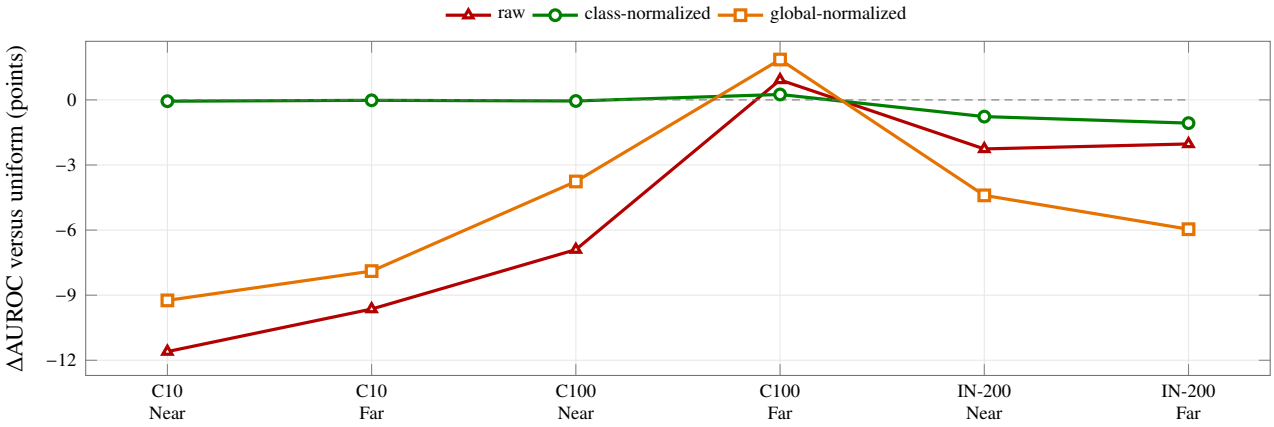

Figure A.5: Effective-rank anchor weighting relative to the uniform static field. Raw and global normalization usually reduce AUROC; class-wise normalization mostly returns to the uniform baseline. The isolated CIFAR-100 Far-OOD gain does not transfer to its Near-OOD counterpart or ImageNet-200.

Appendix A.8. Validation analysis of the static reduction

The staged analysis uses 5000 ImageNet-1K ID and 1763 OpenImage-O validation examples. Table 4 reports stage performance; the following diagnostics quantify adjacent transitions and approximation error.

Table A.13: Full-split transition deltas in the static-reduction sequence. Each entry is the mean change over the three detector seeds and is written as AUROC/FPR@95.

Transition	ΔAUROC	ΔFPR@95
S1 → S2	+13.91	−25.90
S2 → S3	+3.18	−7.83
S3 → S4	0.00	0.00
S4 → S5	+0.95	−3.25
S5 → S6	+0.83	−3.78
MSP → S6	+7.26	−17.68

The radial-centroid bound audit evaluates S2 to S3. All six seed/population combinations satisfy the bound, with zero violations and negative maximum bound excess. The empirical class-level MAE ranges from 0.0340 to 0.0429, and maximum observed errors range from 0.209 to 0.295. Mean class bounds are approximately 0.685–0.692, with maxima 0.974–1.042. The bound identifies the dependence on class dispersion and directional coherence, but is conservative in this experiment. The uniform Gaussian scores lie in $[0, 1]$, whereas the global 2ϵ from Eq. (17) exceeds the largest possible score margin. Consequently, the global error bound does not certify any pairwise ranking here. The observed ranking improvement is an empirical result, not a consequence of a tight error guarantee.

Appendix A.9. Complete local baseline matrix

Tables A.14 and A.15 expose the complete local baseline matrix used to locate the static controls. Each entry is mean AUROC/FPR@95 under the same OpenOOD loaders and checkpoints. ImageNet-1K KNN and ViM were not evaluated and are marked with a dash. The best AUROC and FPR@95 in each task column are boldfaced independently.

Table A.14: Complete Near-OOD local comparison including CTM, CADRef, and LogitGap. Each entry is AUROC/FPR@95 under the unified OpenOOD v1.5 protocol. FPR@95 is the ID rejection rate at 95% OOD recall (OOD positive).

Method	CIFAR-10	CIFAR-100	IN-200	IN-1K
MSP	88.03/48.17	80.27/54.80	83.34/54.84	76.02/65.67
Energy	87.58/61.33	80.91/55.60	82.50/60.22	75.89/68.57
MaxLogit	87.52/61.32	81.05/55.46	82.90/59.75	76.46/67.85
GEN	88.20/53.67	81.31/54.41	83.71/55.11	76.85/65.30
ReAct	87.11/63.55	80.77/56.39	81.87/62.46	77.38/66.74
Scale	82.55/80.46	80.99/55.68	84.84/57.30	**81.36**/**59.77**
KNN	**90.64**/**34.00**	80.18/61.23	81.57/60.18	–
ViM	88.68/44.85	74.98/62.63	78.68/59.19	–
ASH	75.27/86.79	78.19/65.72	82.38/64.88	78.17/63.33
DICE	78.40/70.12	79.38/57.95	81.78/61.87	73.07/72.45
SHE	81.54/79.64	78.95/59.06	80.18/66.79	73.78/73.03
RMDS	86.20/47.20	80.15/55.47	82.57/54.02	76.54/65.24
RankFeat	79.46/60.89	61.88/80.58	56.92/92.04	50.99/91.83
CADRef	88.97/39.46	79.63/57.18	**85.57**/**52.97**	79.62/61.57
LogitGap	88.72/52.34	**81.37**/**54.36**	83.37/57.63	76.70/66.47
Centroid-only	89.35/38.16	80.12/59.07	82.57/60.11	74.34/69.41
Centroid–MSP	89.27/38.15	80.35/58.75	83.33/58.62	76.29/67.93
CTM	89.49/37.29	80.16/59.18	83.13/58.21	75.69/67.86

Table A.15: Complete Far-OOD local comparison including CTM, CADRef, and LogitGap. Each entry is AUROC/FPR@95 under the unified OpenOOD v1.5 protocol. FPR@95 is the ID rejection rate at 95% OOD recall (OOD positive).

Method	CIFAR-10	CIFAR-100	IN-200	IN-1K
MSP	90.73/31.72	77.76/58.71	90.13/35.44	85.23/51.46
Energy	91.21/41.70	79.77/56.58	90.86/34.85	89.47/38.38
MaxLogit	91.10/41.67	79.67/56.73	91.11/34.03	89.58/38.21
GEN	91.35/34.73	79.68/56.71	91.36/32.11	89.77/35.62
ReAct	90.42/44.90	80.39/54.21	92.31/28.50	93.67/26.31
Scale	86.39/67.54	81.42/54.08	93.98/26.47	**96.53**/**16.53**
KNN	92.96/**24.28**	82.40/53.64	93.16/27.27	–
ViM	**93.48**/25.06	81.70/50.73	91.26/27.20	–
ASH	78.49/79.03	80.58/59.20	93.90/27.30	95.74/19.53
DICE	84.29/51.71	80.01/56.25	90.80/36.51	90.95/41.82
SHE	85.32/66.48	76.92/64.12	89.81/42.16	90.92/41.43
RMDS	91.30/29.41	82.92/52.82	88.06/32.45	86.95/39.65
RankFeat	75.87/57.44	67.10/69.45	38.22/97.72	53.93/87.15
CADRef	91.51/28.19	**83.49**/**50.33**	**95.50**/**19.58**	96.27/16.85
LogitGap	91.88/33.41	79.74/55.95	91.24/32.17	89.69/36.02
Centroid-only	91.62/28.24	82.52/52.43	94.18/23.72	94.50/24.54
Centroid–MSP	91.58/28.18	81.70/52.88	94.02/24.17	94.85/23.60
CTM	91.68/27.76	82.62/52.50	93.89/24.52	95.20/21.68

Table A.16: Recent and centroid references on ImageNet-1K/DenseNet-121. CADRef, LogitGap, and CTM are single fixed-model estimates; centroid-only and fixed fusion show mean±standard deviation over three sampled setup banks on the same model. Boldface compares the displayed methods. FPR@95 uses OOD as the positive class, as defined in Section 6.

Group	Method	AUROC	FPR@95
Near	CADRef	**76.13**	**66.60**
	LogitGap	75.49	67.42
	CTM	72.91	71.61
	Centroid-only	71.18 ± 0.35	73.29 ± 0.56
	Centroid–MSP	73.88 ± 0.31	70.80 ± 0.88
Far	CADRef	**95.32**	**21.46**
	LogitGap	89.14	39.69
	CTM	94.02	25.88
	Centroid-only	93.01 ± 0.53	29.59 ± 1.63
	Centroid–MSP	93.95 ± 0.43	26.97 ± 1.46

Table A.17: Additional CTM precision–recall metrics on ImageNet-1K/DenseNet-121. AUROC and FPR@95 are reported in Table A.20. The ID accuracy is 74.86%.

OOD dataset	AUPR-IN	AUPR-OUT
SSB-hard	63.73	64.98
NINCO	96.83	36.02
iNaturalist	98.65	82.65
Textures	99.67	87.78
OpenImage-O	95.99	76.59

Appendix A.10. Complete-coverage ranking

For each ID/group pair, all available methods are ranked by AUROC and FPR@95, including KNN and ViM on the six tasks they cover. Ties receive the minimum rank. Table A.18 reports mean ranks for the 16 scoring configurations with complete eight-task

coverage, including centroid-only. Ranks use the two-decimal means displayed in the complete baseline matrices. The composite averages the AUROC and FPR@95 ranks.

Table A.18: Mean ranks for scoring configurations with complete eight-task coverage, ordered by AUROC rank. Per-task rankings include all available configurations, including KNN and ViM where evaluated; only complete-coverage configurations are listed. Ranks use the two-decimal means displayed in Tables A.14 and A.15, with minimum ranks for ties. Lower is better.

Configuration	AUROC rank	FPR rank	Composite
CADRef	**4.000**	**3.250**	**3.625**
Centroid–MSP	5.750	6.375	6.063
CTM	6.000	6.000	6.000
Scale	6.000	7.625	6.813
Centroid-only	6.750	7.250	7.000
LogitGap	6.875	7.750	7.313
GEN	7.375	7.750	7.563
ReAct	9.125	10.375	9.750
MaxLogit	10.000	10.375	10.188
RMDS	10.250	6.750	8.500
ASH	10.375	12.375	11.375
Energy	10.625	11.625	11.125
MSP	11.250	9.125	10.188
DICE	13.625	13.500	13.563
SHE	14.625	15.375	15.000
RankFeat	17.250	16.250	16.750

Appendix A.11. Per-dataset OOD detection results

Tables A.19 and A.20 report the named OOD datasets underlying the Near-/Far-OOD aggregates in the main text. Each cell is AUROC/FPR@95. CIFAR-10, CIFAR-100, and ImageNet-200 values are mean±standard deviation over three official model seeds. ImageNet-1K uses a fixed pretrained backbone: MSP and Scale are therefore single runs, while the centroid–MSP baseline reports variation over three independently sampled centroid banks. Within each dataset row, the highest AUROC and the lowest FPR@95 are boldfaced independently.

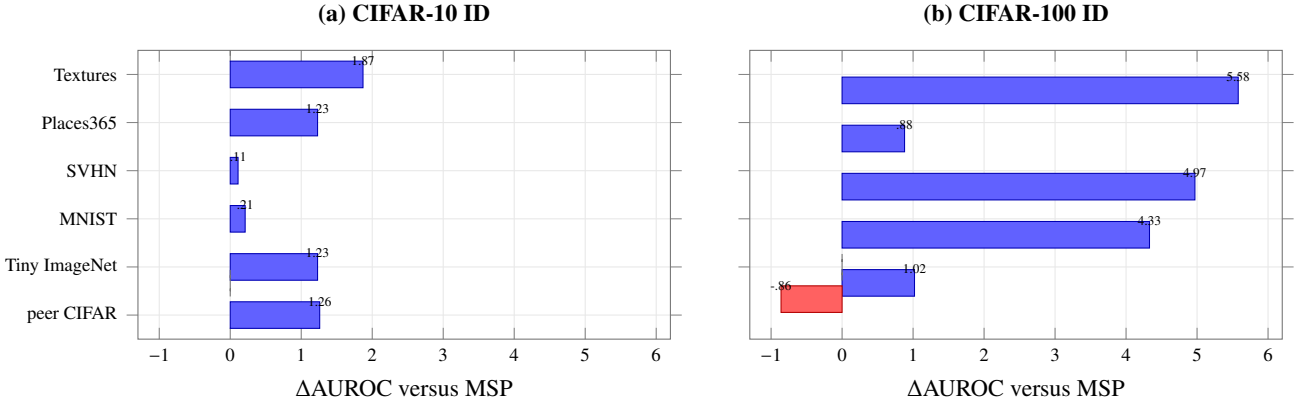

Figure A.6: Per-dataset AUROC change of the centroid–MSP baseline relative to MSP on CIFAR. Blue bars indicate improvement and red indicates degradation. Gains are broad but heterogeneous; peer CIFAR under CIFAR-100 is the only negative case. Exact AUROC/FPR@95 values are reported in Table A.19.

Table A.19: Per-dataset CIFAR results under the unified OpenOOD v1.5 protocol. Each entry is AUROC/FPR@95; all values are mean±standard deviation over three official ResNet-18 seeds.

ID	OOD dataset	Type	MSP	Scale	CTM	Centroid–MSP
CIFAR-10	CIFAR-100	Near	87.19 ± .40/53.09 ± 6.02	81.27 ± .91/81.78 ± 4.87	**88.66 ± .17/40.79 ± .91**	88.45 ± .18/41.96 ± 1.41
	Tiny ImageNet	Near	88.87 ± .24/43.25 ± 3.64	83.84 ± .02/79.14 ± 4.95	**90.32 ± .30/33.79 ± 1.09**	90.10 ± .33/34.33 ± 1.43
	MNIST	Far	92.63 ± 1.92/**23.64 ± 7.13**	90.58 ± 3.62/48.69 ± 18.84	92.82 ± .49/23.84 ± 2.14	**92.84 ± .45**/23.72 ± 1.75
	SVHN	Far	91.46 ± .50/**25.80 ± 2.01**	84.63 ± 3.74/70.54 ± 10.67	91.45 ± .28/26.72 ± 1.17	**91.57 ± .60**/27.05 ± 1.79
	Places365	Far	88.92 ± .57/42.47 ± 4.68	86.41 ± 2.35/70.51 ± 7.32	**90.52 ± .30/33.46 ± .70**	90.15 ± .32/34.77 ± .65
	Textures	Far	89.89 ± .87/34.97 ± 5.71	83.93 ± .42/80.40 ± 4.80	**91.93 ± .26/27.03 ± .33**	91.76 ± .28/27.16 ± .63
CIFAR-100	CIFAR-10	Near	78.47 ± .08/**58.90 ± 1.14**	**79.26 ± .12**/59.11 ± 1.03	77.25 ± .21/68.36 ± .35	77.61 ± .14/67.82 ± .71
	Tiny ImageNet	Near	82.07 ± .20/50.69 ± .43	82.71 ± .18/52.25 ± .89	83.08 ± .25/50.00 ± .29	**83.09 ± .23/49.68 ± .16**
	MNIST	Far	76.08 ± 2.28/57.24 ± 5.72	80.27 ± 1.42/51.64 ± 4.37	**81.32 ± 2.20/50.19 ± 5.15**	80.41 ± 2.07/50.47 ± 4.80
	SVHN	Far	78.42 ± 1.09/59.08 ± 3.10	84.45 ± 1.33/**49.27 ± 3.18**	**84.56 ± 1.59**/49.97 ± 4.04	83.39 ± .85/50.43 ± 2.68
	Places365	Far	79.23 ± .35/**56.62 ± 1.06**	**80.47 ± .28**/56.98 ± .83	80.43 ± .50/58.80 ± 1.04	80.11 ± .37/59.00 ± 1.13
	Textures	Far	77.32 ± .86/61.90 ± 1.59	80.50 ± .86/58.44 ± 2.41	**84.17 ± .66/51.05 ± 2.28**	82.90 ± .65/51.63 ± 1.91

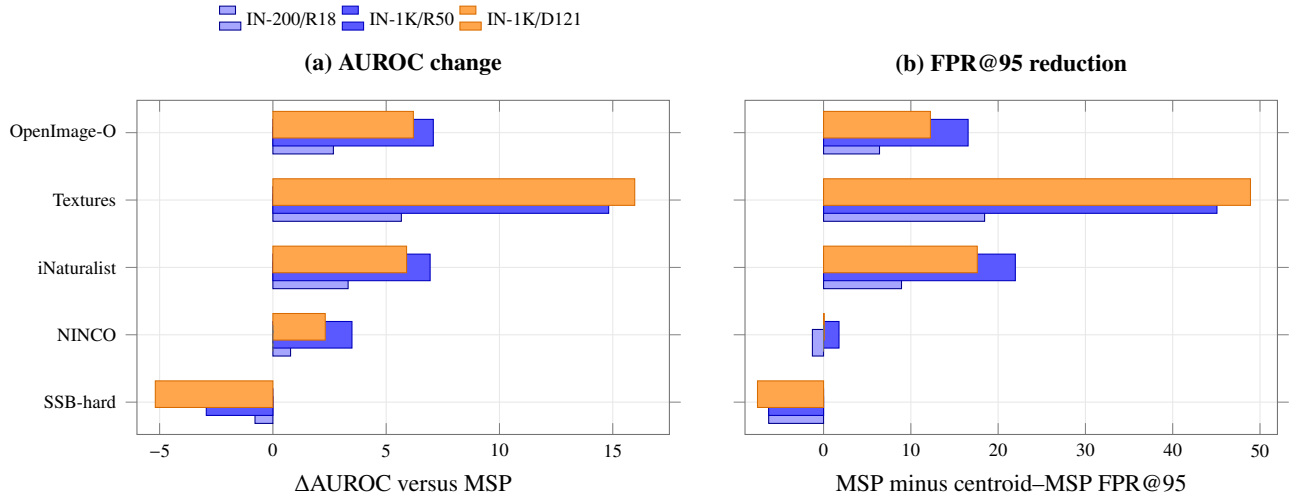

Figure A.7: Per-dataset change of the centroid–MSP baseline relative to MSP on the ImageNet benchmarks. Positive values are favorable in both panels. Grouping the three backbone/scale settings within each dataset exposes the shared pattern: large Far-OOD gains, especially on Textures, and a consistent loss on SSB-hard. Exact values appear in Table A.20.

Table A.20: Per-dataset ImageNet results. Each entry is AUROC/FPR@95. ImageNet-200 uses three official ResNet-18 seeds. ImageNet-1K MSP, Scale, and CTM use one fixed backbone, whereas Centroid–MSP varies three centroid banks.

ID/backbone	OOD dataset	Type	MSP	Scale	CTM	Centroid–MSP
IN-200/R18	SSB-hard	Near	80.38 ± .04/**66.00 ± .14**	**82.08 ± .33**/67.39 ± .78	79.43 ± .25/71.32 ± .29	79.59 ± .24/72.29 ± .36
	NINCO	Near	86.29 ± .13/**43.67 ± .93**	**87.59 ± .36**/47.21 ± 1.55	86.83 ± .17/45.10 ± 1.31	87.07 ± .16/44.95 ± 1.35
	iNaturalist	Far	92.80 ± .30/26.49 ± .91	95.79 ± .56/18.41 ± 2.30	95.90 ± .18/18.70 ± .99	**96.12 ± .33/17.56 ± .68**
	Textures	Far	88.36 ± .16/44.58 ± .86	94.11 ± .08/29.80 ± .64	**94.42 ± .07/25.13 ± .85**	94.03 ± .20/26.13 ± .87
	OpenImage-O	Far	89.24 ± .02/35.24 ± .27	**92.04 ± .30**/31.20 ± 1.32	91.34 ± .12/29.74 ± .70	91.91 ± .06/**28.83 ± .27**
IN-1K/R50	SSB-hard	Near	72.09/74.48	**77.35/67.73**	68.53/80.07	69.15 ± .20/80.77 ± .30
	NINCO	Near	79.95/56.87	**85.37/51.81**	82.85/55.65	83.44 ± .25/55.10 ± .89
	iNaturalist	Far	88.41/43.35	**98.02/9.50**	95.83/19.33	95.35 ± .56/21.38 ± 2.36
	Textures	Far	82.43/60.90	97.63/**11.90**	**97.75**/12.71	97.25 ± .09/15.83 ± .78
	OpenImage-O	Far	84.86/50.14	**93.95/28.18**	92.01/33.00	91.94 ± .46/33.58 ± 1.34
IN-1K/D121	SSB-hard	Near	71.11/75.52	**72.75/75.26**	64.95/82.91	65.92 ± .24/83.09 ± .26
	NINCO	Near	79.53/58.58	**83.56/55.78**	80.86/60.31	81.84 ± .39/58.51 ± 1.51
	iNaturalist	Far	88.61/42.78	**97.25/13.94**	94.80/23.09	94.51 ± .77/25.17 ± 3.08
	Textures	Far	81.28/64.32	96.17/19.47	**97.49/13.44**	97.25 ± .12/15.41 ± .93
	OpenImage-O	Far	83.88/52.59	**92.74/33.02**	89.78/41.10	90.08 ± .58/40.35 ± 1.73

Appendix A.12. Validation mechanism, sample efficiency, and calibration budget

The following tables and plots give the numerical detail for Section 6.5. Pointwise best weights maximize mean AUROC on the designated validation holdout retrospectively; the evaluated rule remains fixed at $\alpha = 0.8$.

Table A.21: Exact values for the controlled nested class-count validation study. α^* maximizes mean holdout AUROC at each class count. The two Δ columns evaluate the single pre-test choice $\alpha = 0.8$ against MSP ($\alpha = 0$) and centroid-only ($\alpha = 1$); positive values favor the fixed mixture. Results use one nested class sequence.

Classes	Pointwise α^*	Δ vs. MSP	Δ vs. centroid	$0.8 = \alpha^*$
10	1.0	+0.64	−1.00	–
25	0.9	+2.39	−0.14	–
50	**0.8**	+2.37	+0.33	✓
100	**0.8**	+2.76	+0.67	✓
200	**0.8**	+4.74	+0.45	✓
500	**0.8**	+5.45	+0.69	✓
1000	0.7	+6.45	+0.86	–

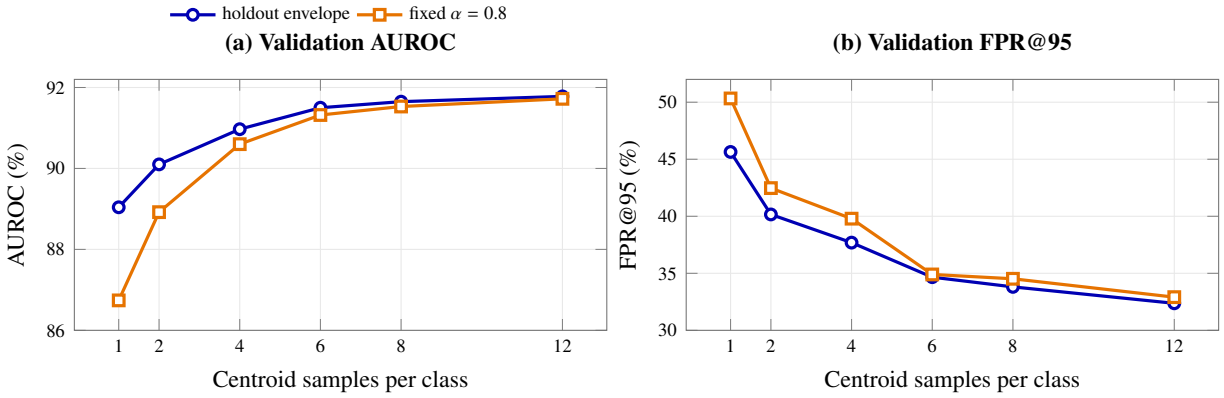

Figure A.8: ImageNet-1K validation-only sample efficiency over three centroid seeds. Six samples per class capture most of the 12-sample performance, and the fixed $\alpha = 0.8$ curve nearly meets the retrospective holdout envelope from six samples onward. The blue curve maximizes mean holdout AUROC over α separately at each bank size; panel (b) uses the corresponding weights.

Table A.22: ImageNet-1K validation-only calibration-budget sensitivity of the centroid–MSP baseline. Columns vary the number of ID calibration images per class; each entry is AUROC/FPR@95. One ID image per class is held out for evaluation, calibration subsets are nested within the other four, and all OOD validation examples are evaluated. The centroid bank is fixed. Positive affine standardization leaves either component's rankings unchanged.

Method	1 image	2 images	3 images	4 images
Fixed $\alpha = 0.8$	92.185/31.17	92.184/31.13	92.185/31.13	92.183/31.13

Table A.23: Validation-holdout mechanism audit on the five natural dataset/backbone targets. "Corr." is Pearson correlation between MSP and centroid scores on the designated OOD validation holdout. Score cells are AUROC/FPR@95; values aggregate three centroid seeds where applicable.

Target	Classes	OOD corr.	Best α	MSP	Fixed $\alpha = 0.8$	Centroid-only
CIFAR-10/R18	10	.880 ± .002	1.0	88.65/42.73	89.72/36.40	89.80/36.53
CIFAR-100/R18	100	.828 ± .003	1.0	83.62/46.53	84.99/45.07	85.00/44.67
IN-200/R18	200	.677 ± .016	1.0	87.81/38.33	90.75/30.94	90.77/31.44
IN-1K/R50	1000	.410 ± .015	.8	85.27/47.13	**91.72/32.91**	90.86/36.89
IN-1K/D121	1000	.315 ± .003	.6	83.77/51.90	**90.31/39.16**	88.81/42.57

Recomputing the 12-sample centroids reproduces the archived checkpoint within 4.47×10^{-8} maximum absolute error for each of the three seeds.

Appendix A.13. Extended batch-one efficiency comparison

Table A.24 reports the complete batch-one latency ratios underlying the abbreviated discussion in Section 6. All measurements use the same RTX 4090 D, include the backbone forward pass, exclude data loading, and use MSP measured in the corresponding timing session as ratio one. CTM and the centroid–MSP baseline are measured in one paired runner on identical preloaded inputs.

Appendix A.14. Implementation details

For the small-bank static control, normalized feature directions are averaged within each class and renormalized to obtain μ_c. The query statistic is $G(x) = \max_c q(x)^\top \mu_c$. If $m(x)$ is MSP and $\widetilde{m}, \widetilde{G}$ denote pooled ID-standardized components, the reported fixed control is $(1 - \alpha)\widetilde{m} + \alpha\widetilde{G}$ with $\alpha = 0.8$.

Table A.24: End-to-end batch-one latency relative to MSP. Lower is better. Per-image latency measures online inference; setup-sample counts for CTM and the centroid–MSP baseline are described in the accompanying resource audit.

Method	CIFAR-10	CIFAR-100	IN-200	IN-1K
CTM	**0.994**	1.043	1.037	1.016
MSP	1.000	**1.000**	1.000	**1.000**
Centroid–MSP	1.041	1.085	1.090	1.039
SHE	1.045	1.180	**0.776**	1.069
RMDS	1.137	1.146	1.103	1.509
Scale	1.141	1.122	1.158	1.516
DICE	1.146	1.120	1.138	1.670
ASH	1.149	1.115	1.207	1.735
RankFeat	3.073	3.045	4.741	4.505

Training banks and candidate classes.. All feature directions are normalized before anchor or centroid construction. The original trajectory settings below are distinct from the mechanism audits. For the original CIFAR trajectories, 80 anchors per class are sampled from the extracted training features using a seeded random permutation. All classes participate in both bandwidth fitting and inference. For ImageNet, the setup bank contains 12 training examples per class, and five are selected by a seeded random permutation as field anchors. The normalized mean of all 12 feature directions defines the class centroid used for candidate selection. At inference, the 20 classes with the largest cosine similarities to the initial query are retained for both force computation and trajectory scoring. This candidate set remains fixed as the query moves. The centroid–MSP baseline uses all class centroids, without candidate truncation.

The score-aligned mechanism audit instead uses five anchors from 12 setup features per class on every benchmark, fits on that setup bank, and uses $\min(C, 20)$ initial candidates. The stability probe uses normalized means of all 12-feature setup banks on every benchmark, compares all class centroids, and does not evaluate an anchor kernel during the probe.

Bandwidth fitting.. The reported trajectory grids, anchor-weighting study, and five-kernel comparison fit bandwidths using zero-step affinities. With the backbone, anchors, and anchor weights fixed, the objective is

$$\mathcal{L}(\sigma) = -\frac{1}{N}\sum_{i=1}^{N} \log \frac{\exp\{KA_{y_i}(q_i)\}}{\sum_{c\in\mathcal{J}_i^{\text{fit}}} \exp\{KA_c(q_i)\}}, \tag{A.6}$$

where (q_i, y_i) are ID fitting features and labels. On CIFAR, $\mathcal{J}_i^{\text{fit}}$ contains every class and fitting uses the extracted training features. On ImageNet, fitting uses the setup bank and $\mathcal{J}_i^{\text{fit}}$ contains the true class and the 19 other classes with highest centroid cosine. The true class is included only because labels are available during fitting; inference uses the 20 highest-scoring classes without this intervention.

The factor K matches the implementation's summed-anchor affinities. It does not change the ordering of a fixed field, but it matters in the fitting loss and force magnitude. We optimize log-bandwidths with Adam at an initial learning rate of 10^{-3} and cosine learning-rate decay. Bandwidths are initialized to 1.0 on CIFAR and 0.5 on ImageNet and constrained to $[0.05, 4.0]$. CIFAR uses 25 epochs with batches of 64; ImageNet uses 10 epochs with batches of 32. No trajectory is differentiated through in this fitting objective, even when inference uses a nonzero step count. The static centroid–MSP baseline requires neither bandwidth fitting nor integration.

Score calibration and fusion.. The aligned audit uses a seeded floor-half split within each ID class and a separate floor-half OOD split for tuning, with the complements held out. Current-winner $T \in \{1, 3, 5, 10\}$ is selected by mean ImageNet-1K tuning AUROC over three banks, breaking ties by lower FPR@95 and then smaller T. The selected $T = 5$ is transferred without retuning. The no-switch audit uses the full validation caches, rather than only the evaluation holdouts. Per seed, the ImageNet-1K selection and holdout counts are respectively 2000/881 and 3000/882 (ID/OOD). Target holdouts contain 500/500 for CIFAR-10, 500/486 for CIFAR-100, and 600/882 for ImageNet-200; no target-side selection is performed for this study.

The probe uses a separate partition with calibration seed zero: two ID examples per class for fit, one for selection, and the remainder for the holdout. A separate OOD permutation has split endpoints $\lfloor .35N_{\text{out}} \rfloor$ and $\lfloor .60N_{\text{out}} \rfloor$; OOD-fit examples are not used for calibration. On ResNet-50 source selection data, $\eta \in \{0, .01, .02, .04, .08, .16\}$ is chosen by mean AUROC over three banks, breaking ties by smaller η. The selected .04 is fixed for all other settings; target selection subsets are unused. Probe statistics and AUROC changes are measured on the same holdout, so their agreement is descriptive rather than an independent prediction test. Per seed, ImageNet-1K uses 2000 ID-fit examples, 1000/440 selection examples, and 2000/706 holdout examples (ID/OOD), for both backbones; only ResNet-50 is used to select η. CIFAR-10, CIFAR-100, and ImageNet-200 use 20, 200, and 400 ID-fit examples, respectively, with holdout counts of 970/400, 700/389, and 400/706. ID-fit examples standardize the probe scores; training-feature banks are separate.

For each benchmark and backbone, the two component scores are standardized using the same deterministic, class-stratified subset of ID validation data. The stored statistics are the pooled means and population standard deviations; inference denominators are floored at 10^{-12}. Selecting the maximum-cosine rule and $\alpha = 0.8$ uses an ImageNet-1K ID/OOD tuning split for the static baseline, with α scanned on $\{0, .05, \ldots, 1\}$ and only the ID tuning part used for standardization. The rule is checked on the complementary holdout. The probe inherits this fixed α; its separate selection split is used to choose η, as described above. These choices remain fixed during transfer, while the ID calibration statistics are recomputed for each target. In the static bridge, S1 already uses all classes; S2 replaces the bandwidths by their per-bank median and uses uniform anchor weights, S3 normalizes each anchor mean to obtain a radial centroid, and S4 applies the equivalent maximum-cosine score. Thus scale sharing and candidate expansion are explicit modeling changes, not consequences of setting $T = 0$.

Appendix A.15. Reproducibility notes

The experiments were executed with Python 3.10.8, PyTorch 2.1.2, torchvision 0.16.2, CUDA 12.1, and an NVIDIA RTX 4090 D. The archived evaluation code records OpenOOD commit `3c35632ee91b54b09d1f085d04f94744cece7d0b`; the CTM completion metadata records official reference commit `3587259bd6a69abd6b4103cb7311ffaa0857d60f`. Dataset preparation validates archive integrity, image-list counts, and missing paths. The ImageNet-1K ResNet-50 file has SHA256 `06 76ba61b6795bbe1773cffd859882e5e297624d384b6993f7c9e683e722fb8a`; the pinned DenseNet-121 file has SHA256 `42bd2c384fae0b346f977fe9feed1bcfdbd6ea4727478c6c76177d725b218e89`. The extended audit covers 50 full-dataset ASH/DICE/SHE/RMDS/RankFeat runs. RankFeat uses full exact SVD, and RMDS streaming statistics match direct tensor computation within 2.84×10^{-14} in float64. The CADRef/LogitGap audit adds 22 full-dataset method/model runs: 11 runs per method across the five dataset/backbone targets. The two fixed ImageNet-1K backbones are not counted as independent training seeds. CADRef processes all ID training images for setup, while fixed LogitGap is training-free; neither uses OOD data for setup or parameter selection. The accompanying CADRef/LogitGap export includes raw Near-/Far-OOD rows, mean and standard-deviation summaries, and completion metadata. The expanded rank table is regenerated from the displayed two-decimal baseline means so that its tie convention is reproducible from the paper itself. The score-aligned mechanism study stores its ImageNet-1K selection gate, locked cross-benchmark transfer gate, pair-decomposition rows, and all stepwise winner traces. The transfer gate has SHA256 `1199cf9e330f4e2f623ce7409193306bfc38672daa5bfb1d6c2471624563ec2b`; the winner-switch explanation gate has SHA256 `f75a83c4a94edd36713ca5e505a01f8a4df28fafd6ae0bf210a8327bd747d9a2`. Both studies are validation-only and record that no Near-/Far-OOD test loader was accessed. The boundary-conditioned stability-probe transfer uses the previously saved validation caches and formal detectors only. Its locked gate has SHA256 `c824da104d547c804166 82c39e12b55041af190b0312534bd19787f35b06a6bb` and rejects benchmark-test access. The accompanying result archive has SHA256 `a2ed6b0dfed5310b343dec71dc7dd5a3d7538d45a3e4100cd77b2d954db34401`. The complete result matrix, per-run configurations, completion markers, and summary CSV files are retained in the accompanying reproducibility archive. The main text and supplementary tables report the formal full-dataset aggregate experiment families; smoke tests, exploratory score-rule candidates, and duplicate per-seed raw rows remain available in the audit archive.

References

Arnold, V. I., Vogtmann, K., and Weinstein, A. (1989). *Mathematical methods of classical mechanics*, volume 60. Springer.

Cayton, L. (2008). Algorithms for manifold learning. Technical report, Department of Computer Science & Engineering, UC San Diego. https://escholarship.org/uc/item/8969r8tc.

Djurisic, A., Bozanic, N., Ashok, A., and Liu, R. (2023). Extremely simple activation shaping for out-of-distribution detection. In *The Eleventh International Conference on Learning Representations*. https://openreview.net/forum?id=ndYXTEL6cZz.

ElAraby, M., Sahoo, S., Pequignot, Y., Novello, P., and Paull, L. (2025). GROOD: Gradient-aware out-of-distribution detection. *Transactions on Machine Learning Research*.

Fefferman, C., Mitter, S., and Narayanan, H. (2016). Testing the manifold hypothesis. *Journal of the American Mathematical Society*, 29(4):983–1049. https://doi.org/10.1090/jams/852.

Gu, J., Qiao, Z., and Li, Z. (2025). Gradient short-circuit: Efficient out-of-distribution detection via feature intervention. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 457–466.

Hendrycks, D., Basart, S., Mazeika, M., Zou, A., Kwon, J., Mostajabi, M., Steinhardt, J., and Song, D. (2022). Scaling out-of-distribution detection for real-world settings. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 8759–8773. PMLR. https://proceedings.mlr.press/v162/hendrycks22a.html.

Hendrycks, D. and Gimpel, K. (2017). A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *International Conference on Learning Representations*. https://openreview.net/forum?id=Hkg4TI9xl.

Isaac-Medina, B. K. S., Che, M., Gaus, Y. F. A., Akcay, S., and Breckon, T. P. (2025). FEVER-OOD: Free energy vulnerability elimination for robust out-of-distribution detection. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 4529–4538.

Ke, A., Chen, W., Feng, C., Cao, Y., Xie, X., Zhou, S. K., and Feng, L. (2025). Prototype-based optimal transport for out-of-distribution detection. In Kwok, J., editor, *Proceedings of the Thirty-Fourth International Joint Conference on Artificial Intelligence, IJCAI-25*, pages 3018–3026. International Joint Conferences on Artificial Intelligence Organization. Main Track.

Lee, K., Lee, K., Lee, H., and Shin, J. (2018). A simple unified framework for detecting out-of-distribution samples and adversarial attacks. In Bengio, S., Wallach, H., Larochelle, H., Grauman, K., Cesa-Bianchi, N., and Garnett, R., editors, *Advances in Neural Information Processing Systems*, volume 31. Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2018/file/abdeb6f575ac5c6676b747bca8d09cc2-Paper.pdf.

Liang, J., Hou, R., Hu, M., Chang, H., Shan, S., and Chen, X. (2025). Revisiting logit distributions for reliable out-of-distribution detection. In *Advances in Neural Information Processing Systems*, volume 38.

Liang, S., Li, Y., and Srikant, R. (2017). Principled detection of out-of-distribution examples in neural networks. *CoRR*, abs/1706.02690. http://arxiv.org/abs/1706.02690.

Ling, Z., Chang, Y., Zhao, H., Zhao, X., Chow, K., and Deng, S. (2025). CADRef: Robust out-of-distribution detection via class-aware decoupled relative feature leveraging. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 4968–4977.

Liu, W., Wang, X., Owens, J., and Li, Y. (2020). Energy-based out-of-distribution detection. In Larochelle, H., Ranzato, M., Hadsell, R., Balcan, M., and Lin, H., editors, *Advances in Neural Information Processing Systems*, volume 33, pages 21464–21475. Curran Associates, Inc. https://proceedings.neurips.cc/paper_files/paper/2020/file/f5496252609c43eb8a3d147ab9b9c006-Paper.pdf.

Liu, X., Lochman, Y., and Zach, C. (2023). Gen: Pushing the limits of softmax-based out-of-distribution detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 23946–23955. https://openaccess.thecvf.com/content/CVPR2023/html/Liu_GEN_Pushing_the_Limits_of_Softmax-Based_Out-of-Distribution_Detection_CVPR_2023_paper.html.

Nguyen, A., Bertrand, A., Le Hégarat-Mascle, S., Aldea, E., El-Korso, M. N., Lustrat, R., and Florin, F. (2026). Fisher–rao sensitivity for out-of-distribution detection in deep neural networks. In *The Fourteenth International Conference on Learning Representations*.

Nguyen, N.-H., Nguyen, H.-Q., Ta, T.-A., Nguyen-Tang, T., Doan, K. D., and Hoang, T.-T. (2023). A cosine similarity-based method for out-of-distribution detection. In *ICML 2023 Workshop on Spurious Correlations, Invariance, and Stability*. https://openreview.net/forum?id=QuXApqb8Yl.

Park, J., Jung, Y. G., and Teoh, A. B. J. (2023). Nearest neighbor guidance for out-of-distribution detection. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 1686–1695.

Ren, J., Fort, S., Liu, J. Z., Roy, A. G., Padhy, S., and Lakshminarayanan, B. (2021). A simple fix to mahalanobis distance for improving near-ood detection. *CoRR*, abs/2106.09022. https://arxiv.org/abs/2106.09022.

Seleznova, M., Chou, H.-H., Mayrink Verdun, C., and Kutyniok, G. (2026). GradPCA: Leveraging NTK alignment for reliable out-of-distribution detection. In *The Fourteenth International Conference on Learning Representations*.

Song, Y., Sebe, N., and Wang, W. (2022). RankFeat: Rank-1 feature removal for out-of-distribution detection. In *Advances in Neural Information Processing Systems*, volume 35, pages 17885–17898. https://proceedings.neurips.cc/paper_files/paper/2022/hash/71c9eb0913e6c7fda3afd69c914b1a0c-Abstract-Conference.html.

Sun, Y., Guo, C., and Li, Y. (2021). React: Out-of-distribution detection with rectified activations. In *Advances in Neural Information Processing Systems*, volume 34. https://proceedings.neurips.cc/paper/2021/hash/01894d6f048493d2cacde3c579c315a3-Abstract.html.

Sun, Y. and Li, Y. (2022). DICE: Leveraging sparsification for out-of-distribution detection. In *Computer Vision – ECCV 2022*, pages 691–708. https://www.ecva.net/papers/eccv_2022/papers_ECCV/papers/136840680.pdf.

Sun, Y., Ming, Y., Zhu, X., and Li, Y. (2022). Out-of-distribution detection with deep nearest neighbors. In Chaudhuri, K., Jegelka, S., Song, L., Szepesvari, C., Niu, G., and Sabato, S., editors, *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 20827–20840. PMLR. https://proceedings.mlr.press/v162/sun22d.html.

Wang, H., Li, Z., Feng, L., and Zhang, W. (2022). Vim: Out-of-distribution with virtual-logit matching. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4921–4930. https://openaccess.thecvf.com/content/CVPR2022/html/Wang_ViM_Out-of-Distribution_With_Virtual-Logit_Matching_CVPR_2022_paper.html.

Wei, T. and Liu, M. (2026). Uncertainty-aware prototype refinement for out-of-distribution detection. *Neurocomputing*, 699:134310.

Xia, G. and Bouganis, C.-S. (2024). Augmenting the softmax with additional confidence scores for improved selective classification with out-of-distribution data. *International Journal of Computer Vision*, 132(9):3714–3752.

Xu, K., Chen, R., Franchi, G., and Yao, A. (2024). Scaling for training time and post-hoc out-of-distribution detection enhancement. In Kim, B., Yue, Y., Chaudhuri, S., Fragkiadaki, K., Khan, M., and Sun, Y., editors, *International Conference on Learning Representations*, volume 2024, pages 16316–16329. https://proceedings.iclr.cc/paper_files/paper/2024/file/46d943bc6a15a57c923829efc0db7c7a-Paper-Conference.pdf.

Zhang, J., Fu, Q., Chen, X., Du, L., Li, Z., Wang, G., Liu, X., Han, S., and Zhang, D. (2023a). Out-of-distribution detection based on in-distribution data patterns memorization with modern hopfield energy. In *The Eleventh International Conference on Learning Representations*. https://openreview.net/forum?id=KkazG4lgKL.

Zhang, J., Yang, J., Wang, P., Wang, H., Lin, Y., Zhang, H., Sun, Y., Du, X., Li, Y., Liu, Z., Chen, Y., and Li, H. (2023b). Openood v1.5: Enhanced benchmark for out-of-distribution detection. *arXiv preprint arXiv:2306.09301*. https://arxiv.org/abs/2306.09301.

Zohrabi, R., Hasani, H., Baghshah, M., Rohrbach, A., Rohrbach, M., and Rohban, M. H. (2025). Spurious-aware prototype refinement for reliable out-of-distribution detection. In *Advances in Neural Information Processing Systems*, volume 38.