FlyHash for LLM Agent Memory: An Empirical Validation and a New Application

Mohan Bheemadolu*[1] and Chrisogonas O. Odhiambo[2]

[1]Independent Researcher. Email: mohan.bmdl@gmail.com
[2]Email: chrisogonas@gmail.com

arXiv:2611.05791v1 [cs.AI] 4 Nov 2026

Abstract

Large language model (LLM) agents increasingly rely on external memory stores to retain information across conversations, typically via dense vector embeddings searched with approximate nearest-neighbor methods. The fruit fly's olfactory circuit computes a locality-sensitive hash of sensory input using a strikingly different mechanism — sparse random projection followed by winner-take-all competition — and prior work (Dasgupta et al., 2017) showed this scheme, dubbed FlyHash, preserves neighbor structure better per bit than conventional random-hyperplane hashing (SimHash) on several benchmarks. We apply FlyHash, to our knowledge for the first time, to LLM agent memory retrieval, and ask three questions a practitioner actually needs answered: does the quality advantage hold under a correctly matched-bit protocol (it does, robustly, across 10 seeds and four datasets spanning two modalities, with one specific, previously undocumented comparison pitfall identified and corrected along the way); does that advantage translate into real memory and latency savings, including against the dense-vector indexing (HNSW) practitioners default to today (no: at equal bits FlyHash's codes are larger and slower to compare than SimHash's, and HNSW dominates both hash methods at ordinary scale; a sparse-set comparison we introduce recovers most but not all of a naive latency gap, leaving hashing's practical case narrower and more specific than the quality result alone implies); and does the advantage survive genuine, non-circular agent-memory relevance labels rather than the nearest-neighbor-by- embedding proxy this literature otherwise relies on (it does). We report all three answers precisely, including the two that narrow rather than flatter the headline result, and release our implementation, a working agent-memory demo, and a full account of the methodological pitfalls we found and corrected.

1 Introduction

Large language model agents that operate over long or repeated interactions increasingly need memory: a way to store facts, preferences, or prior context and retrieve the relevant pieces on demand (Chhikara et al., 2025). The dominant approach embeds each memory as a dense vector and retrieves by approximate nearest-neighbor search over an index such as FAISS (Johnson et al., 2021; see also the RAG survey of Zhao et al., 2024), though hybrid retrieval architectures combining dense and lexical signals are an active and growing alternative. This works well but scales in memory and compute with the size and dimensionality of the stored embeddings, motivating interest in more

*Code: https://github.com/mohanji/flyhash-agent-memory
Archive (DOI): https://doi.org/10.5281/zenodo.23903302

compact representations — in particular, hashing schemes that compress an embedding into a short binary code while preserving enough neighbor structure to remain useful for retrieval.

One such scheme, introduced by Dasgupta, Stevens, and Navlakha (2017) and directly inspired by neurobiology, computes a locality-sensitive hash the way the fruit fly's olfactory circuit is believed to encode odors: roughly 50 receptor-neuron activations are randomly and sparsely projected into approximately 2,000 "Kenyon cells," of which only the top ∼5% most active are allowed to "fire," producing a sparse binary code. The original paper showed that this scheme — which we refer to as FlyHash, following subsequent literature — preserves nearest-neighbor structure better, at a matched number of output bits, than SimHash (Charikar, 2002), the standard random-hyperplane hashing baseline, on several benchmark datasets including SIFT, GloVe, and MNIST.

Despite this result and several follow-up papers extending the algorithm (surveyed in Section 2), we are not aware of any prior work applying FlyHash-style hashing specifically to LLM agent memory retrieval — a setting that is, in several respects, a natural fit: agent memories are frequently decomposed into short, numerous, semantically embedded units retrieved by similarity (Hu et al., 2026; Du, 2026; Zeng et al., 2024) — the regime this paper's evaluation targets, though we note that recent work increasingly favors structured or hierarchical retrieval over flat similarity search for longer-horizon tasks (Cao et al., 2026; Hu et al., 2026).

This paper is organized around three research questions, each answered precisely rather than in the aggregate, because — as we show below — the answers are not all the same shape:

- **RQ1.** Under a protocol that correctly holds FlyHash's and SimHash's *active* bit counts equal (Section 3.2), does FlyHash's reported per-bit retrieval-quality advantage over SimHash hold for data representative of LLM agent memory, robustly across seeds and datasets? (Sections 4.1–4.2; answer: yes, with one specific, previously undocumented comparison pitfall that must be avoided to get there — Appendix A.)
- **RQ2.** Does that quality advantage translate into an actual resource advantage — lower memory, lower query latency — in a realistic deployment, including against the dense-vector indexing methods (HNSW) practitioners use by default today? (Sections 4.3–4.4, 4.6; answer: not by default, and not relative to HNSW at ordinary scale — hashing's resource case is narrower than the quality result alone would suggest.)
- **RQ3.** Does the quality advantage hold under genuine, non-circular agent-memory relevance labels, rather than the nearest-neighbor-by-embedding proxy used everywhere else in this literature (including the original FlyHash paper)? (Section 4.7; answer: yes.)

In service of these questions, the paper's contributions are: (1) the first application, to our knowledge, of FlyHash to LLM agent memory retrieval, with a working end-to-end system (Sections 3.5, 4.5); (2) the matched-active-bit protocol itself, stated precisely enough to prevent a specific conflation we discovered by reproducing a discrepancy with a published result (Section 3.2, Appendix A); (3) statistical and cross-dataset validation beyond the original paper — 10 seeds, four datasets across two modalities, a 36-configuration hyperparameter sweep, and a paired statistical analysis (Sections 4.1–4.2); and (4) an honest, multi-angle resource and deployment analysis that does not stop at the quality result — brute-force memory/latency costs, a sparse-set comparison, a direct HNSW systems baseline, and a relevance-label upgrade (Sections 4.3–4.4, 4.6–4.7).

The remainder of this paper is organized as follows. Section 2 situates this work relative to FlyHash's original paper and its algorithmic follow-ups. Section 3 describes the method, evaluation protocol, datasets, and statistical procedure. Section 4 presents results. Section 5 discusses their implications. Section 6 states limitations and future work. Appendix A gives a full account of the reproducibility investigation that motivated the corrected protocol in Section 3.2.

2 Related Work

The original algorithm. Dasgupta, Stevens, and Navlakha (2017) introduced the algorithm this paper builds on, showing that a neurally-plausible circuit — sparse random projection into a much higher-dimensional space, followed by winner-take-all sparsification — functions as a locality-sensitive hash competitive with, and often superior to, dense random-hyperplane hashing (SimHash; Charikar, 2002) at a matched number of output bits. Their evaluation covered SIFT descriptors, GloVe word vectors, and MNIST images. This paper directly reuses their algorithm and their choice of evaluation datasets (GloVe and MNIST; we do not evaluate on SIFT — see Section 6), applying it to a problem domain they did not consider.

DenseFly. Sharma and Navlakha (2018) extended the base algorithm with "DenseFly," which densifies FlyHash's sparse output as a post-processing step to further improve similarity search accuracy. This is a genuine algorithmic improvement over plain FlyHash that we do not implement here; we discuss it as natural future work in Section 6.

BioHash. Ryali, Hopfield, Grinberg, and Krotov (2020) replaced FlyHash's *random* sparse projection with a *learned* one, trained with a biologically-plausible local learning rule, and reported improved retrieval accuracy over the random-projection version across several datasets. Their paper also established mean Average Precision (mAP) as the standard evaluation metric for this line of work, which we adopt in Section 3.2. We note for transparency that BioHash's learned projection is reported to outperform the plain FlyHash evaluated in this paper; our contribution is not a better hash function, but an application and validation of the simpler, random-projection version, which requires no training data or learning procedure — a practically relevant trade-off for a lightweight agent-memory component.

FlyNN. Most recently, work applying the fly olfactory circuit to federated nearest-neighbor classification ("FlyNN," AAAI 2022) illustrates the breadth of problems this biological mechanism has been adapted to beyond pure hashing, extending it to a federated learning setting.

Existing implementations. A general-purpose `FlyHash` package is available on PyPI (Huang, MIT license), providing a NumPy-compatible implementation of the base algorithm for small-to-medium input dimensionalities. Its existence confirms that an implementation of the algorithm itself is not a novel contribution; our implementation exists to support the specific experiments in this paper rather than to serve as a general-purpose library.

Graph-based approximate nearest-neighbor search. A different line of work achieves fast approximate retrieval not by compressing each vector into a short code, but by building a navigable graph over the full-precision (or lightly quantized) vectors and searching it via greedy graph traversal. HNSW (Hierarchical Navigable Small World graphs; Malkov and Yashunin, 2016/2020) is the dominant practical method in this family and underlies most production vector-search systems in use today (e.g., FAISS's `IndexHNSW`, and the default indexes in several vector databases including Milvus, Pinecone, Weaviate, and pgvector; Idrees et al., 2026). HNSW and hashing schemes like FlyHash and SimHash target different points in the retrieval design space: HNSW optimizes for sublinear query time and high recall at the cost of storing richer per-item structure (full or near-full vectors plus graph edges), while FlyHash and SimHash optimize for extreme per-record compactness — a short binary code compared via cheap Hamming distance — at the cost of the brute-force, linear-scan comparison this paper evaluates throughout. Our main results (Sections 4.1–4.4) isolate hashing-function quality by comparing FlyHash against SimHash at matched bit budgets, rather than against HNSW, since the two are different design points (compact-code quality versus indexed graph search); we separately report a direct, explicitly-labeled systems comparison against HNSW in Section 4 for readers asking the practical "which would I actually deploy" question.

LLM memory architectures and retrieval-augmented generation. A large and rapidly

growing body of work designs memory systems for LLM agents directly, distinct from the hashing-algorithm literature surveyed above: production-oriented long-term memory frameworks (Chhikara et al., 2025), structural memory representations spanning chunks, triples, and summaries (Zeng et al., 2024), and recent surveys of agent memory mechanisms and evaluation (Du, 2026). This literature's central recent finding — that flat, independent similarity retrieval over text snippets can suffer from redundancy and relational blindness at longer horizons, motivating structured or hierarchical retrieval instead (Hu et al., 2026; Cao et al., 2026) — is complementary to this paper's contribution: we evaluate hashing methods within the flat similarity-based retrieval regime that remains a widely used baseline (e.g., Mem0-style systems), rather than proposing or evaluating a structured memory architecture; combining FlyHash-style hashing with structured retrieval is natural future work (Section 6).

Vector quantization and alternative hashing families. A separate line of work compresses dense vectors via quantization rather than binary hashing. Product Quantization (Jégou, Douze, & Schmid, 2010) and its optimized variants (Ge et al., 2013) achieve higher retrieval accuracy than binary hashing at comparable memory budgets by preserving continuous distance estimates via asymmetric lookup tables rather than collapsing to Hamming codes; recent methods such as RaBitQ extend this advantage to extreme, 1-bit-per-dimension compression (Gao & Long, 2024). We do not evaluate against this family here, since our focus is a head-to-head comparison between two hashing paradigms under a matched-bit protocol, but it represents a stronger baseline than either hash method at most operating points — consistent with the HNSW result in Section 4.6. Within the hashing family itself, more sophisticated locality-sensitive schemes such as Cross-Polytope LSH achieve a better asymptotic query exponent than random-hyperplane hashing (Laarhoven, 2020), but this advantage only materializes at large scale and high ambient dimension; at the moderate dataset sizes evaluated here (Table 1), hyperplane-based SimHash remains competitive or superior in wall-clock performance due to lower per-hash computational overhead (Cai, 2016; Laarhoven, 2020), motivating our choice of standard SimHash, rather than Cross-Polytope LSH, as the baseline.

Sparse lexical retrieval. A different line of "sparse retrieval" work — BM25, and learned sparse lexical models such as SPLADE (Formal, Lassance, Piwowarski, & Clinchant, 2024) — operates directly on text via sparse term-weight vectors over a vocabulary, which is a distinct notion of sparsity from the sparse *binary codes* FlyHash derives from dense embeddings via winner-take-all competition. These lexical methods are a complementary, not competing, approach, well-suited to exact keyword matching rather than semantic similarity over learned embeddings, and are outside this paper's scope of embedding-based hashing methods.

This paper's position. None of the above apply FlyHash-style hashing to LLM agent memory retrieval, and none report the depth of statistical validation (multi-seed variance, multi-dataset, multi-modality, hyperparameter ablation) presented in Section 4. The gap this paper fills is at the intersection of a new, currently relevant application and a level of empirical rigor — including a documented and corrected methodological pitfall — that the base algorithm's literature does not yet provide.

3 Method

3.1 FlyHash and SimHash

Both hashing schemes take a dense input vector $x \in \mathbb{R}^d$ and produce a binary output code.

SimHash (Charikar, 2002) draws a dense random matrix $W \in \mathbb{R}^{k \times d}$ with entries sampled i.i.d.

from a standard normal distribution, and outputs

$$h_{\text{sim}}(x) = \text{sign}(Wx) \in \{0,1\}^k.$$

Every one of the k output bits is thresholded independently; the code is dense in the sense that each bit is set to 1 with probability approximately 1/2.

FlyHash (Dasgupta et al., 2017) instead draws a *sparse* random projection matrix $P \in \mathbb{R}^{m \times d}$, where each row (each "Kenyon cell") is connected to only a random subset of $\lceil s \cdot d \rceil$ input dimensions (s is the *sparsity*, fixed at 0.10 throughout this paper, following the fly's own connectivity, where each Kenyon cell integrates inputs from roughly 5–12% of input channels; Caron, Ruta, Abbott, & Axel, 2013), with weights drawn uniformly from $\{-1, +1\}$ on those connections and zero elsewhere. The projection Px is then passed through a k-winners-take-all operation: only the k largest of the m projected values are set to 1, all others to 0:

$$h_{\text{fly}}(x) = \text{WTA}_k(Px) \in \{0,1\}^m, \qquad \|h_{\text{fly}}(x)\|_0 = k.$$

We refer to m as the **pool size** (the total number of Kenyon cells, fixed at 2,000 in our experiments unless otherwise stated) and k as the **bit budget** (the number of winning, active bits). This notational distinction — k versus m — is deliberately explicit, because conflating the two is the source of the reproducibility discrepancy described in Section 3.2 and Appendix A.

3.2 Evaluation protocol

Metric. Following BioHash (Ryali et al., 2020), we evaluate retrieval quality with mean Average Precision at R (mAP@R). For a query q, let the *relevant set* be its R true nearest neighbors under exact Euclidean distance on the original (unhashed) embeddings. Average Precision for q is

$$\text{AP}(q) = \frac{1}{R} \sum_{\ell=1}^{R} \text{Precision}(\ell) \cdot \text{Rel}(\ell),$$

where $\text{Rel}(\ell) = 1$ if the item ranked ℓ-th by Hamming distance under the hash code is in the relevant set, and $\text{Precision}(\ell)$ is the fraction of relevant items among the top ℓ ranked results. mAP@R is the mean of $\text{AP}(q)$ over all queries. Unlike plain recall@k, mAP@R rewards ranking true neighbors earlier, not merely including them somewhere in the top R.

Matched-active-bit (quality-per-active-bit) protocol. The central methodological point of this paper concerns what it means to compare FlyHash and SimHash "at the same number of bits." There are at least three distinct, non-interchangeable notions of parity one could hold fixed:

1. **Equal active/output bits** — the number of bits that are actually compared by Hamming distance per query (k for both methods). This is the notion we use throughout Sections 4.1–4.2, and we call it the **matched-active-bit** protocol to avoid the vaguer, overloaded term "fair comparison."

2. **Equal physical storage** — the number of bytes each method's code occupies at rest, which depends on *representation* (dense vs. packed vs. sparse-index), not just k (Sections 4.3–4.4).

3. **Equal query-time/resource budget** — the number of bytes moved or compared per query, which can differ from storage when a method's on-disk and in-memory representations differ (Sections 4.3–4.4, 4.6).

Conflating these is the specific pitfall this paper is built around. Concretely, under the matched-active-bit protocol, for each bit budget k we evaluate FlyHash with pool size m fixed at a large value (2,000, unless stated otherwise) and the winner-take-all fraction set to k/m, and we evaluate SimHash with exactly k output bits — so both methods produce the same number of *active* bits, isolating the quality of the hash *function* from either method's absolute output size. This is not the naive way to run the comparison: a naive implementation might instead fix FlyHash's *total* output width (i.e., set m to whatever "bits" value is being swept) and hold the winner-take-all *percentage* fixed (e.g., 5%, matching the fly's own biology) — in which case FlyHash's *effective* number of active bits at a given nominal "bits" value is only a small fraction of that value, while SimHash is given the full nominal count as genuinely active bits. This conflation of k (notion 1) with m silently and substantially disadvantages FlyHash, and is distinct from the storage/resource conflations addressed later (notions 2 and 3, Sections 4.3–4.4). We discovered the first pitfall by reproducing a discrepancy with a published result; the full diagnostic process, including hypotheses we initially pursued and later ruled out, is given in Appendix A.

3.3 Datasets

We evaluate on four independent real datasets spanning two modalities, summarized in Table 1.

Table 1: Datasets used for the main comparison (Section 4.1).

Dataset	Modality	Records	Dim.	Role
GloVe (glove-wiki-gigaword-100) [9]	Text (word vectors)	5,000	100	Primary dataset; discovery of the fair-comparison pitfall (Sec. 3.2, App. A); statistical (Sec. 3.4) and hyperparameter (Sec. 4.2) validation
MiniLM (all-MiniLM-L6-v2) [10]	Text (sentence embeddings)	100	384	Validates the exact embedding type used in the agent memory demonstration (Sec. 3.5)
scikit-learn `digits` [12]	Image	1,797	64 (8×8)	Genuine, if small, image dataset used where full MNIST was not reachable during development
MNIST (via OpenML) [11]	Image	5,000 (of 70,000)	784 (28×28)	Full-scale image benchmark used in the original FlyHash paper

All embeddings are L2-normalized prior to hashing where applicable (GloVe, MiniLM). MNIST and `digits` pixel intensities are scaled to $[0, 1]$.

Data availability. All four datasets are public and loaded programmatically, and none of them is redistributed in this paper's code release: GloVe via `gensim`'s model downloader (`https://nlp.stanford.edu/projects/glove/`); MiniLM sentence embeddings via the `sentence-transformers` library, which downloads the pretrained checkpoint from `https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2`; scikit-learn's bundled `digits` dataset via `sklearn.datasets.load_digits()`; and MNIST via OpenML (`https://www.openml.org/d/554`), fetched with `sklearn.datasets.fetch_openml()`. Exact loading code is in the repository linked on the title page. The labeled agent-memory benchmark of Section 4.7 (6,000 memories, 150 queries, relevance labels) is released as JSONL files in the same repository (`datasets/agent_memory_v1/`), together with an offline demo, tests, and a `REPRODUCING.md` guide mapping each table to a command. The release is archived on Zenodo (release v1.0; DOI 10.5281/zenodo.23903302, which resolves to the latest version).

3.4 Statistical protocol

Both FlyHash's sparse projection matrix and SimHash's dense hyperplane matrix are randomly initialized, so a single evaluation reflects one particular random draw, not the expected behavior of the method. For our main results (Section 4.1) and hyperparameter ablation (Section 4.2), we repeat every evaluation across N independent random seeds (10 for the main results, 3 for the larger hyperparameter sweep, given its greater number of configurations), varying *only* the projection matrices' initialization across seeds while holding the dataset and the sampled set of queries fixed. This isolates projection-matrix randomness as the source of measured variance, rather than conflating it with variance from query sampling. We report the mean and standard deviation of mAP@R across seeds, and, as a simple and intentionally conservative robustness check, whether the gap between FlyHash's and SimHash's mean mAP exceeds twice their combined standard deviation ($\Delta > 2\sqrt{\sigma_{\text{fly}}^2 + \sigma_{\text{sim}}^2}$). We present this as a sanity check against attributing an observed gap to projection-matrix noise, not as a substitute for a formal statistical significance test.

Because FlyHash and SimHash are evaluated on the *same* 10 seeds (same sampled queries, same train/query split, differing only in which method's projection matrix that seed initializes), the two methods' mAP values at a given bit budget are naturally paired, not independent samples — each seed contributes one paired difference $d_i = \text{mAP}_{\text{fly},i} - \text{mAP}_{\text{sim},i}$. For our primary GloVe result (the first block of Table 2, and Table 3 below) we additionally report a paired analysis: the mean paired difference with both a t-interval and a bootstrap (10,000 resamples) 95% confidence interval, the paired effect size (Cohen's d_z), and two nonparametric tests that do not assume normality — the Wilcoxon signed-rank test and a sign-flip permutation test (10,000 permutations). This is a strictly stronger statement than the "robust" heuristic above, which compares marginal variances and discards the pairing.

3.5 Agent memory integration

To validate the practical use case motivating this work, we implement a minimal LLM agent memory system: a plain text `remember(text)` / `recall(query, k)` interface wraps a sentence-transformer embedding model and a FlyHash-backed store, so that an agent interacts only with text, not embeddings or hash codes directly. We wire this into a live tool-use loop with a general-purpose LLM exposing two tools — `save_memory`, called when the model judges the user has shared a fact worth remembering, and `recall_memory`, called when answering a question that may depend on a previously shared fact. In a representative interaction, a user states two facts in one turn ("My favorite language is Python and I'm allergic to peanuts"), asks an unrelated question in the next turn, and then asks "What should I avoid eating?" in a third turn; the agent correctly answers "peanuts" by calling `recall_memory`, retrieving the relevant fact via FlyHash-based Hamming-distance ranking over stored memory embeddings from several turns earlier. This demonstrates the full pipeline functioning end-to-end, independent of the benchmark-style evaluation in Section 4.

4 Results

4.1 Main result: corrected comparison across four datasets

Table 2 reports mAP@R (mean $\pm$ standard deviation over 10 random seeds) for FlyHash and SimHash at five bit budgets, on each of the four datasets described in Section 3, under the corrected protocol of Section 3.2. "Robust" indicates whether the gap between the two methods' means exceeds twice their combined standard deviation (Section 3.4).

Table 2: Main results: mAP (mean ± std over 10 seeds) for FlyHash vs. SimHash across four datasets and five bit budgets.

Dataset	k (bits)	FlyHash	SimHash	Robust
GloVe, mAP@100	16	**0.196 ± 0.007**	0.047 ± 0.006	Yes
	32	**0.314 ± 0.009**	0.096 ± 0.009	Yes
	64	**0.449 ± 0.008**	0.185 ± 0.008	Yes
	128	**0.585 ± 0.006**	0.316 ± 0.006	Yes
	256	**0.705 ± 0.006**	0.469 ± 0.007	Yes
MiniLM, mAP@20	16	**0.310 ± 0.013**	0.214 ± 0.026	Yes
	32	**0.408 ± 0.018**	0.266 ± 0.022	Yes
	64	**0.537 ± 0.030**	0.355 ± 0.024	Yes
	128	**0.656 ± 0.016**	0.466 ± 0.027	Yes
	256	**0.753 ± 0.013**	0.587 ± 0.016	Yes
`digits`, mAP@50	16	**0.456 ± 0.015**	0.119 ± 0.031	Yes
	32	**0.593 ± 0.010**	0.236 ± 0.039	Yes
	64	**0.699 ± 0.011**	0.397 ± 0.023	Yes
	128	**0.748 ± 0.005**	0.535 ± 0.021	Yes
	256	**0.806 ± 0.007**	0.664 ± 0.009	Yes
MNIST, mAP@100	16	**0.276 ± 0.009**	0.064 ± 0.014	Yes
	32	**0.382 ± 0.010**	0.133 ± 0.012	Yes
	64	**0.469 ± 0.011**	0.237 ± 0.014	Yes
	128	**0.527 ± 0.008**	0.351 ± 0.009	Yes
	256	**0.563 ± 0.004**	0.450 ± 0.010	Yes

t-interval and bootstrap CIs agree closely at every bit budget, and Cohen's d_z is an order of magnitude past the conventional "large" threshold of 0.8 throughout. The Wilcoxon and permutation p-values sit at $1/2^{10} \approx 0.00195$, the smallest value either test can report with $n = 10$ paired seeds — FlyHash wins on all 10 seeds at every budget, so this is the resolution floor of the test, not a computational artifact. This paired analysis, not the "gap $> 2\times$combined std" heuristic alone, is what we consider the primary evidence of a robust effect; the heuristic is retained in Table 2 as a quick, consistent summary across all four datasets.

FlyHash outperforms SimHash at every bit budget, on every dataset, with every gap classified robust — 20 independent (dataset, bit-budget) comparisons, zero exceptions. A consistent secondary pattern is also visible across all four datasets: the *relative* advantage ($\text{mAP}_{\text{fly}}/\text{mAP}_{\text{sim}}$) is largest at the smallest bit budget and shrinks as k grows — for example, on GloVe, FlyHash is $4.2\times$ SimHash at $k = 16$ but only $1.5\times$ at $k = 256$. Figure 1 plots this ratio against bit budget for all four datasets at once, making the shared pattern visible across modalities. We return to the significance of this pattern in Section 5.

4.2 Hyperparameter ablation

FlyHash has two hyperparameters left fixed throughout Section 4.1: sparsity s (fraction of input dimensions each Kenyon cell samples from; 0.10 by default, following the fly's own connectivity) and pool size m (2,000 by default). To check whether the Section 4.1 result depends on these specific choices, we sweep $s \in \{0.05, 0.10, 0.20\}$ and $m \in \{500, 1000, 2000, 4000\}$ — 12 combinations — at three representative bit budgets ($k \in \{16, 64, 256\}$), on GloVe, with 3 seeds per configuration (fewer than Section 4.1's 10, given the larger number of configurations swept here).

Table 3: Paired statistical analysis, GloVe: mean FlyHash-minus-SimHash mAP difference per seed (10 seeds).

k (bits)	Mean diff	95% CI (t)	95% CI (bootstrap)	Cohen's d_z	Wilcoxon p	Perm. p
16	0.1485	[0.1421, 0.1549]	[0.1437, 0.1539]	16.6	1.95e-03	2.30e-03
32	0.2179	[0.2073, 0.2284]	[0.2095, 0.2267]	14.8	1.95e-03	2.30e-03
64	0.2642	[0.2550, 0.2734]	[0.2568, 0.2719]	20.6	1.95e-03	2.30e-03
128	0.2691	[0.2643, 0.2738]	[0.2654, 0.2732]	40.8	1.95e-03	2.30e-03
256	0.2354	[0.2284, 0.2424]	[0.2299, 0.2414]	24.1	1.95e-03	2.30e-03

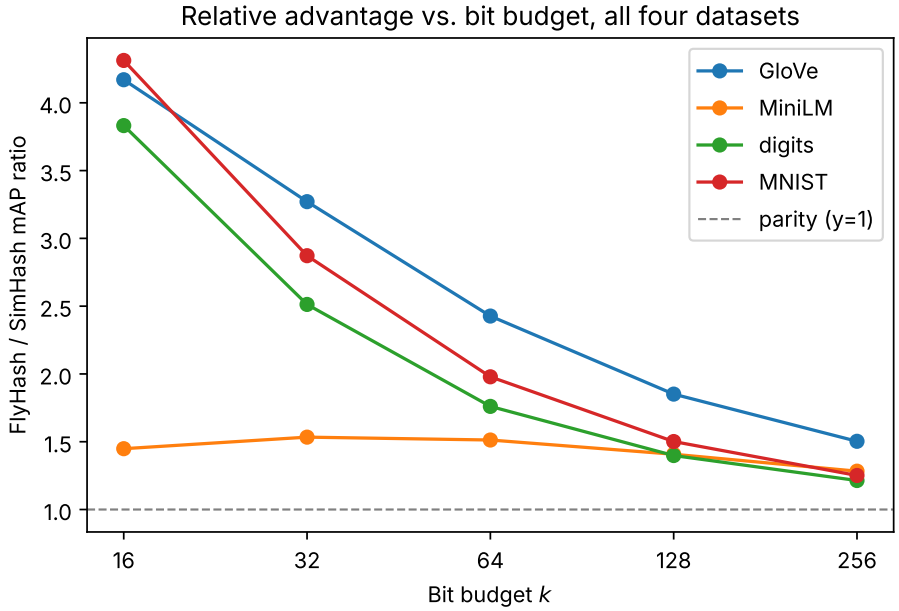

Figure 1: FlyHash/SimHash mAP ratio vs. bit budget, all four datasets. Every dataset's curve is monotonically decreasing and stays above the parity line ($y = 1$) across the full range tested, with GloVe and MNIST showing the steepest initial advantage and MiniLM the flattest curve (a smaller relative gap that still never crosses parity).

All 36 configurations exceed the SimHash baseline at their respective bit budget — no combination of sparsity and pool size reverses the result. Two secondary patterns are visible: sparsity somewhat above the fly's own biological ratio (0.20 vs. ≈0.10–0.12) performs marginally better at larger bit budgets, and increasing pool size beyond roughly 8× the bit budget yields diminishing or slightly negative returns. Neither pattern threatens the main result; both are minor tuning observations. Figure 2 renders Table 4 as heatmaps, which makes both secondary patterns easier to see at a glance than the raw numbers.

4.3 Memory and query latency: quality gains do not imply resource gains

The results above measure retrieval *quality* only. We separately measure memory footprint and query latency on GloVe (5,000 records), comparing FlyHash and SimHash at matched k against an exact (unhashed) Euclidean nearest-neighbor baseline, using brute-force comparison for all three (no approximate index on either side).

Table 4: FlyHash mAP@100 (mean ± std over 3 seeds) across the sparsity × pool-size grid, at three bit budgets. SimHash's mAP@100 at the same k is given as the baseline for each sub-table.

$k = 16$ (SimHash baseline: 0.046 ± 0.004)				
s	$m = 500$	$m = 1000$	$m = 2000$	$m = 4000$
0.05	0.169±0.009	0.169±0.003	0.178±0.006	0.186±0.002
0.10	0.200±0.001	0.199±0.006	0.193±0.010	0.189±0.008
0.20	0.198±0.003	0.205±0.003	0.200±0.005	0.204±0.001
$k = 64$ (SimHash baseline: 0.177 ± 0.004)				
s	$m = 500$	$m = 1000$	$m = 2000$	$m = 4000$
0.05	0.434±0.014	0.415±0.006	0.379±0.005	0.378±0.006
0.10	0.446±0.004	0.456±0.003	0.451±0.009	0.442±0.014
0.20	0.453±0.005	0.472±0.001	0.472±0.003	0.472±0.004
$k = 256$ (SimHash baseline: 0.469 ± 0.003)				
s	$m = 500$	$m = 1000$	$m = 2000$	$m = 4000$
0.05	0.588±0.007	0.667±0.003	0.684±0.007	0.632±0.011
0.10	0.607±0.004	0.683±0.002	0.707±0.006	0.689±0.007
0.20	0.609±0.003	0.686±0.002	0.712±0.003	0.719±0.002

Table 5: Packed memory and mean query latency (50 queries, brute-force ranking against all 5,000 records). Exact float32 baseline: 400 bytes/record, 0.78 ms/query.

k (bits)	SimHash bytes/record	FlyHash (dense) ms/query	SimHash ms/query
16	2	14.05	0.19
32	4	12.76	0.29
64	8	12.93	0.45
128	16	14.15	0.84
256	32	13.17	2.17

FlyHash's code is not itself k bits wide — it is m bits wide (2,000 in our default configuration), of which exactly k are active. Our default dense representation stores and compares this full m-bit array regardless of k, which makes FlyHash's packed size a near-constant ∼250 bytes/record independent of k — larger than SimHash's genuinely k-bit code at every budget tested except $k \gtrsim 250$ — and its brute-force query time a near-constant ∼13 ms/query, 10–70× slower than SimHash at the same k. This is a real structural property of the dense representation, not a bug: mAP depends only on *which* k bits are active, so it is entirely unaffected, but memory and speed under this representation clearly are.

4.4 A sparse-set comparison: partial recovery

Because every FlyHash code has exactly k active bits, the Hamming distance between two such codes reduces to $2k - 2|A \cap B|$, where A and B are their sets of active indices. Ranking by Hamming distance is therefore identical to ranking by descending intersection size, which we compute as a sparse-sparse matrix product over only the active indices, rather than a dense operation over the full m-wide array.

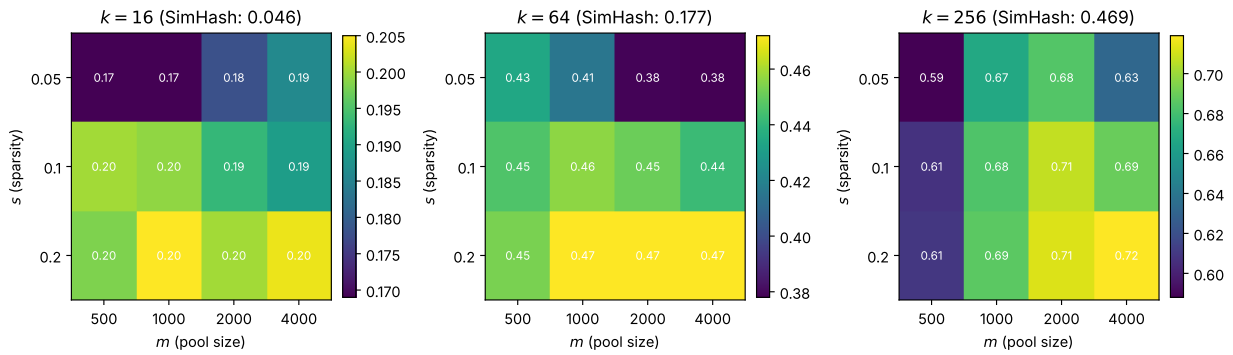

Figure 2: FlyHash mAP@100 across the sparsity × pool-size grid, at $k = 16, 64, 256$. Each panel's SimHash baseline is given in its title; every cell in every panel is above that baseline. The rightward brightening within the $k = 256$ panel visualizes the pool-size effect, and the top-to-bottom brightening within each panel visualizes the sparsity effect.

Table 6: Sparse vs. dense FlyHash comparison, and memory as an explicit index list (uint16 per active index).

k (bits)	Sparse-idx bytes/rec	SimHash bytes/rec	Dense ms/query	Sparse ms/query	SimHash ms/query
16	32	2	15.77	0.71	0.40
32	64	4	14.44	0.73	0.29
64	128	8	14.54	1.59	0.45
128	256	16	15.03	3.39	0.84
256	512	32	13.90	8.55	1.79

The sparse comparison is 10–22× faster than the dense comparison and, unlike the dense version, its latency correctly scales with k rather than remaining flat. It does not, however, close the gap to SimHash entirely: SimHash's small dense $k \times k$ matrix multiplication remains 2–5× faster at every budget tested, and explicit index-list storage, while now also correctly scaling with k, remains ∼16× larger than SimHash's bit-packed code, since storing an explicit position costs a full integer where bit-packing costs a single bit. We report this as a genuine, partial improvement — not a solved problem — and discuss the remaining gap in Section 6.

4.5 Agent memory demonstration

The end-to-end system described in Section 3.5 functions correctly: in the representative interaction described there, the agent retrieves a fact shared several turns earlier via FlyHash-based Hamming ranking over MiniLM sentence embeddings, without that fact appearing anywhere in the immediate conversational context. This confirms the pipeline is functional for its intended use case, independent of the benchmark numbers in Sections 4.1–4.4.

As one measured data point for the full pipeline's latency (embedding included, not just the isolated hash comparison of Sections 4.3–4.4), we timed a single representative query against the 6,000-record corpus of Section 4.7, averaged over 30 repeats on consumer hardware, using the sparse-set comparison method of Section 4.4 for the hash-ranking step (the same method behind Table 6's speed numbers, consistent with the rest of the paper rather than an unrelated, unoptimized implementation):

Table 7: End-to-end pipeline latency, one representative query (30 repeats).

Stage	ms/query	% of total
Sentence embedding (MiniLM)	14.1	68.8%
FlyHash hash + sparse rank	6.4	31.2%
Total end-to-end	**20.5**	100%

With the hash-ranking step measured consistently with Table 6, sentence embedding — not FlyHash's hash comparison — is the dominant cost of the full pipeline, by roughly 2:1. This is a useful finding in its own right: it means that, at this corpus size, FlyHash's own contribution to end-to-end latency is modest relative to the (hash-independent) embedding step, and that further latency optimization in a real deployment would pay off more from the embedding stage than the hash comparison. Embedding time is identical regardless of which hash is used downstream, so this number has no bearing on the FlyHash-vs-SimHash comparisons elsewhere in the paper — it exists only to put a real figure on the full pipeline's latency rather than leaving it unmeasured. We still report it as a single data point on one machine, not a benchmark: no seed averaging and no variation in corpus size or hardware. A proper system-level benchmark remains future work (Section 6).

4.6 Practical systems comparison: FlyHash/SimHash vs. a graph-based ANN index

Sections 4.1–4.4 hold the active-bit budget k fixed and compare FlyHash against SimHash — both compact hash codes compared by Hamming distance. This answers "which hash is better at a given bit budget," but not "would I actually use either of these instead of a standard production ANN index." We report that second, separate question here using FAISS's `IndexHNSWFlat` ($M = 32$, efConstruction= 200) over full-precision GloVe vectors (5,000 records, 50 queries, mAP@100), explicitly as a systems comparison rather than a matched-bit one: HNSW has no "bits" parameter comparable to k, so its quality/speed/memory trade-off is governed by efSearch instead, and the two tables below should be read within themselves, not row-matched against each other.

Table 8: FAISS HNSW, varying efSearch (exact L2 baseline: 0.909 ms/query, 400 bytes/record).

efSearch	mAP@100	ms/query	bytes/record
16	0.7700	0.029	671.5
64	0.9862	0.060	671.5
128	0.9986	0.092	671.5

Table 9: FlyHash (sparse-set) and SimHash, matched active-bit k, same dataset/queries.

k (bits)	FlyHash mAP	FlyHash ms/q	FlyHash B/rec	SimHash mAP	SimHash ms/q
16	0.198	0.66	32	0.048	0.19
32	0.312	1.04	64	0.095	0.22
64	0.440	1.62	128	0.185	0.26
128	0.582	4.46	256	0.325	0.50
256	0.697	10.13	512	0.476	0.70

At this scale, HNSW dominates both hash methods on quality and latency simultaneously — mAP above 0.98 and sub-0.1ms queries, at only $\sim 1.6\times$ the raw exact-storage memory cost (671.5 vs. 400 bytes/record) — while FlyHash and SimHash's only genuine advantage is extreme compactness well below exact storage (as little as 2–32 bytes/record). The honest conclusion is that a practitioner with no unusual memory constraint and a dataset at this scale should use HNSW, not a hash code; FlyHash/SimHash's relevance is specifically to the regime HNSW does not target — severe, explicit storage budgets (e.g., many independent memory stores on resource-constrained agents) — and we revise our framing in Section 5 accordingly rather than letting Sections 4.1–4.4 imply a general retrieval recommendation.

4.7 A realistic, labeled agent-memory benchmark

Every result above uses nearest-neighbor-by-embedding as a stand-in for "relevance" — two items are treated as relevant to each other if an unhashed embedding considers them close, which is circular with respect to the embedding itself and does not test anything about whether hashing preserves *true* query-to-memory relevance. We address this with a separate benchmark using constructed relevance labels that do not derive from any embedding distance.

Established agent-memory benchmarks such as LoCoMo (Maharana et al., 2024) and LongMemEval (Wu et al., 2024) evaluate end-to-end conversational QA accuracy over multi-session dialogue, which conflates retrieval quality with the downstream LLM's reasoning over retrieved context (Wu et al., 2024, explicitly separate indexing, retrieval, and reading for this reason). Our evaluation instead requires ground-truth relevance labels for the retrieval step in isolation — the quantity FlyHash and SimHash actually compete on — which these QA-style benchmarks do not provide directly. We therefore construct a dataset with explicit, non-circular relevance labels anchored to named personas below, in the same spirit as recent calls for benchmark designs that disentangle retrieval correctness from answer correctness (Wu et al., 2024; Hao et al., 2026).

Construction. We generate 150 synthetic named personas, each with 10 independent facts (favorite color, allergy, job title, hometown, and six others), each fact expressed via 3–4 paraphrase templates that name the persona directly (e.g., "Jordan's favorite color is teal." / "If you ask Jordan, teal is their favorite color."), giving 6,000 memory sentences total. Each query is a differently-worded question that also names a target persona (e.g., "Which color does Jordan like best?"), and its ground-truth relevant set is exactly that persona's own paraphrases for that fact (3–4 sentences out of 6,000) — a label fixed by construction, not by embedding similarity. Naming every fact and question to a specific persona is what makes the corpus both genuinely confusable (thousands of near-identical same-template statements about *other* people crowd every query) and answerable from text content alone (the one token that correctly disambiguates the right answer — the name — is actually present in the query), which an earlier, unnamed version of this design was not: without a persona name tying question to answer, no text-similarity method, hashed or not, can identify *whose* stated value is correct among many textually identical same-template statements, and we do not report numbers from that earlier, ill-posed version.

Embedding. We embed both memories and queries with TF-IDF followed by truncated SVD (256 components), fit on the memory corpus — an offline, network-free stand-in for a real sentence-transformer embedding, consistent with this project's practice elsewhere (Section 6).

FlyHash beats SimHash robustly at every budget on this genuinely-labeled task, with absolute mAP values well above chance (random ranking over 6,000 items with $\sim$4 relevant would score far below 0.1) but, expectedly, lower than Table 2's nearest-neighbor-pseudo-label mAP — this is a harder, more realistic task by construction, and the two numbers are not meant to be compared directly. This result is, to our knowledge, the first time FlyHash's advantage over SimHash has

Table 10: FlyHash vs. SimHash, standard mean Average Precision against constructed relevance labels (150 queries, 6,000 memories, 10 seeds).

k (bits)	FlyHash mAP (labeled)	SimHash mAP (labeled)	Robust
16	**0.195 ± 0.014**	0.038 ± 0.010	Yes
32	**0.271 ± 0.017**	0.105 ± 0.011	Yes
64	**0.324 ± 0.012**	0.212 ± 0.016	Yes
128	**0.365 ± 0.009**	0.298 ± 0.011	Yes
256	**0.392 ± 0.009**	0.345 ± 0.011	Yes

been shown against actual, non-circular relevance labels for an agent-memory-shaped task rather than an embedding-similarity proxy.

5 Discussion

Why is the advantage largest at low bit budgets? FlyHash's sparse expansion into a much higher-dimensional space before selecting winners appears to preserve more of the input's neighborhood structure per output bit than SimHash's direct dense projection, particularly when very few bits are available to encode that structure. As the bit budget grows, both methods have enough capacity to represent more of the relevant structure, and the relative gap narrows — though FlyHash's absolute advantage never disappears across any budget we tested (Table 2). This pattern is practically significant: real deployments of compact hashing are motivated precisely by wanting *few* bits per record (that is the entire point of hashing over storing full embeddings), so the regime where FlyHash's advantage is largest is also the regime a real system would actually operate in.

A quality advantage is not automatically a resource advantage. Section 4.3 is, to our knowledge, an underappreciated point in this line of work: a hash function can preserve neighbor structure strictly better than a baseline while, in a naive implementation, being simultaneously worse on memory and speed — because "better structure per bit" and "cheap to store and compare" are different properties that do not automatically follow from each other. Section 4.4 shows that a principled fix targeting the comparison's sparsity recovers much, but not all, of one axis (speed) while leaving the other (memory) only partially improved. We think this distinction — quality-per-bit versus deployed resource cost — is worth making explicit in any evaluation of a hashing scheme derived from a sparse-coding argument, not only in this paper.

Practical implication for agent memory. An agent adopting FlyHash today obtains a genuine, robustly validated retrieval-quality improvement per bit over SimHash, particularly in the low-bit-budget regime most relevant to a resource-constrained memory system, and (Section 4.7) that advantage survives contact with genuine, non-circular relevance labels, not just an embedding-similarity proxy. It should not, on the basis of this paper, expect a memory or latency improvement over a well-implemented SimHash baseline without further engineering (Section 6), nor should it expect either hash method to beat a standard production ANN index such as HNSW at ordinary scale (Section 4.6) — HNSW wins on quality and latency simultaneously there, at a modest memory premium over exact storage. Which of these properties matters most depends on the deployment: a system already committed to compact hash codes, bottlenecked on retrieval accuracy at a fixed, small bit budget, benefits from FlyHash as a drop-in SimHash replacement; a system with no unusual storage constraint should likely use HNSW instead of either hash; and a system bottlenecked specifically on raw per-record storage at a scale where even HNSW's overhead is unaffordable is the narrow regime where FlyHash/SimHash-style hashing is the right tool at all.

6 Limitations and Future Work

- **Memory and speed do not fully match the quality advantage, even after a targeted attempt to close the gap.** Section 4.4's sparse-set comparison recovers most of the speed disadvantage but leaves a meaningful memory disadvantage; a genuinely sublinear index (e.g., banded or multi-probe LSH, or an inverted index verified sublinear at scale — an earlier inverted-index attempt in our own implementation was found not to be sublinear at the scales we tested) remains unbuilt. A more compact encoding of the active-index list (delta-coding of sorted indices, variable-length or entropy coding, or reducing pool size at some cost to quality) is a concrete, identified next step we did not pursue here.

- **Our HNSW comparison (Section 4.6) is a single-scale, single-configuration snapshot.** We report one dataset (GloVe, 5,000 records) and one HNSW configuration ($M = 32$, efConstruction= 200, three efSearch values); HNSW's relative advantage over hashing shrinks as per-record storage budgets tighten and may shrink or reverse at very different corpus scales, which we have not swept. A combined approach (e.g., FlyHash codes as a cheap pre-filter ahead of an HNSW-indexed re-ranking stage, in the spirit of how FAISS combines quantization with HNSW) is a natural direction we have not explored.

- **SIFT1M, part of the original paper's evaluation, is not included.** We judged that four independent datasets across two modalities already provide strong evidence the result is not dataset-specific, and that a fifth, necessarily subsampled dataset (SIFT1M's full 1M-vector scale is impractical to evaluate exhaustively at our compute budget) would add breadth rather than materially new evidence. We report this as a deliberate scope decision, not an oversight.

- **We do not implement DenseFly's densification or BioHash's learned projection**, both of which are reported improvements over the plain, random-projection FlyHash evaluated here (Section 2). Combining either with the agent-memory application is natural future work.

- **We compare two hashing paradigms (sparse random-projection vs. dense random-hyperplane) but do not benchmark against vector quantization methods** (e.g., Product Quantization, RaBitQ), which the literature indicates generally outperform hashing in accuracy-per-bit (Jégou et al., 2011; Gao & Long, 2024; Section 2); a comparison against quantization is natural future work alongside the HNSW baseline already included.

- **The real-text (MiniLM) evaluation uses a small, 100-sentence corpus**, chosen to avoid duplicate-embedding artifacts (see Appendix A) rather than to represent realistic agent-memory scale. A larger, more diverse real-text corpus would strengthen this specific result.

- **We report only a single data point for the full end-to-end agent memory system's latency** (Section 4.5: 20.5 ms/query on one machine, one query, 30 repeats), not a proper benchmark — no seed averaging and no sweep over corpus size or hardware. A dedicated, multi-seed system-level benchmark is future work.

- **The realistic, labeled agent-memory benchmark (Section 4.7) uses templated synthetic text, not naturally-occurring language, and a TF-IDF+SVD embedding as a stand-in for a real sentence-transformer.** Real sentence diversity is considerably messier than any fixed template set, and real deployments would more plausibly scope memory retrieval by session or user ID than rely on a name mentioned in free text. A rerun with real human-written memory logs and a real sentence-transformer embedding (on a machine with internet access, the

same gap already noted for MiniLM/MNIST elsewhere in this paper) would strengthen this specific result.

7 Conclusion

Our three research questions have three different shapes of answer, and we state the claim no wider than that. FlyHash beats SimHash at matched active-bit counts, robustly, on every dataset and seed we tested, including under genuine agent-memory relevance labels rather than an embedding-similarity proxy (RQ1, RQ3) — this is a real, well-supported result, and one previously undocumented comparison pitfall (k-vs-m conflation) had to be corrected to see it. But the practical deployment question (RQ2) does not resolve as cleanly: FlyHash's quality advantage over SimHash does not translate into a memory or latency advantage under a naive comparison, a sparse-set comparison closes most but not all of that gap, and a standard production index (HNSW) outperforms both hash methods on every axis at ordinary scale, leaving genuinely extreme storage constraints as hashing's specific niche rather than a general deployment recommendation. We think this — a validated hash-quality result, a corrected methodological pitfall, and a resource picture that narrows rather than flatters the headline claim — is a more useful contribution to reproducible progress here than a report that stopped at the first, more flattering result.

References

[1] Dasgupta, S., Stevens, C. F., & Navlakha, S. (2017). A neural algorithm for a fundamental computing problem. *Science*, 358(6364), 793–796.

[2] Caron, S. J. C., Ruta, V., Abbott, L. F., & Axel, R. (2013). Random convergence of olfactory inputs in the Drosophila mushroom body. *Nature*, 497(7447), 113–117. `https://doi.org/10.1038/nature12063`

[3] Charikar, M. S. (2002). Similarity estimation techniques from rounding algorithms. *Proceedings of the 34th Annual ACM Symposium on Theory of Computing (STOC '02)*, 380–388, Montréal, Canada. ACM.

[4] Sharma, J., & Navlakha, S. (2018). Improving Similarity Search with High-dimensional Locality-sensitive Hashing. *arXiv:1812.01844*.

[5] Ryali, C., Hopfield, J., Grinberg, L., & Krotov, D. (2020). Bio-Inspired Hashing for Unsupervised Similarity Search. *Proceedings of the 37th International Conference on Machine Learning (ICML)*, PMLR 119, 8295–8306.

[6] Ram, P., & Sinha, K. (2022). Federated Nearest Neighbor Classification with a Colony of Fruit-Flies. *Proceedings of the AAAI Conference on Artificial Intelligence*, 36(7), 8036–8044. *arXiv:2112.07157*.

[7] Huang, T. (2023). FlyHash [Python package, version 1.1.1]. PyPI. `https://pypi.org/project/FlyHash/`

[8] Malkov, Y. A., & Yashunin, D. A. (2020). Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 42(4), 824–836.

[9] Pennington, J., Socher, R., & Manning, C. D. (2014). GloVe: Global Vectors for Word Representation. *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 1532–1543, Doha, Qatar. Association for Computational Linguistics. https://doi.org/10.3115/v1/D14-1162

[10] Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP-IJCNLP)*, 3982–3992, Hong Kong, China. Association for Computational Linguistics. https://doi.org/10.18653/v1/D19-1410. Model checkpoint: https://huggingface.co/sentence-transformers/all-MiniLM-L6-v2

[11] LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324. Accessed via OpenML: https://www.openml.org/d/554

[12] Pedregosa, F., Varoquaux, G., Gramfort, A., et al. (2011). Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830. digits dataset documentation: https://scikit-learn.org/stable/datasets/toy_dataset.html#optical-recognition-of-handwritten-digits-dataset

[13] Chhikara, P., Khant, D., Aryan, S., Singh, T., & Yadav, D. (2025). Mem0: Building Production-Ready AI Agents with Scalable Long-Term Memory. *arXiv:2504.19413*.

[14] Zhao, P., Zhang, H., Yu, Q., Wang, Z., Geng, Y., Fu, F., Yang, L., Zhang, W., Jiang, J., & Cui, B. (2024). Retrieval-Augmented Generation for AI-Generated Content: A Survey. *arXiv:2402.19473*.

[15] Johnson, J., Douze, M., & Jégou, H. (2021). Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3), 535–547.

[16] Du, P. (2026). Memory for Autonomous LLM Agents: Mechanisms, Evaluation, and Emerging Frontiers. *arXiv:2603.07670*.

[17] Zeng, R., Fang, J., Liu, S., & Meng, Z. (2024). On the Structural Memory of LLM Agents. *arXiv:2412.15266*.

[18] Hu, Z., Zhu, Q., Zhao, R., Liang, D., Yan, H., He, Y., & Gui, L. (2026). Beyond RAG for Agent Memory: Retrieval by Decoupling and Aggregation. *arXiv:2602.02007*.

[19] Cao, S., He, J., & Tan, F. (2026). HiGMem: A Hierarchical and LLM-Guided Memory System for Long-Term Conversational Agents. *arXiv:2604.18349*.

[20] Idrees, Y. A., Neiroukh, H., & Al Badarneh, A. (2026). Vector Databases in the Age of AI: Foundations, Architectures, and Emerging Applications. *The 17th International Conference on Information and Communication Systems (ICICS 2026)*, Irbid, Jordan. ACM. https://doi.org/10.1145/3812734.3813407

[21] Jégou, H., Douze, M., & Schmid, C. (2011). Product Quantization for Nearest Neighbor Search. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 33, 117–128.

[22] Ge, T., He, K., Ke, Q., & Sun, J. (2013). Optimized Product Quantization for Approximate Nearest Neighbor Search. *2013 IEEE Conference on Computer Vision and Pattern Recognition*, 2946–2953.

[23] Gao, J., & Long, C. (2024). RaBitQ: Quantizing High-Dimensional Vectors with a Theoretical Error Bound for Approximate Nearest Neighbor Search. *Proceedings of the 2024 ACM SIGMOD International Conference on Management of Data (SIGMOD '24)*. arXiv:2405.12497.

[24] Laarhoven, T. (2020). Polytopes, lattices, and spherical codes for the nearest neighbor problem. *47th International Colloquium on Automata, Languages, and Programming (ICALP 2020)*, 76:1–76:14.

[25] Cai, D. (2016). A Revisit of Hashing Algorithms for Approximate Nearest Neighbor Search. *arXiv:1612.07545*.

[26] Formal, T., Lassance, C., Piwowarski, B., & Clinchant, S. (2024). Towards Effective and Efficient Sparse Neural Information Retrieval. *ACM Transactions on Information Systems*.

[27] Maharana, A., Lee, D.-H., Tulyakov, S., Bansal, M., Barbieri, F., & Fang, Y. (2024). Evaluating Very Long-Term Conversational Memory of LLM Agents. *arXiv:2402.17753*.

[28] Wu, D., Wang, H., Yu, W., Zhang, Y., Chang, K.-W., & Yu, D. (2024). LongMemEval: Benchmarking Chat Assistants on Long-Term Interactive Memory. *arXiv:2410.10813*.

[29] Hao, X., Zhang, Z., Lin, Z., Sun, Y., Guo, Z., Zhang, X., Liang, Y., Xiong, F., & Li, Z. (2026). MemOps: Benchmarking Lifecycle Memory Operations in Long-Horizon Conversations. *arXiv:2607.12893*.

A Reproducibility Investigation

While reproducing the original paper's reported result on GloVe word vectors, we initially found FlyHash losing to SimHash on every dataset we tried — synthetic clustered embeddings, TF-IDF text vectors, real sentence embeddings, and real GloVe vectors — directly contradicting the published finding on GloVe. We treated this discrepancy as a signal that our comparison, not the algorithm, was most likely at fault, and pursued it systematically rather than accepting or explaining away the contradiction.

Hypotheses ruled out. We first tested whether L2-normalizing inputs before hashing affected the result; it did not, meaningfully. We then tested whether using the algorithm's exact original weight definition (true binary $\{0, 1\}$ projection weights, rather than the signed $\{-1, +1\}$ weights we had implemented) affected the result; it did not, and we quantify this claim across all five main bit budgets below rather than resting it on the single seed we checked at the time. We then tested whether our evaluation metric — recall@k by overlap, rather than mAP@R — was responsible; switching to mAP@R produced the same qualitative picture. Each ruled-out hypothesis narrowed the search space without resolving the discrepancy.

Quantifying the weight-mode claim. The "it did not, meaningfully" judgment above was originally made on a single seed. Re-run across all 10 seeds and all five main bit budgets (GloVe), with SimHash included as a reference:

Both weight modes beat SimHash by a wide margin at every budget tested. The signed-vs-binary difference is not zero and grows with k (0.004 at $k = 16$ to 0.058 at $k = 256$), so the honest claim is not "no difference" but "small relative to the FlyHash-vs-SimHash effect at every budget we tested, with a modest upward trend at higher k" — the main results in Section 4.1 use the signed variant throughout, and this table bounds how much of that result could plausibly be an artifact of that choice.

Table 11: Signed vs. binary FlyHash projection weights, GloVe, 10 seeds.

k (bits)	Signed $\{-1, +1\}$ mAP	Binary $\{0, 1\}$ mAP	|diff|	SimHash mAP (ref.)
16	0.196 ± 0.007	0.191 ± 0.005	0.0042	0.047 ± 0.006
32	0.314 ± 0.009	0.298 ± 0.007	0.0156	0.096 ± 0.009
64	0.449 ± 0.008	0.421 ± 0.005	0.0277	0.185 ± 0.008
128	0.585 ± 0.006	0.540 ± 0.005	0.0449	0.316 ± 0.006
256	0.705 ± 0.006	0.647 ± 0.005	0.0577	0.469 ± 0.007

The actual cause. We resolved the discrepancy by comparing our implementation's parameterization directly against a third-party reference implementation of the original algorithm (`github.com/dataplayer12/Fly-LSH`). This revealed that our comparison held FlyHash's *total* hash dimensionality equal to SimHash's bit count while separately fixing FlyHash's winner-take-all *percentage* — silently making FlyHash's *effective* number of active bits, at any given nominal "bits" value, a small fraction of what SimHash was being given as its full, genuinely active bit count. This is precisely the k-versus-m conflation described in Section 3.2. Correcting the protocol to hold k (not m) equal between the two methods, with m fixed independently at a large value, immediately reproduced the original paper's reported pattern: FlyHash winning decisively on GloVe at every bit budget, and — as it turned out — on every other dataset we subsequently tested as well (Section 4.1).

A hypothesis retracted. Prior to finding this cause, we had observed that FlyHash appeared to outperform SimHash specifically on data with correlated, non-negative features (e.g., TF-IDF text vectors) but not on isotropic synthetic data, and had proposed that FlyHash specifically benefits from such data properties. Once the actual comparison bug was corrected, this pattern disappeared: FlyHash outperformed SimHash on the corrected synthetic isotropic data just as decisively as on text data. We state plainly that our earlier hypothesis was incorrect — it was constructed to explain an artifact of a flawed comparison, not a genuine property of either the data or the algorithm — and we retract it here rather than omit it from this account.

We document this investigation in this level of detail because the underlying pitfall — conflating a hash's *effective* bit count with its *nominal* parameter dimensionality — seems easy to make and is, to our knowledge, not explicitly warned against elsewhere in this literature. We hope stating it precisely, alongside the specific hypotheses that seemed plausible but were not the actual cause, helps others attempting a similar comparison avoid the same error.