arXiv:2611.07477v1 [cs.DC] 5 Nov 2026

KOSMOS: SELF-SCHEDULING PERSISTENT KERNELS FOR COMMUNICATION-COMPUTE OVERLAP

Alexander Magro [1] **Aristotle Martin** [1] **Ye Wang** [1] **Wen Chen** [1]

ABSTRACT

Tensor and sequence parallelism (TP+SP) and expert parallelism (EP) put collectives on the critical path of transformer layers. State-of-the-art fused kernels mitigate this with fixed roles or schedules, but this rigidity leaves large idle windows. Predicting a communication-computation split does not remove them either; a predetermined split cannot follow data readiness. We present KOSMOS, a library for self-scheduling persistent kernels whose workgroups use queues ordered by consumer demand to switch between communication and computation roles as data become ready, with no scheduler, per-call cost model, or tuned split. KOSMOS fuses every TP+SP collective of a transformer layer with its GEMM, and runs an entire expert layer, from router to fold, in one launch per direction. Evaluated on AMD Instinct™ MI350X and AMD Instinct™ MI355X GPUs, KOSMOS speeds up TP+SP end-to-end training in Megatron-LM by 1.10–1.28× in BF16 and 1.10–1.25× in MXFP8. KOSMOS accelerates end-to-end training of DeepSeek-V3 and Qwen3-235B proxies by 1.29–1.83× over DeepEP.

1 INTRODUCTION

Training frontier large language models (LLMs) requires parallelism strategies that trade memory for communication. Memory reduction strategies such as tensor and sequence parallelism (TP+SP) and expert parallelism (EP) in Mixture-of-Experts (MoE) put collectives on the critical path. TP+SP wraps every transformer block in all-gather (AG) and reduce-scatter (RS) collectives, whereas EP in MoE models adds all-to-alls (A2A) around each expert layer. To keep matrix cores and links saturated, overlapping communication and computation is critical to training throughput.

Fusing communication into general matrix multiplication (GEMM) avoids hardware-scheduler contention and moves data at tile granularity, overlapping with the GEMM's tile-based computation. Prior fused kernels dedicate a fixed number of compute units (CUs) to communication (Zheng et al., 2025a; Awad et al., 2025; Sul et al., 2026a) or assign each workgroup (WG) its tasks with a scheduler (Aimuyo et al., 2025; Cheng et al., 2026). However, communication demand is not fixed: In dense collectives it shifts as the kernel progresses, and in MoE it shifts with each batch's routing. Predicting a split of communication and computation WGs does not solve this issue. A planned split from a cost model before each call is slower than a fixed split in our measurements (Table 2). Even adapting splits at runtime, as mKernel does with a kernel controller, leaves their kernel 1.18× slower than the best fixed split (Mao et al., 2026).

We present KOSMOS (Kernel-Overlapped Self-scheduling Megakernels On Symmetric-memory), a library of persistent kernels in which WGs dynamically change roles as the kernel runs. A WG is either a mover, which mainly communicates, or a builder, which mainly computes GEMM tiles. A WG selects its role from data readiness, with no scheduler or per-call cost model. Scheduling is implicit: Data is produced in the order consumers claim it. Builders start on local data while movers gather; reductions are produced, ordered by groups of tiled rows (band-major), so every rank's movers start after one step, and dispatch delivers experts in the order each GEMM consumes them.

The switching policy is set by each collective-GEMM pair (Figure 1). In AG+GEMM and MoE dispatch, movers become builders once their transfers are done; in GEMM+RS (Figure 1(A)), builders join the reduction when the GEMM work queue is exhausted; in MoE combine (Figure 1(B)) WGs switch both ways, pushing whenever data is ready and building otherwise. The policy is designed around the shape of the demand: Two-way switching is slower for TP+SP and A2A dispatch collectives whose demand is monotonic; one-way switching leaves combine gaps larger than needed.

KOSMOS provides broad coverage, including, to our knowledge, the first readiness-driven TP+SP fused kernels and MoE megakernels for training. Our contributions are:

1. **Self-scheduling persistent kernels**, with producer and consumer queues co-designed per collective around data readiness. WGs may then switch roles from local

[1]Advanced Micro Devices Inc., Austin, TX, USA. Correspondence to: Alexander Magro <alex.magro@amd.com>.

Preprint.

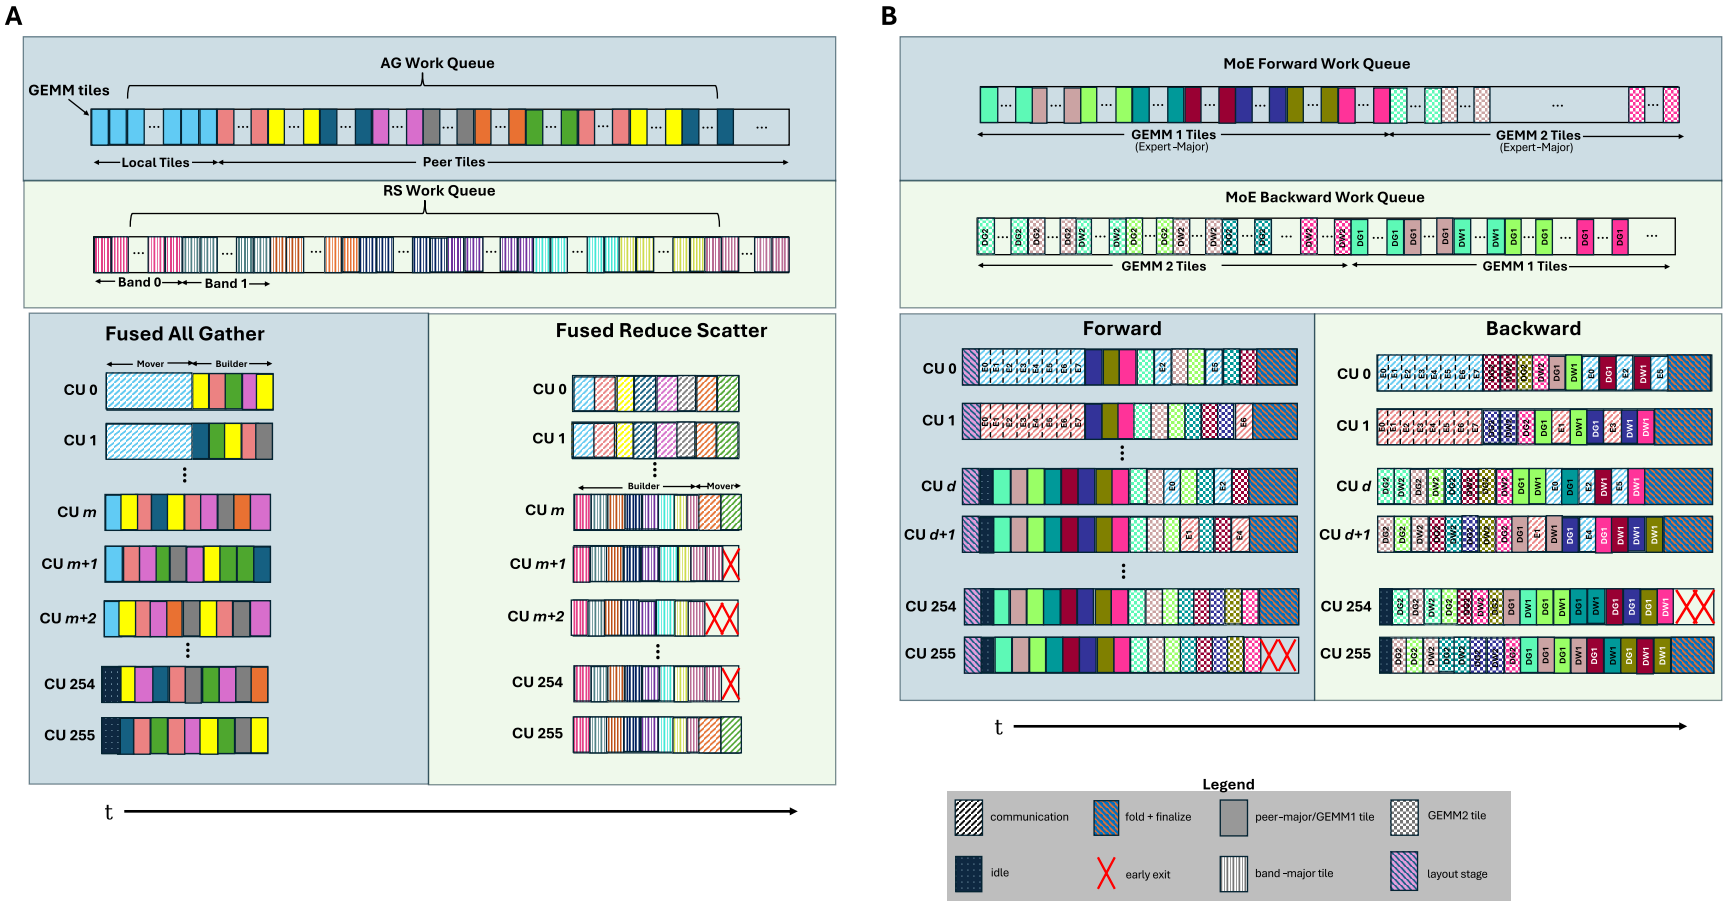

Figure 1. CU timelines and work queues on eight GPUs for (A) fused AG and RS and (B) the MoE forward and backward. Each CU runs one persistent WG; queues hold GEMM output tiles, and hatched segments are transfers (color = rank or expert). (A) The AG queue holds local tiles first, then peer tiles interleaved row by row; the RS queue is band-major (pattern = band). AG movers (CUs 0–1 of m) become builders once their transfers finish; in RS, builders become movers when the GEMM queue empties. (B) The forward queue is expert-major, GEMM1 before GEMM2; the backward queue holds GEMM2 then GEMM1 gradient tiles, interleaving dG and dW. All CUs first route and plan (purple); dispatchers push, then build; combine movers (CUs d, $d+1$) build until a combine unit is ready, then interleave pushes with GEMM2 tiles. Backward: dY is dispatched first and the dX combine overlaps the GEMM1 gradients.

tile-level checks. Switching lets a shape-only heuristic replace per-shape tuning: within 1% of, and up to 15% faster than, the best single fixed split. A fixed split in turn beats a split predicted per call.

2. **A characterization of where idle time falls** in fused communication kernels and a policy rule that reduces it: one-way switching for monotonic demand, two-way switching for interleaved demand.

3. **Complete kernel coverage for training:** persistent fused kernels for every TP+SP collective, forward and backward, in BF16 and MXFP8, and the first role-switching MoE layer with a single launch per direction, from router to fold, including gradients, evaluated end-to-end in Megatron-LM on two GPU models.

4. **Communications optimized for bytes/CU/s:** collectives for fully connected meshes that saturate the bandwidth while maintaining fine grained access, and measurements showing that on power-capped GPUs the CUs lent to communication cost the GEMM little.

We demonstrate KOSMOS on MI350-series GPUs and expect many of its findings to generalize to other systems: On AMD Instinct™ MI350X and AMD Instinct™ MI355X GPUs, KOSMOS speeds up end-to-end Llama 3 training by up to 1.28× over Transformer Engine (TE), and its MoE layers speed up training by up to 1.83× over DeepEP.

2 Background and Motivation

2.1 Parallel Strategies

2.1.1 Tensor and Sequence Parallelism

Under TP+SP (Shoeybi et al., 2020; Korthikanti et al., 2023), each transformer block contains a column-parallel and a row-parallel GEMM pair (Figure 5). Activations outside the pair are sequence-sharded. An AG assembles the sequence before a column-parallel GEMM, and an RS after the row-parallel GEMM reduces sequence shards. The backward pass adds three collectives in BF16: AGs of the output gradients and the saved input, and an RS of the input gradient. MXFP8 gathers output gradients twice, quantized along both dimensions. Each GEMM pair runs five or six collectives per training step, each on the critical path unless overlapped.

2.1.2 Fully Sharded Data Parallelism

Fully sharded data parallelism (FSDP) shards only the model states, allowing collectives to gather model states for the next layer to run off of the critical path. Intra-node FSDP (Zhao et al., 2023; Rajbhandari et al., 2020) is often been preferred to TP+SP on AMD Instinct nodes (AMD, n.d.e; Anthony et al., 2026). However, for large-scale distributed training, FSDP is often paired with TP+SP, which then shards activation memory across sequence, head, and channel dimensions, which dominates at scale.

2.1.3 Expert Parallelism

In an MoE layer with EP, the experts are distributed throughout the world size. A router GEMM matches k experts to each token, then an A2A dispatch sends activations to the selected experts' GEMMs. An A2A combine then returns the expert outputs to the token's original rank, with results folded by token. Because routing is decided per pass, transfer volume and destinations change every step, with skew towards GPUs hosting hot experts. The more skewed the data, the more completion times diverge.

2.2 MXFP8 Training

MXFP8 (Rouhani et al., 2023) stores each element in FP8 and shares one power-of-two (E8M0) scale across each block of 32 consecutive elements, so a tensor moves roughly half the bytes of BF16. A block is quantized in the contracting direction. In training, as each matrix and its transpose are used, we quantize both row-wise and column-wise. Requantizing one copy from the other compounds rounding error, so both must be quantized from the higher-precision source. MXFP8 has been shown to be as accurate in training as BF16 (Mishra et al., 2025), with twice the theoretical GEMM throughput on MI350-series GPUs.

2.3 Motivation and Design Requirements

For TP+SP, the common overlap strategy pipelines chunks of data through transport and compute on two streams. This is generally inefficient, because the two kernels compete for the same CUs without coordinating. In an eight-GPU node in the MI350-series, each GPU has a theoretical bidirectional peak of 1008 GB/s of bandwidth (AMD, 2025a), split between seven 144 GB/s XGMI links. A single peer is reachable at 1/7th of the aggregate. Full bandwidth usage requires sending to and receiving from all peers simultaneously. Host initiated direct memory access engines cannot provide this fan-out at fine granularity: They lag for small transfers because of their launch and synchronization cost, and they cannot perform reductions needed by TP+SP or EP. The tile-granular overlap therefore requires in-kernel communication. This leads to our first requirement: **R1: One kernel, one stream:** Communication and computation must share the GPU within one kernel.

In LLM training with power-capped GPUs, GEMM throughput decreases only slightly when CUs are capped. Both MI350X and MI355X (AMD, 2025a;b) run at their power limit of 1000W and 1400W, respectively, with all 256 CUs active, so when a GEMM runs on fewer CUs, the remaining

ones clock higher (Figure 3). On MI350X, our GEMMs keep an average 93% of their full-GPU throughput on 224 CUs (87.5% of the GPU), and 88% on 192 CUs (75% of the GPU) (Figure 2). Reserving CUs for communication, therefore, costs the GEMM less than the CU count suggests. **R2: Use the power headroom:** Collectives are power cheap, so CUs should be allocated to communication up to the point where the transport saturates.

Overlapping a collective with its GEMM inside one kernel means dividing the GPU's CUs between communication and computation. Most fused kernels make that division once, leading to windows where WGs sit idle on CUs. For AG, this occurs when the gather has finished before compute is complete. On RS, we have two windows: (1) where communicators wait for the first data to send, and (2) at the end when compute WGs have finished their tiles but the reductions are still occurring. In A2A, dispatches act much like AG with front-loaded communication, but combine traffic is interleaved with each batched GEMM, leaving windows scattered throughout the kernel. Figure 4 shows these windows in traces of fused kernels that keep a fixed split of communication WGs. Retuning the split may reduce idle time, but does not solve this structural problem, as communication demand changes as the fused kernel progresses. **R3: Communicators budgeted adaptively:** A fixed mover count leaves either matrix cores or interconnects idle, depending on the phase of the kernel. The number of active movers should follow the work that is available.

MoE layers are a chain of steps whose sizes change from iteration to iteration. Executed as separate kernels, every step requires host overhead, routing decisions must leave the device, and no step can overlap another. Additionally, WGs do not necessarily start in their ID order (Appendix Figure A.1). With cross-WG dependencies (R1, R3), a resident consumer can wait on a producer that has not started; if waiting WGs fill every CU, the producer never spawns and the kernel deadlocks. **R4: Persistent and resident:** Kernels must be persistent with all WGs resident. When CUs are needed for a second stream of work, kernels must adapt to a WG cap so the whole grid stays resident.

One key corollary follows. **R5: Self-scheduling:** While kernel design may guide WGs towards certain work, kernels cannot rely on a host round trip (breaks persistence, R4), a separate scheduling kernel (a second kernel, R1), or a scheduler WG (a CU lost to communication or compute, R2). Each WG must decide its own work only from the state it observes (adapting as demand changes, R3).

3 Overview of KOSMOS

3.1 KOSMOS Design

KOSMOS implements each fused operation as a single persistent kernel that launches one WG per CU, with each WG resident for the kernel's lifetime (R1, R4). Each WG starts in one of two roles: A **mover** transports data to and from remote peers, and a **builder** computes GEMM tiles. Depending on the collective, either role can be the producer or the consumer of the other. WGs coordinate only through in-kernel queues and readiness flags, without the host or a scheduler (R5). Rather than launching more WGs as work appears, KOSMOS keeps this fixed resident set and lets WGs change roles as the work changes (R3). Role switching is *one-way* or *two-way*. In one-way switching, a WG takes on another role for the rest of the kernel once its own work has finished. In AG and MoE dispatch, the movers become builders after their transfers are complete. In GEMM+RS, builders join the reduce once the GEMM tiles run out (Figure 1(A)). In two-way switching, a WG with no work available in its current role works in the other role *temporarily*, and returns when its own work appears: In the MoE combine, a set of WGs push any combine unit that is ready and builds GEMM tiles otherwise (Figure 1(B)).

3.1.1 Readiness and Queue Order

Without a scheduler, two things decide how a WG finds work: the unit at which data is declared ready, and the order of the queues it claims from. A WG claims a tile by atomically taking the next entry from a shared queue; it then computes that tile. KOSMOS operates in bands, chunks, tiles, and steps. A tile row is a set of tiles spanning 256 rows. A rank's chunk is the tokens/TP-size output rows it owns after RS. A band is a group of tile rows within a chunk, and each step produces every chunk's next band.

Work is tracked at the finest grain the GEMM allows. In gathering collectives, the unit is a 256-row tile row of gathered input. In reducing collectives, it is a band of rows that is sized to the output. Every output tile has exactly one completion point, so a consumer can start work on it as soon as it is finished and enqueued.

Two rules fix the order of work: **Upstream before downstream:** A WG claims every tile of an earlier step before any tile of a later one. A WG waits only on a tile claimed earlier in queue order, to avoid deadlocks or stalls. For example, in MoE forward, GEMM2 tiles are claimed only once every GEMM1 tile has been claimed (but not necessarily finished), so a waiting tile only ever depends on the work already claimed by a resident WG (Figure 1(B)). Additionally, a two-way WG never blocks on a push (as it builds instead), so waits are always satisfied. **Produce in the order consumers need:** If consumers are left waiting, we lose

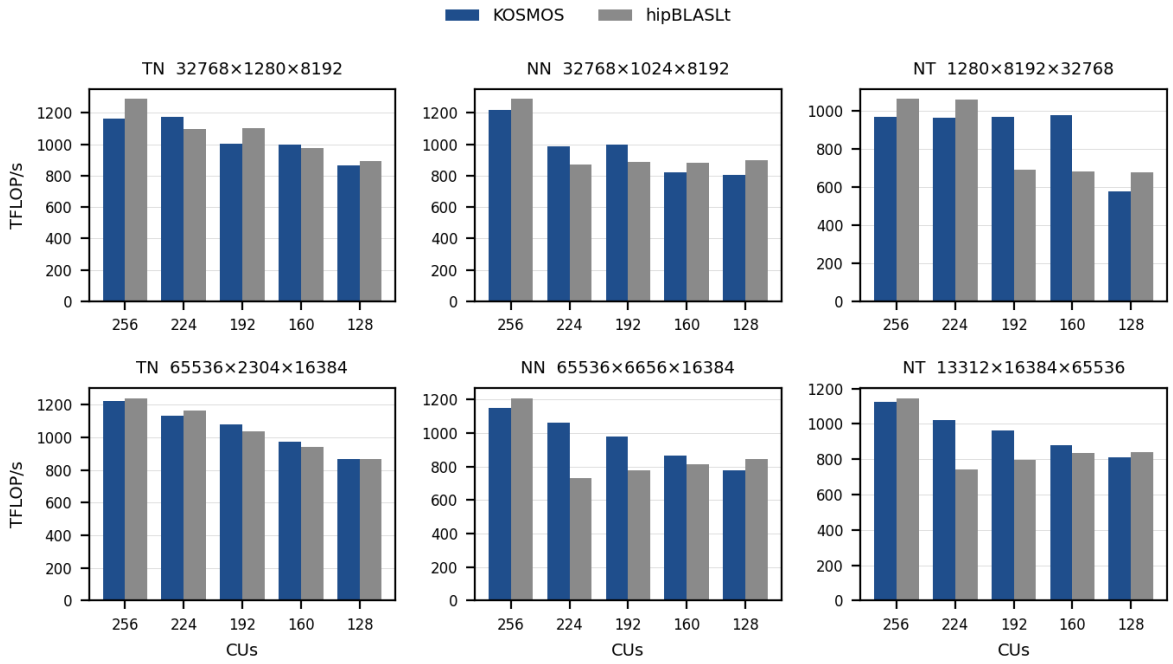

Figure 2. GEMM-only throughput (TFLOPS) of KOSMOS GEMM kernels and hipBLASLt under a CU cap, communication off. Top: small shapes (Llama 3 70B, MBS 4); bottom: large shapes (Llama 3 405B, MBS 8); TN, NN, and NT layouts, titled M×N×K. Measured on MI350X.

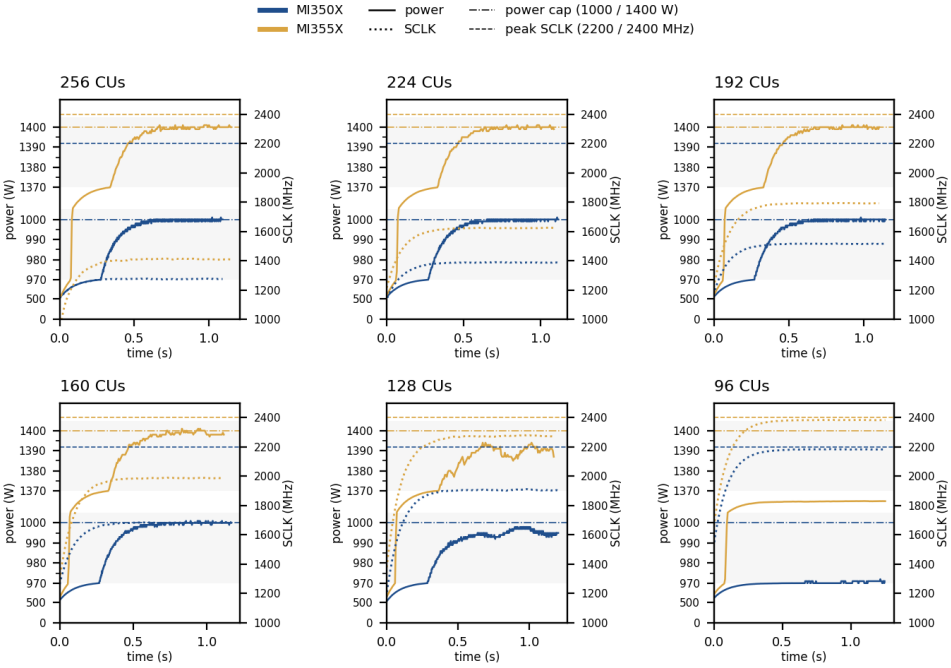

Figure 3. Power and shader clock (SCLK) values over a Llama 3 70B-FC1-MBS8 GEMM stress test of about 1.1–1.3 s. Both MI350X and MI355X remain at their 1000W and 1400W power caps until dropping to 128 CUs. Both GPUs reach their SCLK peak around 96 CUs.

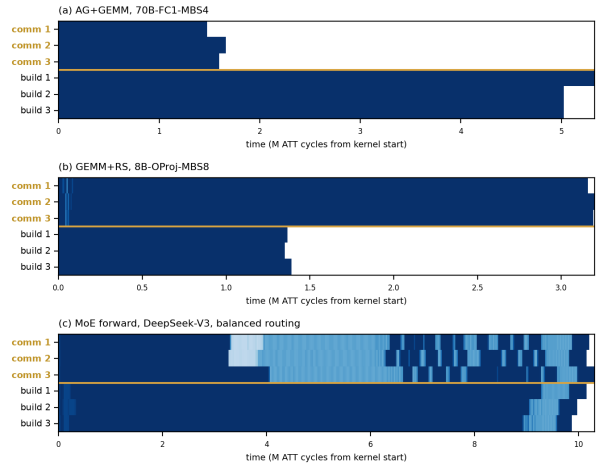

Figure 4. Advanced Thread Trace (ATT) timeline of three communication and three computation WGs for an AG+GEMM from Llama 70B FC1, with a microbatch size (MBS) of 4; a GEMM+RS from Llama 8B with MBS 8; and MoE forward for DeepSeek-V3 with balanced routing. *Dark Blue:* working. *Light Blue:* spinning on a flag or sleeping. *White:* exited.

overlap. In AG+GEMM, builders work on the local shard before claiming other work while movers gather the first shards. In GEMM+RS, builders must produce band-major; finishing the first band of every rank's chunk first allows every rank's reducers to start after one step. In MoE forward, dispatch pushes experts in the order GEMM1 consumes them, and GEMM2 works on a row block so the combine can push to its destination as soon as GEMM2 finishes it.

3.1.2 *Coexistence under a CU Cap*

When other work must run concurrently, such as inter-node FSDP, KOSMOS kernels can be CU capped by calling frameworks to allocate CUs for a second stream of work. Then, AG+GEMM and MoE kernels shrink the builder count and RS kernels shrink both roles proportionally. Every WG of the capped grid then remains resident (R4).

3.2 TP+SP Kernels

As shown in Figure 5, TP+SP requires six GEMMs and five or six collectives. KOSMOS names GEMMs by their layout in BLAS col-major convention: "T" for "transpose," "N" for "no transpose." For BF16, row-parallel wgrad is able to reuse the gathered dY from row-parallel dgrad. In MXFP8 training, dY^T is also gathered. To overlap the increased communication, both AGs are fused along with the pair of backward GEMMs for the row-parallel layer into a single kernel. Thus, KOSMOS supports five persistent kernels for both data types (Figure 5).

3.2.1 *Fused AG*

In the AG kernels (forward AG+TN, backward AG+NN), the movers pull the gathered input directly. Each mover is pinned by its WG index to one slice of one source GPU shard, and copies it one 256-row tile row at a time from the peer into a local buffer. The mover then performs a release-add on a local arrival counter for that peer and tile row. Builders compute the local shard first, then claim peer tiles in gather order. A builder spins on the counters in gather order and performs a single acquire once the count is complete. Once their slice has been transported, the movers perform a one-way switch to the builders.

In MXFP8, the row-parallel backward fuses one more GEMM: Training accuracy requires the weight gradient to use dY quantized along the token dimension, so a second column-wise quantized copy of dY must also be gathered. KOSMOS gathers both quantized copies of dY, and computes the data gradient (NN) and the weight gradient (NT) in one kernel to minimize idle time.

3.2.2 *Fused RS*

The RS kernel (forward TN+RS) uses a "one-shot" pull, derived from the two-shot all-reduce (Thakur et al., 2005). Builders write their partial tiles into a local staging buffer in band-major order. Movers own a band of the output; they pull it from every rank, sum in FP32 in place in a fixed order, and round once to BF16 before storing directly to the output. Every partial crosses the fabric once and is rounded once, which has less accumulation error than RCCL's ring, and is bitwise reproducible (Appendix Table A.6). Arrival is detected in the data itself: The staging buffer is filled with a NaN sentinel, so the landing of the data is the signal in and of itself. A BF16 signal NaN is used to avoid hangs that could arise from a matrix fused multiply-add-produced arithmetic NaN. The first builders (up to saturation) that find no GEMM tile left become movers to assist with the reduction (one-way switching).

3.2.3 *Bulk Kernels*

In the backward pass, two collectives have no data dependency on the GEMM they run beside. The AG of X is needed only by the weight gradient, so it can run alongside the data gradient GEMM, and the RS of dX is not needed by the weight gradient. This becomes a resource optimization problem, so KOSMOS defines a heuristic that uses GB/s/CU and TFLOPS/CU curves to balance WGs. KOSMOS then employs one-way role-switching for movers and builders in bulk fusions to minimize any remaining idle time.

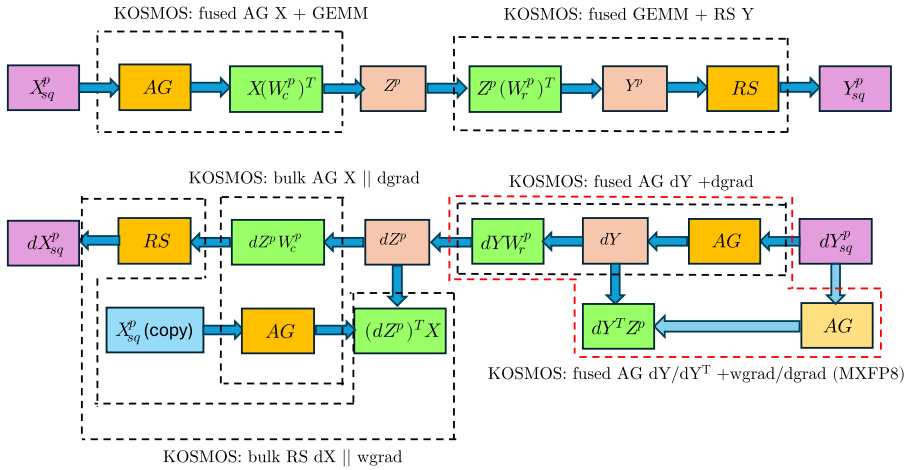

Figure 5. Overview of TP+SP kernel fusion in KOSMOS. Top: forward. Bottom: backward. Each dashed box is one KOSMOS kernel in BF16 or MXFP8; the red box is used only for MXFP8.

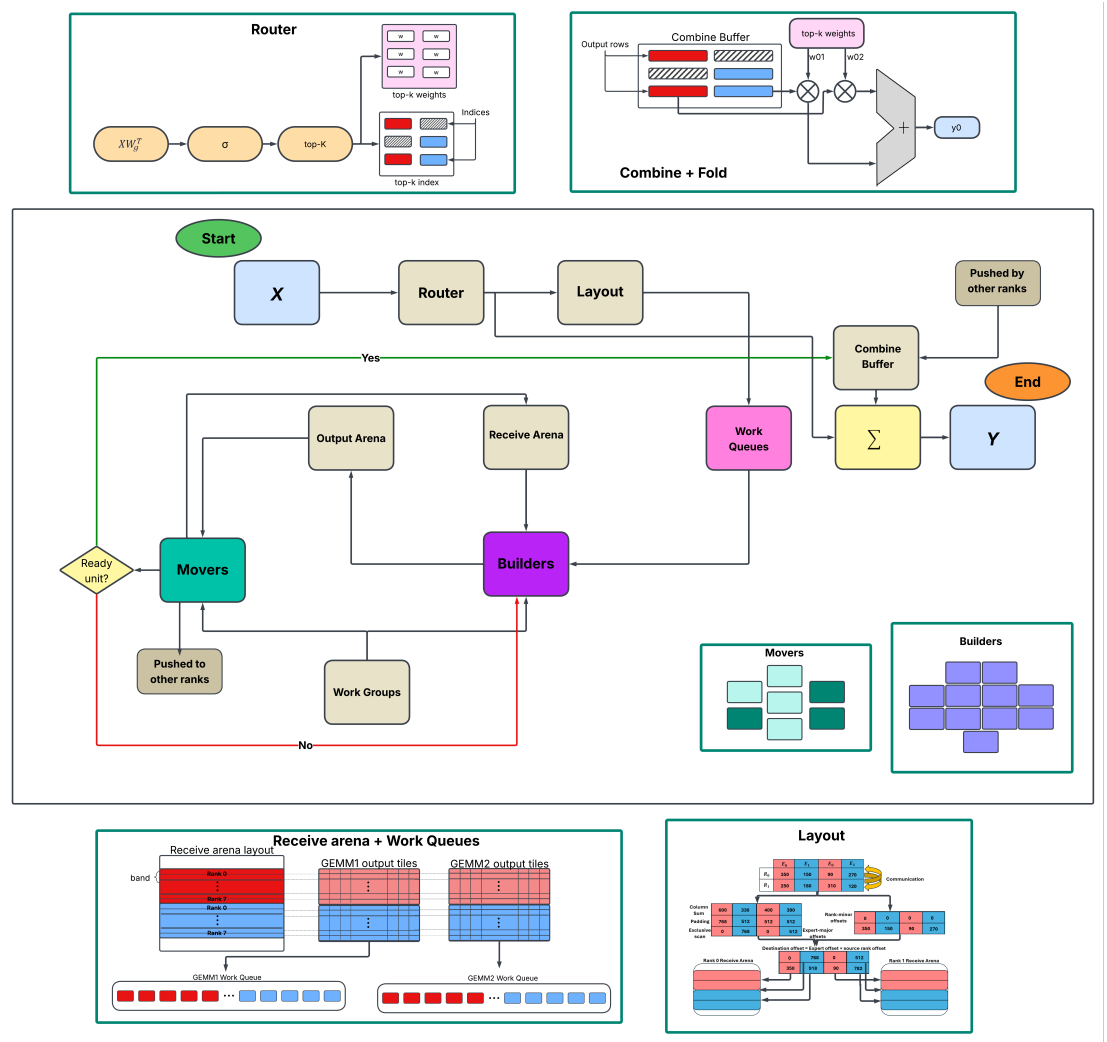

Figure 6. The KOSMOS single-launch MoE forward kernel. Red and blue: local experts; grey hatching: remote experts. Green: movers (light: dispatchers, dark: combiners); purple: builders. The flow runs from Start to End; stages are the forward-pass steps of Section 3.3.1.

3.3 MoE Megakernels

An MoE layer in EP routes each token to k experts. Then, tokens are dispatched to the GPUs that own those experts. The per-expert token counts change every step, so they are unknown until the layer runs. To avoid host overhead and enable overlap across steps, KOSMOS launches all of the steps below in a single kernel launch (Figure 6).

3.3.1 Forward Pass

1. **Route.** All WGs compute the router logits $XW_r^\top$ in FP32 tiles. The last tile to finish a row block of tokens routes that block: score, (group-limited) top-k, dispatch weights, and aux-loss inputs. The last block to be routed reduces the aux loss. The routing never leaves the device and does not need a separate phase.

2. **Plan.** Builders sort the routing into per-expert send lists with a stable counting sort, keeping dW bitwise reproducible, since row order sets the wgrad's accumulation order and MXFP8 scale groups. Only per-expert counts cross ranks. Each rank waits for all peers' counts, then builds the layout and tile queues on the device. This wait is the forward's only cross-rank barrier.

3. **Dispatch ‖ GEMM1.** A crew of 16-96 movers (Appendix Table B.3) pushes token rows directly into their final, expert-grouped position in each owner's receive arena, which at EP 8 is slightly faster than pulling (Appendix Table A.7). Push lists are destination-minor, so consecutive movers target different GPUs to keep every link busy. Per-(source, expert) flags at the destination track readiness. In MXFP8, mover quantizes their rows and pad each group to 32 rows before pushing. Each GEMM1 tile, with SwiGLU fused into its epilogue, waits only for the source whose rows it reads, and stores the pre-activations for the backward. In MXFP8, an epilogue quantizes the SwiGLU output row-wise for GEMM2 and column-wise for wgrad2.

4. **GEMM2 ‖ combine.** A GEMM2 tile starts as soon as GEMM1 has finished its row block. Up to 128 dual-role movers (Appendix Table B.3) push each (expert, destination) unit back to its token owner once the unit's row blocks are done, and claim GEMM tiles between pushes (two-way switching). The combine always pushes in BF16, even in MXFP8 training.

5. **Fold.** Once every row owed to this rank has arrived, each token's k rows are weighted by its FP32 router weights and summed in FP32.

3.3.2 Backward Pass

1. **Dispatch.** A crew of 32-64 movers (Appendix Table B.3) pushes the router-weighted dY to the expert owners. In MXFP8, dY is pushed in BF16 and quantized on the receiver in both the row-wise (dgrad2) and column-wise (wgrad2) layouts. wgrad1 also needs the column-wise X, which the movers push after dY in link time the backward leaves idle; for narrow experts where the backward is communication-bound, it is sent with the forward dispatch instead (Appendix Table B.3). After their pushes, the movers compute the routing-weight gradients $dw_j = \langle dY, Y_j \rangle$ from the forward's combine buffer and back-propagate it with the aux-loss gradient through the score function to the logit gradient.

2. **One queue.** dgrad2, with the SwiGLU backward fused into its epilogue, is interleaved with wgrad2 of the previous expert. The router's dX_r and dW_r tiles are queued after dgrad2/wgrad2 to hide the tail associated with the last experts to land.

3. **Combine ‖ wgrad1.** Up to 128 dual-role movers (Appendix Table B.3) push dX row blocks back to their token owners as dgrad1 completes them, while builders run wgrad1. In BF16, wgrad1 is interleaved into dgrad1; in MXFP8, wgrad1 uses the column-wise X, which arrives late when sent in the backward. All of dgrad1 runs first to cover its arrival; wgrad1 then overlaps the combine.

4. **Fold.** Each token's k gradient rows and the router's dX_r are summed in FP32 into dX.

3.4 Implementation Details

KOSMOS is written in C++20 and HIP using eight waves and 512 threads per WG; local data share and register usage are kept within single occupancy. GEMM tiles use HipKittens (Hu et al., 2026), a C++ domain-specific language for AMD GPUs, with 256x256 output tiles for all expert and TP+SP GEMMs; the router GEMM uses a 64x128 tile as it only has as many columns as experts. Shape-dependent parameters (mover crew size, transport parameters, GEMM parameters) are fixed with closed-form rules during the first training iteration (Appendix Table B.3).

4 EVALUATION

4.1 Experimental Setup

4.1.1 System Specifications

Two single-node systems were used, with eight GPUs (MI350X or MI355X (AMD, 2025a;b)) each, fully connected over XGMI. Both nodes contained two AMD Zen 5 CPUs with 64 cores each. While MI350X runs at a peak clock of 2.2 GHz at up to 1000W, MI355X runs at a peak of 2.4 GHz at up to 1400W. MI350X offers 2.31 PFLOPS

of peak dense BF16 matrix multiplication, while MI355X is slightly higher at 2.52 PFLOPS. The peak performance of MXFP8 is 4.62 PFLOPS and 5.03 PFLOPS, respectively.

The benchmarks were run on ROCm 7.14 (AMD, 2026) with RCCL 2.30.4 (AMD, n.d.c), PyTorch (Paszke et al., 2019) 2.12, ROCm Transformer Engine (AMD, n.d.f) 2.18, and ROCm Megatron-LM (AMD, n.d.d) commit b0695a0.

4.1.2 Integration

The microbenchmarks were run with a standalone harness in KOSMOS. For TP+SP, a RCCL collective and a GEMM are run serially as a baseline (hipBLASLt (AMD, n.d.a) in BF16, a communication-free HipKittens GEMM in MXFP8). Each configuration is run with 24 shapes arising from the Llama3 (Grattafiori et al., 2024) class of models. The 8B, 70B, and 405B model sizes are used with a micro-batch size between 1 and 8 and sequence length 8192. MoE runs one layer, both forward and backward, in a standalone harness. Layers of nine MoE models are used: DeepSeek V3 (DeepSeek-AI, 2025), DeepSeek V4 Pro (DeepSeek-AI, 2026), Qwen3-30B and Qwen3-235B (Yang et al., 2025), Qwen3.5-397B (Qwen Team, 2026), Kimi K2.7 Code (Moonshot AI, 2026), GLM 5.2 (Z.ai, 2026), and Llama 4 Scout and Maverick (Meta AI, 2025). The MoE models use sequence length 4096 and MBS 2. KOSMOS is tested against baselines with balanced and skewed routing to simulate different training scenarios.

End-to-end runs are executed using ROCm Megatron-LM. KOSMOS TP+SP kernels are transplanted into TE and run against Megatron-LM using the default ROCm TE path. The 405B model is restricted to 24 layers to fit on a single node. For MoE, the KOSMOS adapter replaces Megatron's router, token dispatcher, grouped multilayer perceptron (MLP), and shared expert if present. Due to resource constraints, a proxy Qwen-235B model with 24 layers and a DeepSeek-V3 proxy with no dense layers and nine MoE layers are considered. TP+SP models use a global batch size (GBS) of 64, and MoE models use a GBS of 16.

4.1.3 Numerics & Validation

KOSMOS kernels are bitwise run-to-run deterministic: tile claiming changes which WG computes a tile, never the reduction or fold order. Outputs are validated against RCCL+hipBLASLt with ROCm TE's tolerances (AMD, n.d.f), re-poisoned every iteration, and checked twice to catch races. Microbenchmarks run and discard 10 warm-up iterations before 25 timed iterations. Each is run three times with the variant order shuffled, and we report the median of the slowest rank; the median run-to-run spread is 0.4% (max 2.0%). End-to-end results report Megatron-LM's printed iteration time and throughput. Each configuration is run three times for 20 iterations, with the median result presented. MoE TFLOPS count actual routed FLOPs.

4.2 Ablations

4.2.1 Data Readiness Mechanisms

Table 1 isolates the design decisions behind KOSMOS, changing one mechanism at a time with everything else held fixed. The largest effects come from the order in which data is produced and consumed: local-first AG queues consumed in gather order, band-major RS production, and destination-minor A2A push lists. KOSMOS is optimized for bandwidth per CU: its 8-wave movers saturate XGMI links with half the CUs of a 4-wave kernel (Appendix Table A.4). Fusing the router and shared expert into the MoE launch is a small but consistent gain. Readiness-driven switching only pays off if readiness arrives where WGs look next: Thus, the order of production is as important as the dynamic role switching it enables.

4.2.2 Role-Switching Policy

To measure the impact of KOSMOS' derived role switching policies, we compare them to several alternatives (Table 2). For TP+SP, as implied by the one-way relationship of AG/RS and their GEMMs, we found that one-way role switching overall outperformed two-way role switching. One-way role switching is never more than 1% slower than the best single fixed split, and up to 1.15× faster. For AG, communication-bound shapes generally do not benefit from role switching, as movers are active until the very last step: While MLP layers improve by a geometric mean of 1.05×, attention layers average 1.00×. Per-shape results for both GPUs are in Appendix Figures A.2 and A.3.

For MoE, we show the benefits of our two-way role switch in KOSMOS over alternatives. KOSMOS uses an interleaved two-way role switching policy, where WGs decide their role after every tile. We compare this to a policy where dispatch is executed with one-way role switching and combine uses a dedicated WG crew, following UniEP (Zheng et al., 2026), and to a static two-way planned steal, following Weave (Huang et al., 2026), where movers finish dispatch, compute a predetermined number of tiles from a runtime cost model, then return to combine. We find the two-way role switch not only beats a fixed split in all 36 cells on both MI350X and MI355X, but is the best policy on average, with a geometric mean speedup of 1.04×/1.04× on the forward pass and 1.06×/1.08× on the backward pass (MI350X/MI355X). This gain is up to 1.14×/1.16×. Appendix Figures A.4 and A.5 show the speedups per shape on the MI350X and MI355X, respectively.

Table 1. Design mechanisms ablated in isolation: slowdown of the alternative relative to KOSMOS. Full results in Appendix A.

MECHANISM	ALTERNATIVE	ALTERNATIVE VS KOSMOS
AG QUEUE ORDER	NO LOCAL-FIRST, CHUNK-MAJOR	1.17–1.40× SLOWER
RS PRODUCTION ORDER	ROW-MAJOR	1.41–1.45× SLOWER
A2A PUSH ORDER	DESTINATION-MAJOR	1.12–2.03× SLOWER
8-WAVE MOVERS	4 WAVES PER WG (7–56 WGS)	1.23–1.96× SLOWER
ROUTER FUSION	SEPARATE ROUTER KERNELS	1.01–1.02× SLOWER
SHARED-EXPERT FUSION	SEPARATE KERNELS	1.02–1.04× SLOWER

Table 2. Role-policy ablation: geometric-mean speedup over a fixed communication crew. TP+SP: 24 shapes per kernel, fixed crew of 56 WGs. MoE: nine models, fixed crew is the measured peak per shape, also used by the dedicated combine crew (following UniEP (Zheng et al., 2026)). Per-step planned split following Weave (Huang et al., 2026). All policies are implemented in KOSMOS. *Shipped KOSMOS policy; bold: fastest per row and GPU.

	MI350X				MI355X			
	One-way switch	Two-way switch	Dedicated combine	Planned split	One-way switch	Two-way switch	Dedicated combine	Planned split
TP+SP AG+TN	**1.032***	0.965	–	–	**1.033***	0.932	–	–
TP+SP TN+RS	**1.031***	1.024	–	–	**1.018***	1.012	–	–
TP+SP AG+NN	**1.018***	0.951	–	–	**1.017***	0.918	–	–
MoE fwd, balanced	–	**1.029***	1.018	0.985	–	**1.027***	1.008	0.982
MoE fwd, skewed	–	**1.047***	1.029	0.989	–	**1.051***	1.019	0.914
MoE bwd, balanced	–	**1.052***	1.012	0.973	–	**1.067***	1.009	0.836
MoE bwd, skewed	–	**1.064***	1.021	0.963	–	**1.093***	1.009	0.752

4.3 TP+SP Speedups

4.3.1 TP+SP Microbenchmarks

Figure 7 compares each KOSMOS kernel to the same RCCL collective and hipBLASLt/HipKittens GEMM run without overlap. KOSMOS is faster on all 480 combinations of shape, kernel, precision, and GPU. In BF16, we see a geometric mean speedup of 1.38× on MI350X and 1.41× on MI355X. In MXFP8, the geometric mean is 1.36× on MI350X and 1.39× on MI355X. The MXFP8 GEMMs have twice the peak throughput of BF16, while an MXFP8 all-gather moves 33/64 of the BF16 bytes and the reduce-scatters move BF16 in both precisions, which leaves less compute to hide. Compared against unfused KOSMOS communication and GEMM kernels run without overlap gives similar geometric mean speedups: For BF16, the speedup is 1.39×/1.42×, and for MXFP8 the speedup is 1.37×/1.40×. This shows that most of the benefit over the baselines is due to the overlap mechanisms, rather than the components.

4.3.2 TP+SP End-to-End Training

Table 3 reports Llama 3 training in Megatron with TP 8, where KOSMOS replaces Transformer Engine's TP+SP collectives and GEMMs. KOSMOS is faster for every model size, precision, and GPU model.

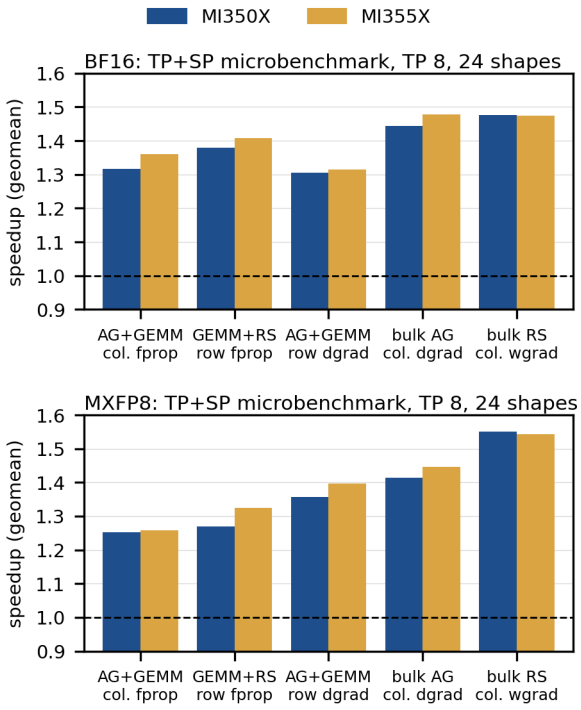

Figure 7. TP+SP microbenchmark: KOSMOS speedup over the baseline, geometric mean over 24 Llama 3 shapes per config.

Table 3. Llama 3 TP+SP end-to-end training, TP 8, sequence 8192, GBS 64: KOSMOS speedup over ROCm Transformer Engine (range over MBS 1–8, same MBS in both arms) and TFLOPS/GPU at each arm's best MBS (TE → KOSMOS). [r]Full recompute in both arms; 70B at MBS 4 and 8 runs all 80 layers, and 405B runs 24 of its 126 layers. Per-MBS results: Appendix Table C.1.

		MI350X		MI355X	
Model	Prec.	Speedup	TFLOPS	Speedup	TFLOPS
8B	BF16	1.20–1.27	474→596	1.24–1.28	478→612
	MXFP8	1.21–1.25	580→718	1.22–1.24	597→730
70B	BF16	1.15–1.21	664→769	1.22–1.28	654→834
	MXFP8	1.19–1.22	893→1067	1.21–1.23	907→1119
405B	BF16	1.10–1.16	570[r]→645[r]	1.17–1.22	624[r]→757[r]
	MXFP8	1.10–1.14	862[r]→985[r]	1.16–1.18	947[r]→1101[r]

4.4 Mixture-of-Experts Speedups

4.4.1 *MoE Microbenchmarks*

Figure 8 compares one KOSMOS MoE layer against RCCL A2A with TE's grouped GEMMs. The figure covers nine models, both forward and backward, with balanced and skewed routing. KOSMOS is faster in all 144 cases, 1.28–3.39×. Gains shrink as GEMM compute increases and top-k decreases: Qwen3-30B (narrow experts, top-8) gains up to 3.39×, while Llama 4 (wide experts, top-1) gains 1.28–1.92×. The forward pass with balanced routing gains the most, and backward pass with skewed routing is sped up the least. We also run KOSMOS components as separate launches; they are 1.25–1.51× slower than the fused layer on geometric mean which we attribute to kernel fusion.

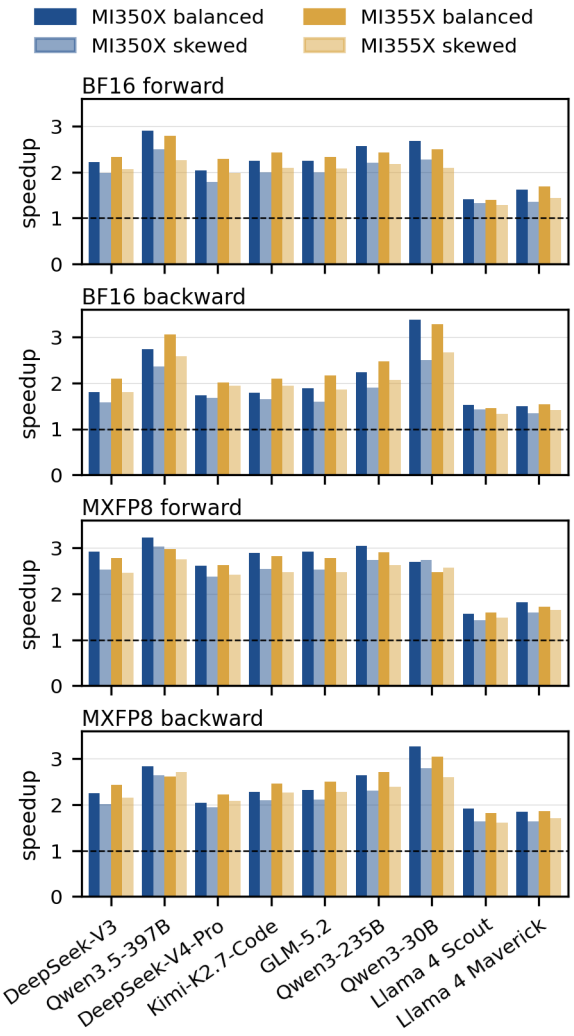

Figure 8. MoE layer speedups over baselines for KOSMOS. Bars are per-model speedup for each direction, precision, and routing.

4.4.2 *MoE Proxy Training*

Table 4 reports end-to-end training of two MoE proxy models. For both, balanced refers to random routing with a uniform distribution, so each expert has an equal expected number of tokens. Skewed results refer to a pure top-k routing defined by each model. For DeepSeek V3 with Megatron-generated mock data, this was seen to be up to 1.3× the data of balanced for the largest peer, whereas Qwen saw skew up to 4.2× in some iterations. KOSMOS relies on Megatron's capacity factor flag to size shared buffers accordingly. A greater skew falls back to the default Megatron MoE layer. KOSMOS is faster than Megatron-LM with ROCm's DeepEP (Zhao et al., 2025a; AMD, 2026) in every configuration, by 1.29–1.83× across GPU models.

5 Related Work

Stream-level overlap. Frameworks (Li et al., 2020; NVIDIA, n.d.b) run NCCL/RCCL (NVIDIA, n.d.a; AMD, n.d.c) collectives on separate streams, where they hold CUs for the whole transfer and compete with the GEMM. CoCoNet (Jangda et al., 2022) adds device-side locks, chunking (Wang et al., 2022; Liang et al., 2025) shrinks the GEMMs, and copy engines (Zheng et al., 2025b) or new hardware (Pati et al., 2024) cannot reduce or signal per tile.

Fused communication kernels. FLUX (Chang et al., 2024) fuses RS into a GEMM epilogue and AG into its prologue; ParallelKittens (Sul et al., 2026a), TileLink (Zheng et al., 2025a), and Iris (Awad et al., 2025) fuse TP+SP collectives per tile with a tuned number of communication streaming multiprocessors (SMs) or WGs. For MoE, COMET (Zhang et al., 2025), Mixture-of-Kittens (Sul et al., 2026b), and the MegaMoE kernels of DeepSeek-V4 (Zhao et al., 2025b; DeepSeek-AI, 2026) and Primus-Turbo (AMD, n.d.b; Zhang et al., 2026) fix communication resources at build time or per shape, before routing is known; Mixture-of-Kittens (a

Table 4. MoE proxy training, EP 8: KOSMOS speedup over DeepEP, and TFLOPS/GPU (DeepEP → KOSMOS). Balanced: uniform random routing; skewed: the model's top-k router. Full results: Appendix Table C.2.

		MI350X		MI355X	
Routing	Prec.	Speedup	TFLOPS	Speedup	TFLOPS
Qwen3-235B					
balanced	BF16	1.70	359→610	1.74	387→673
	MXFP8	1.82	372→677	1.83	400→731
skewed	BF16	1.31	249→327	1.29	275→355
	MXFP8	1.57	245→384	1.54	268→411
DeepSeek-V3					
balanced	BF16	1.31	423→556	1.33	480→637
	MXFP8	1.39	446→619	1.39	500→696
skewed	BF16	1.33	414→551	1.36	465→631
	MXFP8	1.40	437→613	1.41	488→689

second backward megakernel) and Primus-Turbo (fused operators reused in backward) compute the router outside the kernel. FlashMoE (Aimuyo et al., 2025) and MPK (Cheng et al., 2026) are persistent inference kernels driven by scheduler warps. FLUX, ParallelKittens, COMET, and Mixture-of-Kittens are NVIDIA-only, so Section 4.2 instead reproduces the UniEP and Weave policies inside the KOSMOS kernels.

Adaptive communication budgets. UniEP (Zheng et al., 2026) orders one MoE task queue communication-first, with a split set by a cost model and tuner; Weave (Huang et al., 2026) sizes its communication SMs with an in-kernel cost model and lets them steal a precomputed number of GEMM tiles; mKernel (Mao et al., 2026) adapts the split with an in-kernel controller but trails its best fixed split. In KOSMOS, roles follow data readiness alone, with no scheduler, task graph, per-call cost model, or steal budget, and one launch per direction covers every TP+SP collective with its GEMM, and the whole MoE layer, router and gradients included.

6 Conclusion

Current state-of-the-art overlaps communication in TP+SP and EP training by running it on a separate stream, or by giving it a fixed or predicted share of the GPU inside a fused kernel. We show that on CDNA4 GPUs, both approaches are sub-optimal. Concurrent streams may contend without coordination, and a fixed split strands CUs whenever the demand for communication changes. KOSMOS addresses both by running communication and computation in one self-scheduling persistent kernel per operation or layer. Its WGs switch between moving and building as the work available changes, with in-kernel queues rather than a separate scheduler. Because collectives are limited by CUs up to saturation, giving CUs to communication costs little GEMM throughput, if the kernel adapts. On MI350X and MI355X nodes, KOSMOS speeds up dense model training by up to 1.28× (BF16) and 1.25× (MXFP8), and MoE training by up to 1.74× (BF16) and 1.83× (MXFP8) over DeepEP.

KOSMOS kernels are persistent only for intranode transports: XGMI with hipIPC memory is currently required. Remote direct memory access support and likely a hierarchical set of movers would carry the design to multinode. Thus, end-to-end results are limited to a single-node proxy model. We have evaluated with quasi-FSDP and pipeline parallelism stress tests, but complete multinode verification is future work. Finally, our structure requires one completion point per output tile. This disqualifies GEMM algorithms, such as stream-K (Osama et al., 2023), that compute a single output tile using multiple WGs. Other workloads built on fine-grained collectives, such as all-reduce as used in TP-only training and inference; SUMMA (Van De Geijn & Watts, 1997) and sparse distributed GEMMs; 3D fast Fourier transforms and spectral methods; and stencil solvers could use the same role-switching kernel design.

Overlapping collectives with computation is no longer a question of how many CUs to allocate to communication, but of when. KOSMOS answers this with self-scheduling kernels that produce in the order consumers need, and lets data readiness drive role switches.

References

Aimuyo, O., Oh, B., and Singh, R. FlashMoE: Fast distributed MoE in a single kernel. In Belgrave, D., Zhang, C., Lin, H., Pascanu, R., Koniusz, P., Ghassemi, M., and Chen, N. (eds.), *Advances in Neural Information Processing Systems*, volume 38, Main Conference, pp. 100676–100699. Curran Associates, Inc., 2025. doi: 10.52202/085713-3368. URL https://proceedings.neurips.cc/paper_files/paper/2025/file/918d938bd209e5b56072777366f8a211-Paper-Conference.pdf.

AMD. AMD Instinct MI350X GPU datasheet. https://www.amd.com/content/dam/amd/en/documents/instinct-tech-docs/product-briefs/amd-instinct-mi350x-gpu-brochure.pdf, 2025a. Product page: https://www.amd.com/en/products/accelerators/instinct/mi350/mi350x.html. Accessed 2026-10.

AMD. AMD Instinct MI355X GPU datasheet. https://www.amd.com/content/dam/amd/en/documents/instinct-tech-docs/product-briefs/amd-instinct-mi355x-gpu-brochure.pdf, 2025b. Product page: https://www.amd.com/en/products/accelerators/instinct/mi350/mi355x.html. Accessed 2026-10.

AMD. AMD ROCm Software, including RCCL and hipBLASLt. https://rocm.docs.amd.com/, 2026. RCCL: https://github.com/ROCm/rocm-systems/tree/develop/projects/rccl; hipBLASLt: https://github.com/ROCm/rocm-libraries/tree/develop/projects/hipblaslt. Accessed 2026-10.

AMD. hipBLASLt. https://github.com/ROCm/rocm-libraries/tree/develop/projects/hipblaslt, n.d.a. Docs: https://rocm.docs.amd.com/projects/hipBLASLt/en/latest/. Accessed 2026-10.

AMD. Mega MoE (Primus-Turbo documentation). https://github.com/AMD-AGI/Primus-Turbo/blob/main/docs/README_Mega_MoE.md, n.d.b. Accessed 2026-10.

AMD. RCCL: ROCm Communication Collectives Library. https://github.com/ROCm/rocm-systems/tree/develop/projects/rccl, n.d.c. Docs: https://rocm.docs.amd.com/projects/rccl/en/latest/. Accessed 2026-10.

AMD. Megatron-LM for ROCm. https://github.com/ROCm/Megatron-LM, n.d.d. Accessed 2026-10.

AMD. Training a model with Megatron-LM for ROCm. https://rocm.docs.amd.com/en/docs-6.4.1/how-to/rocm-for-ai/training/benchmark-docker/megatron-lm.html, n.d.e. ROCm 6.4.1 documentation. Accessed 2026-10.

AMD. Transformer Engine on ROCm and AMDGPU. https://github.com/ROCm/TransformerEngine, n.d.f. Accessed 2026-10.

Anthony, Q. et al. Training foundation models on a full-stack AMD platform: Compute, networking, and system design. In *ISC High Performance 2026 Research Paper Proceedings (41st International Conference)*, pp. 1–18, 2026. doi: 10.23919/ISC.2026.11520488.

Awad, M., Osama, M., and Potter, B. Iris: First-class Multi-GPU programming experience in Triton, 2025. URL https://arxiv.org/abs/2511.12500.

Chang, L.-W. et al. FLUX: Fast software-based communication overlap on GPUs through kernel fusion, 2024. URL https://arxiv.org/abs/2406.06858.

Cheng, X. et al. MPK: A compiler and runtime for Mega-Kernelizing tensor programs. In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pp. 1909–1926, Seattle, WA, July 2026. USENIX Association. ISBN 978-1-939133-55-7. URL https://www.usenix.org/conference/osdi26/presentation/cheng.

DeepSeek-AI. DeepSeek-V3 technical report, 2025. URL https://arxiv.org/abs/2412.19437.

DeepSeek-AI. DeepSeek-V4: Towards highly efficient million-token context intelligence, 2026. URL https://arxiv.org/abs/2606.19348.

Grattafiori, A. et al. The Llama 3 herd of models, 2024. URL https://arxiv.org/abs/2407.21783.

Hu, W. et al. HipKittens: Fast and furious AMD kernels. In Chowdhery, A. and Jia, Z. (eds.), *Proceedings of Machine Learning and Systems*, volume 8, pp. 1431–1443. MLSys, 2026. URL https://proceedings.mlsys.org/paper_files/paper/2026/file/bc75fa9843a7905bbed9d83895a88f7f-Paper-Conference.pdf.

Huang, Z. et al. Weave: Fine-grained dynamic SM scheduling in an MoE megakernel for compute-communication overlap, 2026. URL https://arxiv.org/abs/2609.21483.

Jangda, A. et al. Breaking the computation and communication abstraction barrier in distributed machine learning workloads. In *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, ASPLOS '22, pp. 402–416, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450392051. doi: 10.1145/3503222.3507778. URL https://doi.org/10.1145/3503222.3507778.

Korthikanti, V. A. et al. Reducing activation recomputation in large Transformer models. In Song, D., Carbin, M., and Chen, T. (eds.), *Proceedings of Machine Learning and Systems*, volume 5, pp. 341–353. Curan, 2023. URL https://proceedings.mlsys.org/paper_files/paper/2023/file/80083951326cf5b35e5100260d64ed81-Paper-mlsys2023.pdf.

Li, S. et al. PyTorch distributed: Experiences on accelerating data parallel training, 2020. URL https://arxiv.org/abs/2006.15704.

Liang, W. et al. TorchTitan: One-stop PyTorch native solution for production ready LLM pretraining. In Yue, Y., Garg, A., Peng, N., Sha, F., and Yu, R. (eds.), *International Conference on Learning Representations*, volume 2025, pp. 92161–92181, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/file/e6231c5f46598cfd09ff1970524e0436-Paper-Conference.pdf.

Mao, Z. et al. mKernel: Fast Multi-GPU, multi-node fused kernels, 2026. URL https://arxiv.org/abs/2609.13585.

Meta AI. The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation. https://ai.meta.com/blog/llama-4-multimodal-intelligence/, 2025. Accessed 2026-10.

Mishra, A., Stosic, D., Layton, S., and Micikevicius, P. Recipes for pre-training LLMs with MXFP8, 2025. URL https://arxiv.org/abs/2506.08027.

Moonshot AI. Kimi K2.7 Code. https://huggingface.co/moonshotai/Kimi-K2.7-Code, 2026. Model card. Accessed 2026-10.

NVIDIA. NVIDIA Collective Communications Library (NCCL). https://github.com/NVIDIA/nccl, n.d.a. Accessed 2026-10.

NVIDIA. Transformer Engine. https://github.com/NVIDIA/TransformerEngine, n.d.b. Accessed 2026-10.

Osama, M., Merrill, D., Cecka, C., Garland, M., and Owens, J. D. Stream-K: Work-centric parallel decomposition for dense matrix-matrix multiplication on the GPU. In *Proceedings of the 28th ACM SIGPLAN Annual Symposium on Principles and Practice of Parallel Programming*, PPoPP '23, pp. 429–431, New York, NY, USA, 2023. Association for Computing Machinery. ISBN 9798400700156. doi: 10.1145/3572848.3577479. URL https://doi.org/10.1145/3572848.3577479.

Paszke, A. et al. PyTorch: An imperative style, high-performance deep learning library. In Wallach, H., Larochelle, H., Beygelzimer, A., d'Alché-Buc, F., Fox, E., and Garnett, R. (eds.), *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL https://proceedings.neurips.cc/paper_files/paper/2019/file/bdbca288fee7f92f2bfa9f7012727740-Paper.pdf.

Pati, S., Aga, S., Islam, M., Jayasena, N., and Sinclair, M. D. T3: Transparent tracking & triggering for fine-grained overlap of compute & collectives. In *Proceedings of the 29th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 2*, ASPLOS '24, pp. 1146–1164, New York, NY, USA, 2024. Association for Computing Machinery. ISBN 9798400703850. doi: 10.1145/3620665.3640410. URL https://doi.org/10.1145/3620665.3640410.

Qwen Team. Qwen3.5: Towards native multimodal agents. https://qwen.ai/blog?id=qwen3.5, 2026. Accessed 2026-10.

Rajbhandari, S., Rasley, J., Ruwase, O., and He, Y. ZeRO: Memory optimizations toward training trillion parameter models. In *SC20: International Conference for High Performance Computing, Networking, Storage and Analysis*, pp. 1–16, 2020. doi: 10.1109/SC41405.2020.00024.

Rouhani, B. D. et al. Microscaling data formats for deep learning, 2023. URL https://arxiv.org/abs/2310.10537.

Shoeybi, M., Patwary, M., Puri, R., LeGresley, P., Casper, J., and Catanzaro, B. Megatron-LM: Training multi-billion parameter language models using model parallelism, 2020. URL https://arxiv.org/abs/1909.08053.

Sul, S. H., Arora, S., Spector, B. F., and Ré, C. ParallelKittens: Systematic and practical simplification of Multi-GPU AI kernels. In Chowdhery, A. and Jia, Z. (eds.), *Proceedings of Machine Learning and Systems*, volume 8, pp. 1076–1089. MLSys, 2026a. URL https://proceedings.mlsys.org/paper_files/paper/2026/file/ff997469ac66cf893c4183efeb22212a-Paper-Conference.pdf.

Sul, S. H., Brown, N., Wildermuth, H., Lin, W., Cassano, F., and Ré, C. Mixture-of-Kittens: MoE megakernel for NVL72s, 2026b. URL https://arxiv.org/abs/2609.36070.

Thakur, R., Rabenseifner, R., and Gropp, W. Optimization of collective communication operations in MPICH. *The International Journal of High Performance Computing Applications*, 19(1):49–66, 2005. doi: 10.1177/1094342005051521. URL https://doi.org/10.1177/1094342005051521.

Van De Geijn, R. A. and Watts, J. SUMMA: Scalable universal matrix multiplication algorithm. *Concurrency: Practice and Experience*, 9(4):255–274, 1997.

Wang, S. et al. Overlap communication with dependent computation via decomposition in large deep learning models. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 1*, ASPLOS 2023, pp. 93–106, New York, NY, USA, 2022. Association for Computing Machinery. ISBN 9781450399159. doi: 10.1145/3567955.3567959. URL https://doi.org/10.1145/3567955.3567959.

Yang, A. et al. Qwen3 technical report, 2025. URL https://arxiv.org/abs/2505.09388.

Z.ai. GLM-5.2: Built for long-horizon tasks. https://z.ai/blog/glm-5.2, 2026. Accessed 2026-10.

Zhang, L. et al. Enabling DeepSeek-V4-Flash training on AMD Instinct MI355X GPUs with Primus. https://rocm.blogs.amd.com/software-tools-optimization/primus-deepseek-v4/README.html, 2026. ROCm Blogs. Accessed 2026-10.

Zhang, S. et al. COMET: Fine-grained computation-communication overlapping for Mixture-of-Experts. In Zaharia, M., Joshi, G., and Lin, Y. (eds.), *Proceedings of Machine Learning and Systems*, volume 7. MLSys, 2025. URL https://proceedings.mlsys.org/paper_files/paper/2025/file/e27ea0cd50b798ff8942caf9203f0992-Paper-Conference.pdf.

Zhao, C. et al. DeepEP: an efficient expert-parallel communication library. https://github.com/deepseek-ai/DeepEP, 2025a.

Zhao, C. et al. DeepGEMM: clean and efficient BLAS kernel library on GPU. https://github.com/deepseek-ai/DeepGEMM, 2025b.

Zhao, Y. et al. PyTorch FSDP: Experiences on scaling fully sharded data parallel, 2023. URL https://arxiv.org/abs/2304.11277.

Zheng, S., Zheng, X., Chang, L.-W., and Zhai, J. UniEP: Unified expert-parallel MegaKernel MoE for LLM training. In *Proceedings of the 35th International Symposium on High-Performance Parallel and Distributed Computing*, HPDC '26, pp. 387–401, New York, NY, USA, 2026. Association for Computing Machinery. ISBN 9798400726408. doi: 10.1145/3806645.3807818. URL https://doi.org/10.1145/3806645.3807818.

Zheng, S. et al. TileLink: Generating efficient compute-communication overlapping kernels using tile-centric primitives. In Zaharia, M., Joshi, G., and Lin, Y. (eds.), *Proceedings of Machine Learning and Systems*, volume 7. MLSys, 2025a. URL https://proceedings.mlsys.org/paper_files/paper/2025/file/c6ee784cbe46d854843e4c883a3321ef-Paper-Conference.pdf.

Zheng, S. et al. Triton-distributed: Programming overlapping kernels on distributed AI systems with the Triton compiler, 2025b. URL https://arxiv.org/abs/2504.19442.

A ABLATIONS

A.1 WGs out of Order

While workgroups of a persistent grid are all resident, they do not start in ID order. The dispatcher has tendencies to fill sorted by XCD rather than ID, but this is not guaranteed either. With oversubscription, the order depends on per-WG work. Variable work per WG can lead to highly unpredictable start order.

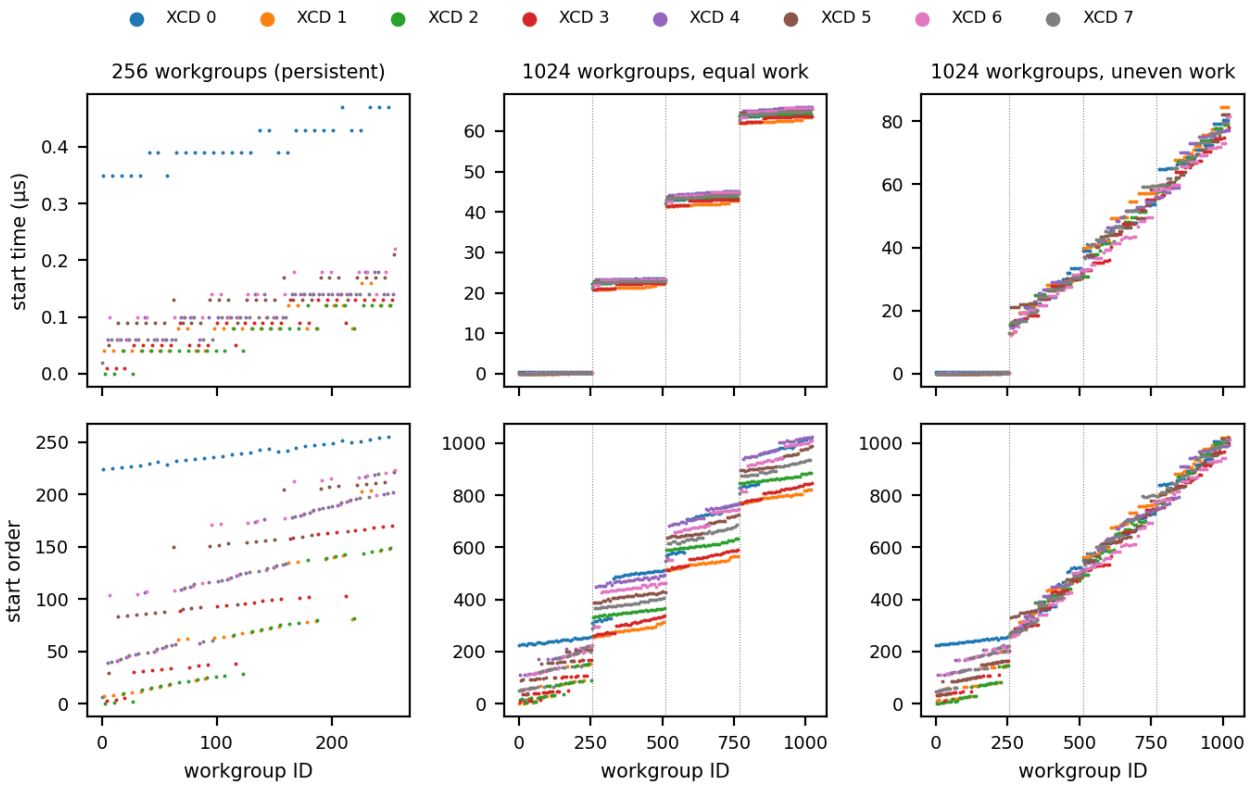

Figure A.1. Workgroup launch order and time sorted plotted against WG ID number. For persistent kernels, all WGs are resident but order is not guaranteed, and instead leans towards an accelerator complex die (XCD) ID-sorted launch order. With oversubscription, 1024 WGs exhibit different behavior based on the workload assigned to each WG.

A.2 Role Switching Policies

These figures give the per-shape results behind Table 2. For TP+SP, one-way switching is never significantly slower than the best fixed split, but up to 15% faster. For MoE, two-way switching beats a fixed split in all 72 cells, and is faster than a dedicated split in 63 of 72 cells and within 2% in the rest. Planned splits appear to lose where routing is skewed within a step, as a fixed crew size is determined before data arrives.

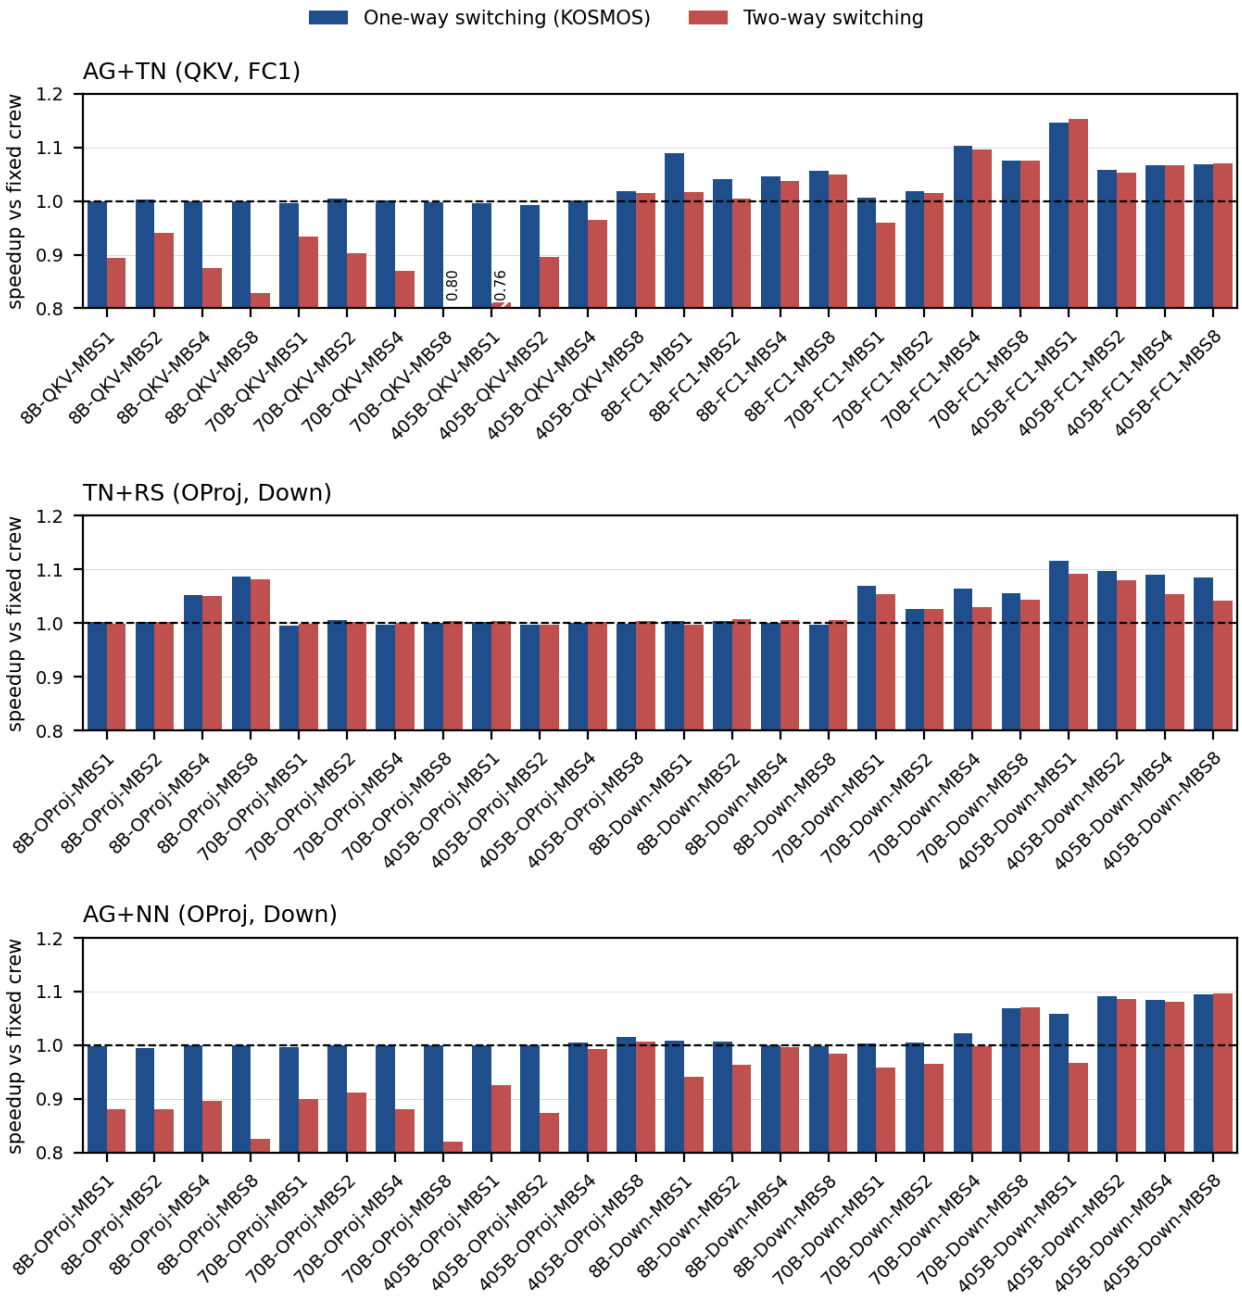

Figure A.2. MI350X: speedup of one-way and two-way role switching over a best single fixed split of 56 communication WGs (8 per peer) for the three fused formats in training: AG+TN GEMM, TN GEMM+RS, and AG+NN GEMM. Switching arms use the KOSMOS heuristic crew (Table B.3).

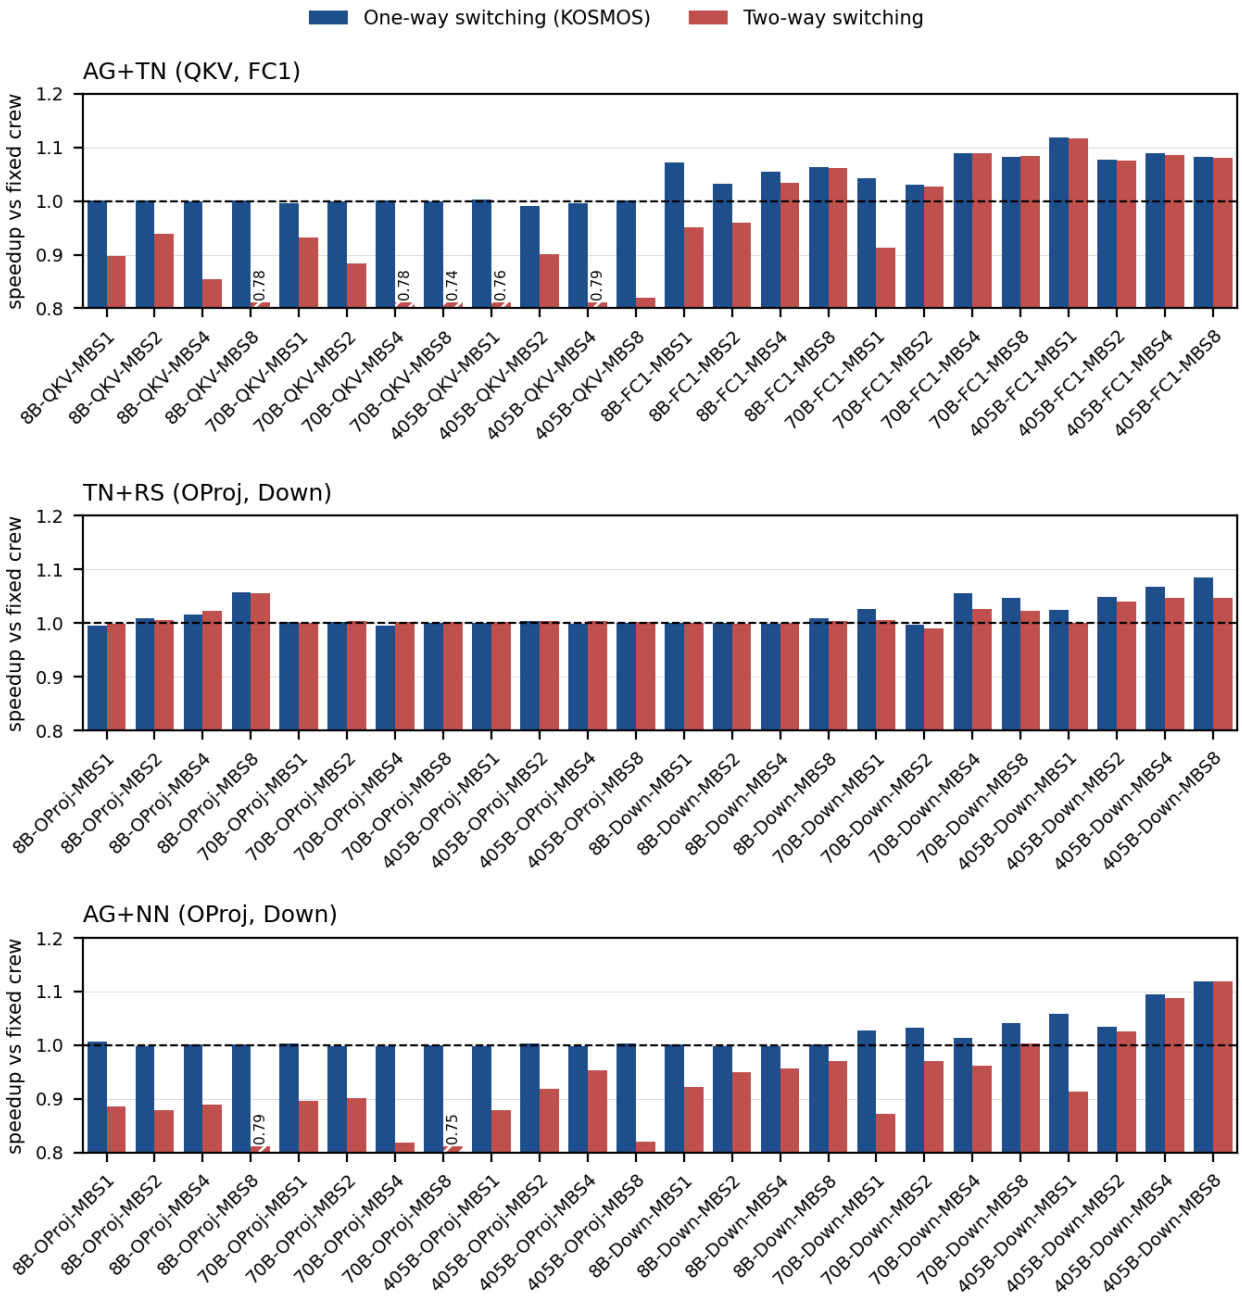

Figure A.3. MI355X: speedup of one-way and two-way role switching over a best single fixed split of 56 communication WGs (8 per peer) for the three fused formats in training: AG+TN GEMM, TN GEMM+RS, and AG+NN GEMM. Switching arms use the KOSMOS heuristic crew (Table B.3).

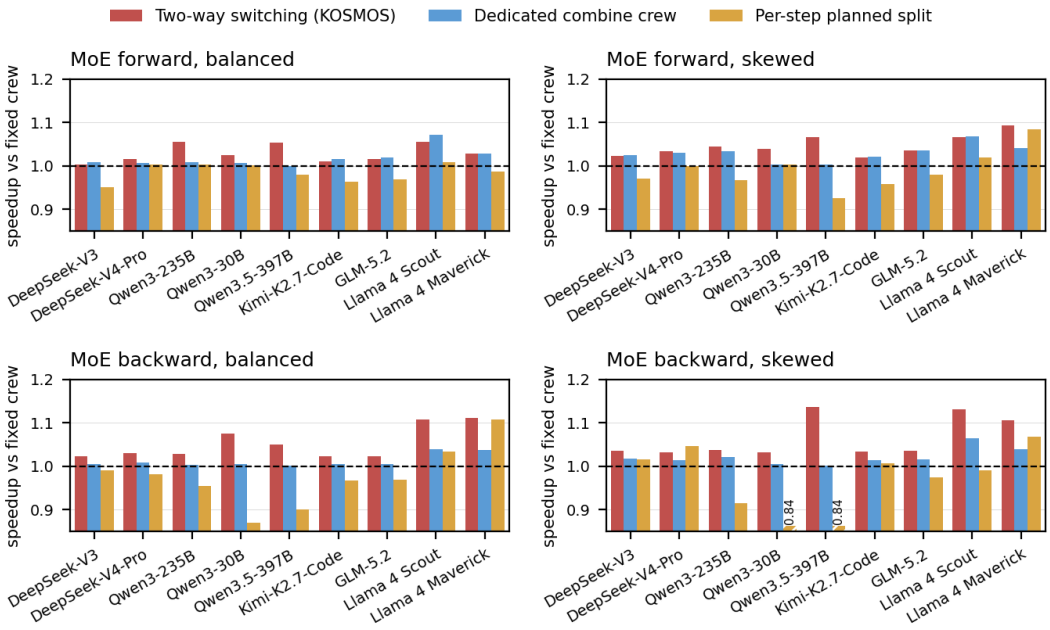

Figure A.4. MI350X: speedup over a fixed split of different role-switching policies in MoE. KOSMOS two-way role switching with data readiness role switching, a UniEP-style one-way role switching for dispatch with a dedicated combine crew, and a Weave-style planned split. Crew size for non-KOSMOS policies is the smallest that reaches 98% of the peak dispatch and combine bandwidth (24, 32, or 48 WGs)

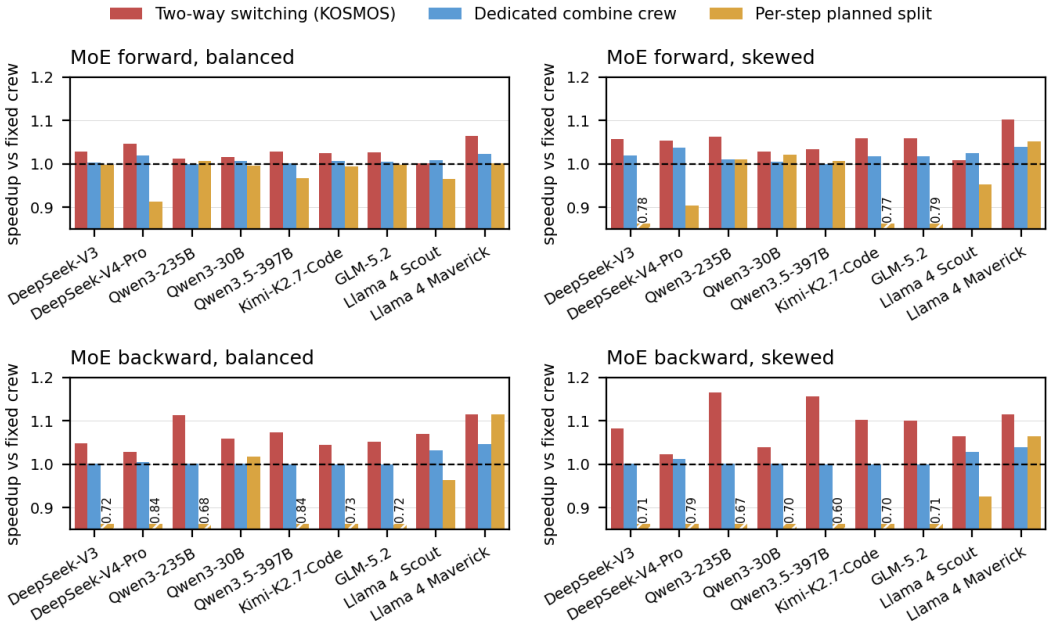

Figure A.5. MI355X: speedup over a fixed split of different role-switching policies in MoE. KOSMOS two-way role switching with data readiness role switching, a UniEP-style one-way role switching for dispatch with a dedicated combine crew, and a Weave-style planned split. Crew size for non-KOSMOS policies is the smallest that reaches 98% of the peak dispatch and combine bandwidth (16, 24, 32, or 48 WGs)

A.3 Fused vs Separate Router and Shared Expert

Running the router or the shared expert as separate kernels is 1.01–1.02× and 1.02–1.04× slower, respectively, than the fused launch (Table A.1).

Table A.1. Fusing the router and the shared expert into the MoE launch. *Fused* (KOSMOS): one launch. *Separate*: The router (GEMM, top-k, aux loss) or the shared expert runs as its own kernels (hipBLASLt for the shared-expert GEMMs) on the same stream. *Without*: the same MoE layer with routing precomputed (router rows) or routed experts only (shared-expert rows).

Component	Model	Pass	Fused (ms)	Separate (ms)	Without (ms)	Speedup
Router	dsv3	forward	**7.557**	7.612	7.168	1.007
Router	dsv3	backward	**13.773**	14.005	13.711	1.017
Router	qwen235	forward	**3.622**	3.679	3.398	1.016
Router	qwen235	backward	**5.808**	5.918	5.752	1.019
Shared expert	dsv3	forward	**7.655**	7.990	7.218	1.044
Shared expert	dsv3	backward	**15.533**	15.797	13.865	1.017

A.4 Local First AG Queue vs Naive Ordering

A naive rank-order queue without local-first is 1.17–1.40× slower than the shipped queue (Table A.2).

Table A.2. AG consumption order, configs 1 (AG+TN) and 3 (AG+NN), TP=8, atomic tile claim (queue order is the schedule). *Shipped*: local shard first, then peer tiles in the order movers gather them. *Chunk-major*: local first, peers chunk by chunk (mismatches the gather order). *Naive*: chunk-major in rank order, no local-first. Time: slowest rank, median of 3 runs × 25 calls; slowdown vs. shipped in parentheses.

Config	Shape	Shipped (ms)	Chunk-major	Naive
TN	70B-FC1-MBS8	**7.579**	8.040 (+6.1%)	8.873 (+17.1%)
TN	405B-QKV-MBS8	**5.552**	7.314 (+31.7%)	7.742 (+39.4%)
NN	70B-Down-MBS8	**4.038**	4.821 (+19.4%)	5.109 (+26.5%)
NN	405B-OProj-MBS8	**5.184**	7.022 (+35.5%)	7.240 (+39.7%)

A.5 Band-Major vs. Row-Major RS Queue

Row-major production is 1.41–1.45× slower than band-major production (Table A.3).

Table A.3. RS production order, config 2 (GEMM → RS), TP=8. Band-major (shipped) produces band b of every rank's chunk before band b+1, so reducers start on the first band; row-major produces each chunk completely before the next.

Shape	Band-major (ms)	Row-major (ms)	Row / band
8B-OProj-MBS8	**1.366**	1.930	1.41×
70B-Down-MBS8	**4.156**	6.038	1.45×

A.6 8 vs. 4 Waves

With four waves per WG, transport is 1.23–1.96× slower measured by bytes/CU/s at 7–56 WGs (Table A.4).

Table A.4. Mover waves vs. request depth, 1 GiB transfer (70B-MBS8), TP=8. The shipped 8-wave mover keeps one 16-byte remote load in flight per wave (two waves per single instruction, multiple data unit). 4-wave variants: same loop (depth 1), two loads in flight per wave (depth 2), or load/store interleave (pipelined, AG only). (Bold: fastest)

Op	WGs	8 waves (shipped)	4 waves	4 waves, depth 2	4 waves, pipelined
AG	7	**9.772**	19.131	10.525	19.460
AG	14	**5.056**	9.791	5.451	9.920
AG	28	**2.734**	5.044	2.899	5.063
AG	56	**2.217**	2.727	2.220	2.729
AG	112	2.278	**2.230**	2.273	2.237
RS	7	**15.900**	30.378	18.490	–
RS	14	**7.988**	15.205	9.263	–
RS	28	**4.068**	7.627	4.683	–
RS	56	**2.288**	3.878	2.513	–
RS	112	2.325	**2.236**	2.259	–

A.7 Dst-Minor vs. Dst-Major Push

Destination-major push lists are 1.12–2.03× slower than destination-minor lists (Table A.5).

Table A.5. Push order for EP transport, dsv3, EP=8, routed experts only (router and shared expert excluded). *Destination-minor* (KOSMOS): Consecutive WGs push to different destinations, one destination per WG. *Destination-major*: Consecutive WGs push to the same destination. Same kernels and data.

Routing	Phase	Destination-minor (ms)	Destination-major (ms)	Slowdown
Balanced	Forward	**8.148**	16.532	2.03×
Balanced	Backward	**13.789**	20.567	1.49×
Skewed	Forward	**18.686**	24.841	1.33×
Skewed	Backward	**30.077**	33.653	1.12×

A.8 RS Error vs. RCCL Ring, against a FP64 Reference

While KOSMOS accumulates all eight partials in FP32 and rounds to BF16 or MXFP8 once, ring algorithms often keep BF16 partials and round at every hop. The result is 2.21× lower reduction error for KOSMOS, and 1.71-1.73× lower error against a FP64 GEMM. KOSMOS is also bitwise run-to-run deterministic.

Table A.6. RS error, TP=8. Every arm reduces the same BF16 partials. *Reduction*: vs. the FP64 sum of the partials. *end-to-end*: vs. an FP64 GEMM (32,768-element sample).

	Reduction			end-to-end		
Shape	KOSMOS	RCCL	Ratio	KOSMOS	RCCL	Ratio
8B-OProj-MBS8	1.67×10^{-3}	3.68×10^{-3}	2.21	2.34×10^{-3}	4.04×10^{-3}	1.73
70B-Down-MBS8	1.67×10^{-3}	3.68×10^{-3}	2.21	2.35×10^{-3}	4.02×10^{-3}	1.71

A.9 Push vs Pull

While pushing rows into the destination's arena is slightly faster than pulling them this requires more memory than a pull algorithm. As such, for larger EP sizes it may be required to use a pull algorithm instead.

Table A.7. Push vs. pull EP transport in the full MoE layer (EP=8). Push (KOSMOS): The producer writes rows into the destination's buffer. Pull: The consumer reads them from the source in the forward dispatch; combine and the backward push in both arms, so the backward rows are a control. Speedup: pull over push time.

Model	Routing	Phase	Push (ms)	Pull (ms)	Speedup
dsv3	balanced	forward	**8.117**	8.158	1.005
dsv3	balanced	backward	13.783	**13.725**	0.996
dsv3	balanced	step	21.043	**21.025**	0.999
dsv3	skewed	forward	**18.628**	19.004	1.020
dsv3	skewed	backward	30.055	**30.052**	1.000
dsv3	skewed	step	**45.144**	45.506	1.008
qwen235	balanced	forward	**3.783**	3.849	1.017
qwen235	balanced	backward	5.710	**5.699**	0.998
qwen235	balanced	step	**9.276**	9.322	1.005

B ADDITIONAL KOSMOS DETAILS

B.1 Fused Communication Framework Comparison

Table B.1. Fused communication–computation systems. *Scope*: what one launch fuses (a collective–GEMM *pair* or a whole MoE *layer*). *Crew sized by*: how the number of communication units is chosen. *Roles/Switch trigger*: whether and when a unit changes between communication and computation. *Rtr/wgrad*: router and weight gradients inside the launch.

System	Scope	MoE launches Train	MoE launches F/B	Crew sized by	Roles	Switch trigger	Switches /unit	Rtr	wgrad	TP+SP
FLUX (Chang et al., 2024)	pair	✓	–	–	own tile	–	0	–	–	✓
TileLink (Zheng et al., 2025a)	pair	✗	–	launch config	fixed	–	0	–	–	✓
ParallelKittens (Sul et al., 2026a)	pair	✗	–	launch config	fixed	–	0	✗	✗	✓
Iris (Awad et al., 2025)	pair	✗	–	launch config	fixed	–	0	–	–	△
COMET (Zhang et al., 2025)	pair	△	∼2 / ?	offline, per shape	fixed	–	0	✗	?	✗
Primus MegaMoE (AMD, n.d.b)	pair	✓	3 / 5[a]	autotuned per shape	fixed	–	0	✗	part.[a]	✗
UniEP (Zheng et al., 2026)	pair	✓	3 / 6[a]	perf-model tuner	one-way	comm tasks exhausted, queue	1	✗	✗[a]	✗
FlashMoE (Aimuyo et al., 2025)	layer†	✗	1 / –	scheduler	sched. tasks	scheduler warp	per task	✓	–	✗
DeepSeek MegaMoE (Zhao et al., 2025b)	layer†	✗	1+aux / –	launch config	fixed warps	–	0	✗	–	✗
Weave (Huang et al., 2026)	layer†	✗	1+aux / –	prologue cost model[b]	planned phases	precomputed steal count[c]	$\leq$2	✗	–	✗
Mixture-of-Kittens (Sul et al., 2026b)	layer	✓	1+aux / 1+aux	launch config	fixed	–	0	✗[d]	✓	✗
MPK (Cheng et al., 2026)	model	✗	1 / –	scheduler	sched. tasks	scheduler SMs	per task	✓	–	✗
KOSMOS	pair + layer	✓	1 / 1	shape model	one-way + two-way	**readiness at every claim**[e]	unbounded	✓[f]	✓	✓

†Forward pass only (inference). [a]From public code. [b]Per layer and per GPU, from routing counts and calibrated bandwidth/throughput curves; fixed for the rest of the kernel. [c]Communication SMs dispatch, steal a precomputed number of GEMM tiles, then combine; all SMs join the final combine. [d]Router runs before the launch; routing-score gradient fused. [e]Movers hand off when their share is done (AG, dispatch, bulk); combine movers check for ready work before every tile claim, in forward and backward. [f]Router GEMM, top-k and aux loss forward; router gradients (logits, dX, dW_r) backward.

B.2 Supported TP+SP Configs

Table B.2. TP+SP KOSMOS kernels per collective-GEMM pair (†: MXFP8 only).

Pass	Collective	GEMM	Kernel
fwd	AG X	col-parallel	fused AG+TN
fwd	RS Y	row-parallel	fused TN+RS
bwd	AG dY	row-parallel dgrad	fused AG+NN
bwd	AG dY, AG $dY^\top$ †	row-parallel dgrad + wgrad	fused AG+NN(+NT) †
bwd	AG X	col-parallel dgrad	bulk AG ‖ NN
bwd	RS dX	col-parallel wgrad	bulk RS ‖ NT

B.3 Heuristics

Table B.3. Per-shape heuristics, evaluated once per shape on the host. Rules were fitted at TP 8 and EP 8; at other TP sizes mover count is scaled. Mover counts are per remote peer; MoE crews are WGs per GPU. All counts are then fitted to the CU cap so that every mover runs beside at least one builder.

Parameter	Kernel	Input	Rule
Movers per peer	AG+GEMM, col. fprop	gathered width N	BF16: 6 if N is at least 3584, else 8. MXFP8: 8.
	AG+GEMM, row dgrad	gathered width N	BF16: 6 if N is at least 4608, else 8. MXFP8: 8.
	bulk AG	contraction K	BF16: 2, 4, 6 if K is at least 7168, 3584, 2304; else 8. MXFP8: 2, 4 if K is at least 3584, 1792; else 8.
Reducers per peer	GEMM+RS, row fprop	local K, waves w (tiles per cap)	K at least 4096: 4 (MXFP8 5). K at least 3072: 6 (MXFP8 8 if w is at least 16, else 7). K at most 768: 14, 10, 8 if w is at least 16, at least 8, or less. Otherwise 8 (MXFP8 10 if w is at least 16).
	bulk RS	output width n	BF16: 2, 4 if n is at least 8192, 3584; else 8. MXFP8: 8 if n is 5120 to 8191, else 4.
Dispatch crew	MoE forward	expert width I	BF16: 32, 24, 16 if I is below 1024, below 3072, or larger. MXFP8: 24576/I rounded to a multiple of 8 between 16 and 48, doubled if the forward also sends the column-wise input. Minimum over ranks.
Combine crew	MoE forward	CU cap	The smaller of 128 and the cap, rounded down to a multiple of 8.
Dispatch crew	MoE backward	total bytes, total GEMM work	256 WGs times push time over push plus GEMM time, rounded up to a multiple of 8, between 32 and 64.
Combine crew	MoE backward	total bytes, total GEMM work	256 WGs times push time over GEMM time, rounded up to a multiple of 8, between 24 and 128 (MXFP8 32 and 128).
Push burst depth	MoE dispatch, combine	hidden size H	Largest of 8, 7, 6, 4 that divides H/512.
Split-K factor	AG+GEMM row dgrad, bulk AG	output tiles	BF16: 4 or 2 if the FP32 partials fit the workspace (at most 64 MiB), else 1. MXFP8: 1.
	bulk RS	output tiles	BF16: 4 if at most 64 tiles and the partials fit the workspace, else 1. MXFP8: 1.
XCD bucketing	AG+GEMM (all)	tiles, tile rows	On if at most 1024 tiles and at least 5 tile rows.

C FULL PER-SHAPE AND PER-GPU TABLES BEHIND THE SUMMARY FIGURES

Table C.1. Llama 3 TP+SP end-to-end training, TP 8, sequence 8192, GBS 64: median over three runs of Megatron-LM's printed numbers (each the mean over iterations 3–19). Speedup is TE baseline ms/iter over KOSMOS ms/iter. OOM: out of memory. [r]Full recompute in both arms. Layer-reduced proxies: The 405B model runs 24 of its 126 layers with full recompute. 70B at MBS 4 and 8 runs with full recompute. Three runs per point, 20 iterations each.

			BF16				MXFP8			
Model	MBS	Layers	TE baseline ms/iter	KOSMOS ms/iter	KOSMOS TFLOPS	Speedup	TE baseline ms/iter	KOSMOS ms/iter	KOSMOS TFLOPS	Speedup
MI350X										
8B	1	32	10056.2	7913.1	426.6	1.271	8708.3	6967.7	484.9	1.250
8B	2	32	7727.6	6435.7	524.4	1.201	6729.2	5454.8	619.1	1.234
8B	4	32	7119.4	5811.7	580.7	1.225	5939.6	4902.3	688.5	1.212
8B	8	32	7117.2	5658.0	596.4	1.258	5813.0	4702.4	717.6	1.236
70B	1	80	47371.0	41279.4	713.3	1.148	37514.8	30841.8	954.7	1.216
70B	2	80	44371.2	38308.3	768.6	1.158	32956.2	27603.6	1066.6	1.194
70B	4	80[r]	62232.3	51524.3	571.4	1.208	43348.7	36146.8	814.5	1.199
70B	8	80[r]	62377.7	51415.7	572.6	1.213	43643.7	36258.4	812.0	1.204
405B	1	24[r]	56475.7	51236.1	628.1	1.102	37706.9	34361.3	936.5	1.097
405B	2	24[r]	57459.5	49923.4	644.6	1.151	37438.2	32840.9	979.9	1.140
405B	4	24[r]	57912.5	49899.2	644.9	1.161	37347.7	32671.8	985.0	1.143
405B	8	24[r]	58049.1	49998.7	643.6	1.161	37655.0	OOM	–	–
MI355X										
8B	1	32	9996.1	7915.3	427.0	1.263	8647.1	6971.3	485.5	1.240
8B	2	32	7876.5	6351.8	531.3	1.240	6664.8	5368.7	628.9	1.241
8B	4	32	7220.7	5709.9	591.0	1.265	5893.8	4833.4	698.5	1.219
8B	8	32	7059.0	5515.6	611.8	1.280	5649.1	4619.6	730.5	1.223
70B	1	80	48695.2	39918.9	737.6	1.220	37191.5	30146.4	976.7	1.234
70B	2	80	45029.7	35307.9	833.9	1.275	32447.1	26317.1	1118.8	1.233
70B	4	80[r]	59942.3	47738.9	616.7	1.256	42303.4	34680.9	848.9	1.220
70B	8	80[r]	58645.8	46949.1	627.1	1.249	41509.8	34187.4	861.2	1.214
405B	1	24[r]	54176.8	46203.8	696.5	1.173	37137.2	31882.3	1009.4	1.165
405B	2	24[r]	52786.9	43727.3	735.9	1.207	35251.2	29995.2	1072.9	1.175
405B	4	24[r]	52032.3	42743.1	752.9	1.217	34474.2	29407.8	1094.3	1.172
405B	8	24[r]	51565.3	42513.7	756.9	1.213	33989.8	29226.6	1101.1	1.163

Table C.2. MoE proxy training, EP 8: median over three runs of Megatron-LM's printed numbers, 20 iterations each. KOSMOS speedup is the DeepEP ms/iter over KOSMOS's. DeepEP: Megatron-LM with ROCm's DeepEP. Balanced: uniform random routing; skewed: the model's top-k router.

			BF16			MXFP8		
Model	Routing	Backend	ms/iter	TFLOPS	KOSMOS speedup	ms/iter	TFLOPS	KOSMOS speedup
MI350X								
Qwen3-235B	balanced	**KOSMOS**	545.1	610	–	491.4	677	–
		DeepEP	925.6	359	1.698	894.7	372	1.821
Qwen3-235B	skewed	**KOSMOS**	1022.6	327	–	868.1	384	–
		DeepEP	1337.7	249	1.308	1359.2	245	1.566
DeepSeek-V3	balanced	**KOSMOS**	612.7	556	–	550.1	619	–
		DeepEP	804.8	423	1.314	762.8	446	1.387
DeepSeek-V3	skewed	**KOSMOS**	617.1	551	–	555.4	613	–
		DeepEP	822.7	414	1.333	778.5	437	1.402
MI355X								
Qwen3-235B	balanced	**KOSMOS**	494.3	673	–	454.9	731	–
		DeepEP	859.7	387	1.739	831.1	400	1.827
Qwen3-235B	skewed	**KOSMOS**	941.6	355	–	810.5	411	–
		DeepEP	1211.8	275	1.287	1244.2	268	1.535
DeepSeek-V3	balanced	**KOSMOS**	534.3	637	–	488.7	696	–
		DeepEP	709.7	480	1.328	680.0	500	1.391
DeepSeek-V3	skewed	**KOSMOS**	539.7	631	–	493.8	689	–
		DeepEP	732.1	465	1.356	697.9	488	1.413

Table C.3. TP+SP microbenchmark, AG+GEMM, column fprop, TP 8, MI350X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-FC1-MBS1	8192	13312	16384	3.994	3.320	1.203	1.188	1.903	1.651	1.152	1.100
405B-QKV-MBS1	8192	2304	16384	1.128	0.926	1.218	1.366	0.740	0.521	1.420	1.343
70B-FC1-MBS1	8192	7168	8192	1.091	0.995	1.097	1.072	0.613	0.539	1.137	1.134
70B-QKV-MBS1	8192	1280	8192	0.516	0.469	1.101	1.028	0.308	0.281	1.097	1.044
8B-FC1-MBS1	8192	3584	4096	0.502	0.310	1.619	1.241	0.259	0.209	1.242	1.180
8B-QKV-MBS1	8192	768	4096	0.377	0.234	1.611	1.041	0.192	0.165	1.167	1.040
405B-FC1-MBS2	16384	13312	16384	8.108	6.908	1.174	1.137	3.943	3.476	1.134	1.086
405B-QKV-MBS2	16384	2304	16384	2.170	1.569	1.383	1.434	1.349	0.883	1.529	1.413
70B-FC1-MBS2	16384	7168	8192	1.963	1.734	1.132	1.165	1.092	0.966	1.132	1.181
70B-QKV-MBS2	16384	1280	8192	0.993	0.756	1.313	1.262	0.583	0.442	1.318	1.240
8B-FC1-MBS2	16384	3584	4096	0.766	0.544	1.409	1.334	0.444	0.346	1.282	1.265
8B-QKV-MBS2	16384	768	4096	0.468	0.398	1.174	1.036	0.269	0.248	1.083	1.048
405B-FC1-MBS4	32768	13312	16384	16.169	13.525	1.196	1.123	8.418	7.455	1.129	1.067
405B-QKV-MBS4	32768	2304	16384	4.268	2.832	1.507	1.515	2.493	1.589	1.568	1.463
70B-FC1-MBS4	32768	7168	8192	4.664	3.530	1.321	1.237	2.174	1.894	1.148	1.190
70B-QKV-MBS4	32768	1280	8192	1.832	1.353	1.354	1.288	1.066	0.786	1.357	1.277
8B-FC1-MBS4	32768	3584	4096	1.348	0.938	1.436	1.460	0.822	0.619	1.328	1.295
8B-QKV-MBS4	32768	768	4096	0.873	0.693	1.260	1.194	0.489	0.411	1.189	1.177
405B-FC1-MBS8	65536	13312	16384	31.274	26.651	1.173	1.131	16.921	14.983	1.129	1.085
405B-QKV-MBS8	65536	2304	16384	9.082	5.578	1.628	1.612	4.993	3.116	1.603	1.485
70B-FC1-MBS8	65536	7168	8192	10.035	7.678	1.307	1.257	5.006	4.238	1.181	1.159
70B-QKV-MBS8	65536	1280	8192	3.528	2.532	1.394	1.344	2.074	1.479	1.402	1.317
8B-FC1-MBS8	65536	3584	4096	2.751	1.829	1.505	1.536	1.621	1.225	1.324	1.280
8B-QKV-MBS8	65536	768	4096	1.668	1.289	1.294	1.219	0.919	0.766	1.200	1.199
geometric mean						*1.316*	*1.249*			*1.252*	*1.205*

Table C.4. TP+SP microbenchmark, AG+GEMM, column fprop, TP 8, MI355X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-FC1-MBS1	8192	13312	16384	3.705	2.929	1.265	1.239	1.958	1.625	1.205	1.202
405B-QKV-MBS1	8192	2304	16384	1.200	0.959	1.252	1.431	0.754	0.524	1.438	1.419
70B-FC1-MBS1	8192	7168	8192	1.183	0.958	1.235	1.261	0.656	0.532	1.233	1.251
70B-QKV-MBS1	8192	1280	8192	0.540	0.486	1.113	1.031	0.308	0.291	1.059	1.051
8B-FC1-MBS1	8192	3584	4096	0.499	0.308	1.621	1.328	0.258	0.210	1.229	1.192
8B-QKV-MBS1	8192	768	4096	0.337	0.243	1.386	1.042	0.190	0.169	1.125	1.030
405B-FC1-MBS2	16384	13312	16384	7.427	5.848	1.270	1.208	3.735	3.176	1.176	1.159
405B-QKV-MBS2	16384	2304	16384	2.335	1.639	1.424	1.544	1.340	0.910	1.472	1.491
70B-FC1-MBS2	16384	7168	8192	2.289	1.682	1.361	1.345	1.226	1.000	1.227	1.247
70B-QKV-MBS2	16384	1280	8192	1.008	0.800	1.260	1.318	0.580	0.455	1.276	1.270
8B-FC1-MBS2	16384	3584	4096	0.795	0.536	1.482	1.481	0.461	0.347	1.328	1.319
8B-QKV-MBS2	16384	768	4096	0.483	0.421	1.149	1.039	0.271	0.257	1.053	1.045
405B-FC1-MBS4	32768	13312	16384	14.025	11.265	1.245	1.229	7.511	6.332	1.186	1.154
405B-QKV-MBS4	32768	2304	16384	4.648	2.924	1.590	1.581	2.525	1.664	1.518	1.538
70B-FC1-MBS4	32768	7168	8192	4.490	3.206	1.401	1.371	2.393	1.923	1.245	1.271
70B-QKV-MBS4	32768	1280	8192	1.926	1.433	1.344	1.347	1.050	0.816	1.288	1.316
8B-FC1-MBS4	32768	3584	4096	1.534	0.943	1.626	1.624	0.847	0.620	1.366	1.360
8B-QKV-MBS4	32768	768	4096	0.891	0.734	1.215	1.186	0.497	0.430	1.155	1.172
405B-FC1-MBS8	65536	13312	16384	27.347	21.930	1.247	1.234	14.561	12.447	1.170	1.157
405B-QKV-MBS8	65536	2304	16384	9.179	5.396	1.701	1.668	4.892	3.205	1.526	1.545
70B-FC1-MBS8	65536	7168	8192	9.132	6.475	1.410	1.347	4.776	3.738	1.278	1.266
70B-QKV-MBS8	65536	1280	8192	3.703	2.686	1.378	1.359	2.036	1.551	1.313	1.319
8B-FC1-MBS8	65536	3584	4096	2.968	1.808	1.641	1.624	1.661	1.250	1.329	1.349
8B-QKV-MBS8	65536	768	4096	1.720	1.375	1.251	1.205	0.926	0.802	1.155	1.182
geometric mean						*1.360*	*1.323*			*1.258*	*1.255*

Table C.5. TP+SP microbenchmark, GEMM+RS, row fprop, TP 8, MI350X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-Down-MBS1	8192	16384	6656	2.015	1.703	1.183	1.252	1.382	1.253	1.103	1.396
405B-OProj-MBS1	8192	16384	2048	1.140	0.819	1.392	1.460	0.967	0.766	1.261	1.367
70B-Down-MBS1	8192	8192	3584	0.752	0.528	1.424	1.316	0.576	0.441	1.307	1.306
70B-OProj-MBS1	8192	8192	1024	0.541	0.375	1.441	1.297	0.462	0.381	1.215	1.220
8B-Down-MBS1	8192	4096	1792	0.354	0.234	1.509	1.317	0.271	0.222	1.217	1.238
8B-OProj-MBS1	8192	4096	512	0.293	0.205	1.431	1.246	0.239	0.210	1.138	1.203
405B-Down-MBS2	16384	16384	6656	4.411	3.646	1.210	1.291	2.728	2.401	1.136	1.464
405B-OProj-MBS2	16384	16384	2048	2.252	1.581	1.424	1.487	1.882	1.417	1.328	1.381
70B-Down-MBS2	16384	8192	3584	1.389	1.002	1.386	1.453	1.105	0.792	1.395	1.560
70B-OProj-MBS2	16384	8192	1024	0.971	0.712	1.364	1.371	0.877	0.718	1.221	1.333
8B-Down-MBS2	16384	4096	1792	0.593	0.381	1.555	1.477	0.502	0.384	1.309	1.327
8B-OProj-MBS2	16384	4096	512	0.503	0.356	1.413	1.309	0.440	0.379	1.161	1.240
405B-Down-MBS4	32768	16384	6656	9.581	7.571	1.265	1.287	5.713	4.763	1.199	1.538
405B-OProj-MBS4	32768	16384	2048	4.385	3.087	1.421	1.528	3.674	2.723	1.349	1.398
70B-Down-MBS4	32768	8192	3584	2.837	1.880	1.509	1.601	2.167	1.436	1.509	1.650
70B-OProj-MBS4	32768	8192	1024	1.886	1.446	1.304	1.298	1.700	1.350	1.259	1.280
8B-Down-MBS4	32768	4096	1792	1.118	0.714	1.567	1.607	0.960	0.711	1.350	1.469
8B-OProj-MBS4	32768	4096	512	0.864	0.698	1.238	1.307	0.837	0.721	1.161	1.257
405B-Down-MBS8	65536	16384	6656	18.697	15.043	1.243	1.284	11.873	9.270	1.281	1.597
405B-OProj-MBS8	65536	16384	2048	9.350	6.196	1.509	1.602	7.322	5.297	1.382	1.434
70B-Down-MBS8	65536	8192	3584	6.004	4.158	1.444	1.527	4.271	2.894	1.476	1.566
70B-OProj-MBS8	65536	8192	1024	3.621	2.921	1.240	1.260	3.313	2.629	1.260	1.278
8B-Down-MBS8	65536	4096	1792	2.171	1.444	1.504	1.548	1.863	1.341	1.390	1.434
8B-OProj-MBS8	65536	4096	512	1.704	1.367	1.247	1.350	1.622	1.361	1.191	1.306
geometric mean						*1.380*	*1.390*			*1.271*	*1.379*

Table C.6. TP+SP microbenchmark, GEMM+RS, row fprop, TP 8, MI355X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-Down-MBS1	8192	16384	6656	2.164	1.795	1.206	1.271	1.514	1.184	1.279	1.525
405B-OProj-MBS1	8192	16384	2048	1.194	0.840	1.421	1.433	1.016	0.770	1.320	1.338
70B-Down-MBS1	8192	8192	3584	0.807	0.559	1.445	1.378	0.609	0.442	1.377	1.383
70B-OProj-MBS1	8192	8192	1024	0.539	0.393	1.371	1.307	0.476	0.383	1.242	1.212
8B-Down-MBS1	8192	4096	1792	0.353	0.246	1.432	1.308	0.275	0.229	1.200	1.218
8B-OProj-MBS1	8192	4096	512	0.281	0.219	1.283	1.220	0.244	0.220	1.107	1.173
405B-Down-MBS2	16384	16384	6656	4.309	3.320	1.298	1.344	2.973	2.198	1.352	1.603
405B-OProj-MBS2	16384	16384	2048	2.404	1.530	1.571	1.545	1.986	1.454	1.366	1.374
70B-Down-MBS2	16384	8192	3584	1.520	1.035	1.468	1.506	1.181	0.792	1.491	1.572
70B-OProj-MBS2	16384	8192	1024	0.995	0.735	1.354	1.329	0.914	0.721	1.268	1.286
8B-Down-MBS2	16384	4096	1792	0.620	0.401	1.545	1.493	0.517	0.388	1.334	1.324
8B-OProj-MBS2	16384	4096	512	0.483	0.383	1.261	1.267	0.452	0.384	1.176	1.225
405B-Down-MBS4	32768	16384	6656	8.764	6.467	1.355	1.351	5.808	4.365	1.331	1.596
405B-OProj-MBS4	32768	16384	2048	4.688	2.949	1.590	1.549	3.893	2.776	1.402	1.400
70B-Down-MBS4	32768	8192	3584	3.047	1.911	1.594	1.630	2.315	1.458	1.588	1.588
70B-OProj-MBS4	32768	8192	1024	1.972	1.456	1.354	1.308	1.777	1.393	1.276	1.274
8B-Down-MBS4	32768	4096	1792	1.154	0.740	1.559	1.579	1.002	0.722	1.388	1.417
8B-OProj-MBS4	32768	4096	512	0.892	0.735	1.214	1.295	0.870	0.744	1.169	1.234
405B-Down-MBS8	65536	16384	6656	16.904	12.565	1.345	1.380	11.491	8.499	1.352	1.615
405B-OProj-MBS8	65536	16384	2048	9.329	5.796	1.610	1.612	7.697	5.408	1.423	1.433
70B-Down-MBS8	65536	8192	3584	5.955	3.731	1.596	1.613	4.547	2.861	1.589	1.595
70B-OProj-MBS8	65536	8192	1024	3.815	2.906	1.313	1.278	3.472	2.719	1.277	1.285
8B-Down-MBS8	65536	4096	1792	2.280	1.495	1.525	1.510	1.959	1.379	1.420	1.425
8B-OProj-MBS8	65536	4096	512	1.768	1.440	1.228	1.335	1.695	1.433	1.183	1.294
geometric mean						*1.408*	*1.404*			*1.324*	*1.384*

Table C.7. TP+SP microbenchmark, AG+GEMM, row dgrad, TP 8, MI350X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-Down-MBS1	8192	6656	16384	2.172	1.976	1.099	1.127	2.314	1.840	1.258	1.163
405B-OProj-MBS1	8192	2048	16384	1.069	0.940	1.137	1.070	1.171	0.871	1.345	1.269
70B-Down-MBS1	8192	3584	8192	0.743	0.666	1.115	1.074	0.806	0.582	1.384	1.377
70B-OProj-MBS1	8192	1024	8192	0.511	0.421	1.214	1.056	0.660	0.382	1.728	1.240
8B-Down-MBS1	8192	1792	4096	0.411	0.273	1.506	1.025	0.465	0.269	1.730	1.386
8B-OProj-MBS1	8192	512	4096	0.368	0.220	1.674	1.152	0.312	0.238	1.307	1.101
405B-Down-MBS2	16384	6656	16384	4.209	3.696	1.139	1.215	4.916	4.165	1.180	1.145
405B-OProj-MBS2	16384	2048	16384	2.052	1.529	1.343	1.299	2.267	1.742	1.301	1.264
70B-Down-MBS2	16384	3584	8192	1.430	1.160	1.233	1.195	1.531	1.088	1.407	1.397
70B-OProj-MBS2	16384	1024	8192	0.897	0.764	1.175	1.055	1.074	0.680	1.579	1.442
8B-Down-MBS2	16384	1792	4096	0.575	0.421	1.366	1.256	0.737	0.455	1.619	1.513
8B-OProj-MBS2	16384	512	4096	0.463	0.380	1.217	1.096	0.484	0.397	1.220	1.083
405B-Down-MBS4	32768	6656	16384	9.128	7.253	1.258	1.260	10.103	8.679	1.164	1.181
405B-OProj-MBS4	32768	2048	16384	3.983	2.788	1.429	1.415	4.554	3.665	1.243	1.205
70B-Down-MBS4	32768	3584	8192	2.595	2.070	1.253	1.351	2.993	2.160	1.386	1.417
70B-OProj-MBS4	32768	1024	8192	1.731	1.351	1.281	1.180	1.952	1.329	1.469	1.469
8B-Down-MBS4	32768	1792	4096	1.112	0.736	1.512	1.423	1.295	0.843	1.536	1.515
8B-OProj-MBS4	32768	512	4096	0.820	0.694	1.181	1.058	0.864	0.722	1.197	1.184
405B-Down-MBS8	65536	6656	16384	18.304	14.027	1.305	1.274	20.140	17.485	1.152	1.205
405B-OProj-MBS8	65536	2048	16384	8.486	5.205	1.630	1.629	9.607	7.415	1.296	1.316
70B-Down-MBS8	65536	3584	8192	5.649	4.066	1.389	1.417	6.696	5.148	1.301	1.329
70B-OProj-MBS8	65536	1024	8192	3.390	2.511	1.350	1.252	3.682	2.670	1.379	1.467
8B-Down-MBS8	65536	1792	4096	2.049	1.370	1.496	1.501	2.418	1.670	1.448	1.551
8B-OProj-MBS8	65536	512	4096	1.566	1.281	1.222	1.130	1.641	1.395	1.176	1.201
geometric mean						*1.305*	*1.220*			*1.357*	*1.302*

Table C.8. TP+SP microbenchmark, AG+GEMM, row dgrad, TP 8, MI355X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-Down-MBS1	8192	6656	16384	2.229	1.919	1.162	1.220	2.438	1.782	1.368	1.267
405B-OProj-MBS1	8192	2048	16384	1.152	0.966	1.192	1.124	1.233	0.907	1.360	1.310
70B-Down-MBS1	8192	3584	8192	0.782	0.630	1.241	1.209	0.827	0.569	1.453	1.462
70B-OProj-MBS1	8192	1024	8192	0.516	0.435	1.184	1.096	0.667	0.396	1.684	1.260
8B-Down-MBS1	8192	1792	4096	0.396	0.279	1.416	1.086	0.457	0.262	1.746	1.457
8B-OProj-MBS1	8192	512	4096	0.328	0.226	1.450	1.073	0.311	0.248	1.255	1.101
405B-Down-MBS2	16384	6656	16384	4.190	3.409	1.229	1.268	4.806	3.604	1.333	1.268
405B-OProj-MBS2	16384	2048	16384	2.238	1.622	1.380	1.336	2.403	1.858	1.293	1.305
70B-Down-MBS2	16384	3584	8192	1.505	1.119	1.345	1.352	1.615	1.046	1.544	1.519
70B-OProj-MBS2	16384	1024	8192	0.931	0.804	1.157	1.092	1.101	0.720	1.528	1.445
8B-Down-MBS2	16384	1792	4096	0.596	0.438	1.360	1.299	0.732	0.445	1.646	1.565
8B-OProj-MBS2	16384	512	4096	0.458	0.401	1.140	1.070	0.499	0.414	1.206	1.092
405B-Down-MBS4	32768	6656	16384	8.425	6.467	1.303	1.307	9.360	7.264	1.289	1.286
405B-OProj-MBS4	32768	2048	16384	4.366	2.924	1.493	1.465	4.724	3.674	1.286	1.276
70B-Down-MBS4	32768	3584	8192	2.910	2.067	1.408	1.390	3.157	1.981	1.594	1.595
70B-OProj-MBS4	32768	1024	8192	1.815	1.433	1.267	1.209	2.023	1.420	1.424	1.464
8B-Down-MBS4	32768	1792	4096	1.137	0.766	1.485	1.455	1.304	0.828	1.574	1.584
8B-OProj-MBS4	32768	512	4096	0.835	0.741	1.128	1.059	0.878	0.759	1.156	1.171
405B-Down-MBS8	65536	6656	16384	16.544	12.430	1.331	1.338	18.272	14.587	1.253	1.298
405B-OProj-MBS8	65536	2048	16384	8.603	5.359	1.605	1.577	9.597	7.376	1.301	1.318
70B-Down-MBS8	65536	3584	8192	5.655	3.903	1.449	1.443	6.479	4.341	1.492	1.500
70B-OProj-MBS8	65536	1024	8192	3.516	2.675	1.314	1.281	3.833	2.846	1.347	1.470
8B-Down-MBS8	65536	1792	4096	2.162	1.453	1.487	1.464	2.454	1.646	1.491	1.655
8B-OProj-MBS8	65536	512	4096	1.613	1.370	1.177	1.115	1.687	1.472	1.146	1.197
geometric mean						*1.314*	*1.255*			*1.398*	*1.360*

Table C.9. TP+SP microbenchmark, bulk AG, column dgrad, TP 8, MI350X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-FC1-MBS1	8192	16384	13312	3.794	3.386	1.121	1.257	1.813	1.639	1.106	1.145
405B-QKV-MBS1	8192	16384	2304	1.178	0.669	1.761	1.743	0.665	0.391	1.701	1.722
70B-FC1-MBS1	8192	8192	7168	1.055	0.879	1.200	1.606	0.598	0.478	1.251	1.355
70B-QKV-MBS1	8192	8192	1280	0.548	0.324	1.694	1.582	0.319	0.197	1.621	1.631
8B-FC1-MBS1	8192	4096	3584	0.493	0.254	1.942	1.511	0.253	0.173	1.468	1.583
8B-QKV-MBS1	8192	4096	768	0.364	0.176	2.062	1.342	0.185	0.121	1.524	1.419
405B-FC1-MBS2	16384	16384	13312	8.014	6.959	1.152	1.354	4.015	3.670	1.094	1.133
405B-QKV-MBS2	16384	16384	2304	2.232	1.393	1.603	1.647	1.279	0.751	1.702	1.770
70B-FC1-MBS2	16384	8192	7168	1.991	1.815	1.097	1.577	1.139	0.904	1.260	1.380
70B-QKV-MBS2	16384	8192	1280	1.008	0.602	1.676	1.578	0.576	0.355	1.620	1.547
8B-FC1-MBS2	16384	4096	3584	0.753	0.493	1.529	1.569	0.430	0.286	1.500	1.698
8B-QKV-MBS2	16384	4096	768	0.499	0.323	1.546	1.388	0.296	0.205	1.443	1.498
405B-FC1-MBS4	32768	16384	13312	16.086	13.411	1.199	1.374	8.535	7.590	1.125	1.183
405B-QKV-MBS4	32768	16384	2304	4.425	2.824	1.567	1.660	2.552	1.478	1.727	1.811
70B-FC1-MBS4	32768	8192	7168	4.335	3.719	1.166	1.639	2.289	1.666	1.374	1.537
70B-QKV-MBS4	32768	8192	1280	1.920	1.218	1.576	1.556	1.094	0.710	1.542	1.553
8B-FC1-MBS4	32768	4096	3584	1.382	0.933	1.480	1.720	0.797	0.533	1.495	1.675
8B-QKV-MBS4	32768	4096	768	0.908	0.617	1.472	1.403	0.531	0.385	1.378	1.357
405B-FC1-MBS8	65536	16384	13312	31.460	26.567	1.184	1.390	17.226	14.820	1.162	1.221
405B-QKV-MBS8	65536	16384	2304	9.445	5.904	1.600	1.680	5.114	3.031	1.687	1.773
70B-FC1-MBS8	65536	8192	7168	9.304	7.733	1.203	1.636	4.825	3.955	1.220	1.330
70B-QKV-MBS8	65536	8192	1280	3.791	2.456	1.544	1.569	2.164	1.445	1.498	1.521
8B-FC1-MBS8	65536	4096	3584	2.661	1.918	1.387	1.638	1.549	1.076	1.440	1.649
8B-QKV-MBS8	65536	4096	768	1.751	1.242	1.410	1.386	1.002	0.754	1.328	1.364
geometric mean						*1.443*	*1.528*			*1.414*	*1.480*

Table C.10. TP+SP microbenchmark, bulk AG, column dgrad, TP 8, MI355X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-FC1-MBS1	8192	16384	13312	3.519	3.047	1.155	1.345	1.937	1.629	1.189	1.232
405B-QKV-MBS1	8192	16384	2304	1.246	0.673	1.850	1.776	0.689	0.401	1.718	1.748
70B-FC1-MBS1	8192	8192	7168	1.168	0.897	1.302	1.619	0.647	0.506	1.279	1.376
70B-QKV-MBS1	8192	8192	1280	0.558	0.339	1.645	1.511	0.320	0.203	1.577	1.581
8B-FC1-MBS1	8192	4096	3584	0.493	0.269	1.838	1.498	0.252	0.169	1.495	1.570
8B-QKV-MBS1	8192	4096	768	0.347	0.187	1.850	1.336	0.180	0.127	1.423	1.405
405B-FC1-MBS2	16384	16384	13312	7.095	6.006	1.181	1.405	3.800	3.120	1.218	1.242
405B-QKV-MBS2	16384	16384	2304	2.406	1.394	1.727	1.725	1.341	0.792	1.693	1.745
70B-FC1-MBS2	16384	8192	7168	2.264	1.709	1.325	1.689	1.238	0.911	1.358	1.467
70B-QKV-MBS2	16384	8192	1280	1.042	0.646	1.613	1.518	0.584	0.374	1.560	1.517
8B-FC1-MBS2	16384	4096	3584	0.792	0.489	1.620	1.625	0.446	0.286	1.556	1.792
8B-QKV-MBS2	16384	4096	768	0.502	0.344	1.459	1.344	0.294	0.212	1.388	1.476
405B-FC1-MBS4	32768	16384	13312	13.862	11.584	1.197	1.421	7.633	6.282	1.215	1.253
405B-QKV-MBS4	32768	16384	2304	4.701	2.777	1.693	1.732	2.651	1.561	1.699	1.762
70B-FC1-MBS4	32768	8192	7168	4.385	3.456	1.269	1.651	2.411	1.758	1.371	1.463
70B-QKV-MBS4	32768	8192	1280	2.000	1.292	1.548	1.521	1.126	0.712	1.581	1.582
8B-FC1-MBS4	32768	4096	3584	1.504	0.928	1.621	1.725	0.840	0.521	1.613	1.812
8B-QKV-MBS4	32768	4096	768	0.935	0.662	1.411	1.344	0.532	0.403	1.320	1.345
405B-FC1-MBS8	65536	16384	13312	26.956	22.603	1.193	1.460	14.849	12.191	1.218	1.282
405B-QKV-MBS8	65536	16384	2304	9.374	5.497	1.705	1.742	5.235	3.045	1.719	1.775
70B-FC1-MBS8	65536	8192	7168	8.699	6.880	1.264	1.667	4.793	3.504	1.368	1.461
70B-QKV-MBS8	65536	8192	1280	3.880	2.596	1.495	1.568	2.228	1.465	1.521	1.545
8B-FC1-MBS8	65536	4096	3584	2.950	1.875	1.573	1.685	1.641	1.070	1.533	1.721
8B-QKV-MBS8	65536	4096	768	1.787	1.318	1.356	1.355	1.025	0.764	1.342	1.375
geometric mean						*1.479*	*1.546*			*1.447*	*1.511*

Table C.11. TP+SP microbenchmark, bulk RS, column wgrad, TP 8, MI350X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-FC1-MBS1	13312	16384	8192	3.971	3.256	1.220	1.589	2.262	1.653	1.369	1.505
405B-QKV-MBS1	2304	16384	8192	1.145	0.660	1.735	1.793	1.011	0.625	1.619	1.516
70B-FC1-MBS1	7168	8192	8192	1.242	0.771	1.611	1.610	0.825	0.552	1.496	1.472
70B-QKV-MBS1	1280	8192	8192	0.559	0.336	1.663	1.558	0.479	0.333	1.436	1.328
8B-FC1-MBS1	3584	4096	8192	0.444	0.238	1.864	1.870	0.321	0.183	1.752	1.670
8B-QKV-MBS1	768	4096	8192	0.330	0.190	1.730	1.509	0.299	0.179	1.674	1.544
405B-FC1-MBS2	13312	16384	16384	8.329	7.155	1.164	1.443	4.827	3.696	1.306	1.453
405B-QKV-MBS2	2304	16384	16384	2.151	1.279	1.682	1.821	1.949	1.193	1.633	1.525
70B-FC1-MBS2	7168	8192	16384	2.386	1.513	1.577	1.681	1.567	1.052	1.489	1.431
70B-QKV-MBS2	1280	8192	16384	1.014	0.619	1.638	1.585	0.886	0.613	1.447	1.329
8B-FC1-MBS2	3584	4096	16384	0.763	0.494	1.545	1.722	0.590	0.335	1.760	1.644
8B-QKV-MBS2	768	4096	16384	0.525	0.364	1.441	1.249	0.552	0.325	1.699	1.568
405B-FC1-MBS4	13312	16384	32768	16.695	14.465	1.154	1.453	10.298	7.812	1.318	1.383
405B-QKV-MBS4	2304	16384	32768	4.193	2.573	1.630	1.885	3.780	2.339	1.616	1.524
70B-FC1-MBS4	7168	8192	32768	4.731	3.673	1.288	1.488	3.151	2.129	1.480	1.404
70B-QKV-MBS4	1280	8192	32768	1.922	1.195	1.608	1.602	1.713	1.174	1.458	1.336
8B-FC1-MBS4	3584	4096	32768	1.396	0.967	1.444	1.759	1.115	0.622	1.792	1.675
8B-QKV-MBS4	768	4096	32768	0.932	0.651	1.433	1.276	1.037	0.608	1.706	1.590
405B-FC1-MBS8	13312	16384	65536	32.763	29.801	1.099	1.419	20.770	16.225	1.280	1.346
405B-QKV-MBS8	2304	16384	65536	9.099	5.821	1.563	1.774	7.988	4.630	1.725	1.551
70B-FC1-MBS8	7168	8192	65536	10.576	8.094	1.307	1.417	6.813	4.902	1.390	1.326
70B-QKV-MBS8	1280	8192	65536	3.724	2.384	1.562	1.613	3.315	2.294	1.445	1.339
8B-FC1-MBS8	3584	4096	65536	2.695	1.956	1.378	1.798	2.188	1.207	1.813	1.703
8B-QKV-MBS8	768	4096	65536	1.736	1.226	1.416	1.302	2.017	1.162	1.735	1.611
geometric mean						*1.476*	*1.582*			*1.551*	*1.486*

Table C.12. TP+SP microbenchmark, bulk RS, column wgrad, TP 8, MI355X: per iteration the slowest rank, median. Base: the RCCL collective and the GEMM run back to back on one stream with no overlap (BF16: torch.mm (hipBLASLt); MXFP8: a communication-free HipKittens MXFP8 GEMM). Speedup is base over KOSMOS; vs unf. is KOSMOS's communication and GEMM WGs run back to back over the fused kernel.

				BF16				MXFP8			
Shape	M	N	K	Base ms	KOSMOS ms	Speedup	vs unf.	Base ms	KOSMOS ms	Speedup	vs unf.
405B-FC1-MBS1	13312	16384	8192	3.620	3.015	1.201	1.586	2.323	1.612	1.441	1.602
405B-QKV-MBS1	2304	16384	8192	1.226	0.691	1.774	1.925	1.065	0.667	1.598	1.501
70B-FC1-MBS1	7168	8192	8192	1.228	0.841	1.460	1.637	0.852	0.533	1.598	1.581
70B-QKV-MBS1	1280	8192	8192	0.566	0.357	1.583	1.518	0.490	0.356	1.375	1.289
8B-FC1-MBS1	3584	4096	8192	0.446	0.234	1.907	1.926	0.323	0.194	1.664	1.594
8B-QKV-MBS1	768	4096	8192	0.304	0.206	1.471	1.432	0.298	0.194	1.536	1.472
405B-FC1-MBS2	13312	16384	16384	7.231	6.188	1.169	1.513	4.638	3.348	1.385	1.491
405B-QKV-MBS2	2304	16384	16384	2.377	1.361	1.747	1.914	2.058	1.276	1.613	1.512
70B-FC1-MBS2	7168	8192	16384	2.329	1.671	1.393	1.613	1.622	1.008	1.610	1.560
70B-QKV-MBS2	1280	8192	16384	1.042	0.659	1.582	1.552	0.918	0.658	1.394	1.291
8B-FC1-MBS2	3584	4096	16384	0.805	0.505	1.594	1.819	0.597	0.354	1.686	1.617
8B-QKV-MBS2	768	4096	16384	0.509	0.383	1.328	1.277	0.556	0.354	1.570	1.499
405B-FC1-MBS4	13312	16384	32768	14.160	12.013	1.179	1.508	9.387	6.753	1.390	1.463
405B-QKV-MBS4	2304	16384	32768	4.596	2.625	1.751	1.902	4.026	2.484	1.621	1.528
70B-FC1-MBS4	7168	8192	32768	4.537	3.368	1.347	1.586	3.215	1.983	1.621	1.558
70B-QKV-MBS4	1280	8192	32768	2.010	1.278	1.573	1.563	1.778	1.262	1.409	1.298
8B-FC1-MBS4	3584	4096	32768	1.546	0.993	1.556	1.848	1.137	0.660	1.723	1.635
8B-QKV-MBS4	768	4096	32768	0.926	0.690	1.341	1.302	1.056	0.656	1.610	1.512
405B-FC1-MBS8	13312	16384	65536	27.920	23.735	1.176	1.515	18.367	13.417	1.369	1.435
405B-QKV-MBS8	2304	16384	65536	9.101	5.212	1.746	1.943	8.139	4.912	1.657	1.554
70B-FC1-MBS8	7168	8192	65536	9.363	6.846	1.368	1.520	6.592	4.424	1.490	1.425
70B-QKV-MBS8	1280	8192	65536	3.918	2.509	1.562	1.562	3.466	2.478	1.399	1.305
8B-FC1-MBS8	3584	4096	65536	3.025	1.967	1.538	1.836	2.231	1.278	1.745	1.653
8B-QKV-MBS8	768	4096	65536	1.764	1.301	1.355	1.326	2.057	1.262	1.629	1.524
geometric mean						*1.474*	*1.617*			*1.543*	*1.492*

Table C.13. MoE microbenchmark, BF16, forward pass, EP 8: per call the slowest rank, median. Baseline: RCCL A2A with TE's grouped GEMMs. Speedup is baseline over KOSMOS; vs KOSMOS serial is the same KOSMOS kernels as separate launches over the fused layer.

Routing	Model	Baseline ms	KOSMOS ms	KOSMOS TFLOPS	Speedup	vs KOSMOS serial
MI350X						
balanced	DeepSeek-V3	17.041	7.670	752.6	2.222	1.463
	Qwen3.5-397B	11.026	3.791	543.8	2.908	1.610
	DeepSeek-V4-Pro	17.443	8.531	761.2	2.045	1.314
	Kimi-K2.7-Code	18.219	8.117	711.1	2.244	1.438
	GLM-5.2	14.753	6.562	754.0	2.248	1.476
	Qwen3-235B	8.742	3.405	726.6	2.568	1.568
	Qwen3-30B	4.158	1.555	397.7	2.674	1.508
	Llama 4 Scout	3.358	2.381	865.8	1.410	1.099
	Llama 4 Maverick	4.373	2.710	760.8	1.614	1.212
	geometric mean				*2.164*	*1.400*
skewed	DeepSeek-V3	31.357	15.792	365.5	1.986	1.324
	Qwen3.5-397B	19.649	7.848	262.7	2.504	1.402
	DeepSeek-V4-Pro	30.492	17.083	380.1	1.785	1.286
	Kimi-K2.7-Code	32.022	16.005	360.7	2.001	1.335
	GLM-5.2	27.324	13.640	362.7	2.003	1.333
	Qwen3-235B	16.608	7.547	327.8	2.201	1.373
	Qwen3-30B	7.093	3.126	197.9	2.269	1.216
	Llama 4 Scout	7.586	5.737	359.4	1.322	1.078
	Llama 4 Maverick	7.679	5.651	364.8	1.359	1.223
	geometric mean				*1.898*	*1.282*
MI355X						
balanced	DeepSeek-V3	15.856	6.806	848.1	2.330	1.592
	Qwen3.5-397B	11.215	4.008	514.4	2.799	1.570
	DeepSeek-V4-Pro	16.249	7.091	915.8	2.292	1.479
	Kimi-K2.7-Code	17.165	7.081	815.2	2.424	1.571
	GLM-5.2	13.696	5.882	841.2	2.329	1.591
	Qwen3-235B	8.421	3.463	714.4	2.432	1.601
	Qwen3-30B	4.103	1.644	376.1	2.495	1.456
	Llama 4 Scout	3.183	2.287	901.6	1.392	1.158
	Llama 4 Maverick	4.090	2.423	850.8	1.688	1.277
	geometric mean				*2.199*	*1.469*
skewed	DeepSeek-V3	30.078	14.522	397.5	2.071	1.379
	Qwen3.5-397B	19.235	8.481	243.1	2.268	1.314
	DeepSeek-V4-Pro	29.068	14.683	442.3	1.980	1.394
	Kimi-K2.7-Code	30.608	14.620	394.8	2.094	1.401
	GLM-5.2	26.230	12.631	391.7	2.077	1.376
	Qwen3-235B	15.958	7.338	337.1	2.175	1.388
	Qwen3-30B	7.048	3.368	183.6	2.093	1.198
	Llama 4 Scout	6.574	5.123	402.4	1.283	1.115
	Llama 4 Maverick	7.143	4.966	415.2	1.438	1.241
	geometric mean				*1.911*	*1.308*

Table C.14. MoE microbenchmark, BF16, backward pass, EP 8: per call the slowest rank, median. Baseline: RCCL A2A with TE's grouped GEMMs. Speedup is baseline over KOSMOS; vs KOSMOS serial is the same KOSMOS kernels as separate launches over the fused layer.

Routing	Model	Baseline ms	KOSMOS ms	KOSMOS TFLOPS	Speedup	vs KOSMOS serial
MI350X						
balanced	DeepSeek-V3	25.939	14.337	805.2	1.809	1.237
	Qwen3.5-397B	16.245	5.918	696.7	2.745	1.479
	DeepSeek-V4-Pro	29.207	16.807	772.8	1.738	1.187
	Kimi-K2.7-Code	26.821	14.994	770.0	1.789	1.249
	GLM-5.2	23.205	12.294	804.9	1.888	1.252
	Qwen3-235B	13.479	6.040	819.1	2.232	1.408
	Qwen3-30B	6.068	1.793	690.0	3.385	1.703
	Llama 4 Scout	6.675	4.383	940.8	1.523	1.111
	Llama 4 Maverick	9.309	6.198	665.2	1.502	1.113
	geometric mean				*1.996*	*1.293*
skewed	DeepSeek-V3	48.415	30.535	378.1	1.586	1.196
	Qwen3.5-397B	28.235	11.970	344.4	2.359	1.467
	DeepSeek-V4-Pro	57.860	34.486	376.6	1.678	1.127
	Kimi-K2.7-Code	50.908	30.922	373.4	1.646	1.220
	GLM-5.2	42.471	26.627	371.6	1.595	1.186
	Qwen3-235B	24.938	13.154	376.1	1.896	1.356
	Qwen3-30B	10.382	4.152	297.9	2.501	1.574
	Llama 4 Scout	14.418	10.122	407.3	1.424	1.101
	Llama 4 Maverick	15.242	11.350	363.3	1.343	1.095
	geometric mean				*1.744*	*1.248*
MI355X						
balanced	DeepSeek-V3	24.104	11.493	1004.5	2.097	1.403
	Qwen3.5-397B	15.765	5.162	798.8	3.054	1.647
	DeepSeek-V4-Pro	27.974	13.883	935.6	2.015	1.270
	Kimi-K2.7-Code	25.323	12.049	958.2	2.102	1.398
	GLM-5.2	21.445	9.930	996.6	2.160	1.410
	Qwen3-235B	12.770	5.176	956.0	2.467	1.559
	Qwen3-30B	5.834	1.777	696.0	3.283	1.763
	Llama 4 Scout	5.651	3.888	1060.4	1.453	1.162
	Llama 4 Maverick	8.226	5.327	774.1	1.544	1.140
	geometric mean				*2.169*	*1.403*
skewed	DeepSeek-V3	46.070	25.534	452.1	1.804	1.297
	Qwen3.5-397B	27.346	10.560	390.5	2.590	1.578
	DeepSeek-V4-Pro	55.182	28.365	457.9	1.945	1.228
	Kimi-K2.7-Code	49.450	25.540	452.0	1.936	1.331
	GLM-5.2	40.206	21.712	455.8	1.852	1.325
	Qwen3-235B	23.386	11.339	436.3	2.062	1.460
	Qwen3-30B	10.072	3.772	328.0	2.670	1.702
	Llama 4 Scout	11.729	8.816	467.7	1.330	1.152
	Llama 4 Maverick	13.823	9.801	420.7	1.410	1.135
	geometric mean				*1.909*	*1.345*

Table C.15. MoE microbenchmark, MXFP8, forward pass, EP 8: per call the slowest rank, median. Baseline: RCCL A2A around Transformer Engine MXFP8 grouped GEMMs. Speedup is baseline over KOSMOS; vs KOSMOS serial is the same KOSMOS kernels as separate launches over the fused layer.

Routing	Model	TE MXFP8 ms	KOSMOS ms	KOSMOS TFLOPS	Speedup	vs KOSMOS serial
MI350X						
balanced	DeepSeek-V3	15.554	5.326	1083.9	2.921	1.498
	Qwen3.5-397B	10.323	3.202	643.9	3.224	1.464
	DeepSeek-V4-Pro	14.399	5.516	1177.4	2.611	1.424
	Kimi-K2.7-Code	15.916	5.514	1046.9	2.887	1.505
	GLM-5.2	13.561	4.637	1067.1	2.925	1.483
	Qwen3-235B	8.661	2.846	869.1	3.043	1.418
	Qwen3-30B	4.285	1.591	388.6	2.693	1.421
	Llama 4 Scout	2.731	1.749	1178.8	1.562	1.111
	Llama 4 Maverick	3.450	1.899	1085.9	1.817	1.240
	geometric mean				*2.566*	*1.390*
skewed	DeepSeek-V3	29.825	11.790	489.6	2.530	1.375
	Qwen3.5-397B	19.397	6.396	322.4	3.033	1.389
	DeepSeek-V4-Pro	27.322	11.497	564.9	2.377	1.364
	Kimi-K2.7-Code	30.105	11.811	488.7	2.549	1.394
	GLM-5.2	25.839	10.215	484.4	2.530	1.386
	Qwen3-235B	16.442	6.002	412.2	2.739	1.416
	Qwen3-30B	7.770	2.839	217.9	2.737	1.320
	Llama 4 Scout	5.586	3.905	527.9	1.430	1.080
	Llama 4 Maverick	6.052	3.794	543.3	1.595	1.211
	geometric mean				*2.328*	*1.322*
MI355X						
balanced	DeepSeek-V3	14.836	5.325	1084.0	2.786	1.489
	Qwen3.5-397B	10.013	3.364	612.8	2.976	1.439
	DeepSeek-V4-Pro	13.527	5.140	1263.5	2.632	1.476
	Kimi-K2.7-Code	15.154	5.363	1076.4	2.826	1.515
	GLM-5.2	12.950	4.664	1060.9	2.777	1.477
	Qwen3-235B	8.353	2.870	862.0	2.911	1.442
	Qwen3-30B	4.174	1.687	366.7	2.474	1.388
	Llama 4 Scout	2.713	1.698	1214.2	1.598	1.136
	Llama 4 Maverick	3.163	1.838	1121.5	1.721	1.242
	geometric mean				*2.468*	*1.395*
skewed	DeepSeek-V3	28.632	11.643	495.8	2.459	1.380
	Qwen3.5-397B	18.822	6.836	301.6	2.753	1.342
	DeepSeek-V4-Pro	26.344	10.897	595.9	2.418	1.399
	Kimi-K2.7-Code	28.936	11.726	492.3	2.468	1.390
	GLM-5.2	24.851	10.066	491.5	2.469	1.391
	Qwen3-235B	15.819	6.031	410.2	2.623	1.406
	Qwen3-30B	7.634	2.964	208.7	2.576	1.287
	Llama 4 Scout	5.361	3.620	569.6	1.481	1.093
	Llama 4 Maverick	5.682	3.455	596.7	1.645	1.256
	geometric mean				*2.276*	*1.323*

Table C.16. MoE microbenchmark, MXFP8, backward pass, EP 8: per call the slowest rank, median. Baseline: RCCL A2A around Transformer Engine MXFP8 grouped GEMMs. Speedup is baseline over KOSMOS; vs KOSMOS serial is the same KOSMOS kernels as separate launches over the fused layer.

Routing	Model	TE MXFP8 ms	KOSMOS ms	KOSMOS TFLOPS	Speedup	vs KOSMOS serial
MI350X						
balanced	DeepSeek-V3	23.684	10.519	1097.5	2.251	1.430
	Qwen3.5-397B	14.637	5.149	800.8	2.843	1.631
	DeepSeek-V4-Pro	25.560	12.512	1038.1	2.043	1.316
	Kimi-K2.7-Code	25.133	11.021	1047.6	2.281	1.476
	GLM-5.2	20.706	8.911	1110.5	2.324	1.452
	Qwen3-235B	12.346	4.681	1056.9	2.637	1.531
	Qwen3-30B	5.557	1.699	728.0	3.270	1.508
	Llama 4 Scout	5.197	2.704	1524.5	1.921	1.317
	Llama 4 Maverick	8.289	4.477	920.9	1.851	1.290
	geometric mean				*2.343*	*1.435*
skewed	DeepSeek-V3	44.307	21.928	526.5	2.021	1.387
	Qwen3.5-397B	26.233	9.945	414.6	2.638	1.691
	DeepSeek-V4-Pro	45.664	23.454	553.8	1.947	1.355
	Kimi-K2.7-Code	45.325	21.599	534.5	2.098	1.472
	GLM-5.2	38.530	18.280	541.3	2.108	1.434
	Qwen3-235B	23.169	10.053	492.2	2.305	1.520
	Qwen3-30B	10.217	3.659	338.1	2.792	1.442
	Llama 4 Scout	10.547	6.415	642.7	1.644	1.225
	Llama 4 Maverick	12.913	7.873	523.7	1.640	1.172
	geometric mean				*2.101*	*1.403*
MI355X						
balanced	DeepSeek-V3	21.941	9.032	1278.2	2.429	1.586
	Qwen3.5-397B	13.748	5.246	786.0	2.621	1.599
	DeepSeek-V4-Pro	23.446	10.559	1230.1	2.220	1.431
	Kimi-K2.7-Code	23.152	9.394	1228.9	2.465	1.628
	GLM-5.2	19.171	7.663	1291.3	2.502	1.607
	Qwen3-235B	11.553	4.268	1159.3	2.707	1.681
	Qwen3-30B	5.281	1.732	714.0	3.048	1.488
	Llama 4 Scout	4.766	2.616	1576.4	1.822	1.311
	Llama 4 Maverick	7.377	3.968	1039.0	1.859	1.308
	geometric mean				*2.378*	*1.510*
skewed	DeepSeek-V3	41.350	19.153	602.8	2.159	1.548
	Qwen3.5-397B	24.885	9.154	450.4	2.718	1.860
	DeepSeek-V4-Pro	42.728	20.454	635.0	2.089	1.469
	Kimi-K2.7-Code	42.354	18.686	617.8	2.267	1.645
	GLM-5.2	36.189	15.886	622.9	2.278	1.607
	Qwen3-235B	21.759	9.105	543.4	2.390	1.654
	Qwen3-30B	9.708	3.730	331.6	2.603	1.412
	Llama 4 Scout	9.539	5.921	696.3	1.611	1.253
	Llama 4 Maverick	11.684	6.847	602.2	1.706	1.250
	geometric mean				*2.173*	*1.510*