arXiv:2611.12690v1 [stat.ME] 9 Nov 2026

Average versus expected causal effects under interference: Modeling assumptions and estimation bottlenecks

Vishesh Karwa[*†] VISHESH.KARWA@TEMPLE.EDU
Edoardo M. Airoldi[*†] EDOARDO.AIROLDI@TEMPLE.EDU
Donald B. Rubin[*†‡] RUBIN@STAT.HARVARD.EDU

[*] *Department of Statistics, Operations, and Data Science, Temple University*
[†] *Data Science Institute, Temple University, Philadelphia, PA 19122, USA*
[‡] *Department of Statistics, Harvard University, Cambridge, MA 02138, USA*

Abstract

Causal inference under treatment interference requires assumptions to reduce the exponentially large number of potential outcomes. We develop hierarchical finite-population models for potential outcomes by separating these assumptions into three layers: *interference neighborhoods*, *exposure mappings*, and *outcome models*. Given an interference neighborhood and exposure mapping, we show that potential outcomes admit a nonparametric decomposition into direct, spillover, and treatment–spillover interaction components. This decomposition provides a common coordinate system for nonparametric, semiparametric, and parametric models of finite-population potential outcomes. For a given potential-outcome model, we define two classes of estimands. Average causal effects (ACEs), such as direct and total treatment effects, contrast potential outcomes under fixed treatment–exposure conditions. Expected causal effects (ECEs), such as expected direct and indirect effects, contrast expectations of potential outcomes under a specified allocation policy. For both classes, we bound the variance of the canonical Horvitz–Thompson estimator by three terms: a diagonal weight term, a term from overlapping interference neighborhoods, and a residual design-dependence term. Our upper and lower bounds show that under *graph-unaware* designs, nonparametric design-unbiased estimation of ACEs can require sample sizes growing exponentially in the interference degree, whereas polynomial growth can suffice for design-aligned ECEs. Under neighborhood interference, we construct *graph-aware* independent-set designs with variance bounds polynomial in the maximum degree for direct and total treatment effects on every graph, whereas graph cluster randomization can incur worst-case variance exponential in the degree.

Keywords: causal inference, interference, potential outcomes, network experiments, randomization inference

Contents

1 Introduction **3**

2 Hierarchical models of potential outcomes **6**
2.1 Interference graph 7
2.2 Exposure mapping 8
2.3 A non-parametric decomposition of the causal response surface 9

2.4 Outcome model . . . 10
2.5 Canonical model classes . . . 11

3 Defining causal estimands under network interference 13
3.1 Average causal effects under network interference . . . 13
3.2 Expected causal effects under network interference . . . 15
3.3 Algebraic representation of ACEs and ECEs . . . 16
3.4 Relevant and Nuisance Potential Outcomes . . . 18
3.5 Conditional expected effects . . . 19

4 Estimation of ACEs and ECEs 19
4.1 Dependence from interference and design . . . 20
4.2 ACE estimation . . . 22
4.3 ECE estimation . . . 25

5 Design-Based Estimation of Average Causal Effects 27
5.1 Independent set designs for the direct treatment effect . . . 28
5.2 Independent set designs for the total treatment effect . . . 31
5.3 Independent set designs for indirect effects . . . 33
5.4 Cluster randomized designs . . . 36

6 Implications for Designs, Estimators, and Models 36
6.1 Classical designs and estimators target expected causal effects . . . 36
6.2 Positivity and overlap for average causal effects . . . 37
6.3 Model-assisted estimation of average causal effects . . . 37

7 Numerical illustration 37

8 Discussion 41

A Additional examples and extensions for Section 2 43
A.1 Asymmetric and bipartite interference . . . 43
A.2 Vector-valued and continuous exposures . . . 43
A.3 Nonlinear and threshold outcome restrictions . . . 43
A.4 Non-anonymous exposure . . . 44

B Proofs for Section 2 44
B.1 Proof of Theorem 1 . . . 44

C Proofs and details for Section 3 45
C.1 Mixed causal effects: EIC, EIT, and ETE . . . 45
C.2 Proof of Proposition 3 . . . 46
C.3 Proof of Proposition 5 . . . 47
C.4 Proof of Proposition 25 . . . 47
C.5 Proof of Corollary 4 . . . 48
C.6 Proof of Proposition 2 . . . 48
C.7 Symmetric exposure mappings . . . 49

D Proofs for Section 4 **50**
D.1 Proof of Lemma 8 . . . 50
D.2 A generic design-dependence bound . . . 50
D.3 Residual design dependence under complete randomization . . . 51
D.4 Proof of Proposition 10 . . . 52
D.5 Proof of Proposition 11 . . . 54
D.6 Proof of Proposition 12 . . . 55
D.7 Weighted dependence bound for graph-aware ECEs . . . 56
D.8 Proof of Proposition 29 . . . 57
D.9 Proof of Proposition 13 . . . 58
D.10 Proof of Corollary 14 . . . 59
D.11 Heterogeneous ACE/ECE comparison under classical designs . . . 60
D.12 Proof of Corollary 30 . . . 61
D.13 Auxiliary calculation: DTE and TTE table rates . . . 62
D.14 Auxiliary calculation: exposure compression and propensity scores . . . 64
D.15 Auxiliary calculation: outcome-model span condition . . . 65

E Proofs for Section 5 **65**
E.1 Proof of Proposition 15 . . . 65
E.2 Proof of Lemma 16 . . . 65
E.3 Proof of Proposition 18 . . . 66
E.4 Proof of Corollary 19 . . . 66
E.5 Proof of Lemma 20 . . . 67
E.6 Proof of Proposition 21 . . . 67
E.7 Proof of Corollary 22 . . . 67
E.8 Proof of Proposition 24 . . . 68

1 Introduction

Under treatment interference, the potential outcome of one unit may depend on the treatment assignments of other units in the finite population. Consequently, the table of science is exponentially large. Moreover, unlike the classical case of no interference, there is no unique canonical causal estimand like the average treatment effect. Thus causal inference under interference requires assumptions on the structure of the table of science, together with a specification of a causal estimand and a design and estimator used to learn the estimand.

A large literature has emerged for causal inference under interference across settings such as peer effects (Manski, 1993; Bramoullé et al., 2009; Graham, 2008), vaccination studies (Hudgens and Halloran, 2008), spatial environments (Zigler et al., 2012; Giffin et al., 2023), and online and network experiments (Aral and Walker, 2011). These approaches use a variety of assumptions to reduce the number of potential outcomes per unit and define different causal estimands along with designs and estimators to estimate them. However, two layers of the causal inference problem are often specified jointly, sometimes even implicitly. The first concerns the modeling assumptions on potential outcomes. Assumptions under interference involve three distinct choices: which units' treatment assignments in the

finite population may affect a given unit's potential outcome, how those assignments are summarized or aggregated, and what functional restrictions, if any, are imposed on the resulting potential outcomes. These choices correspond to *interference neighborhoods*, *exposure mappings*, and *outcome models* respectively. These three modeling assumptions are often conflated in the literature, but they play different mathematical roles.

The second layer concerns the definition of causal effects and its interaction with design and estimation. Some causal estimands under interference are fixed potential-outcome contrasts: they compare outcomes under two specified treatment allocations or treatment–exposure conditions. Other estimands are allocation-averaged contrasts: they average potential outcomes under a specified allocation distribution, often corresponding to a policy implementation. We treat the allocation distribution that is part of the definition of an estimand as distinct from the design used to estimate it.

Without separating these layers, it can be difficult to determine whether differences across methods arise from different assumptions on the potential-outcome table, different causal estimands, different allocation policies, different experimental designs, or different estimators. In this paper, we develop a unified framework for finite population causal inference under interference that makes these distinctions explicit. We consider the Table of Science as the primitive finite-population object. We develop a set of hierarchical modeling assumptions on the Table of Science in the form of interference neighborhoods, exposure mappings, and outcome models. The first two are exclusion restrictions: they reduce the Table of Science by declaring which assignment coordinates can matter and which assignments are equivalent. On the other hand, an outcome model is a functional restriction on the potential outcomes. Once an interference neighborhood and exposure mapping are specified, the full-assignment potential outcome $Y_i(\mathbf{z})$ is represented by a lower-dimensional *reduced potential outcome* $Y_i(z, e)$, where $\mathbf{z}$ is the vector of treatment allocations, z_i is unit i's treatment, and e_i is its *exposure*. A central consequence of this representation is a nonparametric linear decomposition of the reduced potential outcome $Y_i(z, e)$ given by

$$Y_i(z, e) = A_i(z) + B_i(e) + zC_i(e),$$

where A_i captures the direct treatment component, B_i the spillover component, and C_i the treatment–spillover interaction. This decomposition is a reparameterization of the potential outcomes. It gives a common language for specifying nonparametric, semiparametric, and parametric finite population models of potential outcomes under interference. For example, a nonparametric model can leave the exposure as the full interference neighborhood and leave A_i, B_i, C_i unrestricted. Semiparametric and parametric outcome models impose additional restrictions such as additivity, homogeneity, linearity in exposure, or parametric forms for the spillover and interaction components. All potential-outcome models used in the interference literature, including arbitrary interference, linear-in-means models, stratified interference, binary exposure, spatial decay, and threshold or contagion models, can be represented within this hierarchy by appropriate choices of the interference neighborhood, exposure mapping, and outcome-model restrictions.

The framework also separates two natural classes of causal estimands under interference, assuming that the interference neighborhood and exposure mappings are correctly specified. Average causal effects (ACEs), such as the direct and total treatment effects, are contrasts of fixed potential outcomes under specific treatment allocations or treatment–

exposure combinations. They are properties of the finite population and do not involve an allocation distribution. In contrast, expected causal effects (ECEs), such as the expected direct effect (EDE) and expected indirect effect (EIE), are contrasts of expectations of potential outcomes under allocation distributions–sometimes also referred to as policies. The distributions entering the two arms of the contrast may be the same, different, or induced by a common allocation rule; in all cases, the allocation distribution used to average potential outcomes is part of the causal question itself. The experimental design remains distinct from this allocation distribution, though it is usually informed by it.

We show that there is a fundamental separation between the two classes for nonparametric estimation. For ACEs, the design must reveal the particular treatment–exposure potential outcomes that define the target estimand. Potential outcomes revealed by the design but outside this set are nuisance for that estimand: without further outcome-model restrictions on the causal response surface, observing them gives no information about the fixed target potential outcomes. For ECEs, by contrast, the exposure states generated by the target allocation distribution enter the estimand. The design does not need to isolate a set of fixed treatment-exposure conditions; it must reproduce the allocation distribution's averaging weights.

We illustrate this through design-based variance bounds for canonical weighted unbiased Horvitz–Thompson estimators. We show that, for both classes of estimands, the variance of the canonical estimator decomposes into a diagonal weight term, an interference-dependence term from overlapping neighborhoods, and a residual design-dependence term induced by the randomization design. The difference lies in the weight term. ACE weights involve design-induced inverse propensity scores for fixed relevant treatment–exposure conditions. Under classical graph-unaware designs, these propensity scores can be exponentially small in the degree of the interference graph, and matching lower bounds show that this obstruction applies to all unbiased estimators. By contrast, ECE weights are allocation policy-to-design importance weights; when the design is aligned with the allocation policy, the diagonal term is mild and the dominant cost is the dependence among observed weighted outcomes. Consequently we show that, under the same classical designs, fixed treatment-exposure ACE estimation can require sample sizes growing exponentially in the degree of the interference graph, whereas ECE estimation is limited by polynomial dependence on the degree of an appropriate dependence graph.

These results have direct practical implications for design and estimation. Designs for nonparametric ACE estimation must be aligned with the treatment–exposure combinations that define the target estimand. Potential outcomes that do not appear in the definition of the estimand are *nuisance* potential outcomes for nonparametric estimation. Thus, nonparametric estimation of ACEs requires specialized *graph-aware* designs that place sufficient probability on revealing the relevant potential outcomes. We construct such designs from independent sets of the interference graph: a distance-one independent-set design targets the direct treatment effect and a distance-two independent-set design targets the total treatment effect. We show that in each case, these graph aware designs replace the exponentially small propensity of the relevant potential outcome by a focal probability that is polynomial in the relevant local graph size. Sparse seeding of the distance-one independent set additionally estimates the influencer spillover effect—the average effect of treated influencers on untreated peers under the any-treated exposure map—at constant-order focal propen-

sities on locally tree-like regular graphs (Section 5.3). The resulting estimators have mean squared error polynomial in the graph degree scale for broad graph classes. For the total effect, cluster randomized designs can be sharper than the independent-set construction when closed neighborhoods are contained in few clusters, though this requires the clustering to be aligned with the interference graph. Outcome-model restrictions provide another route for ACE estimation by relating nuisance potential outcomes to target potential outcomes, but this route is no longer fully nonparametric and is subject to misspecification of the outcome model. By contrast, for ECEs, the allocation policy is part of the estimand itself, and all potential outcomes with positive probability under the policy enter the estimand. Thus, ECEs can be estimated nonparametrically under classical Bernoulli designs when the target allocation policy is Bernoulli.

The paper is organized as follows. Section 2 introduces the Table of Science and the modeling hierarchy. Section 2.3 gives the non-parametric decomposition of Potential Outcomes and shows how common model classes arise as special cases. Section 3 defines ACEs and ECEs. Section 4 develops the generic estimation bounds and the ACE/ECE bottleneck comparison. Section 6 draws implications for classical estimators, specialized designs, and model-assisted estimation. Section 5 develops estimand-aligned independent-set designs that recover polynomial rates for average causal effects—including a sparse-seeding design that estimates the influencer spillover effect at constant-order propensities on bounded-degree graphs—together with a comparison against cluster randomized designs, and Section 8 concludes.

2 Hierarchical models of potential outcomes

Consider a finite population of n units, indexed by $i \in \{1, \ldots, n\}$, and a binary treatment assignment $z_i \in \{0, 1\}$ for each unit i. The vector of treatment assignments is denoted by $\mathbf{z} = (z_1, \ldots, z_n)$, with $\mathbf{z} \in \{0, 1\}^n$. Under arbitrary treatment interference, the potential outcome of a unit i depends on the entire assignment vector. We write

$$Y_i(\mathbf{z}), \qquad \mathbf{z} \in \{0, 1\}^n.$$

This denotes the *full-assignment potential outcome* of unit i under treatment allocation $\mathbf{z}$. The collection of all full-assignment potential outcomes

$$\mathbb{T} = \{Y_i(\mathbf{z}) : i = 1, \ldots, n,\ \mathbf{z} \in \{0, 1\}^n\}$$

is the *Table of Science* for the study. Crucially, $\mathbb{T}$ is a fixed, full $n \times 2^n$ table dictated by nature, existing prior to any modeling assumptions. Let $\mathbf{Z}$ denote the random treatment assignment generated by a design $p(\mathbf{Z})$. The observed outcome for unit i is

$$Y_i^{\text{obs}} = Y_i(\mathbf{Z}),$$

and in any realized experiment, exactly one column of $\mathbb{T}$ is observed.

Since each unit has exponentially many counterfactuals relative to what is observable, modeling assumptions are *necessary* to relate unobserved potential outcomes to observed ones. We specify finite population models of potential outcomes through three distinct hierarchical components: an interference graph, an exposure mapping, and an outcome model.

Such assumptions constrain the entries of $\mathbb{T}$ in two different ways. The interference graph and exposure mappings are *exclusion restrictions*. The interference graph specifies which treatment coordinates can enter a unit's full-assignment potential outcome and the exposure mapping specifies how those coordinates aggregate. When combined, the interference graph and the exposure mapping partition the 2^n columns of $\mathbb{T}$ into equivalence classes for each row and require $\mathbb{T}$ to be constant within each class. On the other hand, *Outcome models* are a different kind of constraint: they do not merge columns of $\mathbb{T}$, but instead impose functional constraints on the values $\mathbb{T}$ takes across distinct potential outcomes, potentially linking entries within a row or across rows (units). The classical no-interference setting is the most restrictive exclusion restriction, forcing each row to take only two distinct values across all 2^n columns. Formally, each exclusion restriction is specified by an equivalence relation $\sim_i$ on $\{0,1\}^n$ that $\mathbb{T}$ must respect, meaning $\mathbf{z} \sim_i \mathbf{z}' \Rightarrow Y_i(\mathbf{z}) = Y_i(\mathbf{z}')$; $\mathbb{T}$ is then fully determined by its values on the class representatives.

2.1 Interference graph

The first restriction specifies the set of units in the finite popluation whose treatment assignments effect $Y_i(\mathbf{z})$.

Assumption 1 (Interference exclusion restriction) *For each unit i, there exists a subset $N_i \subseteq \{1,\dots,n\} \setminus \{i\}$ such that, for any two treatment allocations $\mathbf{z}$ and $\mathbf{z}'$,*

$$z_i = z_i' \quad \textit{and} \quad z_{N_i} = z'_{N_i} \quad \Longrightarrow \quad Y_i(\mathbf{z}) = Y_i(\mathbf{z}').$$

Formally, Assumption 1 is the equivalence relation

$$\mathbf{z} \sim_i^G \mathbf{z}' \iff z_i = z_i' \text{ and } \mathbf{z}_{N_i} = \mathbf{z}'_{N_i},$$

which partitions $\{0,1\}^n$ into 2^{d_i+1} equivalence classes for row i, where $d_i = |N_i|$. Row i of $\mathbb{T}$ takes at most 2^{d_i+1} distinct values—one per class—and the potential outcome can be written as

$$Y_i(\mathbf{z}) := Y_i(z_i, \mathbf{z}_{N_i}), \tag{2.1}$$

depending only on i's own treatment and the treatments of units in N_i. The neighborhoods $\{N_i\}_{i=1}^n$ may overlap, be asymmetric, or vary in size across units. For a fixed Table of Science, there is a minimal neighborhood

$$N_i^\star = \left\{ j \neq i : \exists \mathbf{z}, \mathbf{z}' \text{ differing only in } z_j \text{ such that } Y_i(\mathbf{z}) \neq Y_i(\mathbf{z}') \right\}.$$

The collection $\{N_i^\star\}$ can be encoded via a minimal functional *interference graph* $G^\star = (V, E^\star)$ on n nodes. The vertices of the graph V correspond to the n units in the finite population. A directed edge from node j to node i exists if the treatment of unit j can affect the potential outcome of unit i. In the remainder of the paper, we work with a specified interference graph $G = (V, E)$ and the corresponding neighborhoods $\{N_i\}_{i=1}^n$, under Assumption 1. Issues arising from random or misspecified interference graphs are discussed in the remark below. We use the degree to denote the in-degree of this graph G. Setting $N_i = \{1,\dots,n\} \setminus \{i\}$ recovers the complete interference case in which all 2^n entries of each row of $\mathbb{T}$ are distinct (the restriction is vacuous), while $N_i = \emptyset$ yields the classical SUTVA case in which each row takes only two distinct values.

It is useful to distinguish the interference graph from an observed *social graph*. The interference graph is an abstract representation of restrictions on the potential-outcome table. It may or may not coincide with a social graph, be an induced subgraph of it, or be a complete graph with the strength of interference modeled through an exposure function. In social-network settings, a natural choice is to set N_i to i's neighbors in the social graph (one-hop interference), or to all units within graph distance k (multi-hop interference). Partial interference, where units can be partitioned into disjoint groups and interference is assumed to occur only within groups, is another special case (Sobel, 2006; Hudgens and Halloran, 2008). Spatial and bipartite interference models can also be represented by appropriate choices of G; see Appendix A.

Remark 1 (Random and misspecified interference) *The minimal functional interference graph $G^\star$ may not be known in applications, and the specified interference graph G may itself be noisy or random. The graph G used in analysis should therefore be understood as an assumed interference graph. If the assumed interference neighborhoods N_i satisfies $N_i^\star \subseteq N_i$, the exclusion restriction remains valid but the representation is nonminimal. If, on the other hand, N_i omits units in $N_i^\star$, the interference graph is misspecified. Another approach is to treat the interference graph as random, and model the observed graph using a distribution; see Bhattacharya et al. (2020), Li and Wager (2022), and Ohnishi et al. (2025) for examples. For doubly robust estimators under a misspecified G, see Weinstein and Nevo (2026).*

Dependence graphs The interference graph should not be confused with the dependence graph used in large-sample analysis of design based estimators. The interference graph is a restriction on the arguments of the potential-outcomes. A dependence graph is a probabilistic object attached to a particular collection of random variables, such as observed outcomes or estimator summands. Under independent Bernoulli assignment, this dependence may be induced by overlapping interference neighborhoods. Under dependent randomization designs, such as complete or cluster randomization, the assignment mechanism can induce additional dependence even when the interference neighborhoods do not overlap. We define the dependence graph formally in Section 4.1, where it is used for variance and consistency results.

Under Assumption 1, row i of $\mathbb{T}$ takes at most 2^{d_i+1} distinct values; $\mathbb{T}$ itself remains an $n \times 2^n$ table, but is fully determined by specifying these 2^{d_i+1} values per row. Even so, 2^{d_i+1} can still be large; the exposure mapping, introduced next, provides a further exclusion restriction.

2.2 Exposure mapping

An *exposure mapping* (Aronow and Samii, 2017) summarizes neighborhood assignments into an exposure value,

$$e_i = f_i(\mathbf{z}_{N_i}) \in \mathcal{E}_i. \tag{2.2}$$

Assumption 2 (Exposure mapping) *For each unit i, there exists an exposure mapping $f_i : \{0,1\}^{|N_i|} \to \mathcal{E}_i$ such that, for any two neighborhood assignments $\mathbf{z}_{N_i}$ and $\mathbf{z}'_{N_i}$,*

$$f_i(\mathbf{z}_{N_i}) = f_i(\mathbf{z}'_{N_i}) \quad \Longrightarrow \quad Y_i(z_i, \mathbf{z}_{N_i}) = Y_i(z_i, \mathbf{z}'_{N_i}) \quad \textit{for } z_i \in \{0,1\}.$$

Formally, Assumption 2 is the equivalence relation

$$\mathbf{z}_{N_i} \sim_i^f \mathbf{z}'_{N_i} \iff f_i(\mathbf{z}_{N_i}) = f_i(\mathbf{z}'_{N_i}),$$

which coarsens $\sim_i^G$ by merging neighborhood assignments that produce the same exposure value. The function f_i encodes *how* the treatment assignments of unit's in N_i affect unit i's potential outcome. In the main text, we focus on the case where exposure values are one-dimensional and discrete, with $|\mathcal{E}_i| = K_i$ exposure levels. Vector-valued and continuous exposures are discussed in Appendix A. We use the symbols $e_i = 0$ to denote no exposure and $e_i = 1$ to denote full exposure, with the exposure mapping specifying what these labels mean in context.

Combined, Assumptions 1 and 2 define the relation

$$\mathbf{z} \sim_i \mathbf{z}' \iff z_i = z'_i \text{ and } f_i(\mathbf{z}_{N_i}) = f_i(\mathbf{z}'_{N_i}),$$

with $2K_i$ equivalence classes, where $K_i = |\mathcal{E}_i|$. Row i of $\mathbb{T}$ is thus fully determined by its $2K_i$ values $Y_i(z_i, e_i)$, one for each (z_i, e_i) pair. We call each such value a *reduced potential outcome*. The collection of these distinct reduced potential outcomes,

$$\mathbb{S} = \{Y_i(z, e) : z \in \{0, 1\},\ e \in \mathcal{E}_i,\ i = 1, \ldots, n\},$$

is the *causal response surface* induced by the specified interference neighborhoods and exposure mappings. Thus $Y_i(\mathbf{z})$ denotes a full-assignment potential outcome in $\mathbb{T}$, whereas $Y_i(z, e)$ denotes the corresponding reduced potential outcome in $\mathbb{S}$. When there is no interference, $Y_i(z)$ is the classical own-treatment potential outcome, the special case with no exposure argument. The non-parametric case corresponds to the identity mapping $e_i = \mathbf{z}_{N_i}$, giving $K_i = 2^{d_i}$ and recovering $\sim_i^G$.

Common examples of exposure mappings include symmetric or anonymous exposure, where f_i depends only on the number or fraction of treated neighbors (Hudgens and Halloran, 2008; Manski, 2013); binary exposure, where $\mathcal{E}_i = \{0, 1\}$; additive exposure of the form $f_i(\mathbf{z}_{N_i}) = \sum_{j \in N_i} h_j(z_j)$; and weighted or non-anonymous exposure, where f_i incorporates neighbor identity, covariates, or graph distance.

2.3 A non-parametric decomposition of the causal response surface

Assumption 1 concerns *who* can interfere with unit i, Assumption 2 concerns the *form* of that interference. Under Assumptions 1 and 2, $\mathbb{T}$ is fully determined by the causal response surface $\mathbb{S}$, comprising $2\sum_i K_i$ values. This response surface admits a natural non-parametric linear representation. This representation is an identity, and, without loss of generality, decomposes the potential outcome of each unit into three components: (a) a component depending only on a unit's own treatment, (b) a component depending only on a unit's exposure, and (c) the interaction between (a) and (b).

Theorem 1 (Non-parametric decomposition) *Under Assumptions 1 and 2, there exist functions $A_i(z)$, $B_i(e_i)$ and $C_i(e_i)$ such that, for each unit i and each $(z_i, e_i) \in \{0, 1\} \times \mathcal{E}_i$, a potential outcome $Y_i(z, e) \in \mathbb{S}$ can be written as*

$$Y_i(z_i, e_i) = A_i(z_i) + B_i(e_i) + z_i C_i(e_i) = \alpha_i + \beta_i z_i + B_i(e_i) + z_i C_i(e_i), \tag{2.3}$$

with the normalization $B_i(0) = C_i(0) = 0$.

The decomposition has a nice interpretation: the A_i term captures the part of the potential outcome depending only on unit i's own treatment, B_i term caputres the additive exposure component, and C_i the interaction between own-treatment and exposure (synergy or competition). These components define the corresponding direct, spillover, and interaction causal effects studied in Section 3.

The representation is a full-rank reparameterization of the causal response surface. In particular, it makes no assumptions on the potential outcomes and parametrizes the $2K_i$ distinct values of row i using $2K_i$ linearly independent parameters. For example, under the identity exposure mapping, the number of exposure levels is $K_i = 2^{d_i}$ and the number of parameters in the representation $Y_i(\mathbf{z}) = \alpha_i + \beta_i z_i + B_i(z_{N_i}) + z_i C_i(z_{N_i})$ for each unit i is 2^{d_i+1}. Similarly, under arbitrary interference, $K_i = 2^{n-1}$ and Theorem 1 gives the non-parametric model $Y_i(\mathbf{z}) = \alpha_i + \beta_i z_i + B_i(\mathbf{z}_{-i}) + z_i C_i(\mathbf{z}_{-i})$, which has 2^n parameters for unit i.

2.4 Outcome model

An outcome model imposes functional constraints on the $2K_i$ distinct values of the causal response surface via specifying semi-parametric or parametric restrictions on $\{A_i, B_i, C_i\}$ in the decomposition (2.3); it does not operate by further merging assignment columns of $\mathbb{T}$. Such constraints reduce the number of free parameters and may arise from scientific knowledge (e.g., a linear response) or be used to express sharp null hypotheses. Unlike exclusion restrictions, which operate within a single row of $\mathbb{T}$ before forming $\mathbb{S}$, outcome models constrain the distinct potential outcomes of $\mathbb{S}$ and can also link entries across rows — homogeneity assumptions, for instance, require β_i to be common across units.

Additive effects. A natural and commonly used restriction is that the direct and indirect effects enter additively, without interaction. In terms of the decomposition in (2.3), this corresponds to

$$C_i(e) = 0 \quad \text{for all } i \text{ and all } e \in \mathcal{E}_i.$$

The potential outcomes then simplify to $Y_i(z, e) = A_i(z) + B_i(e)$, and each unit has $K_i + 2$ free parameters instead of $2K_i$.

Constant (homogeneous) effects. In some settings, it may be reasonable to assume that treatment and exposure effects do not vary across units. This can be written as

$$A_i(z) = A(z), \qquad B_i(e) = B(e), \qquad C_i(e) = C(e), \quad \text{for all } i.$$

This links entries across rows of $\mathbb{T}$: potential outcomes from different units are constrained to share the same effect functions, reducing the parameter count from $O(nK_i)$ to $O(K)$.

Linear effects in exposure. When exposure levels are ordered (e.g., counts of treated neighbors), it may be plausible that the interference effect depends linearly on e. One simple example is

$$B_i(e) = \gamma_i\, e, \qquad C_i(e) = \theta_i\, e,$$

which yields $Y_i(z, e) = \alpha_i + \beta_i z_i + \gamma_i e + \theta_i z e$. If $\alpha_i, \beta_i, \gamma_i$ and θ_i are further assumed constant across units, the model reduces to a four-parameter linear dose-response formulation.

Sharp null hypotheses. Outcome models can also be used to express sharp null hypotheses for randomization testing. The linear decomposition allows these to be stated precisely. For instance, a sharp null of no treatment effect under interference is

$$\beta_i = 0 \text{ and } B_i(e_i) = C_i(e_i) = 0 \text{ for all } e_i, i$$

Similarly, the condition

$$Y_i(1, e) = Y_i(0, e) \quad \text{for all } i \text{ and all } e \in \mathcal{E}_i,$$

is equivalent to the hypothesis $A_i(1) + C_i(e_i) = A_i(0)$ and implies sharp null of no treatment effect only under additivity (i.e. $C_i(e_i) = 0$). More targeted nulls, such as absence of spillover effects or no interaction effects, can be expressed by setting subsets of $\{B_i(e), C_i(e)\}$ to zero. See Rosenbaum (2007); Aronow (2012); Pouget-Abadie et al. (2019); Athey et al. (2018); Basse et al. (2019); Puelz et al. (2022); Rigdon and Hudgens (2015) for construction of statistical tests under various such hypotheses.

Relation to estimation. Outcome modeling assumptions are not required to define causal estimands. However, they can be used for model-assisted design and estimation. In the absence of such assumptions, the parameters $\{A_i, B_i, C_i\}$ vary freely across units and exposure levels. Imposing them constrains these parameters, so that a response observed under one treatment-exposure combination can inform missing responses under other combinations. If correctly specified, this model-assisted approach can significantly reduce variance at either the design or analysis stage.

2.5 Canonical model classes

Different choices of the interference neighborhood N_i, the exposure mapping f_i, and the outcome model restrictions imposed on (A_i, B_i, C_i) lead to different potential-outcome models. The decomposition in Theorem 1 encompasses the full range of models used in the literature. A *non-parametric* model imposes no restrictions on A_i, B_i, C_i beyond the interference neighborhood and exposure mapping; the functions themselves are unrestricted. A *semi-parametric* model imposes qualitative restrictions on A_i, B_i, C_i—such as additivity ($C_i \equiv 0$), symmetry, or monotonicity. A *parametric* model specifies finite-dimensional functional forms for A_i, B_i, C_i, for example linear functions of exposure or homogeneous effects across units. We present some commonly used examples below.

Example 1 (Arbitrary interference) *Let $N_i = \{1, \ldots, n\} \setminus \{i\}$, and let $f_i(\mathbf{z}_{N_i}) = \mathbf{z}_{N_i}$. Then*

$$Y_i(\mathbf{z}) = Y_i(z_i, \mathbf{z}_{-i}) = \alpha_i + \beta_i z_i + B_i(\mathbf{z}_{-i}) + z_i C_i(\mathbf{z}_{-i}).$$

Row i of $\mathbb{T}$ takes 2^n distinct values; Assumption 1 places no constraint.

Example 2 (Neighborhood interference) *Let $N_i \subsetneq \{1, \ldots, n\} \setminus \{i\}$ be a specified proper subset, and let $f_i(\mathbf{z}_{N_i}) = \mathbf{z}_{N_i}$. Then*

$$Y_i(\mathbf{z}) = Y_i(z_i, \mathbf{z}_{N_i}) = \alpha_i + \beta_i z_i + B_i(\mathbf{z}_{N_i}) + z_i C_i(\mathbf{z}_{N_i}).$$

Row i of $\mathbb{T}$ takes 2^{d_i+1} distinct values, where $d_i = |N_i|$; the exponent reduces from $n-1$ to the local neighborhood size.

Example 3 (No interference) *Let $N_i = \emptyset$. Then $Y_i(\mathbf{z}) = Y_i(z_i)$, which is the classical no-interference case. Each unit has two own-treatment potential outcomes.*

Example 4 (Count or fractional exposure mappings) *Let $e_i = f_i(\mathbf{z}_{N_i}) = \sum_{j \in N_i} z_j$ denote the number of treated neighbors of unit i, or alternatively let $e_i = \sum_{j \in N_i} z_j / d_i$ denote the proportion of treated neighbors of unit i. Then*

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + B_i(e_i) + z_i C_i(e_i), \tag{2.4}$$

where $e_i \in \{0, 1, \ldots, d_i\}$ in the count case and $e_i \in \{0, 1/d_i, 2/d_i, \ldots, 1\}$ in the fractional case. Under this model, the number of free parameters per unit is $2(d_i + 1)$. If additivity is imposed, $C_i(e_i) = 0$ for all e_i, and the number of free parameters per unit reduces to $d_i + 2$.

Example 5 (Linear-in-means model) *Starting from equation (2.4), let $B_i(e_i) = \gamma_i e_i$ and $C_i(e_i) = \theta_i e_i$. Then*

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + \gamma_i e_i + \theta_i z_i e_i. \tag{2.5}$$

If the interaction term is zero, the model becomes

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + \gamma_i e_i. \tag{2.6}$$

This additive linear outcome model is the linear-in-means model used in the econometrics literature (Manski, 1993; Bramoullé et al., 2009; Graham, 2008).

Example 6 (2×2 exposure mapping) *Let $f_i(\mathbf{z}_{N_i}) = 0$ if all elements of $\mathbf{z}_{N_i}$ are 0 and $f_i(\mathbf{z}_{N_i}) = 1$ if at least one element of $\mathbf{z}_{N_i}$ is 1. This gives the two-by-two response-surface model, or binary exposure mapping,*

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + \gamma_i e_i + \theta_i z_i e_i, \tag{2.7}$$

where $e_i \in \{0, 1\}$ and $z_i \in \{0, 1\}$. Each unit has four distinct potential outcomes. Setting $\theta_i = 0$ gives the additive binary exposure mapping,

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + \gamma_i e_i, \tag{2.8}$$

with three free parameters per unit.

Example 7 (Weighted exposure mappings) *In non-anonymous interference, the exposure depends on the identity or covariates of the neighbors. For example, one may specify*

$$e_i = \sum_{j \in N_i} w_{ij} z_j,$$

and then write

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + B_i(e_i) + z_i C_i(e_i).$$

The additive version of this model has been used for total treatment effects under network interference (Yu et al., 2022). Additional non-anonymous and spatial examples are given in Appendix A.

The three canonical non-parametric representations—arbitrary interference, neighborhood interference, and exposure mapping (illustrated by examples 4–7 above)—are compared in Section 4 with respect to their estimation bottlenecks.

3 Defining causal estimands under network interference

Under treatment interference, there are many non-equivalent ways to define causal effects, and the appropriate estimand depends on the scientific question of interest. Causal estimands in the framework of Section 2 are summaries of the causal response surface $Y_i(z, e)$. We study two broad classes: *average causal effects* (ACEs), which are contrasts of fixed potential outcomes and do not involve any distribution over treatment allocations; and *expected causal effects* (ECEs), which are contrasts of expectations of potential outcomes under a specified allocation policy.

3.1 Average causal effects under network interference

Average causal effects (ACEs) are defined as contrasts between two fixed potential outcomes. Once an interference neighborhood and exposure mapping have been specified, the treatment–exposure space for unit i is

$$\mathcal{X}_i = \{0, 1\} \times \mathcal{E}_i.$$

For an assignment $\mathbf{z}$, let

$$D_i(\mathbf{z}) = (z_i, e_i(\mathbf{z})) \in \mathcal{X}_i$$

denote the treatment–exposure combination realized by unit i. A generic ACE compares two fixed treatment–exposure combinations

$$d_i^+ = (z_i^+, e_i^+), \qquad d_i^- = (z_i^-, e_i^-), \qquad d_i^+, d_i^- \in \mathcal{X}_i.$$

Let $d^+ = \{d_i^+\}_{i=1}^n$ and $d^- = \{d_i^-\}_{i=1}^n$. The corresponding ACE is defined as

$$\tau_{\text{ACE}}(d^+, d^-) \;=\; \frac{1}{n} \sum_{i=1}^{n} \Big(Y_i(d_i^+) - Y_i(d_i^-) \Big). \tag{3.1}$$

This definition directly generalizes the average treatment effect under no-interference assumptions, and does not require assumptions on the functional form of the response surface. Different choices of d_i^+ and d_i^- correspond to different causal questions about treatment and interference.

This definition also includes arbitrary interference as a special case: take $N_i = \{1, \ldots, n\} \setminus \{i\}$ and use the identity exposure mapping $e_i = \mathbf{z}_{-i}$. Then $D_i(\mathbf{z}) = (z_i, \mathbf{z}_{-i})$, and equation (3.1) defines ACEs directly in terms of the full treatment vector.

We now give a few examples of ACEs. Recall that the exposure levels $e_i = 0$ and $e_i = 1$ are special values defined to represent the situations when a unit is not exposed or fully exposed respectively.

Direct treatment effect. The direct treatment effect compares the average outcome under treatment versus control when exposure is held at the baseline level of 0:

$$\tau_{\text{DTE}} \;=\; \tau_{\text{ACE}}(d^+, d^-) \;=\; \frac{1}{n} \sum_{i=1}^{n} \Big(Y_i(1, 0) - Y_i(0, 0) \Big). \tag{3.2}$$

where $d_i^+ = (1, 0)$ and $d_i^- = (0, 0)$ for all i. The direct treatment effect is the interference analogue of the average treatment effect in the no-interference case. It measures the average effect of treatment on a unit when the exposure is fixed at the baseline level.

Average indirect effects. For two exposure levels $e_a, e_b \in \mathcal{E}$ and own treatment $z \in \{0,1\}$, define the average indirect effect

$$\tau_{\mathrm{IE}}(z; e_a, e_b) \;=\; \frac{1}{n}\sum_{i=1}^{n}\Big(Y_i(z, e_b) - Y_i(z, e_a)\Big), \tag{3.3}$$

which holds the unit's own treatment fixed and varies only the exposure regime between the two levels. The pair (e_a, e_b) is part of the estimand: spillover questions can concern any two exposure levels, and the baseline e_a need not be the no-exposure level. The special case $e_a = 0$ recovers the level-specific indirect effects $\tau_{\mathrm{IEC}}(e) = \tau_{\mathrm{IE}}(0; 0, e)$ and $\tau_{\mathrm{IET}}(e) = \tau_{\mathrm{IE}}(1; 0, e)$ on control and treated units, respectively. Writing $\tau_{\mathrm{IEC}} = \tau_{\mathrm{IEC}}(e_i^{\mathbf{1}})$ and $\tau_{\mathrm{IET}} = \tau_{\mathrm{IET}}(e_i^{\mathbf{1}})$ for the full-exposure versions at the all-treated level $e = e_i^{\mathbf{1}}$,

$$\tau_{\mathrm{IEC}} \;=\; \frac{1}{n}\sum_{i=1}^{n}\Big(Y_i(0,1) - Y_i(0,0)\Big), \tag{3.4}$$

$$\tau_{\mathrm{IET}} \;=\; \frac{1}{n}\sum_{i=1}^{n}\Big(Y_i(1,1) - Y_i(1,0)\Big). \tag{3.5}$$

A natural two-level exposure pair arises from the *any-treated* exposure map

$$U_i = \mathbb{1}\Big\{\textstyle\sum_{j\in N_i} Z_j \geq 1\Big\}, \tag{3.6}$$

which records whether at least one neighbor of unit i is treated. With $e_a = 0$ and $e_b = 1$, the contrast $\tau_{\mathrm{IE}}(z; 0, 1)$ measures the average effect of moving a unit with own treatment z from an environment free of treated neighbors to one with at least one treated neighbor. Held at $z = 0$, this is the *influencer spillover effect* estimated in influencer-seeding experiments in social-network settings: the effect on an average untreated unit of the presence of a treated influencer among its neighbors (Paluck et al., 2016; Aral and Walker, 2012; Bond et al., 2012). Section 5.3 shows that this two-level estimand is estimable at constant-order focal propensities on bounded-degree graphs by a sparse-seeding design, in contrast to the exponentially small propensities of its extreme full-exposure special cases.

By Proposition 3 below, $\tau_{\mathrm{IE}}(0; e_a, e_b) = n^{-1}\sum_i \{B_i(e_b) - B_i(e_a)\}$ captures the additive exposure component, while $\tau_{\mathrm{IE}}(1; e_a, e_b)$ adds the treatment–exposure interaction; the two coincide under additivity ($C_i \equiv 0$).

Remark 2 *We use the term average indirect effect for the fixed-exposure ACE estimands above. This is different from the average indirect effect introduced in Hu et al. (2022), which we define as the expected indirect effect as it belongs to the expected causal effects class.*

Total treatment effect. The total treatment effect compares the fully treated and fully exposed condition to the fully untreated and unexposed condition:

$$\tau_{\mathrm{TTE}} \;=\; \tau_{\mathrm{ACE}}(d^+, d^-) \;=\; \frac{1}{n}\sum_{i=1}^{n}\Big(Y_i(1,1) - Y_i(0,0)\Big). \tag{3.7}$$

where $d_i^+ = (1,1)$ and $d_i^- = (0,0)$ for all i.

3.2 Expected causal effects under network interference

Average causal effects compare two fixed potential outcomes. In many cases, however, the exposure of a unit is random under a policy of scientific interest. Expected causal effects are defined as contrasts of allocation-averaged responses, where entries of the causal response surface are averaged over one or more allocation distributions specified by the estimand. The target allocation policy is part of the causal question itself. A generic expected causal effect contrasts random treatment–exposure combinations induced by an allocation policy. Let $\mathcal{Q}^+$ and $\mathcal{Q}^-$ denote two allocation policies. For each unit i, let

$$D_i^+(\mathbf{Z}), \qquad D_i^-(\mathbf{Z})$$

be two corresponding random elements of $\mathcal{X}_i$. A generic ECE is any contrast of the form

$$\tau_{\mathrm{ECE}}(D^+, D^-; \mathcal{Q}^+, \mathcal{Q}^-) = \frac{1}{n}\sum_{i=1}^{n}\left[\mathbb{E}_{\mathcal{Q}^+} Y_i\{D_i^+(\mathbf{Z})\} - \mathbb{E}_{\mathcal{Q}^-} Y_i\{D_i^-(\mathbf{Z})\}\right]. \tag{3.8}$$

The distributions $\mathcal{Q}^+, \mathcal{Q}^-$ can be different from the experimental design P, but are often induced by P. For example, a study may use a randomized saturation design with several treatment probabilities to estimate the expected effect of a single target rollout policy, such as treating 70 percent of units; the rollout policy defines the estimand, while the saturation design is chosen to provide overlap and precision. For the special case where $\mathcal{Q}^+ = \mathcal{Q}^- = \mathcal{Q}$, we write $\tau_{\mathrm{ECE}}(D^+, D^-; \mathcal{Q})$. We now consider some examples of expected causal effects used in the literature. Let $P_{\mathcal{Q}}(E_i = e)$ denote the marginal distribution of the exposure level E_i under allocation policy $\mathcal{Q}$. The expected response of unit i under own treatment $z \in \{0, 1\}$ is

$$\mu_i(z; \mathcal{Q}) \;=\; \sum_{e \in \mathcal{E}} Y_i(z, e)\, P_{\mathcal{Q}}(E_i = e). \tag{3.9}$$

Expected (average) direct effect. The expected direct effect compares the average outcome under treatment versus control, averaging over the exposure distribution induced by the allocation policy:

$$\tau_{\mathrm{EDE}}(\mathcal{Q}) = \mathbb{E}_{\mathcal{Q}}\left[\frac{1}{n}\sum_{i=1}^{n}\left(Y_i(1, E_i) - Y_i(0, E_i)\right)\right] \;=\; \frac{1}{n}\sum_{i=1}^{n}\Big(\mu_i(1; \mathcal{Q}) - \mu_i(0; \mathcal{Q})\Big). \tag{3.10}$$

This is the one-contrast special case of (3.8) with $D_i^+(\mathbf{Z}) = (1, E_i)$, $D_i^-(\mathbf{Z}) = (0, E_i)$, and common policy $\mathcal{Q}$. The expected average direct effect is also known as the Average Direct Effect (ADE) in the literature, see Sävje et al. (2017). We use the term expected direct effect to emphasize that it is an ECE and not an ACE.

Expected (average) indirect effect. The expected indirect effect compares the expected outcome of unit i when unit j's treatment is changed, keeping all other treatment assignments the same. For $j \in N_i$, define

$$E_i^{+j} = f_i(\mathbf{Z}_{N_i \setminus \{j\}}, Z_j = 1), \qquad E_i^{-j} = f_i(\mathbf{Z}_{N_i \setminus \{j\}}, Z_j = 0),$$

as the exposure of unit i when neighbor j is treated or untreated, holding all others fixed. Then,

$$\tau_{\text{EIE}}(\mathcal{Q}) = \frac{1}{n}\sum_{i=1}^{n}\sum_{j\in N_i}\mathbb{E}_{\mathcal{Q}}\Big[Y_i\big(Z_i, E_i^{+j}\big) - Y_i\big(Z_i, E_i^{-j}\big)\Big]. \tag{3.11}$$

This is an aggregation of neighbor-specific ECE contrasts. For each ordered pair (i, j) with $j \in N_i$, define

$$D_{ij}^{+}(\mathbf{Z}) = (Z_i, E_i^{+j}), \qquad D_{ij}^{-}(\mathbf{Z}) = (Z_i, E_i^{-j}).$$

The expected indirect effect is also known as the Average Indirect Effect. The arbitrary interference version of this estimand was introduced in Hu et al. (2022).

Proposition 2 (Policy derivative representation) *Under a Bernoulli policy $\mathcal{Q}_\pi$ with $\mathbf{Z} \sim \prod_i \text{Bern}(\pi_i)$, writing $\mu(\pi) = n^{-1}\sum_i \mathbb{E}_\pi[Y_i(Z_i, e_i(\mathbf{Z}_{N_i}))]$ for the expected population outcome, the total policy derivative $\tau_{\text{TPE}}(\pi) = \sum_j \partial\mu(\pi)/\partial\pi_j$ decomposes as $\tau_{\text{TPE}}(\pi) = \tau_{\text{EDE}}(\pi) + \tau_{\text{EIE}}(\pi)$, recovering the Average Direct Effect / Average Indirect Effect decomposition of Hu et al. (2022) under arbitrary interference as a special case. Under symmetric exposure mappings, EIE has a simpler finite-difference form; details are in Appendix C.6.*

Expected policy effect. The general definition also allows the two arms of the contrast to be averaged under different allocation policies. Let $\mathcal{Q}^+$ and $\mathcal{Q}^-$ be two policies, for example Bernoulli policies with treatment probabilities p_+ and p_-. The expected policy effect is

$$\tau_{\text{policy}}(\mathcal{Q}^+, \mathcal{Q}^-) = \frac{1}{n}\sum_{i=1}^{n}\left[\mathbb{E}_{\mathcal{Q}^+}Y_i\{D_i(\mathbf{Z})\} - \mathbb{E}_{\mathcal{Q}^-}Y_i\{D_i(\mathbf{Z})\}\right]. \tag{3.12}$$

This is the special case of (3.8) with $D_i^+(\mathbf{Z}) = D_i^-(\mathbf{Z}) = D_i(\mathbf{Z})$, but with $\mathcal{Q}^+ \neq \mathcal{Q}^-$. It represents questions such as comparing an allocation policy with 50 percent treatment probability to one with 20 percent treatment probability.

Mixed effects that combine policy-averaged and fixed-baseline potential outcomes (EIC, EIT, ETE) are defined in Appendix C.1.

3.3 Algebraic representation of ACEs and ECEs

We now represent the two classes of estimands using the non-parametric linear decomposition of the causal response surface.

Proposition 3 (Representation of fixed and expected effects) *Under the non-parametric linear decomposition of Theorem 1,*

$$Y_i(z, e) = \alpha_i + \beta_i z + B_i(e) + zC_i(e), \qquad B_i(0) = C_i(0) = 0,$$

the average causal effects can be written as

$$\begin{aligned}
\tau_{\mathrm{DTE}} &= \frac{1}{n}\sum_i \beta_i, \\
\tau_{\mathrm{IE}}(z; e_a, e_b) &= \frac{1}{n}\sum_i \{B_i(e_b) - B_i(e_a) + z\,(C_i(e_b) - C_i(e_a))\}, \\
\tau_{\mathrm{IEC}}(e) &= \frac{1}{n}\sum_i B_i(e), \\
\tau_{\mathrm{IET}}(e) &= \frac{1}{n}\sum_i \{B_i(e) + C_i(e)\}, \\
\tau_{\mathrm{TTE}} &= \frac{1}{n}\sum_i \{\beta_i + B_i(1) + C_i(1)\}.
\end{aligned}$$

For an allocation policy $\mathcal{Q}$, let $\bar{B}_{\mathcal{Q}} = n^{-1}\sum_i \mathbb{E}_{\mathcal{Q}}[B_i(E_i)]$ and $\bar{C}_{\mathcal{Q}} = n^{-1}\sum_i \mathbb{E}_{\mathcal{Q}}[C_i(E_i)]$. Then

$$\tau_{\mathrm{EDE}}(\mathcal{Q}) = \frac{1}{n}\sum_i \beta_i + \bar{C}_{\mathcal{Q}}.$$

The mixed effects EIC, EIT, and ETE satisfy analogous expressions: $\tau_{\mathrm{EIC}}(\mathcal{Q}) = \bar{B}_{\mathcal{Q}}$, $\tau_{\mathrm{EIT}}(\mathcal{Q}) = \bar{B}_{\mathcal{Q}} + \bar{C}_{\mathcal{Q}}$, and $\tau_{\mathrm{ETE}}(\mathcal{Q}) = n^{-1}\sum_i \beta_i + \bar{B}_{\mathcal{Q}} + \bar{C}_{\mathcal{Q}}$; see Appendix C.1.

These representations connect the estimands directly to the components of the nonparametric response-surface decomposition. While fixed-exposure ACEs select values of B_i and C_i at specific exposure levels, ECEs average these same components over the exposure distribution induced by the allocation policy. Specifically, τ_{DTE} isolates the direct effect component β_i, whereas τ_{IEC} and τ_{IET} capture the pure indirect and indirect-plus-interference components, respectively. Notably, the expected direct effect encapsulates both the pure direct component and the treatment–interference interaction component, the latter being 0 under additivity. This yields the following corollary:

Corollary 4 *Under additivity of the causal response surface, $\tau_{\mathrm{EDE}}(\mathcal{Q}) = \tau_{\mathrm{DTE}}$ for every allocation policy $\mathcal{Q}$.*

Without additivity, the expected direct effect contains the policy-weighted interaction component $\bar{C}_{\mathcal{Q}}$ and is policy-dependent: estimators that target EDE do not target the fixed-exposure DTE. The total treatment effect also has a decomposition:

Proposition 5 (ACE total effect decomposition) *The total treatment effect decomposes as*

$$\tau_{\mathrm{TTE}} = \tau_{\mathrm{DTE}} + \tau_{\mathrm{IET}}.$$

Under additivity ($C_i \equiv 0$), $\tau_{\mathrm{IET}} = \tau_{\mathrm{IEC}}$, so $\tau_{\mathrm{TTE}} = \tau_{\mathrm{DTE}} + \tau_{\mathrm{IEC}}$.

Decompositions involving the mixed effects EIC, EIT, and ETE are given in Appendix C.1.

3.4 Relevant and Nuisance Potential Outcomes

Another distinction between average causal effects (ACEs) and expected causal effects (ECEs) ties to which entries of $\mathbb{S}$ define the estimand. In a non-parametric framework, no outcome model is imposed on the distinct potential outcomes: observing one element of $\mathbb{S}$ provides no information about another distinct element of $\mathbb{S}$. Thus, potential outcomes that do not enter the definition of an estimand are uninformative for its non-parametric estimation, regardless of how often an experimental design reveals them. We formalize this distinction by partitioning the causal response surface relative to a target estimand.

Definition 6 (Relevant and nuisance potential outcomes) *Fix a causal estimand τ, defined as a functional of the causal response surface $\mathbb{S}$. A potential outcome $Y_i(z,e)$ is relevant for τ if the functional defining τ depends on $Y_i(z,e)$. The relevant set is*

$$\mathcal{R}_\tau = \{Y_i(z,e) \in \mathbb{S} : \tau \text{ depends on } Y_i(z,e)\}.$$

The complement $\mathcal{R}_\tau^c = \mathbb{S} \setminus \mathcal{R}_\tau$ is the set of nuisance potential outcomes for τ.

For ACEs, the relevant set $\mathcal{R}_\tau$ is typically a sparse subset of the causal response surface. For example, a direct treatment effect depends only on

$$\{Y_i(1,0), Y_i(0,0)\}_{i=1}^n,$$

whereas a total treatment effect depends only on

$$\{Y_i(1,1), Y_i(0,0)\}_{i=1}^n.$$

All other entries of $\mathbb{S}$ are nuisance for non-parametric estimation of these estimands. For example, if a design exposes a unit to an intermediate exposure level e, revealing $Y_i(1,e)$, that observation provides no non-parametric information about target potential outcomes such as $Y_i(1,1)$ or $Y_i(1,0)$. Without an outcome model relating distinct potential outcomes of $\mathbb{S}$ across treatment or exposure levels, observations from nuisance potential outcomes cannot reduce variance for the target estimand. They may instead waste probability mass, contribute bias, or require outcome modeling assumptions that go beyond non-parametric estimation.

For ECEs, the allocation policy is part of the definition of the estimand. Any treatment–exposure combination that is assigned positive probability under the allocation policy is included in the relevant set $\mathcal{R}_\tau$ of an ECE. For example, the relevant set for the expected direct effect $\tau_{\text{EDE}}(\mathcal{Q})$ is the set of $Y_i(1,e)$ and $Y_i(0,e)$ for each exposure level e in the support of the exposure distribution induced by policy $\mathcal{Q}$. When the policy induces positive probability on every treatment–exposure combination in the specified representation, the relevant set is the entire causal response surface.

This difference in relevant potential outcomes leads to different non-parametric estimation problems, as discussed in Section 4. For ACEs, the design must reveal the relevant potential outcomes. Non-parametric estimation thus requires the propensity scores for those combinations to be bounded away from zero. When the probability of revealing the relevant potential outcomes is rare under the design, inverse-probability weights become unstable

and variance can be large. Consequently, non-parametric estimation of ACEs often requires specialized designs that place sufficient probability mass on the relevant potential outcomes; such designs are discussed in Section 6. For ECEs, the estimand instead averages over the treatment–exposure combinations induced by the policy. The relevant support is therefore determined by the policy itself. In this case, the main variance bottleneck is not the probability of observing a narrow set of fixed treatment–exposure entries, but the overlap between the observed assignment design and the policy. Importance-weighted HT estimators use outcomes across the support of the policy, with precision governed by policy-to-design weights, network dependence, and the extent of support overlap.

3.5 Conditional expected effects

The two classes introduced in Sections 3.1 and 3.2 also clarify conditional estimands that have appeared in the literature on partial interference. Hudgens and Halloran (2008) introduced causal effects under partial interference based on conditional expectations of potential outcomes. In their setting, units are partitioned into disjoint interference groups, and causal effects are defined by conditioning on the treatment assignment within the group. Formally, one considers conditional potential outcomes of the form

$$\mathbb{E}\big[Y_i(z,e) \,\big|\, Z_i = z\big],$$

where the expectation is taken with respect to the distribution of exposure in the group, conditional on the individual treatment. Estimands are then defined as contrasts between such conditional expectations. These estimands can be viewed as *conditional expected causal effects*, which differ from the marginal expected effects defined in Section 3.2. Conditional estimands depend on the randomization scheme within groups, and are most natural under two-stage designs with partial interference in which treatment probabilities are assigned at the group level. In arbitrary network interference, or under designs where a unit's treatment and exposure are dependent, conditioning on Z_i also changes the exposure distribution. The resulting estimand remains well defined, but its interpretation depends on the conditional allocation distribution; see Sävje et al. (2017) for a discussion.

4 Estimation of ACEs and ECEs

In this section, we study design-unbiased estimators of ACEs and ECEs. In particular, under a fixed experimental design P, we derive upper bounds for the variance of canonical Horvitz–Thompson (HT) estimators as well as lower bounds on the design variance of any unbiased estimator. Finally, we establish the conditions under which these estimators achieve design consistency. Under a growing sequence of finite populations indexed by n, an estimator $\hat{\tau}_n$ is design-MSE consistent for τ_n if

$$\mathbb{E}_{P_n}\{(\hat{\tau}_n - \tau_n)^2\} \to 0.$$

The estimation machinery applies to any choice of interference graph and exposure mapping. As examples, we study estimation under the three examples introduced in Section 2: arbitrary interference $(Z_i, \mathbf{Z}_{-i})$, neighborhood interference $(Z_i, \mathbf{Z}_{N_i})$, and exposure mappings $(Z_i, e_i(\mathbf{Z}_{N_i}))$. Throughout this section, no outcome modeling assumptions are imposed.

Once the interference neighborhood and exposure mapping have been fixed, non-parametric estimation treats the distinct potential outcomes of $\mathbb{S}$ as unrestricted. Observing one distinct potential outcome $Y_i(z, e)$ does not, without an outcome model, identify or reduce uncertainty about another distinct potential outcome $Y_i(z', e')$.

4.1 Dependence from interference and design

Consider a generic estimator $\hat{\tau}_n$ that is a function of random variables $X_1, \ldots, X_n$. In the design-based setting, the potential-outcome table is fixed, so any observed-outcome term $Y_i^{obs} = Y_i(\mathbf{Z})$ is a measurable function of the treatment assignment vector. Thus each X_i can be written generally as $X_i = g_i(\mathbf{Z})$. The variance bounds in the sequel use dependence graphs for the random variables $\{X_i\}$.

Definition 7 (Dependence graph and average dependence) *Consider random variables $\{X_i\}_{i=1}^n$. A graph $G_{\text{dep}} = (V, E_{\text{dep}})$ is a dependence graph of $\{X_i\}_{i=1}^n$ if the following holds: whenever two disjoint sets of vertices $A, B \subseteq V$ have no edge between them, the collections $\{X_i : i \in A\}$ and $\{X_j : j \in B\}$ are independent. We write $c_{ij} = 1$ when i and j are adjacent in this dependence graph and $c_{ij} = 0$ otherwise. The dependence degree of i is*

$$c_i = \sum_{j \neq i} c_{ij},$$

and the average dependence degree is

$$c_{\text{avg}} = \frac{1}{n} \sum_{i=1}^{n} c_i = \frac{1}{n} \sum_{i=1}^{n} \sum_{j \neq i} c_{ij}.$$

Remark 3 (Estimator dependence versus interference neighborhoods) *The dependence graph is an object associated with the estimator $\hat{tau}_n$ and the random variables X_i, not directly with the causal estimand. If X_i depends only on a subvector of the treatment assignment, we may write $X_i = g_i(\mathbf{Z}_{A_i})$, where $A_i \subseteq \{1, \ldots, n\}$ is an assignment-dependence set for that estimator variable. The set A_i includes the treatment assignment coordinates used by the observed outcome and by any random weights or normalizing factors. When an interference graph with neighborhoods N_i is specified, write $\mathcal{C}_i = \{i\} \cup N_i$ for the closed interference neighborhood; this includes arbitrary interference as the special case $N_i = \{1, \ldots, n\} \setminus \{i\}$. For the local HT/IPW estimators below, the observed outcome has the form $Y_i(Z_i, \mathbf{Z}_{N_i})$ and the weights are also a function of $(Z_i, \mathbf{Z}_{N_i})$. Hence $A_i \subseteq \mathcal{C}_i$. Ratio estimators such as Hajek estimators may introduce additional assignment dependence because their weights depend on the full treatment assignment vector $\mathbf{Z}$.*

Typically, the variables X_i and X_j can be dependent for two reasons: their treatment-assignment index sets A_i and A_j can overlap because of interference, and the randomization design can directly induce dependence among treatment assignments, even when A_i and A_j do not overlap. Below, we separate these two sources into interference-overlap dependence and residual design dependence.

Under independent treatment assignment, there is no residual design dependence. The remaining dependence is induced by overlap of the assignment coordinates on which the

variables X_i depend: if $X_i = g_i(\mathbf{Z}_{A_i})$, then $\mathbb{1}\{A_i \cap A_j \neq \emptyset\}$ gives a valid dependence edge under independent assignment. The following lemma records the closed-neighborhood specialization used for the local HT/IPW estimators studied below.

Lemma 8 (Dependence induced by local estimator variables) *Suppose the treatment assignments are mutually independent and X_i is a function of $\mathbf{Z}_{\mathcal{C}_i}$, where $\mathcal{C}_i = \{i\} \cup N_i$ is the closed interference neighborhood of unit i. Then a valid dependence graph of $\{X_i\}$ is obtained by setting*

$$c_{ij} = \mathbb{1}\{\mathcal{C}_i \cap \mathcal{C}_j \neq \emptyset\}, \qquad i \neq j.$$

If the interference graph is undirected, $d_i = |N_i|$, and $B_2(i)$ denotes the two-hop interference neighborhood of i, including i, then

$$c_i = |B_2(i)| - 1$$

and

$$b_{2,\mathrm{avg}} := \frac{1}{n}\sum_{i=1}^{n}\{|B_2(i)|-1\} = c_{\mathrm{avg}}, \qquad \bar{d} \leq b_{2,\mathrm{avg}} \leq \min\left\{n-1, \frac{1}{n}\sum_{i=1}^{n} d_i^2\right\}, \qquad \bar{d} = \frac{1}{n}\sum_{i=1}^{n} d_i.$$

Thus, for independent assignment on an undirected interference graph, the average dependence degree of estimators that use local weights is controlled by the second moment of the degree distribution of the interference graph rather than by the average degree alone. For bounded-degree and sparse Erdős–Renyi interference graphs with bounded expected degree, $c_{\mathrm{avg}} = O(1)$. On locally tree-like d-regular graphs, $c_{\mathrm{avg}} \asymp d^2$. On a star graph, $\bar{d} = O(1)$ but $c_{\mathrm{avg}} \asymp n$, because all leaves have overlapping closed neighborhoods through the hub. On complete graphs, $c_{\mathrm{avg}} = n - 1$.

Under dependent randomization designs, two random variables X_i and X_j can be dependent even when their closed interference neighborhoods do not overlap. We measure this residual design-induced dependence by the following coefficient.

Definition 9 (Design dependence coefficient) *For two assignment-index sets $A, B \subseteq \{1, \ldots, n\}$, define*

$$\rho_P(A, B) = \sup_{\|f\|_\infty \leq 1,\, \|g\|_\infty \leq 1} \left|\mathrm{Cov}_P\{f(\mathbf{Z}_A), g(\mathbf{Z}_B)\}\right|,$$

where the supremum is over all measurable functions of the indicated subvectors.

For the local HT/IPW variance bounds below, we use

$$\rho_{ij} = \rho_P(\mathcal{C}_i, \mathcal{C}_j),$$

and

$$r_{\mathrm{avg}} = \frac{1}{n}\sum_{i=1}^{n}\sum_{j \neq i}(1 - c_{ij})\rho_{ij}.$$

Here ρ_{ij} is the pairwise design-dependence coefficient between the two closed-neighborhood assignment subvectors, while r_{avg} averages only the residual design dependence among pairs

whose closed neighborhoods do not overlap. If the design is independent assignment and $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$, then $\rho_{ij} = 0$. Complete randomization, cluster randomization, paired designs, and other dependent designs can have $\rho_{ij} > 0$ even for disjoint closed neighborhoods. These quantities separate overlap of assignment coordinates from dependence induced directly by the randomization design. The variance bounds below therefore separate three sources of variation: the diagonal scale of the HT or importance weights, dependence generated by overlapping closed interference neighborhoods, and residual dependence generated by the randomization design.

4.2 ACE estimation

We first consider nonparametric estimation of an ACE

$$\tau_{\mathrm{ACE}}(d^+, d^-) = \frac{1}{n}\sum_{i=1}^{n}\{Y_i(d_i^+) - Y_i(d_i^-)\}.$$

Write $\pi_i(d) = P\{D_i(\mathbf{Z}) = d\}$ for the design-induced propensity that unit i is observed in treatment–exposure combination d. The generic Horvitz–Thompson estimator is

$$\hat{\tau}^{HT}_{\mathrm{ACE}}(d^+, d^-) = \frac{1}{n}\sum_{i=1}^{n} Y_i^{obs}\left\{\frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^+\}}{\pi_i(d_i^+)} - \frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^-\}}{\pi_i(d_i^-)}\right\}.$$

This is the exposure-mapping Horvitz–Thompson estimator of Aronow and Samii (2017), written here for the generic ACE notation. For the variance bound, let

$$H_i = \frac{1}{\pi_i(d_i^+)} + \frac{1}{\pi_i(d_i^-)}.$$

Proposition 10 (ACE Horvitz–Thompson estimator) *Assume $\pi_i(d_i^+) > 0$ and $\pi_i(d_i^-) > 0$ for all i. Then $\hat{\tau}^{HT}_{\mathrm{ACE}}(d^+, d^-)$ is design-unbiased for $\tau_{\mathrm{ACE}}(d^+, d^-)$ for every potential-outcome table. If, in addition, $|Y_i(d)| \le M$ and the ith summand depends on the assignment only through $\mathbf{Z}_{\mathcal{C}_i}$, then*

$$\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{ACE}}(d^+, d^-)\} \le \frac{M^2}{n^2}\left\{\sum_{i=1}^{n} H_i + \sum_{i=1}^{n}\sum_{j\neq i} c_{ij}H_iH_j + \sum_{i=1}^{n}\sum_{j\neq i}(1 - c_{ij})\rho_{ij}H_iH_j\right\},$$

and if $H_i \le H$ for all i, then

$$\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{ACE}}(d^+, d^-)\} \le M^2\left\{\frac{H}{n} + H^2\frac{c_{\mathrm{avg}} + r_{\mathrm{avg}}}{n}\right\}.$$

Consequently, along a sequence of finite populations, the estimator is design-MSE consistent if $H + H^2(c_{\mathrm{avg}} + r_{\mathrm{avg}}) = o(n)$.

The proof is given in Appendix D. This proposition separates three sources of variation in the HT estimator: the inverse propensity scores of the relevant potential outcomes $Y_i(d_i^+)$ and $Y_i(d_i^-)$ (H_i), the dependence generated by overlapping closed interference neighborhoods (c_{avg}), and residual dependence from the randomization design (r_{avg}). The next proposition shows that the inverse-propensity part is not merely an artifact of the HT estimator.

Proposition 11 (Inverse propensity scores for ACEs) *Fix the design P_n and the induced treatment–exposure combinations $D_i(\mathbf{Z})$. Suppose that $\hat{\tau}_n$ is design-unbiased for $\tau_{\mathrm{ACE}}(d^+, d^-)$ for every bounded potential-outcome table. Then*

$$\sup_{|Y|\le M} \mathrm{Var}_{P_n}(\hat{\tau}_n) \ge \frac{M^2}{n^2} \max_i \left\{ \frac{1}{\pi_i(d_i^+)} - 1,\ \frac{1}{\pi_i(d_i^-)} - 1 \right\}.$$

The proof is given in Appendix D. This coordinate-wise bound demonstrates that no unbiased estimator of a fixed-exposure ACE can avoid the smallest propensity scores among the relevant treatment–exposure combinations. Below, we show that the diagonal term in the upper bound of Proposition 10 scales identically, and is usually the dominant term when compared to off-diagonal network interference and design dependence terms, giving a matching lower bound for regular graph-unaware designs.

Estimation of DTE and TTE under classical designs. We now examine when the Horvitz–Thompson estimator can consistently estimate the direct and total treatment effects under standard graph-unaware designs. Proposition 11 provides a matching lower bound: no unbiased estimator can avoid the inverse-propensity scale H, so the variance conditions we derive are necessary for any design-unbiased estimator, not just the HT. The analysis covers independent designs (Bernoulli, heterogeneous Bernoulli) and exchangeable designs (complete randomization) alike. Proposition 31 (Appendix D.13) shows they share a key property: any specific treatment configuration on a units has probability $\exp\{-\Theta(a)\}$, so its inverse propensity scales as $\exp\{\Theta(a)\}$. Here a is the number of units the exposure condition constrains — d_i+1 for neighborhood-interference targets, n for arbitrary-interference targets. This exponential-in-a scaling is the central bottleneck for fixed-exposure ACE estimation and is the only design feature needed below. The variance bound of Proposition 10 formalizes this bottleneck: under independent assignment,

$$\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{ACE}}\} \le M^2\big\{H/n + H^2 c_{\mathrm{avg}}/n\big\},$$

where $H = \max_i(1/\pi_i(d_i^+)+1/\pi_i(d_i^-))$ and c_{avg} is the average dependence degree (Lemma 8). We write $\mathcal{V}_n(H, c) = H/n + H^2 c/n$ for this scale, suppressing M. An additional residual dependence term appears under exchangeable designs (Appendix D.3), but it does not change the conclusions.

Tables 1 and 2 summarize H and the matching lower bound for the DTE and TTE under different representations. The column "HT H" records the inverse-propensity scale and "LB H" the matching lower bound. The variance scale column combines H with the dependence degree $b_{2,\mathrm{avg}}$ (which equals c_{avg} under independent assignment). When H is exponential, the polynomial dependence is negligible by comparison. For the exposure-mapping rows, let $S_i = \sum_{j\in N_i} Z_j$ denote the number of treated neighbors and $U_i = \mathbb{1}\{S_i \ge 1\}$ the binary any-treated indicator.

Direct treatment effect. The relevant potential outcomes for the DTE are $Y_i(1, 0_{N_i})$ and $Y_i(0, 0_{N_i})$ (Definition 6). Under neighborhood interference, the treated arm $(1, 0_{N_i})$ requires the unit to be treated and all its neighbors to be untreated. Under Bernoulli assignment,

$$\pi_i(1, 0_{N_i}) = p(1-p)^{d_i}, \qquad \pi_i(0, 0_{N_i}) = (1-p)^{d_i+1},$$

Table 1: Direct-effect ACEs under Bernoulli assignment. The rows show the DTE target under arbitrary interference, neighborhood interference, and count/fraction or binary exposure representations.

Representation	Target conditions	HT H	LB H	HT variance scale
Arbitrary interference	$(1, \mathbf{0}_{-i})$ vs. $(0, \mathbf{0}_{-i})$	$e^{\Theta(n)}$	$e^{\Theta(n)}$	$\mathcal{V}_n(e^{\Theta(n)}, n)$
Neighborhood interference	$(1, \mathbf{0}_{N_i})$ vs. $(0, \mathbf{0}_{N_i})$	$e^{\Theta(d_{\max})}$	$e^{\Theta(d_{\max})}$	$\mathcal{V}_n(e^{\Theta(d_{\max})}, b_{2,\text{avg}})$
Count/fraction exposure	$(Z_i, S_i) = (1, 0)$ vs. $(0, 0)$	$e^{\Theta(d_{\max})}$	$e^{\Theta(d_{\max})}$	$\mathcal{V}_n(e^{\Theta(d_{\max})}, b_{2,\text{avg}})$
Binary any-treated exposure	$(Z_i, U_i) = (1, 0)$ vs. $(0, 0)$	$e^{\Theta(d_{\max})}$	$e^{\Theta(d_{\max})}$	$\mathcal{V}_n(e^{\Theta(d_{\max})}, b_{2,\text{avg}})$

Table 2: Total-effect ACEs under Bernoulli assignment. Exposure-mapping rows show representative treated-exposure versus no-exposure targets.

Representation	Target conditions	HT H	LB H	HT variance scale
Arbitrary interference	$(1, \mathbf{1}_{-i})$ vs. $(0, \mathbf{0}_{-i})$	$e^{\Theta(n)}$	$e^{\Theta(n)}$	$\mathcal{V}_n(e^{\Theta(n)}, n)$
Neighborhood interference	$(1, \mathbf{1}_{N_i})$ vs. $(0, \mathbf{0}_{N_i})$	$e^{\Theta(d_{\max})}$	$e^{\Theta(d_{\max})}$	$\mathcal{V}_n(e^{\Theta(d_{\max})}, b_{2,\text{avg}})$
Count/fraction exposure	$(Z_i, S_i) = (1, d_i)$ vs. $(0, 0)$	$e^{\Theta(d_{\max})}$	$e^{\Theta(d_{\max})}$	$\mathcal{V}_n(e^{\Theta(d_{\max})}, b_{2,\text{avg}})$
Binary any-treated exposure	$(Z_i, U_i) = (1, 1)$ vs. $(0, 0)$	$e^{\Theta(d_{\max})}$	$e^{\Theta(d_{\max})}$	$\mathcal{V}_n(e^{\Theta(d_{\max})}, b_{2,\text{avg}})$

so $H_i = \Theta(e^{\gamma d_i})$ with $\gamma = -\log(1-p)$ and $H = e^{\Theta(d_{\max})}$. The same exponential scale applies under count, fraction, and binary any-treated exposure mappings for the DTE, since all encode the no-exposure condition as $S_i = 0$ or $U_i = 0$.

Total treatment effect. The relevant potential outcomes for the TTE are $Y_i(1, \mathbf{1}_{N_i})$ and $Y_i(0, \mathbf{0}_{N_i})$ (Definition 6). The treated arm requires all neighbors to be treated and the control arm requires all neighbors to be untreated, so H is exponential in $d_{\max}$ for every representation in Table 2. Crucially, even the binary any-treated exposure mapping does not remove this scale: the control arm $(0, 0)$ requires all neighbors to be untreated, with propensity $(1-p)^{d_i+1}$ decaying exponentially with d_i.

Implications for practice. Under classical graph-unaware designs, the relevant potential outcomes for the DTE—$Y_i(1, 0)$ and $Y_i(0, 0)$ (Definition 6)—occur with probability $\exp\{-\Theta(d_i)\}$, leading to an inverse-propensity barrier exponential in the neighborhood degree. For the DTE,

$$H \asymp e^{\gamma d_{\max}}, \qquad \gamma = -\log(1-p),$$

and the variance bound vanishes only when $d_{\max} \lesssim \frac{1}{2\gamma} \log n$, i.e., when the degree grows at most logarithmically in n. This condition is necessary for any unbiased estimator, not just the HT (Proposition 11). For the TTE, the exponential barrier applies across *all* exposure mappings (Table 2) because both relevant potential outcomes—$Y_i(1, \mathbf{1})$ and $Y_i(0, \mathbf{0})$—require every neighbor to be in a specific treatment state.

To overcome this exponential barrier induced by classical designs, we need specialized graph aware designs to estimate ACEs. Section 6 develops such designs: independent-set designs isolate the relevant potential outcomes for the DTE on ego units, while cluster randomized designs increase the probability of the all-treated and all-control conditions needed for the TTE.

4.3 ECE estimation

We now consider estimation of ECEs. Following the definition in Section 3.2, let $\mathcal{Q}^+, \mathcal{Q}^-$ be two allocation policies and $D_i^+(\mathbf{Z}), D_i^-(\mathbf{Z}) \in \mathcal{X}_i$ the corresponding random treatment–exposure combinations. Define the policy propensities scores as

$$\lambda_i^+(d) = \mathcal{Q}^+\{D_i^+(\mathbf{Z}) = d\}, \qquad \lambda_i^-(d) = \mathcal{Q}^-\{D_i^-(\mathbf{Z}) = d\},$$

so that the ECE can be written as

$$\tau_{\mathrm{ECE}} = \frac{1}{n}\sum_{i=1}^{n}\sum_{d\in\mathcal{X}_i} Y_i(d)\{\lambda_i^+(d) - \lambda_i^-(d)\}.$$

Its Horvitz–Thompson estimator has the same form as the ACE estimator of Section 4.2, with the fixed inverse propensity scores of the relevant potential outcomes replaced by importance weights that are ratios of the policy propensity scores to the design propensity scores,

$$h_i^+(\mathbf{z}) = \frac{\lambda_i^+\{D_i(\mathbf{z})\}}{\pi_i^P\{D_i(\mathbf{z})\}}, \qquad h_i^-(\mathbf{z}) = \frac{\lambda_i^-\{D_i(\mathbf{z})\}}{\pi_i^P\{D_i(\mathbf{z})\}},$$

with the convention $0/0 = 0$. Define the diagonal scale

$$V_i = \sum_{d:\pi_i^P(d)>0} \frac{\{\lambda_i^+(d) - \lambda_i^-(d)\}^2}{\pi_i^P(d)},$$

and the weight bound $H_i = \sup_{\mathbf{z}}\{|h_i^+(\mathbf{z})| + |h_i^-(\mathbf{z})|\}$.

Proposition 12 (ECE importance-weighted estimator) *Suppose $\lambda_i^+(d) = \lambda_i^-(d) = 0$ whenever $\pi_i^P(d) = 0$. Then*

$$\hat{\tau}_{\mathrm{ECE}}^{HT} = \frac{1}{n}\sum_{i=1}^{n}\{h_i^+(\mathbf{Z}) - h_i^-(\mathbf{Z})\}Y_i^{obs}$$

is design-unbiased for τ_{ECE} for every potential-outcome table. If $|Y_i(\mathbf{z})| \le M$ and the ith summand depends on the assignment only through $\mathbf{Z}_{\mathcal{C}_i}$, then

$$\mathrm{Var}(\hat{\tau}_{\mathrm{ECE}}^{HT}) \le \frac{M^2}{n^2}\left\{\sum_{i=1}^{n} V_i + \sum_{i=1}^{n}\sum_{j\neq i} c_{ij}H_iH_j + \sum_{i=1}^{n}\sum_{j\neq i}(1-c_{ij})\rho_{ij}H_iH_j\right\}.$$

If $V_i \le V$ and $H_i \le H$ for all i, then

$$\mathrm{Var}(\hat{\tau}_{\mathrm{ECE}}^{HT}) \le M^2\left\{\frac{V}{n} + H^2\frac{c_{\mathrm{avg}} + r_{\mathrm{avg}}}{n}\right\},$$

and the estimator is design-MSE consistent if $V + H^2(c_{\mathrm{avg}} + r_{\mathrm{avg}}) = o(n)$.

The proof is given in Appendix D. The bound has the same three terms as the ACE bound of Proposition 10: a diagonal term, an interference-dependence term, and a residual design-dependence term. What changes is the diagonal scale. For an ACE, the policy is a point mass at the relevant potential outcomes $Y_i(d_i^{\pm})$. The importance weights reduce to inverse propensities, and V_i and H_i are of the order of $1/\pi_i(d_i^{\pm})$. For an ECE, V_i and H_i are governed by the ratios $\lambda_i^{\pm}(d)/\pi_i^P(d)$, and these are large only when the allocation policy places probability on combinations that the design rarely produces. When the allocation policy is the design itself, the ratios are bounded, the diagonal term is $O(1/n)$, and the variance is governed by the dependence terms alone. The next two results make this concrete for the expected direct and indirect effects under Bernoulli policy and design.

Expected direct effect. Under Bernoulli assignment with treatment probabilities p_i, consider the EDE whose policy is the same Bernoulli allocation. The importance weight reduces to the classical Horvitz–Thompson weight

$$h_i^{\mathrm{EDE}}(\mathbf{Z}) = \frac{Z_i}{p_i} - \frac{1 - Z_i}{1 - p_i},$$

so $\hat{\tau}^{HT}_{\mathrm{EDE}}$ is the usual difference in weighted means and does not require the interference graph for computation. The weight scales are $V_i = H_i = 1/p_i + 1/(1-p_i)$, which are $O(1)$ when the treatment probabilities are bounded away from zero and one. No inverse propensity of a rare exposure combination appears in the bound.

Proposition 13 (Variance bound for the EDE estimator) *Let $\mathcal{Q}_p$ be a Bernoulli allocation policy, with treatment probabilities bounded away from zero and one. Assume Bernoulli assignment, and $\sup_{\mathbf{z}} |Y_i(\mathbf{z})| \le M$. Then $\hat{\tau}^{HT}_{\mathrm{EDE}}$ is design-unbiased for $\tau_{\mathrm{EDE}}(\mathcal{Q}_p)$, and*

$$\mathrm{Var}_{P_n}(\hat{\tau}^{HT}_{\mathrm{EDE}}) \le C_p M^2 \frac{1 + c_{\mathrm{avg}}}{n},$$

where C_p depends only on the bounds on the treatment probabilities. In particular, $\hat{\tau}^{HT}_{\mathrm{EDE}}$ is design-MSE consistent if $c_{\mathrm{avg}} = o(n)$.

The proof is given in Appendix D. This result should be read against the DTE rows of Table 1. The two estimators use the same design and the same interference graph. The DTE estimator must isolate the treatment–exposure combination $(1, 0_{N_i})$, which the design produces with probability $p(1-p)^{d_i}$, so its inverse-propensity scale is $H = e^{\Theta(d_{\max})}$ and its full variance scale from Proposition 10 is $H/n + H^2 c_{\mathrm{avg}}/n$. The EDE estimator uses every observed outcome, and its variance scale is $(1 + c_{\mathrm{avg}})/n$. On a d-regular locally tree-like interference graph the comparison is $(e^{2\gamma d} d^2)/n$ against d^2/n: the DTE is consistent only in the logarithmic-degree regime $e^{2\gamma d} d^2 = o(n)$, i.e., $d \lesssim \frac{1}{2\gamma} \log n$, while the EDE HT estimator is consistent whenever $d = o(\sqrt{n})$. The price of interference for the EDE is not the rarity of any treatment-exposure combination but the dependence among the summands, measured by c_{avg}.

Expected indirect effect. Graph-aware ECEs that contrast neighbor treatments have larger importance weights. The expected indirect effect sums the weight over neighbors, $h_i^{\mathrm{EIE},G}(\mathbf{Z}) = \sum_{j \in N_i} \{Z_j/p_j - (1-Z_j)/(1-p_j)\}$, so $H_i \le C_p d_i$ when the treatment probabilities are bounded away from zero and one. The weights grow with the degree, but linearly rather than exponentially.

Corollary 14 (Variance bound for the graph-aware EIE estimator) *Assume the true potential outcomes satisfy neighborhood interference with known neighborhoods N_i. Assume Bernoulli assignment with treatment probabilities bounded away from zero and one, and assume* $\sup_{\mathbf{z}} |Y_i(\mathbf{z})| \le M$. *Then $\hat{\tau}^G_{\mathrm{EIE}}$ is design-unbiased for $\tau_{\mathrm{EIE}}(\pi)$, and*

$$\mathrm{Var}(\hat{\tau}^G_{\mathrm{EIE}}) \le \frac{C_p M^2}{n^2} \left\{ \sum_i d_i^2 + \sum_i \sum_{j \ne i} c_{ij} d_i d_j \right\},$$

where $d_i = |N_i|$ and $c_{ij} = \mathbb{1}\{\mathcal{C}_i \cap \mathcal{C}_j \ne \emptyset\}$. If $d_i \le d_{\max}$ for all i, then

$$\mathrm{Var}(\hat{\tau}^G_{\mathrm{EIE}}) \le C_p M^2 d_{\max}^2 \frac{1 + c_{\mathrm{avg}}}{n}.$$

The proof is given in Appendix D. The variance scale $d_{\max}^2(1 + c_{\mathrm{avg}})/n$ is polynomial in the degrees: on a d-regular locally tree-like graph it is of order d^4/n, so the EIE is consistently estimable whenever $d = o(n^{1/4})$. As for the EDE, no exponential barrier due to the propensity score appears, and the dependence term dominates the bound.

Appendix D.11 extends the DTE/EDE comparison to arbitrary heterogeneous degree sequences, where it is governed by two graph functionals: exponential moments of the degree sequence for fixed-exposure ACEs, and the average dependence degree for the EDE.

The variance bounds of Section 4 show that classical graph-unaware designs are matched to expected causal effects and pay an exponential cost for average causal effects. Section 5 constructs designs that remove this cost, and Section 6 collects the resulting implications for classical estimators, model-assisted estimation, and the interpretation of positivity under interference.

5 Design-Based Estimation of Average Causal Effects

When classical estimators fail for average causal effects under interference, the failure arises because the realized assignment does not reveal the relevant potential outcomes with sufficient probability. Section 4 shows that under classical graph-unaware designs the target treatment–exposure combinations have inverse propensities that are exponential in the interference degree, and that this cost is unavoidable for any unbiased estimator. One remedy is to modify the experimental design so that those potential outcomes are observed with non-negligible probability. We study designs that control the assignment structure rather than impose an outcome model. The common device is to use the interference graph to create positive correlation within neighborhoods, so that the event that all of a unit's neighbors are held at a common treatment occurs as a single correlated event rather than as a product of independent events. This replaces the exponential propensity $(1-p)^{d_i}$ of a classical design by a propensity that decays only polynomially in the degree d_i, and the variance is reduced accordingly. Throughout, we assume that the interference neighborhood is known and that an exposure mapping is specified, and we take the potential outcomes to be uniformly bounded, $\sup_{i,z,e} |Y_i(z,e)| \le M$.

Suppose the target average causal effect has the form

$$\tau = \frac{1}{n} \sum_{i=1}^{n} \{Y_i(d_i^+) - Y_i(d_i^-)\},$$

where d_i^+ and d_i^- are two target treatment–exposure combinations. If the design gives positive probability to both combinations,

$$\pi_i(d_i^+) = \Pr\{(Z_i, E_i) = d_i^+\} > 0, \qquad \pi_i(d_i^-) = \Pr\{(Z_i, E_i) = d_i^-\} > 0,$$

then the treatment–exposure Horvitz–Thompson estimator

$$\hat{\tau}^{ZE}_{\mathrm{HT}} = \frac{1}{n}\sum_{i=1}^{n}\left\{\frac{\mathbb{1}\{(Z_i, E_i) = d_i^+\}Y_i^{\mathrm{obs}}}{\pi_i(d_i^+)} - \frac{\mathbb{1}\{(Z_i, E_i) = d_i^-\}Y_i^{\mathrm{obs}}}{\pi_i(d_i^-)}\right\}$$

is unbiased for τ. The role of the design is therefore to make the relevant treatment–exposure combinations occur with positive and reasonably large propensity scores.

Proposition 15 (Treatment–exposure HT unbiasedness for ACEs) *Suppose the observed outcome satisfies $Y_i^{\mathrm{obs}} = Y_i(Z_i, E_i)$. Let*

$$\tau = \frac{1}{n}\sum_{i=1}^{n}\{Y_i(d_i^+) - Y_i(d_i^-)\},$$

and suppose $\pi_i(d_i^+) > 0$ and $\pi_i(d_i^-) > 0$ for all i. Then $\hat{\tau}^{ZE}_{\mathrm{HT}}$ is design-unbiased for τ.

The proof is given in Appendix E.

We develop two graph-aware constructions based on independent sets. The first targets the direct treatment effect and uses a distance-one independent set; the second targets the total treatment effect and uses a distance-two independent set. In each case a randomized greedy construction gives focal inclusion probabilities in closed form, and the resulting Horvitz–Thompson estimator has mean squared error that is polynomial, rather than exponential, in the interference degree. We then discuss cluster randomized designs, whose polynomial behavior is conditional on the clustering being aligned with the graph.

5.1 Independent set designs for the direct treatment effect

The direct treatment effect

$$\tau^D = \frac{1}{n}\sum_{i=1}^{n}\{Y_i(1, \mathbf{0}_{N_i}) - Y_i(0, \mathbf{0}_{N_i})\}$$

contrasts the treated and control potential outcomes of a unit whose neighbors are all held at control. Both arms require the no-exposure condition $Z_{N_i} = \mathbf{0}$; they differ only in the unit's own treatment. Consequently two units i and j that share a neighbor $k \in N_i \cap N_j$ impose the same requirement $Z_k = 0$, so there is no conflict between their target conditions. A distance-one independent set, in which selected units are mutually non-adjacent but may share neighbors, therefore suffices.

We construct a random independent set of *ego* units by a randomized greedy rule, and randomize treatment among the egos while holding their neighbors at control. Units outside the independent set act as *alter* units that buffer the egos from interference.

Design (IS-1).

(1) Draw $U_1, \ldots, U_n$ independently from the uniform distribution on $[0, 1]$.

(2) Let $\mathcal{I} = \{i : U_i < U_k \text{ for all } k \in N_i\}$ be the set of units whose draw is the strict minimum in their closed neighborhood $\{i\} \cup N_i$. The set $\mathcal{I}$ is an independent set almost surely.

(3) For each ego $i \in \mathcal{I}$, draw $B_i \sim \text{Bernoulli}(q)$ independently, set $Z_i = B_i$, and set $Z_k = 0$ for all $k \in N_i$.

(4) Assign any unit not fixed by step (3) a treatment by a prespecified rule that does not overwrite the constraints of step (3).

The construction is graph-aware: whether i is an ego depends only on the draws in its closed neighborhood. If two egos i and j share a neighbor k, step (3) sets $Z_k = 0$ from both, so the assignment is consistent regardless of how many egos claim k. Because the event $\{i \in \mathcal{I}\}$ is determined within $\{i\} \cup N_i$, the inclusion probability is available in closed form.

Lemma 16 (Inclusion and reveal probabilities under IS-1) *Under the design IS-1,*

$$\Pr(i \in \mathcal{I}) = \frac{1}{d_i + 1}.$$

Let $F_i^{(1)} = \{i \in \mathcal{I},\ B_i = 1\}$ *and* $F_i^{(0)} = \{i \in \mathcal{I},\ B_i = 0\}$ *be the focal reveal events. On* $F_i^{(1)}$ *the assignment satisfies* $(Z_i, Z_{N_i}) = (1, \mathbf{0})$*, and on* $F_i^{(0)}$ *it satisfies* $(Z_i, Z_{N_i}) = (0, \mathbf{0})$*, so*

$$\rho_i^{(1)} := \Pr(F_i^{(1)}) = \frac{q}{d_i + 1}, \qquad \rho_i^{(0)} := \Pr(F_i^{(0)}) = \frac{1 - q}{d_i + 1}.$$

The proof is given in Appendix E. The inclusion probability follows from the symmetry of the $d_i + 1$ independent draws in the closed neighborhood, and the reveal probabilities follow from the independent treatment coin B_i. The focal reveal probabilities are exact and do not require simulation. They are stated for the focal events rather than for the raw exposure events $\{Z_i = s, Z_{N_i} = \mathbf{0}\}$, because non-ego units may also happen to satisfy the no-exposure condition depending on the leftover rule in step (4); the estimator below uses only the focal events, whose probabilities are known exactly. Taking $q = \frac{1}{2}$ gives the symmetric values $\rho_i^{(1)} = \rho_i^{(0)} = 1/\{2(d_i + 1)\}$.

The focal Horvitz–Thompson estimator for the direct treatment effect is

$$\hat{\tau}_{\text{IS}}^{D} = \frac{1}{n} \sum_{i=1}^{n} \left\{ \frac{\mathbb{1}\{F_i^{(1)}\} Y_i^{\text{obs}}}{\rho_i^{(1)}} - \frac{\mathbb{1}\{F_i^{(0)}\} Y_i^{\text{obs}}}{\rho_i^{(0)}} \right\},$$

which at $q = \frac{1}{2}$ reduces to

$$\hat{\tau}_{\text{IS}}^{D} = \frac{2}{n} \sum_{i=1}^{n} (d_i + 1) \left\{ \mathbb{1}\{F_i^{(1)}\} Y_i^{\text{obs}} - \mathbb{1}\{F_i^{(0)}\} Y_i^{\text{obs}} \right\}.$$

Proposition 17 (Unbiasedness of the IS-1 estimator) *Under the design IS-1 with $q \in (0,1)$, the estimator $\hat{\tau}^D_{\mathrm{IS}}$ is design-unbiased for the direct treatment effect, $\mathbb{E}[\hat{\tau}^D_{\mathrm{IS}}] = \tau^D$.*

The proof is the same Horvitz–Thompson calculation as Proposition 15, applied to the focal reveal events rather than to the raw treatment–exposure events. On $F_i^{(1)}$ and $F_i^{(0)}$ the observed outcome equals the corresponding direct-effect potential outcome, and the focal probabilities $\rho_i^{(1)}, \rho_i^{(0)}$ are positive and known.

The variance of the estimator is controlled by the closed-neighborhood inclusion probabilities and by the dependence among egos that lie within distance two of one another. Write $\delta_i = Y_i(1, \mathbf{0}_{N_i}) - Y_i(0, \mathbf{0}_{N_i})$ and $p_i = \Pr(i \in \mathcal{I}) = 1/(d_i + 1)$.

Proposition 18 (Variance of the IS-1 estimator) *Under the design IS-1 with $q = \frac{1}{2}$, the estimator satisfies the exact identity*

$$n^2 \operatorname{Var}(\hat{\tau}^D_{\mathrm{IS}}) = \sum_{i=1}^{n} \operatorname{Var}(W_i) + \sum_{i \neq j:\ \operatorname{dist}_G(i,j) \leq 2} \delta_i \delta_j \left(\frac{\Pr(i, j \in \mathcal{I})}{p_i p_j} - 1 \right),$$

where $W_i = 2(d_i + 1)\mathbb{1}\{i \in \mathcal{I}\}\{Y_i(1, \mathbf{0}_{N_i})B_i - Y_i(0, \mathbf{0}_{N_i})(1 - B_i)\}$. For non-adjacent i, j sharing $s_{ij} = |N_i \cap N_j|$ common neighbors, $\Pr(i, j \in \mathcal{I})/(p_i p_j) = (d_i + d_j + 2)/(d_i + d_j + 2 - s_{ij})$, while $\Pr(i, j \in \mathcal{I}) = 0$ when i and j are adjacent. Consequently, uniformly over $|Y_i(\cdot)| \leq M$,

$$\operatorname{Var}(\hat{\tau}^D_{\mathrm{IS}}) \leq \frac{4M^2}{n^2} \sum_{i=1}^{n} (d_i + 1) + \frac{4M^2}{n^2} \left[2|E| + \sum_{i \neq j:\ i \not\sim j,\ s_{ij} \geq 1} \frac{s_{ij}}{d_i + d_j + 2 - s_{ij}} \right].$$

The proof is given in Appendix E. The first term is a diagonal contribution of order $n^{-2} \sum_i (d_i + 1)$; the second collects the dependence among egos whose closed neighborhoods overlap, that is, egos within graph distance two.

Corollary 19 (Polynomial rate for the direct treatment effect) *If G is d-regular and locally tree-like, in the sense that every pair of non-adjacent vertices has at most a bounded number of common neighbors ($s_{ij} = O(1)$), then*

$$\operatorname{Var}(\hat{\tau}^D_{\mathrm{IS}}) = O\left(\frac{M^2 d}{n} \right).$$

For a general d-regular graph, the second-moment bound gives the cruder rate

$$\operatorname{Var}(\hat{\tau}^D_{\mathrm{IS}}) = O\left(\frac{M^2 d^2}{n} \right),$$

since on any d-regular graph the crude bound evaluates to $4M^2(n + nd + nd^2)/n^2$. For an arbitrary graph, the crude second-moment bound gives

$$\operatorname{Var}(\hat{\tau}^D_{\mathrm{IS}}) \leq \frac{4M^2}{n^2} \left[n + \sum_i d_i + \sum_i d_i^2 \right] = O\left(\frac{M^2 (1 + \bar{d} + \overline{d^2})}{n} \right), \qquad \overline{d^2} = \frac{1}{n} \sum_i d_i^2.$$

The proof is given in Appendix E. Under a Bernoulli design the same estimand requires the no-exposure condition, whose propensity $p(1-p)^{d_i}$ is exponentially small in d_i, and the resulting variance is exponential in the average degree by Proposition 11. The IS-1 design replaces this exponential cost by dependence that is polynomial in the degree sequence. The bound of Corollary 19 holds for arbitrary, not necessarily regular, interference graphs: it is of order $M^2(1+\bar{d}+\overline{d^2})/n$, so this estimator is consistent whenever $\overline{d^2} = o(n)$. This is the same second-moment condition on the degree sequence under which the expected direct effect is consistent by the generic dependence-graph bound. On d-regular graphs that are locally tree-like, in the sense that non-adjacent vertices share at most a bounded number of common neighbors, the sharper rate $O(M^2\bar{d}/n)$ holds. On a general d-regular graph the wedge sum $\sum_{i\neq j:\ i\not\sim j} s_{ij}/(d_i+d_j+2-s_{ij})$ can be as large as $\Theta(nd^2)$ when closed neighborhoods overlap heavily, so the crude rate $O(M^2d^2/n)$ is the best that the second-moment bound certifies in general. Figure 3 illustrates the effect: on a fixed family of graphs, the error of the independent set estimator remains nearly flat in the degree, whereas under a Bernoulli design it grows exponentially. For irregular graphs with high-degree hubs the crude bound can be larger, and a hub-aware modification of the focal sampling may control it, though we do not pursue this refinement here. For extreme hub graphs such as a star, this particular randomized local-minimum IS-1 design has a non-vanishing second-moment bound; this indicates a limitation of the simple construction rather than a universal impossibility theorem.

5.2 Independent set designs for the total treatment effect

The total treatment effect

$$\tau^T = \frac{1}{n}\sum_{i=1}^{n}\{Y_i(1, e_i^{\mathbf{1}}) - Y_i(0,0)\}, \qquad e_i^{\mathbf{1}} = f_i(\mathbf{1}_{N_i}),$$

contrasts the all-treated and all-control regimes on the closed neighborhood $\{i\}\cup N_i$. Unlike the direct treatment effect, the two arms now impose opposite requirements on a shared neighbor: if egos i and j share a neighbor k and i is on the treated side while j is on the control side, then i requires $Z_k = 1$ and j requires $Z_k = 0$, a direct conflict. Avoiding such conflicts requires that egos have disjoint closed neighborhoods, that is, that they be at graph distance at least three. This is a distance-two independent set, obtained by running the randomized greedy rule on the radius-two balls.

For a unit i let $B_i^{(2)} = \{k : \mathrm{dist}_G(i,k) \le 2\}$ denote its radius-two ball and $b_i = |B_i^{(2)}|$ its size. Since the neighborhood is contained in the ball,

$$d_i + 1 \le b_i \le 1 + d_i + \sum_{j\in N_i}(d_j - 1) \le 1 + d_i d_{\max}.$$

On d-regular locally tree-like graphs this gives $b_i = 1 + d + d(d-1) = d^2 + 1$.

Design (IS-2).

(1) Draw $U_1, \ldots, U_n$ independently from the uniform distribution on $[0,1]$.

(2) Let $\mathcal{I}_2 = \{i : U_i < U_k \text{ for all } k \in B_i^{(2)} \setminus \{i\}\}$ be the set of units whose draw is the strict minimum in their radius-two ball. The set $\mathcal{I}_2$ is a distance-two independent set almost surely.

(3) For each ego $i \in \mathcal{I}_2$, draw $R_i \sim \text{Bernoulli}(\frac{1}{2})$ independently, and set $Z_k = R_i$ for all $k \in \{i\} \cup N_i$.

(4) Assign any unit not fixed by step (3) an independent $\text{Bernoulli}(\frac{1}{2})$ treatment.

Because egos are at distance at least three, their closed neighborhoods are disjoint and step (3) is free of conflicts. Since the outcome of an ego depends only on the assignment of its closed neighborhood, which is entirely fixed in step (3), the leftover assignments in step (4) do not affect the revealed outcomes of egos.

Lemma 20 (Inclusion and reveal probabilities under IS-2) *Under the design IS-2,*

$$\Pr(i \in \mathcal{I}_2) = \frac{1}{b_i}, \qquad \rho_i^{(1)} := \Pr(i \in \mathcal{I}_2, R_i = 1) = \frac{1}{2b_i}, \qquad \rho_i^{(0)} := \Pr(i \in \mathcal{I}_2, R_i = 0) = \frac{1}{2b_i}.$$

On these focal events the assignment reveals the corresponding all-treated or all-control potential outcome. In particular, if the radius-two balls are polynomially bounded in the interference degree scale, the inverse focal probabilities are polynomial as well.

The proof is given in Appendix E. The focal Horvitz–Thompson estimator for the total treatment effect is

$$\hat{\tau}_{\text{IS}}^T = \frac{2}{n} \sum_{i \in \mathcal{I}_2} b_i \left\{ \mathbb{1}\{R_i = 1\} Y_i^{\text{obs}} - \mathbb{1}\{R_i = 0\} Y_i^{\text{obs}} \right\}.$$

Proposition 21 (Unbiasedness and variance of the IS-2 estimator) *Under the design IS-2, the estimator $\hat{\tau}_{\text{IS}}^T$ is design-unbiased for the total treatment effect. The unbiasedness is with respect to the focal reveal events $(i \in \mathcal{I}_2, R_i = 1)$ and $(i \in \mathcal{I}_2, R_i = 0)$, each of which implies the corresponding treatment–exposure condition. Writing $\Delta_i = Y_i(1, e_i^{\mathbf{1}}) - Y_i(0,0)$ and $q_i = 1/b_i$, it satisfies the exact identity*

$$n^2 \operatorname{Var}(\hat{\tau}_{\text{IS}}^T) = \sum_{i=1}^n \operatorname{Var}(V_i) + \sum_{i \neq j:\ B_i^{(2)} \cap B_j^{(2)} \neq \emptyset} \Delta_i \Delta_j \left(\frac{\Pr(i, j \in \mathcal{I}_2)}{q_i q_j} - 1 \right),$$

where $V_i = 2b_i \mathbb{1}\{i \in \mathcal{I}_2\}\{Y_i(1, e_i^{\mathbf{1}}) R_i - Y_i(0,0)(1 - R_i)\}$. If $\text{dist}_G(i,j) \leq 2$, then $\Pr(i, j \in \mathcal{I}_2) = 0$; if $\text{dist}_G(i,j) \in \{3,4\}$ and $s = |B_i^{(2)} \cap B_j^{(2)}| \geq 1$, then the ratio equals $(b_i + b_j)/(b_i + b_j - s)$. Consequently, uniformly over $|Y_i(\cdot)| \leq M$,

$$\operatorname{Var}(\hat{\tau}_{\text{IS}}^T) \leq \frac{4M^2}{n^2} \sum_{i=1}^n b_i + \frac{4M^2}{n^2} \sum_{i \neq j:\ \text{dist}_G(i,j) \leq 2} 1 + \frac{4M^2}{n^2} \sum_{i \neq j:\ \text{dist}_G(i,j) \in \{3,4\},\ s_{ij}^{(2)} \geq 1} \frac{s_{ij}^{(2)}}{b_i + b_j - s_{ij}^{(2)}}, \qquad s_{ij}^{(2)} = |B_i^{(2)} \cap B_j^{(}$$

The proof is given in Appendix E.

Corollary 22 (Polynomial rate for the total treatment effect) *On d-regular graphs that are locally tree-like, the diagonal and distance-two exclusion terms are of order M^2d^2/n, and the positive radius-two overlap term contributes at most the same order (distance-three partners contribute M^2d/n and distance-four partners M^2d^2/n), so*

$$\mathrm{Var}(\hat{\tau}^T_{\mathrm{IS}}) = O\left(\frac{M^2d^2}{n}\right).$$

For an arbitrary graph, the universal bound is $\mathrm{Var}(\hat{\tau}^T_{\mathrm{IS}}) = O(M^2\overline{b^2}/n)$ *with* $\overline{b^2} = n^{-1}\sum_i b_i^2$.

Since $d_i \le d_{\max}$ and $b_i \le 1 + d_i d_{\max}$, the bounds in Corollaries 19 and 22 give variance bounds of order $O(M^2(1+d_{\max}^2)/n)$ for IS-1 and $O(M^2(1+d_{\max}^4)/n)$ for IS-2 on every interference graph. The proof is given in Appendix E. Under a Bernoulli design the total treatment effect requires the all-treated condition on the closed neighborhood, whose propensity p^{d_i+1} is again exponentially small in d_i. On locally tree-like d-regular graphs, the IS-2 design attains the polynomial rate $O(M^2d^2/n)$, the same degree order as the expected direct effect under a Bernoulli design established in Proposition 13. On arbitrary graphs the rate is governed by the radius-two ball sizes b_i. The exponential separation between average and expected causal effects under classical designs is therefore not intrinsic to the estimands: with a design aligned to the estimand, the total treatment effect can be estimated at polynomial rates. The direct treatment effect is easier still, since Corollary 19 gives the smaller rate $O(M^2\bar{d}/n)$ on locally tree-like regular graphs.

5.3 Independent set designs for indirect effects

Indirect (spillover) effects ask how the outcome of a unit is affected by the treatment of others, its own treatment held fixed. A recurring target in network experiments is the effect of treating a *selected set of influencers* on an average unit: seeding experiments assign treatment to a sparse set of influential units and measure the response of otherwise untreated peers exposed to a treated influencer (Paluck et al., 2016; Aral and Walker, 2012; Aral et al., 2013; Bond et al., 2012). Under the any-treated exposure map (3.6), which records $U_i = 1$ whenever at least one neighbor of i is treated, this target is the two-level average indirect effect $\tau_{\mathrm{IE}}(z;0,1) = n^{-1}\sum_i\{Y_i(z,1) - Y_i(z,0)\}$ of Section 3.1: a comparison of the no-exposure level $e_a = 0$ and the any-treated level $e_b = 1$ with the unit's own treatment z fixed ($z = 0$ for an untreated peer). Estimators of spillover effects under two-stage and independent-set randomization are developed in Baird et al. (2018); Basse and Feller (2018); Cai et al. (2024), and two-level exposure contrasts of this form are the fixed-exposure analogue of the allocation-averaged estimand class studied by Viviano (2026). The uncompressed special cases contrast the fully exposed and unexposed neighborhood regimes, with the unit's own treatment held fixed:

$$\tau_{\mathrm{IEC}} = \frac{1}{n}\sum_{i=1}^{n}\{Y_i(0,e_i^{\mathbf{1}}) - Y_i(0,0)\}, \qquad \tau_{\mathrm{IET}} = \frac{1}{n}\sum_{i=1}^{n}\{Y_i(1,e_i^{\mathbf{1}}) - Y_i(1,0)\}.$$

Both arms of an indirect contrast fix an extreme neighborhood regime, and the exponential barrier of Section 4 applies exactly as it does to the direct and total effects. Under a Bernoulli design with probability p, the unexposed arm has propensity $\Pr(Z_i = z^\dagger)(1-p)^{d_i}$,

the any-treated arm has propensity $\Pr(Z_i = z^\dagger)(1-(1-p)^{d_i})$, and the full-exposure arm has propensity $\Pr(Z_i = z^\dagger)\, p^{d_i}$, where $z^\dagger$ is the own-treatment level of the estimand and all but the any-treated propensity are exponentially small in the degree. Every indirect effect that fixes a rare exposure regime is therefore exponentially costly under a graph-unaware design.

The distance-two independent set construction of Section 5.2 removes this barrier for the indirect effects with a single modification. On the ego set $\mathcal{I}_2$, hold each ego's own treatment at the level $z^\dagger$ required by the estimand ($z^\dagger = 0$ for τ_{IEC} and for control units; $z^\dagger = 1$ for τ_{IET}), and assign the ego's neighborhood the common value $R_i \in \{0,1\}$ by an independent fair coin. An ego then reveals $(z^\dagger, e_i^{\mathbf{1}})$ when $R_i = 1$ and $(z^\dagger, 0)$ when $R_i = 0$, each with focal probability $1/(2b_i)$.

Proposition 23 (Independent set designs for indirect effects) *Under the modified distance-two independent set design above, the focal Horvitz–Thompson estimator is design-unbiased, with focal reveal probabilities $1/(2b_i)$ for both arms, for the any-treated indirect effects $\tau_{\mathrm{IE}}(0;0,1)$ and $\tau_{\mathrm{IE}}(1;0,1)$ under the exposure map (3.6), and for the full-exposure versions τ_{IEC} ($z^\dagger = 0$) and τ_{IET} ($z^\dagger = 1$). Its variance satisfies the bound of Proposition 21, and on locally tree-like d-regular graphs*

$$\mathrm{Var}(\hat{\tau}_{\mathrm{IS}}) = O\left(\frac{M^2\overline{d^2}}{n}\right),$$

matching the rate of the total treatment effect and of the expected direct effect on this graph class.

The proof is identical to that of Proposition 21: fixing the ego's own treatment changes neither the inclusion probabilities nor the covariance structure, and only the pair of revealed treatment–exposure arms differs. No modification of the design is needed for the coarsened two-level estimand. Because the all-treated arm realizes the any-treated level $U_i = 1$ and the all-control arm realizes $U_i = 0$ under (3.6), the polarized design estimates $\tau_{\mathrm{IE}}(z^\dagger;0,1)$ with the same focal probabilities $1/(2b_i)$. More generally, any monotone exposure map that takes the value 1 on the all-treated vector and 0 on the all-control vector yields the same two realized levels, so the same estimator applies to the corresponding coarsened two-level indirect effect. The coarsening is therefore free for the polarized design: it changes only the exposure interpretation, not the assignment, and it can only shorten the distance between the two levels of the contrast.

Sparse seeding with the distance-one independent set. Influencer-seeding experiments treat a sparse, well-separated set of units and measure the responses of their non-treated alters. The distance-one independent set $\mathcal{I}$ of Section 5.1 is exactly such a set: egos are mutually non-adjacent, and on typical graphs each alter is adjacent to a controlled number of egos. Treating each ego with an independent coin $B_i \sim \mathrm{Bernoulli}(q)$ and setting all non-ego units to control yields the *sparse seeding design*: the treated influencers are the egos with $B_i = 1$, and an alter's any-treated indicator (3.6) equals 1 precisely when at least one of its neighbors is a treated ego. Let $K_i = |\mathcal{I} \cap N_i|$ be the number of ego neighbors of

unit i. Because I is an independent set, an ego has no ego neighbors; the unit i consequently reveals $(0,1)$ on the event

$$F_i^{(1)} = \Big\{ \textstyle\sum_{j\in N_i} B_j \mathbb{1}\{j \in \mathcal{I}\} \geq 1 \Big\} \qquad \text{i.e.} \qquad \{Z_i = 0,\ U_i = 1\},$$

and reveals $(0,0)$ on $F_i^{(0)} = \{Z_i = 0\} \setminus F_i^{(1)}$, with $F_i^{(1)} \cup F_i^{(0)} = \{Z_i = 0\}$. The focal probabilities are

$$\rho_i^{(1)} = \mathbb{E}\big[1 - (1-q)^{K_i}\big], \qquad \rho_i^{(0)} = \mathbb{E}\big[(1-q)^{K_i}\big] - \frac{q}{d_i + 1}. \tag{5.1}$$

Both are positive for every unit with degree at least one.

Proposition 24 (Sparse seeding estimates the influencer spillover effect) *Under the sparse seeding design with $q \in (0,1)$ and no isolated units, the focal Horvitz–Thompson estimator*

$$\hat{\tau}_{\text{seed}} = \frac{1}{n} \sum_{i=1}^{n} \Big\{ \frac{\mathbb{1}\{F_i^{(1)}\} Y_i^{\text{obs}}}{\rho_i^{(1)}} - \frac{\mathbb{1}\{F_i^{(0)}\} Y_i^{\text{obs}}}{\rho_i^{(0)}} \Big\}$$

is design-unbiased for the influencer spillover effect $\tau_{\text{IE}}(0;0,1)$ under the any-treated exposure map (3.6)*. On a locally tree-like d-regular graph, $K_i \Rightarrow \text{Poisson}(1)$ and, as $d \to \infty$,*

$$\rho_i^{(1)} \to 1 - e^{-q}, \qquad \rho_i^{(0)} \to e^{-q}:$$

both arms have constant focal probability, independent of n and of the degree, and the inverse-propensity weights $1/\rho_i^{(1)} + 1/\rho_i^{(0)}$ are $O(1)$. The covariance structure is inherited from the ego draws and has the radius-two range of Proposition 18, so the variance satisfies the generic bound of Proposition 10; on locally tree-like d-regular graphs the diagonal term is $O(M^2/n)$, removing the inverse-propensity barrier entirely for this estimand.

The proof is given in Appendix E. The constant-order focal probabilities contrast sharply with the focal probability $1/(2b_i) = O(1/d^2)$ of the full-exposure design of Proposition 23: the coarsening to the any-treated level $e_b = 1$ is what converts the exponentially rare all-treated event into a condition–"at least one treated neighbor"–that has constant probability under sparse seeding. The hub caveat is that the guarantee is tied to the ego-neighborhood count K_i being $O(1)$ in expectation. On a star, for example, the center faces up to d simultaneously eligible ego candidates among its leaves and $\rho_c^{(0)} = \int_0^1 (1-qu)^d \, du = O(1/(qd))$: the unexposed arm of the hub decays polynomially rather than staying constant, so the invertible focal weights grow linearly in the hub degree. The constant-propensity claim is therefore restricted to graph families in which ego-neighborhoods are small in expectation, such as locally tree-like d-regular graphs and, more generally, graphs of bounded degree. Thus the distance-one construction of Section 5.1, run as a seeding experiment, estimates the influencer spillover effect at constant focal propensities on exactly the bounded-degree families where the intervention is well-defined. In every case the exponential separation between average and expected causal effects under classical designs is a consequence of graph-unaware assignment, not a property of the estimand.

5.4 Cluster randomized designs

A cluster randomized design partitions the units into disjoint clusters $\mathcal{C} = \{C_1, \ldots, C_K\}$ and assigns a common treatment to all units in a cluster, drawing the cluster-level indicators $W_1, \ldots, W_K$ according to a specified scheme. A unit i reveals $Y_i(1, e_i^{\mathbf{1}})$ when its cluster and all clusters meeting its neighborhood are treated, and reveals $Y_i(0, 0)$ when they are all controlled. Provided the induced propensities $\pi_i(1, e_i^{\mathbf{1}})$ and $\pi_i(0, 0)$ are positive, the treatment–exposure Horvitz–Thompson estimator of Proposition 15 is unbiased for the total treatment effect; explicit expressions for these propensities under graph cluster randomization are given by Ugander et al. (2013).

The polynomial behavior of cluster randomization is, however, conditional on the clustering being aligned with the graph. Let m_i denote the number of distinct clusters met by the closed neighborhood $\{i\} \cup N_i$. Under independent cluster-level assignment with probability $\frac{1}{2}$, the all-treated condition requires all m_i of these clusters to be treated, so

$$\pi_i(1, e_i^{\mathbf{1}}) = \left(\tfrac{1}{2}\right)^{m_i},$$

and the inverse propensity is 2^{m_i}, exponential in the number of clusters the neighborhood spans rather than in the degree. When the clustering places each closed neighborhood inside a bounded number of clusters, $m_i = O(1)$ and the propensity is bounded below by a constant. Low-boundary or low-conductance partitions can have this property in clustered graphs, but the relevant design quantity is the neighborhood cluster count m_i itself. On graphs for which every practical balanced clustering splits many neighborhoods across a growing number of clusters, m_i grows with the degree, and the exponential cost reappears. Cluster randomization therefore trades an exponential dependence on degree for an exponential dependence on the neighborhood boundary, and removes the estimation barrier only when the chosen clustering is aligned with the interference graph. The independent set design IS-2 gives a graph-agnostic polynomial upper bound for the total treatment effect, while cluster designs can be sharper on graphs with strong community structure.

6 Implications for Designs, Estimators, and Models

The variance bounds of Section 4 have several implications for the design and analysis of experiments under interference. We collect them here.

6.1 Classical designs and estimators target expected causal effects

Bernoulli and completely randomized designs, analyzed with the difference-in-means estimator, are the tools most commonly used in practice. That these classical tools estimate an expected causal effect rather than an average causal effect under interference is well known (Hudgens and Halloran, 2008; Sävje et al., 2017): the difference-in-means estimator targets the expected direct effect, obtained by averaging the unit-level direct effect over the exposure induced by the design. Under a Bernoulli design, the difference-in-means estimator is design-unbiased for the expected direct effect; by Proposition 13, its variance is of order $(1 + c_{\text{avg}})/n$, and its computation does not require the knowledge of interference graph. However, the expected direct effect is policy dependent, and it mixes direct treatment effect

and interference effects, as shown in Proposition 3 which uses the non-parametric linear decomposition.

The same design and estimator are ill-suited to non-parametrically estimate average causal effects. The direct treatment effect is a contrast of the potential outcomes $Y_i(1, \mathbf{0}_{N_i})$ and $Y_i(0, \mathbf{0}_{N_i})$, and under a Bernoulli design the cell $(1, \mathbf{0}_{N_i})$ is observed with probability $p(1-p)^{d_i}$. By Proposition 11, the variance of any graph unaware design-unbiased estimator of the direct treatment effect is then exponential in the interference degree, a similar result holds for the total treatment effect. Moreover, the difference-in-means estimator is not consistent for the direct treatment effect: its probability limit is the expected direct effect, so it carries a bias equal to the expected treatment–exposure interaction $\overline{C}$, which is zero only under additivity. Classical estimators are therefore appropriate for expected causal effects such as the expected direct effect, but not for average causal effects such as the direct and total treatment effects.

6.2 Positivity and overlap for average causal effects

Changing the exposure mapping does not provide an escape. As shown in Section 4, the direct treatment effect requires the no-exposure potential outcome under any mapping, so its propensity remains exponentially small in the degree. The difficulty is a failure of positivity for the relevant treatment–exposure combination, the analogue under interference of the overlap condition in observational studies. On interference graphs with high-degree units this failure is severe, and a fixed-exposure average causal effect can be identified only on the subpopulation of units whose neighborhoods are small enough to provide adequate overlap.

6.3 Model-assisted estimation of average causal effects

A second strategy is available when one is willing to impose outcome model restrictions on the potential outcomes. The non-parametric decomposition $Y_i(z, e) = A_i(z) + B_i(e) + zC_i(e)$ suggests a natural class of working models: a restriction on the spillover and interaction components B_i and C_i relates the nuisance potential outcomes to the relevant ones, so that units that reveal nuisance potential outcomes also become informative for the estimand. The simplest such restriction is additivity, $C_i \equiv 0$, under which the difference-in-means estimator is consistent for the direct treatment effect, since the bias $\overline{C}$ then vanishes. More generally, a working model can be used within a model-assisted estimator, such as a generalized regression estimator, that retains a design-based target while borrowing information across treatment–exposure cells. Such estimators trade the exponential cost of the inverse propensity for sensitivity to misspecification of the working model. Model-assisted estimators of this kind under interference are developed in Karwa and Airoldi (2018).

7 Numerical illustration

We illustrate the estimation gap between ACEs and ECEs on simulated experiments. The interference graph, the exposure mapping, and the potential outcomes are fixed in advance, and the only source of randomness is the treatment assignment.

We draw the interference graph from two random graph model families. The first is the Erdős–Rényi graph $G(n, c/n)$, which is sparse and has degrees concentrated about their mean; it represents limited and roughly homogeneous interference. The second is the Barabási–Albert graph, whose degree distribution is heavy-tailed and contains a few high-degree hubs; it represents the heterogeneous interference common in social and economic networks. We use $n = 2000$ units in the Erdős–Rényi experiments and $n = 3000$ in the Barabási–Albert experiment.

The exposure mapping is the fraction of treated neighbors, $e_i = d_i^{-1} \sum_{j \in N_i} z_j$, with $e_i = 0$ when $d_i = 0$. Potential outcomes follow the decomposition $Y_i(z, e) = \mu_i + A_i(z) + B_i(e) + zC_i(e)$, with $A_i(z) = \alpha_i z$, $B_i(e) = \beta_i e$, and $C_i(e) = \gamma_i e$. The unit-level coefficients are drawn independently once, with $\mu_i \sim \text{Unif}(-0.2, 0.2)$, $\alpha_i \sim \text{Unif}(0.3, 0.7)$, $\beta_i \sim \text{Unif}(0.2, 0.6)$, and $\gamma_i \sim \text{Unif}(0, 0.4)$, and then held fixed. Thus the response surface is bounded and heterogeneous across units. The interaction coefficients have positive mean, so the direct and expected direct effects differ. The figures focus on three estimands: the direct treatment effect, which contrasts $Y_i(1, 0)$ and $Y_i(0, 0)$; the expected direct effect under the Bernoulli allocation policy; and the expected indirect effect under the same policy. Because the treated-neighbor count is binomial conditional on the graph under a Bernoulli design, each target is available in closed form and estimation error is measured against the exact finite-population value.

Each effect is estimated by the corresponding Horvitz–Thompson estimator of Section 4: treatment–exposure weights for the direct treatment effect, fixed own-treatment weights for the expected direct effect, and graph-aware weights for the expected indirect effect. The usual difference-in-means estimator appears separately in Figure 2; it is the ratio form of the own-treatment weighted estimator. We consider two designs: a graph-unaware Bernoulli design with treatment probability one half, used in all but one experiment, and an independent-set design aligned with the direct treatment effect. Each reported error is the root-mean-square error over independent draws of the assignment; Figures 1–3 use Monte Carlo repetitions, while Figure 4 reports the exact probability that the relevant potential outcome is observed at least once in 10^4 Bernoulli assignments. The Horvitz–Thompson estimators are design-unbiased by construction, and we confirmed that the empirical bias of each is zero up to Monte Carlo error.

Figure 1 reports the root-mean-square error against the mean interference degree on the Erdős–Rényi graphs; the vertical axis is logarithmic, so a straight line is exponential growth. The direct treatment effect (red circles) follows the inverse-propensity scale for observing $Y_i(1, 0)$ or $Y_i(0, 0)$ (the red dashed line, Proposition 10). The expected direct effect (blue squares) is horizontal, as Proposition 13 predicts, and stays near its dependence scale (the blue dashed line). The expected indirect effect (green triangles) rises only polynomially, well below the direct effect. The three estimands use the same graph, the same design, and the same outcomes. The separation is not only asymptotic. At bounded degree the RMSE of both classes decreases at the parametric rate $n^{-1/2}$, but the constant differs sharply: for the direct treatment effect the effective sample size is of order $np(1-p)^d$, since only assignments revealing the relevant potential outcomes contribute, whereas the expected direct effect uses every unit and its constant grows only polynomially in the degree. Computed at a single sample size, Figure 1 shows that this gap in constants is already large at the sizes typical of experiments.

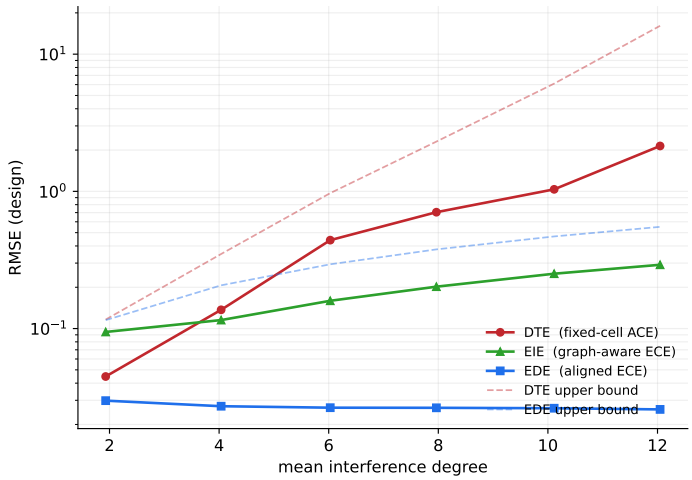

Figure 1: Design root-mean-square error of the Horvitz–Thompson estimator under a Bernoulli design ($p = 1/2$) against the mean interference degree, on Erdős–Rényi graphs ($n = 2000$); the vertical axis is logarithmic. The direct treatment effect (fixed-exposure ACE, red) follows the exponential inverse-propensity scale (red dashed; Proposition 10); the expected direct effect (aligned ECE, blue) is horizontal; the expected indirect effect (graph-aware ECE, green) grows polynomially.

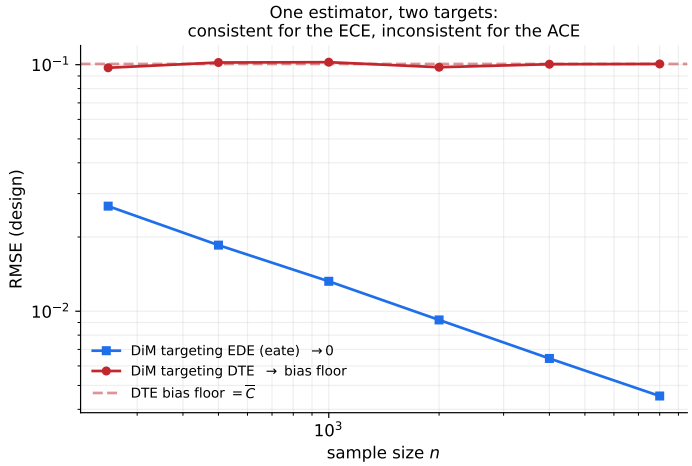

Figure 2: The difference-in-means estimator against two targets as the sample grows (Erdős–Rényi, mean degree 8, $p = 1/2$). Its error for the expected direct effect decays at the parametric rate; its error for the direct treatment effect settles at the policy-averaged interaction floor $\bar{C}_\pi$ set by the treatment–exposure interaction that own-treatment contrasts cannot remove.

Figure 2 follows a single estimator, the difference in means, against two targets as the sample grows. Its error for the expected direct effect (blue) decays as a straight line of

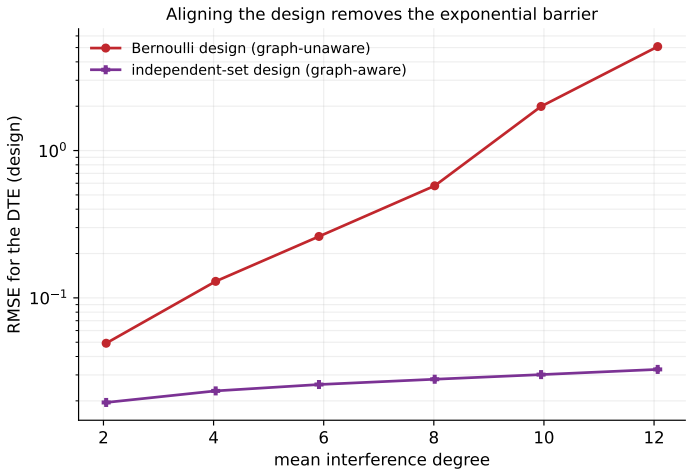

Figure 3: Estimand-aligned design for the direct treatment effect (Erdős–Rényi, $n = 2000$, $p = 1/2$). Under a Bernoulli design the Horvitz–Thompson error grows exponentially in the mean interference degree; under an independent-set design on the same graphs, which reveals $Y_i(1,0)$ and $Y_i(0,0)$ for ego units, it stays nearly flat.

slope $-1/2$ on the logarithmic axes, the parametric rate; its error for the direct treatment effect (red) flattens at a positive floor, the horizontal dashed line at $\bar{C}_\pi$. The estimand class, not the estimator, governs consistency: the same estimator is consistent for the effect the design averages over and inconsistent for the fixed-exposure effect on the same data. The decomposition $Y_i(z,e) = A_i(z) + B_i(e) + zC_i(e)$ identifies the floor exactly. It is the policy-averaged treatment–exposure interaction $\bar{C}_\pi$, equal to $p\,n^{-1}\sum_i \gamma_i$ in this simulation, and is the part of the response surface that own-treatment contrasts cannot remove; under additivity, $C_i \equiv 0$, the floor vanishes and the difference in means is consistent for the direct effect as well.

Figure 3 illustrates the first route around the exponential barrier for estimating ACEs. The independent-set design treats mutually non-adjacent ego units and holds their neighborhoods at control, so each ego reveals one of the two potential outcomes that define the direct treatment effect, $Y_i(1,0)$ or $Y_i(0,0)$. On the same Erdős–Rényi graphs as Figure 1, the error under the Bernoulli design (red) grows exponentially in the degree, while the error under the independent-set design (purple) remains nearly flat.

Figure 4 turns to a heavy-tailed Barabási–Albert graph. The direct-effect potential outcome for a high-degree unit is observed with probability that decays exponentially in the unit's degree, so the inverse propensity of the highest-degree unit exceeds 10^{70} (left panel) and the probability of seeing the relevant potential outcome at least once in 10^4 assignment draws is essentially zero (right panel). The fixed-exposure ACE is not merely expensive to estimate on such graphs; at feasible sample sizes it is effectively non-estimable for the hubs, while the expected direct effect uses every unit. Truncating or stabilizing the extreme inverse-propensity weights can reduce variance, but only by changing the estimator and introducing bias for the fixed finite-population target unless additional modeling assumptions are imposed. Heterogeneity in the degree distribution sharpens the separation rather than

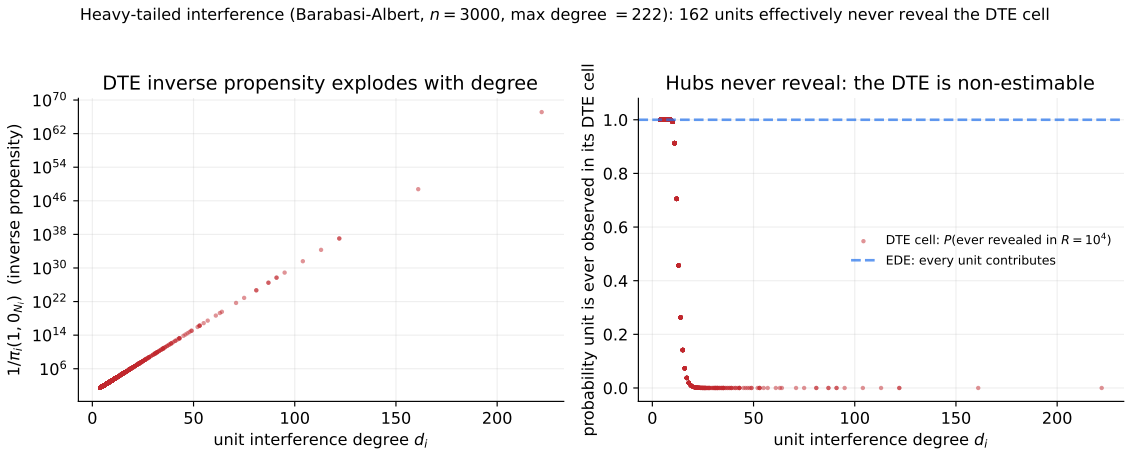

Figure 4: Heavy-tailed interference (Barabási–Albert, $n = 3000$). Left: the inverse propensity of the direct-effect exposure condition grows exponentially in a unit's degree, exceeding 10^{70} for the highest-degree unit. Right: the probability that a unit ever reveals the relevant direct-effect potential outcome, in 10^4 assignment draws, collapses to zero for high-degree units, whereas the expected direct effect uses every unit.

softening it: the cost of a fixed-exposure ACE is governed by the exponential moment of the degree distribution, which the hubs dominate.

8 Discussion

For principled causal inference, the potential-outcomes framework keeps three ingredients separate: restrictions on the potential outcomes, the definition of a causal estimand, and the design and estimator used to learn it. Without interference, the first ingredient is a single clean exclusion restriction: each unit has two potential outcomes. Interference is not the negation of that restriction. It is a family of possible restrictions. One has to specify which units' treatment assignments can affect a given unit's potential outcome, how those assignments enter the potential outcome, and whether the response surface has additional structure such as additivity, homogeneity, or linearity. The non-parametric decomposition of potential outcomes given in this paper provides a unified language to specify any finite-population model of potential outcomes under interference.

Moreover, without interference there is a single canonical finite-population average treatment effect. With interference there are multiple causal effects of interest, such as direct effects, spillover effects, total effects, and policy effects. We classify estimands under interference into two broad categories: average causal effects, which contrast fixed potential outcomes, and expected causal effects, which average potential outcomes under an allocation policy. This distinction matters because the same name can refer to different functionals of the potential-outcome table. For example, a direct effect may be a fixed contrast holding neighbors' treatment fixed, as in the direct treatment effect. It may also average over the neighbor treatments generated by an allocation policy, as in the expected direct effect. These estimands coincide under additivity, but not in general.

This perspective also clarifies why several results in the literature that appear to point in different directions are compatible. Partial-interference analyses specify groups and exposure within groups, and the resulting effects are tied to the group-level allocation policy (Hudgens and Halloran, 2008). Exposure-mapping approaches target average causal effects by contrasting fixed treatment–exposure combinations (Aronow and Samii, 2017). Some estimators remain valid under unknown interference because their targets average over the exposure distribution induced by a policy or design; in the terminology of this paper, the average direct effect studied by Sävje et al. (2017) is an expected direct effect. Other approaches recover average causal effects under unknown network structure by imposing outcome-model restrictions; for example, Yu et al. (2022) target a total treatment effect using additivity and linearity. Estimation of these effects is not free. The graph assumption is replaced either by a policy-averaged target or by an outcome-model assumption on the potential outcomes.

The same distinction governs design. Classical graph-unaware designs are often well matched to estimate policy-averaged expected effects, because they generate the exposure distribution being averaged over. They are poorly matched to estimate fixed-exposure average effects when the defining treatment–exposure combinations are rare. In observational studies without interference, positivity requires treatment variation within the covariate strata that define the target comparison. Under interference, the analogous strata are treatment–exposure cells. A Bernoulli design gives abundant information about typical exposure states, but a direct treatment effect ACE may require the atypical cell in which a treated unit has no treated neighbors. Similarly, a total treatment effect ACE may require all-treated and all-control neighborhoods. These cells can have exponentially small probability in the interference degree under classical designs. Fixed-exposure ACEs therefore require either estimand-aligned designs or outcome-model restrictions. Coarsening the spillover boundary can collapse this cost: under the any-treated exposure map the influencer spillover effect is estimated with constant focal propensities by sparse seeding on locally tree-like regular graphs (Section 5.3), while the uncompressed full-exposure versions retain the exponential barrier and need the distance-two design.

The exponential-versus-polynomial gap implies that the estimand and the design cannot be chosen separately. An investigator who targets a fixed-exposure average causal effect but implements a classical design will incur large estimation errors. The estimand and a design that reveals it must be chosen together, at the planning stage. When the design cannot be controlled, as in observational studies, the situation is more severe: average causal effects are then essentially out of reach without strong outcome models or bounded or logarithmic interference degrees, whereas expected causal effects aligned with the assignment mechanism remain estimable under the conditions that make them estimable by design.

A central problem is that the graph and exposure mapping are rarely known at the design stage. If the specified graph and exposure mapping are correct, a fixed-exposure average causal effect is a contrast on the reduced potential outcomes. If the exposure mapping is misspecified, the same contrast may average over assignment features that the mapping omits; this distinction between a target defined under a working exposure mapping and the corresponding target in the full Table of Science is emphasized by Sävje (2024). The same issue arises for the interference graph. If the graph omits relevant treatment coordinates, a fixed-exposure contrast in the reduced representation averages over those

omitted coordinates in the full Table of Science. It is then no longer the original fixed-exposure ACE, but a coarsened estimand defined by the misspecified representation. This is not only a matter of bias in an estimator. The target itself has changed.

Several problems follow. Experiments should be designed so that conclusions are stable to plausible changes in the interference graph and exposure mapping. When network data, pretreatment covariates, or repeated experiments are used to learn these objects, the uncertainty in that learning step should enter the final inference. Variance estimation must account for randomization dependence and interference dependence at the same time. Model-assisted estimators add another layer: the fitted response model can reduce variance or recover information from nuisance potential outcomes, but its use must be separated from the design-based correction it supports. These are the questions that arise after the estimand, the interference model, and the design have been put in the same language.

Appendix A. Additional examples and extensions for Section 2

A.1 Asymmetric and bipartite interference

While many models assume symmetric interference, interference may be directional. For instance, in social hierarchies or upstream–downstream systems, the treatment of one unit can affect another without a reciprocal effect. In these cases, the in-degree represents a unit's exposure to external interference, while the out-degree measures that unit's reach within the network.

The framework also extends to bipartite interference. In such cases, the graph $G = (U \cup V, E)$ consists of a set of treatment units U and outcome units V, where edges only exist from U to V. This structure is useful when the units being treated differ from the units whose outcomes are measured, as in environmental and service-provider settings (Zigler and Papadogeorgou, 2021).

A.2 Vector-valued and continuous exposures

The exposure function may be multidimensional. For example, in a social network one may use

$$e_i = \big(\text{number of treated friends},\ \text{number of treated friends-of-friends}\big).$$

In settings where interference is governed by physical proximity, the exposure mapping can be defined as a continuous function of geographic distance. For example, N_i can represent all units within a relevant radius, with

$$e_i = \sum_{j \in N_i} K(\rho_{ij}; \phi) z_j,$$

where ρ_{ij} is the distance between units i and j, and $K(\cdot)$ is a kernel function parameterized by ϕ (Giffin et al., 2023).

A.3 Nonlinear and threshold outcome restrictions

The decomposition in Theorem 1 also accommodates nonlinear restrictions on the spillover function B_i and the interaction function C_i. For example, under additivity ($C_i = 0$), one

may posit a concave interference function such as $B_i(e_i) = \gamma \log(1 + e_i)$, a convex function such as $B_i(e_i) = \gamma e_i^2$, or a sigmoidal function

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + \frac{L}{1 + \exp\{-k(e_i - e_0)\}}.$$

Such restrictions arise in settings where the first few treated neighbors have large effects, or where spillovers occur only after a threshold has been reached (Aral and Walker, 2011; Leskovec et al., 2007; De Giorgi et al., 2010).

Threshold and contagion models can be represented by binary exposure mappings such as

$$e_i = \mathbb{I}\left(\sum_{j \in N_i} z_j \geq K\right),$$

which are often used in models of collective behavior and complex contagion (Granovetter, 1978; Centola, 2010). Additional structural restrictions can then be imposed on $B_i(e_i)$ and $C_i(e_i)$ depending on whether individual treatment and exposure are substitutes, complements, or interact in some other way.

A.4 Non-anonymous exposure

In non-anonymous interference, the effect of neighbors' treatments depends on their identities or covariates. One example is

$$e_i = \sum_{j \in N_i} w(X_j) z_j,$$

where X_j are neighbor attributes. The potential outcome can then be written as

$$Y_i(z_i, e_i) = \alpha_i + \beta_i z_i + B_i\left(\sum_{j \in N_i} w(X_j) z_j\right) + z_i C_i\left(\sum_{j \in N_i} w(X_j) z_j\right).$$

This includes weighted and spatial exposure mappings, including environmental settings where exposure depends on distance or direction (Zigler et al., 2012; Patacchini et al., 2017; Adhikari and Zheleva, 2025).

Appendix B. Proofs for Section 2

B.1 Proof of Theorem 1

Restatement of Theorem 1. For any interference neighborhood N_i and exposure mapping $e_i = f(\mathbf{Z}_{N_i})$, there exist functions $A_i(z)$, $B_i(e)$, and $C_i(e)$ such that, for each unit i and each $(z, e) \in \{0, 1\} \times \mathcal{E}_i$,

$$Y_i(z, e) = A_i(z) + B_i(e) + z C_i(e),$$

with the normalization $B_i(0) = C_i(0) = 0$.

Proof Fix a unit i. Once the interference neighborhood and exposure mapping are specified, the potential outcomes for unit i are indexed by $(z, e) \in \{0, 1\} \times \mathcal{E}_i$. Define

$$A_i(z) = Y_i(z, 0), \qquad B_i(e) = Y_i(0, e) - Y_i(0, 0),$$

and

$$C_i(e) = Y_i(1, e) - Y_i(1, 0) - Y_i(0, e) + Y_i(0, 0).$$

Then $B_i(0) = 0$ and $C_i(0) = 0$. If $z = 0$, then

$$A_i(0) + B_i(e) = Y_i(0, 0) + Y_i(0, e) - Y_i(0, 0) = Y_i(0, e).$$

If $z = 1$, then

$$\begin{aligned} A_i(1) + B_i(e) + C_i(e) &= Y_i(1, 0) + Y_i(0, e) - Y_i(0, 0) \\ &\quad + Y_i(1, e) - Y_i(1, 0) - Y_i(0, e) + Y_i(0, 0) \\ &= Y_i(1, e). \end{aligned}$$

Thus $Y_i(z, e) = A_i(z) + B_i(e) + zC_i(e)$ for every z and e.

The representation is also a full-rank reparametrization. For each unit i, there are $2K_i$ distinct potential outcomes. The parametrization contains

$$2 + (K_i - 1) + (K_i - 1) = 2K_i$$

free parameters:

$$\{A_i(0), A_i(1)\}, \qquad \{B_i(e) : e \neq 0\}, \qquad \{C_i(e) : e \neq 0\}.$$

The formulas above give the inverse map from the potential-outcome table to these parameters, so the map is one-to-one and full rank. ■

Appendix C. Proofs and details for Section 3

C.1 Mixed causal effects: EIC, EIT, and ETE

Mixed causal effects contrast expected potential outcomes under a policy against fixed baseline responses. They can be viewed as expected analogues of the ACE indirect effects.

Expected exposure effects. Expected exposure effects capture the policy-averaged indirect effect, holding the own treatment fixed. The expected exposure effect on control units is

$$\tau_{\mathrm{EIC}}(\mathcal{Q}) \;=\; \frac{1}{n}\sum_{i=1}^{n}\sum_{e\in\mathcal{E}}\Big(Y_i(0, e) - Y_i(0, 0)\Big)\, P_{\mathcal{Q}}(E_i = e) = \frac{1}{n}\sum_i \mu_i(0; \mathcal{Q}) - \frac{1}{n}\sum_i Y_i(0, 0), \tag{C.1}$$

and on treated units is

$$\tau_{\mathrm{EIT}}(\mathcal{Q}) \;=\; \frac{1}{n}\sum_{i=1}^{n}\sum_{e\in\mathcal{E}}\Big(Y_i(1, e) - Y_i(1, 0)\Big)\, P_{\mathcal{Q}}(E_i = e) = \frac{1}{n}\sum_i \mu_i(1; \mathcal{Q}) - \frac{1}{n}\sum_i Y_i(1, 0). \tag{C.2}$$

These satisfy $\tau_{\mathrm{EIC}}(\mathcal{Q}) = \mathbb{E}_{\mathcal{Q}}[\tau_{\mathrm{IEC}}(E)]$ and $\tau_{\mathrm{EIT}}(\mathcal{Q}) = \mathbb{E}_{\mathcal{Q}}[\tau_{\mathrm{IET}}(E)]$.

Expected total effect. The expected total effect compares the expectation under full treatment to the baseline condition:

$$\tau_{\mathrm{ETE}}(\mathcal{Q}) = \frac{1}{n}\sum_{i=1}^{n}\sum_{e\in\mathcal{E}}\Big(Y_i(1,e) - Y_i(0,0)\Big)\,P_{\mathcal{Q}}(E_i = e) = \frac{1}{n}\sum_i \mu_i(1;\mathcal{Q}) - \frac{1}{n}\sum_i Y_i(0,0). \tag{C.3}$$

By Proposition 3, these satisfy $\tau_{\mathrm{EIC}}(\mathcal{Q}) = \bar{B}_{\mathcal{Q}}$, $\tau_{\mathrm{EIT}}(\mathcal{Q}) = \bar{B}_{\mathcal{Q}} + \bar{C}_{\mathcal{Q}}$, and $\tau_{\mathrm{ETE}}(\mathcal{Q}) = n^{-1}\sum_i \beta_i + \bar{B}_{\mathcal{Q}} + \bar{C}_{\mathcal{Q}}$, and they decompose as $\tau_{\mathrm{ETE}}(\mathcal{Q}) = \tau_{\mathrm{EDE}}(\mathcal{Q}) + \tau_{\mathrm{EIC}}(\mathcal{Q}) = \tau_{\mathrm{DTE}} + \tau_{\mathrm{EIT}}(\mathcal{Q})$.

Proposition 25 (ECE mixed-effect decomposition) *For any allocation policy $\mathcal{Q}$,*

$$\tau_{\mathrm{ETE}}(\mathcal{Q}) = \tau_{\mathrm{EDE}}(\mathcal{Q}) + \tau_{\mathrm{EIC}}(\mathcal{Q}) = \tau_{\mathrm{DTE}} + \tau_{\mathrm{EIT}}(\mathcal{Q}).$$

C.2 Proof of Proposition 3

Restatement. Under the non-parametric linear representation of Theorem 1,

$$Y_i(z,e) = \alpha_i + \beta_i z + B_i(e) + zC_i(e), \qquad B_i(0) = C_i(0) = 0,$$

the ACEs and ECEs have the representations stated in Proposition 3.

Proof For the fixed-exposure effects, substitute the representation into each contrast. The direct treatment effect is

$$Y_i(1,0) - Y_i(0,0) = \{\alpha_i + \beta_i + B_i(0) + C_i(0)\} - \{\alpha_i + B_i(0)\} = \beta_i.$$

Similarly,

$$Y_i(0,e) - Y_i(0,0) = B_i(e), \qquad Y_i(1,e) - Y_i(1,0) = B_i(e) + C_i(e),$$

and

$$Y_i(1,1) - Y_i(0,0) = \beta_i + B_i(1) + C_i(1).$$

Averaging these identities over i gives the ACE formulas.

For the expected effects, take expectations with respect to the exposure distribution induced by the allocation policy. Since

$$Y_i(1,E_i) - Y_i(0,E_i) = \beta_i + C_i(E_i),$$

we have

$$\tau_{\mathrm{EDE}}(\mathcal{Q}) = \frac{1}{n}\sum_i \beta_i + \frac{1}{n}\sum_i \mathbb{E}_{\mathcal{Q}}[C_i(E_i)].$$

The remaining three formulas follow from

$$Y_i(0,E_i) - Y_i(0,0) = B_i(E_i),$$

$$Y_i(1,E_i) - Y_i(1,0) = B_i(E_i) + C_i(E_i),$$

and

$$Y_i(1,E_i) - Y_i(0,0) = \beta_i + B_i(E_i) + C_i(E_i).$$

■

C.3 Proof of Proposition 5

Restatement. The total treatment effect satisfies

$$\tau_{\mathrm{TTE}} = \tau_{\mathrm{DTE}} + \tau_{\mathrm{IET}}.$$

Under additivity, $C_i(e) = 0$ for all i, e, so

$$\tau_{\mathrm{TTE}} = \tau_{\mathrm{DTE}} + \tau_{\mathrm{IET}} = \tau_{\mathrm{DTE}} + \tau_{\mathrm{IEC}}.$$

Proof By Proposition 3,

$$\tau_{\mathrm{TTE}} = \frac{1}{n}\sum_i \{\beta_i + B_i(1) + C_i(1)\},$$

while

$$\tau_{\mathrm{DTE}} + \tau_{\mathrm{IET}} = \frac{1}{n}\sum_i \beta_i + \frac{1}{n}\sum_i \{B_i(1) + C_i(1)\}.$$

The two expressions are equal. If $C_i(e) = 0$ for all i, e, then $\tau_{\mathrm{IET}} = \tau_{\mathrm{IEC}}$, which gives the additive form. ■

C.4 Proof of Proposition 25

Restatement. For any allocation policy $\mathcal{Q}$,

$$\tau_{\mathrm{ETE}}(\mathcal{Q}) = \tau_{\mathrm{EDE}}(\mathcal{Q}) + \tau_{\mathrm{EIC}}(\mathcal{Q}) = \tau_{\mathrm{DTE}} + \tau_{\mathrm{EIT}}(\mathcal{Q}).$$

Proof Using the notation of Proposition 3,

$$\tau_{\mathrm{ETE}}(\mathcal{Q}) = \frac{1}{n}\sum_i \beta_i + \bar{B}_{\mathcal{Q}} + \bar{C}_{\mathcal{Q}}.$$

Also,

$$\tau_{\mathrm{EDE}}(\mathcal{Q}) + \tau_{\mathrm{EIC}}(\mathcal{Q}) = \left\{\frac{1}{n}\sum_i \beta_i + \bar{C}_{\mathcal{Q}}\right\} + \bar{B}_{\mathcal{Q}},$$

and

$$\tau_{\mathrm{DTE}} + \tau_{\mathrm{EIT}}(\mathcal{Q}) = \frac{1}{n}\sum_i \beta_i + \bar{B}_{\mathcal{Q}} + \bar{C}_{\mathcal{Q}}.$$

Both expressions equal $\tau_{\mathrm{ETE}}(\mathcal{Q})$. ■

C.5 Proof of Corollary 4

Restatement. Under additivity of the causal response surface, $\tau_{\mathrm{EDE}}(\mathcal{Q}) = \tau_{\mathrm{DTE}}$ for every allocation policy $\mathcal{Q}$.

Proof Under additivity, $C_i(e) = 0$ for all units and exposure levels. Therefore $\bar{C}_{\mathcal{Q}} = n^{-1}\sum_i \mathbb{E}_{\mathcal{Q}}[C_i(E_i)] = 0$, and Proposition 3 gives

$$\tau_{\mathrm{EDE}}(\mathcal{Q}) = \frac{1}{n}\sum_i \beta_i = \tau_{\mathrm{DTE}}.$$

■

C.6 Proof of Proposition 2

Restatement. Under a Bernoulli assignment, the exposure mapping versions of EDE and EIE decompose the derivative of the expected population outcome:

$$\tau_{\mathrm{TPE}}(\pi) = \tau_{\mathrm{EDE}}(\pi) + \tau_{\mathrm{EIE}}(\pi).$$

Proof Let the treatment vector be $\mathbf{Z} = (Z_1, \ldots, Z_n)$, where $Z_i \sim \mathrm{Bernoulli}(\pi_i)$ are independent. Let $e_i = f(\mathbf{Z}_{N_i})$ be the exposure mapping for unit i. The expected population outcome under the policy $\pi = (\pi_1, \ldots, \pi_n)$ is

$$\mu(\pi) = \frac{1}{n}\sum_{i=1}^{n} \mathbb{E}_\pi[Y_i(Z_i, e_i)].$$

By the multivariate chain rule,

$$\tau_{\mathrm{TPE}}(\pi) = \sum_{j=1}^{n} \frac{\partial}{\partial \pi_j}\mu(\pi) = \frac{1}{n}\sum_{i=1}^{n}\left\{ \frac{\partial}{\partial \pi_i}\mathbb{E}_\pi[Y_i(Z_i, e_i)] + \sum_{j \in N_i} \frac{\partial}{\partial \pi_j}\mathbb{E}_\pi[Y_i(Z_i, e_i)] \right\},$$

where the partial derivative is zero for $j \notin N_i \cup \{i\}$.

For the own-treatment derivative, the independence of Z_i from e_i under Bernoulli assignment gives

$$\frac{\partial}{\partial \pi_i}\mathbb{E}_\pi[Y_i(Z_i, e_i)] = \mathbb{E}_\pi[Y_i(1, e_i) - Y_i(0, e_i)]\,.$$

Averaging over i gives $\tau_{\mathrm{EDE}}(\pi)$.

For $j \in N_i$, define

$$e_i^{+j} = e_i(Z_{N_i \setminus j}, z_j = 1), \qquad e_i^{-j} = e_i(Z_{N_i \setminus j}, z_j = 0).$$

Again using the product structure of the Bernoulli design,

$$\frac{\partial}{\partial \pi_j}\mathbb{E}_\pi[Y_i(Z_i, e_i)] = \mathbb{E}_\pi\left[Y_i(Z_i, e_i^{+j}) - Y_i(Z_i, e_i^{-j})\right].$$

Summing over $j \in N_i$ and averaging over i gives $\tau_{\mathrm{EIE}}(\pi)$. Combining the own-treatment and neighbor-treatment terms proves

$$\tau_{\mathrm{TPE}}(\pi) = \tau_{\mathrm{EDE}}(\pi) + \tau_{\mathrm{EIE}}(\pi).$$

■

C.7 Symmetric exposure mappings

Proposition 26 (EIE under symmetric exposure) *Suppose the assignment is Bernoulli with common probability π, and suppose the exposure mapping is symmetric. If*

$$e_i(\mathbf{Z}_{N_i}) = \sum_{j \in N_i} Z_j, \qquad d_i = |N_i|,$$

then the expected indirect effect can be written as

$$\tau_{\mathrm{EIE}}(\pi) = \frac{1}{n} \sum_{i=1}^{n} d_i \, \mathbb{E}_\pi \Big[Y_i(Z_i, \widetilde{E}_i + 1) - Y_i(Z_i, \widetilde{E}_i) \Big],$$

where $\widetilde{E}_i$ is the exposure generated by the remaining $d_i - 1$ neighbors. If instead

$$e_i(\mathbf{Z}_{N_i}) = \frac{1}{d_i} \sum_{j \in N_i} Z_j,$$

then

$$\tau_{\mathrm{EIE}}(\pi) = \frac{1}{n} \sum_{i=1}^{n} d_i \, \mathbb{E}_\pi \left[Y_i \left(Z_i, \widetilde{E}_i + \frac{1}{d_i} \right) - Y_i \left(Z_i, \widetilde{E}_i \right) \right],$$

where $\widetilde{E}_i$ is the fractional exposure generated by the remaining $d_i - 1$ neighbors.

Proof For the count exposure, fix i and $j \in N_i$. Since the Bernoulli design has a common treatment probability and the exposure mapping is symmetric, the distribution of the exposure generated by $N_i \setminus \{j\}$ is the same for every neighbor j. Moreover,

$$e_i^{+j} = \widetilde{E}_i + 1, \qquad e_i^{-j} = \widetilde{E}_i.$$

Substituting these two expressions into the definition of $\tau_{\mathrm{EIE}}(\pi)$ and summing the identical neighbor contributions over $j \in N_i$ gives

$$\tau_{\mathrm{EIE}}(\pi) = \frac{1}{n} \sum_{i=1}^{n} d_i \, \mathbb{E}_\pi \Big[Y_i(Z_i, \widetilde{E}_i + 1) - Y_i(Z_i, \widetilde{E}_i) \Big].$$

For fractional exposure, changing Z_j from zero to one increases exposure by $1/d_i$. Thus

$$e_i^{+j} = \widetilde{E}_i + \frac{1}{d_i}, \qquad e_i^{-j} = \widetilde{E}_i,$$

and the same averaging argument gives the second formula. ■

Appendix D. Proofs for Section 4

D.1 Proof of Lemma 8

Proof If there is no dependence-graph edge between two disjoint vertex sets A and B, then $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$ for every $i \in A$ and $j \in B$. Hence the assignment coordinates in $\cup_{i\in A}\mathcal{C}_i$ are disjoint from the assignment coordinates in $\cup_{j\in B}\mathcal{C}_j$. Since treatment assignments are mutually independent, the two assignment subvectors are independent. Because each X_i is a function of $\mathbf{Z}_{\mathcal{C}_i}$, the collections $\{X_i : i \in A\}$ and $\{X_j : j \in B\}$ are independent. Thus the graph defined by $c_{ij} = \mathbb{1}\{\mathcal{C}_i \cap \mathcal{C}_j \neq \emptyset\}$ is a valid dependence graph.

Now suppose the interference graph is undirected. Then $\mathcal{C}_i \cap \mathcal{C}_j \neq \emptyset$ if and only if i and j are at graph distance at most two. Hence $c_i = |B_2(i)| - 1$. Since every neighbor of i lies in $B_2(i)$, $c_i \geq d_i$, and averaging gives $c_{\text{avg}} \geq \bar{d}$. For the upper bound, the vertices within graph distance at most two from i, excluding i, consist of the d_i neighbors of i and the additional vertices reached through those neighbors. Thus

$$c_i \leq d_i + \sum_{\ell\in N_i}(d_\ell - 1) = \sum_{\ell\in N_i} d_\ell.$$

Averaging over i and using the symmetry of the undirected graph gives

$$c_{\text{avg}} \leq \frac{1}{n}\sum_{i=1}^{n}\sum_{\ell\in N_i} d_\ell = \frac{1}{n}\sum_{\ell=1}^{n} d_\ell^2.$$

The bound $c_{\text{avg}} \leq n-1$ is immediate because each unit has at most $n-1$ other units in its dependence neighborhood. ■

D.2 A generic design-dependence bound

The following auxiliary result separates dependence generated by overlapping assignment coordinates from dependence induced by the randomization design, using the design-dependence coefficient in Definition 9.

Proposition 27 (Interference and design dependence bound) *Consider a generic design-based estimator of the form*

$$\hat{\Psi} = \frac{1}{n}\sum_{i=1}^{n}\psi_i(\mathbf{Z}),$$

where $\psi_i(\mathbf{Z})$ is the weighted observed-outcome contribution of unit i. Suppose that ψ_i depends on the assignment only through $\mathbf{Z}_{\mathcal{C}_i}$, and suppose $|\psi_i(\mathbf{Z})| \leq B_i$. Define

$$\ell_{ij} = \mathbb{1}\{\mathcal{C}_i \cap \mathcal{C}_j \neq \emptyset\}, \qquad i \neq j.$$

Define

$$B_{2,n}^2 = \frac{1}{n}\sum_{i=1}^{n} B_i^2,$$

$$L_B = \frac{1}{n}\sum_{i=1}^{n}\sum_{j\neq i}\ell_{ij}B_iB_j,$$

and

$$R_B = \frac{1}{n}\sum_{i=1}^{n}\sum_{j\neq i}(1-\ell_{ij})\rho_{ij}^{P}B_iB_j.$$

Then, under an arbitrary design P,

$$\mathrm{Var}_P\left(\frac{1}{n}\sum_{i=1}^{n}\psi_i(\mathbf{Z})\right) \leq \frac{1}{n}\{B_{2,n}^2 + L_B + R_B\}.$$

Proof Write $\hat{\Psi} = n^{-1}\sum_i \psi_i(\mathbf{Z})$. The variance expansion gives

$$\mathrm{Var}_P(\hat{\Psi}) = \frac{1}{n^2}\left[\sum_i \mathrm{Var}_P\{\psi_i(\mathbf{Z})\} + \sum_i\sum_{j\neq i}\mathrm{Cov}_P\{\psi_i(\mathbf{Z}),\psi_j(\mathbf{Z})\}\right].$$

Since $|\psi_i| \leq B_i$, $\mathrm{Var}_P(\psi_i) \leq B_i^2$. If $\ell_{ij} = 1$, then

$$|\mathrm{Cov}_P\{\psi_i(\mathbf{Z}),\psi_j(\mathbf{Z})\}| \leq \mathbb{E}_P|\psi_i(\mathbf{Z})\psi_j(\mathbf{Z})| \leq B_iB_j.$$

If $\ell_{ij} = 0$, then $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$. When $B_iB_j > 0$, set $f(\mathbf{Z}_{\mathcal{C}_i}) = \psi_i(\mathbf{Z})/B_i$ and $g(\mathbf{Z}_{\mathcal{C}_j}) = \psi_j(\mathbf{Z})/B_j$. These functions have sup-norm at most one, and therefore

$$|\mathrm{Cov}_P\{\psi_i(\mathbf{Z}),\psi_j(\mathbf{Z})\}| \leq \rho_{ij}^{P}B_iB_j.$$

If $B_iB_j = 0$, the same bound is immediate. Substituting these three bounds into the variance expansion proves the result. ∎

D.3 Residual design dependence under complete randomization

The residual design-dependence term can be bounded explicitly for complete randomization. Suppose the design assigns exactly m of the n units to treatment. If $A, B \subseteq \{1,\ldots,n\}$ are disjoint assignment-coordinate sets with $|A| = a$ and $|B| = b$, then

$$\rho_P(A,B) \leq C\min\left\{1, \frac{ab}{n-a-b}\right\},$$

for a universal constant C, whenever $a+b<n$. Consequently, for local HT/IPW estimators whose ith summand depends on $\mathcal{C}_i = \{i\}\cup N_i$,

$$r_{\mathrm{avg}} \leq \frac{C}{n}\sum_{i=1}^{n}\sum_{j\neq i}(1-c_{ij})\min\left\{1, \frac{|\mathcal{C}_i|\,|\mathcal{C}_j|}{n-|\mathcal{C}_i|-|\mathcal{C}_j|}\right\}.$$

In particular, if $\max_i|\mathcal{C}_i| \leq n/4$, then

$$r_{\mathrm{avg}} \leq \frac{C}{n^2}\sum_{i=1}^{n}\sum_{j\neq i}(1-c_{ij})|\mathcal{C}_i|\,|\mathcal{C}_j| \leq C\left(\frac{1}{n}\sum_{i=1}^{n}|\mathcal{C}_i|\right)^2.$$

Thus on a d-regular local interference graph, $r_{\text{avg}} = O(d^2)$, matching the order of $c_{\text{avg}} \asymp d^2$. For bounded-degree graphs, $r_{\text{avg}} = O(1)$, so the residual design-dependence contribution $H\sqrt{r_{\text{avg}}/n}$ has the same $n^{-1/2}$ scaling as the Bernoulli dependence term, up to the relevant propensity scale H.

Proof It suffices to bound the dependence between two disjoint subvectors. For a fixed configuration x_A with $s = \sum_{\ell \in A} x_\ell$, the conditional distribution of $\mathbf{Z}_B$ under complete randomization is sampling without replacement from a population of size $n - a$ containing $m - s$ treated labels. The marginal distribution of $\mathbf{Z}_B$ is sampling without replacement from a population of size n containing m treated labels.

Couple the two draws sequentially over the b coordinates in B. After $k < b$ coordinates in B have been exposed and r treated labels have appeared among them, the two conditional probabilities that the next coordinate is treated are

$$\frac{m - s - r}{n - a - k} \qquad \text{and} \qquad \frac{m - r}{n - k}.$$

Their absolute difference is at most $Ca/(n - a - b)$, uniformly over admissible s, r, k. The union bound over the b sequential draws gives a total-variation bound of order $ab/(n-a-b)$, truncated at one. For any $\|g\|_\infty \le 1$, the difference between the conditional and marginal expectations of $g(\mathbf{Z}_B)$ is at most twice this total variation distance. Taking the supremum also over $\|f\|_\infty \le 1$ gives the displayed bound for $\rho_P(A, B)$.

Substituting $A = \mathcal{C}_i$, $B = \mathcal{C}_j$ for pairs with $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$ gives the first bound on r_{avg}. The simplified bound follows from $|\mathcal{C}_i| + |\mathcal{C}_j| \le n/2$, and the final regular and bounded-degree statements follow by setting $|\mathcal{C}_i| = d + 1$ or $\max_i |\mathcal{C}_i| = O(1)$. ■

D.4 Proof of Proposition 10

Restatement. Assume $\pi_i(d_i^+) > 0$ and $\pi_i(d_i^-) > 0$ for all i, and define

$$\hat{\tau}^{HT}_{\text{ACE}}(d^+, d^-) = \frac{1}{n} \sum_{i=1}^{n} Y_i^{obs} \left\{ \frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^+\}}{\pi_i(d_i^+)} - \frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^-\}}{\pi_i(d_i^-)} \right\}.$$

Then $\hat{\tau}^{HT}_{\text{ACE}}(d^+, d^-)$ is design-unbiased for $\tau_{\text{ACE}}(d^+, d^-)$. If $|Y_i(d)| \le M$ and the ith summand depends on the assignment only through $\mathbf{Z}_{\mathcal{C}_i}$, then the variance is bounded as stated.

Proof Because $Y_i^{obs} = Y_i\{D_i(\mathbf{Z})\}$, the expectation of the first weighted term is

$$\mathbb{E}\left[Y_i^{obs} \frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^+\}}{\pi_i(d_i^+)} \right] = \frac{1}{\pi_i(d_i^+)} \mathbb{E}\left[Y_i\{D_i(\mathbf{Z})\} \mathbb{1}\{D_i(\mathbf{Z}) = d_i^+\} \right] = Y_i(d_i^+).$$

The same calculation gives

$$\mathbb{E}\left[Y_i^{obs} \frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^-\}}{\pi_i(d_i^-)} \right] = Y_i(d_i^-).$$

Therefore

$$\mathbb{E}\{\hat{\tau}^{HT}_{\text{ACE}}(d^+, d^-)\} = \frac{1}{n} \sum_{i=1}^{n} \{Y_i(d_i^+) - Y_i(d_i^-)\} = \tau_{\text{ACE}}(d^+, d^-).$$

For the variance bound, write

$$X_i^{ACE} = Y_i\{D_i(\mathbf{Z})\}\left\{\frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^+\}}{\pi_i(d_i^+)} - \frac{\mathbb{1}\{D_i(\mathbf{Z}) = d_i^-\}}{\pi_i(d_i^-)}\right\}.$$

Then $\hat{\tau}^{HT}_{\mathrm{ACE}}(d^+, d^-) = n^{-1}\sum_i X_i^{ACE}$. Since $|Y_i(d)| \le M$,

$$|X_i^{ACE}| \le M\left\{\frac{1}{\pi_i(d_i^+)} + \frac{1}{\pi_i(d_i^-)}\right\} = MH_i.$$

Because the two target events in the ACE contrast are disjoint,

$$\mathbb{E}\{(X_i^{ACE})^2\} \le M^2\left\{\frac{1}{\pi_i(d_i^+)} + \frac{1}{\pi_i(d_i^-)}\right\} = M^2H_i.$$

Thus $\mathrm{Var}(X_i^{ACE}) \le M^2H_i$. If $c_{ij} = 1$, then

$$|\operatorname{Cov}(X_i^{ACE}, X_j^{ACE})| \le \mathbb{E}|X_i^{ACE}X_j^{ACE}| \le M^2H_iH_j.$$

If $c_{ij} = 0$, then the closed interference neighborhoods are disjoint. When $H_iH_j > 0$, the functions $X_i^{ACE}/(MH_i)$ and $X_j^{ACE}/(MH_j)$ have sup-norm at most one, so Definition 9 gives

$$|\operatorname{Cov}(X_i^{ACE}, X_j^{ACE})| \le M^2\rho_{ij}H_iH_j.$$

The same bound is immediate when $H_iH_j = 0$. Combining the diagonal and off-diagonal terms gives

$$\begin{aligned}\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{ACE}}(d^+, d^-)\} &= \frac{1}{n^2}\left\{\sum_{i=1}^n \mathrm{Var}(X_i^{ACE}) + \sum_{i=1}^n\sum_{j\neq i}\operatorname{Cov}(X_i^{ACE}, X_j^{ACE})\right\} \\ &\le \frac{M^2}{n^2}\left\{\sum_{i=1}^n H_i + \sum_{i=1}^n\sum_{j\neq i} c_{ij}H_iH_j + \sum_{i=1}^n\sum_{j\neq i}(1-c_{ij})\rho_{ij}H_iH_j\right\}.\end{aligned}$$

If $H_i \le H$ for all i, this becomes

$$\begin{aligned}\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{ACE}}(d^+, d^-)\} &\le \frac{M^2}{n^2}\left\{nH + H^2\sum_{i=1}^n\sum_{j\neq i} c_{ij} + H^2\sum_{i=1}^n\sum_{j\neq i}(1-c_{ij})\rho_{ij}\right\} \\ &= M^2\left\{\frac{H}{n} + H^2\frac{c_{\mathrm{avg}} + r_{\mathrm{avg}}}{n}\right\}.\end{aligned}$$

■

D.5 Proof of Proposition 11

Restatement. Fix the design P_n and the induced treatment–exposure combinations $D_i(\mathbf{Z})$. Suppose that $\hat{\tau}_n$ is design-unbiased for $\tau_{\text{ACE}}(d^+, d^-)$ for every bounded potential-outcome table. Then

$$\sup_{|Y|\le M} \text{Var}_{P_n}(\hat{\tau}_n) \ge \frac{M^2}{n^2} \max_i \left\{ \frac{1}{\pi_i(d_i^+)} - 1,\ \frac{1}{\pi_i(d_i^-)} - 1 \right\}.$$

Proof Fix a unit i and a target state d_i^+. Let

$$A_i = \{D_i(\mathbf{Z}) = d_i^+\}, \qquad \pi_i(d_i^+) = P_n(A_i).$$

Construct two potential-outcome tables Y^+ and Y^- that agree everywhere except at the coordinate $Y_i(d_i^+)$, where

$$Y_i^+(d_i^+) = M, \qquad Y_i^-(d_i^+) = -M.$$

All other potential outcomes are the same in the two tables. If this coordinate enters the estimand with coefficient c_i, then

$$\Delta = \tau(Y^+) - \tau(Y^-) = 2Mc_i.$$

Let

$$W(Z) = \hat{\tau}_n(Z, Y^{+,\text{obs}}) - \hat{\tau}_n(Z, Y^{-,\text{obs}}).$$

By design unbiasedness, $\mathbb{E}[W] = \Delta$. On the event A_i^c, the two tables produce identical observed data, and hence $W = 0$. Therefore

$$\Delta = \mathbb{E}\{W \mathbb{1}(A_i)\}.$$

By Cauchy–Schwarz,

$$\Delta^2 \le \mathbb{E}(W^2) P_n(A_i) = \mathbb{E}(W^2)\pi_i(d_i^+),$$

so $\mathbb{E}(W^2) \ge \Delta^2/\pi_i(d_i^+)$. Since $\mathbb{E}(W) = \Delta$,

$$\text{Var}(W) = \mathbb{E}(W^2) - \Delta^2 \ge \Delta^2 \left\{ \frac{1}{\pi_i(d_i^+)} - 1 \right\}.$$

On the other hand,

$$W - \Delta = \{\hat{\tau}_n(Y^+) - \tau(Y^+)\} - \{\hat{\tau}_n(Y^-) - \tau(Y^-)\}.$$

Thus

$$\text{Var}(W) \le 4 \max\left\{ \text{Var}_{P_n}(\hat{\tau}_n; Y^+), \text{Var}_{P_n}(\hat{\tau}_n; Y^-) \right\}.$$

Combining the last two displays gives

$$\sup_{|Y|\le M} \text{Var}_{P_n}(\hat{\tau}_n; Y) \ge \frac{\Delta^2}{4} \left\{ \frac{1}{\pi_i(d_i^+)} - 1 \right\} = c_i^2 M^2 \left\{ \frac{1}{\pi_i(d_i^+)} - 1 \right\}.$$

For an average ACE, the coefficient of each arm is $c_i = \pm 1/n$. Applying the same argument to either d_i^+ or d_i^-, and then taking the maximum over units and arms, gives the claimed bound. ■

D.6 Proof of Proposition 12

Restatement. Suppose $\lambda_i^+(d) = \lambda_i^-(d) = 0$ whenever $\pi_i^P(d) = 0$, and define $h_i^+(\mathbf{z})$ and $h_i^-(\mathbf{z})$ as in the proposition. Let

$$V_i = \sum_{d \in \mathcal{X}_i:\, \pi_i^P(d) > 0} \frac{\{\lambda_i^+(d) - \lambda_i^-(d)\}^2}{\pi_i^P(d)}.$$

Let

$$\hat{\tau}_{\mathrm{ECE}}^{HT} = \frac{1}{n} \sum_{i=1}^{n} \{h_i^+(\mathbf{Z}) - h_i^-(\mathbf{Z})\} Y_i^{obs}.$$

Then $\hat{\tau}_{\mathrm{ECE}}^{HT}$ is design-unbiased for

$$\tau_{\mathrm{ECE}} = \frac{1}{n} \sum_{i=1}^{n} \sum_{d \in \mathcal{X}_i} Y_i(d) \{\lambda_i^+(d) - \lambda_i^-(d)\}.$$

If $|Y_i(\mathbf{z})| \le M$ and

$$X_i = \{h_i^+(\mathbf{Z}) - h_i^-(\mathbf{Z})\} Y_i \{D_i(\mathbf{Z})\}$$

depends on the assignment only through $\mathbf{Z}_{\mathcal{C}_i}$, then

$$\mathrm{Var}(\hat{\tau}_{\mathrm{ECE}}^{HT}) \le \frac{M^2}{n^2} \left\{ \sum_{i=1}^{n} V_i + \sum_{i=1}^{n} \sum_{j \ne i} c_{ij} H_i H_j + \sum_{i=1}^{n} \sum_{j \ne i} (1 - c_{ij}) \rho_{ij} H_i H_j \right\}.$$

Proof Because $Y_i^{obs} = Y_i(\mathbf{Z})$, the expectation of the ith summand is

$$\mathbb{E}[\{h_i^+(\mathbf{Z}) - h_i^-(\mathbf{Z})\} Y_i^{obs}] = \sum_{\mathbf{z} \in \Omega} p(\mathbf{z}) \{h_i^+(\mathbf{z}) - h_i^-(\mathbf{z})\} Y_i \{D_i(\mathbf{z})\}.$$

Grouping allocations by the treatment–exposure combination they induce gives the plus-arm contribution

$$\begin{aligned} &\sum_{\mathbf{z} \in \Omega} p(\mathbf{z}) h_i^+(\mathbf{z}) Y_i \{D_i(\mathbf{z})\} \\ &\quad = \sum_{d \in \mathcal{X}_i:\, \pi_i^P(d) > 0} Y_i(d) \lambda_i^+(d) \frac{\sum_{\mathbf{z} \in \Omega:\, D_i(\mathbf{z}) = d} p(\mathbf{z})}{\pi_i^P(d)} \\ &\quad = \sum_{d \in \mathcal{X}_i:\, \pi_i^P(d) > 0} Y_i(d) \lambda_i^+(d) = \sum_{d \in \mathcal{X}_i} Y_i(d) \lambda_i^+(d), \end{aligned}$$

where the last equality uses the positivity condition. The same calculation for the minus arm gives

$$\mathbb{E}\{h_i^-(\mathbf{Z}) Y_i^{obs}\} = \sum_{d \in \mathcal{X}_i} Y_i(d) \lambda_i^-(d).$$

Summing over i and dividing by n proves design unbiasedness:

$$\mathbb{E}(\hat{\tau}_{\mathrm{ECE}}^{HT}) = \frac{1}{n} \sum_{i=1}^{n} \sum_{d \in \mathcal{X}_i} Y_i(d) \{\lambda_i^+(d) - \lambda_i^-(d)\} = \tau_{\mathrm{ECE}}.$$

For the variance bound, write

$$X_i = \{h_i^+(\mathbf{Z}) - h_i^-(\mathbf{Z})\}Y_i\{D_i(\mathbf{Z})\}.$$

Then

$$\hat{\tau}_{\text{ECE}}^{HT} = \frac{1}{n}\sum_{i=1}^{n} X_i.$$

By boundedness, $|X_i| \le MH_i$, so

$$\text{Var}(X_i) \le \mathbb{E}(X_i^2) \le M^2 \sum_{d\in\mathcal{X}_i:\,\pi_i^P(d)>0} \frac{\{\lambda_i^+(d) - \lambda_i^-(d)\}^2}{\pi_i^P(d)} = M^2 V_i.$$

If $c_{ij} = 1$, Cauchy–Schwarz gives

$$|\operatorname{Cov}(X_i, X_j)| \le \mathbb{E}|X_i X_j| \le M^2 H_i H_j.$$

If $c_{ij} = 0$, then the closed interference neighborhoods are disjoint. When $H_i H_j > 0$, the functions $X_i/(MH_i)$ and $X_j/(MH_j)$ have sup-norm at most one, so Definition 9 gives

$$|\operatorname{Cov}(X_i, X_j)| \le M^2 \rho_{ij} H_i H_j.$$

The same bound is immediate when $H_i H_j = 0$. Therefore

$$\begin{aligned}\text{Var}(\hat{\tau}_{\text{ECE}}^{HT}) &= \frac{1}{n^2}\left\{\sum_{i=1}^{n}\text{Var}(X_i) + \sum_{i=1}^{n}\sum_{j\neq i}\operatorname{Cov}(X_i, X_j)\right\} \\ &\le \frac{M^2}{n^2}\left\{\sum_{i=1}^{n} V_i + \sum_{i=1}^{n}\sum_{j\neq i} c_{ij}H_iH_j + \sum_{i=1}^{n}\sum_{j\neq i}(1-c_{ij})\rho_{ij}H_iH_j\right\}.\end{aligned}$$

If $V_i \le V$ and $H_i \le H$ for all i, this becomes

$$\begin{aligned}\text{Var}(\hat{\tau}_{\text{ECE}}^{HT}) &\le \frac{M^2}{n^2}\left\{nV + H^2\sum_{i=1}^{n}\sum_{j\neq i} c_{ij} + H^2\sum_{i=1}^{n}\sum_{j\neq i}(1-c_{ij})\rho_{ij}\right\} \\ &= M^2\left\{\frac{V}{n} + H^2\frac{c_{\text{avg}} + r_{\text{avg}}}{n}\right\}.\end{aligned}$$

■

D.7 Weighted dependence bound for graph-aware ECEs

Definition 28 (Weighted summand and dependence complexity) *Write* $h_i(\mathbf{Z}) = h_i^+(\mathbf{Z}) - h_i^-(\mathbf{Z})$ *and* $\psi_i(\mathbf{Z}) = h_i(\mathbf{Z})Y_i\{D_i(\mathbf{Z})\}$. *For the canonical ECE estimator* $\hat{\tau}_h = n^{-1}\sum_i \psi_i(\mathbf{Z})$, *define*

$$\omega_{ij}(Y) = \operatorname{Cov}\{\psi_i(\mathbf{Z}), \psi_j(\mathbf{Z})\}, \qquad i \neq j,$$

and

$$\Gamma_h(M) = \sup_{|Y|\le M} \left[\sum_{i=1}^{n} \mathrm{Var}\{\psi_i(\mathbf{Z})\} + \sum_{i=1}^{n} \sum_{j\neq i} \omega_{ij}(Y) \right].$$

The supremum is over all bounded potential-outcome tables satisfying $|Y_i(\mathbf{z})| \le M$.

Proposition 29 (Weighted dependence bound for graph-aware ECEs) *Assume independent Bernoulli assignment. Suppose the potential outcomes satisfy neighborhood interference with known neighborhoods* N_i*, and suppose* $h_i(\mathbf{Z})$ *depends only on* $\mathbf{Z}_{\mathcal{C}_i}$*. Then a valid dependence graph for* $\psi_i(\mathbf{Z}) = h_i(\mathbf{Z})Y_i(\mathbf{Z})$ *is given by*

$$c_{ij} = \mathbb{1}\{\mathcal{C}_i \cap \mathcal{C}_j \neq \emptyset\}, \qquad i \neq j.$$

Moreover,

$$\sup_{|Y|\le M} \mathrm{Var}(\hat{\tau}_h) = \frac{\Gamma_h(M)}{n^2}.$$

In particular, if $|h_i(\mathbf{Z})| \le H_i$*, then*

$$\mathrm{Var}(\hat{\tau}_h) \le \frac{M^2}{n^2} \left\{ \sum_{i=1}^{n} H_i^2 + \sum_{i=1}^{n} \sum_{j\neq i} c_{ij} H_i H_j \right\}.$$

This result separates two roles of the dependence graph. The unweighted graph gives a conservative upper bound by replacing every nonzero covariance by its maximal possible size. The sharper object is the weighted dependence complexity $\Gamma_h(M)$, which keeps the covariance weights induced by the design contrast and the potential-outcome table.

D.8 Proof of Proposition 29

Restatement. Assume independent Bernoulli assignment. Suppose the potential outcomes satisfy neighborhood interference with known neighborhoods N_i, and suppose $h_i(\mathbf{Z})$ depends only on $\mathbf{Z}_{\mathcal{C}_i}$, where $\mathcal{C}_i = \{i\} \cup N_i$. Then a valid dependence graph for $\psi_i(\mathbf{Z}) = h_i(\mathbf{Z})Y_i(\mathbf{Z})$ is given by

$$c_{ij} = \mathbb{1}\{\mathcal{C}_i \cap \mathcal{C}_j \neq \emptyset\}, \qquad i \neq j.$$

Moreover,

$$\sup_{|Y|\le M} \mathrm{Var}(\hat{\tau}_h) = \frac{\Gamma_h(M)}{n^2}.$$

If $|h_i(\mathbf{Z})| \le H_i$, then

$$\mathrm{Var}(\hat{\tau}_h) \le \frac{M^2}{n^2} \left\{ \sum_{i=1}^{n} H_i^2 + \sum_{i=1}^{n} \sum_{j\neq i} c_{ij} H_i H_j \right\}.$$

Proof By neighborhood interference, $Y_i(\mathbf{Z})$ is a function only of $\mathbf{Z}_{\mathcal{C}_i}$. By assumption, $h_i(\mathbf{Z})$ is also a function only of $\mathbf{Z}_{\mathcal{C}_i}$. Hence

$$\psi_i(\mathbf{Z}) = h_i(\mathbf{Z})Y_i(\mathbf{Z})$$

is a function only of $\mathbf{Z}_{\mathcal{C}_i}$. If $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$, then $\mathbf{Z}_{\mathcal{C}_i}$ and $\mathbf{Z}_{\mathcal{C}_j}$ are independent under independent Bernoulli assignment. Therefore $\psi_i(\mathbf{Z})$ and $\psi_j(\mathbf{Z})$ are independent. The graph with edge indicators $c_{ij} = \mathbb{1}\{\mathcal{C}_i \cap \mathcal{C}_j \neq \emptyset\}$ is therefore a valid dependence graph.

For a fixed potential-outcome table Y,

$$\hat{\tau}_h = \frac{1}{n}\sum_{i=1}^{n} \psi_i(\mathbf{Z}),$$

so

$$\mathrm{Var}(\hat{\tau}_h) = \frac{1}{n^2}\left[\sum_{i=1}^{n} \mathrm{Var}\{\psi_i(\mathbf{Z})\} + \sum_{i=1}^{n}\sum_{j\neq i} \mathrm{Cov}\{\psi_i(\mathbf{Z}), \psi_j(\mathbf{Z})\}\right].$$

Taking the supremum over all bounded potential-outcome tables gives

$$\sup_{|Y|\leq M} \mathrm{Var}(\hat{\tau}_h) = \frac{\Gamma_h(M)}{n^2}$$

by Definition 28.

Finally, if $|Y_i(\mathbf{Z})| \leq M$ and $|h_i(\mathbf{Z})| \leq H_i$, then $|\psi_i(\mathbf{Z})| \leq MH_i$. Thus

$$\mathrm{Var}\{\psi_i(\mathbf{Z})\} \leq M^2 H_i^2.$$

If $c_{ij} = 0$, then ψ_i and ψ_j are independent and the covariance is zero. If $c_{ij} = 1$, then Cauchy–Schwarz gives

$$|\,\mathrm{Cov}\{\psi_i(\mathbf{Z}), \psi_j(\mathbf{Z})\}| \leq \{\mathrm{Var}(\psi_i)\,\mathrm{Var}(\psi_j)\}^{1/2} \leq M^2 H_i H_j.$$

Substituting these bounds into the variance expansion gives the stated inequality. ■

D.9 Proof of Proposition 13

Restatement. Assume the true potential outcomes satisfy neighborhood interference with neighborhoods N_i. Assume Bernoulli assignment with treatment probabilities bounded away from zero and one, let $\mathcal{Q}_p$ denote this Bernoulli allocation policy, and assume $\sup_{\mathbf{z}} |Y_i(\mathbf{z})| \leq M$. Then $\hat{\tau}^{HT}_{\mathrm{EDE}}$ is design-unbiased for $\tau_{\mathrm{EDE}}(\mathcal{Q}_p)$, and

$$\mathrm{Var}_{P_n}(\hat{\tau}^{HT}_{\mathrm{EDE}}) \leq C_p M^2 \frac{1 + c_{\mathrm{avg}}}{n}.$$

In particular, $\hat{\tau}^{HT}_{\mathrm{EDE}}$ is design-MSE consistent if $c_{\mathrm{avg}} = o(n)$.

Proof Write the estimator as $n^{-1}\sum_i X_i$, where

$$X_i = Y_i^{obs}\left\{Z_i/p_i - (1 - Z_i)/(1 - p_i)\right\}.$$

First condition on Z_{-i}. Since the assignment is Bernoulli, Z_i is independent of Z_{-i}, and hence

$$\mathbb{E}(X_i \mid Z_{-i}) = Y_i(1, Z_{-i}) - Y_i(0, Z_{-i}).$$

Taking expectations over Z_{-i} gives

$$\mathbb{E}(\hat{\tau}^{HT}_{\mathrm{EDE}}) = \frac{1}{n}\sum_{i=1}^{n}\{\mathbb{E}[Y_i(1, Z_{-i})] - \mathbb{E}[Y_i(0, Z_{-i})]\} = \tau_{\mathrm{EDE}}(\mathcal{Q}_p).$$

The boundedness of the potential outcomes and the positivity of the treatment probabilities imply $|X_i| \leq C_p M$. If $\mathcal{C}_i \cap \mathcal{C}_j = \emptyset$, then X_i and X_j depend on disjoint sets of independent treatment assignments and are therefore independent. Thus only pairs with $c_{ij} = 1$ contribute to the covariance sum. This gives

$$\mathrm{Var}(\hat{\tau}^{HT}_{\mathrm{EDE}}) \leq \frac{C_p M^2}{n^2}\left(n + \sum_i \sum_{j \neq i} c_{ij}\right) = C_p M^2 \frac{1 + c_{\mathrm{avg}}}{n}.$$

■

D.10 Proof of Corollary 14

Restatement. Assume the true potential outcomes satisfy neighborhood interference with known neighborhoods N_i. Assume Bernoulli assignment with treatment probabilities bounded away from zero and one, and assume $\sup_{\mathbf{z}} |Y_i(\mathbf{z})| \leq M$. Then $\hat{\tau}^{G}_{\mathrm{EIE}}$ is design-unbiased for $\tau_{\mathrm{EIE}}(\pi)$, and

$$\mathrm{Var}(\hat{\tau}^{G}_{\mathrm{EIE}}) \leq \frac{C_p M^2}{n^2}\left\{\sum_i d_i^2 + \sum_i \sum_{j \neq i} c_{ij} d_i d_j\right\}.$$

Proof The graph-aware EIE estimator is

$$\hat{\tau}^{G}_{\mathrm{EIE}} = \frac{1}{n}\sum_{i=1}^{n} Y_i^{obs} \sum_{j \in N_i}\left\{\frac{Z_j}{p_j} - \frac{1 - Z_j}{1 - p_j}\right\}.$$

Its design unbiasedness follows by applying the ECE importance-weighting argument to each local neighbor contrast and summing over $j \in N_i$; equivalently, it follows from the policy-derivative calculation in Proposition 2. To obtain the variance bound, write

$$h_i^{\mathrm{EIE},G}(\mathbf{Z}) = \sum_{j \in N_i}\left\{\frac{Z_j}{p_j} - \frac{1 - Z_j}{1 - p_j}\right\}.$$

If the treatment probabilities are bounded away from zero and one, then there is a constant C_p such that

$$|h_i^{\mathrm{EIE},G}(\mathbf{Z})| \leq C_p d_i.$$

Moreover $h_i^{\mathrm{EIE},G}(\mathbf{Z})$ depends only on $\mathbf{Z}_{N_i}$, and hence on $\mathbf{Z}_{\mathcal{C}_i}$. Proposition 29, with $H_i = C_p d_i$, gives

$$\mathrm{Var}(\hat{\tau}^{G}_{\mathrm{EIE}}) \leq \frac{C_p M^2}{n^2}\left\{\sum_i d_i^2 + \sum_i \sum_{j \neq i} c_{ij} d_i d_j\right\},$$

Table 3: Generic comparison of the ACE inverse-propensity scale and the EDE dependence scale under Bernoulli assignment. The ACE column corresponds to a direct-effect target combination such as $(1, 0_{N_i})$. The EDE column gives the order of c_{avg}.

Graph sequence	ACE scale $\mathcal{A}_n$	EDE dependence c_{avg}
Bounded degree	$O(n^{-1})$	$O(1)$
Second-neighborhood regular, degree d	$e^{\gamma d}/n$	d^2
Sparse Erdős–Renyi, mean degree λ	approximately $e^{\gamma\lambda}/n$	approximately $\lambda + \lambda^2$
Dense block model, block size s	approximately $e^{\gamma s}/n$	approximately s
Star graph	dominated by hub term $e^{\gamma n}/n^2$	n
Complete graph	$e^{\gamma n}/n$	n

after absorbing constants depending only on the treatment-probability bounds into C_p. If $d_i \le d_{\max}$, then

$$\sum_i d_i^2 + \sum_i \sum_{j \neq i} c_{ij} d_i d_j \le d_{\max}^2 \left\{ n + \sum_i \sum_{j \neq i} c_{ij} \right\} = n d_{\max}^2 (1 + c_{\text{avg}}),$$

which gives the final display. ∎

D.11 Heterogeneous ACE/ECE comparison under classical designs

The DTE/EDE comparison of Section 4.3 extends to arbitrary heterogeneous degree sequences, where it is governed by two graph functionals. For a fixed-exposure ACE, the relevant scale is the inverse propensity-score moment

$$\mathcal{A}_n(d^+) = \frac{1}{n^2} \sum_{i=1}^{n} \pi_i(d_i^+)^{-1};$$

for the direct-effect entry $(1, 0_{N_i})$ under Bernoulli assignment, $\pi_i(1, 0_{N_i}) = p(1-p)^{d_i}$, so $\mathcal{A}_n(1, 0) \asymp n^{-2} \sum_i \exp(\gamma d_i)$ is an exponential moment of the degree sequence. For the EDE under the same design, the relevant scale is the average dependence degree, $\mathcal{B}_n = c_{\text{avg}}/n$.

The comparison is design-specific. It applies to classical designs such as Bernoulli assignment, and complete randomization with a fixed treatment fraction has the same exponential propensity-score behavior for local fixed-exposure combinations, up to finite-population corrections. It does not rule out designs that are constructed to increase the propensity scores of the ACE target combinations, such as the independent-set designs.

The table is meant to show why a single average degree condition is not adequate for heterogeneous graphs. By Jensen's inequality,

$$\frac{1}{n} \sum_{i=1}^{n} \exp(\gamma d_i) \ge \exp(\gamma \bar{d}),$$

so large average degree is enough to make the ACE inverse-propensity scale large. However, small average degree is not sufficient to make ACE estimation easy, because hubs can dominate the exponential moment. For EDE estimation, average degree is also not the right quantity: a star graph has $\bar{d} = O(1)$ but $c_{\mathrm{avg}} \asymp n$.

Corollary 30 (Heterogeneous ACE/ECE separation under Bernoulli assignment) *Consider a sequence of undirected interference graphs with degrees $d_i = |N_i|$, and suppose treatment is assigned independently with a common probability p bounded away from zero and one. For the direct-effect ACE entry $(1, 0_{N_i})$, there are constants $0 < \gamma_1 \le \gamma_2 < \infty$, depending only on p, such that*

$$\frac{1}{n^2}\sum_{i=1}^{n}\exp(\gamma_1 d_i) \lesssim \frac{1}{n^2}\sum_{i=1}^{n}\pi_i(1,0_{N_i})^{-1} \lesssim \frac{1}{n^2}\sum_{i=1}^{n}\exp(\gamma_2 d_i).$$

The canonical ACE Horvitz–Thompson variance is bounded by

$$\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{ACE}}\} \le \frac{C_p M^2}{n^2}\left\{\sum_{i=1}^{n}\exp(\gamma_2 d_i) + \sum_{i=1}^{n}\sum_{j\neq i} c_{ij}\exp\{\gamma_2(d_i+d_j)\}\right\},$$

where c_{ij} is the interference-overlap graph for the ACE summands. By contrast, the canonical EDE Horvitz–Thompson estimator satisfies

$$\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{EDE}}\} \le C_p M^2 \frac{1 + c_{\mathrm{avg}}}{n}.$$

Thus, under the same Bernoulli design and for arbitrary heterogeneous graph sequences, the ACE benchmark involves exponential moments of the degree sequence and their dependence-weighted analogues, whereas the EDE benchmark involves only the average degree of the dependence graph.

The regular-graph comparison of Section 4.3 is recovered as a special case. More generally, Table 3 should be read through the two graph functionals appearing in the corollary: exponential degree moments for ACE target combinations and average dependence degree for EDE importance-weighted estimators.

D.12 Proof of Corollary 30

Restatement. Under Bernoulli assignment with common treatment probability bounded away from zero and one, the direct-effect ACE benchmark involves exponential moments of the degree sequence and their dependence-weighted analogues, whereas the EDE benchmark involves only the average dependence degree.

Proof For the direct-effect entry $(1, 0_{N_i})$ under Bernoulli assignment,

$$\pi_i(1, 0_{N_i}) = p(1-p)^{d_i}.$$

Since p and $1-p$ are bounded away from zero and one, there are constants $0 < \gamma_1 \le \gamma_2 < \infty$, depending only on p, such that

$$\exp(\gamma_1 d_i) \le \pi_i(1, 0_{N_i})^{-1} \le \exp(\gamma_2 d_i)$$

after absorbing fixed multiplicative constants into the comparison constants. Summing over i gives the inverse-propensity display.

For the ACE HT upper bound, Proposition 10 gives

$$\mathrm{Var}\{\hat{\tau}^{HT}_{\mathrm{ACE}}\} \leq \frac{M^2}{n^2}\left\{\sum_i H_i + \sum_i \sum_{j\neq i} c_{ij} H_i H_j\right\}.$$

Here the residual design-dependence term is zero because treatment is assigned independently. For the direct-effect contrast with no treated neighbors, both arm propensities are bounded below by constants times $\exp(-\gamma_2 d_i)$, so $H_i \leq C_p \exp(\gamma_2 d_i)$. Substitution gives the stated dependence-weighted exponential bound. The EDE bound is exactly Proposition 13. ■

D.13 Auxiliary calculation: DTE and TTE table rates

Proposition 31 (DTE and TTE propensities under graph-unaware designs) *Let A_i denote the assignment coordinates that define the target treatment–exposure combination for unit i. For arbitrary interference, $A_i = \{1, \ldots, n\}$, while under neighborhood interference $A_i = \mathcal{C}_i = \{i\} \cup N_i$. Write $a_i = |A_i|$.*

(i) Suppose the design assigns treatments independently, with $\epsilon \leq P(Z_k = 1) \leq 1 - \epsilon$ for all k. Then every fixed assignment pattern on A_i has probability $\exp\{-\Theta(a_i)\}$, with constants depending only on ϵ. Consequently, the DTE no-exposure conditions and the TTE full-exposure/no-exposure conditions have inverse propensities $\exp\{\Theta(n)\}$ under arbitrary interference and $\exp\{\Theta(d_i)\}$ under neighborhood interference.

(ii) Suppose the design is exchangeable, so that conditional on $T_n = \sum_k Z_k = m$ all assignments with m treated units are equally likely. Assume that, for some $\epsilon > 0$, $T_n/n \in [\epsilon, 1-\epsilon]$ almost surely and that $a_i \leq \epsilon n/2$. Then every fixed assignment pattern on A_i has probability $\exp\{-\Theta(a_i)\}$, with constants depending only on ϵ. Thus the same $\exp\{\Theta(d_i)\}$ inverse-propensity order holds for neighborhood DTE and TTE conditions, including the no-exposure and full-exposure count/fraction exposure representations. Fixed-size complete randomization is a special case when the treatment fraction is bounded away from zero and one.

In both cases, the inverse-propensity lower bound in Proposition 11 applies to any design-unbiased non-parametric estimator of the corresponding DTE or TTE. The exchangeable case covers local treatment–exposure combinations. It does not cover arbitrary-interference TTE/DTE targets under fixed nontrivial treatment fractions, because such designs put zero probability on the global all-treated or all-control assignments.

Proof [Proof of Proposition 31] For the independent-design claim, fix a set A with $|A| = a$ and a pattern $z_A \in \{0,1\}^a$. If $\epsilon \leq P(Z_k = 1) \leq 1 - \epsilon$ for all k, then

$$\epsilon^a \leq P(\mathbf{Z}_A = z_A) \leq (1-\epsilon)^a.$$

Thus every fixed pattern on A has probability $\exp\{-\Theta(a)\}$. For arbitrary interference, the DTE and TTE target conditions impose fixed patterns on all n assignments. Under neighborhood interference, the DTE no-exposure conditions $(1, 0_{N_i})$ and $(0, 0_{N_i})$, and the TTE full-exposure/no-exposure conditions $(1, 1_{N_i})$ and $(0, 0_{N_i})$, impose fixed patterns on $\mathcal{C}_i = \{i\} \cup N_i$, whose size is $d_i + 1$. The inverse propensities therefore have the stated exponential orders.

For the exchangeable-design claim, condition on $T_n = m$. If a fixed pattern on A contains t treated units, then

$$P(\mathbf{Z}_A = z_A \mid T_n = m) = \frac{(m)_t (n-m)_{a-t}}{(n)_a},$$

where $(x)_k = x(x-1)\cdots(x-k+1)$. On the event $m/n \in [\epsilon, 1-\epsilon]$ and for $a \le \epsilon n/2$, each factor in this product is bounded below by $\epsilon/2$ and above by $1-\epsilon$, after adjusting constants depending only on ϵ. Hence

$$(\epsilon/2)^a \le P(\mathbf{Z}_A = z_A \mid T_n = m) \le (1-\epsilon)^a.$$

The same bounds hold unconditionally because the support of T_n is contained in $[\epsilon n, (1-\epsilon)n]$. Applying this with $A = \mathcal{C}_i$ gives the neighborhood DTE and TTE orders. Complete randomization with $m/n \in [\epsilon, 1-\epsilon]$ is the special case in which $T_n = m$ is fixed. The last statement follows immediately from Proposition 11. ∎

Restatement. The orders in Tables 1 and 2 follow from the propensities of the target treatment–exposure combinations under Bernoulli assignment. Complete randomization with a fixed treatment fraction has the same exponential order for the no-exposure and full-exposure fixed entries, up to finite-population corrections.

Proof Assume Bernoulli assignment with common treatment probability $p \in (0,1)$ bounded away from zero and one, and write $d_i = |N_i|$. For a fixed target contrast, the HT weight scale H_i is the larger reciprocal propensity among the two target arms, up to constants. The uniform table scale is $H = \max_i H_i$. The lower-bound column has the same propensity order by Proposition 11, and the HT variance-scale column follows by substituting H and the relevant dependence scale into $\mathcal{V}_n(H, c) = H/n + H^2 c/n$.

For the direct-effect rows, arbitrary interference with no treated units except i has target propensities

$$p(1-p)^{n-1} \qquad \text{and} \qquad (1-p)^n,$$

so $H = e^{\Theta(n)}$. Under neighborhood interference, the target conditions $(1, 0_{N_i})$ and $(0, 0_{N_i})$ have propensities

$$p(1-p)^{d_i} \qquad \text{and} \qquad (1-p)^{d_i+1},$$

so $H = e^{\Theta(d_{\max})}$. The count and fractional no-exposure rows impose $S_i = 0$, equivalently $F_i = 0$, and therefore have the same order.

For the binary any-treated direct-effect row, the DTE no-exposure condition is $U_i = 0$, equivalently $S_i = 0$, so

$$P\{Z_i = 1, U_i = 0\} = p(1-p)^{d_i}, \qquad P\{Z_i = 0, U_i = 0\} = (1-p)^{d_i+1}.$$

Thus $H = e^{\Theta(d_{\max})}$. For a typical count exposure $k_i \asymp pd_i$,

$$P\{S_i = k_i\} = \binom{d_i}{k_i} p^{k_i}(1-p)^{d_i-k_i} = \Theta(d_i^{-1/2})$$

uniformly for k_i in the central binomial range, by Stirling's approximation. Multiplying by the constant arm probabilities p or $1-p$ gives $H = \Theta(\sqrt{d_{\max}})$.

For the total-effect rows, arbitrary interference compares all treated with all control, giving propensities p^n and $(1-p)^n$, hence $H = e^{\Theta(n)}$. Under neighborhood interference, the target conditions $(1, 1_{N_i})$ and $(0, 0_{N_i})$ have propensities

$$p^{d_i+1} \qquad \text{and} \qquad (1-p)^{d_i+1},$$

so $H = e^{\Theta(d_{\max})}$. The count and fractional full-versus-none row imposes $S_i = d_i$ on the treated arm and $S_i = 0$ on the control arm, equivalently $F_i = 1$ versus $F_i = 0$, and has the same order. For the binary any-treated version, the treated arm has probability $p\{1-(1-p)^{d_i}\}$, but the control arm with no exposure has probability $(1-p)^{d_i+1}$. The no-exposure arm therefore still gives $H = e^{\Theta(d_{\max})}$.

Finally, under complete randomization with m treated units, the direct no-exposure target propensities are

$$\frac{\binom{n-d_i-1}{m-1}}{\binom{n}{m}}, \qquad \frac{\binom{n-d_i-1}{m}}{\binom{n}{m}},$$

whenever the binomial coefficients are well defined; uniformly over units this requires $1 \le m \le n - d_{\max} - 1$ for both direct-effect arms. If $m/n \to p \in (0,1)$, these are the finite-population analogues of $p(1-p)^{d_i}$ and $(1-p)^{d_i+1}$; for example, when $d_i = o(n)$ they are equal to the Bernoulli propensities times $\exp\{O(d_i^2/n)\}$. For the full-exposure treated arm in the TTE row, the corresponding propensity is

$$\frac{\binom{n-d_i-1}{m-d_i-1}}{\binom{n}{m}},$$

which is the finite-population analogue of p^{d_i+1}. Uniform positivity for the full-versus-none TTE rows requires $d_{\max} + 1 \le m \le n - d_{\max} - 1$. Thus complete randomization has the same exponential dependence on d_i for local fixed no-exposure and full-exposure entries, up to the usual finite-population correction, but the variance bound also includes the residual design-dependence term r_{avg}. The arbitrary-interference rows are not covered by fixed non-trivial complete randomization, since such a design does not assign positive probability to both all-treated and all-control global regimes. ■

D.14 Auxiliary calculation: exposure compression and propensity scores

Restatement. For an exposure mapping $e = f_i(u)$,

$$\pi_i^f(z, e) = \sum_{u \in f_i^{-1}(e)} \pi_i^{\text{id}}(z, u),$$

and hence

$$\pi_i^f(z,e) \geq \pi_i^{\mathrm{id}}(z,u) \qquad \text{for every } u \in f_i^{-1}(e).$$

Under balanced Bernoulli assignment,

$$\pi_i^f(z,e) = 2^{-(|N_i|+1)}\,|f_i^{-1}(e)|.$$

Proof The event $\{Z_i = z, f_i(Z_{N_i}) = e\}$ is the disjoint union of the events $\{Z_i = z, Z_{N_i} = u\}$ over $u \in f_i^{-1}(e)$. Summing their probabilities gives the first equality and the inequality. Under balanced Bernoulli assignment, each full local configuration (z,u) has probability $2^{-(|N_i|+1)}$. ■

D.15 Auxiliary calculation: outcome-model span condition

Restatement. Suppose an outcome model restricts $\theta_i = X_i\gamma_i$. Let O_i be the set of local states with positive propensity score under the design. A linear functional $\ell_i^\top\theta_i$ is determined by the potential outcomes in the observable states O_i for every γ_i if and only if

$$\ell_i^\top X_i \in \mathrm{rowspan}(X_{i,O_i}).$$

Proof The observable states determine the vector $X_{i,O_i}\gamma_i$. The functional $\ell_i^\top\theta_i = \ell_i^\top X_i\gamma_i$ is determined by these observable values for every γ_i if and only if it is a linear function of $X_{i,O_i}\gamma_i$. This is equivalent to the row vector $\ell_i^\top X_i$ lying in the row span of X_{i,O_i}. ■

Appendix E. Proofs for Section 5

E.1 Proof of Proposition 15

Proof By stability, $\mathbb{1}\{(Z_i,E_i) = d_i^+\}Y_i^{\mathrm{obs}} = \mathbb{1}\{(Z_i,E_i) = d_i^+\}Y_i(d_i^+)$, and since the potential outcomes are fixed under the design,

$$\mathbb{E}\left[\frac{\mathbb{1}\{(Z_i,E_i) = d_i^+\}Y_i^{\mathrm{obs}}}{\pi_i(d_i^+)}\right] = \frac{\pi_i(d_i^+)}{\pi_i(d_i^+)}\,Y_i(d_i^+) = Y_i(d_i^+).$$

The same calculation gives $Y_i(d_i^-)$ for the second arm. Taking the difference and averaging over i yields $\mathbb{E}[\hat{\tau}_{\mathrm{HT}}^{ZE}] = \tau$. ■

E.2 Proof of Lemma 16

Proof The event $\{i \in \mathcal{I}\}$ equals $\{U_i < U_k\ \forall k \in N_i\}$, that is, U_i is the strict minimum among the $d_i + 1$ independent uniform draws in $\{i\} \cup N_i$. By exchangeability each draw is equally likely to be the minimum, so $\Pr(i \in \mathcal{I}) = 1/(d_i+1)$. The coin B_i is drawn independently of the draws U, so

$$\rho_i^{(1)} = \Pr(i \in \mathcal{I})\Pr(B_i = 1) = \frac{q}{d_i+1}, \qquad \rho_i^{(0)} = \Pr(i \in \mathcal{I})\Pr(B_i = 0) = \frac{1-q}{d_i+1}.$$

On $F_i^{(1)}$, step (3) sets $Z_i = 1$ and $Z_k = 0$ for all $k \in N_i$, so $(Z_i, Z_{N_i}) = (1, \mathbf{0})$; on $F_i^{(0)}$ it sets $Z_i = 0$ and $Z_k = 0$ for all $k \in N_i$, giving $(Z_i, Z_{N_i}) = (0, \mathbf{0})$. ■

E.3 Proof of Proposition 18

Proof Write $\hat{\tau}_{\mathrm{IS}}^{D} = n^{-1} \sum_i W_i$ with $W_i = 2(d_i + 1)\mathbb{1}\{i \in \mathcal{I}\}\{Y_i(1, \mathbf{0}_{N_i})B_i - Y_i(0, \mathbf{0}_{N_i})(1 - B_i)\}$. Conditional on $\mathcal{I}$, the coins B_i are independent, and $\mathbb{E}[W_i \mid \mathcal{I}] = \mathbb{1}\{i \in \mathcal{I}\}\,\delta_i/p_i$; taking expectations gives $\mathbb{E}[W_i] = \delta_i$. For $i \neq j$, independence of B_i, B_j given $\mathcal{I}$ yields

$$\mathbb{E}[W_i W_j] = \frac{\Pr(i, j \in \mathcal{I})}{p_i p_j}\,\delta_i \delta_j, \qquad \mathrm{Cov}(W_i, W_j) = \delta_i \delta_j \left(\frac{\Pr(i, j \in \mathcal{I})}{p_i p_j} - 1 \right).$$

If the closed neighborhoods of i and j are disjoint, the events $\{i \in \mathcal{I}\}$ and $\{j \in \mathcal{I}\}$ depend on disjoint sets of draws, so $\Pr(i, j \in \mathcal{I}) = p_i p_j$ and the covariance vanishes; this occurs whenever $\mathrm{dist}_G(i, j) > 2$. Summing the diagonal variances and these covariances gives the stated exact identity.

For the pairwise ratio, if $i \sim j$ then i and j cannot both be closed-neighborhood minima, so $\Pr(i, j \in \mathcal{I}) = 0$. If $i \not\sim j$ share $s_{ij} = |N_i \cap N_j|$ neighbors, set $a = d_i + 1$, $b = d_j + 1$, $s = s_{ij}$. Conditioning on $U_i = x$ and $U_j = y$, both are minima iff the $a - 1 - s$ private neighbors of i exceed x, the $b - 1 - s$ private neighbors of j exceed y, and the s shared neighbors exceed $\max(x, y)$:

$$\Pr(i, j \in \mathcal{I}) = \int_0^1 \int_0^1 (1 - x)^{a-1-s} (1 - y)^{b-1-s} (1 - \max\{x, y\})^s \, dx \, dy = \frac{a + b}{ab(a + b - s)}.$$

Since $p_i p_j = 1/(ab)$, the ratio is $(a + b)/(a + b - s) = (d_i + d_j + 2)/(d_i + d_j + 2 - s_{ij})$, and the excess over 1 is $s_{ij}/(d_i + d_j + 2 - s_{ij})$.

For the bound, $\mathbb{E}[W_i^2] = 2(d_i+1)\{Y_i(1, \mathbf{0}_{N_i})^2 + Y_i(0, \mathbf{0}_{N_i})^2\} \leq 4M^2(d_i+1)$, so $\mathrm{Var}(W_i) \leq 4M^2(d_i + 1)$. Taking absolute values in the covariance sum and using $|\delta_i \delta_j| \leq 4M^2$ with the pairwise ratios (contribution 1 for each of the $2|E|$ ordered adjacent pairs, and $s_{ij}/(d_i + d_j + 2 - s_{ij})$ for each non-adjacent overlapping pair) gives the displayed inequality. ■

E.4 Proof of Corollary 19

Proof On a d-regular graph, $\sum_i (d_i + 1) = n(d+1)$ and $2|E| = nd$. Each adjacent pair contributes at most 1, and each non-adjacent overlapping pair contributes $s_{ij}/(2d + 2 - s_{ij}) \leq 1$ since $s_{ij} \leq d - 1$, while each ordered wedge $\sum_k d_k(d_k - 1) = nd(d - 1)$ counts at most one non-adjacent overlapping pair. The number of non-adjacent overlapping pairs is therefore $O(nd^2)$, and the weighted sum of the pairwise ratios is $O(nd^2)$; this yields the cruder rate $O(M^2 d^2/n)$ after dividing by n^2. If the graph is locally tree-like, $s_{ij} = O(1)$, the pairwise ratio satisfies $s_{ij}/(2d + 2 - s_{ij}) = O(1/d)$ (since the denominator is at least $d + 3$), and the weighted sum over the $O(nd^2)$ overlapping pairs is $O(nd)$; dividing by n^2 gives $O(M^2 d/n)$. For a general graph, drop the denominators using $d_i + d_j + 2 - s_{ij} \geq 1$ and

count ordered wedges: $\sum_{i\neq j:\ i\not\sim j} s_{ij} \le \sum_k d_k(d_k-1)$. Adding $2|E| = \sum_i d_i$ and the diagonal $\sum_i (d_i+1) = n + \sum_i d_i$ bounds the bracket by $n + \sum_i d_i + \sum_i d_i^2$, giving the stated $O(M^2(1+\bar{d}+\overline{d^2})/n)$. ■

E.5 Proof of Lemma 20

Proof The event $\{i \in \mathcal{I}_2\}$ is $\{U_i < U_k \ \forall k \in B_i^{(2)} \setminus \{i\}\}$, i.e. U_i is the strict minimum among the b_i independent draws in $B_i^{(2)}$; by exchangeability this has probability $1/b_i$. On $\{i \in \mathcal{I}_2, R_i = 1\}$, step (3) sets the entire closed neighborhood to 1, giving $(Z_i, Z_{N_i}) = (1, \mathbf{1})$, and R_i is independent of the draws, so $\rho_i^{(1)} = 1/(2b_i)$; the control side is symmetric. The lower bound $d_i + 1 \le b_i$ holds because the closed neighborhood is contained in the radius-two ball. The upper bound follows by counting vertices reached from i in zero, one, or two steps: at most $1 + d_i + \sum_{j\in N_i}(d_j - 1) \le 1 + d_i d_{\max}$. ■

E.6 Proof of Proposition 21

Proof Unbiasedness follows directly from the focal reveal events. On $\{i \in \mathcal{I}_2, R_i = 1\}$ the observed outcome equals $Y_i(1, e_i^{\mathbf{1}})$, and on $\{i \in \mathcal{I}_2, R_i = 0\}$ it equals $Y_i(0,0)$; each event has probability $1/(2b_i)$. Writing $\hat{\tau}_{\mathrm{IS}}^T = n^{-1}\sum_i V_i$, the argument of Proposition 18 applies with p_i replaced by $q_i = 1/b_i$ and δ_i by Δ_i: conditioning on $\mathcal{I}_2$ gives $\mathbb{E}[V_i \mid \mathcal{I}_2] = \mathbb{1}\{i \in \mathcal{I}_2\}\Delta_i/q_i$ and, for $i \neq j$, $\mathrm{Cov}(V_i, V_j) = \Delta_i\Delta_j(\Pr(i,j \in \mathcal{I}_2)/(q_iq_j) - 1)$. If $B_i^{(2)} \cap B_j^{(2)} = \emptyset$ the covariance vanishes. If $\mathrm{dist}_G(i,j) \le 2$, then i and j cannot both be strict minima in their radius-two balls, so $\Pr(i,j \in \mathcal{I}_2) = 0$. If $\mathrm{dist}_G(i,j) \in \{3,4\}$, the same integral as before with ball sizes (b_i, b_j) and overlap $s = s_{ij}^{(2)}$ gives $\Pr(i,j \in \mathcal{I}_2)/(q_iq_j) = (b_i + b_j)/(b_i + b_j - s)$. Finally $\mathbb{E}[V_i^2] = 2b_i\{Y_i(1, e_i^{\mathbf{1}})^2 + Y_i(0,0)^2\} \le 4M^2 b_i$, and taking absolute values in the covariance sum gives the bound: distance-one and distance-two pairs contribute at most one each, and distance-three and distance-four overlapping pairs contribute $s_{ij}^{(2)}/(b_i + b_j - s_{ij}^{(2)})$. ■

E.7 Proof of Corollary 22

Proof On a d-regular locally tree-like graph, $b_i = 1 + d + d(d-1) = d^2 + 1$, so the diagonal term is $4M^2 n^{-2}\sum_i b_i = O(M^2 d^2/n)$. The number of ordered pairs at distance at most two is $O(nd^2)$, so the distance-two exclusion term is also $O(M^2 d^2/n)$. For the positive radius-two overlap term, each unit has at most $d(d-1)^2$ distance-three partners, each with overlap $s = 2$ and denominator of order d^2, contributing $O(nd)$; each unit has $O(d^4)$ distance-four partners, each with overlap $s = 1$ and denominator of order d^2, contributing $O(nd^2)$, which is the same order as the diagonal and exclusion terms. Thus the total is $O(M^2 d^2/n)$. For a general graph, the diagonal and distance-two exclusion terms are bounded by $O(\sum_i b_i)$, and wedge-counting in the radius-two overlap graph bounds the positive overlap sum by $\sum_k b_k(b_k - 1)$, giving $O(M^2\overline{b^2}/n)$. ■

E.8 Proof of Proposition 24

Proof Unbiasedness. On $F_i^{(1)} = \{Z_i = 0, U_i = 1\}$ the observed outcome equals $Y_i(0,1)$, and on $F_i^{(0)} = \{Z_i = 0, U_i = 0\}$ it equals $Y_i(0,0)$. Since $\rho_i^{(b)} > 0$ for every unit of degree at least one,

$$\mathbb{E}\Big[\frac{\mathbb{1}\{F_i^{(1)}\}Y_i^{\text{obs}}}{\rho_i^{(1)}}\Big] = Y_i(0,1), \qquad \mathbb{E}\Big[\frac{\mathbb{1}\{F_i^{(0)}\}Y_i^{\text{obs}}}{\rho_i^{(0)}}\Big] = Y_i(0,0),$$

so $\mathbb{E}[\hat{\tau}_{\text{seed}}] = n^{-1}\sum_i\{Y_i(0,1) - Y_i(0,0)\} = \tau_{\text{IE}}(0;0,1)$.

Focal probabilities. Because $\mathcal{I}$ is an independent set, no ego has an ego neighbor. Consequently the event $\{U_i = 1\}$—at least one treated neighbor—forces $i \notin \mathcal{I}$ and hence $Z_i = 0$. Writing $K_i = |\mathcal{I} \cap N_i|$, the number of treated ego neighbors is a binomial thinning of K_i by the independent coins B_j, so $\Pr(U_i = 1) = \mathbb{E}[1 - (1-q)^{K_i}]$, which gives $\rho_i^{(1)}$. Further, $Z_i = 0$ holds when either $i \notin \mathcal{I}$ or $i \in \mathcal{I}$ with $B_i = 0$, so

$$\Pr(Z_i = 0) = 1 - q\,\Pr(i \in \mathcal{I}) = 1 - \frac{q}{d_i + 1},$$

and $\rho_i^{(0)} = \Pr(Z_i = 0) - \rho_i^{(1)} = \mathbb{E}[(1-q)^{K_i}] - \frac{q}{d_i+1}$, as claimed in (5.1).

Poisson limit on locally tree-like d-regular graphs. Suppose the induced subgraph on the radius-two ball of i is a tree, as on locally tree-like d-regular graphs. Given $U_i = u$, the ego events of the d neighbors of i are determined by disjoint sets of draws and are therefore conditionally independent, each with common probability

$$p(u) = \Pr\Big(U_j < u,\ U_j < U_\ell\ \forall\, \ell \in N_j \setminus \{i\}\Big) = \int_0^u (1-x)^{d-1}\,dx = \frac{1-(1-u)^d}{d}.$$

Hence $K_i \mid U_i = u \sim \text{Binomial}(d, p(u))$ and, for $s \in [0,1]$,

$$\mathbb{E}[s^{K_i}] = \int_0^1 \Big(1 - (1-s)\,p(u)\Big)^d du = \int_0^1 \Big(1 - \frac{1-s}{d}\big(1-(1-u)^d\big)\Big)^d du \xrightarrow{d\to\infty} e^{-(1-s)},$$

by dominated convergence, since $d\,p(u) = 1 - (1-u)^d \to 1$ for $u \in (0,1)$. Thus $K_i \Rightarrow \text{Poisson}(1)$. Taking $s = 1 - q$ gives $\mathbb{E}[(1-q)^{K_i}] \to e^{-q}$; hence $\rho_i^{(1)} \to 1 - e^{-q}$, and since $q/(d+1) \to 0$, $\rho_i^{(0)} \to e^{-q}$.

Star center. For the center c of a star, the leaves are degree-one units: given $U_c = u$, each leaf is an ego exactly when its single draw falls below u, so the leaf ego events are independent Bernoulli(u) and

$$\mathbb{E}[(1-q)^{K_c}] = \int_0^1 (1-qu)^d\,du = \frac{1-(1-q)^{d+1}}{q(d+1)} = O\Big(\frac{1}{qd}\Big),$$

so $\rho_c^{(0)}$ has the same polynomial order and $\rho_c^{(1)} \to 1$. In general $K_i \le d_i$, so $\mathbb{E}[(1-q)^{K_i}] \ge (1-q)^{d_i}$ and $\rho_i^{(0)} \ge (1-q)^{d_i} - q/(d_i+1)$: when the ego-neighborhood count K_i concentrates on values of order d_i the unexposed arm is no larger than the classical $(1-q)^{d_i}$ scale. The constant-order guarantee is therefore tied to units whose ego-neighborhood is $O(1)$ in expectation.

Variance. Write $\hat{\tau}_{\text{seed}} = n^{-1}\sum_i W_i$. Since $|Y_i| \le M$,

$$\mathbb{E}[W_i^2] = \mathbb{E}\Big[\frac{\mathbb{1}\{F_i^{(1)}\}(Y_i^{\text{obs}})^2}{(\rho_i^{(1)})^2} + \frac{\mathbb{1}\{F_i^{(0)}\}(Y_i^{\text{obs}})^2}{(\rho_i^{(0)})^2}\Big] \le M^2\Big(\frac{1}{\rho_i^{(1)}} + \frac{1}{\rho_i^{(0)}}\Big),$$

so the diagonal term is at most $M^2 n^{-2}\sum_i(1/\rho_i^{(1)} + 1/\rho_i^{(0)}) = O(M^2/n)$ on locally tree-like d-regular graphs. The focal event $F_i^{(b)}$ is a function of the draws U in the radius-two ball of i and of the coins of the egos in N_i, so for $\text{dist}_G(i,j) \ge 5$ the focal events of i and j are independent and $\text{Cov}(W_i, W_j) = 0$. The remaining covariance contributions come from the $O(nd^4)$ pairs at distance at most four and obey the same radius-two overlap anatomy as in the proof of Proposition 18; combined with the diagonal term this is exactly the structure of the generic ACE variance bound of Proposition 10. ■

References

Shishir Adhikari and Elena Zheleva. Inferring individual direct causal effects under heterogeneous peer influence. *Machine Learning*, 114(4):113, 2025.

Sinan Aral and Dylan Walker. Creating social contagion through viral product design: A randomized trial of peer influence in networks. *Management Science*, 57(9):1623–1639, 2011.

Sinan Aral and Dylan Walker. Identifying influential and susceptible members of social networks. *Science*, 337(6092):337–341, 2012.

Sinan Aral, Lev Muchnik, and Arun Sundararajan. Engineering social contagions: Optimal network seeding in the presence of homophily. *Network Science*, 1(2):125–153, 2013.

Peter M Aronow. A general method for detecting interference between units in randomized experiments. *Sociological Methods & Research*, 41(1):3–16, 2012.

Peter M. Aronow and Cyrus Samii. Estimating average causal effects under general interference. *The Annals of Applied Statistics*, 11(4):1912–1947, 2017.

Susan Athey, Dean Eckles, and Guido W Imbens. Exact p-values for network interference. *Journal of the American Statistical Association*, 113(521):230–240, 2018.

Sarah Baird, J Aislinn Bohren, Craig McIntosh, and Berk Özler. Optimal design of experiments in the presence of interference. *Review of Economics and Statistics*, 100(5): 844–860, 2018.

Guillaume W. Basse and Avi Feller. Design of experiments in the presence of interference. *Annual Review of Statistics and Its Application*, 5:45–69, 2018. doi: 10.1146/annurev-statistics-031017-100449.

Guillaume W Basse, Avi Feller, and Panos Toulis. Randomization tests of causal effects under interference. *Biometrika*, 106(2):487–494, 2019.

Rohit Bhattacharya, Daniel Malinsky, and Ilya Shpitser. Causal inference under interference and network uncertainty. In *Uncertainty in Artificial Intelligence*, pages 1028–1038. PMLR, 2020.

Robert M Bond, Christopher J Fariss, Jason J Jones, Adam DI Kramer, Cameron Marlow, Jaime E Settle, and James H Fowler. A 61-million-person experiment in social influence and political mobilization. *Nature*, 489(7415):295–298, 2012.

Yann Bramoullé, Habiba Djebbari, and Bernard Fortin. Identification of peer effects through social networks. *Journal of Econometrics*, 150(1):41–55, 2009.

Chencheng Cai, Xu Zhang, and Edoardo M. Airoldi. Independent-set design of experiments for estimating treatment and spillover effects under network interference. In *International Conference on Learning Representations (ICLR)*, 2024.

Damon Centola. The spread of behavior in an online social network experiment. *Science*, 329(5996):1194–1197, 2010.

Giacomo De Giorgi, Michele Pellizzari, and Silvia Redaelli. Identification of social interactions through identifying restrictions. *American Economic Journal: Applied Economics*, 2(2):275–291, 2010. doi: 10.1257/app.2.2.275.

Andrew Giffin, BJ Reich, Shu Yang, and AG Rappold. Generalized propensity score approach to causal inference with spatial interference. *Biometrics*, 79(3):2220–2231, 2023.

Bryan S Graham. Identifying social interactions through conditional variance restrictions. *Econometrica*, 76(3):643–660, 2008.

Mark Granovetter. Threshold models of collective behavior. *American Journal of Sociology*, 83(6):1420–1443, 1978.

Yuchen Hu, Shuangning Li, and Stefan Wager. Average direct and indirect causal effects under interference. *Biometrika*, 109(4):1165–1172, 2022.

Michael G Hudgens and M Elizabeth Halloran. Toward causal inference with interference. *Journal of the American Statistical Association*, (482):832–842, 2008.

Vishesh Karwa and Edoardo M Airoldi. A systematic investigation of classical causal inference strategies under mis-specification due to network interference. *arXiv preprint arXiv:1810.08259*, 2018.

Jure Leskovec, Lada A Adamic, and Bernardo A Huberman. The dynamics of viral marketing. *ACM Transactions on the Web (TWEB)*, 1(1):5–es, 2007. doi: 10.1145/1232722.1232727.

Shuangning Li and Stefan Wager. Random graph asymptotics for treatment effect estimation under network interference. *The Annals of Statistics*, 50(4):2334–2358, 2022.

Charles F Manski. Identification of endogenous social effects: The reflection problem. *The Review of Economic Studies*, 60(3):531–542, 1993.

Charles F Manski. Identification of treatment response with social interactions. *The Econometrics Journal*, 16(1):S1–S23, 2013.

Yuki Ohnishi, Bikram Karmakar, and Arman Sabbaghi. Degree of interference: A general framework for causal inference under interference. *Journal of Machine Learning Research*, 26(120):1–37, 2025.

Elizabeth Levy Paluck, Hana Shepherd, and Peter M. Aronow. Changing climates of conflict: A social network experiment in 56 schools. *Proceedings of the National Academy of Sciences*, 113(3):566–571, 2016.

Eleonora Patacchini, Edoardo Rainone, and Yves Zenou. Heterogeneous peer effects in education. *Journal of Economic Behavior & Organization*, 134:190–227, 2017.

Jean Pouget-Abadie, Guillaume Saint-Jacques, Martin Saveski, Weitao Duan, Souvik Ghosh, Ya Xu, and Edoardo M Airoldi. Testing for arbitrary interference on experimentation platforms. *Biometrika*, 106(4):929–940, 2019.

David Puelz, Guillaume Basse, Avi Feller, and Panos Toulis. A graph-theoretic approach to randomization tests of causal effects under general interference. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 84(1):174–204, 2022.

Joseph Rigdon and Michael G Hudgens. Exact confidence intervals in the presence of interference. *Statistics & probability letters*, 105:130–135, 2015.

Paul R Rosenbaum. Interference between units in randomized experiments. *Journal of the American Statistical Association*, 102(477), 2007.

Fredrik Sävje. Causal inference with misspecified exposure mappings: separating definitions and assumptions. *Biometrika*, 111(1):1–15, 2024.

Fredrik Sävje, Peter M Aronow, and Michael G Hudgens. Average treatment effects in the presence of unknown interference. *arXiv preprint arXiv:1711.06399*, 2017.

Michael E Sobel. What do randomized studies of housing mobility demonstrate? causal inference in the face of interference. *Journal of the American Statistical Association*, 101 (476):1398–1407, 2006.

Johan Ugander, Brian Karrer, Lars Backstrom, and Jon Kleinberg. Graph cluster randomization: Network exposure to multiple universes. In *Proceedings of the 19th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 329–337. ACM, 2013.

Davide Viviano. Experimental design under network interference. *Journal of Political Economy: Microeconomics*, 2026. First version 2020; arXiv:2003.08421.

Bar Weinstein and Daniel Nevo. Causal inference with misspecified network interference structure. *Biometrics*, 82(1):ujag023, 2026.

Christina Lee Yu, Edoardo M Airoldi, Christian Borgs, and Jennifer T Chayes. Estimating the total treatment effect in randomized experiments with unknown network structure. *Proceedings of the National Academy of Sciences*, 119(44):e2208975119, 2022.

Corwin M Zigler and Georgia Papadogeorgou. Bipartite causal inference with interference. *Statistical Science*, 36(1):109–123, 2021.

Corwin M Zigler et al. Estimating causal effects of air quality regulations. *Biostatistics*, 13 (2):289–302, 2012.