Correct First, Fast Second: What a Local 7B Code Model Can and Cannot Do for GPU Kernel Optimization

Jui-Chien Tsou Po-Yi Wu
Department of Computer Science and Information Engineering
National Taiwan University
{b11902044, b11902127}@csie.ntu.edu.tw

arXiv:2611.12079v1 [cs.LG] 9 Nov 2026

***Abstract*—We study how much the *interaction strategy* around a fixed, locally hosted 7B-parameter code model (Qwen2.5-Coder-7B-Instruct) affects its ability to replace PyTorch operators with custom CUDA kernels. On 30 KernelBench tasks (levels 1–3, first 10 each, evaluated on a V100), we compare four regimes of increasing structure: zero-shot prompting, profiling-guided prompting, iterative refinement on execution feedback, and a multi-agent loop (analysis, retrieval, generation, evaluation, feedback triage) in which the same model fills four roles around a non-LLM evaluator. Three findings stand out. (1) Correctness, not speed, is the binding constraint: even *correct* generated kernels average 0.97–0.99× the speed of eager PyTorch, because a 7B model cannot out-write cuBLAS/cuDNN on dense GEMM and convolution; the only real wins are fused memory-bound epilogues. (2) The best strategy depends on difficulty: profiling-guided prompting is strongest overall (43% correct, 27% faster than baseline), but the agent loop is the *only* method that solves any level-2 fusion problem (3/10 vs. 0/10 for all others). (3) Naïve agentic scaffolding is actively harmful—our first agent scored 3% correct, *below* zero-shot's 30%—and the decisive repairs were matching the agent's ambition to the model's competence and spending compute on best-of-n sampling under a real execution oracle (3%→33% correct from n=1 to n=8), not on more reasoning steps.**

***Index Terms*—GPU kernel optimization, large language models, CUDA, agentic systems, KernelBench, test-time compute**

I. Introduction

Writing fast GPU kernels is expert work: a correct, performant CUDA kernel must respect memory coalescing, shared-memory capacity, occupancy, and numerical semantics simultaneously. Large language models (LLMs) promise to automate parts of this work, and benchmarks such as KernelBench [1] formalize the task: given a PyTorch reference module `Model`, produce a functionally identical `ModelNew` that replaces operators with custom CUDA and runs faster.

Most published results on this task use frontier-scale proprietary models. This paper asks a question closer to what a practitioner with fixed local hardware faces: *holding a single mid-size open-weight model constant, how much does the prompting and agent scaffolding around it matter—and where exactly does it help or hurt?* The model is the controlled variable; the four treatment conditions are zero-shot prompting, profiling-guided prompting, iterative refinement, and a multi-agent loop in which the *same* model plays four roles (code analyzer, retrieval researcher, kernel generator, feedback triager) around a real compile-and-run evaluator.

Our contributions: (i) a controlled 1×4 comparison on 30 KernelBench tasks with compile / correctness / speedup metrics and a per-difficulty breakdown; (ii) evidence that no single strategy dominates—guided prompting wins overall and on whole networks, while only the agent solves fused multi-operator problems; (iii) a root-cause analysis of *why* the agent initially scored below zero-shot, and a principled repair (competence-matched ambition plus best-of-n sampling under an execution oracle) validated by a coherent component-ablation study and a test-time-compute scaling curve; (iv) an infrastructure lesson—LLM-generated kernels can deadlock the GPU, so unattended evaluation requires hard per-candidate timeouts.

II. Background and Related Work

Benchmark and dataset. KernelBench [1] contains PyTorch reference modules in difficulty levels: level 1 holds single operators (GEMM variants, convolutions, elementwise ops), level 2 holds short *fused* operator sequences (e.g. `ConvTranspose2d→BiasAdd→Clamp→Scale`), and level 3 holds full architectures (MLPs, LeNet-5, AlexNet, ...). Each task ships `get_inputs()` generators; a candidate passes if its outputs match the reference within tolerance across randomized trials. We use the **first 10 problems of each of levels 1–3** (30 tasks): the subset keeps the full study — four methods plus eight agentic configurations (the component ablations and a best-of-n sweep) — tractable on shared GPUs, while still spanning three qualitatively different difficulty regimes. The headline metric is $fast_1$: the fraction of tasks both correct *and* faster than eager PyTorch [1].

Baseline. The performance baseline is eager PyTorch, whose matmul/convolution dispatch to highly tuned cuBLAS/cuDNN kernels—a strong baseline that, as we show, dominates what a 7B model can hand-write for compute-bound operators.

Model and agent tooling. All generation uses Qwen2.5-Coder-7B-Instruct [2] served locally via HuggingFace Transformers (fp16, multi-GPU workers; one GPU generates while a second evaluates). No proprietary model or API is used anywhere. The agent framework is ∼1,500 lines of Python orchestrating the roles described in §III; retrieval uses BM25

over a small curated corpus of CUDA optimization notes (tiling, coalescing, epilogue fusion, `torch.utils.cpp_extension` usage).

Related work. Our design draws on retrieval-augmented generation [3], self-refinement from execution feedback [4]–[6], and repeated sampling with a verifier, whose effectiveness on code is well documented [7], [8]. Our results echo [8] in a harsher setting: with a weak generator and a real compile-and-run verifier, sampling breadth beats reasoning depth.

III. Methodology

All four methods share the same base task prompt (Appendix A), model, decoding budget (4096 new tokens), and evaluator; they differ only in the information and control flow around generation.

A. The four interaction regimes

Zero-shot. The raw KernelBench prompt; greedy decoding; one attempt. This is the capability floor.

Guided. We profile each *reference* module with the PyTorch profiler on the target GPU and inject its top operators by CUDA time into the prompt, followed by a roofline-style rule we added after the pilot study (§VI): *do not reimplement matmul/conv—cuBLAS/cuDNN already won; fuse the cheap memory-bound epilogue instead*. One attempt, no loop.

Iterative. Up to three turns: generate, compile/check/time on the GPU, feed a distilled error/timing report back, regenerate. The submitted kernel is the best turn under a $(\textit{correct}, \textit{compiled}, \textit{speedup})$ ranking.

Agentic. A five-stage loop — four roles played by the same model under different system prompts, plus a non-LLM Evaluator that closes the loop (Fig. 1; Appendix B):

1) *Code Analyzer* – names the operations, bottlenecks, and fusion opportunities; emits retrieval queries.
2) *RAG Researcher* – retrieves from the documentation corpus and writes a cited optimization brief.
3) *Kernel Generator* – writes `ModelNew`; samples n candidates per turn (best-of-n).
4) *Evaluator* (non-LLM) – compiles, checks correctness, times each candidate; keeps the best per turn.
5) *Feedback Analyzer* – converts raw evaluator output into a diagnosis and a directive for the next turn.

The loop runs up to three turns. Every component can be disabled independently, which yields the ablation study in §V-C.

B. How the agent was developed

A natural question for any agent system is how its instructions were arrived at. Our final design is the outcome of one major iteration: an initial agent that failed badly, and the redesign that repaired it (full chronology in Appendix D). We report both, because the failure is itself informative.

Pilot (v1): 3% correct, 33% compile—below zero-shot. The v1 generator system prompt *mandated* custom kernels ("replace PyTorch operators ... using `load_inline` or Triton"), and the analyzer/researcher briefs recommended tiling, shared memory, and warp reductions on every task. The 7B model attempted ambitious hand-tiled CUDA it could not compile—e.g. calling `load_inline(sources=...)` (a keyword that does not exist) and mismatching `PYBIND11_MODULE` names, yielding `PyInit` import errors. The feedback loop could not recover because each retry stayed in the same over-ambitious mode; tellingly, a single-turn variant *beat* the full loop.

Repairs (v2). Four changes, each motivated by an observed failure: (1) *Competence-matched prompting*: every agent role is instructed to reach a correct, compiling kernel first—reuse `torch` ops for the heavy compute, fuse only the memory-bound epilogue—and to escalate to hand-written tiling only once a correct baseline exists; the prompts also pin the exact `load_inline` signature the model kept hallucinating. (2) *Best-of-n sampling*: $n{=}4$ candidates per turn at temperature 0.7, ranked by the real evaluator. (3) *Monotone turn selection*: the submitted kernel is the best turn by $(\textit{correct}, \textit{compiled}, \textit{speedup})$, so a later broken retry can never displace an earlier working kernel (this bug had been silently degrading the iterative baseline too). (4) *Hard evaluation timeouts*: each candidate is evaluated in a child process killed after 240 s, because generated kernels can deadlock the GPU—one such kernel stalled a production run for 14 hours (§VI-D).

IV. Experimental Setup

Hardware/software. Generation and evaluation run on nodes with 8× NVIDIA V100-SXM2-32GB (Volta, no fp32 tensor cores), CUDA 12.3, PyTorch 2.x, GCC 10. Workers pair GPUs: one serves the model, one evaluates kernels, so generation never competes with timing.

Measurement. For each candidate: compile via `torch.utils.cpp_extension`; correctness over 5 trials with randomized inputs (must match the reference within KernelBench's default tolerance); timing via CUDA events over 10 trials after warm-up, in fp32. The speed baseline is the eager-PyTorch reference time measured on the same GPU.

Success criteria. We report: *compile rate*; *correctness rate*; $\textit{fast}_1$ = correct **and** faster than eager PyTorch (the success definition for this study); and *geometric-mean speedup* over correct samples. A task counts once per (method, task) cell; multi-turn/multi-sample methods submit exactly one final kernel chosen by the evaluator, never by the LLM's own judgment.

Reproducibility. One command reruns the full matrix; the report, tables, and figures regenerate from raw evaluation JSON via a single script. Completed cells are skipped, so the study is incrementally re-runnable. Code and artifacts are released (see Code and Data Availability).

V. Results

A. Headline comparison

Table I: profiling-guided prompting is the best single strategy, lifting correctness 30%→43% and *doubling* fast_1 (13%→27%) over zero-shot at identical compute. Iterative refinement does not help this model: blind "retry with the

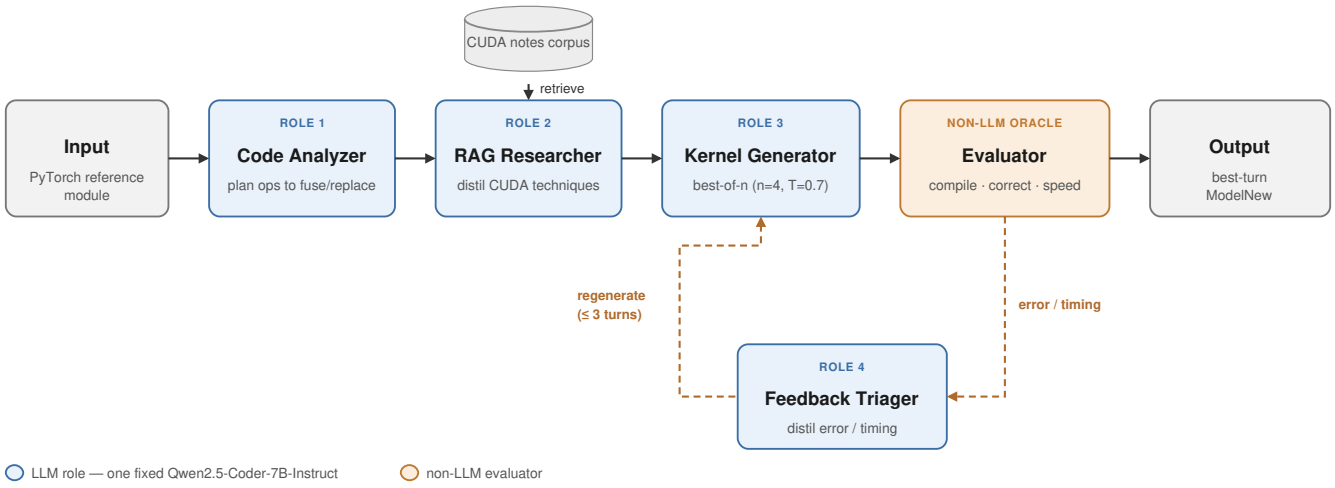

Fig. 1. The agentic loop. Four roles are the same Qwen2.5-Coder-7B-Instruct under different system prompts (blue); a non-LLM Evaluator (orange) compiles, checks, and times each candidate and closes the feedback loop. Generation is best-of-n (n=4); the loop runs at most three turns.

TABLE I
FOUR METHODS ON 30 KERNELBENCH TASKS (LEVELS 1–3). FAST$_1$ = CORRECT AND FASTER THAN EAGER PYTORCH. GEO-MEAN SPEEDUP IS OVER CORRECT SAMPLES ONLY.

Method	Compile%	Correct%	fast$_1$%	Geo-mean
Zero-shot	70	30	13	0.98×
Guided	**77**	**43**	**27**	0.99×
Iterative	67	30	10	0.98×
Agentic	73	30	10	0.97×

TABLE II
CORRECT SOLUTIONS OUT OF 10, PER DIFFICULTY LEVEL.

Method	L1 (single op)	L2 (fused seq.)	L3 (network)
Zero-shot	9	0	0
Guided	**10**	0	**3**
Iterative	9	0	0
Agentic	5	**3**	1

error log" rarely converts failures. The repaired agent reaches parity in aggregate—but with a very different distribution of successes.

B. Difficulty decides the winner

Table II is the central qualitative result. On L1, simpler is better (zero-shot/guided: 9–10/10) and the agent's machinery *hurts* (5/10): its analyzer talks the model into unnecessary custom kernels where a one-line torch op passes. On L2, the agent is the **only** method that scores at all (3/10 vs. 0/10 everywhere else): explicit fusion analysis plus an execution-grounded loop is exactly what multi-operator fusion rewards. On L3, guided profiling uniquely succeeds (3/10) by pointing the model at the dominant layer. No method dominates across difficulty.

TABLE III
AGENTIC COMPONENT ABLATION. EVERY VARIANT FLIPS ONE COMPONENT OFF AGAINST THE SAME BEST-OF-4 BACKBONE; Δ IS CORRECTNESS CHANGE VS. THE FULL SYSTEM.

Variant	Correct%	Compile%	fast$_1$%	Δ
Agentic (full, n=4)	30	73	10	—
− RAG researcher	20	77	7	−10
− Code analyzer	20	63	10	−10
− Feedback analyzer	**37**	73	10	**+7**
− Refinement loop	20	70	10	−10
− Best-of-n (n=1)	3	23	0	**−27**

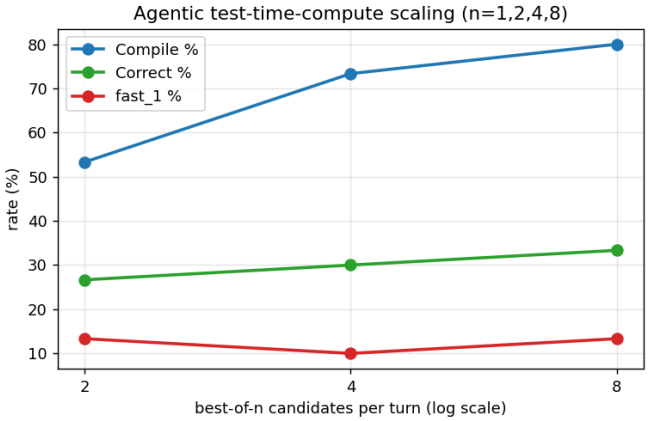

Fig. 2. Test-time-compute scaling: full agent, varying only candidates per turn ($n \in \{1, 2, 4, 8\}$; evaluator keeps the best). Correctness: 3/27/30/33%; compile: 23/53/73/80%. The first doubling is transformative; returns then diminish.

C. Component ablation and test-time compute

Table III and Fig. 2 decompose the agent. **Best-of-n is the load-bearing component by a wide margin**: removing it collapses the system to 3% correct / 23% compile—*worse*

TABLE IV
FAILURE CATEGORY COUNTS OVER 30 TASKS (CORRECT EXCLUDED; "SLOW" = CORRECT BUT NOT FASTER THAN BASELINE).

Method	compile	halluc. API	shape	wrong out	slow
Zero-shot	7	4	5	4	5
Guided	5	6	3	3	5
Iterative	9	3	4	4	6
Agentic	8	8	**0**	5	6

than the unrepaired pilot. The analyzer, retrieval, and the loop each contribute $\approx +10$pp, but only on top of the sampling backbone. The **feedback analyzer is mildly harmful** ($+7$pp when removed): its rewritten diagnoses sometimes over-specify the next attempt (e.g. directing the model to keep refactoring an already-correct design, Appendix C.3), whereas the raw compiler error leaves the model freer to make the small fix. The scaling curve is steep then saturating: $n{=}1{\to}2$ buys $+24$pp correctness, $n{=}4{\to}8$ only $+3$pp for double the evaluation cost.

D. Error taxonomy

Table IV: the agent eliminates *shape mismatch* entirely (its analyzer's explicit shape/fusion reasoning works) but doubles *hallucinated API* (8 cases)—steered toward writing more custom code, the model more often invents symbols, e.g. including the non-public header `ATen/native/cudnn/utils.h` or calling `F.batch_norm` from C++. The residual bottleneck is API grounding, not planning.

VI. DISCUSSION AND ANALYSIS

A. Root cause 1: the cuBLAS/cuDNN ceiling (why "slow" happens)

Across *all* methods, geometric-mean speedup of correct kernels is 0.97–0.99×: correct kernels are on average slightly *slower* than eager PyTorch. This is structural, not incidental. The L1 GEMM tasks are large (4096^2 matrices), firmly compute-bound, and PyTorch dispatches them to cuBLAS kernels with architecture-tuned tiling and double-buffering; a naïve `load_inline` kernel from a 7B model achieves a fraction of that throughput, and even its "correct" attempts only tie by ultimately delegating to `torch.matmul`. The genuine wins are exactly where theory predicts custom code can win: *memory-bound* epilogues and reductions. The agent's best speedup (1.23×) came on the skinny-K GEMM (`7_Matmul_with_small_K`), where cuBLAS is launch/utilization-bound; guided's wins came from fusing bias/activation/clamp chains after convolutions, cutting kernel launches and global-memory round-trips. Consequently the realistic objective for a small model is **correct fused kernels at $\sim$1×**, and methods separate on correctness, which is what our metrics show.

B. Root cause 2: why every single-shot method scores 0/10 on level 2

L2 tasks chain 3–5 operators. A passing `ModelNew` must either write one fused kernel (hard: the pilot showed the model cannot reliably compile tiled multi-op CUDA) or carefully replicate the reference semantics while fusing the easy parts (subtle: e.g. `2_ConvTranspose2d_BiasAdd_...` applies a *learned* bias whose shape broadcasts differently than the conv's own bias). Zero-shot/guided/iterative fail through one of two doors: ambitious fused CUDA that does not compile, or torch-level rewrites that silently change semantics (*wrong output*, Table IV). The agent's analyzer names the exact operator chain and its dataflow before generation, and the loop gets a second chance after a semantic miss—together worth 3/10 where everything else gets 0. Its own L2 successes are honest about the model's competence ceiling: the representative solution (Appendix C.1) keeps `ConvTranspose2d`/`MaxPool` as torch ops and fuses `clamp+mean` into one pass—a correct fusion at 1.01×, not a heroic kernel.

C. Root cause 3: why the pilot agent lost to zero-shot

The v1 failure was not randomness; it was a *systematic mismatch between the scaffold's ambition and the model's competence*. The scaffold pushed every task toward custom tiled CUDA; the model could not compile it; the loop reinforced rather than repaired (each diagnosis suggested more optimization, not retreat); and the selection logic could discard an earlier compiling kernel for a later broken one. Each repair in §III-B targets one of those mechanisms. The ablation confirms the diagnosis quantitatively: with the sampling backbone removed ($n{=}1$), v2 falls right back to 3%—i.e. the prompts alone do not fix a weak generator; *breadth under a real verifier* does. This is the paper's central lesson for agent design: for small models on verifiable tasks, allocate compute to candidates evaluated by an oracle before allocating it to reasoning about any single candidate.

D. Root cause 4: infrastructure—generated kernels can hang the GPU

Two production runs stalled mid-study: a sampled kernel deadlocked at runtime (an infinite loop in device code), and since KernelBench evaluates in-process and a Python signal cannot interrupt a running CUDA kernel, one candidate wedged a worker for 14 hours. We now evaluate every candidate in a child process with a hard 240 s wall-clock kill, scoring timeouts as failures. Anyone running LLM-generated GPU code unattended needs this.

E. Limitations

30 tasks, one model, one seed: differences of 1–2 problems are within noise, so we foreground large effects (3%→30% repair; the 0-vs-3 L2 split; the best-of-n curve) over small gaps. Greedy decoding makes the three non-sampling methods near-deterministic; sampling variance is characterized by the best-of-n sweep itself. Results are fp32 on Volta—on GPUs with tensor cores the library-dominance ceiling may shift. Findings about competence-matched scaffolding are claimed for small models only.

VII. Conclusion

Holding a local 7B code model fixed, the interaction strategy around it changes KernelBench outcomes as much as any prompt detail—but not monotonically with sophistication. Profiling-guided prompting is the best overall (43% correct, 27% faster-than-baseline); a multi-agent loop is uniquely able to solve fused level-2 problems but over-engineers trivial ones; and an unrepaired agent is worse than no agent. The decisive design choices were matching ambition to competence and buying correctness with best-of-n sampling under a real execution oracle—test-time compute, not deeper reasoning. Speed remains bounded by cuBLAS/cuDNN: the deliverable a small model can realistically own is the correct fused memory-bound kernel at parity, which is precisely where these methods differentiate. Future work: difficulty-aware routing (guided for L1/L3, agentic for L2 would reach 53% correct on this set), API-grounded decoding to attack the hallucination bottleneck, and fine-tuning on the preference/reflection trajectories the agent already logs.

Code and Data Availability

The agent framework, experiment harness, and all raw artifacts (per-task prompts, generations, evaluation logs, and the figures/tables in this paper) are available at https://github.com/pango0/KernelBench-Agentic.

Acknowledgment

Experiments were conducted on the Taiwania 2 GPU cluster.

References

[1] A. Ouyang, S. Guo, S. Arora, A. L. Zhang, W. Hu, C. Ré, and A. Mirhoseini, "KernelBench: Can LLMs write efficient GPU kernels?," *arXiv:2502.10517*, 2025.

[2] B. Hui *et al.*, "Qwen2.5-Coder technical report," *arXiv:2409.12186*, 2024.

[3] P. Lewis *et al.*, "Retrieval-augmented generation for knowledge-intensive NLP tasks," in *NeurIPS*, 2020.

[4] A. Madaan *et al.*, "Self-Refine: Iterative refinement with self-feedback," in *NeurIPS*, 2023.

[5] X. Chen, M. Lin, N. Schärli, and D. Zhou, "Teaching large language models to self-debug," in *ICLR*, 2024.

[6] N. Shinn, F. Cassano, E. Berman, A. Gopinath, K. Narasimhan, and S. Yao, "Reflexion: Language agents with verbal reinforcement learning," in *NeurIPS*, 2023.

[7] Y. Li *et al.*, "Competition-level code generation with AlphaCode," *Science*, vol. 378, no. 6624, 2022.

[8] B. Brown *et al.*, "Large language monkeys: Scaling inference compute with repeated sampling," *arXiv:2407.21787*, 2024.

Appendix A
Input Prompt (verbatim)

Base task prompt shared by all four methods, shown for level-1 problem 1 (`1_Square_matrix_multiplication`); the architecture block changes per task.

```
You write custom CUDA operators to replace the pytorch
     operators in the given architecture to get speedups.

You have complete freedom to choose the set of operators
     you want to replace. You may make the decision to
     replace some operators with custom CUDA operators and
     leave others unchanged. You may replace multiple
     operators with custom implementations, consider
     operator fusion opportunities (combining multiple
     operators into a single kernel, for example, combining
      matmul+relu), or algorithmic changes (such as online
     softmax). You are only limited by your imagination.

You are given the following architecture:

import torch
import torch.nn as nn

class Model(nn.Module):
    """
    Simple model that performs a single square matrix
     multiplication (C = A * B)
    """
    def __init__(self):
        super(Model, self).__init__()

    def forward(self, A: torch.Tensor, B: torch.Tensor) ->
     torch.Tensor:
        return torch.matmul(A, B)

N = 2048 * 2

def get_inputs():
    A = torch.rand(N, N)
    B = torch.rand(N, N)
    return [A, B]

def get_init_inputs():
    return []

Note: The kernels should be optimized for FP32 (32-bit
     floating point) precision.

Optimize the architecture named Model with custom CUDA
     operators! Name your optimized output architecture
     ModelNew. Output the new code in codeblocks. Please
     generate real code, NOT pseudocode, make sure the code
      compiles and is fully functional. Just output the new
      model code, no other text, and NO testing code!
```

The *guided* method appends a profiling block before the final instruction (excerpt):

```
Here is GPU profiling information for the reference PyTorch
      implementation (measured on the target GPU):

Total CUDA time (reference forward): <X> ms
Top operators by CUDA time (self time when available):
  1. <op>: <t> ms
  ...

Use this profiling data to guide your optimization: focus
     on the highest-cost operators, consider kernel fusion
     where multiple hot ops appear in sequence, and target
     memory-bound vs compute-bound behavior implied by the
     dominant kernels.

Important: PyTorch already dispatches matmul/linear and
     convolution to highly tuned cuBLAS/cuDNN kernels -- a
     hand-written CUDA replacement for a plain GEMM or conv
      will almost always be SLOWER, not faster. Do not
     reimplement those from scratch. Instead, keep using
     the torch operator for the heavy compute and win by
     (1) fusing the cheap memory-bound epilogue around it (
     bias + activation + scale + clamp into one pass) and
     (2) replacing pure elementwise / reduction /
     normalization chains, where a custom kernel can
     actually cut memory traffic. Correctness first: only
     optimize an operator if you can keep the result
     numerically identical to the reference.
```

Appendix B
System Prompts (verbatim, final v2)

All four agent roles run on the same model.

A. Code Analyzer

```
You are a senior GPU performance engineer. You analyse
     PyTorch reference modules and plan how to replace
     their operators with fast custom CUDA/HIP kernels. You
      are precise, concrete, and you never write the final
     kernel here.
```

B. RAG Researcher

```
You are a research assistant for GPU kernel engineering.
     You are given retrieved documentation excerpts and
     must distil ONLY the techniques that are directly
     useful for the task. Ground every claim in the
     excerpts; never invent APIs. Cite the source tag (e.g.
      [memory_optimization.md#2]) after each point.
```

C. Kernel Generator (the competence-matched v2 prompt)

```
You are an expert CUDA/HIP kernel engineer. Your FIRST
     priority is a ModelNew that compiles and matches the
     reference numerically; speed comes only after that.
     PyTorch already dispatches matmul/linear and
     convolution to highly tuned cuBLAS/cuDNN -- do NOT
     reimplement a plain GEMM or conv as a custom kernel,
     it will be slower. Keep the torch op for the heavy
     compute and add a custom kernel only where it clearly
     helps: fusing the memory-bound epilogue (bias/
     activation/scale/clamp) or replacing pure elementwise
     / reduction / normalization chains. When you do write
     a torch.utils.cpp_extension kernel, use load_inline(
     name=..., cpp_sources=..., cuda_sources=..., functions
     =[...]) -- note load_inline takes cpp_sources/
     cuda_sources, NOT a 'sources=' argument, and the
     pybind names in 'functions' must exactly match the C++
      functions you define. Prefer one simple, correct
     fused kernel over several ambitious ones. You output
     ONLY a single ```python``` code block that defines
     class ModelNew(nn.Module) with the same interface as
     Model, plus get_inputs() and get_init_inputs(). No
     prose, no tests.
```

For comparison, the **v1 generator prompt** that produced the failed pilot (3% correct):

```
You are an expert CUDA/HIP kernel engineer. You replace
     PyTorch operators with correct, high-performance
     custom kernels using torch.utils.cpp_extension (
     load_inline) or Triton. You output ONLY a single ```
     python``` code block that defines class ModelNew(nn.
     Module) with the same interface as Model, plus
     get_inputs() and get_init_inputs(). No prose, no tests.
```

D. Feedback Analyzer

```
You are a debugging and performance triage expert for GPU
     kernels. Given an implementation and its evaluation
     result, you produce a crisp diagnosis and a concrete,
     prioritised fix list for the next attempt. You do NOT
     write the full kernel; you direct the fix.
```

Appendix C
Representative Agent Output

A. A correct level-2 solution (the only method to solve any L2 task)

Task `10_ConvTranspose2d_MaxPool_Hardtanh_Mean_Tanh`; final turn; correct, 1.01×. Note the competence-matched shape of the solution: torch ops retained for conv/pool, `clamp+mean` fused into one pass.

```
import torch
import torch.nn as nn
from torch.utils.cpp_extension import load_inline

class ModelNew(nn.Module):
    def __init__(self, in_channels, out_channels,
     kernel_size, stride, padding,
                 maxpool_kernel_size, maxpool_stride,
     hardtanh_min, hardtanh_max):
        super(ModelNew, self).__init__()
        self.conv_transpose = nn.ConvTranspose2d(
     in_channels, out_channels,
            kernel_size, stride=stride, padding=padding)
        self.maxpool = nn.MaxPool2d(kernel_size=
     maxpool_kernel_size,
            stride=maxpool_stride)
        self.hardtanh = nn.Hardtanh(min_val=hardtanh_min,
     max_val=hardtanh_max)

    def forward(self, x):
        x = self.conv_transpose(x)
        x = self.maxpool(x)
        x = self.fused_hardtanh_mean(x)
        x = torch.tanh(x)
        return x

    @staticmethod
    def fused_hardtanh_mean(x):
        return torch.mean(torch.clamp(x, min=-1, max=1),
     dim=(2, 3), keepdim=True)

def get_inputs():
    return [torch.rand(batch_size, in_channels, height,
     width)]

def get_init_inputs():
    return [in_channels, out_channels, kernel_size, stride,
      padding,
            maxpool_kernel_size, maxpool_stride,
     hardtanh_min, hardtanh_max]
```

B. A representative failure (hallucinated API, −RAG variant)

Compiler output excerpt; the model invented a private ATen header and called a Python API from C++:

```
fuse_relu_bn/cuda.cu:7:10: fatal error: ATen/native/cudnn/
     utils.h: No such file or directory
    7 | #include <ATen/native/cudnn/utils.h>

main.cpp:11:22: error: 'F' was not declared in this scope
   11 |  auto bn_output = F.batch_norm(x, weight, bias,
     running_mean, running_var, true, eps, 0.0);
main.cpp:19:56: error: 'fuse_relu_bn_cuda' was not declared
      in this scope
   19 | m.def("fuse_relu_bn_cuda", torch::
     wrap_pybind_function(fuse_relu_bn_cuda), ...);
```

C. A Feedback Analyzer diagnosis (illustrating its failure mode)

Produced after the already-correct turn 0 of the L2 task above; the directive pushes further refactoring of a working design—the over-specification behavior that makes removing this agent *improve* correctness (Table III):

```
CATEGORY: performance
ROOT_CAUSE: The primary bottleneck identified is the `
     fused_hardtanh_mean` function, which involves both
     clamping and mean reduction operations. These
     operations are memory-bound and can benefit from being
      fused with subsequent operations to reduce the
     overall computational overhead.
DIRECTIVE: Refactor the ModelNew class to use the combined
     `fused_hardtanh_mean_tanh` function instead of the
     separate `fused_hardtanh_mean` function, and verify if
      the performance improves. Additionally, investigate
     any potential optimizations for the Convolution
     Transpose and Max Pooling operations.
```

Appendix D
Agent Modifications

Chronological record of every modification, each motivated by an observed failure.

C1. Generator/analyzer prompt inversion ("correct first"). v1 mandated custom `load_inline`/Triton kernels for every

task; v2 (Appendix B.3) makes correctness the stated priority, forbids reimplementing plain GEMM/conv, directs fusion to memory-bound epilogues, and pins the exact `load_inline` signature the model repeatedly hallucinated. The analyzer's section instructions were rewritten the same way (flag matmul/conv as "keep the torch op"). *Effect (with C2): 3%→30% correct, 33%→73% compile.*

C2. Best-of-n sampling under the execution oracle. The generator samples $n{=}4$ candidates per turn (temperature 0.7); the *evaluator*—not the LLM—ranks them by $(\textit{correct}, \textit{compiled}, \textit{speedup})$ and keeps the best. The ablation (Table III) shows this is the single most important component (-27pp when removed); the sweep (Fig. 2) maps $n \in \{1, 2, 4, 8\}$.

C3. Monotone final-kernel selection. Previously the submitted kernel could be the *last* generated turn even if an earlier turn compiled and the last did not. Selection now maximizes $(\textit{correct}, \textit{compiled}, \textit{speedup}, \textit{turn})$ over all turns, so refinement can never make the submission worse. (The same fix was applied to the iterative baseline, for fairness.)

C4. Hard per-candidate evaluation timeout. Each candidate is compiled/checked/timed in a child process killed after 240 s (`agentic/evaluator.py:evaluate_kernel_safe`). Motivation: a generated kernel deadlocked the GPU and—because a Python signal cannot interrupt a running CUDA kernel—wedged a worker for 14 h. Timeouts are scored as ordinary failures.

C5. Optional reference fallback (kept OFF for all reported numbers). A flag can ship the unmodified reference (`ModelNew = Model`) when no turn compiles—useful as a deployment floor, but disabled throughout this study because a passthrough is trivially "correct" at $1\times$ and would inflate benchmark correctness.

C6. Coherent ablation infrastructure. All agentic variants share one best-of-4 backbone configuration, so each ablation flips exactly one component; the report generator renders the component table, the best-of-n sweep, and the scaling figure directly from raw evaluation JSON.