arXiv:2611.11899v1 [cs.DS] 9 Nov 2026

How Far Must a Test Explore? Certifying Composable Read–Write Interfaces from Their Own Random Edits

Jiaxuan Guo GUOJX@STANFORD.EDU
Stanford University

Xingyu Ren
The Chinese University of Hong Kong

Jingxin Yang
Stanford University

Youran Sun
Haizhao Yang
University of Maryland, College Park

Abstract

Random tests of stateful code check laws, like reading back a written value, at states that random writes reach. The laws hold at every state exactly when the code behaves like a record with hidden data, so a pass is taken as evidence. This can fail: the laws can hold at every sample while the code is far from any record. The laws are local, but behaving like a record is global. Self-testing and stability theorems bridge this gap when the query distribution is controlled by the tester or preserved by the tested maps, but here it is generated by the tested writes, which need not be invertible.

We ask how far a test must explore. An interface has m binary variables, each with a read and two writes; in a record, reads return their bits and writes set only their bits. Samples are the ends of geometric runs of mean ℓ of uniformly random writes from a start distribution; a tester adds its own write chains of depth d. The distance is the least, over records with hidden data, of the largest probability that a read or write differs from it at a sample.

From one start state, law checks at $O(m^2 \log m/\epsilon)$ samples certify, at fixed confidence, any threshold ϵ after geometric runs of every positive mean, while runs of fixed positive length can pass at distance $1/2$. For the worst start distribution, take $\epsilon = m^{-\beta}$ with fixed $0 < \beta \le 1$. For $d \le cm$ with fixed $c < 1/2$, a polynomial budget of samples and writes exists exactly when $\ell = \Omega(m)$, where law checks suffice. With no runs, the least depth for this budget is $(1/2 + o(1))m$. Counting samples alone, in the worst case at $\epsilon = 1/(32m)$ and $d \le m/16$, the critical mean for a polynomial budget is $\Theta(m \log(m/d)/\log m)$, by exhaustive search. So one start state needs only law checks, and the worst case at these thresholds needs exploration linear in m if writes are counted.

Keywords: property testing; self-testing; algebraic laws; read–write interfaces; random walks

1. Introduction

Random tests check stateful code in the way Claessen and Hughes (2002) check queues: they run random sequences of operations from a start state and check laws or a model at the reached states. For m binary variables, each with a read and two writes, the laws say that a written value is read back (*success*), that writing back the current value changes nothing (*write-back*), that of two writes to one variable the second wins (*overwrite*), that a write to one variable keeps the others (*preservation*), and that writes to different variables commute (*commutation*) (Foster et al., 2016). For one variable, success, write-back and overwrite are the PutGet, GetPut and PutPut laws of a very well-behaved

total lens (Foster et al., 2007), which the Haskell lens library checks with QuickCheck (Kmett and Charles, 2026).

The laws hold at every state exactly when the code behaves like a record of m bits with *hidden data*, which no read reveals and no write changes (Proposition 9), so a pass is taken as evidence of this behavior. A one-bit record may have the states $A0$, $A1$, $B0$ and $B1$: reads and writes act on the bit, and the letter is hidden data. But a pass checks only the reached states: each law is local, checked at one state with at most two writes in a row, while behaving like a record is a global property of how all states fit together.

Self-testing and stability theorems justify this interpretation in other settings. There, rare failures of local laws imply closeness to an exact solution of the laws (Blum et al., 1993; Rubinfeld and Sudan, 1996; Becker and Mosheiff, 2021; Becker and Chapman, 2023). Their analyses control the distribution at which the laws are checked. The homomorphism test of Blum et al. (1993) checks $f(xy) = f(x)f(y)$ at a uniformly random pair of group elements x, y, so each of x, y and xy is uniform. In stability theorems, the tested permutations preserve the counting measure.

In random tests of stateful software, even in a record, setting a variable to 0 maps two states that differ only in this variable to the same state, so the image of a sample under a write need not be distributed like a sample. The writes under test also generate the sample distribution.

We ask how far beyond its samples a test must explore before a pass certifies that an interface, an implementation of these reads and writes, is close to a record, and we call such a record a *composable* interface (as in database view updates, Bancilhon and Spyratos, 1981). A sample is the end of a *run*: a geometric number of uniformly random writes from a state drawn from a *start distribution*. We measure exploration by the mean ℓ of the runs and the *depth* d of the tester's own chains of writes from each sample, and closeness by the distance of Section 2.

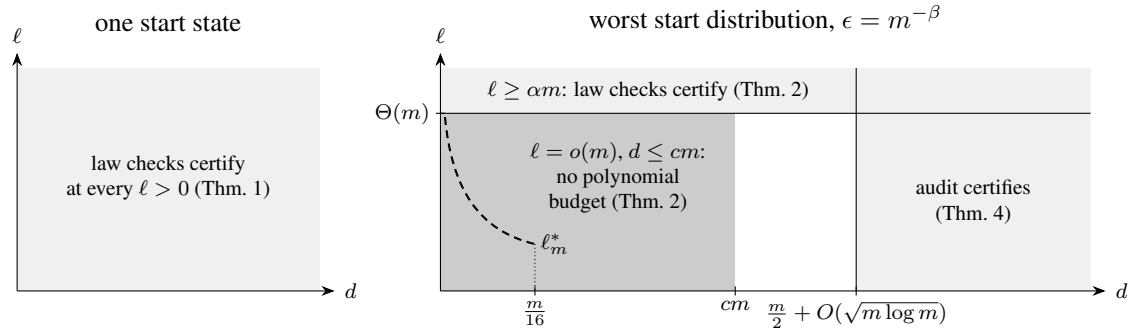

Figure 1: The depth $d \geq 2$ and the run mean ℓ, schematic and not to scale, for fixed $0 < \beta \leq 1$, $\alpha > 0$ and $1/16 < c < 1/2$; Theorem 2 holds for every $0 < c < 1/2$. Budgets count samples and writes. The dashed curve $\ell^*_m = m \log(m/d)/\log m$ counts samples alone, at $\epsilon = 1/(32m)$: polynomially many samples do not suffice below its order and suffice at its order (Thm. 5). The boundary for $d - m/2 = o(m)$ and $0 < \ell = o(m)$ is open.

From one start state, law checks at the ends of geometric runs certify closeness at every positive mean, while for $m \geq 2$ runs of a fixed positive length can pass at distance $1/2$ (Theorem 1). A geometric run can stop after any of its writes, so the sample distribution puts weight on every prefix of a run.

For the worst start distribution at polynomially small thresholds, exploration linear in m is necessary. At each threshold $\epsilon = m^{-\beta}$ with fixed $0 < \beta \leq 1$, a polynomial budget of samples and writes forces $\ell = \Omega(m)$ as long as $d \leq cm$ for a fixed $c < 1/2$ (Theorem 2). Either kind of exploration suffices. Law checks after runs of linear mean meet the budget (Theorem 2). So does an *audit* with no runs, which walks from each sample to a target label and checks that flips rejoin this walk, at depth about $m/2$, the least depth at these thresholds (Theorem 4).

At $\epsilon = 1/(32m)$ and $d \leq m/16$, counting samples alone lowers the necessary mean by a factor of at most $\log m$. A tester that exhausts the ball of radius d around each sample needs only $\ell = \Theta(m \log(m/d)/\log m)$, at a cost in writes exponential in d (Theorem 5). So the two budgets separate: at $d = \lfloor m/16 \rfloor$ and $\ell = m/\log m$, polynomially many samples certify at every threshold $m^{-\beta}$, and polynomially many samples and writes do not. Figure 1 draws these regions.

1.1. Related work

Self-testing and stability. In a composable interface, the flips, which write the opposite value of one variable, commute with each other, each flip undoes itself, and flipping a nonempty set of variables changes every state (Proposition 9). So testing composability resembles testing stability for the group $\{0,1\}^m$ (Chapman and Lubotzky, 2025). Chapman et al. (2026) ask how far unitary matrices $U_1, \ldots, U_m$ that nearly satisfy the relations of the flips, $U_i^2 = I$ and $U_iU_j = U_jU_i$, are from unitary matrices $V_1, \ldots, V_m$ of the same size that satisfy them exactly. For $n \times n$ matrices, they take as squared distance the sum of the squared moduli of the entries of the difference, divided by n. Suppose that, for a relation drawn as an equal mixture of a uniform $U_i^2 = I$ and a uniform $U_iU_j = U_jU_i$ with $i \neq j$, its two sides are at mean squared distance at most η. Then some exact solution has U_i and V_i at mean squared distance $O(\min\{1, m\eta\})$ for a uniform i, and this order is best possible. Their lower bound uses permutations that flip an extra bit conditionally on another bit, close to the twisted interface of Section 4. Our setting differs in two ways. A flip of a general interface need not be a bijection, and the writes under test generate the distribution under which failures are counted.

Testing from samples. Distribution-free testers measure the distance under an unknown distribution from which they sample, and they may also query the tested function at any point (Halevy and Kushilevitz, 2003). In the vertex-distribution-free model of Goldreich (2019), testers may be assumed, at a constant-factor cost, to query only vertices that they have sampled or received as answers. Our testers obey the same restriction for states. In the huge-object model of Goldreich and Ron (2023), a tester draws samples of long strings and probes each at locations of its choice, and the tester with the fewest samples need not be the one with the fewest queries. Our two budgets are these two measures for probes that are writes, and Section 5 separates them at the thresholds $m^{-\beta}$ for the worst start distribution. Local exploration from randomly selected vertices is a standard method in bounded-degree graph testing (Goldreich and Ron, 2002; Bender and Ron, 2002), and our depth is such an exploration radius, bounded as a function of m alone. Identity testers that decide from a single trajectory whether a symmetric Markov chain equals a known reference chain also face states that the trajectory rarely visits (Daskalakis et al., 2018; Cherapanamjeri and Bartlett, 2019); here the chain is generated by the interface under test.

Conformance testing. Conformance testing of finite-state machines decides from input and output sequences whether a black box behaves like a given specification machine, and the length of

the sequences is a central quantity, bounded in terms of the number of states (Vasilevskii, 1973; Chow, 1978; Lee and Yannakakis, 1996). Here the laws are the only specification, and exploration is measured against the number of variables, not of states.

2. Interfaces, laws, runs and distances

Interfaces and laws. An *interface* $I = (r, T)$ supplies readers and setters on a finite set S of states, for $m \geq 1$ binary variables. A *reader* $r_i \colon S \to \{0,1\}$ reports the value of variable i; a *setter* $T_{i,a} \colon S \to S$ is the operation intended to write $a \in \{0,1\}$ to it. The *label* $r(x) = (r_i(x))_{i \in [m]}$ collects all readings at x. The *flip* of variable i is $F_i(x) = T_{i,1-r_i(x)}x$, the write of the opposite value.

A *composable* interface $\tilde{I}$ acts like a record with hidden data. Formally, a bijection $\phi \colon S \to \{0,1\}^m \times \Xi$ gives $\tilde{r}_i(x) = u_i$ and $\phi(\tilde{T}_{i,a}x) = (u[i \mapsto a], z)$ when $\phi(x) = (u, z)$. Here $u[i \mapsto a]$ replaces coordinate i of u by a. Fixing z gives a *cube*, the 2^m states $\phi^{-1}(\{0,1\}^m \times \{z\})$, one of each label, such as $\{A0, A1\}$ in the one-bit record of Section 1. A composable interface on S exists exactly when 2^m divides $|S|$, with $M = |S|/2^m$ cubes.

Fixing the variables and values in a law gives a *law instance*. The following laws hold at every state x of a composable interface, for all $i \neq j$ and a, b, with maps composed from right to left.

- *success*, $r_i(T_{i,a}x) = a$; *write-back*, $T_{i,r_i(x)}x = x$; *overwrite*, $T_{i,a}T_{i,b}x = T_{i,a}x$;
- *preservation*, $r_j(T_{i,a}x) = r_j(x)$; *commutation*, $T_{i,a}T_{j,b}x = T_{j,b}T_{i,a}x$.

Conversely, satisfying all five laws everywhere makes an interface composable (Proposition 9; Power and Shkaravska, 2004).

Runs and samples. A *random edit* applies a uniformly random setter. It sends a distribution μ to μP, where $P = (2m)^{-1} \sum_{i,a} T_{i,a}$. The *start distribution* ν specifies where runs start: a *run* draws its *start state* from ν, makes L independent random edits, and returns its final state as a sample. For geometric run lengths of mean $\ell > 0$, put $q = \ell/(1+\ell)$ and $\Pr(L = s) = (1-q)q^s$, $s \geq 0$. The *sample distribution* is

$$\mu_\ell = (1-q) \sum\nolimits_{s \geq 0} q^s \, \nu P^s, \tag{1}$$

and $\mu_0 = \nu$. For measures, $\mu \leq \mu'$ means $\mu(X) \leq \mu'(X)$ for every $X \subseteq S$. The measure μ_ℓ satisfies $\mu_\ell = (1-q)\nu + q\,\mu_\ell P$, hence $\mu_\ell P \leq (1 + 1/\ell)\mu_\ell$. More generally, μ is λ-*subinvariant* if $\mu P \leq \lambda \mu$, for $\lambda \geq 1$: one random edit increases the probability of any set by at most a factor λ.

Violation rate and distances. Fix a probability measure μ on S. Let $\mathcal{N}$ be the set of states at which some law instance fails. The *violation rate* is $\gamma_\mu(I) = \mu(\mathcal{N})$, the probability that checking all law instances at one sample rejects. The *disagreement* of a candidate $\tilde{I}$ with I is the largest probability that a reader or setter differs,

$$R_\mu(I, \tilde{I}) = \max_{i,a} \max \left\{ \mu\{r_i \neq \tilde{r}_i\},\ \mu\{T_{i,a} \neq \tilde{T}_{i,a}\} \right\}. \tag{2}$$

The *distance* $D_\mu(I)$ is the least disagreement over composable interfaces on S, which exist when 2^m divides $|S|$. The distance $D_\mu^\infty(I)$ is the infimum of the disagreement over composable interfaces on finite sets $S' \supseteq S$, measured under μ on S with an output outside S counted as a disagreement, so $D_\mu^\infty \leq D_\mu$. The theorems below hold for both distances: an upper bound speaks of D_μ only when 2^m divides $|S|$, and every lower bound holds for both. For D_μ, the audit of Section 4 adds

a *bridge check*: it draws a uniform label and, when the depth tests allow, attempts to reach a state with that label and checks, through intermediate states, that the audit's walk sends this state to the same state as the sample (Appendix C.2). The distance on S needs this check because a composable interface on S itself cannot add states to supply missing labels. Subscripts ℓ mean evaluation at μ_ℓ: we write γ_ℓ, D_ℓ and D_ℓ^∞. All logarithms are natural.

Definition 1 (Testers)

(i) Access. *A* tester *holds states only as opaque handles: the samples of μ that it draws and the outputs of its* setter calls*, also called writes. On any held state it may call readers and setters, and it may compare two held states for equality. Its* transcript *lists its draws, calls and comparisons, with the reader outputs and comparison outcomes, and numbers each handle by its first appearance.*

(ii) Depth. *A held state has* depth j *if the shortest chain of the tester's setter calls that leads from a sample to this state has j calls. A setter call has depth one more than its argument, determined when the call is made. A tester has depth d on an input if all its setter calls on this input have depth at most d.*

(iii) Law checks. *A* law-check tester *evaluates law instances at samples and rejects only when one of them fails. It may decide how many samples to draw and which instances to evaluate from the outcomes so far. It calls readers and setters only on samples and on outputs of the setter calls of these evaluations, so its depth is* 2*.*

(iv) One-sided error. *A tester has* one-sided error *if it accepts every composable interface with probability* 1*.*

Depth measures how far calls reach from a sample, not how many calls there are. A call has depth one more than its argument even if it returns a state that the tester already holds, such as its argument. In a composable interface, writing the current value of a variable is such a call. Later calls use the shortest chain of calls that has reached the returned state. For instance, take a sample x with $r_1(x) = 0$ and call $T_{1,1}$ on x, then $T_{1,0}$, then any setter, each on the previous output. In a composable interface, the first output is a new state of depth 1 and the second is x again, so the three calls have depths 1, 2 and 1. If the first two outputs are new states, the calls have depths 1, 2 and 3.

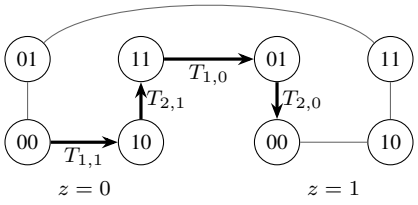

Figure 2: The eight states (u, z) of the interface defined below, each shown by its label u_1u_2, in two layers by the hidden bit z. Lines join states that differ by one flip. The thick path is $(00, 0) \to (10, 0) \to (11, 0) \to (01, 1) \to (00, 1)$.

Certification and budgets. A tester *certifies* $D_\mu < \epsilon$ with confidence $1 - \delta$ if it accepts every composable interface with probability at least $1 - \delta$ and rejects every interface with $D_\mu \geq \epsilon$ with probability at least $1 - \delta$; define this for D_μ^∞ in the same way. Each theorem below certifies the

distance under the sample distribution μ_ℓ of its own mean, and Section 6 compares the measures of different means.

A budget counts samples, or samples and setter calls. It does not count the writes that generate a sample, ℓ in expectation (Lemma 7), and lower bounds leave reader calls, comparisons and computation free. For sequences $\ell_m \geq 0$ and $d_m \geq 2$ and thresholds ϵ_m, a *polynomial budget* is a constant A and, for all large m, a tester of depth at most d_m on samples of μ_{ℓ_m} that certifies at ϵ_m with confidence $2/3$ for all interfaces and start distributions, and uses at most m^A counted operations in expectation on every composable interface. A quantity is *superpolynomial* if, for each A, it exceeds m^A for all large m. A bound $D_\mu \leq K\gamma_\mu$ yields a law-check tester from $\lceil K \log(1/\delta)/\epsilon \rceil$ samples (Appendix A).

Reads and state identity. Success and preservation speak only about reads, while write-back, overwrite and commutation compare states. The interface of Figure 2 separates the two. Its states are $S = \{0,1\}^2 \times \{0,1\}$ with $r(u,z) = u$, and its setters are $T_{1,v}(u,z) = (u[1 \mapsto v], z \oplus (v \oplus u_1)u_2)$ and $T_{2,v}(u,z) = (u[2 \mapsto v], z)$. All its reads are correct: after any sequence of writes, the label is the one that a record predicts. But the thick path of the figure returns to label 00 with the hidden bit toggled, and a composable interface that kept these four writes would return to the state where it started. So the distance is at least $1/8$ under the uniform measure, which is stationary (Proposition 42).

For $\ell > 0$, along independent uniform random writes from a sample of μ_ℓ, reads alone are controlled by the violation rate, with no factor of m (Proposition 41). So the exploration question below asks about *state identity*: whether two write sequences from the same sampled state that predict the same label reach the same state. In the interface of Figure 2, commutation fails at every state, so one law check finds the failure. The instance of Section 4 hides it.

3. One start state

Theorem 1 (One start state) *Let $m \geq 2$.*

(a) Let the start distribution be a point mass $\nu = \delta_o$. There is a coefficient $K = O(m^2 \log m)$, explicit in Appendix G.8, such that for every interface and every $\ell > 0$, $D_\ell \leq \min\{1, K\gamma_\ell\}$ (with D_ℓ^∞ in place of D_ℓ if $2^m \nmid |S|$; Section 2). So, for $0 < \epsilon, \delta < 1$, law checks at $\lceil K \log(1/\delta)/\epsilon \rceil$ independent samples of μ_ℓ certify $D_\ell < \epsilon$ with confidence $1 - \delta$.

(b) For every integer $k \geq 1$, there are a set S of states with $2^m \mid |S|$, an interface on S and a point mass ν such that the sample distribution νP^k of runs of exactly k writes has $\gamma_{\nu P^k} = 0$ and $D^\infty_{\nu P^k} = D_{\nu P^k} = 1/2$.

Why part (b) holds. Figure 3 shows the core, for $m = 2$, with samples uniform on x and w. All law instances hold at x and at w. Take a composable interface, even with additional states, that keeps the four bold entries $T_{1,1}x = y$, $T_{2,1}x = v'$, $T_{2,0}w = y$ and $T_{1,0}w = v$. A setter of a composable interface keeps the cube of its argument, changes only its own variable, and leaves its argument unchanged exactly when the variable already has the written value. So x, y and w lie in one cube, drawn in Figure 3(c). There $y = T_{1,1}x = T_{2,0}w$ gets the label 10, by the values these setters write, and as $T_{1,1}$ changes x and $T_{2,0}$ changes w, x and w get the labels 00 and 11. Then $v' = T_{2,1}x$ and $v = T_{1,0}w$ both get the label 01, which is impossible as $v' \neq v$. So the distance is at least $1/2$. The law instances at x never reach v, and those at w never reach v', so no law instance compares these two states. For $k = 1$, add a start state o with $T_{i,0}o = x$ and $T_{i,1}o = w$: every run of exactly one write ends at x or w (Appendix G.9).

(a)

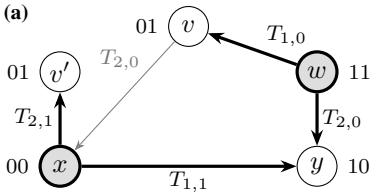

(b)

state	r_1r_2	$T_{1,0}$	$T_{1,1}$	$T_{2,0}$	$T_{2,1}$	laws failing
x	00	x	$\boldsymbol{y}$	x	$\boldsymbol{v'}$	none
w	11	$\boldsymbol{v}$	w	$\boldsymbol{y}$	w	none
y	10	x	y	y	w	commutation
v'	01	v'	w	x	v'	overwrite
v	01	v	w	x	v	overwrite

(c)

01: v', v	11: w
00: x	10: y

Figure 3: An interface with two variables whose square of labels has a split corner, reached at two different states v' and v. (a) Labels r_1r_2 of the states x, y, v', v, w; the samples are the shaded states x and w. Thick arrows are the bold entries of (b), and the gray arrow is $T_{2,0}v = x$. (b) The images under each $T_{i,a}$ and the laws failing at each state. (c) The cube that a composable interface keeping the thick arrows needs, each cell giving a label and its states; the shaded cell holds both v' and v.

Why part (a) holds. In a composable interface, a run from o stays in one cube, and its end state is determined by the *predicted label* B, obtained by applying the run's writes to $r(o)$ as coordinate writes. The proof shows an approximate version for a general interface in four steps (Appendix G).

- *The sample distribution weights every prefix of a run.* A geometric run makes its $(s+1)$-st write with probability q^{s+1}, from a state distributed as νP^s. So by (1), the expected number of writes that a run makes at states of a set X is $\ell\mu_\ell(X)$, and the probability that a run makes a write at a state of $\mathcal{N}$ is at most $\ell\gamma_\ell$. Runs of exactly k writes have no such property: for $k = 1$ above, write-back fails at o, but every run ends at x or w.
- *Runs with one predicted label usually end at one state.* A run *meets a failure with* a variable j if a law instance that involves a setter of j fails at one of its states. Otherwise the last write to j can be moved to the front and the other writes to j deleted, so these writes matter only through B_j. For instance, let a run make the writes $T_{1,1}$, $T_{2,0}$, $T_{1,0}$ and $T_{2,1}$ in this order, and take $j = 1$. Commutation at $T_{1,1}o$ and overwrite at o give $T_{1,0}T_{2,0}T_{1,1}o = T_{2,0}T_{1,0}T_{1,1}o = T_{2,0}T_{1,0}o$, so the run ends at $T_{2,1}T_{2,0}T_{1,0}o$. Now let writes arrive in continuous time at rate 1, and stop the run at an independent exponential time ι of mean ℓ, its *stopping time*, which gives the same sample distribution. Given B and ι, the *histories* of the variables, the times and values of the run's writes to each variable, are independent. So, by the Efron–Stein inequality (van Handel, 2016), which bounds a variance by resampling one independent variable at a time, two independent runs with the same B and ι end at different states with probability at most the sum over j of the probability that resampling the history of j alone, with the same B_j, changes the end state. The j-th term compares the run with the run whose history of j is resampled, and these two end at different states only if one of them meets a failure with j.
- *The stopping time matters little.* A cube has one state of each label, so given B the end state must also vary little with ι. Given B, the conditional histories at different stopping times can be coupled so that the write counts of the variables are pure-birth processes with rates at most $2/m$. At a birth the history of that variable gains a write and may redraw its previous last value, and the end state is unchanged if the old run meets no failure with that variable (Appendix G.2). Appendix G.3 controls the law of ι given B through an integral-kernel bound, which turns this rate into a bound on the variation over ι.

- *One state per label forms a cube.* For each label, choose among the states that read this label the one where runs with this predicted label most often end, weighting a run by $1/(1-\mathrm{e}^{-\iota/m})$, one over the probability, given ι, that it writes a given variable, and choose o for $r(o)$. These states have distinct labels, so they fit into one cube, and filling the other positions, including labels that no state reads, gives the composable interface. A setter $T_{j,a}$ at a sample X acts like resampling the history of j with last value a, and the weight offsets the factor by which this conditioning can raise the probability of a history (Appendix G.4). If neither run meets a failure with j, then $T_{j,a}X$ is the end of the new run, whose predicted label is $B[j \mapsto a]$, so $T_{j,a}X$ is usually the chosen state of this label, which is the output of the composable setter.

4. Many start states

Theorem 2 (Worst start distribution) *Fix $0 < \beta \le 1$ and $0 < c < 1/2$, and let $\epsilon_m = m^{-\beta}$ and $2 \le d_m \le cm$. For both distances, a polynomial budget of samples and setter calls at ϵ_m exists if and only if $\ell_m = \Omega(m)$.*

(a) Necessity. *If $\ell_m = o(m)$, every tester of depth at most d_m that certifies at ϵ_m with confidence $2/3$ for all interfaces and start distributions uses a superpolynomial expected number of samples and setter calls on some composable interface and start distribution, even if it learns the length of each run.*

(b) Sufficiency. *Let $m \ge 2$, $\alpha > 0$ and $K_\alpha = 4\alpha(2\mathrm{e})^{2/\alpha}m^{1+2/\alpha}$. If $\ell \ge \alpha m$, then for all $0 < \epsilon, \delta < 1$, law checks at $\lceil K_\alpha \log(1/\delta)/\epsilon \rceil$ samples of μ_ℓ certify with confidence $1-\delta$, at depth 2 and with at most $2m^2 + 2m$ setter calls per sample. Also $D_\mu \le K_\alpha \gamma_\mu$ for every $(1 + 1/(\alpha m))$-subinvariant measure μ.*

The necessity extends beyond these thresholds. For a fixed $C > 1$, the conclusion of (a) holds for $1/(32m) \le \epsilon_m \le (\log m)^{-C}$ and $\ell_m = o(m \log(1/\epsilon_m)/\log m)$ (Proposition 33(b)); a matching upper bound at this scale is open. So if $(\log m)^C \le \ell_m = o(m)$, there is no polynomial budget at the threshold $1/\ell_m$ either.

Why the necessity holds. The proof builds one family of interfaces at distance above ϵ_m, and the worst start distribution comes from this family (Appendix E).

- *One twist.* Split $8J$ variables into two blocks of $4J$, with values u' and u''. In the *twisted interface*, a setter of the first block also toggles a hidden bit when it moves $|u'|$ across $2J$ while $|u''| \ge 2J$. So commutation between the blocks fails only where both Hamming weights are near $2J$ (Figure 4).
- *No composable interface, even with additional states, can match the twist.* The loop of Figure 4 writes each coordinate it touches once to 1 and once back to 0, and $|u'|$ crosses $2J$ upwards while $|u''| = J$ and downwards while $|u''| = 3J$. So it returns to its start label with the hidden bit toggled, and a composable interface that kept every write would return to the start state. Averaging over such loops, for $\ell = o(m)$ and every composable interface, the expected number of setters of the two blocks that differ from it at a sample is at least a positive constant.
- *A frozen block hides the twist.* Each start state has a *frozen block* of weight J or $3J$. Suppose that its threshold function, $f_1 = \mathbf{1}\{|u'| \ge 2J\}$ or $f_2 = \mathbf{1}\{|u''| \ge 2J\}$, keeps its initial value at every label that the tester queries from a sample. If the first block is frozen, or the second with $f_2 = 0$, no setter toggles the hidden bit. If the second block is frozen with $f_2 = 1$, a setter of

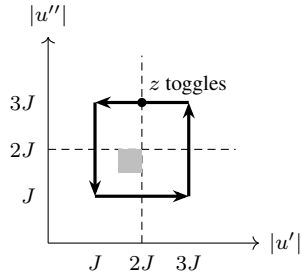

Figure 4: The Hamming weights of the two blocks of $4J$ variables of one group, with the thresholds $2J$ (dashed). Dark square, enlarged: $\{2J-1, 2J\}^2$, where commutation between the blocks can fail. Thick: the start states, each with a frozen block of weight J or $3J$, and the loop along them in the order bottom, right, top, left. The hidden bit z toggles only where the top edge crosses the threshold (dot). Appendix E.1 writes out $J = 1$.

the first block toggles the hidden bit exactly when it changes f_1, so the hidden bit plus f_1 stays constant. Either way, up to this renaming the tester sees a composable interface (Lemma 28).

- *The distance.* For one twist at a random place, dividing the expected number above among all $2m$ setters gives a disagreement lower bound of only order $1/m$. So the instance uses $N_g = \lfloor m/(8J) \rfloor$ disjoint *groups* of $8J$ variables, each with its own twist, hidden bit and frozen block. Their expected numbers add up before the division, which gives a lower bound of order N_g/m, or $1/J$. With $J = \lfloor 1/(512\epsilon_m) \rfloor$, this lower bound exceeds ϵ_m (Lemma 27). At thresholds above order $1/m$, this J is much smaller than m, and the instance hides these small blocks.
- *Copies separate samples.* A state of the instance consists of a copy, a label in $\{0,1\}^m$ and one hidden bit for each group. Each of 2^m *copies* of the states, closed under the setters, draws its own groups and blocks, and polynomially many samples lie in distinct copies except with exponentially small probability. On this event, no state is reached from two samples, and one sample says nothing about where the blocks of another lie.
- *A mask keeps the tester from aiming at a block.* Each copy also draws a *mask*, a random vector added to the labels before the weights are computed. After averaging over the mask, the label of a sample is uniform and independent of the groups, the blocks, and the start label and edits relative to the mask, given the run length. On the composable interface the answers depend only on the labels of the samples and the setters called. So the coordinates that the tester flips are independent of the groups and the frozen blocks, however it chooses them (Lemma 30).
- *The margin: half the dimension.* Flipping d coordinates of a sample flips $4Jd/m$ coordinates of a frozen block in the mean. At a start state of weight J, a quarter of the coordinates of this block are ones, so each flip in this block raises its weight by $1/2$ in the mean, and the mean weight rises to $J + 2Jd/m$. Call the distance from the mean weight to the threshold the *margin*. Then flips multiply the margin by $1 - 2d/m$, for every J. This is a positive constant for $d \le cm$ with $c < 1/2$, and the margin vanishes only when d is about $m/2$.
- *Time: long runs, and where* $\log(1/\epsilon)$ *comes from.* A run lowers the margin as well. Each write to a group replaces a coordinate by a fair bit, so $\hat{t}$ such writes multiply the margin by $(1 - 1/(8J))^{\hat{t}}$, and the margin falls to the fluctuations of the weight after order $J \log J$ of them. The number of writes to one group in a run is geometric with mean $\ell_0 = 8J\ell/m$, so with

$q_0 = \ell_0/(1+\ell_0)$ it exceeds $\hat{t}$ with probability $q_0^{\hat{t}+1} \le \exp[-(\hat{t}+1)/(1+\ell_0)]$. For $\ell > 0$ and $\hat{t}$ of order $J \log J$, this exponent has order $\min\{J \log J, m \log J/\ell\}$, so the tail is superpolynomially small for $\ell = o(m \log J/\log m)$ once $J \log J/\log m \to \infty$. As J has order $1/\epsilon_m$, this is the scale stated after Theorem 2, and at $\epsilon_m = m^{-\beta}$ it reads $\ell = o(m)$. These bounds hold state by state and group by group, so a union over them needs a polynomial budget (Lemma 31).

Why long runs suffice. A *word* is a sequence of setters. The *path* of x under $W = E_t \cdots E_1$ with $E_s = T_{i_s,a_s}$ is $x_0 = x$, $x_s = E_s x_{s-1}$, and the *fiber* of x is $W^{-1}(Wx)$. In a composable interface, a word that writes every variable sends two states to the same state exactly when they lie in one cube, so fibers serve as candidate cubes. With the *old value* $b_s = r_{i_s}(x_{s-1})$, step s *passes* if

$$r(x_s) = r(x_{s-1})[i_s \mapsto a_s], \qquad T_{i_s,b_s} x_s = x_{s-1}. \tag{3}$$

Where all steps pass, $x \mapsto (r(x), Wx)$ is injective, so a fiber fits into a cube at its labels.

Take a word of $t = \lceil m \log(2m) \rceil$ random setters, conditioned on writing every variable, an event of probability at least $1/2$. Before conditioning, the state after s random edits has distribution $\mu P^s \le \lambda^s \mu$, so the path meets $\mathcal{N}$ with probability at most $2 \sum_{s<t} \lambda^s \gamma_\mu$. For μ_ℓ, $\lambda = 1 + 1/\ell$, and λ^t is polynomial exactly when $\ell \gtrsim m$. If the path avoids $\mathcal{N}$, commutation moves any setter along it until the first write to its variable absorbs it, so every setter at x stays in the fiber (Lemma 13). With added states, each fiber takes a cube of its own, and this finishes the proof for D^∞.

On the same states there are only M cubes, so the proof keeps the M heaviest fibers and bounds the probability that this loses (Lemma 11). Here the μ-mass of a fiber $W^{-1}(z)$ counts only its states whose path avoids $\mathcal{N}$, and if weights $\kappa(z) \in [0,1]$ on the fibers of positive mass sum to at most M, this loses at most the mean of $1 - \kappa(Wx)$ for $x \sim \mu$ besides the paths that meet $\mathcal{N}$.

Lemma 3 (Capacity) *A transition of $T_{i,a}$ at x is* kept *if it passes (3), with $T_{i,a}x = x$ when $a = r_i(x)$. The* capacity $\kappa^\circ(z)$ *of a state z is the probability that a second random word, drawn like the first, makes only kept transitions along the path of z. For every label b, the capacities of the states with label b sum to at most M.*

If the transition of $T_{i,a}$ from x to $y \ne x$ is kept, then the transition of $T_{i,r_i(x)}$ from y back to x is kept. The random word has the same law as its reverse, as its setters are independent and uniform before the conditioning, and reversal keeps the event of writing every variable. So a walk along kept transitions leads from x to y with the same probability as from y to x, and this symmetry replaces the invariance of the counting measure, which permutations have and general setters lack. By this symmetry, the sum of these capacities equals the sum over all states x of the probability that such a walk from x ends at label b, at most 2^{-m} each as the written label is uniform (Appendix D.2). Passing paths end at the label that W writes, so Lemma 11 applies, losing at most the mean of $1 - \kappa^\circ(Wx)$ over the covering word W and $x \sim \mu$. As a transition fails to be kept only where a law other than commutation fails, this mean is at most $2\lambda^t$ times the bound $2 \sum_{s<t} \lambda^s \gamma_\mu$ on paths that meet $\mathcal{N}$. Adding the earlier path loss gives K_α (Appendix D.4).

An audit replaces the run by its own walk from each sample. Fix a *target label* $\hat{a}$ and put $W_{\hat{a}} = T_{m,\hat{a}_m} \cdots T_{1,\hat{a}_1}$. The *path check* follows the path of a sample x under $W_{\hat{a}}$ and rejects if a step fails (3). Each step that writes a new value first checks the label of its output and then calls the *undo*, the setter of the old value, whose output must be the state before the step. The *flip check* rejects if write-back fails at x, if a flip $y_i = F_i x$ has a wrong label, or if the path of y_i fails (3) or

does not join the path of x at step i. In a composable interface, step i is the first write to variable i and erases the flip, so the two paths meet there. Passing both checks puts every setter at x in the fiber of x with the updated label (Lemma 15).

The calls of these checks have depth at most $|r(x) \oplus \hat{a}| + 2$, the Hamming distance from the label of x to the target label plus two (Lemma 16). A passed step that writes the current value returns its argument, so the path of x gains depth only at the coordinates where $r(x)$ differs from $\hat{a}$. The flip adds one more, and the undo of a step adds one more. A uniform target label lies at distance about $m/2$ from the label of a sample, so depth about $m/2$ covers all but a small fraction of the rounds, and the other rounds accept without checking. Algorithm 1 in Appendix C lists the calls of one round.

On the interface of Figure 3, with samples uniform on x and w, take $\hat{a} = 00$ and $d = 3$, so $W_{\hat{a}} = T_{2,0}T_{1,0}$. From w, step 1 calls $T_{1,0}w = v$, and its undo $T_{1,1}v = w$ passes. Step 2 calls $T_{2,0}v = x$, but its undo $T_{2,1}x = v'$ is not v, so the path check rejects w, although every law instance holds at w.

Theorem 4 (Audit) *Let μ be a probability measure on S and $0 < \epsilon, \delta < 1$. Rounds of Algorithm 1 have one-sided error.*

(a) For $d = m + 1$ and any fixed $\hat{a}$, a round has depth at most $m + 1$, makes at most $m^2 + 5m$ setter calls and rejects with probability at least D_μ^∞, so $\lceil \log(1/\delta)/\epsilon \rceil$ rounds certify. With the bridge check of Appendix C.2, the same holds for D_μ, with at most $2m^2 + 6m$ setter calls per round.

(b) Let $2 \le d \le m + 1$ and $\tau_j = \Pr\{\mathrm{Bin}(m, 1/2) > d - j\}$. For a uniform $\hat{a}$, a round has depth at most d and rejects with a probability p such that $D_\mu^\infty \le p + \tau_2$; with the bridge check, $D_\mu \le p + (m + 2)\tau_2$. We call the term added to p the tail term*; if it is at most $\epsilon/2$, $\lceil 2\log(1/\delta)/\epsilon \rceil$ rounds certify.*

(c) Let $\ell = 0$ and $\epsilon_m = m^{-\beta}$ with fixed $0 < \beta \le 1$. The least depth with a polynomial budget of samples and setter calls at ϵ_m is $(1/2 + o(1))m$, for both distances.

By Hoeffding's inequality, depth $m/2 + O(\sqrt{m \log m})$ makes the tail term of (b) at most $\epsilon_m/2$. This is the upper bound in (c), and the lower bound is the instance of Theorem 2 at mean zero (Appendix E.7).

5. Counting samples alone

Theorem 5 (Samples alone) *Let $2 \le d_m \le \lfloor m/16 \rfloor$ and $\ell_m^* = m \log(m/d_m)/\log m$. For both distances, a polynomial budget of samples at $\epsilon_m = 1/(32m)$ exists if and only if $\ell_m = \Omega(\ell_m^*)$.*

(a) Necessity. *If $\ell_m = o(\ell_m^*)$, every tester of depth at most d_m that certifies at $\epsilon_m = 1/(32m)$ with confidence $2/3$ for all interfaces and start distributions uses a superpolynomial expected number of samples on some composable interface and start distribution, whatever its number of setter calls.*

(b) Sufficiency. *If $\ell_m \ge b\ell_m^*$ with a constant $b > 0$, then for all $0 < \epsilon, \delta < 1$, evaluating all law instances at the ends of all words of length at most $d_m - 2$ from each sample certifies from a number of samples polynomial in m, $1/\epsilon$ and $\log(1/\delta)$. The setter calls of this tester can be exponential in d_m. For $m \ge 96$, $d = \lfloor m/16 \rfloor$ and $\ell \ge \varsigma m/\log m$ with $\varsigma > 0$, $\lceil 140 m^{2+8/\varsigma} \log(1/\delta)/\epsilon \rceil$ samples suffice.*

Why Theorem 5 holds. For sufficiency, a random word need not write every variable. If it misses at most $d-2$ of them, the exhaustive check at its end stands in for the missing writes, as all words of length at most $d-2$ on the missing variables behave like writes in a record (Theorem 35). Writing all but $d-2$ variables takes about $t = 2m\log(m/d)$ random writes, against $m\log m$ for all of them. The factor $\lambda^t = (1+1/\ell)^t$ that moves failures along the word to the samples is polynomial exactly when $t/\ell \lesssim \log m$, which is $\ell \gtrsim \ell^*_m$.

For necessity, take the instance of Theorem 2 with one group of $8\lfloor m/8 \rfloor$ variables and no mask. The transcripts stay equal as long as the threshold function of the frozen block is constant on the whole ball of radius d around a sample. For a run with at most $\hat{t}$ writes to the group, this holds with high probability if the margin $J(1-1/(8J))^{\hat{t}}$ exceeds d by more than the fluctuations of the weight, which allows $\hat{t}$ of order $m\log(m/d)$. The writes to this group, which holds almost all variables, exceed $\hat{t}$ with probability at most $\exp[-\hat{t}/(1+\ell)]$, so polynomially many samples include such a run only if $\hat{t}/(1+\ell) \lesssim \log m$, the same scale $\ell \gtrsim \ell^*_m$ (Appendix E.5).

The two theorems differ in how the tester's writes can lower the margin. Exhausting the ball reaches the labels that subtract d from the margin, at a cost in setter calls exponential in d. Writes chosen without knowing the mask only multiply the margin by about $1-2d/m$, so with a polynomial budget the twist appears only once the margin falls to the fluctuations, after order $m\log m$ writes.

The two budgets separate. Take $d = \lfloor m/16 \rfloor$ and $\ell = m/\log m$. For $m \geq 96$, part (b) with $\varsigma = 1$ says that $\lceil 140 m^{10}\log(1/\delta)/\epsilon \rceil$ samples certify at every ϵ, so polynomially many samples certify at each $\epsilon_m = m^{-\beta}$. As $\ell = o(m)$, Theorem 2(a) with $c = 1/16$ says that at each such ϵ_m polynomially many samples and setter calls do not.

6. Discussion

Which measure a certificate speaks about. Law checks after runs certify μ_ℓ, and the audit with no runs certifies ν, so these are two different certificates. A two-state interface shows the difference: with target label 0 and depth 2, every audit round passes, and $D_\nu = 0$, while after one random write the distance is $1/2$ (Proposition 43).

A tester that may change the mean of its runs can still certify the distance under the original μ_ℓ. With $L = \max\{\ell, m\}$, $\mu_\ell \leq \frac{1+L}{1+\ell}\mu_L$, so law checks at mean L certify the distance under μ_ℓ from polynomially many samples (Proposition 44). For $\ell > 0$, a distance certified under μ_ℓ bounds how often independent random edits from a sample leave the path of one composable interface (Corollary 10).

How large the worst start distribution must be. With added states and no runs, depth one and a budget polynomial in m and N suffice for a start distribution of support at most N, as a tester can compare states across samples (Proposition 45). So at mean zero, the necessity of Theorem 2 for D^∞ needs start distributions of superpolynomial support.

Larger thresholds. At a fixed threshold, the question for law checks reduces to one quantity. Call a state *locally exact to depth* d if every word of length at most d from it ends at the label that coordinate writes predict, and two such words end at one state exactly when these labels agree. Let a_d be the supremum of $D^\infty_\nu(I)$ over finite interfaces I of any number of variables and distributions ν on such states. Law checks at each fixed positive mean certify D^∞ at each fixed positive threshold from polynomially many samples if and only if $a_d \to 0$ (Theorem 39). Without runs, a sample

bound uniform over interfaces and start distributions requires depth $\Omega(\min\{m, 1/\epsilon\})$ at $0 < \epsilon \leq 1/512$ and confidence above $1/2$, unconditionally (Theorem 34).

Tightness. For $\epsilon \leq 1/16$ and $\delta \leq 1/3$, the number of samples of the audit is optimal up to a constant factor (Proposition 18). At $\ell = \alpha m$ the coefficient of law checks is at least of order $m^{1/(2\alpha)}$ (Proposition 22), so the degree of the law-check budget obtained from this coefficient grows with $1/\alpha$.

Scope. The results assume binary variables, uniformly random writes in the runs, geometric run lengths, exact comparison of states, and a tester that is not given the start states or write sequences of the sampling runs (Definition 1(i)). The lower bounds apply to testers that stay local: the audit of depth $m + 1$ is always available.

Open problems. First, whether $a_d \to 0$. Second, the condition on the mean for a polynomial budget at depths with $d_m - m/2 = o(m)$. Third, the orders of the coefficients in Theorem 1 and in law checks after long runs. Fourth, how the exploration needed at positive means varies between one start state and the worst start distribution, for instance at support polynomial in m.

References

François Bancilhon and Nicolas Spyratos. Update semantics of relational views. *ACM Transactions on Database Systems*, 6(4):557–575, 1981.

Oren Becker and Michael Chapman. Stability of approximate group actions: uniform and probabilistic. *Journal of the European Mathematical Society*, 25(9):3599–3632, 2023.

Oren Becker and Jonathan Mosheiff. Abelian groups are polynomially stable. *International Mathematics Research Notices*, 2021(20):15574–15632, 2021.

Michael A. Bender and Dana Ron. Testing properties of directed graphs: Acyclicity and connectivity. *Random Structures & Algorithms*, 20(2):184–205, 2002.

Manuel Blum, Michael Luby, and Ronitt Rubinfeld. Self-testing/correcting with applications to numerical problems. *Journal of Computer and System Sciences*, 47(3):549–595, 1993.

Michael Chapman and Alexander Lubotzky. Stability of homomorphisms, coverings and cocycles I: Equivalence. *European Journal of Mathematics*, 11:Article 13, 2025.

Michael Chapman, Thomas Vidick, and Henry Yuen. Efficiently stable presentations from error-correcting codes. *Discrete Analysis*, 2026.

Yeshwanth Cherapanamjeri and Peter L. Bartlett. Testing symmetric Markov chains without hitting. In *Proceedings of the Thirty-Second Conference on Learning Theory*, pages 758–785, 2019.

Tsun S. Chow. Testing software design modeled by finite-state machines. *IEEE Transactions on Software Engineering*, SE-4(3):178–187, 1978.

Koen Claessen and John Hughes. Testing monadic code with QuickCheck. In *Proceedings of the 2002 ACM SIGPLAN Workshop on Haskell*, pages 65–77, 2002.

Constantinos Daskalakis, Nishanth Dikkala, and Nick Gravin. Testing symmetric Markov chains from a single trajectory. In *Proceedings of the 31st Conference on Learning Theory (COLT)*, pages 385–409, 2018.

J. Nathan Foster, Michael B. Greenwald, Jonathan T. Moore, Benjamin C. Pierce, and Alan Schmitt. Combinators for bidirectional tree transformations: A linguistic approach to the view-update problem. *ACM Transactions on Programming Languages and Systems*, 29(3):Article 17, 2007.

Simon Foster, Frank Zeyda, and Jim Woodcock. Unifying heterogeneous state-spaces with lenses. In *Theoretical Aspects of Computing – ICTAC 2016*, pages 295–314. Springer, 2016.

Oded Goldreich. Testing graphs in vertex-distribution-free models. In *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing (STOC)*, pages 527–534, 2019.

Oded Goldreich and Dana Ron. Property testing in bounded degree graphs. *Algorithmica*, 32(2): 302–343, 2002.

Oded Goldreich and Dana Ron. Testing distributions of huge objects. *TheoretiCS*, 2:Article 12, 2023.

Shirley Halevy and Eyal Kushilevitz. Distribution-free property testing. In *Approximation, Randomization, and Combinatorial Optimization: Algorithms and Techniques (APPROX 2003 and RANDOM 2003)*, pages 302–317, 2003.

Chris A. J. Klaassen. On an inequality of Chernoff. *The Annals of Probability*, 13(3):966–974, 1985.

Edward Kmett and Oliver Charles. *lens-properties: QuickCheck Properties for lens*, 4.11.1 edition, 2026. URL https://hackage-content.haskell.org/package/lens-5.3.6/lens-5.3.6.tar.gz. Source bundled with lens 5.3.6.

David Lee and Mihalis Yannakakis. Principles and methods of testing finite state machines—a survey. *Proceedings of the IEEE*, 84(8):1090–1123, 1996.

John Power and Olha Shkaravska. From comodels to coalgebras: State and arrays. *Electronic Notes in Theoretical Computer Science*, 106:297–314, 2004.

Ronitt Rubinfeld and Madhu Sudan. Robust characterizations of polynomials with applications to program testing. *SIAM Journal on Computing*, 25(2):252–271, 1996.

Adrien Saumard. Weighted Poincaré inequalities, concentration inequalities and tail bounds related to Stein kernels in dimension one. *Bernoulli*, 25(4B):3978–4006, 2019.

Ramon van Handel. Probability in high dimension. APC 550 lecture notes, Princeton University, 2016. URL https://web.math.princeton.edu/~rvan/APC550.pdf. Version of December 21, 2016.

M. P. Vasilevskii. Failure diagnosis of automata. *Cybernetics*, 9(4):653–665, 1973.

Appendix A. Preliminaries

Lemma 6 (Properties of the sample distribution) *For all $\ell > 0$ and integers $s \geq 0$, (i) $\mu_\ell = (1-q)\nu + q\,\mu_\ell P$, (ii) $\mu_\ell P^s \leq q^{-s}\mu_\ell$ and (iii) $\nu \leq (1+\ell)\,\mu_\ell$.*

Put $t_m = \lceil m \log(2m) \rceil$, and let h_t be the probability that t independent uniform draws cover $[m]$.

Lemma 7 (Samples and coverage) *For every $\ell > 0$, a sample as in Section 2 has distribution μ_ℓ and expected cost $\mathbb{E}L = \ell$ setter calls. Also $h_t \geq 1 - m(1-1/m)^t \geq 1 - m\mathrm{e}^{-t/m}$, $h_t > 0$ for $t \geq m$, and $h_{t_m} \geq 1/2$. Finally $t_m \leq m\log(2m) + 1$, and $t_m \geq m$ for $m \geq 2$.*

Proof of Lemmas 6 and 7 Write $\mu = \mu_\ell$. Given $L = s$, the returned state has distribution νP^s, so the sample has distribution (1), and $\mathbb{E}L = \sum_{s\geq 1} \Pr(L \geq s) = \sum_{s\geq 1} q^s = q/(1-q) = \ell$. Reindexing the series gives (i) of Lemma 6, $q\mu P = (1-q)\sum_{s\geq 0} q^{s+1}\nu P^{s+1} = \mu - (1-q)\nu$. Dropping either nonnegative term of (i) gives (iii), $\nu \leq (1-q)^{-1}\mu = (1+\ell)\mu$, and $\mu P \leq q^{-1}\mu$. By positivity of P, induction with $\mu P^{s+1} \leq q^{-s}\mu P \leq q^{-(s+1)}\mu$ gives (ii). For coverage, a union bound over the variables absent from the t draws gives the first inequality, and $1 - 1/m \leq \mathrm{e}^{-1/m}$ the second. The first m draws are distinct with positive probability, so $h_t > 0$ when $t \geq m$. For $m = 1$ coverage is certain; for $m \geq 2$ and $t = t_m \geq m\log(2m)$, $m\mathrm{e}^{-t/m} \leq 1/2$. The ceiling gives $t_m \leq m\log(2m) + 1$, and $\log(2m) \geq 1$ gives $t_m \geq m$ for $m \geq 2$. ■

Lemma 8 (Subinvariant measures) *Let $\lambda \geq 1$. If μ is λ-subinvariant, then $\mu P^s \leq \lambda^s \mu$ for all $s \geq 0$.*

Proof As P is nonnegative, applying it to both sides of $\mu P^s \leq \lambda^s \mu$ gives $\mu P^{s+1} \leq \lambda^s \mu P \leq \lambda^{s+1}\mu$, and induction gives the claim. ■

Proposition 9 (Characterization) *An interface satisfies the five laws at every state if and only if it is composable.*

Proof A composable interface satisfies the five laws by direct evaluation. Conversely, suppose the laws hold everywhere. By success and preservation, the flip F_i changes exactly reader i. Overwrite and write-back give $F_i^2 x = T_{i,r_i(x)} T_{i,1-r_i(x)} x = T_{i,r_i(x)} x = x$. For $i \neq j$, preservation makes the values chosen by F_i and F_j in F_iF_jx and F_jF_ix equal to the values chosen at x, and commutation of the corresponding setters gives $F_iF_j = F_jF_i$. These involutions define an action of $\{0,1\}^m$ under bitwise addition. Applying $u \in \{0,1\}^m$ changes the label by u, so nonzero elements move every state and every orbit has 2^m points, one per label. With $\mathcal{Z}$ the set of orbits, $x \mapsto (r(x), [x])$ is the required bijection; by write-back, $T_{i,a}x$ is x if $a = r_i(x)$ and F_ix otherwise. ■

Law checks at samples. The five laws of Section 2 have $4m^2+3m$ law instances: $2m$ for success, m for write-back, $4m$ for overwrite, $2m(m-1)$ for preservation and $2m(m-1)$ for commutation, one for each pair $i<j$ and values a,b. Let x be a state with label u, and put $y_i = T_{i,1-u_i}x$. The $2m^2+2m$ setter calls $T_{i,a}x$, $T_{i,a}y_i$ and $T_{i,a}y_j$, for all $i \neq j$ and a, give every law instance at x as a reading or a comparison of their outputs, given write-back at x; for instance, overwrite of i compares $T_{i,a}y_i$ with $T_{i,a}x$.

Let K be a coefficient with $D_\mu \le K\gamma_\mu$ for all μ in a class. Evaluating all law instances at $n = \lceil K\log(1/\delta)/\epsilon \rceil$ independent samples of μ and rejecting when one fails has one-sided error, by Proposition 9. If $D_\mu \ge \epsilon$, each sample shows a failure with probability $\gamma_\mu \ge \epsilon/K$, so all n samples pass with probability at most $\mathrm{e}^{-n\epsilon/K} \le \delta$. The same holds for D_μ^∞. Conversely, a law-check tester rejects with probability at most γ_μ times its expected number of samples, as the decision to draw a sample does not depend on that sample.

What a distance bound gives. A small distance controls sequences of random edits from the samples. Along a sequence, the same setters are applied under I and under a closest composable interface.

Corollary 10 (Sequences of random edits) *Let $\ell > 0$, let $\tilde I$ be a composable interface on S with $R_{\mu_\ell}(I,\tilde I) = D_\ell$, and let $k \ge 1$ be an integer. Draw $x \sim \mu_\ell$ and independent uniformly random setters $(i_1,a_1),\dots,(i_k,a_k)$, independent of x, and put $X_0 = \tilde X_0 = x$, $X_s = T_{i_s,a_s}X_{s-1}$ and $\tilde X_s = \tilde T_{i_s,a_s}\tilde X_{s-1}$ for $s = 1,\dots,k$. Then the probability that $X_s \neq \tilde X_s$ for some $s \le k$ is at most $\ell\,(q^{-k}-1)\,D_\ell$, and the probability that this happens or $r(X_k) \neq \tilde r(\tilde X_k)$ is at most $(\ell\,(q^{-k}-1) + mq^{-k})\,D_\ell$. For $k \le \ell$, these bounds are at most $\mathrm{e}k\,D_\ell$ and $\mathrm{e}(k+m)\,D_\ell$.*

Proof At the first s with $X_s \neq \tilde X_s$, the setters T_{i_s,a_s} and $\tilde T_{i_s,a_s}$ differ at $X_{s-1} = \tilde X_{s-1}$. Here X_{s-1} has distribution $\mu_\ell P^{s-1} \le q^{-(s-1)}\mu_\ell$ by Lemma 6(ii), and (i_s,a_s) is uniform and independent of it, so averaging over the $2m$ setters bounds this probability by $q^{-(s-1)}D_\ell$. Summing over s with $\sum_{s<k} q^{-s} = \ell(q^{-k}-1)$ gives the first bound. If $X_s = \tilde X_s$ for every $s \le k$ but $r(X_k) \neq \tilde r(\tilde X_k)$, then $r_i(X_k) \neq \tilde r_i(X_k)$ for some i. As X_k has distribution $\mu_\ell P^k \le q^{-k}\mu_\ell$, this has probability at most $q^{-k}\sum_i \mu_\ell\{r_i \neq \tilde r_i\} \le mq^{-k}D_\ell$, which gives the second bound. For $k \le \ell$, $q^{-k} = (1+1/\ell)^k \le \mathrm{e}^{k/\ell} \le \mathrm{e}$ and $\ell(q^{-k}-1) \le \ell(\mathrm{e}^{k/\ell}-1) \le k\mathrm{e}^{k/\ell} \le \mathrm{e}k$, by $\mathrm{e}^y - 1 \le y\mathrm{e}^y$ for $y \ge 0$. ∎

Appendix B. Completion and reversible words

This appendix holds the tools behind the upper bounds of Theorems 2, 4 and 5. In a composable interface, a word that writes every variable sends two states to the same state exactly when they lie in one cube. So the fibers of such a word serve as candidate cubes in a general interface. With added states, each fiber can take a cube of its own. On the same states there are only M cubes, so the proofs keep the M heaviest fibers and bound the probability that this loses.

Lemma 11 (Completion by heavy fibers) *Let $|S| = M2^m$, let μ be a probability measure on S, let W be any map $S \to S$, let $G \subseteq S$ be a set on which $x \mapsto (r(x), Wx)$ is injective, and let $A \subseteq G$. Put $w_z = \mu(A \cap W^{-1}(z))$ for $z \in W(S)$, let $\kappa\colon W(S) \to [0,1]$ satisfy $\sum_{z:\,w_z>0}\kappa(z) \le M$, and put $\Delta = \sum_z w_z(1-\kappa(z))$. There are a composable interface $\tilde I$ on S and a set $G' \subseteq G$ with $\mu(A \setminus G') \le \Delta$ such that $\tilde r = r$ on G', and $\tilde T_{i,a}x = T_{i,a}x$ whenever $x \in G'$, $y = T_{i,a}x \in G$, $Wy = Wx$ and $r(y) = r(x)[i \mapsto a]$.*

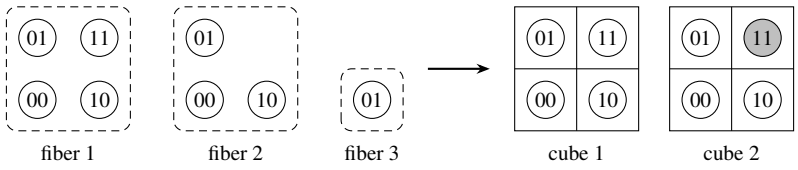

Figure 5: Completion by heavy fibers for $m = 2$ and $M = 2$ cubes. Left: three fibers of a word, with each state shown by its label. Right: the two cubes when the first two fibers carry the most probability. The state of the third fiber (gray), whose label was 01, fills the empty position and reads 11 there.

Proof Let $\Omega_z = G \cap W^{-1}(z)$. By injectivity, the states of Ω_z have distinct labels. Let Z_* be a set of at most M points z with nonempty Ω_z and largest w_z, containing all of them if there are at most M, and let G' be the union of the sets Ω_z with $z \in Z_*$. If at most M sets Ω_z are nonempty, then $G' = G \supseteq A$. Otherwise $|Z_*| = M$; let $w_* = \min_{z \in Z_*} w_z$. If $w_* = 0$, then $w_z = 0$ for $z \notin Z_*$. If $w_* > 0$, then

$$\sum_{z \in Z_*} w_z - \sum_z w_z \kappa(z) = \sum_{z \in Z_*} w_z (1 - \kappa(z)) - \sum_{z \notin Z_*} w_z \kappa(z) \geq w_* \Big(M - \sum_{z:\, w_z > 0} \kappa(z) \Big) \geq 0.$$

In both cases $\sum_{z \in Z_*} w_z \geq \sum_z w_z \kappa(z)$, so $\mu(A \backslash G') \leq \sum_z w_z - \sum_{z \in Z_*} w_z \leq \Delta$. Place each set Ω_z with $z \in Z_*$ in its own cube, at the positions given by the labels of its states. For $|Z_*| = M' \leq M$, these cubes have $M' 2^m - |G'| \leq |S| - |G'|$ empty positions, which we fill with states outside G'. The remaining $(M - M')2^m$ states form the other cubes in any way, which defines $\tilde{I}$. At $x \in G'$ the readers agree with $\tilde{I}$. If $y = T_{i,a} x \in G$, $Wy = Wx$ and $r(y) = r(x)[i \mapsto a]$, then y is the state of Ω_{Wx} with this label, which is the position of $\tilde{T}_{i,a} x$. ■

The first choice of κ is the fraction of labels that occur in a fiber, $\kappa(z) = |r(W^{-1}(z))|/2^m$. The fibers are disjoint and each has at least as many states as distinct labels, so these fractions sum to at most M over all z. A uniformly random label occurs in the fiber $W^{-1}(z)$ with probability $\kappa(z)$, so Δ is the probability that $x \sim \mu$ lies in A and a uniformly random label occurs nowhere in the fiber $W^{-1}(Wx)$. This fiber is a subset of S and may contain states outside G.

Figure 5 shows how the sum bounds the loss. There the second fiber misses a quarter of the labels, the third misses three quarters, and the second carries at least as much probability as the third. So the average of $1 - \kappa$ over the fibers, which counts a quarter of the probability of the second fiber and three quarters of that of the third, is at least the probability of the third. This is the part that the composable interface drops. Appendix D uses another choice of κ.

Reversible words. Let $W = E_t \cdots E_1$ be a word with $E_s = T_{i_s, a_s}$, put $U_s = E_s \cdots E_1$ and $U_0 = \mathrm{Id}$, and let $x_s = U_s x$ be the path of x. The *reversible set* G_W is the set of states x whose path passes (3) at every step $s = 1, \ldots, t$.

Lemma 12 (Reversible set) *The map $x \mapsto (r(x), Wx)$ is injective on G_W.*

Proof Let $x \in G_W$. The first equalities of (3) determine each label $r(x_s)$, and so each old value b_s, from $r(x)$ and the word. Starting at Wx, the setters $T_{i_t, b_t}, \ldots, T_{i_1, b_1}$ undo the steps in reverse order by the second equalities and end at x. ■

Lemma 13 (Inserted setters) *Let $W = E_t \cdots E_1$ write every variable, and let x be a state with $x_j \notin \mathcal{N}$ for $0 \le j < t$. Then $x \in G_W$, and every setter $T_{i,a}$ satisfies $T_{i,a}x \in G_W$, $WT_{i,a}x = Wx$ and $r(T_{i,a}x) = r(x)[i \mapsto a]$.*

Proof A state $z \notin \mathcal{N}$ and a setter $T_{k,e}$ pass (3): success and preservation give the label, and with $v = r_k(z)$, write-back handles $e = v$, while overwrite and write-back give $T_{k,v}T_{k,e}z = T_{k,v}z = z$ for $e \neq v$. So $x \in G_W$. Fix $T = T_{i,a}$ and let τ be the first position with $i_\tau = i$. The label claim is success and preservation at x. Let $y_s = U_s T x$.

For $j < \tau$, let $z = x_{j-1}$, $z' = x_j$ and $E_j = T_{k,e}$ with $k \neq i$. Both z and z' lie outside $\mathcal{N}$, as $j < t$. Commutation at z gives $E_j T z = T z'$, so $y_j = T x_j$ by induction. Success and preservation at z and at z' give $r(Tz) = r(z)[i \mapsto a]$ and $r(Tz') = r(z')[i \mapsto a] = r(Tz)[k \mapsto e]$, which is the first equality of (3) for this step of the path of Tx. For the second, let $v = r_k(Tz) = r_k(z)$. If $e = v$, write-back at z gives $z' = z$, so $T_{k,v}Tz' = E_j Tz = Tz' = Tz$. If $e \neq v$, commutation at z', overwrite at z and write-back at z give $T_{k,v}Tz' = TT_{k,v}T_{k,e}z = TT_{k,v}z = Tz$.

At $j = \tau$, let $z = x_{\tau-1}$ and $E_\tau = T_{i,e}$. Overwrite at z gives $E_\tau Tz = T_{i,e}z = x_\tau$, so the two paths meet, and $T_{i,a}x_\tau = T_{i,a}T_{i,e}z = T_{i,a}z = Tz$, where $a = r_i(Tz)$ by success. Success and preservation at z give $r(x_\tau) = r(Tz)[i \mapsto e]$. So this step passes (3). From step τ on, the path of Tx is that of x, and $x \in G_W$. Hence $Tx \in G_W$ and $WTx = Wx$. ■

For a word W and $A \subseteq S$, let Δ_A be the probability that $x \sim \mu$ lies in A and a uniformly random label, independent of x, occurs nowhere in the fiber $W^{-1}(Wx)$. This is the Δ of Lemma 11 for the first choice of κ.

Lemma 14 (Bounds from a reversible set) *Let W be a word, let $A \subseteq G_W$, and suppose every $x \in A$ and every setter $T_{i,a}$ satisfy $T_{i,a}x \in G_W$, $WT_{i,a}x = Wx$ and $r(T_{i,a}x) = r(x)[i \mapsto a]$. Then $D_\mu^\infty \le \mu(S \setminus A)$, with no condition on $|S|$. If $2^m \mid |S|$, then $D_\mu \le \mu(S \setminus A) + \Delta_A$. In both cases one composable interface agrees with all readers and setters at once outside an event of at most this probability.*

Proof For D_μ^∞: by Lemma 12, the states of G_W in one fiber have distinct labels. Additional states complete each such set, and each state of $S \setminus G_W$, to a cube of its own, with the states of S at the positions of their labels. Coordinate setters on these cubes define a composable interface on a finite set containing S, whose readers agree with r on S. For $x \in A$, $T_{i,a}x$ is the state of the cube of x with the label $r(x)[i \mapsto a]$, which is the output of the composable setter. For D_μ: apply Lemma 11 with $G = G_W$ and the first choice of κ, by Lemma 12. Its composable interface agrees with I on $A \cap G'$. ■

Appendix C. Proof of Theorem 4(a) and (b)

Throughout, $\hat{a}$ is a target label, $W_{\hat{a}} = T_{m,\hat{a}_m} \cdots T_{1,\hat{a}_1}$, $U_j = T_{j,\hat{a}_j} \cdots T_{1,\hat{a}_1}$ and $U_0 = \mathrm{Id}$. A *checked step* from a state z under $T_{j,\hat{a}_j}$ calls $z' = T_{j,\hat{a}_j}z$ and checks $r(z') = r(z)[j \mapsto \hat{a}_j]$. With $b = r_j(z)$, it then compares z' with z if $b = \hat{a}_j$, and otherwise calls $T_{j,b}z'$ and compares the output with z. Each checked step makes at most two calls, and a passed one satisfies (3). A round stops at its first failed check. On the path of y_i, step i compares its output with the state $x_i = U_i x$ of the path of x before it calls the undo. Part (c) is proved in Appendix E.7.

C.1. Path and flip checks

Lemma 15 (Checked setters) *Let x pass the path check and the flip checks of a round with target label $\hat{a}$. Then $x \in G_{W_{\hat{a}}}$, and every setter $T_{i,e}$ satisfies $T_{i,e}x \in G_{W_{\hat{a}}}$, $W_{\hat{a}}T_{i,e}x = W_{\hat{a}}x$ and $r(T_{i,e}x) = r(x)[i \mapsto e]$.*

Proof The path check gives $x \in G_{W_{\hat{a}}}$. If $e = r_i(x)$, the check of write-back gives $T_{i,e}x = x$, and the claims follow. Otherwise $T_{i,e}x = y_i$, and the check of its label gives the third claim. The path of y_i satisfies (3) at steps $1, \dots, i$ and reaches the same state as the path of x at step i. From step i on, the two paths coincide, so the path of y_i satisfies (3) at the later steps as well. Hence $y_i \in G_{W_{\hat{a}}}$ and $W_{\hat{a}}y_i = W_{\hat{a}}x$. ■

Lemma 16 (Depth of the path and flip checks) *Let $\varrho = |r(x) \oplus \hat{a}|$. The calls of the path and flip checks on x have depth at most $\min\{\varrho + 2, m + 1\}$ when they are made, and there are at most $m^2 + 5m$ of them.*

Proof Let ϱ_j be the number of $h \le j$ with $r_h(x) \neq \hat{a}_h$. On the path of x, a passed step that writes the old value is a self-loop that the tester has seen, so x_j has depth at most ϱ_j when step $j + 1$ is made. Hence the calls of the path check have depth at most $\varrho + 1 \le m + 1$. Let $j < i$. Before step j of the path of y_i, the flip and the earlier passed steps give a chain of length at most $1 + \varrho_{j-1}$, so this step and its undo have depth at most $\varrho_j + 2$. As $\varrho_j \le \min\{\varrho, m - 1\}$, this is at most $\min\{\varrho + 2, m + 1\}$. Step i itself has depth at most $\varrho_{i-1} + 2$. It is followed by a comparison with x_i, so its undo acts on a state of depth at most ϱ_i and has depth at most $\varrho_i + 1$. These bounds use only comparisons that passed before the call. The path check makes at most $2m$ calls, the write-backs and flips $2m$, and the paths of the states y_i at most $\sum_i 2i = m(m + 1)$, in total $m^2 + 5m$. ■

C.2. The bridge check

On the same states, the M heaviest fibers lose at most the probability that x passes the checks above and a uniform label u is missing from its fiber (Lemma 14). The *bridge check* pays for this loss by showing that a state of label u lies in the fiber of x. It runs after the path and flip checks of Algorithm 1, as follows.

1. Draw a uniform label u, and let $u^{(j)} = (\hat{a}_1, \dots, \hat{a}_j, u_{j+1}, \dots, u_m)$ for $0 \le j \le m$.

Algorithm 1: One audit round with target label $\hat{a}$ and depth cap d. A failed check rejects.

draw $x \sim \mu$; put $x_0 = x$
if $d \le m$ *and* $|r(x) \oplus \hat{a}| > d - 2$ **then** accept the round `// first test`
for $s = 1, \dots, m$ **do** `// path check`
  put $b_s = r_s(x_{s-1})$; call $x_s = T_{s,\hat{a}_s}x_{s-1}$ and check its label; if $\hat{a}_s = b_s$, compare x_s with x_{s-1}; otherwise call the undo $T_{s,b_s}x_s$ and compare its output with x_{s-1}
for $i = 1, \dots, m$ **do** `// flip check`
  check $T_{i,r_i(x)}x = x$; put $y_i = F_i x$; check $r(y_i) = r(x)[i \mapsto 1 - r_i(x)]$
  for $s = 1, \dots, i$ **do**
    do step s of the path of y_i as above; if $s = i$, compare its output with x_i before any undo

2. If some $|u^{(j)} \oplus r(x)| > d-1$, accept the round (the *second test*).
3. For $j = 0, \ldots, m-1$: from x, write the values of $u^{(j)}$ where it differs from $r(x)$, to get w_j, and check $r(w_j) = u^{(j)}$.
4. Put $w_m = W_{\hat{a}}x$, and check $T_{j,\hat{a}_j} w_{j-1} = w_j$ for $j = 1, \ldots, m$.

Applying $W_{\hat{a}}$ to w_0 directly could take depth $2m$. The bridge check instead reaches each intermediate state w_j from x by a word of length at most m and checks one setter between neighbors, so its depth is at most $m+1$.

Lemma 17 (Bridge check) *If the bridge check of a round passes, then w_0 reads u and lies in the fiber of x under $W_{\hat{a}}$. The bridge check makes at most m^2+m calls, and with a depth cap d, its calls have depth at most d.*

Proof The label check gives $r(w_0) = u^{(0)} = u$, and the comparisons give $W_{\hat{a}}w_0 = w_m = W_{\hat{a}}x$. Each w_j with $j < m$ takes at most m calls, and there are m comparisons with a call each. The word that produces w_j has length $|u^{(j)} \oplus r(x)| \le d-1$, so its calls have depth at most $d-1$, whatever their outputs are, and a bridge call on w_{j-1} has depth at most d. For a depth cap $d = m+1$ the length is at most $m = d-1$ for every u and $\hat{a}$. ■

C.3. Proof of parts (a) and (b)

Proof of Theorem 4(a) and (b) *One-sided error.* In a composable interface every path satisfies (3), writing the current value fixes the state, and a flip changes only coordinate i. The paths of x and y_i reach the same state at step i, the first write to i. Each w_j is the state of the cube of x with the label $u^{(j)}$, and $T_{j,\hat{a}_j}$ takes $u^{(j-1)}$ to $u^{(j)}$. So every check passes, and a round that stops early accepts.

Depth and calls. Lemma 16 bounds the depth by $m+1$ without the first test of Algorithm 1, and by d with it. Lemma 17 bounds the bridge check. So a round has depth at most d, with at most m^2+5m calls without the bridge check and at most $2m^2+6m$ with it.

Part (a). Let $d = m+1$, fix $\hat{a}$, and let A be the set of states that pass the path and flip checks. By Lemma 15, Lemma 14 applies with $W = W_{\hat{a}}$. So $D_\mu^\infty \le \mu(S \setminus A)$, which is the rejection probability of a round without the bridge check. For D_μ, the label u is uniform and independent of x. If $x \in A$ and u occurs nowhere in the fiber of x, the bridge check rejects, by Lemma 17. So a round rejects with probability at least $\mu(S \setminus A) + \Delta_A \ge D_\mu$. If the distance is at least ϵ, then $\lceil \log(1/\delta)/\epsilon \rceil$ independent rounds all pass with probability at most δ.

Part (b). Let $2 \le d \le m$; for $d = m+1$ the bounds follow from part (a). For a fixed target label $\hat{a}$, let $A(\hat{a})$ be the set of states that pass the first test of Algorithm 1 and the path and flip checks. By Lemma 14, $D_\mu^\infty \le \mu(S \setminus A(\hat{a}))$. Averaged over the uniform target label, the probability that the first test ends the round is τ_2, as $r(x) \oplus \hat{a}$ is uniform. So $D_\mu^\infty \le p + \tau_2$. For D_μ, Lemma 14 gives $D_\mu \le \mu(S \setminus A(\hat{a})) + \Delta_{A(\hat{a})}$. If $x \in A(\hat{a})$ and u occurs nowhere in the fiber of x, then either the second test ends the round or the bridge check rejects, by Lemma 17. For fixed x and j, the label $u^{(j)}$ is uniform, as $\hat{a}$ and u are, so the second test ends the round with probability at most $(m+1)\tau_1$. Rejections by the two kinds of checks are disjoint events, so $D_\mu \le p + \tau_2 + (m+1)\tau_1 \le p + (m+2)\tau_2$. If the tail is at most $\epsilon/2$ and the distance at least ϵ, then $p \ge \epsilon/2$, and $\lceil 2\log(1/\delta)/\epsilon \rceil$ independent rounds all pass with probability at most δ. ■

C.4. Lower bounds for every tester

Proposition 18 (Every tester) *Let $0 < \epsilon \le 1/16$, $0 < \delta \le 1/3$ and $N_{\delta,\epsilon} = \lceil \log(1/\delta)/(2\epsilon) \rceil + 2$. Let a tester, whenever μ is the uniform measure on a set of $2^m N_{\delta,\epsilon}$ states, accept composable interfaces and reject interfaces with $D_\mu^\infty > \epsilon$, each with probability at least $1-\delta$. Then it uses, in expectation, at least $\log(1/\delta)/(48\epsilon)$ samples and at least $(m/(48\epsilon))\log(1/\delta)$ setter calls on some composable interface on such a set, even if reader calls and comparisons are free.*

Proof of Proposition 18 *Instances.* Let $N_0 = \lfloor 2\epsilon N_{\delta,\epsilon} \rfloor + 1$ and $\tilde\epsilon = N_0/N_{\delta,\epsilon}$. As $\log(1/\delta) \ge \log 3 > 1$ and $N_{\delta,\epsilon} \ge \log(1/\delta)/(2\epsilon)$, $2\epsilon < \tilde\epsilon \le 2\epsilon + 1/N_{\delta,\epsilon} \le 4\epsilon \le 1/4$ and $1 \le N_0 \le N_{\delta,\epsilon}$. The set of states is $\{0,1\}^m \times [N_{\delta,\epsilon}]$ with cubes $\{0,1\}^m \times \{z\}$, the measure ν is uniform, and $r_i(u,z) = u_i$. The null input has $T_{i,a}(u,z) = (u[i \mapsto a], z)$. A negative input, indexed by a variable j_* and an N_0-set $\mathcal{Z}_0 \subseteq [N_{\delta,\epsilon}]$, makes both setters of j_* the identity on the cubes with $z \in \mathcal{Z}_0$. On each cube the random-edit kernel averages kernels that resample one coordinate of u uniformly, one of them possibly replaced by the identity, so ν is stationary for every input. Identity setters satisfy write-back, overwrite and preservation and commute with setters that keep the cube, so only success fails, on the states of the cubes with $z \in \mathcal{Z}_0$ and $u_{j_*} = 1-a$. At each state of such a cube, a composable interface, on the set of states or on any finite set containing it, moves the state by exactly one of its two setters of variable j_*, while both original setters fix the state. So the two disagreements sum to at least $\tilde\epsilon$. As a composable interface on the same states, the null input differs from the negative input only in the two setters of j_*, each with probability $\tilde\epsilon/2$, so $D_\nu^\infty = D_\nu = \tilde\epsilon/2 \in (\epsilon, 2\epsilon]$.

Avoidance probability. For an integer $0 \le v \le N_{\delta,\epsilon}$, let $\mathrm{Av}(v)$ be the probability that a uniformly random N_0-subset of $[N_{\delta,\epsilon}]$ misses a given v-subset, and let $\overline{\mathrm{Av}}$ be its linear interpolation on $[0, N_{\delta,\epsilon}]$, extended by zero. The decrement $\mathrm{Av}(v) - \mathrm{Av}(v+1) \ge 0$ is the probability of missing v given points and containing another one, which does not increase in v. The slope 0 of the zero extension is at least the last slope of the interpolation. So $\overline{\mathrm{Av}}$ is nonincreasing and convex. For integers $0 \le v \le \lfloor N_{\delta,\epsilon}/2 \rfloor$, each factor of $\mathrm{Av}(v) = \prod_{v'<v}(1 - N_0/(N_{\delta,\epsilon} - v'))$ satisfies $N_0/(N_{\delta,\epsilon} - v') \le 2\tilde\epsilon \le 1/2$, so $\mathrm{Av}(v) \ge (1-2\tilde\epsilon)^v \ge \mathrm{e}^{-4\tilde\epsilon v}$, as $-\log(1-2\tilde\epsilon) \le 4\tilde\epsilon$ for $2\tilde\epsilon \le 1/2$. By convexity of the exponential, $\overline{\mathrm{Av}}(v) \ge \mathrm{e}^{-4\tilde\epsilon v}$ also for real $0 \le v \le \lfloor N_{\delta,\epsilon}/2 \rfloor$.

Coupling. Draw j_* uniformly from $[m]$ and $\mathcal{Z}_0$ uniformly among the N_0-subsets of $[N_{\delta,\epsilon}]$, independently of the tester's randomness and of the samples, and execute the tester on the null input; $\Pr_0$ and $\mathbb{E}_0$ refer to this execution. Let n_T and n_S be its numbers of setter calls and samples. Let $\mathcal{V}_i$ be the set of cubes in which it calls a setter of variable i, so $\sum_i |\mathcal{V}_i| \le n_T$, and let $\mathcal{V}_0$ be the set of cubes that contain a sample, so $|\mathcal{V}_0| \le n_S$. Couple the execution on the negative input $(j_*, \mathcal{Z}_0)$, using the same randomness and samples. Setters keep the cube of their argument, so every held state lies in a cube of $\mathcal{V}_0$. If $\mathcal{Z}_0$ misses $\mathcal{V}_{j_*}$, or if it misses $\mathcal{V}_0$, then every setter call has the same output on both inputs, and, inductively, the transcripts and decisions coincide, as readers and comparisons agree on equal states. Conditional on the null execution, $\mathcal{Z}_0$ misses $\mathcal{V}_{j_*}$ with probability $m^{-1}\sum_i \mathrm{Av}(|\mathcal{V}_i|) \ge \overline{\mathrm{Av}}(m^{-1}\sum_i |\mathcal{V}_i|) \ge \overline{\mathrm{Av}}(n_T/m)$, by convexity and monotonicity of $\overline{\mathrm{Av}}$, and it misses $\mathcal{V}_0$ with probability $\mathrm{Av}(|\mathcal{V}_0|) \ge \overline{\mathrm{Av}}(n_S)$. Let acc be the event that the null input is accepted, so $\Pr_0(\mathrm{acc}) \ge 1-\delta$. Soundness on every negative input, averaged over $(j_*, \mathcal{Z}_0)$, gives $\delta \ge \mathbb{E}_0[\mathbf{1}_{\mathrm{acc}}\,\overline{\mathrm{Av}}(n_T/m)] \ge \Pr_0(\mathrm{acc})\,\overline{\mathrm{Av}}(\bar n_T)$ for $\bar n_T = \mathbb{E}_0[n_T \mid \mathrm{acc}]/m$, by conditional Jensen, and in the same way $\delta \ge \Pr_0(\mathrm{acc})\,\overline{\mathrm{Av}}(\bar n_S)$ for $\bar n_S = \mathbb{E}_0[n_S \mid \mathrm{acc}]$. Both hold trivially if the expectations are infinite.

Extracting $\log(1/\delta)$. We claim $\bar{n}_T \geq \log((1-\delta)/\delta)/(4\tilde{\epsilon})$. As $1-\delta < 1$ and $\tilde{\epsilon} > 2\epsilon$, the right side is at most $\log(1/\delta)/(8\epsilon) \leq (N_{\delta,\epsilon}-1)/2 \leq \lfloor N_{\delta,\epsilon}/2 \rfloor$. So the claim holds if $\bar{n}_T > \lfloor N_{\delta,\epsilon}/2 \rfloor$, and otherwise it follows from $\delta \geq (1-\delta)\mathrm{e}^{-4\tilde{\epsilon}\bar{n}_T}$. The same argument gives $\bar{n}_S \geq \log((1-\delta)/\delta)/(4\tilde{\epsilon})$. As $1-\delta \geq 2/3$, $\tilde{\epsilon} \leq 4\epsilon$ and $\log((1-\delta)/\delta) \geq \log(1/\delta)/2$ by $(1-\delta)^2 \geq \delta$, we get $\mathbb{E}_0 n_T \geq \Pr_0(\mathrm{acc})\, m\bar{n}_T \geq (m/(48\epsilon))\log(1/\delta)$ and $\mathbb{E}_0 n_S \geq \Pr_0(\mathrm{acc})\,\bar{n}_S \geq \log(1/\delta)/(48\epsilon)$. ■

Appendix D. Proof of the sufficiency in Theorem 2

Throughout Appendices D.1–D.4, $|S| = M2^m$, $\lambda \geq 1$, μ is a λ-subinvariant probability measure on S and $\gamma = \gamma_\mu$. Let $\mathcal{U} \subseteq \mathcal{N}$ be the set of states at which success, write-back, overwrite or preservation fails. For $v \geq m$, put $\Phi_v(\lambda) = h_v^{-1}\sum_{k<v}\lambda^k$, with h_v as in Lemma 7. For a composable interface $\tilde{I}$, let $\mathcal{E}_\mu(\tilde{I})$ be the probability that some reader or setter of I differs from that of $\tilde{I}$ at $x \sim \mu$, so $D_\mu \leq \mathcal{E}_\mu(\tilde{I})$.

Theorem 19 (Parametric form) *For $m \geq 1$ and all integers $t, t' \geq m$, some composable interface $\tilde{I}$ on S has*

$$\mathcal{E}_\mu(\tilde{I}) \leq \min\Big\{1,\ \Phi_t(\lambda)\gamma + \frac{\lambda^t}{h_t}\Phi_{t'}(\lambda)\mu(\mathcal{U})\Big\}.$$

The proof fixes a word, takes its fibers as candidate cubes, and applies Lemma 11. To pay for the labels missing from a fiber, it uses a function κ° defined once from the interface alone: the capacity function of Appendix D.2.

D.1. Retained steps

For a setter $T_{i,a}$, let $\Pi_{i,a}$ be the kernel with row zero or one defined as follows. If $a = r_i(x)$, it keeps the transition $x \to T_{i,a}x$ when $T_{i,a}x = x$. If $a \neq r_i(x)$, it keeps $x \to y = T_{i,a}x$ when $r(y) = r(x)[i \mapsto a]$ and $T_{i,r_i(x)}y = x$. Otherwise the row of x is zero. Put $\Pi_i = (\Pi_{i,0} + \Pi_{i,1})/2$ and $\Pi = m^{-1}\sum_i \Pi_i$. Each kept transition passes (3).

Each Π_i is symmetric: if $x \to y$ is a kept transition with $y \neq x$, then the labels of x and y differ exactly at i, and $y \to x$ is kept with the setter of the old value, so both entries are $1/2$. Entrywise, $\Pi_i \leq (T_{i,0} + T_{i,1})/2$ and $\Pi \leq P$. For $x \notin \mathcal{U}$, write-back keeps the setter of the current value, and success, preservation, overwrite and write-back keep the other, so $\Pi_i\mathbf{1}(x) = 1$ for all i. Hence

$$0 \leq \mathbf{1} - \Pi\mathbf{1} \leq \mathbf{1}_{\mathcal{U}}. \tag{4}$$

D.2. The capacity function

Draw t' independent uniform coordinates $i_1, \ldots, i_{t'}$, condition on the event that they cover $[m]$, and put $\Upsilon = \mathbb{E}[\Pi_{i_1}\cdots\Pi_{i_{t'}} \mid \text{cover}]$ and $\kappa^\circ = \Upsilon\mathbf{1}$. The product applies Π_{i_1} first. The transpose of a product reverses the word, and given coverage a uniform word has the same distribution as its reverse. So Υ is symmetric, and $0 \leq \kappa^\circ \leq 1$.

Fix a start state and a covering word of coordinates, and attach independent fair values to its steps. The label that these values write is uniform, as each coordinate takes its last value, and a path

that survives all kept transitions ends at this label. So $\Upsilon(x, r^{-1}(b)) \le 2^{-m}$ for every state x and label b. By symmetry,

$$\sum_{z:\, r(z)=b} \kappa^\circ(z) = \sum_{x \in S} \Upsilon(x, r^{-1}(b)) \le |S| 2^{-m} = M. \tag{5}$$

By (4) and $\Pi \le P$, $\mu(\mathbf{1} - \Pi^{t'}\mathbf{1}) = \sum_{k<t'} \mu\Pi^k(\mathbf{1} - \Pi\mathbf{1}) \le \sum_{k<t'} \mu P^k(\mathcal{U}) \le \sum_{k<t'} \lambda^k \mu(\mathcal{U})$. Now $\Pi^{t'}\mathbf{1}$ is the expectation of $\Pi_{i_1} \cdots \Pi_{i_{t'}}\mathbf{1}$ over the coordinates. Splitting it on coverage, and bounding the product by $\mathbf{1}$ when the coordinates do not cover $[m]$, gives $\Pi^{t'}\mathbf{1} \le h_{t'}\kappa^\circ + (1 - h_{t'})\mathbf{1}$, so $\mathbf{1} - \kappa^\circ \le h_{t'}^{-1}(\mathbf{1} - \Pi^{t'}\mathbf{1})$. Integrating against μ gives

$$\mu(\mathbf{1} - \kappa^\circ) \le \Phi_{t'}(\lambda)\mu(\mathcal{U}). \tag{6}$$

D.3. Proof of Theorem 19

Proof of Theorem 19 Fix a word $W = E_t \cdots E_1$ that writes every variable, with path $x_s = U_s x$, and let $A_W = \{x : x_j \notin \mathcal{N} \text{ for } 0 \le j < t\}$. By Lemma 13, each $x \in A_W$ lies in G_W, and each of its setters leads to a state of G_W in its fiber with the updated label. On G_W the map $x \mapsto (r(x), Wx)$ is injective (Lemma 12). Every $x \in G_W$ has the same end label, the label b_W that W writes. So the points z with $w_z = \mu(A_W \cap W^{-1}(z)) > 0$ are distinct states with label b_W, and (5) gives $\sum_{z:\, w_z > 0} \kappa^\circ(z) \le M$. Lemma 11 with $G = G_W$, $A = A_W$ and $\kappa = \kappa^\circ$ gives a composable interface that agrees with all readers and setters of I on $A_W \cap G'$, so

$$\mathcal{E}_\mu(\tilde{I}) \le \mu(S \setminus A_W) + \sum\nolimits_z w_z(1 - \kappa^\circ(z)). \tag{7}$$

Now draw the steps of W as t independent uniform setters conditioned on writing every variable, with $x \sim \mu$ independent of them. Before conditioning, x_k has distribution $\mu P^k \le \lambda^k \mu$, so

$$\mathbb{E}[\mu(S \setminus A_W) \mid \text{cover}] \le h_t^{-1} \sum\nolimits_{k<t} \mu P^k(\mathcal{N}) \le \Phi_t(\lambda)\gamma.$$

The last term of (7) is the expectation of $\mathbf{1}_{A_W}(x)(1 - \kappa^\circ(Wx))$ over x. Dropping $\mathbf{1}_{A_W}$ and the coverage from the numerator, and using (6),

$$\mathbb{E}\Big[\sum\nolimits_z w_z(1 - \kappa^\circ(z)) \,\Big|\, \text{cover}\Big] \le h_t^{-1}\mu P^t(\mathbf{1} - \kappa^\circ) \le \frac{\lambda^t}{h_t}\mu(\mathbf{1} - \kappa^\circ) \le \frac{\lambda^t}{h_t}\Phi_{t'}(\lambda)\mu(\mathcal{U}).$$

So some word attains at most the sum of these two bounds. ∎

D.4. The coefficients

Let $m \ge 2$, $t = t' = t_m$, so $h_t \ge 1/2$ (Lemma 7), and recall $\mu(\mathcal{U}) \le \gamma$.

The case $\lambda \le 1 + 1/(\alpha m)$. Such a μ is also $(1+1/(\alpha m))$-subinvariant, so take $\lambda = 1 + 1/(\alpha m)$, and put $\bar{\lambda} = \lambda^t \ge 1$. Then $\Phi_t(\lambda) = \alpha m(\bar{\lambda} - 1)/h_t$, and Theorem 19 gives

$$\mathcal{E}_\mu(\tilde{I}) \le \Phi_t(\lambda)\Big(1 + \frac{\bar{\lambda}}{h_t}\Big)\gamma \le 2\alpha m(\bar{\lambda} - 1)(1 + 2\bar{\lambda})\gamma \le 4\alpha m\,\bar{\lambda}^2\gamma, \qquad \bar{\lambda} \le \mathrm{e}^{t/(\alpha m)} \le (2\mathrm{e}m)^{1/\alpha}, \tag{8}$$

as $t \le m\log(2m) + 1$ and $1/m \le 1$. This is at most $K_\alpha\gamma$ with $K_\alpha = 4\alpha(2\mathrm{e})^{2/\alpha}m^{1+2/\alpha}$.

D.5. Law checks, for every set of states

Lemma 20 (Padding) *Let I be an interface on a finite set S, with no condition on $|S|$, and μ a measure on S. Add $(-|S|) \bmod 2^m$ states with arbitrary readings and setters that stay among the added states, and extend μ by zero to $\bar{\mu}$. The new interface $\bar{I}$ on $\bar{S}$ satisfies $\gamma_{\bar{\mu}}(\bar{I}) = \gamma_\mu(I)$, $\bar{\mu}$ is λ-subinvariant if μ is, and $D^\infty_\mu(I) \le D_{\bar{\mu}}(\bar{I})$. Extending a start distribution by zero extends its sample distributions by zero.*

Proof The set S is closed under the setters, and the laws at its states involve only states of S. Added states have mass zero. A composable interface on $\bar{S}$ is a candidate for $D^\infty_\mu(I)$ with the same disagreements under μ. Let $\bar{P}$ be the random-edit kernel of $\bar{I}$. As S is closed under the setters, $\bar{P}$ acts on S as P, so $\bar{\mu}\bar{P}$ is the extension of μP by zero. Hence $\mu P \le \lambda\mu$ gives $\bar{\mu}\bar{P} \le \lambda\bar{\mu}$: on S this is the hypothesis, and on the added states both sides vanish. By induction, if $\bar{\nu}$ extends a start distribution ν by zero, then $\bar{\nu}\bar{P}^s$ extends νP^s by zero for every $s \ge 0$. So, by (1), the sample distribution of $\bar{\nu}$ extends μ_ℓ by zero for every $\ell \ge 0$. ■

Corollary 21 (Long runs) *Let $m \ge 2$, $\alpha > 0$, $\ell \ge \alpha m$ and $K_\alpha = 4\alpha(2\mathrm{e})^{2/\alpha} m^{1+2/\alpha}$. Then $D_\ell \le K_\alpha \gamma_\ell$ when $2^m \mid |S|$, and $D^\infty_\ell \le K_\alpha \gamma_\ell$ for every finite S. So evaluating all law instances at $\lceil K_\alpha \log(1/\delta)/\epsilon \rceil$ samples of μ_ℓ certifies these distances, with depth 2 and $2m^2 + 2m$ setter calls per sample.*

Proof By Lemma 6(i), μ_ℓ is $(1 + 1/\ell)$-subinvariant, so (8) and Lemma 20 give the two bounds, and Appendix A gives the tester. ■

Corollary 21 and (8) are part (b) of Theorem 2.

D.6. The coefficient at mean αm

For $\alpha > 0$, put

$$C_\alpha = 24 \cdot 8^{1/(2\alpha)}\, \Gamma\big(1 + \tfrac{1}{2\alpha}\big) + \Big(\frac{32(1 + 1/(2\alpha))}{\mathrm{e}}\Big)^{1+1/(2\alpha)},$$

$b_\alpha = 1 - \mathrm{e}^{-1/(8\alpha+1)}$ and $c_\alpha = b_\alpha/(128\, C_\alpha\, 16^{1/(2\alpha)})$.

Proposition 22 (Coefficient at mean αm) *For $m \ge 32$ and each $\alpha > 0$, some interface and start distribution have $D^\infty_{\alpha m} \ge c_\alpha m^{1/(2\alpha)} \gamma_{\alpha m} > 0$.*

Proof Let $m \ge 32$ and take the unmasked twisted interface I^- of Appendix E.2 with $J = \lfloor m/8 \rfloor \ge 4$ and $N_c = 1$, which does not depend on α. Let $\ell = \alpha m$. By Lemma 27, $D^\infty_\ell \ge ((J + 1)/(128J^2))(1 - q_0^J)$ with $q_0 = 8\alpha J/(1 + 8\alpha J)$. As $\log(1/q_0) \ge 1/(1 + 8\alpha J)$, $q_0^J \le \mathrm{e}^{-1/(8\alpha+1)}$, so $D^\infty_\ell \ge b_\alpha/(128J)$. It remains to bound $\gamma_\ell = \mu_\ell(\mathcal{N})$, where $\mathcal{N}$ is given by Lemma 23 and $\gamma_\ell > 0$, as μ_ℓ has full support: for $\ell > 0$ a run reaches every label with positive probability, and by Lemma 24 its hidden bit is a fair bit independent of the final label.

Let the edits arrive at the times of a Poisson process of rate 1 and stop at an independent time ι of density $\ell^{-1}\mathrm{e}^{-t/\ell}$; this gives μ_ℓ (Appendix G.1). Given $\iota = t$, the numbers of edits of the variables are independent and Poisson of mean t/m, with fair written values. So each final bit is its initial bit

shifted by an independent Bernoulli variable of parameter $(1-y)/2$, where $y = \mathrm{e}^{-t/m}$, and $\mathrm{e}^{-\iota/m}$ has density $\alpha^{-1}y^{1/\alpha-1}$ on $(0,1)$. Given the start label and y, the two blocks are independent.

For a sum Z of N_1 independent Bernoulli variables, Hoeffding's inequality gives $\Pr(Z - \mathbb{E}Z \geq a) \leq \mathrm{e}^{-2a^2/N_1}$ for $a \geq 0$, and the same for the lower tail. If all parameters lie in $[1/4, 3/4]$, then $\Pr(Z = h) \leq \sqrt{2\pi/(3N_1)}$ for every integer h: Fourier inversion gives $\Pr(Z = h) \leq (2\pi)^{-1}\int_{-\pi}^{\pi} |\mathbb{E}\mathrm{e}^{\mathrm{i}\theta Z}|d\theta$, each factor has squared modulus $1 - 4p(1-p)\sin^2(\theta/2) \leq \mathrm{e}^{-(3/4)\sin^2(\theta/2)}$, $|\sin(\theta/2)| \geq |\theta|/\pi$ on $[-\pi, \pi]$, and a Gaussian integral finishes the bound.

Let $0 < y \leq 1/2$. The parameters lie in $[1/4, 3/4]$, so a block of $4J$ bits lies in $\{2J-1, 2J\}$ with probability at most $2/\sqrt{J}$. A start label has a frozen block, of initial weight J or $3J$, whose final weight Z has mean $2J \mp Jy$. Split its bits into two parts of $2J$, with sums Z_1 and Z_2, which are independent, and let $a = |h - \mathbb{E}Z|$. If $Z = h$, then Z_1 or Z_2 deviates from its mean by at least $a/2$ in the direction of $h - \mathbb{E}Z$. By Hoeffding's inequality with $N_1 = 2J$, Z_1 does so with probability at most $\mathrm{e}^{-2(a/2)^2/(2J)} = \mathrm{e}^{-a^2/(4J)}$. Given such a value of Z_1, the event $Z = h$ fixes Z_2, which by the point bound with $N_1 = 2J$ has probability at most $\sqrt{\pi/(3J)} \leq 2/\sqrt{J}$. The same holds with the two parts exchanged, and adding the two cases gives $\Pr(Z = h) \leq (4/\sqrt{J})\mathrm{e}^{-a^2/(4J)}$. For $h \in \{2J-1, 2J\}$, $a \geq (Jy-1)_+$, and $(x-1)_+^2 \geq x^2/2 - 1$ gives $\Pr(Z \in \{2J-1, 2J\}) \leq (12/\sqrt{J})\mathrm{e}^{-Jy^2/8}$, as $\mathrm{e}^{1/(4J)} < 3/2$. With the other block, $\Pr(\mathcal{N} \mid y) \leq (24/J)\mathrm{e}^{-Jy^2/8}$. For $1/2 < y < 1$, the frozen block has to deviate from its mean by at least $J/2 - 1 \geq J/4$ in a fixed direction, so $\Pr(\mathcal{N} \mid y) \leq \mathrm{e}^{-J/32}$. Integrating over y, with $v = Jy^2/8$,

$$\gamma_\ell \leq \frac{24}{\alpha J}\int_0^\infty y^{1/\alpha-1}\mathrm{e}^{-Jy^2/8}\,dy + \mathrm{e}^{-J/32} = 24\cdot 8^{1/(2\alpha)}\,\Gamma\big(1+\tfrac{1}{2\alpha}\big)J^{-1-1/(2\alpha)} + \mathrm{e}^{-J/32}.$$

As $x^{e'}\mathrm{e}^{-x/32} \leq (32e'/\mathrm{e})^{e'}$ for $x, e' > 0$, taking $e' = 1 + 1/(2\alpha)$ gives $\gamma_\ell \leq C_\alpha J^{-1-1/(2\alpha)}$. So $D_\ell^\infty/\gamma_\ell \geq (b_\alpha/(128C_\alpha))J^{1/(2\alpha)} \geq c_\alpha m^{1/(2\alpha)}$, as $J \geq m/16$. ∎

Appendix E. The twisted rectangle

This appendix proves the necessity in Theorems 2 and 5, part (c) of Theorem 4, and the depth bound at fixed thresholds of Section 6. All of them use one family of interfaces.

E.1. A small instance

The smallest member of the family has $J = 1$. Take $m = 8$, blocks $B' = \{1,2,3,4\}$ and $B'' = \{5,6,7,8\}$, and states (v, z) with $v \in \{0,1\}^8$ and a hidden bit z. Readers read v, and v' and v'' are the restrictions of v to the blocks. A setter of B'' replaces its coordinate. A setter of B' replaces its coordinate, and it also toggles z when it moves $|v'|$ across the threshold, between weights 1 and 2, while $|v''| \geq 2$. Commutation between the blocks fails exactly where both weights lie in $\{1,2\}$, and all other law instances hold everywhere (Lemma 23).

Start at the label with ones at variables 1 and 5, so at weights $(|v'|, |v''|) = (1,1)$. The loop has four segments, with eight writes in all.

- Bottom: write 1 to variables 2 and 3. The weights go $(1,1) \to (2,1) \to (3,1)$, and z stays, as $|v''| = 1$.
- Right: write 1 to variables 6 and 7. The weights go $(3,1) \to (3,2) \to (3,3)$.
- Top: write 0 to variables 3 and 2. The weights go $(3,3) \to (2,3) \to (1,3)$, and the second step toggles z, as $|v''| = 3$.

- Left: write 0 to variables 7 and 6. The weights go $(1,3) \to (1,2) \to (1,1)$.

So the loop returns to its start label with z toggled, and a composable interface that kept all eight writes would return to the start state (Lemma 25). The start distribution of Appendix E.2 draws the arguments of such loops, with random choices of the variables, so each of the 16 setters of the blocks carries at most $8J/(J+1) = 4$ times the start distribution. Every loop has a disagreement, so $1 \le 4 \cdot 16 \cdot R$ for the disagreement R of any candidate, and at $\ell = 0$ the distance is at least $1/64$ (Lemma 27). For $J = 1$ some start states meet the set where commutation fails, and for $J \ge 2$ the frozen block keeps all of them outside it.

E.2. The instances

Let $J, N_g, N_c \ge 1$ be integers and $n = 8J$, with $nN_g \le m$. The states are the triples $(\sigma, v, z) \in [N_c] \times \{0,1\}^m \times \{0,1\}^{N_g}$, so there are $N_c 2^{m+N_g}$ of them, and the readers read v. A *hidden assignment* ω gives each copy σ a *mask* $e_\sigma \in \{0,1\}^m$ and N_g disjoint *groups* $V_{\sigma,g} \subseteq [m]$ of n variables each, and it splits each group into two *blocks* $B'_{\sigma,g}$ and $B''_{\sigma,g}$ of $4J$ variables each. The other variables of the copy are *passive*. The *relative label* of a state is $u = v \oplus e_\sigma$, with restrictions u'_g and u''_g to the blocks of group g, and $f_{g,1} = \mathbf{1}\{|u'_g| \ge 2J\}$ and $f_{g,2} = \mathbf{1}\{|u''_g| \ge 2J\}$ are functions of σ and v. The *unmasked* assignment has $N_g = 1$, $e_\sigma = 0^m$, $B'_{\sigma,1} = [4J]$ and $B''_{\sigma,1} = [8J] \setminus [4J]$ for all σ, and then we drop the index g.

The setters of I^+ replace one coordinate of v and keep σ and z, so I^+ is composable. In the *twisted interface* I^-_ω, a setter of a variable $i \in B'_{\sigma,g}$ maps (σ, v, z) to

$$\big(\sigma,\ v[i \mapsto a],\ z \oplus [f_{g,1}(\sigma, v[i \mapsto a]) \oplus f_{g,1}(\sigma, v)]\, f_{g,2}(\sigma, v)\, \mathbf{e}_g\big), \tag{9}$$

with $\mathbf{e}_g$ the g-th unit vector, and the other setters are those of I^+.

The start distribution ν_ω draws σ uniformly and then, independently for each group, one of four *sides*, each with probability $1/4$: $|u'_g|$ uniform on $\{J, \dots, 3J-1\}$ and $|u''_g| = J$; $|u''_g|$ uniform on $\{J, \dots, 3J-1\}$ and $|u'_g| = 3J$; $|u'_g|$ uniform on $\{J+1, \dots, 3J\}$ and $|u''_g| = 3J$; $|u''_g|$ uniform on $\{J+1, \dots, 3J\}$ and $|u'_g| = J$. Given the weights, the blocks are uniform and independent. The relative labels of the passive variables and z are fair bits, independent of the rest. The *frozen block* B^{f}_g of a start state in group g is the block whose weight the side fixes at J or $3J$; the side also fixes the value of its threshold function.

Lemma 23 (Laws) *In I^-_ω, success, write-back, overwrite and preservation hold everywhere, and $\mathcal{N}$ is exactly the set of states with $|u'_g|, |u''_g| \in \{2J-1, 2J\}$ for some group g.*

Proof Every setter changes the label as written, and writing the current value changes neither v nor any threshold function. Between two writes to $B'_{\sigma,g}$ the value of $f_{g,2}$ does not change, so the changes of z telescope, and overwrite and commutation within this block hold. A setter of a group reads and changes only the labels of this group and the bit z_g, so setters of different groups commute, and so do passive setters with all others. Writes to $B''_{\sigma,g}$ commute with each other. Commutation of a variable i of $B'_{\sigma,g}$ with a variable j of $B''_{\sigma,g}$ fails in the bit z_g by $[f_{g,1}(v[i \mapsto a]) \oplus f_{g,1}(v)][f_{g,2}(v[j \mapsto b]) \oplus f_{g,2}(v)]$. This vanishes unless both weights lie in $\{2J-1, 2J\}$, and at those weights suitable writes cross both thresholds. ∎

Fix a copy σ and a group g. The *rest* of a state is its label outside $V_{\sigma,g}$ with its hidden bits other than z_g. Let ν_g be the distribution of the group's label and z_g under ν_ω given σ, and let P_g apply a uniformly random setter of the group.

Lemma 24 (Common sample distribution) *Let $\ell \geq 0$, $\ell_0 = n\ell/m$ and $q_0 = \ell_0/(1+\ell_0)$. The interfaces I^+ and I^-_ω have the same sample distribution μ_ℓ from ν_ω, and at a sample the hidden vector is a fair vector independent of the label and of the edits of the run. The number L_g of edits of group g in a run satisfies $\Pr(L_g \geq j) = q_0^j$. Given $L_g = j$, these edits are independent and uniform on the setters of the group, and with $\Lambda_{\sigma,g,j}$ the conditional distribution of the final rest,*

$$\mu_\ell(\cdot \mid \sigma) = (1-q_0) \sum\nolimits_{j \geq 0} q_0^j \, (\nu_g P_g^j) \otimes \Lambda_{\sigma,g,j}. \tag{10}$$

Proof Each edit lies in the group with probability n/m, so for $0 \leq y \leq 1$

$$\mathbb{E} y^{L_g} = \sum\nolimits_{j \geq 0} (1-q) q^j (1 - n/m + (n/m)y)^j = \frac{1-q}{1 - q(1 - n/m + (n/m)y)} = \frac{1}{1 + \ell_0 - \ell_0 y}.$$

So L_g is geometric with mean ℓ_0. A random edit can first decide whether it lies in the group and then choose its variable and value, and the choices inside the group are independent of these decisions, of the length of the run and of the choices outside the group. The setters of the group act only on its labels and z_g, the other setters act only on the rest, and the start state draws these two parts independently given σ. This gives (10) for each of the two interfaces. They move labels in the same way, and the hidden vector is its fair initial value plus an amount determined by the labels and the edits, so it stays a fair vector independent of the labels and the edits. So the two sample distributions are equal. ■

E.3. The distance

Lemma 25 (Forcing walk) *Let $w_0, \ldots, w_{n_1}$ satisfy $w_t = T_{i_t,a_t} w_{t-1}$, let every variable that occurs be written exactly twice, first with a value a and then with $1-a$, let each first write change the state, and let $w_{n_1} \neq w_0$. Then every composable interface $\tilde{I}$ on a finite set $S' \supseteq S$ disagrees with one of these calls.*

Proof If $\tilde{I}$ kept all of them, it would follow the same states. Before the first write to a variable i, the reading of $\tilde{I}$ at i is that of w_0. As the first write, of a, changes the state, this reading is $1-a$, and the last write restores it. Variables that do not occur never change, and neither does the hidden data of $\tilde{I}$. So $\tilde{I}$ would return to w_0, a contradiction. The argument does not use the readers of I. ■

Lemma 26 (One group) *Let $\ell \geq 0$, fix ω, a copy σ and a group g, and let $\tilde{I}$ be a composable interface on a finite set $S' \supseteq S$, with $\mathcal{D}_{i,a}$ the set of states of S at which $T_{i,a}$ of I^-_ω and $\tilde{T}_{i,a}$ differ. Then*

$$1 - q_0^J \leq \frac{8J}{J+1} \sum\nolimits_{i \in V_{\sigma,g},\, a} \mu_\ell(\mathcal{D}_{i,a} \mid \sigma). \tag{11}$$

Proof *A random loop.* Fix the rest of the state, and draw z_g uniformly. Let π', π'' be independent uniform permutations of $B'_{\sigma,g}$ and $B''_{\sigma,g}$, and let X_t, Y_t be the relative block labels with ones at the first t elements. From the relative label (X_J, Y_J), the word Q writes the relative value 1 to $\pi'(J+1), \ldots, \pi'(3J)$, then to $\pi''(J+1), \ldots, \pi''(3J)$, then the relative value 0 to $\pi'(3J), \ldots, \pi'(J+1)$, and then to $\pi''(3J), \ldots, \pi''(J+1)$. A relative value $\bar{a}$ at i is the setter $T_{i,\bar{a} \oplus e_{\sigma,i}}$. During the first segment $f_{g,2} = 0$, the second and fourth segments do not change z_g, and during the third segment

$f_{g,2} = 1$ while $f_{g,1}$ falls from 1 to 0 once. So Q returns to its label with z_g toggled, and Lemma 25 applies to its $8J$ calls. Let $\mathrm{occ}_{i,a}(X)$ be the expected number of calls of $T_{i,a}$ on the walk whose argument lies in X. Then $1 \le \sum_{i \in V_{\sigma,g,a}} \mathrm{occ}_{i,a}(\mathcal{D}_{i,a})$.

Occupation. In the first segment, the call at weight t has the argument with relative label (X_t, Y_J). The other segments are alike. The current hidden bit is a fair bit independent of the permutations. So this argument has the distribution of ν_g given the first side and the weight t, an event of probability $1/(8J)$. Given X_t, the next variable written to the relative value 1 is uniform on the $4J - t \ge J + 1$ relative zeros of X_t, and the next one written to 0 in the third segment is uniform on the $t \ge J + 1$ relative ones. At a fixed state, (i, a) is one relative value. So, state by state, $\mathrm{occ}_{i,a} \le (8J/(J + 1))\, \nu_g$.

Noisy loops. Fix a word W of $j \le J - 1$ setters of the group, and let L_W be its action on labels: a coordinate that W does not write keeps its value, and a written one takes its last written value. Let $\hat{\mathsf{L}}_W$ apply L_W to the label and keep the hidden bit. As the hidden bit of ν_g is fair and independent of the label, and W adds to it an amount determined by the label, $\nu_g W = (\hat{\mathsf{L}}_W)_* \nu_g$. Apply L_W to all labels of the loop Q, skip its calls to variables that W writes, and start from a fresh fair hidden bit; call the result the *noisy loop*. By induction over the positions, the label at each position is the image under L_W of the label of Q there. Let $\mathrm{occ}^W_{i,a}$ be the occupation measure of the noisy loop. If W writes i, then $\mathrm{occ}^W_{i,a} = 0$. Otherwise the argument of the noisy loop has the distribution of $\hat{\mathsf{L}}_W$ applied to the argument of Q, so $\mathrm{occ}^W_{i,a} = (\hat{\mathsf{L}}_W)_* \mathrm{occ}_{i,a} \le (8J/(J + 1))(\hat{\mathsf{L}}_W)_* \nu_g = (8J/(J + 1))\, \nu_g W$.

As $j \le J - 1$, every low end of the loop has weight at most $J + j \le 2J - 1$ and every high end weight at least $3J - j \ge 2J + 1$. Every remaining write changes the relative label by one. So in the third segment the weight of the first block falls past $2J$ exactly once while $f_{g,2} = 1$, and the noisy loop returns to its label with the hidden bit toggled. Every first write of a variable changes the state. So Lemma 25 applies again.

Summing. Averaging over j independent uniform setters of the group and over the rest drawn from $\Lambda_{\sigma,g,j}$ gives, for $0 \le j < J$,

$$1 \le \frac{8J}{J + 1} \sum\nolimits_{i \in V_{\sigma,g,\,a}} \left[(\nu_g P_g^j) \otimes \Lambda_{\sigma,g,j}\right](\mathcal{D}_{i,a}).$$

Summing with coefficients $(1 - q_0)q_0^j$ and using (10) gives (11). ∎

Lemma 27 (Distance) *For every ω and $\ell \ge 0$,*

$$D^{\infty}_{\mu_\ell}(I^-_\omega) \ge \frac{N_g(J + 1)}{16mJ}(1 - q_0^J).$$

If $N_g = \lfloor m/n \rfloor$, this is at least $(1 - q_0^J)/(256J)$. If $N_g = 1$ and the group is the same in all copies, then $D^{\infty}_{\mu_\ell}(I^-_\omega) \ge ((J + 1)/(128J^2))(1 - q_0^J) \ge (1 - q_0^J)/(128J)$.

Proof Let $R = R_{\mu_\ell}(I^-_\omega, \tilde{I})$ for a candidate $\tilde{I}$. The groups of one copy are disjoint, so summing (11) over g gives $N_g(1 - q_0^J) \le (8J/(J + 1)) \sum_{i \in [m],a} \mu_\ell(\mathcal{D}_{i,a} \mid \sigma)$. The right side now sums over fixed setters, so we may average over the uniform copy σ, which gives $N_g(1 - q_0^J) \le (8J/(J + 1)) \cdot 2mR$. As $\lfloor m/n \rfloor \ge m/(16J)$, the second claim follows. A sum over one group alone could not be bounded by $16JR$ in this way, as the group varies with the copy. If it does not vary, averaging (11) over σ gives $1 - q_0^J \le (8J/(J + 1)) \cdot 16JR$. ∎

E.4. Renaming

Couple the executions of a tester on I^+ and I^-_ω by the same internal randomness and the same samples, drawn from the common distribution μ_ℓ together with the latent sides and edits. The hidden vector of a sample is a fresh fair vector, which is legal by Lemma 24.

Lemma 28 (Renaming) *Suppose that the samples drawn lie in distinct copies and that, in the execution on I^+, for each sample and each group g, one of $f_{g,1}, f_{g,2}$ is constant on the labels of all states held in the copy of this sample. Then the two transcripts coincide.*

Proof Fix a sample $w_* = (\sigma, v_*, z_*)$. In the execution on I^+ all states held in copy σ have hidden vector z_*. If $f_{g,1}$ is constant on their labels, put $b_g = 0$. Otherwise let b_g be the constant value of $f_{g,2}$. Put

$$\mathrm{ren}(\sigma, v, z_*) = \Big(\sigma, v, z_* \oplus \sum\nolimits_g b_g[f_{g,1}(\sigma, v) \oplus f_{g,1}(\sigma, v_*)]\, \mathbf{e}_g\Big).$$

Then ren fixes w_*, keeps labels and is injective. By (9), $T^-_{i,a}\mathrm{ren}(w) = \mathrm{ren}(T^+_{i,a}w)$ whenever w and $T^+_{i,a}w$ are held: for $i \in B'_{\sigma,g}$ the change of z in (9) is exactly that of ren, and the other setters change neither z nor any $f_{g,1}$. The union of these maps over the samples is injective, as states of distinct copies differ. By induction over the steps of the tester, each state it holds on I^-_ω is the image of the one it holds on I^+, so all readings, comparisons, depths and choices agree. ■

E.5. Unmasked: a ball around each sample

Lemma 29 (Unmasked transcripts) *Let ω be unmasked, let $d \geq 0$ and $\hat{t} \geq 1$ be integers with $J_{\hat{t}} = J(1-1/n)^{\hat{t}} > d$, and put $\hat{p} = q_0^{\hat{t}+1} + \exp[-(J_{\hat{t}} - d)^2/(2\hat{t})]$. The transcripts of a tester of depth at most d with at most s samples of μ_ℓ on I^+ and on I^-_ω differ in total variation by at most $s\hat{p} + s(s-1)/(2N_c)$, whatever its number of setter calls.*

Proof Fix a side, condition on the complete initial label of the frozen block, and let Z_j be the weight of this block after j edits of the group. A coordinate is untouched with probability $(1-1/n)^j$, and otherwise its last value is a fair bit, so $\mathbb{E}Z_j = 2J \mp J(1-1/n)^j$ on a side with initial weight J or $3J$. Changing one edit changes Z_j by at most two, so the bounded-difference inequality (van Handel, 2016) gives $\Pr(Z_j - \mathbb{E}Z_j \geq y) \leq \mathrm{e}^{-y^2/(2j)}$, and the same for the lower tail. Call a sample good if $Z_j \leq 2J - d - 1$ on the low side and $Z_j \geq 2J + d$ on the high side. Then the threshold function of the frozen block is constant on the Hamming ball of radius d around the label of the sample. For $1 \leq j \leq \hat{t}$, a failure needs a deviation of at least $J_{\hat{t}} - d$ in a fixed direction. For $j = 0$ the sample is good, as $J \geq J_{\hat{t}} > d$. So a sample is good except with probability $\hat{p}$. A held state on I^+ differs from its sample in at most d coordinates, as a write changes one coordinate. Lemma 28 finishes the proof, and two of s samples share a copy with probability at most $s(s-1)/(2N_c)$. ■

E.6. Masked: every query separately

For polynomial budgets the tester queries few labels, so a mask replaces the union over the whole ball by a union over the queries. In this subsection $N_g = \lfloor m/n \rfloor$, and ω is *uniform*: for each copy, independently, a uniform mask on $\{0,1\}^m$, uniform disjoint groups and uniform splits into blocks.

We let the tester learn the length of each run, which only helps it. Write each edit $T_{i,a}$ of a sample in copy σ as the pair $(i, \bar{a})$ with relative value $\bar{a} = a \oplus e_{\sigma,i}$.

The *latent data* $\mathcal{K}$ of a sample in a new copy are the groups and blocks of this copy, the sides, the initial relative label and the list of pairs of its edits. The tester sees the visible label $v^\star$ of the sample and the length of its run, and it holds the hidden vector only behind a handle. Hidden from it are $\mathcal{K}$ and the mask.

Lemma 30 (Independence) *Generate the copies of the first s potential samples in advance, condition on the event that they are distinct, and condition on the table of their run lengths.*

(i) The visible label of each sample is uniform and independent of its latent data, and its hidden vector is a fair vector independent of both. The data of different samples are independent.

(ii) The whole execution of the tester on I^+ is a function of the visible labels of the samples, the table and the internal randomness, and it does not depend on the mask or the latent data.

(iii) So, for each held state, its sample and the set $Y \subseteq [m]$ of coordinates where its label differs from the label of this sample are independent of the latent data of all samples. If the depth is at most d, then $|Y| \le d$.

Proof (i) Fix the length of a run in a new copy σ. The groups, the blocks, the sides and the initial relative label do not depend on e_σ. Given these and e_σ, the pairs are independent and uniform on $[m] \times \{0,1\}$, as a is a fair bit independent of i. So $\mathcal{K}$ is independent of e_σ. The relative end label $u^\star$ is a function of $\mathcal{K}$, as each coordinate keeps its initial relative value or takes the last relative value written to it. The visible label is $v^\star = e_\sigma \oplus u^\star$. So, for each v and each event E, $\Pr(v^\star = v, \mathcal{K} \in E) = \sum_y \Pr(\mathcal{K} \in E, u^\star = y) \Pr(e_\sigma = v \oplus y) = 2^{-m} \Pr(\mathcal{K} \in E)$: $v^\star$ is uniform and independent of $\mathcal{K}$. The hidden vector is a fair vector independent of $\mathcal{K}$ and e_σ by Lemma 24. Different samples lie in different copies, whose data are independent.

(ii) In a copy of I^+, a setter replaces one coordinate of the label and keeps the hidden vector. So the end points of paths and their equalities depend only on the root label and the setters called. By induction over the steps of the tester, each reading, each comparison and each choice is a function of the visible labels, the table and the internal randomness.

(iii) This follows from (i) and (ii). A write changes one coordinate of the label, so a state of depth at most d has $|Y| \le d$. ■

Lemma 31 (Masked transcripts) *Let $0 \le d < m/2$ and $\hat{t} \ge 0$ be integers, $\eta = 1 - 2d/m$, and*

$$\zeta = \exp[-J\eta^2/2] + \exp\big[-J^2\eta^2(1-1/n)^{2\hat{t}}/(8(4J+\hat{t}))\big]. \tag{12}$$

A tester of depth at most d with at most s samples of μ_ℓ and at most s_{h} samples and setter calls in total has transcripts on I^+ and on I^-_ω that differ in total variation by at most $sN_g q_0^{\hat{t}+1} + s_{\mathrm{h}} N_g \zeta + s(s-1)/(2N_c)$. This holds with free reader calls, comparisons and computation, and also if the tester learns the length of each run.

Proof *One query, one group.* Fix Y with $|Y| \le d$ and a group g, and condition on $L_g = j \le \hat{t}$. This keeps the law of the groups, as the number of edits of a set of n variables has the same distribution for each such set. So the frozen block B_g^{f} is a uniform $4J$-subset of $[m]$. Let $\mathrm{ov} = |Y \cap B_g^{\mathrm{f}}|$, so $\mathbb{E}\mathrm{ov} \le 4Jd/m = 2J(1-\eta)$. Reveal the $4J$ elements of B_g^{f} one at a time. Each step has

conditional range at most one, as swapping two elements changes ov by at most one, so the Azuma–Hoeffding inequality (van Handel, 2016) gives $\Pr(\mathrm{ov} - \mathbb{E}\mathrm{ov} \geq y) \leq \mathrm{e}^{-2y^2/(4J)}$. With $y = J\eta$, $\Pr(\mathrm{ov} > 2J(1-\eta/2)) \leq \mathrm{e}^{-J\eta^2/2}$, and otherwise $1 - \mathrm{ov}/(2J) \geq \eta/2$.

Let Z be the relative weight of B_g^{f} at the queried label, and let $\mathbb{E}Z$ be its conditional expectation given B_g^{f}, the side, Y and $L_g = j$. Given $L_g = j$, the edits of the group are independent and uniform on its pairs. Each initial bit has mean $1/4$ or $3/4$, a coordinate is untouched with probability $\theta_j = (1-1/n)^j$, and otherwise its last relative value is fair. Flipping the ov coordinates of Y in B_g^{f} gives

$$\mathbb{E}Z = 2J \mp J(1 - \mathrm{ov}/(2J))\theta_j,$$

with the sign of the side. Now fix B_g^{f} and the side. The initial label is the first J or $3J$ elements of a uniform permutation of B_g^{f}. Swapping two of them changes at most two initial bits, and changing one pair at most two final bits. So, given these data, Z is a function with at most $4J + j$ martingale increments of conditional range at most two, and $\Pr(Z - \mathbb{E}Z \geq y) \leq \mathrm{e}^{-y^2/(2(4J+j))}$, also for the lower tail. If $1 - \mathrm{ov}/(2J) \geq \eta/2$, crossing the threshold $2J$ needs a deviation of at least $J\eta\theta_{\hat{t}}/2$ in a fixed direction. Averaging over B_g^{f} and the side, the threshold function of B_g^{f} at the queried label differs from its value on the side with probability at most ζ.

All queries and groups. Let the execution on I^+ be *good* if the first s potential samples lie in distinct copies, each drawn sample has $L_g \leq \hat{t}$ for every group, and at each of the at most s_{h} held states and for every group, the threshold function of the frozen block of its sample has its value on the side. The union bound runs over the unconditioned execution on I^+ and uses only the distribution of each L_g, not a joint law of the groups. By Lemma 30(iii), the query bound applies to each held state and each group. This gives the stated bound on the probability that the execution is not good, and Lemma 28 applies to good executions. ■

Lemma 32 (Expected cost) *Let a tester accept a first input with probability at least $2/3$ and a second input with probability at most $1/3$, where an input may be a mixture of instances. Fix the operations that it counts. The* truncation at t *of an execution is the whole execution before its $(t+1)$-st counted operation, with the free operations in between and the outcome if it stops there. Let $\hat{p} \geq 0$, and suppose that for every t the truncations at t on the two inputs differ in total variation by at most $t\hat{p} + t(t-1)/(2N_c)$. Then the tester uses at least $\lfloor \min\{1/(12\hat{p}), \sqrt{N_c/6}\} \rfloor/6$ counted operations in expectation on the first input, with $1/0 = \infty$.*

Proof Let $t = \lfloor \min\{1/(12\hat{p}), \sqrt{N_c/6}\} \rfloor$, so the truncations at t differ by at most $1/6$. Let E be the event that the tester accepts before its $(t+1)$-st counted operation, which the truncation at t determines, and let E' be the event that it makes more than t of them. With $\Pr_+$ and $\Pr_-$ for the first and the second input, $\Pr_-(E) \leq 1/3$, so $\Pr_+(E) \leq 1/2$, and $\Pr_+(E) \geq 2/3 - \Pr_+(E')$. So $\Pr_+(E') \geq 1/6$, and the expected number is at least $t/6$. ■

Lemmas 29 and 31 supply this hypothesis. The tester stopped before its $(t+1)$-st counted operation has at most t counted operations, and by Lemma 28 the coupling keeps its free operations and its outcome.

E.7. Necessity

Proposition 33 (Necessity) *Fix either distance. Consider testers that certify this distance at ϵ_m with confidence $2/3$ on samples of μ_{ℓ_m}, for all interfaces and start distributions.*

(a) Let $\epsilon_m = 1/(32m)$, $2 \le d_m \le \lfloor m/16 \rfloor$, $\ell_m^* = m \log(m/d_m)/\log m$ *and* $\ell_m = o(\ell_m^*)$. *Every such tester of depth at most* d_m *uses a superpolynomial expected number of samples on* I^+, *for the unmasked start distribution.*

(b) Fix $C > 1$ *and* $0 < c < 1/2$. *Let* $1/(32m) \le \epsilon_m \le (\log m)^{-C}$, $2 \le d_m \le cm$ *and* $\ell_m = o(m \log(1/\epsilon_m)/\log m)$. *Every such tester of depth at most* d_m, *even if it learns the length of each run, uses a superpolynomial expected number of samples and setter calls on* I^+, *for some start distribution* ν_ω.

(c) Let $1/(32m) \le \epsilon_m = o(1/\log m)$, $J = \lfloor 1/(512\epsilon_m) \rfloor$, $\ell_m = 0$ *and* $2 \le d_m \le \lfloor m/2 - w_m \rfloor$ *with* $w_m > 0$ *and* $Jw_m^2/(m^2 \log m) \to \infty$. *The conclusion of (b) holds.*

Proof (a) Take the unmasked assignment with $J = \lfloor m/8 \rfloor$, $n = 8J$ and $N_c = 2^m$. As $\ell_m^* \le m$, $\ell_m = o(m)$, so $J/(1+\ell_0) \to \infty$ and $q_0^J \to 0$. Lemma 27 gives $D_{\ell_m}^\infty \ge (1 - o(1))/(128J) \ge (1 - o(1))/(16m) > \epsilon_m$ for large m and every ω. So a tester as in the statement accepts (I^+, ν_ω) with probability at least $2/3$ and (I_ω^-, ν_ω) with probability at most $1/3$.

Put $d^+ = \max\{d_m, n^{3/4}\}$ and $\hat{t} = \lfloor (n-1)\log(J/(3d^+/2)) \rfloor$. As $d_m \le J/2 = n/16$ and $n^{3/4} \le n/16$ for large m, the logarithm is at least $\log(4/3)$, so $\hat{t} \ge 1$. As $\log(1 - 1/n) \ge -1/(n-1)$, $J_{\hat{t}} \ge 3d^+/2$, so $J_{\hat{t}} - d_m \ge d^+/2$. As $\hat{t} \le n \log n/4$ and $d^+ \ge n^{3/4}$, the second term of $\hat{p}$ in Lemma 29 is at most $\mathrm{e}^{-(d^+)^2/(8\hat{t})} \le \mathrm{e}^{-\sqrt{n}/(2\log n)}$. Let $a_0 = 1 - \log 12/\log 16 > 0$. If $d_m \ge n^{3/4}$, the logarithm in $\hat{t}$ is $\log(n/(12d_m)) \ge a_0 \log(n/d_m)$, as $n/d_m \ge 16$. Otherwise it is at least $(\log n)/8 \ge \log(n/d_m)/8$ for $n \ge 12^8$. So $\hat{t}$ is at least a positive constant times $m \log(m/d_m)$, as $n/m \to 1$. This gives $(\hat{t}+1)/((1+\ell_0)\log m) \to \infty$, since $\ell_0 \le \ell_m = o(\ell_m^*)$ and $\ell_m^* \to \infty$. As $q_0 \le \mathrm{e}^{-1/(1+\ell_0)}$, $\hat{p} = m^{-\omega(1)}$. Lemma 32, counting samples, gives the claim.

(b) Take $J = \lfloor 1/(512\epsilon_m) \rfloor$, $n = 8J$, $N_g = \lfloor m/n \rfloor$ and $N_c = 2^m$. For large m, $J \ge (\log m)^C/1024$ and $\log J \sim \log(1/\epsilon_m)$, and $J \le m/16$ as $\epsilon_m \ge 1/(32m)$, so $N_g \ge 2$. As $\log(1/\epsilon_m) \le \log(32m)$, $\ell_m = o(m)$, so

$$q_0^J \le \exp\Big[-\frac{J}{1 + 8J\ell_m/m}\Big] = \exp\Big[-\frac{1}{1/J + 8\ell_m/m}\Big] \to 0.$$

Lemma 27 gives $D_{\ell_m}^\infty \ge (1 - o(1))/(256J) \ge (2 - o(1))\epsilon_m > \epsilon_m$ for large m and every ω. So a tester as in the statement accepts (I^+, ν_ω) with probability at least $2/3$ and (I_ω^-, ν_ω) with probability at most $1/3$.

Let $a_C = (C-1)/(4C) > 0$ and $\hat{t} = \lfloor a_C(n-1)\log J \rfloor$ in Lemma 31. As $\log(1 - 1/n) \ge -1/(n-1)$, $(1-1/n)^{\hat{t}} \ge J^{-a_C}$. With $\eta_0 = 1 - 2c > 0$ and $\eta \ge \eta_0$, the exponent of the second term of (12) is at least $\eta_0^2 J^{1-2a_C}/(8(4 + 8a_C \log J))$. Here $C(1-2a_C) = (C+1)/2 > 1$, and $y^{1-2a_C}/(4 + 8a_C \log y)$ increases for large y. So, with $J \ge (\log m)^C/1024$, this exponent divided by $\log m$ is at least a positive constant times $(\log m)^{(C-1)/2}/\log\log m$, which tends to infinity. The first term of (12) is superpolynomially small as well, so $\zeta = m^{-\omega(1)}$. Next,

$$\frac{n \log J}{(1 + n\ell_m/m)\log m} = \Big(\frac{\log m}{n \log J} + \frac{\ell_m \log m}{m \log J}\Big)^{-1} \to \infty,$$

so $q_0^{\hat{t}+1} \le \exp[-(\hat{t}+1)/(1+\ell_0)] = m^{-\omega(1)}$. As $N_g \le m$, $\hat{p} = N_g(q_0^{\hat{t}+1} + \zeta) = m^{-\omega(1)}$. Lemma 32, counting samples and setter calls, with the mixtures over ω of the instances (I^+, ν_ω) and of the instances (I_ω^-, ν_ω) as the two inputs, gives the bound on the first mixture, hence for some ω.

(c) Take $n = 8J$, $N_g = \lfloor m/n \rfloor$ and $N_c = 2^m$, so $J \le m/16$. As $q_0 = 0$, Lemma 27 gives $D_\nu^\infty \ge 1/(256J) \ge 2\epsilon_m$. As $d_m \le m/2 - w_m$, $\eta = 1 - 2d_m/m \ge 2w_m/m > 0$, so $d_m < m/2$ and Lemma 31 applies. Take $\hat{t} = 0$ there, so $\zeta \le 2\exp[-J\eta^2/32]$. Squaring $\eta \ge 2w_m/m$, whose two sides are positive, gives $\zeta \le 2\exp[-Jw_m^2/(8m^2)] = m^{-\omega(1)}$. Conclude as in (b). ■

Proof of Theorem 2 Part (b) is Corollary 21 and (8). For part (a), $\epsilon_m = m^{-\beta}$ lies in the range of Proposition 33(b) for $C = 2$ and large m, and $\log(1/\epsilon_m)/\log m = \beta$, so its hypothesis is $\ell_m = o(m)$. If $\ell_m \ge \alpha m$ for all large m, part (b) gives a polynomial budget at depth $2 \le d_m$. If ℓ_m is not $\Omega(m)$, some subsequence has $\ell_m/m \to 0$, and part (a) along it excludes a polynomial budget.

For the statement after the theorem, let $(\log m)^C \le \ell_m = o(m)$ and $\epsilon_m = 1/\ell_m$, which lies in the range of Proposition 33(b) for large m. Fix $y > 0$. As $t/\log t$ increases for $t \ge \mathrm{e}$ and $\ell_m \le ym$ for large m, $\ell_m \log m/(m \log \ell_m) \le y \log m/\log(ym) \to y$. So $\ell_m = o(m \log(1/\epsilon_m)/\log m)$. ■

Proof of Theorem 4(c) *Upper bound.* By Hoeffding's inequality, $\tau_2 \le \exp[-2(d - 2 - m/2)^2/m]$ for $d \ge 2 + m/2$. So $d = 2 + \lceil m/2 + \sqrt{(m/2)\log(2(m+2)/\epsilon_m)} \rceil$ gives $(m+2)\tau_2 \le \epsilon_m/2$, and $d = m/2 + O(\sqrt{m \log m})$. If this value exceeds $m+1$, take $d = m+1$ and part (a). Part (b) gives certification of both distances with $\lceil 2\log 3/\epsilon_m \rceil$ rounds of at most $2m^2 + 6m$ setter calls, which is a polynomial budget.

Lower bound. Take $w_m = m^{1-\beta/2}\log m$, which is positive and $o(m)$, in Proposition 33(c). For large m, $J \ge m^\beta/1024$, so $Jw_m^2/(m^2 \log m) \ge \log m/1024 \to \infty$. So depths at most $m/2 - w_m$ have no polynomial budget, for both distances. ■

E.8. Samples from the start distribution

Let $\ell = 0$, so that a tester draws samples of the start distribution $\mu = \nu$. A tester is an (ϵ, δ)-*tester* if it certifies the distance under μ at ϵ with confidence $1 - \delta$ for every interface on a finite set and every start distribution μ.

Theorem 34 (Samples from the start distribution) *For both distances, at $\ell = 0$: let $0 < \delta < 1/2$, $m \ge 32$ and $0 < \epsilon \le 1/512$. No (ϵ, δ)-tester has a bound on its number of samples, or on its expected number, that depends only on m, ϵ and δ, if its depth is at most $\lfloor \min\{m/8, 1/(128\epsilon)\} \rfloor$, which is at least $\min\{m, 1/\epsilon\}/256$.*

Proof of Theorem 34 Let $y = \min\{m/8, 1/(128\epsilon)\}$. The ranges of m and ϵ give $y \ge 4$, so $J = \lfloor y \rfloor \ge y/2 \ge \min\{m, 1/\epsilon\}/256$, and $m \ge 8J$. Take the unmasked instance with $\ell = 0$ and any N_c. As $J \le 1/(128\epsilon)$, Lemma 27 gives $D_\nu(I^-) \ge D_\nu^\infty(I^-) \ge 1/(128J) \ge \epsilon$.

Fix a start state with block labels u_0', u_0'', hidden bit z_0 and $b = f_2(u_0'')$, and on the states of its copy with hidden bit z_0 whose block labels lie within Hamming distance J of (u_0', u_0'') put $\mathrm{ren}(\sigma, v, z_0) = (\sigma, v, z_0 \oplus b[f_1(u') \oplus f_1(u_0')])$. This map fixes the start state, keeps the readers and is injective. Setters of B_σ'' and of passive variables do not change f_1, and I^- treats them as I^+ does. A setter of B_σ' whose two labels lie in this ball changes f_1 only while $f_2 = b$. Indeed, if $|u_0''| \in \{J, 3J\}$ and $f_2(u'') \ne b$, then u'' is at Hamming distance at least J from u_0'', so $u' = u_0'$ at both ends and f_1 does not change. If $|u_0'| \in \{J, 3J\}$ and f_1 changes, one end is at Hamming distance at least J from u_0', so $u'' = u_0''$ and $f_2 = b$. Every start state has such a block. So $T_{i,a}^- \mathrm{ren}(w) = \mathrm{ren}(T_{i,a}^+ w)$ for every state w of this domain whose block labels lie within Hamming distance $J - 1$ of (u_0', u_0''). A

setter call of depth at most J from the start state acts on a state reached by at most $J-1$ calls. So, as in Lemma 28, the two transcripts of a tester of depth at most J coincide when its samples lie in distinct copies, and two of s samples share a copy with probability at most $s(s-1)/(2N_c)$.

Let $t = \lfloor \sqrt{N_c(1-2\delta)} \rfloor$, let E be the event that the tester accepts before it draws a $(t+1)$-st sample, and let E' be the event that it draws more than t samples. Then $\Pr_-(E) \le \delta$ and $\Pr_+(E) \ge 1 - \delta - \Pr_+(E')$. The coupling then gives $\Pr_+(E') \ge 1 - 2\delta - t(t-1)/(2N_c) \ge (1-2\delta)/2$. Hence the expected number of samples under I^+ is at least $t(1-2\delta)/2$, which grows without bound with N_c. A bound on the number of samples that depends only on m, ϵ and δ would bound this expectation. ■

Appendix F. Proof of Theorem 5

The necessity in Theorem 5 follows from Proposition 33(a). If ℓ_m is not $\Omega(\ell^*_m)$, some subsequence has $\ell_m/\ell^*_m \to 0$, and the proposition, applied along it, excludes a polynomial budget of samples. This appendix proves the sufficiency. A random word need not write every variable: short words from its end can stand in for the few variables that it misses.

For an integer $\bar{d} \ge 0$, let $\mathcal{N}_{\bar{d}}$ be the set of states from which some word of length at most $\bar{d}$ reaches $\mathcal{N}$, so $\mathcal{N}_0 = \mathcal{N}$. Let $H_0 = 0$ and $H_j = \sum_{v \le j} 1/v$.

Theorem 35 (Partial coverage) *Let $m \ge 2$, $2 \le d \le m$, $\bar{d} = d-2$, $\lambda \ge 1$, let μ be λ-subinvariant and $t = \lceil 2m(H_m - H_{\bar{d}}) \rceil$. Then $D^\infty_\mu \le \min\{1, 2(t+1)\lambda^t \mu(\mathcal{N}_{\bar{d}})\}$, and if $2^m \mid |S|$, then $D_\mu \le \min\{1, (t+1)(t+4)\lambda^t \mu(\mathcal{N}_{\bar{d}})\}$.*

The tester. Exhausting all words of length at most $\bar{d}$ from a sample and evaluating all law instances at their ends has depth at most $\bar{d}+2 = d$. It rejects with probability $\mu(\mathcal{N}_{\bar{d}})$ and accepts every composable interface. So, if K is the coefficient of Theorem 35, $\lceil K \log(1/\delta)/\epsilon \rceil$ samples certify as in Appendix A. The number of setter calls of a sample can be exponential in d.

Proof of Theorem 35 *The word.* Draw t independent uniform setters $E_1, \dots, E_t$, with prefixes $U_k = E_k \cdots E_1$. Let $\mathcal{H}$ be the event that at most $\bar{d}$ variables are not written, and let R_0 be the set of these variables. On $\mathcal{H}$, let W_0 write 0 once to each variable of R_0 in increasing order, and put $W = W_0 U_t$, which writes every variable. The number of draws needed to collect $m - \bar{d}$ distinct variables has mean $m(H_m - H_{\bar{d}})$, so Markov's inequality gives $\Pr(\mathcal{H}) \ge 1/2$.

The distance with added states. Let A_W be the set of states with $U_k x \notin \mathcal{N}$ for $k < t$ and $U_t x \notin \mathcal{N}_{\bar{d}}$. Every step of W_0 acts on a state within $\bar{d}-1$ writes of $U_t x$, so the path of $x \in A_W$ under W avoids $\mathcal{N}$ before its last step. By Lemma 13, Lemma 14 applies to A_W, and $D^\infty_\mu \le \mu(S \setminus A_W)$. For $x \sim \mu$ independent of the setters, $U_k x$ has distribution $\mu P^k \le \lambda^k \mu$, and $\mathcal{N} \subseteq \mathcal{N}_{\bar{d}}$. So, dropping $\mathcal{H}$ from the numerator and dividing by $\Pr(\mathcal{H}) \ge 1/2$,

$$\mathbb{E}[\mu(S \setminus A_W) \mid \mathcal{H}] \le 2\Big[\sum\nolimits_{k<t} \lambda^k + \lambda^t\Big]\mu(\mathcal{N}_{\bar{d}}) \le 2(t+1)\lambda^t \mu(\mathcal{N}_{\bar{d}}). \tag{13}$$

A local rule. Let $z \notin \mathcal{N}_{\bar{d}}$ and $|R_0| \le \bar{d}$. Let W_g write the values $g \in \{0,1\}^{R_0}$ to the variables of R_0, once each, in increasing order. Each suffix of W_g acts on a state within $\bar{d}-1$ writes of z, so these states lie outside $\mathcal{N}$. Commuting a setter $T_{i,e}$ with $i \in R_0$ through W_g and absorbing it by overwrite gives $T_{i,e} W_g z = W_{g'} z$, where g' is g with coordinate i replaced by e. In particular

$W_{g''}W_g z = W_{g''}z$ for all g, g''. Commuting a setter E of a variable outside R_0 through W_g gives $EW_g z = W_g E z$.

A second word. Let $g_1, \dots, g_t$ be fair bits and $g_0 \in \{0,1\}^{R_0}$ a fair vector, independent of the rest. Put $V_k = T_{i_k, g_k}$ and $V = V_1 \cdots V_t$, which applies V_t first. For $1 \le a \le k \le t$ put $z_{a,k} = U_{a-1} V_{k+1} \cdots V_t x$, and for $0 \le j \le t$ put $x'_j = U_j V_{j+1} \cdots V_t x$. The positions used by $z_{a,k}$ are disjoint, so before conditioning $z_{a,k}$ has distribution $\mu P^{t-k+a-1}$, and x'_j has distribution μP^t. Consider the event $\mathcal{G}$ that all $z_{a,k} \notin \mathcal{N}$ and all $x'_j \notin \mathcal{N}_{\bar{d}}$. Summing the $t(t+1)/2$ and the $t+1$ terms,

$$\Pr(\mathcal{G}^c \mid \mathcal{H}) \le 2\Big[\frac{t(t+1)}{2} + t + 1\Big]\lambda^t \mu(\mathcal{N}_{\bar{d}}). \tag{14}$$

A witness. Assume $\mathcal{G}$. Fix k and let j_k be the first position with $i_{j_k} = i_k$. Commutation at $z_{a,k}$ for $a < j_k$ and overwrite at $z_{j_k,k}$ give, for $j \ge k$, $U_j V_k V_{k+1} \cdots V_t x = U_j V_{k+1} \cdots V_t x$. Removing $V_1, \dots, V_j$ in turn gives $U_j V x = x'_j$. Let $y = W_{g_0} V x$, and let $\hat{u}$ be the label predicted by applying the writes of V and W_{g_0} to $r(x)$. Given W and x, $\hat{u}$ is uniform: V applies V_1 last, so each written variable takes the bit g_k of V at its first position k in W, and each variable of R_0 a coordinate of g_0. The steps of V act at $z_{1,k} \notin \mathcal{N}$ and those of W_{g_0} within $\bar{d} - 1$ writes of $x'_0 = Vx \notin \mathcal{N}_{\bar{d}}$, so $r(y) = \hat{u}$. The steps of U_t avoid R_0 and act at $U_{j-1} V x = x'_{j-1} \notin \mathcal{N}_{\bar{d}}$, so the local rule moves W_{g_0} through them: $U_t y = W_{g_0} U_t V x = W_{g_0} U_t x$. As $x'_t = U_t x \notin \mathcal{N}_{\bar{d}}$, the local rule gives $W_0 W_{g_0} U_t x = W_0 U_t x$. Hence $Wy = Wx$ and $r(y) = \hat{u}$.

The distance on the same states. By Lemma 14, $D_\mu \le \mu(S \setminus A_W) + \Delta_{A_W}$. Given W, the term Δ_{A_W} is at most the probability that $\hat{u}$ occurs nowhere in the fiber of x, which needs $\mathcal{G}^c$. Averaging given $\mathcal{H}$ and adding (13) and (14) gives $2[t(t+1)/2 + 2t + 2]\lambda^t \mu(\mathcal{N}_{\bar{d}}) = (t+1)(t+4)\lambda^t \mu(\mathcal{N}_{\bar{d}})$. ■

Proof of the sufficiency in Theorem 5 Let $2 \le d \le m/16$, $\bar{d} = d - 2$, and $\ell \ge b\ell^*_m$ with $\ell^*_m = m \log(m/d)/\log m$ and a constant $b > 0$. For $\bar{d} \ge 1$, integral comparison gives $H_m - H_{\bar{d}} \le \log(m/\bar{d})$, and $\bar{d} \ge d/3$ for $d \ge 3$, so $\log(m/\bar{d}) \le \log 3 + \log(m/d) \le 2\log(m/d)$. For $\bar{d} = 0$, $H_m \le 1 + \log m \le 2\log(m/2)$ for $m \ge 32$. So $t \le 5m \log(m/d)$, and with $\lambda = 1 + 1/\ell$, $\lambda^t \le \mathrm{e}^{t/\ell} \le m^{5/b}$. As $(t+1)(t+4)$ is polynomial in m, the tester above uses a polynomial number of samples, for both distances; for D^∞, no condition on $|S|$ is needed.

The explicit bound. Let $m \ge 96$, $d = \lfloor m/16 \rfloor$ and $\ell \ge \varsigma m/\log m$ with $\varsigma > 0$. Then $\bar{d} \ge m/32 \ge 1$, so $H_m - H_{\bar{d}} \le \log 32$ and $t \le 2m \log 32 + 1 \le 8m$, and $(t+1)(t+4) \le 140m^2$. Also $\lambda^t \le \mathrm{e}^{t/\ell} \le m^{8/\varsigma}$. So $\lceil 140 m^{2+8/\varsigma} \log(1/\delta)/\epsilon \rceil$ samples certify both distances. ■

Appendix G. Proof of Theorem 1

Route. Part (a) follows the four steps of Section 3, and each subsection gives one ingredient.

- Appendix G.1 sets up continuous time and predicted labels, and gives two events on which the history of a variable matters only through B_j: (16) moves the last write to the front, and strongly good words allow a whole history to be replaced. Their failure probabilities are charged to γ in (17) and (18).
- Appendix G.2 couples the histories given B as the time grows, which bounds the derivative in time of the end state given B in (19).
- Appendix G.3 proves a bound on the density of the time given a label (Lemma 36) and turns a derivative into a variance bound (20).

- Appendix G.4 introduces the weight ϖ, which is what a setter pays when it resamples a history. It bounds the error from resampling one history at a time (23) and from the time (25), and it chooses a state $f(b)$ for each label $b \neq r(o)$.
- Appendix G.5 treats the label $r(o)$, where the weight is not integrable at time 0 and $f(r(o)) = o$.
- Appendix G.6 places the chosen states in one cube and proves the bound (28) for $\ell \geq 2$.
- Appendix G.7 treats means at most 2, where a Stein kernel replaces the coupling and the integral kernel.
- Appendix G.8 assembles the cases, with means of at least $2m$ taken from (8), and Appendix G.9 proves part (b).

Throughout Appendices G.1–G.6, $m \geq 2$, the start distribution is $\nu = \delta_o$, $\ell \geq 2$, $q = \ell/(1+\ell)$, $\mu = \mu_\ell$, $\gamma = \gamma_\mu$, $\alpha = \ell/m$ (this ratio, not the free constant α of Theorem 2), and all law instances hold at o. Each law instance then fails with probability at most γ at $x \sim \mu$, and this is the only use of γ.

G.1. Histories and two good events

Continuous time. Let the edits arrive at the times of a Poisson process of rate 1, each a uniformly random setter, and let $\rho_t = \mathrm{e}^{-t} \sum_{n \geq 0} (t^n/n!)\, \nu P^n$ be the distribution of the state at time t. Stop at an independent time ι with density $\ell^{-1}\mathrm{e}^{-t/\ell}$. The number L of edits before ι satisfies

$$\Pr(L = n) = \frac{1}{\ell\, n!} \int_0^\infty \mathrm{e}^{-(1+1/\ell)t} t^n \, dt = \frac{\ell^n}{(1+\ell)^{n+1}} = (1-q)q^n,$$

by integration by parts and induction on n. So the state X at time ι has distribution μ; this holds for every start distribution and every $\ell > 0$. For a nonnegative function φ on states, this gives $\mathbb{E} \int_0^\iota \rho_s \varphi \, ds = \int_0^\infty \mathrm{e}^{-s/\ell} \rho_s \varphi \, ds = \ell \mu \varphi$.

Predicted labels. Let $b^0 = r(o)$. Applying the same edits to the vector b^0 as coordinate writes gives the *predicted label* B. Split by variable, the edits form independent Poisson processes of rate $1/m$, with independent fair written values. So, with $w = |b \oplus b^0|$,

$$\Pr(B = b \mid \iota = t) = 2^{-m}(1 - \mathrm{e}^{-t/m})^w (1 + \mathrm{e}^{-t/m})^{m-w}, \tag{15}$$

which is positive for $t > 0$. Given $B = b$ and $\iota = t$, the histories of the variables, with their arrival times and written values, are still independent, as the event $B = b$ is the intersection of one event for each variable.

Moving the last write to the front. Fix a variable j. For an edit $E = T_{i,a}$, let $I_j(x, E)$ indicate that, for $b = 0$ or $b = 1$, the equality $T_{j,b}Ex = ET_{j,b}x$ fails if $i \neq j$, or $T_{j,b}T_{j,a}x = T_{j,b}x$ fails if $i = j$. Let $\bar{I}_j(x) = (2m)^{-1} \sum_{i,a} I_j(x, T_{i,a})$, so $\mu \bar{I}_j \leq 2\gamma$. Let $W = E_n \cdots E_1$ be a word, $U_s = E_s \cdots E_1$ and $U_0 = \mathrm{Id}$, and let W_{-j} delete the writes to j from W. If $I_j(U_{s-1}o, E_s) = 0$ for every s before the last write to j, then

$$Wo = W_{-j} T_{j,B_j} o. \tag{16}$$

Indeed, if W writes j, let s_* be its last write to j, with value $b = B_j$, and put $T_* = T_{j,b}$. For $s < s_*$, induction gives $T_* U_s o = U_s^{(-j)} T_* o$, where $U_s^{(-j)}$ deletes the writes to j: an edit of another variable commutes with T_* at $U_{s-1}o$, and a write to j is overwritten by T_*. So $U_{s_*} o = T_* U_{s_*-1} o =$

$U_{s_*-1}^{(-j)} T_* o$, and the later edits do not write j. If W does not write j, then $B_j = b_j^0$, and write-back at o gives $T_{j,b_j^0} o = o$. If $I_j(U_{s-1}o, E_s) = 0$ for all s, the same induction gives $T_{j,a} W o = W_{-j} T_{j,a} o$ for $a = 0, 1$.

Let $\bar{p}_j(t)$ be the probability that the word of the edits up to time t fails the condition of (16), and $p_j(t)$ the probability that it fails the stronger condition for all s. An edit arriving at time s acts at a state of distribution ρ_s, and a later write to j occurs with probability $1 - \mathrm{e}^{-(t-s)/m}$, independently of this state and this edit. So

$$\bar{p}_j(t) \le \int_0^t \rho_s \bar{I}_j \, [1 - \mathrm{e}^{-(t-s)/m}] \, ds, \qquad p_j(t) \le \int_0^t \rho_s \bar{I}_j \, ds. \tag{17}$$

Strongly good words. Call a word *strongly good for* j if the four overwrite equalities $T_{j,a} T_{j,b} x = T_{j,a} x$ hold at every prefix state, including the end state, and every edit of another variable commutes with $T_{j,0}$ and $T_{j,1}$ at the state where it acts. If a word is strongly good for j, then replacing the whole history of j by another one with the same value B_j keeps the end state. To see this, scan the old writes to j, the new writes to j and the common other edits in their merged order, with x the old prefix state, y the new prefix state and e the new predicted value of j. The invariant $y = T_{j,e} x$ holds at the start by write-back at o. An old write to j changes only x, and overwrite at x keeps the invariant. A new write to j changes y and e, and overwrite at x keeps it. Another edit commutes with $T_{j,e}$ at x. At the end, $y = T_{j,B_j} X$. Induction along the old word, with commutation for other edits and $T_{j,a} T_{j,a} x = T_{j,a} x$ for writes to j, gives $T_{j,B_j} X = X$. So $y = X$. A new write may come after all old edits, which is why the end state is included.

Let I_j° indicate that one of the four overwrite equalities fails at a state, and let $\bar{I}_j^*$ be I_j° plus the average over the edit of the indicator that this edit, if it does not write j, fails to commute with $T_{j,0}$ or $T_{j,1}$. Then $\mu I_j^\circ \le 4\gamma$ and $\mu \bar{I}_j^* \le 6\gamma$. Let $p_j^*(t)$ be the probability that the word up to time t is not strongly good for j, and $p_j^*(t \mid b)$ the same given $B = b$. Then $p_j^*(t) \le \int_0^t \rho_s \bar{I}_j^* \, ds + \rho_t I_j^\circ$, so

$$\mathbb{E}\, p_j^*(\iota) \le (6\ell + 4)\gamma. \tag{18}$$

G.2. A pure-birth coupling given the label

Fix a variable, let $y = t/m$, $p_n = \mathrm{e}^{-y} y^n / n!$ and $R_n = \Pr\{\mathrm{Pois}(y) > n\}$. Given its predicted value and t, the number of writes to this variable has the masses $p_n/(1-p_0)$ for $n \ge 1$ if the value differs from the initial one, and $2p_0/(1+p_0)$ for $n = 0$ and $p_n/(1+p_0)$ for $n \ge 1$ otherwise. Differentiating the distribution functions $1 - R_n/(1 \mp p_0)$ in t, using $dp_0/dy = -p_0$ and $dR_n/dy = p_n$, shows that these masses are the marginals of a pure-birth process with rates

$$\frac{1}{m}\Big[1 - \frac{p_0 R_n}{(1-p_0)p_n}\Big], \qquad \frac{1}{m}\Big[1 + \frac{p_0 R_n}{(1+p_0)p_n}\Big] \ (n \ge 1), \qquad \frac{1}{m(1+p_0)} \ (n = 0),$$

in the two cases. Let $N(t)$ be the state of this process at time t. Its entrance law is $N(0+) = 1$ when the predicted value differs from the initial one, and $N(0+) = 0$ when they are equal. As $R_n/p_n = \sum_{v \ge 1} y^v n!/(n+v)! \le \mathrm{e}^y - 1$, all rates lie in $[0, 2/m]$. Given $n \ge 1$ writes, the normalized arrival times are independent uniform points of $[0, 1]$, the last value is the given one, and the earlier values are fair. At a birth, insert an independent uniform point. If it is not the new last point, give it a fair value. If it is, redraw the old last value as a fair bit and give the new point the given value. For $n = 0$ there is no old last value, so the first birth skips the redraw. This keeps the

law of the history. Multiplying by t does not change the order of the normalized points. So running this process independently for each variable couples the histories given $B = b$ for all $t > 0$, and a birth changes one history only. If the word before the birth is strongly good for that variable, the end state does not change.

Let $g_{b,x}(t) = \Pr(X = x \mid B = b, \iota = t)$ and $v_b(t) = 1 - \sum_x g_{b,x}(t)^2$. As the rates are at most $2/m$, the probability that the end state changes in $(t, t+h]$ is at most $(2/m)\int_t^{t+h} \sum_j p_j^*(s \mid b)\, ds$. These conditional probabilities are quotients of power series that converge uniformly with their derivatives on compact subsets of $(0, \infty)$. So

$$\sum\nolimits_x |g'_{b,x}(t)| \le \frac{4}{m} \sum\nolimits_j p_j^*(t \mid b). \tag{19}$$

G.3. An integral kernel of the time given the label

Lemma 36 (Integral kernel) *Let $n \ge 1$ and $0 \le \bar{k} \le n$ be integers and $\ell' > 0$. Let ψ be the probability density on $(0, \infty)$ proportional to $\mathrm{e}^{-t/\ell'}(1 - \mathrm{e}^{-t/n})^{\bar{k}}(1 + \mathrm{e}^{-t/n})^{n-\bar{k}}$, with distribution function Ψ. Then $\Psi(t)(1 - \Psi(t))/\psi(t) \le 200(1 + \ell'/n)\min\{t, n\}$ for all $t > 0$.*

Proof Let $\mathcal{S} = (\log \psi)'$, so

$$\mathcal{S}(t) = -\frac{1}{\ell'} + \frac{\bar{k}}{n(\mathrm{e}^{t/n} - 1)} - \frac{n - \bar{k}}{n(\mathrm{e}^{t/n} + 1)}, \qquad \mathcal{S}(t) \ge -\frac{1}{\ell'} - \frac{1}{2}.$$

With $a' = n/\ell'$ and $y = \mathrm{e}^{-t/n}$, $n(1 - y^2)\mathcal{S}(t) = (a' + n)y^2 + (2\bar{k} - n)y - a'$. The right side has roots $y_+ \in (0, 1]$ and $-y_- < 0$, with $(a' + n)y_+y_- = a'$; $y_+ = 1$ exactly when $\bar{k} = 0$. So ψ has a unique mode at $t_* = -n \log y_+$: an interior maximum if $\bar{k} \ge 1$, and the boundary point $t_* = 0$ if $\bar{k} = 0$, when ψ decreases on $(0, \infty)$. With $v = \mathrm{e}^{-(t - t_*)/n}$,

$$\mathcal{S}(t) = \frac{1}{\ell'}(v - 1)\frac{1 + (y_+/y_-)v}{1 - y_+^2 v^2},$$

where the last fraction is at least 1. Hence $\mathcal{S}(t_* + z) \le -(1 - \mathrm{e}^{-z/n})/\ell'$ for $z > 0$, and $\mathcal{S}(t_* - z) \ge (\mathrm{e}^{z/n} - 1)/\ell'$ for $0 < z < t_*$.

The scale n. For $t \le t_*$, moving left by at most n the density ratio is at most 1, and beyond that the logarithmic derivative is at least $(\mathrm{e} - 1)/\ell' \ge 1/\ell'$. So $\Psi(t)/\psi(t) \le n + \ell'$. For $t \ge t_*$, moving right by at most n the ratio is at most 1, and beyond that the logarithmic derivative is at most $-(1 - \mathrm{e}^{-1})/\ell' \le -1/(2\ell')$. So $(1 - \Psi(t))/\psi(t) \le n + 2\ell'$. As $\Psi(1 - \Psi) \le \min\{\Psi, 1 - \Psi\}$, always $\Psi(1 - \Psi)/\psi \le n + 2\ell' \le 2(1 + \ell'/n)n$. This gives the claim for $t \ge n/64$, as then $n \le 64 \min\{t, n\}$.

The scale t. Let $t < n/64$. If $t \le t_*$, ψ does not decrease on $(0, t]$, so $\Psi(t)/\psi(t) \le t$. Let $t > t_*$, and let $q' = \ell'/(1 + \ell')$. If $t \le q'$, then for $0 < z < t$, $\psi(z)/\psi(t) \le \mathrm{e}^{(1/\ell' + 1/2)t} \le \mathrm{e}$, so $\Psi(t)/\psi(t) \le 3t$.

Let $t > q'$ and $\ell' \ge 2$, so $t \ge 2/3$. For $0 < z \le \bar{k}/2 \le n/2$, $\mathrm{e}^{z/n} - 1 < 2z/n$ gives $\mathcal{S}(z) > -1/\ell' - 1/2 + \bar{k}/(2z) \ge 0$. So $t_* \ge \bar{k}/2$, and $\bar{k} \le 2t$. Let $z_0 = 32t < n/2$. For $z_0 \le z \le n$, $\mathrm{e}^{z/n} - 1 \ge z/n$ and $\mathrm{e}^{z/n} + 1 < 4$ give $\mathcal{S}(z) \le \bar{k}/z - (1 - \bar{k}/n)/4 \le 1/16 - 31/128 < -1/16$. So ψ does not increase on $[t, z_0]$, it falls at rate at least $1/16$ on $[z_0, n]$, and $\psi(n)/\psi(t) \le \mathrm{e}^{-n/32}$. With the bound $n + 2\ell'$ at n and $n\mathrm{e}^{-n/32} \le 32$,

$$\frac{1 - \Psi(t)}{\psi(t)} \le 32t + 16 + 32(1 + 2\ell'/n) \le (104 + 96\ell'/n)t.$$

Let $t > q'$ and $\ell' < 2$. For $0 < z \le \bar{k}\ell'/(2+\ell') < n/2$, the same estimate gives $\mathcal{S}(z) > 0$, so $\bar{k} \le (2/\ell' + 1)t$. For $z \ge 8t$, $\mathcal{S}(z) \le -1/\ell' + \bar{k}/z \le -(6-\ell')/(8\ell') \le -1/(2\ell')$. So $(1-\Psi(t))/\psi(t) \le 8t + 2\ell' \le 14t$, as $\ell' \le (1+\ell')t < 3t$. ∎

For a positive probability density ψ with distribution function Ψ and a locally absolutely continuous vector g of probabilities, the inequality $\|y - y'\|_2^2 \le \|y - y'\|_1$ for probability vectors gives

$$\sum\nolimits_x \operatorname{Var}_\psi g_x = \int_{s<t} \|g(t) - g(s)\|_2^2 \, \psi(s)\psi(t)\, ds\, dt \le \int_0^\infty \|g'(z)\|_1 \Psi(z)(1-\Psi(z))\, dz. \quad (20)$$

G.4. Weighted errors and the chosen states

Let $\varpi(t) = 1/(1-\mathrm{e}^{-t/m})$, which is at least 1, and $\Omega_\ell = \ell + (m/q)\log(1+\ell)$. For every nonnegative φ on states,

$$\mathbb{E}\Big[\varpi(\iota) \int_0^\iota \rho_s \varphi \, ds\Big] \le \Omega_\ell \, \mu\varphi. \quad (21)$$

Indeed, $\varpi(t) \le 1 + m/t$. The constant part gives $\ell\mu\varphi$. For the other part, expand $\rho_s\varphi$ over the number k of edits. The coefficient of $\nu P^k \varphi$ is $(m/\ell)\int_{1/\ell}^\infty [v(1+v)^{k+1}]^{-1} dv$, after writing $\mathrm{e}^{-t/\ell}/t = \int_{1/\ell}^\infty \mathrm{e}^{-tv} dv$ and exchanging integrals. The substitution $(1+v)^{-1} = y$ turns it into $m(1-q)q^k \sum_{v\ge 0} q^v/(k+v+1) \le (1-q)q^k (m/q)\log(1+\ell)$, as $q/\ell = 1-q$.

One history at a time. Let $V(t) = \sum_b \Pr(B = b \mid \iota = t)\, v_b(t)$. For independent variables $\omega_1, \dots, \omega_m$ and a vector-valued function f of them,

$$\mathbb{E}\|f - \mathbb{E}f\|^2 \le \sum\nolimits_j \mathbb{E} \operatorname{Var}_{\omega_j}(f \mid \omega_{-j}) = \tfrac{1}{2} \sum\nolimits_j \mathbb{E}\|f - f^{(j)}\|^2, \quad (22)$$

where $f^{(j)}$ resamples ω_j. This is the tensorization of the variance, also called the Efron–Stein inequality (van Handel, 2016). Take $f = (\mathbf{1}_{X=x})_x$ and resample the history of j given B_j. By (16) the two end states agree unless one of the two words fails the condition there, and the two words have the same distribution. So $V(t) \le 2\sum_j \bar{p}_j(t)$. As $\varpi(t)[1 - \mathrm{e}^{-(t-s)/m}] \le 1$, (17) gives

$$\mathbb{E}[\varpi(\iota) V(\iota)] \le 4m\ell\gamma. \quad (23)$$

Along a run, the reading of j stays equal to B_j unless success or preservation fails. So, by (21),

$$\mathbb{E}[\varpi(\iota)\mathbf{1}_{r(X)\ne B}] \le m\Omega_\ell\gamma. \quad (24)$$

The time given the label. Let $b \ne b^0$. By (15), the density of ι given $B = b$ under the measure $\varpi(\iota)\, d\Pr$ is proportional to $\mathrm{e}^{-t/\ell}(1-\mathrm{e}^{-t/m})^{w-1}(1+\mathrm{e}^{-t/m})^{m-w}$; call it ψ_b, with distribution function Ψ_b. The substitution $n = m-1$, $\bar{k} = w-1$, $z = nt/m$, $\ell' = n\ell/m$ reduces it to Lemma 36, and gives $\Psi_b(1-\Psi_b)/\psi_b \le 200(1+\alpha)\min\{t, m\}$. Let $Z_b = \mathbb{E}[\varpi(\iota)\mathbf{1}_{B=b}]$. With (19), (20), $\varpi(t)\min\{t,m\}/m \le 2$ and (18),

$$\begin{aligned}\sum\nolimits_{b\ne b^0} Z_b \sum\nolimits_x \operatorname{Var}_{\psi_b} g_{b,x} &\le \frac{800(1+\alpha)}{m} \mathbb{E}\Big[\varpi(\iota)\min\{\iota, m\} \sum\nolimits_j p_j^*(\iota \mid B)\Big] \\ &\le 1600(1+\alpha)m(6\ell+4)\gamma.\end{aligned} \quad (25)$$

Put $\bar{g}_{b,x} = \mathbb{E}_{\psi_b} g_{b,x}$. The law of total variance gives $1 - \sum_x \bar{g}_{b,x}^2 = \mathbb{E}_{\psi_b} v_b + \sum_x \operatorname{Var}_{\psi_b} g_{b,x}$. For $b \ne b^0$ with $S_b = r^{-1}(b)$ nonempty, let $f(b) \in S_b$ have the largest $\bar{g}_{b,x}$. The largest $\bar{g}_{b,x}$ over all x is at least $\sum_x \bar{g}_{b,x}^2$ and exceeds $\bar{g}_{b,f(b)}$ by at most the mass of $\bar{g}_b$ outside S_b. So $1 - \bar{g}_{b,f(b)} \le 1 - \sum_x \bar{g}_{b,x}^2 + \sum_{x\notin S_b} \bar{g}_{b,x}$, also when S_b is empty, with the convention $\bar{g}_{b,f(b)} = 0$.

G.5. The initial label

Put $f(b^0) = o$. Let $v_0(t) = \ell^{-1}\mathrm{e}^{-t/\ell}\Pr(B = b^0 \mid t)$ and $\varpi_0 = \varpi v_0$, which is not integrable at 0. Given $B = b^0$ and t, the probability of no edits tends to 1 as $t \downarrow 0$, so $g_{b^0,o}(t) \to 1$. Let $e_0 = \int_0^\infty \varpi_0(t)[1 - g_{b^0,o}(t)]\,dt$. The bound of Appendix G.2 on the probability that the end state changes gives $1 - g_{b^0,o}(t) \le (2/m)\int_0^t \sum_j p_j^*(s \mid b^0)\,ds$. So, by an exchange of integrals,

$$e_0 \le \frac{2}{m}\int_0^\infty \sum\nolimits_j p_j^*(s \mid b^0)\int_s^\infty \varpi_0(t)\,dt\,ds.$$

For $t \le m$, the logarithmic derivative of v_0 is at most $-1/(\mathrm{e}^{t/m}+1) \le -1/4$, and beyond m at most $-1/\ell$. With $\varpi(t) \le 1 + m/t$ and $\varpi(t) \le 2$ for $t \ge m$, for $s \le m$,

$$\frac{1}{v_0(s)}\int_s^\infty \varpi_0(t)\,dt \le 4 + m[1 + \log(1 + 4/s)] + 2\ell \le m[6 + 2\alpha + \log^+(1/s)],$$

using $\int_0^\infty \mathrm{e}^{-z}/(z + s/4)\,dz \le \log(1 + 4/s) + 1$, $4 \le 2m$ and $\log 5 < 2$. For $s \ge m$ the left side is at most 2ℓ.

Also $\mathbb{E}[\log^+(1/\iota)\,p_j^*(\iota)] \le 18\gamma$. For the path term, $\int_s^1 \mathrm{e}^{-t/\ell}\log(1/t)\,dt \le 1 \le 2\mathrm{e}^{-s/\ell}$ for $0 < s < 1$, so its expectation is at most $2\mu\bar{I}_j^* \le 12\gamma$. For the end state, given k edits the density of ι is proportional to $\mathrm{e}^{-\lambda' t}t^k$ with $\lambda' = 1 + 1/\ell \le 3/2$, and ι is independent of the word. So $\mathbb{E}[\log^+(1/\iota) \mid L = k] \le \lambda'^{k+1}/(k!(k+1)^2) \le 3/2$, the bound being largest at $k = 0$, and this term is at most $(3/2)\mu I_j^\circ \le 6\gamma$. With (18),

$$e_0 \le (12 + 4\alpha)m(6\ell + 4)\gamma + 36m\gamma. \tag{26}$$

G.6. One cube

The states $f(b)$ have distinct labels. Place them in one cube, each at its label, fill the other positions of this cube with other states, which exist as 2^m divides $|S|$, and split the remaining states into cubes in any way. This defines a composable interface $\tilde{I}$ on S that keeps the readers at every state $f(b)$. Let $\mathcal{G} = \{(f(b), b) : S_b \ne \varnothing\}$ and $e_* = \mathbb{E}[\varpi(\iota)\mathbf{1}_{(X,B)\notin\mathcal{G}}]$. By the last paragraph of Appendix G.4 and (23), (25), (24) and (26),

$$e_* \le m\{4\ell + [1612 + 1604\alpha](6\ell + 4) + 36 + \Omega_\ell\}\gamma. \tag{27}$$

Fix a setter $T_{j,a}$. Given $\iota = t$, keep the histories of the other variables and resample the history of j conditioned on its last value being a. The new history has density at most $2\varpi(t)$ with respect to the old one, by (15). If both words satisfy the stronger condition of Appendix G.1, then the new end state is $Y = W_{-j}T_{j,a}o = T_{j,a}X$. If moreover (X, B) and $(Y, B[j \mapsto a])$ lie in $\mathcal{G}$, then $T_{j,a}X = f(B[j \mapsto a]) = T_{j,a}X$. By (21) and (17), $\mathbb{E}[\varpi(\iota)p_j(\iota)] \le 2\Omega_\ell\gamma$. So each setter disagrees with probability at most $e_* + 2e_* + 2\Omega_\ell\gamma + 4\Omega_\ell\gamma$, and each reader with probability at most e_*. With (27), $6\ell + 4 \le 8\ell$, $108 \le 54\ell$, $q \ge 2/3$ and $3m + 6 \le 6m$,

$$\begin{aligned} D_\mu &\le m\{12\ell + [4836 + 4812\alpha](6\ell + 4) + 108\}\gamma + (3m + 6)\Omega_\ell\gamma \\ &\le m^2[38760\alpha + 38496\alpha^2 + 9\log(1 + \ell)]\gamma. \end{aligned} \tag{28}$$

G.7. Short means

Let $0 < \ell < 4$, and keep the other notation. For an integer $0 \le w \le m$, let ψ_w° be the density on $(0,\infty)$ proportional to $\mathrm{e}^{-t/\ell}(1-\mathrm{e}^{-t/m})^w(1+\mathrm{e}^{-t/m})^{m-w}$, the density of ι given $|B \oplus b^0| = w$, with mean $\bar{t}_w$. Define ϑ_w, the Stein kernel of ψ_w° (Saumard, 2019), by $\vartheta_w(t)\psi_w^\circ(t) = \int_0^t (\bar{t}_w - z)\psi_w^\circ(z)\,dz$.

Lemma 37 (Short means) *For $0 < \ell < 4$, $\vartheta_w(t) \le \xi_\ell\, t$ for all $t > 0$, with $\xi_\ell = \ell/(1-\ell/4)$.*

Proof Fix w and drop it from the notation, and write $\xi = \xi_\ell$. *The mean.* Write $\psi^\circ(t)$ as a positive constant times $t^w \mathrm{e}^{-t/\ell} g(t)$ with $g(t) = ((1-\mathrm{e}^{-t/m})/t)^w(1+\mathrm{e}^{-t/m})^{m-w}$. Both factors of g do not increase: the derivative of $(1-\mathrm{e}^{-t/m})/t$ has the numerator $(1+t/m)\mathrm{e}^{-t/m} - 1 \le 0$. Let τ, τ' be independent Gamma variables with shape $w+1$ and scale ℓ, of mean $(w+1)\ell$. Then $2\,\mathrm{Cov}(\tau, g(\tau)) = \mathbb{E}[(\tau-\tau')(g(\tau)-g(\tau'))] \le 0$, and g is bounded and positive, so $\bar{t} = \mathbb{E}[\tau g(\tau)]/\mathbb{E}g(\tau) \le (w+1)\ell$.

A sign change. Let $\Psi^\circ(t) = \xi t\,\psi^\circ(t) - \int_0^t (\bar{t} - z)\psi^\circ(z)\,dz$ and $\beta_0(t) = \xi + t - \bar{t} + \xi t\mathcal{S}(t)$, with $\mathcal{S}$ the logarithmic derivative of ψ°. Then $(\Psi^\circ)' = \beta_0\psi^\circ$. Let $\varphi_0(y) = y/(\mathrm{e}^y - 1)$ and $\varphi_1(y) = y/(\mathrm{e}^y + 1)$, so that $t\mathcal{S}(t) = -t/\ell + w\varphi_0(t/m) - (m-w)\varphi_1(t/m)$. The function φ_0 does not increase, as the numerator $\mathrm{e}^y(1-y) - 1$ of its derivative vanishes at 0 and has derivative $-y\mathrm{e}^y \le 0$. Also $-\varphi_1'(y) \le 1/4$ for $y > 0$: this is equivalent to $4y \le \mathrm{e}^y + 6 + 5\mathrm{e}^{-y}$, and $\mathrm{e}^y + 6 - 4y \ge 7 - 3y + y^2/2 = (y-3)^2/2 + 5/2 > 0$. Hence

$$\beta_0'(t) = 1 - \frac{\xi}{\ell} + \frac{\xi}{m}\big[w\,\varphi_0'(t/m) - (m-w)\varphi_1'(t/m)\big] \le 1 - \frac{\xi}{\ell} + \frac{\xi}{4} = 0.$$

As $t\mathcal{S}(t) \to w$ for $t \downarrow 0$, $\beta_0(0+) = \xi(w+1) - \bar{t} \ge 0$, by the bound on the mean and $\xi > \ell$. As $\mathcal{S}(t) \to -1/\ell$ and $t(\mathcal{S}(t) + 1/\ell) \to 0$ for $t \to \infty$, $\beta_0(t) = \xi - \bar{t} + (1-\xi/\ell)t + o(1) \to -\infty$. So β_0 is first nonnegative and then nonpositive.

Conclusion. As $t\psi^\circ(t) \to 0$ at 0 and at ∞, and $\int_0^\infty (\bar{t} - z)\psi^\circ(z)\,dz = 0$, $\Psi^\circ(0+) = \Psi^\circ(\infty) = 0$. As Ψ° first increases and then decreases, $\Psi^\circ \ge 0$, which is $\vartheta(t) \le \xi t$. ∎

Lemma 38 (Bound for short means) *For $0 < \ell \le 2$ and $m \ge 2$, if all law instances hold at o, then $D_\mu \le 296 m^2 \gamma$.*

Proof Write $\psi = \psi_w^\circ$, $\vartheta = \vartheta_w$ and $\xi = \xi_\ell$, with $\xi \le 4$ as $\ell \le 2$. For every bounded smooth function φ on $(0,\infty)$,

$$\begin{aligned} \mathrm{Var}_\psi\,\varphi &= \int_{y<y'} (\varphi(y') - \varphi(y))^2 \psi(y)\psi(y')\,dy\,dy' \\ &\le \int_{y<y'} (y'-y) \int_y^{y'} \varphi'(t)^2\,dt\,\psi(y)\psi(y')\,dy\,dy' = \int_0^\infty \varphi'(t)^2 \vartheta(t)\psi(t)\,dt, \end{aligned} \tag{29}$$

by the Cauchy–Schwarz inequality and an exchange of integrals, with $\bar{t}$ the mean of ψ and

$$\int_0^t \int_t^\infty (y'-y)\psi(y)\psi(y')\,dy'\,dy = \int_0^t (\bar{t} - y)\psi(y)\,dy.$$

This is the weighted Poincaré inequality (Klaassen, 1985; Saumard, 2019). By Lemma 37, the right side is at most $\xi \int_0^\infty t\varphi'(t)^2\psi(t)\,dt$.

Fix b and t, and put $y = t/m$. Let L_i be the number of edits of variable i. If $b_i \neq b_i^0$, then L_i is Poisson with mean y, conditioned to be positive, and its variance is $y/(1-\mathrm{e}^{-y}) - y^2\mathrm{e}^{-y}/(1-\mathrm{e}^{-y})^2 \le y$. If $b_i = b_i^0$, its variance is $y/(1+\mathrm{e}^{-y}) + y^2\mathrm{e}^{-y}/(1+\mathrm{e}^{-y})^2 \le y + y^2\mathrm{e}^{-y} \le 2y$. These numbers are conditionally independent, so $L_t = \sum_i L_i$ has $\mathrm{Var}(L_t \mid b, t) \le 2t$. Given $\iota = t$, L_t is Poisson with mean t, and given the numbers of edits and B, the end state does not depend on t. So $\Pr(X = x, B = b \mid \iota = t) = \mathrm{e}^{-t}\sum_{n\ge 0}\chi_{b,x,n}t^n/n!$ with numbers $\chi_{b,x,n}$ that do not depend on t, namely the probabilities of $X = x$ and $B = b$ given n edits. As $t\,(t^n/n!)' = n\,t^n/n!$, the quotient rule for $g_{b,x}$ gives $g'_{b,x}(t) = t^{-1}\mathrm{Cov}(\mathbf{1}_{X=x}, L_t \mid b, t)$. By the Cauchy–Schwarz inequality,

$$\sum\nolimits_x g'_{b,x}(t)^2 \le \frac{\mathrm{Var}(L_t \mid b, t)}{t^2}\sum\nolimits_x g_{b,x}(t)(1-g_{b,x}(t)) \le \frac{2}{t}\,v_b(t).$$

Let $p_{b,x} = \Pr(X = x \mid B = b)$. By the law of total variance, $1 - \sum_x p_{b,x}^2 = \mathbb{E}[v_b(\iota) \mid B = b] + \sum_x \mathrm{Var}(g_{b,x}(\iota) \mid B = b) \le (1+2\xi)\mathbb{E}[v_b(\iota) \mid B = b]$. Averaging over b, $V \le 2\sum_j \bar{p}_j$ and $\mathbb{E}\bar{p}_j(\iota) \le 2\ell\gamma$ give

$$\sum\nolimits_b \Pr(B = b)\Big(1 - \sum\nolimits_x p_{b,x}^2\Big) \le 4m\ell(1+2\xi)\gamma.$$

For each label b with nonempty S_b, choose $f(b) \in S_b$ with the largest $p_{b,x}$, and let π_0 be the joint distribution of (X, B) and $\mathcal{G} = \{(f(b), b) : S_b \neq \varnothing\}$. The inequality at the end of Appendix G.4, with $p_{b,x}$ in place of $\bar{g}_{b,x}$, gives $1 - p_{b,f(b)} \le 1 - \sum_x p_{b,x}^2 + \sum_{x\notin S_b} p_{b,x}$, also when S_b is empty, with the convention $p_{b,f(b)} = 0$. Average this over b. The bound above takes care of the first two terms, and the last term averages to $\Pr(r(X) \neq B) \le m\ell\gamma$. So $\pi_0(\mathcal{G}^c) \le m\ell(5+8\xi)\gamma$. Place the states $f(b)$ in one cube as in Appendix G.6. Let $\mathcal{T}_{i,a}(x, b) = (T_{i,a}x, b[i \mapsto a])$. The geometric history gives $\pi_0 = (1-q)\delta_{(o,b^0)} + (q/(2m))\sum_{i,a}\pi_0\mathcal{T}_{i,a}$, so $\pi_0\mathcal{T}_{i,a} \le (2m/q)\pi_0$. If $(X, B) \in \mathcal{G}$ and $\mathcal{T}_{i,a}(X, B) \in \mathcal{G}$, then $T_{i,a}X = \tilde{T}_{i,a}X$. So each setter disagrees with probability at most $(1+2m/q)\pi_0(\mathcal{G}^c)$. With $\ell/q = 1+\ell$, $D_\mu \le m[2m + (2m+1)\ell](5+8\xi)\gamma$. As $\ell \le 2$ and $\xi \le 4$, this is at most $8m^2 \cdot 37\gamma = 296m^2\gamma$. ∎

G.8. Proof of Theorem 1(a)

The proof treats four cases: means $\ell \ge 2m$ by (8), means $\ell < 2m$ with a law instance failing at o by Lemma 6(iii), and, with all law instances holding at o, means $\ell \le 2$ by Lemma 38 and means $2 \le \ell \le 2m$ by (28).

Let $K = m^2[231504 + 9\log(1+2m)]$.

Proof of Theorem 1(a) As $m \ge 2$, the intervals $(0, 2]$, $[2, 2m]$ and $[2m, \infty)$ cover every mean $\ell > 0$. We show $D_\ell \le K\gamma_\ell$ in four cases. With $D_\ell \le 1$, this gives $D_\ell \le \min\{1, K\gamma_\ell\}$.

Means $\ell \ge 2m$. By Lemma 6(i), μ_ℓ is $(1+1/\ell)$-subinvariant, hence $(1+1/(2m))$-subinvariant, whether or not the laws hold at o. So (8) with $\alpha = 2$ gives $\bar{\lambda} \le (2\mathrm{e}m)^{1/2}$ and $D_\ell \le 8m\bar{\lambda}^2\gamma_\ell \le 16\mathrm{e}m^2\gamma_\ell$. As $16\mathrm{e} < 48$, this coefficient is at most K.

Means $\ell < 2m$ with a law instance failing at o. Then $o \in \mathcal{N}$, and Lemma 6(iii) gives $\gamma_\ell \ge \mu_\ell(\{o\}) \ge 1/(1+\ell)$. So $D_\ell \le 1 \le (1+\ell)\gamma_\ell \le (1+2m)\gamma_\ell$, and $1 + 2m \le 2m^2 \le K$ as $m \ge 2$.

Means $\ell \le 2$ with all law instances holding at o. Lemma 38 gives $D_\ell \le 296m^2\gamma_\ell$, and $296m^2 \le K$.

Means $2 \le \ell \le 2m$ with all law instances holding at o. These are the assumptions of Appendices G.1–G.6, with $\alpha = \ell/m \le 2$ and $\log(1+\ell) \le \log(1+2m)$. So (28) gives

$$D_\ell \le m^2[2 \cdot 38760 + 4 \cdot 38496 + 9\log(1+2m)]\gamma_\ell = K\gamma_\ell.$$

For D^∞ on every finite S, apply this to the interface of Lemma 20, which keeps the start state, the law instances at it and γ_ℓ. The tester follows from Appendix A. ■

G.9. Proof of Theorem 1(b)

Proof of Theorem 1(b) *A core.* Let S° consist of eight distinct states $x, w, y, v', v, \bar{x}, \bar{y}, \bar{w}$. For $m = 2$, the table in Figure 3 gives the readers and setters at x, w, y, v' and v. The states $\bar{x}, \bar{y}, \bar{w}$ read $(0, 0)$, and every setter fixes them. Let $\nu^\circ = (\delta_x + \delta_w)/2$, and call x and w its *atoms*.

Four entries that no composable interface keeps. We claim that no composable interface $\tilde{I}$, on S° or on a finite set containing S°, has $\tilde{T}_{1,1}x = y$, $\tilde{T}_{2,1}x = v'$, $\tilde{T}_{2,0}w = y$ and $\tilde{T}_{1,0}w = v$. In a composable interface, every setter keeps the cube of its argument. So the first and third entries put x, y and w in one cube, and the other two entries put v' and v in it as well. By the first and third entries, the label of y in $\tilde{I}$ starts with $(1, 0)$. As $x \neq y$, the first entry changes x, so the label of x starts with $(0, 0)$. As $w \neq y$, the label of w starts with $(1, 1)$ in the same way. So $v' = \tilde{T}_{2,1}x$ and $v = \tilde{T}_{1,0}w$ both have labels that start with $(0, 1)$. The four setters keep all other coordinates, so the labels of v' and v agree, and $v' = v$, a contradiction. Hence every composable interface, on S° or on a larger set, disagrees with I at x or at w, and $D^\infty_{\nu^\circ} \geq 1/2$.

The laws hold at x and w. Success, write-back and preservation at x and w follow from the table. Overwrite involves two writes to one variable. At x, these stay in $\{x, y\}$ for variable 1 and in $\{x, v'\}$ for variable 2, and at w, in $\{v, w\}$ and in $\{y, w\}$. On each of these pairs, the two setters of the variable act as an exact binary write, so overwrite holds. For commutation, both orders of $T_{1,a}$ and $T_{2,b}$ give the same state:

x	$b=0$	$b=1$
$a=0$	x	v'
$a=1$	y	w

w	$b=0$	$b=1$
$a=0$	x	v
$a=1$	y	w

.

So all $4m^2 + 3m = 22$ law instances hold at x and w, and $\gamma_{\nu^\circ} = 0$.

A composable interface at distance $1/2$. Give x, y, v', w the labels $00, 10, 01, 11$ in one cube, and $\bar{x}, \bar{y}, v, \bar{w}$ the same labels in a second cube. This composable interface keeps the readers at x and w. At x and w, its setters differ from those of I only in $T_{1,0}w = v' \neq v$. So $D_{\nu^\circ} = 1/2$.

More variables. For $m \geq 2$, let $S^\circ_m = S^\circ \times \{0, 1\}^{m-2}$. The first two variables act on the first component as above, the other variables are coordinate readers and setters of the second component, and labels concatenate. Let $\nu^\circ = (\delta_{(x,0^{m-2})} + \delta_{(w,0^{m-2})})/2$. The laws that involve only the first two variables hold as for $m = 2$, and those that involve only the others hold for coordinate readers and setters. Setters of the first two variables and of the others act on different components, so they commute and keep each other's readings. Hence $\gamma_{\nu^\circ} = 0$. The four entries at the atoms keep the second component, so the argument above gives $D^\infty_{\nu^\circ} \geq 1/2$. The product of the composable interface above with the interface that reads and replaces the coordinates of $\{0, 1\}^{m-2}$ changes only $T_{1,0}$ at $(w, 0^{m-2})$ on the support of ν°, so $D_{\nu^\circ} = 1/2$. The core S°_m has 2^{m+1} states.

One start state. Let $k \geq 1$, let $o_0 = (x, 0^{m-2})$ and $o_1 = (w, 0^{m-2})$ be the atoms of ν°, and add k extra states $o'_0, \ldots, o'_{k-1}$ with arbitrary readings. Every setter maps o'_j to o'_{j+1} for $j < k - 1$, and $T_{i,a}$ maps o'_{k-1} to o_a. With ν the point mass at o'_0, the last of k random edits reaches o_a when it writes a, so $\nu P^k = \nu^\circ$. Add $(-k) \bmod 2^m$ further states, with arbitrary readings and fixed by every setter, so that 2^m divides $|S|$. As S°_m is closed under the setters, the law instances at the atoms

involve only states of S_m°, and $\gamma_{\nu P^k} = \gamma_{\nu^\circ} = 0$. The lower bound concerns only the entries at the atoms, so it holds as before. The added states form full cubes, so the composable interface above extends to S with disagreement $1/2$. Thus $D^\infty_{\nu P^k} = D_{\nu P^k} = 1/2$. ■

Appendix H. Fixed thresholds: reduction to one quantity

For an integer $d \geq 1$, a state x is *locally exact to depth* d if every word of length at most d that starts at x ends at a state whose label is the one predicted by coordinate writes from $r(x)$, and two such words end at the same state exactly when their predicted labels agree. Let A_d be the set of these states, and let

$$a_d = \sup\{D^\infty_\nu(I) : I \text{ a finite interface of any number of variables, } \nu(A_d) = 1\}.$$

So $a_d \in [0, 1]$ does not increase with d. The *local exactness check* of depth d reads the end label of every word of length at most d from a sample, checks it against the predicted label, and checks for every two of these words that their ends are equal exactly when their predicted labels agree. It rejects with probability $\mu(S \setminus A_d)$.

Theorem 39 (Reduction) *For every interface and μ, $D^\infty_\mu \leq \mu(S \setminus A_d) + \mu(A_d)a_d$. The following are equivalent: (1) $a_d \to 0$; (2) for every fixed $\ell > 0$, every sequence of interfaces and start distributions with $m \to \infty$ and $\gamma_\ell = m^{-\omega(1)}$ has $D^\infty_\ell \to 0$; (3) the same for one fixed $\ell > 0$; (4) for all fixed $\ell > 0$ and $0 < \eta, \delta < 1$, law checks at samples of μ_ℓ certify $D^\infty_{\mu_\ell} < \eta$ from a number of samples polynomial in m, with degree and coefficient depending on ℓ, η, δ only.*

Lemma 40 (Passive embedding) *Let I be an interface with n_1 variables on S, and let $\nu(A_d) = 1$ with $d \geq 2$. For all $m \geq n_1$ and $\ell > 0$ there is an interface $\bar{I}$ with m variables and a start distribution with $D^\infty_\ell(\bar{I}) \geq D^\infty_\nu(I)/(1 + \ell)$ and $\gamma_\ell(\bar{I}) \leq (\ell n_1/(m + \ell n_1))^{d-1}$.*

Proof Let $\bar{S} = S \times \{0, 1\}^{m-n_1}$, $\bar{r}(s, v) = (r(s), v)$ and $\bar{\nu} = \nu \otimes \delta_{0^{m-n_1}}$. The first n_1 setters act on s, the others replace coordinates of v. Then $\bar{\mu}_\ell \geq \bar{\nu}/(1+\ell)$. A composable interface with m variables, with the last $m - n_1$ readers and setters forgotten, is a composable interface with n_1 variables: each m-dimensional cube splits into 2^{m-n_1} cubes by the last coordinates. Its set of states contains the slice $S \times \{0^{m-n_1}\}$, so it is a candidate for $D^\infty_\nu(I)$. This gives the distance bound. The number L_a of edits of the first n_1 variables is geometric with mean $\ell n_1/m$, as in Lemma 24. If $L_a \leq d - 2$, both sides of each law instance on these variables are words of length at most $L_a + 2 \leq d$ from the start state, which lies in A_d, so the law holds at the end of the run. Passive coordinates are exact and commute with all writes. So $\gamma_\ell \leq \Pr(L_a \geq d - 1)$. ■

Proof of Theorem 39 *The inequality.* If $\mu(A_d) = 0$, the right side is 1 and each disagreement probability is at most 1. Otherwise, let ν be μ conditioned on A_d. For every composable interface and every disagreement set X, $\mu(X) \leq \mu(S \setminus A_d) + \mu(A_d)\nu(X)$. Taking the largest X and the infimum over candidates gives the inequality. As $A_{d+1} \subseteq A_d$, a_d does not increase.

From laws to local exactness. If no state within $d-1$ writes of x lies in $\mathcal{N}$, then $x \in A_d$. Indeed, along a word of length at most d, success and preservation give the predicted labels. Commutation sorts the writes by variable, overwrite deletes all but the last write of each variable, and write-back at

x deletes a write of the current value of x. Each commutation or overwrite applies to a state reached by at most $d-2$ writes, and write-back applies at x after commutation moves the write to the front. So two words with the same predicted label reduce to the same word, and words with different predicted labels end at different labels. If μ is λ-subinvariant, $\mu(S \setminus A_d) \le \gamma_\mu \sum_{k<d} (2m\lambda)^k$, as the push-forwards of μ under all $(2m)^k$ words of length k sum to $(2m)^k \mu P^k \le (2m\lambda)^k \mu$.

(1) implies (2). Fix d and ℓ. As μ_ℓ is $(1+1/\ell)$-subinvariant, $\gamma_\ell = m^{-\omega(1)}$ gives $\mu_\ell(S \setminus A_d) \to 0$, so $\limsup D_\ell^\infty \le a_d$. Let $d \to \infty$. Choosing a fixed $\ell > 0$ in (2) gives (3).

(3) implies (1). Suppose $a_\infty = \lim_d a_d > 0$ and put $\eta = a_\infty/2$. For each $j \ge 2$, choose an interface I_j with n_j variables and ν_j with $\nu_j(A_j) = 1$ and $D_{\nu_j}^\infty(I_j) \ge \eta$. Fix $\ell > 0$ and choose increasing integers $m_j \ge \max\{j, n_j^2, (\ell n_j)^2, 4\}$. For $m \ge m_2$, let $j(m) = \max\{j : m_j \le m\}$, which tends to infinity, and embed $I_{j(m)}$ by Lemma 40. Then $D_\ell^\infty \ge \eta/(1+\ell)$ and $\gamma_\ell \le (\ell n_{j(m)}/m)^{j(m)-1} \le m^{-(j(m)-1)/2} = m^{-\omega(1)}$, against (3).

(1) implies (4). Fix ℓ, η, δ, choose d with $a_d \le \eta/2$ and put $\lambda = 1 + 1/\ell$. If $D_{\mu_\ell}^\infty \ge \eta$, the inequality gives $\mu_\ell(S \setminus A_d) \ge \eta/2$, so $\gamma_\ell \ge \eta/(2\sum_{k<d}(2m\lambda)^k)$. So $\lceil (2/\eta) \sum_{k<d} (2m\lambda)^k \log(1/\delta) \rceil$ samples suffice.

(4) implies (2). If (2) fails, there are a fixed $\ell > 0$, a fixed $\eta > 0$ and a sequence with $D_\ell^\infty \ge \eta$ and $\gamma_\ell = m^{-\omega(1)}$. A law-check tester with polynomially many samples rejects these with probability $o(1)$. ∎

Appendix I. Reads, measures and small supports

This appendix proves Propositions 41–45.

I.1. Reads along random writes

Let $\ell > 0$ and let $k \ge 1$ be an integer. Draw $X_0 \sim \mu_\ell$ and independent uniformly random setters $(i_1, a_1), \dots, (i_k, a_k)$, independent of X_0. Put $X_s = T_{i_s,a_s} X_{s-1}$, $B_0 = r(X_0)$ and $B_s = B_{s-1}[i_s \mapsto a_s]$, the label that a record predicts.

Proposition 41 (Reads) *For every interface and start distribution, the probability that $r(X_s) \ne B_s$ for some $s \le k$ is at most $\ell(q^{-k} - 1)\gamma_\ell$, which is at most $\mathrm{e}k\gamma_\ell$ for $k \le \ell$.*

Proof Take the first s with $r(X_s) \ne B_s$. As $r(X_{s-1}) = B_{s-1}$, the setter does not update the label as written at X_{s-1}, so success or preservation fails there, and $X_{s-1} \in \mathcal{N}$. Here X_{s-1} has distribution $\mu_\ell P^{s-1} \le q^{-(s-1)} \mu_\ell$ by Lemma 6(ii). Summing over s with $\sum_{s<k} q^{-s} = \ell(q^{-k} - 1)$ gives the bound. The last estimate of Corollary 10 gives $\ell(q^{-k} - 1) \le \mathrm{e}k$ for $k \le \ell$. ∎

Each step checks a whole label at once, so no union over the m variables appears.

I.2. Correct reads at constant distance

For $m \ge 2$, let $S = \{0,1\}^m \times \{0,1\}$ and $r(u, z) = u$, and let

$$T_{1,v}(u, z) = (u[1 \mapsto v],\ z \oplus (v \oplus u_1)u_2), \qquad T_{i,v}(u, z) = (u[i \mapsto v],\ z) \quad (i \ne 1). \tag{30}$$

For $m = 2$ this is the interface of Figure 2.

Proposition 42 (Eight states) *Let μ be the uniform measure on S. Success, preservation, write-back and overwrite hold at every state, so after any sequence of writes the label is the one that coordinate writes predict. Commutation fails at every state. Also $\mu P = \mu$ and $D_\mu^\infty \geq 1/8$.*

Proof Every setter updates the label as written, which gives success and preservation. For $v = u_1$ the change of z in (30) vanishes, which gives write-back. Writing w and then v to variable 1 changes z by $(w \oplus u_1)u_2 \oplus (v \oplus w)u_2 = (v \oplus u_1)u_2$, the change of writing v alone, which gives overwrite for this variable; for the other variables it is direct. At (u, z), the two orders of $T_{1,1-u_1}$ and $T_{2,1-u_2}$ give the same label, and their values of z differ by $(1 \oplus u_1 \oplus u_1)(u_2 \oplus 1 \oplus u_2) = 1$.

Each flip is an involution, as $F_1(u, z) = (u \oplus \mathbf{e}_1, z \oplus u_2)$ and the other flips change one coordinate of u. So each flip preserves μ. By write-back, the two setters of variable i average to $(\mathrm{Id} + F_i)/2$, so $\mu P = \mu$.

From $x = (u, z)$, apply F_1, F_2, F_1, F_2. Each of the two coordinates is first written with the opposite value and then with its initial value. The label returns to u, and z changes by $u_2 \oplus (1-u_2) = 1$. Let $\tilde{I}$ be a composable interface on a finite set $S' \supseteq S$. By Lemma 25, it disagrees with one of the four calls. Draw $x \sim \mu$. As the flips are involutions, the argument of each call is uniform. A call of variable i uses one of its two setters, so it disagrees with probability at most the sum of their disagreements, at most $2R_\mu(I, \tilde{I})$. So $1 \leq 8R_\mu(I, \tilde{I})$. ■

I.3. An audit at mean zero certifies the start distribution

Let $m = 1$, $S = \{x, y\}$, $\nu = \delta_x$, $r(x) = 0$ and $r(y) = 1$, and let $T_0 x = x$, $T_1 x = y$, $T_0 y = x$ and $T_1 y = x$.

Proposition 43 (Two states) *Every audit round with target label $\hat{a} = 0$ and depth $m + 1 = 2$ passes, with or without its bridge check, and $D_\nu = D_\nu^\infty = 0$. After one random write, $D_{\nu P} = D_{\nu P}^\infty = 1/2$.*

Proof *The round.* The path check has $T_0 x = x$, with the correct label, and the comparison with x passes. Write-back at x is $T_0 x = x$. The flip is $y_1 = T_1 x = y$, of label 1. Its path has $T_0 y = x$, of label 0, which equals $x_1 = x$, and the undo $T_1 x = y$ returns to y. For the bridge check, $u^{(1)} = a = 0$, and w_0 is x if $u = 0$ and y if $u = 1$, with the correct label. In both cases $T_0 w_0 = x = w_1$.

The distances. The record with labels $r(x) = 0$ and $r(y) = 1$, whose setters T_0 and T_1 always output x and y, agrees with the interface at x, so both distances under ν are 0. Now $\nu P = (\delta_x + \delta_y)/2$. This record differs only in $T_1 y$, which gives the upper bound $1/2$. A candidate with disagreement below $1/2$ would keep all readers and setters at both atoms. As $T_1 y = x$ and $r(x) = 0$, it would violate success, on the same states or with added states. ■

I.4. Lengthening the runs

Proposition 44 (Longer runs) *Let $m \geq 2$, $\ell \geq 0$ and $L \geq \ell$ with $L > 0$. Then $\mu_\ell \leq \frac{1+L}{1+\ell}\mu_L$, so $D_{\mu_\ell}^\infty \leq \frac{1+L}{1+\ell} D_{\mu_L}^\infty$, and also $D_{\mu_\ell} \leq \frac{1+L}{1+\ell} D_{\mu_L}$ when $2^m \mid |S|$. Let $L = \max\{\ell, m\}$. Then $\frac{1+L}{1+\ell} \leq 1+m$, and law checks at depth 2 at $\lceil 16\mathrm{e}^2 m^3 \frac{1+L}{1+\ell} \log(1/\delta)/\epsilon \rceil$ samples of μ_L certify $D_{\mu_\ell} < \epsilon$ when $2^m \mid |S|$, and $D_{\mu_\ell}^\infty < \epsilon$ for every finite S.*

Proof Put $q_\ell = \ell/(1+\ell)$. As $q_\ell \le q_L$, comparing (1) term by term gives $(1-q_\ell)q_\ell^s \le ((1-q_\ell)/(1-q_L))(1-q_L)q_L^s$, and $(1-q_\ell)/(1-q_L) = (1+L)/(1+\ell)$. The bound on the distances follows for each candidate and each disagreement set. For the tester, Corollary 21 with $\alpha = 1$ gives $D_{\mu_L} \le 16\mathrm{e}^2 m^3 \gamma_{\mu_L}$ when $2^m \mid |S|$, and $D^\infty_{\mu_L} \le 16\mathrm{e}^2 m^3 \gamma_{\mu_L}$ for every finite S. So a distance of at least ϵ under μ_ℓ gives $\gamma_{\mu_L} \ge \epsilon(1+\ell)/(16\mathrm{e}^2 m^3(1+L))$, and law checks reject as in Appendix A. If $\ell \ge m$ then $L = \ell$, and otherwise $L = m$, which gives the bound on the ratio. ■

I.5. Small supports with added states

Proposition 45 (Small support) *Let $m, N \ge 1$, $0 < \epsilon, \delta < 1$, and let μ have support of size at most N. A tester of depth at most one with one-sided error, $\lceil (N/\epsilon)\log(N/\delta)\rceil$ samples and at most $2mN$ setter calls certifies $D^\infty_\mu < \epsilon$ with confidence $1-\delta$. It needs only the bound N.*

Proof *The tester.* Draw the samples and remove repetitions by comparison, which leaves a set B_N. For each $x \in B_N$, call all $T_{i,a}x$ and read the labels of their outputs. Reject if $T_{i,r_i(x)}x \ne x$ or $r(T_{i,a}x) \ne r(x)[i \mapsto a]$ for some i, a. Among all held states, join x and $T_{i,a}x$ by an edge, and reject if a connected component has two distinct states of one label. The calls act only on samples, so the depth is at most one, and there are at most $2mN$ of them. A composable interface passes, as each setter keeps the cube and a cube has one state of each label.

A pass gives a composable interface. Suppose the tester passes. Each component has distinct labels. Give each component its own cube, with its states at their labels, and fill the other positions with added states. Give each state of S that is not held a cube of its own in the same way. This is a composable interface on a finite set containing S. It keeps the readers on S, and for $x \in B_N$ all setters at x, as $T_{i,a}x$ lies in the cube of x with the updated label. So $D^\infty_\mu \le \mu(S \setminus B_N)$.

Heavy atoms are sampled. Let A_N be the support of μ and $H_N = \{x \in A_N : \mu(x) \ge \epsilon/N\}$. With $s = \lceil (N/\epsilon)\log(N/\delta)\rceil$ samples, a state of H_N is missed with probability at most $(1-\epsilon/N)^s \le \delta/N$, so $H_N \subseteq B_N$ except with probability at most δ. On this event, $\mu(S \setminus B_N) \le \mu(A_N \setminus H_N) < \epsilon$. So an interface with $D^\infty_\mu \ge \epsilon$ passes with probability at most δ. ■