Beyond Interaction Scores: Selecting Feature Interactions via Independent Replication

Jincheng Gao
HKUST(GZ)
jgao627@connect.hkust-gz.edu.cn

Shuo Sun
HKUST(GZ)
shuosun@hkust-gz.edu.cn

arXiv:2611.13765v1 [cs.LG] 10 Nov 2026

***Abstract*—Learning feature interactions is essential for accurate prediction and interpretable modeling, yet a high interaction score may reflect fitting noise rather than a reproducible pattern. Ranking methods score feature pairs but leave the cutoff to the user; testing methods set the cutoff but rely on repeated model refits or a model of the joint feature distribution. We present LENS, a model-agnostic framework for selecting feature interactions through independent replication. LENS trains two models with the same learning algorithm on disjoint data and accesses them only through predictions. For each feature pair, it removes main effects from the joint-effect surface to isolate an interaction pattern, ranks pairs by its strength, and tests whether the second model reproduces it at a meaningful strength. A single multiple-testing correction determines which pairs to retain, using only two model fits regardless of the number of features. We prove that a replication ratio alone can retain arbitrarily weak patterns shared by both fits, whereas a strength floor restores selection consistency, and we characterize when normalization in Friedman's H^2 reverses interaction rankings. On 16 synthetic designs, LENS matches leading model-specific ranking methods (e.g., TreeSHAP) and recovers 91% of material interactions, against 45% for the knockoff-based DIAMOND, without selecting false pairs; with 1,225 candidate pairs and no interactions, it makes substantially fewer false discoveries. On large-scale French motor insurance data, it identifies interactions documented in the actuarial literature and recovers 90% of the deviance gain obtained using all pairs. On two credit datasets, LENS selects few pairs while maintaining or slightly improving predictive performance over the additive baseline.**

***Index Terms*—interpretable machine learning, feature interactions, generalized additive models, hypothesis testing**

I. Introduction

Two features interact when the effect of one depends on the value of the other. Learning insightful feature interactions is of vital importance for various forecasting and modeling tasks [1]–[3]. Fig. 1 illustrates one example in insurance pricing, where the effect of vehicle age can vary across brands [4]. The elevated claim frequency associated with a particular combination of vehicle age and brand motivates an explicit interaction term in an otherwise additive tariff. Practitioners can use a flexible predictive model to identify candidate interactions, then add selected interactions as explicit terms in an interpretable model [4]–[7]. The usefulness of this approach depends on choosing reliable interactions. Some patterns learned by the predictive model may arise from noise in its training sample and disappear when the model is trained on new data. This raises the question of which learned interactions are reproducible enough to retain.

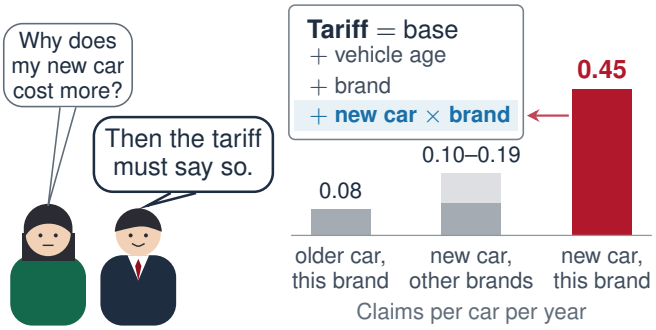

Fig. 1: Demonstration example on the importance of feature interaction in insurance pricing.

Selecting these interactions is challenging because p features give rise to $p(p-1)/2$ candidate pairs. Existing approaches fall into two broad groups, summarized in Fig. 3. Ranking methods assign a score to each pair using model internals, such as tree structures or neural network weights [8]–[10], or predictions alone [11]. They prioritize candidates but leave users to choose a cutoff, often by retaining the top k pairs [7], [9], [12]. A high score, however, can reflect fitting noise, so the ranking alone cannot determine which pairs have sufficient evidence for inclusion. Statistical testing methods address this selection problem, but typically require repeated model refits [11], [13]–[15] or a model of the joint feature distribution [16]. These requirements incur substantial computational cost or introduce distributional modeling assumptions, especially when many candidate pairs must be considered.

Independent replication offers a way to assess whether a learned interaction persists beyond a single training sample. If a pattern reflects a stable dependency captured by a learning algorithm, a model trained by the same algorithm on independent data should reproduce it. We build on this idea in LENS, a model-agnostic framework for selecting feature interactions through independent replication. LENS trains a discovery model and a confirmation model on disjoint data and accesses both only through predictions. It uses two model fits regardless of the number of features and performs the remaining computations through prediction queries.

Fig. 2 presents the procedure. For each feature pair, LENS constructs a joint-effect surface and removes the main effects of the two features to isolate their interaction pattern. The

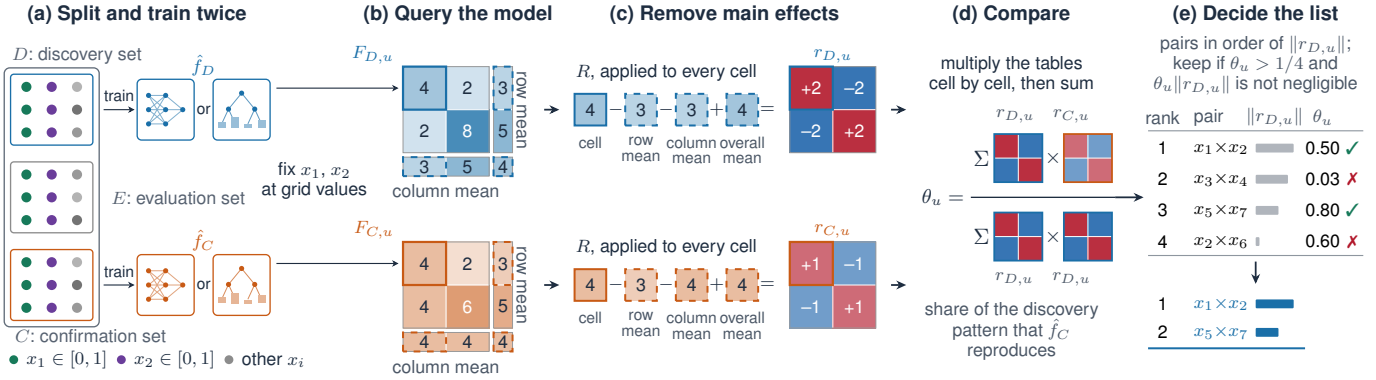

Fig. 2: Overview of LENS. (a) One data set is split into D, C, and E, and two models of any kind are trained on D and C. (b) With x_1 and x_2 fixed at grid values and the other features taken from real rows (D for $\hat{f}_D$, E for $\hat{f}_C$), the mean prediction fills the table F_u. (c) Removing the row and column means leaves the residual table r_u. (d) θ_u is the share of $r_{D,u}$ that $r_{C,u}$ reproduces. (e) Pairs ranked by $\|r_{D,u}\|$ are kept if the strength that $\hat{f}_C$ reproduces, $\theta_u\|r_{D,u}\|$, is significantly above both $\|r_{D,u}\|/4$ and a floor set by the typical pair strength and the spread of the predictions, after one Benjamini–Yekutieli correction over all pairs; $x_3 \times x_4$ is strong but not reproduced, and $x_2 \times x_6$ is reproduced but negligible. Red positive, blue negative.

Method	Models trained	Sets the list	No model of $p(x)$	Few fits
Ranking H^2 [11], TreeSHAP [8], NID [9]	one fit	✗	✓	✓
Refit tests permutation [13], [14], iLOCO [15]	many refits	✓	✓	✗
Knockoff tests DIAMOND [16]	one fit + a model of $p(x)$	✓	✗	✓
LENS	two fits, any p	✓	✓	✓

Fig. 3: How LENS differs from existing methods. Sets the list: the method decides which pairs to keep. No model of $p(x)$: it needs no model of the joint law of the features. Few fits: the number of fits does not grow with the number of pairs.

energy of the discovery pattern provides a score for ranking the pairs. Using a separate evaluation set, LENS then tests whether the confirmation model reproduces that pattern at a meaningful strength. A single multiple-testing correction determines which pairs to retain, potentially returning an empty set. The distinction between strength and replication is visible in the French motor insurance data. In the first split, VehAge×VehBrand is the strongest pair in the discovery model, and the confirmation model reproduces 78% of its pattern. VehAge×Region ranks fifth, but only 3% of its pattern is reproduced. LENS retains the former and rejects the latter.

Replication alone, however, is insufficient. Comparing independent fits removes the contribution of independent fitting variation in expectation, but both fits can share a systematic pattern induced by the learning algorithm. A replication ratio can consequently approach one even when the shared pattern is too weak to matter. We prove that a selection rule based on this ratio alone can retain arbitrarily weak shared patterns. LENS therefore also requires the reproduced interaction strength to exceed a floor. Under stated assumptions, this floor restores consistent selection of interactions whose limiting strength exceeds the threshold. Applying the Benjamini–Yekutieli correction [17] across all pairs further yields asymptotic false discovery rate control for the replication hypotheses under arbitrary dependence between pairs.

Our analysis also explains why the choice of ranking score matters. Friedman's H^2 divides interaction energy by the total variation of the joint-effect surface [11]. This normalization makes a pair's score depend on the magnitude of its main effects. A substantial interaction accompanied by strong main effects can receive a small score, while a weak interaction on a nearly constant surface can receive a large one. We characterize exactly when this normalization reverses the ranking of two pairs and establish consistency of the unnormalized interaction energy with a deterministic error bound. LENS uses this energy to rank candidates. On the same fitted XGBoost models, removing the normalization raises Recall@15 from 0.48 to 0.88, comparable to the tree-specific TreeSHAP score at 0.87. The ranking and replication test thus address complementary requirements of interaction discovery, quantifying the strength of a candidate and assessing whether an independent fit reproduces it sufficiently.

We evaluate LENS on 16 synthetic designs and three real-world datasets, covering diverse interaction structures, interaction-free settings, correlated features, and two model families. Its ranking accuracy matches leading model-specific methods. On the F1–F10 benchmarks, LENS with XGBoost recovers 91% of material interactions, those that move predictions by at least 2% of their spread, without selecting false

pairs, compared with 45% for DIAMOND with the same model family. In interaction-free settings with 1,225 candidate pairs, LENS makes substantially fewer false discoveries. On large-scale French motor insurance data, it identifies interactions documented in the actuarial literature and recovers 90% of the deviance improvement obtained by adding all candidate pairs to the additive model. On two credit datasets, it selects few pairs and matches or slightly improves on the additive baseline. Our contributions are threefold:

- **Interaction selection through independent replication.** We introduce LENS, a model-agnostic framework that ranks feature pairs and tests whether their interaction patterns replicate at a meaningful strength. It requires only two model fits and prediction access, with one multiple-testing correction determining the selected set.
- **Theory for replication and interaction ranking.** We prove that replication ratios alone can retain arbitrarily weak shared patterns and establish selection consistency with a strength floor under stated assumptions. We characterize ranking reversals caused by H^2 normalization, establish consistency of the unnormalized score with a deterministic error bound, and prove asymptotic false discovery rate control for the replication hypotheses.
- **Evaluation of ranking, selection, and predictive utility.** Experiments on synthetic and real-world data demonstrate competitive ranking accuracy, recovery of material interactions with few false discoveries, and the predictive utility of selected interactions in additive models.

The code and every run behind the tables and figures are available at https://github.com/JinchengGao-Infty/LENS.

II. Problem Formulation

We are given n rows with p features and a learning algorithm, such as gradient-boosted trees or a neural network, that fits a model $\hat{f} : \mathbb{R}^p \to \mathbb{R}$. The output of $\hat{f}$ is on the link scale, such as a log-rate or a logit. This section defines when a fitted model has a pair interaction and states the problem.

A. The interaction of a pair

A model takes all p features at once, so the interaction of a pair $u = (j, k)$ is defined on the model's joint effect of x_j and x_k: its partial-dependence surface [11]

$$F_u(a, b) = \mathbb{E}\big[\hat{f}(X \text{ with } X_j = a,\ X_k = b)\big], \qquad (1)$$

the mean prediction when the two features are set to a and b and the other features keep their distribution in the data. The pair has no interaction when the surface is additive, $F_u(a, b) = c + \alpha(a) + \beta(b)$: the effect of one feature does not depend on the other. Any departure from this form is an interaction.

Example 1. The left of Table I is additive, with $c = 100$, $\alpha = (0, -10, -30)$, and $\beta = (0, 20)$. Raising x_2 from 0 to 1 adds 20 in every row, so the effect of x_2 does not depend on x_1. The effects need not be straight lines: α is not. The right of the table changes one cell. Raising x_2 now adds 60 in the last row, so the effect of x_2 depends on x_1, and no choice of c, α, and β can produce this table.

TABLE I: Example 1. Left: an additive table. Right: a table with an interaction. The columns δ give the effect of raising x_2 from 0 to 1.

	Additive			Interaction		
	$x_2 = 0$	$x_2 = 1$	δ	$x_2 = 0$	$x_2 = 1$	δ
$x_1 = 1$	100	120	20	100	120	20
$x_1 = 2$	90	110	20	90	110	20
$x_1 = 3$	70	90	20	70	**130**	**60**

TABLE II: Notation

Symbol	Meaning
$u = (j, k)$, P	a pair of features; number of pairs, $p(p-1)/2$
D, C, E	discovery, confirmation, and evaluation parts
$F_u^{(i)}$, F_u	context surface of row i; its average
R	projection that removes both main effects
$r_{D,u} = RF_{D,u}$	residual table of the discovery model
$r_{C,u} = RF_{C,u}$	residual table of the confirmation model
$h_{C,u,i} = RF_{C,u}^{(i)}$	residual table of the confirmation model in context i
S_u, a_u	ranking score $\|r_{D,u}\|^2$; discovery strength $\|r_{D,u}\|$
$A_{u,i}$, θ_u	replication ratio in context i; its mean
$\mu_u = \theta_u a_u$	reproduced strength along the discovery pattern
s_D, σ_D	RMS strength over all pairs; spread of $\hat{f}_D$
η, κ, α	relative margin, scale floor, BY level
Δ_u, M_u	interaction and main-effect energy of a surface

B. The selection problem

A fitted model has some interaction on nearly every pair, even when the data have none, because it also fits noise. We call an interaction real when a model trained by the same algorithm on independent rows learns it again, at a strength that matters. A pair whose interaction is not real is a false discovery, and $\mathcal{H}_0$ denotes the set of such pairs; Eq. (6) states this hypothesis exactly, and Theorem 2 controls errors on it. The target is thus defined through the learning algorithm: an interaction learned on every sample is real for the model, even when the generating function puts it on a correlated pair.

Problem. Given the data and the learning algorithm, return a set $\widehat{G}$ of the $P = p(p-1)/2$ pairs, ordered by interaction strength, such that the false discovery rate $\mathbb{E}\big[|\widehat{G} \cap \mathcal{H}_0| / \max(|\widehat{G}|, 1)\big]$, conditional on the fitted models, is asymptotically at most α. The size of $\widehat{G}$ is decided by the data, and $\widehat{G}$ may be empty. The procedure should query the models only through their predictions and train a number of models fixed in p.

III. The LENS Method

LENS splits the data once into a discovery part D, a confirmation part C, and an evaluation part E, and trains two models of the same kind (Fig. 2(a)): $\hat{f}_D$ on D and $\hat{f}_C$ on C. The evaluation rows have no role in training; they supply the contexts in Eq. (2) for the test. LENS then measures every pair's interaction in each model (Section III-A), ranks the pairs by their strength in $\hat{f}_D$, tests whether $\hat{f}_C$ reproduces them, and decides the list with one correction over all pairs. Algorithm 1 summarizes the procedure; Table II lists the notation.

A. Measuring an interaction

To compute Eq. (1) for a pair $u = (j, k)$, LENS evaluates it on a grid of values for each feature: quantiles with equal weights for a continuous feature, levels weighted by their frequencies for a categorical one. Both are computed from the features of D, before $\hat{f}_C$ is trained. Write w_a and v_b for the weights of grid values a and b. The context surface of row x_i is the table

$$F_u^{(i)}(a, b) = \hat{f}\big(x_i \text{ with } x_j = a,\ x_k = b\big). \tag{2}$$

We set the two features to each pair of grid values and keep the rest of the row. Averaging over rows gives the partial-dependence surface $F_u = \frac{1}{m}\sum_{i=1}^{m} F_u^{(i)}$. Because the other features come from real rows, their correlations are preserved.

This table still contains what each feature does alone. To keep only the interaction, we remove these main effects: from every cell we subtract the mean of its row and the mean of its column, and add back the overall mean,

$$(RF)(a, b) = F(a, b) - \bar{F}(a, \cdot) - \bar{F}(\cdot, b) + \bar{F}(\cdot, \cdot), \tag{3}$$

where the means use the weights w_a and v_b. The residual table $r_u = RF_u$ (Fig. 2(c)) is the classical two-way interaction of the surface. To measure it, write $\langle F, F' \rangle = \sum_{a,b} w_a v_b F(a, b) F'(a, b)$ for the weighted average of the cell-by-cell products. The energy of the residual table is $\Delta_u = \|r_u\|^2 = \langle r_u, r_u \rangle$. An additive table $c + \alpha(a) + \beta(b)$ has $RF = 0$, so $\Delta_u = 0$ exactly when the surface is additive on the grid, the definition of Section II-A.

B. Ranking

The larger the residual table, the stronger the interaction. For every pair u, LENS computes the residual table of the discovery model, $r_{D,u} = RF_{D,u}$, with contexts drawn from D, and ranks pairs by its energy,

$$S_u = \|r_{D,u}\|^2. \tag{4}$$

It is the numerator of Friedman's H^2, computed on the grid rather than at the data points. H^2 divides it by the variance of the whole surface; Inglis et al. already drop this division [18], and Section VI adds an error bound for the ranking and the exact condition under which the division reverses it. The ranking orders the list; the test chooses the pairs.

C. The replication test

A positive S_u does not show that $\hat{f}_D$ learned something stable: a model fitted to noise also has small interactions. LENS asks whether the pattern $r_{D,u}$ reappears in the independently trained $\hat{f}_C$ (Fig. 2(d)). For each evaluation row $i \in E$, let $h_{C,u,i} = RF_{C,u}^{(i)}$ be the residual table of the confirmation model on that row, and define

$$A_{u,i} = \frac{\langle r_{D,u}, h_{C,u,i} \rangle}{\|r_{D,u}\|^2}, \qquad \theta_u = \mathbb{E}\big[A_{u,i}\big]. \tag{5}$$

$A_{u,i}$ compares the two tables cell by cell, and its mean θ_u measures how much of the pattern of $\hat{f}_D$ the model $\hat{f}_C$ contains. It is 1 when $\hat{f}_C$ reproduces the pattern exactly and 0 when $\hat{f}_C$ has nothing of it; a different pattern in $\hat{f}_C$ does not count. It can also exceed 1 or turn negative. Because R and $\langle \cdot, \cdot \rangle$ are linear, the mean of $A_{u,i}$ over the rows of E equals $\langle r_{D,u}, r_{C,u} \rangle / \|r_{D,u}\|^2$, where $r_{C,u} = RF_{C,u}$ is the residual table of $\hat{f}_C$ with contexts from E (Fig. 2(d)); the values per row are needed only for the standard error of the test.

TABLE III: Terms of the Boundary in Eq. (6)

Term	A pair must reproduce	This rules out
ηa_u	a quarter of its own strength	a strong pair that barely replicates
ηs_D	a quarter of the typical strength	a weak pair that replicates a negligible pattern
$\kappa \sigma_D$	2% of the spread of the predictions	a model that learned no interaction, where every table shrinks

The two models are trained on disjoint data, so their fitting errors are independent and cancel on average when the tables are compared. Only what both models learned survives. A pattern that only one model learned gives θ_u near 0; a pattern that both learned gives θ_u near the ratio of the two strengths.

A percentage, however, is blind to size. If both models carry the same faint pattern, θ_u is near one however faint the pattern is, and Section VI-C shows that a rule based on θ_u alone keeps such pairs even as the pattern vanishes. LENS therefore also asks how large the reproduced pattern is. Let $a_u = \|r_{D,u}\|$ be the strength of pair u in $\hat{f}_D$, and let $\mu_u = \theta_u a_u$ be the strength that $\hat{f}_C$ reproduces. LENS tests

$$H_{0,u}:\ \mu_u \le \max(\eta a_u, \eta s_D, \kappa \sigma_D), \quad \eta = 0.25,\ \kappa = 0.02, \tag{6}$$

where s_D is the root-mean-square strength over all pairs and σ_D is the standard deviation of the predictions of $\hat{f}_D$. Table III gives the role of each term; all are fixed before $\hat{f}_C$ is trained.

To test this with the rows of E, divide by a_u: the hypothesis becomes $\theta_u \le b_u$ with $b_u = \max(\eta a_u, \eta s_D, \kappa \sigma_D)/a_u$. Each row gives one value $A_{u,i}$, so this is a one-sample test. With $\bar{A}_u$ and s_u the mean and standard deviation of these values over the m rows,

$$Z_u = \frac{\sqrt{m}\,(\bar{A}_u - b_u)}{s_u}, \qquad p_u = 1 - \Phi(Z_u). \tag{7}$$

Two cases need no division: a pair with $a_u = 0$ has nothing to replicate and gets $p_u = 1$; if $s_u = 0$, $p_u = \mathbf{1}\{\bar{A}_u \le b_u\}$.

Example 2. Take two pairs in the first split of the French motor insurance data of Section V. On the log-rate scale, the discovery model has $\sigma_D = 0.591$ and $s_D = 0.040$, so $\eta s_D = 0.010$ and $\kappa \sigma_D = 0.012$, a 1.2% change in the claim rate. VehAge×VehBrand is the strongest pair, with $a_u = 0.151$; its own term dominates, and the boundary is $\eta a_u = 0.038$. The confirmation model reproduces $\mu_u = 0.118$ along the discovery pattern ($\theta_u = 0.78$), with a standard error of 0.0002 over 179,000 evaluation rows, so $Z_u \approx 370$. VehAge×Region is the fifth strongest pair, with $a_u = 0.060$ and boundary 0.015. The confirmation model reproduces $\mu_u = 0.0019$ ($\theta_u = 0.03$), so $Z_u \approx -270$ and $p_u = 1$. Both decisions lie far from the boundary. With this many evaluation rows, the standard errors

Algorithm 1 LENS

Require: Data split into D, C, E; a learning algorithm; grids; margins η, κ; level α.

Train $\hat{f}_D$ on D.
for every pair u **do**
  Compute $r_{D,u} = RF_{D,u}$ with contexts from D, and $S_u = \|r_{D,u}\|^2$.
end for
Set $s_D = (P^{-1}\sum_u S_u)^{1/2}$, σ_D the standard deviation of $\hat{f}_D$ over the contexts, and $b_u = \max(\eta a_u, \eta s_D, \kappa\sigma_D)/a_u$ with $a_u = \|r_{D,u}\|$.
Freeze all $r_{D,u}$ and b_u. Train $\hat{f}_C$ on C.
for every pair u **do**
  For each $i \in E$, compute $h_{C,u,i} = RF^{(i)}_{C,u}$ and $A_{u,i}$ (Eq. (5)).
  Compute p_u (Eq. (7)).
end for
Apply BY at level α to $\{p_u\}$ over all pairs.
return Rejected pairs, ordered by S_u; possibly none.

are small, and the list is set by where the boundary lies rather than by the level of the test (Section IV-D).

D. The decision: one correction over all pairs

Testing many pairs at once lets some pass by chance, so LENS applies one correction to the complete family. It computes p_u for all $p(p-1)/2$ pairs and applies the Benjamini–Yekutieli procedure [17] at level $\alpha = 0.05$. BY controls the false discovery rate under arbitrary dependence between the p-values, which matters here because all p-values share the same two models and the same evaluation rows. The rejected pairs, ordered by S_u, form $\widehat{G}$. If no p-value passes, $\widehat{G} = \varnothing$: LENS reports that the model has no interaction that replicates.

E. Cost

LENS trains two models, whatever the number of features. Everything else is prediction. A pair with G_j and G_k grid values needs G_jG_k predictions per context, so the ranking costs $m_D\sum_u G_jG_k$ predictions and the test costs $m_E\sum_u G_jG_k$, where m_D and m_E are the numbers of discovery and evaluation contexts. With 12 grid points per feature, 50 features give 1,225 pairs and 144 grid cells per pair; the procedure takes about four minutes on one GPU (Section IV-E). A refit test with empirical rank p-values multiplies the training cost by the number of refits, which must grow with P. After B refits, a rank p-value cannot fall below $1/(B+1)$, while the Benjamini–Yekutieli procedure at level 0.05 needs p-values below $0.05k/(mc_m)$ to keep k of m pairs, where $c_m = \sum_{i\le m} 1/i$. So B must exceed $mc_m/(0.05k) - 1$. Keeping three interactions among the 1,225 pairs of 50 features takes about 63,000 refits: about five days on the same GPU, not four minutes.

Implementation. LENS predicts all grid cells of a pair in one batched call, and the pairs are independent, so they can be spread over devices. Each context reduces to the scalar $A_{u,i}$, so the test stores only m_E numbers per pair. The procedure has no tuning loop: the grids, η, κ, and α are set once, and no step uses the outcome of the test to change an earlier step.

IV. Experiments on Synthetic Data

This section evaluates LENS on synthetic data with known interactions, and Section V applies it to real data. Section IV-A describes the models, designs, baselines, and metrics. Section IV-B first evaluates the residual energy as a ranking, against scores built for one model family and against H^2. Section IV-C then evaluates the lists that LENS returns, including designs without interactions and with correlated features, as the candidate pairs grow from 45 to 1,225. Section IV-D studies how the lists depend on the margins η and κ and on the level α, and Section IV-E reports the cost.

A. Setup

Models. The main model is XGBoost with 1,000 trees, learning rate 0.05, and depth 6, with early stopping on 10% of its training rows. To show that LENS needs only predictions, we repeat the synthetic study with the multilayer perceptron of DIAMOND [16], which has layers of 140, 100, 60, and 20 ELU units. LENS splits the data into equal discovery, confirmation, and evaluation parts, so each of its models sees a third of the rows. It uses 12 grid points per feature, 3,000 discovery contexts, every evaluation row as a test context, $\eta = 0.25$, $\kappa = 0.02$, and $\alpha = 0.05$. We fixed κ on three pilot seeds that the results below do not use.

Synthetic data. Each design has $n = 20{,}000$ rows of standard normal features, Gaussian noise with standard deviation 0.25, and $p \in \{10, 30, 50\}$ features, that is, 45, 435, or 1,225 candidate pairs. All features beyond the active ones are irrelevant. *Sparse* adds $1.5x_0x_1 + x_2x_3 + 0.75x_4x_5$ to the main effects $\sin x_6 + 0.5x_7^2 + 0.5x_8$. *Additive* keeps only these main effects. *Additive-corr* has main effects in x_0, x_1, and x_4, and makes x_2 and x_3 noisy copies of x_0 and x_1 with correlation 0.9. *Pair-corr* has the interaction x_0x_1, small main effects, and the same copies, so (x_0, x_3), (x_2, x_1), and (x_2, x_3) are proxies of the true pair and count as false. We also use F1–F10, the ten functions of the NID and DIAMOND benchmarks [9], [16], generated by DIAMOND's code with $p = 30$, $n = 20{,}000$, and no noise. Their true pairs follow DIAMOND's convention, which splits higher-order terms into pairs. All designs use 20 seeds.

Methods that decide. We compare LENS with five methods that return a list, each run with its published procedure. DIAMOND [16] trains a model on the features and their knockoffs and keeps the pairs that pass its knockoff filter at $q = 0.2$. We run its official code with its XGBoost model and with its MLP. Its lists barely change at $q = 0.05$ (FDP within 0.08), and on F1–F10, its own benchmark, its FDP of 0.07–0.10 stays below q, as its paper reports. H^2 with a refit null follows Friedman and Popescu [11]: 100 refits on responses simulated from an additive fit, then BY at 0.05 over all pairs [14]. With rank p-values, BY kept no pair in any of the 440 runs, as the count in Section III-E predicts, so we give H^2

TABLE IV: Ranking accuracy on twelve functions with known interactions. The last block uses reference implementations.

Method	Model	Recall@15	AP
LENS	XGBoost	**0.88**	**0.90**
TreeSHAP [8]	XGBoost	0.87	0.88
H^2 [11]	XGBoost	0.48	0.38
LENS	MLP	**0.95**	**0.96**
Integrated Hessians [10]	MLP	0.94	0.95
NID [9]	MLP	0.67	0.62
TreeSHAP-IQ [19]	-	0.87	0.87
Archipelago [20]	-	0.77	0.63
SHAP-IQ [21]	-	0.72	0.71
Faith-Shap [22]	-	0.56	0.49
hstats [23]	-	0.46	0.38
iRF [24]	-	0.44	0.22

a normal approximation to its null distribution. iLOCO [15] enters in both of its versions. iLOCO-Split refits the model without x_j, without x_k, and without both, for every pair; iLOCO-MP uses 10,000 minipatches of 20% of the rows and features. Both apply BY at 0.05 to normal p-values. FAST [5] ranks the pairs, and the validation loss of the resulting additive model with pair terms (GA^{2}M) chooses how many to keep from $K \in \{0, 1, 2, 3, 5, 8, 12, 20\}$.

Metrics. A ranking is scored by Recall@15, the share of the true pairs that rank in the top 15, and by average precision (AP), which averages the precision at the rank of each true pair and so scores the whole ranking. A list is scored against the generating function by its size, its false discovery proportion (FDP), the share of its pairs that are false, and its power, the share of the true pairs it contains. On the additive designs every pair is false, so we report the share of runs that return any pair. This FDP counts against the generating function, a different target from the replication hypotheses of Theorem 2: a proxy that both fits learn replicates but counts as false. Each metric is averaged over seeds. On real data, where the true pairs are unknown, a list is scored by the test loss of an additive model using it (Section V).

B. Ranking accuracy

The residual energy ranks interactions as accurately as scores that read the model's internals, and removing the denominator of H^2 is what makes the difference. Table IV reports Recall@15 and AP over twelve functions with known interactions: Friedman's benchmark, a smooth design, and F1–F10, each with $n = 10{,}000$ and 32 seeds. The model-specific scores run on the same fitted models as LENS. On XGBoost, LENS reaches a Recall@15 of 0.88, against 0.87 for TreeSHAP. On a softplus network it reaches 0.95, against 0.94 for Integrated Hessians and 0.67 for NID. H^2, computed on the same XGBoost models and grids, reaches 0.48.

The failure of H^2 comes from its denominator: the stronger the main effects of a pair, the lower H^2 ranks it (Proposition 3). We recorded both scores for every pair in 96 of these runs. Of the false pairs in the top 15 of H^2, 88% join two irrelevant features. Over 353,744 comparisons of a true pair with a false pair from the same run, the energy orders 99.4% correctly, and H^2 reverses 7.5% of these. In the reversed comparisons, the median ratio M/Δ is 736 for the true pair and 103 for the false one, which is Eq. (9). The true interactions sit on features with large main effects, and the denominator dilutes them.

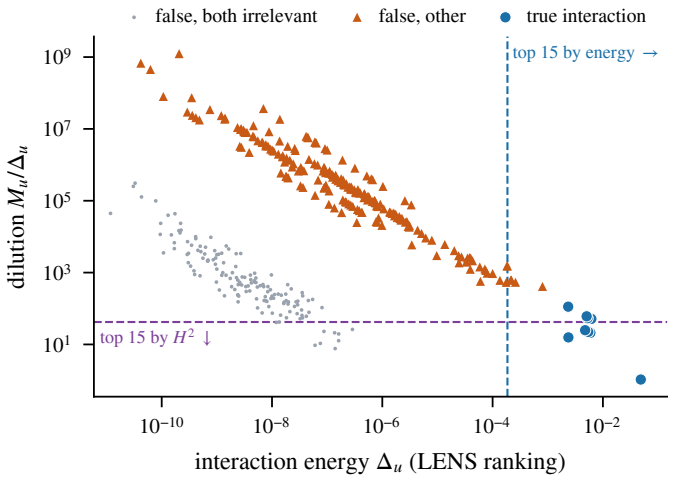

Fig. 4: The two scores in one run of F4 ($p = 30$, XGBoost). Each point is a pair; $H^2 = 1/(1 + M_u/\Delta_u)$ ranks by height alone and the energy by horizontal position alone. Dashed lines mark the top 15 of each score. The 71 pairs with $\Delta_u < 10^{-11}$, which the model treats additively, are omitted; none is true.

Fig. 4 shows the mechanism in one run of F4 with 435 pairs. Each pair is placed by its energy Δ_u, which orders the LENS ranking, and by its dilution M_u/Δ_u. Since $H^2 = 1/(1+M_u/\Delta_u)$, H^2 orders the pairs by the vertical axis alone. The ten true pairs have the ten largest energies, and the top 15 by energy contains all of them. They also carry large main effects, with M_u/Δ_u up to 112, so H^2 admits only five of them to its top 15. The other ten places go to pairs of two irrelevant features. Their energies are 10^{-8} to 3×10^{-7}, four to five orders of magnitude below the weakest true pair, but their faint surfaces are as non-additive in shape as the true ones, with M_u/Δ_u between 8 and 42. The other false pairs contain at least one relevant feature. They carry its main effect but almost no interaction, so they lie high on a line of slope -1, below the true pairs in both scores.

C. Deciding which pairs to keep

LENS is the only method whose lists are both complete and clean on the sparse design, and it returns no false pair on the benchmark functions (Table V). On the sparse design it returns exactly the three true interactions in 19 of 20 runs, and one extra pair in the twentieth; the same holds with 45 and 435 candidate pairs. With the neural network, it returns exactly the three in all 60 runs. DIAMOND finds 93% of the true pairs, and about one pair in six on its lists is false. The refit and leave-out tests find all three pairs but bury them:

TABLE V: Results on synthetic data, means over 20 seeds. The four designs use $p = 50$ and F1–F10 use $p = 30$. Red marks the best value in each column.

		Additive	Additive-corr	Sparse		Pair-corr		F1–F10	
Method	Model	FDP↓	FDP↓	FDP↓	power↑	FDP↓	power↑	FDP↓	power↑
LENS	XGBoost	**0.10**	0.05	**0.01**	1.00	**0.67**	1.00	0.00	**0.62**
LENS	MLP	0.00	**0.15**	0.00	1.00	0.03	1.00	0.00	**0.66**
DIAMOND	XGBoost	0.90	0.90	0.18	0.93	0.35	0.55	0.07	0.31
DIAMOND	MLP	0.65	0.90	0.14	0.93	0.12	0.20	0.10	0.29
H^2, normal refit null	XGBoost	1.00	1.00	0.96	1.00	0.99	1.00	0.72	0.96
iLOCO-Split	XGBoost	0.45	0.15	0.98	1.00	0.99	0.85	0.67	0.60
iLOCO-MP	trees	1.00	1.00	0.99	1.00	1.00	1.00	0.99	0.30
FAST	EBM	0.30	0.60	0.85	1.00	0.48	1.00	0.51	0.90

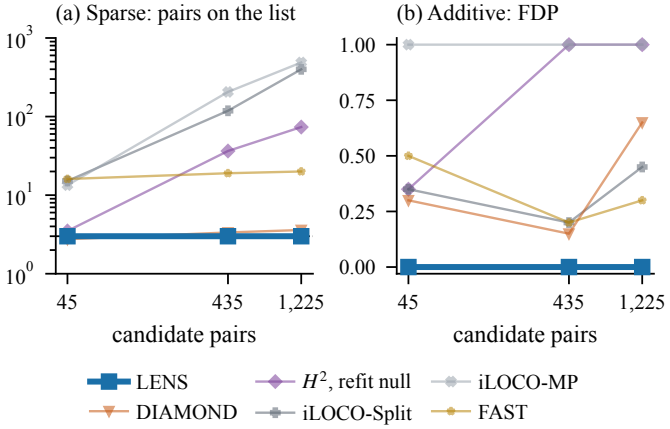

Fig. 5: The list of LENS keeps its length as the candidate pairs grow from 45 to 1,225; 20 seeds per point, with the same neural network for LENS and DIAMOND. (a) Mean list size on Sparse, which has three true pairs (dashed line). (b) FDP on Additive.

the list of H^2 has 74 pairs and that of iLOCO-Split 399, and more than 95% of each is false. FAST chooses its largest K in every run. These false pairs are not interactions of the model. Of those of H^2, 83% join two irrelevant features: their H^2 is a ratio of two near-zero energies, high in some fits and low in others, so the normal approximation to its null passes many of them (Section IV-B). So do 67% of those of iLOCO-Split. It tests, as designed, the loss contrast of the refits it has fitted, and for such a pair this contrast is a difference between four refits, which is not zero; the differences between fits thus count as interaction. FAST chooses K by validation loss, which rewards any pair term that lowers the loss. Of its false pairs, 83% contain one feature with a main effect, whose remaining curvature the pair term helps fit. Fig. 5 shows how the lists change as the candidate pairs grow from 45 to 1,225. LENS keeps its list length; the refit and leave-out tests do not. Appendix A reports every design at every number of features, for every method and decision rule.

Without interactions, LENS returns a pair in 7 of the 120 XGBoost runs on the two additive designs, never more than one. With the neural network it returns none on independent features, and one pair in 9 of 60 runs on correlated features. DIAMOND with XGBoost returns pairs in 90% of the runs at $p = 50$, and in 40% already at $p = 10$. In 96% of its false pairs, both features have main effects: the fitted model spreads a little of the nonlinear main effects $\sin x_6$ and $0.5x_7^2$ over their pair, and the knockoff copies, which carry no signal, cannot match it. The few false pairs of LENS on Additive are of the same kind. H^2 and iLOCO-MP return pairs in every run. For iLOCO-MP the leave-out score is not zero even when the function is additive. Removing a feature changes how often every other feature appears in the remaining minipatches, and with it the shrinkage of the ensemble; even 100,000 minipatches did not shorten its lists.

On F1–F10, LENS returns no false pair on any function with either model, and its power is 0.62 with XGBoost and 0.66 with the network. The pairs it misses are weak. The benchmark labels a pair as true when its features share a term of the formula, whatever the size of their interaction. We computed the strength of every true pair from the generating function on the same grids, in units of the standard deviation of the function. Of the 104 true pairs, 40 fall below 2%, and 38 of them come from splitting a higher-order term. The other two have no interaction at all, because their terms are additive on the unit cube. Against the 64 pairs above 2%, LENS finds 91% with XGBoost and all of them with the network, while DIAMOND finds 45–47% (Fig. 6). The 2% cut equals κ, but the gap does not depend on it: above 1%, LENS finds 76–81% and DIAMOND 38%, and over all true pairs Table V shows the same factor of two. Both LENS models lie beyond the Pareto front of the baselines: every baseline has a higher FDP, and none with an FDP below 0.5 finds more than half of the material pairs. By strength, the cut is sharp. With the network, LENS detects every pair above the floor and none below it; with XGBoost it misses some pairs just above the floor. DIAMOND misses a third of the strongest pairs.

FAST and H^2 find all material pairs, but half or more of their lists are false. The scale floor certifies the interactions that move predictions by at least 2% of their spread, and the weak pairs fall below it as intended.

The row $\kappa = 0$ of Table VI shows why the scale floor is needed. Without it, LENS on XGBoost finds every material pair of F1–F10, but its FDP rises to 0.24, and on the additive designs it returns a list in every run with 435 or more pairs.

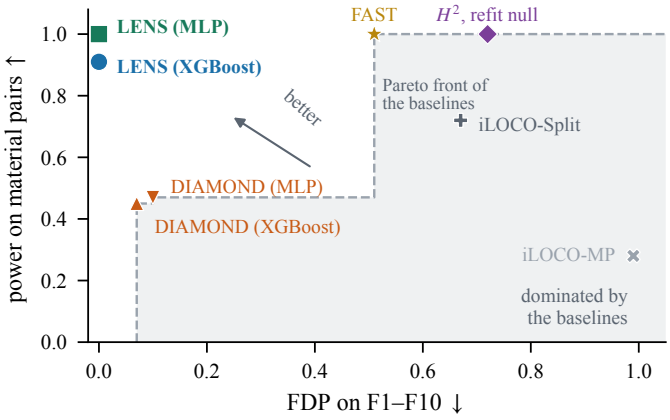

Fig. 6: FDP and power on the material pairs of F1–F10, means over the ten functions and 20 seeds. The dashed line is the Pareto front of the baselines; LENS lies beyond it.

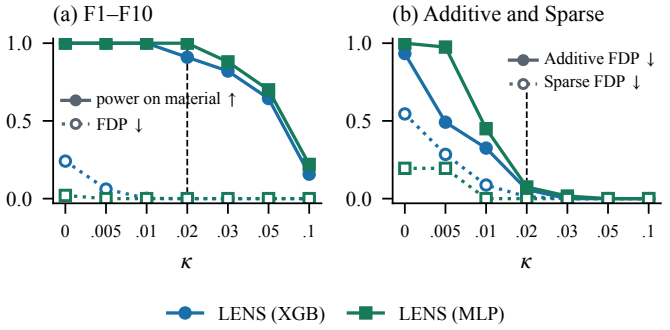

Fig. 7: The scale floor trades power for clean null lists. Sparse uses $p = 50$. The dashed line marks the default $\kappa = 0.02$.

The neural network shows the failure of Section VI-C in full. Its tables are smooth, so the two fits reproduce each other's artifacts closely, and on the additive design the relative terms keep 16, 146, and 403 pairs on average at $p = 10$, 30, and 50. Yet on F1–F10 the same rule is nearly exact, with an FDP of 0.02. The relative terms judge each pair against the others, so they fail exactly when the model has no interaction to compare with. The false pairs they keep on the additive designs reproduce a median of 0.08% (XGBoost) and 0.16% (network) of the prediction spread, and 99% of them fall below the floor. The true pairs of Sparse reproduce a median of 24%. For the network, the scale floor removes the failure and changes nothing on F1–F10 except the weak pairs.

The second model is needed as much as the floor. In the settings of Table V, keeping the pairs whose discovery strength a_u clears the same floor, without $\hat{f}_C$, returns about three false pairs in every additive XGBoost run, doubles the Sparse list to 6.1 pairs (FDP 0.51), and raises the FDP on F1–F10 from 0 to 0.06, at the same power. With the network it returns pairs in half of the Additive-corr runs, against 15% for LENS, and its Pair-corr FDP is 0.18, against 0.03. Requiring both fits to clear the floor, without comparing their tables, still returns pairs in every additive XGBoost run and gives a Sparse FDP of 0.49: the second fit helps only through the comparison. The floor removes what is too weak to matter, and replication what only one fit learned. Appendix B removes each ingredient of LENS in turn, on synthetic and real data: without the comparison the false pairs return on every design and the credit lists grow three- to fifteenfold at a higher loss; the energy ranking cannot set the length of the list by itself; and without the test the lists barely change at these sample sizes but lose the error control of Theorem 2.

Correlated proxies remain the hard case for every method. On Pair-corr, LENS with XGBoost returns the true pair with two proxies in 19 of 20 runs, an FDP of 0.67. With the network it returns the true pair alone in 56 of 60 runs. DIAMOND misses the true pair in 45–80% of the runs, and 12–35% of its pairs are false. The lists of the refit and leave-out tests are almost entirely false. As Section VI-C explains, the XGBoost proxies are interactions of the fitted model, reproduced at about a seventh of the true pair's strength, so a test of the model's interactions correctly reports them.

The test itself is conservative. In a simulation that draws the ratios $A_{u,i}$ directly, with every null pair at or below its boundary b_u, the false discovery rate stays below 0.007 at a nominal level of 0.05 across 24 settings: 36 or 253 pairs, 128 to 2,048 contexts, Gaussian, t_5, and log-normal noise, and independent and correlated pairs. Most of the slack comes from BY, whose bound holds under any dependence.

D. Sensitivity to the margins and the level

The list depends on the scale floor κ, little on the relative margin η, and hardly at all on the level α. Table VI re-decides every run under other settings. This needs no new fits, because the reproduced strength $\hat{\mu}_u = a_u \bar{A}_u$ and its standard error do not depend on the margins or the level; at the defaults the recomputed lists match all 440 stored XGBoost lists.

The scale floor. The floor trades power for clean null lists, and the trade is cheap because false and true pairs lie far apart: without the floor, the false pairs reproduce about 0.1% of the prediction spread and the true pairs of Sparse 24%. Below $\kappa = 0.02$, the null designs leak. At $\kappa = 0.01$, XGBoost returns a pair in a third of the additive runs and the network in 45%, and the network keeps the proxies of Pair-corr. Above 0.02, the lists stay clean but lose the weaker true pairs: power on the material pairs of F1–F10 falls from 0.91 to 0.82 at $\kappa = 0.03$ and to 0.64 at 0.05. The default is the smallest value in the grid at which both models return a pair in fewer than one additive run in ten and recover the sparse design exactly in at least 95% of runs. We chose it on pilot seeds, and the table shows that the choice carries over to the evaluation seeds. At $\kappa = 0.1$, the XGBoost proxies of Pair-corr fall below the floor as well, but so do 84% of the material pairs of F1–F10: a floor removes a proxy only by giving up the true interactions that are as weak as the proxy. Fig. 7 plots the trade-off. On real data, the list shrinks steadily with κ, and on the credit data, where the black box adds little, it is nearly empty at 0.03.

TABLE VI: Sensitivity of LENS, one parameter at a time from the defaults ($\kappa = 0.02$, $\eta = 0.25$, $\alpha = 0.05$). F1–F10 power counts the material pairs. Additive FDP pools $p = 10, 30, 50$; Sparse and Pair-corr use $p = 50$, with power 1.00 throughout.

	F1–F10 FDP↓		F1–F10 power↑		Additive FDP↓		Sparse FDP↓		Pair-corr FDP↓		List size		
Setting	XGB	MLP	XGB	MLP	XGB	MLP	XGB	MLP	XGB	MLP	freMTPL2	HELOC	UCI
$\kappa = 0$	0.24	0.02	1.00	1.00	0.93	1.00	0.54	0.20	0.78	0.90	13.0	23.2	35.2
$\kappa = 0.005$	0.06	0.00	1.00	1.00	0.49	0.97	0.28	0.20	0.71	0.90	13.0	19.0	33.0
$\kappa = 0.01$	0.00	0.00	1.00	1.00	0.33	0.45	0.09	0.00	0.70	0.36	13.0	10.4	17.0
$\kappa = 0.02$ (LENS)	0.00	0.00	0.91	1.00	0.06	0.07	0.01	0.00	0.67	0.03	12.0	2.0	5.2
$\kappa = 0.03$	0.00	0.00	0.82	0.88	0.00	0.02	0.00	0.00	0.67	0.00	9.0	0.2	3.6
$\kappa = 0.05$	0.00	0.00	0.64	0.70	0.00	0.00	0.00	0.00	0.56	0.00	5.2	0.0	1.0
$\kappa = 0.1$	0.00	0.00	0.16	0.22	0.00	0.00	0.00	0.00	0.03	0.00	1.2	0.0	0.4
$\eta = 0$	0.00	0.00	0.91	1.00	0.06	0.07	0.01	0.00	0.67	0.03	12.4	2.0	5.2
$\eta = 0.1$	0.00	0.00	0.91	1.00	0.06	0.07	0.01	0.00	0.67	0.03	12.4	2.0	5.2
$\eta = 0.5$	0.00	0.00	0.91	1.00	0.06	0.07	0.00	0.00	0.64	0.03	6.4	1.4	4.4
$\alpha = 0.01$	0.00	0.00	0.91	1.00	0.06	0.07	0.01	0.00	0.67	0.03	12.0	2.0	5.0
$\alpha = 0.2$	0.00	0.00	0.91	1.00	0.06	0.07	0.01	0.00	0.67	0.03	12.0	2.0	5.2

The relative margin. On the synthetic designs, η between 0 and 0.5 changes almost nothing. Weak pairs face the scale floor, which is above ηs_D, and strong true pairs reproduce far more than half of their strength. On freMTPL2, raising η to 0.5 halves the list, from 12 pairs to 6.4. There the real interactions reproduce between half and all of their strength (Table VIII), so η decides how complete a replication must be. The default asks for a quarter, which lies on the flat part: lowering η to 0 changes the list by at most half a pair on average.

The level. Changing α from 0.01 to 0.2 changes no synthetic list, and on the real data it removes one pair from one split of UCI credit at $\alpha = 0.01$. Example 2 explains why: with thousands of evaluation rows, the standard error of $\hat{\mu}_u$ is small compared with the distance of most pairs from their boundary, so the decision is made by the hypothesis, not by the level. With few evaluation rows the level matters, and the simulation of Section IV-C shows that the test then stays conservative, down to 128 rows. In practice, the choice that matters is where the boundary lies, which is why LENS states it on the link scale, where an analyst can set it from the size of an effect that matters to the application.

E. Cost

Two fits are enough for a test over any number of pairs, and a refit test with rank p-values cannot match this at any affordable number of fits. A refit test draws its p-values from its refits, so they cannot fall below $1/(B+1)$. With 100 refits, BY kept no pair of H^2 in any of the 440 runs (Section IV-A). Each refit takes 7 seconds at $p = 50$, so the 63,000 refits that Section III-E requires would take about five days on one RTX 5090. LENS draws its p-values from the evaluation rows instead, and two fits give p-values far below any refit floor: the strongest pair on freMTPL2 has $Z_u \approx 370$ (Example 2).

The time of LENS is set by prediction queries, which grow linearly with the pairs and run in parallel. A run on Sparse takes 36, 63, and 252 seconds at $p = 10$, 30, and 50, so 1,225 pairs are tested in about four minutes. At $p = 50$ LENS makes about 1.7×10^9 predictions, $1{,}225 \times 144 \times 9{,}600$, at about 7 million per second including both fits. The refit methods are slower even with far too few refits for BY: H^2 takes 35, 197, and 730 seconds, and iLOCO-Split, which refits for every pair, 39, 667, and 1,838. Methods without a refit null finish in 92 (DIAMOND), 139 (iLOCO-MP), and 144 seconds (FAST) at $p = 50$, and their lists are those of Table V; DIAMOND also needs a generator for the joint law of the features. The cost carries over to real data. On freMTPL2, with 180,000 policies per part, a run takes 3.3–4.5 minutes, mostly for the test of 179,000 evaluation rows; on credit data, under three.

V. Experiments on Real Data

We apply LENS to one insurance dataset, where actuaries have reported interactions, and to two credit datasets, where the black box adds little to an additive model.

A. Setup

Data. freMTPL2 has 678,013 French motor policies with nine rating factors, so 36 pairs. We model claim counts with a Poisson XGBoost and an offset of log exposure, and cap counts at 4 and exposure at 1, as is common for this dataset; the four categorical factors use XGBoost's native splits. LENS therefore tests interactions on the log claim rate, the scale on which the tariff is additive. HELOC has 10,459 credit applications and UCI credit has 30,000 card holders, each with 23 features and 253 pairs; both use a logistic XGBoost. Each dataset is split five times. A fifth of the rows is held out as a test set, and the rest is divided into the three parts of LENS.

Measuring a list. An additive model of boosted depth-one trees is fitted on the discovery part. Each pair on a list enters as one term: its discovery residual table, frozen as a lookup function, times a coefficient. The coefficients are fitted jointly on the confirmation part, and we report the loss on the test set. This is a generalized linear model (GLM) with fixed interaction shapes, a deliberately simple use of the list; a GA^{2}M would refit the shapes. We compare the LENS list with all pairs and with the top K pairs by energy for a fixed K, 7 on freMTPL2 and 10 on the credit data, which is what a ranking without a test leaves to the user. With 20,000 policies the LENS list has nearly this length (Section V-B).

TABLE VII: Test loss on real data, mean over five splits: Poisson deviance on freMTPL2, log loss on credit data. Red: lowest loss among the lists. No floor: LENS with $\kappa = 0$.

	freMTPL2		HELOC		UCI credit	
List	pairs	loss↓	pairs	loss↓	pairs	loss↓
None (additive)	0	0.3128	0	0.5508	0	0.4312
LENS	12.0	**0.3095**	2.0	0.5508	5.2	0.4303
No floor	13.0	0.3095	23.2	0.5541	35.2	0.4318
Top K by energy	7	0.3108	10	0.5518	10	0.4303
All pairs	36	0.3091	253	0.5915	253	0.4421
Black box		0.3055		0.5633		0.4325

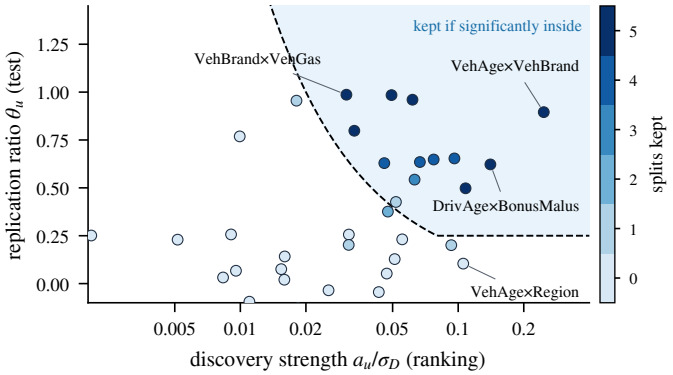

Fig. 8: The 36 pairs of freMTPL2, averaged over five splits: discovery strength, which orders the ranking, against the replication ratio, which the test uses. The dashed line is the boundary of the null hypothesis, $\theta_u = \max(\eta, \kappa\sigma_D/a_u)$; here ηs_D lies below the scale floor. Color gives the number of splits in which LENS keeps the pair.

B. Insurance data

Agreement with the actuarial literature. The interactions LENS keeps on freMTPL2 are those that actuaries have reported. Eleven pairs enter its list in at least four of the five splits. They include all five interactions that Schelldorfer and Wüthrich identified [4]: VehPower×VehAge, VehPower×VehBrand, VehAge×VehBrand, VehAge×VehGas, and DrivAge×BonusMalus. They also include all five pairs to which the boosted model of Havrylenko and Heger gives a nonzero H [6], which adds DrivAge×Region, BonusMalus×VehBrand, and VehAge×BonusMalus. The two studies share only two pairs, and the stable list contains the pairs of both. The remaining three pairs, VehPower×DrivAge, VehBrand×VehGas, and DrivAge×VehBrand, appear in neither study. Table VIII lists every pair that LENS keeps in some split. The stable pairs are not simply the strongest: VehPower×DrivAge and VehBrand×VehGas rank 20th on average, but the confirmation model reproduces 80% and 99% of their patterns, and both exceed the floor in every split.

The test also explains the pairs that the literature disputes (Fig. 8). Havrylenko and Heger's neural network ranks the pairs of BonusMalus and VehAge with the regional factors high, and their AIC-based selection step recommends

TABLE VIII: Pairs kept on freMTPL2 in at least one of five splits, then the two strongest pairs never kept. Kept: splits that keep the pair; rank (by energy), a_u/σ_D, and θ_u: means over splits; range: lowest and highest θ_u. S&W [4], H&H [6]: reported.

Pair	kept	rank	a_u/σ_D	θ_u	range	S&W	H&H
VehAge×VehBrand	5	1.0	0.247	0.90	0.78 to 1.06	✓	
DrivAge×BonusMalus	5	2.2	0.141	0.62	0.43 to 0.89	✓	✓
DrivAge×Region	5	4.4	0.108	0.50	0.31 to 0.67		✓
VehAge×VehGas	5	11.4	0.062	0.96	0.47 to 1.45	✓	✓
VehPower×VehAge	5	13.0	0.049	0.98	0.55 to 1.46	✓	
VehPower×DrivAge	5	19.6	0.033	0.80	0.54 to 1.35		
VehBrand×VehGas	5	20.0	0.031	0.99	0.59 to 1.59		
VehAge×BonusMalus	4	5.2	0.096	0.65	0.23 to 1.19		✓
BonusMalus×VehBrand	4	8.6	0.077	0.65	0.08 to 1.00		✓
DrivAge×VehBrand	4	9.2	0.067	0.63	0.22 to 1.20		
VehPower×VehBrand	4	14.6	0.046	0.63	0.41 to 0.96	✓	
VehAge×DrivAge	3	10.8	0.063	0.54	0.20 to 1.03		
DrivAge×Density	2	15.4	0.048	0.38	−0.13 to 0.82		
BonusMalus×Region	1	5.4	0.093	0.20	0.00 to 0.44		
BonusMalus×Density	1	14.2	0.052	0.43	−0.12 to 1.30		
VehAge×Region	0	4.6	0.105	0.10	0.03 to 0.21		
VehBrand×Region	0	11.8	0.055	0.23	0.12 to 0.33		

BonusMalus×Region [6]. Our discovery models agree on the strength: VehAge×Region is the fourth strongest pair on average and is in the top five in four of five splits, so a top-5 cutoff, as in the pipeline of Lindholm and Palmquist [7], would keep it. The confirmation model, however, reproduces at most 21% of its pattern, and its table never significantly lowers the Poisson loss on held-out policies. LENS never keeps it, and keeps BonusMalus×Region in one split.

What the tables say. The residual tables are themselves readable. In the first split, the strongest interaction sits in one cell. New cars of brand B12, which make up 64% of the new cars in the portfolio, have a residual of $+0.85$ on the log rate, so their predicted claim rate is 2.3 times what the row and column effects give. The raw data agree (Fig. 1): new B12 cars have 0.45 claims per policy-year, against 0.08 for older B12 cars and 0.10 to 0.19 for new cars of other brands. This one cell holds 64% of the pair's energy, and the confirmation model reproduces it at $+0.70$. DrivAge×BonusMalus has the shape that motivates this term in tariffs. In both tables, a bonus-malus level of 95 raises the rate more for drivers under 36 than the additive model allows, and less for drivers over 43, while the level 60 shows the opposite trend. In VehAge×BonusMalus, the bonus-malus scale is flatter for new cars and steeper for cars older than 14 years, again in both tables. The dropped pair looks different. VehAge×Region puts 62% of its energy into the row of new cars, where the discovery model gives each region its own correction, of either sign. The confirmation model has corrections of the same size in that row, but they do not line up with the first: the two tables agree in R21, a region with 0.4% of the policies, and disagree in most others, so the confirmation model reproduces 3% of the pattern ($\theta_u = 0.03$). Each model has learned a region effect for new cars from the few hundred to few thousand new cars per region in its part of the data, and the two effects differ. This is the pattern the test is designed to reject, and an energy ranking alone cannot

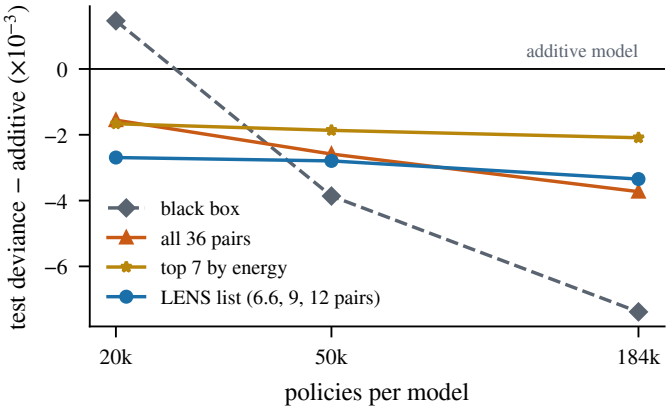

Fig. 9: freMTPL2 test deviance relative to the additive model, by the number of policies used to fit each model (mean over five splits). Lower is better. The legend gives the mean length of the LENS list at each size.

tell it apart from the kept pairs.

Accuracy. The list carries most of the gain of all pairs into the additive model (Table VII). The 12 pairs that LENS keeps on average close 45% of the deviance gap between the additive model and the black box, and 90% of what all 36 pairs close (85–96% across splits). The 7 strongest pairs close only 28%, so the length that the test sets matters. The value of the list is clearest when data are scarce. With 50,000 policies per model, the 9 pairs that LENS keeps match all 36 on average. With 20,000, the black box (0.3173) is worse than the additive model (0.3159), yet the 6.6 pairs extracted from it lower the deviance to 0.3132, below both and below all 36 pairs (0.3143), in every split. Fig. 9 shows the three sizes together. The black box gains the most from more data, all 36 pairs gain steadily, and the LENS list gains while staying short. With few policies the two fits can confirm only the strongest interactions, so the list holds 6.6 pairs; with more, it grows to 9 and 12. It grows by adding pairs: 73% of the pairs kept with 20,000 policies are kept with all of them, and 82% of those kept with 50,000. The test thus adapts the length of the list to the evidence, which a fixed K cannot do. The top 7 by energy does worse than the LENS list at every size. At 20,000 policies the two lists have nearly the same length but share only two pairs on average. The pairs that only the top 7 contains have replication ratios between -0.23 and 0.23; BonusMalus×Region is among the four strongest pairs in every split, yet its ratio never exceeds 0.23. With little data, the strongest patterns of one fit are largely its own noise. At the same length, the energy ranking alone is worse than LENS at every size and in every split (Table XVII).

C. Credit data

When the black box has little to give, LENS returns a short list, and the scale floor is what keeps it short. On both credit datasets the black box is no better than the additive model. LENS keeps 2 pairs on HELOC and 5 on UCI credit. On HELOC the log loss is unchanged; on UCI credit it drops by 0.0009, in four of five splits, and the additive model with its pairs beats the black box. Without the floor, the test keeps 23 and 35 pairs, and these lists raise the log loss, on HELOC in every split. Here the floor $\kappa\sigma_D$ is four to five times ηs_D, whereas on freMTPL2 it exceeds ηs_D by only 12–28% and removes one pair on average. The short lists vary across splits, with a Jaccard similarity of 0.05 and 0.21, as expected when no interaction is strong; only PAY_0×BILL_AMT1 on UCI credit enters in every split. The lists of UCI credit concentrate on one feature: PAY_0, the repayment status of the latest month, appears in 8 of the 16 distinct pairs and in 16 of the 26 selections over five splits. On HELOC, the two pairs kept twice both involve a summary risk score or the share of trades never delinquent, ExternalRiskEstimate×NumRevolvingTradesWBalance and PercentTradesNeverDelq×NetFractionRevolvingBurden. A short list that changes from split to split is the honest answer when the interactions are weak; the long lists without the floor change as much and add loss.

VI. Theory

The experiments show what LENS does; this section explains why. All results concern the interaction projection R shared by the ranking and the test. Section VI-A shows what H^2 measures and when it reverses the order of two pairs, and that the energy ranks consistently while H^2 does not. Section VI-B shows what the cross inner product of two independent fits removes and what it keeps, and states the error control of the replication test. Section VI-C shows why the test needs a floor on the reproduced strength. All proofs are given in full.

A. Why the energy ranks and H^2 does not

Let $\mathcal{A}$ be the additive tables. Write a table as $F = \mu+\alpha+\beta+r$, where $\mu = \sum_{a,b} w_a v_b F(a,b)$ is its weighted grand mean, $\alpha(a) = \sum_b v_b F(a,b) - \mu$ and $\beta(b) = \sum_a w_a F(a,b) - \mu$ are its centered weighted row and column means, and $r = RF$. Let $M = \|\alpha\|^2 + \|\beta\|^2$ be the main-effect energy and $\Delta = \|r\|^2$ the interaction energy.

Proposition 1 (Decomposition): The tables μ, α, β, and r are pairwise orthogonal, r is orthogonal to every additive table, and $\|F-\mu\|^2 = M + \Delta$. So R is the orthogonal projection onto the complement of $\mathcal{A}$, and

$$\boxed{\Delta = \min_{A\in\mathcal{A}} \|F - A\|^2, \qquad H^2 = \frac{\Delta}{M+\Delta} = \sin^2\vartheta,} \tag{8}$$

where ϑ is the angle between $F-\mu$ and the centered additive tables.

Proof: Because the weights are a product, the inner product of a table that depends only on a with a table that depends only on b factors into two weighted means. Thus $\langle \alpha, \beta\rangle = \big(\sum_a w_a \alpha(a)\big)\big(\sum_b v_b \beta(b)\big) = 0$, and the constant μ is orthogonal to α, β, and r, each of which has weighted mean zero. By Eq. (3), the weighted row means of r vanish: $\sum_b v_b r(a,b) = \sum_b v_b F(a,b) - \sum_{b'} v_{b'} F(a,b') - \mu + \mu = 0$, and so do its column means. Hence for any additive table $c + \alpha'(a) + \beta'(b)$, $\langle r, \alpha'\rangle = \sum_a w_a \alpha'(a) \sum_b v_b r(a,b) = 0$,

and likewise $\langle r, \beta' \rangle = \langle r, c \rangle = 0$. Since $F - r = \mu + \alpha + \beta$ is additive and r is orthogonal to all additive tables, r is the orthogonal projection of F onto the complement of $\mathcal{A}$, which gives the first identity in Eq. (8). Pythagoras gives $\|F - \mu\|^2 = \|\alpha\|^2 + \|\beta\|^2 + \|r\|^2$. Finally, the projection of $F - \mu$ onto the centered additive tables is $\alpha + \beta$, because the remainder r is orthogonal to them, so $\cos^2 \vartheta = \|\alpha+\beta\|^2 / \|F - \mu\|^2 = M/(M + \Delta)$. ∎

The energy is the distance of a surface to the nearest additive table; H^2 is only the share of its variation that is not additive, so it measures the direction of a surface and ignores its size. A surface with almost no variation can still point away from the additive tables and score near one. The decomposition is exact because the grid carries product weights; it does not require the features to be independent in the data, since the background is averaged before the projection.

Proposition 2 (Invariance to main effects): Let $\tilde{f} = \hat{f} + b_0 + \sum_{\ell=1}^{p} b_\ell(x_\ell)$ for any functions b_ℓ. Replacing $\hat{f}_D$ or $\hat{f}_C$ by such an $\tilde{f}$ leaves every ranking score S_u, the root-mean-square strength s_D, and every ratio $A_{u,i}$ unchanged. In contrast, for a fixed residual $r \neq 0$ and an additive table A_0 with weighted mean zero, $H^2(r + A_0) = \|r\|^2 / (\|r\|^2 + \|A_0\|^2)$, which tends to zero as A_0 grows.

Proof: Fix a pair $u = (j, k)$ and a context row i. The context surface of $\tilde{f}$ is

$$\tilde{F}_u^{(i)}(a, b) = F_u^{(i)}(a, b) + b_j(a) + b_k(b) + b_0 + \textstyle\sum_{\ell \neq j,k} b_\ell(x_{i\ell}).$$

The added table is a function of a plus a function of b plus a constant, so it lies in $\mathcal{A}$ and $R\tilde{F}_u^{(i)} = RF_u^{(i)}$. This gives $\tilde{h}_{C,u,i} = h_{C,u,i}$ for every evaluation row. R is linear, so it commutes with the average over contexts, and $\tilde{r}_{D,u} = r_{D,u}$. The scores S_u, the strengths a_u and s_D, and the ratios $A_{u,i}$ are functions of these tables alone. For the second claim, $r + A_0$ has grand mean zero, its centered main effects sum to A_0, and its residual is r; Proposition 1 gives $M = \|A_0\|^2$ and $\Delta = \|r\|^2$. ∎

Adding main effects changes no energy and no ratio $A_{u,i}$, but can push any H^2 toward zero. Adding an unrelated main effect to the model therefore changes neither the order of the pairs nor any relative threshold of the test. It moves only the floor $\kappa\sigma_D$, which is measured in the spread of the predictions by design. The next result gives the exact condition under which H^2 reverses two pairs.

Proposition 3 (Reversal): If $\Delta_u > \Delta_v > 0$, then H^2 ranks v above u if and only if

$$\boxed{M_u/\Delta_u > M_v/\Delta_v.} \tag{9}$$

Proof: By Proposition 1,

$$H_u^2 - H_v^2 = \frac{\Delta_u M_v - \Delta_v M_u}{(M_u + \Delta_u)(M_v + \Delta_v)}.$$

The denominator is positive, so $H_v^2 > H_u^2$ exactly when $\Delta_v M_u > \Delta_u M_v$. Dividing by $\Delta_u \Delta_v > 0$ gives Eq. (9). ∎

H^2 ranks a weak interaction between features with small main effects above a strong interaction between features with large main effects. If every pair had the same main-effect energy, H^2 and Δ would order the pairs identically; the reversal is caused by differences in M between pairs. The following example shows that this happens even for a model that converges to the truth.

Example 1: Let $x_1, \ldots, x_4$ be independent and uniform on $\{-1, 1\}$, let the true function be $f^\star = 10x_1 + x_1x_2$, and let the fitted models be $\hat{f}_n = f^\star + \varepsilon_n x_3 x_4$ with $\varepsilon_n \downarrow 0$. Then $\hat{f}_n \to f^\star$ uniformly. With grids $\{-1, 1\}$ and equal weights, the surface of the true pair $(1, 2)$ is $10a + ab$, because the term $\varepsilon_n x_3 x_4$ averages to zero, so $\Delta = 1$, $M = 100$, and $H^2 = 1/101$. The surface of the spurious pair $(3, 4)$ is $\varepsilon_n cd$, because $10x_1 + x_1x_2$ averages to zero, so $\Delta = \varepsilon_n^2$, $M = 0$, and $H^2 = 1$. For every n, H^2 ranks the spurious pair first, while the energy ranks the true pair first once $\varepsilon_n < 1$.

The root cause is that H^2 has no limit at the zero surface. For a unit additive table A and a unit interaction table B, $H^2(\epsilon A) = 0$ and $H^2(\epsilon B) = 1$ for every $\epsilon > 0$, although both tables tend to zero. No value assigned to $H^2(0)$ makes it continuous there. A spurious surface that vanishes as the fit improves therefore keeps the score its shape dictates, whereas $\Delta(\epsilon B) = \epsilon^2$ vanishes.

Let $\mathcal{S}$ be the pairs whose target surface $F_u^\star$ is not additive and $\mathcal{N}$ the others, and write $\hat{F}_u$ for the estimated surface; the family is finite. A pair belongs to $\mathcal{S}$ only if its averaged surface is non-additive: a term such as $x_j x_k x_\ell$ with $\mathbb{E}x_\ell = 0$ is non-additive in every context but averages out, which leaves (j, k) in $\mathcal{N}$.

Theorem 1 (Energy ranks consistently): Let $d = \min_{u \in \mathcal{S}} \|RF_u^\star\|$ and $e = \max_u \|\hat{F}_u - F_u^\star\|$. Then

$$\boxed{e < d/2 \implies \min_{u \in \mathcal{S}} \hat{\Delta}_u > \max_{v \in \mathcal{N}} \hat{\Delta}_v.} \tag{10}$$

If $e \to 0$ in probability, all interacting pairs rank above all non-interacting pairs with probability tending to one, and Recall@$K \to 1$ for every $K \geq |\mathcal{S}|$.

Proof: R is an orthogonal projection, so $\|RG\| \leq \|G\|$ for every table G. By the reverse triangle inequality,

$$\left| \sqrt{\hat{\Delta}_u} - \|RF_u^\star\| \right| \leq \|R(\hat{F}_u - F_u^\star)\| \leq \|\hat{F}_u - F_u^\star\| \leq e.$$

For $v \in \mathcal{N}$, $RF_v^\star = 0$, so $\hat{\Delta}_v \leq e^2$. For $u \in \mathcal{S}$, $\sqrt{\hat{\Delta}_u} \geq d - e > e$ when $e < d/2$, so $\hat{\Delta}_u \geq (d - e)^2 > e^2$. The family is finite and every pair in $\mathcal{S}$ has a nonzero residual, so $d > 0$, and $\Pr(e < d/2) \to 1$ when $e \to 0$ in probability. On this event, the $|\mathcal{S}|$ highest energies belong to $\mathcal{S}$, so every list of length $K \geq |\mathcal{S}|$ contains all of $\mathcal{S}$. ∎

Any estimate whose surfaces are within half of the weakest true interaction ranks every interacting pair above every other pair. The bound is deterministic, so it holds in finite samples. Surface errors enter the energy of a non-interacting pair only at second order. Nothing in the bound requires the surfaces of non-interacting pairs to vanish: a pair with large main effects and no interaction, such as $x_j + x_k$, has $RF^\star = 0$ and is covered. H^2 has no such guarantee. Let

$h^\star = \min_{u\in\mathcal{S}} \Delta_u^\star/(M_u^\star + \Delta_u^\star)$ be the H^2 of the weakest interacting pair in the target.

Proposition 4 (H^2 does not rank consistently): Suppose the estimated surfaces converge on $\mathcal{S}$, and for some $\delta, \pi > 0$, $\liminf \Pr\big(\max_{v\in\mathcal{N}} \hat{H}_v^2 > h^\star + \delta\big) \geq \pi$. Then the probability that H^2 ranks all interacting pairs above all non-interacting pairs is at most $1 - \pi$ in the limit.

Proof: For $u \in \mathcal{S}$, $\Delta_u^\star > 0$, so the denominator $M_u^\star + \Delta_u^\star$ is positive and H^2 is continuous at $F_u^\star$. Hence $\hat{H}_u^2 \to h_u^\star$ in probability, and over the finite set $\mathcal{S}$, $\min_{u\in\mathcal{S}} \hat{H}_u^2 \to h^\star$. Let C_m be the event that H^2 ranks correctly and E_m the event in the assumption. On $C_m \cap E_m$, $\min_{\mathcal{S}} \hat{H}_u^2 > \max_{\mathcal{N}} \hat{H}_v^2 > h^\star + \delta$, an event whose probability tends to zero. Therefore $\Pr(C_m) \leq \Pr(C_m \cap E_m) + 1 - \Pr(E_m)$, and $\limsup \Pr(C_m) \leq 1 - \pi$. ∎

By the scale argument above, the condition depends only on the shapes of the spurious surfaces, not on their size. If a spurious surface is $\varepsilon_n G$ for a fixed shape G with $H^2(G) > h^\star + \delta$, the condition holds however fast ε_n vanishes, because $H^2(\varepsilon_n G) = H^2(G)$; Example 1 is this case with $\pi = 1$. Two forces make the condition easy to meet: main effects dilute the true pairs and lower $h^\star$, and the number of pairs of irrelevant features grows quadratically. Section IV-B measured both on fitted XGBoost models.

B. What replication removes

The ranking uses the self inner product of a table and the test uses the cross inner product of two tables. The difference is what the second fit adds. Let $c_u = \mathbb{E}[h_{C,u,i}]$ be the confirmation table averaged over evaluation contexts, and let $\rho_D = \mathbb{E}[r_{D,u} \mid \mathcal{G}]$ and $\rho_C = \mathbb{E}[c_u \mid \mathcal{G}]$ be the expected tables of the learner given the grids and weights $\mathcal{G}$, over training samples and training randomness. The grids are computed from D before $\hat{f}_C$ is trained, so given them the two fits remain independent.

Proposition 5 (Cross products cancel variance, not bias): If $\hat{f}_D$ and $\hat{f}_C$ are trained on independent samples with independent randomness, then, with all expectations given $\mathcal{G}$,

$$\boxed{\begin{aligned}&\mathbb{E}\langle r_{D,u}, c_u\rangle = \langle \rho_D, \rho_C\rangle,\\ &\mathbb{E}\|r_{D,u}\|^2 = \|\rho_D\|^2 + \mathbb{E}\|r_{D,u} - \rho_D\|^2.\end{aligned}} \tag{11}$$

Proof: Given $\mathcal{G}$, the table $r_{D,u}$ is a function of the discovery part and its contexts, and c_u is a function of the confirmation fit and the evaluation contexts. These are independent of D, and $\mathcal{G}$ is a function of D, so given $\mathcal{G}$ the two tables are independent and the expectation of each cell product factors: $\mathbb{E}[r_{D,u}(a,b)c_u(a,b) \mid \mathcal{G}] = \rho_D(a,b)\rho_C(a,b)$. The weights $w_a v_b$ are fixed by $\mathcal{G}$, so summing gives the first identity. The second is the bias–variance decomposition of $\|r_{D,u}\|^2$, cell by cell. ∎

Comparing two independent fits cancels their variance in expectation but not the bias they share: independent fitting noise adds nothing to the cross product on average, but a systematic pattern of the learner does. The self inner product carries the variance of the fit: a model trained on data without any interaction still has $\mathbb{E}\|r_{D,u}\|^2 > 0$, which is why every pair receives a positive score from one model. If either learner has an expected table of zero, the cross product has mean zero. What survives is the agreement of the two expected tables. When the data contain an interaction $r_u^\star$ and the learner is consistent, both expected tables approach $r_u^\star$, and the cross product approaches $\|r_u^\star\|^2$. But a learner can have a nonzero expected table where the data have no interaction, from the systematic part of its error, and two fits of the same learner share it. Section VI-C shows that this shared part is what a purely relative threshold cannot remove.

For the test, condition on everything fixed before the evaluation rows are used: both fitted models, the grids, the tables $r_{D,u}$, and the thresholds b_u. These are functions of the discovery and confirmation parts and of constants fixed in advance. The ratios $A_{u,1}, \ldots, A_{u,m}$ of a pair are then functions of independent evaluation rows.

Theorem 2 (False discovery rate): If the $A_{u,i}$ are i.i.d. with finite positive variance, then at the boundary $\theta_u = b_u$, $\sqrt{m}(\bar{A}_u - b_u)/s_u \to N(0,1)$ in distribution, and for every $t \in (0,1)$, $\limsup_{m\to\infty} \Pr(p_u \leq t) \leq t$ whenever $\theta_u \leq b_u$. Let the family contain P pairs, m_0 of which satisfy $\theta_u \leq b_u$, and let V and R_{BY} be the false and total rejections of BY at level α, with critical values $c_r = r\alpha/(PL_P)$ and $L_P = \sum_{i=1}^P 1/i$. For arbitrarily dependent p-values,

$$\boxed{\mathbb{E}\left[\frac{V}{R_{\mathrm{BY}} \vee 1}\right] \leq \frac{\alpha m_0}{P} + m_0\varepsilon_m,} \tag{12}$$

where $\varepsilon_m = \max_u \max_{r\leq P}[\Pr(p_u \leq c_r) - c_r]_+$ over the null pairs, and $\varepsilon_m \to 0$ as $m \to \infty$.

Proof: Pointwise validity. Let σ_u^2 be the variance of $A_{u,i}$. The central limit theorem gives $\sqrt{m}(\bar{A}_u - \theta_u)/\sigma_u \to N(0,1)$, the law of large numbers gives $s_u \to \sigma_u$ in probability, and Slutsky's theorem gives $\sqrt{m}(\bar{A}_u - \theta_u)/s_u \to N(0,1)$. At the boundary this is Z_u, so $p_u = 1 - \Phi(Z_u)$ converges in distribution to a uniform variable and $\Pr(p_u \leq t) \to t$. Inside the null, $\theta_u < b_u$, write

$$Z_u = \frac{\sqrt{m}(\bar{A}_u - \theta_u)}{s_u} - \frac{\sqrt{m}(b_u - \theta_u)}{s_u}.$$

The first term is bounded in probability and the second tends to $+\infty$, so $Z_u \to -\infty$, $p_u \to 1$ in probability, and $\Pr(p_u \leq t) \to 0$.

False discovery rate. We follow the argument of Benjamini and Yekutieli [17], keeping the marginal excess. Set $c_0 = 0$, $B_1 = [0, c_1]$, and $B_r = (c_{r-1}, c_r]$ for $r \geq 2$. BY rejects the R_{BY} smallest p-values, where R_{BY} is the largest r with $p_{(r)} \leq c_r$. If a null pair u is rejected and $p_u \in B_r$, then $r \leq R_{\mathrm{BY}}$. Hence, writing $F_u(t) = \Pr(p_u \leq t)$,

$$\begin{aligned}\mathbb{E}\left[\frac{\mathbf{1}\{u \text{ rejected}\}}{R_{\mathrm{BY}} \vee 1}\right] &\leq \sum_{r=1}^{P} \frac{\Pr(p_u \in B_r)}{r}\\ &= \sum_{r=1}^{P-1} \frac{F_u(c_r)}{r(r+1)} + \frac{F_u(c_P)}{P},\end{aligned}$$

where the equality is summation by parts. This step uses only the marginal law of p_u, which is why any dependence between pairs is allowed. Bound $F_u(c_r) \le c_r + \varepsilon_m$. The coefficients satisfy $\sum_{r<P} c_r/(r(r+1)) + c_P/P = \alpha(L_P - 1 + 1)/(PL_P) = \alpha/P$ and $\sum_{r<P} 1/(r(r+1)) + 1/P = 1$, so each null pair contributes at most $\alpha/P + \varepsilon_m$. Summing over the m_0 null pairs gives Eq. (12). By the first part, $F_u(c_r) \to c_r$ for a null pair at its boundary and $F_u(c_r) \to 0$ for one inside its null; the family and its critical values are finite, so $\varepsilon_m \to 0$. ∎

Conditional on the two fits, the replication test asymptotically controls the false discovery rate under any dependence between pairs. Every p-value uses the same two models and evaluation rows, so this dependence must be allowed. The guarantee is conditional on the two fits, so it holds for the models that the user actually has, not on average over models the user might have trained. The family is fixed while the number of evaluation rows grows; a family that grows with the data would need uniform bounds on the normal approximation, which we do not pursue.

C. Where replication alone fails

Theorem 2 holds for any threshold fixed before evaluation, including the purely relative one, $b_u = \eta$. The threshold decides which pairs the list contains as the fits improve.

Proposition 6 (A floor gives consistent selection): Suppose that for every pair $r_{D,u}$ and c_u converge in probability to a common table $r_u^\dagger$, possibly zero, and that $\bar{A}_u - \theta_u \to 0$ and $s_u/\sqrt{m} \to 0$ in probability. (a) With the relative threshold $b_u = \eta < 1$, LENS keeps every pair with $r_u^\dagger \neq 0$ with probability tending to one, however small $\|r_u^\dagger\|$ is. (b) With the threshold $\max(\eta a_u, \tau_u)$ on μ_u, where τ_u is fixed before evaluation and converges to $\tau_u^\dagger > 0$, and no $\|r_u^\dagger\|$ equals $\tau_u^\dagger$, LENS selects with probability tending to one

$$\boxed{\widehat{G} = \{u : \|r_u^\dagger\| > \tau_u^\dagger\}.} \tag{13}$$

The floor $\max(\eta s_D, \kappa\sigma_D)$ of Eq. (6) has a positive limit whenever the limit model is not constant on the contexts, because σ_D converges to its standard deviation there.

Proof: (a) For a pair with $r_u^\dagger \neq 0$, by the definition of θ_u and the Cauchy–Schwarz inequality,

$$|\theta_u - 1| = \frac{|\langle r_{D,u}, c_u - r_{D,u}\rangle|}{\|r_{D,u}\|^2} \le \frac{\|c_u - r_{D,u}\|}{\|r_{D,u}\|} \to 0,$$

because the numerator tends to zero and the denominator to $\|r_u^\dagger\| > 0$. Then $\bar{A}_u - \eta \to 1 - \eta > 0$, and since $s_u/\sqrt{m} \to 0$, $Z_u \to +\infty$ and $p_u \to 0$. For a fixed family, the smallest BY critical value $c_1 = \alpha/(PL_P)$ is a positive constant. With probability tending to one, every such p_u is below c_1, and the step-up procedure rejects every p-value below c_1.

(b) The statistic Z_u equals $(\hat{\mu}_u - t_u)/(a_u s_u/\sqrt{m})$, where $\hat{\mu}_u = a_u \bar{A}_u$ and $t_u = \max(\eta a_u, \tau_u)$, so its sign is that of the margin $\hat{\mu}_u - t_u$. Suppose first that $r_u^\dagger \neq 0$. Then $a_u \to \|r_u^\dagger\|$, and $\mu_u = \langle r_{D,u}, c_u\rangle/\|r_{D,u}\| \to \|r_u^\dagger\|$, so the margin converges to $g_u = \|r_u^\dagger\| - \max(\eta\|r_u^\dagger\|, \tau_u^\dagger)$. If $\|r_u^\dagger\| > \tau_u^\dagger$, then $g_u > 0$ because also $\|r_u^\dagger\| > \eta\|r_u^\dagger\|$; if $\|r_u^\dagger\| < \tau_u^\dagger$, then $g_u \le \|r_u^\dagger\| - \tau_u^\dagger < 0$. Suppose next that $r_u^\dagger = 0$. If $r_{D,u} = 0$, then $p_u = 1$. Otherwise $|\mu_u| \le \|c_u\| \to 0$ by the Cauchy–Schwarz inequality, whatever the direction of $r_{D,u}$, while $t_u \ge \tau_u \to \tau_u^\dagger$, so the margin tends to $-\tau_u^\dagger < 0$. With vanishing standard errors, a positive limit margin gives $p_u \to 0$ and a negative one gives $p_u \to 1$. The first kind is rejected as in part (a); the second is not, because the largest critical value $c_P = \alpha/L_P$ is below one. ∎

A threshold on the share alone keeps every pattern that both fits share, however faint; a floor on the reproduced strength keeps exactly the pairs whose interaction exceeds it. Part (a) asks only that the two tables agree relative to their size, so it also covers tables that vanish. If $r_{D,n} = c_n = \varepsilon_n B$ for a unit table B and $\varepsilon_n \downarrow 0$, then $\theta_n = 1$ for every n, and if the context tables scale with ε_n, the whole ratio test does not depend on ε_n. Shared faint patterns arise without any true interaction; they are the shared bias of Proposition 5. Let x_1, x_2 be uniform on $\{-1, 1\}$ and the true log-rate $ax_1 + bx_2$. A three-leaf tree that splits x_2 only where $x_1 = 1$ fits the residual table $(b/2)x_1x_2$, and it does so on every sample. Correlated features add a second source: when the data never visit some grid cells, a model is free to extrapolate there, and the same learner tends to extrapolate the same way on two samples. A model that converges to an additive function can thus hand over a list that does not shrink.

The relative floor ηs_D removes pairs that are weak compared with the rest of the family. It cannot remove patterns that shrink in every pair at once, however: if every table is $\varepsilon_n B_u$, then $a_u = s_D = \varepsilon_n$ and the floor shrinks with them. This is the case of a model that has learned no interaction, where every table is an artifact of the learner. The scale floor $\kappa\sigma_D$ does not shrink with the tables, so part (b) covers this case too.

No floor separates a true interaction from its proxies when features are strongly correlated. If x_3 is a noisy copy of x_1, a model of x_1x_2 can route part of the interaction through the pair (x_3, x_2), and the fitted surface of that pair is then non-additive in every fit. LENS tests the interactions of the model, so it reports the proxy when both fits route the interaction the same way.

VII. Discussion

The price of a split. Each LENS model sees a third of the data, so the list is the one that a third of the data can confirm. On freMTPL2 the list grows from 6.6 to 12 pairs as the data per model grow from 20,000 to 184,000 policies, and most of the pairs kept with less data stay (Section V), so a split mainly costs the weakest pairs. The three parts need not be equal. The evaluation part only supplies contexts for an average, and Section IV-D shows that the standard errors are rarely what decides, so it can be much smaller than the two training parts. It must be fresh, however: the conditioning argument of Section VI-B needs contexts that are independent of both fits. Rotating the roles of the parts would use the data more fully. Multi-split methods combine p-values across random splits for a fixed set of hypotheses [25]; here the hypotheses themselves change with the rotation, because the discovery directions and

boundaries do, so keeping Theorem 2 across rotations needs a new argument. We use one split because its error control follows directly from the conditioning.

Setting the boundary. The floor $\kappa\sigma_D$ is a size of effect, so an analyst can replace it with a tolerance from the application. In a motor tariff, a floor of 0.01 on the log claim rate says that the reproduced interaction must move the claim rate by about 1% over the grid to be worth a term.

Two fits of one learner. LENS replicates across samples, not across learners. Both fits come from the same algorithm, which is what lets them share the systematic part of their errors (Proposition 5), and what the floor must then remove. Fitting the confirmation model with a different learner would cancel some of that shared part. It would also change the question from "did this model learn a stable interaction" to "do two kinds of model agree on it". Our experiments hint at what this would do. On Pair-corr, XGBoost routes part of the true interaction through the proxies and the network almost does not; a cross-learner test would likely drop the proxies, together with any true interaction that only one learner represents. We have not evaluated this variant.

What the averaged surface misses. LENS tests the partial-dependence surface, the average of the context surfaces. A three-way term $x_j x_k x_\ell$ with $\mathbb{E}x_\ell = 0$ leaves every pair of $\{j, k, \ell\}$ additive on average, and an interaction whose sign changes across subgroups can cancel in the same way. The test inherits this from partial dependence. Extrapolation to grid cells that the data never visit is a related concern. A pattern that one fit invents there must replicate in the other fit and exceed the floor. But a pattern that the learner extrapolates the same way on both samples can pass both, which is one source of the shared bias of Section VI-C. The same projection extends to triples, with three-way tables in place of two-way ones and $p(p-1)(p-2)/6$ candidates, which keeps the two fits and multiplies the queries by the third grid. We leave this extension to future work.

VIII. Related Work

False discoveries in data management. The database community has long asked how many results mined in bulk are real. Interactive exploration tests many hypotheses implicitly and so invites false discoveries [26]. In association rule mining, the permutation correction is the most powerful but also the most costly [27], and FSR reduces the resampled datasets that significant pattern mining needs [28]. CleanML adopts the Benjamini–Yekutieli procedure for the same reason as LENS: its tests are dependent [29]. Closest to our setting, Slice Finder tests the data slices on which a trained model underperforms [30], [31]. These methods test patterns in the data or in the loss of a model. LENS asks the same question of the interactions a model has learned.

Ranking interactions. Model-specific scores read one model family: TreeSHAP for tree ensembles [8], NID for network weights [9], and Integrated Hessians for differentiable networks [10]. Model-agnostic scores include H^2 [11], built on functional ANOVA [32], Archipelago [20], and SHAP-IQ [21]. Other methods build interactions into the model: ARM-Net learns cross features for structured data [33], and FAST adds pairs to an additive model one at a time [5]. None of them tests the pairs, so the length of the list is left to the user. Others have also found H^2 unreliable [18], [34], [35].

Testing and replication. Refit tests build a null distribution by refitting the model [11], [13], [14], but an empirical rank p-value from B refits cannot fall below $1/(B+1)$ [36], so a correction over many pairs needs many refits. DIAMOND controls the false discovery rate with model-X knockoffs, which must reproduce the joint law of the features [16], [37]; ARDA uses random features as a related baseline for single features [38]. iLOCO compares the loss after leaving out features and pairs [15]. For single features, sample-splitting tests rest on the same fact as LENS: fitting noise in one model is independent of a model trained on separate data [39], [40]. Statistics has formal notions of replicability across studies [41]–[43]. LENS applies replication to the interactions of a fitted model.

IX. Conclusion

We presented LENS, a model-agnostic framework that selects feature interactions through independent replication: it asks whether a second model, trained on independent data, reproduces the interaction pattern of the first. It needs two fits for any number of pairs, and one multiple-testing correction sets the length of the list. Our theory shows why the unnormalized energy ranks consistently where H^2 does not, and why consistent selection needs a floor on the reproduced strength. The list is clean and finds most material interactions on synthetic benchmarks, agrees with the actuarial literature and recovers most of the gain of all pairs on insurance data, and keeps few pairs without raising the loss on credit data. A natural next step is three-way interactions.

AI-Generated Content Acknowledgement

Claude (Anthropic) assisted with language editing throughout the manuscript and with developing the experimental code. GPT Image (OpenAI) produced a layout draft for Fig. 2, which we redrew as a vector figure. The authors reviewed and revised all content and take full responsibility for it.

References

[1] H.-T. Cheng, L. Koc, J. Harmsen, T. Shaked, T. Chandra, H. Aradhye, G. Anderson, G. Corrado, W. Chai, M. Ispir, R. Anil, Z. Haque, L. Hong, V. Jain, X. Liu, and H. Shah, "Wide & deep learning for recommender systems," in *Proceedings of the Workshop on Deep Learning for Recommender Systems*, 2016, pp. 7–10.

[2] Y. Luo, M. Wang, H. Zhou, Q. Yao, W.-W. Tu, Y. Chen, W. Dai, and Q. Yang, "AutoCross: Automatic feature crossing for tabular data in real-world applications," in *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2019, pp. 1936–1945.

[3] H. J. Cordell, "Detecting gene–gene interactions that underlie human diseases," *Nature Reviews Genetics*, vol. 10, no. 6, pp. 392–404, 2009.

[4] J. Schelldorfer and M. V. Wüthrich, "Nesting classical actuarial models into neural networks," SSRN 3320525, 2019.

[5] Y. Lou, R. Caruana, J. Gehrke, and G. Hooker, "Accurate intelligible models with pairwise interactions," in *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2013, pp. 623–631.

[6] Y. Havrylenko and J. Heger, "Detection of interacting variables for generalized linear models via neural networks," *European Actuarial Journal*, vol. 14, no. 2, pp. 551–580, 2024.

[7] M. Lindholm and J. Palmquist, "Black-box guided generalised linear model building with non-life pricing applications," *Annals of Actuarial Science*, vol. 18, no. 3, pp. 675–691, 2024.

[8] S. M. Lundberg, G. Erion, H. Chen, A. DeGrave, J. M. Prutkin, B. Nair, R. Katz, J. Himmelfarb, N. Bansal, and S.-I. Lee, "From local explanations to global understanding with explainable AI for trees," *Nature Machine Intelligence*, vol. 2, no. 1, pp. 56–67, 2020.

[9] M. Tsang, D. Cheng, and Y. Liu, "Detecting statistical interactions from neural network weights," in *Proceedings of the International Conference on Learning Representations*, 2018.

[10] J. D. Janizek, P. Sturmfels, and S.-I. Lee, "Explaining explanations: Axiomatic feature interactions for deep networks," *Journal of Machine Learning Research*, vol. 22, no. 104, pp. 1–54, 2021.

[11] J. H. Friedman and B. E. Popescu, "Predictive learning via rule ensembles," *The Annals of Applied Statistics*, vol. 2, no. 3, pp. 916–954, 2008.

[12] R. Henckaerts, K. Antonio, and M.-P. Côté, "When stakes are high: Balancing accuracy and transparency with model-agnostic interpretable data-driven surrogates," *Expert Systems with Applications*, vol. 202, 2022, Art. no. 117230.

[13] E. Lampa, L. Lind, P. M. Lind, and A. Bornefalk-Hermansson, "The identification of complex interactions in epidemiology and toxicology: A simulation study of boosted regression trees," *Environmental Health*, vol. 13, 2014, Art. no. 57.

[14] T. Welchowski, K. O. Maloney, R. Mitchell, and M. Schmid, "Techniques to improve ecological interpretability of black-box machine learning models," *Journal of Agricultural, Biological and Environmental Statistics*, vol. 27, no. 1, pp. 175–197, 2022.

[15] C. Little, L. Zheng, and G. Allen, "iLOCO: Distribution-free inference for feature interactions," *arXiv preprint arXiv:2502.06661*, 2025.

[16] W. Chen, Y. Jiang, W. S. Noble, and Y. Y. Lu, "Error-controlled non-additive interaction discovery in machine learning models," *Nature Machine Intelligence*, vol. 7, no. 9, pp. 1541–1554, 2025.

[17] Y. Benjamini and D. Yekutieli, "The control of the false discovery rate in multiple testing under dependency," *The Annals of Statistics*, vol. 29, no. 4, pp. 1165–1188, 2001.

[18] A. Inglis, A. Parnell, and C. B. Hurley, "Visualizing variable importance and variable interaction effects in machine learning models," *Journal of Computational and Graphical Statistics*, vol. 31, no. 3, pp. 766–778, 2022.

[19] M. Muschalik, F. Fumagalli, B. Hammer, and E. Hüllermeier, "Beyond TreeSHAP: Efficient computation of any-order Shapley interactions for tree ensembles," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 38, 2024, pp. 14 388–14 396.

[20] M. Tsang, S. Rambhatla, and Y. Liu, "How does this interaction affect me? Interpretable attribution for feature interactions," in *Advances in Neural Information Processing Systems*, vol. 33, 2020, pp. 6147–6159.

[21] F. Fumagalli, M. Muschalik, P. Kolpaczki, E. Hüllermeier, and B. Hammer, "SHAP-IQ: Unified approximation of any-order Shapley interactions," in *Advances in Neural Information Processing Systems*, vol. 36, 2023, pp. 11 515–11 551.

[22] C.-P. Tsai, C.-K. Yeh, and P. Ravikumar, "Faith-Shap: The faithful Shapley interaction index," *Journal of Machine Learning Research*, vol. 24, no. 94, pp. 1–42, 2023.

[23] M. Mayer, "hstats: Interaction statistics," R package, CRAN, 2023.

[24] S. Basu, K. Kumbier, J. B. Brown, and B. Yu, "Iterative random forests to discover predictive and stable high-order interactions," *Proceedings of the National Academy of Sciences*, vol. 115, no. 8, pp. 1943–1948, 2018.

[25] N. Meinshausen, L. Meier, and P. Bühlmann, "p-values for high-dimensional regression," *Journal of the American Statistical Association*, vol. 104, no. 488, pp. 1671–1681, 2009.

[26] Z. Zhao, L. De Stefani, E. Zgraggen, C. Binnig, E. Upfal, and T. Kraska, "Controlling false discoveries during interactive data exploration," in *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 2017, pp. 527–540.

[27] G. Liu, H. Zhang, and L. Wong, "Controlling false positives in association rule mining," *Proceedings of the VLDB Endowment*, vol. 5, no. 2, pp. 145–156, 2011.

[28] L. Pellegrina and F. Vandin, "Efficient discovery of significant patterns with few-shot resampling," *Proceedings of the VLDB Endowment*, vol. 17, no. 10, pp. 2668–2680, 2024.

[29] P. Li, S. X. Rao, J. L. Blase, Y. Zhang, X. Chu, and C. Zhang, "CleanML: A study for evaluating the impact of data cleaning on ML classification tasks," in *Proceedings of the IEEE International Conference on Data Engineering*, 2021, pp. 13–24.

[30] Y. Chung, T. Kraska, N. Polyzotis, K. H. Tae, and S. E. Whang, "Slice Finder: Automated data slicing for model validation," in *Proceedings of the IEEE International Conference on Data Engineering*, 2019, pp. 1550–1553.

[31] Y. Chung, T. Kraska, N. Polyzotis, K. H. Tae, and S. E. Whang, "Automated data slicing for model validation: A big data–AI integration approach," *IEEE Transactions on Knowledge and Data Engineering*, vol. 32, no. 12, pp. 2284–2296, 2020.

[32] G. Hooker, "Discovering additive structure in black box functions," in *Proceedings of the ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2004, pp. 575–580.

[33] S. Cai, K. Zheng, G. Chen, H. V. Jagadish, B. C. Ooi, and M. Zhang, "ARM-Net: Adaptive relation modeling network for structured data," in *Proceedings of the ACM SIGMOD International Conference on Management of Data*, 2021, pp. 207–220.

[34] B. M. Greenwell, B. C. Boehmke, and A. J. McCarthy, "A simple and effective model-based variable importance measure," *arXiv preprint arXiv:1805.04755*, 2018.

[35] T. Welchowski and D. Edelmann, "Interaction difference hypothesis test for prediction models," *Machine Learning and Knowledge Extraction*, vol. 6, no. 2, pp. 1298–1322, 2024.

[36] B. Phipson and G. K. Smyth, "Permutation p-values should never be zero: Calculating exact p-values when permutations are randomly drawn," *Statistical Applications in Genetics and Molecular Biology*, vol. 9, no. 1, 2010, Art. no. 39.

[37] E. Candès, Y. Fan, L. Janson, and J. Lv, "Panning for gold: 'Model-X' knockoffs for high dimensional controlled variable selection," *Journal of the Royal Statistical Society: Series B*, vol. 80, no. 3, pp. 551–577, 2018.

[38] N. Chepurko, R. Marcus, E. Zgraggen, R. C. Fernandez, T. Kraska, and D. Karger, "ARDA: Automatic relational data augmentation for machine learning," *Proceedings of the VLDB Endowment*, vol. 13, no. 9, pp. 1373–1387, 2020.

[39] B. Dai, X. Shen, and W. Pan, "Significance tests of feature relevance for a black-box learner," *IEEE Transactions on Neural Networks and Learning Systems*, vol. 35, no. 2, pp. 1898–1911, 2024.

[40] W. Tansey, V. Veitch, H. Zhang, R. Rabadan, and D. M. Blei, "The holdout randomization test for feature selection in black box models," *Journal of Computational and Graphical Statistics*, vol. 31, no. 1, pp. 151–162, 2022.

[41] Y. Benjamini and R. Heller, "Screening for partial conjunction hypotheses," *Biometrics*, vol. 64, no. 4, pp. 1215–1222, 2008.

[42] M. Bogomolov and R. Heller, "Discovering findings that replicate from a primary study of high dimension to a follow-up study," *Journal of the American Statistical Association*, vol. 108, no. 504, pp. 1480–1492, 2013.

[43] R. Heller, M. Bogomolov, and Y. Benjamini, "Deciding whether follow-up studies have replicated findings in a preliminary large-scale omics study," *Proceedings of the National Academy of Sciences*, vol. 111, no. 46, pp. 16 262–16 267, 2014.

APPENDIX A
COMPLETE SYNTHETIC RESULTS

Tables IX–XII give every design at every number of features, for every method and every variant of its decision rule that we ran. Table V in the body reports the default rule of each method at $p = 50$. LENS appears with its default boundary ($\kappa = 0.02$), with a smaller scale floor ($\kappa = 0.01$), without the scale floor ($\kappa = 0$, the boundary $\max(\eta a_u, \eta s_D)$), and with the purely relative boundary $b_u = \eta$ of Proposition 6(a). DIAMOND appears at both levels of its knockoff filter, with its own XGBoost, with the XGBoost settings of LENS, and with its MLP. H^2 appears with the rank p-values of the refit test, which keep no pair, and with the normal approximation used in the body. Tables XIII–XV break F1–F10 down by function.

TABLE IX: Sparse design: list size, FDP, and power, means over 20 seeds, for $p = 10$, 30, and 50 (45, 435, and 1,225 candidate pairs).

		$p = 10$			$p = 30$			$p = 50$		
Method	Model	size	FDP	power	size	FDP	power	size	FDP	power
LENS (default, $\kappa = 0.02$)	XGBoost	3.0	0.01	1.00	3.0	0.01	1.00	3.0	0.01	1.00
LENS, $\kappa = 0.01$	XGBoost	3.0	0.01	1.00	3.2	0.06	1.00	3.4	0.09	1.00
LENS, $\kappa = 0$	XGBoost	3.0	0.01	1.00	3.5	0.12	1.00	7.2	0.54	1.00
LENS, $b_u = \eta$	XGBoost	8.3	0.61	1.00	79.0	0.96	1.00	233.8	0.99	1.00
LENS (default, $\kappa = 0.02$)	MLP	3.0	0.00	1.00	3.0	0.00	1.00	3.0	0.00	1.00
LENS, $\kappa = 0.01$	MLP	3.0	0.00	1.00	3.0	0.00	1.00	3.0	0.00	1.00
LENS, $\kappa = 0$	MLP	3.0	0.00	1.00	3.0	0.00	1.00	3.9	0.20	1.00
DIAMOND, $q = 0.2$	XGBoost	2.8	0.03	0.90	3.5	0.16	0.82	4.0	0.18	0.93
DIAMOND, $q = 0.05$	XGBoost	2.8	0.03	0.90	3.1	0.14	0.82	3.9	0.17	0.93
DIAMOND, $q = 0.2$	XGBoost†	2.9	0.04	0.90	3.6	0.19	0.82	4.0	0.17	0.93
DIAMOND, $q = 0.2$	MLP	2.8	0.01	0.90	3.4	0.16	0.82	3.6	0.14	0.93
DIAMOND, $q = 0.05$	MLP	2.8	0.01	0.90	3.0	0.14	0.82	3.6	0.14	0.93
H^2, normal refit null	XGBoost	3.5	0.11	1.00	36.5	0.91	1.00	73.5	0.96	1.00
H^2, rank p-values	XGBoost	0.0	0.00	0.00	0.0	0.00	0.00	0.0	0.00	0.00
iLOCO-Split, $\alpha = 0.05$	XGBoost	15.2	0.76	1.00	118.5	0.92	1.00	399.1	0.98	1.00
iLOCO-Split, $\alpha = 0.2$	XGBoost	19.2	0.82	1.00	136.9	0.94	1.00	437.8	0.98	1.00
iLOCO-MP, $\alpha = 0.05$	trees	13.4	0.77	1.00	205.8	0.99	1.00	487.6	0.99	1.00
FAST	EBM	16.1	0.78	1.00	19.0	0.83	1.00	20.0	0.85	1.00

† DIAMOND with the XGBoost settings of LENS.

TABLE X: Pair-corr design: list size, FDP, and power, means over 20 seeds, for $p = 10$, 30, and 50 (45, 435, and 1,225 candidate pairs).

		$p = 10$			$p = 30$			$p = 50$		
Method	Model	size	FDP	power	size	FDP	power	size	FDP	power
LENS (default, $\kappa = 0.02$)	XGBoost	3.0	0.67	1.00	3.0	0.67	1.00	3.0	0.67	1.00
LENS, $\kappa = 0.01$	XGBoost	3.0	0.67	1.00	3.2	0.69	1.00	3.4	0.70	1.00
LENS, $\kappa = 0$	XGBoost	3.0	0.67	1.00	3.5	0.71	1.00	4.9	0.78	1.00
LENS, $b_u = \eta$	XGBoost	8.6	0.88	1.00	76.8	0.99	1.00	240.8	1.00	1.00
LENS (default, $\kappa = 0.02$)	MLP	1.1	0.03	1.00	1.1	0.05	1.00	1.1	0.03	1.00
LENS, $\kappa = 0.01$	MLP	1.1	0.03	1.00	2.1	0.40	1.00	2.4	0.36	1.00
LENS, $\kappa = 0$	MLP	1.1	0.03	1.00	2.8	0.58	1.00	12.3	0.90	1.00
DIAMOND, $q = 0.2$	XGBoost	0.5	0.11	0.20	1.3	0.24	0.35	2.0	0.35	0.55
DIAMOND, $q = 0.05$	XGBoost	0.5	0.11	0.20	1.2	0.23	0.35	1.8	0.35	0.55
DIAMOND, $q = 0.2$	XGBoost†	0.5	0.12	0.20	1.7	0.30	0.40	1.6	0.31	0.50
DIAMOND, $q = 0.2$	MLP	0.6	0.13	0.25	1.4	0.27	0.35	0.6	0.12	0.20
DIAMOND, $q = 0.05$	MLP	0.6	0.13	0.25	1.2	0.27	0.35	0.6	0.12	0.20
H^2, normal refit null	XGBoost	9.1	0.88	1.00	54.3	0.98	1.00	87.4	0.99	1.00
H^2, rank p-values	XGBoost	0.0	0.00	0.00	0.0	0.00	0.00	0.0	0.00	0.00
iLOCO-Split, $\alpha = 0.05$	XGBoost	5.2	0.60	1.00	91.2	0.97	1.00	181.6	0.99	0.85
iLOCO-Split, $\alpha = 0.2$	XGBoost	6.5	0.65	1.00	118.2	0.98	1.00	226.8	0.99	0.90
iLOCO-MP, $\alpha = 0.05$	trees	19.0	0.95	1.00	256.9	1.00	1.00	588.6	1.00	1.00
FAST	EBM	9.6	0.71	1.00	9.3	0.48	1.00	8.9	0.48	1.00

† DIAMOND with the XGBoost settings of LENS.

TABLE XI: Additive design: list size and the share of runs that return any pair (every pair is false), means over 20 seeds, for $p = 10$, 30, and 50 (45, 435, and 1,225 candidate pairs).

		$p = 10$		$p = 30$		$p = 50$	
Method	Model	size	any	size	any	size	any
LENS (default, $\kappa = 0.02$)	XGBoost	0.1	0.05	0.1	0.05	0.1	0.10
LENS, $\kappa = 0.01$	XGBoost	0.1	0.15	0.2	0.20	0.2	0.25
LENS, $\kappa = 0$	XGBoost	0.9	0.65	7.5	1.00	17.2	1.00
LENS, $b_u = \eta$	XGBoost	8.3	1.00	76.5	1.00	219.6	1.00
LENS (default, $\kappa = 0.02$)	MLP	0.0	0.00	0.0	0.00	0.0	0.00
LENS, $\kappa = 0.01$	MLP	0.1	0.05	0.1	0.05	0.2	0.25
LENS, $\kappa = 0$	MLP	16.1	1.00	145.8	1.00	403.2	1.00
DIAMOND, $q = 0.2$	XGBoost	1.1	0.40	2.2	0.85	2.6	0.90
DIAMOND, $q = 0.05$	XGBoost	1.1	0.40	2.2	0.85	2.6	0.90
DIAMOND, $q = 0.2$	XGBoost†	1.3	0.50	2.2	0.80	2.5	0.85
DIAMOND, $q = 0.2$	MLP	0.7	0.30	0.4	0.15	1.8	0.65
DIAMOND, $q = 0.05$	MLP	0.7	0.30	0.4	0.15	1.8	0.65
H^2, normal refit null	XGBoost	0.6	0.35	8.0	1.00	25.1	1.00
H^2, rank p-values	XGBoost	0.0	0.00	0.0	0.00	0.0	0.00
iLOCO-Split, $\alpha = 0.05$	XGBoost	1.4	0.35	1.2	0.20	9.3	0.45
iLOCO-Split, $\alpha = 0.2$	XGBoost	3.0	0.60	5.8	0.40	21.7	0.50
iLOCO-MP, $\alpha = 0.05$	trees	21.0	1.00	264.4	1.00	609.5	1.00
FAST	EBM	5.0	0.50	1.1	0.20	1.6	0.30

† DIAMOND with the XGBoost settings of LENS.

TABLE XII: Additive-corr design: list size and the share of runs that return any pair (every pair is false), means over 20 seeds, for $p = 10$, 30, and 50 (45, 435, and 1,225 candidate pairs).

		$p = 10$		$p = 30$		$p = 50$	
Method	Model	size	any	size	any	size	any
LENS (default, $\kappa = 0.02$)	XGBoost	0.1	0.05	0.1	0.05	0.1	0.05
LENS, $\kappa = 0.01$	XGBoost	0.5	0.40	0.6	0.45	0.5	0.50
LENS, $\kappa = 0$	XGBoost	1.9	0.95	10.1	1.00	20.2	1.00
LENS, $b_u = \eta$	XGBoost	9.4	1.00	92.4	1.00	228.4	1.00
LENS (default, $\kappa = 0.02$)	MLP	0.2	0.20	0.1	0.10	0.1	0.15
LENS, $\kappa = 0.01$	MLP	1.4	0.90	1.1	0.75	1.1	0.70
LENS, $\kappa = 0$	MLP	13.4	1.00	137.4	1.00	377.7	1.00
DIAMOND, $q = 0.2$	XGBoost	1.4	0.40	2.3	0.80	3.0	0.90
DIAMOND, $q = 0.05$	XGBoost	1.4	0.40	2.3	0.80	2.9	0.90
DIAMOND, $q = 0.2$	XGBoost†	1.2	0.35	2.4	0.80	3.0	0.90
DIAMOND, $q = 0.2$	MLP	0.6	0.15	1.1	0.45	3.1	0.90
DIAMOND, $q = 0.05$	MLP	0.5	0.15	1.1	0.45	3.1	0.90
H^2, normal refit null	XGBoost	0.4	0.20	9.7	0.95	23.6	1.00
H^2, rank p-values	XGBoost	0.0	0.00	0.0	0.00	0.0	0.00
iLOCO-Split, $\alpha = 0.05$	XGBoost	0.3	0.25	2.1	0.15	25.4	0.15
iLOCO-Split, $\alpha = 0.2$	XGBoost	1.6	0.40	4.6	0.30	33.1	0.30
iLOCO-MP, $\alpha = 0.05$	trees	22.2	1.00	247.7	1.00	599.3	1.00
FAST	EBM	4.3	0.60	2.2	0.50	2.5	0.60

† DIAMOND with the XGBoost settings of LENS.

TABLE XIII: F1–F10 ($p = 30$): FDP per function, means over 20 seeds.

Method	Model	F1	F2	F3	F4	F5	F6	F7	F8	F9	F10	mean
LENS (default, $\kappa = 0.02$)	XGBoost	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
LENS (default, $\kappa = 0.02$)	MLP	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
DIAMOND, $q = 0.2$	XGBoost	0.11	0.13	0.11	0.09	0.11	0.03	0.05	0.03	0.02	0.07	0.07
DIAMOND, $q = 0.2$	MLP	0.16	0.18	0.14	0.13	0.16	0.05	0.03	0.02	0.02	0.07	0.10
H^2, normal refit null	XGBoost	0.76	0.78	0.82	0.77	0.70	0.83	0.31	0.70	0.72	0.83	0.72
iLOCO-Split, $\alpha = 0.05$	XGBoost	0.72	0.64	0.76	0.76	0.49	0.70	0.63	0.61	0.62	0.77	0.67
iLOCO-MP, $\alpha = 0.05$	trees	0.98	0.98	0.99	0.99	0.99	0.98	1.00	0.99	0.98	0.99	0.99
FAST	EBM	0.51	0.54	0.46	0.47	0.57	0.68	0.35	0.38	0.40	0.70	0.51

TABLE XIV: F1–F10 ($p = 30$): power per function, means over 20 seeds.

Method	Model	F1	F2	F3	F4	F5	F6	F7	F8	F9	F10	mean
LENS (default, $\kappa = 0.02$)	XGBoost	0.54	0.55	1.00	1.00	0.88	0.25	0.21	0.67	0.13	1.00	0.62
LENS (default, $\kappa = 0.02$)	MLP	0.64	0.64	1.00	1.00	0.88	0.25	0.21	0.75	0.27	1.00	0.66
DIAMOND, $q = 0.2$	XGBoost	0.29	0.29	0.40	0.35	0.76	0.12	0.14	0.26	0.06	0.47	0.31
DIAMOND, $q = 0.2$	MLP	0.29	0.28	0.16	0.20	0.75	0.12	0.15	0.28	0.10	0.59	0.29
H^2, normal refit null	XGBoost	0.94	0.93	1.00	1.00	0.88	1.00	0.84	1.00	1.00	1.00	0.96
iLOCO-Split, $\alpha = 0.05$	XGBoost	0.57	0.55	0.94	0.95	0.76	0.81	0.02	0.19	0.56	0.64	0.60
iLOCO-MP, $\alpha = 0.05$	trees	0.40	0.38	0.17	0.15	0.46	0.62	0.04	0.15	0.38	0.27	0.30
FAST	EBM	0.81	0.80	1.00	1.00	0.88	0.80	0.93	1.00	0.80	1.00	0.90

TABLE XV: F1–F10 ($p = 30$): list size per function, means over 20 seeds.

Method	Model	F1	F2	F3	F4	F5	F6	F7	F8	F9	F10	mean
LENS (default, $\kappa = 0.02$)	XGBoost	6.0	6.0	9.0	10.0	7.0	2.0	3.0	8.0	2.0	6.0	5.9
LENS (default, $\kappa = 0.02$)	MLP	7.0	7.0	9.0	10.0	7.0	2.0	3.0	9.0	4.0	6.0	6.4
DIAMOND, $q = 0.2$	XGBoost	3.6	3.7	4.5	4.2	7.0	1.0	2.1	3.3	0.9	3.2	3.4
DIAMOND, $q = 0.2$	MLP	4.0	3.9	2.4	3.1	7.3	1.1	2.2	3.5	1.6	4.0	3.3
H^2, normal refit null	XGBoost	43.2	46.5	50.3	45.2	24.5	48.6	17.5	41.1	56.0	36.2	40.9
iLOCO-Split, $\alpha = 0.05$	XGBoost	34.4	35.8	67.7	52.5	19.4	39.1	3.4	18.1	27.7	38.1	33.6
iLOCO-MP, $\alpha = 0.05$	trees	253.2	251.7	220.3	225.2	249.9	272.4	226.3	255.1	270.4	260.8	248.5
FAST	EBM	18.8	19.6	17.6	19.2	17.4	20.0	20.0	19.6	20.0	20.0	19.2

APPENDIX B
WHICH PART OF LENS DOES THE WORK

LENS combines four ingredients: a ranking by the energy a_u^2, a second model, a comparison of the two fits along the discovery pattern, and a test. Each of the following rules removes one of them and keeps the boundary of Eq. (6).

- *Top-$|\widehat{G}|$ by energy* removes the decision: it keeps the $|\widehat{G}|$ pairs of highest energy, as many as LENS keeps. It isolates which pairs the test chooses, at a fixed length.
- *$\hat{f}_D$ only* removes the second model: it keeps every pair whose discovery strength a_u exceeds the floor $\max(\eta s_D, \kappa\sigma_D)$.
- *Both fits above the floor, no comparison* uses the second model but not the comparison: it keeps a pair when both a_u and the strength $\|r_{C,u}\|$ of the confirmation table exceed the same floor, whatever the direction of the two tables.
- *$\hat{\mu}_u$ above the boundary, no test* keeps the comparison and drops the test: it keeps a pair when the reproduced strength $\hat{\mu}_u = a_u \bar{A}_u$ exceeds $\max(\eta a_u, \eta s_D, \kappa\sigma_D)$.

All four are computed from the same two fits as LENS, so the comparison differs only in the rule.

TABLE XVI: Ablations on synthetic data with XGBoost, means over 20 seeds. Each rule removes one ingredient of LENS and keeps the boundary of Eq. (6). For the additive designs, the share of runs that return any pair (every pair is false); otherwise list size, FDP, and power.

	Additive	Additive-corr	Sparse			Pair-corr			F1–F10 ($p = 30$)		
Rule	any	any	size	FDP	power	size	FDP	power	size	FDP	power
$p = 10$											
LENS (default)	0.05	0.05	3.0	0.01	1.00	3.0	0.67	1.00			
$\hat{\mu}_u$ above the boundary, no test	0.05	0.05	3.0	0.01	1.00	3.0	0.67	1.00			
Top-$\lvert\widehat{G}\rvert$ by energy	0.05	0.05	3.0	0.01	1.00	3.0	0.67	1.00			
Both fits above the floor, no comparison	1.00	1.00	4.3	0.27	1.00	4.1	0.76	1.00			
$\hat{f}_D$ only: a_u above the floor	1.00	1.00	5.0	0.37	1.00	5.1	0.80	1.00			
$p = 30$											
LENS (default)	0.05	0.05	3.0	0.01	1.00	3.0	0.67	1.00	5.9	0.00	0.62
$\hat{\mu}_u$ above the boundary, no test	0.05	0.05	3.0	0.01	1.00	3.0	0.67	1.00	5.9	0.00	0.62
Top-$\lvert\widehat{G}\rvert$ by energy	0.05	0.05	3.0	0.01	1.00	3.0	0.67	1.00	5.9	0.00	0.62
Both fits above the floor, no comparison	1.00	1.00	5.8	0.48	1.00	4.0	0.75	1.00	6.3	0.04	0.64
$\hat{f}_D$ only: a_u above the floor	1.00	1.00	6.0	0.50	1.00	4.5	0.77	1.00	6.5	0.06	0.64
$p = 50$											
LENS (default)	0.10	0.05	3.0	0.01	1.00	3.0	0.67	1.00			
$\hat{\mu}_u$ above the boundary, no test	0.10	0.05	3.1	0.03	1.00	3.1	0.67	1.00			
Top-$\lvert\widehat{G}\rvert$ by energy	0.10	0.05	3.0	0.01	1.00	3.0	0.67	1.00			
Both fits above the floor, no comparison	1.00	1.00	5.9	0.49	1.00	4.0	0.75	1.00			
$\hat{f}_D$ only: a_u above the floor	1.00	1.00	6.1	0.51	1.00	4.2	0.76	1.00			

TABLE XVII: Selection rules on real data: mean list size and test loss over five splits (Poisson deviance for freMTPL2 at three training sizes per model; log loss for credit). Each list is refitted on top of the same additive model f_0.

	freMTPL2 (rows per model)						Credit			
	20,000		50,000		184,420		HELOC		UCI	
Rule	size	loss	size	loss	size	loss	size	loss	size	loss
Additive f_0	–	0.31587	–	0.31447	–	0.31285	–	0.55084	–	0.43118
Black box $\hat{f}_D$	–	0.31747	–	0.31002	–	0.30536	–	0.56082	–	0.43097
LENS (default)	6.6	0.31318	9.0	0.31168	12.0	0.30950	2.0	0.55076	5.2	0.43028
LENS, $\kappa = 0$	6.6	0.31318	9.0	0.31168	13.0	0.30945	23.2	0.55408	35.2	0.43177
LENS, $b_u = \eta$	11.4	0.31340	14.8	0.31160	17.2	0.30938	55.4	0.55590	67.0	0.43266
$\hat{\mu}_u$ above the boundary, no test	6.6	0.31318	9.0	0.31168	12.2	0.30950	2.6	0.55038	5.4	0.43010
Top-$\lvert\widehat{G}\rvert$ by energy	6.6	0.31414	9.0	0.31251	12.0	0.31019	2.0	0.55066	5.2	0.43024
$\hat{f}_D$ only: a_u above the floor	25.8	0.31382	26.4	0.31166	23.4	0.30921	30.2	0.55507	41.0	0.43173
Both fits above the floor, no comparison	22.2	0.31368	22.6	0.31164	21.6	0.30922	16.8	0.55345	15.4	0.43098
All pairs	36.0	0.31432	36.0	0.31189	36.0	0.30912	253.0	0.59148	253.0	0.44213

The comparison removes what one fit learned. Both rules without the comparison return a pair in every run of the two additive designs, where every pair is false; LENS does so in at most one run in ten (Table XVI). On Sparse their lists grow from 3.0 to 4.3–6.1 pairs with an FDP of 0.27–0.51, against 0.01 for LENS, and on F1–F10 their FDP is 0.04–0.06, against 0, for a power gain of 0.02. Using the second model only as a second threshold barely helps: *both fits above the floor* is almost as long as $\hat{f}_D$ *only*, because a model fitted to noise leaves energy in many pairs; the two fits agree on how strong a pair is, not on its pattern. The same holds on real data (Table XVII). On freMTPL2 the two rules keep 22–26 of the 36 pairs, two to four times the LENS list. With 20,000 policies per model their loss is higher than that of LENS; with more data it is lower by at most 3×10^{-4}, the gain of adding nearly every pair. On the credit data they keep 15–41 pairs, against 2.0 and 5.2, and their loss is higher than that of LENS on both data sets; on HELOC it is higher than that of the additive model.

The ranking alone cannot set the list. Cut at the length of the LENS list, the energy ranking matches LENS in size, FDP, and power on every synthetic design: there the true pairs are also the strongest. The length, however, comes from LENS, and on real data the strongest pairs are not always the reproduced ones. On freMTPL2 the top-$\lvert\widehat{G}\rvert$ list has a higher test deviance than the LENS list at every training size (0.31414 against 0.31318, 0.31251 against 0.31168, and 0.31019 against 0.30950) and in each of the 15 splits. On the credit data the two lists have the same loss to within 10^{-4}.

The test makes the list error-controlled at no cost in accuracy. Keeping every pair whose point estimate $\hat{\mu}_u$ clears the boundary gives the LENS list in all but a few runs: it differs only at $p = 50$ on Sparse and Pair-corr (3.1 pairs each; FDP 0.03 against 0.01 on Sparse), on full freMTPL2 (12.2 pairs against 12.0), and on the credit data (2.6 against 2.0 and 5.4 against 5.2), with losses within 4×10^{-4} of those of LENS. With the evaluation sets used here the standard errors are small, so the list is set by the boundary, as the insensitivity to α in Section IV-D also shows. The pairs the point rule adds are those within sampling error of the boundary, including the false pairs at $p = 50$ on Sparse. The test is what keeps them out and what gives the list the error control of Theorem 2; the point rule has no such guarantee, and nothing bounds its errors when the evaluation set is small.

Summary. The second model and the comparison remove the patterns that only one fit learned, the boundary sets how long the list is, and the test makes the list error-controlled. Removing the comparison brings back false pairs on every design and longer, worse lists on real data; removing the test changes almost nothing on these sample sizes but forfeits the guarantee.