arXiv:2611.14311v1 [cs.AI] 10 Nov 2026

Multimodal Abstention That Preserves Acquired Correct Answers

Jinyuan Zhang
202621116012480@stu.hubu.edu.cn

Abstract

Adding abstention to a task-adapted multimodal model can discard correct answers, including capabilities acquired during adaptation. We formulate abstention repair as maximizing correct abstention on unanswerable questions under a budget on newly lost correct answers. We propose ReFrain (Regression-Free Abstention), which combines a linear probe on first-token logits, threshold calibration on held-out image clusters, and a deterministic veto that replaces rejected outputs with the displayed "none of the above" option while preserving all other responses. "Regression-free" denotes zero loss on the calibration set at the zero-budget target. All abstention methods in our main comparison start from the same answering-SFT checkpoint; we separately track lost correct answers, repaired errors, and retained SFT gains. On Qwen2.5-VL-3B, ReFrain at zero budget retains all 57 SFT-gained answers with no losses and improves correct abstention by 7.30 percentage points. At a 0.5% calibration budget, it reaches 74.05% correct abstention with 8 losses and retains 55 gains; CLIP-UP EmbLoRA reaches 64.65% with 62 losses, 41 repairs, and 40 retained gains. The same probe at its default threshold loses 298 correct answers. With models and thresholds frozen, confirmation on 228 independent question clusters yields 66.05% correct abstention and no losses at the 0.5% point, exceeding EmbLoRA by 32.72 percentage points. Ordinary refusal SFT achieves higher correct abstention at greater regression cost. Code and reproduction materials: https://github.com/ModeLiberaL/ReFrain.

1 Introduction

After a multimodal model has been adapted to its task, it still needs the ability to abstain: to stop guessing on inputs where the evidence is insufficient or the question is invalid. This order is close to practice: the model first learns to do the task, and its behavior when it cannot answer reliably is constrained afterwards. The order also determines how the cost of the abstention stage should be counted. Adaptation has already made the model answer a set of questions correctly, some of which adaptation has only just corrected; if the abstention stage turns these answers back into errors, overall accuracy may change very little while the specific capabilities that users can rely on have changed. We call adding abstention to an already-adapted answer model *abstention repair*, and we measure its cost against the records that this model answered correctly.

MM-UPD turns part of the answering boundary into a measurable task: the correct answer is absent from the candidate set (absent answer detection, AAD), the candidate set is incompatible with the question (incompatible answer set detection, IASD), or the image is incompatible with the question (incompatible visual question detection, IVQD) [16]. In these cases, stronger domain knowledge does not necessarily make the current input solvable. For example, knowing the general properties of an object cannot replace the corresponding evidence that is missing from the current image, and knowing the answer to the open-ended question does not make a candidate set without the correct answer yield a correct option. Joint evaluation on Standard and UPD therefore requires both valid answering and correct abstention.

Existing approaches differ in where the abstention decision is made: continuing to train the model on abstention examples [10], injecting a visual–textual matching signal into the input [22], perturbing the input

several times and deciding by the stability of the outputs [33], or training a classifier that abstains directly at the output [34]. TAPS separates accuracy training from post-hoc selection and compares this approach with abstention training at matched coverage [32]. We study the subsequent upgrade of an already-adapted answerer: in our main comparison, all abstention schemes start from the same checkpoint obtained by supervised fine-tuning (SFT) on answering, the *answer-SFT checkpoint*, and we record which answers each scheme changes relative to that checkpoint. There are two reasons. First, net accuracy lets losses and corrections cancel: at our common starting point, the Standard accuracy of CLIP-UP EmbLoRA is only 0.37 percentage points below the starting point, and this difference contains 62 lost original correct answers and 41 corrections. Second, only with the same starting point can the answers changed by each scheme be compared record by record, including the answers that answer SFT has just corrected.

Accordingly, we state abstention repair as a constrained problem: maximize the correct-abstention rate on unanswerable questions, subject to the original-correct loss not exceeding a budget ε. The original-correct loss is the negative flip rate of the model-update literature [28, 30]. The constrained form follows the existing formulation of regression-free model updates, in which the negative flip rate is constrained to be zero or at most a given constant [28]. The accompanying shared-checkpoint protocol specifies that all schemes use the same checkpoint, fit only on the training set, select operating points only on the calibration set, and neither tune thresholds nor select checkpoints on the evaluation set; besides net metrics, losses, corrections, and the retention of SFT-gained records are reported record by record.

Under this formulation we propose ReFrain (Regression-Free Abstention). It has three steps: a linear probe on the full-vocabulary logits of the first token gives an answerability score; on held-out image clusters, the largest threshold that meets the budget on the original-correct loss is selected; and when the score is below the threshold, the output is replaced with the "none of the above" (NOTA) option actually displayed in the question, while all other outputs stay unchanged. The scorer is the probe of LVLM-LP. Replacing the output directly on a classifier's decision also has precedents: LVLM-LP generates a template answer instead when its probe judges the input unanswerable, and the appendix of CLIP-UP gives a baseline that outputs a fixed refusal string after classification [22, 34]. For post-hoc abstention baselines that threshold a confidence score on a frozen model, the threshold has been selected by an aggregate metric on a validation set [7]. ReFrain calibrates the threshold against the original-correct loss of the adapted checkpoint; our controls compare this choice with the classifier's default threshold while holding the score and executor fixed. Executing by veto only, without rewriting, gives one-sidedness: an original correct answer can be lost only by being vetoed and is never changed to another substantive option; the correct-abstention rate on unanswerable questions is not below that of the starting point; and both the loss and the correct-abstention rate are monotone in the threshold. Calibration therefore needs only labeled answerable records, and changing the budget requires no retraining. "Regression-free" in the name refers to the calibration target at $\varepsilon = 0$: no original correct answer is lost on the calibration set. This is an empirical constraint on a finite calibration set. For a new image cluster that is exchangeable with the K calibration clusters, a standard order-statistic argument bounds the probability of any loss at this target by $1/(K + 1)$ (Property 3); the realized loss on the evaluation set is reported separately.

On the shared Qwen2.5-VL-3B [3] answer-SFT checkpoint B1, ReFrain loses no original correct answer on the 5,626 Standard records of the evaluation set, retains all 57 SFT-gained records, and raises the correct-abstention rate from 54.63% to 61.94%. ReFrain (ε=0.5%) loses 8, with a correct-abstention rate of 74.05%. The same probe, abstaining at its default probability threshold, has a correct-abstention rate of 92.16% and loses 298 original correct answers. From the same starting point, CLIP-UP EmbLoRA loses 62 and corrects 41, with a correct-abstention rate of 64.65%; plain abstention SFT has a correct-abstention rate of 87.47%, loses 140 and corrects 33, and does not pass the joint criterion (Section 3.6) on the evaluation set. With calibration and the executor fixed and only the scorer replaced, the two logits probes have correct-abstention rates of 74.05% and 75.40% at the 0.5% budget, and the other seven scorer configurations range from 56.59% to 67.20%. On two original base models without SFT, the zero-budget point of the same pipeline loses 2 and

4 original correct answers on the evaluation set, with correct-abstention rates 2.04 and 2.85 percentage points above EmbLoRA.

After the models and thresholds were frozen, we ran a single independent confirmation on AI2D images [12] that were not used in the development of this project. Neither operating point lost an original correct answer on the 912 Standard records, and the correct-abstention rate rose by 20.10 and 44.36 percentage points, respectively. The trained branches both lose and correct answers on the confirmation set: EmbLoRA loses 7 and corrects 13; plain abstention SFT loses 4 and corrects 5, passes the joint criterion, and has a correct-abstention rate 4.90 percentage points above ReFrain (ε=0.5%). ReFrain does not correct the errors of the answer model, which is the price of one-sidedness: on the confirmation set, the point estimates of the net Standard change of the two trained branches are positive, with intervals that contain zero.

This paper makes three contributions. First, problem and protocol: we state abstention repair after answer adaptation as a problem constrained by the original-correct loss, give a shared-checkpoint protocol and a taxonomy of schemes by where the abstention decision is made, and report the retention of SFT-gained records as a separate readout. Second, method: ReFrain and its one-sidedness properties, which reduce output-level abstention repair to a frontier parameterized by the threshold, on which the budget selects a point. Third, evidence: on the same checkpoint we compare ReFrain, two parameter-update schemes, and the configurations obtained by replacing the scorer; we repeat the same pipeline on two original base models; and we run a single confirmation on independent images after freezing. The main numbers are reported as direct paired differences within the same starting point; results across base models measure how the same pipeline performs when it is fitted separately on each.

2 Related Work

Unanswerable questions and selective prediction. UPD concerns inputs that are unsolvable by construction, and limits blanket abstention by evaluating each Standard question together with its unsolvable counterpart [16]. Reliable VQA studies the trade-off between answer coverage and the risk of error, and ReCoVERR further examines how to reduce unnecessary abstention [18, 25]. Together they show that the practical value of abstention depends both on stopping wrong outputs and on keeping valid answers.

Answer training followed by selection. TAPS combines accuracy-oriented post-training with a frozen answerer and an external confidence estimator [32]. Its main implementation uses GRPO for accuracy training and a hidden-state correctness probe for selection; its matched-coverage analysis separates improvements in answering from improvements in rejection. Our main experiment instead compares subsequent abstention schemes from one shared answer-SFT checkpoint, calibrates veto thresholds against its original-correct loss, and tracks the retention of records newly corrected by SFT. This protocol measures the cost of adding abstention after task adaptation, complementing TAPS's comparison of alternative training routes.

Answerability and uncertainty signals. LVLM-LP trains probes on the information in token distributions [34]; VL-Uncertainty estimates the instability of outputs under visual and textual perturbations [33]; CLIP-UP provides a visual–textual matching signal through structured CLIP correlation vectors [22]. These methods represent different sources of information and are the direct points of comparison for understanding the differences among the operating points studied here.

Correctness probes and confidence fusion. LRP compares hidden-state and attention probes in OCR question answering [31]; BICR trains a correctness scorer on the difference between the states under the real image and under a black image [13]; ACR fuses the maximum softmax probability (MSP) with a correctness prediction through a learnable gate [20]. These works provide different scoring mechanisms for a frozen

answerer. We plug multiple-choice adaptations of BICR and ACR into the calibrate and veto steps of ReFrain as scorers and compare them on the same B1 (Sections 4.5–4.6).

Trained abstention and the evidence boundary. CARA studies the detection of insufficient context [15]; MAD-MED trains with weighted abstention examples [10]; SAVOR adds calibration and abstention to a reinforcement-learning reward [4]; Variational VQA improves selective prediction through the variance of a parameter posterior [26]. They cover different stages—evidence recognition, generation policy, and uncertainty-aware training—and give a broader background for applications of abstention.

From signal to abstention action. CLIP-UP injects the correlation signal into the multimodal model, in an Emb and an EmbLoRA form; its appendix also gives a classify-then-abstain baseline: a logistic regression on the correlation vector decides whether the input is answerable, a fixed refusal string is output when the input is classified as unanswerable, and the model answers otherwise. The paper reports that this classifier labels only 46.27% of Circular Standard questions as answerable [22]. LVLM-LP also describes probe-guided decoding: when the probe judges the input unanswerable, a template answer is generated instead [34]. Dr. No steers sparse-autoencoder features on a frozen model to induce abstention, and compares on MM-UPD a set of post-hoc abstention baselines that threshold a confidence score, with thresholds selected on a validation split to maximize mean Dual accuracy [7]. The classify-then-abstain baseline, probe-guided decoding, and the thresholded post-hoc baselines belong, like ReFrain, to output-level override (Section 3.3): the answer model is frozen, and a scalar score decides whether the output is replaced. These methods and TAPS share the external-selection architecture with ReFrain. Our experiments compare scorers and threshold rules within this architecture, using the original-correct loss of the frozen answer model to select the operating point. Within the output-level override family we fix the executor and replace the scorer and the threshold source separately; schemes in the other families keep their own training and execution and are compared as wholes. The two kinds of comparison answer questions about information quality and about end-to-end effect, respectively.

Constrained threshold selection. Minimizing one type of error subject to an upper bound on the other is the setting of Neyman–Pearson classification [19]. This setting, and finite-sample methods such as conformal risk control [2], have been used to calibrate abstention thresholds. For language models, what is constrained is the rate at which wrong answers are not intercepted [17], or the rates at which unanswerable questions are answered and answerable ones are answered wrongly [23]; for vision–language models, what is constrained is the error or hallucination risk on the answered part [24, 29]. These works constrain erroneous outputs, and wrongly abstaining on correct answers is the side that is minimized or only reported. Our constraint runs in the opposite direction: what is constrained is the fraction of records answered correctly by the frozen answer model that are vetoed, and what is maximized is the correct-abstention rate on unanswerable inputs. Our calibration rule is an empirical constraint on a finite calibration set. At a nonzero budget it gives no population-level guarantee; at the zero budget, the order-statistic arguments of these two settings bound the probability of a loss on a new exchangeable image cluster (Property 3). The realized loss on the evaluation set is reported separately, with paired intervals.

Regression-free updates. The model-update literature calls a sample that the old model answers correctly and the new model answers wrongly a negative flip, measures regression by the negative flip rate, and reduces negative flips by training the new model with distillation or by using model ensembles [28, 30]. Among these, Xie et al. [28] write the regression-free model update as a constrained optimization problem, constraining the negative flip rate to be zero or at most a given constant, and then relax it into joint training with a distillation term. Another line keeps the predictions of the old model at inference time and chooses between, or fuses, the old and new predictions sample by sample [14, 21]. For language models, negative flips are also used to

measure the regression caused by model version updates and by preference alignment [5, 36]. We adopt the negative flip rate as the measure together with the constrained form, and change the object to the addition of abstention: the reference model is a checkpoint that has completed answer adaptation, and the objective is the correct-abstention rate on unanswerable inputs. ReFrain neither trains a new answer model nor chooses between two answer models; there is only one answer model, the only alternative output is abstention, and the veto threshold is calibrated against the original-correct loss on held-out image clusters. In the post-training setting, we additionally condition on the records that the previous stage newly made correct, to characterize how a capability is retained after it is acquired.

3 Problem, Protocol, and Method

3.1 Problem: Abstention Repair after Answer SFT

Records and correctness. A record x consists of an image, a question, and a set of candidate options, among which a "none of the above" (NOTA) option is displayed. S denotes the set of Standard records, whose correct answer is a substantive option and on which a wrongful abstention counts as an error. U denotes the set of UPD records, on which an abstention that conforms to the task definition counts as correct; selecting the displayed NOTA option is such an abstention. $c_g(i) \in \{0, 1\}$ indicates whether system g is correct on record i.

Answer model. Let B0 be the original multimodal answer model and B1 the model after a fixed answer SFT. Our main setting takes $f = \mathrm{B1}$ and freezes its parameters. This corresponds to a common order of use: the model is first adapted to the task and only afterwards given the ability to abstain, and the cost of the abstention stage is counted against the records that this model answered correctly. Section 4.1 additionally takes f to be each of two original base models, to examine the same pipeline from starting points without SFT.

Abstention repair schemes. An abstention repair scheme takes f, a training set T, and a calibration set CM, and outputs a repaired system g. T is used for fitting or training and CM for selecting an operating point or a model; the evaluation set takes part in no selection.

Loss and correction. The original-correct loss and the correction rate of g relative to f are

$$L(g; f) = \frac{1}{|S|} \sum_{i \in S} c_f(i)\,[1 - c_g(i)], \qquad F(g; f) = \frac{1}{|S|} \sum_{i \in S} [1 - c_f(i)]\, c_g(i). \tag{1}$$

The net change in accuracy satisfies $\Delta\mathrm{Acc} = F - L$. L is the fraction of records that f answers correctly and g answers wrongly, that is, the negative flip rate of the model-update literature [28, 30]; F is the fraction in the opposite direction.

Constrained problem. Given a budget $\varepsilon \geq 0$, abstention repair solves

$$\max_{g} \ \mathrm{Acc}_U(g) \quad \text{s.t.} \quad L(g; f) \leq \varepsilon, \tag{2}$$

where Acc_U is the correct-abstention rate on U. $\varepsilon = 0$ requires that no original correct answer be lost and is the preset primary operating point; $\varepsilon = 0.5\%$ is a preset additional operating point. Writing the regression-free model update as a constrained optimization problem, with the negative flip rate constrained to be zero or at most a given constant, has a precedent [28]; we adopt this form, restrict the update to adding

abstention to an already-adapted answer model, and take the objective to be the correct-abstention rate on unanswerable records.

The constraint is placed on L rather than on net accuracy because the net value lets losses and corrections cancel. One trained branch in Section 4.2 is only 0.37 percentage points below f on Standard, and this difference contains 62 lost original correct answers and 41 corrections. For a model that is already in use, what users encounter are these record-level regressions.

SFT-gained records. Define $R = \{i \in S : c_{\mathrm{B0}}(i) = 0,\ c_{\mathrm{B1}}(i) = 1\}$, the records that answer SFT has just corrected. The gain retention rate is

$$\mathrm{Retain}(g; R) = \frac{\sum_{i \in R} c_g(i)}{|R|}. \tag{3}$$

It answers whether the abstention stage concentrates its losses on the capability gained in the previous stage. L, F, and Retain have different denominators and meanings and are listed separately in the results.

3.2 Shared-Checkpoint Protocol

Answer SFT uses the 677 Standard records in T, trains for three epochs, and saves a single final checkpoint, B1. All abstention repair schemes start from this one B1.

T has 1,354 records. The calibration set CM has 4,567 records in 183 image clusters, of which 2,406 are Standard. The evaluation set E has 10,678 records in 416 image clusters, of which 5,626 are Standard and 5,052 are UPD. The image-cluster split isolates training, calibration, and evaluation, and the image cluster is the resampling unit of the paired intervals.

All schemes follow the same set of rules (Table 1). When the starting point is an original base model (Section 4.1), the same protocol applies, and scorers and thresholds are fitted separately on each base model.

Table 1: Rules of the shared-checkpoint protocol.

Resource	Rule
Answering starting point	The same B1 checkpoint
Fitting or training	T only
Selection of operating point or model	CM only
Evaluation	E and the confirmation set I (Section 3.7); no threshold tuning or checkpoint selection on them
Parsing and decoding	The same parser, the same NOTA mapping, greedy decoding
Readouts	All readouts of Section 3.6

Each compared scheme carries a fidelity label. A *full reproduction* uses the official code and all components. A *mechanism port* implements the mechanism of the original paper in our setting and discloses the hyperparameters that the paper does not give. A *partial implementation* lacks a component of the original method and does not enter the main tables under the name of the original method.

3.3 A Taxonomy of Abstention Repair Schemes

By whether the answer model is modified and by where the abstention decision is made, existing approaches fall into six families (Table 2).

Within the output-level override family there are two layers of comparison. The first fixes the score and changes the source of the threshold: the same score is executed at the default threshold of the original method,

Table 2: Families of abstention repair schemes. "Modifies f" states whether the parameters of f are changed; "Can rewrite" states whether a correct answer of f can be changed to another substantive option.

Family	What it does	Modifies f	Can rewrite	Representative methods
Output-level override	Freezes f; a scalar score decides whether the output is replaced with NOTA	No	No	ReFrain (this paper); the LVLM-LP probe abstaining at its default threshold [34], the CLIP-UP correlation-vector classifier [22], BICR [13], ACR [20]
Parameter update	Continues training f on abstention examples	Yes	Yes	Plain abstention SFT, CLIP-UP EmbLoRA (which also injects CLIP correlation vectors) [22], MAD-MED [10], SAVOR [4], VL-Calibration [27], Reliable-V-RL [11]
Input or representation injection	Adds a signal to the input embeddings or to intermediate activations	Only the added projection or steering parameters	Yes	CLIP-UP Emb [22], Dr. No [7]
Inference-time sampling	Perturbs the input several times and decides by the consistency of the outputs	No	Depends on the execution	VL-Uncertainty [33]
Learned routing	Trains the model to output an answer-or-abstain decision before answering	Yes	Yes	Rep2Act [8]
Prompting	Changes the instruction only	No	Yes	The additional-instruction setting of MM-UPD [16]

or calibrated by the rule of Section 3.4. The second fixes calibration and the executor and replaces the score. Schemes in the other families are compared as wholes, together with their own training and execution.

On the same B1 we evaluate the scorers of the output-level override family (Sections 4.3–4.6) and two schemes of the parameter-update family (Section 4.2). CLIP-UP Emb, from the input-injection family, and the inference-time sampling method are evaluated on the original base models (Section 4.1). Dr. No, the learned-routing family, and the prompting family have no same-setting results yet (Section 6).

3.4 ReFrain

ReFrain (Regression-Free Abstention) belongs to the output-level override family and consists of three steps: score, calibrate, and veto (Algorithm 1). It changes neither the parameters of f nor the substantive answers of f.

Algorithm 1 ReFrain

Input: frozen answer model f; training set T with Standard/UPD labels; the Standard records of CM and the correctness of f on them; budget ε

Output: g_{τ_ε}

1: **Score.** Fit a scorer s on T; a low score indicates that abstention is more warranted.
2: **Calibrate.** $\tau_\varepsilon \leftarrow \max\{\tau \in \mathcal{T} : \widehat{L}_{\mathsf{CM}}(g_\tau; f) \le \varepsilon\}$, where $\mathcal{T} = \{-\infty, +\infty\} \cup \{s(x) : x \in \mathsf{T}\}$.
3: **Veto.** $g_{\tau_\varepsilon}(x) \leftarrow n(x)$ if $s(x) < \tau_\varepsilon$; otherwise $g_{\tau_\varepsilon}(x) \leftarrow f(x)$.

Score. The default scorer is a linear probe. It reads the full-vocabulary logits $\phi_f(x)$ at the position where f begins to generate its answer and outputs an answerability score $s(x) = w^\top \phi_f(x) + b$. The probe is trained

on T with Standard as the positive class and UPD as the negative class; the regularization strength is selected by five-fold cross-validation over image clusters within T, with mean binary cross-entropy as the criterion. The feature and the form of the probe follow LVLM-LP [34]. $\phi_f(x)$ comes from the same forward pass that produces the answer, so scoring adds no forward passes.

Calibrate. The candidate threshold set is $\mathcal{T} = \{-\infty, +\infty\} \cup \{s(x) : x \in \mathsf{T}\}$, including the no-veto and all-veto endpoints. For a budget ε,

$$\tau_\varepsilon = \max\{\tau \in \mathcal{T} : \ \widehat{L}_{\mathsf{CM}}(g_\tau; f) \le \varepsilon\}, \tag{4}$$

where $\widehat{L}_{\mathsf{CM}}$ is the empirical original-correct loss on CM, whose denominator is all 2,406 Standard records of CM. $\varepsilon = 0.5\%$ allows at most 12 calibration losses, and the default scorer attains 11. The threshold is determined directly by how many original correct answers may be lost, and not by the default probability threshold of a classifier, an aggregate metric, or an error risk. Under veto-only execution, this rule is the empirical form of Neyman–Pearson threshold selection: the constrained error is the veto of a record that f answered correctly. Selecting a threshold by the empirical error rate does not guarantee with high probability that the population error rate stays within the same level [19]; we therefore report the constraint on the calibration set and the realized loss on the evaluation set separately. For $\varepsilon = 0$, Property 3 states the level that does hold for a new image cluster under an exchangeability assumption.

Veto. Let $n(x)$ be the NOTA letter actually displayed in the question, and

$$g_\tau(x) = \begin{cases} n(x), & s(x) < \tau, \\ f(x), & \text{otherwise.} \end{cases} \tag{5}$$

The system has only a veto: it can replace the output of f with NOTA, but not with another substantive option. Records that are not vetoed keep the original output of f. Replacing the output on a classifier's decision has precedents: the probe-guided decoding of LVLM-LP generates a template answer instead when the input is judged unanswerable, and the classification baseline in the appendix of CLIP-UP outputs a fixed refusal string [22, 34]. What ReFrain substitutes is the NOTA option displayed in the question.

The scorer is a replaceable component. The first step can be replaced with any scalar score, with calibration and veto unchanged. We use this to compare, under the same calibration and executor, a probe with a fixed regularization strength ($C = 1$), the maximum softmax probability, the negative entropy, the CLIP correlation-vector classifier, and the scores of BICR and ACR (Sections 4.3–4.6). Of these, BICR and ACR predict whether the answer of f is correct, whereas the probes and the CLIP classifier predict whether the input is answerable.

Name and notation. In what follows, ReFrain without annotation denotes the default configuration: the cross-validated probe at $\varepsilon = 0$. "Regression-free" in the name refers to this calibration target: not a single record that f answered correctly is lost on the calibration set. It is an empirical constraint on a finite calibration set. For a new image cluster that is exchangeable with the calibration clusters, Property 3 bounds the probability of any loss; for data from other sources there is no guarantee, and the realized loss on the evaluation set is reported separately. The operating point at $\varepsilon = 0.5\%$ is written ReFrain (ε=0.5%) and is the relaxed version. When the scorer is replaced, we write ReFrain (scorer name). The names of the original methods, such as LVLM-LP, CLIP-UP, BICR, and ACR, are used only for results obtained with the original method's own threshold and execution.

3.5 Properties

The following two properties depend only on the veto rule $g_\tau(x) \in \{f(x), n(x)\}$ and hold for any scorer.

Property 1 (One-sidedness)**.** (a) For every Standard record, $c_g \le c_f$. Hence $F(g_\tau; f) = 0$ and $\Delta\mathrm{Acc} = -L$: a veto does not correct an error of f, nor does it change a correct answer of f to another substantive option; the only cost is the original correct answers that are vetoed. (b) For every UPD record, $c_g \ge c_f$; the correct-abstention rate is not below that of f at any threshold. (c) As τ increases, the set of vetoed records only grows, and $L(g_\tau; f)$ and $\mathrm{Acc}_U(g_\tau)$ are both non-decreasing in τ. The whole method is therefore a frontier parameterized by τ, on which the budget ε selects a point (Section 4.7).

Part (a) holds because NOTA is always a wrong answer on a Standard record; part (b) holds because the displayed NOTA option is always a correct answer on a UPD record. Parameter-update schemes have neither (a) nor (b): they can lose and correct Standard answers at the same time, and can fall below f on some UPD records.

Property 2 (The loss is determined by one score distribution)**.**

$$L(g_\tau; f) = \frac{1}{|S|} \sum_{i \in S} c_f(i)\, \mathbf{1}[s(x_i) < \tau]. \tag{6}$$

The loss equals the Standard accuracy of f multiplied by the fraction of the Standard records answered correctly by f whose score is below τ. Calibration is therefore the selection of a point on a monotone curve, and it needs only answerable records with their original task labels and the correctness of f on them; unanswerable examples are used only to train the scorer. Changing the budget only requires selecting the threshold again, not retraining. At $\varepsilon = 0$, τ_0 is the largest candidate threshold that does not exceed the lowest score among the CM Standard records that f answers correctly. When $f = \mathrm{B1}$, the gain retention rate is the restriction of the same quantity to R:

$$\mathrm{Retain}(g_\tau; R) = 1 - \frac{|\{i \in R :\ s(x_i) < \tau\}|}{|R|}. \tag{7}$$

The third property concerns new data and, unlike the first two, needs a sampling assumption. For an image cluster k, let m_k be the lowest score among its Standard records that f answers correctly, with $m_k = +\infty$ if there is none. By Property 2, g_τ loses an original correct answer in cluster k exactly when $m_k < \tau$, and the zero-budget rule gives $\tau_0 \le \min_{k \le K} m_k$ over the K image clusters of CM.

Property 3 (Zero-budget loss on a new image cluster)**.** Suppose that f and the scorer s are fixed without reference to the image clusters of CM or to the new image cluster, and that, conditional on the training set and on everything else that determines f and s, the K image clusters of CM and the new image cluster are exchangeable. (a) The probability that g_{τ_0} loses at least one original correct answer in the new cluster is at most $1/(K+1)$. (b) If, under the same conditioning, the clusters are moreover i.i.d., then for every $\alpha \in (0, 1)$, with probability at least $1 - (1-\alpha)^K$ over the calibration clusters, the probability of such a loss in a new cluster is at most α.

Both parts follow from a standard order-statistic argument and are stated for completeness, not as a new result. Part (a) is the case of conformal risk control in which the target level is $1/(K+1)$, so that the rule requires zero empirical risk [2]; part (b) is the bound for the extreme order statistic in the Neyman–Pearson umbrella algorithm [19]. Proofs are given in Appendix A. With $K = 183$, part (a) gives $1/184 \approx 0.54\%$, and part (b) gives $\alpha \approx 1.62\%$ at the 95% level. The probability in part (a) is over the joint draw of the calibration clusters and the new cluster and is not conditional on the realized calibration set; part (b) is the conditional statement.

The scope of Property 3 is narrow. It is stated for image clusters and not for records; it bounds the probability for a single new cluster and is not simultaneous over the clusters of an evaluation set; it covers only $\varepsilon = 0$ and says nothing about ReFrain (ε=0.5%); and it does not apply to data from another source, such as the confirmation set. In the protocol of Section 3.2, f and the default scorer are fitted on T only. For E, exchangeability with the image clusters of CM is an assumption about how the data were partitioned; it holds by design when image clusters are assigned to CM and E uniformly at random. Under it, the expected number of the 416 image clusters of E in which an original correct answer is lost is at most $416/184 \approx 2.26$; the observed number is 0 for f = B1 (Section 4.2). Both sets were also used while the method was developed, so the first condition is not verified for the design choices; the observed number is descriptive, a consistency check and not a test of the bound.

Cost. One-sidedness means that ReFrain gives up correction, and its Standard accuracy cannot be higher than that of f. Schemes that can correct the errors of f can be higher than f: on the confirmation set, the net Standard point estimates of the two trained branches are both above f, with intervals that contain zero (Section 4.8). The gain in the correct-abstention rate equals the fraction of UPD records that f answered wrongly and whose score is below τ, so the ceiling on abstention coverage is set by how well the scorer separates the two kinds of records.

3.6 Evaluation Protocol

Readouts. We report Standard accuracy, the UPD correct-abstention rate, Circular Dual, the number of original correct answers lost, the number of corrections, and the gain retention rate. Circular aggregates over the permutations of the candidates; Dual requires that the Standard side and the UPD side of the same question both satisfy their conditions.

Joint criterion. Relative to the common starting point f, the lower bound of the paired 95% interval of ΔStandard is above -1 percentage point, and at the same time the lower bound for ΔUPD is above zero. The intervals use 2,000 paired bootstrap resamples over image clusters [6]. An advantage of one scheme over another is tested separately with direct paired differences. Passing the criterion relative to f and being better than another scheme are reported separately.

Native view. Each scheme is reported at its own default operating point, with all readouts and the joint criterion.

3.7 Independent Confirmation of the Frozen Setting

The confirmation set I comes from the non-test images of the complete official AI2D image collection [12].[1] Its construction excluded the images of the project's earlier MM-UPD and CLIP-UP data, the complete official AI2D test images, and the images reserved by a parallel study, and also excluded exact and near-duplicate images. The original questions keep their official answers; AAD removes the correct option, IASD replaces the option set, and IVQD replaces the image with a mismatched foreign image. After a review of text and image semantics by fixed rules, 228 of 300 candidate clusters were kept: 96 AAD, 89 IASD, and 43 IVQD. Standard records rotate over all four substantive options, AAD records over three, and the other UPD records over four; NOTA is fixed in the last position. The final set has 271 images, 912 Standard records, and 816 UPD records. Each IVQD foreign image is used in only one cluster and is resampled together with all records of the original question.

[1] https://allenai.org/data/diagrams

B1, the EmbLoRA and plain abstention SFT checkpoints, the probe coefficients, and the two CM thresholds were frozen before inference on the confirmation set; the thresholds are -10.226138449867452 and -6.183404341319426. The confirmation set takes no part in training or calibration, and neither the selection nor the semantic review read the predictions of any branch under evaluation. B0 only generates Standard answers, to identify the SFT-gained records. All comparisons keep the original parser, greedy decoding, and the joint criterion relative to B1; the two operating points and the four paired comparisons between methods were fixed before inference. The results are reported separately from E, with both the UPD micro-average and the task macro-average. Details of the construction are given in Appendix E.

4 Results

Unless stated otherwise, the numbers in this section are measured on E, differences between rates are given in percentage points (pp), and intervals are paired 95% intervals over image clusters (Section 3.6). In the tables, UPD is the correct-abstention rate on UPD records, "Lost" and "Corrected" count Standard records relative to the starting point f, and "Retained" counts SFT-gained records.

4.1 Two Original Base Models under the Same Protocol

Table 3 takes f to be each of two original base models without answer SFT, Qwen2.5-VL-3B and InternVL3-8B [35]. The zero budget is a constraint on CM. On E, the zero-budget points on the two base models lose 2 and 4 original correct Standard answers, and their correct-abstention rates are 2.04 and 2.85 pp above EmbLoRA. The probe and the threshold are fitted separately on each base model. These losses are consistent with the scope of Property 3, which bounds the probability of a loss in a new image cluster and does not state that the loss on the evaluation set is zero.

Table 3: Two original base models without answer SFT, on E. Each scheme is fitted and calibrated separately on each base model, and the joint criterion is relative to the base model without an abstention stage. ReFrain takes f to be the base model. VL-Uncertainty is reported with its native execution and with its score under the common executor.

Original base model	Scheme	Standard (%)	UPD (%)	Joint criterion
Qwen2.5-VL-3B	No abstention stage	86.72	56.43	reference
	ReFrain	86.69	67.85	passed
	CLIP-UP EmbLoRA	86.76	65.82	passed
	CLIP-UP Emb	86.21	68.01	not passed
	VL-Uncertainty (native execution)	80.16	75.36	not passed
	VL-Uncertainty score, common executor	86.70	56.63	passed
InternVL3-8B	No abstention stage	92.18	74.33	reference
	ReFrain	92.11	80.84	passed
	CLIP-UP EmbLoRA	92.09	77.99	passed
	CLIP-UP Emb	80.48	53.35	not passed

CLIP-UP [22] behaves differently across base models and forms of attachment. EmbLoRA passes the joint criterion on both base models; Emb degrades markedly on InternVL3-8B under the common prompt. This observation corresponds to the present combination of input, injection, and training settings, and we do not read it as a statement about the original method in its own setting (Appendix C). The native execution of VL-Uncertainty [33] and its score under the common executor give clearly different operating points and are therefore reported separately.

4.2 A Common Starting Point after Answer SFT

Answer SFT corrects 57 answers that B0 got wrong and loses 29 that B0 got right, a net gain of 28 correct answers. The 57 SFT-gained records lie in 28 image clusters. Table 4 gives the results of the abstention branches that start from B1.

Table 4: Abstention branches from the common starting point B1, on E. Lost and corrected answers are counted relative to B1 among the 5,626 Standard records; retained records are out of the 57 SFT-gained records.

Branch	Standard (%)	UPD (%)	Circular Dual (%)	Lost / corrected	Retained
B0	86.72	56.43	40.26	—	—
B1	87.22	54.63	38.86	0 / 0	57 / 57
ReFrain	87.22	61.94	46.22	0 / 0	57 / 57
ReFrain (ε=0.5%)	87.08	74.05	55.73	8 / 0	55 / 57
CLIP-UP EmbLoRA	86.85	64.65	47.82	62 / 41	40 / 57
Plain abstention SFT	85.32	87.47	66.41	140 / 33	23 / 57

Both ReFrain operating points and EmbLoRA pass the joint criterion relative to B1. ReFrain keeps every original correct answer of B1 and raises the correct-abstention rate by 7.30 pp, with a 95% interval of [6.21, 8.37]. Its correct-abstention rate is 2.71 pp below that of EmbLoRA, with a direct paired interval of [−4.07, −1.32], which shows a trade-off between stricter retention and abstention coverage.

The evaluation-set loss of ReFrain (ε=0.5%) is 8 of 5,626 Standard records, of which 2 are SFT-gained records and 6 are other records that B1 answered correctly. Relative to EmbLoRA, its correct-abstention rate is 9.40 pp higher, with an interval of [7.77, 11.05]; it loses 54 fewer original correct answers, a difference of 0.96 pp with an interval of [0.52, 1.47]. The point estimate of its Standard accuracy is 0.23 pp higher, with an interval of [−0.37, 0.87].

The gain retention rate is 96.49% against 70.18%, a direct difference of 26.32 pp with an interval of [7.35, 49.06]. EmbLoRA loses 17 SFT-gained records and 45 other correct records, and corrects 41 records that B1 answered wrongly. The corrections are a real gain of this branch; the two kinds of transition together make up net accuracy.

Plain abstention SFT starts from the same B1, uses the same 1,354 training records and three epochs of training, and takes the LoRA configuration and the learning rate of EmbLoRA. Its target is the correct option or the NOTA letter in the common multiple-choice format. This branch has a higher correct-abstention rate and a higher Circular Dual than the other branches in Table 4; it loses 34 SFT-gained records and 106 other answers that B1 got right, and corrects 33 old errors. Its Standard accuracy falls by 1.90 pp relative to B1, with a 95% interval of [−2.84, −1.01], so it does not pass the accuracy part of the joint criterion.

Relative to plain abstention SFT, ReFrain (ε=0.5%) is 1.76 pp higher on Standard, with an interval of [0.89, 2.68], and 13.42 pp lower in correct-abstention rate, with an interval of [−14.88, −11.90]. The two offer, respectively, stronger retention of correct behavior and higher abstention coverage. Figure 1 shows the abstention gain, the transitions of correct answers, and the retention of SFT-gained records together.

After SFT, the correct-abstention rate of B1 is 1.80 pp below that of B0. Relative to B0, the final correct-abstention rates of ReFrain, ReFrain (ε=0.5%), and EmbLoRA are higher by 5.50, 17.62, and 8.21 pp. The final gain of the abstention stage therefore exceeds the recovery of the drop in abstention caused by this SFT.

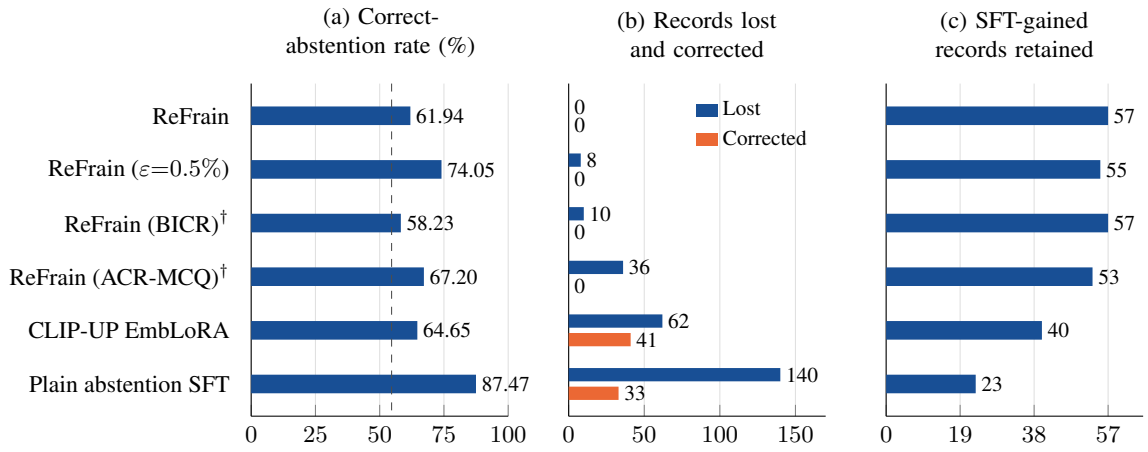

Figure 1: Results on the same B1, on E. (a) Correct-abstention rate on UPD records; the dashed line marks B1 (54.63%). (b) Original correct answers lost and old errors corrected, relative to B1. (c) SFT-gained records retained, out of 57. †The scorer is replaced with that of BICR or of ACR-MCQ (standardized features) and calibrated at the 0.5% budget under the common executor (Sections 4.5 and 4.6).

4.3 Scorers under the Same Calibration and Executor

Table 5 replaces the scorer and keeps the calibration rule and the executor. The ordering of the probe with fixed $C = 1$ and the cross-validated probe changes with the budget. At $\varepsilon = 0$, the correct-abstention rate of the cross-validated probe is 2.06 pp higher, with an interval of [1.58, 2.58]; at $\varepsilon = 0.5\%$, that of the $C = 1$ probe is 1.35 pp higher, with an interval of [0.86, 1.84], and this probe loses 5 more original correct answers. The two probes retain the same number of SFT-gained records.

Table 5: Scorers under the same calibration and executor on B1, on E. MSP is the maximum softmax probability. Every row is ReFrain with the named scorer.

Scorer	ε (%)	Standard (%)	UPD (%)	Lost	Retained
Cross-validated probe (default)	0	87.22	61.94	0	57 / 57
Probe with fixed $C = 1$	0	87.22	59.88	0	57 / 57
Cross-validated probe (default)	0.5	87.08	74.05	8	55 / 57
Probe with fixed $C = 1$	0.5	86.99	75.40	13	55 / 57
MSP	0.5	86.97	59.34	14	54 / 57
Negative entropy	0.5	87.17	57.30	3	57 / 57

At these operating points, the learned logits probes have higher correct-abstention rates than MSP and the negative entropy. The common framework can reuse different scorers, and the choice of regularization is a choice of operating point. This table is a comparison of components: the threshold calibration and the executor of the $C = 1$ probe are those of this paper, and the probe at its default probability threshold is listed in Table 7.

4.4 A Common-Executor Control for CLIP Correlation Features

Appendix F.2 of CLIP-UP gives a mechanism that abstains directly after a classification of the correlation vector [22]. We use the official Structure-CLIP weights [9], multiply the normalized image embedding element-wise with each of the four question–candidate text embeddings, and concatenate the products; the

NOTA candidate is excluded. The original paper does not give the complete fitting hyperparameters of this classifier, so we report two mechanism ports: an unstandardized fit with $C = 1$, and a fit selected by cross-validation over image clusters within T. The cross-validation chooses whether to standardize the features and a value of $C \in \{0.01, 0.1, 1, 10, 100\}$, and it selects the unstandardized fit with $C = 100$. The threshold rule and the executor are the same as for the logits probe (Table 6).

Table 6: CLIP correlation-vector classifier on B1, on E, at its default threshold and under the calibration and executor of ReFrain. The calibrated rows are ReFrain with the CLIP classifier as the scorer.

Classifier fit	Operating point	Standard (%)	UPD (%)	Lost	Retained	Joint criterion
Fixed $C = 1$	Default threshold ($p = 0.5$)	48.99	92.40	2151	30 / 57	not passed
	Calibrated, $\varepsilon = 0$	87.22	55.09	0	57 / 57	passed
	Calibrated, $\varepsilon = 0.5\%$	86.70	57.46	29	55 / 57	not passed
Cross-validated	Default threshold ($p = 0.5$)	66.57	85.25	1162	51 / 57	not passed
	Calibrated, $\varepsilon = 0$	87.22	54.75	0	57 / 57	not passed
	Calibrated, $\varepsilon = 0.5\%$	87.08	56.59	8	57 / 57	passed

At the common 0.5% budget, the cross-validated CLIP classifier and the cross-validated logits probe both lose exactly 8 Standard answers; the correct-abstention rate of the logits probe is 17.46 pp higher, with a 95% interval of [15.71, 19.25]. The CLIP classifier retains all 57 SFT-gained records and the logits probe retains 55, which shows that the same total loss can correspond to different sets of lost records. The classifier operating points that are not calibrated against the loss have high correct-abstention rates and veto a large number of original correct answers; after calibration, accuracy is retained, but the additional abstention gain is small. At the low-loss operating points that we tested, this supports the internal logits score over the CLIP correlation features.

Table 7 puts the default threshold next to the common calibration for the two cross-validated classifiers, all computed on the same B1.

Table 7: Default threshold against common calibration for the two cross-validated classifiers on B1, on E.

Feature	Decision rule	Standard (%)	UPD (%)	Lost	Retained	Joint criterion
CLIP correlation vector	Default threshold ($p = 0.5$)	66.57	85.25	1162	51 / 57	not passed
	Calibrated, $\varepsilon = 0.5\%$	87.08	56.59	8	57 / 57	passed
First-token logits	Default threshold ($p = 0.5$)	81.92	92.16	298	45 / 57	not passed
	Calibrated, $\varepsilon = 0.5\%$	87.08	74.05	8	55 / 57	passed

The default classification threshold does not retain the original accuracy with either feature; the common calibration against the loss brings both to operating points that pass the joint criterion. Calibration also lowers abstention coverage, and the logits probe keeps more of the gain in correct-abstention rate at the calibrated operating point.

4.5 A Multiple-Choice Adaptation of Adaptive Confidence Fusion

ACR fuses MSP with a correctness signal and a gate signal that are produced by two MLPs [20]. We adapt its inputs to the final-prompt hidden state of B1 and the logits of the options A–E, keep the answer model frozen and the 15 epochs of training, and report one variant with raw features and one with features standardized using statistics of T only. We call this adaptation ACR-MCQ. Both variants use the same loss calibration and NOTA executor (Table 8).

Table 8: ACR-MCQ on B1, on E. The calibrated rows are ReFrain with the ACR-MCQ score as the scorer.

Features	Operating point	Standard (%)	UPD (%)	Lost	Retained
Raw	Default threshold ($p = 0.5$)	85.76	67.00	82	25 / 57
	Calibrated, $\varepsilon = 0$	87.22	55.56	0	57 / 57
	Calibrated, $\varepsilon = 0.5\%$	86.97	59.32	14	54 / 57
Standardized	Default threshold ($p = 0.5$)	82.28	82.21	278	23 / 57
	Calibrated, $\varepsilon = 0$	87.22	57.98	0	57 / 57
	Calibrated, $\varepsilon = 0.5\%$	86.58	67.20	36	53 / 57

With raw inputs, the gate equals 1 on the whole training set, and the deployed behavior equals MSP over the option space A–E; standardization on T makes the gate vary with the input, and the correctness AUROC on E rises from 0.806 to 0.913. At the 0.5% calibration budget, the standardized variant has higher abstention coverage and also a larger evaluation-set loss. Relative to this operating point, ReFrain (ε=0.5%) is 0.50 pp higher on Standard, with a 95% interval of [0.21, 0.81], and 6.85 pp higher in correct-abstention rate, with an interval of [5.32, 8.32].

4.6 A Visual-Contrast Correctness Probe

BICR trains a shared MLP on the final-prompt states under the real image and under a black image, with the correctness of the actual answer of B1 as the label and a joint objective of a weighted binary cross-entropy, a Brier term, and a ranking loss on positive examples [13]. Within T we make an 80/20 split by image cluster, run 50 trials of an Optuna search [1], and then train the selected configuration on the full T for 79 epochs. CM is used only to calibrate the common loss threshold (Table 9).

Table 9: BICR on B1, on E. The calibrated rows are ReFrain with the BICR score as the scorer.

Operating point	Standard (%)	UPD (%)	Lost	Retained	Joint criterion
Default threshold ($p = 0.5$)	66.81	95.19	1148	22 / 57	not passed
Calibrated, $\varepsilon = 0$	87.22	55.48	0	57 / 57	passed
Calibrated, $\varepsilon = 0.5\%$	87.04	58.23	10	57 / 57	passed

The correctness AUROC of BICR on E is 0.904. At the common 0.5% budget it retains all SFT-gained records; the correct-abstention rate of ReFrain (ε=0.5%) is 15.82 pp higher, with a 95% interval of [14.25, 17.41], and the difference on Standard is 0.04 pp in favor of ReFrain (ε=0.5%), with an interval of [−0.14, 0.20]. The ranking of overall correctness and the abstention gain at low loss are different readouts. BICR and ACR predict whether the answer is correct, whereas our probe and the CLIP classifier predict whether the input is answerable, with Standard against UPD as the classes; the differences among these complete scorers therefore combine features, targets, and fitting procedures.

4.7 Budget Curves and Failure Types

With the scorer fixed, we read the same threshold rule at CM budgets of 0%, 0.1%, 0.25%, 0.5%, and 1%. The correct-abstention rate of the default scorer on E is then 61.94%, 66.83%, 71.81%, 74.05%, and 82.90%, and the number of original correct answers lost is 0, 0, 5, 8, and 70. A wider budget increases abstention coverage and also the realized loss. The curve is the single-parameter frontier of Property 1(c); it shows the trade-off among operating points, and the preset primary and additional operating points stay unchanged (Figure 2). At the 1% budget, the realized loss rate on E, 1.24%, is above the nominal budget.

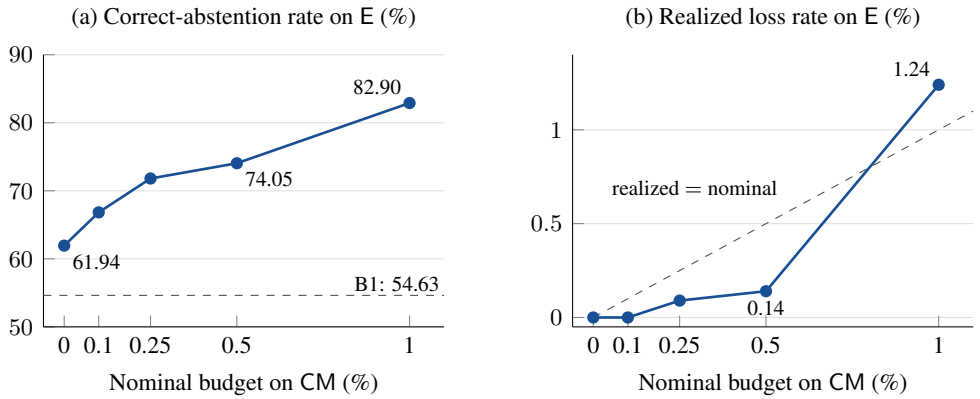

Figure 2: Calibration budget against abstention gain and evaluation-set loss for the default scorer on B1. The horizontal axis is the nominal budget on CM. (a) Correct-abstention rate on E; the dashed line marks B1. (b) Realized original-correct loss rate on E; on the dashed line the realized loss rate equals the nominal budget. Empirical calibration does not force the evaluation points to lie below this line.

Table 10 separates the comparison by failure type. Relative to EmbLoRA, the advantage of ReFrain (ε=0.5%) in correct-abstention rate is concentrated in AAD and IASD. On IVQD, EmbLoRA has a higher correct-abstention rate and corrects 9 errors of B1, so that its Standard accuracy on this subtask is also higher; ReFrain (ε=0.5%) retains all 13 SFT-gained records and loses 1 original correct answer. Different failure types give different trade-offs among answering, abstention, and retention.

Table 10: Results by failure type on B1, on E.

Task	Scheme	Standard (%)	UPD (%)	Lost	Retained
AAD	ReFrain (ε=0.5%)	88.32	60.47	3	20 / 21
	CLIP-UP EmbLoRA	88.01	38.80	25	14 / 21
	Plain abstention SFT	86.39	80.39	58	7 / 21
IASD	ReFrain (ε=0.5%)	86.31	85.38	4	22 / 23
	CLIP-UP EmbLoRA	85.87	77.35	32	15 / 23
	Plain abstention SFT	84.24	94.09	70	8 / 23
IVQD	ReFrain (ε=0.5%)	86.14	68.36	1	13 / 13
	CLIP-UP EmbLoRA	86.67	76.83	5	11 / 13
	Plain abstention SFT	85.61	82.65	12	8 / 13

4.8 Frozen Confirmation on Independent Images

On the confirmation set, the Standard accuracy of B1 is 91.67% and its correct-abstention rate is 21.69%. Both frozen operating points of ReFrain keep all 836 original correct answers of B1, an original-correct loss of 0 of 912; the correct-abstention rate rises by 20.10 pp, with a 95% interval of [16.11, 24.64], and by 44.36 pp, with an interval of [39.30, 50.00], respectively. Both operating points pass the joint criterion (Table 11).

Relative to EmbLoRA, the correct-abstention rate of ReFrain is 8.46 pp higher, with an interval of [3.58, 13.73], and that of ReFrain (ε=0.5%) is 32.72 pp higher, with an interval of [27.13, 38.53]. Both operating points lose 7 fewer original correct answers, a reduction of 0.77 pp in the loss rate, with an interval of [0.22, 1.43]. EmbLoRA also corrects 13 old errors, which puts its net Standard accuracy 0.66 pp above that of ReFrain; the interval of the Standard difference, ReFrain minus EmbLoRA, is [−1.87, 0.44].

Table 11: Frozen confirmation on I. UPD is given as the micro-average over the 816 UPD records and as the macro-average over the failure types. The joint criterion is relative to B1; lost and corrected answers are counted among the 912 Standard records, and retained records are out of the 8 SFT-gained records.

Branch	Standard (%)	UPD, micro (%)	UPD, macro (%)	Lost / corrected	Retained	Joint criterion
B1	91.67	21.69	21.29	0 / 0	8 / 8	reference
ReFrain	91.67	41.79	37.95	0 / 0	8 / 8	passed
ReFrain (ε=0.5%)	91.67	66.05	60.10	0 / 0	8 / 8	passed
CLIP-UP EmbLoRA	92.32	33.33	33.33	7 / 13	6 / 8	passed
Plain abstention SFT	91.78	70.96	66.75	4 / 5	6 / 8	passed

Plain abstention SFT also passes the joint criterion in this round. Relative to ReFrain (ε=0.5%), its correct-abstention rate is 4.90 pp higher, with an interval of [1.34, 8.49], and it loses 4 original correct answers and corrects 5 old errors. The original-correct loss rate of ReFrain is 0.44 pp lower, with an interval of [0.11, 0.88]; the net difference on Standard is -0.11 pp, with an interval of [-0.88, 0.55]. The independent confirmation thus continues the trade-off between retention and abstention coverage: ReFrain keeps all original correct answers, plain abstention SFT has higher abstention coverage, and EmbLoRA brings more additional corrections.

By task, the correct-abstention rates of ReFrain (ε=0.5%) on AAD and IASD are 60.07% and 86.52%, above the 8.68% and 48.88% of EmbLoRA. On IVQD, ReFrain (ε=0.5%) is at 33.72%, and EmbLoRA and plain abstention SFT are at 42.44% and 49.42%, in agreement with the relative trade-off on E. Relative to B0, B1 gains 8 records and loses 4 on the confirmation set; the SFT-gained records lie in 7 clusters, ReFrain retains all of them at both operating points, and each of the two trained branches retains 6. Only counts are reported for this small set. The complete joint criterion, the four direct comparisons, Circular Dual, and the per-task values are given in Appendix E.

5 Discussion

The retention requirement of abstention. When net accuracies are close, abstention branches can still change different numbers of original correct answers. The difference between ReFrain (ε=0.5%) and EmbLoRA on B1 shows that the decomposition into losses and corrections complements the net metric. For applications that continue to rely on existing answers, a deployer can explicitly choose a stricter retention point, or accept a small loss for higher abstention coverage.

The frozen setting on new images. On E and on the independent confirmation set I, both preset operating points of ReFrain give a low original-correct loss and additional abstention. Plain abstention SFT loses 140 of 5,626 on E and 4 of 912 on I, and passes the joint criterion on the latter; the retention cost of the trained branches varies with the evaluation population. On I, EmbLoRA reaches a higher net Standard accuracy through corrections, while ReFrain obtains more correct abstentions and keeps all original correct answers, which again shows the difference between net accuracy and the retention of behavior.

Post-training changes the abstention operating point. The same probe pipeline has a lower correct-abstention rate on B1 than on B0: 67.85% becomes 61.94% at the zero budget, and 80.56% becomes 74.05% at the 0.5% budget. The abstention layer therefore needs to be evaluated again after answer adaptation. The drop in the operating point alone does not yet separate the effects of a change in ranking, a change in calibration, and the reuse of the training set; the score distributions on T and CM and the ranking at low false-positive rates could separate these explanations further.

Choice of compared methods. CLIP-UP provides a route of training and injection that is aimed directly at UPD; LVLM-LP is the direct source of the internal logits score; VL-Uncertainty provides a perturbation-based uncertainty control; MSP and the negative entropy are simple scorers that need no additional supervised learning. The correlation-vector classifier in Appendix F.2 of CLIP-UP is the closest component control; BICR adds a visual-contrast correctness probe, and ACR-MCQ adds a learnable confidence fusion. The comparison is organized by task, mechanism, and reproducibility, and does not merge numbers that papers report on different datasets into a single ranking.

6 Limitations

The shared-SFT setting has a single training seed, and the training set has 677 records. The net Standard gain of SFT is a point estimate whose interval contains zero; the 57 SFT-gained records lie in 28 image clusters, and the interval of the difference in gain retention rate is wide. The image-cluster intervals of this paper measure sampling variation under fixed checkpoints and do not include the randomness of retraining.

E is a historically exposed evaluation population, and I provides a confirmation on independent images after the method was frozen. I covers only AI2D science diagrams, the three failure types have different sample sizes, and the questions and transformations were reviewed by rules and by a multimodal model, without annotation by two human annotators. That the images were not used in the development of this project does not exclude that the base model saw public AI2D data during pretraining or instruction tuning. The abstention baseline of B1 is lower on I than on E, so the absolute gains on the two populations cannot be read directly as a change in the ability of the method; we interpret the results through the paired comparisons within each set. On I there are only 8 SFT-gained records in 7 clusters, and their retention counts do not support a new statistical conclusion.

The preceding SFT and the probe share the Standard records of T. The budget on CM is an empirical constraint. Property 3 covers only the zero budget and a single new image cluster that is exchangeable with the calibration clusters; the zero loss of the two operating points on the confirmation set is not covered by it and gives no population-level guarantee. For E, the exchangeability of its image clusters with those of CM is an assumption about how the data were partitioned. The results on the two original base models support fitting the pipeline separately on each base model; the results on retention after post-training currently come only from Qwen2.5-VL-3B.

In an experiment with variations of the frozen prompt, part of the gain shrinks; these runs are not reported here. The direct retention comparison for abstention by parameter update covers EmbLoRA and plain abstention SFT; the two differ in target text and prompt format, so plain abstention SFT is not a causal ablation that removes only the CLIP signal. The CLIP classifier result is a mechanism port under the common prompt, fitted over a disclosed finite grid of candidates. BICR keeps its core objective and search scheme and is adapted to the common B1, data, and executor. ACR-MCQ changes the inputs and the head width of the original audio-visual classification model to the setting of this task, and the standardized adaptation is reported in separate rows. The added controls were run after the main results had been observed and are extension experiments. Rep2Act [8], Reliable-V-RL [11], VL-Calibration [27], Dr. No [7], and the additional-instruction prompting of MM-UPD [16] have no results in the same setting; Appendix C gives the state of each.

7 Conclusion

On the present multimodal abstention task, a constraint on the loss of correct answers, combined with deterministic execution, turns an existing answerability score into favorable operating points between abstention and retention. Record-level evaluation from a shared answer-SFT starting point shows that the

abstention stage can keep all or most of the SFT-gained answers while adding correct abstentions. The confirmation on independent images, run after the models and thresholds were frozen, further shows that the two preset operating points raise the correct-abstention rate while keeping all original correct answers. Losses, corrections, and gain retention, together with net accuracy, show the actual trade-offs of the different abstention routes.

References

[1] Takuya Akiba, Shotaro Sano, Toshihiko Yanase, Takeru Ohta, and Masanori Koyama. Optuna: A Next-generation Hyperparameter Optimization Framework, 2019. URL https://arxiv.org/abs/1907.10902.

[2] Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal risk control. In *International Conference on Learning Representations (ICLR)*, 2024. URL https://openreview.net/forum?id=33XGfHLtZg. arXiv:2208.02814.

[3] Shuai Bai, Keqin Chen, Xuejing Liu, Jialin Wang, Wenbin Ge, Sibo Song, Kai Dang, Peng Wang, Shijie Wang, Jun Tang, Humen Zhong, Yuanzhi Zhu, Mingkun Yang, Zhaohai Li, Jianqiang Wan, Pengfei Wang, Wei Ding, Zheren Fu, Yiheng Xu, Jiabo Ye, Xi Zhang, Tianbao Xie, Zesen Cheng, Hang Zhang, Zhibo Yang, Haiyang Xu, and Junyang Lin. Qwen2.5-VL technical report, 2025. URL https://arxiv.org/abs/2502.13923.

[4] Zixiu Ding, Zilin Zhao, Yingjie He, Xinlang Kang, Guansu Wang, and Wei Zhang. SAVOR: Self-Aware Visual Grounding via Confidence-Calibrated Reinforcement Learning for Multimodal Hallucination Mitigation, 2026. URL https://arxiv.org/abs/2609.16601. Accepted at ICONIP 2026 according to the arXiv record.

[5] Jessica Echterhoff, Fartash Faghri, Raviteja Vemulapalli, Ting-Yao Hu, Chun-Liang Li, Oncel Tuzel, and Hadi Pouransari. MUSCLE: A Model Update Strategy for Compatible LLM Evolution, 2024. URL https://arxiv.org/abs/2407.09435.

[6] Bradley Efron. Bootstrap methods: Another look at the jackknife. *The Annals of Statistics*, 7(1):1–26, 1979. doi: 10.1214/aos/1176344552.

[7] Farid Hasanov, Riccardo Renzulli, Bahram Zonooz, Mykola Pechenizkiy, and Decebal Constantin Mocanu. Dr. No: Feature Steering for Abstention in Vision-Language Models, 2026. URL https://arxiv.org/abs/2611.03862.

[8] Jialuo He and Huangxun Chen. From knowing to abstaining: Bridging the representation-action gap in vision-language models, 2026. URL https://arxiv.org/abs/2609.32653.

[9] Yufeng Huang, Jiji Tang, Zhuo Chen, Rongsheng Zhang, Xinfeng Zhang, Weijie Chen, Zeng Zhao, Zhou Zhao, Tangjie Lv, Zhipeng Hu, and Wen Zhang. Structure-CLIP: Towards Scene Graph Knowledge to Enhance Multi-modal Structured Representations, 2023. URL https://arxiv.org/abs/2305.06152.

[10] Sofia Jamil, Chinmay Purushottam Bhat, Sooraj Veer R, and Sriparna Saha. *Learn To Say NO!* incorporating model abstention in medical vision-language models for pharmacovigilance. In *Proceedings of the 15th International Joint Conference on Natural Language Processing and the 5th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*,

pages 834–848, Hengqin, China, November 2026. Association for Computational Linguistics. URL https://aclanthology.org/2026.aacl-long.40/.

[11] Paul Jian, Julian Wu, Leo Wei, Dylan Chen, Garvin Guo, Qingpei Guo, and Jianing Li. To answer or not to answer? reliability-oriented reinforcement learning for visual language models. In *Proceedings of the 2026 Conference on Empirical Methods in Natural Language Processing*, pages 53194–53220, Budapest, Hungary, October 2026. Association for Computational Linguistics. doi: 10.18653/2026.emnlp-main.2342. URL https://aclanthology.org/2026.emnlp-main.2342/.

[12] Aniruddha Kembhavi, Mike Salvato, Eric Kolve, Minjoon Seo, Hannaneh Hajishirzi, and Ali Farhadi. A Diagram Is Worth A Dozen Images, 2016. URL https://arxiv.org/abs/1603.07396.

[13] Reza Khanmohammadi, Erfan Miahi, Simerjot Kaur, Charese Smiley, Ivan Brugere, Kundan S Thind, and Mohammad M. Ghassemi. Grounded or guessing? LVLM confidence estimation via blind-image contrastive ranking. In *Proceedings of the 15th International Joint Conference on Natural Language Processing and the 5th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4645–4721, Hengqin, China, November 2026. Association for Computational Linguistics. URL https://aclanthology.org/2026.aacl-long.210/.

[14] Yi-An Lai, Elman Mansimov, Yuqing Xie, and Yi Zhang. Improving Prediction Backward-Compatiblility in NLP Model Upgrade with Gated Fusion, 2023. URL https://arxiv.org/abs/2302.02080. Accepted to Findings of EACL 2023 according to the arXiv record.

[15] Junzhang Liu, Zhecan Wang, Hammad Ayyubi, Haoxuan You, Chris Thomas, Rui Sun, Shih-Fu Chang, and Kai-Wei Chang. Detecting Multimodal Situations with Insufficient Context and Abstaining from Baseless Predictions, 2024. URL https://arxiv.org/abs/2405.11145.

[16] Atsuyuki Miyai, Jingkang Yang, Jingyang Zhang, Yifei Ming, Qing Yu, Go Irie, Yixuan Li, Hai Helen Li, Ziwei Liu, and Kiyoharu Aizawa. Unsolvable problem detection: Robust understanding evaluation for large multimodal models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 6497–6540, Vienna, Austria, July 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.acl-long.326. URL https://aclanthology.org/2025.acl-long.326/.

[17] Fan Nie, Xiaotian Hou, Shuhang Lin, James Zou, Huaxiu Yao, and Linjun Zhang. FactTest: Factuality Testing in Large Language Models with Finite-Sample and Distribution-Free Guarantees, 2024. URL https://arxiv.org/abs/2411.02603.

[18] Tejas Srinivasan, Jack Hessel, Tanmay Gupta, Bill Yuchen Lin, Yejin Choi, Jesse Thomason, and Khyathi Chandu. Selective "selective prediction": Reducing unnecessary abstention in vision-language reasoning. In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 12935–12948, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.findings-acl.767. URL https://aclanthology.org/2024.findings-acl.767/.

[19] Xin Tong, Yang Feng, and Jingyi Jessica Li. Neyman-Pearson (NP) classification algorithms and NP receiver operating characteristics (NP-ROC). *Science Advances*, 4(2):eaao1659, 2018. URL https://arxiv.org/abs/1608.03109.

[20] Dinh Phu Tran, Jihoon Jeong, Saad Wazir, Seongah Kim, Thao Do, Cem Subakan, and Daeyoung Kim. Knowing When to Answer: Adaptive Confidence Refinement for Reliable Audio-Visual Question Answering, 2026. URL https://arxiv.org/abs/2602.04924.

[21] Frederik Träuble, Julius von Kügelgen, Matthäus Kleindessner, Francesco Locatello, Bernhard Schölkopf, and Peter Gehler. Backward-Compatible Prediction Updates: A Probabilistic Approach, 2021. URL https://arxiv.org/abs/2107.01057.

[22] Ben Vardi, Oron Nir, and Ariel Shamir. CLIP-UP: CLIP-based unanswerable problem detection for visual question answering. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, pages 5898–5908, March 2026. URL https://openaccess.thecvf.com/content/WACV2026/html/Vardi_CLIP-UP_CLIP-Based_Unanswerable_Problem_Detection_for_Visual_Question_Answering_WACV_2026_paper.html. arXiv:2501.01371.

[23] Benedikt J. Wagner. Two Axes of LLM Abstention: Answer Correctness and Question Answerability, 2026. URL https://arxiv.org/abs/2607.08456.

[24] Zhiyuan Wang, Aniri, Tianlong Chen, Yue Zhang, Heng Tao Shen, Xiaoshuang Shi, and Kaidi Xu. LEC: Linear Expectation Constraints for Selection-Conditioned Risk Control in Selective Prediction and Routing Systems, 2025. URL https://arxiv.org/abs/2512.01556. Accepted at ICML 2026 according to the arXiv record.

[25] Spencer Whitehead, Suzanne Petryk, Vedaad Shakib, Joseph Gonzalez, Trevor Darrell, Anna Rohrbach, and Marcus Rohrbach. Reliable visual question answering: Abstain rather than answer incorrectly. In *Computer Vision – ECCV 2022*, Lecture Notes in Computer Science, pages 148–166. Springer Nature Switzerland, 2022. doi: 10.1007/978-3-031-20059-5_9. arXiv:2204.13631.

[26] Tobias Jan Wieczorek, Nathalie Daun, Mohammad Emtiyaz Khan, and Marcus Rohrbach. Variational Visual Question Answering for Uncertainty-Aware Selective Prediction. *Transactions on Machine Learning Research*, 2026. URL https://arxiv.org/abs/2505.09591.

[27] Wenyi Xiao, Xinchi Xu, and Leilei Gan. VL-Calibration: Decoupled Confidence Calibration for Large Vision-Language Models Reasoning, 2026. URL https://arxiv.org/abs/2604.09529. Accepted at ACL 2026 according to the arXiv record.

[28] Yuqing Xie, Yi an Lai, Yuanjun Xiong, Yi Zhang, and Stefano Soatto. Regression Bugs Are In Your Model! Measuring, Reducing and Analyzing Regressions In NLP Model Updates, 2021. URL https://arxiv.org/abs/2105.03048. Accepted at ACL 2021 according to the arXiv record.

[29] Jian Xu, Yanning Wu, Delu Zeng, John Paisley, and Qibin Zhao. Look Again Before You Abstain: Budgeted Conformal Evidence Acquisition for Reliable Vision-Language Model, 2026. URL https://arxiv.org/abs/2606.16667.

[30] Sijie Yan, Yuanjun Xiong, Kaustav Kundu, Shuo Yang, Siqi Deng, Meng Wang, Wei Xia, and Stefano Soatto. Positive-congruent training: Towards regression-free model updates. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 14299–14308, June 2021. doi: 10.1109/CVPR46437.2021.01407.

[31] Jihan Yao, Achin Kulshrestha, Nathalie Rauschmayr, Reed Roberts, Banghua Zhu, Yulia Tsvetkov, and Federico Tombari. Reading Between the Lines: Abstaining from VLM-Generated OCR Errors via Latent Representation Probes, 2025. URL https://arxiv.org/abs/2511.19806.

[32] Junwei Zhang, Yifei Wang, Wanli Yang, Hexiang Tan, Rui Tang, Xin Chen, Jingang Wang, Du Su, and Fei Sun. Strong yet Controllable LLM Abstention without Learning to Abstain, 2026. URL https://arxiv.org/abs/2610.16318.

[33] Ruiyang Zhang, Hu Zhang, and Zhedong Zheng. VL-Uncertainty: Detecting Hallucination in Large Vision-Language Model via Uncertainty Estimation, 2024. URL https://arxiv.org/abs/2411.11919.

[34] Qinyu Zhao, Ming Xu, Kartik Gupta, Akshay Asthana, Liang Zheng, and Stephen Gould. The First to Know: How Token Distributions Reveal Hidden Knowledge in Large Vision-Language Models?, 2024. URL https://arxiv.org/abs/2403.09037. Accepted at ECCV 2024 according to the arXiv record.

[35] Jinguo Zhu, Weiyun Wang, Zhe Chen, Zhaoyang Liu, Shenglong Ye, Lixin Gu, Hao Tian, Yuchen Duan, Weijie Su, Jie Shao, Zhangwei Gao, Erfei Cui, Xuehui Wang, Yue Cao, Yangzhou Liu, Xingguang Wei, Hongjie Zhang, Haomin Wang, Weiye Xu, Hao Li, Jiahao Wang, Nianchen Deng, Songze Li, Yinan He, Tan Jiang, Jiapeng Luo, Yi Wang, Conghui He, Botian Shi, Xingcheng Zhang, Wenqi Shao, Junjun He, Yingtong Xiong, Wenwen Qu, Peng Sun, Penglong Jiao, Han Lv, Lijun Wu, Kaipeng Zhang, Huipeng Deng, Jiaye Ge, Kai Chen, Limin Wang, Min Dou, Lewei Lu, Xizhou Zhu, Tong Lu, Dahua Lin, Yu Qiao, Jifeng Dai, and Wenhai Wang. InternVL3: Exploring Advanced Training and Test-Time Recipes for Open-Source Multimodal Models, 2025. URL https://arxiv.org/abs/2504.10479.

[36] Mingye Zhu, Yi Liu, Quan Wang, Junbo Guo, and Zhendong Mao. FlipGuard: Defending Preference Alignment against Update Regression with Constrained Optimization, 2024. URL https://arxiv.org/abs/2410.00508. Accepted at EMNLP 2024 according to the arXiv record.

A Proofs

Properties 1 and 2 follow directly from the veto rule of Eq. (5), as explained after their statements in Section 3.5. This appendix proves Property 3 and adds three remarks on its scope.

Setup. An image cluster is a finite set of records, and the score s is real-valued. For an image cluster k, $m_k \in (-\infty, +\infty]$ is the lowest score among its Standard records that f answers correctly, with $m_k = +\infty$ if there is no such record. Clusters $1, \ldots, K$ are the image clusters of CM, cluster $K+1$ is the new image cluster, and $M_K = \min_{k \le K} m_k$.

The zero-budget threshold. By Property 2, g_τ loses a Standard record of CM that f answers correctly exactly when the score of the record is below τ. Hence $\widehat{L}_{\mathsf{CM}}(g_\tau; f) = 0$ if and only if $\tau \le M_K$, and the rule of Eq. (4) gives

$$\tau_0 \le M_K. \tag{8}$$

The candidate set includes $-\infty$, which is always feasible and vetoes no record. The proofs below use only Eq. (8) and therefore hold for any candidate set $\mathcal{T}$ with an always-feasible no-veto endpoint. In the same way, g_{τ_0} loses an original correct answer in the new cluster exactly when $m_{K+1} < \tau_0$, which implies $m_{K+1} < M_K$.

Proof of part (a). For $i = 1, \ldots, K+1$, let A_i be the event that $m_i < m_j$ for all $j \neq i$. The events A_i are pairwise disjoint, so $\sum_i \Pr[A_i] \le 1$. Under the conditioning of the property, f and s are fixed, each m_k is the same function of cluster k, and the clusters are exchangeable; the values $m_1, \ldots, m_{K+1}$ are therefore exchangeable, and all $\Pr[A_i]$ are equal. Hence $\Pr[A_{K+1}] \le 1/(K+1)$. A loss in the new cluster implies $m_{K+1} < M_K$, which is the event A_{K+1}. If the lowest value is attained by more than one cluster, including the case in which several clusters have $m_k = +\infty$, no A_i occurs, which only lowers the probability of a loss. The bound holds under the conditioning and therefore also after averaging over it. □

Proof of part (b). Under the conditioning, $m_1, \ldots, m_{K+1}$ are i.i.d. For $t \in [-\infty, +\infty]$, let $G(t) = \Pr[m_1 < t]$. G is non-decreasing and left-continuous, with $G(-\infty) = 0$ and $G(+\infty) = \Pr[m_1 < +\infty]$. Because the clusters are i.i.d., conditioning on the calibration clusters does not change the distribution of the new cluster; the conditional probability of a loss in the new cluster is therefore $G(\tau_0)$, and $G(\tau_0) \leq G(M_K)$ by Eq. (8). It remains to show that $\Pr[G(M_K) > \alpha] \leq (1-\alpha)^K$. If $G(+\infty) \leq \alpha$, then $G(M_K) \leq \alpha$ always. Otherwise, let $t^* = \sup\{t \in \mathbb{R} : G(t) \leq \alpha\}$. This set is non-empty because $G(t) \to 0$ as $t \to -\infty$, and it is bounded above because $G(t) \to G(+\infty) > \alpha$ as $t \to +\infty$; hence t^* is finite. By left-continuity, $G(t^*) \leq \alpha$. Since $G(t) > \alpha$ for every $t > t^*$, $\Pr[m_1 \leq t^*] = \lim_{t \downarrow t^*} G(t) \geq \alpha$. If $G(M_K) > \alpha$, then $M_K > t^*$, that is, $m_k > t^*$ for all $k \leq K$; this event has probability $(1 - \Pr[m_1 \leq t^*])^K \leq (1-\alpha)^K$. □

For the level quoted in Section 3.5, setting $(1-\alpha)^K = 0.05$ with $K = 183$ gives $\alpha = 1 - 0.05^{1/183} \approx 1.62\%$.

Remark 1: several evaluation clusters. Consider an evaluation set with N image clusters, each of which is exchangeable with the K calibration clusters in the sense of the property. By part (a) and the linearity of expectation, the expected number of evaluation clusters in which an original correct answer is lost is at most $N/(K+1)$; for E this is $416/184 \approx 2.26$. The loss events of different evaluation clusters share the threshold τ_0; the marginal bound alone does not establish concentration of this count around its expectation. The probability that at least one of the N clusters has a loss is not small. If all $N + K$ clusters are exchangeable, $\tau_0 = M_K$, and the values m_k are almost surely distinct, a loss occurs in some evaluation cluster exactly when the lowest of the $N + K$ values belongs to an evaluation cluster, which has probability $N/(N+K) = 416/599 \approx 69.4\%$.

Remark 2: budgets above zero. The same argument gives a statement at the level of clusters for a budget of r calibration records. If g_τ loses at most r records of CM, then at most r calibration clusters have $m_k < \tau$, so the threshold does not exceed the $(r+1)$-th smallest of $m_1, \ldots, m_K$. A loss in the new cluster then requires that at most r of the other K values be less than or equal to m_{K+1}. At most $r + 1$ of the $K + 1$ clusters can have this property at the same time, because the cluster with the largest value among them has all the others at or below it; by exchangeability, the probability is at most $(r+1)/(K+1)$. For $\varepsilon = 0.5\%$, $r = 12$, and the bound is $13/184 \approx 7.07\%$. This is the probability that a new cluster contains at least one lost record; it is loose, and it is not a bound on the loss rate over records. We therefore state Property 3 for the zero budget only and base no claim about ReFrain (ε=0.5%) on this remark.

Remark 3: what the property does not cover. The bound is marginal over image clusters and is not conditional on a failure type or on any other subgroup. It concerns the loss only and is silent on utility: a rule that never vetoes satisfies it, and the gain in correct-abstention rate at τ_0 is an empirical matter. It is tied to the frozen f, the frozen scorer, the prompt, the option format, and the NOTA mapping $n(\cdot)$; if any of these changes after calibration, the threshold has to be calibrated again.

B Calibration Budget and Evaluation-Set Loss

On B1, the zero-budget operating point of every scorer has an original-correct loss of 0 on both CM and E. Table 12 gives the operating points at the 0.5% budget; the denominators are the 2,406 Standard records of CM and the 5,626 Standard records of E.

The loss rates on E of the CLIP classifier with fixed $C = 1$ and of the standardized ACR-MCQ are above the nominal calibration budget. The empirical constraint on the calibration set and the realized value on the

Table 12: Original-correct loss at the 0.5% calibration budget on B1, on the calibration set and on the evaluation set. Every row is ReFrain with the named scorer.

Scorer	Losses on CM	Loss rate on CM (%)	Losses on E	Loss rate on E (%)
Cross-validated probe (default)	11	0.46	8	0.14
Probe with fixed $C = 1$	12	0.50	13	0.23
MSP	7	0.29	14	0.25
Negative entropy	8	0.33	3	0.05
CLIP classifier, fixed $C = 1$	12	0.50	29	0.52
CLIP classifier, cross-validated	12	0.50	8	0.14
ACR-MCQ, raw features	7	0.29	14	0.25
ACR-MCQ, standardized features	10	0.42	36	0.64
BICR	11	0.46	10	0.18

evaluation set therefore have to be reported separately. Table 13 gives the complete counts for the five budgets of Section 4.7.

Table 13: The default scorer on B1 at five calibration budgets, on E.

Nominal budget on CM (%)	Correct-abstention rate on E (%)	Losses on E	Loss rate on E (%)
0	61.94	0	0.00
0.1	66.83	0	0.00
0.25	71.81	5	0.09
0.5	74.05	8	0.14
1	82.90	70	1.24

C Implementation and Fidelity of the Compared Methods

This appendix gives, for each compared method, its fidelity label in the sense of Section 3.2 and the points at which our implementation differs from the original. All methods use the data roles of Table 1, the common option prompt, and the common parser.

LVLM-LP probe. Mechanism port. The feature, namely the full-vocabulary logits at the first answer position, and the linear form of the probe follow LVLM-LP and its official repository [34]. The regularization strength is selected by cross-validation over image clusters within T, and the fit with fixed $C = 1$ is reported as a second scorer. Table 7 lists the probe at its default probability threshold; all calibrated rows are ReFrain.

CLIP-UP. On InternVL3-8B, a base model supported by the original work, we use the released code and the training recipe of the authors [22], train on T, and evaluate under the common option prompt. On Qwen2.5-VL-3B, for both B0 and B1, the implementation is a mechanism port, because the released code does not cover this model family. The projected correlation token is placed after the vision-end token, whereas the official InternVL implementation places it before the image-end token. Training, fitting, and calibration are done separately on each base model. The Emb result on InternVL3-8B in Table 3 is a result of transferring the method to the present protocol; it is not a verdict on the numbers of the original paper, which were obtained with its own data and prompts.

VL-Uncertainty. Mechanism port of the official repository [33], with the model interface changed from Qwen2-VL to Qwen2.5-VL; we report it on the original Qwen2.5-VL-3B only. Each record is answered under five perturbed versions of the input, which pair a blurred image with a rephrased question, and the uncertainty is the base-2 entropy of the answer clusters over the five samples. In the native execution, a record with an uncertainty of at least 1.0 receives the displayed abstention letter, and every other record receives the initial answer, which is sampled at temperature 0.1. Under the common executor, the uncertainty is used as the score, and the cached greedy answer of the base model is kept on the answered side. Three deviations from the official code are disclosed. The full question and hint are rephrased, where the official code rephrases the first line only. Answers are clustered by the frozen option parser of the protocol, where the official code takes the first integer. Records whose questions have an identical prefix share their rephrasings.

MSP and negative entropy. Both are computed from the output distribution of f and need no fitting; they have no threshold of their own and appear only with the calibration and executor of ReFrain.

CLIP correlation-vector classifier. Mechanism port of the classifier in Appendix F.2 of CLIP-UP, with the official Structure-CLIP weights [9]. The original paper does not give the complete fitting hyperparameters, so two fits are disclosed (Section 4.4).

BICR. Mechanism port. The core objective and the search scheme of BICR are kept [13], and the method is adapted to B1, to T, and to the common executor.

ACR-MCQ. An adaptation of ACR, which was proposed for an audio-visual classification model [20]. The inputs and the width of the heads are changed to the setting of this task (Section 4.5), and we therefore report it under a separate name.

Plain abstention SFT. This branch is defined in this paper (Section 4.2) and has no original implementation.

Methods without results in the same setting. Rep2Act trains a routing head together with rationale targets that are grounded in the image [8]. We have implemented the head, the four training objectives, and the two-stage training; T has no such rationales, so a run on it would be a partial implementation, and none is reported. Reliable-V-RL trains with rewards from a grader model that scores responses against rubrics written for each question [11]. We have implemented the reward and the training; the rubrics, the contradiction statements from which they are written, and the grader model are not available for T, and no run is reported. For VL-Calibration [27], the official source has been adapted and its data checks pass, but training has not been run. Dr. No needs a sparse autoencoder that matches the answer model [7]; there is none for B1. The additional-instruction prompting of MM-UPD [16] has not been run on B1.

D Loss of SFT-Gained and of Persistently Correct Records

On E, B1 has 57 SFT-gained records and a further 4,850 Standard records that both B0 and B1 answer correctly, which we call persistently correct. Table 14 compares the rates at which the abstention stage loses the two kinds of records; the intervals use 2,000 paired bootstrap resamples over image clusters.

EmbLoRA and plain abstention SFT lose SFT-gained records at a higher rate, with intervals above zero; the interval of the difference for ReFrain (ε=0.5%) contains zero. The set of SFT-gained records is itself selected by correctness before and after SFT and differs in difficulty from the persistently correct set; these statistics describe a conditional difference in behavior on the present sample.

Table 14: Loss of SFT-gained records and of persistently correct records in the abstention stage, on E.

Scheme	SFT-gained records lost (%)	Persistently correct records lost (%)	Difference [95% interval] (pp)
ReFrain (ε=0.5%)	2 / 57 (3.51)	6 / 4,850 (0.12)	3.39 [−0.22, 12.25]
CLIP-UP EmbLoRA	17 / 57 (29.82)	45 / 4,850 (0.93)	28.90 [12.82, 49.88]
Plain abstention SFT	34 / 57 (59.65)	106 / 4,850 (2.19)	57.46 [37.68, 79.47]

E Construction of the Independent Confirmation and Full Comparisons

Images and duplicates. Among the original AI2D questions, we use only those that need no added letter overlay on the image, and we keep the official answers and the original images. Candidate images are matched, for exact and near duplicates, against the images used earlier in this project, all official AI2D test images, and the images of Geometry3K; the near-duplicate criterion is a Hamming distance of at most 4 between 64-bit dHash values. Different confirmation clusters share neither original images nor replacement images.

Semantic review. The review of semantics uses fixed rules and a multimodal model. Its inputs are limited to the original question, the options, the image, and the transformation; they do not include the answers, scores, or correctness of the branches under evaluation. The 300 candidate clusters that entered the review were screened as in Table 15, and the original correct answers were not modified. All decisions were recorded, and the final sample was fixed before inference.

Table 15: Semantic screening of the candidate clusters of I. Numbers in parentheses are counts of excluded clusters.

Type	Reviewed	Excluded	Final	Reasons for exclusion
AAD	100	4	96	Options overlap in meaning, a meta-option cannot be judged, the stem reveals the answer, or the original question is unclear
IASD	100	11	89	Nested option labels (3); an "NA" meta-option (3); still answerable or ambiguous after replacement (3); a synonymous correct answer (1); confused concepts in the original question (1)
IVQD	100	57	43	Insufficient dependence on the image or an ambiguous question (55); the original answer is not unique on review of the image (2)

Records and resampling unit. Every Standard question keeps four substantive options and is rotated over all circular shifts; AAD rotates the three options that remain after the correct one is removed, and IASD and IVQD rotate four options. The NOTA text is "None of the above" for AAD and IASD and "The image and question are irrelevant." for IVQD, and it is fixed in the last position. The unit of resampling consists of all Standard and UPD records of an original question together with its unique foreign image, which gives 228 clusters. Inference uses a batch size of 1, greedy decoding, BF16 precision, and at most 1,024 new tokens.

Full comparisons. The models and thresholds are those of the frozen configuration, and I is not used to select methods or operating points. Table 16 gives the joint criterion with its intervals, and Table 17 reports all four direct comparisons that were fixed before inference. Table 18 gives the results by failure type.

Table 16: Joint criterion on l, relative to B1. The lost count is out of 912 Standard records; "Invalid" counts outputs that the parser cannot map to an option.

Branch	ΔStandard [95% interval] (pp)	ΔUPD [95% interval] (pp)	Lost (%)	Circular Dual (%)	Invalid
B1	reference	reference	0 (0.00)	10.08	0
ReFrain	0.00 [0.00, 0.00]	20.10 [16.11, 24.64]	0 (0.00)	23.93	—
ReFrain (ε=0.5%)	0.00 [0.00, 0.00]	44.36 [39.30, 50.00]	0 (0.00)	44.94	—
CLIP-UP EmbLoRA	0.66 [−0.44, 1.87]	11.64 [8.30, 15.17]	7 (0.77)	20.05	1
Plain abstention SFT	0.11 [−0.55, 0.88]	49.26 [44.00, 54.63]	4 (0.44)	52.82	0

Table 17: The four direct comparisons on l that were fixed before inference, as differences with paired 95% intervals over image clusters. A positive ΔStandard or ΔUPD means that ReFrain is higher; a positive reduction in original-correct loss means that ReFrain loses fewer answers.

Comparison	ΔStandard (pp)	ΔUPD (pp)	Reduction in original-correct loss (pp)
ReFrain − CLIP-UP EmbLoRA	−0.66 [−1.87, 0.44]	8.46 [3.58, 13.73]	0.77 [0.22, 1.43]
ReFrain − plain abstention SFT	−0.11 [−0.88, 0.55]	−29.17 [−34.10, −24.26]	0.44 [0.11, 0.88]
ReFrain (ε=0.5%) − CLIP-UP EmbLoRA	−0.66 [−1.87, 0.44]	32.72 [27.13, 38.53]	0.77 [0.22, 1.43]
ReFrain (ε=0.5%) − plain abstention SFT	−0.11 [−0.88, 0.55]	−4.90 [−8.49, −1.34]	0.44 [0.11, 0.88]

On l, the Standard accuracy of B0 is 91.23% and that of B1 is 91.67%; SFT gains 8 records and loses 4. The SFT-gained records lie in 7 clusters, below the threshold of 10 clusters that was set in advance for reporting an interval.

Table 18: Results by failure type on l. The numbers of Standard and UPD records are 384 and 288 for AAD, 356 and 356 for IASD, and 172 and 172 for IVQD.

Task	Branch	Standard (%)	UPD (%)	Lost / corrected
AAD	B1	92.19	11.11	0 / 0
	ReFrain	92.19	28.47	0 / 0
	ReFrain (ε=0.5%)	92.19	60.07	0 / 0
	CLIP-UP EmbLoRA	94.01	8.68	1 / 8
	Plain abstention SFT	92.97	63.19	1 / 4
IASD	B1	95.79	29.49	0 / 0
	ReFrain	95.79	60.96	0 / 0
	ReFrain (ε=0.5%)	95.79	86.52	0 / 0
	CLIP-UP EmbLoRA	95.22	48.88	4 / 2
	Plain abstention SFT	95.22	87.64	2 / 0
IVQD	B1	81.98	23.26	0 / 0
	ReFrain	81.98	24.42	0 / 0
	ReFrain (ε=0.5%)	81.98	33.72	0 / 0
	CLIP-UP EmbLoRA	82.56	42.44	2 / 3
	Plain abstention SFT	81.98	49.42	1 / 1